

C

K19G

2026-09-30

Table of contents

C Programming	62
Who this is for	62
How to use this book	62
How to compile examples	63
Workbook (content/99-workbook)	63
Related books	63
Modules (15 + workbook + appendices)	64
Prerequisites	64
Getting started	65
 Intro	 66
 Comprehensive C Programming Syllabus	 67
Course Overview	67
Learning Objectives	67
Prerequisites	67
Course Structure	67
Module 1: Foundations and Environment Setup	67
1.1 Introduction to C Programming	68
1.2 Development Environment Setup	68
1.3 First C Program	68
1.4 Basic I/O	68
Module 2: Data Types and Variables	69
2.1 Fundamental Data Types	69
2.2 Variables and Constants	69
2.3 Type Conversions	69
2.4 Operators and Expressions	69
Module 3: Control Flow	69
3.1 Conditional Statements	70
3.2 Loops	70
3.3 Program Flow Control	70
Module 4: Functions and Modular Programming	70
4.1 Function Fundamentals	70
4.2 Advanced Function Concepts	71
4.3 Modular Programming	71
4.4 Standard Library Functions	71
Module 5: Arrays and Strings	71
5.1 Arrays	71
5.2 Strings	72

5.3 Advanced Array Concepts	72
Module 6: Pointers and Memory Management	72
6.1 Pointer Fundamentals	72
6.2 Pointers and Arrays	72
6.3 Dynamic Memory Management	73
6.4 Advanced Pointer Concepts	73
6.5 Common Pointer Pitfalls	73
Module 7: Structures and User-Defined Types	73
7.1 Structures	73
7.2 Unions and Enumerations	73
7.3 Advanced Structure Concepts	74
7.4 Abstract Data Types	74
Module 8: File I/O and System Programming	74
8.1 File Operations	74
8.2 File I/O Functions	74
8.3 Advanced File Operations	74
8.4 System-Level Programming	75
Module 9: Advanced C Features and Modern Standards	75
9.1 C99 Features	75
9.2 C11 Features	75
9.3 C17 and C23 Features	75
9.4 Modern C Development Practices (2024)	76
9.5 Advanced Preprocessor	76
Module 10: Data Structures and Algorithms	77
10.1 Linear Data Structures	77
10.2 Tree Structures	77
10.3 Hash Tables	77
10.4 Graph Algorithms	77
10.5 Sorting and Searching	77
Module 11: Network Programming	78
11.1 Network Fundamentals	78
11.2 Socket Programming	78
11.3 Advanced Network Programming	78
11.4 Network Protocols	78
11.5 Security Considerations	78
Module 12: Embedded Systems Programming	79
12.1 Embedded Systems Basics	79
12.2 Low-Level Programming	79
12.3 Embedded C Techniques	79
12.4 Real-Time Programming	79
12.5 Hardware Interfaces	79
Module 13: Performance Optimization and Profiling	80
13.1 Performance Analysis	80
13.2 Optimization Techniques	80
13.3 Memory Optimization	80
13.4 Advanced Optimization	80
Module 14: Testing and Debugging	80
14.1 Testing Methodologies	81

14.2 Debugging Techniques	81
14.3 Error Handling	81
14.4 Code Quality	81
Module 15: Advanced Topics and Capstone	81
15.1 Multi-threading and Concurrency	81
15.2 Advanced System Programming	82
15.3 Domain-Specific Applications	82
15.4 Capstone Project	82
Assessment Methods	82
Continuous Assessment (60%)	82
Major Projects (25%)	82
Examinations (15%)	83
Recommended Resources	83
Essential Books	83
Advanced and Specialized References	83
Current Standards and Specifications	83
Modern Online Resources (2024)	83
Development Tools (2024 Current Versions)	84
Embedded and IoT Development (Current)	84
Security and Best Practices	84
Course Policies	84
Prerequisites Reinforcement	84
Academic Integrity	84
Accessibility and Support	85
Learning Outcomes Verification	85

Foundations 86

Introduction to C Programming 87

What is C?	87
A short history (enough to orient you)	87
Why learn C?	87
From source to process: how a C program runs	88
What each stage does	89
Watching the stages with GCC	89
Memory model intro: text, data, stack, heap	90
Hello, World — anatomy line by line	91
Small but important details	92
Compile and run	92
Full mini-programs	92
1. Temperature converter (°C °F)	92
2. Maximum of three integers	93
3. Simple calculator CLI	94
4. Exit status practice	95
Common misconceptions	96
“C is basically the same as C++”	96
“If it compiles, it is correct”	96

“The compiler is just a translator; any compiler is fine”	97
“I need to master every standard library function first”	97
“void main() is fine”	97
How this book is organized	97
What to practice first	98
Applications of C (why the language stays relevant)	98
Summary	99
Exercises	99
Development Environment Setup	101
Introduction	101
Installing Compilers	101
GCC (GNU Compiler Collection)	101
Clang	102
Microsoft Visual C++ (MSVC)	102
Modern IDEs and Editors	103
Visual Studio Code	103
CLion	104
Code::Blocks	105
Command-Line Tools	105
Make	105
CMake	105
Git	106
Debugging Tools	106
GDB (GNU Debugger)	106
Valgrind	107
AddressSanitizer	107
Package Management	108
vcpkg	108
Conan	108
Cloud Development Environments	108
GitHub Codespaces	108
Gitpod	109
Cross-Platform Development	109
Windows Subsystem for Linux (WSL)	109
Docker	109
Real Workflows: From One File to a Small Project	110
Single-file recipe	110
Multi-file build recipes	111
Full program: tiny CLI calculator (single file)	112
GCC vs Clang (What Actually Differs)	114
Side-by-side on the same bug	114
Useful shared flags	114
macOS note	115
Make: Multi-File Builds Without Typing Everything	115
Minimal Makefile for the util/main project	115
Full program + Make: word count of argv strings	116

CMake Mini-Project (Practical One-Pager)	116
Configure, build, run	117
Optional: compile_commands.json for editors	117
Editor + Debugger One-Pager	118
Terminal editor + GDB (always works)	118
LLDB (macOS default)	118
VS Code (minimum useful setup)	119
GDB against a real bug (complete mini-lab)	119
Sanitizers and Memory Tools (Environment Checklist)	120
Modern Development Practices	120
Version Control Workflow	120
Project Structure	121
Continuous Integration	121
Verify Your Setup (Smoke Tests)	121
Exercises	122
Summary	123
Your First C Program	124
Introduction	124
The classic “Hello, World!”	124
Line-by-line	124
Create, compile, run	125
What the compiler does (short version)	125
Structure of a C program	126
Comments	126
More complete single-file examples	127
Greeting with command-line arguments	127
Sum of two integers (with validation)	127
Multiple helpers in one file	128
Print a small banner	128
Multi-file tiny program	128
Build it	129
Why the include guard?	130
Header vs source rules	130
Common compile errors (with explanations)	130
Missing semicolon	130
Missing <code>#include <stdio.h></code>	130
Wrong main spelling or signature	131
Case sensitivity	131
Undeclared function (multi-file)	131
Multiple definition of a function	131
Forgetting to link a file	131
Format / type mismatch (warning)	131
Using <code>=</code> instead of <code>==</code> in a condition (logic bug; may still compile)	132
Unterminated string	132
Tips for reading diagnostics	132
Lab A — Personal info printer	132
Lab B — Unit converter	133

Variations worth knowing	135
Implicit <code>return 0</code> in <code>main</code> (C99+)	135
Prototype before <code>main</code>	135
Using <code>stderr</code> for errors	135
Beginner mistakes checklist	136
Summary	136
Exercises	136
Basic Input and Output	138
Introduction	138
The <code><stdio.h></code> library	138
Output with <code>printf</code>	139
Prototype (conceptual)	139
Minimal example	139
Format specifier table	139
Runnable demo — mixed types	140
Width, precision, flags	140
<code>fprintf</code> and streams	141
<code>snprintf</code> — safe string formatting	141
Input with <code>scanf</code>	142
Prototype (conceptual)	142
Address operator	142
Specifier table for <code>scanf</code> (common cases)	142
Full demo — reading several fields	143
Buffer caveats with <code>scanf</code>	144
Leftover newlines	144
<code>%s</code> stops at whitespace	144
Always bound string width	144
Failed conversions leave bad data	145
Clearing the rest of a line	145
Character I/O: <code>getchar</code> and <code>putchar</code>	145
Echo a character	146
Copy <code>stdin</code> to <code>stdout</code> until EOF	146
Count characters and lines	146
Reading lines safely with <code>fgets</code>	147
Prototype	147
Basic pattern	147
Why prefer <code>fgets</code> for text lines?	148
Mixing <code>fgets</code> and <code>scanf</code>	148
Detecting truncated lines	149
<code>sscanf</code> — parse from a string	150
Checking return values	150
<code>printf</code>	150
<code>scanf</code>	151
Mini-project 1 — Menu-driven CLI	151
Mini-project 2 — Simple invoice printer	153
Practical calculator (formatted I/O review)	155
Best practices (quick reference)	156

Summary	156
Exercises	156
Module 1: Foundations and Environment Setup	158
Overview	158
Learning goals	158
Chapters (map)	158
Suggested path	159
Tooling cheat sheet	159
Prerequisites	159
Self-check before Module 2	160
Estimated time	160
Next module	160
Resources	160
Data Types	161
Fundamental Data Types	162
Introduction	162
Basic Data Types in C	162
Integer Types	162
Typical Sizes (LP64 Linux, e.g. x86_64)	162
Unsigned Variants	163
Worked Demo: sizeof of Major Types	163
Limits from <code>limits.h</code>	164
Floating-Point Types and <code>float.h</code>	165
Classic Precision Surprise	166
Character Types	167
Boolean Type (C99+)	167
Fixed-Width Types (<code>stdint.h</code>) — Prefer These When Width Matters	168
Exact-Width	168
Printing Portably: <code>inttypes.h</code>	168
When to Use What	169
Example: Portable Binary Header	169
Signed Overflow Is Undefined Behavior	171
Safer Patterns Summary	173
Type Specifiers and Qualifiers	173
Type Promotion (Preview)	174
Practical Combined Explorer	175
Best Practices	175
Exercises	176
Exercise 1 — ABI report	176
Exercise 2 — limits-driven clamp	176
Exercise 3 — safe multiply	176
Exercise 4 — protocol fields	176
Exercise 5 — float.h scavenger	176
Exercise 6 — UB sanitizer lab	177

Summary	177
Variables and Constants	178
Introduction	178
Variables	178
What is a Variable?	178
Variable Declaration	178
Variable Initialization	178
Variable Assignment	179
Variable Naming Rules	179
Variable Naming Conventions	180
Constants	180
Literal Constants	180
#define Preprocessor Directive	181
const Qualifier	181
enum (Enumeration)	182
Storage Classes	182
auto	182
register	183
static	183
extern	184
Scope and Lifetime	185
Scope	185
Lifetime	186
Type Qualifiers	186
const	186
volatile	187
restrict (C99+)	187
Practical Examples	188
Variable Declaration and Usage	188
Constants Comparison	188
Storage Class Demonstration	189
Best Practices	190
1. Initialize Variables	190
2. Use Meaningful Names	190
3. Use const for Fixed Values	190
4. Limit Global Variables	191
5. Use Appropriate Storage Classes	191
Summary	191
Type Conversions	192
Introduction	192
Implicit Type Conversions (Coercion)	192
Usual Arithmetic Conversions	192
Integer Promotion	193
Conversion Hierarchy	193
Boolean Context Conversions	194

Explicit Type Conversions (Casting)	195
Cast Operator Syntax	195
Common Casting Scenarios	195
Pointer Casting	196
Potential Pitfalls and Data Loss	197
Truncation	197
Overflow	197
Precision Loss	198
Safe Conversion Practices	198
Checking for Overflow	198
Using Standard Library Functions	199
Conversion Functions	199
Mathematical Conversions	199
String Conversions	200
Practical Examples	200
Temperature Converter with Type Safety	200
Data Type Range Checker	201
Mixed Type Calculator	202
Best Practices	203
1. Be Explicit About Conversions	203
2. Check for Data Loss	203
3. Use Appropriate Functions for String Conversion	203
4. Understand Truncation vs Rounding	204
5. Be Careful with Signed/Unsigned Conversions	204
Summary	204
Operators and Expressions	205
Introduction	205
Categories of Operators	205
Arithmetic Operators	205
Basic Arithmetic Operations	206
Division and Modulus Behavior	206
Relational Operators	207
Relational Operations	207
Logical Operators	207
Logical Operations	208
Bitwise Operators	208
Bitwise Operations	209
Assignment Operators	210
Assignment Operations	210
Increment and Decrement Operators	211
Increment/Decrement Examples	211
Conditional (Ternary) Operator	212
Conditional Operator Examples	212
Special Operators	213
sizeof Operator	213
Address-of (&) and Dereference (*) Operators	213

Operator Precedence and Associativity	214
Precedence Levels (Highest to Lowest)	214
Precedence Examples	214
Practical Examples	215
Mathematical Expression Evaluator	215
Bit Manipulation Utility	216
Conditional Logic Demonstrator	217
Best Practices	218
1. Use Parentheses for Clarity	218
2. Be Careful with Side Effects	218
3. Use Meaningful Variable Names in Complex Expressions	218
4. Be Aware of Integer Division	219
5. Use Bitwise Operations for Bit Manipulation	219
Summary	219
Module 2: Data Types and Variables	221
Overview	221
Learning Objectives	221
Chapters	221
Key Concepts Covered	221
Prerequisites	222
Tools and Technologies	222
Estimated Time to Complete	222
Assessment	222
Next Module	222
Additional Resources	223
Control Flow	224
Conditional Statements	225
Introduction	225
The if Statement	225
Basic Syntax	225
Simple if Statement	225
if-else Statement	226
if-else if-else Chain	226
The switch Statement	227
Basic Syntax	227
Simple switch Example	227
Intentional Fall-Through	228
Nested Conditionals	229
Nested Condition Pitfalls	229
The Dangling else (Else-Bind Ambiguity)	230
Conditional (Ternary) Expressions	231
Ternary Best Practices	232
Best Practices for Conditionals	233
1. Use Braces Consistently	233

2. Handle All Cases in switch	233
3. Prefer Early Returns Over Deep Nesting	234
4. Name Complex Conditions	234
Common Pitfalls	234
1. Assignment vs Comparison	234
2. Floating-Point Comparisons	235
3. Missing break in switch	235
Worked Example 1: Full Grading System	235
Worked Example 2: Leap Year (Rules + Nested Form)	237
Worked Example 3: Simple ATM Menu with switch	239
Short-Circuit Evaluation	241
Exercises	242
Exercise 1 — Triangle classifier	242
Exercise 2 — Rock–paper–scissors	243
Exercise 3 — BMI category	243
Exercise 4 — Rewrite dangling-else trap	243
Exercise 5 — ATM extension	243
Summary	243

Loops 245

Introduction	245
The while Loop	245
Basic Syntax	245
Simple while Loop	245
while with User Input	246
The do-while Loop	246
Basic Syntax	247
Simple do-while	247
Menu Pattern (classic do-while use)	247
The for Loop	248
Basic Syntax	248
Simple for Loop	248
for with Arrays	248
Decision Guide: for vs while vs do-while	249
Same Task Three Ways	249
When while Beats for	250
Loop Control: break and continue	251
break	251
continue	251
Nested Loops	252
Multiplication Table Pattern	252
Nested break/continue Scope	252
Nested Loop Patterns Cheatsheet	253
Infinite Loops (Servers, Embedded, Tools)	254
Forms	254
Pattern: Event Loop with break Conditions	254
Pattern: Embedded-Style Superloop	255
Pattern: Retry with Cap (avoid true infinite hang)	256

Worked Example 1: Prime Check	257
Worked Example 2: FizzBuzz	258
Worked Example 3: Multiplication Table	259
Worked Example 4: Running Average from Input	260
Performance Notes	261
Best Practices	262
Common Pitfalls	262
Infinite loop by forgetting the update	262
Off-by-one	262
Floating loop variable	263
More Practical Examples	263
Number Patterns	263
Array Stats and Reverse	264
Exercises	265
Exercise 1 — Collatz steps	265
Exercise 2 — Nested loop: prime pairs	265
Exercise 3 — Choose the right loop	265
Exercise 4 — FizzBuzz customization	265
Exercise 5 — Server-style loop	265
Exercise 6 — Running variance (optional stretch)	266
Summary	266
Advanced Control Flow	267
Introduction	267
The goto Statement	267
Basic Syntax	267
Simple goto Example	267
goto for Error Handling	268
goto in Switch Statements	269
Function Calls as Control Flow	270
Basic Function Calls	270
Recursive Function Calls	270
Function Pointers for Dynamic Control Flow	271
Error Handling Patterns	272
Return Code Pattern	272
errno and perror for System Errors	273
Setjmp and Longjmp for Non-Local Jumps	274
Advanced Loop Control Techniques	275
Loop Labels and Break to Outer Loops	275
State Machines	275
Best Practices for Advanced Control Flow	277
1. Use goto Sparingly and Structured	277
2. Avoid Deep Nesting	277
3. Use Function Pointers for Strategy Pattern	278
Common Pitfalls and How to Avoid Them	279
1. goto Leading to Spaghetti Code	279
2. Infinite Recursion	279
3. Longjmp Across Function Boundaries	280

Practical Examples	281
Resource Management with goto	281
Command-Line Parser with State Machine	282
Complex Error Handling System	284
Summary	286
Module 3: Control Flow	287
Overview	287
Learning Objectives	287
Chapters	287
Key Concepts Covered	287
Prerequisites	288
Tools and Technologies	288
Estimated Time to Complete	288
Assessment	288
Next Module	288
Additional Resources	288
Functions	289
Function Fundamentals	290
Introduction	290
What is a Function?	290
Function Definition	290
Basic Syntax	290
Simple Function Example	291
Function with No Parameters	291
Function with No Return Value (void)	291
Function Declaration (Prototypes)	292
Syntax	292
Function Prototype Example	292
Function Parameters	293
Parameter Passing by Value	293
Multiple Parameters	293
Return Statements	294
Returning Values	294
Early Returns	295
Function Scope and Lifetime	295
Local Variables	296
Static Local Variables	296
Function Naming Conventions	297
Common Naming Styles	297
Best Practices for Functions	297
1. Single Responsibility Principle	297
2. Appropriate Function Size	298
3. Meaningful Parameter Names	299

Common Pitfalls and How to Avoid Them	300
1. Missing Return Statements	300
2. Using Uninitialized Variables	300
3. Function Declaration vs Definition Mismatch	301
Practical Examples	301
Temperature Converter Functions	301
Mathematical Utility Functions	302
String Processing Functions	304
Summary	305
Advanced Functions	306
Introduction	306
Recursive Functions	306
Basic Recursion Concepts	306
Simple Recursion Example	306
Fibonacci Sequence	307
Recursive Array Processing	307
Variable Argument Functions	308
Basic Variable Argument Functions	308
Formatted Print Function	310
Inline Functions (C99)	311
Basic Inline Functions	311
Performance Considerations	311
Function Pointers	312
Basic Function Pointers	312
Function Pointer Arrays	313
Callback Functions	314
Higher-Order Functions	315
Function Factory Pattern	315
Best Practices for Advanced Functions	316
1. Recursion Best Practices	316
2. Variable Argument Functions Best Practices	317
3. Function Pointer Best Practices	318
Common Pitfalls and How to Avoid Them	319
1. Infinite Recursion	319
2. Variable Argument Function Errors	319
3. Function Pointer Errors	320
Practical Examples	321
Recursive Directory Traversal	321
Mathematical Expression Evaluator	322
Event Handler System	323
Summary	324
Modular Programming	325
Introduction	325
Header Files and Source Files	325
Header Files (.h)	325
Source Files (.c)	325

Example Layout	326
Include Guards	326
Traditional Include Guards	326
#pragma once (Non-standard, Widely Supported)	326
Compilation Units and Linking	326
Linkage: extern , static , and No Linkage	327
External Linkage (Default for File-Scope Functions)	327
Internal Linkage (static at File Scope)	328
static vs extern at a Glance	329
No Linkage	329
Complete Multi-Module Project: string_stats	329
Project Tree	329
strstats.h — Public Interface	330
strstats.c — Implementation	330
main.c — Driver	332
Manual Build	333
Full Makefile: Debug, Release, and -MMD Dependencies	333
Variant with build/ Directory	334
Creating Libraries	335
Static Library (.a)	335
Shared Library (.so on Linux / macOS)	335
Static vs Shared	336
Header Design Rules	336
Recommended Include Order	337
Forward Declarations	337
Opaque Pointers / Incomplete Types	338
vector.h	338
vector.c	338
main_vector.c	340
Alternate Multi-Module Example: 2D Vector Math	340
vec2.h	341
vec2.c	341
main_vec2.c	342
Callback Interfaces (Flexible Modules)	342
Best Practices	343
Exercises	343
Summary	344
Standard Library Functions	345
Introduction	345
String Functions (<string.h>)	345
String Copying	345
String Concatenation	345
String Comparison	346
String Searching	346
String Tokenization	346
String Length	346
Memory Functions (<string.h>)	346

Memory Copying	346
Memory Setting	346
Memory Comparison	347
Memory Searching	347
Mathematical Functions (<math.h>)	347
Basic Mathematical Functions	347
Hyperbolic Functions	348
Floating-Point Functions	348
Classification and Comparison Functions	348
Character Classification (<ctype.h>)	348
Character Testing	348
Character Conversion	349
Time Functions (<time.h>)	349
Time Types	349
Time Functions	349
Input/Output Functions (<stdio.h>)	350
File Operations	350
Formatted I/O	350
Character I/O	350
Direct I/O	350
File Positioning	351
Error Handling	351
Memory Management (<stdlib.h>)	351
Dynamic Memory Allocation	351
Utility Functions	351
Random Number Generation	352
Environment Functions	352
Searching and Sorting	352
Integer Arithmetic	352
Advanced Standard Library Features (C23)	353
New String Functions	353
Enhanced Memory Functions	353
Practical Examples	353
String Manipulation Example	353
Mathematical Functions Example	354
Time Functions Example	355
Memory Management Example	355
Full Programs: Standard Library in Practice	356
Program 1 — Safe string toolbox (strtool.c)	356
Program 2 — Math lab: vectors and roots (mathlab.c)	359
Program 3 — Character classifier report (ctype_report.c)	360
Program 4 — Time stamps and elapsed work (timelab.c)	361
Program 5 — qsort + bsearch + stdlib utilities (sortlab.c)	362
Program 6 — Mini word frequency with string.h + heap (wfreq.c)	364
Best Practices for Using Standard Library Functions	366
1. Error Checking	366
2. Buffer Safety	366
3. Memory Management	366

4. Locale Awareness	366
5. Thread Safety	367
6. Link math explicitly	367
Exercises	367
Summary	368
Module 4: Functions and Modular Programming	369
Overview	369
Learning Objectives	369
Chapters	369
Key Concepts Covered	369
Prerequisites	370
Tools and Technologies	370
Estimated Time to Complete	370
Assessment	370
Next Module	370
Additional Resources	371
 Arrays and Strings	 372
Arrays	373
Introduction	373
Array Declaration and Initialization	373
Basic Declaration	373
Initialization	373
Constant Arrays	374
Access and Indexing	374
No Automatic Bounds Checking	374
Traversal	374
sizeof and Array Length	375
Passing Arrays to Functions	375
Equivalent Parameter Declarations	375
Full Program: Size Parameter Pattern	376
Optional: Static Array Size Hints (C99+)	376
Multidimensional Arrays and Row-Major Layout	377
2D Declaration	377
Row-Major Storage	377
Passing 2D Arrays	377
Worked Algorithms (Full Programs)	378
1. Reverse an Array In Place	378
2. Find Minimum and Maximum	379
3. Frequency Count	380
4. Small Matrix Multiply ($2 \times 3 \cdot 3 \times 2 \rightarrow 2 \times 2$)	381
5. Histogram of Grades	382
6. Linear Search and Selection Sort (Together)	383
Common Bugs: Wrong vs Right	384
Off-by-One in Loops	384

Forgetting That Arrays Decay	385
Using Uninitialized Local Arrays	385
Copying with =	386
Returning a Pointer to a Local Array	386
Practical Examples	386
Temperature Analysis	386
Student Scores (Parallel Arrays)	387
Mini-Lab: Array Toolkit	387
Exercises	388
Summary	389
Strings	390
Introduction	390
String Representation	390
Character Arrays and Literals	390
Memory Layout	390
Literals Must Not Be Modified	391
Declaration Patterns	391
Standard Library Overview	391
Safe String Handling Patterns	391
Prefer <code>snprintf</code> over bare <code>strcpy</code> / <code>strcat</code>	391
Building Paths or Messages Safely	392
<code>strncpy</code> vs <code>snprintf</code> — Know the Trap	393
Bounded Copy Helper (Recommended Pattern)	393
Safe Concatenation	394
Input: Prefer <code>fgets</code> + Strip Newline	395
Implement <code>strlen</code> / <code>strcpy</code> / <code>strcmp</code> Yourself	395
Full Learning Program	395
Tokenizing and CSV Parsing	397
<code>strtok</code> vs <code>strtok_r</code>	397
Full Program: CSV Line with <code>strtok_r</code>	397
Full Program: Manual CSV Parse (No <code>strtok</code>)	398
Simple Key=Value Line	400
Buffer Overflows: Awareness and Safe Alternatives	401
Fragile Pattern (Do Not Ship)	401
Safe Alternatives Checklist	401
Full Program: Safe Line Reader and Copy	402
Compiler Help	403
More String Techniques	403
Remove Newline, Reverse, Case	403
Palindrome Check	404
Word Count	405
Character Class Statistics	405
Pitfalls Recap	406
Exercises	406
Summary	407

Advanced Arrays	408
Introduction	408
Multi-dimensional Arrays	408
Two-dimensional Arrays	408
Accessing Multi-dimensional Arrays	408
Multi-dimensional Array Functions	409
Three-dimensional and Higher-dimensional Arrays	410
Three-dimensional Arrays	410
Higher-dimensional Arrays	410
Variable-length Arrays (VLAs) - C99	410
Basic VLA Declaration	410
Multi-dimensional VLAs	411
VLA Function Parameters	412
Array of Arrays vs. Multi-dimensional Arrays	412
Array of Arrays	412
Multi-dimensional Arrays	413
Jagged Arrays	413
Creating Jagged Arrays	413
Jagged Array Functions	414
Array Algorithms and Operations	415
Matrix Operations	415
Dynamic Array Resizing	416
Array Performance Considerations	416
Memory Layout	416
Cache Performance	417
Advanced Array Techniques	418
Array Pointers	418
Array of Pointers	418
Compound Literals (C99)	418
Practical Examples	419
Sudoku Validator	419
Image Processing Example	420
Summary	422
Module 5: Arrays and Strings	423
Overview	423
Learning Objectives	423
Chapters	423
Key Concepts Covered	423
Prerequisites	424
Tools and Technologies	424
Estimated Time to Complete	424
Assessment	424
Next Module	424
Additional Resources	425

Pointers 426

Pointer Fundamentals 427

Introduction	427
What Is a Pointer?	427
Visual Step-by-Step: <code>&</code> and <code>*</code>	427
Step 1 — Create an <code>int</code>	427
Step 2 — Point at it	428
Step 3 — Read through the pointer	428
Step 4 — Write through the pointer	428
Full Demo Program	428
Declaration and Initialization	429
Null Pointers	429
Swap via Pointers (Full Program)	430
Why Values Alone Fail	430
Correct Swap with Pointers	430
<code>const</code> and Pointers: Three Combinations	431
Full Program	431
Why <code>const int *</code> Appears in APIs	432
Note on <code>const int *</code> vs <code>int const *</code>	432
<code>void *</code> with Careful Casting	432
Full Program: Generic Swap of Bytes	432
Rules of Thumb for <code>void *</code>	433
Pointer Arithmetic (Brief)	434
Pointers vs Arrays: Equivalence and Differences	434
Where They Look the Same	434
Where They Differ	434
Full Comparison Program	435
Out-Parameters: Return Status Separately from Data	435
Full Program	436
Design Tips	437
Pointers and Functions (More)	438
Returning a Pointer into a Caller-Owned Array	438
Unsafe: Returning Address of Local	438
Pointer Comparison and Size	439
Best Practices	439
Practical Examples	439
Pointer-Based String Reverse	439
Sum with Pointer Walking	440
Exercises	441
Summary	441

Pointers and Arrays 442

Introduction	442
Array-Pointer Equivalence	442
Arrays as Pointers	442
Important Distinction	442

Pointer Subscripting	443
Using Pointers as Arrays	443
Array Access Methods	443
Pointer Arithmetic with Arrays	444
Incrementing Pointers	444
Pointer Difference	444
Multi-dimensional Arrays and Pointers	444
Two-dimensional Arrays	444
Row-wise Traversal	445
Arrays as Function Parameters	446
Array Decay to Pointers	446
Multi-dimensional Array Parameters	446
Dynamic Arrays with Pointers	447
Allocating Dynamic Arrays	447
Resizing Dynamic Arrays	448
Pointer Arrays vs. Arrays of Pointers	449
Array of Pointers	449
Pointer to Array	450
Advanced Array-Pointer Techniques	450
Pointer Arithmetic for Multi-dimensional Arrays	450
Function Pointers for Array Operations	451
Practical Examples	452
Matrix Transpose Using Pointers	452
String Array Sorting	453
Dynamic 2D Array	453
Summary	455
Dynamic Memory Management	456
Introduction	456
Memory Layout in C Programs	456
Program Memory Sections	456
Stack vs. Heap	456
Dynamic Memory Allocation Functions	457
malloc()	457
calloc()	457
realloc()	458
free()	459
Memory Allocation Strategies	460
Single Element Allocation	460
Array Allocation	460
2D Array Allocation	461
Memory Management Best Practices	463
Error Checking	463
Preventing Dangling Pointers	463
Avoiding Memory Leaks	463
Memory Reallocation Safety	463
Common Memory Management Patterns	464
Dynamic String Handling	464

Dynamic Array with Resizing	465
Memory Debugging Tools	466
Valgrind (Linux/Unix)	466
AddressSanitizer (GCC/Clang)	467
Static Analysis Tools	467
Common Memory Errors and Solutions	467
Memory Leaks	467
Dangling Pointers	468
Buffer Overflows	468
Advanced Memory Management (C11 and Later)	468
aligned_alloc()	468
Practical Examples	469
Memory Pool Implementation	469
Smart Pointer-like Implementation	472
Summary	473
Advanced Pointer Concepts	474
Introduction	474
Pointers to Pointers	474
Basic Concept	474
Practical Applications	474
Function Pointers	476
Basic Declaration and Usage	476
Function Pointer Arrays	477
Callback Functions	477
Function Pointers as Parameters	478
Void Pointers	479
Basic Concept	479
Generic Functions with Void Pointers	480
Void Pointer Arrays	481
Const Correctness with Pointers	482
Pointer to Constant	482
Constant Pointer	482
Constant Pointer to Constant	483
Const Correctness in Function Parameters	483
Advanced Pointer Techniques	484
Pointer Arithmetic with Structs	484
Flexible Array Members (C99)	485
Restricted Pointers (C99)	486
Practical Examples	487
Generic Linked List Implementation	487
Event System with Function Pointers	489
Summary	491
Common Pointer Pitfalls	492
Introduction	492
Dangling Pointers	492
What Are Dangling Pointers?	492

Common Causes	492
Prevention Strategies	493
Memory Leaks	494
What Are Memory Leaks?	494
Common Causes	495
Prevention Strategies	496
Buffer Overflows	497
What Are Buffer Overflows?	497
Common Causes	497
Prevention Strategies	498
Uninitialized Pointers	499
What Are Uninitialized Pointers?	499
Common Causes	499
Prevention Strategies	500
Pointer Arithmetic Errors	502
Common Mistakes	502
Prevention Strategies	503
NULL Pointer Dereferencing	503
What Is NULL Pointer Dereferencing?	503
Common Causes	503
Prevention Strategies	504
Type Safety Issues	504
Common Problems	504
Prevention Strategies	505
Debugging Pointer Issues	506
Using Debugging Tools	506
Static Analysis Tools	506
Best Practices for Pointer Safety	506
1. Initialize Pointers	506
2. Check for NULL	507
3. Validate Array Bounds	507
4. Use Safe String Functions	507
5. Match Allocation with Deallocation	507
Practical Example: Safe Dynamic Array	508
Summary	510

Module 6: Pointers and Memory Management 511

Overview	511
Learning Objectives	511
Chapters	511
Key Concepts Covered	511
Prerequisites	512
Tools and Technologies	512
Estimated Time to Complete	512
Assessment	512
Next Module	513
Additional Resources	513

Structures 514

Structures 515

Introduction	515
Structure Declaration and Definition	515
Basic Structure Declaration	515
Structure Variable Declaration	515
Structure Initialization	516
Designated Initializers (C99)	516
Nested Structure Initialization	517
Accessing Structure Members	517
Dot Operator (.)	517
Arrow Operator (->)	518
Arrays of Structures	519
Declaration and Initialization	519
Dynamic Array of Structures	519
Nested Structures	520
Basic Nested Structures	520
Self-Referential Structures	521
Structure Assignment and Comparison	522
Structure Assignment	522
Structure Comparison	523
Passing Structures to Functions	523
Pass by Value	523
Pass by Reference (Pointer)	524
Structure Size and Memory Layout	525
Calculating Structure Size	525
Memory Alignment	526
Practical Examples	526
Student Management System	526
Point Structure with Operations	528
Summary	529

Unions and Enumerations 530

Introduction	530
Unions	530
Union Declaration and Basics	530
Union Memory Layout	531
Practical Union Usage	531
Anonymous Unions (C11)	533
Enumerations	534
Basic Enumeration Declaration	534
Explicit Enumeration Values	535
Enumeration with Negative Values	536
Enumeration Size and Type Specification (C23)	537
Combining Unions and Enumerations	538
Tagged Union Pattern	538
Configuration Structure with Unions	540

Practical Examples	542
Network Packet Header	542
Mathematical Expression Evaluator	544
Summary	547
Advanced Structure Concepts	548
Introduction	548
Bit Fields	548
Basic Bit Field Declaration	548
Practical Bit Field Usage	549
Bit Field Limitations and Considerations	551
Structure Packing and Alignment	552
Default Structure Alignment	552
Controlling Alignment with Pragma Pack	553
Using _Alignas (C11)	554
Flexible Array Members (C99)	554
Basic Flexible Array Member	554
Flexible Array in String Structure	555
Self-Referential Structures	557
Linked List Implementation	557
Binary Tree Implementation	558
Advanced Structure Techniques	560
Structure with Function Pointers (C Object-Oriented Approach)	560
Structure with Nested Anonymous Structures (C11)	562
Practical Examples	563
Configuration Manager with Advanced Structures	563
Memory Pool with Flexible Array Members	566
Summary	568
Abstract Data Types	570
Introduction	570
Information Hiding Principles	570
Separation of Interface and Implementation	570
Opaque Pointers	573
Creating Opaque Pointers	573
Constructor and Destructor Patterns	575
Resource Acquisition Is Initialization (RAII) in C	575
Generic Abstract Data Types	577
Using Void Pointers for Generic ADTs	577
Interface Design in C	582
Designing Clean APIs	582
Practical Examples	587
Complete ADT Implementation: Dynamic Array	587
Summary	592
Module 7: Structures and User-Defined Types	593
Overview	593
Learning Objectives	593

Chapters	593
Key Concepts Covered	593
Prerequisites	594
Tools and Technologies	594
Estimated Time to Complete	594
Assessment	594
Next Module	595
Additional Resources	595

File I/O 596

File Operations 597

File Pointers and the FILE Structure	597
Opening Files: <code>fopen</code>	597
Modes	597
Minimal open / close	598
Text vs Binary Mode (Caveats)	598
Text mode (" <code>r</code> ", " <code>w</code> ", ...)	598
Binary mode (" <code>rb</code> ", " <code>wb</code> ", ...)	599
Closing and Buffering	599
Error Handling: <code>fopen</code> , <code>ferror</code> , <code>errno</code>	600
Check every open	600
Distinguish EOF from error	600
Pattern: cleanup on partial failure	601
Full Program 1: Copy a File	602
Full Program 2: Simplified <code>wc</code> (Lines, Words, Bytes)	603
Full Program 3: Append Log Lines with Timestamps	606
File Size and Status Helpers	608
Best Practices	609
Exercises	610
Exercise 1 — Safe truncate-free write	610
Exercise 2 — Byte-at-a-time vs block copy	610
Exercise 3 — Extend <code>mywc</code>	610
Exercise 4 — Rotating log	610
Exercise 5 — Error path drill	610
Exercise 6 — Binary round-trip	610
Summary	611

File I/O Functions 612

Introduction	612
Character I/O Functions	612
<code>fgetc()</code> and <code>getc()</code>	612
<code>fputc()</code> and <code>putc()</code>	613
<code>ungetc()</code>	613
Examples	613
Line I/O Functions	614
<code>fgets()</code>	614

fputs()	614
Examples	615
Formatted I/O Functions	615
fprintf()	616
fscanf()	616
Examples	616
Binary I/O Functions	617
fwrite()	617
fread()	617
Examples	618
Practical Examples	619
CSV File Processor	619
Binary File Analyzer	620
Error Handling and Best Practices	622
Comprehensive Error Checking	622
Full Programs: Character, Line, Formatted, Binary	623
Program 1 — ROT13 filter (rot13.c)	623
Program 2 — Line numberer with long-line handling (nlnum.c)	624
Program 3 — CSV-ish records with fprintf / fscanf (scores_io.c)	625
Program 4 — Binary record dump / load (binrec.c)	627
Program 5 — ungetc simple number peek (peek_num.c)	629
Program 6 — Block copy with progress (bcopy.c)	631
Error Handling Patterns (Quick Reference)	632
Summary	633
Exercises	633

Advanced File Operations 635

Introduction	635
File Positioning	635
fseek()	635
ftell()	635
rewind()	636
Examples	636
Random Access Files	637
Structured Data Example	637
Temporary Files	640
tmpfile()	640
tmpnam()	640
mkstemp() (POSIX)	640
Examples	641
File Locking and Concurrent Access	642
flockfile() and funlockfile()	642
ftrylockfile()	642
Example with File Locking	642
File Status and Information	643
fileno()	644
fstat() (POSIX)	644
Example with File Information	644

Buffering Strategies	645
setvbuf()	645
Example with Custom Buffering	645
Practical Examples	646
File Merger Utility	646
File Splitter Utility	648
Summary	649
System-Level Programming	650
Introduction	650
Command-Line Arguments	650
argc and argv	650
Examples	651
Practical Command-Line Processor	651
Environment Variables	652
getenv()	653
Examples	653
Setting Environment Variables	653
Process Management	654
exit()	654
atexit()	654
Examples	655
Process Exit Status	656
Signal Handling	656
signal()	656
Common Signals	657
Examples	657
Advanced Signal Handling	658
System Calls and Process Creation	659
system()	659
Examples	660
Practical Examples	660
Configuration File Processor	660
Process Monitor	663
Full Programs: CLI, Env, Signals, Processes	664
Program 1 — Robust argument parser (<code>opts.c</code>)	664
Program 2 — Environment-driven config (<code>envcfg.c</code>)	666
Program 3 — <code>atexit</code> cleanup and <code>exit</code> status (<code>cleanup_demo.c</code>)	667
Program 4 — Graceful SIGINT shutdown (<code>sigint_loop.c</code>)	668
Program 5 — <code>fork</code> + <code>exec</code> + <code>wait</code> mini runner (<code>run_cmd.c</code>)	669
Program 6 — Echo lines with PID and env prefix (<code>line_echo.c</code>)	671
Exercises	671
Summary	672
Module 8: File I/O and System Programming	673
Overview	673
Learning Objectives	673
Chapters	673

Key Concepts Covered	673
Prerequisites	674
Tools and Technologies	674
Estimated Time to Complete	674
Assessment	674
Next Module	675
Additional Resources	675

Modern C 676

C99 Features 677

Introduction	677
Variable-Length Arrays (VLAs)	677
Basic VLA Declaration	677
Multi-dimensional VLAs	678
VLA Function Parameters	678
Designated Initializers	679
Array Designated Initializers	679
Structure Designated Initializers	680
Compound Literals	681
Array Compound Literals	681
Structure Compound Literals	682
New Data Types	683
long long Type	683
_Bool Type	683
Inline Functions	684
restrict Keyword	685
Flexible Array Members	686
Practical Examples	687
Matrix Library Using C99 Features	687
Summary	690

C11 Features 691

Introduction	691
Static Assertions	691
_Static_assert	691
Practical Static Assertion Examples	692
Anonymous Structures and Unions	693
Anonymous Structures	693
Anonymous Unions	693
Thread-Local Storage	694
Basic Thread-Local Storage	695
Thread-Local Storage with Initialization	696
Bounds-Checking Functions (Annex K)	697
Safe String Functions	697
Safe Memory Functions	698

Generic Selection	699
Basic Generic Selection	699
Advanced Generic Selection	700
Alignment Support	701
Alignment Specifier	702
Practical Examples	702
Thread-Safe Logger with Static Assertions	702
Summary	704
C17 and C23 Features	706
Introduction	706
C17: Technical Corrigendum	706
Defect Fixes	706
Example of C17 Improvements	706
C23: typeof Operators	707
Basic typeof Usage	707
typeof with Complex Expressions	708
typeof_unqual for Unqualified Types	708
Binary Literals	709
Basic Binary Literals	709
Bit Manipulation with Binary Literals	709
Enhanced Preprocessor Directives	710
#ifdef and #ifndef	711
#warning Directive	711
#embed Directive (Conceptual)	712
Attribute Syntax	712
Standard Attributes	712
Extended Integer Types	713
__BitInt Example (Conceptual)	713
Enhanced __Static_assert	714
Enhanced Static Assertions	714
New Standard Library Functions	715
strdup and strndup	715
Enhanced aligned_alloc	716
Practical Examples	716
Modern Configuration System	716
Summary	718
Modern C Development Practices	720
Introduction	720
DevOps and C: CI/CD Pipelines	720
GitHub Actions for C Projects	720
GitLab CI for C Projects	721
Containerized Development	722
Docker for C Development	722
Multi-stage Docker Build	723
Docker Compose for Development	724

Cross-compilation	725
CMake Cross-compilation	725
Cross-compilation Toolchain File	726
WebAssembly (WASM)	727
Emscripten Setup	727
Simple C to WASM Example	727
HTML Integration	728
Cloud-native C Applications	728
Simple HTTP Server in C	728
Dockerfile for Cloud Deployment	731
Supply Chain Security	731
Conan Package Manager	732
Vulnerability Scanning	732
Modern Testing Practices	733
Unity Testing Framework	733
Property-Based Testing with CMocka	734
Fuzzing with AFL	735
Documentation as Code	737
Doxygen Configuration	737
Example Documentation	738
Practical Examples	739
Modern C Project Structure	739
Complete CMakeLists.txt	739
Summary	740
Advanced Preprocessor	742
Introduction	742
Macro Definitions and Function-like Macros	742
Advanced Macro Techniques	742
Generic Programming with Macros	743
Conditional Compilation	745
Platform-Specific Compilation	745
Version-Specific Code	747
Pragma Directives	748
Common Pragma Directives	749
Compiler-Specific Pragmas	750
Predefined Macros	752
System and Compiler Information	752
Token Manipulation and Advanced Techniques	754
Deferred Macro Expansion	754
X-Macros Pattern	755
C23 Enhanced Preprocessor Features	757
<code>#elifdef</code> and <code>#elifndef</code>	757
<code>#warning</code> Directive	758
Digit Separators in Preprocessor	758
Practical Examples	759
Configuration System Using Preprocessor	759
Advanced Error Handling with Preprocessor	761

Summary	763
Module 9: Advanced C Features and Modern Standards	765
Overview	765
Learning Objectives	765
Chapters	765
Key Concepts Covered	766
Prerequisites	766
Tools and Technologies	766
Estimated Time to Complete	767
Assessment	767
Next Module	767
Additional Resources	767
 Data Structures	 768
Linear Data Structures	769
Introduction	769
Dynamic Arrays (Vectors)	769
Basic Dynamic Array Implementation	769
Generic Dynamic Array	772
Linked Lists	774
Singly Linked List	775
Doubly Linked List	779
Stacks	782
Array-based Stack	782
Linked List-based Stack	784
Queues	787
Circular Buffer Queue	787
Linked List-based Queue	790
Practical Examples	793
Expression Evaluation Using Stacks	793
Summary	796
 Tree Structures	 797
Introduction	797
Binary Trees	797
Basic Binary Tree Implementation	797
Binary Search Trees (BST)	801
BST Implementation	801
AVL Trees	805
AVL Tree Implementation	806
Red-Black Trees	812
Red-Black Tree Properties	812
B-Trees	817
B-Tree Implementation (Simplified)	817

Practical Examples	821
File System Directory Tree	821
Summary	824
Hash Tables	826
Introduction	826
Hash Function Design	826
Basic Hash Functions	826
Collision Resolution Strategies	828
Chaining	828
Open Addressing	832
Dynamic Hash Tables	837
Resizable Hash Table with Chaining	837
Hash Table Performance Analysis	843
Load Factor and Performance	843
Practical Examples	847
Symbol Table Implementation	847
Cache Implementation	852
Summary	857
Graph Algorithms	858
Introduction	858
Graph Representations	858
Adjacency Matrix	858
Adjacency List	860
Graph Traversal Algorithms	864
Breadth-First Search (BFS)	864
Depth-First Search (DFS)	868
Shortest Path Algorithms	872
Dijkstra's Algorithm	872
Minimum Spanning Tree Algorithms	879
Prim's Algorithm	879
Practical Examples	886
Social Network Analysis	886
Summary	891
Sorting and Searching	893
Introduction	893
Comparison-Based Sorting Algorithms	893
Bubble Sort	893
Selection Sort	894
Insertion Sort	895
Merge Sort	896
Quick Sort	898
Heap Sort	900
Non-Comparison Sorting Algorithms	902
Counting Sort	902
Radix Sort	903

Bucket Sort	905
Searching Algorithms	907
Linear Search	907
Binary Search	908
Interpolation Search	910
Advanced Searching Techniques	911
Exponential Search	911
Jump Search	912
Selection Algorithms	914
Quickselect (Finding kth Smallest Element)	914
Performance Comparison and Analysis	916
Sorting Algorithm Performance Test	916
Practical Examples	922
Phone Book Implementation	922
Summary	925
Module 10: Data Structures and Algorithms	927
Overview	927
Learning Objectives	927
Chapters	927
Key Concepts Covered	927
Prerequisites	928
Tools and Technologies	928
Estimated Time to Complete	929
Assessment	929
Next Module	929
Additional Resources	929
Network	930
Network Fundamentals	931
Computer Network Basics	931
Topologies (conceptual)	931
Network types by scope	931
OSI and TCP/IP Models	932
OSI (7 layers)	932
TCP/IP (practical 4 layers)	932
IP Addressing and Ports	932
IPv4	932
IPv6 (awareness)	933
Ports	933
TCP vs UDP (what sockets expose)	933
Sockets and Address Structures	933
Headers	933
Structures	934
Endianness: Why Byte Order Matters	934
Program: inspect endianness and multi-byte layout	934

htons / htonl / ntohs / ntohl	936
Demo program	936
inet_pton and inet_ntop	937
Full conversion + sockaddr_in setup demo	937
Filling sockaddr_in checklist	939
Minimal UDP Echo: Server and Client	939
Server (udp_echo_server.c)	939
Client (udp_echo_client.c)	941
Build and run (two terminals)	942
Optional: see the packets	943
Network Errors (practical)	943
Hostname Lookup (brief, modern)	944
Protocol Quick Reference	944
Exercises	944
Exercise 1 — Endianness report	944
Exercise 2 — Port round-trip	944
Exercise 3 — inet_pton validation	945
Exercise 4 — Bind failure	945
Exercise 5 — Echo timeout	945
Exercise 6 — Multi-message server log	945
Exercise 7 — Binary payload	945
Summary	945
Socket Programming	946
Introduction	946
Socket Basics	946
Socket Creation	946
Example Socket Creation	946
Socket Address Structures	947
IPv4 Address Structure	947
Generic Socket Address Structure	947
Server-Side Socket Programming	947
Binding a Socket	948
Listening for Connections	948
Accepting Connections	948
Complete TCP Server Example	948
Client-Side Socket Programming	950
Connecting to a Server	950
Complete TCP Client Example	950
UDP Socket Programming	951
UDP Server Example	951
UDP Client Example	953
Advanced Socket Options	954
Socket Options	954
Common Socket Options	954
Example: Setting Socket Options	954
Error Handling in Socket Programming	955
Common Socket Errors	955

Robust Error Handling Example	955
Practical Examples	956
Simple Echo Server	956
Simple Echo Client	958
Summary	960
Advanced Networking	961
Introduction	961
Non-blocking I/O	961
Setting Non-blocking Mode	961
Non-blocking Socket Example	961
I/O Multiplexing	963
select()	963
Example: Using select() for Multiple Clients	963
poll()	966
Example: Using poll()	966
Concurrent Server Architectures	969
Threading Approach	969
Forking Approach	971
Advanced Socket Options	974
Socket Buffer Sizes	974
Timeout Options	975
Signal-Driven I/O	976
Practical Examples	977
High-Performance Echo Server	977
Summary	981
Network Protocols	982
Introduction	982
HTTP Protocol Implementation	982
HTTP Request Structure	982
HTTP Response Structure	982
Simple HTTP Client	983
Simple HTTP Server	985
FTP Protocol Implementation	988
FTP Client Implementation	988
DNS Protocol Implementation	991
Simple DNS Resolver	991
SMTP Protocol Implementation	994
Practical Examples	998
Protocol Analyzer	998
Summary	999
Network Security	1001
Introduction	1001
Security Fundamentals	1001
CIA Triad	1001
Common Security Threats	1001

Cryptography Basics	1001
Symmetric Encryption	1002
Hash Functions	1002
Secure Socket Layer (SSL/TLS)	1003
Installing OpenSSL	1003
Simple HTTPS Client with OpenSSL	1004
Authentication and Authorization	1007
Simple Authentication System	1007
Secure Coding Practices	1009
Input Validation	1009
Buffer Overflow Prevention	1011
Firewall and Access Control	1012
Simple Packet Filter	1012
Secure Communication Patterns	1014
Mutual Authentication	1014
Practical Examples	1016
Secure File Transfer	1016
Summary	1019
Module 11: Network Programming	1021
Overview	1021
Learning Objectives	1021
Chapters	1021
Prerequisites	1021
Key Concepts	1022
Embedded	1023
Embedded Basics	1024
Introduction	1024
What Are Embedded Systems?	1024
Characteristics	1024
Types (rough scale)	1025
Microcontroller vs Microprocessor	1025
Typical Components	1025
Development Environments and Cross-Compilation	1025
High-Level Toolchain Notes (<code>arm-none-eabi-gcc</code>)	1025
Other Tools You Will Meet	1026
Memory Organization	1027
Startup Sketch (bare metal idea)	1027
Section Placement Reminders	1027
Bit Manipulation Helpers	1028
<code>volatile</code> — Why Embedded Needs It	1029
Packed Structs	1031
Memory-Mapped Register Pattern as Struct Overlay	1032
Software Model: Button Debounce State Machine	1034

Classic Patterns (Still Useful)	1036
LED Blink Superloop (conceptual MMIO)	1036
Interrupt Sketch	1037
Timing	1037
Practical Host Lab: Mini “Firmware” Module	1037
Exercises	1039
Exercise 1 — Bit library	1039
Exercise 2 — Packed vs unpacked	1039
Exercise 3 — Overlay UART	1039
Exercise 4 — Debounce FSM extension	1039
Exercise 5 — Cross-compile dry run	1039
Exercise 6 — Volatile experiment	1040
Summary	1040
Low-Level Programming	1041
Introduction	1041
Memory-Mapped I/O	1041
Direct Register Access	1041
Bit Manipulation Techniques	1042
Advanced Bit Operations	1043
Bit Fields	1043
Inline Assembly	1044
Basic Inline Assembly	1044
Advanced Inline Assembly	1045
Memory Management	1046
Custom Memory Allocator	1046
Memory Sections and Linker Scripts	1047
Hardware Abstraction Layer (HAL)	1047
GPIO HAL Implementation	1047
Interrupt Management	1051
Nested Vectored Interrupt Controller (NVIC)	1051
Practical Examples	1052
UART Driver Implementation	1052
Summary	1054
Embedded C Techniques	1055
Introduction	1055
Compiler Extensions and Attributes	1055
GCC Attributes	1055
IAR and Keil Extensions	1056
Optimization Techniques	1057
Compiler Optimization Levels	1057
Manual Optimization Techniques	1058
Fixed-Point Arithmetic	1059
Volatile Keyword and Memory-Mapped I/O	1060
Interrupt Service Routines	1060
Memory-Constrained Programming	1061
Stack Management	1062

Code Size Optimization	1063
Real-Time Considerations	1064
Timing Analysis	1064
Hardware-Specific Features	1065
ARM Cortex-M Specific Features	1065
Practical Examples	1066
Efficient Sensor Data Processing	1066
Power Management	1068
Summary	1069
Real-time Programming	1070
Introduction	1070
Real-time System Characteristics	1070
Types of Real-time Systems	1070
Real-time Requirements	1070
Task Scheduling	1071
Rate Monotonic Scheduling (RMS)	1071
Earliest Deadline First (EDF) Scheduling	1072
Timing and Synchronization	1073
Timer Management	1073
Semaphore Implementation	1075
Interrupt Handling in Real-time Systems	1077
Interrupt Latency Management	1078
Resource Management	1080
Memory Pool Management	1080
Priority Inheritance	1081
Timing Analysis and Deadline Monitoring	1083
Deadline Monitoring	1083
Practical Examples	1086
Real-time Control System	1086
Real-time Data Acquisition System	1088
Summary	1091
Hardware Interfaces	1092
Introduction	1092
GPIO (General Purpose Input/Output)	1092
GPIO Configuration and Control	1092
SPI (Serial Peripheral Interface)	1096
SPI Driver Implementation	1096
I2C (Inter-Integrated Circuit)	1100
I2C Driver Implementation	1100
UART (Universal Asynchronous Receiver-Transmitter)	1108
UART Driver Implementation	1108
Practical Examples	1112
Sensor Interface Example	1112
Summary	1117

Module 12: Embedded Systems Programming	1118
Overview	1118
Learning Objectives	1118
Chapters	1118
Prerequisites	1118
Key Concepts	1119
 Performance	 1120
Performance Analysis	1121
Introduction	1121
Performance Metrics	1121
Time Complexity Metrics	1121
Memory Usage Metrics	1122
CPU Usage Metrics	1123
Profiling Tools	1124
GNU gprof	1124
Valgrind Callgrind	1125
perf (Linux Performance Events)	1126
Custom Profiling Framework	1126
Benchmarking Framework	1129
Cache Performance Analysis	1132
Memory Access Pattern Analysis	1133
Practical Examples	1135
Comprehensive Performance Analysis Tool	1135
Summary	1138
 Optimization Techniques	 1140
Introduction	1140
Algorithmic Optimization	1140
Time Complexity Analysis	1140
Search Algorithm Optimization	1143
Loop Optimization	1146
Loop Unrolling	1146
Loop Fusion and Fission	1148
Data Structure Optimization	1150
Array vs Linked List	1150
Compiler Optimization Techniques	1153
Compiler Flags and Pragmas	1153
Mathematical Optimization	1155
Fast Mathematical Operations	1155
Practical Examples	1157
Image Processing Optimization	1157
Summary	1160
 Memory Optimization	 1161
Introduction	1161

Memory Hierarchy and Performance	1161
Memory Latency Comparison	1161
Cache Behavior Impact	1161
Memory Allocation Optimization	1163
Custom Memory Allocators	1163
Stack Allocation vs Heap Allocation	1166
Data Structure Optimization	1167
Structure of Arrays vs Array of Structures	1167
Memory Alignment	1169
Memory Layout Optimization	1170
Data Locality	1170
Memory Pooling	1172
Memory Access Patterns	1175
Sequential vs Random Access	1175
Practical Memory Optimization Techniques	1177
Reducing Memory Fragmentation	1177
Memory Prefetching	1178
Memory Debugging and Analysis	1179
Simple Memory Tracker	1180
Summary	1183
Advanced Optimization	1184
Introduction	1184
Compiler Optimization	1184
Optimization Levels	1184
Compiler Hints and Pragmas	1185
Profile-Guided Optimization (PGO)	1186
Vectorization	1186
Manual Vectorization with Intrinsics	1186
Auto-Vectorization	1189
Parallel Processing	1190
OpenMP for Shared-Memory Parallelism	1190
Thread-Level Parallelism with Pthreads	1192
Architecture-Specific Optimizations	1194
Cache Optimization Techniques	1194
Branch Prediction Optimization	1196
Memory-Level Parallelism	1197
Advanced Profiling and Analysis	1198
Optimization Decision Making	1199
Practical Optimization Workflow	1200
Summary	1201
Module 13: Performance Optimization	1202
Overview	1202
Learning Objectives	1202
Chapters	1202
Prerequisites	1202
Key Concepts	1203

Testing	1204
Testing Methodologies	1205
Why Testing Matters	1205
Types of Testing (C-focused)	1205
Unit testing characteristics	1205
Style 1: <code>assert.h</code> Unit Tests	1206
Style 2: Minimal Custom <code>TEST</code> Macro Harness	1206
Style 3: Unity-Style Suite That Actually Builds	1209
Option A — Self-contained “unity-like” runner (no download)	1209
Option B — Real Unity (vendored)	1211
Test-Driven Development (TDD) Walkthrough: <code>is_leap_year()</code>	1212
Step 0 — Scaffold (no real logic yet)	1212
Step 1 — RED: first failing test	1212
Step 2 — GREEN: minimal code that passes	1213
Step 3 — RED: century rule	1213
Step 4 — GREEN: handle centuries	1213
Step 5 — RED: year 2000 must be leap	1214
Step 6 — GREEN: full Gregorian rule	1214
Step 7 — REFACTOR: keep tests green	1214
Alternate TDD target: <code>clamp()</code>	1214
Boundary Value Testing	1215
Integers	1215
Strings	1216
Table-Driven Tests in C	1217
Code Coverage with <code>gcov</code> / <code>lcov</code>	1218
Instrument and run	1218
Sanitizers: Address and Undefined Behavior	1219
AddressSanitizer (ASan)	1219
UndefinedBehaviorSanitizer (UBSan)	1219
Combined (common for tests)	1220
Integration and System Testing (brief)	1220
Black-Box vs White-Box (quick map)	1220
FIRST and AAA	1221
Best Practices	1221
Exercises	1221
Exercise 1 — Custom harness for <code>clamp</code>	1221
Exercise 2 — TDD <code>int max3(int a, int b, int c)</code>	1222
Exercise 3 — Table-driven string trim	1222
Exercise 4 — Coverage gap hunt	1222
Exercise 5 — Sanitizer lab	1222
Exercise 6 — Boundary design	1222
Summary	1222
Debugging Techniques	1224
The Debugging Process	1224
GDB: Session-Style Walkthroughs	1224
Program under study	1225

Transcript 1 — break, run, print, next, continue	1225
Transcript 2 — backtrace on a crash	1226
Transcript 3 — watchpoints	1227
Handy GDB extras	1228
printf Debugging and DEBUG_LOG Macros	1228
Rubber duck and binary-search debugging	1230
Valgrind Memcheck: Intentional Leak and Reading the Report	1231
Example program with a real leak and a bad read	1231
How to read typical output	1232
Common Memcheck flags	1232
strace and ltrace (brief)	1233
strace — system calls	1233
ltrace — library calls	1233
Core Dumps: Capture and Analyze	1233
Enable dumps	1233
Produce a core	1234
Analyze with GDB	1234
Common Failure Modes	1235
Segmentation faults	1235
Logic errors	1235
Memory leaks and corruption	1235
Race conditions (threads)	1235
Static Analysis (quick)	1236
Debugging Best Practices	1236
Exercises	1236
Exercise 1 — GDB fix of <code>sum_range</code>	1236
Exercise 2 — <code>DEBUG_LOG</code> toggle	1236
Exercise 3 — Valgrind leak	1237
Exercise 4 — Invalid write	1237
Exercise 5 — Core dump	1237
Exercise 6 — <code>strace</code> detective	1237
Summary	1237
Error Handling	1239
Introduction to Error Handling	1239
Why Error Handling is Important	1239
Error Handling Strategies in C	1239
Return Codes	1239
Global Error Variables	1241
Advanced Error Handling Techniques	1243
Setjmp and Longjmp	1243
Resource Acquisition Is Initialization (RAII)	1244
Error Handling Best Practices	1246
Check All Return Values	1246
Use Assertions for Debugging	1246
Provide Meaningful Error Messages	1247
Implement Error Logging	1248

Error Handling in Different Contexts	1249
File I/O Error Handling	1249
Memory Management Error Handling	1251
Conclusion	1252
Code Quality	1253
Introduction to Code Quality	1253
Dimensions of Code Quality	1253
Coding Standards and Style Guidelines	1253
Naming Conventions	1254
Code Formatting	1254
Commenting Best Practices	1254
Code Review Practices	1255
Benefits of Code Reviews	1255
Code Review Checklist	1255
Example Code Review	1255
Static Analysis Tools	1256
GCC with Warnings	1256
Cppcheck	1256
Clang Static Analyzer	1257
Code Formatting and Style Tools	1257
Clang-Format	1257
AStyle (Artistic Style)	1257
Test Coverage Analysis	1257
GCC with Coverage	1258
Coverage Metrics	1258
Performance Analysis	1258
Profiling with gprof	1258
Profiling with perf	1259
Security Best Practices	1259
Input Validation	1259
Memory Safety	1260
Secure Coding Practices	1260
Documentation and Code Comments	1260
Self-Documenting Code	1261
API Documentation	1261
Refactoring and Code Improvement	1262
Code Smells to Avoid	1262
Refactoring Example	1262
Full Programs: Quality Gates You Can Run Locally	1264
Program 1 — Strict-warning “clean” module (clamp_lib)	1264
Program 2 — Warning archaeology (smell.c → clean.c)	1265
Program 3 — Complexity reduction: nested ifs → early return	1266
Program 4 — Static assertions and size contracts	1268
Program 5 — Local quality script (check.sh + target)	1268
Program 6 — Review checklist applied to a PR-sized change	1269
Continuous Integration and Quality Gates	1271
CI Pipeline Components	1271

Example CI Configuration	1271
Local gate CI gate	1272
Exercises	1272
Conclusion	1273
Module 14: Testing	1274
Learning goals	1274
Chapters (map)	1274
Suggested order	1275
Tooling cheat sheet	1275
How this module connects	1276
Module self-check	1276
Estimated time	1276
Next	1276
Advanced	1277
Module 15.1: Multithreading in C	1278
Introduction	1278
Learning Objectives	1278
1. Fundamentals of Multithreading	1278
1.1 Processes vs Threads	1278
1.2 Thread Lifecycle	1278
1.3 Thread Attributes	1279
2. Creating and Managing Threads	1279
2.1 Basic Thread Creation	1279
2.2 Thread Arguments and Return Values	1280
3. Thread Synchronization	1281
3.1 Race Conditions	1281
3.2 Mutexes	1281
3.3 Condition Variables	1282
4. Advanced Threading Concepts	1284
4.1 Thread-Specific Data	1284
4.2 Thread Cancellation	1285
5. Performance Considerations	1286
5.1 Thread Pool Implementation	1286
6. Debugging Multithreaded Applications	1289
6.1 Common Issues	1289
6.2 Debugging Tools	1289
7. Best Practices	1289
7.1 Design Principles	1289
7.2 Performance Optimization	1290
Summary	1290
Module 15.2: System Programming in C	1291
Introduction	1291
Learning Objectives	1291

1. System Calls and APIs	1291
1.1 What are System Calls?	1291
1.2 Common System Call Categories	1291
1.3 Error Handling in System Calls	1291
2. Process Management	1292
2.1 Process Creation with fork()	1292
2.2 Process Execution with exec()	1293
2.3 Process Groups and Sessions	1293
3. Advanced File System Operations	1294
3.1 File Metadata with stat()	1294
3.2 Directory Operations	1295
3.3 Hard Links and Symbolic Links	1296
4. Inter-Process Communication (IPC)	1297
4.1 Pipes	1297
4.2 Named Pipes (FIFOs)	1299
4.3 Shared Memory	1300
5. Signal Handling	1301
5.1 Basic Signal Handling	1301
5.2 Advanced Signal Handling with sigaction()	1302
6. Memory Management	1303
6.1 Memory Mapping with mmap()	1303
6.2 Memory Protection	1304
7. Security Considerations	1306
7.1 Secure Programming Practices	1306
7.2 Example of Secure File Handling	1306
8. Best Practices	1307
8.1 Error Handling	1307
8.2 Resource Management	1307
8.3 Portability	1307
Summary	1308

Module 15.3: Domain-Specific Applications in C 1309

Introduction	1309
Learning Objectives	1309
1. Characteristics of Domain-Specific Applications	1309
1.1 Domain Requirements	1309
1.2 Design Considerations	1309
2. Embedded Systems Applications	1310
2.1 Microcontroller Programming	1310
2.2 Real-Time Operating System Integration	1311
3. Scientific Computing Applications	1314
3.1 Numerical Libraries Integration	1314
3.2 Fast Fourier Transform (FFT)	1318
4. Networking Applications	1321
4.1 Protocol Implementation	1321
4.2 HTTP Client Implementation	1325
5. Best Practices for Domain-Specific Applications	1331
5.1 Domain Analysis	1331

5.2 Design Patterns	1331
5.3 Testing Strategies	1331
Summary	1332
Module 15.4: Capstone Project	1333
Introduction	1333
Learning Objectives	1333
Project Requirements	1333
1. Technical Requirements	1333
2. Project Scope	1334
3. Deliverables	1334
Project Ideas	1334
1. System Monitoring Tool	1334
2. Custom Database Engine	1339
3. Network File Transfer System	1343
Project Development Process	1351
1. Planning Phase	1351
2. Implementation Phase	1352
3. Testing and Debugging Phase	1352
4. Deployment and Presentation Phase	1352
Evaluation Criteria	1352
Technical Excellence (40%)	1352
Design and Architecture (25%)	1353
Documentation (20%)	1353
Testing (10%)	1353
Presentation (5%)	1353
Best Practices	1353
1. Code Quality	1353
2. Project Management	1353
3. Collaboration	1354
Resources and Support	1354
Tools and Libraries	1354
Learning Resources	1354
Conclusion	1354
Workbook	1355
Examples	1356
Foundations	1357
Data Types	1358
Control Flow	1359

Functions	1360
Arrays and Strings	1361
Pointers	1362
Structures	1363
File I/O	1364
Modern C	1365
Data Structures	1366
Network	1367
Embedded	1368
Performance	1369
Testing	1370
Advanced	1371
Exercises	1372
Foundations	1373
Module 1: Foundations Exercises	1374
Exercise 1: Basic Program Structure	1374
Exercise 2: Personal Information Display	1374
Exercise 3: Simple Calculator Input	1374
Exercise 4: Character Analysis	1374
Exercise 5: Temperature Converter	1375
Exercise 6: String Formatting Practice	1375
Exercise 7: Input Buffer Management	1375
Exercise 8: Comprehensive Demo Program	1375

Solutions and Tips	1376
Exercise 1 Solution Example:	1376
Exercise 3 Solution Example:	1376
Common Pitfalls to Avoid:	1377
Compilation Tips:	1377
Data Types	1378
Module 2: Data Types and Variables Exercises	1379
Exercise 1: Data Type Explorer	1379
Exercise 2: Variable Declaration Challenge	1379
Exercise 3: Constants Comparison	1379
Exercise 4: Storage Class Demonstrator	1380
Exercise 5: Type Conversion Safety	1380
Exercise 6: Operator Expression Builder	1380
Exercise 7: Bit Manipulation Toolkit	1380
Exercise 8: Data Type Selection Advisor	1381
Exercise 9: Comprehensive Calculator	1381
Exercise 10: Real-World Application	1381
Solutions and Tips	1381
Exercise 1 Solution Example:	1381
Exercise 3 Solution Example:	1382
Exercise 6 Solution Example:	1383
Exercise 7 Solution Example:	1383
Common Pitfalls to Avoid	1384
Compilation Tips	1385
Control Flow	1386
Module 3: Control Flow Exercises	1387
Exercise 1: Grade Calculator with Multiple Conditions	1387
Exercise 2: Pattern Generator	1387
Exercise 3: Menu-Driven Calculator	1387
Exercise 4: Prime Number Analyzer	1387
Exercise 5: Resource Manager with goto	1388
Exercise 6: Function Pointer Calculator	1388
Exercise 7: State Machine Implementation	1388
Exercise 8: Error Handling Framework	1388
Exercise 9: Complex Loop Optimization	1388
Exercise 10: Comprehensive Control Flow Application	1389
Solutions and Tips	1389
Exercise 1 Solution Example:	1389
Exercise 2 Solution Example:	1391
Exercise 3 Solution Example:	1392
Exercise 5 Solution Example:	1395
Common Pitfalls to Avoid	1397

Compilation Tips	1397
Functions	1398
Module 4: Functions and Modular Programming Exercises	1399
Exercise 1: Basic Function Implementation	1399
Exercise 2: Recursive Functions	1399
Exercise 3: Variable Argument Functions	1399
Exercise 4: Function Pointers and Callbacks	1400
Exercise 5: Modular Programming with Multiple Files	1400
Exercise 6: Inline Functions and Macro Functions	1400
Exercise 7: Standard Library Function Practice	1400
Exercise 8: Advanced Modular Programming	1401
Solutions and Tips	1401
Exercise 1 Solution Example:	1401
Exercise 2 Solution Example:	1403
Exercise 3 Solution Example:	1405
Exercise 4 Solution Example:	1407
Common Pitfalls to Avoid	1410
Compilation Tips	1410
Arrays and Strings	1411
Module 5: Arrays and Strings Exercises	1412
Exercise 1: Array Manipulation Functions	1412
Exercise 2: Matrix Operations	1412
Exercise 3: String Processing Library	1412
Exercise 4: Dynamic Array Implementation	1413
Exercise 5: Multi-dimensional Array Challenges	1413
Exercise 6: String Search and Replace	1413
Exercise 7: Command Line Argument Processing	1413
Exercise 8: Advanced Array and String Integration	1414
Solutions and Tips	1414
Exercise 1 Solution Example:	1414
Exercise 3 Solution Example:	1415
Common Pitfalls to Avoid:	1416
Performance Tips:	1416
Pointers	1418
Module 6: Pointers Exercises	1419
Exercise 1: Basic Pointer Operations	1419
Exercise 2: Pointers and Arrays	1419
Exercise 3: Dynamic Memory Allocation	1419
Exercise 4: Pointer to Pointers	1420
Exercise 5: Function Pointers	1420

Exercise 6: Advanced Pointer Techniques	1420
Exercise 7: Pointer Debugging and Safety	1420
Exercise 8: Comprehensive Pointer Application	1421
Solutions and Tips	1421
Exercise 1 Solution Example:	1421
Exercise 3 Solution Example:	1422
Common Pitfalls to Avoid:	1423
Best Practices:	1424
Structures	1425
Module 7: Structures, Unions, and Enums Exercises	1426
Exercise 1: Basic Structure Implementation	1426
Exercise 2: Array of Structures	1426
Exercise 3: Linked List Implementation	1426
Exercise 4: Unions and Bit Fields	1427
Exercise 5: Enumerations and Type Definitions	1427
Exercise 6: Advanced Structure Techniques	1427
Exercise 7: Structure Memory Management	1427
Exercise 8: Comprehensive Structure Application	1428
Solutions and Tips	1428
Exercise 1 Solution Example:	1428
Exercise 3 Solution Example:	1429
Common Pitfalls to Avoid:	1431
Best Practices:	1431
File I/O	1432
Module 8: File I/O Exercises	1433
Exercise 1: Basic File Operations	1433
Exercise 2: Text File Processing	1433
Exercise 3: Binary File Operations	1433
Exercise 4: File System Operations	1433
Exercise 5: Random Access Files	1434
Exercise 6: Advanced File I/O Techniques	1434
Exercise 7: Error Handling and Recovery	1434
Exercise 8: Comprehensive File Management System	1434
Solutions and Tips	1435
Exercise 1 Solution Example:	1435
Exercise 3 Solution Example:	1437
Common Pitfalls to Avoid:	1439
Best Practices:	1439

Modern C 1440

Module 9: Modern C Features Exercises 1441

Exercise 1: C99 Features Implementation	1441
Exercise 2: C11 Standard Features	1441
Exercise 3: C17/C23 New Features	1441
Exercise 4: Type Safety and Modern Practices	1442
Exercise 5: Advanced Preprocessor Techniques	1442
Exercise 6: Memory Management and Allocation	1442
Exercise 7: Error Handling and Diagnostics	1442
Exercise 8: Comprehensive Modern C Application	1443
Solutions and Tips	1443
Exercise 1 Solution Example:	1443
Exercise 2 Solution Example:	1444
Common Pitfalls to Avoid:	1446
Best Practices:	1446

Data Structures 1447

Module 10: Data Structures Exercises 1448

Exercise 1: Linear Data Structures Implementation	1448
Exercise 2: Linked List Variations	1448
Exercise 3: Tree Data Structures	1448
Exercise 4: Hash Tables and Associative Arrays	1449
Exercise 5: Graph Representations and Algorithms	1449
Exercise 6: Advanced Data Structures	1449
Exercise 7: Data Structure Performance Analysis	1449
Exercise 8: Comprehensive Data Structure Application	1450
Solutions and Tips	1450
Exercise 1 Solution Example:	1450
Exercise 3 Solution Example:	1453
Common Pitfalls to Avoid:	1456
Best Practices:	1456

Network 1457

Module 11: Network Programming Exercises 1458

Exercise 1: Basic Socket Programming	1458
Exercise 2: HTTP Client Implementation	1458
Exercise 3: Multi-client Server	1458
Exercise 4: Network Protocol Implementation	1459
Exercise 5: Advanced Networking Concepts	1459
Exercise 6: Network Security Implementation	1459
Exercise 7: Network Debugging and Monitoring	1459
Exercise 8: Comprehensive Network Application	1460
Solutions and Tips	1460
Exercise 1 Solution Example:	1460

Exercise 2 Solution Example:	1464
Common Pitfalls to Avoid:	1467
Best Practices:	1467
Embedded	1468
Module 12: Embedded Systems Programming Exercises	1469
Exercise 1: Bit Manipulation and Register Operations	1469
Exercise 2: Memory-Mapped I/O Implementation	1469
Exercise 3: Interrupt Service Routines	1469
Exercise 4: Real-Time Constraints Implementation	1470
Exercise 5: Low-Power Programming Techniques	1470
Exercise 6: Hardware Abstraction Layer (HAL)	1470
Exercise 7: Embedded Debugging and Monitoring	1470
Exercise 8: Comprehensive Embedded Application	1471
Solutions and Tips	1471
Exercise 1 Solution Example:	1471
Exercise 2 Solution Example:	1473
Common Pitfalls to Avoid:	1475
Best Practices:	1475
Performance	1476
Module 13: Performance Optimization Exercises	1477
Exercise 1: Algorithm Complexity Analysis	1477
Exercise 2: Memory Access Patterns Optimization	1477
Exercise 3: Compiler Optimization Techniques	1477
Exercise 4: Data Structure Optimization	1478
Exercise 5: Parallel Processing Implementation	1478
Exercise 6: SIMD and Vectorization	1478
Exercise 7: Memory Management Optimization	1478
Exercise 8: Comprehensive Performance Optimization	1479
Solutions and Tips	1479
Exercise 1 Solution Example:	1479
Exercise 2 Solution Example:	1482
Common Pitfalls to Avoid:	1484
Best Practices:	1484
Testing	1486
Module 14: Testing and Quality Assurance Exercises	1487
Exercise 1: Unit Testing Framework Implementation	1487
Exercise 2: Mocking and Stubbing Techniques	1487
Exercise 3: Integration Testing Implementation	1487
Exercise 4: Code Coverage Analysis	1488
Exercise 5: Static Analysis and Linting	1488

Exercise 6: Fuzz Testing Implementation	1488
Exercise 7: Performance and Stress Testing	1488
Exercise 8: Comprehensive Testing Strategy	1489
Solutions and Tips	1489
Exercise 1 Solution Example:	1489
Exercise 2 Solution Example:	1492
Common Pitfalls to Avoid:	1496
Best Practices:	1496
Advanced	1497
Module 15: Advanced Topics Exercises	1498
Exercise 1: Metaprogramming with Macros	1498
Exercise 2: Generic Programming Techniques	1498
Exercise 3: Coroutines and Cooperative Multitasking	1498
Exercise 4: Reflection and Introspection	1499
Exercise 5: Template Metaprogramming Simulation	1499
Exercise 6: Memory Model and Concurrency	1499
Exercise 7: Language Interoperability	1499
Exercise 8: Comprehensive Advanced Application	1500
Solutions and Tips	1500
Exercise 1 Solution Example:	1500
Exercise 2 Solution Example:	1502
Common Pitfalls to Avoid:	1505
Best Practices:	1506
Projects	1507
Advanced	1508
Advanced Projects	1509
Project 1: Concurrent Web Server	1509
Description	1509
Learning Objectives	1509
Requirements	1509
Implementation Steps	1509
Sample Code Structure	1510
Common Pitfalls to Avoid	1512
Best Practices	1512
Project 2: Memory Allocator	1513
Description	1513
Learning Objectives	1513
Requirements	1513
Implementation Steps	1513
Sample Code Structure	1513
Common Pitfalls to Avoid	1515

Best Practices	1516
Project 3: Real-time Data Processing System	1516
Description	1516
Learning Objectives	1516
Requirements	1516
Implementation Steps	1517
Sample Code Structure	1517
Common Pitfalls to Avoid	1519
Best Practices	1520
Project 4: Embedded System Simulator	1520
Description	1520
Learning Objectives	1520
Requirements	1520
Implementation Steps	1520
Sample Code Structure	1521
Common Pitfalls to Avoid	1523
Best Practices	1524
Tips for Success	1524

Beginner 1525

Beginner Projects 1526

Project 1: Personal Information Manager	1526
Description	1526
Learning Objectives	1526
Requirements	1526
Implementation Steps	1526
Sample Code Structure	1527
Common Pitfalls to Avoid	1528
Best Practices	1528
Project 2: Simple Calculator	1528
Description	1528
Learning Objectives	1529
Requirements	1529
Implementation Steps	1529
Sample Code Structure	1529
Common Pitfalls to Avoid	1530
Best Practices	1530
Project 3: Number Guessing Game	1530
Description	1530
Learning Objectives	1530
Requirements	1531
Implementation Steps	1531
Sample Code Structure	1531
Common Pitfalls to Avoid	1532
Best Practices	1532

Project 4: Temperature Converter	1533
Description	1533
Learning Objectives	1533
Requirements	1533
Implementation Steps	1533
Sample Code Structure	1533
Common Pitfalls to Avoid	1534
Best Practices	1535
Tips for Success	1535
Capstone	1536
Capstone Projects	1537
Project 1: Distributed Chat Application	1537
Description	1537
Learning Objectives	1537
Requirements	1537
Implementation Steps	1538
Sample Code Structure	1538
Key Challenges	1542
Evaluation Criteria	1542
Project 2: Operating System Kernel	1542
Description	1542
Learning Objectives	1543
Requirements	1543
Implementation Steps	1544
Sample Code Structure	1544
Key Challenges	1548
Evaluation Criteria	1548
Project 3: Database Management System	1549
Description	1549
Learning Objectives	1549
Requirements	1549
Implementation Steps	1550
Sample Code Structure	1550
Key Challenges	1555
Evaluation Criteria	1555
Project 4: Game Engine	1555
Description	1555
Learning Objectives	1555
Requirements	1556
Implementation Steps	1556
Sample Code Structure	1557
Key Challenges	1563
Evaluation Criteria	1563
Tips for Capstone Success	1563

Intermediate	1564
Intermediate Projects	1565
Project 1: Student Grade Management System	1565
Description	1565
Learning Objectives	1565
Requirements	1565
Implementation Steps	1565
Sample Code Structure	1566
Common Pitfalls to Avoid	1568
Best Practices	1568
Project 2: Simple Text Editor	1568
Description	1568
Learning Objectives	1568
Requirements	1568
Implementation Steps	1569
Sample Code Structure	1569
Common Pitfalls to Avoid	1570
Best Practices	1571
Project 3: Library Management System	1571
Description	1571
Learning Objectives	1571
Requirements	1571
Implementation Steps	1571
Sample Code Structure	1572
Common Pitfalls to Avoid	1574
Best Practices	1574
Project 4: Simple Database Engine	1574
Description	1574
Learning Objectives	1575
Requirements	1575
Implementation Steps	1575
Sample Code Structure	1575
Common Pitfalls to Avoid	1577
Best Practices	1578
Tips for Success	1578
Appendices	1579
Appendix A: Additional Resources	1580
Books and Publications	1580
Core C Programming	1580
Systems Programming	1580
Data Structures and Algorithms	1580
Online Resources	1581
Official Documentation	1581
Tutorials and Learning Platforms	1581

Reference Sites	1582
Development Tools	1582
Compilers	1582
Debugging Tools	1582
Profiling Tools	1583
Static Analysis Tools	1583
Libraries and Frameworks	1583
Standard Libraries	1583
Networking Libraries	1584
Data Structure Libraries	1584
Communities and Forums	1584
Online Communities	1584
Mailing Lists	1585
Conferences and Events	1585
Major Conferences	1585
Open Source Projects	1585
Learning from Real Code	1585
Practice Platforms	1586
Coding Challenges	1586
Standards Organizations	1586
Key Organizations	1586
Additional Learning Resources	1586
Video Courses	1586
Podcasts	1587
Appendix B: C Language Reference	1588
C Keywords	1588
Standard C Keywords (C17/C18)	1588
Deprecated Keywords	1588
Data Types and Sizes	1588
Basic Data Types	1588
Type Specifiers	1589
Type Qualifiers	1589
Operators	1589
Arithmetic Operators	1589
Relational Operators	1589
Logical Operators	1590
Bitwise Operators	1590
Assignment Operators	1590
Other Operators	1591
Standard Library Functions	1591
Input/Output Functions (<stdio.h>)	1591
String Functions (<string.h>)	1592
Memory Management Functions (<stdlib.h>)	1592
Mathematical Functions (<math.h>)	1592
Utility Functions (<stdlib.h>)	1593
Time Functions (<time.h>)	1594

Preprocessor Directives	1594
Common Directives	1594
Standard Macros	1594
Format Specifiers	1594
printf Format Specifiers	1594
scanf Format Specifiers	1595
Format Modifiers	1595
Escape Sequences	1596
Standard Escape Sequences	1596
Octal and Hexadecimal Escapes	1596
Storage Classes	1596
Storage Class Specifiers	1596
Error Handling	1597
errno Values (<errno.h>)	1597
Assertions (<assert.h>)	1597
Signal Handling (<signal.h>)	1597
Signal Types	1597
Signal Functions	1597
File Operations Modes	1597
fopen Modes	1597
Binary Mode Suffix	1598
Standard Streams	1598
Predefined Streams	1598
Limits and Constants	1598
Implementation Limits (<limits.h>)	1598
Floating-Point Limits (<float.h>)	1599
Thread Support (C11) (<threads.h>)	1599
Thread Functions	1599
Mutex Functions	1599
Condition Variable Functions	1599
Atomic Operations (C11) (<stdatomic.h>)	1600
Atomic Types	1600
Atomic Operations	1600
Generic Selection (C11) (<stdnoreturn.h>)	1600
_Generic Keyword	1600
Alignment (C11) (<stdalign.h>)	1601
Alignment Specifiers	1601
Appendix C: Exercise Solutions	1602
Module 1: Foundations	1602
Exercise 1.1: Hello World Variations	1602
Exercise 1.2: Basic Calculator	1602
Module 2: Data Types	1603
Exercise 2.1: Size and Range Calculator	1603
Exercise 2.2: Temperature Converter	1604
Module 3: Control Flow	1605
Exercise 3.1: Prime Number Checker	1605
Exercise 3.2: Pattern Printer	1606

Module 4: Functions	1607
Exercise 4.1: Mathematical Functions Library	1607
Exercise 4.2: Recursive Functions	1610
Module 5: Arrays and Strings	1611
Exercise 5.1: Array Operations	1611
Exercise 5.2: String Manipulation	1613
Module 6: Pointers	1615
Exercise 6.1: Pointer Arithmetic	1615
Exercise 6.2: Dynamic Memory Allocation	1616
Module 7: Structures	1618
Exercise 7.1: Student Database	1618
Exercise 7.2: Complex Number Operations	1621
Module 8: File I/O	1624
Exercise 8.1: Text File Processor	1624
Exercise 8.2: Binary File Operations	1627
Module 9: Modern C Features	1630
Exercise 9.1: C99 Features Demonstration	1630
Exercise 9.2: C11 Features Demonstration	1631
Module 10: Data Structures	1633
Exercise 10.1: Linked List Implementation	1633
Exercise 10.2: Binary Search Tree Implementation	1637
Module 11: Network Programming	1641
Exercise 11.1: Simple TCP Client-Server	1641
Exercise 11.2: HTTP Client	1644
Module 12: Embedded Systems Programming	1646
Exercise 12.1: Bit Manipulation Library	1646
Exercise 12.2: Memory-Mapped I/O Simulation	1647
Module 13: Performance Optimization	1650
Exercise 13.1: Algorithm Complexity Analysis	1650
Exercise 13.2: Memory Access Patterns	1652
Module 14: Testing and Quality Assurance	1655
Exercise 14.1: Unit Testing Framework	1655
Exercise 14.2: Mocking Framework	1658
Module 15: Advanced Topics	1662
Exercise 15.1: Multithreading Example	1662
Exercise 15.2: System Programming Example	1663

C Programming

A modular C library: foundations through systems topics, with a separate workbook for extra drills and projects.

This book uses **C17** as the default dialect. Examples prefer:

```
gcc -std=c17 -Wall -Wextra -Wpedantic -o prog prog.c
./prog
```

Clang is fine as a drop-in for learning:

```
clang -std=c17 -Wall -Wextra -Wpedantic -o prog prog.c
```

Who this is for

Path A — Beginner (language first)

Start at Module 1 and move in order through Module 5 before jumping ahead. Type every full program. Treat warnings as errors in practice (**-Werror** once you are comfortable). Goal: confident control flow, functions, arrays/strings, and a first honest encounter with pointers.

Path B — Systems-oriented (already know some C syntax)

Skim Module 1 for the compile/link mental model and warning flags, then prioritize Modules 6–8 (pointers, structures, file/system I/O). Use Modules 11–15 when you need sockets, embedded patterns, performance, testing, or threads. Still run Module 1’s multi-file and I/O labs if those habits are rusty.

Both paths should enable warnings early and avoid treating “it compiled” as “it is correct.” Undefined behavior is introduced early for a reason.

How to use this book

1. Read a section, then **type** the programs (do not only read them).
2. Compile with **-std=c17 -Wall -Wextra -Wpedantic**.
3. Break one thing on purpose; read the diagnostic; fix it.
4. Do the chapter exercises before moving on.
5. Use the workbook when you want more reps or a larger project—not as a substitute for chapter work.

How to compile examples

From any directory where you saved a single-file example:

```
gcc -std=c17 -Wall -Wextra -Wpedantic -o hello hello.c
./hello
```

Multi-file:

```
gcc -std=c17 -Wall -Wextra -Wpedantic -o app main.c util.c
# or
gcc -std=c17 -Wall -Wextra -Wpedantic -c main.c
gcc -std=c17 -Wall -Wextra -Wpedantic -c util.c
gcc -o app main.o util.o
```

Useful extras:

Flags	Purpose
-g -O0	Debug builds
-O2	Optimized builds
-fsanitize=address,undefined	Runtime checks (GCC/Clang, host-dependent)

Workbook demos live under `content/99-workbook/examples/` (per module). Some multi-file folders include a `Makefile`—read it, then still know how to issue the raw `gcc` line yourself.

Workbook (content/99-workbook)

Area	Path	Use for
Examples	<code>content/99-workbook/examples/</code>	Questions → demos aligned to modules
Exercises	<code>content/99-workbook/exercises/</code>	Drills → problem sets
Projects	<code>content/99-workbook/projects/</code>	Design → capstone project prompts

Suggested rhythm: finish the module chapters → attempt workbook exercises for that module → pick a project only when the drills feel mechanical.

Appendices under `content/100-appendices/` hold resources, quick reference material, and solution notes.

Related books

Book	Overlap
Linux	Toolchain, make, and systems environment
Go	Contrasting memory and concurrency models
Maths	Discrete math and complexity behind algorithms chapters

Modules (15 + workbook + appendices)

1. **Foundations and Environment** — history and process model, toolchain, first programs, console I/O
 2. **Data Types and Variables** — types, storage, conversions, operators
 3. **Control Flow** — decisions, loops, non-local flow
 4. **Functions and Modularity** — design, linkage, library overview
 5. **Arrays and Strings** — contiguous data and text
 6. **Pointers and Memory** — addresses, dynamic allocation, pitfalls
 7. **Structures and Types** — struct, union, enum, simple ADTs
 8. **File I/O and System Programming** — streams and system-facing I/O
 9. **Modern C Standards** — C99 through C23 and modern practice
 10. **Data Structures and Algorithms** — lists, trees, tables, graphs, sorting
 11. **Network Programming** — sockets and protocols
 12. **Embedded Systems** — constrained and hardware-oriented C
 13. **Performance Optimization** — measure, then improve
 14. **Testing and Debugging** — quality, debug, error handling
 15. **Advanced Topics and Capstone** — concurrency, systems topics, capstone
- Plus **99-workbook** (examples, exercises, projects) and **appendices**.

Prerequisites

- Comfort using a text editor and a terminal

- Ability to install a C compiler (GCC, Clang, or MSVC/MinGW on Windows)
- No prior programming required for Path A; Path B assumes basic syntax familiarity

Getting started

1. [Introduction to C](#) — mental model and first mini-programs
2. [Development Environment](#) — install and verify the toolchain
3. [Your First C Program](#) — compile, multi-file, labs
4. [Basic I/O](#) — `printf` / `scanf` / `fgets`

Then continue through the modules in order (Path A) or jump to pointers and systems modules (Path B) once foundations are solid.

Intro

Comprehensive C Programming Syllabus

Course Overview

This comprehensive C programming course is designed to take students from complete beginners to advanced practitioners capable of writing efficient, maintainable, and professional C code. The syllabus draws from classic texts like Kernighan & Ritchie's "The C Programming Language," modern best practices, and real-world applications in systems programming, embedded development, and network programming.

Learning Objectives

By the end of this course, students will be able to:

- Write efficient and maintainable C programs
- Understand memory management and pointer manipulation
- Develop system-level applications
- Create network applications and embedded systems
- Apply modern C development practices and tooling
- Debug and optimize C programs
- Work with various C standards (C89, C99, C11, C17, C23)

Prerequisites

- Basic understanding of computer systems
- Familiarity with command-line interfaces
- Mathematical reasoning skills
- No prior programming experience required

Course Structure

The course is divided into 12 comprehensive modules, progressing from fundamentals to specialized applications.

Module 1: Foundations and Environment Setup

Duration: 1-2 weeks

1.1 Introduction to C Programming

- History and evolution of C (1972-present)
- Why C remains crucial in modern programming (2024 perspective)
- **Current C standards: C89/C90, C99, C11, C17, C23 (latest - ISO/IEC 9899:2024)**
- Relationship to other languages (C++, Rust, Go, Zig influence)
- Modern applications: Operating Systems, Embedded Systems, IoT, WebAssembly, System Tools
- C in cloud computing and microservices
- Performance-critical applications and real-time systems

1.2 Development Environment Setup

- **Modern compilers supporting C23:** GCC 14+, Clang 18+, MSVC 2022+
- Cross-platform development: MinGW-w64, WSL2, Docker containers
- **Modern IDEs and editors:** VS Code with C/C++ extensions, CLion 2024, Neovim/Vim with LSP
- **Cloud development:** GitHub Codespaces, Gitpod, online compilers
- Command-line tools and build systems: Make, CMake 3.25+, Ninja, Meson
- **Modern version control:** Git workflows, GitHub/GitLab CI/CD
- **Advanced debugging tools:** GDB 13+, LLDB, Valgrind, AddressSanitizer, ThreadSanitizer
- **Package management:** vcpkg, Conan 2.0, CPM.cmake
- **Static analysis:** Clang Static Analyzer, PVS-Studio, Cppcheck, SonarQube

1.3 First C Program

- Structure of a C program
- The `main()` function
- Comments and code style
- Compilation process: preprocessing, compilation, linking
- Understanding object files and executables

1.4 Basic I/O

- `printf()` and format specifiers
- `scanf()` and input handling
- Character I/O with `getchar()` and `putchar()`
- Introduction to streams

Projects: - Hello World variations - Simple calculator - Number guessing game

Module 2: Data Types and Variables

Duration: 2 weeks

2.1 Fundamental Data Types

- Integer types: `char`, `short`, `int`, `long`, `long long`
- Floating-point types: `float`, `double`, `long double`
- Boolean type (`_Bool`, `<stdbool.h>`)
- Size and range considerations
- Fixed-width integer types (`<stdint.h>`)

2.2 Variables and Constants

- Variable declaration and initialization
- Scope and lifetime
- Storage classes: `auto`, `static`, `extern`, `register`
- Constants: `const` keyword, `#define`, `enum`
- Volatile qualifier

2.3 Type Conversions

- Implicit conversions and promotion rules
- Explicit casting
- Common pitfalls and warnings

2.4 Operators and Expressions

- Arithmetic operators and precedence
- Relational and logical operators
- Bitwise operators
- Assignment operators
- Increment/decrement operators
- Conditional operator
- `sizeof` operator

Projects: - Temperature conversion utility - Bitwise operations calculator - Simple arithmetic expression evaluator

Module 3: Control Flow

Duration: 2 weeks

3.1 Conditional Statements

- `if`, `else if`, `else` statements
- Nested conditionals
- `switch` statements and `break`
- Best practices for conditional logic

3.2 Loops

- `while` and `do-while` loops
- `for` loops and variants
- Nested loops
- Loop control: `break`, `continue`
- Infinite loops and when to use them

3.3 Program Flow Control

- `goto` statement (when and why to avoid)
- Function calls as control flow
- Error handling patterns

Projects: - Menu-driven program - Pattern printing programs - Prime number generator - Simple text-based games (Tic-tac-toe, Hangman)

Module 4: Functions and Modular Programming

Duration: 2-3 weeks

4.1 Function Fundamentals

- Function declaration and definition
- Parameter passing (by value)
- Return values and `return` statement
- Function prototypes and headers
- Scope rules and name hiding

4.2 Advanced Function Concepts

- Recursive functions
- Variable argument functions (`<stdarg.h>`)
- Inline functions (C99)
- Function pointers and callbacks
- Higher-order functions

4.3 Modular Programming

- Header files and source files
- Include guards and `#pragma once`
- Compilation units and linking
- Static vs. external linkage
- Creating and using libraries

4.4 Standard Library Functions

- String functions (`<string.h>`)
- Mathematical functions (`<math.h>`)
- Character classification (`<ctype.h>`)
- Time functions (`<time.h>`)
- Random number generation (`<stdlib.h>`)

Projects: - Mathematical library implementation - Text processing utilities - Recursive algorithms (factorial, Fibonacci, tree traversal) - Function pointer demonstration programs

Module 5: Arrays and Strings

Duration: 2-3 weeks

5.1 Arrays

- Array declaration and initialization
- Multi-dimensional arrays
- Array bounds and buffer overflows
- Variable-length arrays (C99)
- Arrays as function parameters

5.2 Strings

- Character arrays and null termination
- String literals and string constants
- String input/output functions
- String manipulation functions
- Common string operations

5.3 Advanced Array Concepts

- Array of arrays vs. multi-dimensional arrays
- Jagged arrays using pointers
- Dynamic arrays (preview to pointers)
- Array algorithms: searching, sorting

Projects: - Text analyzer (word count, character frequency) - Simple sorting algorithms implementation - String manipulation library - Command-line argument processor

Module 6: Pointers and Memory Management

Duration: 3-4 weeks

6.1 Pointer Fundamentals

- Understanding memory addresses
- Pointer declaration and initialization
- Dereferencing and address-of operators
- Null pointers and pointer validation
- Pointer arithmetic

6.2 Pointers and Arrays

- Array-pointer equivalence
- Pointer subscripting
- Multi-dimensional arrays via pointers
- Function parameters: arrays vs. pointers

6.3 Dynamic Memory Management

- `malloc()`, `calloc()`, `realloc()`, `free()`
- Memory leaks and double-free errors
- Memory debugging tools
- Best practices for dynamic allocation

6.4 Advanced Pointer Concepts

- Pointers to pointers
- Function pointers advanced usage
- Void pointers and generic programming
- Const correctness with pointers

6.5 Common Pointer Pitfalls

- Dangling pointers
- Buffer overflows
- Segmentation faults
- Memory alignment issues

Projects: - Dynamic array implementation - Linked list data structure - Memory allocator simulation - Generic sorting function using function pointers

Module 7: Structures and User-Defined Types

Duration: 2-3 weeks

7.1 Structures

- Structure declaration and definition
- Structure initialization and assignment
- Nested structures
- Arrays of structures
- Pointers to structures

7.2 Unions and Enumerations

- Union declaration and usage
- Anonymous unions (C11)
- Enumeration types
- Enumeration with explicit values

7.3 Advanced Structure Concepts

- Bit fields
- Structure packing and alignment
- Flexible array members (C99)
- Self-referential structures

7.4 Abstract Data Types

- Information hiding principles
- Opaque pointers
- Constructor/destructor patterns
- Interface design in C

Projects: - Student record management system - Binary tree implementation - Stack and queue data structures - Simple database using structures

Module 8: File I/O and System Programming

Duration: 2-3 weeks

8.1 File Operations

- File pointers and FILE structure
- Opening and closing files
- Text vs. binary modes
- Error handling in file operations

8.2 File I/O Functions

- Character I/O: `fgetc()`, `fputc()`, `ungetc()`
- Line I/O: `fgets()`, `fputs()`
- Formatted I/O: `fprintf()`, `fscanf()`
- Binary I/O: `fread()`, `fwrite()`

8.3 Advanced File Operations

- File positioning: `fseek()`, `ftell()`, `rewind()`
- Random access files
- Temporary files
- File locking and concurrent access

8.4 System-Level Programming

- Command-line arguments
- Environment variables
- Exit status and `atexit()`
- Signal handling basics
- Process creation (system calls overview)

Projects: - File encryption/decryption utility - Log file analyzer - Configuration file parser - Simple text editor

Module 9: Advanced C Features and Modern Standards

Duration: 2-3 weeks

9.1 C99 Features

- Variable-length arrays
- Designated initializers
- Compound literals
- `restrict` keyword
- `inline` functions
- New data types (`long long`, `_Bool`)

9.2 C11 Features

- Anonymous structures and unions
- Static assertions (`_Static_assert`)
- Improved Unicode support
- Thread-local storage
- Bounds-checking functions (Annex K)

9.3 C17 and C23 Features

- C17 (ISO/IEC 9899:2018): Technical corrigendum with bug fixes and clarifications
- **C23 (ISO/IEC 9899:2024): Latest Standard - Official Release**
 - `typeof` and `typeof_unqual` operators
 - `_BitInt(N)` extended integer types
 - Binary integer constants (e.g., `0b1010`)
 - `#elifdef` and `#elifndef` preprocessor directives
 - `#warning` directive
 - `#embed` directive for binary resource inclusion

- `[[...]]` attribute syntax support
- `[[noreturn]]`, `[[maybe_unused]]`, `[[deprecated]]` attributes
- `[[nodiscard]]`, `[[fallthrough]]` attributes
- `_Decimal128`, `_Decimal32`, `_Decimal64` decimal floating types
- Improved `_Static_assert` (message now optional)
- Enhanced integer constant expressions
- `auto` type inference (limited scope)
- Improved Unicode support (`char8_t` via `__STDC_UTF_8__`)
- New standard library functions in `<stdbit.h>`
- Memory alignment improvements
- `strdup()`, `strndup()` functions standardized
- `aligned_alloc()` improvements
- Enhanced time utilities
- Better support for embedded systems

9.4 Modern C Development Practices (2024)

- **DevOps and C:** CI/CD pipelines for C projects
- **Containerized development:** Docker for C development environments
- **Cross-compilation:** Modern toolchains and target support
- **WebAssembly (WASM):** Compiling C to WebAssembly with Emscripten
- **Cloud-native C:** Microservices and serverless functions in C
- **Supply chain security:** Dependency management and vulnerability scanning
- **Documentation as code:** Doxygen, Sphinx integration
- **Modern testing:** Fuzzing, property-based testing, mutation testing

9.5 Advanced Preprocessor

- Macro definitions and function-like macros
- Conditional compilation
- Pragma directives
- Predefined macros
- Token pasting and stringification

Projects: - Generic data structure library using macros - Configuration system with conditional compilation - Cross-platform compatibility layer - C23 feature detection and compatibility testing - Modern C development workflow setup

Note on C23 Compiler Support (2024): - **GCC 14+:** Partial C23 support, `typeof`, binary literals, `#elifdef` - **Clang 18+:** Growing C23 support, attributes, enhanced preprocessor - **MSVC 2022:** Limited C23 features, focus on C11/C17 compliance - **Industry adoption:** Gradual migration, embedded systems leading adoption - **Compatibility strategies:** Feature detection macros, fallback implementations

Module 10: Data Structures and Algorithms

Duration: 3-4 weeks

10.1 Linear Data Structures

- Dynamic arrays (vectors)
- Linked lists (singly, doubly, circular)
- Stacks and queues
- Deques and priority queues

10.2 Tree Structures

- Binary trees and binary search trees
- AVL trees and red-black trees
- B-trees and B+ trees
- Tries and suffix trees

10.3 Hash Tables

- Hash function design
- Collision resolution strategies
- Dynamic resizing
- Performance analysis

10.4 Graph Algorithms

- Graph representations
- Breadth-first and depth-first search
- Shortest path algorithms
- Minimum spanning trees

10.5 Sorting and Searching

- Comparison-based sorting algorithms
- Non-comparison sorting (radix, counting)
- Binary search variations
- Selection algorithms

Projects: - Complete data structure library - Graph algorithms visualization - Sorting algorithm performance comparison - Hash table implementation with benchmarks

Module 11: Network Programming

Duration: 3-4 weeks

11.1 Network Fundamentals

- OSI and TCP/IP models
- Sockets concept
- IPv4 and IPv6 addressing
- Ports and protocols

11.2 Socket Programming

- Socket creation and configuration
- TCP client-server programming
- UDP programming
- Address structures and conversion functions

11.3 Advanced Network Programming

- Non-blocking I/O
- I/O multiplexing (`select()`, `poll()`, `epoll()`)
- Multi-threading in network programs
- Signal handling in network applications

11.4 Network Protocols

- HTTP client implementation
- FTP basics
- DNS resolution
- Custom protocol design

11.5 Security Considerations

- Buffer overflow prevention
- Input validation
- Encryption basics (with OpenSSL)
- Secure coding practices

Projects: - Echo server and client - HTTP web server - Chat application - File transfer utility - Network monitoring tools

Module 12: Embedded Systems Programming

Duration: 3-4 weeks

12.1 Embedded Systems Basics

- Microcontroller architecture
- Memory models (Harvard vs. Von Neumann)
- Real-time constraints
- Resource limitations

12.2 Low-Level Programming

- Direct memory access
- Hardware registers manipulation
- Interrupt service routines
- Device drivers basics

12.3 Embedded C Techniques

- Volatile variables usage
- Bit manipulation for hardware control
- Fixed-point arithmetic
- Code optimization for space and speed

12.4 Real-Time Programming

- Task scheduling concepts
- Inter-task communication
- Synchronization primitives
- Real-time operating systems (RTOS) basics

12.5 Hardware Interfaces

- GPIO programming
- Serial communication (UART, SPI, I2C)
- Analog-to-digital conversion
- PWM and timing control

Projects: - LED control system - Sensor data acquisition - Motor control application - Communication protocol implementation - Embedded web server (if hardware permits)

Module 13: Performance Optimization and Profiling

Duration: 2 weeks

13.1 Performance Analysis

- Profiling tools (gprof, perf, Instruments)
- Time and space complexity analysis
- Benchmarking methodologies
- Performance bottleneck identification

13.2 Optimization Techniques

- Algorithmic optimizations
- Data structure selection
- Cache-friendly programming
- Compiler optimizations

13.3 Memory Optimization

- Memory alignment and padding
- Memory pool allocation
- Stack vs. heap usage
- Memory-mapped files

13.4 Advanced Optimization

- SIMD instructions
- Loop optimizations
- Function inlining strategies
- Profile-guided optimization

Projects: - Performance comparison suite - Memory usage analyzer - Optimized mathematical library - Real-world application optimization

Module 14: Testing and Debugging

Duration: 2 weeks

14.1 Testing Methodologies

- Unit testing in C
- Integration testing
- Test-driven development
- Testing frameworks (Unity, CUnit, Check)

14.2 Debugging Techniques

- GDB usage and advanced features
- Static analysis tools (lint, cppcheck, clang-analyzer)
- Dynamic analysis (Valgrind, AddressSanitizer)
- Debugging embedded systems

14.3 Error Handling

- Error codes and return values
- Exception simulation in C
- Logging and tracing
- Graceful degradation

14.4 Code Quality

- Code review practices
- Static analysis integration
- Continuous integration
- Documentation standards

Projects: - Comprehensive test suite for previous projects - Debugging exercise with intentional bugs - Code quality assessment tools

Module 15: Advanced Topics and Capstone

Duration: 3-4 weeks

15.1 Multi-threading and Concurrency

- POSIX threads (pthreads)
- Thread synchronization (mutexes, semaphores)
- Thread-safe programming
- Lock-free data structures

15.2 Advanced System Programming

- Inter-process communication
- Shared memory
- Memory mapping
- System call wrappers

15.3 Domain-Specific Applications

- Compiler construction basics
- Operating system components
- Database engine components
- Graphics programming

15.4 Capstone Project

- Project planning and design
- Implementation phases
- Testing and documentation
- Code review and presentation

Capstone Project Ideas: - Mini operating system kernel - Database management system - Network protocol implementation - Embedded control system - Compiler for a simple language - Real-time data processing system

Assessment Methods

Continuous Assessment (60%)

- Weekly programming assignments (30%)
- Module projects (20%)
- Code review participation (10%)

Major Projects (25%)

- Mid-term project (Module 6-8)
- Capstone project (Module 15)

Examinations (15%)

- Practical coding exams
 - Theoretical understanding tests
-

Recommended Resources

Essential Books

1. **Kernighan, Brian W. & Ritchie, Dennis M.** - “The C Programming Language” (2nd Edition) - Still the definitive classic
2. **King, K.N.** - “C Programming: A Modern Approach” (2nd Edition)
3. **Harbison, Samuel P. & Steele, Guy L.** - “C: A Reference Manual” (5th Edition)
4. **Gustedt, Jens** - “Modern C” (2019) - Covers C11/C17 with modern practices

Advanced and Specialized References

5. **Stevens, W. Richard & Rago, Stephen A.** - “Advanced Programming in the UNIX Environment” (3rd Edition, 2013)
6. **Stevens, W. Richard** - “UNIX Network Programming” (3rd Edition)
7. **Barr, Michael & Massa, Anthony** - “Programming Embedded Systems” (2nd Edition)
8. **Seacord, Robert C.** - “Effective C” (2020) - Modern secure coding practices
9. **van der Linden, Peter** - “Expert C Programming” - Classic advanced techniques

Current Standards and Specifications

10. **ISO/IEC 9899:2024** - C23 Standard (Latest Official Release)
11. **ISO/IEC 9899:2018** - C17 Standard
12. **ISO/IEC 9899:2011** - C11 Standard
13. **POSIX.1-2017** - IEEE Std 1003.1-2017

Modern Online Resources (2024)

- **C23 Documentation:** Official ISO documentation and compiler support matrices
- **cppreference.com/c** - Comprehensive C reference with C23 updates
- **GNU C Library Documentation** - Latest glibc documentation
- **Clang/LLVM Documentation** - Modern compiler features and C23 support
- **GCC Documentation** - C23 implementation status and extensions
- **CERT C Coding Standard** - Updated security guidelines
- **Linux man pages** - System programming references

Development Tools (2024 Current Versions)

- **Compilers:** GCC 14.x, Clang 18.x, MSVC 2022 (17.8+)
- **IDEs:** CLion 2024.x, VS Code with C/C++ Extension Pack
- **Build Systems:** CMake 3.28+, Meson 1.3+, Bazel for large projects
- **Debugging:** GDB 14.x, LLDB 18.x, Intel Inspector, ARM Development Studio
- **Static Analysis:** Clang Static Analyzer, PVS-Studio 2024, PC-lint Plus
- **Dynamic Analysis:** Valgrind 3.22+, AddressSanitizer, Dr. Memory
- **Profiling:** Intel VTune 2024, AMD Prof, Google Performance Tools
- **Version Control:** Git 2.43+, GitHub Advanced Security features
- **CI/CD:** GitHub Actions, GitLab CI, Azure DevOps with C/C++ support

Embedded and IoT Development (Current)

- **ARM Cortex Development:** ARM Keil MDK, STM32CubeIDE 1.14+
- **RISC-V Tools:** SiFive Freedom Studio, PlatformIO
- **IoT Platforms:** ESP-IDF 5.x, Zephyr RTOS 3.x, FreeRTOS 11.x
- **Real-time Systems:** QNX 8.0, VxWorks 23.09

Security and Best Practices

- **OWASP C/C++ Security Guidelines** (2023 updates)
 - **NIST Cybersecurity Framework** - C development guidelines
 - **MISRA C:2023** - Latest automotive safety standard
 - **CERT C Secure Coding Standard** - Updated for modern threats
-

Course Policies

Prerequisites Reinforcement

- Mathematical foundation review sessions
- Command-line tutorial for beginners
- Basic computer architecture overview

Academic Integrity

- Code plagiarism detection
- Proper attribution for reference code
- Collaboration guidelines

Accessibility and Support

- Office hours schedule
 - Peer programming sessions
 - Online forum for questions
 - Additional resources for struggling students
-

Learning Outcomes Verification

Upon completion, students should demonstrate:

1. Technical Proficiency

- Write efficient, readable, and maintainable C code
- Debug complex programs using various tools
- Optimize code for performance and memory usage

2. System Understanding

- Understand memory management and pointer manipulation
- Implement system-level functionality
- Work with hardware interfaces in embedded systems

3. Professional Skills

- Follow coding standards and best practices
- Conduct code reviews effectively
- Document code appropriately
- Work collaboratively on coding projects

4. Problem-Solving Abilities

- Analyze complex problems and design solutions
- Choose appropriate data structures and algorithms
- Handle edge cases and error conditions

5. Specialized Knowledge

- Network programming capabilities
- Embedded systems development
- Multi-threaded programming
- Cross-platform development awareness

This comprehensive syllabus provides a thorough foundation in C programming while preparing students for real-world development challenges across various domains including systems programming, embedded development, and network applications.

Foundations

Introduction to C Programming

What is C?

C is a general-purpose, procedural programming language developed in the early 1970s by Dennis Ritchie at Bell Labs. It was created to rewrite the UNIX operating system in a portable high-level language, and it still sits at the foundation of modern systems software.

C is often called a “middle-level” language: it has high-level control structures and types, yet it exposes memory addresses, integer sizes, and hardware-friendly representations. That combination is why kernels, drivers, embedded firmware, databases, and language runtimes are still written in C (or in languages that interoperate closely with C).

You will write source files ending in `.c`, compile them into machine code, and run the resulting process. This chapter explains what that pipeline does, how a running C program is laid out in memory, and how to read a simple program line by line.

A short history (enough to orient you)

Era	Milestone
1969–1972	C evolves from B (and BCPL) on early UNIX at Bell Labs
1973	UNIX kernel rewritten in C
1978	Kernighan & Ritchie publish <i>The C Programming Language</i> (K&R C)
1989 / 1990	ANSI C (C89), then ISO C (C90)
1999	C99: <code>//</code> comments, <code>long long</code> , VLAs, mixed declarations
2011	C11: atomics, threads API, better Unicode support
2018	C17: mostly clarifications and defect fixes of C11
2024	C23: binary literals, <code>typeof</code> , attributes, more library work

This book defaults to **C17** for portable examples (`-std=c17`). When a feature needs C23 or is C99-only, the text says so.

Why learn C?

- **Control:** you decide layout, lifetimes, and how memory is used.
- **Performance:** little runtime magic between your code and the CPU.
- **Systems literacy:** APIs for operating systems and libraries are often C (or C-shaped).

- **Transferable mental model:** C++, Rust, Go, Zig, and many VMs make more sense after you understand C's model of memory, linking, and calling conventions.
- **Career relevance:** kernels, embedded, networking, game engines, databases, security tooling, and high-performance services still need C-fluent engineers.

C is not the only systems language. It is still the lingua franca of low-level interfaces.

From source to process: how a C program runs

A `.c` file is not what the CPU executes. The build tools turn text into a binary image; the operating system loads that image as a **process**.

`source.c`

Preprocess `#include`, `#define`, conditional compilation

`.i` (optional intermediate)

Compile `C` → assembly (or straight to object on some toolchains)

`.s`

Assemble assembly → machine code in a relocatable object

`.o`

Link resolve symbols, pull in `libc`, produce executable

`a.out` / `hello`

Run OS loads program → process with its own address space

What each stage does

1. **Preprocess** — text-level expansion. `#include <stdio.h>` is replaced by the contents of that header (plus nested includes). Macros like `#define N 10` become substitutions.
2. **Compile** — the compiler checks types and syntax, then generates assembly (or an internal IR that becomes object code).
3. **Assemble** — assembly becomes a **relocatable object file** (`.o`): machine instructions and data with unresolved external references (for example, `printf`).
4. **Link** — the linker stitches object files and libraries together, assigns final addresses, and produces an **executable**.
5. **Run** — the OS creates a process, maps the executable (and shared libraries) into memory, sets up the stack, and jumps to the program entry point (which eventually calls `main`).

Watching the stages with GCC

Save this as `hello.c`:

```
#include <stdio.h>

int main(void) {
    printf("Hello, World!\n");
    return 0;
}
```

Recommended compile line (use this habit throughout the book):

```
gcc -std=c17 -Wall -Wextra -Wpedantic -o hello hello.c
./hello
```

Inspect stages individually:

```
# Preprocessed source (huge; useful when macros go wrong)
gcc -std=c17 -E hello.c -o hello.i

# Assembly listing
gcc -std=c17 -S hello.c -o hello.s

# Object file without linking
gcc -std=c17 -c hello.c -o hello.o

# Link the object into an executable
gcc -o hello hello.o
```

You do not need to memorize assembly. You do need the mental model: **source** → **translation unit** → **object** → **executable** → **process**.

Memory model intro: text, data, stack, heap

When your program runs, the process address space is divided into regions. Details differ by OS and architecture, but this classic picture is enough for foundations:

High addresses

stack	grows downward ← local vars, call frames
(free)	
heap	grows upward ← malloc / free
BSS (uninitialized data)	global/static zeros
data (initialized)	global/static with values
text (code)	machine instructions (often read-only)

Low addresses

Region	Typical contents	Lifetime
text	Compiled instructions	Whole program
data	Initialized globals / statics	Whole program
BSS	Zero-initialized globals / statics	Whole program
heap	malloc / calloc / realloc memory	Until free (or exit)
stack	Locals, parameters, return addresses	Function call

Early chapters use **stack** locals and string literals. Dynamic **heap** allocation appears later with pointers. The important habit now: every object has a **storage duration** (automatic, static, or allocated) and a **scope** (where its name is visible). Confusing those two is a common source of bugs.

Minimal illustration of where things live (you are not expected to print addresses yet—this is orientation):

```
#include <stdio.h>
#include <stdlib.h>

int g_count = 1;      /* data (initialized global) */
int g_zero;           /* BSS (zero-initialized global) */
```

```

int main(void) {
    int local = 42;          /* stack */
    static int sticky = 7; /* data/BSS, survives across calls */
    int *p = malloc(sizeof *p); /* heap (if malloc succeeds) */

    if (p == NULL) {
        fprintf(stderr, "malloc failed\n");
        return 1;
    }
    *p = 99;

    printf("local=%d sticky=%d *p=%d g_count=%d g_zero=%d\n",
           local, sticky, *p, g_count, g_zero);

    free(p);
    return 0;
}

gcc -std=c17 -Wall -Wextra -Wpedantic -o mem_regions mem_regions.c
./mem_regions

```

Hello, World — anatomy line by line

```

#include <stdio.h>

int main(void) {
    printf("Hello, World!\n");
    return 0;
}

```

Line	Role
<code>#include <stdio.h></code>	Preprocessor pulls in declarations for standard I/O, including <code>printf</code> . Angle brackets mean “search system include paths.”
<code>int main(void)</code>	Program entry from the C language’s point of view. <code>int</code> is the return type. <code>(void)</code> means “no parameters” (preferred over empty <code>()</code> in C).
<code>{ ... }</code>	Function body. Execution of <code>main</code> starts at the first statement.
<code>printf("Hello, World!\n");</code>	Call into the C library to write to standard output. <code>\n</code> is a newline. Statement ends with <code>;</code> .

Line	Role
<code>return 0;</code>	Exit status for the process. By convention, 0 means success; non-zero means failure.

Small but important details

- **C is case-sensitive:** `Printf` is not `printf`.
- **Headers declare; libraries define:** including `<stdio.h>` is not optional if you call `printf`—without a declaration, modern compilers reject or warn about implicit declarations.
- **main forms:** `int main(void)` or `int main(int argc, char *argv[])` are the portable forms. Returning from `main` is equivalent to calling `exit` with that status.
- **In C99+,** falling off the end of `main` without a `return` is treated as returning 0. Still write `return 0;` while you are learning—clarity beats cleverness.

Compile and run

```
gcc -std=c17 -Wall -Wextra -Wpedantic -o hello hello.c
./hello
echo $?      # shell shows exit status: 0 on success
```

Full mini-programs

These programs are complete and compilable. Type them yourself; do not only read them.

1. Temperature converter (°C °F)

```
#include <stdio.h>

int main(void) {
    int choice;
    double value;
    double result;

    printf("Temperature converter\n");
    printf("  1) Celsius to Fahrenheit\n");
    printf("  2) Fahrenheit to Celsius\n");
    printf("Choice: ");

    if (scanf("%d", &choice) != 1) {
```

```

        fprintf(stderr, "Invalid choice input\n");
        return 1;
    }

    printf("Value: ");
    if (scanf("%lf", &value) != 1) {
        fprintf(stderr, "Invalid number\n");
        return 1;
    }

    if (choice == 1) {
        result = value * 9.0 / 5.0 + 32.0;
        printf("%.2f C = %.2f F\n", value, result);
    } else if (choice == 2) {
        result = (value - 32.0) * 5.0 / 9.0;
        printf("%.2f F = %.2f C\n", value, result);
    } else {
        fprintf(stderr, "Choice must be 1 or 2\n");
        return 1;
    }

    return 0;
}

gcc -std=c17 -Wall -Wextra -Wpedantic -o temp_convert temp_convert.c
./temp_convert

```

Notes:

- `scanf("%lf", &value)` is required for double (%f is for float with `scanf`).
- Always check that `scanf` returned the number of items you expected.

2. Maximum of three integers

```

#include <stdio.h>

int main(void) {
    int a, b, c;
    int max;

    printf("Enter three integers: ");
    if (scanf("%d %d %d", &a, &b, &c) != 3) {
        fprintf(stderr, "Expected three integers\n");
        return 1;
    }
}

```

```

    max = a;
    if (b > max) {
        max = b;
    }
    if (c > max) {
        max = c;
    }

    printf("max(%d, %d, %d) = %d\n", a, b, c, max);
    return 0;
}

```

```

gcc -std=c17 -Wall -Wextra -Wpedantic -o max3 max3.c
./max3
# sample input: 9 -2 7

```

This is deliberately free of arrays and functions so you can focus on comparison logic. Later modules rewrite the same idea with loops and helpers.

3. Simple calculator CLI

```

#include <stdio.h>

int main(void) {
    double left, right, result;
    char op;

    printf("Enter expression (e.g. 3.5 * 2): ");
    if (scanf("%lf %c %lf", &left, &op, &right) != 3) {
        fprintf(stderr, "Could not parse expression\n");
        return 1;
    }

    switch (op) {
        case '+':
            result = left + right;
            break;
        case '-':
            result = left - right;
            break;
        case '*':
            result = left * right;
            break;
        case '/':
            if (right == 0.0) {

```

```

        fprintf(stderr, "Division by zero\n");
        return 1;
    }
    result = left / right;
    break;
default:
    fprintf(stderr, "Unsupported operator '%c'\n", op);
    return 1;
}

printf("%.10g %c %.10g = %.10g\n", left, op, right, result);
return 0;
}

```

```

gcc -std=c17 -Wall -Wextra -Wpedantic -o calc calc.c
./calc
# sample input: 12.5 / 2.5

```

%.10g prints a readable floating-point form without forcing fixed decimals. Division by zero is checked for / only—floating-point edge cases (NaN, overflow) come later.

4. Exit status practice

```

#include <stdio.h>
#include <stdlib.h>

int main(void) {
    int code;

    printf("Enter exit code (0-255 recommended): ");
    if (scanf("%d", &code) != 1) {
        return EXIT_FAILURE;
    }

    if (code == 0) {
        printf("Reporting success\n");
        return EXIT_SUCCESS;
    }

    fprintf(stderr, "Reporting failure with code %d\n", code);
    return code;
}

```

```
gcc -std=c17 -Wall -Wextra -Wpedantic -o exitdemo exitdemo.c
./exitdemo
echo $?
```

Common misconceptions

“C is basically the same as C++”

No. C++ accepts a large amount of C-looking code, but they are different languages:

- C has no classes, references, templates, or RAII.
- Valid C is not always valid C++ (`void *` assignments, designated initializers historically differed, `restrict`, compound literals, etc.).
- Compile C with a C compiler and `-std=c17`. Do not learn C by pasting into a C++ project and “fixing until it builds.”

“If it compiles, it is correct”

C allows **undefined behavior** (UB): the standard places no requirements on what the program may do. Classic examples include signed integer overflow, using an uninitialized value, buffer overruns, and dangling pointers. The program may crash, appear to work, or corrupt data later—and the behavior can change with optimization level.

```
/* Illustrating a class of bug - do NOT treat the result as meaningful */
#include <stdio.h>

int main(void) {
    int x;
    /* x is indeterminate; reading it is undefined behavior */
    /* printf("%d\n", x); */
    printf("Uncommenting the printf above would be undefined behavior.\n");
    return 0;
}
```

Enable warnings and treat them seriously:

```
gcc -std=c17 -Wall -Wextra -Wpedantic -O2 -o prog prog.c
```

Later: AddressSanitizer (`-fsanitize=address`), undefined-behavior sanitizer (`-fsanitize=undefined`), and tools like Valgrind.

“The compiler is just a translator; any compiler is fine”

Compilers differ in diagnostics, optimization, and default dialects. A program that “works” with loose defaults may fail under stricter flags or another vendor. Standardize on:

```
gcc -std=c17 -Wall -Wextra -Wpedantic
```

Add `-Werror` in CI when you want warnings to fail the build. Clang is an excellent alternative with often clearer messages; the same flags apply in spirit (`clang -std=c17 -Wall -Wextra -Wpedantic`).

“I need to master every standard library function first”

You need a core loop: edit → compile with warnings → run → fix. Learn I/O, types, control flow, functions, then pointers and memory. The library is large; this book introduces pieces when you have a use for them.

“void main() is fine”

Not for hosted (normal OS) programs. Use `int main(void)` or `int main(int argc, char *argv[])`.

How this book is organized

The C book under `books/C/` is a modular library, not a 90-day calendar:

Module	Path	Focus
1 Foundations	<code>content/01-foundations/</code>	Tooling, first programs, console I/O
2 Data types	<code>content/02-data-types/</code>	Types, variables, operators, conversions
3 Control flow	<code>content/03-control-flow/</code>	<code>if</code> , loops, <code>switch</code> , jumps
4 Functions	<code>content/04-functions/</code>	Design, scope, modularity, libc overview
5 Arrays & strings	<code>content/05-arrays-strings/</code>	Contiguous data and text
6 Pointers	<code>content/06-pointers/</code>	Addresses, dynamic memory, pitfalls
7 Structures	<code>content/07-structures/</code>	<code>struct</code> , <code>union</code> , <code>enum</code> , simple ADTs
8 File I/O	<code>content/08-file-io/</code>	Streams and system-facing I/O
9 Modern C	<code>content/09-modern-c/</code>	C99–C23 and modern practice

Module	Path	Focus
10 Data structures	content/10-data-structures/	Lists, trees, tables, graphs, sorting
11 Network	content/11-network/	Sockets and protocols
12 Embedded	content/12-embedded/	Bare-metal-oriented techniques
13 Performance	content/13-performance/	Measurement and optimization
14 Testing	content/14-testing/	Debug, test, harden
15 Advanced	content/15-advanced/	Threads, systems topics, capstone
Workbook	content/99-workbook/	Extra demos, exercises, projects
Appendices	content/100-appendices/	Resources, reference, solution notes

What to practice first

1. Install a compiler and confirm `gcc --version` (or `clang --version`).
2. Compile Hello World with `-std=c17 -Wall -Wextra -Wpedantic`.
3. Type the temperature converter and calculator by hand; break them on purpose and read the diagnostics.
4. Complete the exercises at the end of this chapter before racing ahead to pointers.
5. Use the workbook ([content/99-workbook/examples/01-foundations/](#)) as optional extra drills—not a substitute for typing chapter programs.

Recommended daily loop for Module 1:

```
read a section → type one full program → compile with warnings → run →
change one thing → explain the new behavior out loud or in a comment
```

Applications of C (why the language stays relevant)

- **Operating systems and runtimes:** kernels, process loaders, language VMs.
- **Embedded and IoT:** microcontrollers, sensors, firmware with tight RAM/flash budgets.
- **Infrastructure:** databases, web servers, container tooling, networking stacks.
- **Performance-critical domains:** games, HFT-adjacent tooling, scientific kernels, codecs.
- **Interoperability:** foreign-function interfaces almost always speak C ABIs.

Relationship sketch (not a full language survey):

Language	Relation to C
C++	Historically grew from C; different language and type system
Rust / Go / Zig	Modern systems languages; often call C libraries via FFI
Java / C#	Syntax influenced by C-family; managed runtimes

Language	Relation to C
Python / scripting	Many interpreters and extension modules are C under the hood

Summary

- C is a systems-oriented language with a small core and a large ecosystem of libraries and platforms.
- A program is **preprocessed, compiled, assembled, linked**, then loaded as a **process**.
- Running processes have **text, data/BSS, heap, and stack** regions; know what lives where at a high level.
- Hello World teaches includes, **main**, library calls, and exit status.
- Prefer **gcc -std=c17 -Wall -Wextra -Wpedantic**, check I/O return values, and treat undefined behavior as a first-class hazard.
- Work module by module; use the **workbook** for extra practice after the chapter programs.

Next: set up (or verify) your toolchain in *Development Environment Setup*, then deepen compilation and multi-file builds in *Your First C Program*.

Exercises

Do these with **-std=c17 -Wall -Wextra -Wpedantic**. Fix every warning.

1. **Hello variants** — Print three lines: your name, today's focus topic, and the value of `sizeof(int)`.
2. **Stage tour** — For `hello.c`, produce `.i`, `.s`, and `.o` artifacts. Open the `.s` file and find the string `Hello` (or a reference to it). You do not need to understand every instruction.
3. **Exit codes** — Write a program that returns 0 if the user enters an even integer and 1 otherwise. Inspect `$?` in the shell.
4. **Swap print** — Read two integers and print them in reverse order without using a third variable *for output only* is optional; using a temporary is fine. Focus on correct **scanf** checks.
5. **Fahrenheit table** — Print a table of Celsius from 0 to 100 in steps of 10 with the Fahrenheit equivalent (no user input required).
6. **Max of three** — **self-test** — Extend the max program to also print the minimum.
7. **Calculator harden** — Reject invalid operators and division by zero; print errors to **stderr**.
8. **Comment audit** — Take any mini-program above and add brief comments describing each stage of execution in **main** (not line-by-line noise).
9. **Broken on purpose** — Remove `#include <stdio.h>` and compile. Record the diagnostic. Restore it. Then misspell `main` as `mian` and record what happens.

10. **Mental model** — In 5–8 sentences, explain the difference between an object file (.o) and an executable, and between stack storage and heap storage.

Optional stretch: time yourself rewriting the temperature converter from memory after a short break. Accuracy and clean compile matter more than speed.

Development Environment Setup

Introduction

Setting up a proper development environment is crucial for effective C programming. This chapter walks through installing compilers, choosing editors, configuring debuggers, and — most importantly — **real day-to-day workflows**: multi-file builds, Make, a mini CMake project, and GCC vs Clang diagnostics.

Default compile line used throughout this book:

```
gcc -std=c17 -Wall -Wextra -Wpedantic -o prog prog.c
```

Clang is a fine drop-in for learning:

```
clang -std=c17 -Wall -Wextra -Wpedantic -o prog prog.c
```

Installing Compilers

GCC (GNU Compiler Collection)

GCC is the most widely used C compiler, especially on Linux and macOS systems.

Linux Installation

```
# Ubuntu/Debian
sudo apt update
sudo apt install build-essential

# Fedora/RHEL/CentOS
sudo dnf install gcc gcc-c++ make

# Arch Linux
sudo pacman -S gcc make
```

macOS Installation

Install Xcode Command Line Tools:

```
xcode-select --install
```

Or use Homebrew:

```
brew install gcc
```

Windows Installation

1. Install MinGW-w64: <https://www.mingw-w64.org/>
2. Or use MSYS2: <https://www.msys2.org/>
3. Or install Visual Studio with C/C++ support

Clang

Clang is a modern compiler known for excellent error messages and fast compilation.

Installation

```
# Ubuntu/Debian
sudo apt install clang

# macOS (via Homebrew)
brew install llvm

# Windows (via MSYS2)
pacman -S mingw-w64-x86_64-clang
```

Microsoft Visual C++ (MSVC)

For Windows development, MSVC is Microsoft's official C/C++ compiler.

Installation

1. Download Visual Studio Community (free): <https://visualstudio.microsoft.com/>
2. During installation, select "Desktop development with C++" workload

Modern IDEs and Editors

Visual Studio Code

A lightweight but powerful editor with excellent C/C++ support.

Setup Steps

1. Download and install VS Code: <https://code.visualstudio.com/>
2. Install the C/C++ extension by Microsoft
3. Install additional helpful extensions:
 - C/C++ Extension Pack
 - Code Runner
 - GitLens

Configuration

Create a `.vscode` directory in your project with these files:

`tasks.json` (for building):

```
{
  "version": "2.0.0",
  "tasks": [
    {
      "type": "cppbuild",
      "label": "C/C++: gcc build active file",
      "command": "/usr/bin/gcc",
      "args": [
        "-fdiagnostics-color=always",
        "-g",
        "${file}",
        "-o",
        "${fileDirname}/${fileBasenameNoExtension}",
        "-lm"
      ],
      "options": {
        "cwd": "${fileDirname}"
      },
      "problemMatcher": ["$gcc"],
      "group": {
        "kind": "build",
        "isDefault": true
      },
      "detail": "Compiler: /usr/bin/gcc"
    }
  ]
}
```

```
    ]
}
```

launch.json (for debugging):

```
{
  "version": "0.2.0",
  "configurations": [
    {
      "name": "C/C++: gcc build and debug active file",
      "type": "cppdbg",
      "request": "launch",
      "program": "${fileDirname}/${fileBasenameNoExtension}",
      "args": [],
      "stopAtEntry": false,
      "cwd": "${fileDirname}",
      "environment": [],
      "externalConsole": false,
      "MIMode": "gdb",
      "setupCommands": [
        {
          "description": "Enable pretty-printing for gdb",
          "text": "-enable-pretty-printing",
          "ignoreFailures": true
        }
      ],
      "preLaunchTask": "C/C++: gcc build active file",
      "miDebuggerPath": "/usr/bin/gdb"
    }
  ]
}
```

CLion

A full-featured IDE specifically designed for C/C++ development.

Features

- Intelligent code completion
- Integrated debugger
- Built-in terminal
- Version control integration
- CMake support

Installation

1. Download from JetBrains: <https://www.jetbrains.com/clion/>
2. Install with a free student license if you're a student
3. Configure toolchain to point to your compiler

Code::Blocks

A free, open-source C/C++ IDE.

Installation

```
# Ubuntu/Debian
sudo apt install codeblocks

# macOS (via Homebrew)
brew install --cask codeblocks

# Windows
Download from http://www.codeblocks.org/
```

Command-Line Tools

Make

A build automation tool that automatically determines which pieces of a program need to be recompiled.

Installation

```
# Usually comes with build-essential on Linux
sudo apt install make

# macOS (via Xcode Command Line Tools)
xcode-select --install

# Windows (via MSYS2)
pacman -S make
```

CMake

A cross-platform build system generator.

Installation

```
# Ubuntu/Debian
sudo apt install cmake

# macOS (via Homebrew)
brew install cmake

# Windows
Download from https://cmake.org/download/
```

Git

Version control system essential for modern development.

Installation

```
# Ubuntu/Debian
sudo apt install git

# macOS
xcode-select --install

# Windows
Download from https://git-scm.com/
```

Debugging Tools

GDB (GNU Debugger)

The standard debugger for GCC.

Installation

```
# Ubuntu/Debian
sudo apt install gdb

# macOS
# Comes with Xcode Command Line Tools

# Windows (via MSYS2)
pacman -S mingw-w64-x86_64-gdb
```

Basic Usage

```
# Compile with debug symbols
gcc -g -o program program.c

# Start debugging
gdb program

# Common GDB commands
(gdb) run                # Run the program
(gdb) break main         # Set breakpoint at main function
(gdb) step               # Step into function
(gdb) next               # Step over function
(gdb) print variable_name # Print variable value
(gdb) continue           # Continue execution
(gdb) quit               # Exit GDB
```

Valgrind

A powerful tool for detecting memory leaks and errors.

Installation

```
# Ubuntu/Debian
sudo apt install valgrind

# macOS
brew install valgrind

# Note: Valgrind has limited support on macOS
```

Usage

```
# Compile with debug symbols
gcc -g -o program program.c

# Run with Valgrind
valgrind --leak-check=full ./program
```

AddressSanitizer

A fast memory error detector built into GCC and Clang.

Usage

```
# Compile with AddressSanitizer
gcc -fsanitize=address -g -o program program.c

# Run normally - errors will be reported automatically
./program
```

Package Management

vcpkg

Microsoft's C++ package manager that works with C as well.

Installation

```
git clone https://github.com/Microsoft/vcpkg.git
cd vcpkg
./bootstrap-vcpkg.sh # Linux/macOS
./bootstrap-vcpkg.bat # Windows
```

Conan

A decentralized C/C++ package manager.

Installation

```
pip install conan
```

Cloud Development Environments

GitHub Codespaces

Browser-based development environment with full VS Code features.

Setup

1. Create a GitHub repository
2. Go to repository settings
3. Enable Codespaces
4. Create a dev container configuration

Gitpod

Ready-to-code development environments in the cloud.

Setup

1. Create a Gitpod account
2. Add a `.gitpod.yml` file to your repository:

```
image: gitpod/workspace-full
tasks:
  - init: echo "Initializing project"
    command: echo "Starting development environment"
vscode:
  extensions:
    - ms-vscode.cpptools
```

Cross-Platform Development

Windows Subsystem for Linux (WSL)

Run Linux development tools on Windows.

Installation

1. Open PowerShell as Administrator
2. Run: `wsl --install`
3. Restart your computer
4. Install Ubuntu from Microsoft Store

Docker

Containerize your development environment.

Example Dockerfile

```
FROM gcc:latest

WORKDIR /app
COPY . .

RUN gcc -o program program.c

CMD ["/program"]
```

Usage

```
# Build the container
docker build -t c-program .

# Run the program
docker run c-program
```

Real Workflows: From One File to a Small Project

Installation is only half the job. This section is the **daily loop** you will use for the rest of the book.

Single-file recipe

```
# write hello.c, then:
gcc -std=c17 -Wall -Wextra -Wpedantic -o hello hello.c
./hello
echo $?          # exit status: 0 means success
```

Debug-friendly variant:

```
gcc -std=c17 -Wall -Wextra -g -O0 -o hello hello.c
gdb ./hello      # or lldb on macOS
```

Release-ish variant:

```
gcc -std=c17 -Wall -Wextra -O2 -o hello hello.c
```

Multi-file build recipes

Suppose you have:

util.h
util.c
main.c

util.h

```
#ifndef UTIL_H
#define UTIL_H

int add(int a, int b);
void greet(const char *name);

#endif
```

util.c

```
#include "util.h"
#include <stdio.h>

int add(int a, int b) {
    return a + b;
}

void greet(const char *name) {
    if (name == NULL) {
        name = "(null)";
    }
    printf("Hello, %s!\n", name);
}
```

main.c

```
#include <stdio.h>
#include "util.h"

int main(void) {
    greet("C developer");
    printf("2 + 3 = %d\n", add(2, 3));
    return 0;
}
```

Recipe A — one-shot link

```
gcc -std=c17 -Wall -Wextra -Wpedantic -o app main.c util.c
./app
```

Recipe B — separate compile, then link (scales better)

```
gcc -std=c17 -Wall -Wextra -Wpedantic -c main.c -o main.o
gcc -std=c17 -Wall -Wextra -Wpedantic -c util.c -o util.o
gcc -o app main.o util.o
./app
```

When you change only `util.c`, recompile that object and re-link:

```
gcc -std=c17 -Wall -Wextra -Wpedantic -c util.c -o util.o
gcc -o app main.o util.o
```

Recipe C — show each stage (learning aid)

```
gcc -std=c17 -E util.c -o util.i      # preprocess
gcc -std=c17 -S util.c -o util.s      # assembly
gcc -std=c17 -c util.c -o util.o      # object
nm util.o                             # symbols defined/referenced
```

Full program: tiny CLI calculator (single file)

Type this, compile it, break it on purpose, fix warnings.

```
/* file: calc.c */
#include <stdio.h>
#include <stdlib.h>
#include <string.h>

static void usage(const char *prog) {
    fprintf(stderr, "Usage: %s <add|sub|mul|div> <a> <b>\n", prog);
}

int main(int argc, char *argv[]) {
    double a, b, result;
    char *end;

    if (argc != 4) {
```



```

        usage(argv[0]);
        return EXIT_FAILURE;
    }

    a = strtod(argv[2], &end);
    if (end == argv[2] || *end != '\0') {
        fprintf(stderr, "invalid number: %s\n", argv[2]);
        return EXIT_FAILURE;
    }
    b = strtod(argv[3], &end);
    if (end == argv[3] || *end != '\0') {
        fprintf(stderr, "invalid number: %s\n", argv[3]);
        return EXIT_FAILURE;
    }

    if (strcmp(argv[1], "add") == 0) {
        result = a + b;
    } else if (strcmp(argv[1], "sub") == 0) {
        result = a - b;
    } else if (strcmp(argv[1], "mul") == 0) {
        result = a * b;
    } else if (strcmp(argv[1], "div") == 0) {
        if (b == 0.0) {
            fprintf(stderr, "division by zero\n");
            return EXIT_FAILURE;
        }
        result = a / b;
    } else {
        usage(argv[0]);
        return EXIT_FAILURE;
    }

    printf("%.10g\n", result);
    return EXIT_SUCCESS;
}

```

```

gcc -std=c17 -Wall -Wextra -Wpedantic -o calc calc.c
./calc add 2 3
./calc div 10 4
./calc div 1 0    # should fail cleanly

```

GCC vs Clang (What Actually Differs)

Both implement the C language and understand the same core flags. Prefer **either** for learning; use both occasionally so diagnostics look familiar.

Topic	GCC	Clang
Typical binary	<code>gcc</code>	<code>clang</code>
Error style	Compact, classic	Often more “caret” detail / notes
Sanitizers	Excellent (<code>-fsanitize=...</code>)	Excellent (often slightly clearer reports)
Static analysis extras	<code>-fanalyzer</code> (GCC)	<code>clang --analyze</code> / clang-tidy ecosystem
Default on macOS	Homebrew <code>gcc</code> or Xcode’s <code>clang</code> as <code>gcc</code> alias	Xcode CLT is Clang under the hood

Side-by-side on the same bug

```
/* file: warn_me.c - deliberately sloppy */
#include <stdio.h>

int main(void) {
    int x;
    printf("%d\n", x);    /* uninitialized use */
    return 0;
}
```

```
gcc -std=c17 -Wall -Wextra -Wpedantic -o warn_me warn_me.c
clang -std=c17 -Wall -Wextra -Wpedantic -o warn_me warn_me.c
```

Compare messages. Fix:

```
int x = 0;
```

Useful shared flags

```
# Treat warnings as errors once you are comfortable
gcc -std=c17 -Wall -Wextra -Wpedantic -Werror -o prog prog.c

# Extra useful warnings (optional)
gcc -std=c17 -Wall -Wextra -Wpedantic -Wshadow -Wconversion -o prog prog.c

# Sanitizers (Linux/macOS; great for learning memory bugs)
```

```
gcc -std=c17 -Wall -Wextra -g -O1 -fsanitize=address,undefined -o prog prog.c
```

macOS note

gcc --version on stock macOS often prints **Apple Clang**. For real GCC:

```
brew install gcc
# then use gcc-14 / g++-14 (version may differ)
```

Make: Multi-File Builds Without Typing Everything

Minimal Makefile for the util/main project

```
# Makefile
CC      = gcc
CFLAGS  = -std=c17 -Wall -Wextra -Wpedantic
LDFLAGS =
TARGET  = app
OBJS     = main.o util.o

.PHONY: all clean run

all: $(TARGET)

$(TARGET): $(OBJS)
    $(CC) $(LDFLAGS) -o $@ $(OBJS)

%.o: %.c
    $(CC) $(CFLAGS) -c $< -o $@

# Rebuild if headers change (simple approach)
main.o: main.c util.h
util.o: util.c util.h

run: $(TARGET)
    ./$$(TARGET)

clean:
    rm -f $(OBJS) $(TARGET)
```

```
make
make run
make clean
```

Full program + Make: word count of argv strings

wcargs.c

```
/* file: wcargs.c */
#include <stdio.h>
#include <string.h>

int main(int argc, char *argv[]) {
    int i;
    size_t total = 0;

    for (i = 1; i < argc; i++) {
        size_t n = strlen(argv[i]);
        printf("%zu\t%s\n", n, argv[i]);
        total += n;
    }
    printf("%zu\ttotal chars (no spaces between args)\n", total);
    return 0;
}

# one-file Makefile variant
CC = gcc
CFLAGS = -std=c17 -Wall -Wextra -Wpedantic
wcargs: wcargs.c
    $(CC) $(CFLAGS) -o $@ $<

make wcargs
./wcargs hello world
```

CMake Mini-Project (Practical One-Pager)

CMake generates native build files (Make, Ninja, Xcode, VS). For a tiny C project:

```
mini_cmake/
  CMakeLists.txt
  include/
```

```
    util.h
src/
    util.c
    main.c
```

CMakeLists.txt

```
cmake_minimum_required(VERSION 3.16)
project(mini_app C)

set(CMAKE_C_STANDARD 17)
set(CMAKE_C_STANDARD_REQUIRED ON)
set(CMAKE_C_EXTENSIONS OFF)

add_compile_options(-Wall -Wextra -Wpedantic)

add_executable(mini_app
    src/main.c
    src/util.c
)

target_include_directories(mini_app PRIVATE include)
```

Copy the same `util.h` / `util.c` / `main.c` content from the multi-file section (put `util.h` under `include/`, sources under `src/`, and use `#include "util.h"`).

Configure, build, run

```
cd mini_cmake
cmake -S . -B build
cmake --build build
./build/mini_app
```

Debug build:

```
cmake -S . -B build -DCMAKE_BUILD_TYPE=Debug
cmake --build build
```

Out-of-source build/ keeps sources clean. Delete build/ anytime and reconfigure.

Optional: `compile_commands.json` for editors

```
cmake -S . -B build -DCMAKE_EXPORT_COMPILE_COMMANDS=ON
# many editors read build/compile_commands.json (or a symlink at repo root)
```

Editor + Debugger One-Pager

You do not need a heavy IDE. You need: **edit** → **build** → **run under a debugger**.

Terminal editor + GDB (always works)

```
# 1. edit files in vim/nano/emacs/VS Code
# 2. build with symbols
gcc -std=c17 -Wall -Wextra -g -O0 -o app main.c util.c

# 3. debug
gdb ./app
```

Essential GDB session:

```
(gdb) break main
(gdb) run
(gdb) next          # step over
(gdb) step          # step into
(gdb) print a
(gdb) backtrace     # or: bt
(gdb) continue
(gdb) quit
```

LLDB (macOS default)

```
clang -std=c17 -Wall -Wextra -g -O0 -o app main.c util.c
lldb ./app
```

```
(lldb) b main
(lldb) run
(lldb) n
(lldb) p a
(lldb) bt
(lldb) c
(lldb) q
```

VS Code (minimum useful setup)

1. Install C/C++ extension (Microsoft) or clangd-based setup.
2. Ensure gcc/clang and gdb/lldb are on PATH.
3. Use the earlier `tasks.json` / `launch.json`, or just open an integrated terminal and run `make + gdb`.
4. Prefer a multi-file task that builds the whole project, not only the active file, once you leave single-file labs.

GDB against a real bug (complete mini-lab)

```
/* file: off_by_one.c */
#include <stdio.h>

static int sum_first_n(const int *a, int n) {
    int total = 0;
    int i;
    /* BUG: i <= n reads one past the end when n is the count */
    for (i = 0; i <= n; i++) {
        total += a[i];
    }
    return total;
}

int main(void) {
    int data[5] = {1, 2, 3, 4, 5};
    printf("sum = %d\n", sum_first_n(data, 5));
    return 0;
}
```

```
gcc -std=c17 -Wall -Wextra -g -O0 -o off_by_one off_by_one.c
gdb ./off_by_one
```

```
(gdb) break sum_first_n
(gdb) run
(gdb) print n
(gdb) print i
(gdb) next
# watch i go to 5 while valid indices are 0..4
```

Fix the loop to `i < n`, rebuild, confirm `sum = 15`.

Sanitizers and Memory Tools (Environment Checklist)

```
# Address + undefined behavior (recommended learning default)
gcc -std=c17 -Wall -Wextra -g -O1 -fsanitize=address,undefined -o app main.c util.c
./app

# Valgrind Memcheck (Linux; limited/absent on modern macOS)
gcc -std=c17 -Wall -Wextra -g -O0 -o app main.c util.c
valgrind --leak-check=full ./app
```

Keep a personal alias if you like:

```
alias c17='gcc -std=c17 -Wall -Wextra -Wpedantic'
alias c17san='gcc -std=c17 -Wall -Wextra -g -O1 -fsanitize=address,undefined'
```

Modern Development Practices

Version Control Workflow

1. Initialize a Git repository: `git init`
2. Create a `.gitignore` for C projects:

```
# Compiled objects
*.o
*.obj

# Executables
*.exe
*.out
a.out
app
calc

# Debug files
*.dSYM/
core
core.*

# Build directories
build/
cmake-build-*/
```



```
compile_commands.json
```

```
# Editor
.vscode/
.idea/
```

3. Commit early, commit small: “compiles with -Wall” is a good checkpoint.

Project Structure

```
project/
  src/           # Source files
  include/       # Header files
  tests/         # Test files
  docs/          # Documentation
  build/         # Build output (gitignored)
  CMakeLists.txt # or Makefile
  README.md
```

Continuous Integration

```
# .github/workflows/c-ci.yml
name: C CI

on: [push, pull_request]

jobs:
  build:
    runs-on: ubuntu-latest
    steps:
      - uses: actions/checkout@v4
      - name: Compile
        run: gcc -std=c17 -Wall -Wextra -Wpedantic -Werror -o app src/*.c -I include
      - name: Run
        run: ./app
```

Verify Your Setup (Smoke Tests)

Run these once after installing tools:

```
# Compilers
gcc --version
```

```

clang --version    # if installed

# Build tools
make --version
cmake --version

# Debuggers
gdb --version      # or: lldb --version

# Sanitizer smoke
printf '%s\n' '#include <stdio.h>' 'int main(void){puts("ok");}' > t.c
gcc -std=c17 -Wall -Wextra -fsanitize=address,undefined -o t t.c && ./t

```

Expected: versions print without error; ./t prints ok.

Exercises

Compile every program with:

```
gcc -std=c17 -Wall -Wextra -Wpedantic -o <name> <sources...>
```

1. **Hello toolchain** — Compile the same `hello.c` with both `gcc` and `clang` (if available). Note any diagnostic differences when you introduce an unused variable.
 2. **Multi-file link** — Create `add.h` / `add.c` / `main.c`. Build with Recipe A and Recipe B. Delete `add.o` and rebuild only what is needed.
 3. **Makefile** — Write a `Makefile` that builds the multi-file app, supports `make clean` and `make run`, and rebuilds when a header changes.
 4. **CMake mini** — Recreate the `mini_cmake` layout. Configure an out-of-source `build/`, build, and run. Then change `util.c`, rebuild, and confirm only the needed objects recompile (watch timestamps / build log).
 5. **Debugger lab** — Use GDB or LLDB on `off_by_one.c`. Set a breakpoint, print `i` and `a[i]`, fix the bug, and show the correct sum.
 6. **Sanitizer lab** — Write a program that does `int *p = NULL; *p = 1;`. Run it normally (crash), then under AddressSanitizer. Paste the first few lines of the ASan report into a comment in your fixed file.
 7. **calc CLI** — Extend `calc.c` with a `mod` operation using integers (`strtol`) and reject non-integers cleanly.
 8. **gitignore** — Initialize git in a throwaway folder, build an app, confirm artifacts are ignored, and make one commit of sources only.
-

Summary

A solid C environment is:

1. **Compiler** — GCC and/or Clang with `-std=c17 -Wall -Wextra`
2. **Editor** — anything that edits text reliably; IDE optional
3. **Debugger** — GDB/LLDB with `-g -O0` builds
4. **Build tools** — one-shot `gcc`, then `Make`, then `CMake` as projects grow
5. **Sanitizers / Valgrind** — treat memory and UB as first-class
6. **Git** — ignore build products; commit sources

Master the multi-file and debugger loops now; every later module assumes you can edit, build, run, and inspect without friction.

Your First C Program

Introduction

This chapter moves from “what C is” to a working edit–compile–run loop. You will write complete programs, read compiler diagnostics, split a tiny program across multiple files, and finish a short lab (personal info printer + unit converter).

Default compile line for every example:

```
gcc -std=c17 -Wall -Wextra -Wpedantic -o prog prog.c
```

The classic “Hello, World!”

```
#include <stdio.h>

int main(void) {
    printf("Hello, World!\n");
    return 0;
}
```

Line-by-line

Piece	Meaning
<code>#include <stdio.h></code>	Bring in declarations for standard I/O (<code>printf</code> , <code>scanf</code> , ...).
<code>int main(void)</code>	Entry point; returns an <code>int</code> status; no parameters.
<code>printf(...)</code>	Library call writing to standard output.
<code>\n</code>	Newline character in the string.
<code>return 0;</code>	Report success to the operating system / shell.

Create, compile, run

```
# create hello.c with your editor, then:
gcc -std=c17 -Wall -Wextra -Wpedantic -o hello hello.c
./hello
```

On Windows (MinGW / similar), the executable may be `hello.exe`; run `hello` or `.\hello.exe` depending on your shell.

What the compiler does (short version)

`hello.c` → preprocess → compile → assemble → link → `hello`

Stage	Flag to stop here	Output
Preprocess	<code>-E</code>	Preprocessed source (<code>.i</code>)
Compile to assembly	<code>-S</code>	Assembly (<code>.s</code>)
Assemble only	<code>-c</code>	Object file (<code>.o</code>)
Full link	(default)	Executable

```
gcc -std=c17 -E hello.c -o hello.i
gcc -std=c17 -S hello.c -o hello.s
gcc -std=c17 -c hello.c -o hello.o
gcc -o hello hello.o
```

Useful flags you will reuse:

```
gcc -std=c17 -Wall -Wextra -Wpedantic -g -O0 -o hello hello.c # debug-friendly
gcc -std=c17 -Wall -Wextra -Wpedantic -O2 -o hello hello.c   # optimized release-ish
```

Flag	Role
<code>-std=c17</code>	Language dialect
<code>-Wall -Wextra -Wpedantic</code>	Strong warnings
<code>-g</code>	Debug symbols (GDB/LLDB)
<code>-O0 / -O2</code>	Optimization level
<code>-o name</code>	Output file name

Structure of a C program

```
/* 1. Preprocessor directives */
#include <stdio.h>
#define APP_NAME "demo"

/* 2. Optional global declarations / prototypes */
void greet(const char *who);

/* 3. main - required for hosted programs */
int main(void) {
    greet("World");
    return 0;
}

/* 4. Other function definitions */
void greet(const char *who) {
    printf("Hello, %s!\n", who);
}
```

Order rules of thumb:

1. Includes and macros at the top.
2. Prototypes before first use (or define functions before `main`).
3. One clear `main`.
4. Keep globals rare; prefer locals and parameters.

Comments

```
// Single-line comment (C99+)

/*
 * Multi-line comment
 * useful for file headers or longer notes
 */
```

Prefer comments that explain **why** or non-obvious constraints. Restating `i++` as “increment `i`” adds noise.

More complete single-file examples

Greeting with command-line arguments

```
#include <stdio.h>

int main(int argc, char *argv[]) {
    if (argc < 2) {
        printf("Usage: %s <name>\n", argv[0]);
        return 1;
    }

    printf("Hello, %s!\n", argv[1]);
    printf("Argument count: %d\n", argc);
    return 0;
}

gcc -std=c17 -Wall -Wextra -Wpedantic -o greeter greeter.c
./greeter Ada
./greeter          # should print usage and exit non-zero
```

`argv[0]` is usually the program name as invoked. `argc` is the count of elements in `argv`.

Sum of two integers (with validation)

```
#include <stdio.h>

int main(void) {
    int a, b;

    printf("Enter two integers: ");
    if (scanf("%d %d", &a, &b) != 2) {
        fprintf(stderr, "error: expected two integers\n");
        return 1;
    }

    printf("%d + %d = %d\n", a, b, a + b);
    return 0;
}
```

Multiple helpers in one file

```
#include <stdio.h>

static int square(int x) {
    return x * x;
}

static int max2(int a, int b) {
    return (a > b) ? a : b;
}

int main(void) {
    int x = 5;
    int y = 12;

    printf("square(%d) = %d\n", x, square(x));
    printf("max(%d, %d) = %d\n", x, y, max2(x, y));
    return 0;
}
```

`static` on a function at file scope means “internal linkage”: the name is not exported to the linker for other `.c` files. That keeps helpers private to the translation unit.

Print a small banner

```
#include <stdio.h>

int main(void) {
    puts("*****");
    puts("*   My First C Program   *");
    puts("*****");
    return 0;
}
```

`puts` appends a newline for you and is convenient for simple strings. Prefer `printf` when you need formatting.

Multi-file tiny program

Real projects split declarations (headers) from definitions (source files). Minimal pattern:

`util.h`


```

#ifndef UTIL_H
#define UTIL_H

int add(int a, int b);
void print_sum(int a, int b);

#endif /* UTIL_H */

```

util.c

```

#include <stdio.h>
#include "util.h"

int add(int a, int b) {
    return a + b;
}

void print_sum(int a, int b) {
    printf("%d + %d = %d\n", a, b, add(a, b));
}

```

main.c

```

#include "util.h"

int main(void) {
    print_sum(3, 4);
    return 0;
}

```

Build it

```

# Separate compile + link
gcc -std=c17 -Wall -Wextra -Wpedantic -c util.c -o util.o
gcc -std=c17 -Wall -Wextra -Wpedantic -c main.c -o main.o
gcc -o multifile main.o util.o

# Or one shot
gcc -std=c17 -Wall -Wextra -Wpedantic -o multifile main.c util.c

./multifile

```

Why the include guard?

```
#ifndef UTIL_H
#define UTIL_H
/* declarations */
#endif
```

If two files include `util.h` and a third includes both, guards prevent duplicate declarations in one translation unit.

Header vs source rules

Put in <code>.h</code>	Put in <code>.c</code>
Function prototypes	Function bodies
Shared <code>struct</code> types (later)	<code>static</code> helpers
<code>extern</code> data declarations (rare)	Definitions of globals

Include your own headers with quotes: `#include "util.h"`. System headers use angle brackets: `#include <stdio.h>`.

Common compile errors (with explanations)

Missing semicolon

```
printf("hi\n")    /* error: expected ';' */
```

Fix: statements end with `;`. The diagnostic often points at the *next* line—check the previous statement.

Missing `#include <stdio.h>`

```
int main(void) {
    printf("hi\n");
    return 0;
}
```

Symptom: implicit declaration / unknown identifier `printf`.

Fix: include the proper header.

Wrong main spelling or signature

```
int mian(void) { return 0; }
```

Symptom: linker error undefined reference to main.

Fix: the entry point must be named main.

Case sensitivity

```
Printf("hi\n"); /* wrong */
```

Fix: printf, not Printf.

Undeclared function (multi-file)

main.c calls add but never includes util.h and has no prototype:

Symptom: implicit declaration warning/error, sometimes wrong assumptions about types.

Fix: #include "util.h" or declare int add(int, int); before use.

Multiple definition of a function

Defining int add(int a, int b) { ... } in a header that two .c files include.

Symptom: linker multiple definition of add.

Fix: declare in the header; define in exactly one .c file.

Forgetting to link a file

```
gcc -std=c17 -Wall -Wextra -Wpedantic -o multifile main.c  
# undefined reference to print_sum
```

Fix: include every translation unit: main.c util.c.

Format / type mismatch (warning)

```
int x = 3;  
printf("%f\n", x); /* wrong specifier for int */
```

Fix: use %d for int. Enable -Wall; modern GCC/Clang catch many of these.

Using = instead of == in a condition (logic bug; may still compile)

```
int x = 0;
if (x = 1) { /* assignment; always true here */
    printf("oops\n");
}
```

Mitigation: -Wall often warns about suspicious assignments in conditions. Prefer clear comparisons.

Unterminated string

```
printf("hello\n);
```

Symptom: missing terminating " / stray errors cascading.

Fix: close the string literal; recompile—cascading errors often shrink after the first fix.

Tips for reading diagnostics

1. Start with the **first** error, not the last.
 2. Fix includes and syntax before chasing linker issues.
 3. Recompile after each fix; noise often disappears.
 4. If the message is opaque, reduce to a 10-line file that still fails.
-

Lab A — Personal info printer

Goal: print a short personal profile using variables and formatted output.

```
#include <stdio.h>

int main(void) {
    char name[] = "Ada Lovelace";
    int age = 36;
    char city[] = "London";
    char language[] = "C";
    double years_experience = 2.5;

    printf("=====\n");
    printf("      PERSONAL PROFILE\n");
    printf("=====\n");
    printf("Name:      %s\n", name);
    printf("Age:       %d\n", age);
}
```

```

printf("City:      %s\n", city);
printf("Language:  %s\n", language);
printf("Experience: %.1f years\n", years_experience);
printf("=====\n");

return 0;
}

gcc -std=c17 -Wall -Wextra -Wpedantic -o profile profile.c
./profile

```

Lab stretch:

1. Change the data to your own details.
 2. Add a second `printf` block that re-prints the same data as a single comma-separated line.
 3. Replace hard-coded values with `scanf` / `fgets` input (after the next chapter if needed).
-

Lab B — Unit converter

Goal: convert between a few everyday units with a tiny menu.

```

#include <stdio.h>

int main(void) {
    int choice;
    double input;
    double output;

    printf("Unit converter\n");
    printf("  1) kilometers to miles\n");
    printf("  2) miles to kilometers\n");
    printf("  3) kilograms to pounds\n");
    printf("  4) pounds to kilograms\n");
    printf("Select: ");

    if (scanf("%d", &choice) != 1) {
        fprintf(stderr, "Invalid menu selection\n");
        return 1;
    }

    printf("Value: ");
    if (scanf("%lf", &input) != 1) {
        fprintf(stderr, "Invalid numeric value\n");
    }
}

```

```

        return 1;
    }

    switch (choice) {
    case 1:
        output = input * 0.621371;
        printf("%.4f km = %.4f mi\n", input, output);
        break;
    case 2:
        output = input * 1.60934;
        printf("%.4f mi = %.4f km\n", input, output);
        break;
    case 3:
        output = input * 2.20462;
        printf("%.4f kg = %.4f lb\n", input, output);
        break;
    case 4:
        output = input * 0.453592;
        printf("%.4f lb = %.4f kg\n", input, output);
        break;
    default:
        fprintf(stderr, "Unknown option %d\n", choice);
        return 1;
    }

    return 0;
}

```

```

gcc -std=c17 -Wall -Wextra -Wpedantic -o units units.c
./units

```

Lab stretch:

1. Add Celsius Fahrenheit (reuse logic from the intro chapter).
2. Split conversion formulas into functions in a second file (`convert.c` / `convert.h`) and link them.
3. Loop the menu until the user chooses 0 to quit (preview of control flow).

Variations worth knowing

Implicit return 0 in main (C99+)

```
#include <stdio.h>

int main(void) {
    printf("Hello, World!\n");
}
```

Legal in C99 and later for main only. Prefer an explicit `return` while learning.

Prototype before main

```
#include <stdio.h>

void print_greeting(void);

int main(void) {
    print_greeting();
    return 0;
}

void print_greeting(void) {
    printf("Hello, World!\n");
}
```

Using stderr for errors

```
#include <stdio.h>

int main(void) {
    fprintf(stderr, "This goes to standard error\n");
    printf("This goes to standard output\n");
    return 1;
}
```

```
./prog >out.txt 2>err.txt
```

Beginner mistakes checklist

Mistake	Fix
Missing ;	Terminate statements
Missing header	<code>#include <stdio.h></code> etc.
<code>void main()</code>	Use <code>int main(void)</code>
Calling without prototype	Declare or include header
Ignoring warnings	Compile with <code>-Wall -Wextra -Wpedantic</code>
Wrong file extension / dialect	Save as <code>.c</code> , pass <code>-std=c17</code>
Editing the wrong directory	Confirm path before compile
Running old binary	Rebuild after every edit

Summary

- Every hosted C program needs a **main** and usually standard library headers for I/O.
- Compilation is a pipeline; object files link into executables.
- Multi-file programs separate **declarations** (headers) from **definitions** (sources).
- Learn to read the **first** compiler/linker diagnostic.
- Labs: personal profile printer and unit converter reinforce the full cycle.

Next: *Basic Input and Output* — format strings, **scanf** pitfalls, and safer line input with **fgets**.

Exercises

Compile each with:

```
gcc -std=c17 -Wall -Wextra -Wpedantic -o exN exN.c
```

1. **Banner** — Print a boxed title using only `printf` or `puts` (no user input).
2. **Three-file arithmetic** — Create `mathx.h`, `mathx.c` (`mul`, `div_safe`), and `main.c`. `div_safe` should return 0 and print an error to `stderr` when dividing by zero (or use an `int` status code—your design, document it).
3. **Argc printer** — Print all command-line arguments, one per line, including `argv[0]`.
4. **Warning hunt** — Write a file that triggers at least two different warnings under `-Wall -Wextra`, then fix them. Save both versions in comments at the bottom of the file.
5. **Profile lab upgrade** — Personal info printer that reads age with `scanf` and name with a fixed string or `scanf("%49s", ...)`.
6. **Converter lab upgrade** — Add an option 0 to exit; use a loop if you already know `while/for`, otherwise document that the stretch is deferred.

7. **Explain a log** — Deliberately omit `util.o` from a multi-file link command and paste the linker error into a comment; explain it in one sentence.
8. **Style pass** — Reformat one lab with consistent indentation (4 spaces), braces on the same style throughout, and a one-line file header comment with your name and purpose.

Basic Input and Output

Introduction

Almost every useful program talks to the outside world. In C, console I/O goes through the **standard streams**:

Stream	Pointer	Typical use
Standard input	stdin	Keyboard (or redirected file)
Standard output	stdout	Normal program output
Standard error	stderr	Diagnostics and errors

The declarations live in `<stdio.h>`. This chapter covers formatted I/O (`printf` / `scanf`), character I/O, safer line input with `fgets`, buffer pitfalls, and two mini-projects: a menu-driven CLI and a simple invoice printer.

Compile examples with:

```
gcc -std=c17 -Wall -Wextra -Wpedantic -o prog prog.c
```

The `<stdio.h>` library

```
#include <stdio.h>
```

You get:

- Formatted I/O: `printf`, `fprintf`, `sprintf`, `snprintf`, `scanf`, `fscanf`, `sscanf`
- Character I/O: `getchar`, `putchar`, `fgetc`, `fputc`
- Line I/O: `fgets`, `fputs`
- File operations (later module): `fopen`, `fclose`, ...

Output with printf

Prototype (conceptual)

```
int printf(const char *format, ...);
```

Returns the number of characters written, or a negative value on error.

Minimal example

```
#include <stdio.h>

int main(void) {
    printf("Hello, World!\n");
    return 0;
}
```

Format specifier table

Specifier	Typical argument	Meaning
%d / %i	int	Signed decimal
%u	unsigned int	Unsigned decimal
%ld / %lu	long / unsigned long	Long variants
%lld / %llu	long long / unsigned long long	Long long variants
%f	double (for printf)	Fixed-point floating
%e / %E	double	Scientific notation
%g / %G	double	Compact float format
%c	int (promoted char)	Character
%s	char *	String (must be null-terminated)
%x / %X	unsigned	Hexadecimal
%o	unsigned	Octal
%p	void *	Pointer address
%zu	size_t	Size type (C99+)
%%	—	Literal %

Important: for printf, float is promoted to double, so %f is used for both. For scanf, %f is float * and %lf is double *.

Runnable demo — mixed types

```
#include <stdio.h>

int main(void) {
    int age = 25;
    unsigned int score = 9001u;
    double height = 1.75;
    char initial = 'J';
    char name[] = "Jordan";
    void *self = &age;

    printf("Name:    %s\n", name);
    printf("Initial: %c\n", initial);
    printf("Age:      %d\n", age);
    printf("Score:    %u\n", score);
    printf("Height:   %.2f m\n", height);
    printf("Hex age:  0x%x\n", (unsigned)age);
    printf("Pointer:  %p\n", self);
    printf("Percent:  100%% done\n");

    return 0;
}
```

```
gcc -std=c17 -Wall -Wextra -Wpedantic -o print_types print_types.c
./print_types
```

Width, precision, flags

```
#include <stdio.h>

int main(void) {
    int n = 42;
    double pi = 3.1415926535;
    char word[] = "HelloWorld";

    printf("|%10d|\n", n);      /* right-align in width 10 */
    printf("|%-10d|\n", n);    /* left-align */
    printf("|%010d|\n", n);    /* zero-pad */
    printf("|%.2f|\n", pi);     /* 2 digits after decimal */
    printf("|%8.2f|\n", pi);    /* width 8, 2 decimals */
    printf("|%.5s|\n", word);   /* at most 5 chars from string */

    return 0;
}
```

Pattern	Effect
%10d	Minimum width 10, right-aligned
%-10d	Left-aligned
%010d	Zero-filled to width 10
%.2f	Two digits after decimal point
%.5s	Print at most 5 characters of a string
%#x	Hex with 0x prefix (alternate form)

fprintf and streams

```
#include <stdio.h>

int main(void) {
    printf("normal message\n");
    fprintf(stdout, "also normal\n");
    fprintf(stderr, "something went wrong\n");
    return 1;
}
```

Redirect separately:

```
./prog >out.txt 2>err.txt
```

snprintf — safe string formatting

Prefer `snprintf` over `sprintf` to avoid buffer overruns:

```
#include <stdio.h>

int main(void) {
    char buf[32];
    int age = 30;
    int n = snprintf(buf, sizeof buf, "age=%d", age);

    if (n < 0) {
        fprintf(stderr, "encoding error\n");
        return 1;
    }
    if ((size_t)n >= sizeof buf) {
        fprintf(stderr, "output truncated\n");
    }

    printf("%s\n", buf);
    return 0;
}
```

```
}
```

Input with scanf

Prototype (conceptual)

```
int scanf(const char *format, ...);
```

Returns the number of input **items** assigned, or EOF on input failure before any assignment.

Address operator

scanf needs **pointers** to the objects it will write:

```
int number;
scanf("%d", &number); /* correct */
/* scanf("%d", number); // wrong: passes value, not address */
```

Arrays (and string buffers) already decay to pointers:

```
char name[50];
scanf("%49s", name); /* no & for the array */
```

Specifier table for scanf (common cases)

Specifier	Argument type	Notes
%d	int *	Signed int
%u	unsigned int *	Unsigned
%ld	long *	Long
%f	float *	Float
%lf	double *	Double — easy to forget
%c	char *	Does not skip whitespace unless you force it
%s	char *	Stops at whitespace; always set a max width
%x	unsigned int *	Hex input

Full demo — reading several fields

```
#include <stdio.h>

int main(void) {
    int age;
    double height;
    char initial;
    char name[50];

    printf("Enter age: ");
    if (scanf("%d", &age) != 1) {
        fprintf(stderr, "bad age\n");
        return 1;
    }

    printf("Enter height (meters): ");
    if (scanf("%lf", &height) != 1) {
        fprintf(stderr, "bad height\n");
        return 1;
    }

    printf("Enter middle initial: ");
    if (scanf(" %c", &initial) != 1) { /* leading space skips whitespace */
        fprintf(stderr, "bad initial\n");
        return 1;
    }

    printf("Enter first name (no spaces): ");
    if (scanf("%49s", name) != 1) {
        fprintf(stderr, "bad name\n");
        return 1;
    }

    printf("\n--- Profile ---\n");
    printf("%s %c. age=%d height=%.2f\n", name, initial, age, height);
    return 0;
}
```

Buffer caveats with scanf

Leftover newlines

`scanf("%d", &n)` reads the number but leaves the terminating newline in the input buffer. A following `%c` can read that newline instead of the character you meant.

```
#include <stdio.h>

int main(void) {
    int n;
    char ch;

    printf("Enter an integer: ");
    if (scanf("%d", &n) != 1) {
        return 1;
    }

    printf("Enter a character: ");
    /* Without the leading space, ch often becomes '\n' */
    if (scanf(" %c", &ch) != 1) {
        return 1;
    }

    printf("n=%d ch='%c' (code %d)\n", n, ch, ch);
    return 0;
}
```

Rule of thumb: use a leading space in the format for `%c` when you want to skip leftover whitespace: `scanf(" %c", &ch)`.

%s stops at whitespace

`scanf("%s", buf)` cannot read "Ada Lovelace" as one field—it stops at the space. Use `fgets` for full lines (below).

Always bound string width

```
char name[20];
scanf("%19s", name); /* leave room for '\0' */
```

Unbounded `%s` is a classic buffer overflow.

Failed conversions leave bad data

If the user types `abc` for `%d`, `scanf` returns 0 and does not update the `int` reliably for your purposes—check the return value before using the variable.

Clearing the rest of a line

After a failed or partial read, discard through newline:

```
#include <stdio.h>

static void clear_line(void) {
    int c;
    while ((c = getchar()) != '\n' && c != EOF) {
        /* discard */
    }
}

int main(void) {
    int n;

    printf("Enter integer: ");
    while (scanf("%d", &n) != 1) {
        fprintf(stderr, "Not an integer. Try again: ");
        clear_line();
    }
    clear_line(); /* eat trailing newline after a good read if needed later */

    printf("Got %d\n", n);
    return 0;
}
```

Character I/O: getchar and putchar

```
int getchar(void); /* returns unsigned char cast to int, or EOF */
int putchar(int c);
```

Echo a character

```
#include <stdio.h>

int main(void) {
    int ch;

    printf("Press a key then Enter: ");
    ch = getchar();
    if (ch == EOF) {
        return 1;
    }

    printf("You typed: ");
    putchar(ch);
    putchar('\n');
    return 0;
}
```

Copy stdin to stdout until EOF

```
#include <stdio.h>

int main(void) {
    int ch;

    while ((ch = getchar()) != EOF) {
        putchar(ch);
    }
    return 0;
}
```

```
gcc -std=c17 -Wall -Wextra -Wpedantic -o catish catish.c
```

```
./catish < some_file.txt
```

```
# interactive: type text, end with Ctrl+D (Unix) or Ctrl+Z then Enter (Windows consoles)
```

Count characters and lines

```
#include <stdio.h>

int main(void) {
    int ch;
    unsigned long chars = 0;
```

```

    unsigned long lines = 0;

    while ((ch = getchar()) != EOF) {
        chars++;
        if (ch == '\n') {
            lines++;
        }
    }

    printf("chars=%lu lines=%lu\n", chars, lines);
    return 0;
}

```

Use `int` for the variable that holds `getchar()` so EOF (−1) is distinguishable from a byte value 0xFF.

Reading lines safely with `fgets`

Prototype

```
char *fgets(char *s, int size, FILE *stream);
```

Reads at most `size - 1` characters into `s`, stops early on newline, always null-terminates on success. Returns `s` on success, NULL on EOF/error.

Basic pattern

```

#include <stdio.h>
#include <string.h>

int main(void) {
    char line[128];

    printf("Enter a line: ");
    if (fgets(line, sizeof line, stdin) == NULL) {
        fprintf(stderr, "no input\n");
        return 1;
    }

    /* strip trailing newline if present */
    line[strcspn(line, "\n")] = '\0';
}

```

```

    printf("You said: \"%s\\n\"", line);
    return 0;
}

gcc -std=c17 -Wall -Wextra -Wpedantic -o line_echo line_echo.c
./line_echo

```

Why prefer fgets for text lines?

Approach	Risk
gets	Removed from C11 — never use; unbounded
scanf("%s")	Stops at space; easy to overflow without width
scanf("%[^\n]")	Awkward; still easy to get wrong
fgets	Bounded; keeps spaces; predictable

Mixing fgets and scanf

Leftover newlines still bite. Patterns that work:

Option A — prefer fgets for everything, then parse with sscanf:

```

#include <stdio.h>

int main(void) {
    char line[64];
    int age;

    printf("Age: ");
    if (fgets(line, sizeof line, stdin) == NULL) {
        return 1;
    }
    if (sscanf(line, "%d", &age) != 1) {
        fprintf(stderr, "invalid age\\n");
        return 1;
    }

    printf("age=%d\\n", age);
    return 0;
}

```

Option B — after scanf, drain the line before the next fgets:

```

#include <stdio.h>
#include <string.h>

```

```

static void clear_line(void) {
    int c;
    while ((c = getchar()) != '\n' && c != EOF) { }
}

int main(void) {
    int id;
    char name[64];

    printf("id: ");
    if (scanf("%d", &id) != 1) {
        return 1;
    }
    clear_line();

    printf("full name: ");
    if (fgets(name, sizeof name, stdin) == NULL) {
        return 1;
    }
    name[strcspn(name, "\n")] = '\0';

    printf("id=%d name=\"%s\"\n", id, name);
    return 0;
}

```

Detecting truncated lines

If `fgets` fills the buffer without seeing `\n`, the line was longer than the buffer:

```

#include <stdio.h>
#include <string.h>

int main(void) {
    char line[16];

    printf("Short buffer demo (type a long line):\n");
    if (fgets(line, sizeof line, stdin) == NULL) {
        return 1;
    }

    if (strchr(line, '\n') == NULL) {
        fprintf(stderr, "warning: line truncated\n");
        /* optional: drain the rest of the line */
        int c;
        while ((c = getchar()) != '\n' && c != EOF) { }
    }
}

```

```

    } else {
        line[strcspn(line, "\n")] = '\0';
    }

    printf("stored: \"%s\"\n", line);
    return 0;
}

```

sscanf — parse from a string

```

#include <stdio.h>

int main(void) {
    const char *input = "Jordan 25 1.75";
    char name[32];
    int age;
    double height;

    if (sscanf(input, "%31s %d %lf", name, &age, &height) != 3) {
        fprintf(stderr, "parse failed\n");
        return 1;
    }

    printf("%s is %d years old and %.2f m\n", name, age, height);
    return 0;
}

```

Useful when you already have a full line from `fgets` and want conversion with error checking.

Checking return values

printf

```

int n = printf("Hello\n");
if (n < 0) {
    /* rare on console; more relevant with broken streams */
}

```

scanf

```
int a, b;
int got = scanf("%d %d", &a, &b);
if (got != 2) {
    fprintf(stderr, "expected two integers, got %d\n", got);
    return 1;
}
```

Treat return-value checks as part of correct I/O, not optional polish.

Mini-project 1 — Menu-driven CLI

A small loop that presents options until the user quits. Uses `fgets` + `sscanf` to avoid buffer traps.

```
#include <stdio.h>
#include <string.h>

static int read_int(const char *prompt, int *out) {
    char line[64];

    printf("%s", prompt);
    if (fgets(line, sizeof line, stdin) == NULL) {
        return 0;
    }
    return sscanf(line, "%d", out) == 1;
}

int main(void) {
    char line[64];
    int choice;
    int running = 1;

    while (running) {
        printf("\n==== Mini CLI ==== \n");
        printf("1) Echo a line\n");
        printf("2) Add two integers\n");
        printf("3) Show ASCII code of a character\n");
        printf("0) Quit\n");
        printf("Select: ");

        if (fgets(line, sizeof line, stdin) == NULL) {
            break;
        }
    }
}
```

```

}
if (sscanf(line, "%d", &choice) != 1) {
    printf("Please enter a number from the menu.\n");
    continue;
}

switch (choice) {
case 0:
    running = 0;
    printf("Bye.\n");
    break;
case 1: {
    char text[128];
    printf("Text: ");
    if (fgets(text, sizeof text, stdin) == NULL) {
        running = 0;
        break;
    }
    text[strcspn(text, "\n")] = '\0';
    printf("Echo: %s\n", text);
    break;
}
case 2: {
    int a, b;
    if (!read_int("a = ", &a) || !read_int("b = ", &b)) {
        printf("Invalid integer input.\n");
        break;
    }
    printf("%d + %d = %d\n", a, b, a + b);
    break;
}
case 3: {
    char raw[16];
    printf("Character: ");
    if (fgets(raw, sizeof raw, stdin) == NULL) {
        running = 0;
        break;
    }
    if (raw[0] == '\0' || raw[0] == '\n') {
        printf("No character entered.\n");
        break;
    }
    printf("'%'c' has code %d\n", raw[0], (unsigned char)raw[0]);
    break;
}
default:
    printf("Unknown option.\n");
}

```



```

        break;
    }
}

return 0;
}

gcc -std=c17 -Wall -Wextra -Wpedantic -o menu_cli menu_cli.c
./menu_cli

```

Mini-project 2 — Simple invoice printer

Reads a few line items and prints a formatted invoice on `stdout`.

```

#include <stdio.h>
#include <string.h>

#define MAX_ITEMS 5
#define NAME_LEN 48

int main(void) {
    char customer[64];
    char names[MAX_ITEMS][NAME_LEN];
    int quantities[MAX_ITEMS];
    double prices[MAX_ITEMS];
    int count = 0;
    double subtotal = 0.0;
    const double tax_rate = 0.08;
    char line[128];

    printf("Customer name: ");
    if (fgets(customer, sizeof customer, stdin) == NULL) {
        return 1;
    }
    customer[strcspn(customer, "\n")] = '\0';

    printf("How many line items (1-%d)? ", MAX_ITEMS);
    if (fgets(line, sizeof line, stdin) == NULL) {
        return 1;
    }
    if (sscanf(line, "%d", &count) != 1 || count < 1 || count > MAX_ITEMS) {
        fprintf(stderr, "Invalid item count\n");
        return 1;
    }
}

```

```

}

for (int i = 0; i < count; i++) {
    printf("\nItem %d name: ", i + 1);
    if (fgets(names[i], NAME_LEN, stdin) == NULL) {
        return 1;
    }
    names[i][strcspn(names[i], "\n")] = '\0';

    printf("Quantity: ");
    if (fgets(line, sizeof line, stdin) == NULL) {
        return 1;
    }
    if (sscanf(line, "%d", &quantities[i]) != 1 || quantities[i] < 0) {
        fprintf(stderr, "Invalid quantity\n");
        return 1;
    }

    printf("Unit price: ");
    if (fgets(line, sizeof line, stdin) == NULL) {
        return 1;
    }
    if (sscanf(line, "%lf", &prices[i]) != 1 || prices[i] < 0.0) {
        fprintf(stderr, "Invalid price\n");
        return 1;
    }
}

printf("\n===== \n");
printf("          INVOICE\n");
printf("===== \n");
printf("Customer: %s\n\n", customer);
printf("%-20s %5s %10s %10s\n", "Item", "Qty", "Unit", "Line");
printf("----- \n");

for (int i = 0; i < count; i++) {
    double line_total = quantities[i] * prices[i];
    subtotal += line_total;
    printf("%-20.20s %5d %10.2f %10.2f\n",
        names[i], quantities[i], prices[i], line_total);
}

double tax = subtotal * tax_rate;
double total = subtotal + tax;

printf("----- \n");
printf("%-28s %10.2f\n", "Subtotal", subtotal);

```

```

printf("%-28s %10.2f\n", "Tax (8%)", tax);
printf("%-28s %10.2f\n", "TOTAL", total);
printf("=====\n");

return 0;
}

```

```

gcc -std=c17 -Wall -Wextra -Wpedantic -o invoice invoice.c
./invoice

```

This uses a C99-style `for (int i = 0; ...)` loop (fine under `-std=c17`). Arrays show up formally in a later module; here they are only storage for the demo.

Practical calculator (formatted I/O review)

```

#include <stdio.h>

int main(void) {
    double num1, num2, result;
    char op;
    char line[128];

    printf("Enter expression (e.g. 3.5 * 2): ");
    if (fgets(line, sizeof line, stdin) == NULL) {
        return 1;
    }
    if (sscanf(line, "%lf %c %lf", &num1, &op, &num2) != 3) {
        fprintf(stderr, "Could not parse expression\n");
        return 1;
    }

    switch (op) {
    case '+': result = num1 + num2; break;
    case '-': result = num1 - num2; break;
    case '*': result = num1 * num2; break;
    case '/':
        if (num2 == 0.0) {
            fprintf(stderr, "Division by zero\n");
            return 1;
        }
        result = num1 / num2;
        break;
    default:

```

```

        fprintf(stderr, "Invalid operator '%c'\n", op);
        return 1;
    }

    printf("%.10g %c %.10g = %.10g\n", num1, op, num2, result);
    return 0;
}

```

Best practices (quick reference)

1. Check return values of `scanf`, `fgets`, and parsing helpers.
 2. Bound every string read (`%49s`, `fgets(..., sizeof buf, ...)`).
 3. Prefer `fgets` + `sscanf` for interactive line-oriented tools.
 4. Never use `gets`.
 5. Send errors to `stderr` with `fprintf(stderr, ...)`.
 6. Mind whitespace when mixing `%d` and `%c`, or `scanf` and `fgets`.
 7. Use `snprintf`, not unbounded `sprintf`, when building strings.
 8. Compile with `-std=c17 -Wall -Wextra -Wpedantic`.
-

Summary

- `printf` formats output; width and precision control alignment and detail.
- `scanf` reads tokens; always pass addresses (except arrays), check the return count, and bound `%s`.
- Character I/O with `getchar` / `putchar` is simple and stream-friendly; store results in `int`.
- `fgets` is the default safe way to read a line; strip `\n` with `strcspn` or equivalent.
- Menu CLIs and small printers are good practice for real interactive programs.

Next module: data types and variables—so the values you print and scan have precise sizes, ranges, and conversion rules.

Exercises

Use `-std=c17 -Wall -Wextra -Wpedantic` for all of these.

1. **Specifier drill** — Print the same integer as decimal, octal, and hex on one line.

2. **Column report** — Print three rows of “name / quantity / price” using width specifiers so columns align.
3. **Scanf trap lab** — Demonstrate the leftover-newline bug with `%d` followed by `%c` *without* a leading space, then show the fixed version.
4. **Line reverser (partial)** — Read a line with `fgets` and print its length (excluding newline). Full reverse can wait until arrays/strings if needed; do it now if you already can.
5. **Retry loop** — Keep asking for an integer until `sscanf` after `fgets` succeeds.
6. **Menu extension** — Add a fourth menu option that converts Celsius to Fahrenheit.
7. **Invoice tax flag** — Modify the invoice program to ask whether tax applies (y/n) before printing totals.
8. **stderr routing** — Write a program that prints OK to stdout and BAD to stderr; run with redirects and show both files.
9. **snprintf card** — Build a string “ID=###;NAME=...” into a 64-byte buffer and report truncation if it occurs.
10. **Char histogram (lite)** — Read characters until EOF and count how many were digits ('0'–'9'). Print the count.

Module 1: Foundations and Environment Setup

Overview

This module is the **on-ramp** for the C book: what C is, how to install and drive a toolchain, the edit-compile-run loop, and console I/O. Finish it before jumping to data types and control flow unless you already ship C daily.

Default compile line (memorize it):

```
gcc -std=c17 -Wall -Wextra -Wpedantic -o prog prog.c
```

Clang is fine: replace `gcc` with `clang`.

Learning goals

By the end of this module you should be able to:

- Explain where C sits (systems language, standards C17/C23) and when it is the right tool
- Install **GCC** and/or **Clang**, a debugger (**GDB/LLDB**), and optional **CMake** / **Make**
- Write, compile, and run small programs; read common compiler diagnostics
- Split a program across `.c` / `.h` files and link object files
- Use `printf` / `scanf` / `fgets` carefully (check return values, avoid unbounded `%s`)
- Prefer `-g -O0` when debugging and treat warnings as signal, not noise

Chapters (map)

#	Chapter	What you practice
1	Introduction to C	History, standards, “Hello”, mini converters, mental model of compile/link
2	Development Environment	Install compilers; multi-file recipes ; Make ; mini CMake ; GCC vs Clang; GDB one-pager; sanitizers

#	Chapter	What you practice
3	First C Program	Program structure, flags, multi-file lab, personal-info / converter labs
4	Basic Input/Output	<code>printf/scanf</code> , format specs, <code>fgets</code> , mixing line and token input safely

Suggested path

1. Skim **Intro** so the compile pipeline and `main` form are clear.
2. Do **Environment** *for real* — multi-file build + one debugger session, not only install notes.
3. Type every program in **First Program**.
4. Complete **Basic I/O** exercises before Module 2; I/O bugs dominate early frustration.

Tooling cheat sheet

```
# single file
gcc -std=c17 -Wall -Wextra -Wpedantic -o hello hello.c && ./hello

# multi-file
gcc -std=c17 -Wall -Wextra -Wpedantic -c main.c util.c
gcc -o app main.o util.o

# debug
gcc -std=c17 -Wall -Wextra -g -O0 -o app main.c util.c
gdb ./app      # or: lldb ./app

# memory / UB (Linux/macOS)
gcc -std=c17 -Wall -Wextra -g -O1 -fsanitize=address,undefined -o app app.c
```

Prerequisites

- Basic computer literacy and a terminal
- Ability to install packages (apt, brew, MSYS2, or Xcode CLT)
- No prior C required for Path A (see book landing page)

Self-check before Module 2

- ☐ `hello.c` builds with zero warnings under `-Wall -Wextra`
- ☐ Multi-file program links (`main.c + util.c`)
- ☐ One GDB/LLDB session: breakpoint, `print`, continue
- ☐ Program reads an integer and a line of text without leftover-newline confusion
- ☐ You know how to show the compiler's version and the executable's exit status (`echo $?`)

Estimated time

- Reading: 4–6 hours
- Typing examples + exercises: 6–10 hours
- Optional: CMake mini-project and sanitizer lab from chapter 2

Next module

Module 2: Data Types and Variables — types, sizes, constants, conversions, operators.

Resources

- K&R *The C Programming Language* (classic reference, older dialect)
- GCC / Clang manuals for flag details
- Book appendices for reference tables and solutions pointers

Data Types

Fundamental Data Types

Introduction

Data types define what values you can store, how much memory they use, which operations apply, and how the compiler interprets bits. In C, sizes of some types are **implementation-defined** (within standard minimums), so portable code measures with `sizeof` and uses limits headers or fixed-width types when exact width matters.

This chapter covers integer and floating types, `char` and `bool`, full `sizeof / limits.h / float.h` demos, signed overflow undefined behavior and safer patterns, and a strong emphasis on `stdint.h` fixed-width types—with full programs you can run on Linux.

Compile convention:

```
gcc -std=c17 -Wall -Wextra -o program program.c
./program
```

Basic Data Types in C

1. **Integer types** — whole numbers
2. **Floating-point types** — real numbers
3. **Character types** — `char` (integer family)
4. **Boolean** — `_Bool` / `bool` (C99+)

Integer Types

Typical Sizes (LP64 Linux, e.g. x86_64)

These are **common** on 64-bit Linux with GCC—not guarantees for every platform.

Type	Size (typical)	Range (typical, signed)
<code>char</code>	1 byte	−128..127 or 0..255 if unsigned <code>char</code>

Type	Size (typical)	Range (typical, signed)
short	2 bytes	−32 768..32 767
int	4 bytes	−2 147 483 648..2 147 483 647
long	8 bytes (LP64)	full 64-bit signed range
long long	8 bytes	full 64-bit signed range

Unsigned Variants

Type	Typical max
unsigned char	255
unsigned short	65 535
unsigned int	4 294 967 295
unsigned long	2 −1 on LP64
unsigned long long	2 −1

Unsigned arithmetic wraps modulo 2 (well-defined). Signed overflow is **not** well-defined (see below).

Worked Demo: sizeof of Major Types

Run this on your machine and keep the output—it is your ABI cheatsheet.

```
/* sizeof_tour.c */
#include <stdio.h>
#include <stdbool.h>
#include <stddef.h>
#include <stdint.h>

int main(void) {
    printf("=== sizeof tour (this machine) ===\n");
    printf("char:           %zu\n", sizeof(char));
    printf("signed char:    %zu\n", sizeof(signed char));
    printf("unsigned char: %zu\n", sizeof(unsigned char));
    printf("short:          %zu\n", sizeof(short));
    printf("unsigned short: %zu\n", sizeof(unsigned short));
    printf("int:            %zu\n", sizeof(int));
    printf("unsigned int:   %zu\n", sizeof(unsigned int));
    printf("long:           %zu\n", sizeof(long));
    printf("unsigned long:  %zu\n", sizeof(unsigned long));
    printf("long long:      %zu\n", sizeof(long long));
    printf("unsigned long long: %zu\n", sizeof(unsigned long long));
}
```

```

printf("float:           %zu\n", sizeof(float));
printf("double:          %zu\n", sizeof(double));
printf("long double:     %zu\n", sizeof(long double));
printf("bool / _Bool:    %zu\n", sizeof(bool));
printf("void *:           %zu\n", sizeof(void *));
printf("size_t:           %zu\n", sizeof(size_t));
printf("ptrdiff_t:        %zu\n", sizeof(ptrdiff_t));
printf("int8_t:            %zu\n", sizeof(int8_t));
printf("int16_t:           %zu\n", sizeof(int16_t));
printf("int32_t:           %zu\n", sizeof(int32_t));
printf("int64_t:           %zu\n", sizeof(int64_t));
printf("uint64_t:          %zu\n", sizeof(uint64_t));
printf("intptr_t:          %zu\n", sizeof(intptr_t));
printf("uintptr_t:         %zu\n", sizeof(uintptr_t));
printf("intmax_t:          %zu\n", sizeof(intmax_t));

int x = 0;
printf("sizeof x:        %zu (parens optional for expressions)\n", sizeof x);

return 0;
}

```

```

gcc -std=c17 -Wall -Wextra -o sizeof_tour sizeof_tour.c
./sizeof_tour

```

Sample-shaped output on a typical x86_64 Linux GCC (yours may match):

```

char: 1, short: 2, int: 4, long: 8, long long: 8
float: 4, double: 8, long double: 16
void *: 8

```

On 32-bit ILP32 Linux, long and pointers are often 4 bytes—another reason not to assume.

Limits from limits.h

```

/* limits_demo.c */
#include <stdio.h>
#include <limits.h>

int main(void) {
    printf("=== limits.h (signed) ===\n");
    printf("CHAR_BIT    = %d bits in a byte\n", CHAR_BIT);
    printf("CHAR_MIN    = %d\n", CHAR_MIN);
}

```

```

printf("CHAR_MAX    = %d\n", CHAR_MAX);
printf("SCHAR_MIN   = %d\n", SCHAR_MIN);
printf("SCHAR_MAX    = %d\n", SCHAR_MAX);
printf("SHRT_MIN     = %d\n", SHRT_MIN);
printf("SHRT_MAX     = %d\n", SHRT_MAX);
printf("INT_MIN      = %d\n", INT_MIN);
printf("INT_MAX      = %d\n", INT_MAX);
printf("LONG_MIN     = %ld\n", LONG_MIN);
printf("LONG_MAX     = %ld\n", LONG_MAX);
printf("LLONG_MIN    = %lld\n", LLONG_MIN);
printf("LLONG_MAX    = %lld\n", LLONG_MAX);

printf("\n=== limits.h (unsigned maxima) ===\n");
printf("UCHAR_MAX    = %u\n", UCHAR_MAX);
printf("USHRT_MAX     = %u\n", USHRT_MAX);
printf("UINT_MAX      = %u\n", UINT_MAX);
printf("ULONG_MAX     = %lu\n", ULONG_MAX);
printf("ULLONG_MAX    = %llu\n", ULLONG_MAX);

return 0;
}

```

```

gcc -std=c17 -Wall -Wextra -o limits_demo limits_demo.c
./limits_demo

```

Use these macros for **portable** range checks instead of hard-coded 2147483647.

Floating-Point Types and float.h

Type	Typical size	Approx. decimal digits
float	4	~6–7 (FLT_DIG)
double	8	~15–16 (DBL_DIG)
long double	8–16	extended (LDBL_DIG)

```

/* float_limits_demo.c */
#include <stdio.h>
#include <float.h>

int main(void) {
    printf("=== sizes ===\n");
    printf("float:      %zu\n", sizeof(float));
}

```

```

printf("double:      %zu\n", sizeof(double));
printf("long double: %zu\n", sizeof(long double));

printf("\n=== decimal digits of precision ===\n");
printf("FLT_DIG  = %d\n", FLT_DIG);
printf("DBL_DIG  = %d\n", DBL_DIG);
printf("LDBL_DIG = %d\n", LDBL_DIG);

printf("\n=== min positive normalized / max finite ===\n");
printf("FLT_MIN  = %e  FLT_MAX  = %e\n", FLT_MIN, FLT_MAX);
printf("DBL_MIN  = %e  DBL_MAX  = %e\n", DBL_MIN, DBL_MAX);
printf("LDBL_MIN = %Le LDBL_MAX = %Le\n", LDBL_MIN, LDBL_MAX);

printf("\n=== epsilon (1 + eps != 1) ===\n");
printf("FLT_EPSILON = %e\n", FLT_EPSILON);
printf("DBL_EPSILON = %e\n", DBL_EPSILON);
printf("LDBL_EPSILON = %Le\n", LDBL_EPSILON);

printf("\n=== radix and mantissa bits ===\n");
printf("FLT_RADIX = %d\n", FLT_RADIX);
printf("FLT_MANT_DIG = %d  DBL_MANT_DIG = %d  LDBL_MANT_DIG = %d\n",
      FLT_MANT_DIG, DBL_MANT_DIG, LDBL_MANT_DIG);

return 0;
}

```

```

gcc -std=c17 -Wall -Wextra -o float_limits_demo float_limits_demo.c
./float_limits_demo

```

Classic Precision Surprise

```

#include <stdio.h>

int main(void) {
    double a = 0.1 + 0.2;
    printf("0.1 + 0.2 = %.20f\n", a);
    printf("equals 0.3? %s\n", (a == 0.3) ? "yes" : "no");
    return 0;
}

```

```

gcc -std=c17 -Wall -Wextra -o fp_surprise fp_surprise.c
./fp_surprise

```

Character Types

`char` is an integer type holding a character code (often ASCII/UTF-8 unit). Whether plain `char` is signed is **implementation-defined**.

```
#include <stdio.h>

int main(void) {
    char letter = 'A';
    char from_code = 65;

    printf("Letter: %c\n", letter);
    printf("As int: %d\n", letter);
    printf("from_code: %c\n", from_code);

    signed char sc = -1;
    unsigned char uc = 255;
    printf("signed char -1 -> %d\n", sc);
    printf("unsigned char 255 -> %u\n", uc);
    return 0;
}

gcc -std=c17 -Wall -Wextra -o char_demo char_demo.c
./char_demo
```

Prefer `unsigned char` for raw byte buffers; prefer `signed char` when you need negative small integers explicitly.

Boolean Type (C99+)

```
#include <stdio.h>
#include <stdbool.h>

int main(void) {
    _Bool flag = 1;
    bool ok = true;
    bool bad = false;

    printf("flag=%d ok=%d bad=%d\n", flag, ok, bad);
    if (ok) {
        printf("Task complete path\n");
    }
    return 0;
}
```

```
}
```

Any scalar converts to `_Bool`: zero $\rightarrow$ 0, nonzero $\rightarrow$ 1.

Fixed-Width Types (`stdint.h`) — Prefer These When Width Matters

`int` and `long` change meaning across ABIs. For wire formats, file headers, embedded registers, cryptography, and cross-platform protocols, use **exact-width** types when the implementation provides them (all modern hosted Linux toolchains do for 8/16/32/64).

Exact-Width

Type	Width	Signed range (two's complement)
<code>int8_t</code> / <code>uint8_t</code>	8	$-128..127$ / $0..255$
<code>int16_t</code> / <code>uint16_t</code>	16	$\pm 2^1$ / $0..2^1 - 1$
<code>int32_t</code> / <code>uint32_t</code>	32	$\pm 2^{31}$ / $0..2^{32} - 1$
<code>int64_t</code> / <code>uint64_t</code>	64	$\pm 2^3$ / $0..2 - 1$

Also: `int_leastN_t`, `int_fastN_t`, `intmax_t`, `intptr_t`, `uintptr_t`.

Printing Portably: `inttypes.h`

```
/* stdint_demo.c */
#include <stdio.h>
#include <stdint.h>
#include <inttypes.h>
#include <stddef.h>

int main(void) {
    int8_t    a = -5;
    uint8_t   b = 200;
    int16_t   c = -1000;
    uint16_t  d = 50000;
    int32_t   e = 123456789;
    uint32_t  f = 4000000000u;
    int64_t   g = -123456789012345LL;
    uint64_t  h = 123456789012345ULL;

    printf("int8:    %" PRIi8 "\n", a);
    printf("uint8:   %" PRIu8 "\n", b);
```



```

printf("int16:  %" PRIu16 "\n", c);
printf("uint16: %" PRIu16 "\n", d);
printf("int32:  %" PRIu32 "\n", e);
printf("uint32: %" PRIu32 "\n", f);
printf("int64:  %" PRIu64 "\n", g);
printf("uint64: %" PRIu64 "\n", h);

/* Hex */
printf("uint32 hex: 0x%08" PRIx32 "\n", f);

/* Pointer-sized integers */
void *p = &e;
uintptr_t raw = (uintptr_t)p;
printf("pointer as uintptr_t: 0x%" PRIxPTR "\n", raw);

/* Limits from stdint.h */
printf("INT32_MIN=%" PRIu32 " INT32_MAX=%" PRIu32 "\n",
      INT32_MIN, INT32_MAX);
printf("UINT64_MAX=%" PRIu64 "\n", UINT64_MAX);

return 0;
}

```

```

gcc -std=c17 -Wall -Wextra -o stdint_demo stdint_demo.c
./stdint_demo

```

When to Use What

Need	Choose
Exact layout / protocol field	<code>intN_t</code> / <code>uintN_t</code>
“At least N bits”, speed OK	<code>int_fastN_t</code>
“At least N bits”, small size	<code>int_leastN_t</code>
Object sizes, array indices	<code>size_t</code>
Pointer difference	<code>ptrdiff_t</code>
Generic local counter on host	<code>int</code> is fine
File offsets (POSIX)	often <code>off_t</code> (with feature macros)

Example: Portable Binary Header

```

/* packet_header.c */
#include <stdio.h>
#include <stdint.h>
#include <inttypes.h>

```

```

#include <string.h>

struct PacketHeader {
    uint32_t magic;      /* 0x4B313947 'K19G' style demo */
    uint16_t version;
    uint16_t length;
    uint32_t crc32;
};

int main(void) {
    struct PacketHeader h = {
        .magic = 0x4B313947u,
        .version = 1,
        .length = 42,
        .crc32 = 0xDEADBEEFu
    };

    printf("sizeof header = %zu (expect 12 if no padding)\n", sizeof h);
    printf("magic=%08" PRIx32 " ver=%" PRIu16 " len=%" PRIu16 " crc=%08" PRIx32 "\n",
        h.magic, h.version, h.length, h.crc32);

    /* Serialize explicitly (endianness still your responsibility) */
    uint8_t buf[12];
    memcpy(buf + 0, &h.magic, 4);
    memcpy(buf + 4, &h.version, 2);
    memcpy(buf + 6, &h.length, 2);
    memcpy(buf + 8, &h.crc32, 4);

    printf("bytes:");
    for (size_t i = 0; i < sizeof buf; i++) {
        printf(" %02X", buf[i]);
    }
    printf("\n");
    return 0;
}

gcc -std=c17 -Wall -Wextra -o packet_header packet_header.c
./packet_header

```

Note: `memcpy` of multi-byte integers uses **host endianness**. Network protocols often need explicit big-endian packing.

Signed Overflow Is Undefined Behavior

For **signed** integer types, overflow (and some underflows) are **undefined behavior (UB)** in C. The compiler may assume signed overflow never happens and optimize accordingly—including deleting “impossible” checks.

```
/* overflow_ub_note.c - educational; do not rely on wraparound of signed int */
#include <stdio.h>
#include <limits.h>
#include <stdint.h>
#include <inttypes.h>

int main(void) {
    int x = INT_MAX;
    printf("INT_MAX = %d\n", x);

    /* The expression (x + 1) is UB if it overflows.
     * In practice you might see wrap, a trap, or surprising optimization.
     * Do NOT write production code that expects wrap on signed overflow.
     */
    printf("Naive signed x+1 is UB if evaluated at INT_MAX.\n");

    /* Well-defined wrap: use unsigned */
    unsigned int ux = (unsigned int)INT_MAX;
    printf("unsigned wrap: %u + 1 = %u\n", ux, ux + 1u);

    /* Or use wider type for the math */
    int64_t wide = (int64_t)x + 1;
    printf("widened: %" PRIu64 "\n", wide);

    return 0;
}
```

```
gcc -std=c17 -Wall -Wextra -o overflow_ub_note overflow_ub_note.c
./overflow_ub_note
```

Safer patterns:

```
/* overflow_safe_patterns.c */
#include <stdio.h>
#include <limits.h>
#include <stdint.h>
#include <inttypes.h>
#include <stdbool.h>

/* Return false if a + b would overflow int */
```

```

static bool safe_add_int(int a, int b, int *out) {
    if (b > 0 && a > INT_MAX - b) {
        return false;
    }
    if (b < 0 && a < INT_MIN - b) {
        return false;
    }
    *out = a + b;
    return true;
}

/* Unsigned: wrap is defined; detect wrap if you need saturation/error */
static bool safe_add_u32(uint32_t a, uint32_t b, uint32_t *out) {
    if (a > UINT32_MAX - b) {
        return false; /* would wrap */
    }
    *out = a + b;
    return true;
}

static bool safe_mul_int(int a, int b, int *out) {
    if (a == 0 || b == 0) {
        *out = 0;
        return true;
    }
    /* Use int64_t intermediate when int is 32-bit - common on Linux */
    int64_t r = (int64_t)a * (int64_t)b;
    if (r > INT_MAX || r < INT_MIN) {
        return false;
    }
    *out = (int)r;
    return true;
}

int main(void) {
    int sum;
    if (!safe_add_int(INT_MAX, 1, &sum)) {
        printf("add INT_MAX+1: overflow detected (good)\n");
    }

    if (safe_add_int(100, 23, &sum)) {
        printf("100+23 = %d\n", sum);
    }

    uint32_t u;
    if (!safe_add_u32(UINT32_MAX, 1u, &u)) {
        printf("u32 MAX+1: overflow detected\n");
    }
}

```

```

}

int prod;
if (!safe_mul_int(100000, 100000, &prod)) {
    printf("100000*100000 does not fit in int\n");
}

/* Compiler helpers (GCC/Clang) - practical on Linux */
int r;
if (__builtin_add_overflow(INT_MAX, 1, &r)) {
    printf("__builtin_add_overflow caught INT_MAX+1\n");
}
if (!__builtin_add_overflow(20, 22, &r)) {
    printf("__builtin_add_overflow 20+22 = %d\n", r);
}

return 0;
}

gcc -std=c17 -Wall -Wextra -o overflow_safe_patterns overflow_safe_patterns.c
./overflow_safe_patterns

```

Safer Patterns Summary

Pattern	Notes
Check before op	if (a > INT_MAX - b) style for add
Wider intermediate	int64_t for 32-bit int multiply/add
Unsigned modular math	Defined wrap; still check if wrap is an error
__builtin*_overflow	GCC/Clang; clean and efficient
Sanitizers	gcc -fsanitize=undefined catches overflow at runtime in tests

```

gcc -std=c17 -Wall -Wextra -fsanitize=undefined -o ov_san ov.c
./ov_san

```

Type Specifiers and Qualifiers

Specifier / qualifier	Role
signed / unsigned	Signedness
short / long / long long	Rank / width hints
const	Not assignable through this lvalue
volatile	May change outside program's knowledge
restrict	Pointer aliasing promise (C99)

```
#include <stdio.h>

int main(void) {
    const int max_size = 100;
    long long big = 123456789012345LL;
    unsigned short port = 8080;

    printf("max_size=%d big=%lld port=%hu\n", max_size, big, port);
    return 0;
}
```

Type Promotion (Preview)

char and short promote to int (or unsigned int) in expressions. Mixed int/float/double follow usual arithmetic conversions. Details belong in the type-conversion chapter; for now, know that `sizeof(char + char)` is often `sizeof(int)`.

```
#include <stdio.h>

int main(void) {
    char a = 1, b = 2;
    printf("sizeof(char)=%zu sizeof(a+b)=%zu\n", sizeof(char), sizeof(a + b));
    return 0;
}
```

```
gcc -std=c17 -Wall -Wextra -o promote promote.c
./promote
```

Practical Combined Explorer

```
/* type_explorer.c */
#include <stdio.h>
#include <limits.h>
#include <float.h>
#include <stdint.h>
#include <inttypes.h>
#include <stddef.h>

int main(void) {
    printf("=== Memory Requirements ===\n");
    printf("char %zu | int %zu | long %zu | float %zu | double %zu | void* %zu\n",
        sizeof(char), sizeof(int), sizeof(long),
        sizeof(float), sizeof(double), sizeof(void *));

    int n = 1000;
    printf("\nArray of %d ints: %zu bytes\n", n, (size_t)n * sizeof(int));
    printf("Array of %d int64_t: %zu bytes\n", n, (size_t)n * sizeof(int64_t));

    printf("\n=== Critical limits ===\n");
    printf("INT_MAX=%d UINT_MAX=%u\n", INT_MAX, UINT_MAX);
    printf("LONG_MAX=%ld\n", LONG_MAX);
    printf("SIZE_MAX=%zu\n", (size_t)SIZE_MAX);
    printf("DBL_DIG=%d DBL_EPSILON=%e\n", DBL_DIG, DBL_EPSILON);

    printf("\n=== Fixed width sanity ===\n");
    printf("sizeof(int32_t)=%zu sizeof(int64_t)=%zu\n",
        sizeof(int32_t), sizeof(int64_t));

    return 0;
}

gcc -std=c17 -Wall -Wextra -o type_explorer type_explorer.c
./type_explorer
```

Best Practices

1. Measure with `sizeof`; don't hard-code type sizes.
2. Use `stdint.h` for on-disk, on-wire, and register-facing fields.

3. Print fixed-width values with `inttypes.h` macros (`PRId32`, ...).
 4. Treat **signed overflow as a bug**, not as wraparound.
 5. Prefer `size_t` for sizes and indices into objects.
 6. Compare floats with **tolerances**, not `==`, unless you need exact bit identity.
 7. Enable `-Wall -Wextra`; add `-fsanitize=undefined` in debug builds.
-

Exercises

Exercise 1 — ABI report

Extend `sizeof_tour.c` to also print `CHAR_BIT`, whether plain `char` is signed (`((char)-1 < 0)`), and `sizeof(struct { char c; int i; })` vs sum of members (padding).

```
gcc -std=c17 -Wall -Wextra -o abi_report abi_report.c
./abi_report
```

Exercise 2 — limits-driven clamp

Write `int clamp_to_int(long long x)` that returns `INT_MIN/INT_MAX` if out of range, else `(int)x`. Test with values from `limits_demo` output.

Exercise 3 — safe multiply

Implement `bool safe_mul_i64(int64_t a, int64_t b, int64_t *out)` without invoking UB. Compare against `__builtin_mul_overflow`.

Exercise 4 — protocol fields

Define a message with `uint8_t` `type`, `uint16_t` `id`, `uint32_t` `payload_len` using fixed-width types. Pack into a byte buffer in **little-endian** explicitly (shifts and masks), then unpack and verify round-trip.

Exercise 5 — float.h scavenger

Print `FLT_ROUNDS` and explain what each value means (look up in the standard or man pages). Measure how many times you can add `FLT_EPSILON` to `1.0f` before the sum changes (loop with a counter).

Exercise 6 — UB sanitizer lab

Intentionally evaluate `int x = INT_MAX; x = x + 1;` in a tiny program under:

```
gcc -std=c17 -Wall -Wextra -fsanitize=undefined -O0 -o ub ub.c
./ub
```

Then rewrite using a safe helper and confirm the sanitizer is silent.

Summary

1. **Integer and floating types** — families, typical Linux sizes
2. **sizeof tour** — measure your ABI
3. **limits.h / float.h** — portable ranges and FP characteristics
4. **char / bool** — integer nature and boolean conversion
5. **stdint.h + inttypes.h** — fixed-width types and portable I/O
6. **Signed overflow UB** — detection patterns and builtins
7. **Qualifiers** — `const`, `volatile`, `restrict` (preview)
8. **Exercises** — ABI report, clamp, safe mul, packing, sanitizers

Next: variables, constants, and storage duration—how these types appear in real programs.

Variables and Constants

Introduction

Variables and constants are fundamental building blocks in C programming. Variables are named storage locations in memory that can hold values which may change during program execution, while constants are fixed values that cannot be altered after they are defined.

Understanding how to properly declare, initialize, and use variables and constants is crucial for writing effective C programs. This chapter will cover all aspects of variables and constants, including their declaration, scope, lifetime, and best practices.

Variables

What is a Variable?

A variable is a named memory location that stores data which can be modified during program execution. Every variable in C has: - A name (identifier) - A type (determines what kind of data it can store) - A value (the actual data stored) - A memory address (location in memory)

Variable Declaration

Before using a variable, you must declare it by specifying its type and name:

```
int age;           // Declares an integer variable named 'age'
float height;      // Declares a float variable named 'height'
char initial;      // Declares a character variable named 'initial'
double salary;     // Declares a double variable named 'salary'
```

Variable Initialization

Variables can be initialized (given an initial value) at the time of declaration:

```
int age = 25;           // Initialize with value 25
float height = 5.9f;    // Initialize with value 5.9
char initial = 'J';     // Initialize with character 'J'
double salary = 75000.50; // Initialize with value 75000.50
```

You can also declare multiple variables of the same type in a single statement:

```
int x = 10, y = 20, z = 30;    // Multiple integers with initial values
float a, b, c;                // Multiple floats without initialization
```

Variable Assignment

After declaration, you can assign new values to variables using the assignment operator (=):

```
#include <stdio.h>

int main() {
    int number;           // Declaration
    number = 42;          // Assignment

    printf("Number: %d\n", number);

    number = 100;         // New assignment
    printf("Number: %d\n", number);

    return 0;
}
```

Variable Naming Rules

C has specific rules for naming variables:

1. ****Must begin with a letter (a-z, A-Z) or underscore (_)****
2. **Can contain letters, digits (0-9), and underscores**
3. **Cannot be C keywords (int, float, if, etc.)**
4. **Case-sensitive (age and Age are different variables)**
5. **No spaces or special characters allowed**

Valid Variable Names

```
int age;                // Valid
int _temp;              // Valid (starts with underscore)
int student1;           // Valid (contains digit)
int myVariable;         // Valid (camelCase)
int MAX_SIZE;           // Valid (uppercase convention for constants)
```

Invalid Variable Names

```
int 1age;           // Invalid (starts with digit)
int my-age;         // Invalid (contains hyphen)
int int;            // Invalid (keyword)
int my age;         // Invalid (contains space)
int $amount;        // Invalid (contains special character)
```

Variable Naming Conventions

While C allows various naming styles, following consistent conventions improves code readability:

```
// Snake case (recommended for C)
int student_age;
float account_balance;
char first_name[50];

// Camel case (also acceptable)
int studentAge;
float accountBalance;
char firstName[50];

// Upper case for constants (convention)
const int MAX_STUDENTS = 100;
const float PI = 3.14159f;
```

Constants

Constants are fixed values that cannot be changed during program execution. C provides several ways to define constants.

Literal Constants

Literal constants are values written directly in the code:

```
int main() {
    int age = 25;           // 25 is an integer literal
    float price = 19.99f;   // 19.99 is a floating-point literal
    char grade = 'A';       // 'A' is a character literal
    char name[] = "John";   // "John" is a string literal

    return 0;
}
```

#define Preprocessor Directive

The `#define` directive creates symbolic constants:

```
#include <stdio.h>

#define MAX_SIZE 100
#define PI 3.14159
#define COMPANY_NAME "TechCorp"
#define IS_ACTIVE 1

int main() {
    int array[MAX_SIZE];
    float area = PI * 10 * 10;

    printf("Company: %s\n", COMPANY_NAME);
    printf("Area of circle: %.2f\n", area);

    if (IS_ACTIVE) {
        printf("System is active\n");
    }

    return 0;
}
```

const Qualifier

The `const` keyword creates typed constants:

```
#include <stdio.h>

int main() {
    const int max_size = 100;
    const float pi = 3.14159f;
    const char company_name[] = "TechCorp";

    printf("Max size: %d\n", max_size);
    printf("PI: %.5f\n", pi);
    printf("Company: %s\n", company_name);

    // max_size = 200; // Error: cannot modify const variable

    return 0;
}
```

enum (Enumeration)

Enumerations create named integer constants:

```
#include <stdio.h>

enum Weekday {
    MONDAY = 1,
    TUESDAY,
    WEDNESDAY,
    THURSDAY,
    FRIDAY,
    SATURDAY,
    SUNDAY
};

enum Status {
    INACTIVE = 0,
    ACTIVE = 1,
    PENDING = 2
};

int main() {
    enum Weekday today = WEDNESDAY;
    enum Status user_status = ACTIVE;

    printf("Today is day %d of the week\n", today);
    printf("User status: %d\n", user_status);

    if (user_status == ACTIVE) {
        printf("User is active\n");
    }

    return 0;
}
```

Storage Classes

Storage classes define the scope, visibility, and lifetime of variables and functions.

auto

The `auto` storage class is the default for local variables:

```

#include <stdio.h>

int main() {
    auto int count = 0; // auto is optional
    int number = 10;    // implicitly auto

    printf("Count: %d\n", count);
    printf("Number: %d\n", number);

    return 0;
}

```

register

The **register** storage class suggests that the variable be stored in a CPU register for faster access:

```

#include <stdio.h>

int main() {
    register int counter = 0;
    register int i;

    for (i = 0; i < 1000000; i++) {
        counter += i;
    }

    printf("Counter: %d\n", counter);

    return 0;
}

```

static

The **static** storage class has different meanings depending on context:

Static Local Variables

```

#include <stdio.h>

void counter() {
    static int count = 0; // Retains value between function calls
    count++;
    printf("Count: %d\n", count);
}

```

```

int main() {
    counter(); // Output: Count: 1
    counter(); // Output: Count: 2
    counter(); // Output: Count: 3

    return 0;
}

```

Static Global Variables

```

#include <stdio.h>

static int internal_count = 0; // Only accessible within this file

void increment_count() {
    internal_count++;
}

int get_count() {
    return internal_count;
}

int main() {
    increment_count();
    increment_count();
    printf("Internal count: %d\n", get_count());

    return 0;
}

```

extern

The **extern** storage class declares a variable that is defined in another file:

```

// file1.c
int global_counter = 0; // Definition

// file2.c
#include <stdio.h>

extern int global_counter; // Declaration (defined elsewhere)

void print_counter() {
    printf("Global counter: %d\n", global_counter);
}

```



```
}
```

Scope and Lifetime

Scope

Scope determines where a variable can be accessed in the code:

Local Scope

```
#include <stdio.h>

int main() {
    int local_var = 10; // Local to main function

    if (local_var > 5) {
        int block_var = 20; // Local to this block
        printf("Local var: %d, Block var: %d\n", local_var, block_var);
    }

    // printf("Block var: %d\n", block_var); // Error: out of scope

    return 0;
}
```

Global Scope

```
#include <stdio.h>

int global_var = 100; // Global scope

void function1() {
    printf("Global var in function1: %d\n", global_var);
}

int main() {
    printf("Global var in main: %d\n", global_var);
    function1();

    return 0;
}
```

Lifetime

Lifetime determines how long a variable exists in memory:

```
#include <stdio.h>

int global_var = 100; // Exists for entire program duration

void function() {
    static int static_var = 0; // Initialized only once, exists for program duration
    int local_var = 10;        // Created and destroyed each function call

    static_var++;
    printf("Static var: %d, Local var: %d\n", static_var, local_var);
}

int main() {
    function(); // Static var: 1, Local var: 10
    function(); // Static var: 2, Local var: 10
    function(); // Static var: 3, Local var: 10

    return 0;
}
```

Type Qualifiers

const

The `const` qualifier makes a variable read-only:

```
#include <stdio.h>

int main() {
    const int max_items = 50;
    const float tax_rate = 0.08f;

    // max_items = 100; // Error: cannot modify const variable

    printf("Max items: %d\n", max_items);
    printf("Tax rate: %.2f\n", tax_rate);

    return 0;
}
```

volatile

The `volatile` qualifier tells the compiler that a variable's value may change unexpectedly:

```
#include <stdio.h>

int main() {
    volatile int sensor_value = 0; // May be changed by hardware

    // Compiler won't optimize away reads/writes to volatile variables
    while (sensor_value < 100) {
        // Read sensor value (might be updated by hardware)
        printf("Sensor value: %d\n", sensor_value);
    }

    return 0;
}
```

restrict (C99+)

The `restrict` qualifier is used with pointers to indicate that no other pointer will access the same memory location:

```
#include <stdio.h>

void add_arrays(int n, int *restrict a, int *restrict b, int *restrict c) {
    // Compiler can optimize knowing a, b, c don't overlap
    for (int i = 0; i < n; i++) {
        c[i] = a[i] + b[i];
    }
}

int main() {
    int array1[5] = {1, 2, 3, 4, 5};
    int array2[5] = {10, 20, 30, 40, 50};
    int result[5];

    add_arrays(5, array1, array2, result);

    for (int i = 0; i < 5; i++) {
        printf("%d ", result[i]);
    }
    printf("\n");

    return 0;
}
```

Practical Examples

Variable Declaration and Usage

```
#include <stdio.h>

int main() {
    // Different ways to declare and initialize variables
    int age = 25;
    float height = 5.9f;
    double weight = 150.5;
    char grade = 'A';
    char name[] = "John Doe";

    // Display initial values
    printf("=== Initial Values ===\n");
    printf("Age: %d\n", age);
    printf("Height: %.1f feet\n", height);
    printf("Weight: %.1f pounds\n", weight);
    printf("Grade: %c\n", grade);
    printf("Name: %s\n", name);

    // Modify variables
    age = 26;
    height = 6.0f;
    weight = 155.0;
    grade = 'B';

    // Display modified values
    printf("\n=== Modified Values ===\n");
    printf("Age: %d\n", age);
    printf("Height: %.1f feet\n", height);
    printf("Weight: %.1f pounds\n", weight);
    printf("Grade: %c\n", grade);

    return 0;
}
```

Constants Comparison

```
#include <stdio.h>

#define PI_DEFINE 3.14159
const float PI_CONST = 3.14159f;
```

```

enum Colors {
    RED = 1,
    GREEN = 2,
    BLUE = 4
};

int main() {
    // Using #define constant
    float area1 = PI_DEFINE * 10 * 10;

    // Using const variable
    float area2 = PI_CONST * 15 * 15;

    // Using enum constants
    int selected_colors = RED | BLUE;

    printf("Area using #define: %.2f\n", area1);
    printf("Area using const: %.2f\n", area2);
    printf("Selected colors: %d\n", selected_colors);

    return 0;
}

```

Storage Class Demonstration

```

#include <stdio.h>

int global_counter = 0; // Global variable

void demonstrate_storage_classes() {
    auto int auto_var = 10; // Automatic storage (default)
    register int reg_var = 20; // Register storage
    static int static_var = 0; // Static storage

    global_counter++;
    static_var++;

    printf("Auto var: %d\n", auto_var);
    printf("Register var: %d\n", reg_var);
    printf("Static var: %d\n", static_var);
    printf("Global counter: %d\n", global_counter);
    printf("----\n");
}

int main() {

```

```

    printf("First call:\n");
    demonstrate_storage_classes();

    printf("Second call:\n");
    demonstrate_storage_classes();

    printf("Third call:\n");
    demonstrate_storage_classes();

    return 0;
}

```

Best Practices

1. Initialize Variables

```

// Good practice
int count = 0;
float price = 0.0f;

// Avoid (uninitialized variables)
int count; // Contains garbage value

```

2. Use Meaningful Names

```

// Good
int student_age;
float account_balance;
char first_name[50];

// Avoid
int x;
float y;
char z[50];

```

3. Use const for Fixed Values

```

// Good
const int MAX_STUDENTS = 100;
const float TAX_RATE = 0.08f;

// Avoid

```

```
#define MAX_STUDENTS 100
#define TAX_RATE 0.08
```

4. Limit Global Variables

```
// Prefer local variables when possible
void calculate_area(float radius) {
    const float PI = 3.14159f; // Local constant
    float area = PI * radius * radius;
    printf("Area: %.2f\n", area);
}
```

5. Use Appropriate Storage Classes

```
// Use static for persistent local state
void counter() {
    static int count = 0;
    count++;
    printf("Call #%d\n", count);
}
```

Summary

In this chapter, you've learned about:

1. **Variables:** Declaration, initialization, naming rules, and assignment
2. **Constants:** `#define`, `const`, and `enum` for creating fixed values
3. **Storage Classes:** `auto`, `register`, `static`, and `extern`
4. **Scope and Lifetime:** Local vs global scope, variable lifetime
5. **Type Qualifiers:** `const`, `volatile`, and `restrict`
6. **Best Practices:** Proper variable naming, initialization, and usage

Understanding variables and constants is essential for effective C programming. In the next chapter, we'll explore type conversions and how C handles different data types in expressions.

Type Conversions

Introduction

Type conversion, also known as type casting, is the process of converting a value from one data type to another. C provides both implicit (automatic) and explicit (manual) type conversion mechanisms. Understanding how type conversions work is crucial for writing correct and efficient C programs, as improper conversions can lead to unexpected behavior, data loss, or undefined results.

This chapter will explore both implicit and explicit type conversions, their rules, potential pitfalls, and best practices for safe conversions.

Implicit Type Conversions (Coercion)

Implicit type conversions occur automatically when the compiler converts one data type to another without explicit programmer intervention. These conversions follow specific rules to ensure data integrity while maintaining program functionality.

Usual Arithmetic Conversions

When operators are applied to operands of different types, C performs usual arithmetic conversions to bring them to a common type:

```
#include <stdio.h>

int main() {
    int i = 10;
    float f = 3.14f;
    double d = 2.5;

    // int + float -> float
    float result1 = i + f;
    printf("int + float = %.2f (float)\n", result1);

    // float + double -> double
    double result2 = f + d;
    printf("float + double = %.2f (double)\n", result2);

    // int + double -> double
```



```

    double result3 = i + d;
    printf("int + double = %.2f (double)\n", result3);

    return 0;
}

```

Integer Promotion

Integer promotion is the automatic conversion of smaller integer types to `int` or `unsigned int`:

```

#include <stdio.h>

int main() {
    char c = 'A';
    short s = 100;

    // char and short are promoted to int in expressions
    int result = c + s;
    printf("char + short = %d (int)\n", result);

    // Promotion also happens with unary operators
    int negated = -c;
    printf("Negated char: %d\n", negated);

    return 0;
}

```

Conversion Hierarchy

C follows a specific hierarchy for implicit conversions:

1. **long double** (highest)
2. **double**
3. **float**
4. **unsigned long long**
5. **long long**
6. **unsigned long**
7. **long**
8. **unsigned int**
9. **int** (lowest)

When operands of different types are used together, the lower type is converted to the higher type:

```

#include <stdio.h>

int main() {

```

```

char c = 'X';
short s = 1000;
int i = 50000;
float f = 3.14f;
double d = 2.71828;

// All converted to double
double result = c + s + i + f + d;
printf("Mixed arithmetic result: %.5f\n", result);

return 0;
}

```

Boolean Context Conversions

In conditional expressions, all scalar types are implicitly converted to boolean values:

```

#include <stdio.h>

int main() {
    int number = 42;
    float zero_float = 0.0f;
    char letter = 'A';
    char null_char = '\0';

    // Non-zero values are true, zero values are false
    if (number) {
        printf("Number %d is true\n", number);
    }

    if (!zero_float) {
        printf("Zero float is false\n");
    }

    if (letter) {
        printf("Character '%c' is true\n", letter);
    }

    if (!null_char) {
        printf("Null character is false\n");
    }

    return 0;
}

```

Explicit Type Conversions (Casting)

Explicit type conversions, or casting, allow programmers to manually convert values from one type to another using cast operators.

Cast Operator Syntax

The cast operator has the form `(type)` where `type` is the target data type:

```
#include <stdio.h>

int main() {
    double pi = 3.14159;
    int integer_pi = (int)pi; // Cast double to int

    printf("Original: %.5f\n", pi);
    printf("Casted: %d\n", integer_pi);

    return 0;
}
```

Common Casting Scenarios

Floating-Point to Integer

```
#include <stdio.h>

int main() {
    float temperature = 98.6f;
    int temp_int = (int)temperature;

    printf("Float temperature: %.1f\n", temperature);
    printf("Integer temperature: %d\n", temp_int);

    // Note: Truncation, not rounding
    float negative = -3.7f;
    int neg_int = (int)negative;
    printf("Negative float: %.1f -> int: %d\n", negative, neg_int);

    return 0;
}
```

Integer to Floating-Point

```
#include <stdio.h>

int main() {
    int total = 100;
    int count = 3;

    // Without casting, integer division occurs
    float average_int_div = total / count;
    printf("Integer division: %d / %d = %.1f\n", total, count, average_int_div);

    // With casting, floating-point division occurs
    float average_float_div = (float)total / count;
    printf("Float division: %.1f / %d = %.2f\n", (float)total, count, average_float_div);

    return 0;
}
```

Character to Integer and Vice Versa

```
#include <stdio.h>

int main() {
    char letter = 'A';
    int ascii_value = (int)letter;

    printf("Character: %c\n", letter);
    printf("ASCII value: %d\n", ascii_value);

    int number = 66;
    char character = (char)number;
    printf("Number: %d\n", number);
    printf("Character: %c\n", character);

    return 0;
}
```

Pointer Casting

C allows casting between different pointer types:

```
#include <stdio.h>
```

```

int main() {
    int number = 42;
    int *int_ptr = &number;

    // Cast int pointer to void pointer
    void *void_ptr = (void*)int_ptr;

    // Cast void pointer back to int pointer
    int *new_int_ptr = (int*)void_ptr;

    printf("Original number: %d\n", number);
    printf("Via pointer: %d\n", *new_int_ptr);

    return 0;
}

```

Potential Pitfalls and Data Loss

Truncation

```

#include <stdio.h>

int main() {
    double large_number = 12345.6789;
    int truncated = (int)large_number;

    printf("Original: %.4f\n", large_number);
    printf("Truncated: %d\n", truncated);

    return 0;
}

```

Overflow

```

#include <stdio.h>
#include <limits.h>

int main() {
    long long big_number = 3000000000LL; // Larger than INT_MAX
    int overflowed = (int)big_number;

    printf("Big number: %lld\n", big_number);
    printf("Overflowed: %d\n", overflowed); // Undefined behavior
}

```

```

    return 0;
}

```

Precision Loss

```

#include <stdio.h>

int main() {
    float precise = 0.1f;
    double more_precise = 0.1;

    printf("Float: %.20f\n", precise);
    printf("Double: %.20f\n", more_precise);

    // Converting double to float loses precision
    float less_precise = (float)more_precise;
    printf("Casted double: %.20f\n", less_precise);

    return 0;
}

```

Safe Conversion Practices

Checking for Overflow

```

#include <stdio.h>
#include <limits.h>

int safe_int_cast(long long value) {
    if (value > INT_MAX || value < INT_MIN) {
        printf("Warning: Value %lld cannot fit in int\n", value);
        return (value > 0) ? INT_MAX : INT_MIN;
    }
    return (int)value;
}

int main() {
    long long big_value = 3000000000LL;
    int safe_value = safe_int_cast(big_value);

    printf("Safe value: %d\n", safe_value);

    return 0;
}

```

```
}
```

Using Standard Library Functions

```
#include <stdio.h>
#include <stdlib.h>

int main() {
    char number_str[] = "12345";
    long parsed_number = strtol(number_str, NULL, 10);

    printf("String: %s\n", number_str);
    printf("Parsed long: %ld\n", parsed_number);

    // Safe conversion with range checking
    if (parsed_number <= INT_MAX && parsed_number >= INT_MIN) {
        int converted = (int)parsed_number;
        printf("Converted to int: %d\n", converted);
    } else {
        printf("Number too large for int\n");
    }

    return 0;
}
```

Conversion Functions

Mathematical Conversions

```
#include <stdio.h>
#include <math.h>

int main() {
    double value = -3.7;

    printf("Original: %.1f\n", value);
    printf("Floor: %.1f\n", floor(value));    // Round down
    printf("Ceil: %.1f\n", ceil(value));      // Round up
    printf("Truncate: %.1f\n", trunc(value)); // Truncate towards zero
    printf("Round: %.1f\n", round(value));    // Round to nearest

    return 0;
}
```

String Conversions

```
#include <stdio.h>
#include <stdlib.h>

int main() {
    // String to integer
    char int_str[] = "123";
    int int_val = atoi(int_str);
    printf("atoi(\"%s\") = %d\n", int_str, int_val);

    // String to long
    char long_str[] = "123456789";
    long long_val = atol(long_str);
    printf("atol(\"%s\") = %ld\n", long_str, long_val);

    // String to double
    char double_str[] = "3.14159";
    double double_val = atof(double_str);
    printf("atof(\"%s\") = %.5f\n", double_str, double_val);

    // More robust conversions with error checking
    char *endptr;
    long strtol_val = strtol("456abc", &endptr, 10);
    printf("strtol(\"456abc\") = %ld, stopped at: %s\n", strtol_val, endptr);

    return 0;
}
```

Practical Examples

Temperature Converter with Type Safety

```
#include <stdio.h>

// Safe conversion from Fahrenheit to Celsius
double fahrenheit_to_celsius(double fahrenheit) {
    return (fahrenheit - 32.0) * 5.0 / 9.0;
}

// Safe conversion from Celsius to Fahrenheit
double celsius_to_fahrenheit(double celsius) {
    return (celsius * 9.0 / 5.0) + 32.0;
}
```



```

int main() {
    float user_input;
    int choice;

    printf("Temperature Converter\n");
    printf("1. Fahrenheit to Celsius\n");
    printf("2. Celsius to Fahrenheit\n");
    printf("Enter choice: ");
    scanf("%d", &choice);

    if (choice == 1) {
        printf("Enter temperature in Fahrenheit: ");
        scanf("%f", &user_input);
        double result = fahrenheit_to_celsius((double)user_input);
        printf("%.2f°F = %.2f°C\n", user_input, result);
    } else if (choice == 2) {
        printf("Enter temperature in Celsius: ");
        scanf("%f", &user_input);
        double result = celsius_to_fahrenheit((double)user_input);
        printf("%.2f°C = %.2f°F\n", user_input, result);
    } else {
        printf("Invalid choice\n");
    }

    return 0;
}

```

Data Type Range Checker

```

#include <stdio.h>
#include <limits.h>
#include <float.h>

void check_int_range(long long value) {
    printf("Checking value: %lld\n", value);

    if (value > INT_MAX) {
        printf(" Too large for int (max: %d)\n", INT_MAX);
    } else if (value < INT_MIN) {
        printf(" Too small for int (min: %d)\n", INT_MIN);
    } else {
        printf(" Fits in int range\n");
        int converted = (int)value;
        printf(" Converted value: %d\n", converted);
    }
}

```

```

    printf("\n");
}

int main() {
    long long test_values[] = {
        100LL,
        3000000000LL,
        -3000000000LL,
        INT_MAX,
        INT_MIN,
        (long long)INT_MAX + 1,
        (long long)INT_MIN - 1
    };

    int num_tests = sizeof(test_values) / sizeof(test_values[0]);

    for (int i = 0; i < num_tests; i++) {
        check_int_range(test_values[i]);
    }

    return 0;
}

```

Mixed Type Calculator

```

#include <stdio.h>

int main() {
    int int_val = 10;
    float float_val = 3.14f;
    double double_val = 2.71828;
    char char_val = 'A'; // ASCII 65

    printf("=== Mixed Type Arithmetic ===\n");
    printf("int: %d\n", int_val);
    printf("float: %.2f\n", float_val);
    printf("double: %.5f\n", double_val);
    printf("char: %c (ASCII: %d)\n", char_val, char_val);
    printf("\n");

    // Various combinations
    printf("int + char: %d + %d = %d\n", int_val, char_val, int_val + char_val);
    printf("float + int: %.2f + %d = %.2f\n", float_val, int_val, float_val + int_val);
    printf("double + float: %.5f + %.2f = %.5f\n", double_val, float_val, double_val + float_val);
    printf("char + double: %d + %.5f = %.5f\n", char_val, double_val, char_val + double_val);
}

```

```

    // Explicit casting examples
    printf("\n=== Explicit Casting ===\n");
    printf("double to int: (int)%.5f = %d\n", double_val, (int)double_val);
    printf("int to float: (float)%d = %.1f\n", int_val, (float)int_val);
    printf("float to char: (char)%.2f = %c\n", float_val, (char)float_val);

    return 0;
}

```

Best Practices

1. Be Explicit About Conversions

```

// Good: Clear intent
float average = (float)total / count;

// Less clear: Relies on implicit conversion
float average = total / (float)count;

```

2. Check for Data Loss

```

// Good: Range checking
if (big_value <= INT_MAX && big_value >= INT_MIN) {
    int small_value = (int)big_value;
}

```

3. Use Appropriate Functions for String Conversion

```

// Good: Error checking with strtol
char *endptr;
long value = strtol(string, &endptr, 10);
if (*endptr != '\0') {
    printf("Invalid number format\n");
}

```

4. Understand Truncation vs Rounding

```
#include <math.h>
// For truncation (default casting)
int truncated = (int)3.7; // Results in 3

// For proper rounding
int rounded = (int)round(3.7); // Results in 4
```

5. Be Careful with Signed/Unsigned Conversions

```
// Dangerous conversion
unsigned int u_value = 4000000000U;
int s_value = (int)u_value; // Implementation-defined behavior
```

Summary

In this chapter, you've learned about:

1. **Implicit Conversions:** Automatic type conversions following C's hierarchy
2. **Explicit Conversions:** Manual casting using the cast operator
3. **Integer Promotion:** Automatic promotion of smaller types to `int`
4. **Arithmetic Conversions:** Rules for mixed-type arithmetic operations
5. **Potential Pitfalls:** Truncation, overflow, and precision loss
6. **Safe Practices:** Range checking and proper conversion techniques
7. **Standard Library Functions:** `atoi`, `atof`, `strtol` for string conversions
8. **Best Practices:** Writing safe and clear conversion code

Understanding type conversions is essential for writing robust C programs. In the next chapter, we'll explore operators and expressions, which heavily rely on type conversion rules.

Operators and Expressions

Introduction

Operators are symbols that perform specific operations on operands (variables, constants, or expressions). Expressions are combinations of operators and operands that evaluate to a single value. Understanding operators and how to construct expressions is fundamental to C programming, as they form the building blocks of all program logic.

This chapter will cover all major categories of operators in C, their precedence and associativity rules, and how to write effective expressions.

Categories of Operators

C provides a rich set of operators that can be categorized as follows:

1. **Arithmetic Operators:** Perform mathematical operations
2. **Relational Operators:** Compare values
3. **Logical Operators:** Combine boolean expressions
4. **Bitwise Operators:** Manipulate individual bits
5. **Assignment Operators:** Assign values to variables
6. **Increment/Decrement Operators:** Modify values by one
7. **Conditional Operator:** Ternary selection operator
8. **Special Operators:** sizeof, address-of, dereference, etc.

Arithmetic Operators

Arithmetic operators perform basic mathematical operations:

Operator	Description	Example	Result
+	Addition	5 + 3	8
-	Subtraction	5 - 3	2
*	Multiplication	5 * 3	15
/	Division	5 / 3	1 (integer division)
%	Modulus (remainder)	5 % 3	2

Basic Arithmetic Operations

```
#include <stdio.h>

int main() {
    int a = 10, b = 3;

    printf("a = %d, b = %d\n", a, b);
    printf("a + b = %d\n", a + b);
    printf("a - b = %d\n", a - b);
    printf("a * b = %d\n", a * b);
    printf("a / b = %d\n", a / b); // Integer division
    printf("a %% b = %d\n", a % b); // Modulus

    // Floating-point arithmetic
    float x = 10.0f, y = 3.0f;
    printf("\nFloat arithmetic:\n");
    printf("x = %.1f, y = %.1f\n", x, y);
    printf("x / y = %.2f\n", x / y); // Floating-point division

    return 0;
}
```

Division and Modulus Behavior

```
#include <stdio.h>

int main() {
    // Integer division truncates toward zero
    printf("Integer division:\n");
    printf("7 / 3 = %d\n", 7 / 3); // 2
    printf("-7 / 3 = %d\n", -7 / 3); // -2
    printf("7 / -3 = %d\n", 7 / -3); // -2
    printf("-7 / -3 = %d\n", -7 / -3); // 2

    // Modulus result has the same sign as the dividend
    printf("\nModulus operations:\n");
    printf("7 %% 3 = %d\n", 7 % 3); // 1
    printf("-7 %% 3 = %d\n", -7 % 3); // -1
    printf("7 %% -3 = %d\n", 7 % -3); // 1
    printf("-7 %% -3 = %d\n", -7 % -3); // -1

    return 0;
}
```

Relational Operators

Relational operators compare two values and return a boolean result (1 for true, 0 for false):

Operator	Description	Example	Result
==	Equal to	5 == 3	0 (false)
!=	Not equal to	5 != 3	1 (true)
>	Greater than	5 > 3	1 (true)
<	Less than	5 < 3	0 (false)
>=	Greater than or equal to	5 >= 5	1 (true)
<=	Less than or equal to	5 <= 3	0 (false)

Relational Operations

```
#include <stdio.h>

int main() {
    int x = 10, y = 20;

    printf("x = %d, y = %d\n", x, y);
    printf("x == y: %d\n", x == y);
    printf("x != y: %d\n", x != y);
    printf("x > y: %d\n", x > y);
    printf("x < y: %d\n", x < y);
    printf("x >= y: %d\n", x >= y);
    printf("x <= y: %d\n", x <= y);

    // Floating-point comparison (be careful!)
    float a = 0.1f + 0.2f;
    float b = 0.3f;
    printf("\nFloating-point comparison:\n");
    printf("a = %.20f\n", a);
    printf("b = %.20f\n", b);
    printf("a == b: %d\n", a == b);

    return 0;
}
```

Logical Operators

Logical operators work with boolean values and are used to combine or invert conditions:

Operator	Description	Example	Result
&&	Logical AND	1 && 0	0 (false)
	Logical OR	1 0	1 (true)
!	Logical NOT	!1	0 (false)

Logical Operations

```
#include <stdio.h>

int main() {
    int a = 1, b = 0, c = 1;

    printf("a = %d, b = %d, c = %d\n", a, b, c);
    printf("a && b: %d\n", a && b); // 1 AND 0 = 0
    printf("a || b: %d\n", a || b); // 1 OR 0 = 1
    printf("!a: %d\n", !a);         // NOT 1 = 0
    printf("a && c: %d\n", a && c); // 1 AND 1 = 1
    printf("a || (b && c): %d\n", a || (b && c)); // 1 OR (0 AND 1) = 1

    // Short-circuit evaluation
    printf("\nShort-circuit evaluation:\n");
    int x = 5, y = 0;
    if (y != 0 && x / y > 1) { // Second condition not evaluated
        printf("This won't print\n");
    } else {
        printf("Short-circuit prevented division by zero\n");
    }

    return 0;
}
```

Bitwise Operators

Bitwise operators work at the bit level and are particularly useful in systems programming and embedded development:

Operator	Description	Example	Result
&	Bitwise AND	5 & 3	1
	Bitwise OR	5 3	7
^	Bitwise XOR	5 ^ 3	6
~	Bitwise NOT	~5	-6
<<	Left shift	5 << 1	10
>>	Right shift	5 >> 1	2

Bitwise Operations

```
#include <stdio.h>

void print_binary(unsigned int num) {
    for (int i = 31; i >= 0; i--) {
        printf("%d", (num >> i) & 1);
        if (i % 4 == 0) printf(" ");
    }
    printf("\n");
}

int main() {
    unsigned int a = 12; // Binary: 1100
    unsigned int b = 10; // Binary: 1010

    printf("a = %u, b = %u\n", a, b);
    printf("a in binary: ");
    print_binary(a);
    printf("b in binary: ");
    print_binary(b);

    printf("\nBitwise operations:\n");
    printf("a & b = %u (AND)\n", a & b);
    printf("a | b = %u (OR)\n", a | b);
    printf("a ^ b = %u (XOR)\n", a ^ b);
    printf("~a = %u (NOT)\n", ~a);
    printf("a << 1 = %u (Left shift)\n", a << 1);
    printf("a >> 1 = %u (Right shift)\n", a >> 1);

    // Practical example: Setting, clearing, and toggling bits
    unsigned char flags = 0b00001010; // 10 in decimal
    printf("\nFlag operations:\n");
    printf("Original flags: ");
    print_binary(flags);

    // Set bit 2 (make it 1)
    flags |= (1 << 2);
    printf("After setting bit 2: ");
    print_binary(flags);

    // Clear bit 3 (make it 0)
    flags &= ~(1 << 3);
    printf("After clearing bit 3: ");
    print_binary(flags);
}
```

```

// Toggle bit 1
flags ^= (1 << 1);
printf("After toggling bit 1: ");
print_binary(flags);

return 0;
}

```

Assignment Operators

Assignment operators assign values to variables. C provides several compound assignment operators for convenience:

Operator	Description	Equivalent To
=	Simple assignment	a = b
+=	Add and assign	a = a + b
-=	Subtract and assign	a = a - b
*=	Multiply and assign	a = a * b
/=	Divide and assign	a = a / b
%=	Modulus and assign	a = a % b
&=	Bitwise AND and assign	a = a & b
=	Bitwise OR and assign	a = a b
^=	Bitwise XOR and assign	a = a ^ b
<<=	Left shift and assign	a = a << b
>>=	Right shift and assign	a = a >> b

Assignment Operations

```

#include <stdio.h>

int main() {
    int a = 10;
    printf("Initial value: a = %d\n", a);

    a += 5;    // a = a + 5
    printf("After a += 5: a = %d\n", a);

    a -= 3;    // a = a - 3
    printf("After a -= 3: a = %d\n", a);

    a *= 2;    // a = a * 2
    printf("After a *= 2: a = %d\n", a);
}

```

```

a /= 4;    // a = a / 4
printf("After a /= 4: a = %d\n", a);

a %= 3;    // a = a % 3
printf("After a %= 3: a = %d\n", a);

// Multiple assignments
int x, y, z;
x = y = z = 5; // All variables get value 5
printf("\nMultiple assignment: x = %d, y = %d, z = %d\n", x, y, z);

return 0;
}

```

Increment and Decrement Operators

These operators increase or decrease a variable's value by 1:

Operator	Description	Example	Result
++	Pre-increment	++a	Increment a, then use new value
++	Post-increment	a++	Use current value, then increment a
--	Pre-decrement	--a	Decrement a, then use new value
--	Post-decrement	a--	Use current value, then decrement a

Increment/Decrement Examples

```

#include <stdio.h>

int main() {
    int a = 5, b = 5;
    int result1, result2;

    printf("Initial: a = %d, b = %d\n", a, b);

    // Pre-increment
    result1 = ++a;
    printf("After ++a: a = %d, result1 = %d\n", a, result1);

    // Post-increment
    result2 = b++;
    printf("After b++: b = %d, result2 = %d\n", b, result2);

    // Practical example in loops
}

```

```

printf("\nLoop examples:\n");

// Using pre-increment
printf("Pre-increment loop:\n");
for (int i = 0; i < 3; ++i) {
    printf("i = %d\n", i);
}

// Using post-increment
printf("Post-increment loop:\n");
for (int i = 0; i < 3; i++) {
    printf("i = %d\n", i);
}

// Be careful with complex expressions
int x = 10;
int y = ++x + x++; // Undefined behavior!
printf("\nUndefined behavior example: y = %d\n", y);

return 0;
}

```

Conditional (Ternary) Operator

The conditional operator is C's only ternary operator, taking three operands:

```
condition ? value_if_true : value_if_false
```

Conditional Operator Examples

```

#include <stdio.h>

int main() {
    int a = 10, b = 20;

    // Simple conditional assignment
    int max = (a > b) ? a : b;
    printf("Maximum of %d and %d is %d\n", a, b, max);

    // Conditional string selection
    char *message = (a % 2 == 0) ? "even" : "odd";
    printf("%d is %s\n", a, message);
}

```

```

// Nested conditional operators
int score = 85;
char *grade = (score >= 90) ? "A" :
               (score >= 80) ? "B" :
               (score >= 70) ? "C" :
               (score >= 60) ? "D" : "F";
printf("Score %d gets grade %s\n", score, grade);

// Conditional with function calls (be careful with side effects)
int x = 5, y = 10;
int result = (x > y) ? (x += 5) : (y += 5);
printf("After conditional: x = %d, y = %d, result = %d\n", x, y, result);

return 0;
}

```

Special Operators

sizeof Operator

Returns the size (in bytes) of a data type or variable:

```

#include <stdio.h>

int main() {
    printf("Size of char: %zu bytes\n", sizeof(char));
    printf("Size of int: %zu bytes\n", sizeof(int));
    printf("Size of float: %zu bytes\n", sizeof(float));
    printf("Size of double: %zu bytes\n", sizeof(double));
    printf("Size of pointer: %zu bytes\n", sizeof(int*));

    int array[10];
    printf("Size of array[10]: %zu bytes\n", sizeof(array));
    printf("Size of array element: %zu bytes\n", sizeof(array[0]));

    return 0;
}

```

Address-of (&) and Dereference (*) Operators

Used with pointers (covered in detail in Module 6):

```

#include <stdio.h>

```

```

int main() {
    int number = 42;
    int *ptr = &number; // Address-of operator

    printf("Value of number: %d\n", number);
    printf("Address of number: %p\n", (void*)&number);
    printf("Value of ptr: %p\n", (void*)ptr);
    printf("Value pointed to by ptr: %d\n", *ptr); // Dereference operator

    return 0;
}

```

Operator Precedence and Associativity

Operators in C have specific precedence levels and associativity rules that determine the order of evaluation in expressions:

Precedence Levels (Highest to Lowest)

1. **Postfix:** `()`, `[]`, `->`, `.`, `++`, `--` (Left to Right)
2. **Unary:** `++`, `--`, `+`, `-`, `!`, `~`, `(type)`, `*`, `&`, `sizeof` (Right to Left)
3. **Multiplicative:** `*`, `/`, `%` (Left to Right)
4. **Additive:** `+`, `-` (Left to Right)
5. **Shift:** `<<`, `>>` (Left to Right)
6. **Relational:** `<`, `<=`, `>`, `>=` (Left to Right)
7. **Equality:** `==`, `!=` (Left to Right)
8. **Bitwise AND:** `&` (Left to Right)
9. **Bitwise XOR:** `^` (Left to Right)
10. **Bitwise OR:** `|` (Left to Right)
11. **Logical AND:** `&&` (Left to Right)
12. **Logical OR:** `||` (Left to Right)
13. **Conditional:** `?:` (Right to Left)
14. **Assignment:** `=`, `+=`, `-=`, etc. (Right to Left)
15. **Comma:** `,` (Left to Right)

Precedence Examples

```

#include <stdio.h>

int main() {
    int a = 5, b = 3, c = 2;

    // Multiplication has higher precedence than addition

```

```

int result1 = a + b * c;          // 5 + (3 * 2) = 11
int result2 = (a + b) * c;        // (5 + 3) * 2 = 16

printf("a + b * c = %d\n", result1);
printf("(a + b) * c = %d\n", result2);

// Relational operators have lower precedence than arithmetic
int result3 = a > b + c;          // a > (b + c) = 5 > 5 = 0
int result4 = (a > b) + c;        // (a > b) + c = 1 + 2 = 3

printf("a > b + c = %d\n", result3);
printf("(a > b) + c = %d\n", result4);

// Logical operators have lower precedence than relational
int x = 10, y = 5, z = 15;
int result5 = x > y && y < z;      // (x > y) && (y < z) = 1 && 1 = 1
int result6 = x > y & y < z;       // x > (y & y) < z = 10 > 1 < 15 = 1

printf("x > y && y < z = %d\n", result5);
printf("x > y & y < z = %d\n", result6);

return 0;
}

```

Practical Examples

Mathematical Expression Evaluator

```

#include <stdio.h>

int main() {
    double x, y, z;

    printf("Enter three numbers: ");
    scanf("%lf %lf %lf", &x, &y, &z);

    // Complex mathematical expression
    double result = (x + y) * z / (x - y) + (x * y - z);

    printf("Result of (x + y) * z / (x - y) + (x * y - z):\n");
    printf("= (%.2f + %.2f) * %.2f / (%.2f - %.2f) + (%.2f * %.2f - %.2f)\n",
           x, y, z, x, y, x, y, z);
    printf("= %.2f\n", result);
}

```

```

    return 0;
}

```

Bit Manipulation Utility

```

#include <stdio.h>

// Function to check if a bit is set
int is_bit_set(unsigned int num, int position) {
    return (num >> position) & 1;
}

// Function to set a bit
unsigned int set_bit(unsigned int num, int position) {
    return num | (1 << position);
}

// Function to clear a bit
unsigned int clear_bit(unsigned int num, int position) {
    return num & ~(1 << position);
}

// Function to toggle a bit
unsigned int toggle_bit(unsigned int num, int position) {
    return num ^ (1 << position);
}

void print_bits(unsigned int num) {
    printf("Binary: ");
    for (int i = 31; i >= 0; i--) {
        printf("%d", (num >> i) & 1);
        if (i % 4 == 0) printf(" ");
    }
    printf("\n");
}

int main() {
    unsigned int number = 0b10101010;

    printf("Original number: %u\n", number);
    print_bits(number);

    // Check bit at position 3
    printf("Bit at position 3: %d\n", is_bit_set(number, 3));
}

```



```

// Set bit at position 2
number = set_bit(number, 2);
printf("After setting bit 2: %u\n", number);
print_bits(number);

// Clear bit at position 7
number = clear_bit(number, 7);
printf("After clearing bit 7: %u\n", number);
print_bits(number);

// Toggle bit at position 5
number = toggle_bit(number, 5);
printf("After toggling bit 5: %u\n", number);
print_bits(number);

return 0;
}

```

Conditional Logic Demonstrator

```

#include <stdio.h>

int main() {
    int score;
    printf("Enter student score (0-100): ");
    scanf("%d", &score);

    // Validate input
    if (score < 0 || score > 100) {
        printf("Invalid score!\n");
        return 1;
    }

    // Determine letter grade using conditional operators
    char *letter_grade = (score >= 90) ? "A" :
        (score >= 80) ? "B" :
        (score >= 70) ? "C" :
        (score >= 60) ? "D" : "F";

    // Determine pass/fail status
    char *status = (score >= 60) ? "PASS" : "FAIL";

    // Provide feedback
    char *feedback = (score >= 90) ? "Excellent!" :
        (score >= 80) ? "Good job!" :

```

```

        (score >= 70) ? "Satisfactory" :
        (score >= 60) ? "Needs improvement" : "Failed";

    printf("\nResults:\n");
    printf("Score: %d\n", score);
    printf("Grade: %s\n", letter_grade);
    printf("Status: %s\n", status);
    printf("Feedback: %s\n", feedback);

    return 0;
}

```

Best Practices

1. Use Parentheses for Clarity

```

// Good: Clear intent
int result = (a + b) * c;

// Less clear: Relies on precedence knowledge
int result = a + b * c;

```

2. Be Careful with Side Effects

```

// Dangerous: Undefined behavior
int a = 5;
int result = a++ + ++a;

// Safe: Clear separation of operations
int a = 5;
int temp1 = a++;
int temp2 = ++a;
int result = temp1 + temp2;

```

3. Use Meaningful Variable Names in Complex Expressions

```

// Good: Self-documenting
int base_salary = 50000;
int bonus_percentage = 10;
int years_of_service = 5;
int total_compensation = base_salary + (base_salary * bonus_percentage / 100) * years_of_ser

```

```
// Less clear
int a = 50000;
int b = 10;
int c = 5;
int d = a + (a * b / 100) * c;
```

4. Be Aware of Integer Division

```
// If you want floating-point result
float average = (float)total / count;

// If you want integer result
int average = total / count;
```

5. Use Bitwise Operations for Bit Manipulation

```
// Good for setting flags
flags |= FLAG_READ;

// Good for clearing flags
flags &= ~FLAG_WRITE;

// Good for checking flags
if (flags & FLAG_EXECUTE) {
    // Execute code
}
```

Summary

In this chapter, you've learned about:

1. **Arithmetic Operators:** +, -, *, /, % for mathematical operations
2. **Relational Operators:** ==, !=, >, <, >=, <= for comparisons
3. **Logical Operators:** &&, ||, ! for boolean logic
4. **Bitwise Operators:** &, |, ^, ~, <<, >> for bit manipulation
5. **Assignment Operators:** =, +=, -=, etc. for value assignment
6. **Increment/Decrement Operators:** ++, -- for modifying values by one
7. **Conditional Operator:** ?: for ternary selection
8. **Special Operators:** sizeof, &, * for specific operations
9. **Operator Precedence:** Rules determining evaluation order
10. **Best Practices:** Writing clear and safe expressions

Understanding operators and expressions is crucial for C programming, as they form the foundation for all program logic. With this knowledge, you're now ready to tackle more complex programming constructs like control flow statements in the next module.

Module 2: Data Types and Variables

Overview

This module delves into the fundamental building blocks of C programming: data types and variables. You'll learn about C's rich type system, from basic integer and floating-point types to advanced fixed-width types introduced in modern C standards. The module also covers variable declaration, initialization, scope, constants, and type conversions.

Learning Objectives

By the end of this module, you will be able to:

- Understand and use all fundamental C data types
- Declare, initialize, and manipulate variables effectively
- Work with constants using various methods (`#define`, `const`, `enum`)
- Apply storage classes (`auto`, `static`, `register`, `extern`) appropriately
- Perform safe type conversions and understand implicit/explicit casting
- Use operators effectively to create complex expressions
- Understand operator precedence and associativity
- Apply best practices for data type selection and variable management

Chapters

1. **Fundamental Data Types** - Integer, floating-point, character, and boolean types
2. **Variables and Constants** - Declaration, initialization, scope, and storage classes
3. **Type Conversions** - Implicit and explicit conversions, potential pitfalls
4. **Operators and Expressions** - Arithmetic, relational, logical, bitwise, and assignment operators

Key Concepts Covered

- Integer types (`char`, `short`, `int`, `long`, `long long`) and their unsigned variants
- Floating-point types (`float`, `double`, `long double`)
- Fixed-width integer types (`int8_t`, `int16_t`, `int32_t`, `int64_t`)
- Boolean types (`_Bool`, `bool` with `stdbool.h`)
- Variable declaration, initialization, and naming conventions
- Constants using `#define`, `const`, and `enum`
- Storage classes (`auto`, `static`, `register`, `extern`)
- Scope (local, global) and lifetime of variables
- Type qualifiers (`const`, `volatile`, `restrict`)

- Implicit and explicit type conversions
- Arithmetic, relational, logical, bitwise, and assignment operators
- Operator precedence and associativity
- Bit manipulation techniques
- Safe conversion practices

Prerequisites

- Completion of Module 1: Foundations and Environment Setup
- Basic understanding of C program structure
- Familiarity with basic I/O operations

Tools and Technologies

- C compiler (GCC, Clang, or MSVC)
- Text editor or IDE
- Understanding of number systems (binary, hexadecimal)
- Basic mathematical concepts

Estimated Time to Complete

- Reading: 6-8 hours
- Exercises: 8-10 hours
- Projects: 3-5 hours

Assessment

- Complete all chapter exercises
- Successfully implement variable manipulation programs
- Demonstrate understanding of type conversions
- Pass operators and expressions quiz
- Submit a comprehensive data type demonstration program

Next Module

After completing this module, proceed to **Module 3: Control Flow** to learn about decision-making and looping constructs in C.

Additional Resources

- “The C Programming Language” by Kernighan and Ritchie (Chapter 2)
- C standard documentation (ISO/IEC 9899:2024 for C23)
- Online references for `limits.h` and `stdint.h`
- Bit manipulation tutorials and resources
- Operator precedence tables

Control Flow

Conditional Statements

Introduction

Conditional statements are fundamental control structures that allow programs to make decisions and execute different code paths based on specific conditions. They enable programs to respond dynamically to varying inputs, user interactions, and runtime conditions.

In C, conditional statements provide the foundation for implementing logic, validation, error handling, and complex decision-making processes. This chapter explores `if / else`, `switch`, nested conditions, the ternary operator, common pitfalls, and several full worked programs you can compile on Linux.

Compile flag convention used throughout this book:

```
gcc -std=c17 -Wall -Wextra -o program program.c
./program
```

The if Statement

The `if` statement is the most basic conditional construct in C. It executes a block of code only if a specified condition evaluates to true (non-zero).

Basic Syntax

```
if (condition) {
    // Code to execute if condition is true
}
```

Simple if Statement

```
#include <stdio.h>

int main(void) {
    int age = 20;

    if (age >= 18) {
```

```

        printf("You are eligible to vote.\n");
    }

    return 0;
}

gcc -std=c17 -Wall -Wextra -o if_simple if_simple.c
./if_simple

```

if-else Statement

The `if-else` construct allows you to specify alternative code paths for when the condition is true or false:

```

#include <stdio.h>

int main(void) {
    int number = 15;

    if (number % 2 == 0) {
        printf("%d is even.\n", number);
    } else {
        printf("%d is odd.\n", number);
    }

    return 0;
}

```

if-else if-else Chain

For multiple conditions, chain `if-else if-else` statements. Order matters: the first true branch wins.

```

#include <stdio.h>

int main(void) {
    int score = 85;

    if (score >= 90) {
        printf("Grade: A\n");
    } else if (score >= 80) {
        printf("Grade: B\n");
    } else if (score >= 70) {
        printf("Grade: C\n");
    }
}

```

```

    } else if (score >= 60) {
        printf("Grade: D\n");
    } else {
        printf("Grade: F\n");
    }

    return 0;
}

```

The switch Statement

The `switch` statement provides an alternative to long `if-else` if chains when comparing a single expression against multiple **constant integer** values (including character constants).

Basic Syntax

```

switch (expression) {
    case constant1:
        // Code for constant1
        break;
    case constant2:
        // Code for constant2
        break;
    default:
        // Code for no match
        break;
}

```

Simple switch Example

```

#include <stdio.h>

int main(void) {
    int day = 3;

    switch (day) {
        case 1: printf("Monday\n");    break;
        case 2: printf("Tuesday\n");   break;
        case 3: printf("Wednesday\n"); break;
        case 4: printf("Thursday\n");  break;
        case 5: printf("Friday\n");    break;
        case 6: printf("Saturday\n");  break;
    }
}

```

```

        case 7: printf("Sunday\n");    break;
        default: printf("Invalid day\n"); break;
    }

    return 0;
}

```

Intentional Fall-Through

Without `break`, control falls into the next case. Use this only deliberately, and document it:

```

#include <stdio.h>

int main(void) {
    int month = 2;

    switch (month) {
        case 12:
        case 1:
        case 2:
            printf("Winter\n");
            break;
        case 3:
        case 4:
        case 5:
            printf("Spring\n");
            break;
        case 6:
        case 7:
        case 8:
            printf("Summer\n");
            break;
        case 9:
        case 10:
        case 11:
            printf("Autumn\n");
            break;
        default:
            printf("Invalid month\n");
            break;
    }

    return 0;
}

```

```
gcc -std=c17 -Wall -Wextra -o seasons seasons.c
./seasons
```

Nested Conditionals

Conditionals can be nested to handle multi-factor decisions:

```
#include <stdio.h>
#include <stdbool.h>

int main(void) {
    int age = 25;
    bool has_license = true;

    if (age >= 18) {
        if (has_license) {
            printf("You can drive legally.\n");
        } else {
            printf("You need to get a license.\n");
        }
    } else {
        printf("You are not old enough to drive.\n");
    }

    return 0;
}
```

Nested Condition Pitfalls

Deep nesting is hard to read and easy to get wrong. Prefer early returns or combined boolean expressions when the logic allows.

```
#include <stdio.h>
#include <stdbool.h>

/* Nested - correct but noisy */
static void check_access_nested(int age, bool licensed, bool insured) {
    if (age >= 18) {
        if (licensed) {
            if (insured) {
                printf("Drive allowed.\n");
            } else {
                printf("Need insurance.\n");
            }
        }
    }
}
```

```

        } else {
            printf("Need license.\n");
        }
    } else {
        printf("Too young.\n");
    }
}

/* Flattened with early exits - usually clearer */
static void check_access_flat(int age, bool licensed, bool insured) {
    if (age < 18) {
        printf("Too young.\n");
        return;
    }
    if (!licensed) {
        printf("Need license.\n");
        return;
    }
    if (!insured) {
        printf("Need insurance.\n");
        return;
    }
    printf("Drive allowed.\n");
}

int main(void) {
    check_access_nested(20, true, false);
    check_access_flat(20, true, false);
    return 0;
}

```

The Dangling else (Else-Bind Ambiguity)

In C, an `else` always binds to the **nearest unmatched** `if`. Without braces this can surprise you.

```

#include <stdio.h>

int main(void) {
    int x = 0;
    int y = 1;

    /* DANGEROUS style (shown only to illustrate binding):
     *
     * if (x)

```

```

*      if (y)
*          printf("both\n");
*  else
*      printf("x was false?\n");    // actually binds to if (y)!
*
*  Always use braces:
*/

if (x) {
    if (y) {
        printf("both\n");
    }
} else {
    printf("x was false (this else belongs to if (x))\n");
}

/* Explicit binding for the inner else */
if (x) {
    if (y) {
        printf("both\n");
    } else {
        printf("x true, y false\n");
    }
}

return 0;
}

```

```

gcc -std=c17 -Wall -Wextra -o dangling_else dangling_else.c
./dangling_else

```

Rule: always brace `if` / `else` bodies, even one-liners. It eliminates the ambiguity.

Conditional (Ternary) Expressions

C supports the conditional operator as a shorthand for simple value selection:

```

#include <stdio.h>

int main(void) {
    int a = 10, b = 20;

    int max_if;
    if (a > b) {
        max_if = a;
    }
}

```

```

    } else {
        max_if = b;
    }

    int max_ternary = (a > b) ? a : b;

    printf("Max (if-else): %d\n", max_if);
    printf("Max (ternary): %d\n", max_ternary);

    int age = 17;
    const char *status = (age >= 18) ? "adult" : "minor";
    printf("Status: %s\n", status);

    return 0;
}

```

Ternary Best Practices

Prefer ternary when...	Prefer if/else when...
Selecting a value for assignment/return	The branches have side effects (I/O, mutations)
Expression stays on one readable line	Nested ternaries would appear (a ? b : c ? d : e)
Types of both arms match cleanly	Branches need multiple statements

```

#include <stdio.h>

/* Good: pure value selection */
static int clamp(int v, int lo, int hi) {
    return (v < lo) ? lo : (v > hi) ? hi : v; /* nested: still OK if tiny */
}

/* Better readability for nested selection */
static int clamp_clear(int v, int lo, int hi) {
    if (v < lo) {
        return lo;
    }
    if (v > hi) {
        return hi;
    }
    return v;
}

/* Bad: side effects buried in a ternary */

```



```

static void bad_style(int flag) {
    /* avoid: flag ? printf("yes\n") : printf("no\n"); */
    if (flag) {
        printf("yes\n");
    } else {
        printf("no\n");
    }
}

int main(void) {
    printf("clamp(150, 0, 100) = %d\n", clamp(150, 0, 100));
    printf("clamp_clear(-5, 0, 100) = %d\n", clamp_clear(-5, 0, 100));
    bad_style(1);
    return 0;
}

gcc -std=c17 -Wall -Wextra -o ternary_style ternary_style.c
./ternary_style

```

Best Practices for Conditionals

1. Use Braces Consistently

```

/* Good */
if (condition) {
    do_something();
}

/* Avoid: inconsistent / brace-less style */
if (condition)
    do_something();
else
    do_something_else();

```

2. Handle All Cases in switch

Always include `default` for unexpected values (or document why it is impossible). Always `break` unless fall-through is intentional and commented.

3. Prefer Early Returns Over Deep Nesting

```
int process_data(int data) {
    if (data < 0) {
        return -1;
    }
    if (data > 100) {
        return -2;
    }
    return data * 2;
}
```

4. Name Complex Conditions

```
int is_adult = (age >= 18);
int has_permission = (permission_level > 0);

if (is_adult && has_permission) {
    grant_access();
}
```

Common Pitfalls

1. Assignment vs Comparison

```
#include <stdio.h>

int main(void) {
    int x = 5;

    /* Wrong: assignment; condition is always true here after assign */
    if (x = 10) {
        printf("Assigned and truthy: x = %d\n", x);
    }

    /* Correct comparison */
    if (x == 10) {
        printf("x equals 10\n");
    }

    /* Yoda style some teams use to catch typos: if (10 == x) */
    if (10 == x) {
        printf("Yoda comparison OK\n");
    }
}
```

```

    }

    return 0;
}

```

With `-Wall`, gcc warns about assignment used as a condition unless you add extra parentheses: `if ((x = 10))`.

2. Floating-Point Comparisons

```

#include <stdio.h>
#include <math.h>

int main(void) {
    double a = 0.1 + 0.2;
    double b = 0.3;

    if (a == b) {
        printf("Equal\n");
    } else {
        printf("Not bit-equal (a = %.20f, b = %.20f)\n", a, b);
    }

    const double epsilon = 1e-9;
    if (fabs(a - b) < epsilon) {
        printf("Approximately equal\n");
    }

    return 0;
}

gcc -std=c17 -Wall -Wextra -o fp_cmp fp_cmp.c -lm
./fp_cmp

```

3. Missing break in switch

Fall-through without intent is a classic bug. Enable and respect warnings; some projects use `/* fallthrough */` comments for intentional cases (GCC recognizes them).

Worked Example 1: Full Grading System

Letter grade, pass/fail, GPA-style points, and input validation.

```

/* grade_system.c - interactive grading helper */
#include <stdio.h>

static char letter_grade(float score) {
    if (score >= 90.0f) {
        return 'A';
    }
    if (score >= 80.0f) {
        return 'B';
    }
    if (score >= 70.0f) {
        return 'C';
    }
    if (score >= 60.0f) {
        return 'D';
    }
    return 'F';
}

static float grade_points(char letter) {
    switch (letter) {
        case 'A': return 4.0f;
        case 'B': return 3.0f;
        case 'C': return 2.0f;
        case 'D': return 1.0f;
        default: return 0.0f;
    }
}

static const char *feedback(char letter) {
    switch (letter) {
        case 'A': return "Excellent";
        case 'B': return "Good";
        case 'C': return "Satisfactory";
        case 'D': return "Needs improvement";
        default: return "Failing - retake recommended";
    }
}

int main(void) {
    float score;

    printf("Enter student score (0-100): ");
    if (scanf("%f", &score) != 1) {
        printf("Invalid input (not a number).\n");
        return 1;
    }
}

```

```

if (score < 0.0f || score > 100.0f) {
    printf("Invalid score! Enter a value between 0 and 100.\n");
    return 1;
}

char letter = letter_grade(score);
float gpa = grade_points(letter);

printf("Score : %.1f\n", score);
printf("Grade : %c (%s)\n", letter, feedback(letter));
printf("Points: %.1f / 4.0\n", gpa);

if (score >= 60.0f) {
    printf("Result: PASS\n");
} else {
    printf("Result: FAIL\n");
}

/* Plus/minus band using nested conditions */
if (letter != 'F') {
    int band = (int)score % 10;
    if (band >= 7 || score >= 97.0f) {
        printf("Band : %c+\n", letter);
    } else if (band <= 2 && score < 90.0f) {
        printf("Band : %c-\n", letter);
    } else {
        printf("Band : %c\n", letter);
    }
}

return 0;
}

```

```

gcc -std=c17 -Wall -Wextra -o grade_system grade_system.c
echo "87.5" | ./grade_system

```

Worked Example 2: Leap Year (Rules + Nested Form)

Gregorian rule: divisible by 4, except centuries unless also divisible by 400.

```

/* leap_year.c */
#include <stdio.h>

```

```

#include <stdbool.h>

static bool is_leap_flat(int year) {
    return (year % 4 == 0 && year % 100 != 0) || (year % 400 == 0);
}

static bool is_leap_nested(int year) {
    if (year % 4 == 0) {
        if (year % 100 == 0) {
            if (year % 400 == 0) {
                return true;
            }
            return false;
        }
        return true;
    }
    return false;
}

int main(void) {
    int year;

    printf("Enter a year: ");
    if (scanf("%d", &year) != 1) {
        printf("Invalid input.\n");
        return 1;
    }

    if (year <= 0) {
        printf("Use a positive year (Gregorian-style check).\n");
        return 1;
    }

    bool flat = is_leap_flat(year);
    bool nest = is_leap_nested(year);

    printf("%d is %sa leap year (flat).\n", year, flat ? "" : "not ");
    printf("%d is %sa leap year (nested).\n", year, nest ? "" : "not ");

    /* Spot-check known cases without user input */
    int samples[] = {1900, 2000, 2024, 2023};
    printf("\nSelf-check:\n");
    for (int i = 0; i < 4; i++) {
        int y = samples[i];
        printf("  %d -> %s\n", y, is_leap_flat(y) ? "leap" : "common");
    }
}

```

```

    return 0;
}

```

```

gcc -std=c17 -Wall -Wextra -o leap_year leap_year.c
echo "2000" | ./leap_year

```

Expected self-check: 1900 common, 2000 leap, 2024 leap, 2023 common.

Worked Example 3: Simple ATM Menu with switch

Balance, deposit, withdraw, and PIN check — menu loop with switch.

```

/* atm_menu.c - console ATM simulation */
#include <stdio.h>
#include <stdbool.h>

#define PIN_CODE 1234
#define MAX_TRIES 3

int main(void) {
    double balance = 1000.00;
    int pin;
    int tries = 0;
    bool authenticated = false;

    printf("=== Mini ATM ===\n");
    while (tries < MAX_TRIES) {
        printf("Enter PIN: ");
        if (scanf("%d", &pin) != 1) {
            printf("Invalid input.\n");
            /* drain line */
            int ch;
            while ((ch = getchar()) != '\n' && ch != EOF) {
            }
            tries++;
            continue;
        }
        if (pin == PIN_CODE) {
            authenticated = true;
            break;
        }
        tries++;
        printf("Wrong PIN (%d/%d).\n", tries, MAX_TRIES);
    }
}

```

```

}

if (!authenticated) {
    printf("Card locked. Goodbye.\n");
    return 1;
}

int choice = -1;
do {
    printf("\n--- Menu ---\n");
    printf("1. Check balance\n");
    printf("2. Deposit\n");
    printf("3. Withdraw\n");
    printf("4. Quick $50 withdraw\n");
    printf("0. Exit\n");
    printf("Choice: ");

    if (scanf("%d", &choice) != 1) {
        printf("Please enter a number.\n");
        int ch;
        while ((ch = getchar()) != '\n' && ch != EOF) {
        }
        choice = -1;
        continue;
    }

    switch (choice) {
        case 1:
            printf("Balance: $%.2f\n", balance);
            break;

        case 2: {
            double amount;
            printf("Deposit amount: ");
            if (scanf("%lf", &amount) != 1 || amount <= 0.0) {
                printf("Invalid amount.\n");
                break;
            }
            balance += amount;
            printf("Deposited. New balance: $%.2f\n", balance);
            break;
        }

        case 3: {
            double amount;
            printf("Withdraw amount: ");
            if (scanf("%lf", &amount) != 1 || amount <= 0.0) {

```



```

        printf("Invalid amount.\n");
        break;
    }
    if (amount > balance) {
        printf("Insufficient funds.\n");
    } else {
        balance -= amount;
        printf("Withdrawn. New balance: $%.2f\n", balance);
    }
    break;
}

case 4:
    if (balance < 50.0) {
        printf("Insufficient funds for quick withdraw.\n");
    } else {
        balance -= 50.0;
        printf("Dispensed $50.00. Balance: $%.2f\n", balance);
    }
    break;

case 0:
    printf("Thank you. Session ended.\n");
    break;

default:
    printf("Unknown option.\n");
    break;
}
} while (choice != 0);

return 0;
}

```

```

gcc -std=c17 -Wall -Wextra -o atm_menu atm_menu.c
# Interactive: enter PIN 1234, then try menu options
./atm_menu

```

Short-Circuit Evaluation

&& and || stop evaluating once the result is known. Use this to guard expensive or unsafe operations:

```

#include <stdio.h>

int main(void) {
    int a = 5, b = 0;

    /* Safe: a / b never runs when b == 0 */
    if (b != 0 && a / b > 1) {
        printf("Division is safe and result > 1\n");
    } else {
        printf("Skipped unsafe division or result not > 1\n");
    }

    int calls = 0;
    if (a > 10 && (++calls > 0)) {
        printf("Won't print\n");
    }
    printf("Side-effect calls: %d (short-circuit skipped ++)\n", calls);

    return 0;
}

gcc -std=c17 -Wall -Wextra -o short_circuit short_circuit.c
./short_circuit

```

Exercises

Exercise 1 — Triangle classifier

Read three side lengths. Print whether they form a valid triangle, and if so whether it is equilateral, isosceles, or scalene. Reject non-positive sides.

Solution sketch:

```

/* Valid if a+b>c, a+c>b, b+c>a for all permutations.
 * Then: all equal → equilateral; any two equal → isosceles; else scalene.
 * Use if / else if chain; avoid floating equality issues by using int sides first.
 */

gcc -std=c17 -Wall -Wextra -o triangle triangle.c

```

Exercise 2 — Rock–paper–scissors

Map chars `r/p/s` (or integers 0/1/2) for player and computer. Use `switch` for the winner matrix. Print win/lose/draw.

Solution sketch:

```
/* switch (player) {
 *   case 'r': if (cpu=='s') win; else if (cpu=='p') lose; else draw; break;
 *   ...
 * }
 * Or encode both as ints and use (player - cpu + 3) % 3.
 */
```

Exercise 3 — BMI category

Read weight (kg) and height (m). Compute $\text{BMI} = \text{kg} / \text{m}^2$. Categories: underweight <18.5, normal <25, overweight <30, obese 30. Guard height 0.

Solution sketch:

```
/* if (height <= 0) error;
 * double bmi = weight / (height * height);
 * else-if ladder on thresholds; print bmi with %.1f
 */
```

Exercise 4 — Rewrite dangling-else trap

Given brace-less nested `if/else` that misclassifies, rewrite with braces so the `else` binds to the outer `if`. Add a unit-style table of (x,y) inputs and expected messages.

Exercise 5 — ATM extension

Extend `atm_menu.c` with a transfer option (case 5) that subtracts from balance only if funds suffice, and logs a fake destination account id. Add a daily withdraw limit with a nested check.

Summary

In this chapter you covered:

1. `if` / `if-else` / `else-if` — ordered multi-way branching

2. **switch** — integer/character multi-way with break and intentional fall-through
3. **Nested conditions** — when needed, and how to flatten
4. **Dangling else** — else binds to nearest unmatched if; always use braces
5. **Ternary** — value selection; avoid side effects and deep nesting
6. **Worked programs** — grading system, leap year, ATM menu
7. **Pitfalls** — = vs ==, float compares, missing **break**, short-circuit
8. **Exercises** — triangle, RPS, BMI, else-binding, ATM extensions

Next: loops (**while**, **do-while**, **for**) for repetition and iterative algorithms.

Loops

Introduction

Loops are control structures that allow programs to repeat a block of code multiple times. They are essential for processing collections of data, implementing algorithms, and automating repetitive tasks. C provides three main loop kinds: **while**, **do-while**, and **for**.

This chapter covers syntax, a **for vs while vs do-while** decision guide, nested patterns, infinite loops for servers/embedded-style code, worked examples (primes, FizzBuzz, tables, running average), and exercises.

Compile convention:

```
gcc -std=c17 -Wall -Wextra -o program program.c
./program
```

The while Loop

The **while** loop is a **pre-test** loop: the condition is evaluated before each iteration. The body may run zero times.

Basic Syntax

```
while (condition) {
    // Code to execute repeatedly
}
```

Simple while Loop

```
#include <stdio.h>

int main(void) {
    int count = 1;

    while (count <= 5) {
        printf("Count: %d\n", count);
    }
}
```

```

        count++; /* must update the control variable */
    }

    return 0;
}

gcc -std=c17 -Wall -Wextra -o while_simple while_simple.c
./while_simple

```

while with User Input

```

#include <stdio.h>

int main(void) {
    int number;

    printf("Enter positive numbers (0 to stop):\n");
    if (scanf("%d", &number) != 1) {
        return 1;
    }

    while (number > 0) {
        printf("You entered: %d\n", number);
        if (scanf("%d", &number) != 1) {
            break;
        }
    }

    printf("Loop ended.\n");
    return 0;
}

```

The do-while Loop

The do-while loop is a **post-test** loop: the body runs **at least once**, then the condition is checked.

Basic Syntax

```
do {  
    // Code to execute  
} while (condition); /* note the trailing semicolon */
```

Simple do-while

```
#include <stdio.h>  
  
int main(void) {  
    int count = 1;  
  
    do {  
        printf("Count: %d\n", count);  
        count++;  
    } while (count <= 5);  
  
    return 0;  
}
```

Menu Pattern (classic do-while use)

```
#include <stdio.h>  
  
int main(void) {  
    int choice;  
  
    do {  
        printf("\n=== Menu ===\n");  
        printf("1. Option 1\n");  
        printf("2. Option 2\n");  
        printf("0. Exit\n");  
        printf("Choice: ");  
        if (scanf("%d", &choice) != 1) {  
            int ch;  
            while ((ch = getchar()) != '\n' && ch != EOF) {  
            }  
            choice = -1;  
            continue;  
        }  
  
        switch (choice) {
```

```

        case 1: printf("Option 1\n"); break;
        case 2: printf("Option 2\n"); break;
        case 0: printf("Exiting...\n"); break;
        default: printf("Invalid choice.\n"); break;
    }
} while (choice != 0);

return 0;
}

```

```

gcc -std=c17 -Wall -Wextra -o menu_do menu_do.c
./menu_do

```

The for Loop

The for loop packs initialization, condition, and update in one place. Ideal when iteration count or index range is clear.

Basic Syntax

```

for (initialization; condition; update) {
    // body
}

```

Simple for Loop

```

#include <stdio.h>

int main(void) {
    for (int i = 1; i <= 5; i++) {
        printf("Count: %d\n", i);
    }
    return 0;
}

```

for with Arrays

```

#include <stdio.h>

int main(void) {
    int numbers[] = {10, 20, 30, 40, 50};

```



```

    int size = (int)(sizeof numbers / sizeof numbers[0]);

    for (int i = 0; i < size; i++) {
        printf("numbers[%d] = %d\n", i, numbers[i]);
    }
    return 0;
}

gcc -std=c17 -Wall -Wextra -o for_array for_array.c
./for_array

```

Any of the three `for` clauses may be empty:

```

int i = 0;
for (; i < 10; ) {
    printf("%d\n", i);
    i++;
}

```

Decision Guide: for vs while vs do-while

Situation	Prefer	Why
Known count / index over a range	for	Init, test, and step live together; hard to forget the step
Iterate until a runtime condition (file ends, sentinel input)	while	Condition is the story; body may not run
Must show a menu / prompt at least once	do-while	Post-test guarantees one execution
Event loop / server / firmware “forever”	while (1) or for (;;) 	Exit only via break , return , or signal
Multiple control variables advanced together	for (multi-update)	e.g. <code>for (i = 0, j = n-1; i < j; i++, j--)</code>

Same Task Three Ways

Sum integers from 1 to `n`:

```

#include <stdio.h>

int main(void) {
    int n = 10;
    long sum;

    /* for - clearest for counted range */
    sum = 0;
    for (int i = 1; i <= n; i++) {
        sum += i;
    }
    printf("for:      %ld\n", sum);

    /* while - same logic, control variable outside */
    sum = 0;
    int i = 1;
    while (i <= n) {
        sum += i;
        i++;
    }
    printf("while:    %ld\n", sum);

    /* do-while - awkward if n could be 0; body always runs once */
    sum = 0;
    i = 1;
    if (n >= 1) {
        do {
            sum += i;
            i++;
        } while (i <= n);
    }
    printf("do-while: %ld\n", sum);

    return 0;
}

```

```

gcc -std=c17 -Wall -Wextra -o sum_three sum_three.c
./sum_three

```

When while Beats for

Reading until failure — the “next value” is also the loop condition:

```

#include <stdio.h>

```

```

int main(void) {
    int x;
    long sum = 0;
    int count = 0;

    printf("Enter integers; non-numeric or EOF ends input:\n");
    while (scanf("%d", &x) == 1) {
        sum += x;
        count++;
    }

    if (count == 0) {
        printf("No numbers.\n");
    } else {
        printf("count=%d sum=%ld avg=%.2f\n",
            count, sum, (double)sum / count);
    }
    return 0;
}

```

Loop Control: break and continue

break

Leaves the **innermost** loop immediately.

```

#include <stdio.h>

int main(void) {
    for (int i = 1; i <= 10; i++) {
        if (i == 5) {
            break;
        }
        printf("%d ", i);
    }
    printf("\nDone.\n");
    return 0;
}

```

continue

Skips the rest of the current iteration and proceeds to the next test/update.

```

#include <stdio.h>

int main(void) {
    for (int i = 1; i <= 10; i++) {
        if (i % 2 == 0) {
            continue; /* skip evens */
        }
        printf("%d ", i);
    }
    printf("\n");
    return 0;
}

```

Nested Loops

Multiplication Table Pattern

```

#include <stdio.h>

int main(void) {
    for (int i = 1; i <= 5; i++) {
        for (int j = 1; j <= 5; j++) {
            printf("%4d", i * j);
        }
        printf("\n");
    }
    return 0;
}

```

Nested break/continue Scope

`break` and `continue` affect only the **inner** loop unless you use flags or `goto` (rare) for multi-level exit.

```

#include <stdio.h>

int main(void) {
    for (int i = 1; i <= 3; i++) {
        printf("outer %d\n", i);
        for (int j = 1; j <= 5; j++) {
            if (j == 3) {
                continue; /* next j */
            }
        }
    }
}

```

```

    }
    if (i == 2 && j == 4) {
        break;    /* leave inner only */
    }
    printf("    inner %d\n", j);
}
}
return 0;
}

```

Nested Loop Patterns Cheatsheet

Pattern	Structure	Use
Rectangle	outer rows $\times$ inner cols	grids, matrices, tables
Triangle	inner runs to i or $n-i$	patterns, combinations
Pairwise	for i / for $j = i+1$	prefix unique pairs without double count
2D walk	for r / for c over $a[r][c]$	images, boards

```

#include <stdio.h>

int main(void) {
    int n = 4;

    printf("Triangle:\n");
    for (int i = 1; i <= n; i++) {
        for (int j = 1; j <= i; j++) {
            printf("* ");
        }
        printf("\n");
    }

    printf("\nUnique pairs (i < j):\n");
    for (int i = 0; i < n; i++) {
        for (int j = i + 1; j < n; j++) {
            printf("(%d,%d) ", i, j);
        }
    }
    printf("\n");
    return 0;
}

```

```
}
```

Infinite Loops (Servers, Embedded, Tools)

Infinite loops are normal when the process should run until an explicit stop condition.

Forms

```
while (1) { /* ... */ }
while (true) { /* need stdbool.h */ }
for (;;) { /* idiomatic "forever" */ }
```

Pattern: Event Loop with break Conditions

```
/* event_loop_demo.c - desktop stand-in for a service loop */
#include <stdio.h>
#include <stdbool.h>
#include <string.h>

int main(void) {
    char line[128];
    bool running = true;
    unsigned long ticks = 0;

    printf("Commands: status | quit | panic\n");

    while (running) {
        ticks++;
        printf("> ");
        if (!fgets(line, sizeof line, stdin)) {
            /* EOF: clean shutdown */
            break;
        }

        /* strip newline */
        line[strcspn(line, "\n")] = '\0';

        if (strcmp(line, "quit") == 0) {
            running = false;          /* graceful flag */
        } else if (strcmp(line, "panic") == 0) {
```

```

        printf("Fatal: breaking out immediately.\n");
        break; /* hard exit from loop */
    } else if (strcmp(line, "status") == 0) {
        printf("ticks=%lu running=%d\n", ticks, running);
    } else if (line[0] != '\0') {
        printf("Unknown command: %s\n", line);
    }

    /* In a real server: poll sockets, handle timers, etc. here */
}

printf("Shutdown after %lu loop iterations.\n", ticks);
return 0;
}

gcc -std=c17 -Wall -Wextra -o event_loop_demo event_loop_demo.c
printf 'status\nquit\n' | ./event_loop_demo

```

Pattern: Embedded-Style Superloop

```

#include <stdio.h>
#include <stdbool.h>

/* Simulated "hardware" flags */
static volatile bool g_button_irq;
static volatile bool g_shutdown;

static void poll_inputs(void) {
    /* On bare metal: read GPIO. Here we fake nothing. */
}

static void service_button(void) {
    if (g_button_irq) {
        g_button_irq = false;
        printf("button handled\n");
    }
}

/* Call this many times; break when shutdown requested */
static void superloop_iteration(void) {
    poll_inputs();
    service_button();
    /* feed watchdog, sleep until interrupt, etc. */
}

```

```

int main(void) {
    int simulated_cycles = 0;

    for (;;) {
        superloop_iteration();
        simulated_cycles++;

        /* Host demo stop conditions */
        if (simulated_cycles >= 3) {
            g_shutdown = true;
        }
        if (g_shutdown) {
            break;
        }
    }

    printf("MCU superloop exited cleanly.\n");
    return 0;
}

gcc -std=c17 -Wall -Wextra -o superloop superloop.c
./superloop

```

Pattern: Retry with Cap (avoid true infinite hang)

```

#include <stdio.h>
#include <stdbool.h>

static bool try_connect(int attempt) {
    printf("connect attempt %d\n", attempt);
    return attempt >= 3; /* succeed on 3rd try in this demo */
}

int main(void) {
    const int max_attempts = 5;
    int attempt = 0;
    bool ok = false;

    while (attempt < max_attempts) {
        attempt++;
        if (try_connect(attempt)) {
            ok = true;
            break;
        }
    }
}

```



```
printf(ok ? "Connected.\n" : "Gave up.\n");
return ok ? 0 : 1;
}
```

Worked Example 1: Prime Check

```
/* prime_check.c */
#include <stdio.h>
#include <stdbool.h>
#include <math.h>

static bool is_prime(long n) {
    if (n <= 1) {
        return false;
    }
    if (n <= 3) {
        return true;
    }
    if (n % 2 == 0 || n % 3 == 0) {
        return false;
    }
    /* check 6k ± 1 up to sqrt(n) */
    for (long i = 5; i * i <= n; i += 6) {
        if (n % i == 0 || n % (i + 2) == 0) {
            return false;
        }
    }
    return true;
}

int main(void) {
    long n;
    printf("Enter integer: ");
    if (scanf("%ld", &n) != 1) {
        printf("Invalid input.\n");
        return 1;
    }

    if (is_prime(n)) {
        printf("%ld is prime.\n", n);
    } else {
        printf("%ld is not prime.\n", n);
    }
}
```

```

    }

    printf("Primes up to 50: ");
    for (long i = 2; i <= 50; i++) {
        if (is_prime(i)) {
            printf("%ld ", i);
        }
    }
    printf("\n");
    return 0;
}

gcc -std=c17 -Wall -Wextra -o prime_check prime_check.c
echo 97 | ./prime_check

```

Worked Example 2: FizzBuzz

```

/* fizzbuzz.c */
#include <stdio.h>

int main(void) {
    int limit = 30;

    for (int i = 1; i <= limit; i++) {
        if (i % 15 == 0) {
            printf("FizzBuzz\n");
        } else if (i % 3 == 0) {
            printf("Fizz\n");
        } else if (i % 5 == 0) {
            printf("Buzz\n");
        } else {
            printf("%d\n", i);
        }
    }

    /* Compact variant with continue-style layering */
    printf("--- compact ---\n");
    for (int i = 1; i <= 15; i++) {
        int printed = 0;
        if (i % 3 == 0) {
            printf("Fizz");
            printed = 1;
        }
        if (i % 5 == 0) {
            printf("Buzz");
            printed = 1;
        }
        if (!printed) {
            printf("%d", i);
        }
        printf("\n");
    }
}

```

```

    }
    if (i % 5 == 0) {
        printf("Buzz");
        printed = 1;
    }
    if (!printed) {
        printf("%d", i);
    }
    printf("\n");
}

return 0;
}

gcc -std=c17 -Wall -Wextra -o fizzbuzz fizzbuzz.c
./fizzbuzz

```

Worked Example 3: Multiplication Table

```

/* mult_table.c */
#include <stdio.h>

int main(void) {
    int n;

    printf("Table size (1-12): ");
    if (scanf("%d", &n) != 1 || n < 1 || n > 12) {
        printf("Use an integer from 1 to 12.\n");
        return 1;
    }

    printf("    ");
    for (int c = 1; c <= n; c++) {
        printf("%4d", c);
    }
    printf("\n");

    for (int r = 1; r <= n; r++) {
        printf("%4d", r);
        for (int c = 1; c <= n; c++) {
            printf("%4d", r * c);
        }
    }
}

```

```

        printf("\n");
    }

    return 0;
}

gcc -std=c17 -Wall -Wextra -o mult_table mult_table.c
echo 10 | ./mult_table

```

Worked Example 4: Running Average from Input

Read numbers until a sentinel (or EOF); maintain count and sum for a live average.

```

/* running_avg.c */
#include <stdio.h>

int main(void) {
    double x;
    double sum = 0.0;
    int count = 0;

    printf("Enter numbers (non-numeric or Ctrl-D to finish):\n");

    while (scanf("%lf", &x) == 1) {
        count++;
        sum += x;
        printf("  n=%d  last=%.4f  sum=%.4f  avg=%.4f\n",
               count, x, sum, sum / count);
    }

    if (count == 0) {
        printf("No data.\n");
        return 1;
    }

    printf("Final average of %d values: %.4f\n", count, sum / count);
    return 0;
}

gcc -std=c17 -Wall -Wextra -o running_avg running_avg.c
printf '10 20 30 40\n' | ./running_avg

```

Sentinel variant (stop on a magic value without relying on EOF):

```

/* running_avg_sentinel.c */
#include <stdio.h>

int main(void) {
    double x, sum = 0.0;
    int count = 0;

    printf("Enter numbers; enter -999 to stop:\n");
    while (scanf("%lf", &x) == 1) {
        if (x == -999.0) {
            break;
        }
        count++;
        sum += x;
        printf("    running avg = %.4f\n", sum / count);
    }
    if (count > 0) {
        printf("Final avg = %.4f\n", sum / count);
    }
    return 0;
}

```

```

gcc -std=c17 -Wall -Wextra -o running_avg_sentinel running_avg_sentinel.c
printf '2 4 6 -999\n' | ./running_avg_sentinel

```

Performance Notes

```

#include <stdio.h>

int main(void) {
    int array[1000];
    const int n = (int)(sizeof array / sizeof array[0]);

    /* Hoist invariant size (here already a constant) */
    for (int i = 0; i < n; i++) {
        array[i] = i;
    }

    long sum = 0;
    for (int i = 0; i < n; i++) {
        sum += array[i];
    }
}

```

```

    printf("sum=%ld\n", sum);
    return 0;
}

```

Avoid `float`/`double` as the primary loop counter (precision can skip or overshoot the end test). Use integer counters and derive floats inside the body.

Best Practices

1. **Initialize** control variables clearly; prefer `for (int i = 0; ...)` (C99+).
2. **Name** indices meaningfully in complex loops (`row`, `col`, `student_i`).
3. Prefer **`break`** / **flags** over mutating the loop index mid-body.
4. Match **loop kind** to the problem (see decision guide).
5. Always ensure a **progress condition** (or intentional infinite loop with external exit).

Common Pitfalls

Infinite loop by forgetting the update

```

int i = 0;
while (i < 10) {
    printf("%d\n", i);
    /* missing i++ */
}

```

Off-by-one

```

int a[5];
for (int i = 0; i <= 5; i++) { /* BAD: i==5 is out of bounds */
    a[i] = 0;
}
/* Correct: i < 5 or i < n */

```

Floating loop variable

```
/* Fragile */
for (float x = 0.0f; x != 1.0f; x += 0.1f) { }

/* Robust */
for (int i = 0; i < 10; i++) {
    float x = (float)i * 0.1f;
}
```

More Practical Examples

Number Patterns

```
#include <stdio.h>

int main(void) {
    int rows = 5;

    printf("Right triangle:\n");
    for (int i = 1; i <= rows; i++) {
        for (int j = 1; j <= i; j++) {
            printf("* ");
        }
        printf("\n");
    }

    printf("\nPyramid:\n");
    for (int i = 1; i <= rows; i++) {
        for (int j = 1; j <= rows - i; j++) {
            printf(" ");
        }
        for (int k = 1; k <= 2 * i - 1; k++) {
            printf("*");
        }
        printf("\n");
    }
    return 0;
}
```



```
gcc -std=c17 -Wall -Wextra -o patterns patterns.c
./patterns
```

Array Stats and Reverse

```
#include <stdio.h>

int main(void) {
    int numbers[] = {12, 7, 23, 8, 15, 3, 19, 11, 5, 17};
    int size = (int)(sizeof numbers / sizeof numbers[0]);

    int max = numbers[0];
    for (int i = 1; i < size; i++) {
        if (numbers[i] > max) {
            max = numbers[i];
        }
    }
    printf("Max: %d\n", max);

    long sum = 0;
    for (int i = 0; i < size; i++) {
        sum += numbers[i];
    }
    printf("Average: %.2f\n", (double)sum / size);

    for (int i = 0, j = size - 1; i < j; i++, j--) {
        int tmp = numbers[i];
        numbers[i] = numbers[j];
        numbers[j] = tmp;
    }

    printf("Reversed: ");
    for (int i = 0; i < size; i++) {
        printf("%d ", numbers[i]);
    }
    printf("\n");
    return 0;
}
```

```
gcc -std=c17 -Wall -Wextra -o array_stats array_stats.c
./array_stats
```


Exercises

Exercise 1 — Collatz steps

Read a positive integer n . While $n \neq 1$, if even set $n = n/2$, else $n = 3n+1$. Count steps. Print the sequence.

Sketch: `while (n != 1) { print; if (n%2==0) n/=2; else n=3*n+1; steps++; }` Guard overflow with `long`.

```
gcc -std=c17 -Wall -Wextra -o collatz collatz.c
```

Exercise 2 — Nested loop: prime pairs

Find all pairs (p, q) with $2 \leq p < q \leq 50$ where both are prime and $p + q == 50$ (Goldbach-style for 50).

Sketch: outer p , inner not needed if $q = 50-p$ and $q \geq p$; reuse `is_prime`.

Exercise 3 — Choose the right loop

Implement three programs that read a password until correct:

1. `while` — re-prompt only if first read failed validation
2. `do-while` — always prompt at least once
3. `for` — at most 5 attempts

Document which feels most natural and why.

Exercise 4 — FizzBuzz customization

Accept n , a , b from the command line or stdin. Print $1..n$ with Fizz for multiples of a , Buzz for b , FizzBuzz for both.

Exercise 5 — Server-style loop

Write a `for (;;) loop` that reads lines; commands `inc`, `dec`, `show`, `quit`. Maintain a counter. Reject unknown commands without exiting. Exit only on `quit` or EOF.

Exercise 6 — Running variance (optional stretch)

Extend running average to track sum of squares and print population variance each step:

`var = (sumsq / n) - mean2` (careful with floating error; fine for a learning demo).

Summary

1. **while** — pre-test; zero or more iterations
2. **do-while** — post-test; at least once (menus)
3. **for** — counted ranges and clear index updates
4. **Decision guide** — pick the loop that matches the story
5. **break / continue** — innermost loop only
6. **Nested loops** — tables, patterns, pairs
7. **Infinite loops** — servers/superloops with explicit exit conditions
8. **Worked examples** — primes, FizzBuzz, mult table, running average
9. **Pitfalls** — forgotten updates, off-by-one, float counters

Next: advanced flow control (`goto` rarity, multi-level exit patterns, structured alternatives).

Advanced Control Flow

Introduction

Advanced control flow mechanisms in C go beyond basic conditionals and loops to provide more sophisticated ways to manage program execution. This chapter explores the `goto` statement, function calls as control flow mechanisms, error handling patterns, and advanced techniques for structuring complex programs.

While some of these mechanisms are powerful, they must be used judiciously to maintain code readability and maintainability. Understanding when and how to use these advanced features is crucial for writing professional C code.

The goto Statement

The `goto` statement allows unconditional jumps to labeled statements within the same function. While generally discouraged in modern programming, there are specific scenarios where `goto` can improve code clarity and efficiency.

Basic Syntax

```
goto label;  
...  
label: statement;
```

Simple goto Example

```
#include <stdio.h>  
  
int main() {  
    int choice;  
  
    printf("Enter 1 to continue, 0 to exit: ");  
    scanf("%d", &choice);  
  
    if (choice == 0) {  
        goto end; // Jump to the 'end' label  
    }  
}
```

```

    }

    printf("Continuing with the program...\n");
    // More program logic here

end:
    printf("Program ended.\n");
    return 0;
}

```

goto for Error Handling

One of the most legitimate uses of `goto` is in error handling, particularly when cleaning up resources:

```

#include <stdio.h>
#include <stdlib.h>

int main() {
    FILE *file1 = NULL;
    FILE *file2 = NULL;
    int *array = NULL;

    // Allocate resources
    file1 = fopen("file1.txt", "r");
    if (file1 == NULL) {
        printf("Error opening file1.txt\n");
        goto cleanup;
    }

    file2 = fopen("file2.txt", "w");
    if (file2 == NULL) {
        printf("Error opening file2.txt\n");
        goto cleanup;
    }

    array = (int*)malloc(100 * sizeof(int));
    if (array == NULL) {
        printf("Error allocating memory\n");
        goto cleanup;
    }

    // Main program logic here
    printf("All resources allocated successfully!\n");

    // Normal cleanup (reverse order of allocation)

```

```

    free(array);
    fclose(file2);
    fclose(file1);
    return 0;

cleanup:
    // Error cleanup (free allocated resources)
    if (array != NULL) {
        free(array);
    }
    if (file2 != NULL) {
        fclose(file2);
    }
    if (file1 != NULL) {
        fclose(file1);
    }
    return 1; // Error exit code
}

```

goto in Switch Statements

goto can be used to jump between cases in switch statements:

```

#include <stdio.h>

int main() {
    int state = 1;

    switch (state) {
        case 1:
            printf("State 1\n");
            goto state_3; // Jump to case 3
        case 2:
            printf("State 2\n");
            break;
        case 3:
        state_3:
            printf("State 3\n");
            break;
        default:
            printf("Unknown state\n");
    }

    return 0;
}

```

Function Calls as Control Flow

Function calls are a fundamental control flow mechanism that transfer execution to another part of the program. They enable code reuse, modularity, and structured programming.

Basic Function Calls

```
#include <stdio.h>

void print_hello() {
    printf("Hello from function!\n");
}

int add(int a, int b) {
    return a + b;
}

int main() {
    print_hello(); // Function call transfers control

    int result = add(5, 3); // Function call with return value
    printf("5 + 3 = %d\n", result);

    return 0;
}
```

Recursive Function Calls

Recursive functions call themselves to solve problems that can be broken down into smaller, similar subproblems:

```
#include <stdio.h>

// Factorial calculation using recursion
int factorial(int n) {
    if (n <= 1) {
        return 1; // Base case
    }
    return n * factorial(n - 1); // Recursive call
}

// Fibonacci sequence using recursion
int fibonacci(int n) {
    if (n <= 1) {
        return n; // Base cases
    }
}
```

```

    }
    return fibonacci(n - 1) + fibonacci(n - 2); // Recursive calls
}

int main() {
    int num = 5;

    printf("Factorial of %d: %d\n", num, factorial(num));
    printf("Fibonacci of %d: %d\n", num, fibonacci(num));

    return 0;
}

```

Function Pointers for Dynamic Control Flow

Function pointers allow dynamic selection of which function to call at runtime:

```

#include <stdio.h>

int add(int a, int b) {
    return a + b;
}

int subtract(int a, int b) {
    return a - b;
}

int multiply(int a, int b) {
    return a * b;
}

int divide(int a, int b) {
    if (b != 0) {
        return a / b;
    }
    return 0;
}

int main() {
    int x = 10, y = 5;
    int (*operation)(int, int); // Function pointer

    // Dynamic function selection
    char op;
    printf("Enter operation (+, -, *, /): ");
    scanf(" %c", &op);
}

```

```

switch (op) {
    case '+':
        operation = add;
        break;
    case '-':
        operation = subtract;
        break;
    case '*':
        operation = multiply;
        break;
    case '/':
        operation = divide;
        break;
    default:
        printf("Invalid operation\n");
        return 1;
}

int result = operation(x, y); // Call selected function
printf("Result: %d\n", result);

return 0;
}

```

Error Handling Patterns

C doesn't have built-in exception handling, so error handling is typically done through return codes and careful checking.

Return Code Pattern

```

#include <stdio.h>
#include <stdlib.h>

#define SUCCESS 0
#define ERROR_INVALID_INPUT 1
#define ERROR_MEMORY_ALLOCATION 2
#define ERROR_FILE_OPERATION 3

int process_data(int *data, int size) {
    if (data == NULL || size <= 0) {
        return ERROR_INVALID_INPUT;
    }
}

```



```

    // Process data
    for (int i = 0; i < size; i++) {
        data[i] *= 2;
    }

    return SUCCESS;
}

int main() {
    int data[] = {1, 2, 3, 4, 5};
    int size = sizeof(data) / sizeof(data[0]);

    int result = process_data(data, size);

    switch (result) {
        case SUCCESS:
            printf("Data processed successfully!\n");
            for (int i = 0; i < size; i++) {
                printf("%d ", data[i]);
            }
            printf("\n");
            break;
        case ERROR_INVALID_INPUT:
            printf("Error: Invalid input data\n");
            break;
        case ERROR_MEMORY_ALLOCATION:
            printf("Error: Memory allocation failed\n");
            break;
        case ERROR_FILE_OPERATION:
            printf("Error: File operation failed\n");
            break;
        default:
            printf("Unknown error occurred\n");
    }

    return result;
}

```

errno and perror for System Errors

```

#include <stdio.h>
#include <errno.h>
#include <string.h>

int main() {

```

```

FILE *file = fopen("nonexistent_file.txt", "r");

if (file == NULL) {
    printf("Error opening file: %s\n", strerror(errno));
    perror("fopen"); // Alternative way to print error
    return errno;
}

fclose(file);
return 0;
}

```

Setjmp and Longjmp for Non-Local Jumps

The `setjmp` and `longjmp` functions provide a way to perform non-local jumps, similar to exceptions in other languages:

```

#include <stdio.h>
#include <setjmp.h>
#include <stdlib.h>

jmp_buf jump_buffer;

void risky_function() {
    printf("Entering risky function\n");

    // Simulate an error condition
    int error_condition = 1;
    if (error_condition) {
        printf("Error occurred, jumping back!\n");
        longjmp(jump_buffer, 1); // Jump back to setjmp
    }

    printf("This won't be printed\n");
}

int main() {
    int result = setjmp(jump_buffer);

    if (result == 0) {
        printf("First time through setjmp\n");
        risky_function();
    } else {
        printf("Returned from longjmp with value: %d\n", result);
    }
}

```

```
    return 0;
}
```

Advanced Loop Control Techniques

Loop Labels and Break to Outer Loops

While C doesn't support loop labels like some languages, you can achieve similar functionality using flags or goto:

```
#include <stdio.h>

int main() {
    int found = 0;

    // Search for a value in a 2D array
    int matrix[3][3] = {{1, 2, 3}, {4, 5, 6}, {7, 8, 9}};
    int target = 5;

    for (int i = 0; i < 3 && !found; i++) {
        for (int j = 0; j < 3; j++) {
            if (matrix[i][j] == target) {
                printf("Found %d at position (%d, %d)\n", target, i, j);
                found = 1;
                break; // Breaks inner loop
            }
        }
        // Outer loop will exit due to !found condition
    }

    return 0;
}
```

State Machines

State machines are a powerful pattern for managing complex control flow:

```
#include <stdio.h>

typedef enum {
    STATE_IDLE,
    STATE_PROCESSING,
    STATE_COMPLETE,
    STATE_ERROR
}
```

```

} State;

int main() {
    State current_state = STATE_IDLE;
    int input;

    while (current_state != STATE_COMPLETE && current_state != STATE_ERROR) {
        switch (current_state) {
            case STATE_IDLE:
                printf("Enter command (1=start, 2=quit): ");
                scanf("%d", &input);
                if (input == 1) {
                    current_state = STATE_PROCESSING;
                    printf("Processing started...\n");
                } else if (input == 2) {
                    current_state = STATE_COMPLETE;
                }
                break;

            case STATE_PROCESSING:
                printf("Enter data (0=complete, -1=error): ");
                scanf("%d", &input);
                if (input == 0) {
                    current_state = STATE_COMPLETE;
                    printf("Processing completed!\n");
                } else if (input == -1) {
                    current_state = STATE_ERROR;
                    printf("Processing error!\n");
                }
                break;

            case STATE_COMPLETE:
                printf("Process complete.\n");
                break;

            case STATE_ERROR:
                printf("Process error.\n");
                break;
        }
    }

    return 0;
}

```

Best Practices for Advanced Control Flow

1. Use goto Sparingly and Structured

```
// Good: Structured error handling with goto
int process_files() {
    FILE *file1 = NULL;
    FILE *file2 = NULL;
    int *buffer = NULL;
    int result = -1; // Default error

    file1 = fopen("input.txt", "r");
    if (!file1) goto cleanup;

    file2 = fopen("output.txt", "w");
    if (!file2) goto cleanup;

    buffer = malloc(1024);
    if (!buffer) goto cleanup;

    // Main processing logic
    result = 0; // Success

cleanup:
    if (buffer) free(buffer);
    if (file2) fclose(file2);
    if (file1) fclose(file1);
    return result;
}
```

2. Avoid Deep Nesting

```
// Good: Early returns to reduce nesting
int validate_input(int value) {
    if (value < 0) {
        return -1; // Invalid
    }

    if (value > 100) {
        return -1; // Invalid
    }

    if (value % 2 != 0) {
        return -1; // Invalid
    }
}
```

```

        return 0; // Valid
    }

    // Avoid: Deep nesting
    int validate_input_nested(int value) {
        if (value >= 0) {
            if (value <= 100) {
                if (value % 2 == 0) {
                    return 0; // Valid
                } else {
                    return -1; // Invalid
                }
            } else {
                return -1; // Invalid
            }
        } else {
            return -1; // Invalid
        }
    }
}

```

3. Use Function Pointers for Strategy Pattern

```

#include <stdio.h>

typedef int (*operation_func)(int, int);

int add(int a, int b) { return a + b; }
int subtract(int a, int b) { return a - b; }
int multiply(int a, int b) { return a * b; }

int calculate(int a, int b, operation_func op) {
    return op(a, b);
}

int main() {
    int x = 10, y = 5;

    printf("Add: %d\n", calculate(x, y, add));
    printf("Subtract: %d\n", calculate(x, y, subtract));
    printf("Multiply: %d\n", calculate(x, y, multiply));

    return 0;
}

```

Common Pitfalls and How to Avoid Them

1. goto Leading to Spaghetti Code

```
#include <stdio.h>

int main() {
    // Wrong: Complex goto usage leading to unreadable code
    int i = 0;
    goto start;

loop:
    printf("In loop: %d\n", i);
    i++;
    if (i < 5) goto loop;
    goto end;

start:
    printf("Starting...\n");
    goto loop;

end:
    printf("Ending...\n");

    // Correct: Use proper loop constructs
    printf("Starting...\n");
    for (int j = 0; j < 5; j++) {
        printf("In loop: %d\n", j);
    }
    printf("Ending...\n");

    return 0;
}
```

2. Infinite Recursion

```
#include <stdio.h>

// Wrong: No base case
int factorial_wrong(int n) {
    return n * factorial_wrong(n - 1); // Infinite recursion
}

// Correct: With proper base case
int factorial_correct(int n) {
```

```

        if (n <= 1) return 1; // Base case
        return n * factorial_correct(n - 1);
    }

    int main() {
        printf("Factorial of 5: %d\n", factorial_correct(5));
        return 0;
    }

```

3. Longjmp Across Function Boundaries

```

#include <stdio.h>
#include <setjmp.h>

jmp_buf env;

void function_a() {
    printf("In function A\n");
    longjmp(env, 1); // This is safe within the same function context
}

void function_b() {
    function_a(); // This is also safe
}

int main() {
    if (setjmp(env) == 0) {
        printf("Setting jump point\n");
        function_b();
    } else {
        printf("Jumped back!\n");
    }

    return 0;
}

```


Practical Examples

Resource Management with goto

```
#include <stdio.h>
#include <stdlib.h>
#include <string.h>

typedef struct {
    char *name;
    int *data;
    FILE *file;
} ResourceContainer;

int initialize_resources(ResourceContainer *container) {
    // Allocate name
    container->name = malloc(50);
    if (!container->name) {
        goto cleanup;
    }
    strcpy(container->name, "Resource");

    // Allocate data
    container->data = malloc(100 * sizeof(int));
    if (!container->data) {
        goto cleanup;
    }

    // Open file
    container->file = fopen("data.txt", "w");
    if (!container->file) {
        goto cleanup;
    }

    return 0; // Success
}

cleanup:
// Clean up in reverse order
if (container->file) {
    fclose(container->file);
}
if (container->data) {
    free(container->data);
}
if (container->name) {
    free(container->name);
}
```

```

    }
    return -1; // Error
}

void cleanup_resources(ResourceContainer *container) {
    if (container->file) {
        fclose(container->file);
    }
    if (container->data) {
        free(container->data);
    }
    if (container->name) {
        free(container->name);
    }
    memset(container, 0, sizeof(ResourceContainer));
}

int main() {
    ResourceContainer container = {0};

    if (initialize_resources(&container) == 0) {
        printf("Resources initialized successfully\n");
        printf("Name: %s\n", container.name);
        // Use resources...
        cleanup_resources(&container);
    } else {
        printf("Failed to initialize resources\n");
        return 1;
    }

    return 0;
}

```

Command-Line Parser with State Machine

```

#include <stdio.h>
#include <string.h>

typedef enum {
    STATE_START,
    STATE_OPTION,
    STATE_ARGUMENT,
    STATE_DONE,
    STATE_ERROR
} ParseState;

```

```

typedef struct {
    int verbose;
    char *output_file;
    int count;
} Options;

int parse_arguments(int argc, char *argv[], Options *options) {
    ParseState state = STATE_START;

    for (int i = 1; i < argc && state != STATE_ERROR; i++) {
        switch (state) {
            case STATE_START:
                if (argv[i][0] == '-') {
                    state = STATE_OPTION;
                    i--; // Re-process this argument
                } else {
                    state = STATE_DONE;
                }
                break;

            case STATE_OPTION:
                if (strcmp(argv[i], "-v") == 0) {
                    options->verbose = 1;
                } else if (strcmp(argv[i], "-o") == 0) {
                    state = STATE_ARGUMENT;
                } else if (strcmp(argv[i], "-c") == 0) {
                    state = STATE_ARGUMENT;
                } else {
                    printf("Unknown option: %s\n", argv[i]);
                    state = STATE_ERROR;
                }
                break;

            case STATE_ARGUMENT:
                if (strcmp(argv[i-1], "-o") == 0) {
                    options->output_file = argv[i];
                } else if (strcmp(argv[i-1], "-c") == 0) {
                    options->count = atoi(argv[i]);
                }
                state = STATE_START;
                break;

            case STATE_DONE:
            case STATE_ERROR:
                break;
        }
    }
}

```

```

    }

    return (state == STATE_DONE || state == STATE_START) ? 0 : -1;
}

int main(int argc, char *argv[]) {
    Options options = {0};

    if (parse_arguments(argc, argv, &options) == 0) {
        printf("Parsing successful:\n");
        printf("  Verbose: %s\n", options.verbose ? "yes" : "no");
        printf("  Output file: %s\n", options.output_file ? options.output_file : "not speci");
        printf("  Count: %d\n", options.count);
    } else {
        printf("Parsing failed\n");
        return 1;
    }

    return 0;
}

```

Complex Error Handling System

```

#include <stdio.h>
#include <stdlib.h>
#include <errno.h>
#include <string.h>

typedef enum {
    ERR_SUCCESS = 0,
    ERR_MEMORY,
    ERR_FILE_OPEN,
    ERR_FILE_READ,
    ERR_INVALID_INPUT,
    ERR_UNKNOWN
} ErrorCode;

const char* error_string(ErrorCode code) {
    switch (code) {
        case ERR_SUCCESS: return "Success";
        case ERR_MEMORY: return "Memory allocation failed";
        case ERR_FILE_OPEN: return "File open failed";
        case ERR_FILE_READ: return "File read failed";
        case ERR_INVALID_INPUT: return "Invalid input";
        case ERR_UNKNOWN: return "Unknown error";
    }
}

```

```

        default: return "Unrecognized error";
    }
}

ErrorCode process_file(const char *filename) {
    FILE *file = NULL;
    char *buffer = NULL;
    ErrorCode result = ERR_SUCCESS;

    // Open file
    file = fopen(filename, "r");
    if (!file) {
        result = ERR_FILE_OPEN;
        goto cleanup;
    }

    // Allocate buffer
    buffer = malloc(1024);
    if (!buffer) {
        result = ERR_MEMORY;
        goto cleanup;
    }

    // Read file
    if (fread(buffer, 1, 1024, file) == 0) {
        if (feof(file)) {
            printf("File is empty\n");
        } else {
            result = ERR_FILE_READ;
            goto cleanup;
        }
    }

    printf("File processed successfully\n");

cleanup:
    if (buffer) free(buffer);
    if (file) fclose(file);
    return result;
}

int main() {
    const char *filename = "test.txt";
    ErrorCode result = process_file(filename);

    if (result == ERR_SUCCESS) {
        printf("Operation completed successfully\n");
    }
}

```

```

    } else {
        printf("Error: %s\n", error_string(result));
        if (result == ERR_FILE_OPEN) {
            printf("System error: %s\n", strerror(errno));
        }
    }

    return result;
}

```

Summary

In this chapter, you've learned about:

1. **goto Statement:** When and how to use unconditional jumps
2. **Function Calls:** Basic and recursive function calls as control flow
3. **Function Pointers:** Dynamic function selection and callback mechanisms
4. **Error Handling:** Return codes, `errno`, and advanced error management
5. **setjmp/longjmp:** Non-local jumps for exceptional control flow
6. **Advanced Loop Control:** Techniques for managing complex iterations
7. **State Machines:** Pattern for managing complex program states
8. **Best Practices:** Writing clean, maintainable advanced control flow
9. **Common Pitfalls:** Avoiding typical mistakes with advanced control structures

Advanced control flow mechanisms provide powerful tools for managing complex program execution, but they must be used carefully to maintain code quality. With these techniques, you can implement sophisticated control logic while maintaining readability and reliability.

Module 3: Control Flow

Overview

This module explores the control structures that govern program execution flow in C. You'll learn how to implement decision-making with conditionals, repetition with loops, and advanced control mechanisms. Mastering control flow is essential for creating programs that respond dynamically to inputs and execute complex logic.

Learning Objectives

By the end of this module, you will be able to: - Implement decision-making using if, if-else, and switch statements - Create repetitive code execution using while, do-while, and for loops - Apply loop control statements (break, continue) effectively - Use advanced control flow mechanisms including goto and function calls - Implement error handling patterns and state machines - Write efficient and maintainable control flow code - Avoid common pitfalls in control structure implementation

Chapters

1. **Conditional Statements** - if, if-else, switch, and nested conditionals
2. **Loops** - while, do-while, for loops and nested iterations
3. **Advanced Control Flow** - goto, function pointers, error handling, and state machines

Key Concepts Covered

- Conditional execution with if, if-else, and switch statements
- Loop constructs: while, do-while, and for loops
- Loop control with break and continue statements
- Nested loops and complex iteration patterns
- Infinite loops and loop performance considerations
- The goto statement and when to use it appropriately
- Function calls as control flow mechanisms
- Error handling patterns and return codes
- State machines for complex program logic
- Advanced techniques like setjmp/longjmp
- Best practices for writing clean control flow code

Prerequisites

- Completion of Module 1: Foundations and Environment Setup
- Completion of Module 2: Data Types and Variables
- Basic understanding of operators and expressions
- Familiarity with variable declaration and initialization

Tools and Technologies

- C compiler (GCC, Clang, or MSVC)
- Text editor or IDE
- Understanding of basic programming concepts
- Familiarity with debugging tools

Estimated Time to Complete

- Reading: 6-8 hours
- Exercises: 8-12 hours
- Projects: 4-6 hours

Assessment

- Complete all chapter exercises
- Successfully implement various control flow patterns
- Demonstrate understanding of loop optimization techniques
- Pass control flow programming quiz
- Submit a comprehensive program using multiple control structures

Next Module

After completing this module, proceed to **Module 4: Functions and Modular Programming** to learn about code organization, function design, and modular programming techniques.

Additional Resources

- “The C Programming Language” by Kernighan and Ritchie (Chapter 3)
- C standard documentation (ISO/IEC 9899:2024 for C23)
- Online tutorials on control flow patterns
- Debugging tools documentation (GDB, Valgrind)
- Best practices guides for structured programming

Functions

Function Fundamentals

Introduction

Functions are fundamental building blocks in C programming that allow you to break down complex problems into smaller, manageable pieces. They promote code reusability, modularity, and maintainability by encapsulating specific tasks into named, callable units.

A function in C is a self-contained block of code that performs a specific task and can optionally return a value to the caller. Functions help organize code logically, reduce redundancy, and make programs easier to understand and debug.

This chapter will explore the fundamentals of functions in C, including their definition, declaration, calling mechanisms, and parameter passing.

What is a Function?

A function is a named block of code that performs a specific task and can be called (invoked) from other parts of the program. Functions have several key characteristics:

1. **Name:** A unique identifier used to call the function
2. **Parameters:** Input values passed to the function
3. **Return Type:** The type of value the function returns (or void if none)
4. **Body:** The code that executes when the function is called

Function Definition

A function definition specifies what the function does and how it does it. It includes the function's return type, name, parameters, and body.

Basic Syntax

```
return_type function_name(parameter_list) {  
    // Function body  
    // Local variable declarations  
    // Statements  
    // Return statement (if not void)  
}
```

Simple Function Example

```
#include <stdio.h>

// Function definition
int add(int a, int b) {
    int sum = a + b;
    return sum;
}

int main() {
    int result = add(5, 3); // Function call
    printf("5 + 3 = %d\n", result);
    return 0;
}
```

Function with No Parameters

```
#include <stdio.h>

void print_greeting() {
    printf("Hello, welcome to C programming!\n");
}

int main() {
    print_greeting(); // Function call with no arguments
    return 0;
}
```

Function with No Return Value (void)

```
#include <stdio.h>

void print_line() {
    for (int i = 0; i < 20; i++) {
        printf("-");
    }
    printf("\n");
}

int main() {
    print_line();
    printf("Main Content\n");
    print_line();
}
```

```
    return 0;
}
```

Function Declaration (Prototypes)

A function declaration (also called a prototype) tells the compiler about a function's name, return type, and parameters without providing the implementation. This allows you to call a function before its definition appears in the code.

Syntax

```
return_type function_name(parameter_types);
```

Function Prototype Example

```
#include <stdio.h>

// Function prototypes
int multiply(int x, int y);
void print_result(int value);
double calculate_average(double a, double b);

int main() {
    int product = multiply(4, 5);
    print_result(product);

    double avg = calculate_average(10.5, 20.3);
    printf("Average: %.2f\n", avg);

    return 0;
}

// Function definitions (can appear after main)
int multiply(int x, int y) {
    return x * y;
}

void print_result(int value) {
    printf("Result: %d\n", value);
}

double calculate_average(double a, double b) {
```

```
    return (a + b) / 2.0;
}
```

Function Parameters

Function parameters are variables that receive values when the function is called. C uses “pass by value” for parameter passing, meaning that copies of the actual arguments are passed to the function.

Parameter Passing by Value

```
#include <stdio.h>

void modify_value(int x) {
    x = 100; // Only modifies the local copy
    printf("Inside function: x = %d\n", x);
}

int main() {
    int num = 5;
    printf("Before function call: num = %d\n", num);
    modify_value(num);
    printf("After function call: num = %d\n", num); // Unchanged
    return 0;
}
```

Multiple Parameters

```
#include <stdio.h>

int find_maximum(int a, int b, int c) {
    int max = a;
    if (b > max) max = b;
    if (c > max) max = c;
    return max;
}

void print_student_info(char name[], int age, float gpa) {
    printf("Student: %s\n", name);
    printf("Age: %d\n", age);
    printf("GPA: %.2f\n", gpa);
}
```

```

int main() {
    int max_value = find_maximum(10, 25, 15);
    printf("Maximum value: %d\n", max_value);

    print_student_info("John Doe", 20, 3.75);

    return 0;
}

```

Return Statements

The `return` statement terminates function execution and optionally returns a value to the caller. For functions with a return type other than `void`, a return statement with a value is required.

Returning Values

```

#include <stdio.h>

int square(int n) {
    return n * n;
}

float calculate_area(float radius) {
    const float PI = 3.14159f;
    return PI * radius * radius;
}

char get_grade(int score) {
    if (score >= 90) return 'A';
    else if (score >= 80) return 'B';
    else if (score >= 70) return 'C';
    else if (score >= 60) return 'D';
    else return 'F';
}

int main() {
    printf("Square of 5: %d\n", square(5));
    printf("Area of circle (radius 3): %.2f\n", calculate_area(3.0f));
    printf("Grade for score 85: %c\n", get_grade(85));

    return 0;
}

```

Early Returns

```
#include <stdio.h>

int absolute_value(int n) {
    if (n < 0) {
        return -n; // Early return for negative numbers
    }
    return n; // Return positive numbers as-is
}

int find_factorial(int n) {
    if (n < 0) {
        return -1; // Error case
    }
    if (n == 0 || n == 1) {
        return 1; // Base cases
    }

    int result = 1;
    for (int i = 2; i <= n; i++) {
        result *= i;
    }
    return result;
}

int main() {
    printf("Absolute value of -5: %d\n", absolute_value(-5));
    printf("Absolute value of 5: %d\n", absolute_value(5));

    printf("Factorial of 5: %d\n", find_factorial(5));
    printf("Factorial of -3: %d\n", find_factorial(-3)); // Error case

    return 0;
}
```

Function Scope and Lifetime

Variables declared inside a function have local scope and exist only during the function's execution.

Local Variables

```
#include <stdio.h>

void function_a() {
    int local_var = 10; // Local to function_a
    printf("In function_a: local_var = %d\n", local_var);
}

void function_b() {
    int local_var = 20; // Different variable, local to function_b
    printf("In function_b: local_var = %d\n", local_var);
}

int main() {
    function_a();
    function_b();

    // printf("%d\n", local_var); // Error: local_var not accessible here

    return 0;
}
```

Static Local Variables

Static local variables retain their values between function calls:

```
#include <stdio.h>

void counter() {
    static int count = 0; // Initialized only once
    count++;
    printf("Function called %d time(s)\n", count);
}

int main() {
    counter(); // Output: Function called 1 time(s)
    counter(); // Output: Function called 2 time(s)
    counter(); // Output: Function called 3 time(s)

    return 0;
}
```


Function Naming Conventions

Following consistent naming conventions improves code readability and maintainability:

Common Naming Styles

```
#include <stdio.h>

// Snake case (recommended for C)
int calculate_area_rectangle(int width, int height) {
    return width * height;
}

// Camel case (also acceptable)
int calculateAreaRectangle(int width, int height) {
    return width * height;
}

// Descriptive names
void print_student_information(char student_name[], int student_age) {
    printf("Name: %s, Age: %d\n", student_name, student_age);
}

int main() {
    int area1 = calculate_area_rectangle(5, 10);
    int area2 = calculateAreaRectangle(3, 7);

    printf("Area 1: %d\n", area1);
    printf("Area 2: %d\n", area2);

    print_student_information("Alice Smith", 22);

    return 0;
}
```

Best Practices for Functions

1. Single Responsibility Principle

```
#include <stdio.h>

// Good: Each function has a single, clear purpose
int calculate_rectangle_area(int width, int height) {
```

```

    return width * height;
}

int calculate_rectangle_perimeter(int width, int height) {
    return 2 * (width + height);
}

void print_rectangle_info(int width, int height) {
    printf("Rectangle: %d x %d\n", width, height);
    printf("Area: %d\n", calculate_rectangle_area(width, height));
    printf("Perimeter: %d\n", calculate_rectangle_perimeter(width, height));
}

// Avoid: Function doing multiple unrelated things
/*
int calculate_and_print_rectangle(int width, int height) {
    int area = width * height;
    printf("Area: %d\n", area);
    return area;
}
*/

int main() {
    print_rectangle_info(5, 3);
    return 0;
}

```

2. Appropriate Function Size

```

#include <stdio.h>

// Good: Functions are small and focused
int is_even(int number) {
    return number % 2 == 0;
}

int is_positive(int number) {
    return number > 0;
}

int is_valid_input(int number) {
    return is_positive(number) && is_even(number);
}

// Avoid: Functions that are too long

```

```

/*
int process_complex_logic(int a, int b, int c, int d, int e) {
    // 100+ lines of code doing many different things
}
*/

int main() {
    int test_number = 10;
    if (is_valid_input(test_number)) {
        printf("%d is valid input\n", test_number);
    }

    return 0;
}

```

3. Meaningful Parameter Names

```

#include <stdio.h>

// Good: Descriptive parameter names
double calculate_triangle_area(double base, double height) {
    return 0.5 * base * height;
}

// Avoid: Generic parameter names
/*
double calculate_triangle_area(double b, double h) {
    return 0.5 * b * h;
}
*/

int main() {
    double area = calculate_triangle_area(10.0, 5.0);
    printf("Triangle area: %.2f\n", area);

    return 0;
}

```

Common Pitfalls and How to Avoid Them

1. Missing Return Statements

```
#include <stdio.h>

// Wrong: Missing return statement in non-void function
/*
int calculate_square(int n) {
    int result = n * n;
    // Missing return statement
}
*/

// Correct: Always include return statement
int calculate_square(int n) {
    int result = n * n;
    return result;
}

int main() {
    int square = calculate_square(5);
    printf("Square of 5: %d\n", square);

    return 0;
}
```

2. Using Uninitialized Variables

```
#include <stdio.h>

// Wrong: Using uninitialized local variable
/*
int add_numbers(int a, int b) {
    int sum; // Not initialized
    sum = sum + a + b; // Using uninitialized variable
    return sum;
}
*/

// Correct: Initialize variables properly
int add_numbers(int a, int b) {
    int sum = 0; // Properly initialized
    sum = sum + a + b;
    return sum;
}
```

```

}

int main() {
    int result = add_numbers(3, 4);
    printf("3 + 4 = %d\n", result);

    return 0;
}

```

3. Function Declaration vs Definition Mismatch

```

#include <stdio.h>

// Function prototype
int multiply(int x, int y);

int main() {
    int result = multiply(5, 3);
    printf("5 * 3 = %d\n", result);
    return 0;
}

// Correct: Definition matches prototype
int multiply(int x, int y) {
    return x * y;
}

// Wrong: Definition doesn't match prototype
/*
int multiply(int a, int b) { // Different parameter names are OK
    return a * b;
}
// But different types or number of parameters would be wrong
*/

```

Practical Examples

Temperature Converter Functions

```

#include <stdio.h>

// Function prototypes
double celsius_to_fahrenheit(double celsius);

```

```

double fahrenheit_to_celsius(double fahrenheit);
void print_temperature_conversion(double celsius, double fahrenheit);

// Convert Celsius to Fahrenheit
double celsius_to_fahrenheit(double celsius) {
    return (celsius * 9.0 / 5.0) + 32.0;
}

// Convert Fahrenheit to Celsius
double fahrenheit_to_celsius(double fahrenheit) {
    return (fahrenheit - 32.0) * 5.0 / 9.0;
}

// Print temperature conversion
void print_temperature_conversion(double celsius, double fahrenheit) {
    printf("%.2f°C = %.2f°F\n", celsius, fahrenheit);
    printf("%.2f°F = %.2f°C\n", fahrenheit, celsius_to_fahrenheit(fahrenheit));
}

int main() {
    double c_temp = 25.0;
    double f_temp = 77.0;

    double c_to_f = celsius_to_fahrenheit(c_temp);
    double f_to_c = fahrenheit_to_celsius(f_temp);

    printf("Temperature Conversions:\n");
    print_temperature_conversion(c_temp, f_temp);

    return 0;
}

```

Mathematical Utility Functions

```

#include <stdio.h>

// Function prototypes
int power(int base, int exponent);
int factorial(int n);
int gcd(int a, int b);
int lcm(int a, int b);

// Calculate power (base^exponent)
int power(int base, int exponent) {
    if (exponent < 0) return 0; // Handle negative exponents

```

```

        if (exponent == 0) return 1;

        int result = 1;
        for (int i = 0; i < exponent; i++) {
            result *= base;
        }
        return result;
    }

    // Calculate factorial
    int factorial(int n) {
        if (n < 0) return -1; // Error for negative numbers
        if (n == 0 || n == 1) return 1;

        int result = 1;
        for (int i = 2; i <= n; i++) {
            result *= i;
        }
        return result;
    }

    // Calculate Greatest Common Divisor
    int gcd(int a, int b) {
        while (b != 0) {
            int temp = b;
            b = a % b;
            a = temp;
        }
        return a;
    }

    // Calculate Least Common Multiple
    int lcm(int a, int b) {
        return (a * b) / gcd(a, b);
    }

    int main() {
        printf("Mathematical Utilities:\n");
        printf("2^8 = %d\n", power(2, 8));
        printf("5! = %d\n", factorial(5));
        printf("GCD of 48 and 18 = %d\n", gcd(48, 18));
        printf("LCM of 12 and 18 = %d\n", lcm(12, 18));

        return 0;
    }

```

String Processing Functions

```
#include <stdio.h>
#include <string.h>

// Function prototypes
int string_length(char str[]);
void string_copy(char dest[], char src[]);
int string_compare(char str1[], char str2[]);
void string_reverse(char str[]);

// Calculate string length
int string_length(char str[]) {
    int length = 0;
    while (str[length] != '\0') {
        length++;
    }
    return length;
}

// Copy string
void string_copy(char dest[], char src[]) {
    int i = 0;
    while (src[i] != '\0') {
        dest[i] = src[i];
        i++;
    }
    dest[i] = '\0'; // Null terminate
}

// Compare strings
int string_compare(char str1[], char str2[]) {
    int i = 0;
    while (str1[i] != '\0' && str2[i] != '\0') {
        if (str1[i] != str2[i]) {
            return str1[i] - str2[i];
        }
        i++;
    }
    return str1[i] - str2[i];
}

// Reverse string
void string_reverse(char str[]) {
    int length = string_length(str);
    for (int i = 0; i < length / 2; i++) {
```



```

        char temp = str[i];
        str[i] = str[length - 1 - i];
        str[length - 1 - i] = temp;
    }
}

int main() {
    char str1[100] = "Hello";
    char str2[100] = "World";
    char str3[100];

    printf("String Processing Examples:\n");
    printf("Length of '%s': %d\n", str1, string_length(str1));

    string_copy(str3, str1);
    printf("Copied string: %s\n", str3);

    printf("Comparison '%s' vs '%s': %d\n", str1, str2, string_compare(str1, str2));

    printf("Original: %s\n", str1);
    string_reverse(str1);
    printf("Reversed: %s\n", str1);

    return 0;
}

```

Summary

In this chapter, you've learned about:

1. **Function Basics:** What functions are and why they're important
2. **Function Definition:** How to define functions with proper syntax
3. **Function Declaration:** Using prototypes to declare functions before use
4. **Parameter Passing:** How C passes parameters by value
5. **Return Statements:** Using return to send values back to callers
6. **Scope and Lifetime:** Understanding variable scope in functions
7. **Best Practices:** Writing clean, maintainable functions
8. **Common Pitfalls:** Avoiding typical function-related errors

Functions are essential for organizing code into logical, reusable components. In the next chapter, we'll explore advanced function concepts including recursion, variable arguments, and function pointers.

Advanced Functions

Introduction

Building upon the fundamental concepts of functions, this chapter explores advanced function features in C that provide greater flexibility and power. These advanced features include recursive functions, variable argument functions, inline functions, and function pointers. Understanding these concepts is crucial for writing sophisticated C programs and leveraging the full capabilities of the language.

Advanced functions enable you to implement complex algorithms, create flexible interfaces, and build more efficient code. This chapter will dive deep into each of these advanced topics with practical examples and best practices.

Recursive Functions

Recursion is a powerful programming technique where a function calls itself to solve a problem. Recursive functions are particularly useful for problems that can be broken down into smaller, similar subproblems.

Basic Recursion Concepts

A recursive function must have: 1. **Base Case**: A condition that stops the recursion 2. **Recursive Case**: The function calling itself with modified parameters

Simple Recursion Example

```
#include <stdio.h>

// Calculate factorial using recursion
int factorial(int n) {
    // Base case
    if (n <= 1) {
        return 1;
    }

    // Recursive case
    return n * factorial(n - 1);
}
```

```

}

int main() {
    for (int i = 0; i <= 6; i++) {
        printf("%d! = %d\n", i, factorial(i));
    }

    return 0;
}

```

Fibonacci Sequence

```

#include <stdio.h>

// Calculate Fibonacci number using recursion
int fibonacci(int n) {
    // Base cases
    if (n <= 1) {
        return n;
    }

    // Recursive case
    return fibonacci(n - 1) + fibonacci(n - 2);
}

int main() {
    printf("Fibonacci sequence (first 10 numbers):\n");
    for (int i = 0; i < 10; i++) {
        printf("F(%d) = %d\n", i, fibonacci(i));
    }

    return 0;
}

```

Recursive Array Processing

```

#include <stdio.h>

// Calculate sum of array elements using recursion
int array_sum(int arr[], int size) {
    // Base case
    if (size <= 0) {
        return 0;
    }
}

```

```

        // Recursive case
        return arr[size - 1] + array_sum(arr, size - 1);
    }

    // Find maximum element in array using recursion
    int array_max(int arr[], int size) {
        // Base case
        if (size == 1) {
            return arr[0];
        }

        // Recursive case
        int max_of_rest = array_max(arr + 1, size - 1);
        return (arr[0] > max_of_rest) ? arr[0] : max_of_rest;
    }

    int main() {
        int numbers[] = {3, 7, 2, 9, 1, 8, 4};
        int size = sizeof(numbers) / sizeof(numbers[0]);

        printf("Array: ");
        for (int i = 0; i < size; i++) {
            printf("%d ", numbers[i]);
        }
        printf("\n");

        printf("Sum: %d\n", array_sum(numbers, size));
        printf("Maximum: %d\n", array_max(numbers, size));

        return 0;
    }

```

Variable Argument Functions

C provides the `<stdarg.h>` header for creating functions that can accept a variable number of arguments. This is useful for functions like `printf()` that can handle different numbers of parameters.

Basic Variable Argument Functions

```

#include <stdio.h>
#include <stdarg.h>

```

```

// Function that calculates the sum of variable arguments
int sum(int count, ...) {
    va_list args;
    va_start(args, count);

    int total = 0;
    for (int i = 0; i < count; i++) {
        int value = va_arg(args, int);
        total += value;
    }

    va_end(args);
    return total;
}

// Function that finds maximum among variable arguments
int maximum(int count, ...) {
    if (count <= 0) return 0;

    va_list args;
    va_start(args, count);

    int max = va_arg(args, int);
    for (int i = 1; i < count; i++) {
        int value = va_arg(args, int);
        if (value > max) {
            max = value;
        }
    }

    va_end(args);
    return max;
}

int main() {
    printf("Sum of 3 numbers: %d\n", sum(3, 10, 20, 30));
    printf("Sum of 5 numbers: %d\n", sum(5, 1, 2, 3, 4, 5));

    printf("Maximum of 4 numbers: %d\n", maximum(4, 15, 8, 23, 42));
    printf("Maximum of 3 numbers: %d\n", maximum(3, 100, 50, 75));

    return 0;
}

```

Formatted Print Function

```
#include <stdio.h>
#include <stdarg.h>

// Custom printf-like function
void my_printf(const char *format, ...) {
    va_list args;
    va_start(args, format);

    while (*format != '\0') {
        if (*format == '%' && *(format + 1) != '\0') {
            format++;
            switch (*format) {
                case 'd':
                    printf("%d", va_arg(args, int));
                    break;
                case 'f':
                    printf("%f", va_arg(args, double));
                    break;
                case 'c':
                    printf("%c", va_arg(args, int)); // char is promoted to int
                    break;
                case 's':
                    printf("%s", va_arg(args, char*));
                    break;
                default:
                    printf("%%", *format);
                    break;
            }
        } else {
            printf("%c", *format);
        }
        format++;
    }

    va_end(args);
}

int main() {
    my_printf("Hello, %s! You are %d years old.\n", "World", 25);
    my_printf("Pi is approximately %f\n", 3.14159);
    my_printf("Character: %c\n", 'A');

    return 0;
}
```

Inline Functions (C99)

Inline functions are a C99 feature that suggests to the compiler that the function code should be inserted directly at the call site, potentially reducing function call overhead.

Basic Inline Functions

```
#include <stdio.h>

// Inline function suggestion (C99)
inline int square(int x) {
    return x * x;
}

inline int max(int a, int b) {
    return (a > b) ? a : b;
}

int main() {
    int result1 = square(5);
    int result2 = max(10, 20);

    printf("Square of 5: %d\n", result1);
    printf("Maximum of 10 and 20: %d\n", result2);

    return 0;
}
```

Performance Considerations

```
#include <stdio.h>
#include <time.h>

// Regular function
int regular_square(int x) {
    return x * x;
}

// Inline function
inline int inline_square(int x) {
    return x * x;
}

int main() {
    const int iterations = 100000000;
```

```

    clock_t start, end;

    // Test regular function
    start = clock();
    volatile int sum1 = 0;
    for (int i = 0; i < iterations; i++) {
        sum1 += regular_square(i % 1000);
    }
    end = clock();
    double regular_time = ((double)(end - start)) / CLOCKS_PER_SEC;

    // Test inline function
    start = clock();
    volatile int sum2 = 0;
    for (int i = 0; i < iterations; i++) {
        sum2 += inline_square(i % 1000);
    }
    end = clock();
    double inline_time = ((double)(end - start)) / CLOCKS_PER_SEC;

    printf("Regular function time: %.4f seconds\n", regular_time);
    printf("Inline function time: %.4f seconds\n", inline_time);
    printf("Results: %d, %d\n", sum1, sum2);

    return 0;
}

```

Function Pointers

Function pointers are variables that store the address of a function. They enable dynamic function selection, callbacks, and more flexible program design.

Basic Function Pointers

```

#include <stdio.h>

// Function definitions
int add(int a, int b) {
    return a + b;
}

int subtract(int a, int b) {
    return a - b;
}

```



```

int multiply(int a, int b) {
    return a * b;
}

int divide(int a, int b) {
    if (b != 0) {
        return a / b;
    }
    return 0;
}

int main() {
    // Function pointer declaration
    int (*operation)(int, int);

    // Assign function addresses to pointer
    operation = add;
    printf("5 + 3 = %d\n", operation(5, 3));

    operation = subtract;
    printf("5 - 3 = %d\n", operation(5, 3));

    operation = multiply;
    printf("5 * 3 = %d\n", operation(5, 3));

    operation = divide;
    printf("6 / 3 = %d\n", operation(6, 3));

    return 0;
}

```

Function Pointer Arrays

```

#include <stdio.h>

// Mathematical operations
int add(int a, int b) { return a + b; }
int subtract(int a, int b) { return a - b; }
int multiply(int a, int b) { return a * b; }
int divide(int a, int b) { return b != 0 ? a / b : 0; }

int main() {
    // Array of function pointers
    int (*operations[])(int, int) = {add, subtract, multiply, divide};
    char *operation_names[] = {"Add", "Subtract", "Multiply", "Divide"};
}

```

```

    int x = 10, y = 5;

    printf("Operations on %d and %d:\n", x, y);
    for (int i = 0; i < 4; i++) {
        int result = operations[i](x, y);
        printf("%s: %d\n", operation_names[i], result);
    }

    return 0;
}

```

Callback Functions

```

#include <stdio.h>

// Function that takes a function pointer as parameter
void process_array(int arr[], int size, void (*callback)(int)) {
    for (int i = 0; i < size; i++) {
        callback(arr[i]);
    }
}

// Callback functions
void print_square(int n) {
    printf("%d^2 = %d\n", n, n * n);
}

void print_double(int n) {
    printf("%d * 2 = %d\n", n, n * 2);
}

void print_negative(int n) {
    printf("-%d = %d\n", n, -n);
}

int main() {
    int numbers[] = {1, 2, 3, 4, 5};
    int size = sizeof(numbers) / sizeof(numbers[0]);

    printf("Squares:\n");
    process_array(numbers, size, print_square);

    printf("\nDoubles:\n");
    process_array(numbers, size, print_double);
}

```

```

    printf("\nNegatives:\n");
    process_array(numbers, size, print_negative);

    return 0;
}

```

Higher-Order Functions

Higher-order functions are functions that take other functions as parameters or return functions as results. While C doesn't support returning functions directly, it can work with function pointers to achieve similar functionality.

Function Factory Pattern

```

#include <stdio.h>
#include <stdlib.h>

// Function pointer type for mathematical operations
typedef int (*math_operation)(int, int);

// Function that returns a function pointer
math_operation get_operation(char op) {
    switch (op) {
        case '+': return (math_operation)add;
        case '-': return (math_operation)subtract;
        case '*': return (math_operation)multiply;
        case '/': return (math_operation)divide;
        default: return NULL;
    }
}

// Mathematical operations (defined earlier)
int add(int a, int b) { return a + b; }
int subtract(int a, int b) { return a - b; }
int multiply(int a, int b) { return a * b; }
int divide(int a, int b) { return b != 0 ? a / b : 0; }

int main() {
    char operators[] = {'+', '-', '*', '/'};
    int x = 15, y = 3;

    for (int i = 0; i < 4; i++) {
        math_operation op = get_operation(operators[i]);
        if (op != NULL) {

```

```

        int result = op(x, y);
        printf("%d %c %d = %d\n", x, operators[i], y, result);
    }

    return 0;
}

```

Best Practices for Advanced Functions

1. Recursion Best Practices

```

#include <stdio.h>

// Good: Tail recursion optimization possible
int factorial_tail(int n, int accumulator) {
    if (n <= 1) {
        return accumulator;
    }
    return factorial_tail(n - 1, n * accumulator);
}

int factorial(int n) {
    return factorial_tail(n, 1);
}

// Avoid: Deep recursion that may cause stack overflow
/*
int fibonacci_slow(int n) {
    if (n <= 1) return n;
    return fibonacci_slow(n - 1) + fibonacci_slow(n - 2); // Exponential time!
}
*/

// Better: Iterative approach or memoization
int fibonacci_fast(int n) {
    if (n <= 1) return n;

    int a = 0, b = 1, result;
    for (int i = 2; i <= n; i++) {
        result = a + b;
        a = b;
        b = result;
    }
    return result;
}

```

```

}

int main() {
    printf("Factorial of 5: %d\n", factorial(5));
    printf("Fibonacci of 10: %d\n", fibonacci_fast(10));

    return 0;
}

```

2. Variable Argument Functions Best Practices

```

#include <stdio.h>
#include <stdarg.h>

// Good: Clear count parameter
int sum_with_count(int count, ...) {
    if (count <= 0) return 0;

    va_list args;
    va_start(args, count);

    int total = 0;
    for (int i = 0; i < count; i++) {
        total += va_arg(args, int);
    }

    va_end(args);
    return total;
}

// Better: Sentinel-based approach (like printf)
int sum_with_sentinel(int first, ...) {
    if (first == 0) return 0;

    va_list args;
    va_start(args, first);

    int total = first;
    int value;
    while ((value = va_arg(args, int)) != 0) {
        total += value;
    }

    va_end(args);
    return total;
}

```

```

}

int main() {
    printf("Sum with count: %d\n", sum_with_count(4, 1, 2, 3, 4));
    printf("Sum with sentinel: %d\n", sum_with_sentinel(1, 2, 3, 4, 0));

    return 0;
}

```

3. Function Pointer Best Practices

```

#include <stdio.h>

// Good: Type-safe function pointers
typedef int (*binary_operation)(int, int);

int add(int a, int b) { return a + b; }
int subtract(int a, int b) { return a - b; }

// Good: Function pointer validation
int safe_calculate(binary_operation op, int x, int y) {
    if (op == NULL) {
        printf("Error: Invalid operation\n");
        return 0;
    }
    return op(x, y);
}

int main() {
    binary_operation operations[] = {add, subtract, NULL};
    int x = 10, y = 5;

    for (int i = 0; operations[i] != NULL; i++) {
        int result = safe_calculate(operations[i], x, y);
        printf("Result: %d\n", result);
    }

    // Test with NULL pointer
    int result = safe_calculate(NULL, x, y);
    printf("NULL operation result: %d\n", result);

    return 0;
}

```

Common Pitfalls and How to Avoid Them

1. Infinite Recursion

```
#include <stdio.h>

// Wrong: Missing or incorrect base case
/*
int factorial_wrong(int n) {
    return n * factorial_wrong(n - 1); // Infinite recursion!
}
*/

// Correct: Proper base case
int factorial_correct(int n) {
    if (n < 0) return -1; // Handle negative input
    if (n <= 1) return 1; // Base case
    return n * factorial_correct(n - 1);
}

int main() {
    printf("Factorial of 5: %d\n", factorial_correct(5));
    printf("Factorial of -3: %d\n", factorial_correct(-3));

    return 0;
}
```

2. Variable Argument Function Errors

```
#include <stdio.h>
#include <stdarg.h>

// Wrong: Incorrect argument type retrieval
/*
void print_numbers_wrong(int count, ...) {
    va_list args;
    va_start(args, count);

    for (int i = 0; i < count; i++) {
        double value = va_arg(args, double); // Wrong if ints were passed
        printf("%.2f ", value);
    }

    va_end(args);
    printf("\n");
}
```

```

}
*/

// Correct: Match argument types
void print_numbers_correct(int count, ...) {
    va_list args;
    va_start(args, count);

    for (int i = 0; i < count; i++) {
        int value = va_arg(args, int); // Correct type
        printf("%d ", value);
    }

    va_end(args);
    printf("\n");
}

int main() {
    print_numbers_correct(4, 10, 20, 30, 40);

    return 0;
}

```

3. Function Pointer Errors

```

#include <stdio.h>

int add(int a, int b) { return a + b; }

int main() {
    // Wrong: Incorrect function pointer type
    /*
    int (*wrong_ptr)(int) = add; // Mismatch: add takes 2 params
    */

    // Correct: Matching function pointer type
    int (*correct_ptr)(int, int) = add;

    int result = correct_ptr(5, 3);
    printf("5 + 3 = %d\n", result);

    return 0;
}

```


Practical Examples

Recursive Directory Traversal

```
#include <stdio.h>
#include <string.h>

// Simulate directory structure
typedef struct {
    char name[50];
    int is_directory;
    struct {
        char name[50];
        int is_directory;
    } contents[10];
    int content_count;
} Directory;

// Recursive function to print directory structure
void print_directory(Directory *dir, int depth) {
    // Print indentation
    for (int i = 0; i < depth; i++) {
        printf("  ");
    }

    // Print directory name
    printf("%s%s\n", dir->name, dir->is_directory ? "/" : "");

    // Recursively print contents
    if (dir->is_directory) {
        for (int i = 0; i < dir->content_count; i++) {
            Directory sub_dir;
            strcpy(sub_dir.name, dir->contents[i].name);
            sub_dir.is_directory = dir->contents[i].is_directory;
            // In real implementation, would load actual directory data
            print_directory(&sub_dir, depth + 1);
        }
    }
}

int main() {
    Directory root = {"root", 1, {"documents", 1, {"file1.txt", 0}}, 2};
    print_directory(&root, 0);

    return 0;
}
```

Mathematical Expression Evaluator

```
#include <stdio.h>
#include <stdarg.h>
#include <string.h>

// Function pointer type for operations
typedef double (*operation_func)(double, double);

// Mathematical operations
double add(double a, double b) { return a + b; }
double subtract(double a, double b) { return a - b; }
double multiply(double a, double b) { return a * b; }
double divide(double a, double b) { return b != 0 ? a / b : 0; }

// Calculator with variable arguments
double calculate(const char *operations, double first, ...) {
    va_list args;
    va_start(args, first);

    double result = first;
    int op_index = 0;

    while (operations[op_index] != '\0') {
        double operand = va_arg(args, double);
        operation_func op;

        switch (operations[op_index]) {
            case '+': op = add; break;
            case '-': op = subtract; break;
            case '*': op = multiply; break;
            case '/': op = divide; break;
            default: op = NULL; break;
        }

        if (op != NULL) {
            result = op(result, operand);
        }

        op_index++;
    }

    va_end(args);
    return result;
}

int main() {
```

```

    // Calculate: 10 + 5 * 2 - 3 / 2
    double result = calculate("+-*/", 10.0, 5.0, 2.0, 3.0, 2.0);
    printf("Result: %.2f\n", result);

    return 0;
}

```

Event Handler System

```

#include <stdio.h>
#include <string.h>

#define MAX_HANDLERS 10

// Event handler function pointer type
typedef void (*event_handler)(const char *event_data);

// Event handler registry
typedef struct {
    char event_name[50];
    event_handler handlers[MAX_HANDLERS];
    int handler_count;
} EventHandlerRegistry;

EventHandlerRegistry registry = {0};

// Register event handler
void register_handler(const char *event_name, event_handler handler) {
    if (registry.handler_count < MAX_HANDLERS) {
        registry.handlers[registry.handler_count] = handler;
        registry.handler_count++;
        strcpy(registry.event_name, event_name);
        printf("Handler registered for event: %s\n", event_name);
    }
}

// Trigger event
void trigger_event(const char *event_name, const char *event_data) {
    printf("Triggering event: %s\n", event_name);
    for (int i = 0; i < registry.handler_count; i++) {
        if (registry.handlers[i] != NULL) {
            registry.handlers[i](event_data);
        }
    }
}

```

```

// Sample event handlers
void on_user_login(const char *user_data) {
    printf("User logged in: %s\n", user_data);
}

void on_user_logout(const char *user_data) {
    printf("User logged out: %s\n", user_data);
}

void on_data_update(const char *data) {
    printf("Data updated: %s\n", data);
}

int main() {
    // Register handlers
    register_handler("user_login", on_user_login);
    register_handler("user_logout", on_user_logout);
    register_handler("data_update", on_data_update);

    // Trigger events
    trigger_event("user_login", "Alice");
    trigger_event("data_update", "Profile information");
    trigger_event("user_logout", "Alice");

    return 0;
}

```

Summary

In this chapter, you've learned about:

1. **Recursive Functions:** Functions that call themselves to solve problems
2. **Variable Argument Functions:** Functions that accept a variable number of parameters
3. **Inline Functions:** C99 feature for potential performance optimization
4. **Function Pointers:** Variables that store function addresses
5. **Higher-Order Functions:** Functions that work with other functions
6. **Callback Functions:** Functions passed as parameters to other functions
7. **Best Practices:** Writing safe and efficient advanced functions
8. **Common Pitfalls:** Avoiding typical errors with advanced functions

Advanced functions provide powerful mechanisms for creating flexible, efficient, and maintainable C programs. These concepts are essential for implementing complex algorithms, building extensible systems, and leveraging the full power of the C language. In the next chapter, we'll explore modular programming techniques that build upon these advanced function concepts.

Modular Programming

Introduction

Modular programming organizes code into separate, independent modules that combine into a complete program. In C, you achieve this with header files, source files, and careful linkage. The payoff is reusability, maintainability, and clear separation of concerns.

This chapter walks through a full multi-module project, a production-style Makefile, **static** vs **extern**, static and shared libraries, header design, opaque types, and include hygiene — all with code that compiles under:

```
gcc -std=c17 -Wall -Wextra
```

Header Files and Source Files

Header Files (.h)

Headers declare the **interface** the rest of the program can see:

- Function prototypes
- Type definitions (**struct**, **union**, **enum**, **typedef**)
- Macro definitions
- Global variable **declarations** (**extern**)
- Includes of other headers the interface needs

Source Files (.c)

Source files provide the **implementation**:

- Function definitions
- Global variable **definitions**
- File-scope helpers marked **static**
- Algorithms and data

Example Layout

```
project/  
  include/  
    strstats.h  
  src/  
    strstats.c  
    main.c  
  Makefile
```

Include Guards

Include guards prevent a header from being processed more than once in a translation unit (which would cause redefinition errors).

Traditional Include Guards

```
#ifndef STRSTATS_H  
#define STRSTATS_H  
  
// Header content goes here  
  
#endif /* STRSTATS_H */
```

#pragma once (Non-standard, Widely Supported)

```
#pragma once  
  
// Header content goes here
```

Prefer traditional guards for maximum portability; many codebases use both.

Compilation Units and Linking

Each `.c` file is a **translation unit**. The compiler turns it into an object file (`.o`). The linker then combines object files, resolves symbols, and produces the executable (or library).

```
# Compile each unit separately  
gcc -std=c17 -Wall -Wextra -c strstats.c -o strstats.o  
gcc -std=c17 -Wall -Wextra -c main.c -o main.o
```

```
# Link
gcc strstats.o main.o -o program
```

Pipeline summary:

1. **Compile** — syntax check, generate machine code, emit undefined external references
2. **Link** — match calls to definitions, bind addresses, pull in libraries
3. **Load / run** — OS maps the binary and starts `main`

Linkage: extern, static, and No Linkage

External Linkage (Default for File-Scope Functions)

Symbols with external linkage are visible across translation units. Function names at file scope have external linkage by default. Global variables need an `extern` declaration in a header and exactly one definition in one `.c` file.

```
/* config.h */
#ifndef CONFIG_H
#define CONFIG_H

extern int g_verbose;          /* declaration - many TUs may include this */
int parse_args(int argc, char **argv);

#endif

/* config.c */
#include "config.h"

int g_verbose = 0;             /* definition - exactly once in the program */

int parse_args(int argc, char **argv) {
    for (int i = 1; i < argc; i++) {
        if (argv[i][0] == '-' && argv[i][1] == 'v') {
            g_verbose = 1;
        }
    }
    return 0;
}

/* main.c */
#include <stdio.h>
#include "config.h"

int main(int argc, char **argv) {
```

```

    parse_args(argc, argv);
    if (g_verbose) {
        printf("verbose mode on\n");
    }
    return 0;
}

```

Build:

```

gcc -std=c17 -Wall -Wextra -c config.c -o config.o
gcc -std=c17 -Wall -Wextra -c main.c -o main.o
gcc config.o main.o -o config_demo
./config_demo -v

```

Internal Linkage (static at File Scope)

`static` on a file-scope function or variable restricts it to that translation unit. Other `.c` files cannot call or see it — even if they somehow declare the same name.

```

/* helpers.c */
#include <stdio.h>

/* Only visible inside helpers.c */
static int clamp(int x, int lo, int hi) {
    if (x < lo) return lo;
    if (x > hi) return hi;
    return x;
}

static int call_count = 0; /* private counter */

int public_scale(int x) {
    call_count++;
    return clamp(x, 0, 100);
}

void public_report(void) {
    printf("public_scale called %d times\n", call_count);
}

/* main.c - cannot call clamp(); link would fail if it tried */
#include <stdio.h>

int public_scale(int x);
void public_report(void);

```



```
int main(void) {
    printf("%d\n", public_scale(150)); /* 100 */
    printf("%d\n", public_scale(-5)); /* 0 */
    public_report();
    return 0;
}
```

```
gcc -std=c17 -Wall -Wextra helpers.c main.c -o static_demo
./static_demo
```

static vs extern at a Glance

Keyword / form	Scope	Linkage	Typical use
int foo(void); in header	program	external	public API
static int helper(void) in .c	one .c file	internal	private helper
extern int g; in header	program	external	shared global (use sparingly)
static int g; in .c	one .c file	internal	module-private state
local int x; in a function	block	none	ordinary local

No Linkage

Parameters and locals have **no linkage**. Each function invocation gets its own stack storage; names are not visible outside the block.

Complete Multi-Module Project: string_stats

A small library that analyzes C strings: length, word count, letter/digit counts, and case-folded copy. Public API lives in a header; implementation and private helpers stay in the .c file.

Project Tree

```
string_stats/
  strstats.h
  strstats.c
  main.c
  Makefile
```

strstats.h — Public Interface

```
#ifndef STRSTATS_H
#define STRSTATS_H

#include <stddef.h> /* size_t */

typedef struct {
    size_t length;      /* characters, excluding '\0' */
    size_t words;       /* whitespace-separated tokens */
    size_t letters;
    size_t digits;
    size_t spaces;
    size_t others;
} StrStats;

/* Analyze src; returns 0 on success, -1 if src is NULL. */
int strstats_analyze(const char *src, StrStats *out);

/* Print a human-readable report to stdout. */
void strstats_print(const StrStats *s);

/*
 * Copy src into dest, converting A-Z to a-z.
 * dest_size is the full buffer size (including room for '\0').
 * Returns 0 on success, -1 on NULL args or if dest is too small.
 */
int strstats_to_lower(char *dest, size_t dest_size, const char *src);

#endif /* STRSTATS_H */
```

strstats.c — Implementation

```
#include "strstats.h"

#include <ctype.h>
#include <stdio.h>
#include <string.h>

/* File-private: not declared in the header. */
static int is_word_char(unsigned char c) {
    return !isspace(c);
}

int strstats_analyze(const char *src, StrStats *out) {
```

```

    if (src == NULL || out == NULL) {
        return -1;
    }

    memset(out, 0, sizeof(*out));

    int in_word = 0;
    for (size_t i = 0; src[i] != '\0'; i++) {
        unsigned char c = (unsigned char)src[i];
        out->length++;

        if (isalpha(c)) {
            out->letters++;
        } else if (isdigit(c)) {
            out->digits++;
        } else if (isspace(c)) {
            out->spaces++;
        } else {
            out->others++;
        }

        if (is_word_char(c)) {
            if (!in_word) {
                out->words++;
                in_word = 1;
            }
        } else {
            in_word = 0;
        }
    }
    return 0;
}

void strstats_print(const StrStats *s) {
    if (s == NULL) {
        return;
    }
    printf("length   : %zu\n", s->length);
    printf("words    : %zu\n", s->words);
    printf("letters  : %zu\n", s->letters);
    printf("digits   : %zu\n", s->digits);
    printf("spaces   : %zu\n", s->spaces);
    printf("others   : %zu\n", s->others);
}

int strstats_to_lower(char *dest, size_t dest_size, const char *src) {
    if (dest == NULL || src == NULL || dest_size == 0) {

```

```

        return -1;
    }

    size_t need = strlen(src) + 1; /* include '\0' */
    if (need > dest_size) {
        return -1;
    }

    for (size_t i = 0; src[i] != '\0'; i++) {
        dest[i] = (char)tolower((unsigned char)src[i]);
    }
    dest[need - 1] = '\0';
    return 0;
}

```

main.c — Driver

```

#include <stdio.h>

#include "strstats.h"

int main(void) {
    const char *sample = "Hello, C17 World! 42 cats.";
    StrStats stats;

    if (strstats_analyze(sample, &stats) != 0) {
        fprintf(stderr, "analyze failed\n");
        return 1;
    }

    printf("Input: \"%s\"\n\n", sample);
    strstats_print(&stats);

    char lower[128];
    if (strstats_to_lower(lower, sizeof lower, sample) != 0) {
        fprintf(stderr, "to_lower failed\n");
        return 1;
    }
    printf("\nlower: \"%s\"\n", lower);
    return 0;
}

```

Manual Build

```
gcc -std=c17 -Wall -Wextra -c strstats.c -o strstats.o
gcc -std=c17 -Wall -Wextra -c main.c -o main.o
gcc strstats.o main.o -o string_stats
./string_stats
```

Expected output (abbreviated):

```
Input: "Hello, C17 World! 42 cats."

length  : 26
words   : 5
letters : 15
digits  : 2
...
lower: "hello, c17 world! 42 cats."
```

Full Makefile: Debug, Release, and -MMD Dependencies

A Makefile that scales beyond one file: separate build directories, automatic header dependency tracking, and clean targets.

```
# Makefile for string_stats
CC      := gcc
STD     := -std=c17
WARN    := -Wall -Wextra -Wpedantic
SRCS    := strstats.c main.c
OBJS    := $(SRCS:.c=.o)
DEPS    := $(OBJS:.o=.d)

TARGET := string_stats

# Default = debug
BUILD  ?= debug

ifeq ($(BUILD),release)
    CFLAGS := $(STD) $(WARN) -O2 -DNDEBUG
else
    CFLAGS := $(STD) $(WARN) -g -O0
endif

# -MMD: write .d dependency files; -MP: phony headers so deletes don't break make
DEPFLAGS := -MMD -MP
```

```

.PHONY: all debug release clean run

all: $(TARGET)

debug:
    $(MAKE) BUILD=debug

release:
    $(MAKE) BUILD=release

$(TARGET): $(OBJS)
    $(CC) $(OBJS) -o $@

%.o: %.c
    $(CC) $(CFLAGS) $(DEPFLAGS) -c $< -o $@

run: $(TARGET)
    ./$$(TARGET)

clean:
    rm -f $(OBJS) $(DEPS) $(TARGET)

# Pull in generated dependency fragments (if present)
#include $(DEPS)

```

Usage:

```

make          # debug build
make release  # optimized
make run
make clean

```

Why `-MMD -MP` matters: if you edit `strstats.h`, Make recompiles every `.c` that includes it — without you listing headers by hand.

Variant with build/ Directory

```

CC      := gcc
CFLAGS  := -std=c17 -Wall -Wextra -g -O0 -MMD -MP
SRCS    := strstats.c main.c
BUILD   := build
OBJS    := $(addprefix $(BUILD)/,$(SRCS:.c=.o))
DEPS    := $(OBJS:.o=.d)
TARGET  := $(BUILD)/string_stats

.PHONY: all clean

```

```

all: $(TARGET)

$(TARGET): $(OBJS) | $(BUILD)
    $(CC) $(OBJS) -o $@

$(BUILD)/%.o: %.c | $(BUILD)
    $(CC) $(CFLAGS) -c $< -o $@

$(BUILD):
    mkdir -p $(BUILD)

clean:
    rm -rf $(BUILD)

-include $(DEPS)

```

Creating Libraries

Static Library (.a)

Object code is copied into the final executable at link time. No runtime dependency on the .a file.

```

# Compile objects
gcc -std=c17 -Wall -Wextra -c strstats.c -o strstats.o

# Archive into libstrstats.a
#   r = replace/insert, c = create if needed, s = write symbol index
ar rcs libstrstats.a strstats.o

# Link: -L. looks in current dir; -lstrstats finds libstrstats.a
gcc -std=c17 -Wall -Wextra main.c -L. -lstrstats -o string_stats_static

./string_stats_static

```

Inspect the archive:

```

ar t libstrstats.a          # list members
nm libstrstats.a | head    # symbols

```

Shared Library (.so on Linux / macOS)

Shared libraries are loaded at runtime and can be shared by multiple processes.

```

# Position-independent code is required for shared objects
gcc -std=c17 -Wall -Wextra -fPIC -c strstats.c -o strstats.o

# Build the shared library
gcc -shared -o libstrstats.so strstats.o

# Link the program against it
gcc -std=c17 -Wall -Wextra main.c -L. -lstrstats -o string_stats_shared

# Runtime: loader must find libstrstats.so
# Linux:
export LD_LIBRARY_PATH=.:$LD_LIBRARY_PATH
# macOS:
export DYLD_LIBRARY_PATH=.:$DYLD_LIBRARY_PATH

./string_stats_shared

```

Check linkage:

```

# Linux
ldd ./string_stats_shared
# macOS
otool -L ./string_stats_shared

```

Static vs Shared

Static (.a)	Shared (.so)
Linked Build time when	Load / run time
Binary Larger (code copied in) size	Smaller
UpdatesRebuild app to pick up lib fix	Replace .so, restart app
Deploy Single executable (simpler)	Ship binary + libraries

Header Design Rules

1. **Headers declare; sources define.** Do not put non-inline function bodies or non-static data definitions in headers (except `static inline` helpers when appropriate).
2. **Self-contained.** A header should compile after `#include "foo.h"` alone — include what *it* needs (`stddef.h` for `size_t`, etc.).
3. **Include what you use** in .c files; do not rely on transitive includes from other headers.
4. **Minimize public surface.** Prefer incomplete types / opaque pointers for private structs.
5. **One module, one header** as a default; split only when the API is large.
6. **Document contracts** in comments: preconditions, ownership, error returns.

7. **Avoid using the preprocessor for “namespaces”** beyond a clear prefix (`strstats_`, `vec_`).
8. **Never `#include` a `.c` file** in normal modular code.

Recommended Include Order

A common, readable order inside a `.c` file:

```
/* 1. Related header first - proves the header is self-contained */
#include "strstats.h"

/* 2. Other project headers */
#include "config.h"

/* 3. Standard library headers (alphabetical is fine) */
#include <ctype.h>
#include <stdio.h>
#include <string.h>

/* 4. System / third-party headers */
/* #include <unistd.h> */
```

Why related header first: if `strstats.h` forgets a needed include, the compiler fails when compiling `strstats.c` — not later, deep in some other unit.

Forward Declarations

When you only need a pointer or reference to a type, forward-declare instead of including a heavy header:

```
/* log.h - does not need the full definition of StrStats */
#ifndef LOG_H
#define LOG_H

struct StrStats; /* forward declaration (incomplete type) */

void log_stats(const struct StrStats *s);

#endif

/* log.c - full type only needed where members are accessed */
#include "log.h"
#include "strstats.h"
#include <stdio.h>
```

```

void log_stats(const struct StrStats *s) {
    if (s == NULL) {
        return;
    }
    printf("[log] length=%zu words=%zu\n", s->length, s->words);
}

```

Forward declarations break circular header dependencies and speed up compiles.

Opaque Pointers / Incomplete Types

Hide the real struct layout so callers cannot poke private fields. The header only names the type; the .c file owns the definition.

vector.h

```

#ifndef VECTOR_H
#define VECTOR_H

#include <stddef.h>

typedef struct Vector Vector; /* incomplete - no members visible */

Vector *vector_create(void);
void    vector_destroy(Vector *v);

int      vector_push(Vector *v, double x); /* 0 ok, -1 OOM / NULL */
size_t   vector_size(const Vector *v);
int      vector_get(const Vector *v, size_t i, double *out); /* 0 ok, -1 bad index */
void     vector_clear(Vector *v);

#endif /* VECTOR_H */

```

vector.c

```

#include "vector.h"

#include <stdlib.h>
#include <string.h>

struct Vector {
    double *data;
    size_t  size;
}

```

```

    size_t capacity;
};

Vector *vector_create(void) {
    Vector *v = calloc(1, sizeof *v);
    return v; /* NULL if OOM */
}

void vector_destroy(Vector *v) {
    if (v == NULL) {
        return;
    }
    free(v->data);
    free(v);
}

int vector_push(Vector *v, double x) {
    if (v == NULL) {
        return -1;
    }
    if (v->size == v->capacity) {
        size_t new_cap = (v->capacity == 0) ? 8 : v->capacity * 2;
        double *p = realloc(v->data, new_cap * sizeof *p);
        if (p == NULL) {
            return -1;
        }
        v->data = p;
        v->capacity = new_cap;
    }
    v->data[v->size++] = x;
    return 0;
}

size_t vector_size(const Vector *v) {
    return v ? v->size : 0;
}

int vector_get(const Vector *v, size_t i, double *out) {
    if (v == NULL || out == NULL || i >= v->size) {
        return -1;
    }
    *out = v->data[i];
    return 0;
}

void vector_clear(Vector *v) {
    if (v) {

```

```

        v->size = 0;
    }
}

```

main_vector.c

```

#include <stdio.h>
#include "vector.h"

int main(void) {
    Vector *v = vector_create();
    if (v == NULL) {
        return 1;
    }

    vector_push(v, 1.5);
    vector_push(v, 2.5);
    vector_push(v, 3.0);

    for (size_t i = 0; i < vector_size(v); i++) {
        double x;
        if (vector_get(v, i, &x) == 0) {
            printf("v[%zu] = %.1f\n", i, x);
        }
    }

    /* This would NOT compile - incomplete type:
    printf("%zu\n", v->size);
    */

    vector_destroy(v);
    return 0;
}

gcc -std=c17 -Wall -Wextra vector.c main_vector.c -o vector_demo
./vector_demo

```

Benefits: you can change `struct Vector` internals without breaking callers, and you force all access through a documented API.

Alternate Multi-Module Example: 2D Vector Math

Same patterns, different domain — useful if you prefer numeric code over string stats.

vec2.h

```
#ifndef VEC2_H
#define VEC2_H

typedef struct {
    double x;
    double y;
} Vec2;

Vec2  vec2_add(Vec2 a, Vec2 b);
Vec2  vec2_sub(Vec2 a, Vec2 b);
Vec2  vec2_scale(Vec2 a, double s);
double vec2_dot(Vec2 a, Vec2 b);
double vec2_length(Vec2 a);
void  vec2_print(const char *label, Vec2 a);

#endif
```

vec2.c

```
#include "vec2.h"

#include <math.h>
#include <stdio.h>

Vec2 vec2_add(Vec2 a, Vec2 b) {
    return (Vec2){a.x + b.x, a.y + b.y};
}

Vec2 vec2_sub(Vec2 a, Vec2 b) {
    return (Vec2){a.x - b.x, a.y - b.y};
}

Vec2 vec2_scale(Vec2 a, double s) {
    return (Vec2){a.x * s, a.y * s};
}

double vec2_dot(Vec2 a, Vec2 b) {
    return a.x * b.x + a.y * b.y;
}

double vec2_length(Vec2 a) {
    return sqrt(vec2_dot(a, a));
}
```

```

void vec2_print(const char *label, Vec2 a) {
    printf("%s: (%.3f, %.3f)\n", label ? label : "v", a.x, a.y);
}

```

main_vec2.c

```

#include "vec2.h"
#include <stdio.h>

int main(void) {
    Vec2 a = {3.0, 4.0};
    Vec2 b = {1.0, 0.0};

    vec2_print("a", a);
    vec2_print("b", b);
    vec2_print("a+b", vec2_add(a, b));
    printf("|a| = %.3f\n", vec2_length(a));
    printf("a·b = %.3f\n", vec2_dot(a, b));
    return 0;
}

gcc -std=c17 -Wall -Wextra vec2.c main_vec2.c -lm -o vec2_demo
./vec2_demo

```

Note: `-lm` links the math library for `sqrt`.

Callback Interfaces (Flexible Modules)

Function pointers let a module stay generic while callers plug in behavior:

```

/* map.h */
#ifndef MAP_H
#define MAP_H

typedef int (*map_fn)(int value, void *ctx);

void map_apply(int *arr, int n, map_fn fn, void *ctx);

#endif

/* map.c */
#include "map.h"

```

```

void map_apply(int *arr, int n, map_fn fn, void *ctx) {
    if (arr == NULL || fn == NULL || n <= 0) {
        return;
    }
    for (int i = 0; i < n; i++) {
        arr[i] = fn(arr[i], ctx);
    }
}

/* main_map.c */
#include "map.h"
#include <stdio.h>

static int add_k(int value, void *ctx) {
    int k = *(const int *)ctx;
    return value + k;
}

int main(void) {
    int a[] = {1, 2, 3, 4};
    int k = 10;
    map_apply(a, 4, add_k, &k);
    for (int i = 0; i < 4; i++) {
        printf("%d ", a[i]); /* 11 12 13 14 */
    }
    printf("\n");
    return 0;
}

```

Best Practices

1. **Naming** — Prefix public symbols with the module name (`strstats_`, `vector_`).
2. **Error handling** — Return error codes; document them; let callers decide how to report.
3. **Documentation** — Comment the header: parameters, returns, ownership, thread-safety.
4. **Testing** — Prefer unit-testable modules with minimal I/O in the library layer.
5. **Globals** — Prefer parameters and opaque state over `extern` globals.
6. **Versioning** — Keep header and library in lockstep; document API breaks.
7. **Compiler flags** — Always develop with `-std=c17 -Wall -Wextra` (add `-Wpedantic` when practical).

Exercises

1. **Split a program.** Take any single-file program you have written and split it into `utils.h` / `utils.c` / `main.c`. Add a Makefile with debug/release and `-MMD`.

2. **Extend `string_stats`.** Add `strstats_to_upper` and a function that counts vowels. Keep helpers `static` in the `.c` file.
3. **Static library.** Build `libstrstats.a` and link a second small program that only calls `strstats_analyze`. Confirm with `nm` that symbols come from the archive.
4. **Shared library.** Build `libstrstats.so`, link, and run with the appropriate library path environment variable on your OS.
5. **Opaque counter.** Implement `counter.h` / `counter.c` with `counter_create`, `counter_inc`, `counter_get`, `counter_destroy`. Do not expose the struct layout in the header.
6. **Include audit.** Deliberately break a header's self-containment (remove a needed `#include`), put the related `#include "mod.h"` first in `mod.c`, and observe the compile error. Fix it correctly.
7. **Forward declarations.** Create two modules that would circularly include each other if you `#include` both headers. Resolve the cycle with a forward declaration and move the full include into a `.c` file.
8. **Vector math library.** Add `vec2_normalize` and unit tests (or a small `assert`-based driver). Link with `-lm`.

Summary

Modular C rests on a few durable ideas:

1. **Headers declare interfaces; sources define implementations.**
2. **Include guards** keep headers safe to include repeatedly.
3. **Linkage** (`extern` / `static`) controls who can see a symbol.
4. **Makefiles** with `-MMD` keep multi-file projects rebuild-correct.
5. **Static (`.a`) and shared (`.so`) libraries** package modules for reuse.
6. **Opaque pointers** hide representation and stabilize APIs.
7. **Include order and forward declarations** keep dependencies honest and compiles fast.

Master these patterns and you can grow from single-file experiments to maintainable multi-module programs and libraries without fighting the toolchain.

Standard Library Functions

Introduction

The C standard library is a collection of functions, macros, and type definitions that are part of the C programming language specification. These functions provide essential functionality for input/output operations, string manipulation, mathematical computations, memory management, and more. Understanding and effectively using the standard library is crucial for C programmers as it provides a foundation for building robust applications.

This chapter mixes **API maps** with **complete, compilable programs**. Prefer the programs: type them, break them, and re-run.

```
gcc -std=c17 -Wall -Wextra -o prog prog.c
# math examples need the math library on many systems:
gcc -std=c17 -Wall -Wextra -o prog prog.c -lm
```

String Functions (<string.h>)

String Copying

```
#include <string.h>

char *strcpy(char *dest, const char *src);
char *strncpy(char *dest, const char *src, size_t n);
char *strdup(const char *s);           // C23 standard
char *strndup(const char *s, size_t n); // C23 standard
```

String Concatenation

```
char *strcat(char *dest, const char *src);
char *strncat(char *dest, const char *src, size_t n);
```

String Comparison

```
int strcmp(const char *s1, const char *s2);
int strncmp(const char *s1, const char *s2, size_t n);
int strcoll(const char *s1, const char *s2); // Locale-specific
```

String Searching

```
char *strchr(const char *s, int c); // Find first occurrence
char *strrchr(const char *s, int c); // Find last occurrence
char *strstr(const char *haystack, const char *needle); // Find substring
size_t strcspn(const char *s, const char *reject); // Span of non-matching chars
size_t strspn(const char *s, const char *accept); // Span of matching chars
char *strpbrk(const char *s, const char *accept); // Find any of a set of chars
```

String Tokenization

```
char *strtok(char *str, const char *delim);
```

String Length

```
size_t strlen(const char *s);
size_t strnlen(const char *s, size_t maxlen); // C23 standard
```

Memory Functions (<string.h>)

Memory Copying

```
void *memcpy(void *dest, const void *src, size_t n);
void *memmove(void *dest, const void *src, size_t n);
```

Memory Setting

```
void *memset(void *s, int c, size_t n);
```

Memory Comparison

```
int memcmp(const void *s1, const void *s2, size_t n);
```

Memory Searching

```
void *memchr(const void *s, int c, size_t n);
```

Mathematical Functions (<math.h>)

Basic Mathematical Functions

```
#include <math.h>

double fabs(double x);           // Absolute value
double ceil(double x);          // Ceiling function
double floor(double x);         // Floor function
double round(double x);         // Round to nearest integer
double trunc(double x);         // Truncate towards zero

double sqrt(double x);          // Square root
double cbrt(double x);          // Cube root
double pow(double x, double y); // Power function
double exp(double x);           // Exponential function
double log(double x);           // Natural logarithm
double log10(double x);         // Base-10 logarithm
double log2(double x);          // Base-2 logarithm

double sin(double x);           // Sine function
double cos(double x);           // Cosine function
double tan(double x);           // Tangent function
double asin(double x);          // Arc sine
double acos(double x);          // Arc cosine
double atan(double x);          // Arc tangent
double atan2(double y, double x); // Arc tangent with two arguments
```

Hyperbolic Functions

```
double sinh(double x);           // Hyperbolic sine
double cosh(double x);           // Hyperbolic cosine
double tanh(double x);           // Hyperbolic tangent
```

Floating-Point Functions

```
double fmod(double x, double y); // Floating-point remainder
double remainder(double x, double y); // IEEE remainder
double copysign(double x, double y); // Copy sign
double nextafter(double x, double y); // Next representable value
```

Classification and Comparison Functions

```
int isfinite(double x);           // Check if finite
int isinf(double x);              // Check if infinite
int isnan(double x);              // Check if NaN
int isnormal(double x);           // Check if normal
int signbit(double x);            // Check sign
```

Character Classification (<ctype.h>)

Character Testing

```
#include <ctype.h>

int isalnum(int c);               // Alphanumeric character
int isalpha(int c);               // Alphabetic character
int isdigit(int c);               // Decimal digit
int isxdigit(int c);              // Hexadecimal digit
int islower(int c);               // Lowercase letter
int isupper(int c);               // Uppercase letter
int isblank(int c);               // Blank character (space or tab)
int isspace(int c);               // Whitespace character
int iscntrl(int c);               // Control character
int isprint(int c);               // Printable character
int isgraph(int c);               // Graphic character
int ispunct(int c);               // Punctuation character
```

Character Conversion

```
int tolower(int c);           // Convert to lowercase
int toupper(int c);          // Convert to uppercase
```

Time Functions (<time.h>)

Time Types

```
#include <time.h>

typedef long time_t;           // Calendar time
typedef long clock_t;         // Processor time
struct tm {                   // Broken-down time
    int tm_sec;               // Seconds (0-60)
    int tm_min;               // Minutes (0-59)
    int tm_hour;              // Hours (0-23)
    int tm_mday;              // Day of month (1-31)
    int tm_mon;               // Month (0-11)
    int tm_year;              // Year since 1900
    int tm_wday;              // Day of week (0-6, Sunday = 0)
    int tm_yday;              // Day of year (0-365)
    int tm_isdst;             // Daylight saving time flag
};
```

Time Functions

```
time_t time(time_t *tloc);    // Get current time
double difftime(time_t time1, time_t time0); // Difference between two times
time_t mktime(struct tm *tm); // Convert broken-down time to calendar time
char *asctime(const struct tm *tm); // Convert to string
char *ctime(const time_t *timep); // Convert to string
struct tm *gmtime(const time_t *timep); // Convert to UTC
struct tm *localtime(const time_t *timep); // Convert to local time
size_t strftime(char *s, size_t max, const char *format, const struct tm *tm);
```

Input/Output Functions (<stdio.h>)

File Operations

```
#include <stdio.h>

FILE *fopen(const char *pathname, const char *mode);
int fclose(FILE *stream);
int fflush(FILE *stream);
FILE *freopen(const char *pathname, const char *mode, FILE *stream);
FILE *fdopen(int fd, const char *mode);
FILE *fmemopen(void *buf, size_t size, const char *mode); // C23 standard
```

Formatted I/O

```
int printf(const char *format, ...);
int fprintf(FILE *stream, const char *format, ...);
int sprintf(char *str, const char *format, ...);
int snprintf(char *str, size_t size, const char *format, ...); // C99

int scanf(const char *format, ...);
int fscanf(FILE *stream, const char *format, ...);
int sscanf(const char *str, const char *format, ...);
```

Character I/O

```
int fgetc(FILE *stream);
char *fgets(char *s, int size, FILE *stream);
int fputc(int c, FILE *stream);
int fputs(const char *s, FILE *stream);
int getchar(void);
char *gets(char *s); // Deprecated, unsafe
int putchar(int c);
int puts(const char *s);
int ungetc(int c, FILE *stream);
```

Direct I/O

```
size_t fread(void *ptr, size_t size, size_t nmemb, FILE *stream);
size_t fwrite(const void *ptr, size_t size, size_t nmemb, FILE *stream);
```

File Positioning

```
int fseek(FILE *stream, long offset, int whence);
long ftell(FILE *stream);
void rewind(FILE *stream);
int fgetpos(FILE *stream, fpos_t *pos);
int fsetpos(FILE *stream, const fpos_t *pos);
```

Error Handling

```
void clearerr(FILE *stream);
int feof(FILE *stream);
int ferror(FILE *stream);
void perror(const char *s);
```

Memory Management (<stdlib.h>)

Dynamic Memory Allocation

```
#include <stdlib.h>

void *malloc(size_t size);
void *calloc(size_t nmemb, size_t size);
void *realloc(void *ptr, size_t size);
void free(void *ptr);
void *aligned_alloc(size_t alignment, size_t size); // C11 standard
```

Utility Functions

```
double atof(const char *nptr);           // Convert string to double
int atoi(const char *nptr);               // Convert string to int
long atol(const char *nptr);              // Convert string to long
long long atoll(const char *nptr);        // Convert string to long long (C99)

double strtod(const char *nptr, char **endptr); // Convert string to double
float strtof(const char *nptr, char **endptr);   // Convert string to float (C99)
long double strtold(const char *nptr, char **endptr); // Convert string to long double (C99)
long strtol(const char *nptr, char **endptr, int base);
long long strtoll(const char *nptr, char **endptr, int base); // C99
unsigned long strtoul(const char *nptr, char **endptr, int base);
unsigned long long strtoull(const char *nptr, char **endptr, int base); // C99
```

Random Number Generation

```
int rand(void);           // Generate pseudo-random number
void srand(unsigned int seed); // Seed the random number generator
```

Environment Functions

```
void abort(void);          // Cause abnormal program termination
int atexit(void (*func)(void)); // Register function to be called at exit
int at_quick_exit(void (*func)(void)); // Register function for quick exit (C11)
void exit(int status);     // Terminate program normally
void _Exit(int status);    // Terminate program immediately (C99)
void quick_exit(int status); // Quick program termination (C11)
char *getenv(const char *name); // Get environment variable
int system(const char *command); // Execute system command
```

Searching and Sorting

```
void *bsearch(const void *key, const void *base,
              size_t nmemb, size_t size,
              int (*compar)(const void *, const void *));
void qsort(void *base, size_t nmemb, size_t size,
           int (*compar)(const void *, const void *));
```

Integer Arithmetic

```
int abs(int j);           // Absolute value of int
long labs(long j);        // Absolute value of long
long long llabs(long long j); // Absolute value of long long (C99)
div_t div(int numer, int denom); // Integer division
ldiv_t ldiv(long numer, long denom); // Long integer division
lldiv_t lldiv(long long numer, long long denom); // Long long integer division (C99)
```


Advanced Standard Library Features (C23)

New String Functions

```
// C23 additions
char *strdup(const char *s);
char *strndup(const char *s, size_t n);
size_t strnlen(const char *s, size_t maxlen);
```

Enhanced Memory Functions

```
// Bounds-checking functions (Annex K, optional in C23)
errno_t memcpy_s(void *dest, rsize_t destsz,
                 const void *src, rsize_t count);
errno_t strcpy_s(char *dest, rsize_t destsz,
                 const char *src);
```

Practical Examples

String Manipulation Example

```
#include <stdio.h>
#include <string.h>
#include <stdlib.h>

int main() {
    char str1[50] = "Hello, ";
    char str2[] = "World!";
    char str3[100];

    // String concatenation
    strcpy(str3, str1);
    strcat(str3, str2);
    printf("Concatenated: %s\n", str3);

    // String comparison
    if (strcmp(str1, str2) == 0) {
        printf("Strings are equal\n");
    } else {
        printf("Strings are different\n");
    }
}
```

```

// String searching
char *found = strstr(str3, "World");
if (found) {
    printf("Found 'World' at position: %ld\n", found - str3);
}

// String tokenization
char sentence[] = "This is a sample sentence";
char *token = strtok(sentence, " ");
while (token != NULL) {
    printf("Token: %s\n", token);
    token = strtok(NULL, " ");
}

return 0;
}

```

Mathematical Functions Example

```

#include <stdio.h>
#include <math.h>

int main() {
    double x = 4.0;
    double y = 3.0;

    printf("Square root of %.2f: %.2f\n", x, sqrt(x));
    printf("%.2f raised to the power of %.2f: %.2f\n", x, y, pow(x, y));
    printf("Sine of %.2f radians: %.2f\n", x, sin(x));
    printf("Natural log of %.2f: %.2f\n", x, log(x));

    // Rounding functions
    double value = 3.7;
    printf("Value: %.2f\n", value);
    printf("Ceiling: %.2f\n", ceil(value));
    printf("Floor: %.2f\n", floor(value));
    printf("Round: %.2f\n", round(value));

    return 0;
}

```

Time Functions Example

```
#include <stdio.h>
#include <time.h>

int main() {
    // Get current time
    time_t rawtime;
    struct tm *timeinfo;
    char buffer[80];

    time(&rawtime);
    timeinfo = localtime(&rawtime);

    printf("Current local time: %s", asctime(timeinfo));

    // Format time
    strftime(buffer, sizeof(buffer), "%Y-%m-%d %H:%M:%S", timeinfo);
    printf("Formatted time: %s\n", buffer);

    // Calculate time difference
    time_t start_time = time(NULL);

    // Simulate some work
    for (int i = 0; i < 1000000; i++) {
        // Busy loop
    }

    time_t end_time = time(NULL);
    double elapsed = difftime(end_time, start_time);
    printf("Elapsed time: %.0f seconds\n", elapsed);

    return 0;
}
```

Memory Management Example

```
#include <stdio.h>
#include <stdlib.h>
#include <string.h>

int main() {
    // Dynamic memory allocation
    int n = 5;
    int *arr = (int*)malloc(n * sizeof(int));
```

```

    if (arr == NULL) {
        printf("Memory allocation failed\n");
        return 1;
    }

    // Initialize array
    for (int i = 0; i < n; i++) {
        arr[i] = i + 1;
        printf("arr[%d] = %d\n", i, arr[i]);
    }

    // Resize array
    n = 10;
    arr = (int*)realloc(arr, n * sizeof(int));

    if (arr == NULL) {
        printf("Memory reallocation failed\n");
        return 1;
    }

    // Initialize new elements
    for (int i = 5; i < n; i++) {
        arr[i] = i + 1;
        printf("arr[%d] = %d\n", i, arr[i]);
    }

    // Free memory
    free(arr);
    arr = NULL; // Good practice to avoid dangling pointers

    return 0;
}

```

Full Programs: Standard Library in Practice

The snippets above are reference maps. The programs below are meant to be **typed, compiled, and run**.

Program 1 — Safe string toolbox (strtool.c)

Demonstrates length, bounded copy, search, tokenize, and compare without buffer overflows.

```

/* file: strtool.c */
#include <stdio.h>

```

```

#include <stdlib.h>
#include <string.h>
#include <ctype.h>

/* Copy src into dest; always NUL-terminates if dest_sz > 0.
 * Returns 0 on full copy, -1 if truncated or bad args. */
static int str_copy(char *dest, size_t dest_sz, const char *src) {
    size_t i;

    if (dest == NULL || src == NULL || dest_sz == 0) {
        return -1;
    }
    for (i = 0; i + 1 < dest_sz && src[i] != '\0'; i++) {
        dest[i] = src[i];
    }
    dest[i] = '\0';
    return (src[i] == '\0') ? 0 : -1;
}

static void trim_inplace(char *s) {
    char *start, *end;
    size_t len;

    if (s == NULL) {
        return;
    }
    start = s;
    while (*start != '\0' && isspace((unsigned char)*start)) {
        start++;
    }
    if (*start == '\0') {
        s[0] = '\0';
        return;
    }
    end = start + strlen(start) - 1;
    while (end > start && isspace((unsigned char)*end)) {
        end--;
    }
    len = (size_t)(end - start + 1);
    memmove(s, start, len);
    s[len] = '\0';
}

int main(void) {
    char line[128];
    char key[32];
    char *tok;

```

```

const char *hay = "the quick brown fox jumps over the lazy dog";
const char *needle = "fox";
char *found;

printf("strlen(hay) = %zu\n", strlen(hay));
found = strstr(hay, needle);
if (found != NULL) {
    printf("strstr: \"%s\" at offset %td\n", needle, found - hay);
}

if (str_copy(key, sizeof key, "port=8080#comment") != 0) {
    fprintf(stderr, "key truncated\n");
}
printf("copy: \"%s\"\n", key);

/* strtok is standard C (not reentrant - fine for this demo) */
{
    char mutable[] = "alice,bob,carol";
    printf("tokens:");
    for (tok = strtok(mutable, ","); tok != NULL; tok = strtok(NULL, ",")) {
        printf(" [%s]", tok);
    }
    printf("\n");
}

printf("Enter a line to trim: ");
if (fgets(line, sizeof line, stdin) == NULL) {
    return EXIT_FAILURE;
}

/* strip newline then trim spaces */
line[strcspn(line, "\n")] = '\0';
trim_inplace(line);
printf("trimmed: \"%s\"\n", line);

if (strcmp(line, "quit") == 0) {
    printf("matched quit\n");
} else {
    printf("strcmp vs \"quit\": %d\n", strcmp(line, "quit"));
}

return EXIT_SUCCESS;
}

gcc -std=c17 -Wall -Wextra -o strtool strtool.c
./strtool

```

Notes

- Prefer `strncpy` / manual copies / `snprintf` over unbounded `strcpy/strcat`.
- Cast to `unsigned char` before `isspace` / `ctype.h` macros.
- `strtok` modifies its buffer and keeps hidden state — fine for simple tools; avoid in threaded code.

Program 2 — Math lab: vectors and roots (`mathlab.c`)

```

/* file: mathlab.c */
#include <stdio.h>
#include <stdlib.h>
#include <math.h>
#include <errno.h>

struct Vec2 {
    double x;
    double y;
};

static double vec_len(struct Vec2 v) {
    return hypot(v.x, v.y); /* safer than sqrt(x*x + y*y) for large values */
}

static struct Vec2 vec_norm(struct Vec2 v) {
    double len = vec_len(v);
    struct Vec2 out = {0.0, 0.0};
    if (len > 0.0) {
        out.x = v.x / len;
        out.y = v.y / len;
    }
    return out;
}

static int safe_sqrt(double x, double *out) {
    if (out == NULL) {
        return -1;
    }
    if (x < 0.0) {
        errno = EDOM;
        return -1;
    }
    *out = sqrt(x);
    return 0;
}

```

```

int main(void) {
    struct Vec2 a = {3.0, 4.0};
    struct Vec2 n = vec_norm(a);
    double root;
    double angle;

    printf("|a| = %.6f\n", vec_len(a));
    printf("unit(a) = (%.6f, %.6f)\n", n.x, n.y);

    angle = atan2(a.y, a.x);
    printf("atan2(y,x) = %.6f rad (%.2f deg)\n",
           angle, angle * 180.0 / 3.14159265358979323846);

    if (safe_sqrt(2.0, &root) == 0) {
        printf("sqrt(2) = %.10f\n", root);
    }
    if (safe_sqrt(-1.0, &root) != 0) {
        printf("sqrt(-1) rejected (errno may be EDOM)\n");
    }

    printf("pow(2,10) = %.0f\n", pow(2.0, 10.0));
    printf("floor(3.7)=%.1f ceil(3.2)=%.1f round(3.5)=%.1f\n",
           floor(3.7), ceil(3.2), round(3.5));

    return 0;
}

gcc -std=c17 -Wall -Wextra -o matlab matlab.c -lm
./matlab

```

Program 3 — Character classifier report (ctype_report.c)

```

/* file: ctype_report.c */
#include <stdio.h>
#include <ctype.h>
#include <string.h>

static void classify(int ch) {
    unsigned char c = (unsigned char)ch;
    printf("'%c' (0x%02x):", (c >= 32 && c < 127) ? c : '?', c);
    if (isdigit(c)) printf(" digit");
    if (isalpha(c)) printf(" alpha");
    if (isalnum(c)) printf(" alnum");
    if (isspace(c)) printf(" space");
    if (ispunct(c)) printf(" punct");
}

```



```

        if (isupper(c)) printf(" upper");
        if (islower(c)) printf(" lower");
        printf(" -> upper='%c' lower='%c'\n",
                toupper(c), tolower(c));
    }

int main(int argc, char *argv[]) {
    const char *s;
    size_t i;

    if (argc < 2) {
        fprintf(stderr, "Usage: %s <string>\n", argv[0]);
        return 1;
    }
    s = argv[1];
    for (i = 0; s[i] != '\0'; i++) {
        classify((unsigned char)s[i]);
    }
    return 0;
}

gcc -std=c17 -Wall -Wextra -o ctype_report ctype_report.c
./ctype_report 'Hi 42!'

```

Program 4 — Time stamps and elapsed work (timelab.c)

```

/* file: timelab.c */
#include <stdio.h>
#include <time.h>
#include <string.h>

int main(void) {
    time_t now;
    struct tm local_tm;
    char buf[64];
    clock_t t0, t1;
    volatile double sink = 0.0;
    long i;

    now = time(NULL);
    if (now == (time_t)-1) {
        perror("time");
        return 1;
    }
}

```

```

#if defined(_POSIX_VERSION)
    if (localtime_r(&now, &local_tm) == NULL) {
        perror("localtime_r");
        return 1;
    }
#else
    {
        struct tm *p = localtime(&now);
        if (p == NULL) {
            return 1;
        }
        local_tm = *p;
    }
#endif

    if (strftime(buf, sizeof buf, "%Y-%m-%d %H:%M:%S %Z", &local_tm) == 0) {
        fprintf(stderr, "strftime failed\n");
        return 1;
    }
    printf("local: %s\n", buf);
    printf("asctime: %s", asctime(&local_tm)); /* includes newline */

    t0 = clock();
    for (i = 0; i < 10000000L; i++) {
        sink += 1.0;
    }
    t1 = clock();
    printf("busy loop: %.3f seconds of CPU (clock)\n",
        (double)(t1 - t0) / (double)CLOCKS_PER_SEC);
    printf("sink=%.0f\n", sink);

    return 0;
}

```

```

gcc -std=c17 -Wall -Wextra -o timelab timelab.c
./timelab

```

Program 5 — qsort + bsearch + stdlib utilities (sortlab.c)

```

/* file: sortlab.c */
#include <stdio.h>
#include <stdlib.h>
#include <string.h>

static int cmp_int_asc(const void *a, const void *b) {

```

```

    int x = *(const int *)a;
    int y = *(const int *)b;
    return (x > y) - (x < y);
}

static int cmp_str(const void *a, const void *b) {
    const char *const *sa = a;
    const char *const *sb = b;
    return strcmp(*sa, *sb);
}

int main(void) {
    int nums[] = {42, 7, 19, 3, 100, 19};
    size_t n = sizeof nums / sizeof nums[0];
    int key = 19;
    int *found;
    const char *words[] = {"pear", "apple", "orange", "banana"};
    size_t wn = sizeof words / sizeof words[0];
    size_t i;
    char *end;
    long port;

    qsort(nums, n, sizeof nums[0], cmp_int_asc);
    printf("sorted ints:");
    for (i = 0; i < n; i++) {
        printf(" %d", nums[i]);
    }
    printf("\n");

    found = bsearch(&key, nums, n, sizeof nums[0], cmp_int_asc);
    if (found != NULL) {
        printf("bsearch %d -> index %td\n", key, found - nums);
    }

    qsort(words, wn, sizeof words[0], cmp_str);
    printf("sorted words:");
    for (i = 0; i < wn; i++) {
        printf(" %s", words[i]);
    }
    printf("\n");

    port = strtol("8080", &end, 10);
    if (end != NULL && *end == '\0' && port > 0 && port <= 65535) {
        printf("parsed port: %ld\n", port);
    }

    printf("abs(-12)=%d, RAND_MAX=%d, rand sample=%d\n",

```

```

        abs(-12), RAND_MAX, rand() % 100);

    return 0;
}

gcc -std=c17 -Wall -Wextra -o sortlab sortlab.c
./sortlab

```

Program 6 — Mini word frequency with string.h + heap (wfreq.c)

```

/* file: wfreq.c - counts words from stdin (letters/digits only) */
#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <ctype.h>

#define MAX_WORD 64
#define MAX_ENTRIES 256

struct Entry {
    char word[MAX_WORD];
    int count;
};

static int find_or_add(struct Entry *tab, int *n, const char *w) {
    int i;
    for (i = 0; i < *n; i++) {
        if (strcmp(tab[i].word, w) == 0) {
            tab[i].count++;
            return 0;
        }
    }
    if (*n >= MAX_ENTRIES) {
        return -1;
    }
    if (strlen(w) >= MAX_WORD) {
        return -1;
    }
    strcpy(tab[*n].word, w); /* bounded by check above */
    tab[*n].count = 1;
    (*n)++;
    return 0;
}

static int cmp_entry_desc(const void *a, const void *b) {

```

```

    const struct Entry *ea = a;
    const struct Entry *eb = b;
    if (ea->count != eb->count) {
        return eb->count - ea->count;
    }
    return strcmp(ea->word, eb->word);
}

int main(void) {
    struct Entry table[MAX_ENTRIES];
    int n = 0;
    char word[MAX_WORD];
    int wi = 0;
    int ch;
    int i;

    while ((ch = getchar()) != EOF) {
        if (isalnum((unsigned char)ch)) {
            if (wi + 1 < MAX_WORD) {
                word[wi++] = (char)tolower((unsigned char)ch);
            }
        } else if (wi > 0) {
            word[wi] = '\0';
            if (find_or_add(table, &n, word) != 0) {
                fprintf(stderr, "table full or word too long\n");
                return EXIT_FAILURE;
            }
            wi = 0;
        }
    }

    if (wi > 0) {
        word[wi] = '\0';
        find_or_add(table, &n, word);
    }

    qsort(table, (size_t)n, sizeof table[0], cmp_entry_desc);
    for (i = 0; i < n; i++) {
        printf("%4d  %s\n", table[i].count, table[i].word);
    }
    return 0;
}

```

```

gcc -std=c17 -Wall -Wextra -o wfreq wfreq.c
echo 'To be or not to be' | ./wfreq

```

Best Practices for Using Standard Library Functions

1. Error Checking

Always check return values for error conditions:

```
FILE *file = fopen("data.txt", "r");
if (file == NULL) {
    perror("Error opening file");
    return 1;
}
```

2. Buffer Safety

Use safe versions of functions when available and ensure adequate buffer sizes:

```
char buffer[100];
// Safe approach
if (fgets(buffer, sizeof(buffer), stdin) != NULL) {
    // Process input
}
```

Prefer `snprintf` over `sprintf`:

```
char path[64];
int n = snprintf(path, sizeof path, "/tmp/%s.log", "app");
if (n < 0 || (size_t)n >= sizeof path) {
    /* truncated or error */
}
```

3. Memory Management

Always free dynamically allocated memory:

```
int *ptr = malloc(10 * sizeof(int));
if (ptr != NULL) {
    // Use memory
    free(ptr);
    ptr = NULL; // Prevent dangling pointer
}
```

4. Locale Awareness

Be aware of locale-specific functions:

```
#include <locale.h>

setlocale(LC_ALL, ""); // Use system locale
```

5. Thread Safety

Some standard library functions are not thread-safe. Use reentrant versions when needed:

```
// Thread-safe version (POSIX)
struct tm tm_result;
localtime_r(&rawtime, &tm_result);
```

6. Link math explicitly

On many Linux toolchains, math symbols live in `libm`:

```
gcc -std=c17 -Wall -Wextra -o mathlab mathlab.c -lm
```

Exercises

```
gcc -std=c17 -Wall -Wextra -o exN exN.c    # add -lm when using <math.h>
```

1. **Bounded concat** — Write `int str_cat(char *dest, size_t dest_sz, const char *src)` that appends without overflow and always NUL-terminates. Unit-test with full and almost-full buffers.
2. **strtoul port parser** — Accept a string; accept only 1...65535; reject junk, empty, negatives, and overflow. Return status codes, not just print.
3. **Case-insensitive compare** — Implement `int str_casecmp(const char *a, const char *b)` with `tolower((unsigned char)...) and test "Hello" vs "hello".`
4. **Math: distance** — Given two `Vec2` points, print Euclidean distance with `hypot`. Compare naive `sqrt(dx*dx+dy*dy)` on large coordinates (optional exploration).
5. **qsort structs** — Sort an array of `{name, score}` by score descending, then name ascending.
6. **Time format** — Print “days since epoch” and a `strftime` format `YYYY-MM-DDThh:mm:ss`.
7. **wfreq upgrade** — Read from a filename in `argv[1]` instead of `stdin`; print only the top 10 words.
8. **ctype filter** — Read `stdin`, write only alphanumeric characters and spaces to `stdout` (like a crude sanitizer).

Summary

The C standard library provides a rich set of functions that form the foundation of most C programs. Mastering these functions is essential for effective C programming. Key points to remember:

1. **String Functions:** Essential for text processing and manipulation — prefer bounded copies and `snprintf`
2. **Mathematical Functions:** Provide a wide range of mathematical operations — link with `-lm`
3. **Character Functions:** Enable character classification and conversion — cast to `unsigned char`
4. **Time Functions:** Handle time-related operations and formatting — prefer reentrant APIs when available
5. **I/O Functions:** Enable file and console input/output operations — always check returns
6. **Memory Management:** Provide dynamic memory allocation capabilities — pair every `malloc` with `free`
7. **Utility Functions:** `qsort`, `bsearch`, `strtol`, `abs`, `rand` cover many everyday needs

Understanding and properly using these standard library functions will make your C programs more robust, portable, and efficient. Always consult the documentation for specific behavior and error handling requirements of each function.

Module 4: Functions and Modular Programming

Overview

This module explores the fundamental concepts of functions in C and how to organize code into modular, reusable components. You'll learn about function design, parameter passing, return values, and advanced function concepts like recursion and function pointers. The module also covers modular programming techniques, including header files, compilation units, and library creation.

Learning Objectives

By the end of this module, you will be able to: - Design and implement functions with proper parameters and return values - Understand parameter passing mechanisms (by value) - Create and use recursive functions effectively - Work with variable argument functions and inline functions - Utilize function pointers for callbacks and dynamic dispatch - Organize code into modular components using header and source files - Create and use both static and shared libraries - Understand linkage concepts and scope rules - Apply modular programming principles for maintainable code

Chapters

1. **Function Fundamentals** - Function definition, declaration, parameters, and return values
2. **Advanced Function Concepts** - Recursion, variable arguments, inline functions, and function pointers
3. **Modular Programming** - Header files, compilation units, and linkage concepts
4. **Standard Library Functions** - Commonly used functions from C standard library

Key Concepts Covered

- Function declaration and definition syntax
- Parameter passing by value
- Return values and the return statement
- Function prototypes and forward declarations
- Scope rules and name hiding
- Recursive function design and implementation
- Variable argument functions using `stdarg.h`
- Inline functions for performance optimization
- Function pointers and callback mechanisms

- Header file organization and include guards
- Compilation units and linking process
- Static vs. external linkage
- Creating and using libraries
- Standard library functions overview

Prerequisites

- Completion of Module 1: Foundations and Environment Setup
- Completion of Module 2: Data Types and Variables
- Completion of Module 3: Control Flow
- Understanding of basic programming concepts
- Familiarity with variable declaration and initialization

Tools and Technologies

- C compiler (GCC, Clang, or MSVC)
- Text editor or IDE
- Understanding of basic programming concepts
- Familiarity with debugging tools
- Make or CMake for building multi-file projects

Estimated Time to Complete

- Reading: 8-10 hours
- Exercises: 12-15 hours
- Projects: 6-8 hours

Assessment

- Complete all chapter exercises
- Successfully implement various function types
- Demonstrate understanding of modular programming concepts
- Pass function and modular programming quiz
- Submit a comprehensive program using advanced function features

Next Module

After completing this module, proceed to **Module 5: Arrays and Strings** to learn about data collections and text processing.

Additional Resources

- “The C Programming Language” by Kernighan and Ritchie (Chapter 4)
- C standard documentation (ISO/IEC 9899:2024 for C23)
- Online tutorials on function design patterns
- Debugging tools documentation (GDB, Valgrind)
- Best practices guides for modular programming

Arrays and Strings

Arrays

Introduction

Arrays store multiple values of the same type in a contiguous block of memory. They are the foundation of strings, buffers, tables, and many algorithms. This chapter covers declaration, indexing, traversal, passing arrays to functions, multidimensional layout, worked algorithm programs, common bugs, and a mini-lab.

Compile examples with:

```
gcc -std=c17 -Wall -Wextra program.c -o program
```

Array Declaration and Initialization

Basic Declaration

```
int numbers[10];          /* 10 integers (uninitialized if local) */
float temperatures[7];
char letters[26];
```

Initialization

```
/* Full initializer list */
int numbers[5] = {1, 2, 3, 4, 5};

/* Partial - remaining elements are zero */
int values[10] = {1, 2, 3};

/* All zeros */
int zeros[5] = {0};

/* Designated initializers (C99+) */
int array[10] = {[0] = 1, [5] = 10, [9] = 5};

/* Size inferred from initializer */
int primes[] = {2, 3, 5, 7, 11, 13};
```

Constant Arrays

```
const int DAYS_IN_MONTH[] = {
    31, 28, 31, 30, 31, 30, 31, 31, 30, 31, 30, 31
};
```

Access and Indexing

Indexing is zero-based. Valid indices for `int a[n]` are 0 through `n - 1`.

```
#include <stdio.h>

int main(void) {
    int numbers[5] = {10, 20, 30, 40, 50};

    printf("First: %d\n", numbers[0]);    /* 10 */
    printf("Third: %d\n", numbers[2]);    /* 30 */
    printf("Last:  %d\n", numbers[4]);    /* 50 */

    numbers[1] = 25;
    numbers[3] = numbers[0] + numbers[2]; /* 40 */
    return 0;
}
```

No Automatic Bounds Checking

C does **not** check indices. Out-of-range access is undefined behavior:

```
int array[5] = {1, 2, 3, 4, 5};
array[0] = 10;    /* OK */
array[4] = 50;    /* OK */
/* array[-1] = 5; */ /* UB - do not */
/* array[10] = 100; */ /* UB - do not */
```

Traversal

```
#include <stdio.h>

int main(void) {
    int numbers[5] = {1, 2, 3, 4, 5};
    int size = 5;
```

```

    for (int i = 0; i < size; i++) {
        printf("a[%d] = %d\n", i, numbers[i]);
    }

    for (int i = size - 1; i >= 0; i--) {
        printf("rev a[%d] = %d\n", i, numbers[i]);
    }
    return 0;
}

```

sizeof and Array Length

```

#include <stdio.h>

#define ARRAY_LEN(a) (sizeof(a) / sizeof((a)[0]))

int main(void) {
    int numbers[10];
    printf("len = %zu\n", ARRAY_LEN(numbers)); /* 10 */
    return 0;
}

```

Critical: ARRAY_LEN only works on a real array object, not on a pointer. After an array decays to a pointer (for example as a function parameter), `sizeof` yields the pointer size, not the array size. Always pass an explicit length.

Passing Arrays to Functions

When you pass an array, the callee receives a pointer to the first element. **Always pass the size** (or a sentinel convention you document).

Equivalent Parameter Declarations

These three mean the same thing for the callee — a pointer:

```

void f(int arr[], int n);
void f(int arr[10], int n); /* 10 is ignored by the compiler for type */
void f(int *arr, int n);

```

Full Program: Size Parameter Pattern

```
#include <stdio.h>

void print_ints(const int *arr, size_t n) {
    for (size_t i = 0; i < n; i++) {
        printf("%d%s", arr[i], (i + 1 < n) ? " " : "\n");
    }
}

int sum_ints(const int *arr, size_t n) {
    int s = 0;
    for (size_t i = 0; i < n; i++) {
        s += arr[i];
    }
    return s;
}

/* WRONG idea of "getting length inside the function" */
void wrong_len_demo(int arr[]) {
    /* sizeof(arr) is sizeof(int *), not the caller's array */
    printf("wrong \"length\" calculation: %zu\n",
        sizeof(arr) / sizeof(arr[0]));
}

int main(void) {
    int data[] = {3, 1, 4, 1, 5};
    size_t n = sizeof data / sizeof data[0];

    print_ints(data, n);
    printf("sum = %d\n", sum_ints(data, n));
    wrong_len_demo(data);
    return 0;
}

gcc -std=c17 -Wall -Wextra pass_array.c -o pass_array
./pass_array
```

Optional: Static Array Size Hints (C99+)

```
/* Documents that arr should point to at least n elements;
   some compilers can warn on misuse with the right flags. */
void process(size_t n, int arr[static n]);
```


Multidimensional Arrays and Row-Major Layout

2D Declaration

```
int matrix[3][4];    /* 3 rows, 4 columns */
int m[2][3] = {
    {1, 2, 3},
    {4, 5, 6}
};
```

Row-Major Storage

C stores multidimensional arrays in **row-major** order: all elements of row 0, then row 1, and so on. For `int a[R][C]`, element `a[i][j]` lives at offset:

```
index = i * C + j
```

in a flat view of `R * C` contiguous ints.

```
#include <stdio.h>

int main(void) {
    int a[2][3] = {
        {10, 20, 30},
        {40, 50, 60}
    };

    int *flat = &a[0][0];
    printf("a[1][2] via [i][j] = %d\n", a[1][2]);    /* 60 */
    printf("a[1][2] via flat   = %d\n", flat[1 * 3 + 2]); /* 60 */

    printf("memory order: ");
    for (int k = 0; k < 6; k++) {
        printf("%d ", flat[k]); /* 10 20 30 40 50 60 */
    }
    printf("\n");
    return 0;
}
```

Passing 2D Arrays

The callee needs the column count (or full type) so index arithmetic works:

```

#include <stdio.h>

#define COLS 3

void print_matrix(const int m[][COLS], int rows) {
    for (int i = 0; i < rows; i++) {
        for (int j = 0; j < COLS; j++) {
            printf("%4d", m[i][j]);
        }
        printf("\n");
    }
}

/* Flat form: treat as one vector with explicit rows/cols */
void print_flat(const int *m, int rows, int cols) {
    for (int i = 0; i < rows; i++) {
        for (int j = 0; j < cols; j++) {
            printf("%4d", m[i * cols + j]);
        }
        printf("\n");
    }
}

int main(void) {
    int m[2][COLS] = {{1, 2, 3}, {4, 5, 6}};
    print_matrix(m, 2);
    print_flat(&m[0][0], 2, COLS);
    return 0;
}

```

Worked Algorithms (Full Programs)

1. Reverse an Array In Place

```

#include <stdio.h>

static void reverse(int *a, size_t n) {
    for (size_t i = 0; i < n / 2; i++) {
        int tmp = a[i];
        a[i] = a[n - 1 - i];
        a[n - 1 - i] = tmp;
    }
}

```

```

static void print_arr(const int *a, size_t n) {
    for (size_t i = 0; i < n; i++) {
        printf("%d%s", a[i], (i + 1 < n) ? " " : "\n");
    }
}

int main(void) {
    int a[] = {1, 2, 3, 4, 5};
    size_t n = sizeof a / sizeof a[0];

    printf("before: ");
    print_arr(a, n);
    reverse(a, n);
    printf("after: ");
    print_arr(a, n);
    return 0;
}

```

2. Find Minimum and Maximum

```

#include <stdio.h>
#include <stdbool.h>

/* Returns false if n == 0; otherwise writes *min_out / *max_out. */
static bool minmax(const int *a, size_t n, int *min_out, int *max_out) {
    if (a == NULL || min_out == NULL || max_out == NULL || n == 0) {
        return false;
    }
    int lo = a[0];
    int hi = a[0];
    for (size_t i = 1; i < n; i++) {
        if (a[i] < lo) lo = a[i];
        if (a[i] > hi) hi = a[i];
    }
    *min_out = lo;
    *max_out = hi;
    return true;
}

int main(void) {
    int a[] = {7, -2, 15, 0, 9, 15, -8};
    size_t n = sizeof a / sizeof a[0];
    int lo, hi;

    if (!minmax(a, n, &lo, &hi)) {

```

```

        fprintf(stderr, "empty array\n");
        return 1;
    }
    printf("min = %d, max = %d\n", lo, hi);
    return 0;
}

```

3. Frequency Count

Count how often each value appears in a small range (here: digits 0–9).

```

#include <stdio.h>
#include <string.h>

enum { RANGE = 10 };

static void count_digits(const int *a, size_t n, int freq[RANGE]) {
    memset(freq, 0, RANGE * sizeof freq[0]);
    for (size_t i = 0; i < n; i++) {
        int v = a[i];
        if (v >= 0 && v < RANGE) {
            freq[v]++;
        }
    }
}

int main(void) {
    int data[] = {1, 3, 1, 7, 3, 3, 9, 0, 1};
    size_t n = sizeof data / sizeof data[0];
    int freq[RANGE];

    count_digits(data, n, freq);

    printf("value : count\n");
    for (int v = 0; v < RANGE; v++) {
        if (freq[v] > 0) {
            printf("  %d : %d\n", v, freq[v]);
        }
    }
    return 0;
}

```

4. Small Matrix Multiply ($2 \times 3 \cdot 3 \times 2 \rightarrow 2 \times 2$)

```
#include <stdio.h>

enum { A_ROWS = 2, A_COLS = 3, B_COLS = 2 };

/*
 * C = A * B
 * A is (A_ROWS x A_COLS), B is (A_COLS x B_COLS), C is (A_ROWS x B_COLS)
 * Row-major: C[i][j] = sum_k A[i][k] * B[k][j]
 */
static void matmul(const int A[A_ROWS][A_COLS],
                  const int B[A_COLS][B_COLS],
                  int C[A_ROWS][B_COLS]) {
    for (int i = 0; i < A_ROWS; i++) {
        for (int j = 0; j < B_COLS; j++) {
            int sum = 0;
            for (int k = 0; k < A_COLS; k++) {
                sum += A[i][k] * B[k][j];
            }
            C[i][j] = sum;
        }
    }
}

static void print_2x2(const int M[A_ROWS][B_COLS]) {
    for (int i = 0; i < A_ROWS; i++) {
        for (int j = 0; j < B_COLS; j++) {
            printf("%4d", M[i][j]);
        }
        printf("\n");
    }
}

int main(void) {
    int A[A_ROWS][A_COLS] = {
        {1, 2, 3},
        {4, 5, 6}
    };
    int B[A_COLS][B_COLS] = {
        {7, 8},
        {9, 10},
        {11, 12}
    };
    int C[A_ROWS][B_COLS];

    matmul(A, B, C);
```

```

    printf("A * B =\n");
    print_2x2(C);
    /* Expected:
       58  64
       139 154
    */
    return 0;
}

gcc -std=c17 -Wall -Wextra matmul.c -o matmul
./matmul

```

5. Histogram of Grades

Bucket scores into letter-grade ranges and print a text histogram.

```

#include <stdio.h>
#include <string.h>

enum { N_BUCKETS = 5 };

/* 0: F <60, 1: D 60-69, 2: C 70-79, 3: B 80-89, 4: A 90-100 */
static int grade_bucket(int score) {
    if (score < 60) return 0;
    if (score < 70) return 1;
    if (score < 80) return 2;
    if (score < 90) return 3;
    return 4;
}

static void histogram(const int *scores, size_t n, int buckets[N_BUCKETS]) {
    memset(buckets, 0, N_BUCKETS * sizeof buckets[0]);
    for (size_t i = 0; i < n; i++) {
        int s = scores[i];
        if (s < 0 || s > 100) {
            continue; /* skip invalid */
        }
        buckets[grade_bucket(s)]++;
    }
}

static void print_bar(const char *label, int count) {
    printf("%s (%2d): ", label, count);
    for (int i = 0; i < count; i++) {
        putchar('#');
    }
}

```

```

    }
    putchar('\n');
}

int main(void) {
    int scores[] = {
        95, 82, 77, 61, 54, 88, 91, 73, 66, 100,
        59, 84, 70, 45, 92
    };
    size_t n = sizeof scores / sizeof scores[0];
    int buckets[N_BUCKETS];
    const char *labels[N_BUCKETS] = {"F", "D", "C", "B", "A"};

    histogram(scores, n, buckets);

    printf("Grade histogram (%zu scores)\n", n);
    for (int i = 0; i < N_BUCKETS; i++) {
        print_bar(labels[i], buckets[i]);
    }
    return 0;
}

```

6. Linear Search and Selection Sort (Together)

```

#include <stdio.h>

static int linear_search(const int *a, size_t n, int target) {
    for (size_t i = 0; i < n; i++) {
        if (a[i] == target) {
            return (int)i;
        }
    }
    return -1;
}

static void selection_sort(int *a, size_t n) {
    for (size_t i = 0; i + 1 < n; i++) {
        size_t min_i = i;
        for (size_t j = i + 1; j < n; j++) {
            if (a[j] < a[min_i]) {
                min_i = j;
            }
        }
        if (min_i != i) {
            int t = a[i];

```

```

        a[i] = a[min_i];
        a[min_i] = t;
    }
}

static void print_arr(const int *a, size_t n) {
    for (size_t i = 0; i < n; i++) {
        printf("%d%s", a[i], (i + 1 < n) ? " " : "\n");
    }
}

int main(void) {
    int a[] = {64, 25, 12, 22, 11};
    size_t n = sizeof a / sizeof a[0];

    printf("index of 22: %d\n", linear_search(a, n, 22));
    selection_sort(a, n);
    printf("sorted: ");
    print_arr(a, n);
    return 0;
}

```

Common Bugs: Wrong vs Right

Off-by-One in Loops

```

#include <stdio.h>

int main(void) {
    int a[5] = {1, 2, 3, 4, 5};

    /* WRONG: i <= 5 reads a[5] - out of bounds */
    /* for (int i = 0; i <= 5; i++) printf("%d\n", a[i]); */

    /* RIGHT: i < n */
    for (int i = 0; i < 5; i++) {
        printf("%d\n", a[i]);
    }
    return 0;
}

```


Forgetting That Arrays Decay

```
#include <stdio.h>

/* WRONG: cannot recover length from decayed pointer */
static void bad_average(int arr[]) {
    size_t n = sizeof(arr) / sizeof(arr[0]); /* pointer size! */
    int sum = 0;
    for (size_t i = 0; i < n; i++) {
        sum += arr[i]; /* may read garbage / crash */
    }
    printf("bad avg path used n=%zu\n", n);
    (void)sum;
}

/* RIGHT */
static double good_average(const int *arr, size_t n) {
    if (n == 0) {
        return 0.0;
    }
    long sum = 0;
    for (size_t i = 0; i < n; i++) {
        sum += arr[i];
    }
    return (double)sum / (double)n;
}

int main(void) {
    int a[] = {10, 20, 30, 40};
    size_t n = sizeof a / sizeof a[0];
    printf("avg = %.2f\n", good_average(a, n));
    bad_average(a); /* educational only - do not copy this pattern */
    return 0;
}
```

Using Uninitialized Local Arrays

```
/* WRONG */
int buf[10];
printf("%d\n", buf[0]); /* indeterminate value */

/* RIGHT */
int buf[10] = {0};
/* or fill before read */
```

Copying with =

```
int a[3] = {1, 2, 3};
int b[3];
/* b = a; */ /* ILLEGAL - arrays are not assignable */

/* RIGHT */
for (int i = 0; i < 3; i++) {
    b[i] = a[i];
}
/* or memcpy(b, a, sizeof a); with <string.h> */
```

Returning a Pointer to a Local Array

```
/* WRONG */
int *bad(void) {
    int local[4] = {1, 2, 3, 4};
    return local; /* dangling - storage ends when function returns */
}

/* RIGHT: caller owns the buffer */
void good(int *out, size_t n) {
    for (size_t i = 0; i < n; i++) {
        out[i] = (int)(i + 1);
    }
}
```

Practical Examples

Temperature Analysis

```
#include <stdio.h>

#define DAYS 7

int main(void) {
    float temperatures[DAYS] = {18.5f, 20.0f, 19.2f, 21.5f, 17.8f, 16.0f, 22.1f};

    float sum = 0.0f;
    float max = temperatures[0];
    float min = temperatures[0];

    for (int i = 0; i < DAYS; i++) {
```

```

        sum += temperatures[i];
        if (temperatures[i] > max) max = temperatures[i];
        if (temperatures[i] < min) min = temperatures[i];
    }

    printf("Average: %.2f\n", sum / DAYS);
    printf("Maximum: %.2f\n", max);
    printf("Minimum: %.2f\n", min);
    return 0;
}

```

Student Scores (Parallel Arrays)

```

#include <stdio.h>

#define STUDENTS 4

int main(void) {
    const char *names[STUDENTS] = {"Ada", "Linus", "Grace", "Ken"};
    float grades[STUDENTS] = {94.0f, 88.5f, 97.0f, 91.0f};

    int top = 0;
    float sum = 0.0f;
    for (int i = 0; i < STUDENTS; i++) {
        sum += grades[i];
        if (grades[i] > grades[top]) {
            top = i;
        }
    }

    printf("Top: %s (%.1f)\n", names[top], grades[top]);
    printf("Average: %.2f\n", sum / STUDENTS);
    return 0;
}

```

Mini-Lab: Array Toolkit

Implement a single file `array_lab.c` that:

1. Reads up to 32 integers (first number is `n`, then `n` values), **or** uses a built-in sample if you prefer non-interactive code.
2. Prints the array.
3. Prints min, max, and sum.
4. Prints a frequency table for values in 0..9 (ignore others).

5. Reverses a copy of the array and prints it (leave the original intact).

Skeleton:

```
#include <stdio.h>
#include <string.h>

enum { MAX_N = 32 };

static void print_arr(const int *a, int n) {
    for (int i = 0; i < n; i++) {
        printf("%d%s", a[i], (i + 1 < n) ? " " : "\n");
    }
}

/* Exercise: implement minmax, sum, freq, reverse_into */

int main(void) {
    int sample[] = {1, 4, 1, 5, 9, 2, 6, 5};
    int n = (int)(sizeof sample / sizeof sample[0]);
    int copy[MAX_N];

    printf("data: ");
    print_arr(sample, n);

    /* Fill in the rest for the lab. */
    (void)copy;
    return 0;
}

gcc -std=c17 -Wall -Wextra array_lab.c -o array_lab
./array_lab
```

Exercises

1. **Rotate left by k.** Given an array and k, rotate elements left by k positions in place (or via a temp buffer). Test with $k = 0$, $k = n$, and $k > n$ (normalize with $k \% n$).
2. **Second largest.** Return the second-largest distinct value, or document a clear error if it does not exist.
3. **Merge sorted arrays.** Given two sorted int arrays, produce a third sorted array of length $n + m$.
4. **Saddle point.** In a small matrix, find an element that is the minimum of its row and maximum of its column (if any).

5. **Sparse frequency.** Count frequencies without assuming values are in `0..9` — use a parallel “value / count” table or sort-then-scan.
6. **Bounds-safe wrapper.** Write `int at(const int *a, size_t n, size_t i, int *out)` that returns 0 on success and -1 if `i >= n`, never reading out of range.
7. **Row and column sums.** For a fixed `R×C` matrix, print each row sum and each column sum.
8. **Mini-lab completion.** Finish `array_lab.c` from above and add selection sort on a copy.

Summary

1. Arrays are contiguous same-type storage with zero-based indices.
2. C does not bounds-check — you own every index.
3. Arrays decay to pointers at function boundaries; **pass the length**.
4. Multidimensional arrays are row-major; `a[i][j]` sits at `i * cols + j` in flat form.
5. Classic operations — reverse, min/max, frequency, matrix multiply, histograms — are loops plus careful indexing.
6. Common bugs: off-by-one, `sizeof` on parameters, uninitialized data, illegal array assignment, dangling returns.

Solid array discipline is the base for strings, pointers, and almost every algorithm you will write in C.

Strings

Introduction

In C, a string is a contiguous sequence of characters ending with a null terminator `'\0'`. There is no built-in string type: you work with `char` arrays and `char *`, plus the helpers in `<string.h>` and careful buffer management.

This chapter covers representation, standard functions, **safe** copy/format patterns, DIY `strlen` / `strcpy` / `strcmp`, CSV parsing with `strtok_r` and manual scanners, buffer overflow awareness with safe alternatives, and exercises.

Compile with:

```
gcc -std=c17 -Wall -Wextra program.c -o program
```

String Representation

Character Arrays and Literals

```
char str1[6] = {'H', 'e', 'l', 'l', 'o', '\0'};
char str2[6] = "Hello"; /* same content; size includes '\0' */
char str3[] = "Hello"; /* size 6 inferred */

char *lit = "Hello"; /* pointer to string literal - treat as read-only */
char buf[] = "Hello"; /* mutable copy in your array */
```

Memory Layout

```
char s[] = "Hi";

index:  0    1    2

        'H'  'i'  '\0'

length strlen(s) == 2
storage size sizeof s == 3
```

Literals Must Not Be Modified

```
char *p = "Hello";
/* p[0] = 'h'; */ /* undefined behavior - may crash */

char a[] = "Hello";
a[0] = 'h'; /* OK - a is a mutable array */
```

Prefer `const char *p = "Hello";` so the compiler helps you avoid writes.

Declaration Patterns

```
char name[64] = "Ada"; /* modifiable, fixed capacity */
char empty[32] = ""; /* just '\0' */
char raw[32]; /* uninitialized - do not read as a string yet */
char *msg = "ready"; /* points at literal */

#include <string.h>
memset(raw, 0, sizeof raw); /* explicit clear */
```

Standard Library Overview

```
#include <string.h>

size_t strlen(const char *s);
char *strcpy(char *dst, const char *src); /* unsafe if dst too small */
char *strncpy(char *dst, const char *src, size_t n); /* subtle; see below */
int strcmp(const char *a, const char *b);
int strncmp(const char *a, const char *b, size_t n);
char *strcat(char *dst, const char *src); /* unsafe if dst too small */
char *strncat(char *dst, const char *src, size_t n);
char *strchr(const char *s, int c);
char *strstr(const char *haystack, const char *needle);
```

Also useful: `snprintf` from `<stdio.h>` for bounded formatting into a buffer.

Safe String Handling Patterns

Prefer `snprintf` over bare `strcpy` / `strcat`

`snprintf` always null-terminates when the buffer size is greater than zero, and it tells you how many characters would have been written.

```

#include <stdio.h>
#include <string.h>

int main(void) {
    char buf[16];
    const char *user = "Ada Lovelace";

    int n = snprintf(buf, sizeof buf, "hi %s", user);
    if (n < 0) {
        /* encoding error (rare) */
        return 1;
    }
    if ((size_t)n >= sizeof buf) {
        /* truncated - buf is still a valid C string */
        printf("truncated: \"%s\" (needed %d bytes incl. nul logic)\n", buf, n + 1);
    } else {
        printf("full: \"%s\"\n", buf);
    }
    return 0;
}

gcc -std=c17 -Wall -Wextra snip.c -o snip
./snip

```

Building Paths or Messages Safely

```

#include <stdio.h>

static int join3(char *out, size_t out_sz,
                 const char *a, const char *b, const char *c) {
    int n = snprintf(out, out_sz, "%s/%s/%s", a, b, c);
    if (n < 0) {
        return -1;
    }
    if ((size_t)n >= out_sz) {
        return -2; /* truncated */
    }
    return 0;
}

int main(void) {
    char path[32];
    if (join3(path, sizeof path, "home", "ada", "notes.txt") == 0) {
        printf("%s\n", path);
    }
}

```



```

    return 0;
}

```

strncpy vs snprintf — Know the Trap

strncpy is **not** a “safe strcpy.” It pads with zeros up to `n`, and if `src` length is `n`, it **does not** write a `'\0'`.

```

#include <stdio.h>
#include <string.h>

int main(void) {
    char dest[8];
    const char *src = "OVERFLOW!"; /* longer than dest */

    /* TRAP: dest may be not null-terminated */
    strncpy(dest, src, sizeof dest);
    /* dest has 8 chars, no '\0' - using it as a string is UB */

    /* Repair pattern if you must use strncpy: */
    dest[sizeof dest - 1] = '\0';

    /* BETTER: */
    char dest2[8];
    snprintf(dest2, sizeof dest2, "%s", src); /* always terminated if size > 0 */

    printf("dest2 = \"%s\"\n", dest2);
    return 0;
}

```

Bounded Copy Helper (Recommended Pattern)

```

#include <stdio.h>
#include <string.h>

/* Returns 0 on full copy, -1 on truncation or bad args. Always terminates if dst_sz > 0. */
static int str_copy(char *dst, size_t dst_sz, const char *src) {
    if (dst == NULL || src == NULL || dst_sz == 0) {
        return -1;
    }
    int n = snprintf(dst, dst_sz, "%s", src);
    if (n < 0 || (size_t)n >= dst_sz) {
        return -1;
    }
}

```

```

    return 0;
}

int main(void) {
    char b[8];
    if (str_copy(b, sizeof b, "hello") == 0) {
        printf("ok: %s\n", b);
    }
    if (str_copy(b, sizeof b, "toolong!") != 0) {
        printf("truncated-or-fail: \"%s\"\n", b);
    }
    return 0;
}

```

Safe Concatenation

```

#include <stdio.h>
#include <string.h>

static int str_append(char *dst, size_t dst_sz, const char *src) {
    if (dst == NULL || src == NULL || dst_sz == 0) {
        return -1;
    }
    size_t used = strlen(dst);
    if (used >= dst_sz) {
        return -1; /* dst already not a valid sized string */
    }
    int n = snprintf(dst + used, dst_sz - used, "%s", src);
    if (n < 0 || (size_t)n >= dst_sz - used) {
        return -1;
    }
    return 0;
}

int main(void) {
    char msg[32] = "Hello";
    str_append(msg, sizeof msg, ", ");
    str_append(msg, sizeof msg, "world");
    printf("%s\n", msg);
    return 0;
}

```

Input: Prefer fgets + Strip Newline

```
#include <stdio.h>
#include <string.h>

static void strip_newline(char *s) {
    size_t n = strlen(s);
    if (n > 0 && s[n - 1] == '\n') {
        s[n - 1] = '\0';
    }
}

int main(void) {
    char line[64];

    /* BAD: scanf("%s") - no bounds unless you set a width */
    /* BAD: gets - removed from C; never use */

    printf("name: ");
    if (fgets(line, sizeof line, stdin) == NULL) {
        return 1;
    }
    strip_newline(line);
    printf("hello, %s\n", line);
    return 0;
}
```

If you use `scanf` for a single token, always set a maximum field width: `scanf("%63s", line)` for a 64-byte buffer.

Implement strlen / strcpy / strcmp Yourself

Writing these clarifies how C strings work. Production code should still prefer the standard library and bounded helpers.

Full Learning Program

```
#include <stdio.h>
#include <assert.h>

static size_t my_strlen(const char *s) {
    const char *p = s;
    while (*p != '\0') {
        p++;
    }
}
```

```

    }
    return (size_t)(p - s);
}

/* Classic unbounded copy - destination must be large enough (like strcpy). */
static char *my_strcpy(char *dst, const char *src) {
    char *out = dst;
    while ((*dst++ = *src++) != '\0') {
        /* copy including '\0' */
    }
    return out;
}

/* Like strcmp: <0, 0, >0 */
static int my_strcmp(const char *a, const char *b) {
    while (*a != '\0' && *a == *b) {
        a++;
        b++;
    }
    return (unsigned char)*a - (unsigned char)*b;
}

/* Bounded copy: always writes '\0' if dst_sz > 0. Returns total length of src
   (as snprintf does conceptually) so caller can detect truncation. */
static size_t my_strncpy(char *dst, const char *src, size_t dst_sz) {
    size_t src_len = my_strlen(src);
    if (dst_sz == 0) {
        return src_len;
    }
    size_t copy = (src_len < dst_sz - 1) ? src_len : dst_sz - 1;
    for (size_t i = 0; i < copy; i++) {
        dst[i] = src[i];
    }
    dst[copy] = '\0';
    return src_len;
}

int main(void) {
    const char *s = "Hello";
    assert(my_strlen(s) == 5);
    assert(my_strlen("") == 0);

    char buf[16];
    my_strcpy(buf, s);
    assert(my_strcmp(buf, "Hello") == 0);
    assert(my_strcmp("abc", "abd") < 0);
}

```

```

assert(my_strncmp("abd", "abc") > 0);

char small[4];
size_t need = my_strncpy(small, "Hello", sizeof small);
printf("small=\"%s\" need=%zu truncated=%s\n",
      small, need, (need >= sizeof small) ? "yes" : "no");

printf("all my_string tests passed\n");
return 0;
}

```

```

gcc -std=c17 -Wall -Wextra my_string.c -o my_string
./my_string

```

Notes:

- Cast to unsigned char in compare so high-bit characters sort consistently.
- Unbounded my_strcpy is for learning; use my_strncpy / snprintf in real code.

Tokenizing and CSV Parsing

strtok vs strtok_r

strtok keeps hidden static state — **not reentrant**, awkward in libraries and threads. Prefer **strtok_r** (POSIX) with an explicit context pointer, or parse manually for full control.

Full Program: CSV Line with strtok_r

```

#include <stdio.h>
#include <string.h>

enum { MAX_FIELDS = 16, FIELD_LEN = 64 };

/*
 * Split a mutable CSV line on commas into fields[][].
 * Note: strtok_r mutates the line (writes '\0' over delimiters).
 * Returns field count, or -1 on overflow of MAX_FIELDS.
 */
static int split_csv_strtok_r(char *line, char fields[][FIELD_LEN], int max_fields) {
    int count = 0;
    char *save = NULL;
    char *tok = strtok_r(line, ",", &save);

    while (tok != NULL) {

```

```

        if (count >= max_fields) {
            return -1;
        }
        /* Bounded copy into fixed field buffer */
        snprintf(fields[count], FIELD_LEN, "%s", tok);
        count++;
        tok = strtok_r(NULL, ",", &save);
    }
    return count;
}

int main(void) {
    char line[] = "ada,lovelace,36,math"; /* mutable */
    char fields[MAX_FIELDS][FIELD_LEN];

    int n = split_csv_strtok_r(line, fields, MAX_FIELDS);
    if (n < 0) {
        fprintf(stderr, "too many fields\n");
        return 1;
    }

    for (int i = 0; i < n; i++) {
        printf("field[%d] = \"%s\"\n", i, fields[i]);
    }
    /* Warning: original line is now shredded into tokens */
    return 0;
}

```

```

gcc -std=c17 -Wall -Wextra csv_strtok.c -o csv_strtok
./csv_strtok

```

On platforms without `strtok_r`, use the manual parser below (portable C17).

Full Program: Manual CSV Parse (No `strtok`)

Does not need POSIX; leaves the original string intact if you copy fields out.

```

#include <stdio.h>
#include <string.h>
#include <ctype.h>

enum { MAX_FIELDS = 16, FIELD_LEN = 64 };

static void trim_inplace(char *s) {
    /* leading */
    char *start = s;

```

```

while (*start != '\0' && isspace((unsigned char)*start)) {
    start++;
}
if (start != s) {
    memmove(s, start, strlen(start) + 1);
}
/* trailing */
size_t n = strlen(s);
while (n > 0 && isspace((unsigned char)s[n - 1])) {
    s[--n] = '\0';
}
}

/*
 * Parse comma-separated fields from src into fields.
 * Returns count or -1 if a field would not fit FIELD_LEN or too many fields.
 */
static int split_csv_manual(const char *src,
                           char fields[][FIELD_LEN],
                           int max_fields) {

    int count = 0;
    const char *p = src;

    while (*p != '\0') {
        if (count >= max_fields) {
            return -1;
        }
        /* skip leading spaces in field */
        while (*p != '\0' && *p != ',' && isspace((unsigned char)*p)) {
            p++;
        }
        const char *start = p;
        while (*p != '\0' && *p != ',') {
            p++;
        }
        size_t len = (size_t)(p - start);
        /* trim trailing spaces from the slice */
        while (len > 0 && isspace((unsigned char)start[len - 1])) {
            len--;
        }
        if (len + 1 > FIELD_LEN) {
            return -1;
        }
        memcpy(fields[count], start, len);
        fields[count][len] = '\0';
        count++;
    }
}

```

```

        if (*p == ',') {
            p++; /* consume comma */
        }
    }
    return count;
}

int main(void) {
    const char *line = " ada , lovelace , 36 , math ";
    char fields[MAX_FIELDS][FIELD_LEN];

    int n = split_csv_manual(line, fields, MAX_FIELDS);
    if (n < 0) {
        fprintf(stderr, "parse error\n");
        return 1;
    }
    printf("original still: \"%s\"\n", line);
    for (int i = 0; i < n; i++) {
        printf("field[%d] = \"%s\"\n", i, fields[i]);
    }
    (void)trim_inplace; /* available for whole-string trimming if needed */
    return 0;
}

gcc -std=c17 -Wall -Wextra csv_manual.c -o csv_manual
./csv_manual

```

Simple Key=Value Line

```

#include <stdio.h>
#include <string.h>

/* Parse "key=value" into key/value buffers. Returns 0 on success. */
static int parse_kv(const char *line,
                    char *key, size_t key_sz,
                    char *val, size_t val_sz) {
    const char *eq = strchr(line, '=');
    if (eq == NULL) {
        return -1;
    }
    size_t klen = (size_t)(eq - line);
    if (klen + 1 > key_sz) {
        return -1;
    }
    memcpy(key, line, klen);

```



```

    key[klen] = '\0';

    if (snprintf(val, val_sz, "%s", eq + 1) < 0) {
        return -1;
    }
    if (strlen(eq + 1) >= val_sz) {
        return -1; /* truncated */
    }
    return 0;
}

int main(void) {
    char k[32], v[64];
    if (parse_kv("port=8080", k, sizeof k, v, sizeof v) == 0) {
        printf("key=\"%s\" val=\"%s\"\n", k, v);
    }
    return 0;
}

```

Buffer Overflows: Awareness and Safe Alternatives

A **buffer overflow** happens when code writes past the end of an array. With strings, the usual causes are unbounded copies, missing null terminators, and `scanf("%s")` without a width.

This section is **defensive**. Do not use these ideas to attack systems. The goal is to recognize fragile patterns and replace them with bounded APIs.

Fragile Pattern (Do Not Ship)

```

/* ILLUSTRATION ONLY - unbounded write into a fixed buffer */
void fragile_copy(char *dst, const char *src) {
    while (*src) {
        *dst++ = *src++; /* no length check */
    }
    *dst = '\0';
}

```

If `src` is longer than the space behind `dst`, memory after the buffer is overwritten. That can crash the program or corrupt other data. Treat it as a bug to eliminate.

Safe Alternatives Checklist

Fragile	Prefer
gets	fgets(buf, sizeof buf, stdin)
scanf("%s", buf)	scanf("%63s", buf) or fgets
strcpy(dst, src)	snprintf(dst, dst_sz, "%s", src)
strcat(dst, src)	str_append helper with remaining size
strncpy without forced '\0'	snprintf or copy + dst[sz-1]='\0'
sprintf(dst, ...)	snprintf(dst, dst_sz, ...)
DIY loops without n	index loops with $i + 1 < sz$

Full Program: Safe Line Reader and Copy

```
#include <stdio.h>
#include <string.h>

enum { LINE_MAX = 128 };

static int read_line(char *buf, size_t buf_sz) {
    if (buf == NULL || buf_sz == 0) {
        return -1;
    }
    if (fgets(buf, (int)buf_sz, stdin) == NULL) {
        return -1;
    }
    size_t n = strlen(buf);
    if (n > 0 && buf[n - 1] == '\n') {
        buf[n - 1] = '\0';
    } else {
        /* Line may have been longer than buffer - drain rest of line */
        int c;
        int overflow = (n == buf_sz - 1);
        if (overflow) {
            while ((c = getchar()) != '\n' && c != EOF) {
                /* discard */
            }
            return -2; /* truncated */
        }
    }
    return 0;
}

int main(void) {
    char name[LINE_MAX];
    char greeting[LINE_MAX];

    printf("Enter your name: ");
```

```

int rc = read_line(name, sizeof name);
if (rc == -1) {
    fprintf(stderr, "no input\n");
    return 1;
}
if (rc == -2) {
    fprintf(stderr, "input too long (truncated)\n");
}

if (snprintf(greeting, sizeof greeting, "Hello, %s!", name) < 0) {
    return 1;
}
printf("%s\n", greeting);
return 0;
}

```

Compiler Help

```
gcc -std=c17 -Wall -Wextra -Wformat -Wformat-security program.c -o program
```

Some platforms also offer fortified library functions when optimization is on (`_FORTIFY_SOURCE`). Still write correct sizes yourself — fortification is a backstop, not a design.

More String Techniques

Remove Newline, Reverse, Case

```

#include <stdio.h>
#include <string.h>
#include <ctype.h>

static void strip_newline(char *s) {
    size_t n = strlen(s);
    if (n > 0 && s[n - 1] == '\n') {
        s[n - 1] = '\0';
    }
}

static void reverse_inplace(char *s) {
    size_t n = strlen(s);
    for (size_t i = 0; i < n / 2; i++) {
        char t = s[i];
        s[i] = s[n - 1 - i];
    }
}

```

```

        s[n - 1 - i] = t;
    }
}

static void to_lower_inplace(char *s) {
    for (; *s; s++) {
        *s = (char)tolower((unsigned char)*s);
    }
}

int main(void) {
    char s[] = "AbCdEf";
    to_lower_inplace(s);
    reverse_inplace(s);
    printf("%s\n", s); /* fedcba */
    (void)strip_newline;
    return 0;
}

```

Palindrome Check

```

#include <stdio.h>
#include <string.h>
#include <ctype.h>

static int is_palindrome(const char *str) {
    size_t i = 0;
    size_t j = strlen(str);
    if (j == 0) {
        return 1;
    }
    j--;
    while (i < j) {
        while (i < j && !isalnum((unsigned char)str[i])) i++;
        while (i < j && !isalnum((unsigned char)str[j])) j--;
        if (tolower((unsigned char)str[i]) != tolower((unsigned char)str[j])) {
            return 0;
        }
        i++;
        j--;
    }
    return 1;
}

int main(void) {

```

```

    printf("%d\n", is_palindrome("A man, a plan, a canal: Panama")); /* 1 */
    printf("%d\n", is_palindrome("hello")); /* 0 */
    return 0;
}

```

Word Count

```

#include <stdio.h>
#include <ctype.h>

static int count_words(const char *str) {
    int count = 0;
    int in_word = 0;
    for (size_t i = 0; str[i] != '\0'; i++) {
        if (isspace((unsigned char)str[i])) {
            in_word = 0;
        } else if (!in_word) {
            in_word = 1;
            count++;
        }
    }
    return count;
}

int main(void) {
    printf("%d\n", count_words("  hello  c strings ")); /* 3 */
    return 0;
}

```

Character Class Statistics

```

#include <stdio.h>
#include <ctype.h>
#include <string.h>

static void analyze(const char *str) {
    size_t letters = 0, digits = 0, spaces = 0, others = 0;
    for (size_t i = 0; str[i] != '\0'; i++) {
        unsigned char c = (unsigned char)str[i];
        if (isalpha(c)) letters++;
        else if (isdigit(c)) digits++;
        else if (isspace(c)) spaces++;
        else others++;
    }
}

```

```

        printf("letters=%zu digits=%zu spaces=%zu others=%zu len=%zu\n",
               letters, digits, spaces, others, strlen(str));
    }

    int main(void) {
        analyze("Hello, C17! 42");
        return 0;
    }

```

Pitfalls Recap

1. **Forgetting '\0'** after manual copies.
2. **sizeof on char *** — that is pointer size, not string length.
3. **Modifying string literals.**
4. **strncpy without terminator.**
5. **Using strtok in concurrent or nested tokenization** — prefer `strtok_r` or manual parse.
6. **Assuming `scanf("%s")` is safe** without a width.

Exercises

1. **Bounded library.** Implement `str_copy`, `str_append`, and `str_eq` (boolean compare) with full size parameters and tests.
2. **DIY suite.** Write `my_strlen`, `my_strcmp`, `my_strncmp`, and `my_strlcpy`. Compare results against the C library for a table of sample strings.
3. **CSV grades.** Parse lines like `name,score` from a hardcoded multi-line string; compute average score. Use the manual splitter (no `strtok`).
4. **strtok_r word counter.** Tokenize on whitespace with `strtok_r` and count tokens. Then reimplement without mutating a second copy.
5. **Safe login prompt.** Read username and password with `fgets` into fixed buffers; reject truncated lines; never use `gets` or unbounded `scanf`.
6. **URL path join.** `join_url(char *out, size_t n, const char *base, const char *path)` using `snprintf`, handling trailing/leading slashes cleanly.
7. **Find and replace first occurrence.** Replace the first occurrence of a substring with another (same length for simplicity, or build into a second buffer with size checks).
8. **Overflow hunt.** Take a small program that uses `strcpy`/`sprintf` and rewrite every call site to a bounded form. List each change in comments.

Summary

1. C strings are **char sequences with '\0'** — capacity and length are not the same.
2. Prefer **snprintf**, **fgets**, and explicit sizes over **strcpy** / **gets** / bare **scanf("%s")**.
3. **strncpy** is easy to misuse; force termination or switch to **snprintf**.
4. Implementing **strlen** / **strcpy** / **strcmp** teaches the model; ship bounded wrappers.
5. **strtok_r** or **manual parsing** beats classic **strtok** for reentrancy and clarity.
6. Treat buffer overflows as **bugs to prevent** with correct sizes and safer APIs — never as something to exploit.

With disciplined buffer handling, C string code can be both efficient and robust.

Advanced Arrays

Introduction

Advanced array concepts in C extend beyond basic one-dimensional arrays to include multi-dimensional arrays, variable-length arrays, and specialized array techniques. These concepts are essential for handling complex data structures and implementing efficient algorithms. Understanding advanced arrays enables programmers to work with matrices, tables, and other multi-dimensional data representations.

Multi-dimensional Arrays

Two-dimensional Arrays

Two-dimensional arrays are essentially arrays of arrays, often used to represent matrices or tables:

```
// Declaration
int matrix[3][4]; // 3 rows, 4 columns

// Initialization
int matrix1[2][3] = {{1, 2, 3}, {4, 5, 6}};
int matrix2[2][3] = {1, 2, 3, 4, 5, 6}; // Equivalent to above

// Partial initialization
int matrix3[3][3] = {{1, 2}, {4, 5}}; // Remaining elements set to 0

// Designated initializers (C99)
int matrix4[3][3] = {[0][0] = 1, [1][1] = 5, [2][2] = 9};
```

Accessing Multi-dimensional Arrays

```
int matrix[3][4] = {{1, 2, 3, 4}, {5, 6, 7, 8}, {9, 10, 11, 12}};

// Access individual elements
int element = matrix[1][2]; // Element at row 1, column 2 (value: 7)

// Modify elements
matrix[0][0] = 100;
```



```

// Traverse the array
for (int i = 0; i < 3; i++) {
    for (int j = 0; j < 4; j++) {
        printf("%d ", matrix[i][j]);
    }
    printf("\n");
}

```

Multi-dimensional Array Functions

When passing multi-dimensional arrays to functions, all dimensions except the first must be specified:

```

// Function that takes a 3x4 matrix
void print_matrix(int matrix[][4], int rows) {
    for (int i = 0; i < rows; i++) {
        for (int j = 0; j < 4; j++) {
            printf("%d ", matrix[i][j]);
        }
        printf("\n");
    }
}

// Function with pointer notation (equivalent)
void print_matrix_ptr(int (*matrix)[4], int rows) {
    for (int i = 0; i < rows; i++) {
        for (int j = 0; j < 4; j++) {
            printf("%d ", matrix[i][j]);
        }
        printf("\n");
    }
}

// Usage
int matrix[3][4] = {{1, 2, 3, 4}, {5, 6, 7, 8}, {9, 10, 11, 12}};
print_matrix(matrix, 3);

```

Three-dimensional and Higher-dimensional Arrays

Three-dimensional Arrays

```
// Declaration
int cube[2][3][4]; // 2 layers, 3 rows, 4 columns

// Initialization
int cube1[2][2][2] = {
    {{1, 2}, {3, 4}},
    {{5, 6}, {7, 8}}
};

// Accessing elements
int element = cube1[1][0][1]; // Value: 6

// Traversal
for (int i = 0; i < 2; i++) {
    for (int j = 0; j < 2; j++) {
        for (int k = 0; k < 2; k++) {
            printf("cube[%d][%d][%d] = %d\n", i, j, k, cube1[i][j][k]);
        }
    }
}
```

Higher-dimensional Arrays

While C supports arrays of any dimension, higher-dimensional arrays become increasingly complex:

```
// Four-dimensional array
int hypercube[2][2][2][2];

// Accessing elements
hypercube[0][1][0][1] = 42;

// Generally, it's better to use structures or dynamic allocation for complex multi-dimensional arrays
```

Variable-length Arrays (VLAs) - C99

Basic VLA Declaration

Variable-length arrays allow the size to be determined at runtime:

```

#include <stdio.h>

int main() {
    int size;
    printf("Enter array size: ");
    scanf("%d", &size);

    // Variable-length array
    int array[size]; // Size determined at runtime

    // Initialize and use the array
    for (int i = 0; i < size; i++) {
        array[i] = i * i;
        printf("array[%d] = %d\n", i, array[i]);
    }

    return 0;
}

```

Multi-dimensional VLAs

```

#include <stdio.h>

int main() {
    int rows, cols;
    printf("Enter rows and columns: ");
    scanf("%d %d", &rows, &cols);

    // Multi-dimensional VLA
    int matrix[rows][cols];

    // Initialize the matrix
    for (int i = 0; i < rows; i++) {
        for (int j = 0; j < cols; j++) {
            matrix[i][j] = i * cols + j;
        }
    }

    // Print the matrix
    for (int i = 0; i < rows; i++) {
        for (int j = 0; j < cols; j++) {
            printf("%d ", matrix[i][j]);
        }
        printf("\n");
    }
}

```

```

    return 0;
}

```

VLA Function Parameters

```

// Function that accepts VLA parameters
void process_matrix(int rows, int cols, int matrix[rows][cols]) {
    for (int i = 0; i < rows; i++) {
        for (int j = 0; j < cols; j++) {
            matrix[i][j] *= 2; // Double each element
        }
    }
}

// Usage
int main() {
    int rows = 3, cols = 4;
    int matrix[3][4] = {{1, 2, 3, 4}, {5, 6, 7, 8}, {9, 10, 11, 12}};

    process_matrix(rows, cols, matrix);

    // Print modified matrix
    for (int i = 0; i < rows; i++) {
        for (int j = 0; j < cols; j++) {
            printf("%d ", matrix[i][j]);
        }
        printf("\n");
    }

    return 0;
}

```

Array of Arrays vs. Multi-dimensional Arrays

Array of Arrays

An array of arrays treats each row as a separate array:

```

// Array of pointers to arrays (jagged array)
int *matrix[3];
matrix[0] = malloc(2 * sizeof(int)); // First row: 2 elements
matrix[1] = malloc(4 * sizeof(int)); // Second row: 4 elements
matrix[2] = malloc(3 * sizeof(int)); // Third row: 3 elements

```

```
// Initialize elements
matrix[0][0] = 1; matrix[0][1] = 2;
matrix[1][0] = 3; matrix[1][1] = 4; matrix[1][2] = 5; matrix[1][3] = 6;
matrix[2][0] = 7; matrix[2][1] = 8; matrix[2][2] = 9;

// Free memory
free(matrix[0]);
free(matrix[1]);
free(matrix[2]);
```

Multi-dimensional Arrays

Multi-dimensional arrays have fixed dimensions and contiguous memory layout:

```
// True multi-dimensional array
int matrix[3][4]; // Fixed size: 3 rows, 4 columns

// All elements stored contiguously in memory
// Memory layout: [0][0] [0][1] [0][2] [0][3] [1][0] [1][1] ...
```

Jagged Arrays

Creating Jagged Arrays

Jagged arrays are arrays of arrays where each sub-array can have different lengths:

```
#include <stdio.h>
#include <stdlib.h>

int main() {
    // Create a jagged array
    int rows = 3;
    int *jagged[3];

    // Allocate memory for each row with different sizes
    jagged[0] = malloc(2 * sizeof(int)); // Row 0: 2 elements
    jagged[1] = malloc(4 * sizeof(int)); // Row 1: 4 elements
    jagged[2] = malloc(3 * sizeof(int)); // Row 2: 3 elements

    // Initialize elements
    jagged[0][0] = 1; jagged[0][1] = 2;
    jagged[1][0] = 3; jagged[1][1] = 4; jagged[1][2] = 5; jagged[1][3] = 6;
    jagged[2][0] = 7; jagged[2][1] = 8; jagged[2][2] = 9;
```

```

// Print the jagged array
for (int i = 0; i < rows; i++) {
    int cols = (i == 0) ? 2 : (i == 1) ? 4 : 3;
    for (int j = 0; j < cols; j++) {
        printf("%d ", jagged[i][j]);
    }
    printf("\n");
}

// Free memory
for (int i = 0; i < rows; i++) {
    free(jagged[i]);
}

return 0;
}

```

Jagged Array Functions

```

// Function to create a jagged array
int** create_jagged_array(int rows, int *col_sizes) {
    int **array = malloc(rows * sizeof(int*));

    for (int i = 0; i < rows; i++) {
        array[i] = malloc(col_sizes[i] * sizeof(int));
    }

    return array;
}

// Function to free a jagged array
void free_jagged_array(int **array, int rows) {
    for (int i = 0; i < rows; i++) {
        free(array[i]);
    }
    free(array);
}

// Function to print a jagged array
void print_jagged_array(int **array, int rows, int *col_sizes) {
    for (int i = 0; i < rows; i++) {
        for (int j = 0; j < col_sizes[i]; j++) {
            printf("%d ", array[i][j]);
        }
        printf("\n");
    }
}

```

```
}  
}
```

Array Algorithms and Operations

Matrix Operations

```
#include <stdio.h>  
  
#define ROWS 3  
#define COLS 3  
  
// Matrix addition  
void matrix_add(int a[ROWS][COLS], int b[ROWS][COLS], int result[ROWS][COLS]) {  
    for (int i = 0; i < ROWS; i++) {  
        for (int j = 0; j < COLS; j++) {  
            result[i][j] = a[i][j] + b[i][j];  
        }  
    }  
}  
  
// Matrix multiplication  
void matrix_multiply(int a[ROWS][COLS], int b[COLS][ROWS], int result[ROWS][ROWS]) {  
    for (int i = 0; i < ROWS; i++) {  
        for (int j = 0; j < ROWS; j++) {  
            result[i][j] = 0;  
            for (int k = 0; k < COLS; k++) {  
                result[i][j] += a[i][k] * b[k][j];  
            }  
        }  
    }  
}  
  
// Transpose matrix  
void matrix_transpose(int matrix[ROWS][COLS], int transpose[COLS][ROWS]) {  
    for (int i = 0; i < ROWS; i++) {  
        for (int j = 0; j < COLS; j++) {  
            transpose[j][i] = matrix[i][j];  
        }  
    }  
}
```

Dynamic Array Resizing

```
#include <stdio.h>
#include <stdlib.h>

typedef struct {
    int *data;
    int size;
    int capacity;
} DynamicArray;

// Initialize dynamic array
DynamicArray* create_array(int initial_capacity) {
    DynamicArray *arr = malloc(sizeof(DynamicArray));
    arr->data = malloc(initial_capacity * sizeof(int));
    arr->size = 0;
    arr->capacity = initial_capacity;
    return arr;
}

// Add element to dynamic array
void array_push(DynamicArray *arr, int value) {
    if (arr->size >= arr->capacity) {
        // Double the capacity
        arr->capacity *= 2;
        arr->data = realloc(arr->data, arr->capacity * sizeof(int));
    }

    arr->data[arr->size] = value;
    arr->size++;
}

// Free dynamic array
void free_array(DynamicArray *arr) {
    free(arr->data);
    free(arr);
}
```

Array Performance Considerations

Memory Layout

Arrays have contiguous memory layout, which provides:

- Cache locality for sequential access
- Predictable memory usage
- Efficient indexing operations


```

// Row-major order traversal (efficient)
for (int i = 0; i < rows; i++) {
    for (int j = 0; j < cols; j++) {
        process(matrix[i][j]);
    }
}

// Column-major order traversal (less efficient for C arrays)
for (int j = 0; j < cols; j++) {
    for (int i = 0; i < rows; i++) {
        process(matrix[i][j]);
    }
}

```

Cache Performance

```

#define SIZE 1000

// Good cache performance - sequential access
void good_access_pattern(int matrix[SIZE][SIZE]) {
    for (int i = 0; i < SIZE; i++) {
        for (int j = 0; j < SIZE; j++) {
            matrix[i][j] = i + j;
        }
    }
}

// Poor cache performance - non-sequential access
void poor_access_pattern(int matrix[SIZE][SIZE]) {
    for (int j = 0; j < SIZE; j++) {
        for (int i = 0; i < SIZE; i++) {
            matrix[i][j] = i + j;
        }
    }
}

```

Advanced Array Techniques

Array Pointers

```
// Pointer to array
int (*ptr)[5]; // Pointer to array of 5 integers
int array[5] = {1, 2, 3, 4, 5};
ptr = &array; // Point to the entire array

// Access elements
printf("%d\n", (*ptr)[2]); // Access third element
```

Array of Pointers

```
// Array of pointers
int *ptr_array[5]; // Array of 5 pointers to integers
int values[5] = {10, 20, 30, 40, 50};

// Point each pointer to an element
for (int i = 0; i < 5; i++) {
    ptr_array[i] = &values[i];
}

// Access through pointers
for (int i = 0; i < 5; i++) {
    printf("%d ", *ptr_array[i]);
}
```

Compound Literals (C99)

```
// Compound literal for array
int *ptr = (int[]){1, 2, 3, 4, 5};

// Use in function calls
void process_array(int *arr, int size) {
    for (int i = 0; i < size; i++) {
        printf("%d ", arr[i]);
    }
}

// Call with compound literal
process_array((int[]){10, 20, 30, 40, 50}, 5);
```

Practical Examples

Sudoku Validator

```
#include <stdio.h>
#include <stdbool.h>

#define SIZE 9

bool is_valid_sudoku(int board[SIZE][SIZE]) {
    // Check rows
    for (int i = 0; i < SIZE; i++) {
        bool seen[SIZE + 1] = {false};
        for (int j = 0; j < SIZE; j++) {
            if (board[i][j] != 0) {
                if (seen[board[i][j]]) return false;
                seen[board[i][j]] = true;
            }
        }
    }

    // Check columns
    for (int j = 0; j < SIZE; j++) {
        bool seen[SIZE + 1] = {false};
        for (int i = 0; i < SIZE; i++) {
            if (board[i][j] != 0) {
                if (seen[board[i][j]]) return false;
                seen[board[i][j]] = true;
            }
        }
    }

    // Check 3x3 boxes
    for (int box_row = 0; box_row < 3; box_row++) {
        for (int box_col = 0; box_col < 3; box_col++) {
            bool seen[SIZE + 1] = {false};
            for (int i = 0; i < 3; i++) {
                for (int j = 0; j < 3; j++) {
                    int row = box_row * 3 + i;
                    int col = box_col * 3 + j;
                    if (board[row][col] != 0) {
                        if (seen[board[row][col]]) return false;
                        seen[board[row][col]] = true;
                    }
                }
            }
        }
    }
}
```

```

    }
}

return true;
}

int main() {
    int sudoku[SIZE][SIZE] = {
        {5, 3, 0, 0, 7, 0, 0, 0, 0},
        {6, 0, 0, 1, 9, 5, 0, 0, 0},
        {0, 9, 8, 0, 0, 0, 0, 6, 0},
        {8, 0, 0, 0, 6, 0, 0, 0, 3},
        {4, 0, 0, 8, 0, 3, 0, 0, 1},
        {7, 0, 0, 0, 2, 0, 0, 0, 6},
        {0, 6, 0, 0, 0, 0, 2, 8, 0},
        {0, 0, 0, 4, 1, 9, 0, 0, 5},
        {0, 0, 0, 0, 8, 0, 0, 7, 9}
    };

    if (is_valid_sudoku(sudoku)) {
        printf("Valid Sudoku\n");
    } else {
        printf("Invalid Sudoku\n");
    }

    return 0;
}

```

Image Processing Example

```

#include <stdio.h>
#include <stdlib.h>

#define WIDTH 100
#define HEIGHT 100

// Simple image structure
typedef struct {
    unsigned char pixels[HEIGHT][WIDTH][3]; // RGB pixels
    int width, height;
} Image;

// Convert to grayscale
void grayscale(Image *img) {
    for (int i = 0; i < img->height; i++) {

```

```

    for (int j = 0; j < img->width; j++) {
        // Calculate grayscale value
        unsigned char gray = (img->pixels[i][j][0] * 0.3 +
                               img->pixels[i][j][1] * 0.59 +
                               img->pixels[i][j][2] * 0.11);

        // Set all channels to grayscale value
        img->pixels[i][j][0] = gray;
        img->pixels[i][j][1] = gray;
        img->pixels[i][j][2] = gray;
    }
}

// Apply simple blur filter
void blur(Image *img) {
    // Create a copy for processing
    unsigned char temp[HEIGHT][WIDTH][3];

    for (int i = 1; i < img->height - 1; i++) {
        for (int j = 1; j < img->width - 1; j++) {
            for (int c = 0; c < 3; c++) {
                // Average of 3x3 neighborhood
                int sum = 0;
                for (int di = -1; di <= 1; di++) {
                    for (int dj = -1; dj <= 1; dj++) {
                        sum += img->pixels[i + di][j + dj][c];
                    }
                }
                temp[i][j][c] = sum / 9;
            }
        }
    }

    // Copy back to original image
    for (int i = 1; i < img->height - 1; i++) {
        for (int j = 1; j < img->width - 1; j++) {
            for (int c = 0; c < 3; c++) {
                img->pixels[i][j][c] = temp[i][j][c];
            }
        }
    }
}

```

Summary

Advanced array concepts in C provide powerful tools for handling complex data structures and implementing efficient algorithms. Key points to remember:

1. **Multi-dimensional Arrays:** Arrays of arrays with fixed dimensions and contiguous memory layout
2. **Variable-length Arrays:** C99 feature allowing runtime-determined array sizes
3. **Jagged Arrays:** Arrays of arrays with different sub-array sizes
4. **Memory Layout:** Row-major order storage affects performance characteristics
5. **Array Pointers:** Distinction between pointers to arrays and arrays of pointers
6. **Performance Considerations:** Cache locality and access patterns significantly impact performance
7. **Advanced Techniques:** Compound literals, designated initializers, and dynamic resizing

Understanding these advanced array concepts is essential for efficient C programming, particularly when working with matrices, tables, and complex data structures.

Module 5: Arrays and Strings

Overview

This module explores arrays and strings in C, fundamental data structures for storing and manipulating collections of data. You'll learn about array declaration, initialization, and access, as well as multi-dimensional arrays and advanced array concepts. The module also covers string handling in C, including string literals, character arrays, and the standard library functions for string manipulation.

Learning Objectives

By the end of this module, you will be able to:

- Declare, initialize, and access arrays of various types
- Work with multi-dimensional arrays and jagged arrays
- Understand array bounds and buffer overflow prevention
- Manipulate strings using character arrays and standard library functions
- Implement string searching, comparison, and manipulation algorithms
- Apply arrays and strings to solve practical programming problems
- Understand the relationship between arrays and pointers
- Use variable-length arrays (C99 feature) appropriately

Chapters

1. **Arrays** - Declaration, initialization, access, and basic operations
2. **Strings** - Character arrays, string literals, and string manipulation
3. **Advanced Arrays** - Multi-dimensional arrays, jagged arrays, and advanced concepts

Key Concepts Covered

- Array declaration and initialization syntax
- Array indexing and bounds checking
- Multi-dimensional arrays and initialization
- Array traversal and common operations
- String representation in C
- String literals vs. character arrays
- Standard library string functions
- String searching and manipulation
- Array algorithms (searching, sorting)
- Relationship between arrays and pointers

- Variable-length arrays (C99)
- Array of arrays vs. multi-dimensional arrays
- Jagged arrays using pointers
- Buffer overflow prevention techniques

Prerequisites

- Completion of Module 1: Foundations and Environment Setup
- Completion of Module 2: Data Types and Variables
- Completion of Module 3: Control Flow
- Completion of Module 4: Functions and Modular Programming
- Understanding of basic programming concepts
- Familiarity with variable declaration and initialization

Tools and Technologies

- C compiler (GCC, Clang, or MSVC)
- Text editor or IDE
- Understanding of basic programming concepts
- Familiarity with debugging tools

Estimated Time to Complete

- Reading: 8-10 hours
- Exercises: 12-15 hours
- Projects: 6-8 hours

Assessment

- Complete all chapter exercises
- Successfully implement array and string manipulation algorithms
- Demonstrate understanding of array bounds and string handling
- Pass array and string programming quiz
- Submit a comprehensive program using arrays and strings

Next Module

After completing this module, proceed to **Module 6: Pointers and Memory Management** to learn about memory addressing and dynamic allocation.

Additional Resources

- “The C Programming Language” by Kernighan and Ritchie (Chapter 5)
- C standard documentation (ISO/IEC 9899:2024 for C23)
- Online tutorials on array and string algorithms
- Debugging tools documentation (GDB, Valgrind)
- Best practices guides for secure string handling

Pointers

Pointer Fundamentals

Introduction

A pointer is a variable that stores the **address** of another object, not the object's value itself. Pointers unlock efficient parameter passing, dynamic data structures, and low-level system programming. This chapter builds intuition with step-by-step memory pictures, complete programs, `const` combinations, `void *`, array/pointer relationships, and out-parameters.

Compile with:

```
gcc -std=c17 -Wall -Wextra program.c -o program
```

What Is a Pointer?

Every object in your program lives at a memory address. The address-of operator `&` yields that address; the dereference operator `*` follows a pointer to the object.

```
#include <stdio.h>

int main(void) {
    int x = 42;
    printf("Value of x:  %d\n", x);
    printf("Address of x: %p\n", (void *)&x);
    return 0;
}
```

Always cast to `void *` when printing with `%p`.

Visual Step-by-Step: `&` and `*`

Think of memory as numbered boxes. Variable names label boxes; pointers hold box numbers.

Step 1 — Create an `int`

```
int x = 42;
```

```

name: x

42      address e.g. 0x1000

```

Step 2 — Point at it

```

int *p = &x;    /* p holds the address of x */

```

```

name: p                name: x

0x1000                42      0x1000

```

Step 3 — Read through the pointer

```

int y = *p;    /* y becomes 42 */

```

*p means “the int sitting at the address stored in p.”

Step 4 — Write through the pointer

```

*p = 99;        /* x is now 99 */

```

```

p      x: 99

```

Full Demo Program

```

#include <stdio.h>

int main(void) {
    int x = 42;
    int *p = &x;

    printf("1) x=%d  &x=%p  p=%p  *p=%d\n",
           x, (void *)&x, (void *)p, *p);

    *p = 99;
    printf("2) after *p=99: x=%d  *p=%d\n", x, *p);

    x = 7;
}

```

```

    printf("3) after x=7:   x=%d  *p=%d\n", x, *p);

    int z = 100;
    p = &z;
    printf("4) p now points at z: *p=%d  x still %d\n", *p, x);
    return 0;
}

gcc -std=c17 -Wall -Wextra ptr_steps.c -o ptr_steps
./ptr_steps

```

Declaration and Initialization

```

int *ptr;           /* pointer to int (uninitialized - dangerous) */
char *cptr;
float *fptr;

int x = 42;
int *p1 = &x;       /* initialized */
int *p2 = NULL;     /* safe empty pointer */
int *p3 = 0;        /* also null */

```

Null Pointers

```

#include <stdio.h>

int main(void) {
    int *p = NULL;

    if (p != NULL) {
        printf("%d\n", *p);
    } else {
        printf("Pointer is null - do not dereference\n");
    }
    return 0;
}

```

Never dereference NULL. On most systems that crashes; the language only says the behavior is undefined.

Swap via Pointers (Full Program)

Pass-by-value cannot swap the caller's variables. Pass pointers (addresses) so the function can write back.

Why Values Alone Fail

```
#include <stdio.h>

/* This does NOT swap the caller's x and y */
static void swap_by_value(int a, int b) {
    int t = a;
    a = b;
    b = t;
}

int main(void) {
    int x = 10, y = 20;
    swap_by_value(x, y);
    printf("still x=%d y=%d\n", x, y); /* 10 20 */
    return 0;
}
```

Correct Swap with Pointers

```
#include <stdio.h>

static void swap(int *a, int *b) {
    if (a == NULL || b == NULL) {
        return;
    }
    int temp = *a;
    *a = *b;
    *b = temp;
}

int main(void) {
    int x = 10, y = 20;
    printf("before: x=%d y=%d\n", x, y);

    swap(&x, &y);

    printf("after:  x=%d y=%d\n", x, y); /* 20 10 */
    return 0;
}
```

```
}
```

Picture of the call:

```
main:  x=10  y=20
```

```
swap:   a      b      (a and b hold addresses)
        *a    *b      (values exchanged)
```

```
gcc -std=c17 -Wall -Wextra swap.c -o swap
./swap
```

const and Pointers: Three Combinations

const can protect the **data**, the **pointer**, or **both**.

Declaration	Can change *p?	Can reassign p?
const int *p	no	yes
int *const p	yes	no
const int *const p	no	no
int *p	yes	yes

Full Program

```
#include <stdio.h>

int main(void) {
    int x = 10;
    int y = 20;

    /* 1) Pointer to const int - data is read-only through p */
    const int *p1 = &x;
    printf("p1 -> %d\n", *p1);
    /* *p1 = 99; */          /* ERROR: discards const */
    p1 = &y;                  /* OK: pointer may move */
    printf("p1 now -> %d\n", *p1);

    /* 2) Const pointer to int - pointer frozen, data mutable */
    int *const p2 = &x;
    *p2 = 11;                 /* OK */
    /* p2 = &y; */           /* ERROR: p2 is const */
    printf("x is now %d\n", x);
}
```

```

/* 3) Const pointer to const int - both frozen */
const int *const p3 = &y;
printf("p3 -> %d\n", *p3);
/* *p3 = 0; */          /* ERROR */
/* p3 = &x; */          /* ERROR */

/* Reading is always fine for all three */
printf("read-only views: %d %d %d\n", *p1, *p2, *p3);
return 0;
}

```

Why const int * Appears in APIs

```

/* Caller knows print_ints will not modify the buffer */
void print_ints(const int *a, size_t n);

```

That is a promise to the reader and a constraint the compiler helps enforce (unless you cast const away — don't).

Note on const int * vs int const *

These are the **same**: pointer to const int. Read declarations right-to-left: `const int *p` → “p is a pointer to const int.”

void * with Careful Casting

`void *` is a generic object pointer. You can convert to and from other object pointer types, but you **must not** dereference a `void *` without casting to a concrete type first.

Full Program: Generic Swap of Bytes

```

#include <stdio.h>
#include <stdlib.h>
#include <string.h>

/* Swap n bytes between the objects pointed to by a and b. */
static int memswap(void *a, void *b, size_t n) {
    if (a == NULL || b == NULL) {
        return -1;
    }
    if (n == 0) {
        return 0;
    }

```



```

    }
    unsigned char *p = a;    /* implicit void* → unsigned char* */
    unsigned char *q = b;
    for (size_t i = 0; i < n; i++) {
        unsigned char t = p[i];
        p[i] = q[i];
        q[i] = t;
    }
    return 0;
}

static void print_ints(const char *label, const int *a, size_t n) {
    printf("%s:", label);
    for (size_t i = 0; i < n; i++) {
        printf(" %d", a[i]);
    }
    printf("\n");
}

int main(void) {
    int u = 5, v = 9;
    memswap(&u, &v, sizeof u);
    printf("u=%d v=%d\n", u, v);

    int a[] = {1, 2, 3};
    int b[] = {7, 8, 9};
    memswap(a, b, sizeof a);
    print_ints("a", a, 3);
    print_ints("b", b, 3);

    /* Careful cast example: recover a typed pointer from void* */
    void *raw = a;
    int *typed = (int *)raw;    /* explicit cast for clarity */
    printf("typed[0]=%d\n", typed[0]);
    return 0;
}

gcc -std=c17 -Wall -Wextra void_swap.c -o void_swap
./void_swap

```

Rules of Thumb for void *

1. Use for generic buffers, allocators (`malloc` returns `void *`), and callbacks.
2. Cast to the **actual** type before arithmetic or dereference.
3. Keep size information alongside the pointer when the type is erased.
4. Do not cast `void *` to function pointers (different category in C).

Pointer Arithmetic (Brief)

Arithmetic scales by the size of the pointed-to type:

```
#include <stdio.h>

int main(void) {
    int arr[5] = {10, 20, 30, 40, 50};
    int *p = arr; /* same as &arr[0] */

    printf("*p      = %d\n", *p);      /* 10 */
    printf("*(p+1)  = %d\n", *(p + 1)); /* 20 */
    printf("p[2]    = %d\n", p[2]);    /* 30 - syntactic sugar */

    p++;
    printf("after p++: %d\n", *p);    /* 20 */
    return 0;
}
```

`p[i]` is defined as `*(p + i)`.

Pointers vs Arrays: Equivalence and Differences

Where They Look the Same

In most expression contexts, an array name **decays** to a pointer to its first element:

```
int a[4] = {1, 2, 3, 4};
int *p = a; /* decay */
int *q = &a[0]; /* same address */

/* These all refer to the third element (value 3) */
a[2];
*(a + 2);
p[2];
*(p + 2);
```

Where They Differ

	Array <code>int a[4]</code>	Pointer <code>int *p</code>
Storage	Holds 4 ints	Holds one address
<code>sizeof</code>	<code>4 * sizeof(int)</code>	<code>sizeof(int *)</code>
Assignability	<code>a = ...</code> illegal	<code>p = ...</code> OK
Address	<code>&a</code> is <code>int (*)[4]</code>	<code>&p</code> is <code>int **</code>

	Array <code>int a[4]</code>	Pointer <code>int *p</code>
Lifetime	As declared (e.g. <code>auto</code>)	Points wherever you set

Full Comparison Program

```
#include <stdio.h>

static void show_param(int arr[]) {
    /* arr is really int * here */
    printf("    sizeof(param) = %zu (pointer, not array)\n", sizeof arr);
}

int main(void) {
    int a[4] = {1, 2, 3, 4};
    int *p = a;

    printf("sizeof(a) = %zu\n", sizeof a);    /* e.g. 16 */
    printf("sizeof(p) = %zu\n", sizeof p);    /* e.g. 8 */

    printf("a[2]=%d p[2]=%d\n", a[2], p[2]);

    p = &a[1];    /* OK */
    /* a = p; */ /* ERROR: array is not assignable */

    show_param(a);

    /* &a vs &a[0] */
    printf("&a    = %p (pointer to whole array)\n", (void *)&a);
    printf("&a[0] = %p (pointer to first int)\n", (void *)&a[0]);
    printf("same address, different types/sizeof arithmetic\n");
    return 0;
}
```

Out-Parameters: Return Status Separately from Data

A common C pattern: return an `int` status (0 success, non-zero error) and write results through pointer out-parameters.

Full Program

```
#include <stdio.h>
#include <stdbool.h>

enum {
    OK = 0,
    ERR_NULL = -1,
    ERR_RANGE = -2,
    ERR_EMPTY = -3
};

/* Divide n / d into *quot_out. */
static int safe_div(int n, int d, int *quot_out) {
    if (quot_out == NULL) {
        return ERR_NULL;
    }
    if (d == 0) {
        return ERR_RANGE;
    }
    *quot_out = n / d;
    return OK;
}

/* Find max of a[0..n). */
static int array_max(const int *a, size_t n, int *max_out) {
    if (a == NULL || max_out == NULL) {
        return ERR_NULL;
    }
    if (n == 0) {
        return ERR_EMPTY;
    }
    int m = a[0];
    for (size_t i = 1; i < n; i++) {
        if (a[i] > m) {
            m = a[i];
        }
    }
    *max_out = m;
    return OK;
}

/* Parse a simple non-negative integer from a C string. */
static int parse_u32(const char *s, unsigned *out) {
    if (s == NULL || out == NULL) {
        return ERR_NULL;
    }
}
```

```

    if (*s == '\\0') {
        return ERR_EMPTY;
    }
    unsigned value = 0;
    for (const char *p = s; *p != '\\0'; p++) {
        if (*p < '0' || *p > '9') {
            return ERR_RANGE;
        }
        value = value * 10u + (unsigned)(*p - '0');
    }
    *out = value;
    return OK;
}

int main(void) {
    int q;
    if (safe_div(42, 5, &q) == OK) {
        printf("42/5 = %d\\n", q);
    }
    if (safe_div(42, 0, &q) != OK) {
        printf("division by zero handled\\n");
    }

    int data[] = {3, 9, 2, 9, 5};
    int mx;
    if (array_max(data, 5, &mx) == OK) {
        printf("max = %d\\n", mx);
    }

    unsigned u;
    if (parse_u32("12345", &u) == OK) {
        printf("parsed %u\\n", u);
    }
    if (parse_u32("12x", &u) != OK) {
        printf("reject bad digits\\n");
    }
    return 0;
}

```

```

gcc -std=c17 -Wall -Wextra outparams.c -o outparams
./outparams

```

Design Tips

1. Document who allocates and who frees any pointer you return.
2. Prefer `const T *` for inputs you only read.

3. Check for NULL out-parameters when the API is public.
4. Do not leave `*out` half-written on failure unless you document it; either write only on success or define failure values.

Pointers and Functions (More)

Returning a Pointer into a Caller-Owned Array

```
#include <stdio.h>

/* Returns pointer to first element equal to target, or NULL. */
static int *find_int(int *a, size_t n, int target) {
    for (size_t i = 0; i < n; i++) {
        if (a[i] == target) {
            return &a[i];
        }
    }
    return NULL;
}

int main(void) {
    int nums[] = {10, 50, 30, 40, 20};
    int *hit = find_int(nums, 5, 40);
    if (hit != NULL) {
        printf("found %d at offset %td\n", *hit, hit - nums);
        *hit = 41; /* modify through returned pointer */
    }
    printf("nums[3]=%d\n", nums[3]);
    return 0;
}
```

Safe because the array lives in `main` longer than the pointer is used.

Unsafe: Returning Address of Local

```
/* NEVER do this */
static int *bad(void) {
    int local = 42;
    return &local; /* dangling */
}
```

Pointer Comparison and Size

```
#include <stdio.h>

int main(void) {
    int arr[5] = {10, 20, 30, 40, 50};
    int *p1 = &arr[1];
    int *p2 = &arr[4];

    if (p1 < p2) {
        printf("p1 before p2 in the same array\n");
    }
    printf("element distance: %td\n", p2 - p1); /* 3 */

    printf("sizeof(int*) = %zu\n", sizeof(int *));
    printf("sizeof(char*) = %zu\n", sizeof(char *));
    /* Same size on typical platforms; type still matters for arithmetic */
    return 0;
}
```

Relational comparisons are only defined for pointers into the same array object (or one past the end).

Best Practices

1. **Initialize** every pointer (NULL or a valid address).
2. **Check** for NULL before dereference when the value might be empty.
3. Prefer **const** on pointer parameters that only read data.
4. Do not return addresses of **automatic** locals.
5. Keep **size** next to raw buffers and **void ***.
6. Use **out-parameters** + **status** instead of overloading return values for both data and errors when both matter.
7. Compile with **-Wall -Wextra**; treat warnings as bugs.

Practical Examples

Pointer-Based String Reverse

```
#include <stdio.h>
#include <string.h>

static void reverse_string(char *str) {
    if (str == NULL || *str == '\0') {
```

```

        return;
    }
    char *start = str;
    char *end = str + strlen(str) - 1;
    while (start < end) {
        char temp = *start;
        *start = *end;
        *end = temp;
        start++;
        end--;
    }
}

int main(void) {
    char str[] = "Hello, World!"; /* must be mutable array */
    printf("Original: %s\n", str);
    reverse_string(str);
    printf("Reversed: %s\n", str);
    return 0;
}

```

Sum with Pointer Walking

```

#include <stdio.h>

static int sum_array(const int *arr, size_t n) {
    int sum = 0;
    const int *end = arr + n;
    for (const int *p = arr; p < end; p++) {
        sum += *p;
    }
    return sum;
}

int main(void) {
    int numbers[] = {1, 2, 3, 4, 5};
    size_t n = sizeof numbers / sizeof numbers[0];
    printf("Sum: %d\n", sum_array(numbers, n));
    return 0;
}

```


Exercises

1. **Trace on paper.** For `int x=3; int *p=&x; *p=5; int y=*p;`, draw boxes after each statement.
2. **Swap three ways.** Implement swap of two ints with pointers; of two doubles; and a generic `memswap` with `void *` and `sizeof`.
3. **const experiments.** Write four small snippets that should fail to compile for illegal uses of `const int *`, `int *const`, and `const int *const`. Confirm with `gcc -std=c17 -Wall -Wextra`.
4. **Out-parameter min/max.** `int minmax(const int *a, size_t n, int *min_out, int *max_out)` returning status codes; write a driver that tests empty and null cases.
5. **Find and replace.** Return a pointer to the first space in a mutable string; replace that space with `'_'` through the pointer.
6. **Array vs pointer quiz.** Print `sizeof` for an array and for a pointer parameter that “receives” it. Explain the numbers in a comment.
7. **void * printer.** Write `void print_bytes(const void *p, size_t n)` that prints `n` bytes in hex by casting to `const unsigned char *`.
8. **Pointer distance.** Given two pointers into the same array, print how many elements lie between them. What happens if they are not from the same array? (Do not rely on that case.)

Summary

1. Pointers store **addresses**; `*` accesses the object; `&` produces an address.
2. **Swap and out-parameters** are the classic reasons to pass pointers to functions.
3. **const** can lock data, the pointer variable, or both — use it in APIs.
4. **void *** erases type; cast carefully and track sizes.
5. Arrays **decay** to pointers but are not the same as pointers (`sizeof`, assignment).
6. Always initialize, avoid dangling pointers, and return status separately when needed.

Mastering these fundamentals makes every later C topic — arrays, strings, dynamic memory, data structures — much easier to reason about.

Pointers and Arrays

Introduction

One of the most important relationships in C programming is between pointers and arrays. In fact, arrays and pointers are closely related concepts, with arrays often being treated as pointers in many contexts. Understanding this relationship is crucial for effective C programming, as it enables efficient array manipulation, dynamic memory allocation, and implementation of complex data structures.

Array-Pointer Equivalence

Arrays as Pointers

In most contexts, an array name is treated as a pointer to its first element:

```
int arr[5] = {10, 20, 30, 40, 50};

// These are equivalent:
int *ptr1 = arr;      // Array name decays to pointer
int *ptr2 = &arr[0];  // Explicit address of first element

printf("arr[0] = %d, *ptr1 = %d, *ptr2 = %d\n", arr[0], *ptr1, *ptr2);
```

Important Distinction

While arrays and pointers are often interchangeable, they are not identical:

```
int arr[5] = {1, 2, 3, 4, 5};
int *ptr = arr;

// Arrays have fixed size
sizeof(arr); // Returns 20 bytes (5 * sizeof(int))

// Pointers have size of address
sizeof(ptr); // Returns 8 bytes on 64-bit system (size of pointer)

// Arrays cannot be reassigned
```

```
// arr = ptr; // Error: Cannot assign to array

// Pointers can be reassigned
ptr = &arr[2]; // OK: Pointer can be changed
```

Pointer Subscripting

Using Pointers as Arrays

Pointers can be used with array subscript notation:

```
int arr[5] = {10, 20, 30, 40, 50};
int *ptr = arr;

// These are equivalent:
printf("arr[2] = %d\n", arr[2]);
printf("ptr[2] = %d\n", ptr[2]);
printf("*(arr + 2) = %d\n", *(arr + 2));
printf("*(ptr + 2) = %d\n", *(ptr + 2));
```

Array Access Methods

There are multiple equivalent ways to access array elements:

```
int arr[5] = {10, 20, 30, 40, 50};

// Method 1: Array subscript notation
int value1 = arr[2];

// Method 2: Pointer arithmetic with dereference
int value2 = *(arr + 2);

// Method 3: Using pointer variable
int *ptr = arr;
int value3 = *(ptr + 2);
int value4 = ptr[2];

// All methods produce the same result
printf("%d %d %d %d\n", value1, value2, value3, value4); // 30 30 30 30
```

Pointer Arithmetic with Arrays

Incrementing Pointers

Pointers can be incremented to traverse arrays:

```
int arr[5] = {10, 20, 30, 40, 50};
int *ptr = arr;

// Traverse array using pointer arithmetic
for (int i = 0; i < 5; i++) {
    printf("Element %d: %d\n", i, *ptr);
    ptr++; // Move to next element
}

// Reset pointer to beginning
ptr = arr;

// Alternative traversal
for (int i = 0; i < 5; i++) {
    printf("Element %d: %d\n", i, *(ptr + i));
}
```

Pointer Difference

The difference between two pointers gives the number of elements between them:

```
int arr[10] = {0, 1, 2, 3, 4, 5, 6, 7, 8, 9};
int *ptr1 = &arr[2];
int *ptr2 = &arr[8];

// Number of elements between ptr1 and ptr2
ptrdiff_t distance = ptr2 - ptr1; // 6
printf("Distance: %td\n", distance);

// Check if pointer is within array bounds
if (ptr1 >= arr && ptr1 < arr + 10) {
    printf("ptr1 is within array bounds\n");
}
```

Multi-dimensional Arrays and Pointers

Two-dimensional Arrays

Two-dimensional arrays can be accessed using various pointer techniques:

```

int matrix[3][4] = {
    {1, 2, 3, 4},
    {5, 6, 7, 8},
    {9, 10, 11, 12}
};

// Method 1: Standard array notation
printf("matrix[1][2] = %d\n", matrix[1][2]); // 7

// Method 2: Pointer to array
int (*ptr_to_row)[4] = matrix; // Pointer to array of 4 integers
printf("ptr_to_row[1][2] = %d\n", ptr_to_row[1][2]); // 7

// Method 3: Pointer to first element
int *ptr_to_element = &matrix[0][0];
printf("*(ptr_to_element + 1*4 + 2) = %d\n", *(ptr_to_element + 1*4 + 2)); // 7

```

Row-wise Traversal

```

int matrix[3][4] = {
    {1, 2, 3, 4},
    {5, 6, 7, 8},
    {9, 10, 11, 12}
};

// Traverse row by row
for (int i = 0; i < 3; i++) {
    for (int j = 0; j < 4; j++) {
        printf("%d ", matrix[i][j]);
    }
    printf("\n");
}

// Using pointer arithmetic
int *ptr = &matrix[0][0];
for (int i = 0; i < 3; i++) {
    for (int j = 0; j < 4; j++) {
        printf("%d ", *ptr);
        ptr++;
    }
    printf("\n");
}

```

Arrays as Function Parameters

Array Decay to Pointers

When arrays are passed to functions, they decay to pointers:

```
// These function declarations are equivalent:
void process_array1(int arr[]);
void process_array2(int *arr);
void process_array3(int arr[10]); // Size is ignored

// Implementation
void process_array(int arr[], int size) {
    // sizeof(arr) returns pointer size, not array size!
    for (int i = 0; i < size; i++) {
        printf("Element %d: %d\n", i, arr[i]);
        // Equivalent to: printf("Element %d: %d\n", i, *(arr + i));
    }
}

int main() {
    int numbers[5] = {10, 20, 30, 40, 50};
    process_array(numbers, 5); // Pass array and size
    return 0;
}
```

Multi-dimensional Array Parameters

For multi-dimensional arrays, all dimensions except the first must be specified:

```
// Function accepting 2D array - second dimension must be specified
void process_matrix(int matrix[][4], int rows) {
    for (int i = 0; i < rows; i++) {
        for (int j = 0; j < 4; j++) {
            printf("%d ", matrix[i][j]);
        }
        printf("\n");
    }
}

// Alternative using pointer notation
void process_matrix_ptr(int (*matrix)[4], int rows) {
    for (int i = 0; i < rows; i++) {
        for (int j = 0; j < 4; j++) {
            printf("%d ", matrix[i][j]);
        }
    }
}
```

```

        printf("\n");
    }
}

int main() {
    int matrix[3][4] = {
        {1, 2, 3, 4},
        {5, 6, 7, 8},
        {9, 10, 11, 12}
    };

    process_matrix(matrix, 3);
    return 0;
}

```

Dynamic Arrays with Pointers

Allocating Dynamic Arrays

Pointers enable dynamic array allocation:

```

#include <stdio.h>
#include <stdlib.h>

int main() {
    int size;
    printf("Enter array size: ");
    scanf("%d", &size);

    // Allocate dynamic array
    int *arr = malloc(size * sizeof(int));
    if (arr == NULL) {
        printf("Memory allocation failed\n");
        return 1;
    }

    // Initialize array
    for (int i = 0; i < size; i++) {
        arr[i] = i * i;
    }

    // Use array
    for (int i = 0; i < size; i++) {
        printf("arr[%d] = %d\n", i, arr[i]);
    }
}

```

```

    // Free memory
    free(arr);
    arr = NULL; // Prevent dangling pointer

    return 0;
}

```

Resizing Dynamic Arrays

```

#include <stdio.h>
#include <stdlib.h>

int main() {
    int *arr = malloc(5 * sizeof(int));
    int size = 5;

    // Initialize array
    for (int i = 0; i < size; i++) {
        arr[i] = i + 1;
    }

    printf("Original array: ");
    for (int i = 0; i < size; i++) {
        printf("%d ", arr[i]);
    }
    printf("\n");

    // Resize array
    size = 10;
    arr = realloc(arr, size * sizeof(int));
    if (arr == NULL) {
        printf("Memory reallocation failed\n");
        return 1;
    }

    // Initialize new elements
    for (int i = 5; i < size; i++) {
        arr[i] = (i + 1) * 2;
    }

    printf("Resized array: ");
    for (int i = 0; i < size; i++) {
        printf("%d ", arr[i]);
    }
    printf("\n");
}

```



```
    // Free memory
    free(arr);
    return 0;
}
```

Pointer Arrays vs. Arrays of Pointers

Array of Pointers

An array of pointers stores multiple pointer values:

```
#include <stdio.h>
#include <string.h>
#include <stdlib.h>

int main() {
    // Array of string pointers
    char *strings[4] = {
        "Hello",
        "World",
        "C Programming",
        "Pointers"
    };

    // Print all strings
    for (int i = 0; i < 4; i++) {
        printf("String %d: %s\n", i, strings[i]);
    }

    // Dynamic string array
    char **dynamic_strings = malloc(3 * sizeof(char*));
    dynamic_strings[0] = malloc(6 * sizeof(char));
    strcpy(dynamic_strings[0], "First");

    dynamic_strings[1] = malloc(7 * sizeof(char));
    strcpy(dynamic_strings[1], "Second");

    dynamic_strings[2] = malloc(6 * sizeof(char));
    strcpy(dynamic_strings[2], "Third");

    // Print dynamic strings
    for (int i = 0; i < 3; i++) {
        printf("Dynamic string %d: %s\n", i, dynamic_strings[i]);
    }
}
```

```

    // Free dynamic memory
    for (int i = 0; i < 3; i++) {
        free(dynamic_strings[i]);
    }
    free(dynamic_strings);

    return 0;
}

```

Pointer to Array

A pointer to an array points to the entire array:

```

int arr[5] = {1, 2, 3, 4, 5};

// Pointer to array of 5 integers
int (*ptr_to_array)[5] = &arr;

// Access elements
printf("First element: %d\n", (*ptr_to_array)[0]);
printf("Third element: %d\n", (*ptr_to_array)[2]);

// Increment pointer (moves by entire array size)
ptr_to_array++; // Points to next array of 5 integers

```

Advanced Array-Pointer Techniques

Pointer Arithmetic for Multi-dimensional Arrays

```

int matrix[3][4] = {
    {1, 2, 3, 4},
    {5, 6, 7, 8},
    {9, 10, 11, 12}
};

// Treat 2D array as 1D array
int *ptr = &matrix[0][0];
for (int i = 0; i < 12; i++) {
    printf("%d ", ptr[i]);
    if ((i + 1) % 4 == 0) printf("\n");
}

```

Function Pointers for Array Operations

```
#include <stdio.h>

// Function pointer type for array operations
typedef int (*array_operation)(int);

// Array operation functions
int square(int x) { return x * x; }
int cube(int x) { return x * x * x; }
int double_value(int x) { return x * 2; }

// Apply operation to array
void apply_operation(int *arr, int size, array_operation op) {
    for (int i = 0; i < size; i++) {
        arr[i] = op(arr[i]);
    }
}

int main() {
    int numbers[5] = {1, 2, 3, 4, 5};
    int size = 5;

    printf("Original: ");
    for (int i = 0; i < size; i++) {
        printf("%d ", numbers[i]);
    }
    printf("\n");

    // Apply square operation
    apply_operation(numbers, size, square);

    printf("Squared: ");
    for (int i = 0; i < size; i++) {
        printf("%d ", numbers[i]);
    }
    printf("\n");

    return 0;
}
```

Practical Examples

Matrix Transpose Using Pointers

```
#include <stdio.h>

void transpose_matrix(int *matrix, int *transpose, int rows, int cols) {
    for (int i = 0; i < rows; i++) {
        for (int j = 0; j < cols; j++) {
            // matrix[i][j] becomes transpose[j][i]
            *(transpose + j * rows + i) = *(matrix + i * cols + j);
        }
    }
}

int main() {
    int matrix[3][4] = {
        {1, 2, 3, 4},
        {5, 6, 7, 8},
        {9, 10, 11, 12}
    };

    int transpose[4][3];

    transpose_matrix(&matrix[0][0], &transpose[0][0], 3, 4);

    printf("Original matrix:\n");
    for (int i = 0; i < 3; i++) {
        for (int j = 0; j < 4; j++) {
            printf("%d ", matrix[i][j]);
        }
        printf("\n");
    }

    printf("\nTransposed matrix:\n");
    for (int i = 0; i < 4; i++) {
        for (int j = 0; j < 3; j++) {
            printf("%d ", transpose[i][j]);
        }
        printf("\n");
    }

    return 0;
}
```

String Array Sorting

```
#include <stdio.h>
#include <string.h>

void sort_strings(char *strings[], int count) {
    for (int i = 0; i < count - 1; i++) {
        for (int j = 0; j < count - i - 1; j++) {
            if (strcmp(strings[j], strings[j + 1]) > 0) {
                // Swap pointers
                char *temp = strings[j];
                strings[j] = strings[j + 1];
                strings[j + 1] = temp;
            }
        }
    }
}

int main() {
    char *fruits[] = {"apple", "orange", "banana", "pear", "grape"};
    int count = 5;

    printf("Original order:\n");
    for (int i = 0; i < count; i++) {
        printf("%s ", fruits[i]);
    }
    printf("\n");

    sort_strings(fruits, count);

    printf("Sorted order:\n");
    for (int i = 0; i < count; i++) {
        printf("%s ", fruits[i]);
    }
    printf("\n");

    return 0;
}
```

Dynamic 2D Array

```
#include <stdio.h>
#include <stdlib.h>

int** create_2d_array(int rows, int cols) {
```

```

// Allocate array of row pointers
int **matrix = malloc(rows * sizeof(int*));
if (matrix == NULL) return NULL;

// Allocate each row
for (int i = 0; i < rows; i++) {
    matrix[i] = malloc(cols * sizeof(int));
    if (matrix[i] == NULL) {
        // Clean up on failure
        for (int j = 0; j < i; j++) {
            free(matrix[j]);
        }
        free(matrix);
        return NULL;
    }
}

return matrix;
}

void free_2d_array(int **matrix, int rows) {
    for (int i = 0; i < rows; i++) {
        free(matrix[i]);
    }
    free(matrix);
}

int main() {
    int rows = 3, cols = 4;
    int **matrix = create_2d_array(rows, cols);

    if (matrix == NULL) {
        printf("Memory allocation failed\n");
        return 1;
    }

    // Initialize matrix
    for (int i = 0; i < rows; i++) {
        for (int j = 0; j < cols; j++) {
            matrix[i][j] = i * cols + j;
        }
    }

    // Print matrix
    for (int i = 0; i < rows; i++) {
        for (int j = 0; j < cols; j++) {

```

```
        printf("%d ", matrix[i][j]);
    }
    printf("\n");
}

// Free memory
free_2d_array(matrix, rows);

return 0;
}
```

Summary

The relationship between pointers and arrays is fundamental to C programming. Key points to remember:

1. **Array-Pointer Equivalence:** Arrays decay to pointers in most contexts
2. **Pointer Subscripting:** Pointers can use array notation and vice versa
3. **Pointer Arithmetic:** Operations are scaled by data type size
4. **Function Parameters:** Arrays pass as pointers to functions
5. **Dynamic Arrays:** Pointers enable runtime array allocation
6. **Multi-dimensional Arrays:** Various pointer techniques for 2D arrays
7. **Memory Management:** Proper allocation and deallocation of dynamic arrays

Understanding this relationship enables efficient array manipulation, dynamic memory allocation, and implementation of complex data structures in C.

Dynamic Memory Management

Introduction

Dynamic memory management is a crucial aspect of C programming that allows programs to allocate and deallocate memory at runtime. Unlike static memory allocation where the size and lifetime of variables are determined at compile time, dynamic memory allocation provides flexibility for handling data structures whose size is unknown until the program runs. This capability is essential for implementing complex data structures like linked lists, trees, and dynamic arrays.

Memory Layout in C Programs

Program Memory Sections

A C program's memory is typically organized into several sections:

1. **Text Segment:** Contains executable code
2. **Data Segment:** Contains global and static variables
3. **Stack:** Contains local variables and function call information
4. **Heap:** Contains dynamically allocated memory

Stack vs. Heap

```
// Stack allocation - automatic memory management
void stack_example() {
    int local_array[100]; // Allocated on stack
    // Automatically deallocated when function returns
}

// Heap allocation - manual memory management
void heap_example() {
    int *dynamic_array = malloc(100 * sizeof(int)); // Allocated on heap
    // Must be manually deallocated with free()
    free(dynamic_array);
}
```


Dynamic Memory Allocation Functions

malloc()

The malloc() function allocates a block of memory of specified size:

```
#include <stdio.h>
#include <stdlib.h>

int main() {
    // Allocate memory for 10 integers
    int *arr = malloc(10 * sizeof(int));

    if (arr == NULL) {
        printf("Memory allocation failed\n");
        return 1;
    }

    // Use the allocated memory
    for (int i = 0; i < 10; i++) {
        arr[i] = i * i;
        printf("arr[%d] = %d\n", i, arr[i]);
    }

    // Free the allocated memory
    free(arr);
    arr = NULL; // Prevent dangling pointer

    return 0;
}
```

calloc()

The calloc() function allocates memory and initializes it to zero:

```
#include <stdio.h>
#include <stdlib.h>

int main() {
    // Allocate and zero-initialize memory for 10 integers
    int *arr = calloc(10, sizeof(int));

    if (arr == NULL) {
        printf("Memory allocation failed\n");
        return 1;
    }
}
```

```

// Memory is already zero-initialized
for (int i = 0; i < 10; i++) {
    printf("arr[%d] = %d\n", i, arr[i]); // All zeros
}

// Free the allocated memory
free(arr);
arr = NULL;

return 0;
}

```

realloc()

The `realloc()` function changes the size of previously allocated memory:

```

#include <stdio.h>
#include <stdlib.h>

int main() {
    // Initial allocation
    int *arr = malloc(5 * sizeof(int));
    if (arr == NULL) {
        printf("Initial allocation failed\n");
        return 1;
    }

    // Initialize array
    for (int i = 0; i < 5; i++) {
        arr[i] = i + 1;
    }

    printf("Original array: ");
    for (int i = 0; i < 5; i++) {
        printf("%d ", arr[i]);
    }
    printf("\n");

    // Resize array to hold 10 elements
    arr = realloc(arr, 10 * sizeof(int));
    if (arr == NULL) {
        printf("Reallocation failed\n");
        return 1;
    }
}

```

```

// Initialize new elements
for (int i = 5; i < 10; i++) {
    arr[i] = (i + 1) * 2;
}

printf("Resized array: ");
for (int i = 0; i < 10; i++) {
    printf("%d ", arr[i]);
}
printf("\n");

// Free memory
free(arr);
return 0;
}

```

free()

The `free()` function deallocates previously allocated memory:

```

#include <stdio.h>
#include <stdlib.h>

int main() {
    // Allocate memory
    int *ptr = malloc(100 * sizeof(int));
    if (ptr == NULL) {
        printf("Allocation failed\n");
        return 1;
    }

    // Use memory...
    for (int i = 0; i < 100; i++) {
        ptr[i] = i;
    }

    // Deallocate memory
    free(ptr);
    ptr = NULL; // Prevent dangling pointer

    return 0;
}

```

Memory Allocation Strategies

Single Element Allocation

```
#include <stdio.h>
#include <stdlib.h>

typedef struct {
    int id;
    char name[50];
} Person;

int main() {
    // Allocate single structure
    Person *person = malloc(sizeof(Person));
    if (person == NULL) {
        printf("Allocation failed\n");
        return 1;
    }

    // Initialize structure
    person->id = 1;
    strcpy(person->name, "John Doe");

    printf("ID: %d, Name: %s\n", person->id, person->name);

    // Free memory
    free(person);
    person = NULL;

    return 0;
}
```

Array Allocation

```
#include <stdio.h>
#include <stdlib.h>

int main() {
    int size;
    printf("Enter array size: ");
    scanf("%d", &size);

    // Allocate dynamic array
    int *arr = malloc(size * sizeof(int));
```

```

if (arr == NULL) {
    printf("Allocation failed\n");
    return 1;
}

// Initialize array
for (int i = 0; i < size; i++) {
    arr[i] = i * i;
}

// Use array
printf("Array elements: ");
for (int i = 0; i < size; i++) {
    printf("%d ", arr[i]);
}
printf("\n");

// Free memory
free(arr);
arr = NULL;

return 0;
}

```

2D Array Allocation

```

#include <stdio.h>
#include <stdlib.h>

int** create_2d_array(int rows, int cols) {
    // Allocate array of row pointers
    int **matrix = malloc(rows * sizeof(int*));
    if (matrix == NULL) return NULL;

    // Allocate each row
    for (int i = 0; i < rows; i++) {
        matrix[i] = malloc(cols * sizeof(int));
        if (matrix[i] == NULL) {
            // Clean up on failure
            for (int j = 0; j < i; j++) {
                free(matrix[j]);
            }
            free(matrix);
            return NULL;
        }
    }
}

```

```

    }

    return matrix;
}

void free_2d_array(int **matrix, int rows) {
    for (int i = 0; i < rows; i++) {
        free(matrix[i]);
    }
    free(matrix);
}

int main() {
    int rows = 3, cols = 4;
    int **matrix = create_2d_array(rows, cols);

    if (matrix == NULL) {
        printf("Matrix allocation failed\n");
        return 1;
    }

    // Initialize matrix
    for (int i = 0; i < rows; i++) {
        for (int j = 0; j < cols; j++) {
            matrix[i][j] = i * cols + j;
        }
    }

    // Print matrix
    printf("Matrix:\n");
    for (int i = 0; i < rows; i++) {
        for (int j = 0; j < cols; j++) {
            printf("%d ", matrix[i][j]);
        }
        printf("\n");
    }

    // Free memory
    free_2d_array(matrix, rows);

    return 0;
}

```

Memory Management Best Practices

Error Checking

Always check if memory allocation was successful:

```
int *ptr = malloc(100 * sizeof(int));
if (ptr == NULL) {
    // Handle allocation failure
    fprintf(stderr, "Memory allocation failed\n");
    exit(EXIT_FAILURE);
}
```

Preventing Dangling Pointers

Set pointers to NULL after freeing:

```
free(ptr);
ptr = NULL; // Prevent dangling pointer
```

Avoiding Memory Leaks

Ensure every malloc/calloc/realloc has a corresponding free:

```
// Good practice
int *arr = malloc(10 * sizeof(int));
if (arr != NULL) {
    // Use array...
    free(arr);
    arr = NULL;
}
```

Memory Reallocation Safety

Handle realloc failure properly:

```
int *temp = realloc(arr, new_size * sizeof(int));
if (temp == NULL) {
    // realloc failed, original array still valid
    // Handle error without losing original data
    fprintf(stderr, "Reallocation failed\n");
    // arr is still valid
} else {
    // realloc succeeded
```

```
    arr = temp;
}
```

Common Memory Management Patterns

Dynamic String Handling

```
#include <stdio.h>
#include <stdlib.h>
#include <string.h>

char* duplicate_string(const char *source) {
    if (source == NULL) return NULL;

    size_t len = strlen(source);
    char *copy = malloc((len + 1) * sizeof(char));
    if (copy == NULL) return NULL;

    strcpy(copy, source);
    return copy;
}

char* append_strings(const char *str1, const char *str2) {
    if (str1 == NULL || str2 == NULL) return NULL;

    size_t len1 = strlen(str1);
    size_t len2 = strlen(str2);
    char *result = malloc((len1 + len2 + 1) * sizeof(char));
    if (result == NULL) return NULL;

    strcpy(result, str1);
    strcat(result, str2);

    return result;
}

int main() {
    char *original = "Hello, ";
    char *append = "World!";

    char *copy = duplicate_string(original);
    char *combined = append_strings(original, append);

    if (copy != NULL) {
        printf("Copy: %s\n", copy);
    }
}
```



```

        free(copy);
    }

    if (combined != NULL) {
        printf("Combined: %s\n", combined);
        free(combined);
    }

    return 0;
}

```

Dynamic Array with Resizing

```

#include <stdio.h>
#include <stdlib.h>

typedef struct {
    int *data;
    size_t size;
    size_t capacity;
} DynamicArray;

DynamicArray* create_array(size_t initial_capacity) {
    DynamicArray *arr = malloc(sizeof(DynamicArray));
    if (arr == NULL) return NULL;

    arr->data = malloc(initial_capacity * sizeof(int));
    if (arr->data == NULL) {
        free(arr);
        return NULL;
    }

    arr->size = 0;
    arr->capacity = initial_capacity;
    return arr;
}

void push_back(DynamicArray *arr, int value) {
    if (arr->size >= arr->capacity) {
        // Double the capacity
        size_t new_capacity = arr->capacity * 2;
        int *temp = realloc(arr->data, new_capacity * sizeof(int));
        if (temp == NULL) return; // Handle allocation failure

        arr->data = temp;
    }
}

```

```

        arr->capacity = new_capacity;
    }

    arr->data[arr->size] = value;
    arr->size++;
}

void free_array(DynamicArray *arr) {
    if (arr != NULL) {
        free(arr->data);
        free(arr);
    }
}

int main() {
    DynamicArray *arr = create_array(2);
    if (arr == NULL) {
        printf("Array creation failed\n");
        return 1;
    }

    // Add elements (will trigger reallocation)
    for (int i = 0; i < 10; i++) {
        push_back(arr, i * i);
    }

    // Print elements
    printf("Array elements: ");
    for (size_t i = 0; i < arr->size; i++) {
        printf("%d ", arr->data[i]);
    }
    printf("\n");

    // Clean up
    free_array(arr);

    return 0;
}

```

Memory Debugging Tools

Valgrind (Linux/Unix)

Valgrind is a powerful tool for detecting memory errors:

```
# Compile with debugging information
gcc -g -Wall program.c -o program

# Run with Valgrind
valgrind --leak-check=full ./program
```

AddressSanitizer (GCC/Clang)

AddressSanitizer detects memory errors at runtime:

```
# Compile with AddressSanitizer
gcc -fsanitize=address -g program.c -o program

# Run program (errors will be reported automatically)
./program
```

Static Analysis Tools

Tools like cppcheck can detect potential memory issues:

```
# Static analysis
cppcheck --enable=all program.c
```

Common Memory Errors and Solutions

Memory Leaks

```
// Problem: Memory leak
void memory_leak() {
    int *ptr = malloc(100 * sizeof(int));
    // Missing free(ptr) - memory leak!
}

// Solution: Always free allocated memory
void no_memory_leak() {
    int *ptr = malloc(100 * sizeof(int));
    if (ptr != NULL) {
        // Use memory...
        free(ptr);
        ptr = NULL;
    }
}
```

Dangling Pointers

```
// Problem: Dangling pointer
void dangling_pointer() {
    int *ptr = malloc(sizeof(int));
    *ptr = 42;
    free(ptr);
    printf("%d\n", *ptr); // Undefined behavior!
}

// Solution: Set pointer to NULL after freeing
void no_dangling_pointer() {
    int *ptr = malloc(sizeof(int));
    *ptr = 42;
    free(ptr);
    ptr = NULL;
    // if (ptr != NULL) printf("%d\n", *ptr); // Safe check
}
```

Buffer Overflows

```
// Problem: Buffer overflow
void buffer_overflow() {
    int *arr = malloc(5 * sizeof(int));
    for (int i = 0; i < 10; i++) {
        arr[i] = i; // Writing beyond allocated memory
    }
    free(arr);
}

// Solution: Always check bounds
void no_buffer_overflow() {
    int *arr = malloc(5 * sizeof(int));
    for (int i = 0; i < 5; i++) { // Correct bound
        arr[i] = i;
    }
    free(arr);
}
```

Advanced Memory Management (C11 and Later)

aligned_alloc()

The `aligned_alloc()` function allocates aligned memory (C11):

```

#include <stdio.h>
#include <stdlib.h>

int main() {
    // Allocate 100 bytes aligned to 16-byte boundary
    void *ptr = aligned_alloc(16, 100);
    if (ptr == NULL) {
        printf("Aligned allocation failed\n");
        return 1;
    }

    printf("Allocated address: %p\n", ptr);
    printf("Alignment check: %s\n",
        ((uintptr_t)ptr % 16 == 0) ? "OK" : "FAIL");

    free(ptr);
    return 0;
}

```

Practical Examples

Memory Pool Implementation

```

#include <stdio.h>
#include <stdlib.h>
#include <string.h>

typedef struct {
    void *memory;
    size_t block_size;
    size_t total_blocks;
    size_t free_blocks;
    char *free_list;
} MemoryPool;

MemoryPool* create_pool(size_t block_size, size_t num_blocks) {
    MemoryPool *pool = malloc(sizeof(MemoryPool));
    if (pool == NULL) return NULL;

    pool->memory = malloc(block_size * num_blocks);
    if (pool->memory == NULL) {
        free(pool);
        return NULL;
    }
}

```

```

pool->block_size = block_size;
pool->total_blocks = num_blocks;
pool->free_blocks = num_blocks;

// Initialize free list
pool->free_list = malloc(num_blocks);
if (pool->free_list == NULL) {
    free(pool->memory);
    free(pool);
    return NULL;
}

for (size_t i = 0; i < num_blocks; i++) {
    pool->free_list[i] = 1; // 1 = free, 0 = allocated
}

return pool;
}

void* pool_alloc(MemoryPool *pool) {
    if (pool->free_blocks == 0) return NULL;

    for (size_t i = 0; i < pool->total_blocks; i++) {
        if (pool->free_list[i]) {
            pool->free_list[i] = 0;
            pool->free_blocks--;
            return (char*)pool->memory + (i * pool->block_size);
        }
    }

    return NULL;
}

void pool_free(MemoryPool *pool, void *ptr) {
    if (ptr == NULL || ptr < pool->memory ||
        ptr >= (char*)pool->memory + (pool->total_blocks * pool->block_size)) {
        return;
    }

    size_t index = ((char*)ptr - (char*)pool->memory) / pool->block_size;
    if (index < pool->total_blocks && !pool->free_list[index]) {
        pool->free_list[index] = 1;
        pool->free_blocks++;
    }
}

void destroy_pool(MemoryPool *pool) {

```

```

    if (pool != NULL) {
        free(pool->memory);
        free(pool->free_list);
        free(pool);
    }
}

int main() {
    // Create pool for 10 integers
    MemoryPool *pool = create_pool(sizeof(int), 10);
    if (pool == NULL) {
        printf("Pool creation failed\n");
        return 1;
    }

    // Allocate from pool
    int *numbers[5];
    for (int i = 0; i < 5; i++) {
        numbers[i] = (int*)pool_alloc(pool);
        if (numbers[i] != NULL) {
            *numbers[i] = i * 10;
            printf("Allocated %d at %p\n", *numbers[i], (void*)numbers[i]);
        }
    }

    // Free some allocations
    pool_free(pool, numbers[2]);
    pool_free(pool, numbers[4]);

    // Allocate again (should reuse freed blocks)
    int *new_num = (int*)pool_alloc(pool);
    if (new_num != NULL) {
        *new_num = 999;
        printf("Reallocated %d at %p\n", *new_num, (void*)new_num);
    }

    // Clean up
    destroy_pool(pool);

    return 0;
}

```

Smart Pointer-like Implementation

```
#include <stdio.h>
#include <stdlib.h>

typedef struct {
    void *ptr;
    size_t ref_count;
    void (*destructor)(void*);
} SmartPointer;

SmartPointer* create_smart_pointer(size_t size, void (*destructor)(void*)) {
    SmartPointer *sp = malloc(sizeof(SmartPointer));
    if (sp == NULL) return NULL;

    sp->ptr = malloc(size);
    if (sp->ptr == NULL) {
        free(sp);
        return NULL;
    }

    sp->ref_count = 1;
    sp->destructor = destructor;

    return sp;
}

SmartPointer* smart_pointer_retain(SmartPointer *sp) {
    if (sp != NULL) {
        sp->ref_count++;
    }
    return sp;
}

void smart_pointer_release(SmartPointer *sp) {
    if (sp != NULL) {
        sp->ref_count--;
        if (sp->ref_count == 0) {
            if (sp->destructor != NULL) {
                sp->destructor(sp->ptr);
            } else {
                free(sp->ptr);
            }
            free(sp);
        }
    }
}
```



```

int main() {
    // Create smart pointer for integer
    SmartPointer *sp = create_smart_pointer(sizeof(int), NULL);
    if (sp == NULL) {
        printf("Smart pointer creation failed\n");
        return 1;
    }

    // Use the smart pointer
    *(int*)sp->ptr = 42;
    printf("Value: %d\n", *(int*)sp->ptr);

    // Create another reference
    SmartPointer *sp2 = smart_pointer_retain(sp);
    printf("Reference count: %zu\n", sp->ref_count);

    // Release first reference
    smart_pointer_release(sp);
    printf("Reference count after release: %zu\n", sp2->ref_count);

    // Release second reference (triggers deallocation)
    smart_pointer_release(sp2);

    return 0;
}

```

Summary

Dynamic memory management is a powerful feature of C that provides flexibility for handling data structures whose size is unknown at compile time. Key points to remember:

1. **Memory Allocation Functions:** malloc, calloc, realloc, and free for memory management
2. **Memory Layout:** Understanding stack vs. heap allocation
3. **Error Handling:** Always check for allocation failures
4. **Memory Safety:** Prevent memory leaks, dangling pointers, and buffer overflows
5. **Best Practices:** Set pointers to NULL after freeing, check bounds, and use debugging tools
6. **Advanced Features:** aligned_alloc for aligned memory (C11+)
7. **Debugging Tools:** Valgrind, AddressSanitizer, and static analysis tools

Proper dynamic memory management is essential for writing robust, efficient C programs and avoiding common memory-related bugs.

Advanced Pointer Concepts

Introduction

Advanced pointer concepts in C extend beyond basic pointer usage to include sophisticated techniques like pointers to pointers, function pointers, void pointers, and const correctness. These concepts enable powerful programming paradigms such as generic programming, callback mechanisms, and complex data structures. Understanding these advanced concepts is essential for writing flexible, reusable, and efficient C code.

Pointers to Pointers

Basic Concept

A pointer to a pointer is a variable that stores the address of another pointer:

```
#include <stdio.h>

int main() {
    int x = 42;
    int *ptr = &x;          // Pointer to integer
    int **ptr_to_ptr = &ptr; // Pointer to pointer to integer

    printf("Value of x: %d\n", x);
    printf("Value via ptr: %d\n", *ptr);
    printf("Value via ptr_to_ptr: %d\n", **ptr_to_ptr);

    // Modifying through pointer to pointer
    **ptr_to_ptr = 100;
    printf("After modification, x = %d\n", x);

    return 0;
}
```

Practical Applications

Pointers to pointers are commonly used for:

1. Modifying pointers in functions

2. Dynamic memory allocation for pointers
3. Implementing complex data structures

```
#include <stdio.h>
#include <stdlib.h>

// Function to allocate memory and update pointer
void allocate_memory(int **ptr, size_t size) {
    *ptr = malloc(size * sizeof(int));
    if (*ptr != NULL) {
        for (size_t i = 0; i < size; i++) {
            (*ptr)[i] = (int)i;
        }
    }
}

// Function to swap two pointers
void swap_pointers(int **ptr1, int **ptr2) {
    int *temp = *ptr1;
    *ptr1 = *ptr2;
    *ptr2 = temp;
}

int main() {
    int *arr1 = NULL;
    int *arr2 = NULL;

    // Allocate memory through pointer to pointer
    allocate_memory(&arr1, 5);
    allocate_memory(&arr2, 3);

    if (arr1 != NULL && arr2 != NULL) {
        printf("Array 1: ");
        for (int i = 0; i < 5; i++) {
            printf("%d ", arr1[i]);
        }
        printf("\n");

        printf("Array 2: ");
        for (int i = 0; i < 3; i++) {
            printf("%d ", arr2[i]);
        }
        printf("\n");

        // Swap pointers
        swap_pointers(&arr1, &arr2);
    }
}
```

```

        printf("After swap:\n");
        printf("Array 1: ");
        for (int i = 0; i < 3; i++) {
            printf("%d ", arr1[i]);
        }
        printf("\n");

        // Free memory
        free(arr1);
        free(arr2);
    }

    return 0;
}

```

Function Pointers

Basic Declaration and Usage

Function pointers store the address of functions and can be used to call functions indirectly:

```

#include <stdio.h>

// Function prototypes
int add(int a, int b) { return a + b; }
int subtract(int a, int b) { return a - b; }
int multiply(int a, int b) { return a * b; }

int main() {
    // Function pointer declaration
    int (*operation)(int, int);

    // Assign function addresses
    operation = add;
    printf("5 + 3 = %d\n", operation(5, 3));

    operation = subtract;
    printf("5 - 3 = %d\n", operation(5, 3));

    operation = multiply;
    printf("5 * 3 = %d\n", operation(5, 3));

    return 0;
}

```

Function Pointer Arrays

Function pointers can be stored in arrays for dispatch tables:

```
#include <stdio.h>

// Calculator operations
int add(int a, int b) { return a + b; }
int subtract(int a, int b) { return a - b; }
int multiply(int a, int b) { return a * b; }
int divide(int a, int b) { return b != 0 ? a / b : 0; }

int main() {
    // Array of function pointers
    int (*operations[])(int, int) = {add, subtract, multiply, divide};
    const char *op_names[] = {"+", "-", "*", "/"};

    int a = 10, b = 5;

    // Use function pointer array
    for (int i = 0; i < 4; i++) {
        int result = operations[i](a, b);
        printf("%d %s %d = %d\n", a, op_names[i], b, result);
    }

    return 0;
}
```

Callback Functions

Function pointers enable callback mechanisms:

```
#include <stdio.h>

// Callback function type
typedef void (*callback_func)(int value);

// Function that takes a callback
void process_array(int *arr, size_t size, callback_func callback) {
    for (size_t i = 0; i < size; i++) {
        callback(arr[i]);
    }
}

// Callback functions
void print_value(int value) {
```

```

    printf("Value: %d\n", value);
}

void square_value(int value) {
    printf("Square: %d\n", value * value);
}

int main() {
    int numbers[] = {1, 2, 3, 4, 5};
    size_t size = sizeof(numbers) / sizeof(numbers[0]);

    printf("Printing values:\n");
    process_array(numbers, size, print_value);

    printf("\nPrinting squares:\n");
    process_array(numbers, size, square_value);

    return 0;
}

```

Function Pointers as Parameters

Function pointers are commonly used as parameters for generic algorithms:

```

#include <stdio.h>

// Comparison function type
typedef int (*compare_func)(const void *a, const void *b);

// Generic sorting function
void bubble_sort(void *arr, size_t count, size_t size, compare_func compare) {
    char *array = (char*)arr;

    for (size_t i = 0; i < count - 1; i++) {
        for (size_t j = 0; j < count - i - 1; j++) {
            char *elem1 = array + j * size;
            char *elem2 = array + (j + 1) * size;

            if (compare(elem1, elem2) > 0) {
                // Swap elements
                for (size_t k = 0; k < size; k++) {
                    char temp = elem1[k];
                    elem1[k] = elem2[k];
                    elem2[k] = temp;
                }
            }
        }
    }
}

```

```

    }
}

// Comparison functions
int compare_ints(const void *a, const void *b) {
    int int_a = *(const int*)a;
    int int_b = *(const int*)b;
    return (int_a > int_b) - (int_a < int_b);
}

int compare_strings(const void *a, const void *b) {
    const char *str_a = *(const char**)a;
    const char *str_b = *(const char**)b;
    return strcmp(str_a, str_b);
}

int main() {
    // Sort integers
    int numbers[] = {64, 34, 25, 12, 22, 11, 90};
    size_t num_count = sizeof(numbers) / sizeof(numbers[0]);

    printf("Original numbers: ");
    for (size_t i = 0; i < num_count; i++) {
        printf("%d ", numbers[i]);
    }
    printf("\n");

    bubble_sort(numbers, num_count, sizeof(int), compare_ints);

    printf("Sorted numbers: ");
    for (size_t i = 0; i < num_count; i++) {
        printf("%d ", numbers[i]);
    }
    printf("\n");

    return 0;
}

```

Void Pointers

Basic Concept

Void pointers can point to any data type but cannot be dereferenced directly:

```

#include <stdio.h>

int main() {
    int x = 42;
    float y = 3.14f;
    char z = 'A';

    // Void pointer can point to any type
    void *ptr;

    ptr = &x;
    printf("Integer: %d\n", *(int*)ptr);

    ptr = &y;
    printf("Float: %.2f\n", *(float*)ptr);

    ptr = &z;
    printf("Character: %c\n", *(char*)ptr);

    return 0;
}

```

Generic Functions with Void Pointers

Void pointers enable generic programming:

```

#include <stdio.h>
#include <stdlib.h>
#include <string.h>

// Generic swap function
void generic_swap(void *a, void *b, size_t size) {
    char *ptr_a = (char*)a;
    char *ptr_b = (char*)b;

    for (size_t i = 0; i < size; i++) {
        char temp = ptr_a[i];
        ptr_a[i] = ptr_b[i];
        ptr_b[i] = temp;
    }
}

// Generic copy function
void generic_copy(void *dest, const void *src, size_t size) {
    char *d = (char*)dest;
    const char *s = (const char*)src;

```



```

    for (size_t i = 0; i < size; i++) {
        d[i] = s[i];
    }
}

int main() {
    // Swap integers
    int a = 10, b = 20;
    printf("Before swap: a = %d, b = %d\n", a, b);
    generic_swap(&a, &b, sizeof(int));
    printf("After swap: a = %d, b = %d\n", a, b);

    // Swap doubles
    double x = 3.14, y = 2.71;
    printf("Before swap: x = %.2f, y = %.2f\n", x, y);
    generic_swap(&x, &y, sizeof(double));
    printf("After swap: x = %.2f, y = %.2f\n", x, y);

    // Copy structure
    struct {
        int id;
        char name[20];
    } person1 = {1, "John"}, person2;

    generic_copy(&person2, &person1, sizeof(person1));
    printf("Copied person: ID = %d, Name = %s\n", person2.id, person2.name);

    return 0;
}

```

Void Pointer Arrays

Arrays of void pointers can store pointers to different types:

```

#include <stdio.h>

int main() {
    int x = 42;
    float y = 3.14f;
    char z = 'A';

    // Array of void pointers
    void *array[3] = {&x, &y, &z};

    // Access elements with proper casting

```

```

printf("Integer: %d\n", *(int*)array[0]);
printf("Float: %.2f\n", *(float*)array[1]);
printf("Character: %c\n", *(char*)array[2]);

return 0;
}

```

Const Correctness with Pointers

Pointer to Constant

A pointer to constant prevents modification of the pointed-to value:

```

#include <stdio.h>

int main() {
    int x = 42;
    int y = 100;

    // Pointer to constant - cannot modify value
    const int *ptr_to_const = &x;

    // This is OK - can change where pointer points
    ptr_to_const = &y;

    // This is ERROR - cannot modify value through pointer
    // *ptr_to_const = 50; // Compilation error

    printf("Value: %d\n", *ptr_to_const);

    return 0;
}

```

Constant Pointer

A constant pointer cannot be changed to point to a different location:

```

#include <stdio.h>

int main() {
    int x = 42;
    int y = 100;

    // Constant pointer - cannot change pointer value

```

```

int *const const_ptr = &x;

// This is OK - can modify value through pointer
*const_ptr = 50;

// This is ERROR - cannot change where pointer points
// const_ptr = &y; // Compilation error

printf("Value: %d\n", *const_ptr);

return 0;
}

```

Constant Pointer to Constant

Both the pointer and the pointed-to value are constant:

```

#include <stdio.h>

int main() {
    int x = 42;
    int y = 100;

    // Constant pointer to constant - nothing can be changed
    const int *const const_ptr_to_const = &x;

    // Both of these are ERROR
    // *const_ptr_to_const = 50; // Cannot modify value
    // const_ptr_to_const = &y; // Cannot change pointer

    printf("Value: %d\n", *const_ptr_to_const);

    return 0;
}

```

Const Correctness in Function Parameters

Using const in function parameters improves code safety and clarity:

```

#include <stdio.h>
#include <string.h>

// Function that doesn't modify the string
size_t safe_string_length(const char *str) {
    // str is a pointer to constant - cannot modify the string
}

```

```

    return strlen(str);
}

// Function that modifies the string
void unsafe_string_uppercase(char *str) {
    // str is a pointer to non-constant - can modify the string
    for (int i = 0; str[i] != '\0'; i++) {
        if (str[i] >= 'a' && str[i] <= 'z') {
            str[i] = str[i] - 'a' + 'A';
        }
    }
}

// Function that takes a constant pointer to array
void print_array(const int *arr, size_t size) {
    // arr is a pointer to constant - cannot modify array elements
    for (size_t i = 0; i < size; i++) {
        printf("%d ", arr[i]);
    }
    printf("\n");
}

int main() {
    const char *text = "Hello, World!";
    char mutable_text[] = "hello, world!";

    printf("Length: %zu\n", safe_string_length(text));

    printf("Before: %s\n", mutable_text);
    unsafe_string_uppercase(mutable_text);
    printf("After: %s\n", mutable_text);

    int numbers[] = {1, 2, 3, 4, 5};
    print_array(numbers, 5);

    return 0;
}

```

Advanced Pointer Techniques

Pointer Arithmetic with Structs

Pointer arithmetic can be used with structures:

```

#include <stdio.h>
#include <stdlib.h>

typedef struct {
    int id;
    char name[20];
    float score;
} Student;

int main() {
    // Array of structures
    Student class[3] = {
        {1, "Alice", 95.5},
        {2, "Bob", 87.0},
        {3, "Charlie", 92.5}
    };

    // Pointer to first element
    Student *ptr = class;

    // Access elements using pointer arithmetic
    for (int i = 0; i < 3; i++) {
        printf("Student %d: ID=%d, Name=%s, Score=%.1f\n",
            i + 1, ptr->id, ptr->name, ptr->score);
        ptr++; // Move to next structure
    }

    return 0;
}

```

Flexible Array Members (C99)

Structures can have flexible array members:

```

#include <stdio.h>
#include <stdlib.h>
#include <string.h>

typedef struct {
    int count;
    int data[]; // Flexible array member (C99)
} DynamicArray;

DynamicArray* create_dynamic_array(int count) {
    // Allocate memory for structure + array
    DynamicArray *arr = malloc(sizeof(DynamicArray) + count * sizeof(int));
}

```

```

    if (arr != NULL) {
        arr->count = count;
        // arr->data can now be used as array of 'count' integers
    }
    return arr;
}

int main() {
    DynamicArray *arr = create_dynamic_array(5);
    if (arr != NULL) {
        // Initialize array elements
        for (int i = 0; i < arr->count; i++) {
            arr->data[i] = i * i;
        }

        // Print array elements
        printf("Array elements: ");
        for (int i = 0; i < arr->count; i++) {
            printf("%d ", arr->data[i]);
        }
        printf("\n");

        // Free memory
        free(arr);
    }

    return 0;
}

```

Restricted Pointers (C99)

The `restrict` keyword tells the compiler that a pointer is the only way to access the object:

```

#include <stdio.h>

// Function with restricted pointers
void add_arrays(int *restrict result,
               const int *restrict a,
               const int *restrict b,
               size_t size) {
    // Compiler knows a, b, and result don't overlap
    for (size_t i = 0; i < size; i++) {
        result[i] = a[i] + b[i];
    }
}

```

```

int main() {
    int a[] = {1, 2, 3, 4, 5};
    int b[] = {10, 20, 30, 40, 50};
    int result[5];

    add_arrays(result, a, b, 5);

    printf("Result: ");
    for (int i = 0; i < 5; i++) {
        printf("%d ", result[i]);
    }
    printf("\n");

    return 0;
}

```

Practical Examples

Generic Linked List Implementation

```

#include <stdio.h>
#include <stdlib.h>

typedef struct Node {
    void *data;
    struct Node *next;
} Node;

typedef struct {
    Node *head;
    size_t data_size;
} LinkedList;

LinkedList* create_list(size_t data_size) {
    LinkedList *list = malloc(sizeof(LinkedList));
    if (list != NULL) {
        list->head = NULL;
        list->data_size = data_size;
    }
    return list;
}

void list_append(LinkedList *list, const void *data) {
    Node *new_node = malloc(sizeof(Node));

```

```

    if (new_node == NULL) return;

    new_node->data = malloc(list->data_size);
    if (new_node->data == NULL) {
        free(new_node);
        return;
    }

    // Copy data
    char *src = (char*)data;
    char *dest = (char*)new_node->data;
    for (size_t i = 0; i < list->data_size; i++) {
        dest[i] = src[i];
    }

    new_node->next = NULL;

    // Add to end of list
    if (list->head == NULL) {
        list->head = new_node;
    } else {
        Node *current = list->head;
        while (current->next != NULL) {
            current = current->next;
        }
        current->next = new_node;
    }
}

void list_print_int(LinkedList *list) {
    Node *current = list->head;
    while (current != NULL) {
        printf("%d ", *(int*)current->data);
        current = current->next;
    }
    printf("\n");
}

void list_free(LinkedList *list) {
    Node *current = list->head;
    while (current != NULL) {
        Node *next = current->next;
        free(current->data);
        free(current);
        current = next;
    }
    free(list);
}

```



```

}

int main() {
    LinkedList *int_list = create_list(sizeof(int));

    // Add integers to list
    for (int i = 1; i <= 5; i++) {
        list_append(int_list, &i);
    }

    printf("List contents: ");
    list_print_int(int_list);

    list_free(int_list);

    return 0;
}

```

Event System with Function Pointers

```

#include <stdio.h>
#include <stdlib.h>

// Event types
typedef enum {
    EVENT_CLICK,
    EVENT_KEYPRESS,
    EVENT_MOUSEMOVE,
    EVENT_COUNT
} EventType;

// Event data structures
typedef struct {
    int x, y;
} ClickEvent;

typedef struct {
    char key;
} KeyPressEvent;

typedef struct {
    int x, y;
} MouseMoveEvent;

// Event handler function type

```

```

typedef void (*EventHandler)(void *event_data);

// Event system structure
typedef struct {
    EventHandler handlers[EVENT_COUNT];
} EventSystem;

// Event handler functions
void handle_click(void *data) {
    ClickEvent *event = (ClickEvent*)data;
    printf("Click at (%d, %d)\n", event->x, event->y);
}

void handle_keypress(void *data) {
    KeyPressEvent *event = (KeyPressEvent*)data;
    printf("Key pressed: %c\n", event->key);
}

void handle_mousemove(void *data) {
    MouseMoveEvent *event = (MouseMoveEvent*)data;
    printf("Mouse moved to (%d, %d)\n", event->x, event->y);
}

// Event system functions
EventSystem* create_event_system() {
    EventSystem *system = malloc(sizeof(EventSystem));
    if (system != NULL) {
        for (int i = 0; i < EVENT_COUNT; i++) {
            system->handlers[i] = NULL;
        }
    }
    return system;
}

void register_handler(EventSystem *system, EventType type, EventHandler handler) {
    if (type >= 0 && type < EVENT_COUNT) {
        system->handlers[type] = handler;
    }
}

void trigger_event(EventSystem *system, EventType type, void *event_data) {
    if (type >= 0 && type < EVENT_COUNT && system->handlers[type] != NULL) {
        system->handlers[type](event_data);
    }
}

```

```

void destroy_event_system(EventSystem *system) {
    free(system);
}

int main() {
    EventSystem *events = create_event_system();

    // Register event handlers
    register_handler(events, EVENT_CLICK, handle_click);
    register_handler(events, EVENT_KEYPRESS, handle_keypress);
    register_handler(events, EVENT_MOUSEMOVE, handle_mousemove);

    // Trigger events
    ClickEvent click = {100, 200};
    trigger_event(events, EVENT_CLICK, &click);

    KeyPressEvent keypress = {'A'};
    trigger_event(events, EVENT_KEYPRESS, &keypress);

    MouseMoveEvent mousemove = {50, 75};
    trigger_event(events, EVENT_MOUSEMOVE, &mousemove);

    destroy_event_system(events);

    return 0;
}

```

Summary

Advanced pointer concepts in C provide powerful mechanisms for creating flexible and efficient programs. Key points to remember:

1. **Pointers to Pointers:** Enable modification of pointers in functions and complex data structures
2. **Function Pointers:** Enable callback mechanisms, dispatch tables, and generic algorithms
3. **Void Pointers:** Enable generic programming and storage of different data types
4. **Const Correctness:** Improves code safety and communicates intent to other developers
5. **Advanced Techniques:** Flexible array members, restricted pointers, and pointer arithmetic with structures

These advanced concepts are essential for writing professional C code, implementing complex data structures, and creating reusable libraries. Understanding them enables programmers to leverage the full power of C's pointer system while maintaining code safety and clarity.

Common Pointer Pitfalls

Introduction

Pointers are one of the most powerful features of the C programming language, but they are also one of the most error-prone. Incorrect pointer usage can lead to subtle bugs, program crashes, security vulnerabilities, and undefined behavior. Understanding common pointer pitfalls and how to avoid them is crucial for writing robust, reliable C programs. This chapter explores the most frequent pointer-related errors and provides strategies for preventing them.

Dangling Pointers

What Are Dangling Pointers?

A dangling pointer is a pointer that points to memory that has been deallocated or is no longer valid. Dereferencing a dangling pointer leads to undefined behavior.

Common Causes

1. Premature Deallocation

```
#include <stdio.h>
#include <stdlib.h>

int* create_integer(int value) {
    int *ptr = malloc(sizeof(int));
    if (ptr != NULL) {
        *ptr = value;
    }
    return ptr;
}

int main() {
    int *ptr = create_integer(42);

    // Problem: Using pointer after deallocation
    free(ptr);
    printf("Value: %d\n", *ptr); // Dangling pointer - undefined behavior!
```

```
    return 0;
}
```

2. Returning Pointer to Local Variable

```
#include <stdio.h>

int* get_pointer() {
    int local_var = 42;
    return &local_var; // Problem: Returning pointer to local variable
}

int main() {
    int *ptr = get_pointer();
    printf("Value: %d\n", *ptr); // Dangling pointer - undefined behavior!

    return 0;
}
```

Prevention Strategies

1. Set Pointers to NULL After Freeing

```
#include <stdio.h>
#include <stdlib.h>

int main() {
    int *ptr = malloc(sizeof(int));
    if (ptr != NULL) {
        *ptr = 42;
        printf("Value: %d\n", *ptr);

        free(ptr);
        ptr = NULL; // Prevent dangling pointer

        // Safe check before use
        if (ptr != NULL) {
            printf("Value: %d\n", *ptr);
        } else {
            printf("Pointer is NULL\n");
        }
    }

    return 0;
}
```

```
}
```

2. Avoid Returning Pointers to Local Variables

```
#include <stdio.h>
#include <stdlib.h>

// Solution 1: Use dynamic allocation
int* create_integer_safe(int value) {
    int *ptr = malloc(sizeof(int));
    if (ptr != NULL) {
        *ptr = value;
    }
    return ptr;
}

// Solution 2: Use output parameters
void get_value(int *output) {
    *output = 42; // Modify through parameter
}

int main() {
    // Using dynamic allocation
    int *ptr = create_integer_safe(42);
    if (ptr != NULL) {
        printf("Value: %d\n", *ptr);
        free(ptr);
    }

    // Using output parameter
    int value;
    get_value(&value);
    printf("Value: %d\n", value);

    return 0;
}
```

Memory Leaks

What Are Memory Leaks?

A memory leak occurs when dynamically allocated memory is not properly deallocated, making it unavailable for reuse during program execution.

Common Causes

1. Forgetting to Call free()

```
#include <stdio.h>
#include <stdlib.h>

void memory_leak_example() {
    int *ptr = malloc(100 * sizeof(int));
    if (ptr != NULL) {
        // Use the memory...
        for (int i = 0; i < 100; i++) {
            ptr[i] = i;
        }

        // Problem: Forgot to call free() - memory leak!
        // free(ptr);
    }
}

int main() {
    for (int i = 0; i < 1000; i++) {
        memory_leak_example(); // Leaks 400 bytes each iteration
    }

    return 0;
}
```

2. Losing Pointer References

```
#include <stdio.h>
#include <stdlib.h>

int main() {
    int *ptr1 = malloc(100 * sizeof(int));
    int *ptr2 = malloc(100 * sizeof(int));

    if (ptr1 != NULL && ptr2 != NULL) {
        // Use memory...

        // Problem: Losing reference to ptr1 - memory leak!
        ptr1 = ptr2;

        free(ptr2); // Only frees one block
        // ptr1's memory is leaked
    }
}
```

```

    }

    return 0;
}

```

Prevention Strategies

1. Use RAII-like Patterns

```

#include <stdio.h>
#include <stdlib.h>

typedef struct {
    int *data;
    size_t size;
} ManagedArray;

ManagedArray* create_array(size_t size) {
    ManagedArray *arr = malloc(sizeof(ManagedArray));
    if (arr != NULL) {
        arr->data = malloc(size * sizeof(int));
        if (arr->data == NULL) {
            free(arr);
            return NULL;
        }
        arr->size = size;
    }
    return arr;
}

void destroy_array(ManagedArray *arr) {
    if (arr != NULL) {
        free(arr->data);
        free(arr);
    }
}

int main() {
    ManagedArray *arr = create_array(100);
    if (arr != NULL) {
        // Use array...

        // Always clean up
        destroy_array(arr);
    }
}

```



```
    return 0;
}
```

2. Match Every malloc with free

```
#include <stdio.h>
#include <stdlib.h>

void safe_memory_usage() {
    int *ptr = malloc(100 * sizeof(int));
    if (ptr == NULL) {
        return; // Handle allocation failure
    }

    // Use the memory...
    for (int i = 0; i < 100; i++) {
        ptr[i] = i;
    }

    // Always free allocated memory
    free(ptr);
    ptr = NULL; // Prevent dangling pointer
}
```

Buffer Overflows

What Are Buffer Overflows?

A buffer overflow occurs when more data is written to a buffer than it can hold, potentially corrupting adjacent memory.

Common Causes

1. Array Index Out of Bounds

```
#include <stdio.h>

int main() {
    int arr[5] = {1, 2, 3, 4, 5};

    // Problem: Writing beyond array bounds
    for (int i = 0; i < 10; i++) { // i goes to 9, but array only has 5 elements
```

```

        arr[i] = i * i; // Buffer overflow when i >= 5
    }

    return 0;
}

```

2. String Buffer Overflows

```

#include <stdio.h>
#include <string.h>

int main() {
    char buffer[10];

    // Problem: String too long for buffer
    strcpy(buffer, "This string is too long for the buffer!");
    // Buffer overflow - undefined behavior!

    return 0;
}

```

Prevention Strategies

1. Always Check Array Bounds

```

#include <stdio.h>

int main() {
    int arr[5] = {1, 2, 3, 4, 5};
    int size = sizeof(arr) / sizeof(arr[0]);

    // Safe: Check bounds before access
    for (int i = 0; i < size; i++) {
        arr[i] = i * i;
    }

    // Alternative: Use safer loop
    for (int i = 0; i < 5; i++) {
        arr[i] = i * i;
    }

    return 0;
}

```

2. Use Safe String Functions

```
#include <stdio.h>
#include <string.h>

int main() {
    char buffer[10];

    // Safe: Use strncpy with proper bounds checking
    const char *source = "This string is too long";
    strncpy(buffer, source, sizeof(buffer) - 1);
    buffer[sizeof(buffer) - 1] = '\0'; // Ensure null termination

    printf("Buffer: %s\n", buffer);

    return 0;
}
```

Uninitialized Pointers

What Are Uninitialized Pointers?

An uninitialized pointer contains garbage values and points to random memory locations. Dereferencing such pointers leads to undefined behavior.

Common Causes

1. Using Uninitialized Local Pointers

```
#include <stdio.h>

int main() {
    int *ptr; // Uninitialized pointer

    // Problem: Using uninitialized pointer
    *ptr = 42; // Undefined behavior!

    return 0;
}
```

2. Partial Initialization

```
#include <stdio.h>

int main() {
    int *ptr1, *ptr2;

    ptr1 = malloc(sizeof(int));
    if (ptr1 != NULL) {
        *ptr1 = 42;
    }

    // Problem: ptr2 is uninitialized
    *ptr2 = 100; // Undefined behavior!

    free(ptr1);

    return 0;
}
```

Prevention Strategies

1. Always Initialize Pointers

```
#include <stdio.h>
#include <stdlib.h>

int main() {
    // Good: Initialize to NULL
    int *ptr = NULL;

    // Safe: Check before use
    if (ptr != NULL) {
        *ptr = 42;
    }

    // Or initialize with valid address
    int value = 42;
    int *ptr2 = &value; // Safe initialization

    printf("Value: %d\n", *ptr2);

    return 0;
}
```

2. Use Separate Declarations

```
#include <stdio.h>
#include <stdlib.h>

int main() {
    // Good: Separate declarations for clarity
    int *ptr1 = NULL;
    int *ptr2 = NULL;

    ptr1 = malloc(sizeof(int));
    if (ptr1 != NULL) {
        *ptr1 = 42;
    }

    ptr2 = malloc(sizeof(int));
    if (ptr2 != NULL) {
        *ptr2 = 100;
    }

    // Safe usage
    if (ptr1 != NULL) {
        printf("ptr1: %d\n", *ptr1);
    }

    if (ptr2 != NULL) {
        printf("ptr2: %d\n", *ptr2);
    }

    // Clean up
    free(ptr1);
    free(ptr2);

    return 0;
}
```

Pointer Arithmetic Errors

Common Mistakes

1. Incorrect Pointer Arithmetic

```
#include <stdio.h>

int main() {
    int arr[5] = {1, 2, 3, 4, 5};
    int *ptr = arr;

    // Problem: Incorrect arithmetic
    ptr = ptr + 10; // Points beyond array bounds

    // Undefined behavior when dereferencing
    // printf("Value: %d\n", *ptr);

    return 0;
}
```

2. Pointer Arithmetic with Different Types

```
#include <stdio.h>

int main() {
    int arr[5] = {1, 2, 3, 4, 5};
    char *char_ptr = (char*)arr;

    // Problem: Different pointer arithmetic
    printf("int ptr + 1: %p\n", (void*)(arr + 1));
    printf("char ptr + 1: %p\n", (void*)(char_ptr + 1));
    // These point to different memory locations!

    return 0;
}
```

Prevention Strategies

1. Validate Pointer Arithmetic

```
#include <stdio.h>

int main() {
    int arr[5] = {1, 2, 3, 4, 5};
    int *ptr = arr;
    int size = sizeof(arr) / sizeof(arr[0]);

    // Safe: Validate bounds before arithmetic
    int offset = 2;
    if (offset >= 0 && offset < size) {
        int *new_ptr = ptr + offset;
        printf("Value at offset %d: %d\n", offset, *new_ptr);
    } else {
        printf("Offset out of bounds\n");
    }

    return 0;
}
```

NULL Pointer Dereferencing

What Is NULL Pointer Dereferencing?

Dereferencing a NULL pointer leads to program crashes or undefined behavior.

Common Causes

1. Forgetting to Check Return Values

```
#include <stdio.h>
#include <stdlib.h>

int main() {
    int *ptr = malloc(100 * sizeof(int));

    // Problem: Not checking for NULL
    *ptr = 42; // Crash if malloc failed!

    free(ptr);
}
```

```
    return 0;
}
```

Prevention Strategies

1. Always Check for NULL

```
#include <stdio.h>
#include <stdlib.h>

int main() {
    int *ptr = malloc(100 * sizeof(int));

    // Safe: Check for NULL
    if (ptr != NULL) {
        *ptr = 42;
        printf("Value: %d\n", *ptr);
        free(ptr);
    } else {
        printf("Memory allocation failed\n");
        return 1;
    }

    return 0;
}
```

Type Safety Issues

Common Problems

1. Incorrect Casting

```
#include <stdio.h>

int main() {
    int x = 42;
    int *int_ptr = &x;

    // Problem: Incorrect casting
    char *char_ptr = (char*)int_ptr;
    printf("Char value: %c\n", *char_ptr); // Interprets as character

    return 0;
}
```



```
}
```

2. Function Pointer Mismatches

```
#include <stdio.h>

int add(int a, int b) { return a + b; }
void print_message(void) { printf("Hello\n"); }

int main() {
    // Problem: Wrong function pointer type
    int (*func_ptr)(int, int) = (int (*)(int, int))print_message;

    // Undefined behavior when called
    // int result = func_ptr(5, 3);

    return 0;
}
```

Prevention Strategies

1. Use Proper Typedefs

```
#include <stdio.h>

// Safe: Use typedef for function pointers
typedef int (*binary_op)(int, int);

int add(int a, int b) { return a + b; }
int subtract(int a, int b) { return a - b; }

int main() {
    binary_op operations[] = {add, subtract};

    int result1 = operations[0](10, 5); // Safe: add(10, 5)
    int result2 = operations[1](10, 5); // Safe: subtract(10, 5)

    printf("10 + 5 = %d\n", result1);
    printf("10 - 5 = %d\n", result2);

    return 0;
}
```

Debugging Pointer Issues

Using Debugging Tools

1. Valgrind (Linux/Unix)

```
# Compile with debugging information
gcc -g -Wall program.c -o program

# Run with Valgrind
valgrind --leak-check=full ./program
```

2. AddressSanitizer (GCC/Clang)

```
# Compile with AddressSanitizer
gcc -fsanitize=address -g -Wall program.c -o program

# Run program (errors will be reported automatically)
./program
```

Static Analysis Tools

```
# Static analysis with cppcheck
cppcheck --enable=all program.c

# Static analysis with clang-analyzer
clang --analyze program.c
```

Best Practices for Pointer Safety

1. Initialize Pointers

```
// Good
int *ptr = NULL;

// Good
int value = 42;
int *ptr = &value;

// Bad
int *ptr; // Uninitialized
```

2. Check for NULL

```
int *ptr = malloc(sizeof(int));
if (ptr != NULL) {
    *ptr = 42;
    // Use ptr...
    free(ptr);
    ptr = NULL;
}
```

3. Validate Array Bounds

```
int arr[10];
int size = sizeof(arr) / sizeof(arr[0]);

for (int i = 0; i < size; i++) {
    arr[i] = i;
}
```

4. Use Safe String Functions

```
char buffer[100];
const char *source = "Some text";

// Safe copy
strncpy(buffer, source, sizeof(buffer) - 1);
buffer[sizeof(buffer) - 1] = '\0';
```

5. Match Allocation with Deallocation

```
// For malloc/calloc/realloc
int *ptr = malloc(100 * sizeof(int));
// ... use ptr ...
free(ptr);
ptr = NULL;

// For arrays
int *arr = calloc(100, sizeof(int));
// ... use arr ...
free(arr);
arr = NULL;
```

Practical Example: Safe Dynamic Array

```
#include <stdio.h>
#include <stdlib.h>
#include <string.h>

typedef struct {
    int *data;
    size_t size;
    size_t capacity;
} SafeArray;

SafeArray* safe_array_create(size_t initial_capacity) {
    SafeArray *arr = malloc(sizeof(SafeArray));
    if (arr == NULL) {
        return NULL;
    }

    arr->data = malloc(initial_capacity * sizeof(int));
    if (arr->data == NULL) {
        free(arr);
        return NULL;
    }

    arr->size = 0;
    arr->capacity = initial_capacity;
    return arr;
}

int safe_array_push(SafeArray *arr, int value) {
    if (arr == NULL) {
        return -1; // Error: NULL pointer
    }

    // Resize if needed
    if (arr->size >= arr->capacity) {
        size_t new_capacity = arr->capacity * 2;
        int *temp = realloc(arr->data, new_capacity * sizeof(int));
        if (temp == NULL) {
            return -1; // Error: Reallocation failed
        }
        arr->data = temp;
        arr->capacity = new_capacity;
    }

    arr->data[arr->size] = value;
```

```

    arr->size++;
    return 0; // Success
}

int safe_array_get(SafeArray *arr, size_t index) {
    if (arr == NULL || index >= arr->size) {
        return 0; // Error: Invalid parameters
    }
    return arr->data[index];
}

void safe_array_destroy(SafeArray *arr) {
    if (arr != NULL) {
        free(arr->data);
        free(arr);
    }
}

int main() {
    SafeArray *arr = safe_array_create(2);
    if (arr == NULL) {
        printf("Array creation failed\n");
        return 1;
    }

    // Add elements
    for (int i = 0; i < 10; i++) {
        if (safe_array_push(arr, i * i) != 0) {
            printf("Array push failed\n");
            safe_array_destroy(arr);
            return 1;
        }
    }

    // Print elements
    printf("Array elements: ");
    for (size_t i = 0; i < arr->size; i++) {
        printf("%d ", safe_array_get(arr, i));
    }
    printf("\n");

    // Clean up
    safe_array_destroy(arr);

    return 0;
}

```

Summary

Pointer-related errors are among the most common and dangerous bugs in C programs. Key strategies for avoiding these pitfalls include:

1. **Dangling Pointers:** Always set pointers to NULL after freeing and avoid returning pointers to local variables
2. **Memory Leaks:** Match every malloc/calloc/realloc with a corresponding free
3. **Buffer Overflows:** Always validate array bounds and use safe string functions
4. **Uninitialized Pointers:** Initialize all pointers to NULL or valid addresses
5. **NULL Pointer Dereferencing:** Always check for NULL before dereferencing
6. **Pointer Arithmetic:** Validate bounds and understand type-specific arithmetic
7. **Type Safety:** Use proper casting and typedefs for function pointers

Using debugging tools like Valgrind and AddressSanitizer, along with static analysis tools, can help detect many of these issues during development. Following best practices and defensive programming techniques will lead to more robust and reliable C programs.

Module 6: Pointers and Memory Management

Overview

This module explores pointers and dynamic memory management in C, fundamental concepts for systems programming and efficient resource utilization. You'll learn about memory addresses, pointer arithmetic, dynamic memory allocation, and advanced pointer techniques. The module also covers common pointer pitfalls and best practices for safe memory management.

Learning Objectives

By the end of this module, you will be able to: - Understand memory addresses and pointer fundamentals - Perform pointer arithmetic and manipulation - Allocate and deallocate dynamic memory using malloc/free - Work with arrays and pointers interchangeably - Implement complex data structures using pointers - Apply function pointers for callbacks and dynamic dispatch - Understand void pointers and generic programming - Identify and prevent common pointer-related errors - Use const correctness with pointers - Apply advanced pointer concepts like pointers to pointers

Chapters

1. **Pointer Fundamentals** - Memory addresses, pointer declaration, initialization, and dereferencing
2. **Pointers and Arrays** - Array-pointer equivalence, pointer subscripting, and multi-dimensional arrays
3. **Dynamic Memory Management** - malloc, calloc, realloc, free, and memory debugging
4. **Advanced Pointer Concepts** - Pointers to pointers, function pointers, void pointers, and const correctness
5. **Common Pointer Pitfalls** - Dangling pointers, buffer overflows, segmentation faults, and memory leaks

Key Concepts Covered

- Understanding memory addresses and the & operator
- Pointer declaration, initialization, and dereferencing with *
- Pointer arithmetic and increment/decrement operations
- Array-pointer equivalence and pointer subscripting
- Dynamic memory allocation with malloc, calloc, realloc

- Memory deallocation with free and resource cleanup
- Function pointers and callback mechanisms
- Void pointers and generic programming techniques
- Const correctness with pointers
- Pointers to pointers and multi-level indirection
- Common pointer errors and debugging techniques
- Memory alignment and performance considerations

Prerequisites

- Completion of Module 1: Foundations and Environment Setup
- Completion of Module 2: Data Types and Variables
- Completion of Module 3: Control Flow
- Completion of Module 4: Functions and Modular Programming
- Completion of Module 5: Arrays and Strings
- Understanding of basic programming concepts
- Familiarity with variable declaration and initialization

Tools and Technologies

- C compiler (GCC, Clang, or MSVC)
- Text editor or IDE
- Memory debugging tools (Valgrind, AddressSanitizer)
- Understanding of basic programming concepts
- Familiarity with debugging tools

Estimated Time to Complete

- Reading: 10-12 hours
- Exercises: 15-20 hours
- Projects: 8-10 hours

Assessment

- Complete all chapter exercises
- Successfully implement dynamic data structures
- Demonstrate understanding of memory management concepts
- Pass pointer programming quiz
- Submit a comprehensive program using advanced pointer techniques

Next Module

After completing this module, proceed to **Module 7: Structures and User-Defined Types** to learn about creating custom data types.

Additional Resources

- “The C Programming Language” by Kernighan and Ritchie (Chapter 5)
- C standard documentation (ISO/IEC 9899:2024 for C23)
- Online tutorials on pointer arithmetic and memory management
- Debugging tools documentation (GDB, Valgrind, AddressSanitizer)
- Best practices guides for safe memory management

Structures

Structures

Introduction

Structures (structs) are user-defined data types in C that allow you to group related variables of different types under a single name. They are fundamental for creating complex data representations and implementing abstract data types. Structures enable programmers to model real-world entities and organize data in a logical, maintainable way.

Structure Declaration and Definition

Basic Structure Declaration

A structure is declared using the `struct` keyword followed by the structure name and a list of members enclosed in braces:

```
struct Person {  
    char name[50];  
    int age;  
    float height;  
    char gender;  
};
```

Structure Variable Declaration

Structure variables can be declared in several ways:

```
// Method 1: Declare structure and variables separately  
struct Person {  
    char name[50];  
    int age;  
    float height;  
    char gender;  
};  
  
struct Person person1, person2;  
  
// Method 2: Declare structure and variables together
```

```

struct {
    char name[50];
    int age;
    float height;
    char gender;
} person1, person2;

// Method 3: Use typedef for cleaner syntax
typedef struct {
    char name[50];
    int age;
    float height;
    char gender;
} Person;

Person person1, person2;

```

Structure Initialization

Designated Initializers (C99)

```

typedef struct {
    char name[50];
    int age;
    float height;
    char gender;
} Person;

// Initialize all members
Person person1 = {"John Doe", 25, 5.9, 'M'};

// Designated initializers (C99)
Person person2 = {
    .name = "Jane Smith",
    .age = 30,
    .height = 5.5,
    .gender = 'F'
};

// Partial initialization
Person person3 = {
    .name = "Bob Johnson",
    .age = 35
    // height and gender will be zero-initialized

```

```
};
```

Nested Structure Initialization

```
typedef struct {
    int day;
    int month;
    int year;
} Date;

typedef struct {
    char name[50];
    Date birth_date;
    float salary;
} Employee;

// Nested structure initialization
Employee emp = {
    .name = "Alice Brown",
    .birth_date = {15, 6, 1990},
    .salary = 75000.0
};

// Designated initialization with nested structures
Employee emp2 = {
    .name = "Charlie Wilson",
    .birth_date.day = 20,
    .birth_date.month = 3,
    .birth_date.year = 1985,
    .salary = 80000.0
};
```

Accessing Structure Members

Dot Operator (.)

The dot operator is used to access members of a structure variable:

```
typedef struct {
    char name[50];
    int age;
    float height;
    char gender;
```

```

} Person;

int main() {
    Person person = {"John Doe", 25, 5.9, 'M'};

    // Access members using dot operator
    printf("Name: %s\n", person.name);
    printf("Age: %d\n", person.age);
    printf("Height: %.1f\n", person.height);
    printf("Gender: %c\n", person.gender);

    // Modify members
    person.age = 26;
    person.height = 6.0;
    strcpy(person.name, "John Smith");

    return 0;
}

```

Arrow Operator (->)

The arrow operator is used to access members through a pointer to a structure:

```

typedef struct {
    char name[50];
    int age;
    float height;
    char gender;
} Person;

int main() {
    Person person = {"John Doe", 25, 5.9, 'M'};
    Person *ptr = &person;

    // Access members using arrow operator
    printf("Name: %s\n", ptr->name);
    printf("Age: %d\n", ptr->age);
    printf("Height: %.1f\n", ptr->height);
    printf("Gender: %c\n", ptr->gender);

    // Modify members through pointer
    ptr->age = 26;
    ptr->height = 6.0;
    strcpy(ptr->name, "John Smith");

    return 0;
}

```

```
}
```

Arrays of Structures

Declaration and Initialization

```
typedef struct {
    char name[50];
    int id;
    float gpa;
} Student;

int main() {
    // Array of structures
    Student students[3] = {
        {"Alice", 101, 3.8},
        {"Bob", 102, 3.5},
        {"Charlie", 103, 3.9}
    };

    // Access array elements
    for (int i = 0; i < 3; i++) {
        printf("Student %d: %s, ID: %d, GPA: %.1f\n",
            i + 1, students[i].name, students[i].id, students[i].gpa);
    }

    return 0;
}
```

Dynamic Array of Structures

```
#include <stdio.h>
#include <stdlib.h>
#include <string.h>

typedef struct {
    char name[50];
    int id;
    float gpa;
} Student;

int main() {
    int count = 3;
```

```

// Dynamic array of structures
Student *students = malloc(count * sizeof(Student));
if (students == NULL) {
    printf("Memory allocation failed\n");
    return 1;
}

// Initialize array elements
strcpy(students[0].name, "Alice");
students[0].id = 101;
students[0].gpa = 3.8;

strcpy(students[1].name, "Bob");
students[1].id = 102;
students[1].gpa = 3.5;

strcpy(students[2].name, "Charlie");
students[2].id = 103;
students[2].gpa = 3.9;

// Access array elements
for (int i = 0; i < count; i++) {
    printf("Student %d: %s, ID: %d, GPA: %.1f\n",
        i + 1, students[i].name, students[i].id, students[i].gpa);
}

// Free memory
free(students);

return 0;
}

```

Nested Structures

Basic Nested Structures

```

typedef struct {
    int day;
    int month;
    int year;
} Date;

typedef struct {
    char street[100];

```



```

    char city[50];
    char state[30];
    int zip_code;
} Address;

typedef struct {
    char name[50];
    Date birth_date;
    Address home_address;
    float salary;
} Employee;

int main() {
    Employee emp = {
        .name = "John Doe",
        .birth_date = {15, 6, 1990},
        .home_address = {
            .street = "123 Main St",
            .city = "Anytown",
            .state = "CA",
            .zip_code = 12345
        },
        .salary = 75000.0
    };

    // Access nested structure members
    printf("Name: %s\n", emp.name);
    printf("Birth Date: %d/%d/%d\n",
        emp.birth_date.month, emp.birth_date.day, emp.birth_date.year);
    printf("Address: %s, %s, %s %d\n",
        emp.home_address.street, emp.home_address.city,
        emp.home_address.state, emp.home_address.zip_code);
    printf("Salary: $%.2f\n", emp.salary);

    return 0;
}

```

Self-Referential Structures

Structures can contain pointers to their own type, enabling linked data structures:

```

typedef struct Node {
    int data;
    struct Node *next; // Pointer to same structure type
} Node;

```

```

int main() {
    // Create nodes
    Node node1 = {10, NULL};
    Node node2 = {20, NULL};
    Node node3 = {30, NULL};

    // Link nodes
    node1.next = &node2;
    node2.next = &node3;

    // Traverse linked list
    Node *current = &node1;
    while (current != NULL) {
        printf("Data: %d\n", current->data);
        current = current->next;
    }

    return 0;
}

```

Structure Assignment and Comparison

Structure Assignment

Structures can be assigned to each other if they are of the same type:

```

typedef struct {
    char name[50];
    int age;
    float height;
} Person;

int main() {
    Person person1 = {"John Doe", 25, 5.9};
    Person person2;

    // Structure assignment
    person2 = person1;

    printf("Person 1: %s, %d, %.1f\n", person1.name, person1.age, person1.height);
    printf("Person 2: %s, %d, %.1f\n", person2.name, person2.age, person2.height);

    return 0;
}

```

Structure Comparison

Structures cannot be directly compared using comparison operators. You need to compare members individually:

```
typedef struct {
    char name[50];
    int age;
    float height;
} Person;

int compare_persons(const Person *p1, const Person *p2) {
    // Compare names
    int name_cmp = strcmp(p1->name, p2->name);
    if (name_cmp != 0) return name_cmp;

    // Compare ages
    if (p1->age != p2->age) return p1->age - p2->age;

    // Compare heights
    if (p1->height < p2->height) return -1;
    if (p1->height > p2->height) return 1;

    return 0; // Structures are equal
}

int main() {
    Person person1 = {"John Doe", 25, 5.9};
    Person person2 = {"John Doe", 25, 5.9};

    if (compare_persons(&person1, &person2) == 0) {
        printf("Persons are equal\n");
    } else {
        printf("Persons are different\n");
    }

    return 0;
}
```

Passing Structures to Functions

Pass by Value

When passing structures by value, a copy of the entire structure is made:

```

typedef struct {
    char name[50];
    int age;
    float height;
} Person;

void print_person(Person p) {
    printf("Name: %s\n", p.name);
    printf("Age: %d\n", p.age);
    printf("Height: %.1f\n", p.height);
}

void modify_person(Person p) {
    // Modifications affect only the copy
    p.age = 30;
    strcpy(p.name, "Modified Name");
}

int main() {
    Person person = {"John Doe", 25, 5.9};

    printf("Before modification:\n");
    print_person(person);

    modify_person(person);

    printf("After modification:\n");
    print_person(person); // Original unchanged

    return 0;
}

```

Pass by Reference (Pointer)

Passing pointers to structures is more efficient and allows modification:

```

typedef struct {
    char name[50];
    int age;
    float height;
} Person;

void print_person(const Person *p) {
    printf("Name: %s\n", p->name);
    printf("Age: %d\n", p->age);
    printf("Height: %.1f\n", p->height);
}

```

```

}

void modify_person(Person *p) {
    // Modifications affect the original structure
    p->age = 30;
    strcpy(p->name, "Modified Name");
}

int main() {
    Person person = {"John Doe", 25, 5.9};

    printf("Before modification:\n");
    print_person(&person);

    modify_person(&person);

    printf("After modification:\n");
    print_person(&person); // Original changed

    return 0;
}

```

Structure Size and Memory Layout

Calculating Structure Size

```

#include <stdio.h>

typedef struct {
    char a;    // 1 byte
    int b;     // 4 bytes
    char c;    // 1 byte
} Example;

int main() {
    printf("Size of char: %zu bytes\n", sizeof(char));
    printf("Size of int: %zu bytes\n", sizeof(int));
    printf("Size of Example: %zu bytes\n", sizeof(Example));
    // Note: Size may be larger than sum due to padding

    return 0;
}

```

Memory Alignment

Compilers may add padding to ensure proper alignment:

```
#include <stdio.h>

// Without pragma pack
typedef struct {
    char a;        // 1 byte
    int b;         // 4 bytes (may have 3 bytes padding before)
    char c;        // 1 byte (may have 3 bytes padding after)
} Unpacked;

// With pragma pack (1) - no padding
#pragma pack(1)
typedef struct {
    char a;        // 1 byte
    int b;         // 4 bytes (no padding)
    char c;        // 1 byte (no padding)
} Packed;
#pragma pack()

int main() {
    printf("Size of Unpacked: %zu bytes\n", sizeof(Unpacked));
    printf("Size of Packed: %zu bytes\n", sizeof(Packed));

    return 0;
}
```

Practical Examples

Student Management System

```
#include <stdio.h>
#include <stdlib.h>
#include <string.h>

typedef struct {
    int id;
    char name[50];
    float grades[5];
    float average;
} Student;

void calculate_average(Student *student) {
```

```

    float sum = 0;
    for (int i = 0; i < 5; i++) {
        sum += student->grades[i];
    }
    student->average = sum / 5;
}

void print_student(const Student *student) {
    printf("ID: %d\n", student->id);
    printf("Name: %s\n", student->name);
    printf("Grades: ");
    for (int i = 0; i < 5; i++) {
        printf("%.1f ", student->grades[i]);
    }
    printf("\nAverage: %.2f\n", student->average);
    printf("-----\n");
}

int main() {
    Student students[3] = {
        {1, "Alice", {85.5, 90.0, 78.5, 92.0, 88.5}, 0},
        {2, "Bob", {76.0, 82.5, 79.0, 85.5, 81.0}, 0},
        {3, "Charlie", {95.0, 92.5, 98.0, 90.5, 94.0}, 0}
    };

    // Calculate averages
    for (int i = 0; i < 3; i++) {
        calculate_average(&students[i]);
    }

    // Print student information
    for (int i = 0; i < 3; i++) {
        print_student(&students[i]);
    }

    // Find student with highest average
    int top_student = 0;
    for (int i = 1; i < 3; i++) {
        if (students[i].average > students[top_student].average) {
            top_student = i;
        }
    }

    printf("Top student:\n");
    print_student(&students[top_student]);

    return 0;
}

```

```
}
```

Point Structure with Operations

```
#include <stdio.h>
#include <math.h>

typedef struct {
    double x;
    double y;
} Point;

// Function to create a point
Point create_point(double x, double y) {
    Point p = {x, y};
    return p;
}

// Function to calculate distance between two points
double distance(const Point *p1, const Point *p2) {
    double dx = p1->x - p2->x;
    double dy = p1->y - p2->y;
    return sqrt(dx * dx + dy * dy);
}

// Function to move a point
void move_point(Point *p, double dx, double dy) {
    p->x += dx;
    p->y += dy;
}

// Function to print a point
void print_point(const Point *p) {
    printf("(%.2f, %.2f)", p->x, p->y);
}

int main() {
    Point p1 = create_point(0, 0);
    Point p2 = create_point(3, 4);

    printf("Point 1: ");
    print_point(&p1);
    printf("\n");

    printf("Point 2: ");
```



```

    print_point(&p2);
    printf("\n");

    printf("Distance: %.2f\n", distance(&p1, &p2));

    move_point(&p1, 1, 1);
    printf("After moving Point 1: ");
    print_point(&p1);
    printf("\n");

    printf("New distance: %.2f\n", distance(&p1, &p2));

    return 0;
}

```

Summary

Structures are fundamental to C programming, providing a way to group related data and create complex data types. Key points to remember:

1. **Declaration and Initialization:** Structures can be declared and initialized in various ways
2. **Member Access:** Use dot operator for structure variables, arrow operator for pointers
3. **Arrays of Structures:** Enable management of multiple related data items
4. **Nested Structures:** Allow creation of complex hierarchical data representations
5. **Function Parameters:** Pass by value creates copies, pass by pointer is more efficient
6. **Memory Layout:** Structures may have padding for alignment
7. **Self-Referential Structures:** Enable linked data structures like lists and trees

Understanding structures is essential for implementing complex data structures, managing related data efficiently, and creating well-organized C programs.

Unions and Enumerations

Introduction

Unions and enumerations are two important features in C that complement structures by providing additional ways to define and work with data types. Unions allow different data types to share the same memory location, making them useful for memory optimization and type-generic programming. Enumerations provide a way to define named integer constants, improving code readability and maintainability.

Unions

Union Declaration and Basics

A union is a special data type that allows storing different data types in the same memory location. Unlike structures where each member has its own memory, all members of a union share the same memory space.

```
#include <stdio.h>

union Data {
    int i;
    float f;
    char str[20];
};

int main() {
    union Data data;

    printf("Memory size occupied by data: %zu\n", sizeof(data));

    data.i = 10;
    printf("data.i: %d\n", data.i);

    data.f = 220.5;
    printf("data.f: %f\n", data.f);

    // Note: Previous value of data.i is lost
    printf("data.i after setting data.f: %d\n", data.i);
}
```

```

    strcpy(data.str, "C Programming");
    printf("data.str: %s\n", data.str);

    return 0;
}

```

Union Memory Layout

All members of a union occupy the same memory location, so the size of a union is the size of its largest member:

```

#include <stdio.h>

union Example {
    char c;          // 1 byte
    int i;           // 4 bytes (typically)
    double d;        // 8 bytes (typically)
    char str[20];    // 20 bytes
};

int main() {
    union Example ex;

    printf("Size of char: %zu bytes\n", sizeof(char));
    printf("Size of int: %zu bytes\n", sizeof(int));
    printf("Size of double: %zu bytes\n", sizeof(double));
    printf("Size of char[20]: %zu bytes\n", sizeof(char[20]));
    printf("Size of union Example: %zu bytes\n", sizeof(ex));

    // All members share the same memory address
    printf("Address of ex.c: %p\n", (void*)&ex.c);
    printf("Address of ex.i: %p\n", (void*)&ex.i);
    printf("Address of ex.d: %p\n", (void*)&ex.d);
    printf("Address of ex.str: %p\n", (void*)&ex.str);

    return 0;
}

```

Practical Union Usage

1. Memory Optimization

Unions are useful when you need to store different types of data but only one at a time:

```

#include <stdio.h>
#include <string.h>

typedef enum {
    TYPE_INT,
    TYPE_FLOAT,
    TYPE_STRING
} DataType;

typedef struct {
    DataType type;
    union {
        int i;
        float f;
        char str[50];
    } data;
} Variant;

void print_variant(const Variant *v) {
    switch (v->type) {
        case TYPE_INT:
            printf("Integer: %d\n", v->data.i);
            break;
        case TYPE_FLOAT:
            printf("Float: %.2f\n", v->data.f);
            break;
        case TYPE_STRING:
            printf("String: %s\n", v->data.str);
            break;
    }
}

int main() {
    Variant var1 = {TYPE_INT, .data.i = 42};
    Variant var2 = {TYPE_FLOAT, .data.f = 3.14f};
    Variant var3 = {TYPE_STRING, .data = {.str = "Hello, World!"}};

    print_variant(&var1);
    print_variant(&var2);
    print_variant(&var3);

    return 0;
}

```

2. Hardware Register Representation

Unions are commonly used in embedded programming to represent hardware registers:

```
#include <stdio.h>
#include <stdint.h>

// 32-bit register with bit fields
typedef union {
    uint32_t value;
    struct {
        uint32_t enable : 1;
        uint32_t mode : 2;
        uint32_t reserved : 5;
        uint32_t frequency : 8;
        uint32_t address : 16;
    } bits;
} Register;

int main() {
    Register reg;

    // Set register value
    reg.value = 0x12345678;
    printf("Register value: 0x%08X\n", reg.value);
    printf("Enable: %d\n", reg.bits.enable);
    printf("Mode: %d\n", reg.bits.mode);
    printf("Frequency: %d\n", reg.bits.frequency);
    printf("Address: 0x%04X\n", reg.bits.address);

    // Modify individual bits
    reg.bits.enable = 1;
    reg.bits.mode = 2;
    printf("Modified register: 0x%08X\n", reg.value);

    return 0;
}
```

Anonymous Unions (C11)

C11 introduced anonymous unions, which allow direct access to union members without specifying the union name:

```
#include <stdio.h>

typedef struct {
```

```

    int type;
    union {
        int integer_value;
        float float_value;
        char string_value[20];
    }; // Anonymous union
} DataRecord;

int main() {
    DataRecord record;

    record.type = 1;
    record.integer_value = 42; // Direct access to union member
    printf("Integer value: %d\n", record.integer_value);

    record.type = 2;
    record.float_value = 3.14f; // Direct access to union member
    printf("Float value: %.2f\n", record.float_value);

    return 0;
}

```

Enumerations

Basic Enumeration Declaration

Enumerations provide a way to define named integer constants, making code more readable and maintainable:

```

#include <stdio.h>

// Basic enumeration
enum Color {
    RED,
    GREEN,
    BLUE
};

// Enumeration with typedef
typedef enum {
    MONDAY,
    TUESDAY,
    WEDNESDAY,
    THURSDAY,
    FRIDAY,
}

```

```

        SATURDAY,
        SUNDAY
    } Day;

int main() {
    enum Color favorite_color = BLUE;
    Day today = WEDNESDAY;

    printf("Favorite color: %d\n", favorite_color);
    printf("Today is day: %d\n", today);

    // Using in switch statements
    switch (today) {
        case MONDAY:
            printf("It's Monday\n");
            break;
        case WEDNESDAY:
            printf("It's Wednesday\n");
            break;
        case FRIDAY:
            printf("It's Friday\n");
            break;
        default:
            printf("It's another day\n");
    }

    return 0;
}

```

Explicit Enumeration Values

Enumeration constants can be assigned specific values:

```

#include <stdio.h>

typedef enum {
    HTTP_OK = 200,
    HTTP_BAD_REQUEST = 400,
    HTTP_UNAUTHORIZED = 401,
    HTTP_FORBIDDEN = 403,
    HTTP_NOT_FOUND = 404,
    HTTP_INTERNAL_SERVER_ERROR = 500
} HttpStatusCode;

typedef enum {
    FLAG_NONE = 0,

```

```

    FLAG_READ = 1,
    FLAG_WRITE = 2,
    FLAG_EXECUTE = 4,
    FLAG_DELETE = 8
} FilePermission;

int main() {
    HttpStatusCode status = HTTP_NOT_FOUND;
    FilePermission permissions = FLAG_READ | FLAG_WRITE;

    printf("HTTP Status: %d\n", status);
    printf("Permissions: %d\n", permissions);

    // Check permissions using bitwise operations
    if (permissions & FLAG_READ) {
        printf("Read permission granted\n");
    }

    if (permissions & FLAG_EXECUTE) {
        printf("Execute permission granted\n");
    } else {
        printf("Execute permission denied\n");
    }

    return 0;
}

```

Enumeration with Negative Values

Enumerations can include negative values:

```

#include <stdio.h>

typedef enum {
    TEMPERATURE_FREEZING = 0,
    TEMPERATURE_COLD = 10,
    TEMPERATURE_WARM = 20,
    TEMPERATURE_HOT = 30,
    TEMPERATURE_BOILING = 100
} Temperature;

typedef enum {
    DIRECTION_NORTH = 0,
    DIRECTION_EAST = 90,
    DIRECTION_SOUTH = 180,
    DIRECTION_WEST = 270
}

```



```

} Direction;

// Enumeration with negative values
typedef enum {
    STATUS_ERROR = -1,
    STATUS_UNKNOWN = 0,
    STATUS_SUCCESS = 1
} Status;

int main() {
    Temperature temp = TEMPERATURE_HOT;
    Direction dir = DIRECTION_EAST;
    Status result = STATUS_SUCCESS;

    printf("Temperature: %d°C\n", temp);
    printf("Direction: %d°\n", dir);
    printf("Status: %d\n", result);

    return 0;
}

```

Enumeration Size and Type Specification (C23)

C23 introduces more control over enumeration types:

```

#include <stdio.h>

// Specify underlying type (C23 feature)
typedef enum : unsigned char {
    SMALL_RED,
    SMALL_GREEN,
    SMALL_BLUE
} SmallColor;

typedef enum : long long {
    LARGE_VALUE1 = 1000000000000LL,
    LARGE_VALUE2 = 2000000000000LL
} LargeEnum;

int main() {
    SmallColor color = SMALL_GREEN;
    LargeEnum large = LARGE_VALUE2;

    printf("Size of SmallColor: %zu bytes\n", sizeof(color));
    printf("Size of LargeEnum: %zu bytes\n", sizeof(large));
    printf("SmallColor value: %d\n", color);
}

```

```

    printf("LargeEnum value: %lld\n", large);

    return 0;
}

```

Combining Unions and Enumerations

Tagged Union Pattern

A common pattern combines unions with enumerations to create type-safe variant types:

```

#include <stdio.h>
#include <string.h>

typedef enum {
    VALUE_INTEGER,
    VALUE_FLOAT,
    VALUE_STRING,
    VALUE_BOOLEAN
} ValueType;

typedef struct {
    ValueType type;
    union {
        int integer_value;
        float float_value;
        char string_value[100];
        int boolean_value; // 0 = false, 1 = true
    } data;
} Value;

// Constructor functions
Value create_integer(int value) {
    Value v = {VALUE_INTEGER, .data.integer_value = value};
    return v;
}

Value create_float(float value) {
    Value v = {VALUE_FLOAT, .data.float_value = value};
    return v;
}

Value create_string(const char *value) {
    Value v = {VALUE_STRING};
    strncpy(v.data.string_value, value, sizeof(v.data.string_value) - 1);
}

```

```

        v.data.string_value[sizeof(v.data.string_value) - 1] = '\0';
        return v;
    }

    Value create_boolean(int value) {
        Value v = {VALUE_BOOLEAN, .data.boolean_value = value ? 1 : 0};
        return v;
    }

    // Print function
    void print_value(const Value *v) {
        switch (v->type) {
            case VALUE_INTEGER:
                printf("Integer: %d\n", v->data.integer_value);
                break;
            case VALUE_FLOAT:
                printf("Float: %.2f\n", v->data.float_value);
                break;
            case VALUE_STRING:
                printf("String: %s\n", v->data.string_value);
                break;
            case VALUE_BOOLEAN:
                printf("Boolean: %s\n", v->data.boolean_value ? "true" : "false");
                break;
        }
    }

    int main() {
        Value values[] = {
            create_integer(42),
            create_float(3.14f),
            create_string("Hello, World!"),
            create_boolean(1)
        };

        int count = sizeof(values) / sizeof(values[0]);

        for (int i = 0; i < count; i++) {
            print_value(&values[i]);
        }

        return 0;
    }

```

Configuration Structure with Unions

Using unions and enumerations for flexible configuration:

```
#include <stdio.h>
#include <string.h>

typedef enum {
    CONFIG_TYPE_INT,
    CONFIG_TYPE_FLOAT,
    CONFIG_TYPE_STRING,
    CONFIG_TYPE_BOOL
} ConfigType;

typedef struct {
    char key[50];
    ConfigType type;
    union {
        int int_value;
        float float_value;
        char string_value[100];
        int bool_value;
    } value;
} ConfigItem;

typedef struct {
    ConfigItem *items;
    int count;
    int capacity;
} Config;

Config* create_config() {
    Config *config = malloc(sizeof(Config));
    if (config != NULL) {
        config->items = malloc(10 * sizeof(ConfigItem));
        if (config->items != NULL) {
            config->count = 0;
            config->capacity = 10;
        } else {
            free(config);
            config = NULL;
        }
    }
    return config;
}

void add_int_config(Config *config, const char *key, int value) {
```

```

    if (config->count >= config->capacity) {
        // Resize array
        ConfigItem *new_items = realloc(config->items,
                                         (config->capacity + 10) * sizeof(ConfigItem));

        if (new_items != NULL) {
            config->items = new_items;
            config->capacity += 10;
        } else {
            return; // Allocation failed
        }
    }

    ConfigItem *item = &config->items[config->count];
    strncpy(item->key, key, sizeof(item->key) - 1);
    item->key[sizeof(item->key) - 1] = '\0';
    item->type = CONFIG_TYPE_INT;
    item->value.int_value = value;

    config->count++;
}

void print_config(const Config *config) {
    for (int i = 0; i < config->count; i++) {
        const ConfigItem *item = &config->items[i];
        printf("Key: %s, ", item->key);

        switch (item->type) {
            case CONFIG_TYPE_INT:
                printf("Value: %d (int)\n", item->value.int_value);
                break;
            case CONFIG_TYPE_FLOAT:
                printf("Value: %.2f (float)\n", item->value.float_value);
                break;
            case CONFIG_TYPE_STRING:
                printf("Value: %s (string)\n", item->value.string_value);
                break;
            case CONFIG_TYPE_BOOL:
                printf("Value: %s (bool)\n", item->value.bool_value ? "true" : "false");
                break;
        }
    }
}

int main() {
    Config *config = create_config();
    if (config == NULL) {

```

```

        printf("Failed to create config\n");
        return 1;
    }

    add_int_config(config, "window_width", 800);
    add_int_config(config, "window_height", 600);

    print_config(config);

    // Clean up
    free(config->items);
    free(config);

    return 0;
}

```

Practical Examples

Network Packet Header

Using unions and enumerations for network protocol implementation:

```

#include <stdio.h>
#include <stdint.h>

typedef enum {
    PACKET_TYPE_DATA = 0x01,
    PACKET_TYPE_ACK = 0x02,
    PACKET_TYPE_NACK = 0x03,
    PACKET_TYPE_CONTROL = 0x04
} PacketType;

typedef enum {
    CONTROL_START = 0x01,
    CONTROL_STOP = 0x02,
    CONTROL_RESET = 0x03
} ControlType;

typedef struct {
    uint8_t version : 4;
    uint8_t type : 4;
    uint16_t length;
    uint32_t sequence;
    union {
        struct {

```

```

        uint8_t data[1024];
    } data_packet;

    struct {
        uint32_t ack_sequence;
    } ack_packet;

    struct {
        uint8_t control_type;
        uint8_t reserved[3];
    } control_packet;
} payload;
} NetworkPacket;

void print_packet_info(const NetworkPacket *packet) {
    printf("Packet Version: %d\n", packet->version);
    printf("Packet Type: 0x%02X\n", packet->type);
    printf("Packet Length: %d\n", packet->length);
    printf("Sequence Number: %u\n", packet->sequence);

    switch (packet->type) {
        case PACKET_TYPE_DATA:
            printf("Data Packet\n");
            break;
        case PACKET_TYPE_ACK:
            printf("ACK for sequence: %u\n", packet->payload.ack_packet.ack_sequence);
            break;
        case PACKET_TYPE_CONTROL:
            printf("Control Type: 0x%02X\n", packet->payload.control_packet.control_type);
            break;
        default:
            printf("Unknown Packet Type\n");
    }
}

int main() {
    NetworkPacket packet = {
        .version = 1,
        .type = PACKET_TYPE_CONTROL,
        .length = sizeof(packet.payload.control_packet),
        .sequence = 1000,
        .payload.control_packet = {
            .control_type = CONTROL_RESET,
            .reserved = {0}
        }
    };
};

```

```

    print_packet_info(&packet);

    return 0;
}

```

Mathematical Expression Evaluator

Using unions and enumerations for a simple expression evaluator:

```

#include <stdio.h>
#include <stdlib.h>
#include <string.h>

typedef enum {
    TOKEN_NUMBER,
    TOKEN_OPERATOR,
    TOKEN_VARIABLE
} TokenType;

typedef enum {
    OP_ADD = '+',
    OP_SUBTRACT = '-',
    OP_MULTIPLY = '*',
    OP_DIVIDE = '/',
    OP_ASSIGN = '='
} Operator;

typedef struct {
    TokenType type;
    union {
        double number_value;
        Operator operator_value;
        char variable_name[20];
    } data;
} Token;

typedef struct {
    Token *tokens;
    int count;
    int capacity;
} Expression;

Expression* create_expression() {
    Expression *expr = malloc(sizeof(Expression));
    if (expr != NULL) {
        expr->tokens = malloc(10 * sizeof(Token));
    }
}

```



```

        if (expr->tokens != NULL) {
            expr->count = 0;
            expr->capacity = 10;
        } else {
            free(expr);
            expr = NULL;
        }
    }
    return expr;
}

void add_number_token(Expression *expr, double value) {
    if (expr->count >= expr->capacity) {
        Token *new_tokens = realloc(expr->tokens,
                                     (expr->capacity + 10) * sizeof(Token));

        if (new_tokens != NULL) {
            expr->tokens = new_tokens;
            expr->capacity += 10;
        } else {
            return;
        }
    }

    Token *token = &expr->tokens[expr->count];
    token->type = TOKEN_NUMBER;
    token->data.number_value = value;
    expr->count++;
}

void add_operator_token(Expression *expr, Operator op) {
    if (expr->count >= expr->capacity) {
        Token *new_tokens = realloc(expr->tokens,
                                     (expr->capacity + 10) * sizeof(Token));

        if (new_tokens != NULL) {
            expr->tokens = new_tokens;
            expr->capacity += 10;
        } else {
            return;
        }
    }

    Token *token = &expr->tokens[expr->count];
    token->type = TOKEN_OPERATOR;
    token->data.operator_value = op;
    expr->count++;
}

```

```

void print_expression(const Expression *expr) {
    for (int i = 0; i < expr->count; i++) {
        const Token *token = &expr->tokens[i];

        switch (token->type) {
            case TOKEN_NUMBER:
                printf("%.2f ", token->data.number_value);
                break;
            case TOKEN_OPERATOR:
                printf("%c ", token->data.operator_value);
                break;
            case TOKEN_VARIABLE:
                printf("%s ", token->data.variable_name);
                break;
        }
    }
    printf("\n");
}

int main() {
    Expression *expr = create_expression();
    if (expr == NULL) {
        printf("Failed to create expression\n");
        return 1;
    }

    // Create expression: 3.14 + 2.5 * x
    add_number_token(expr, 3.14);
    add_operator_token(expr, OP_ADD);
    add_number_token(expr, 2.5);
    add_operator_token(expr, OP_MULTIPLY);
    // For simplicity, we'll represent variable as a number
    add_number_token(expr, 0); // Placeholder for variable x

    printf("Expression: ");
    print_expression(expr);

    // Clean up
    free(expr->tokens);
    free(expr);

    return 0;
}

```

Summary

Unions and enumerations are powerful features that extend the capabilities of C programming:

1. **Unions:** Allow different data types to share the same memory location, useful for memory optimization and type-generic programming
2. **Memory Sharing:** All union members occupy the same memory space, with size equal to the largest member
3. **Practical Applications:** Hardware register representation, variant types, and memory-efficient data storage
4. **Anonymous Unions:** C11 feature enabling direct access to union members (C11)
5. **Enumerations:** Provide named integer constants for improved code readability and maintainability
6. **Explicit Values:** Enumerations can have specific values, including negative numbers
7. **Type Safety:** Combining unions with enumerations creates type-safe variant types
8. **Advanced Features:** C23 introduces more control over enumeration types

Understanding unions and enumerations enables programmers to write more efficient, readable, and maintainable C code, particularly in systems programming, embedded development, and protocol implementation.

Advanced Structure Concepts

Introduction

Advanced structure concepts in C extend beyond basic structure usage to include sophisticated features like bit fields, structure packing, flexible array members, and self-referential structures. These concepts enable programmers to optimize memory usage, implement complex data structures, and create efficient representations of real-world entities. Understanding these advanced features is essential for systems programming, embedded development, and performance-critical applications.

Bit Fields

Basic Bit Field Declaration

Bit fields allow you to specify the exact number of bits used for structure members, enabling efficient memory usage for small values:

```
#include <stdio.h>

struct FilePermissions {
    unsigned int read : 1;    // 1 bit
    unsigned int write : 1;   // 1 bit
    unsigned int execute : 1; // 1 bit
    unsigned int reserved : 5; // 5 bits
};

int main() {
    struct FilePermissions perm;

    // Set permissions
    perm.read = 1;
    perm.write = 1;
    perm.execute = 0;
    perm.reserved = 0;

    printf("Size of struct: %zu bytes\n", sizeof(perm));
    printf("Read: %d\n", perm.read);
    printf("Write: %d\n", perm.write);
    printf("Execute: %d\n", perm.execute);
}
```

```

    printf("Reserved: %d\n", perm.reserved);

    return 0;
}

```

Practical Bit Field Usage

1. Hardware Register Representation

Bit fields are commonly used to represent hardware registers with specific bit layouts:

```

#include <stdio.h>
#include <stdint.h>

struct GPIO_Register {
    uint32_t pin0 : 2; // 2 bits for pin 0 configuration
    uint32_t pin1 : 2; // 2 bits for pin 1 configuration
    uint32_t pin2 : 2; // 2 bits for pin 2 configuration
    uint32_t pin3 : 2; // 2 bits for pin 3 configuration
    uint32_t mode : 4; // 4 bits for mode selection
    uint32_t reserved : 20; // Remaining bits reserved
};

int main() {
    struct GPIO_Register gpio = {0};

    // Configure pins
    gpio.pin0 = 1; // Output mode
    gpio.pin1 = 0; // Input mode
    gpio.pin2 = 1; // Output mode
    gpio.pin3 = 2; // Alternate function
    gpio.mode = 5; // Specific mode

    printf("GPIO Register Value: 0x%08X\n", *(uint32_t*)&gpio);
    printf("Pin 0 Configuration: %d\n", gpio.pin0);
    printf("Pin 1 Configuration: %d\n", gpio.pin1);
    printf("Mode: %d\n", gpio.mode);

    return 0;
}

```

2. Network Protocol Headers

Bit fields are useful for implementing network protocol headers with specific bit layouts:

```

#include <stdio.h>
#include <stdint.h>

struct IP_Header {
    uint32_t version : 4;          // 4 bits
    uint32_t ihl : 4;              // 4 bits
    uint32_t tos : 8;              // 8 bits
    uint32_t total_length : 16;    // 16 bits
    uint32_t identification : 16; // 16 bits
    uint32_t flags : 3;            // 3 bits
    uint32_t fragment_offset : 13; // 13 bits
    uint32_t ttl : 8;              // 8 bits
    uint32_t protocol : 8;         // 8 bits
    uint32_t checksum : 16;        // 16 bits
    uint32_t source_ip : 32;       // 32 bits
    uint32_t dest_ip : 32;         // 32 bits
};

int main() {
    struct IP_Header ip_header = {0};

    // Set header values
    ip_header.version = 4;          // IPv4
    ip_header.ihl = 5;              // 5 * 4 = 20 bytes
    ip_header.tos = 0;
    ip_header.total_length = 100;
    ip_header.identification = 12345;
    ip_header.flags = 2;            // Don't fragment
    ip_header.fragment_offset = 0;
    ip_header.ttl = 64;
    ip_header.protocol = 6;         // TCP
    ip_header.source_ip = 0xC0A80101; // 192.168.1.1
    ip_header.dest_ip = 0xC0A80102;  // 192.168.1.2

    printf("IP Header Size: %zu bytes\n", sizeof(ip_header));
    printf("Version: %d\n", ip_header.version);
    printf("Header Length: %d words\n", ip_header.ihl);
    printf("Total Length: %d bytes\n", ip_header.total_length);
    printf("Protocol: %d\n", ip_header.protocol);

    return 0;
}

```

Bit Field Limitations and Considerations

1. Portability Issues

Bit field layout is implementation-defined and may vary between compilers:

```
#include <stdio.h>

struct BitFieldExample {
    unsigned int a : 3;
    unsigned int b : 4;
    unsigned int c : 5;
};

int main() {
    struct BitFieldExample bf = {1, 2, 3};

    printf("Size: %zu bytes\n", sizeof(bf));
    printf("Values: a=%d, b=%d, c=%d\n", bf.a, bf.b, bf.c);

    // Bit layout is implementation-defined
    unsigned int *ptr = (unsigned int*)&bf;
    printf("Raw value: 0x%08X\n", *ptr);

    return 0;
}
```

2. Alignment and Padding

Bit fields may introduce padding for alignment:

```
#include <stdio.h>

struct PackedBits {
    unsigned int a : 1;
    unsigned int b : 1;
    unsigned int c : 1;
};

struct AlignedStruct {
    char x;
    struct PackedBits bits;
    char y;
};

int main() {
```

```

printf("Size of PackedBits: %zu bytes\n", sizeof(struct PackedBits));
printf("Size of AlignedStruct: %zu bytes\n", sizeof(struct AlignedStruct));

return 0;
}

```

Structure Packing and Alignment

Default Structure Alignment

Compilers add padding to structures to ensure proper alignment:

```

#include <stdio.h>

struct Unpacked {
    char a;        // 1 byte
    int b;         // 4 bytes (may have 3 bytes padding before)
    char c;        // 1 byte (may have 3 bytes padding after)
};

#pragma pack(1)
struct Packed {
    char a;        // 1 byte
    int b;         // 4 bytes (no padding)
    char c;        // 1 byte (no padding)
};

#pragma pack()

int main() {
    printf("Size of Unpacked: %zu bytes\n", sizeof(struct Unpacked));
    printf("Size of Packed: %zu bytes\n", sizeof(struct Packed));

    // Show member offsets
    struct Unpacked u = {0};
    printf("Unpacked offsets - a: %td, b: %td, c: %td\n",
        (char*)&u.a - (char*)&u,
        (char*)&u.b - (char*)&u,
        (char*)&u.c - (char*)&u);

    struct Packed p = {0};
    printf("Packed offsets - a: %td, b: %td, c: %td\n",
        (char*)&p.a - (char*)&p,
        (char*)&p.b - (char*)&p,
        (char*)&p.c - (char*)&p);
}

```



```
    return 0;
}
```

Controlling Alignment with Pragma Pack

The `#pragma pack` directive controls structure packing:

```
#include <stdio.h>

// Default alignment
struct Default {
    char a;
    int b;
    char c;
};

// 1-byte alignment (no padding)
#pragma pack(1)
struct Packed1 {
    char a;
    int b;
    char c;
};
#pragma pack()

// 2-byte alignment
#pragma pack(2)
struct Packed2 {
    char a;
    int b;
    char c;
};
#pragma pack()

// 4-byte alignment
#pragma pack(4)
struct Packed4 {
    char a;
    int b;
    char c;
};
#pragma pack()

int main() {
    printf("Default alignment: %zu bytes\n", sizeof(struct Default));
    printf("1-byte packing: %zu bytes\n", sizeof(struct Packed1));
}
```

```

    printf("2-byte packing: %zu bytes\n", sizeof(struct Packed2));
    printf("4-byte packing: %zu bytes\n", sizeof(struct Packed4));

    return 0;
}

```

Using `_Alignas` (C11)

C11 introduces `_Alignas` for explicit alignment control:

```

#include <stdio.h>
#include <stdalign.h>

struct AlignedStruct {
    _Alignas(16) char data[16]; // Align to 16-byte boundary
    int value;
};

int main() {
    struct AlignedStruct s = {0};

    printf("Size: %zu bytes\n", sizeof(s));
    printf("Address of data: %p\n", (void*)s.data);
    printf("Alignment of data: %zu\n", (size_t)s.data % 16);

    return 0;
}

```

Flexible Array Members (C99)

Basic Flexible Array Member

Flexible array members allow structures to have variable-sized arrays at the end:

```

#include <stdio.h>
#include <stdlib.h>
#include <string.h>

struct DynamicArray {
    size_t count;
    int data[]; // Flexible array member (must be last)
};

struct DynamicArray* create_array(size_t count) {

```

```

// Allocate memory for structure + array
struct DynamicArray *arr = malloc(sizeof(struct DynamicArray) +
                                count * sizeof(int));

if (arr != NULL) {
    arr->count = count;
    // arr->data can now be used as array of 'count' integers
}
return arr;
}

int main() {
    struct DynamicArray *arr = create_array(5);
    if (arr != NULL) {
        // Initialize array elements
        for (size_t i = 0; i < arr->count; i++) {
            arr->data[i] = (int)(i * i);
        }

        // Print array elements
        printf("Array elements: ");
        for (size_t i = 0; i < arr->count; i++) {
            printf("%d ", arr->data[i]);
        }
        printf("\n");

        printf("Size of structure: %zu bytes\n", sizeof(struct DynamicArray));
        printf("Total allocated size: %zu bytes\n",
               sizeof(struct DynamicArray) + arr->count * sizeof(int));

        // Free memory
        free(arr);
    }

    return 0;
}

```

Flexible Array in String Structure

Using flexible array members for dynamic strings:

```

#include <stdio.h>
#include <stdlib.h>
#include <string.h>

struct DynamicString {
    size_t length;

```

```

    char data[]; // Flexible array member
};

struct DynamicString* create_string(const char *source) {
    size_t len = strlen(source);
    // Allocate memory for structure + string + null terminator
    struct DynamicString *str = malloc(sizeof(struct DynamicString) +
                                       len + 1);

    if (str != NULL) {
        str->length = len;
        strcpy(str->data, source);
    }
    return str;
}

struct DynamicString* append_strings(const struct DynamicString *s1,
                                    const char *s2) {
    size_t len2 = strlen(s2);
    size_t new_len = s1->length + len2;

    struct DynamicString *result = malloc(sizeof(struct DynamicString) +
                                           new_len + 1);

    if (result != NULL) {
        result->length = new_len;
        strcpy(result->data, s1->data);
        strcat(result->data, s2);
    }
    return result;
}

int main() {
    struct DynamicString *str1 = create_string("Hello, ");
    if (str1 != NULL) {
        printf("String 1: %s (length: %zu)\n", str1->data, str1->length);

        struct DynamicString *str2 = append_strings(str1, "World!");
        if (str2 != NULL) {
            printf("String 2: %s (length: %zu)\n", str2->data, str2->length);
            free(str2);
        }

        free(str1);
    }

    return 0;
}

```

Self-Referential Structures

Linked List Implementation

Self-referential structures are fundamental for implementing linked data structures:

```
#include <stdio.h>
#include <stdlib.h>

struct Node {
    int data;
    struct Node *next;
};

// Function to create a new node
struct Node* create_node(int data) {
    struct Node *node = malloc(sizeof(struct Node));
    if (node != NULL) {
        node->data = data;
        node->next = NULL;
    }
    return node;
}

// Function to insert at beginning
struct Node* insert_at_beginning(struct Node *head, int data) {
    struct Node *new_node = create_node(data);
    if (new_node != NULL) {
        new_node->next = head;
        head = new_node;
    }
    return head;
}

// Function to print the list
void print_list(struct Node *head) {
    struct Node *current = head;
    while (current != NULL) {
        printf("%d -> ", current->data);
        current = current->next;
    }
    printf("NULL\n");
}

// Function to free the list
void free_list(struct Node *head) {
    struct Node *current = head;
```

```

        while (current != NULL) {
            struct Node *next = current->next;
            free(current);
            current = next;
        }
    }

    int main() {
        struct Node *head = NULL;

        // Build the list: 3 -> 2 -> 1 -> NULL
        head = insert_at_beginning(head, 1);
        head = insert_at_beginning(head, 2);
        head = insert_at_beginning(head, 3);

        printf("Linked List: ");
        print_list(head);

        // Free the list
        free_list(head);

        return 0;
    }

```

Binary Tree Implementation

Self-referential structures for binary trees:

```

#include <stdio.h>
#include <stdlib.h>

struct TreeNode {
    int data;
    struct TreeNode *left;
    struct TreeNode *right;
};

// Function to create a new node
struct TreeNode* create_node(int data) {
    struct TreeNode *node = malloc(sizeof(struct TreeNode));
    if (node != NULL) {
        node->data = data;
        node->left = NULL;
        node->right = NULL;
    }
    return node;
}

```

```

}

// Function to insert a node in BST
struct TreeNode* insert(struct TreeNode *root, int data) {
    if (root == NULL) {
        return create_node(data);
    }

    if (data < root->data) {
        root->left = insert(root->left, data);
    } else if (data > root->data) {
        root->right = insert(root->right, data);
    }

    return root;
}

// In-order traversal
void inorder_traversal(struct TreeNode *root) {
    if (root != NULL) {
        inorder_traversal(root->left);
        printf("%d ", root->data);
        inorder_traversal(root->right);
    }
}

// Free the tree
void free_tree(struct TreeNode *root) {
    if (root != NULL) {
        free_tree(root->left);
        free_tree(root->right);
        free(root);
    }
}

int main() {
    struct TreeNode *root = NULL;

    // Insert nodes
    root = insert(root, 50);
    insert(root, 30);
    insert(root, 70);
    insert(root, 20);
    insert(root, 40);
    insert(root, 60);
    insert(root, 80);

```

```

    printf("In-order traversal: ");
    inorder_traversal(root);
    printf("\n");

    // Free the tree
    free_tree(root);

    return 0;
}

```

Advanced Structure Techniques

Structure with Function Pointers (C Object-Oriented Approach)

Using function pointers to implement object-oriented concepts in C:

```

#include <stdio.h>
#include <stdlib.h>
#include <string.h>

// Forward declaration
struct Shape;

// Function pointer types
typedef double (*area_func)(const struct Shape *self);
typedef void (*print_func)(const struct Shape *self);

// Base structure
struct Shape {
    area_func area;
    print_func print;
    char name[20];
};

// Rectangle structure
struct Rectangle {
    struct Shape base;
    double width;
    double height;
};

// Circle structure
struct Circle {
    struct Shape base;
    double radius;
};

```



```

};

// Area calculation functions
double rectangle_area(const struct Shape *self) {
    const struct Rectangle *rect = (const struct Rectangle *)self;
    return rect->width * rect->height;
}

double circle_area(const struct Shape *self) {
    const struct Circle *circle = (const struct Circle *)self;
    return 3.14159 * circle->radius * circle->radius;
}

// Print functions
void rectangle_print(const struct Shape *self) {
    const struct Rectangle *rect = (const struct Rectangle *)self;
    printf("Rectangle: %s, Width: %.2f, Height: %.2f\n",
        self->name, rect->width, rect->height);
}

void circle_print(const struct Shape *self) {
    const struct Circle *circle = (const struct Circle *)self;
    printf("Circle: %s, Radius: %.2f\n",
        self->name, circle->radius);
}

// Constructor functions
struct Rectangle* create_rectangle(const char *name, double width, double height) {
    struct Rectangle *rect = malloc(sizeof(struct Rectangle));
    if (rect != NULL) {
        strncpy(rect->base.name, name, sizeof(rect->base.name) - 1);
        rect->base.name[sizeof(rect->base.name) - 1] = '\0';
        rect->base.area = rectangle_area;
        rect->base.print = rectangle_print;
        rect->width = width;
        rect->height = height;
    }
    return rect;
}

struct Circle* create_circle(const char *name, double radius) {
    struct Circle *circle = malloc(sizeof(struct Circle));
    if (circle != NULL) {
        strncpy(circle->base.name, name, sizeof(circle->base.name) - 1);
        circle->base.name[sizeof(circle->base.name) - 1] = '\0';
        circle->base.area = circle_area;
        circle->base.print = circle_print;
    }
}

```

```

        circle->radius = radius;
    }
    return circle;
}

int main() {
    struct Rectangle *rect = create_rectangle("MyRectangle", 5.0, 3.0);
    struct Circle *circle = create_circle("MyCircle", 4.0);

    if (rect != NULL && circle != NULL) {
        // Polymorphic behavior
        struct Shape *shapes[] = {
            (struct Shape*)rect,
            (struct Shape*)circle
        };

        for (int i = 0; i < 2; i++) {
            shapes[i]->print(shapes[i]);
            printf("Area: %.2f\n", shapes[i]->area(shapes[i]));
        }

        free(rect);
        free(circle);
    }

    return 0;
}

```

Structure with Nested Anonymous Structures (C11)

Using nested anonymous structures for cleaner APIs:

```

#include <stdio.h>

struct Point3D {
    union {
        struct {
            double x, y, z;
        };
        double coords[3];
    };
};

struct Color {
    union {
        struct {

```

```

        unsigned char r, g, b, a;
    };
    unsigned char components[4];
};

};

int main() {
    struct Point3D point = {1.0, 2.0, 3.0};
    struct Color color = {255, 128, 64, 255};

    // Access using named members
    printf("Point: (%.1f, %.1f, %.1f)\n", point.x, point.y, point.z);

    // Access using array
    printf("Coordinates: ");
    for (int i = 0; i < 3; i++) {
        printf("%.1f ", point.coords[i]);
    }
    printf("\n");

    // Access color components
    printf("Color: RGBA(%d, %d, %d, %d)\n",
           color.r, color.g, color.b, color.a);

    // Modify using array access
    for (int i = 0; i < 4; i++) {
        color.components[i] = 128;
    }

    printf("Modified color: RGBA(%d, %d, %d, %d)\n",
           color.r, color.g, color.b, color.a);

    return 0;
}

```

Practical Examples

Configuration Manager with Advanced Structures

Combining advanced structure concepts for a configuration manager:

```

#include <stdio.h>
#include <stdlib.h>
#include <string.h>

```

```

// Configuration entry types
typedef enum {
    CONFIG_INT,
    CONFIG_FLOAT,
    CONFIG_STRING,
    CONFIG_BOOL
} ConfigType;

// Configuration entry with bit fields
struct ConfigEntry {
    unsigned int type : 3;          // 3 bits for type
    unsigned int is_default : 1;    // 1 bit for default flag
    unsigned int is_locked : 1;    // 1 bit for lock flag
    unsigned int reserved : 27;    // 27 bits reserved

    char key[32];
    union {
        int int_value;
        float float_value;
        char *string_value;
        int bool_value;
    } value;
};

// Configuration manager
struct ConfigManager {
    struct ConfigEntry *entries;
    size_t count;
    size_t capacity;
};

struct ConfigManager* create_config_manager() {
    struct ConfigManager *cm = malloc(sizeof(struct ConfigManager));
    if (cm != NULL) {
        cm->entries = malloc(10 * sizeof(struct ConfigEntry));
        if (cm->entries != NULL) {
            cm->count = 0;
            cm->capacity = 10;
        } else {
            free(cm);
            cm = NULL;
        }
    }
    return cm;
}

```

```

int add_int_config(struct ConfigManager *cm, const char *key, int value) {
    if (cm->count >= cm->capacity) {
        struct ConfigEntry *new_entries = realloc(cm->entries,
            (cm->capacity + 10) * sizeof(struct ConfigEntry));
        if (new_entries == NULL) return -1;
        cm->entries = new_entries;
        cm->capacity += 10;
    }

    struct ConfigEntry *entry = &cm->entries[cm->count];
    entry->type = CONFIG_INT;
    entry->is_default = 0;
    entry->is_locked = 0;
    entry->reserved = 0;
    strncpy(entry->key, key, sizeof(entry->key) - 1);
    entry->key[sizeof(entry->key) - 1] = '\0';
    entry->value.int_value = value;

    cm->count++;
    return 0;
}

void print_config(const struct ConfigManager *cm) {
    for (size_t i = 0; i < cm->count; i++) {
        const struct ConfigEntry *entry = &cm->entries[i];
        printf("Key: %s, Type: %d, Default: %d, Locked: %d\n",
            entry->key, entry->type, entry->is_default, entry->is_locked);

        switch (entry->type) {
            case CONFIG_INT:
                printf("  Value: %d (int)\n", entry->value.int_value);
                break;
            // Add other cases as needed
        }
    }
}

void free_config_manager(struct ConfigManager *cm) {
    if (cm != NULL) {
        free(cm->entries);
        free(cm);
    }
}

int main() {
    struct ConfigManager *cm = create_config_manager();
    if (cm == NULL) {

```

```

        printf("Failed to create config manager\n");
        return 1;
    }

    add_int_config(cm, "window_width", 800);
    add_int_config(cm, "window_height", 600);
    add_int_config(cm, "max_fps", 60);

    print_config(cm);

    free_config_manager(cm);
    return 0;
}

```

Memory Pool with Flexible Array Members

Implementing a memory pool using flexible array members:

```

#include <stdio.h>
#include <stdlib.h>
#include <string.h>

struct MemoryBlock {
    size_t size;
    int is_free;
    char data[]; // Flexible array member
};

struct MemoryPool {
    size_t block_size;
    size_t total_blocks;
    struct MemoryBlock blocks[]; // Flexible array member
};

struct MemoryPool* create_memory_pool(size_t block_size, size_t num_blocks) {
    size_t pool_size = sizeof(struct MemoryPool) +
        num_blocks * (sizeof(struct MemoryBlock) + block_size);

    struct MemoryPool *pool = malloc(pool_size);
    if (pool != NULL) {
        pool->block_size = block_size;
        pool->total_blocks = num_blocks;

        // Initialize blocks
        char *block_data = (char*)(pool->blocks);
        for (size_t i = 0; i < num_blocks; i++) {

```

```

        struct MemoryBlock *block = (struct MemoryBlock*)block_data;
        block->size = block_size;
        block->is_free = 1;
        block_data += sizeof(struct MemoryBlock) + block_size;
    }
}
return pool;
}

void* pool_alloc(struct MemoryPool *pool, size_t size) {
    if (size > pool->block_size) return NULL;

    char *block_data = (char*)(pool->blocks);
    for (size_t i = 0; i < pool->total_blocks; i++) {
        struct MemoryBlock *block = (struct MemoryBlock*)block_data;
        if (block->is_free && block->size >= size) {
            block->is_free = 0;
            return block->data;
        }
        block_data += sizeof(struct MemoryBlock) + pool->block_size;
    }
    return NULL;
}

void pool_free(struct MemoryPool *pool, void *ptr) {
    if (ptr == NULL) return;

    char *pool_start = (char*)pool;
    char *ptr_char = (char*)ptr;

    if (ptr_char >= pool_start &&
        ptr_char < pool_start + sizeof(struct MemoryPool) +
            pool->total_blocks * (sizeof(struct MemoryBlock) + pool->block_size)) {

        char *block_data = (char*)(pool->blocks);
        for (size_t i = 0; i < pool->total_blocks; i++) {
            struct MemoryBlock *block = (struct MemoryBlock*)block_data;
            if (ptr_char == block->data) {
                block->is_free = 1;
                return;
            }
            block_data += sizeof(struct MemoryBlock) + pool->block_size;
        }
    }
}
}

```

```

void free_memory_pool(struct MemoryPool *pool) {
    free(pool);
}

int main() {
    struct MemoryPool *pool = create_memory_pool(64, 10);
    if (pool == NULL) {
        printf("Failed to create memory pool\n");
        return 1;
    }

    // Allocate memory from pool
    int *numbers = (int*)pool_alloc(pool, 10 * sizeof(int));
    if (numbers != NULL) {
        for (int i = 0; i < 10; i++) {
            numbers[i] = i * i;
        }

        printf("Numbers: ");
        for (int i = 0; i < 10; i++) {
            printf("%d ", numbers[i]);
        }
        printf("\n");

        // Free memory
        pool_free(pool, numbers);
    }

    free_memory_pool(pool);
    return 0;
}

```

Summary

Advanced structure concepts in C provide powerful mechanisms for creating efficient and sophisticated data representations:

1. **Bit Fields:** Enable precise control over memory usage for small values and hardware register representation
2. **Structure Packing:** Control memory alignment and padding for optimization
3. **Flexible Array Members:** Allow structures to have variable-sized arrays at the end (C99)
4. **Self-Referential Structures:** Enable implementation of linked data structures like lists and trees
5. **Function Pointers in Structures:** Support object-oriented programming concepts in C
6. **Anonymous Structures:** Provide cleaner APIs (C11)

These advanced concepts are essential for systems programming, embedded development, and performance-critical applications where memory efficiency and precise control over data representation are crucial.

Abstract Data Types

Introduction

Abstract Data Types (ADTs) are a fundamental concept in software engineering that separate the interface of a data type from its implementation. In C, ADTs are implemented using structures, function pointers, and information hiding principles to create modular, maintainable, and reusable code. This approach enables programmers to encapsulate data and operations, providing a clean interface while hiding implementation details.

Information Hiding Principles

Separation of Interface and Implementation

Information hiding is achieved by separating the public interface (what users can access) from the private implementation (internal details):

```
// stack.h - Public interface
#ifndef STACK_H
#define STACK_H

#include <stddef.h>

// Opaque pointer - hides implementation details
typedef struct Stack Stack;

// Public functions
Stack* stack_create(size_t capacity);
void stack_destroy(Stack *stack);
int stack_push(Stack *stack, int value);
int stack_pop(Stack *stack);
int stack_peek(const Stack *stack);
int stack_is_empty(const Stack *stack);
int stack_is_full(const Stack *stack);
size_t stack_size(const Stack *stack);

#endif // STACK_H
```

```

// stack.c - Private implementation
#include "stack.h"
#include <stdlib.h>
#include <stdbool.h>

// Actual structure definition (hidden from users)
struct Stack {
    int *data;
    size_t capacity;
    size_t top;
};

Stack* stack_create(size_t capacity) {
    Stack *stack = malloc(sizeof(Stack));
    if (stack != NULL) {
        stack->data = malloc(capacity * sizeof(int));
        if (stack->data != NULL) {
            stack->capacity = capacity;
            stack->top = 0;
        } else {
            free(stack);
            stack = NULL;
        }
    }
    return stack;
}

void stack_destroy(Stack *stack) {
    if (stack != NULL) {
        free(stack->data);
        free(stack);
    }
}

int stack_push(Stack *stack, int value) {
    if (stack == NULL || stack->top >= stack->capacity) {
        return -1; // Error
    }
    stack->data[stack->top] = value;
    stack->top++;
    return 0; // Success
}

int stack_pop(Stack *stack) {
    if (stack == NULL || stack->top == 0) {
        return -1; // Error
    }
}

```

```

    stack->top--;
    return stack->data[stack->top];
}

int stack_peek(const Stack *stack) {
    if (stack == NULL || stack->top == 0) {
        return -1; // Error
    }
    return stack->data[stack->top - 1];
}

int stack_is_empty(const Stack *stack) {
    return (stack == NULL || stack->top == 0);
}

int stack_is_full(const Stack *stack) {
    return (stack != NULL && stack->top >= stack->capacity);
}

size_t stack_size(const Stack *stack) {
    return (stack != NULL) ? stack->top : 0;
}

// main.c - Using the ADT
#include <stdio.h>
#include "stack.h"

int main() {
    Stack *stack = stack_create(10);
    if (stack == NULL) {
        printf("Failed to create stack\n");
        return 1;
    }

    // Use the stack
    stack_push(stack, 10);
    stack_push(stack, 20);
    stack_push(stack, 30);

    printf("Stack size: %zu\n", stack_size(stack));
    printf("Top element: %d\n", stack_peek(stack));

    while (!stack_is_empty(stack)) {
        printf("Popped: %d\n", stack_pop(stack));
    }
}

```

```

    stack_destroy(stack);
    return 0;
}

```

Opaque Pointers

Creating Opaque Pointers

Opaque pointers hide the implementation details of a data structure by only declaring the structure type in the header file:

```

// vector.h - Opaque pointer example
#ifndef VECTOR_H
#define VECTOR_H

#include <stddef.h>

// Forward declaration - no structure definition
typedef struct Vector Vector;

// Public interface
Vector* vector_create(size_t initial_capacity);
void vector_destroy(Vector *vector);
int vector_push_back(Vector *vector, int value);
int vector_pop_back(Vector *vector);
int vector_get(const Vector *vector, size_t index);
int vector_set(Vector *vector, size_t index, int value);
size_t vector_size(const Vector *vector);
size_t vector_capacity(const Vector *vector);

#endif // VECTOR_H

// vector.c - Implementation with opaque pointer
#include "vector.h"
#include <stdlib.h>

// Actual structure definition (private)
struct Vector {
    int *data;
    size_t size;
    size_t capacity;
};

Vector* vector_create(size_t initial_capacity) {

```

```

Vector *vector = malloc(sizeof(Vector));
if (vector != NULL) {
    vector->data = malloc(initial_capacity * sizeof(int));
    if (vector->data != NULL) {
        vector->size = 0;
        vector->capacity = initial_capacity;
    } else {
        free(vector);
        vector = NULL;
    }
}
return vector;
}

void vector_destroy(Vector *vector) {
    if (vector != NULL) {
        free(vector->data);
        free(vector);
    }
}

int vector_push_back(Vector *vector, int value) {
    if (vector == NULL) return -1;

    // Resize if needed
    if (vector->size >= vector->capacity) {
        size_t new_capacity = vector->capacity * 2;
        int *new_data = realloc(vector->data, new_capacity * sizeof(int));
        if (new_data == NULL) return -1;

        vector->data = new_data;
        vector->capacity = new_capacity;
    }

    vector->data[vector->size] = value;
    vector->size++;
    return 0;
}

int vector_pop_back(Vector *vector) {
    if (vector == NULL || vector->size == 0) return -1;
    vector->size--;
    return vector->data[vector->size];
}

int vector_get(const Vector *vector, size_t index) {
    if (vector == NULL || index >= vector->size) return -1;

```

```

    return vector->data[index];
}

int vector_set(Vector *vector, size_t index, int value) {
    if (vector == NULL || index >= vector->size) return -1;
    vector->data[index] = value;
    return 0;
}

size_t vector_size(const Vector *vector) {
    return (vector != NULL) ? vector->size : 0;
}

size_t vector_capacity(const Vector *vector) {
    return (vector != NULL) ? vector->capacity : 0;
}

```

Constructor and Destructor Patterns

Resource Acquisition Is Initialization (RAII) in C

Implementing constructor and destructor patterns for proper resource management:

```

// file_manager.h
#ifndef FILE_MANAGER_H
#define FILE_MANAGER_H

#include <stdio.h>

typedef struct FileManager FileManager;

FileManager* file_manager_create(const char *filename, const char *mode);
void file_manager_destroy(FileManager *fm);
int file_manager_write(FileManager *fm, const char *data);
int file_manager_read(FileManager *fm, char *buffer, size_t size);
long file_manager_get_size(FileManager *fm);

#endif // FILE_MANAGER_H

// file_manager.c
#include "file_manager.h"
#include <stdlib.h>
#include <string.h>

```

```

struct FileManager {
    FILE *file;
    char *filename;
    char *mode;
};

FileManager* file_manager_create(const char *filename, const char *mode) {
    FileManager *fm = malloc(sizeof(FileManager));
    if (fm == NULL) return NULL;

    fm->filename = malloc(strlen(filename) + 1);
    if (fm->filename == NULL) {
        free(fm);
        return NULL;
    }

    fm->mode = malloc(strlen(mode) + 1);
    if (fm->mode == NULL) {
        free(fm->filename);
        free(fm);
        return NULL;
    }

    strcpy(fm->filename, filename);
    strcpy(fm->mode, mode);

    fm->file = fopen(filename, mode);
    if (fm->file == NULL) {
        free(fm->mode);
        free(fm->filename);
        free(fm);
        return NULL;
    }

    return fm;
}

void file_manager_destroy(FileManager *fm) {
    if (fm != NULL) {
        if (fm->file != NULL) {
            fclose(fm->file);
        }
        free(fm->mode);
        free(fm->filename);
        free(fm);
    }
}

```



```

int file_manager_write(FileManager *fm, const char *data) {
    if (fm == NULL || fm->file == NULL || data == NULL) return -1;
    return fputs(data, fm->file) >= 0 ? 0 : -1;
}

int file_manager_read(FileManager *fm, char *buffer, size_t size) {
    if (fm == NULL || fm->file == NULL || buffer == NULL) return -1;
    return fgets(buffer, (int)size, fm->file) != NULL ? 0 : -1;
}

long file_manager_get_size(FileManager *fm) {
    if (fm == NULL || fm->file == NULL) return -1;

    long current_pos = ftell(fm->file);
    fseek(fm->file, 0, SEEK_END);
    long size = ftell(fm->file);
    fseek(fm->file, current_pos, SEEK_SET);

    return size;
}

```

Generic Abstract Data Types

Using Void Pointers for Generic ADTs

Creating generic ADTs that can work with any data type:

```

// generic_list.h
#ifndef GENERIC_LIST_H
#define GENERIC_LIST_H

#include <stddef.h>

typedef struct GenericList GenericList;

// Function pointer types for generic operations
typedef void (*destructor_func)(void *data);
typedef int (*compare_func)(const void *a, const void *b);
typedef void (*print_func)(const void *data);

GenericList* generic_list_create(destructor_func destructor);
void generic_list_destroy(GenericList *list);
int generic_list_append(GenericList *list, void *data);
int generic_list_prepend(GenericList *list, void *data);

```

```

void* generic_list_get(const GenericList *list, size_t index);
int generic_list_remove(GenericList *list, size_t index);
size_t generic_list_size(const GenericList *list);
void generic_list_print(const GenericList *list, print_func printer);

```

```

#endif // GENERIC_LIST_H

```

```

// generic_list.c

```

```

#include "generic_list.h"

```

```

#include <stdlib.h>

```

```

struct Node {
    void *data;
    struct Node *next;
};

struct GenericList {
    struct Node *head;
    size_t size;
    destructor_func destructor;
};

```

```

GenericList* generic_list_create(destructor_func destructor) {
    GenericList *list = malloc(sizeof(GenericList));
    if (list != NULL) {
        list->head = NULL;
        list->size = 0;
        list->destructor = destructor;
    }
    return list;
}

```

```

void generic_list_destroy(GenericList *list) {
    if (list != NULL) {
        struct Node *current = list->head;
        while (current != NULL) {
            struct Node *next = current->next;
            if (list->destructor != NULL) {
                list->destructor(current->data);
            }
            free(current);
            current = next;
        }
        free(list);
    }
}

```

```

int generic_list_append(GenericList *list, void *data) {
    if (list == NULL) return -1;

    struct Node *new_node = malloc(sizeof(struct Node));
    if (new_node == NULL) return -1;

    new_node->data = data;
    new_node->next = NULL;

    if (list->head == NULL) {
        list->head = new_node;
    } else {
        struct Node *current = list->head;
        while (current->next != NULL) {
            current = current->next;
        }
        current->next = new_node;
    }

    list->size++;
    return 0;
}

void* generic_list_get(const GenericList *list, size_t index) {
    if (list == NULL || index >= list->size) return NULL;

    struct Node *current = list->head;
    for (size_t i = 0; i < index && current != NULL; i++) {
        current = current->next;
    }

    return (current != NULL) ? current->data : NULL;
}

int generic_list_remove(GenericList *list, size_t index) {
    if (list == NULL || index >= list->size) return -1;

    struct Node *to_remove;
    if (index == 0) {
        to_remove = list->head;
        list->head = to_remove->next;
    } else {
        struct Node *current = list->head;
        for (size_t i = 0; i < index - 1 && current != NULL; i++) {
            current = current->next;
        }
    }
}

```

```

        if (current == NULL || current->next == NULL) return -1;

        to_remove = current->next;
        current->next = to_remove->next;
    }

    if (list->destructor != NULL) {
        list->destructor(to_remove->data);
    }
    free(to_remove);
    list->size--;

    return 0;
}

size_t generic_list_size(const GenericList *list) {
    return (list != NULL) ? list->size : 0;
}

void generic_list_print(const GenericList *list, print_func printer) {
    if (list == NULL || printer == NULL) return;

    struct Node *current = list->head;
    while (current != NULL) {
        printer(current->data);
        current = current->next;
    }
}

// Example usage with different data types
#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include "generic_list.h"

// Integer destructor
void int_destructor(void *data) {
    free(data);
}

// Integer printer
void int_printer(const void *data) {
    printf("%d ", *(int*)data);
}

// String destructor

```

```

void string_destructor(void *data) {
    free(data);
}

// String printer
void string_printer(const void *data) {
    printf("%s ", (char*)data);
}

int main() {
    // Integer list
    GenericList *int_list = generic_list_create(int_destructor);
    if (int_list != NULL) {
        for (int i = 1; i <= 5; i++) {
            int *value = malloc(sizeof(int));
            *value = i * 10;
            generic_list_append(int_list, value);
        }

        printf("Integer list: ");
        generic_list_print(int_list, int_printer);
        printf("\n");

        generic_list_destroy(int_list);
    }

    // String list
    GenericList *string_list = generic_list_create(string_destructor);
    if (string_list != NULL) {
        char *str1 = malloc(20);
        char *str2 = malloc(20);
        char *str3 = malloc(20);

        strcpy(str1, "Hello");
        strcpy(str2, "World");
        strcpy(str3, "C");

        generic_list_append(string_list, str1);
        generic_list_append(string_list, str2);
        generic_list_append(string_list, str3);

        printf("String list: ");
        generic_list_print(string_list, string_printer);
        printf("\n");

        generic_list_destroy(string_list);
    }
}

```

```

    return 0;
}

```

Interface Design in C

Designing Clean APIs

Creating well-designed interfaces that are easy to use and maintain:

```

// hash_table.h - Well-designed hash table interface
#ifndef HASH_TABLE_H
#define HASH_TABLE_H

#include <stddef.h>
#include <stdbool.h>

typedef struct HashTable HashTable;

// Function pointer types for custom operations
typedef size_t (*hash_func)(const void *key);
typedef int (*compare_func)(const void *key1, const void *key2);
typedef void (*destructor_func)(void *data);

// Error codes
typedef enum {
    HASH_SUCCESS = 0,
    HASH_ERROR_NULL_POINTER = -1,
    HASH_ERROR_MEMORY = -2,
    HASH_ERROR_KEY_NOT_FOUND = -3,
    HASH_ERROR_DUPLICATE_KEY = -4
} HashError;

// Public functions
HashTable* hash_table_create(size_t initial_capacity,
                             hash_func hash_function,
                             compare_func compare_function,
                             destructor_func key_destructor,
                             destructor_func value_destructor);

void hash_table_destroy(HashTable *table);
HashError hash_table_insert(HashTable *table, void *key, void *value);
HashError hash_table_remove(HashTable *table, const void *key);
void* hash_table_lookup(const HashTable *table, const void *key);
bool hash_table_contains(const HashTable *table, const void *key);

```

```

size_t hash_table_size(const HashTable *table);
void hash_table_clear(HashTable *table);

// Utility functions
size_t hash_string(const void *key); // Default string hash function
int compare_strings(const void *key1, const void *key2); // Default string compare

#endif // HASH_TABLE_H

// hash_table.c - Implementation
#include "hash_table.h"
#include <stdlib.h>
#include <string.h>

// Hash table entry
struct Entry {
    void *key;
    void *value;
    struct Entry *next;
};

// Hash table structure
struct HashTable {
    struct Entry **buckets;
    size_t bucket_count;
    size_t size;
    hash_func hash_function;
    compare_func compare_function;
    destructor_func key_destructor;
    destructor_func value_destructor;
};

// Default hash function for strings
size_t hash_string(const void *key) {
    const char *str = (const char*)key;
    size_t hash = 5381;
    int c;

    while ((c = *str++)) {
        hash = ((hash << 5) + hash) + c; // hash * 33 + c
    }

    return hash;
}

// Default compare function for strings

```

```

int compare_strings(const void *key1, const void *key2) {
    return strcmp((const char*)key1, (const char*)key2);
}

HashTable* hash_table_create(size_t initial_capacity,
                             hash_func hash_function,
                             compare_func compare_function,
                             destructor_func key_destructor,
                             destructor_func value_destructor) {
    HashTable *table = malloc(sizeof(HashTable));
    if (table == NULL) return NULL;

    table->buckets = calloc(initial_capacity, sizeof(struct Entry*));
    if (table->buckets == NULL) {
        free(table);
        return NULL;
    }

    table->bucket_count = initial_capacity;
    table->size = 0;
    table->hash_function = hash_function ? hash_function : hash_string;
    table->compare_function = compare_function ? compare_function : compare_strings;
    table->key_destructor = key_destructor;
    table->value_destructor = value_destructor;

    return table;
}

void hash_table_destroy(HashTable *table) {
    if (table != NULL) {
        hash_table_clear(table);
        free(table->buckets);
        free(table);
    }
}

HashError hash_table_insert(HashTable *table, void *key, void *value) {
    if (table == NULL || key == NULL) return HASH_ERROR_NULL_POINTER;

    size_t index = table->hash_function(key) % table->bucket_count;

    // Check for duplicate key
    struct Entry *current = table->buckets[index];
    while (current != NULL) {
        if (table->compare_function(current->key, key) == 0) {
            return HASH_ERROR_DUPLICATE_KEY;
        }
    }
}

```



```

        current = current->next;
    }

    // Create new entry
    struct Entry *new_entry = malloc(sizeof(struct Entry));
    if (new_entry == NULL) return HASH_ERROR_MEMORY;

    new_entry->key = key;
    new_entry->value = value;
    new_entry->next = table->buckets[index];
    table->buckets[index] = new_entry;
    table->size++;

    return HASH_SUCCESS;
}

void* hash_table_lookup(const HashTable *table, const void *key) {
    if (table == NULL || key == NULL) return NULL;

    size_t index = table->hash_function(key) % table->bucket_count;

    struct Entry *current = table->buckets[index];
    while (current != NULL) {
        if (table->compare_function(current->key, key) == 0) {
            return current->value;
        }
        current = current->next;
    }

    return NULL;
}

HashError hash_table_remove(HashTable *table, const void *key) {
    if (table == NULL || key == NULL) return HASH_ERROR_NULL_POINTER;

    size_t index = table->hash_function(key) % table->bucket_count;

    struct Entry *current = table->buckets[index];
    struct Entry *previous = NULL;

    while (current != NULL) {
        if (table->compare_function(current->key, key) == 0) {
            if (previous == NULL) {
                table->buckets[index] = current->next;
            } else {
                previous->next = current->next;
            }
        }
    }
}

```

```

        if (table->key_destructor != NULL) {
            table->key_destructor(current->key);
        }
        if (table->value_destructor != NULL) {
            table->value_destructor(current->value);
        }

        free(current);
        table->size--;
        return HASH_SUCCESS;
    }

    previous = current;
    current = current->next;
}

return HASH_ERROR_KEY_NOT_FOUND;
}

bool hash_table_contains(const HashTable *table, const void *key) {
    return hash_table_lookup(table, key) != NULL;
}

size_t hash_table_size(const HashTable *table) {
    return (table != NULL) ? table->size : 0;
}

void hash_table_clear(HashTable *table) {
    if (table == NULL) return;

    for (size_t i = 0; i < table->bucket_count; i++) {
        struct Entry *current = table->buckets[i];
        while (current != NULL) {
            struct Entry *next = current->next;

            if (table->key_destructor != NULL) {
                table->key_destructor(current->key);
            }
            if (table->value_destructor != NULL) {
                table->value_destructor(current->value);
            }

            free(current);
            current = next;
        }
        table->buckets[i] = NULL;
    }
}

```

```

    }

    table->size = 0;
}

```

Practical Examples

Complete ADT Implementation: Dynamic Array

A complete implementation of a dynamic array ADT with comprehensive error handling:

```

// dynamic_array.h
#ifndef DYNAMIC_ARRAY_H
#define DYNAMIC_ARRAY_H

#include <stddef.h>

typedef struct DynamicArray DynamicArray;

// Error codes
typedef enum {
    DYNARRAY_SUCCESS = 0,
    DYNARRAY_ERROR_NULL_POINTER = -1,
    DYNARRAY_ERROR_OUT_OF_BOUNDS = -2,
    DYNARRAY_ERROR_MEMORY = -3,
    DYNARRAY_ERROR_INVALID_SIZE = -4
} DynamicArrayError;

// Function pointer types
typedef void (*destructor_func)(void *data);
typedef int (*compare_func)(const void *a, const void *b);

// Public interface
DynamicArray* dynamic_array_create(size_t element_size, destructor_func destructor);
void dynamic_array_destroy(DynamicArray *array);
DynamicArrayError dynamic_array_push_back(DynamicArray *array, const void *element);
DynamicArrayError dynamic_array_pop_back(DynamicArray *array, void *element);
DynamicArrayError dynamic_array_get(const DynamicArray *array, size_t index, void *element);
DynamicArrayError dynamic_array_set(DynamicArray *array, size_t index, const void *element);
DynamicArrayError dynamic_array_insert(DynamicArray *array, size_t index, const void *element);
DynamicArrayError dynamic_array_remove(DynamicArray *array, size_t index);
size_t dynamic_array_size(const DynamicArray *array);
size_t dynamic_array_capacity(const DynamicArray *array);
DynamicArrayError dynamic_array_resize(DynamicArray *array, size_t new_size);
DynamicArrayError dynamic_array_clear(DynamicArray *array);

```

```

DynamicArrayError dynamic_array_sort(DynamicArray *array, compare_func compare);

#endif // DYNAMIC_ARRAY_H

// dynamic_array.c
#include "dynamic_array.h"
#include <stdlib.h>
#include <string.h>

#define INITIAL_CAPACITY 8
#define GROWTH_FACTOR 2

struct DynamicArray {
    void *data;
    size_t size;
    size_t capacity;
    size_t element_size;
    destructor_func destructor;
};

DynamicArray* dynamic_array_create(size_t element_size, destructor_func destructor) {
    if (element_size == 0) return NULL;

    DynamicArray *array = malloc(sizeof(DynamicArray));
    if (array == NULL) return NULL;

    array->data = malloc(INITIAL_CAPACITY * element_size);
    if (array->data == NULL) {
        free(array);
        return NULL;
    }

    array->size = 0;
    array->capacity = INITIAL_CAPACITY;
    array->element_size = element_size;
    array->destructor = destructor;

    return array;
}

void dynamic_array_destroy(DynamicArray *array) {
    if (array != NULL) {
        dynamic_array_clear(array);
        free(array->data);
        free(array);
    }
}

```

```

}

DynamicArrayError dynamic_array_push_back(DynamicArray *array, const void *element) {
    if (array == NULL || element == NULL) return DYNARRAY_ERROR_NULL_POINTER;

    if (array->size >= array->capacity) {
        size_t new_capacity = array->capacity * GROWTH_FACTOR;
        void *new_data = realloc(array->data, new_capacity * array->element_size);
        if (new_data == NULL) return DYNARRAY_ERROR_MEMORY;

        array->data = new_data;
        array->capacity = new_capacity;
    }

    void *dest = (char*)array->data + array->size * array->element_size;
    memcpy(dest, element, array->element_size);
    array->size++;

    return DYNARRAY_SUCCESS;
}

DynamicArrayError dynamic_array_pop_back(DynamicArray *array, void *element) {
    if (array == NULL) return DYNARRAY_ERROR_NULL_POINTER;
    if (array->size == 0) return DYNARRAY_ERROR_OUT_OF_BOUNDS;

    if (element != NULL) {
        void *src = (char*)array->data + (array->size - 1) * array->element_size;
        memcpy(element, src, array->element_size);
    }

    array->size--;
    return DYNARRAY_SUCCESS;
}

DynamicArrayError dynamic_array_get(const DynamicArray *array, size_t index, void *element)
    if (array == NULL || element == NULL) return DYNARRAY_ERROR_NULL_POINTER;
    if (index >= array->size) return DYNARRAY_ERROR_OUT_OF_BOUNDS;

    void *src = (char*)array->data + index * array->element_size;
    memcpy(element, src, array->element_size);

    return DYNARRAY_SUCCESS;
}

DynamicArrayError dynamic_array_set(DynamicArray *array, size_t index, const void *element)
    if (array == NULL || element == NULL) return DYNARRAY_ERROR_NULL_POINTER;

```

```

    if (index >= array->size) return DYNARRAY_ERROR_OUT_OF_BOUNDS;

    void *dest = (char*)array->data + index * array->element_size;
    memcpy(dest, element, array->element_size);

    return DYNARRAY_SUCCESS;
}

size_t dynamic_array_size(const DynamicArray *array) {
    return (array != NULL) ? array->size : 0;
}

size_t dynamic_array_capacity(const DynamicArray *array) {
    return (array != NULL) ? array->capacity : 0;
}

DynamicArrayError dynamic_array_clear(DynamicArray *array) {
    if (array == NULL) return DYNARRAY_ERROR_NULL_POINTER;

    if (array->destructor != NULL) {
        for (size_t i = 0; i < array->size; i++) {
            void *element = (char*)array->data + i * array->element_size;
            array->destructor(element);
        }
    }

    array->size = 0;
    return DYNARRAY_SUCCESS;
}

DynamicArrayError dynamic_array_sort(DynamicArray *array, compare_func compare) {
    if (array == NULL || compare == NULL) return DYNARRAY_ERROR_NULL_POINTER;

    // Simple bubble sort implementation
    for (size_t i = 0; i < array->size - 1; i++) {
        for (size_t j = 0; j < array->size - i - 1; j++) {
            void *elem1 = (char*)array->data + j * array->element_size;
            void *elem2 = (char*)array->data + (j + 1) * array->element_size;

            if (compare(elem1, elem2) > 0) {
                // Swap elements
                char *temp = malloc(array->element_size);
                if (temp == NULL) return DYNARRAY_ERROR_MEMORY;

                memcpy(temp, elem1, array->element_size);
                memcpy(elem1, elem2, array->element_size);
            }
        }
    }
}

```

```

        memcpy(elem2, temp, array->element_size);

        free(temp);
    }
}

return DYNARRAY_SUCCESS;
}

// Example usage
#include <stdio.h>
#include <stdlib.h>
#include "dynamic_array.h"

// Integer destructor
void int_destructor(void *data) {
    // No special cleanup needed for integers
}

// Integer compare function
int compare_ints(const void *a, const void *b) {
    int int_a = *(const int*)a;
    int int_b = *(const int*)b;
    return (int_a > int_b) - (int_a < int_b);
}

int main() {
    DynamicArray *array = dynamic_array_create(sizeof(int), int_destructor);
    if (array == NULL) {
        printf("Failed to create dynamic array\n");
        return 1;
    }

    // Add elements
    for (int i = 10; i >= 1; i--) {
        DynamicArrayError err = dynamic_array_push_back(array, &i);
        if (err != DYNARRAY_SUCCESS) {
            printf("Error adding element: %d\n", err);
            break;
        }
    }

    printf("Array size: %zu, capacity: %zu\n",
        dynamic_array_size(array), dynamic_array_capacity(array));
}

```

```

// Print elements
printf("Before sorting: ");
for (size_t i = 0; i < dynamic_array_size(array); i++) {
    int value;
    if (dynamic_array_get(array, i, &value) == DYNARRAY_SUCCESS) {
        printf("%d ", value);
    }
}
printf("\n");

// Sort array
DynamicArrayError err = dynamic_array_sort(array, compare_ints);
if (err == DYNARRAY_SUCCESS) {
    printf("After sorting: ");
    for (size_t i = 0; i < dynamic_array_size(array); i++) {
        int value;
        if (dynamic_array_get(array, i, &value) == DYNARRAY_SUCCESS) {
            printf("%d ", value);
        }
    }
    printf("\n");
}

dynamic_array_destroy(array);
return 0;
}

```

Summary

Abstract Data Types in C provide a powerful mechanism for creating modular, maintainable, and reusable code:

1. **Information Hiding:** Separate interface from implementation using opaque pointers
2. **Opaque Pointers:** Hide structure details by only declaring types in headers
3. **Constructor/Destructor Patterns:** Ensure proper resource management
4. **Generic ADTs:** Use void pointers and function pointers for type flexibility
5. **Clean Interface Design:** Provide clear, consistent APIs with proper error handling
6. **Memory Management:** Implement proper allocation and deallocation strategies

These concepts enable C programmers to write code that is as maintainable and reusable as code written in object-oriented languages, while maintaining C's performance characteristics and low-level control.

Module 7: Structures and User-Defined Types

Overview

This module explores structures and user-defined types in C, essential for creating complex data representations and abstract data types. You'll learn about structure declaration, initialization, and usage, as well as unions, enumerations, and advanced structure concepts. The module also covers information hiding principles and object-oriented programming techniques in C.

Learning Objectives

By the end of this module, you will be able to: - Declare, define, and use structures effectively - Initialize structures using various methods - Work with nested structures and arrays of structures - Understand structure packing and alignment - Implement unions and enumerations appropriately - Apply bit fields for memory-efficient data representation - Create abstract data types using information hiding principles - Design opaque pointers for encapsulation - Implement constructor/destructor patterns in C - Apply best practices for structure design and usage

Chapters

1. **Structures** - Structure declaration, definition, initialization, and usage
2. **Unions and Enumerations** - Union declaration and usage, enumeration types and values
3. **Advanced Structure Concepts** - Bit fields, structure packing, flexible array members, self-referential structures
4. **Abstract Data Types** - Information hiding principles, opaque pointers, constructor/destructor patterns, interface design

Key Concepts Covered

- Structure declaration and member access
- Structure initialization techniques
- Nested structures and arrays of structures
- Pointers to structures and the arrow operator
- Structure packing and memory alignment
- Union declaration and size optimization
- Enumeration types and explicit value assignment
- Bit fields for memory-efficient storage

- Flexible array members (C99)
- Self-referential structures for linked data
- Information hiding and encapsulation
- Opaque pointers for abstraction
- Constructor and destructor patterns
- Interface design in C

Prerequisites

- Completion of Module 1: Foundations and Environment Setup
- Completion of Module 2: Data Types and Variables
- Completion of Module 3: Control Flow
- Completion of Module 4: Functions and Modular Programming
- Completion of Module 5: Arrays and Strings
- Completion of Module 6: Pointers and Memory Management
- Understanding of basic programming concepts
- Familiarity with variable declaration and initialization

Tools and Technologies

- C compiler (GCC, Clang, or MSVC)
- Text editor or IDE
- Memory debugging tools (Valgrind, AddressSanitizer)
- Understanding of basic programming concepts
- Familiarity with debugging tools

Estimated Time to Complete

- Reading: 10-12 hours
- Exercises: 15-20 hours
- Projects: 8-10 hours

Assessment

- Complete all chapter exercises
- Successfully implement complex data structures using structures
- Demonstrate understanding of unions, enumerations, and bit fields
- Pass structure programming quiz
- Submit a comprehensive program using abstract data types

Next Module

After completing this module, proceed to **Module 8: File I/O and System Programming** to learn about file operations and system-level programming.

Additional Resources

- “The C Programming Language” by Kernighan and Ritchie (Chapter 6)
- C standard documentation (ISO/IEC 9899:2024 for C23)
- Online tutorials on data structure implementation
- Debugging tools documentation (GDB, Valgrind, AddressSanitizer)
- Best practices guides for structure design

File I/O

File Operations

File operations let programs persist data, process large datasets, and talk to the rest of the system. In C, the high-level interface lives in `<stdio.h>`: open a stream (`FILE *`), read or write, then close it. This chapter focuses on **Linux/POSIX** practice with complete, buildable programs.

Compile everything with:

```
gcc -std=c17 -Wall -Wextra -o program program.c
```

File Pointers and the FILE Structure

All stdio file work goes through a `FILE *`:

```
#include <stdio.h>

FILE *fp; /* stream handle */
```

The exact layout of `FILE` is implementation-defined. Conceptually it holds buffering state, current position, error/EOF flags, and the underlying file descriptor.

Standard streams (already open):

Stream	Meaning	Typical fd
<code>stdin</code>	Standard input	0
<code>stdout</code>	Standard output	1
<code>stderr</code>	Standard error	2

Opening Files: `fopen`

```
FILE *fopen(const char *filename, const char *mode);
```

Returns a stream on success, or `NULL` on failure (check `errno`).

Modes

Mode	Access	Position	Create?	Truncate?
"r"	read	start	no	no
"w"	write	start	yes	yes
"a"	write	end	yes	no
"r+"	read/write	start	no	no
"w+"	read/write	start	yes	yes
"a+"	read/write	end (writes)	yes	no

Append **b** for binary ("rb", "wb", ...). On Linux, "b" has little effect (no CRLF translation), but **always use it for non-text data** for portability and clarity.

Minimal open / close

```
#include <stdio.h>
#include <stdlib.h>

int main(void) {
    FILE *fp = fopen("data.txt", "r");
    if (fp == NULL) {
        perror("fopen data.txt");
        return EXIT_FAILURE;
    }

    /* ... use fp ... */

    if (fclose(fp) != 0) {
        perror("fclose");
        return EXIT_FAILURE;
    }
    return EXIT_SUCCESS;
}
```

Text vs Binary Mode (Caveats)

Text mode ("r", "w", ...)

- Oriented toward characters and lines
- On some systems (notably Windows), `\n` may be translated to `\r\n` on write and the reverse on read
- On Linux, translation is typically a no-op, but programs that care about **exact bytes** should still use binary mode

Binary mode ("rb", "wb", ...)

- Bytes in, bytes out — no newline translation
- Correct choice for images, compressed data, structs dumped as bytes, executables
- **Caveat:** writing a raw `struct` with `fwrite` is **not** a portable file format (padding, endianness, alignment). Prefer explicit serialization for interchange

```
#include <stdio.h>
#include <stdlib.h>

int main(void) {
    FILE *tf = fopen("note.txt", "w");
    FILE *bf = fopen("nums.bin", "wb");
    int nums[] = {1, 2, 3, 4, 5};

    if (tf == NULL || bf == NULL) {
        perror("fopen");
        return EXIT_FAILURE;
    }

    fprintf(tf, "Hello, World!\n");
    if (fwrite(nums, sizeof nums[0], 5, bf) != 5) {
        fprintf(stderr, "short write or error on nums.bin\n");
    }

    fclose(tf);
    fclose(bf);
    return 0;
}
```

Portability note: Never assume `sizeof(int)` or endianness when exchanging binary files across machines. Use fixed-width types (`uint32_t`) and explicit byte order if the file leaves your process.

Closing and Buffering

```
int fclose(FILE *stream); /* 0 success, EOF failure */
int fflush(FILE *stream); /* flush buffers to OS */
```

- **Always** `fclose` streams you opened (and check the result after writes — buffered data may fail only at flush/close).
- `fflush(stdout)` after prompts when you need the user to see text before a blocking read.
- `setvbuf` can change buffering mode (`_IOFBF`, `_IOLBF`, `_IONBF`).

```

#include <stdio.h>

int main(void) {
    printf("Enter name: ");
    fflush(stdout); /* ensure prompt appears */

    char name[64];
    if (fgets(name, sizeof name, stdin) == NULL) {
        return 1;
    }
    printf("Hi, %s", name);
    return 0;
}

```

Error Handling: fopen, ferror, errno

Check every open

```

#include <stdio.h>
#include <stdlib.h>
#include <errno.h>
#include <string.h>

int main(void) {
    FILE *fp = fopen("missing.txt", "r");
    if (fp == NULL) {
        /* perror uses errno */
        perror("fopen missing.txt");

        /* or format yourself */
        fprintf(stderr, "fopen failed: %s (errno=%d)\n",
                strerror(errno), errno);
        return EXIT_FAILURE;
    }
    fclose(fp);
    return 0;
}

```

Common failure reasons: file not found (ENOENT), permission denied (EACCES), is a directory, too many open files, disk full on write.

Distinguish EOF from error

fgetc returns EOF both at end-of-file **and** on error. Use feof / ferror:


```

#include <stdio.h>
#include <stdlib.h>

int main(void) {
    FILE *fp = fopen("example.txt", "r");
    int ch;

    if (fp == NULL) {
        perror("fopen");
        return EXIT_FAILURE;
    }

    while ((ch = fgetc(fp)) != EOF) {
        putchar(ch);
    }

    if (ferror(fp)) {
        fprintf(stderr, "read error on example.txt\n");
        clearerr(fp);
        fclose(fp);
        return EXIT_FAILURE;
    }
    if (feof(fp)) {
        /* normal end */
    }

    fclose(fp);
    return 0;
}

```

Function	Role
<code>feof(fp)</code>	Non-zero if EOF indicator is set
<code>ferror(fp)</code>	Non-zero if error indicator is set
<code>clearerr(fp)</code>	Clears both indicators
<code>perror(msg)</code>	Prints <code>msg: <strerror(errno)>\n</code> to <code>stderr</code>
<code>strerror(errno)</code>	Thread-aware alternatives exist (<code>strerror_r</code>) on POSIX

Pattern: cleanup on partial failure

When multiple resources open, close what you already opened before returning:

```

FILE *in = fopen(src, "rb");
if (!in) {

```

```

        perror(src);
        return 1;
    }
    FILE *out = fopen(dst, "wb");
    if (!out) {
        perror(dst);
        fclose(in);
        return 1;
    }
    /* ... */
    fclose(out);
    fclose(in);

```

Full Program 1: Copy a File

Binary-safe copy with block reads (faster than byte-at-a-time for large files):

```

/* file: fcopy.c */
#include <stdio.h>
#include <stdlib.h>

#define BUFSZ 8192

int main(int argc, char *argv[]) {
    FILE *src = NULL;
    FILE *dst = NULL;
    unsigned char buf[BUFSZ];
    size_t n;
    int status = EXIT_SUCCESS;

    if (argc != 3) {
        fprintf(stderr, "Usage: %s <source> <destination>\n", argv[0]);
        return EXIT_FAILURE;
    }

    src = fopen(argv[1], "rb");
    if (src == NULL) {
        perror(argv[1]);
        return EXIT_FAILURE;
    }

    dst = fopen(argv[2], "wb");
    if (dst == NULL) {

```

```

        perror(argv[2]);
        fclose(src);
        return EXIT_FAILURE;
    }

    while ((n = fread(buf, 1, sizeof buf, src)) > 0) {
        if (fwrite(buf, 1, n, dst) != n) {
            fprintf(stderr, "write error: %s\n", argv[2]);
            status = EXIT_FAILURE;
            break;
        }
    }

    if (ferror(src)) {
        fprintf(stderr, "read error: %s\n", argv[1]);
        status = EXIT_FAILURE;
    }

    if (fclose(dst) != 0) {
        perror("fclose destination");
        status = EXIT_FAILURE;
    }
    if (fclose(src) != 0) {
        perror("fclose source");
        status = EXIT_FAILURE;
    }

    if (status == EXIT_SUCCESS) {
        fprintf(stderr, "copied %s -> %s\n", argv[1], argv[2]);
    }
    return status;
}

```

```

gcc -std=c17 -Wall -Wextra -o fcopy fcopy.c
./fcopy photo.jpg photo.bak
cmp photo.jpg photo.bak && echo identical

```

Full Program 2: Simplified wc (Lines, Words, Bytes)

Counts lines, words (whitespace-separated tokens), and bytes — like a small `wc`:

```

/* file: mywc.c */
#include <stdio.h>
#include <stdlib.h>
#include <ctype.h>
#include <stdbool.h>

struct counts {
    unsigned long lines;
    unsigned long words;
    unsigned long bytes;
};

static int count_stream(FILE *fp, struct counts *c) {
    int ch;
    bool in_word = false;

    c->lines = c->words = c->bytes = 0;

    while ((ch = fgetc(fp)) != EOF) {
        c->bytes++;

        if (ch == '\n') {
            c->lines++;
        }

        if (isspace((unsigned char)ch)) {
            in_word = false;
        } else if (!in_word) {
            in_word = true;
            c->words++;
        }
    }

    if (ferror(fp)) {
        return -1;
    }
    return 0;
}

static void print_counts(const struct counts *c, const char *name) {
    if (name != NULL) {
        printf("%8lu %8lu %8lu %s\n", c->lines, c->words, c->bytes, name);
    } else {
        printf("%8lu %8lu %8lu\n", c->lines, c->words, c->bytes);
    }
}

```

```

int main(int argc, char *argv[]) {
    struct counts total = {0, 0, 0};
    int i;
    int status = EXIT_SUCCESS;

    if (argc == 1) {
        struct counts c;
        if (count_stream(stdin, &c) != 0) {
            fprintf(stderr, "read error on stdin\n");
            return EXIT_FAILURE;
        }
        print_counts(&c, NULL);
        return EXIT_SUCCESS;
    }

    for (i = 1; i < argc; i++) {
        FILE *fp = fopen(argv[i], "rb");
        struct counts c;

        if (fp == NULL) {
            perror(argv[i]);
            status = EXIT_FAILURE;
            continue;
        }

        if (count_stream(fp, &c) != 0) {
            fprintf(stderr, "read error: %s\n", argv[i]);
            status = EXIT_FAILURE;
            fclose(fp);
            continue;
        }
        fclose(fp);

        print_counts(&c, argv[i]);
        total.lines += c.lines;
        total.words += c.words;
        total.bytes += c.bytes;
    }

    if (argc > 2) {
        print_counts(&total, "total");
    }

    return status;
}

```

```
gcc -std=c17 -Wall -Wextra -o mywc mywc.c
echo -e "hello world\nfoo" > sample.txt
./mywc sample.txt
./mywc sample.txt sample.txt
printf 'a b c' | ./mywc
```

Notes

- Binary mode avoids surprises on platforms with newline translation when counting bytes.
- `isspace((unsigned char)ch)` avoids undefined behavior for negative `char` values.
- Words follow a simple “transition into non-space” definition, similar in spirit to classic `wc`.

Full Program 3: Append Log Lines with Timestamps

Append-only log with ISO-like local timestamps:

```
/* file: logappend.c */
#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <time.h>
#include <errno.h>

static int timestamp_now(char *buf, size_t buflen) {
    time_t t = time(NULL);
    struct tm tm_now;

    if (t == (time_t)-1) {
        return -1;
    }
    if (localtime_r(&t, &tm_now) == NULL) {
        return -1;
    }
    /* YYYY-MM-DD HH:MM:SS */
    if (strftime(buf, buflen, "%Y-%m-%d %H:%M:%S", &tm_now) == 0) {
        return -1;
    }
    return 0;
}

int main(int argc, char *argv[]) {
    const char *path;
```

```

FILE *fp;
char ts[32];
char line[1024];

if (argc != 2) {
    fprintf(stderr, "Usage: %s <logfile>\n", argv[0]);
    fprintf(stderr, "Reads lines from stdin and appends them with timestamps.\n");
    return EXIT_FAILURE;
}
path = argv[1];

fp = fopen(path, "a");
if (fp == NULL) {
    fprintf(stderr, "fopen %s: %s\n", path, strerror(errno));
    return EXIT_FAILURE;
}

while (fgets(line, sizeof line, stdin) != NULL) {
    size_t len = strlen(line);

    /* strip trailing newline for clean log formatting */
    if (len > 0 && line[len - 1] == '\n') {
        line[len - 1] = '\0';
    }

    if (timestamp_now(ts, sizeof ts) != 0) {
        fprintf(stderr, "timestamp failed\n");
        fclose(fp);
        return EXIT_FAILURE;
    }

    if (fprintf(fp, "[%s] %s\n", ts, line) < 0) {
        fprintf(stderr, "write failed: %s\n", strerror(errno));
        fclose(fp);
        return EXIT_FAILURE;
    }

    if (fflush(fp) != 0) {
        perror("fflush");
        fclose(fp);
        return EXIT_FAILURE;
    }
}

if (ferror(stdin)) {
    fprintf(stderr, "error reading stdin\n");
    fclose(fp);
    return EXIT_FAILURE;
}

```

```

    }

    if (fclose(fp) != 0) {
        perror("fclose");
        return EXIT_FAILURE;
    }

    return EXIT_SUCCESS;
}

gcc -std=c17 -Wall -Wextra -o logappend logappend.c
echo "server started" | ./logappend app.log
echo "accepted client" | ./logappend app.log
cat app.log
# [2026-04-01 14:22:10] server started
# [2026-04-01 14:22:11] accepted client

```

`fopen(..., "a")` guarantees writes append even if another process also appends (with the usual caveats about line atomicity for large writes).

File Size and Status Helpers

```

/* file: fsize.c */
#include <stdio.h>
#include <stdlib.h>

int main(int argc, char *argv[]) {
    FILE *fp;
    long size;

    if (argc != 2) {
        fprintf(stderr, "Usage: %s <file>\n", argv[0]);
        return EXIT_FAILURE;
    }

    fp = fopen(argv[1], "rb");
    if (fp == NULL) {
        perror(argv[1]);
        return EXIT_FAILURE;
    }

    if (fseek(fp, 0, SEEK_END) != 0) {
        perror("fseek");
    }
}

```



```

        fclose(fp);
        return EXIT_FAILURE;
    }
    size = ftell(fp);
    if (size < 0) {
        perror("ftell");
        fclose(fp);
        return EXIT_FAILURE;
    }

    printf("%s: %ld bytes\n", argv[1], size);
    fclose(fp);
    return 0;
}

```

For huge files, prefer `fseeko` / `ftello` and `off_t` (define `_FILE_OFFSET_BITS=64` on some 32-bit systems). POSIX `stat()` is often clearer for metadata.

Best Practices

1. Check `fopen`, `fread`/`fwrite` counts, `ferror`, and `fclose` after writes.
 2. **Close** everything you open; use reverse-order cleanup on error paths.
 3. Use `"rb"/"wb"` for non-text or exact byte copies.
 4. Prefer **block I/O** (`fread`/`fwrite` with an 4–64 KiB buffer) over `fgetc` for bulk data.
 5. Use `perror` / `strerror(errno)` so users know *why* open failed.
 6. **fflush** when the next operation needs data on disk or the user needs a prompt.
 7. Do not ignore **short writes** — disks fill up; pipes break.
 8. Treat raw `struct` dumps as non-portable binary unless you control every reader.
-

Exercises

Exercise 1 — Safe truncate-free write

Write `write_all.txt` using mode `"w"`. Then change the program to open with `"r+"` and overwrite only the first line without truncating the rest of the file. Explain when `"w"` would destroy data.

Exercise 2 — Byte-at-a-time vs block copy

Implement file copy with `fgetc/fputc` and again with an 8 KiB buffer. Time both on a ~50 MB file (`time ./fcopy ...`). Record the difference.

Exercise 3 — Extend `mywc`

Add a `-l` flag that prints only the line count (like `wc -l`). Keep the multi-file total behavior.

Exercise 4 — Rotating log

Modify `logappend` so that if the log exceeds 10000 bytes (`ftell` after open with `"a+"` and seek end), it renames `app.log` to `app.log.1` (use `rename(2)`) and opens a fresh log.

Exercise 5 — Error path drill

Point `fcopy` at a source you can read and a destination on a full filesystem or a path without write permission. Confirm you get a useful message and a non-zero exit code, and that the source stream is closed.

Exercise 6 — Binary round-trip

Write an array of ten `uint32_t` values to `data.bin` with `fwrite`, then read them back and print in hex. Re-open the same file in text mode and discuss whether anything appears different on Linux (and what would change on Windows).

Topic	Key API / idea
-------	----------------

Summary

Topic	Key API / idea
Open / close	<code>fopen</code> , <code>fclose</code>
Modes	<code>"r"/"w"/"a"</code> and <code>+</code> / <code>b</code> variants
Text vs binary	Use <code>b</code> for exact bytes; watch portability of structs
Errors	NULL from <code>fopen</code> , <code>errno</code> , <code>ferror</code> vs <code>feof</code>
Bulk I/O	<code>fread</code> / <code>fwrite</code> with a fixed buffer
Practice programs	<code>fcopy</code> , <code>mywc</code> , <code>logappend</code>

Mastering careful open/check/copy/close patterns is the foundation for every later file-processing and system program in C.

File I/O Functions

Introduction

C provides a rich set of functions for performing input and output operations on files. These functions can be categorized into character I/O, line I/O, formatted I/O, and binary I/O functions. Each category serves different purposes and offers varying levels of control and convenience for file processing tasks.

Understanding these functions is crucial for effectively reading from and writing to files, whether you're processing text data, binary data, or structured information.

Compile everything with:

```
gcc -std=c17 -Wall -Wextra -o program program.c
```

Prefer complete programs: create input files, run tools, inspect outputs.

Character I/O Functions

Character I/O functions allow reading and writing files one character at a time. These functions are useful for simple text processing and when fine-grained control over file operations is needed.

fgetc() and getc()

Both functions read a single character from a file:

```
int fgetc(FILE *stream);  
int getc(FILE *stream);
```

Key Points: - Returns the character read as an **unsigned char** cast to **int** - Returns EOF on end of file or error - **fgetc()** is a function, **getc()** may be a macro

fputc() and putc()

Both functions write a single character to a file:

```
int fputc(int char, FILE *stream);
int putc(int char, FILE *stream);
```

Key Points: - Writes the character (cast to `unsigned char`) to the file - Returns the written character on success - Returns EOF on error - `fputc()` is a function, `putc()` may be a macro

ungetc()

Pushes a character back onto the input stream:

```
int ungetc(int char, FILE *stream);
```

Key Points: - Allows “unread” operations for parsing - Only one character can be pushed back - Returns the character on success, EOF on failure

Examples

```
#include <stdio.h>
#include <stdlib.h>
#include <ctype.h>

int main() {
    FILE *input, *output;
    int ch;

    input = fopen("input.txt", "r");
    output = fopen("output.txt", "w");

    if (input == NULL || output == NULL) {
        printf("Error opening files\n");
        exit(EXIT_FAILURE);
    }

    // Convert lowercase to uppercase
    while ((ch = fgetc(input)) != EOF) {
        if (islower(ch)) {
            ch = toupper(ch);
        }
        fputc(ch, output);
    }
}
```

```

// Demonstrate ungetc
ch = fgetc(input);
if (ch != EOF) {
    ungetc(ch, input); // Push character back
    ch = fgetc(input); // Read it again
    printf("Character read twice: %c\n", ch);
}

fclose(input);
fclose(output);

return 0;
}

```

Line I/O Functions

Line I/O functions are designed for reading and writing entire lines of text, making them ideal for processing text files where data is organized in lines.

fgets()

Reads a line from a file:

```
char *fgets(char *str, int count, FILE *stream);
```

Key Points: - Reads at most `count-1` characters - Stops at newline or end of file - Always null-terminates the string - Stores newline in the buffer (if space available) - Returns `str` on success, `NULL` on error or EOF

fputs()

Writes a string to a file:

```
int fputs(const char *str, FILE *stream);
```

Key Points: - Writes string without null terminator - Does not append newline automatically - Returns non-negative value on success - Returns EOF on error

Examples

```
#include <stdio.h>
#include <stdlib.h>
#include <string.h>

int main() {
    FILE *fp;
    char line[256];

    fp = fopen("lines.txt", "w");
    if (fp == NULL) {
        printf("Error opening file for writing\n");
        exit(EXIT_FAILURE);
    }

    // Write lines to file
    fputs("First line\n", fp);
    fputs("Second line\n", fp);
    fputs("Third line\n", fp);

    fclose(fp);

    // Read lines from file
    fp = fopen("lines.txt", "r");
    if (fp == NULL) {
        printf("Error opening file for reading\n");
        exit(EXIT_FAILURE);
    }

    while (fgets(line, sizeof(line), fp) != NULL) {
        // Remove newline if present
        line[strcspn(line, "\n")] = '\0';
        printf("Read line: '%s'\n", line);
    }

    fclose(fp);
    return 0;
}
```

Formatted I/O Functions

Formatted I/O functions provide high-level input and output operations similar to `printf()` and `scanf()`, but operate on files instead of standard streams.

fprintf()

Writes formatted output to a file:

```
int fprintf(FILE *stream, const char *format, ...);
```

Key Points: - Works like `printf()` but writes to a file - Returns number of characters written - Returns negative value on error

fscanf()

Reads formatted input from a file:

```
int fscanf(FILE *stream, const char *format, ...);
```

Key Points: - Works like `scanf()` but reads from a file - Returns number of items successfully read - Returns EOF on error or end of file

Examples

```
#include <stdio.h>
#include <stdlib.h>

typedef struct {
    int id;
    char name[50];
    double salary;
} Employee;

int main() {
    FILE *fp;
    Employee emp;

    // Write structured data to file
    fp = fopen("employees.txt", "w");
    if (fp == NULL) {
        printf("Error opening file for writing\n");
        exit(EXIT_FAILURE);
    }

    fprintf(fp, "%d %s %.2f\n", 1, "John Doe", 50000.0);
    fprintf(fp, "%d %s %.2f\n", 2, "Jane Smith", 60000.0);
    fprintf(fp, "%d %s %.2f\n", 3, "Bob Johnson", 55000.0);

    fclose(fp);
}
```



```

// Read structured data from file
fp = fopen("employees.txt", "r");
if (fp == NULL) {
    printf("Error opening file for reading\n");
    exit(EXIT_FAILURE);
}

printf("Employee Records:\n");
while (fscanf(fp, "%d %s %lf", &emp.id, emp.name, &emp.salary) == 3) {
    printf("ID: %d, Name: %s, Salary: %.2f\n",
           emp.id, emp.name, emp.salary);
}

fclose(fp);
return 0;
}

```

Binary I/O Functions

Binary I/O functions allow reading and writing raw data in binary format, preserving the exact bit patterns of the data. These functions are essential for working with binary files like images, executables, or structured data.

fwrite()

Writes binary data to a file:

```
size_t fwrite(const void *ptr, size_t size, size_t count, FILE *stream);
```

Key Points: - Writes `count` items of `size` bytes each - Returns number of items successfully written - Returns value less than `count` on error

fread()

Reads binary data from a file:

```
size_t fread(void *ptr, size_t size, size_t count, FILE *stream);
```

Key Points: - Reads `count` items of `size` bytes each - Returns number of items successfully read - Returns value less than `count` on error or EOF

Examples

```
#include <stdio.h>
#include <stdlib.h>

typedef struct {
    int id;
    char name[20];
    double value;
} Record;

int main() {
    FILE *fp;
    Record records[] = {
        {1, "Record1", 100.5},
        {2, "Record2", 200.75},
        {3, "Record3", 300.25}
    };
    Record read_records[3];
    size_t written, read;

    // Write binary data
    fp = fopen("records.bin", "wb");
    if (fp == NULL) {
        printf("Error opening file for writing\n");
        exit(EXIT_FAILURE);
    }

    written = fwrite(records, sizeof(Record), 3, fp);
    if (written != 3) {
        printf("Error writing records\n");
    }

    fclose(fp);

    // Read binary data
    fp = fopen("records.bin", "rb");
    if (fp == NULL) {
        printf("Error opening file for reading\n");
        exit(EXIT_FAILURE);
    }

    read = fread(read_records, sizeof(Record), 3, fp);
    if (read != 3) {
        printf("Error reading records\n");
    }
}
```

```

fclose(fp);

// Display read data
printf("Read Records:\n");
for (int i = 0; i < 3; i++) {
    printf("ID: %d, Name: %s, Value: %.2f\n",
           read_records[i].id, read_records[i].name, read_records[i].value);
}

return 0;
}

```

Practical Examples

CSV File Processor

```

#include <stdio.h>
#include <stdlib.h>
#include <string.h>

typedef struct {
    int id;
    char name[50];
    int age;
    double salary;
} Employee;

int main() {
    FILE *input, *output;
    char line[256];
    Employee emp;
    double total_salary = 0;
    int count = 0;

    input = fopen("employees.csv", "r");
    output = fopen("report.txt", "w");

    if (input == NULL || output == NULL) {
        printf("Error opening files\n");
        exit(EXIT_FAILURE);
    }

    // Skip header line
    fgets(line, sizeof(line), input);

```

```

fprintf(output, "Employee Report\n");
fprintf(output, "=====\n");

// Process each line
while (fgets(line, sizeof(line), input) != NULL) {
    // Remove newline
    line[strcspn(line, "\n")] = '\0';

    // Parse CSV line
    if (sscanf(line, "%d,%[^,],%d,%lf",
               &emp.id, emp.name, &emp.age, &emp.salary) == 4) {
        fprintf(output, "ID: %d, Name: %s, Age: %d, Salary: %.2f\n",
               emp.id, emp.name, emp.age, emp.salary);
        total_salary += emp.salary;
        count++;
    }
}

if (count > 0) {
    fprintf(output, "\nAverage Salary: %.2f\n", total_salary / count);
}

fclose(input);
fclose(output);

printf("Report generated successfully\n");
return 0;
}

```

Binary File Analyzer

```

#include <stdio.h>
#include <stdlib.h>

int main(int argc, char *argv[]) {
    FILE *fp;
    unsigned char buffer[16];
    long offset = 0;
    size_t bytes_read;

    if (argc != 2) {
        printf("Usage: %s <filename>\n", argv[0]);
        exit(EXIT_FAILURE);
    }
}

```

```

fp = fopen(argv[1], "rb");
if (fp == NULL) {
    perror("Error opening file");
    exit(EXIT_FAILURE);
}

printf("Offset  Hex Dump                                ASCII\n");
printf("-----  -----                                -----\n");

while ((bytes_read = fread(buffer, 1, sizeof(buffer), fp)) > 0) {
    // Print offset
    printf("%06lx  ", offset);

    // Print hex values
    for (size_t i = 0; i < sizeof(buffer); i++) {
        if (i < bytes_read) {
            printf("%02x ", buffer[i]);
        } else {
            printf("   ");
        }

        // Add extra space after 8 bytes
        if (i == 7) printf(" ");
    }

    printf(" ");

    // Print ASCII representation
    for (size_t i = 0; i < bytes_read; i++) {
        if (buffer[i] >= 32 && buffer[i] <= 126) {
            printf("%c", buffer[i]);
        } else {
            printf(".");
        }
    }

    printf("\n");
    offset += bytes_read;
}

fclose(fp);
return 0;
}

```

Error Handling and Best Practices

Comprehensive Error Checking

```
#include <stdio.h>
#include <stdlib.h>

int safe_file_copy(const char *source, const char *destination) {
    FILE *src, *dst;
    int ch;

    // Open source file
    src = fopen(source, "rb");
    if (src == NULL) {
        perror("Error opening source file");
        return -1;
    }

    // Open destination file
    dst = fopen(destination, "wb");
    if (dst == NULL) {
        perror("Error opening destination file");
        fclose(src);
        return -1;
    }

    // Copy data
    while ((ch = fgetc(src)) != EOF) {
        if (fputc(ch, dst) == EOF) {
            perror("Error writing to destination file");
            fclose(src);
            fclose(dst);
            return -1;
        }
    }

    // Check for read errors
    if (ferror(src)) {
        perror("Error reading source file");
        fclose(src);
        fclose(dst);
        return -1;
    }

    // Close files
    if (fclose(src) != 0) {
```

```

        perror("Error closing source file");
        fclose(dst);
        return -1;
    }

    if (fclose(dst) != 0) {
        perror("Error closing destination file");
        return -1;
    }

    return 0; // Success
}

int main() {
    if (safe_file_copy("source.txt", "destination.txt") == 0) {
        printf("File copied successfully\n");
    } else {
        printf("File copy failed\n");
    }

    return 0;
}

```

Full Programs: Character, Line, Formatted, Binary

Program 1 — ROT13 filter (rot13.c)

Character I/O over stdin/stdout (and optional files via shell redirection).

```

/* file: rot13.c */
#include <stdio.h>
#include <stdlib.h>
#include <ctype.h>

static int rot13_char(int ch) {
    if (ch >= 'a' && ch <= 'z') {
        return 'a' + (ch - 'a' + 13) % 26;
    }
    if (ch >= 'A' && ch <= 'Z') {
        return 'A' + (ch - 'A' + 13) % 26;
    }
    return ch;
}

int main(void) {

```

```

    int ch;
    while ((ch = fgetc(stdin)) != EOF) {
        if (fputc(rot13_char(ch), stdout) == EOF) {
            perror("fputc");
            return EXIT_FAILURE;
        }
    }
    if (ferror(stdin)) {
        perror("fgetc");
        return EXIT_FAILURE;
    }
    return EXIT_SUCCESS;
}

gcc -std=c17 -Wall -Wextra -o rot13 rot13.c
echo 'Hello, World!' | ./rot13
echo 'Hello, World!' | ./rot13 | ./rot13    # identity
./rot13 < input.txt > output.txt

```

Program 2 — Line numberer with long-line handling (nlnum.c)

```

/* file: nlnum.c */
#include <stdio.h>
#include <stdlib.h>
#include <string.h>

int main(int argc, char *argv[]) {
    FILE *in = stdin;
    char buf[64];
    unsigned long lineno = 0;
    int in_line = 0; /* true if we already printed a number for current logical line */

    if (argc > 2) {
        fprintf(stderr, "Usage: %s [file]\n", argv[0]);
        return EXIT_FAILURE;
    }
    if (argc == 2) {
        in = fopen(argv[1], "r");
        if (in == NULL) {
            perror(argv[1]);
            return EXIT_FAILURE;
        }
    }

    while (fgets(buf, sizeof buf, in) != NULL) {

```



```

        if (!in_line) {
            lineno++;
            printf("%6lu\t", lineno);
            in_line = 1;
        }
        fputs(buf, stdout);
        if (strchr(buf, '\n') != NULL) {
            in_line = 0; /* finished a full line */
        }
    }

    if (ferror(in)) {
        perror("read");
        if (in != stdin) {
            fclose(in);
        }
        return EXIT_FAILURE;
    }
    if (in != stdin) {
        fclose(in);
    }
    return EXIT_SUCCESS;
}

```

```

gcc -std=c17 -Wall -Wextra -o nlnum nlnum.c
printf 'short\n' > sample.txt
# a line longer than the buffer still gets one line number:
python3 -c 'print("x"*100)' >> sample.txt
./nlnum sample.txt

```

Program 3 — CSV-ish records with fprintf / fscanf (scores_io.c)

```

/* file: scores_io.c */
#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <errno.h>

struct Score {
    char name[32];
    int points;
};

static int write_scores(const char *path, const struct Score *rows, size_t n) {
    FILE *fp = fopen(path, "w");

```

```

size_t i;

if (fp == NULL) {
    fprintf(stderr, "fopen %s: %s\n", path, strerror(errno));
    return -1;
}

for (i = 0; i < n; i++) {
    if (fprintf(fp, "%s %d\n", rows[i].name, rows[i].points) < 0) {
        perror("fprintf");
        fclose(fp);
        return -1;
    }
}

if (fclose(fp) != 0) {
    perror("fclose");
    return -1;
}

return 0;
}

static int read_scores(const char *path) {
    FILE *fp = fopen(path, "r");
    char name[32];
    int points;
    int nread;

    if (fp == NULL) {
        fprintf(stderr, "fopen %s: %s\n", path, strerror(errno));
        return -1;
    }

    printf("name\tpoints\n");
    printf("----\t-----\n");
    while ((nread = fscanf(fp, "%31s %d", name, &points)) == 2) {
        printf("%s\t%d\n", name, points);
    }

    if (ferror(fp)) {
        perror("fscanf");
        fclose(fp);
        return -1;
    }

    /* nread==EOF or failed match ends loop - OK for this simple format */
    fclose(fp);
    return 0;
}

int main(void) {

```

```

    struct Score rows[] = {
        {"alice", 120},
        {"bob", 95},
        {"carol", 140}
    };

    if (write_scores("scores.txt", rows, 3) != 0) {
        return EXIT_FAILURE;
    }
    if (read_scores("scores.txt") != 0) {
        return EXIT_FAILURE;
    }
    return EXIT_SUCCESS;
}

```

```

gcc -std=c17 -Wall -Wextra -o scores_io scores_io.c
./scores_io
cat scores.txt

```

Caveat: `fscanf("%31s")` stops at whitespace — names with spaces need a different format (or `fgets` + parse).

Program 4 — Binary record dump / load (`binrec.c`)

```

/* file: binrec.c */
#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <stdint.h>
#include <errno.h>

struct Rec {
    uint32_t id;
    int32_t value;
    char tag[8]; /* fixed field; not a portable interchange format */
};

static int save_recs(const char *path, const struct Rec *recs, size_t n) {
    FILE *fp = fopen(path, "wb");
    size_t wrote;

    if (fp == NULL) {
        fprintf(stderr, "fopen %s: %s\n", path, strerror(errno));
        return -1;
    }
}

```

```

wrote = fwrite(recs, sizeof recs[0], n, fp);
if (wrote != n) {
    fprintf(stderr, "short write (%zu of %zu)\n", wrote, n);
    fclose(fp);
    return -1;
}
if (fclose(fp) != 0) {
    perror("fclose");
    return -1;
}
return 0;
}

static int load_and_print(const char *path) {
    FILE *fp = fopen(path, "rb");
    struct Rec r;
    size_t n;

    if (fp == NULL) {
        fprintf(stderr, "fopen %s: %s\n", path, strerror(errno));
        return -1;
    }
    while ((n = fread(&r, sizeof r, 1, fp)) == 1) {
        /* ensure tag is printable for display */
        char tag[9];
        memcpy(tag, r.tag, 8);
        tag[8] = '\0';
        printf("id=%u value=%d tag=\"%s\"\n", r.id, r.value, tag);
    }
    if (ferror(fp)) {
        perror("fread");
        fclose(fp);
        return -1;
    }
    fclose(fp);
    return 0;
}

int main(void) {
    struct Rec data[2];

    memset(data, 0, sizeof data);
    data[0].id = 1;
    data[0].value = 42;
    memcpy(data[0].tag, "alpha", 6);
    data[1].id = 2;
    data[1].value = -7;
}

```

```

memcpy(data[1].tag, "beta", 5);

if (save_recs("recs.bin", data, 2) != 0) {
    return EXIT_FAILURE;
}
if (load_and_print("recs.bin") != 0) {
    return EXIT_FAILURE;
}

/* hex peek of the file */
{
    FILE *fp = fopen("recs.bin", "rb");
    unsigned char buf[16];
    size_t n;
    size_t i;
    if (fp) {
        n = fread(buf, 1, sizeof buf, fp);
        printf("first %zu bytes:", n);
        for (i = 0; i < n; i++) {
            printf(" %02x", buf[i]);
        }
        printf("\n");
        fclose(fp);
    }
}
return EXIT_SUCCESS;
}

```

```

gcc -std=c17 -Wall -Wextra -o binrec binrec.c
./binrec
xxd recs.bin | head

```

Portability note: raw `struct` dumps depend on padding and endianness. Fine for same-machine cache files; use explicit serialization for network or long-term formats.

Program 5 — ungetc simple number peek (`peek_num.c`)

```

/* file: peek_num.c */
#include <stdio.h>
#include <stdlib.h>
#include <ctype.h>

/* Read an integer from stdin; first non-digit is pushed back. */
static int read_int(int *out) {
    int ch;

```

```

    int sign = 1;
    long val = 0;
    int digits = 0;

    ch = fgetc(stdin);
    while (ch != EOF && isspace((unsigned char)ch)) {
        ch = fgetc(stdin);
    }
    if (ch == '+' || ch == '-') {
        sign = (ch == '-') ? -1 : 1;
        ch = fgetc(stdin);
    }
    while (ch != EOF && isdigit((unsigned char)ch)) {
        digits++;
        val = val * 10 + (ch - '0');
        ch = fgetc(stdin);
    }
    if (ch != EOF) {
        if (ungetc(ch, stdin) == EOF) {
            return -1;
        }
    }
    if (digits == 0) {
        return -1;
    }
    *out = (int)(sign * val);
    return 0;
}

int main(void) {
    int a, b;
    int ch;

    printf("Type: <int><op><int> e.g. 12+34\n");
    if (read_int(&a) != 0) {
        fprintf(stderr, "expected integer\n");
        return EXIT_FAILURE;
    }
    ch = fgetc(stdin);
    if (ch != '+' && ch != '-' && ch != '*' ) {
        fprintf(stderr, "expected operator, got %d\n", ch);
        return EXIT_FAILURE;
    }
    if (read_int(&b) != 0) {
        fprintf(stderr, "expected second integer\n");
        return EXIT_FAILURE;
    }
}

```

```

        switch (ch) {
        case '+': printf("%d\n", a + b); break;
        case '-': printf("%d\n", a - b); break;
        case '*': printf("%d\n", a * b); break;
        default: return EXIT_FAILURE;
        }
    return EXIT_SUCCESS;
}

gcc -std=c17 -Wall -Wextra -o peek_num peek_num.c
printf '12+34\n' | ./peek_num

```

Program 6 — Block copy with progress (bcopy.c)

```

/* file: bcopy.c */
#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <errno.h>

#define BUF_SZ 4096

int main(int argc, char *argv[]) {
    FILE *in, *out;
    unsigned char buf[BUF_SZ];
    size_t n;
    unsigned long long total = 0;

    if (argc != 3) {
        fprintf(stderr, "Usage: %s <src> <dst>\n", argv[0]);
        return EXIT_FAILURE;
    }

    in = fopen(argv[1], "rb");
    if (in == NULL) {
        fprintf(stderr, "open %s: %s\n", argv[1], strerror(errno));
        return EXIT_FAILURE;
    }

    out = fopen(argv[2], "wb");
    if (out == NULL) {
        fprintf(stderr, "open %s: %s\n", argv[2], strerror(errno));
        fclose(in);
        return EXIT_FAILURE;
    }
}

```

```

while ((n = fread(buf, 1, sizeof buf, in)) > 0) {
    size_t off = 0;
    while (off < n) {
        size_t w = fwrite(buf + off, 1, n - off, out);
        if (w == 0) {
            perror("fwrite");
            fclose(in);
            fclose(out);
            return EXIT_FAILURE;
        }
        off += w;
    }
    total += n;
}
if (ferror(in)) {
    perror("fread");
    fclose(in);
    fclose(out);
    return EXIT_FAILURE;
}

if (fclose(in) != 0 || fclose(out) != 0) {
    perror("fclose");
    return EXIT_FAILURE;
}
fprintf(stderr, "copied %llu bytes\n", total);
return EXIT_SUCCESS;
}

```

```

gcc -std=c17 -Wall -Wextra -o bcopy bcopy.c
./bcopy scores.txt scores.bak

```

Error Handling Patterns (Quick Reference)

Function	Success	Failure / EOF
fgetc / getc	character as int	EOF — then check ferror vs feof
fputc / putc	character written	EOF
fgets	str pointer	NULL
fputs	non-negative	EOF
fprintf	char count 0	negative
fscanf	count of assigned fields	EOF or short count

Function	Success	Failure / EOF
<code>fread</code> / <code>fwrite</code>	items transferred	short count — check <code>ferror</code>

```

int ch = fgetc(fp);
if (ch == EOF) {
    if (ferror(fp)) {
        perror("read");
    } else {
        /* clean EOF */
    }
}

```

Summary

C's file I/O functions provide comprehensive capabilities for working with files:

1. **Character I/O:** `fgetc()`, `fputc()`, `ungetc()` for single character operations
2. **Line I/O:** `fgets()`, `fputs()` for line-based text processing
3. **Formatted I/O:** `fprintf()`, `fscanf()` for structured data input/output
4. **Binary I/O:** `fwrite()`, `fread()` for raw binary data operations

Each category of functions serves specific purposes and offers different levels of control and convenience. Proper error handling and understanding of text vs. binary modes are essential for robust file processing applications.

Exercises

```
gcc -std=c17 -Wall -Wextra -o exN exN.c
```

1. **Case flip** — Write a filter that toggles case of letters (`Hello` → `hELLO`) using `fgetc/fputc` and `ctype.h`.
2. **nlnum width** — Add `-w N` to control line-number field width (default 6).
3. **TSV writer** — Store three columns (string, int, double) with `fprintf` as tab-separated values; read them back with `fscanf` or `fgets+parse`.
4. **Binary double array** — Write 100 double values with `fwrite`, read them back, print the sum. Confirm file size is `100 * sizeof(double)`.

5. **Short write drill** — Modify `bcopy` to write in 1-byte `fwrite` loops (slow) and compare; still handle short writes correctly.
6. **`ungetc` comment skipper** — Skip C-style `//` comments from `stdin`: when you see `/` then `/`, discard until newline; use `ungetc` when the second char is not `/`.
7. **EOF vs error** — Open a directory as a file (e.g. `fopen(".", "r")` on Linux may fail) or force an error path; print whether you hit `ferror` or clean EOF on a normal file.

Advanced File Operations

Introduction

Advanced file operations in C go beyond basic reading and writing to include sophisticated techniques such as file positioning, random access, temporary file management, and concurrent access control. These operations are essential for building efficient file processing applications, database systems, and other programs that require fine-grained control over file access.

Understanding these advanced concepts enables programmers to create more robust and performant file-handling code, particularly when dealing with large files or complex data structures.

File Positioning

File positioning functions allow you to move the file pointer to specific locations within a file, enabling random access to data. This is particularly useful for databases, configuration files, and other structured data formats.

fseek()

Moves the file pointer to a specified position:

```
int fseek(FILE *stream, long offset, int origin);
```

Parameters: - **stream:** File pointer - **offset:** Number of bytes to move (can be negative) - **origin:** Reference point for the offset

Origin Values: - **SEEK_SET:** Beginning of file - **SEEK_CUR:** Current position - **SEEK_END:** End of file

Return Value: - 0 on success - Non-zero on error

ftell()

Returns the current position of the file pointer:

```
long ftell(FILE *stream);
```

Return Value: - Current file position (in bytes from beginning) - -1L on error

rewind()

Moves the file pointer to the beginning of the file:

```
void rewind(FILE *stream);
```

Examples

```
#include <stdio.h>
#include <stdlib.h>

int main() {
    FILE *fp;
    char buffer[100];

    // Create a test file
    fp = fopen("test.txt", "w");
    if (fp == NULL) {
        perror("Error opening file");
        exit(EXIT_FAILURE);
    }

    fprintf(fp, "Line 1: This is the first line.\n");
    fprintf(fp, "Line 2: This is the second line.\n");
    fprintf(fp, "Line 3: This is the third line.\n");
    fclose(fp);

    // Open for reading and demonstrate positioning
    fp = fopen("test.txt", "r");
    if (fp == NULL) {
        perror("Error opening file");
        exit(EXIT_FAILURE);
    }

    // Read first line
    fgets(buffer, sizeof(buffer), fp);
    printf("First line: %s", buffer);

    // Move to beginning of file
    rewind(fp);
    fgets(buffer, sizeof(buffer), fp);
    printf("First line again: %s", buffer);

    // Move to position 20 from beginning
    fseek(fp, 20, SEEK_SET);
```

```

fgets(buffer, sizeof(buffer), fp);
printf("From position 20: %s", buffer);

// Get current position
long pos = ftell(fp);
printf("Current position: %ld\n", pos);

// Move 10 bytes forward from current position
fseek(fp, 10, SEEK_CUR);
pos = ftell(fp);
printf("Position after moving 10 bytes forward: %ld\n", pos);

fclose(fp);
return 0;
}

```

Random Access Files

Random access files allow direct access to any part of the file without reading through preceding data. This is particularly useful for databases and other applications where quick access to specific records is required.

Structured Data Example

```

#include <stdio.h>
#include <stdlib.h>
#include <string.h>

#define RECORD_SIZE 64

typedef struct {
    int id;
    char name[20];
    double value;
} Record;

// Write records to file
void write_records(const char *filename) {
    FILE *fp;
    Record records[] = {
        {1, "Record1", 100.5},
        {2, "Record2", 200.75},
        {3, "Record3", 300.25},
        {4, "Record4", 400.0},
    };
}

```

```

        {5, "Record5", 500.75}
    };

    fp = fopen(filename, "wb");
    if (fp == NULL) {
        perror("Error opening file");
        exit(EXIT_FAILURE);
    }

    for (int i = 0; i < 5; i++) {
        fwrite(&records[i], sizeof(Record), 1, fp);
    }

    fclose(fp);
}

// Read specific record by ID
int read_record(const char *filename, int record_id, Record *record) {
    FILE *fp;
    long position = (record_id - 1) * sizeof(Record);

    fp = fopen(filename, "rb");
    if (fp == NULL) {
        perror("Error opening file");
        return -1;
    }

    // Position to the specific record
    if (fseek(fp, position, SEEK_SET) != 0) {
        perror("Error seeking to record");
        fclose(fp);
        return -1;
    }

    // Read the record
    if (fread(record, sizeof(Record), 1, fp) != 1) {
        if (feof(fp)) {
            printf("Record not found\n");
        } else {
            perror("Error reading record");
        }
        fclose(fp);
        return -1;
    }

    fclose(fp);
    return 0;
}

```

```

}

// Update specific record by ID
int update_record(const char *filename, int record_id, const Record *new_record) {
    FILE *fp;
    long position = (record_id - 1) * sizeof(Record);

    fp = fopen(filename, "r+b");
    if (fp == NULL) {
        perror("Error opening file");
        return -1;
    }

    // Position to the specific record
    if (fseek(fp, position, SEEK_SET) != 0) {
        perror("Error seeking to record");
        fclose(fp);
        return -1;
    }

    // Write the updated record
    if (fwrite(new_record, sizeof(Record), 1, fp) != 1) {
        perror("Error writing record");
        fclose(fp);
        return -1;
    }

    fclose(fp);
    return 0;
}

int main() {
    Record record;

    // Create test file with records
    write_records("records.dat");

    // Read specific record
    if (read_record("records.dat", 3, &record) == 0) {
        printf("Record 3: ID=%d, Name=%s, Value=%.2f\n",
            record.id, record.name, record.value);
    }

    // Update record
    Record updated = {3, "Updated3", 999.99};
    if (update_record("records.dat", 3, &updated) == 0) {

```

```

        printf("Record 3 updated successfully\n");
    }

    // Read updated record
    if (read_record("records.dat", 3, &record) == 0) {
        printf("Updated Record 3: ID=%d, Name=%s, Value=%.2f\n",
            record.id, record.name, record.value);
    }

    return 0;
}

```

Temporary Files

Temporary files are useful for storing data temporarily during program execution. C provides functions to create temporary files that are automatically managed by the system.

tmpfile()

Creates a temporary binary file that is automatically deleted when closed:

```
FILE *tmpfile(void);
```

Return Value: - File pointer on success - NULL on error

tmpnam()

Generates a unique temporary filename:

```
char *tmpnam(char *str);
```

Parameters: - **str:** Buffer to store filename (if NULL, uses internal static buffer)

Return Value: - Pointer to temporary filename - NULL on error

mkstemp() (POSIX)

Creates and opens a unique temporary file (more secure):

```
int mkstemp(char *template);
```


Examples

```
#include <stdio.h>
#include <stdlib.h>
#include <string.h>

int main() {
    FILE *tmp_fp;
    char tmp_name[L_tmpnam];
    int fd;

    // Method 1: Using tmpfile()
    tmp_fp = tmpfile();
    if (tmp_fp == NULL) {
        perror("Error creating temporary file");
        exit(EXIT_FAILURE);
    }

    fprintf(tmp_fp, "This is temporary data\n");
    fprintf(tmp_fp, "It will be deleted automatically\n");

    // Rewind and read back
    rewind(tmp_fp);
    int ch;
    while ((ch = fgetc(tmp_fp)) != EOF) {
        putchar(ch);
    }

    // File is automatically deleted when closed
    fclose(tmp_fp);

    // Method 2: Using tmpnam()
    if (tmpnam(tmp_name) == NULL) {
        printf("Error generating temporary filename\n");
        exit(EXIT_FAILURE);
    }

    printf("Temporary filename: %s\n", tmp_name);

    tmp_fp = fopen(tmp_name, "w+");
    if (tmp_fp == NULL) {
        perror("Error opening temporary file");
        exit(EXIT_FAILURE);
    }

    fprintf(tmp_fp, "Data in named temporary file\n");
```

```

fclose(tmp_fp);

// Note: tmpnam() doesn't delete the file automatically
// You need to remove it manually
if (remove(tmp_name) != 0) {
    perror("Error deleting temporary file");
} else {
    printf("Temporary file deleted successfully\n");
}

return 0;
}

```

File Locking and Concurrent Access

When multiple processes or threads access the same file simultaneously, file locking mechanisms are necessary to prevent data corruption and ensure consistency.

flockfile() and funlockfile()

Provide explicit locking for FILE objects:

```

void flockfile(FILE *file);
void funlockfile(FILE *file);

```

ftrylockfile()

Attempts to lock a file without blocking:

```

int ftrylockfile(FILE *file);

```

Example with File Locking

```

#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>
#include <pthread.h>

void* writer_thread(void* arg) {
    FILE *fp = (FILE*)arg;
    int thread_id = *(int*)pthread_getspecific(pthread_self());

```

```

    for (int i = 0; i < 5; i++) {
        // Lock the file
        flockfile(fp);

        fprintf(fp, "Writer %d: Line %d\n", thread_id, i);
        fflush(fp);

        // Unlock the file
        funlockfile(fp);

        sleep(1); // Simulate work
    }

    return NULL;
}

int main() {
    FILE *fp;
    pthread_t threads[3];
    int thread_ids[3] = {1, 2, 3};

    fp = fopen("shared.log", "w");
    if (fp == NULL) {
        perror("Error opening file");
        exit(EXIT_FAILURE);
    }

    // Create writer threads
    for (int i = 0; i < 3; i++) {
        pthread_create(&threads[i], NULL, writer_thread, fp);
    }

    // Wait for threads to complete
    for (int i = 0; i < 3; i++) {
        pthread_join(threads[i], NULL);
    }

    fclose(fp);
    printf("All threads completed\n");

    return 0;
}

```

File Status and Information

C provides functions to query file status and obtain information about files without opening them.

fileno()

Returns the file descriptor associated with a FILE pointer:

```
int fileno(FILE *stream);
```

fstat() (POSIX)

Gets file status information:

```
#include <sys/stat.h>
int fstat(int fd, struct stat *buf);
```

Example with File Information

```
#include <stdio.h>
#include <stdlib.h>
#include <sys/stat.h>
#include <time.h>
#include <unistd.h>

void print_file_info(const char *filename) {
    struct stat file_info;

    if (stat(filename, &file_info) != 0) {
        perror("Error getting file information");
        return;
    }

    printf("File: %s\n", filename);
    printf("Size: %ld bytes\n", file_info.st_size);
    printf("Permissions: %o\n", file_info.st_mode & 0777);
    printf("Last modified: %s", ctime(&file_info.st_mtime));
    printf("Number of links: %ld\n", file_info.st_nlink);

    // Check file type
    if (S_ISREG(file_info.st_mode)) {
        printf("Type: Regular file\n");
    } else if (S_ISDIR(file_info.st_mode)) {
        printf("Type: Directory\n");
    } else if (S_ISLNK(file_info.st_mode)) {
        printf("Type: Symbolic link\n");
    }
}
```

```

int main() {
    FILE *fp;

    // Create a test file
    fp = fopen("info_test.txt", "w");
    if (fp == NULL) {
        perror("Error creating test file");
        exit(EXIT_FAILURE);
    }

    fprintf(fp, "This is a test file for information retrieval.\n");
    fprintf(fp, "It contains multiple lines of text.\n");
    fclose(fp);

    // Print file information
    print_file_info("info_test.txt");

    return 0;
}

```

Buffering Strategies

Understanding and controlling buffering can significantly impact file I/O performance.

setvbuf()

Controls buffering behavior:

```

int setvbuf(FILE *stream, char *buffer, int mode, size_t size);

```

Mode Values: - `_IOFBF`: Full buffering - `_IOLBF`: Line buffering - `_IONBF`: No buffering

Example with Custom Buffering

```

#include <stdio.h>
#include <stdlib.h>

int main() {
    FILE *fp;
    char buffer[1024];

    fp = fopen("buffer_test.txt", "w");
    if (fp == NULL) {

```

```

        perror("Error opening file");
        exit(EXIT_FAILURE);
    }

    // Set custom buffer
    if (setvbuf(fp, buffer, _IOFBF, sizeof(buffer)) != 0) {
        printf("Error setting buffer\n");
        fclose(fp);
        exit(EXIT_FAILURE);
    }

    // Write data (will be buffered)
    for (int i = 0; i < 100; i++) {
        fprintf(fp, "Line %d: This is test data.\n", i);
    }

    // Force buffer flush
    fflush(fp);

    fclose(fp);
    printf("Data written with custom buffering\n");

    return 0;
}

```

Practical Examples

File Merger Utility

```

#include <stdio.h>
#include <stdlib.h>

int merge_files(const char *output, const char **inputs, int count) {
    FILE *out_fp, *in_fp;

    // Open output file
    out_fp = fopen(output, "wb");
    if (out_fp == NULL) {
        perror("Error opening output file");
        return -1;
    }

    // Process each input file
    for (int i = 0; i < count; i++) {

```

```

        in_fp = fopen(inputs[i], "rb");
        if (in_fp == NULL) {
            perror("Error opening input file");
            fclose(out_fp);
            return -1;
        }

        // Copy data
        int ch;
        while ((ch = fgetc(in_fp)) != EOF) {
            fputc(ch, out_fp);
        }

        fclose(in_fp);
    }

    fclose(out_fp);
    return 0;
}

int main() {
    const char *input_files[] = {"file1.txt", "file2.txt", "file3.txt"};

    // Create test files
    FILE *fp;
    fp = fopen("file1.txt", "w");
    fprintf(fp, "Content of file 1\n");
    fclose(fp);

    fp = fopen("file2.txt", "w");
    fprintf(fp, "Content of file 2\n");
    fclose(fp);

    fp = fopen("file3.txt", "w");
    fprintf(fp, "Content of file 3\n");
    fclose(fp);

    // Merge files
    if (merge_files("merged.txt", input_files, 3) == 0) {
        printf("Files merged successfully\n");
    } else {
        printf("Error merging files\n");
    }

    return 0;
}

```

File Splitter Utility

```
#include <stdio.h>
#include <stdlib.h>
#include <string.h>

int split_file(const char *input, long chunk_size) {
    FILE *in_fp, *out_fp;
    char out_name[256];
    int ch, chunk_num = 0;
    long bytes_written = 0;

    in_fp = fopen(input, "rb");
    if (in_fp == NULL) {
        perror("Error opening input file");
        return -1;
    }

    while ((ch = fgetc(in_fp)) != EOF) {
        // Open new output file if needed
        if (bytes_written == 0) {
            snprintf(out_name, sizeof(out_name), "%s.part%d", input, chunk_num);
            out_fp = fopen(out_name, "wb");
            if (out_fp == NULL) {
                perror("Error opening output file");
                fclose(in_fp);
                return -1;
            }
            chunk_num++;
        }

        // Write byte to current output file
        fputc(ch, out_fp);
        bytes_written++;

        // Close current output file if chunk is complete
        if (bytes_written >= chunk_size) {
            fclose(out_fp);
            bytes_written = 0;
        }
    }

    // Close last output file if open
    if (bytes_written > 0) {
        fclose(out_fp);
    }
}
```



```

    fclose(in_fp);
    printf("File split into %d parts\n", chunk_num);
    return 0;
}

int main() {
    // Create test file
    FILE *fp = fopen("large_file.txt", "w");
    for (int i = 0; i < 1000; i++) {
        fprintf(fp, "This is line %d of the large test file.\n", i);
    }
    fclose(fp);

    // Split file into 1KB chunks
    if (split_file("large_file.txt", 1024) == 0) {
        printf("File split successfully\n");
    } else {
        printf("Error splitting file\n");
    }

    return 0;
}

```

Summary

Advanced file operations in C provide powerful capabilities for sophisticated file handling:

1. **File Positioning:** `fseek()`, `ftell()`, `rewind()` for random access
2. **Random Access Files:** Direct access to specific file locations
3. **Temporary Files:** `tmpfile()`, `tmpnam()` for temporary data storage
4. **File Locking:** Synchronization for concurrent file access
5. **File Information:** Status and metadata retrieval
6. **Buffering Control:** Performance optimization through custom buffering

These advanced techniques are essential for building robust file processing applications, database systems, and other programs requiring fine-grained control over file operations.

System-Level Programming

Introduction

System-level programming in C involves interacting directly with the operating system to perform tasks such as processing command-line arguments, accessing environment variables, managing process execution, and handling system signals. These capabilities allow C programs to integrate seamlessly with the operating system environment and provide more sophisticated functionality than basic file I/O operations.

Understanding system-level programming concepts is essential for developing robust applications that can adapt to different environments, handle errors gracefully, and interact with system resources effectively.

Environment: Linux/POSIX (macOS mostly fine). Compile with:

```
gcc -std=c17 -Wall -Wextra -o program program.c
```

Some examples use POSIX APIs (`sigaction`, `fork`, `unistd.h`).

Command-Line Arguments

Command-line arguments allow programs to receive input directly from the command line when they are executed. This is a fundamental way for programs to be configurable and flexible.

`argc` and `argv`

The `main()` function can accept two parameters for processing command-line arguments:

```
int main(int argc, char *argv[]);
```

Parameters: - `argc`: Argument count (number of command-line arguments) - `argv`: Argument vector (array of string pointers to the arguments)

Examples

```
#include <stdio.h>
#include <stdlib.h>
#include <string.h>

int main(int argc, char *argv[]) {
    printf("Program name: %s\n", argv[0]);
    printf("Number of arguments: %d\n", argc - 1);

    // Print all arguments
    for (int i = 1; i < argc; i++) {
        printf("Argument %d: %s\n", i, argv[i]);
    }

    return 0;
}
```

Practical Command-Line Processor

```
#include <stdio.h>
#include <stdlib.h>
#include <string.h>

void print_usage(const char *program_name) {
    printf("Usage: %s [options] <input_file> <output_file>\n", program_name);
    printf("Options:\n");
    printf("  -v, --verbose    Verbose output\n");
    printf("  -h, --help       Show this help message\n");
    printf("  -c, --copy       Copy mode (default)\n");
    printf("  -m, --move       Move mode\n");
}

int main(int argc, char *argv[]) {
    int verbose = 0;
    int move_mode = 0;
    const char *input_file = NULL;
    const char *output_file = NULL;

    // Process command-line arguments
    for (int i = 1; i < argc; i++) {
        if (strcmp(argv[i], "-v") == 0 || strcmp(argv[i], "--verbose") == 0) {
            verbose = 1;
        } else if (strcmp(argv[i], "-h") == 0 || strcmp(argv[i], "--help") == 0) {
            print_usage(argv[0]);
        }
    }
}
```

```

        exit(EXIT_SUCCESS);
    } else if (strcmp(argv[i], "-m") == 0 || strcmp(argv[i], "--move") == 0) {
        move_mode = 1;
    } else if (strcmp(argv[i], "-c") == 0 || strcmp(argv[i], "--copy") == 0) {
        move_mode = 0;
    } else if (input_file == NULL) {
        input_file = argv[i];
    } else if (output_file == NULL) {
        output_file = argv[i];
    } else {
        fprintf(stderr, "Error: Too many arguments\n");
        print_usage(argv[0]);
        exit(EXIT_FAILURE);
    }
}

// Validate required arguments
if (input_file == NULL || output_file == NULL) {
    fprintf(stderr, "Error: Input and output files required\n");
    print_usage(argv[0]);
    exit(EXIT_FAILURE);
}

if (verbose) {
    printf("Mode: %s\n", move_mode ? "Move" : "Copy");
    printf("Input file: %s\n", input_file);
    printf("Output file: %s\n", output_file);
}

// Perform the operation (simplified)
if (move_mode) {
    if (verbose) printf("Moving file...\n");
    // In a real implementation, you would use rename()
} else {
    if (verbose) printf("Copying file...\n");
    // In a real implementation, you would copy the file contents
}

printf("Operation completed successfully\n");
return 0;
}

```

Environment Variables

Environment variables provide a way for programs to access system configuration information and user preferences. They are key-value pairs that are inherited from the parent process and can be

accessed by C programs.

getenv()

Retrieves the value of an environment variable:

```
char *getenv(const char *name);
```

Return Value: - Pointer to the value string on success - NULL if the variable is not found

Examples

```
#include <stdio.h>
#include <stdlib.h>

int main() {
    char *home_dir, *path, *user, *shell;

    // Get common environment variables
    home_dir = getenv("HOME");
    path = getenv("PATH");
    user = getenv("USER");
    shell = getenv("SHELL");

    printf("Home directory: %s\n", home_dir ? home_dir : "Not set");
    printf("User: %s\n", user ? user : "Not set");
    printf("Shell: %s\n", shell ? shell : "Not set");
    printf("Path: %s\n", path ? path : "Not set");

    // Check for custom environment variable
    char *debug = getenv("DEBUG");
    if (debug != NULL && strcmp(debug, "1") == 0) {
        printf("Debug mode enabled\n");
    }

    return 0;
}
```

Setting Environment Variables

```
#include <stdio.h>
#include <stdlib.h>

int main() {
```

```

// Set an environment variable
if (setenv("MY_VARIABLE", "Hello, World!", 1) != 0) {
    perror("Error setting environment variable");
    exit(EXIT_FAILURE);
}

// Retrieve and display the variable
char *value = getenv("MY_VARIABLE");
printf("MY_VARIABLE = %s\n", value ? value : "Not set");

// Unset the environment variable
unsetenv("MY_VARIABLE");

// Try to retrieve it again
value = getenv("MY_VARIABLE");
printf("MY_VARIABLE after unset = %s\n", value ? value : "Not set");

return 0;
}

```

Process Management

Process management functions allow programs to control their execution flow, terminate gracefully, and manage exit status codes.

exit()

Terminates the program with a specified exit status:

```
void exit(int status);
```

atexit()

Registers functions to be called when the program exits:

```
int atexit(void (*func)(void));
```

Examples

```
#include <stdio.h>
#include <stdlib.h>

void cleanup_handler1(void) {
    printf("Cleanup handler 1 called\n");
}

void cleanup_handler2(void) {
    printf("Cleanup handler 2 called\n");
}

void cleanup_handler3(void) {
    printf("Cleanup handler 3 called\n");
}

int main() {
    // Register cleanup handlers
    if (atexit(cleanup_handler1) != 0) {
        fprintf(stderr, "Failed to register cleanup handler 1\n");
        exit(EXIT_FAILURE);
    }

    if (atexit(cleanup_handler2) != 0) {
        fprintf(stderr, "Failed to register cleanup handler 2\n");
        exit(EXIT_FAILURE);
    }

    if (atexit(cleanup_handler3) != 0) {
        fprintf(stderr, "Failed to register cleanup handler 3\n");
        exit(EXIT_FAILURE);
    }

    printf("Main function executing\n");

    // Exit with success status
    exit(EXIT_SUCCESS);

    // This code will never be reached
    printf("This won't be printed\n");
}
```

Process Exit Status

```
#include <stdio.h>
#include <stdlib.h>

int main(int argc, char *argv[]) {
    if (argc != 2) {
        fprintf(stderr, "Usage: %s <number>\n", argv[0]);
        exit(EXIT_FAILURE); // Exit with failure status
    }

    int number = atoi(argv[1]);

    if (number < 0) {
        fprintf(stderr, "Error: Negative numbers not allowed\n");
        exit(EXIT_FAILURE); // Exit with failure status
    }

    if (number == 0) {
        printf("Zero is neither positive nor negative\n");
        exit(EXIT_SUCCESS); // Exit with success status
    }

    printf("Number %d is positive\n", number);

    // Normal exit (equivalent to exit(EXIT_SUCCESS))
    return 0;
}
```

Signal Handling

Signals are software interrupts that allow the operating system to notify a process of various events. Signal handling enables programs to respond gracefully to events like user interrupts, termination requests, and other system events.

signal()

Sets up a signal handler:

```
#include <signal.h>
void (*signal(int sig, void (*handler)(int)))(int);
```


Common Signals

Signal	Description	Default Action
SIGINT	Interrupt (Ctrl+C)	Terminate process
SIGTERM	Termination request	Terminate process
SIGKILL	Kill signal	Terminate process (cannot be caught)
SIGSEGV	Segmentation fault	Terminate process
SIGUSR1	User-defined signal 1	Terminate process
SIGUSR2	User-defined signal 2	Terminate process

Examples

```
#include <stdio.h>
#include <stdlib.h>
#include <signal.h>
#include <unistd.h>

volatile sig_atomic_t signal_received = 0;

void signal_handler(int sig) {
    printf("\nReceived signal %d\n", sig);
    signal_received = 1;
}

int main() {
    // Set up signal handlers
    if (signal(SIGINT, signal_handler) == SIG_ERR) {
        perror("Error setting SIGINT handler");
        exit(EXIT_FAILURE);
    }

    if (signal(SIGTERM, signal_handler) == SIG_ERR) {
        perror("Error setting SIGTERM handler");
        exit(EXIT_FAILURE);
    }

    printf("Program running. Press Ctrl+C to interrupt.\n");

    // Main loop
    while (!signal_received) {
        printf("Working...\n");
        sleep(2);
    }

    printf("Cleaning up and exiting gracefully...\n");
}
```

```

        // Perform cleanup operations here
        // Close files, free memory, etc.

        exit(EXIT_SUCCESS);
    }

```

Advanced Signal Handling

```

#include <stdio.h>
#include <stdlib.h>
#include <signal.h>
#include <unistd.h>
#include <string.h>

volatile sig_atomic_t sigint_count = 0;
volatile sig_atomic_t sigterm_received = 0;

void sigint_handler(int sig) {
    sigint_count++;
    printf("\nSIGINT received (%d times)\n", sigint_count);

    if (sigint_count >= 3) {
        printf("Three interrupts received. Exiting...\n");
        exit(EXIT_SUCCESS);
    }
}

void sigterm_handler(int sig) {
    printf("\nSIGTERM received. Cleaning up...\n");
    sigterm_received = 1;
}

void sigsegv_handler(int sig) {
    printf("\nSegmentation fault detected! Saving emergency data...\n");
    // Save critical data before terminating
    exit(EXIT_FAILURE);
}

int main() {
    // Set up signal handlers
    struct sigaction sa_int, sa_term, sa_segv;

    // SIGINT handler
    memset(&sa_int, 0, sizeof(sa_int));
    sa_int.sa_handler = sigint_handler;

```

```

sigemptyset(&sa_int.sa_mask);
sa_int.sa_flags = 0;
sigaction(SIGINT, &sa_int, NULL);

// SIGTERM handler
memset(&sa_term, 0, sizeof(sa_term));
sa_term.sa_handler = sigterm_handler;
sigemptyset(&sa_term.sa_mask);
sa_term.sa_flags = 0;
sigaction(SIGTERM, &sa_term, NULL);

// SIGSEGV handler
memset(&sa_segv, 0, sizeof(sa_segv));
sa_segv.sa_handler = sigsegv_handler;
sigemptyset(&sa_segv.sa_mask);
sa_segv.sa_flags = SA_RESETHAND; // Only handle once
sigaction(SIGSEGV, &sa_segv, NULL);

printf("Advanced signal handling demo\n");
printf("Press Ctrl+C up to 3 times, or send SIGTERM to terminate\n");

// Main loop
while (!sigterm_received) {
    printf("Working... (PID: %d)\n", getpid());
    sleep(3);
}

printf("Program terminating normally\n");
return 0;
}

```

System Calls and Process Creation

System calls provide direct access to operating system services. While C's standard library provides higher-level interfaces, system calls offer more control and are sometimes necessary for specific tasks.

system()

Executes a command through the shell:

```
int system(const char *command);
```

Examples

```
#include <stdio.h>
#include <stdlib.h>

int main() {
    int result;

    // Execute system commands
    printf("Listing current directory:\n");
    result = system("ls -la");
    if (result == -1) {
        perror("Error executing command");
    }

    printf("\nCurrent date and time:\n");
    result = system("date");
    if (result == -1) {
        perror("Error executing command");
    }

    // Conditional execution based on system command result
    printf("\nChecking if 'git' is available:\n");
    result = system("which git > /dev/null 2>&1");
    if (result == 0) {
        printf("Git is available\n");
        system("git --version");
    } else {
        printf("Git is not available\n");
    }

    return 0;
}
```

Practical Examples

Configuration File Processor

```
#include <stdio.h>
#include <stdlib.h>
#include <string.h>

typedef struct {
    char database_host[100];
```

```

    int database_port;
    char log_level[20];
    int max_connections;
} Config;

int load_config(const char *filename, Config *config) {
    FILE *fp;
    char line[256];
    char key[100], value[100];

    // Set default values
    strcpy(config->database_host, "localhost");
    config->database_port = 5432;
    strcpy(config->log_level, "INFO");
    config->max_connections = 100;

    fp = fopen(filename, "r");
    if (fp == NULL) {
        // Try to load from environment variables
        char *env_host = getenv("DB_HOST");
        char *env_port = getenv("DB_PORT");
        char *env_log = getenv("LOG_LEVEL");
        char *env_max = getenv("MAX_CONN");

        if (env_host) strcpy(config->database_host, env_host);
        if (env_port) config->database_port = atoi(env_port);
        if (env_log) strcpy(config->log_level, env_log);
        if (env_max) config->max_connections = atoi(env_max);

        return 0; // Success with environment variables
    }

    // Parse configuration file
    while (fgets(line, sizeof(line), fp) != NULL) {
        // Skip comments and empty lines
        if (line[0] == '#' || line[0] == '\n') continue;

        // Parse key=value pairs
        if (sscanf(line, "%[^=]=%s", key, value) == 2) {
            if (strcmp(key, "database_host") == 0) {
                strcpy(config->database_host, value);
            } else if (strcmp(key, "database_port") == 0) {
                config->database_port = atoi(value);
            } else if (strcmp(key, "log_level") == 0) {
                strcpy(config->log_level, value);
            } else if (strcmp(key, "max_connections") == 0) {
                config->max_connections = atoi(value);
            }
        }
    }
}

```

```

        }
    }
}

fclose(fp);
return 0;
}

void print_config(const Config *config) {
    printf("Configuration:\n");
    printf("  Database Host: %s\n", config->database_host);
    printf("  Database Port: %d\n", config->database_port);
    printf("  Log Level: %s\n", config->log_level);
    printf("  Max Connections: %d\n", config->max_connections);
}

int main(int argc, char *argv[]) {
    Config config;
    const char *config_file = "app.conf";

    // Check for command-line configuration file
    if (argc > 1) {
        config_file = argv[1];
    }

    // Load configuration
    if (load_config(config_file, &config) != 0) {
        fprintf(stderr, "Error loading configuration\n");
        exit(EXIT_FAILURE);
    }

    // Print configuration
    print_config(&config);

    // Use configuration in application
    printf("\nApplication starting with above configuration...\n");

    return 0;
}

```

Process Monitor

```
#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <signal.h>
#include <unistd.h>
#include <sys/wait.h>

volatile sig_atomic_t terminate = 0;

void signal_handler(int sig) {
    printf("\nTermination signal received\n");
    terminate = 1;
}

int main(int argc, char *argv[]) {
    pid_t pid;
    int status;

    if (argc != 2) {
        fprintf(stderr, "Usage: %s <command>\n", argv[0]);
        exit(EXIT_FAILURE);
    }

    // Set up signal handlers
    signal(SIGINT, signal_handler);
    signal(SIGTERM, signal_handler);

    printf("Process monitor started\n");
    printf("Monitoring command: %s\n", argv[1]);

    while (!terminate) {
        // Fork a child process
        pid = fork();

        if (pid == -1) {
            perror("Error forking process");
            exit(EXIT_FAILURE);
        }

        if (pid == 0) {
            // Child process - execute the command
            execl("/bin/sh", "sh", "-c", argv[1], (char *)NULL);
            perror("Error executing command");
            exit(EXIT_FAILURE);
        } else {
            // Parent process - wait for child to finish
            waitpid(pid, &status, 0);
        }
    }
}
```

```

        // Parent process - wait for child
        printf("Started process with PID %d\n", pid);

        // Wait for child to complete
        waitpid(pid, &status, 0);

        if (WIFEXITED(status)) {
            printf("Process exited with status %d\n", WEXITSTATUS(status));
        } else if (WIFSIGNALED(status)) {
            printf("Process terminated by signal %d\n", WTERMSIG(status));
        }

        // Wait before restarting (unless terminating)
        if (!terminate) {
            printf("Restarting in 5 seconds...\n");
            sleep(5);
        }
    }

    printf("Process monitor terminated\n");
    return 0;
}

```

Full Programs: CLI, Env, Signals, Processes

Program 1 — Robust argument parser (opts.c)

```

/* file: opts.c */
#include <stdio.h>
#include <stdlib.h>
#include <string.h>

struct Options {
    const char *input;
    const char *output;
    int verbose;
    int count;
};

static void usage(const char *prog) {
    fprintf(stderr,
        "Usage: %s [-v] [-n N] -i input -o output\n"
        "    -v          verbose\n"
    );
}

```



```

        "  -n N      count (default 1)\n"
        "  -i path   input file\n"
        "  -o path   output file\n",
        prog);
}

int main(int argc, char *argv[]) {
    struct Options opt = {NULL, NULL, 0, 1};
    int i;

    for (i = 1; i < argc; i++) {
        if (strcmp(argv[i], "-v") == 0) {
            opt.verbose = 1;
        } else if (strcmp(argv[i], "-n") == 0) {
            if (i + 1 >= argc) {
                usage(argv[0]);
                return EXIT_FAILURE;
            }
            opt.count = atoi(argv[++i]);
            if (opt.count <= 0) {
                fprintf(stderr, "-n must be positive\n");
                return EXIT_FAILURE;
            }
        } else if (strcmp(argv[i], "-i") == 0) {
            if (i + 1 >= argc) {
                usage(argv[0]);
                return EXIT_FAILURE;
            }
            opt.input = argv[++i];
        } else if (strcmp(argv[i], "-o") == 0) {
            if (i + 1 >= argc) {
                usage(argv[0]);
                return EXIT_FAILURE;
            }
            opt.output = argv[++i];
        } else if (strcmp(argv[i], "-h") == 0 || strcmp(argv[i], "--help") == 0) {
            usage(argv[0]);
            return EXIT_SUCCESS;
        } else {
            fprintf(stderr, "unknown option: %s\n", argv[i]);
            usage(argv[0]);
            return EXIT_FAILURE;
        }
    }

    if (opt.input == NULL || opt.output == NULL) {

```

```

        usage(argv[0]);
        return EXIT_FAILURE;
    }

    if (opt.verbose) {
        fprintf(stderr, "in=%s out=%s count=%d\n",
            opt.input, opt.output, opt.count);
    }
    printf("would process %s -> %s (%d times)\n",
        opt.input, opt.output, opt.count);
    return EXIT_SUCCESS;
}

gcc -std=c17 -Wall -Wextra -o opts opts.c
./opts -v -n 3 -i a.txt -o b.txt
./opts -h
./opts -i a.txt           # should fail with usage

```

Program 2 — Environment-driven config (envcfg.c)

```

/* file: envcfg.c */
#include <stdio.h>
#include <stdlib.h>
#include <string.h>

struct Config {
    const char *home;
    const char *log_level;
    int max_items;
};

static int parse_positive_int(const char *s, int *out) {
    char *end;
    long v;

    if (s == NULL || out == NULL) {
        return -1;
    }
    v = strtol(s, &end, 10);
    if (end == s || *end != '\0' || v <= 0 || v > 1000000L) {
        return -1;
    }
    *out = (int)v;
    return 0;
}

```

```

int main(void) {
    struct Config cfg;
    const char *max_s;

    cfg.home = getenv("HOME");
    if (cfg.home == NULL) {
        cfg.home = ".";
    }

    cfg.log_level = getenv("APP_LOG_LEVEL");
    if (cfg.log_level == NULL) {
        cfg.log_level = "INFO";
    }

    max_s = getenv("APP_MAX_ITEMS");
    if (max_s == NULL) {
        cfg.max_items = 100;
    } else if (parse_positive_int(max_s, &cfg.max_items) != 0) {
        fprintf(stderr, "invalid APP_MAX_ITEMS=%s\n", max_s);
        return EXIT_FAILURE;
    }

    printf("home=%s\n", cfg.home);
    printf("log_level=%s\n", cfg.log_level);
    printf("max_items=%d\n", cfg.max_items);
    return EXIT_SUCCESS;
}

```

```

gcc -std=c17 -Wall -Wextra -o envcfg envcfg.c
./envcfg
APP_LOG_LEVEL=DEBUG APP_MAX_ITEMS=50 ./envcfg
APP_MAX_ITEMS=nope ./envcfg # should fail

```

Program 3 — atexit cleanup and exit status (cleanup_demo.c)

```

/* file: cleanup_demo.c */
#include <stdio.h>
#include <stdlib.h>
#include <string.h>

static FILE *g_log;

static void close_log(void) {
    if (g_log != NULL) {

```

```

        fprintf(g_log, "shutdown\n");
        fclose(g_log);
        g_log = NULL;
        fprintf(stderr, "atexit: log closed\n");
    }
}

int main(int argc, char *argv[]) {
    int fail = (argc > 1 && strcmp(argv[1], "fail") == 0);

    g_log = fopen("app_runtime.log", "w");
    if (g_log == NULL) {
        perror("fopen");
        return EXIT_FAILURE;
    }
    if (atexit(close_log) != 0) {
        fprintf(stderr, "atexit registration failed\n");
        fclose(g_log);
        return EXIT_FAILURE;
    }

    fprintf(g_log, "started\n");
    printf("working...\n");

    if (fail) {
        fprintf(stderr, "failing with exit code 2\n");
        exit(2); /* still runs atexit handlers */
    }

    return EXIT_SUCCESS;
}

```

```

gcc -std=c17 -Wall -Wextra -o cleanup_demo cleanup_demo.c
./cleanup_demo
cat app_runtime.log
./cleanup_demo fail; echo exit=$?

```

Note: `#include <string.h>` is required for `strcmp` — add it if your compiler rejects the program under pedantic modes.

Program 4 — Graceful SIGINT shutdown (sigint_loop.c)

```

/* file: sigint_loop.c */
#include <stdio.h>
#include <stdlib.h>
#include <signal.h>

```

```

#include <unistd.h>
#include <string.h>

static volatile sig_atomic_t g_stop = 0;

static void on_sigint(int signo) {
    (void)signo;
    g_stop = 1;
}

int main(void) {
    struct sigaction sa;
    unsigned long ticks = 0;

    memset(&sa, 0, sizeof sa);
    sa.sa_handler = on_sigint;
    sigemptyset(&sa.sa_mask);
    sa.sa_flags = 0; /* no SA_RESTART: sleep may return early */

    if (sigaction(SIGINT, &sa, NULL) != 0) {
        perror("sigaction");
        return EXIT_FAILURE;
    }

    printf("pid=%d - press Ctrl+C to stop\n", (int)getpid());
    while (!g_stop) {
        printf("tick %lu\n", ticks++);
        sleep(1);
    }
    printf("clean shutdown after %lu ticks\n", ticks);
    return EXIT_SUCCESS;
}

```

```

gcc -std=c17 -Wall -Wextra -o sigint_loop sigint_loop.c
./sigint_loop
# Ctrl+C → clean message

```

Program 5 — fork + exec + wait mini runner (run_cmd.c)

```

/* file: run_cmd.c */
#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <unistd.h>
#include <sys/wait.h>

```

```

#include <errno.h>

int main(int argc, char *argv[]) {
    pid_t pid;
    int status;

    if (argc < 2) {
        fprintf(stderr, "Usage: %s <command> [args...]\n", argv[0]);
        return EXIT_FAILURE;
    }

    pid = fork();
    if (pid < 0) {
        perror("fork");
        return EXIT_FAILURE;
    }

    if (pid == 0) {
        /* child */
        execvp(argv[1], &argv[1]);
        fprintf(stderr, "execvp %s: %s\n", argv[1], strerror(errno));
        _exit(127);
    }

    /* parent */
    if (waitpid(pid, &status, 0) < 0) {
        perror("waitpid");
        return EXIT_FAILURE;
    }

    if (WIFEXITED(status)) {
        printf("child exited %d\n", WEXITSTATUS(status));
        return WEXITSTATUS(status);
    }

    if (WIFSIGNALED(status)) {
        printf("child killed by signal %d\n", WTERMSIG(status));
        return 128 + WTERMSIG(status);
    }

    return EXIT_FAILURE;
}

```

```

gcc -std=c17 -Wall -Wextra -o run_cmd run_cmd.c
./run_cmd echo hello from child
./run_cmd false; echo status=$?
./run_cmd /no/such/cmd; echo status=$?

```

Program 6 — Echo lines with PID and env prefix (line_echo.c)

```
/* file: line_echo.c */
#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <unistd.h>

int main(void) {
    char line[256];
    const char *prefix = getenv("LINE_PREFIX");
    pid_t pid = getpid();

    if (prefix == NULL) {
        prefix = "";
    }

    while (fgets(line, sizeof line, stdin) != NULL) {
        /* strip newline for formatting */
        line[strcspn(line, "\n")] = '\0';
        printf("%s[pid=%d] %s\n", prefix, (int)pid, line);
        fflush(stdout);
    }
    if (ferror(stdin)) {
        perror("stdin");
        return EXIT_FAILURE;
    }
    return EXIT_SUCCESS;
}

gcc -std=c17 -Wall -Wextra -o line_echo line_echo.c
printf 'a\nb\n' | LINE_PREFIX='>>' ./line_echo
```

Exercises

```
gcc -std=c17 -Wall -Wextra -o exN exN.c
```

1. **opts long form** — Accept `--input=path` and `--verbose` in addition to short flags.
2. **Required env** — Fail startup if `APP_TOKEN` is unset or empty; print a clear error to `stderr`.
3. **Double atexit** — Register two `atexit` handlers; show they run in reverse order of registration.

4. **SIGTERM too** — Extend `sigint_loop` to stop on `SIGTERM` as well as `SIGINT` (one handler, two `sigaction` calls).
 5. **Pipeline** — Use `run_cmd` to run `ls` and then `true/false`; document exit codes in a comment.
 6. **Avoid `system()`** — Replace a `system("ls -l")` sketch with `fork/execvp/waitpid` for fixed `argv {"ls", "-l", NULL}`.
 7. **Argv0 banner** — Print only the basename of `argv[0]` (after last `/`) as the program name in usage errors.
 8. **Timeout sketch** — After `fork`, parent uses `alarm(2)` or a timed wait strategy (research `sigtimedwait` or poll loop) to kill a hanging child — advanced stretch.
-

Summary

System-level programming in C provides powerful capabilities for interacting with the operating system:

1. **Command-Line Arguments:** Processing `argc` and `argv` for flexible program configuration
2. **Environment Variables:** Accessing system configuration through `getenv()`, `setenv()`, and `unsetenv()`
3. **Process Management:** Controlling program termination with `exit()` and cleanup with `atexit()`
4. **Signal Handling:** Responding to system events and interrupts gracefully (`sig_atomic_t`, `sigaction`)
5. **Process control:** `fork`, `exec*`, `waitpid` for running and supervising programs
6. **System Calls:** Executing shell commands carefully (`system` is convenient but shells inject risk)

These system-level programming concepts are essential for developing robust, production-ready C applications that can adapt to different environments and handle various runtime conditions effectively.

Module 8: File I/O and System Programming

Overview

This module covers file input/output operations and system-level programming in C. You'll learn how to work with files, handle different file modes, perform various I/O operations, and understand system programming concepts. The module also explores advanced file operations, command-line arguments, environment variables, and basic system calls.

Learning Objectives

By the end of this module, you will be able to: - Open, read, write, and close files using standard library functions - Understand different file modes and their appropriate usage - Perform character, line, formatted, and binary I/O operations - Manipulate file positions using seeking functions - Handle file errors and implement proper error checking - Work with command-line arguments and environment variables - Understand process creation and basic system calls - Implement file locking and concurrent access mechanisms - Apply best practices for file handling and system programming

Chapters

1. **File Operations** - File pointers, opening and closing files, text vs. binary modes, error handling
2. **File I/O Functions** - Character I/O, line I/O, formatted I/O, binary I/O functions
3. **Advanced File Operations** - File positioning, random access files, temporary files, file locking
4. **System-Level Programming** - Command-line arguments, environment variables, exit status, signal handling basics

Key Concepts Covered

- File pointers and the `FILE` structure
- File opening modes (`r`, `w`, `a`, `r+`, `w+`, `a+`, binary variants)
- Text vs. binary file operations
- Character I/O functions (`fgetc()`, `fputc()`, `ungetc()`)
- Line I/O functions (`fgets()`, `fputs()`)
- Formatted I/O functions (`fprintf()`, `fscanf()`)

- Binary I/O functions (`fread()`, `fwrite()`)
- File positioning functions (`fseek()`, `ftell()`, `rewind()`)
- Random access file operations
- Temporary file creation and management
- File locking mechanisms
- Command-line argument processing (`argc`, `argv`)
- Environment variable access (`getenv()`)
- Process exit handling (`exit()`, `atexit()`)
- Signal handling basics
- Error handling in file operations
- Buffering strategies and flushing

Prerequisites

- Completion of Module 1: Foundations and Environment Setup
- Completion of Module 2: Data Types and Variables
- Completion of Module 3: Control Flow
- Completion of Module 4: Functions and Modular Programming
- Completion of Module 5: Arrays and Strings
- Completion of Module 6: Pointers and Memory Management
- Completion of Module 7: Structures and User-Defined Types
- Understanding of basic programming concepts
- Familiarity with variable declaration and initialization

Tools and Technologies

- C compiler (GCC, Clang, or MSVC)
- Text editor or IDE
- File system exploration tools
- Understanding of basic operating system concepts
- Familiarity with command-line interfaces

Estimated Time to Complete

- Reading: 8-10 hours
- Exercises: 12-15 hours
- Projects: 6-8 hours

Assessment

- Complete all chapter exercises
- Successfully implement file processing programs

- Demonstrate understanding of different file I/O techniques
- Pass file I/O and system programming quiz
- Submit a comprehensive file manipulation utility

Next Module

After completing this module, proceed to **Module 9: Advanced C Features and Modern Standards** to learn about modern C language features and standards.

Additional Resources

- “The C Programming Language” by Kernighan and Ritchie (Chapter 7)
- C standard documentation (ISO/IEC 9899:2024 for C23)
- Operating system documentation
- File system documentation
- System programming guides
- POSIX standard documentation

Modern C

C99 Features

Introduction

The C99 standard (ISO/IEC 9899:1999) introduced numerous significant enhancements to the C programming language, making it more powerful, expressive, and suitable for modern programming tasks. These features addressed many limitations of the previous C89/C90 standard and brought C closer to contemporary programming language capabilities.

C99 features include new data types, improved control structures, enhanced initialization syntax, and better support for numerical computing. Understanding these features is essential for writing modern, efficient C code.

Variable-Length Arrays (VLAs)

One of the most significant additions in C99 was Variable-Length Arrays (VLAs), which allow the size of an array to be determined at runtime rather than compile time.

Basic VLA Declaration

```
#include <stdio.h>

int main() {
    int size;

    printf("Enter array size: ");
    scanf("%d", &size);

    // Variable-length array - size determined at runtime
    int array[size];

    // Initialize and use the array
    for (int i = 0; i < size; i++) {
        array[i] = i * i;
        printf("array[%d] = %d\n", i, array[i]);
    }

    return 0;
}
```

Multi-dimensional VLAs

```
#include <stdio.h>

int main() {
    int rows, cols;

    printf("Enter rows and columns: ");
    scanf("%d %d", &rows, &cols);

    // Multi-dimensional variable-length array
    int matrix[rows][cols];

    // Initialize the matrix
    for (int i = 0; i < rows; i++) {
        for (int j = 0; j < cols; j++) {
            matrix[i][j] = i * cols + j;
        }
    }

    // Print the matrix
    for (int i = 0; i < rows; i++) {
        for (int j = 0; j < cols; j++) {
            printf("%4d ", matrix[i][j]);
        }
        printf("\n");
    }

    return 0;
}
```

VLA Function Parameters

```
#include <stdio.h>

// Function that accepts VLA parameters
void process_matrix(int rows, int cols, int matrix[rows][cols]) {
    printf("Processing %dx%d matrix:\n", rows, cols);

    for (int i = 0; i < rows; i++) {
        for (int j = 0; j < cols; j++) {
            matrix[i][j] *= 2; // Double each element
        }
    }
}
```

```

// Function to print a matrix
void print_matrix(int rows, int cols, int matrix[rows][cols]) {
    for (int i = 0; i < rows; i++) {
        for (int j = 0; j < cols; j++) {
            printf("%4d ", matrix[i][j]);
        }
        printf("\n");
    }
}

int main() {
    int rows = 3, cols = 4;

    // Declare and initialize a VLA
    int matrix[rows][cols] = {
        {1, 2, 3, 4},
        {5, 6, 7, 8},
        {9, 10, 11, 12}
    };

    printf("Original matrix:\n");
    print_matrix(rows, cols, matrix);

    // Process the matrix
    process_matrix(rows, cols, matrix);

    printf("\nModified matrix:\n");
    print_matrix(rows, cols, matrix);

    return 0;
}

```

Designated Initializers

C99 introduced designated initializers, which allow you to initialize specific elements of arrays and structures by specifying their indices or member names.

Array Designated Initializers

```

#include <stdio.h>

int main() {
    // Traditional initialization
    int array1[10] = {0, 0, 5, 0, 0, 0, 0, 8, 0, 0};
}

```

```

// Designated initializers - only initialize specific elements
int array2[10] = {[2] = 5, [7] = 8};

// Mixed initialization
int array3[10] = {1, 2, [5] = 10, 11, 12};

printf("Array 1: ");
for (int i = 0; i < 10; i++) {
    printf("%d ", array1[i]);
}
printf("\n");

printf("Array 2: ");
for (int i = 0; i < 10; i++) {
    printf("%d ", array2[i]);
}
printf("\n");

printf("Array 3: ");
for (int i = 0; i < 10; i++) {
    printf("%d ", array3[i]);
}
printf("\n");

return 0;
}

```

Structure Designated Initializers

```

#include <stdio.h>
#include <string.h>

typedef struct {
    int id;
    char name[50];
    double salary;
    int age;
} Employee;

int main() {
    // Traditional initialization (order matters)
    Employee emp1 = {1, "John Doe", 50000.0, 30};

    // Designated initializers (order doesn't matter)
    Employee emp2 = {

```



```

        .id = 2,
        .name = "Jane Smith",
        .salary = 60000.0,
        .age = 28
    };

    // Partial initialization with designated initializers
    Employee emp3 = {
        .id = 3,
        .name = "Bob Johnson"
        // salary and age will be zero-initialized
    };

    printf("Employee 1: ID=%d, Name=%s, Salary=%.2f, Age=%d\n",
        emp1.id, emp1.name, emp1.salary, emp1.age);

    printf("Employee 2: ID=%d, Name=%s, Salary=%.2f, Age=%d\n",
        emp2.id, emp2.name, emp2.salary, emp2.age);

    printf("Employee 3: ID=%d, Name=%s, Salary=%.2f, Age=%d\n",
        emp3.id, emp3.name, emp3.salary, emp3.age);

    return 0;
}

```

Compound Literals

Compound literals allow you to create unnamed objects of any data type, including arrays and structures.

Array Compound Literals

```

#include <stdio.h>
#include <string.h>

void print_array(int size, int arr[size]) {
    for (int i = 0; i < size; i++) {
        printf("%d ", arr[i]);
    }
    printf("\n");
}

int main() {
    // Traditional approach

```

```

    int temp_array[] = {1, 2, 3, 4, 5};
    print_array(5, temp_array);

    // Using compound literal
    print_array(5, (int[]){1, 2, 3, 4, 5});

    // Compound literal with VLA
    int size = 3;
    print_array(size, (int[size]){10, 20, 30});

    return 0;
}

```

Structure Compound Literals

```

#include <stdio.h>
#include <string.h>

typedef struct {
    int x, y;
} Point;

typedef struct {
    Point center;
    double radius;
} Circle;

void print_point(Point p) {
    printf("Point(%d, %d)\n", p.x, p.y);
}

void print_circle(Circle c) {
    printf("Circle(center=(%d, %d), radius=%.2f)\n",
        c.center.x, c.center.y, c.radius);
}

int main() {
    // Traditional approach
    Point p1 = {10, 20};
    print_point(p1);

    // Using compound literal
    print_point((Point){30, 40});

    // Structure with nested structure

```

```

    Circle c1 = {{0, 0}, 5.0};
    print_circle(c1);

    // Using compound literals for nested structures
    print_circle((Circle){{10, 10}, 7.5});

    return 0;
}

```

New Data Types

C99 introduced several new data types to improve numerical computing capabilities.

long long Type

```

#include <stdio.h>
#include <limits.h>

int main() {
    long long big_number = 123456789012345LL;
    unsigned long long huge_number = 1234567890123456789ULL;

    printf("long long value: %lld\n", big_number);
    printf("unsigned long long value: %llu\n", huge_number);

    printf("LLONG_MIN: %lld\n", LLONG_MIN);
    printf("LLONG_MAX: %lld\n", LLONG_MAX);
    printf("ULLONG_MAX: %llu\n", ULLONG_MAX);

    return 0;
}

```

_Bool Type

```

#include <stdio.h>
#include <stdbool.h>

int main() {
    // Using _Bool (C99)
    _Bool flag1 = 1;
    _Bool flag2 = 0;
}

```

```

// Using bool with stdbool.h (recommended)
bool is_valid = true;
bool is_error = false;

printf("_Bool values: %d, %d\n", flag1, flag2);
printf("bool values: %d, %d\n", is_valid, is_error);

// Boolean expressions
int x = 10, y = 20;
bool result = (x < y);
printf("x < y: %s\n", result ? "true" : "false");

return 0;
}

```

Inline Functions

The `inline` keyword suggests to the compiler that a function should be inlined to reduce function call overhead.

```

#include <stdio.h>

// Inline function suggestion
inline int square(int x) {
    return x * x;
}

// Static inline function (more portable)
static inline int cube(int x) {
    return x * x * x;
}

int main() {
    int num = 5;

    printf("Square of %d: %d\n", num, square(num));
    printf("Cube of %d: %d\n", num, cube(num));

    // Inline functions work in loops without function call overhead
    for (int i = 1; i <= 5; i++) {
        printf("Square of %d: %d\n", i, square(i));
    }

    return 0;
}

```

restrict Keyword

The `restrict` keyword is a hint to the compiler that a pointer is the only way to access a particular memory region, allowing for better optimization.

```
#include <stdio.h>

// Without restrict - compiler must be conservative
void add_arrays_normal(int n, int *a, int *b, int *result) {
    for (int i = 0; i < n; i++) {
        result[i] = a[i] + b[i];
    }
}

// With restrict - compiler can optimize more aggressively
void add_arrays_restrict(int n, int *restrict a, int *restrict b, int *restrict result) {
    for (int i = 0; i < n; i++) {
        result[i] = a[i] + b[i];
    }
}

int main() {
    int a[5] = {1, 2, 3, 4, 5};
    int b[5] = {10, 20, 30, 40, 50};
    int result1[5], result2[5];

    add_arrays_normal(5, a, b, result1);
    add_arrays_restrict(5, a, b, result2);

    printf("Normal addition: ");
    for (int i = 0; i < 5; i++) {
        printf("%d ", result1[i]);
    }
    printf("\n");

    printf("Restrict addition: ");
    for (int i = 0; i < 5; i++) {
        printf("%d ", result2[i]);
    }
    printf("\n");

    return 0;
}
```

Flexible Array Members

C99 introduced flexible array members, allowing structures to have arrays of variable size at the end.

```
#include <stdio.h>
#include <stdlib.h>
#include <string.h>

typedef struct {
    int count;
    double data[]; // Flexible array member (must be last)
} DynamicArray;

DynamicArray* create_dynamic_array(int count) {
    // Allocate memory for structure + array
    DynamicArray *arr = malloc(sizeof(DynamicArray) + count * sizeof(double));
    if (arr != NULL) {
        arr->count = count;
    }
    return arr;
}

void free_dynamic_array(DynamicArray *arr) {
    free(arr);
}

int main() {
    int size = 5;
    DynamicArray *arr = create_dynamic_array(size);

    if (arr != NULL) {
        // Initialize the flexible array
        for (int i = 0; i < arr->count; i++) {
            arr->data[i] = i * 1.5;
        }

        // Print the array
        printf("Dynamic array contents: ");
        for (int i = 0; i < arr->count; i++) {
            printf("%.1f ", arr->data[i]);
        }
        printf("\n");

        free_dynamic_array(arr);
    }
}
```

```
    return 0;
}
```

Practical Examples

Matrix Library Using C99 Features

```
#include <stdio.h>
#include <stdlib.h>

// Matrix structure using flexible array member
typedef struct {
    int rows, cols;
    double data[]; // Flexible array member
} Matrix;

// Create matrix
Matrix* create_matrix(int rows, int cols) {
    Matrix *m = malloc(sizeof(Matrix) + rows * cols * sizeof(double));
    if (m != NULL) {
        m->rows = rows;
        m->cols = cols;
    }
    return m;
}

// Free matrix
void free_matrix(Matrix *m) {
    free(m);
}

// Initialize matrix with designated initializers
void init_matrix(Matrix *m, double values[m->rows][m->cols]) {
    for (int i = 0; i < m->rows; i++) {
        for (int j = 0; j < m->cols; j++) {
            m->data[i * m->cols + j] = values[i][j];
        }
    }
}

// Print matrix
void print_matrix(const Matrix *m) {
    for (int i = 0; i < m->rows; i++) {
        for (int j = 0; j < m->cols; j++) {
```

```

        printf("%8.2f ", m->data[i * m->cols + j]);
    }
    printf("\n");
}

}

// Matrix multiplication using VLAs
Matrix* multiply_matrices(const Matrix *a, const Matrix *b) {
    if (a->cols != b->rows) {
        printf("Error: Incompatible matrix dimensions\n");
        return NULL;
    }

    Matrix *result = create_matrix(a->rows, b->cols);
    if (result == NULL) {
        printf("Error: Memory allocation failed\n");
        return NULL;
    }

    // Use VLA for computation
    int rows = a->rows, cols = b->cols, inner = a->cols;
    double a_data[rows][inner], b_data[inner][cols], result_data[rows][cols];

    // Copy data to VLAs for easier access
    for (int i = 0; i < rows; i++) {
        for (int j = 0; j < inner; j++) {
            a_data[i][j] = a->data[i * a->cols + j];
        }
    }

    for (int i = 0; i < inner; i++) {
        for (int j = 0; j < cols; j++) {
            b_data[i][j] = b->data[i * b->cols + j];
        }
    }

    // Perform multiplication
    for (int i = 0; i < rows; i++) {
        for (int j = 0; j < cols; j++) {
            result_data[i][j] = 0;
            for (int k = 0; k < inner; k++) {
                result_data[i][j] += a_data[i][k] * b_data[k][j];
            }
        }
    }
}

```



```

    // Copy result back to matrix
    for (int i = 0; i < rows; i++) {
        for (int j = 0; j < cols; j++) {
            result->data[i * cols + j] = result_data[i][j];
        }
    }

    return result;
}

int main() {
    // Create matrices using compound literals
    Matrix *m1 = create_matrix(2, 3);
    Matrix *m2 = create_matrix(3, 2);

    if (m1 == NULL || m2 == NULL) {
        printf("Error: Memory allocation failed\n");
        return 1;
    }

    // Initialize matrices
    double values1[2][3] = {{1, 2, 3}, {4, 5, 6}};
    double values2[3][2] = {{7, 8}, {9, 10}, {11, 12}};

    init_matrix(m1, values1);
    init_matrix(m2, values2);

    printf("Matrix 1:\n");
    print_matrix(m1);

    printf("\nMatrix 2:\n");
    print_matrix(m2);

    // Multiply matrices
    Matrix *result = multiply_matrices(m1, m2);
    if (result != NULL) {
        printf("\nResult of multiplication:\n");
        print_matrix(result);
        free_matrix(result);
    }

    free_matrix(m1);
    free_matrix(m2);

    return 0;
}

```

Summary

C99 introduced several powerful features that significantly enhanced the C programming language:

1. **Variable-Length Arrays (VLAs):** Runtime-determined array sizes
2. **Designated Initializers:** Initialize specific array elements or structure members
3. **Compound Literals:** Create unnamed objects of any type
4. **New Data Types:** `long long`, `_Bool` for better numerical computing
5. **Inline Functions:** Suggest function inlining for performance
6. **restrict Keyword:** Hint for compiler optimization
7. **Flexible Array Members:** Variable-size arrays at structure end

These features make C99 code more expressive, efficient, and easier to write and maintain. Understanding and utilizing these features is essential for modern C programming.

C11 Features

Introduction

The C11 standard (ISO/IEC 9899:2011) introduced several important enhancements to the C programming language, focusing on improved type safety, concurrency support, and alignment with modern programming practices. These features addressed many of the limitations and safety concerns that had been identified in earlier versions of the language.

C11 features include static assertions for compile-time checking, anonymous structures and unions for more flexible data organization, thread-local storage for multi-threaded applications, and bounds-checking functions for improved security.

Static Assertions

Static assertions allow you to perform compile-time checks that can help catch errors early in the development process. This is particularly useful for validating assumptions about data types, structure sizes, or other compile-time constants.

`_Static_assert`

```
#include <stdio.h>
#include <stdint.h>

// Static assertion to ensure int is 32 bits
_Static_assert(sizeof(int) == 4, "int must be 32 bits");

// Static assertion to ensure pointer alignment
_Static_assert(sizeof(void*) >= sizeof(int), "Pointer must be at least as large as int");

typedef struct {
    uint8_t id;
    uint32_t timestamp;
    char data[256];
} Packet;

// Static assertion to ensure structure packing
_Static_assert(sizeof(Packet) == 261, "Packet structure size must be exactly 261 bytes");
```

```

int main() {
    printf("All static assertions passed!\n");
    printf("Size of int: %zu bytes\n", sizeof(int));
    printf("Size of Packet: %zu bytes\n", sizeof(Packet));

    return 0;
}

```

Practical Static Assertion Examples

```

#include <stdio.h>
#include <stdint.h>
#include <limits.h>

// Ensure char is 8 bits
_Static_assert(CHAR_BIT == 8, "char must be 8 bits");

// Ensure two's complement representation
_Static_assert(-1 == ~0, "System must use two's complement representation");

// Ensure floating-point format
_Static_assert(sizeof(float) == 4, "float must be 32 bits");
_Static_assert(sizeof(double) == 8, "double must be 64 bits");

// Custom assertion macro for better error messages
#define STATIC_ASSERT(condition, message) _Static_assert(condition, message)

// Assert that a structure has no padding
typedef struct {
    uint32_t a;
    uint16_t b;
    uint8_t c;
    uint8_t d;
} CompactStruct;

STATIC_ASSERT(sizeof(CompactStruct) == 8, "CompactStruct should be exactly 8 bytes");

int main() {
    printf("System characteristics validated at compile time\n");
    printf("Size of CompactStruct: %zu bytes\n", sizeof(CompactStruct));

    return 0;
}

```

Anonymous Structures and Unions

C11 introduced anonymous structures and unions, which allow you to embed structures and unions directly within other structures without requiring a member name.

Anonymous Structures

```
#include <stdio.h>

typedef struct {
    int x, y;
} Point;

typedef struct {
    struct { // Anonymous structure
        int width, height;
    };
    Point position;
} Rectangle;

int main() {
    Rectangle rect = {{100, 50}, {10, 20}};

    // Access anonymous structure members directly
    printf("Width: %d, Height: %d\n", rect.width, rect.height);
    printf("Position: (%d, %d)\n", rect.position.x, rect.position.y);

    // Modify anonymous structure members
    rect.width = 200;
    rect.height = 100;

    printf("Modified Width: %d, Height: %d\n", rect.width, rect.height);

    return 0;
}
```

Anonymous Unions

```
#include <stdio.h>

typedef struct {
    enum { INTEGER, FLOAT, STRING } type;
    union { // Anonymous union
        int i;
```

```

        float f;
        char *s;
    };
} Value;

void print_value(const Value *v) {
    switch (v->type) {
        case INTEGER:
            printf("Integer: %d\n", v->i);
            break;
        case FLOAT:
            printf("Float: %.2f\n", v->f);
            break;
        case STRING:
            printf("String: %s\n", v->s);
            break;
    }
}

int main() {
    Value values[3];

    // Initialize different types of values
    values[0].type = INTEGER;
    values[0].i = 42;

    values[1].type = FLOAT;
    values[1].f = 3.14f;

    values[2].type = STRING;
    values[2].s = "Hello, World!";

    // Print all values
    for (int i = 0; i < 3; i++) {
        print_value(&values[i]);
    }

    return 0;
}

```

Thread-Local Storage

C11 introduced the `_Thread_local` storage class specifier, which allows variables to have separate instances for each thread in a multi-threaded program.

Basic Thread-Local Storage

```
#include <stdio.h>
#include <threads.h>
#include <stdlib.h>

// Thread-local variable
_Thread_local int thread_counter = 0;

// Function to be executed by threads
int thread_function(void *arg) {
    int thread_id = *(int*)arg;

    // Each thread has its own copy of thread_counter
    for (int i = 0; i < 5; i++) {
        thread_counter++;
        printf("Thread %d: counter = %d\n", thread_id, thread_counter);
        thrd_sleep(&(struct timespec){.tv_sec=1}, NULL);
    }

    return 0;
}

int main() {
    thrd_t threads[3];
    int thread_ids[3] = {1, 2, 3};

    // Create threads
    for (int i = 0; i < 3; i++) {
        if (thrd_create(&threads[i], thread_function, &thread_ids[i]) != thrd_success) {
            printf("Error creating thread %d\n", i);
            return 1;
        }
    }

    // Wait for threads to complete
    for (int i = 0; i < 3; i++) {
        thrd_join(threads[i], NULL);
    }

    printf("All threads completed\n");
    return 0;
}
```

Thread-Local Storage with Initialization

```
#include <stdio.h>
#include <threads.h>
#include <stdlib.h>
#include <time.h>

// Thread-local variable with initializer
_Thread_local unsigned int thread_seed = 0;

// Function to initialize thread-local seed
void init_thread_seed(void) {
    if (thread_seed == 0) {
        thread_seed = (unsigned int)time(NULL) ^ (unsigned int)(uintptr_t)&thread_seed;
    }
}

int random_thread_function(void *arg) {
    int thread_id = *(int*)arg;

    // Initialize thread-local seed
    init_thread_seed();

    printf("Thread %d seed: %u\n", thread_id, thread_seed);

    // Generate some random numbers
    for (int i = 0; i < 5; i++) {
        int random_num = rand_r(&thread_seed) % 100;
        printf("Thread %d, Random %d: %d\n", thread_id, i, random_num);
        thrd_sleep(&(struct timespec){.tv_sec=0, .tv_nsec=100000000}, NULL);
    }

    return 0;
}

int main() {
    thrd_t threads[3];
    int thread_ids[3] = {1, 2, 3};

    // Create threads
    for (int i = 0; i < 3; i++) {
        if (thrd_create(&threads[i], random_thread_function, &thread_ids[i]) != thrd_success)
            printf("Error creating thread %d\n", i);
        return 1;
    }
}
```



```

    // Wait for threads to complete
    for (int i = 0; i < 3; i++) {
        thrd_join(threads[i], NULL);
    }

    printf("All threads completed\n");
    return 0;
}

```

Bounds-Checking Functions (Annex K)

C11 introduced bounds-checking functions as an optional extension (Annex K) to improve memory safety and prevent buffer overflows. These functions include additional parameters to specify buffer sizes.

Safe String Functions

```

#define __STDC_WANT_LIB_EXT1__ 1
#include <stdio.h>
#include <string.h>
#include <stdlib.h>

#ifdef __STDC_LIB_EXT1__
// Check if bounds-checking functions are available
void safe_string_operations(void) {
    char buffer[20];
    char source[] = "This is a long string that might overflow";
    errno_t result;

    // Safe string copy
    result = strcpy_s(buffer, sizeof(buffer), source);
    if (result != 0) {
        printf("strcpy_s failed: buffer too small\n");
    } else {
        printf("strcpy_s succeeded: %s\n", buffer);
    }

    // Safe string concatenation
    char buffer2[50] = "Hello, ";
    char append[] = "World!";

    result = strcat_s(buffer2, sizeof(buffer2), append);
    if (result != 0) {
        printf("strcat_s failed\n");
    }
}

```

```

    } else {
        printf("strcat_s succeeded: %s\n", buffer2);
    }

    // Safe formatted output
    char formatted[30];
    result = sprintf_s(formatted, sizeof(formatted), "Number: %d, String: %s", 42, "test");
    if (result < 0) {
        printf("sprintf_s failed\n");
    } else {
        printf("sprintf_s succeeded: %s\n", formatted);
    }
}
#else
void safe_string_operations(void) {
    printf("Bounds-checking functions not available\n");
}
#endif

int main() {
    safe_string_operations();
    return 0;
}

```

Safe Memory Functions

```

#define __STDC_WANT_LIB_EXT1__ 1
#include <stdio.h>
#include <string.h>
#include <stdlib.h>

#ifdef __STDC_LIB_EXT1__
void safe_memory_operations(void) {
    char source[] = "Hello, World!";
    char destination[20];
    errno_t result;

    // Safe memory copy
    result = memcpy_s(destination, sizeof(destination), source, strlen(source) + 1);
    if (result != 0) {
        printf("memcpy_s failed\n");
    } else {
        printf("memcpy_s succeeded: %s\n", destination);
    }
}

```

```

// Safe memory set
char buffer[10];
result = memset_s(buffer, sizeof(buffer), 'A', sizeof(buffer) - 1);
if (result != 0) {
    printf("memset_s failed\n");
} else {
    buffer[sizeof(buffer) - 1] = '\0';
    printf("memset_s succeeded: %s\n", buffer);
}
}
#else
void safe_memory_operations(void) {
    printf("Bounds-checking functions not available\n");
}
#endif

int main() {
    safe_memory_operations();
    return 0;
}

```

Generic Selection

C11 introduced the `_Generic` keyword, which allows for type-generic programming by selecting different implementations based on the type of an expression.

Basic Generic Selection

```

#include <stdio.h>

// Generic print function using _Generic
#define print(x) _Generic((x), \
    int: print_int, \
    float: print_float, \
    double: print_double, \
    char*: print_string, \
    default: print_unknown \
)(x)

void print_int(int i) {
    printf("Integer: %d\n", i);
}

void print_float(float f) {

```

```

    printf("Float: %.2f\n", f);
}

void print_double(double d) {
    printf("Double: %.2f\n", d);
}

void print_string(char *s) {
    printf("String: %s\n", s);
}

void print_unknown(void *p) {
    printf("Unknown type: %p\n", p);
}

int main() {
    int i = 42;
    float f = 3.14f;
    double d = 2.71828;
    char *s = "Hello, World!";

    print(i);
    print(f);
    print(d);
    print(s);

    return 0;
}

```

Advanced Generic Selection

```

#include <stdio.h>
#include <math.h>

// Generic absolute value function
#define abs_val(x) _Generic((x), \
    int: abs, \
    long: labs, \
    long long: llabs, \
    float: fabsf, \
    double: fabs, \
    long double: fabsl \
)(x)

// Generic square function

```

```

#define square(x) _Generic((x), \
    int: square_int, \
    float: square_float, \
    double: square_double \
)(x)

int square_int(int x) {
    return x * x;
}

float square_float(float x) {
    return x * x;
}

double square_double(double x) {
    return x * x;
}

int main() {
    int i = -5;
    long l = -1000000L;
    float f = -3.14f;
    double d = -2.71828;

    printf("Absolute values:\n");
    printf("abs(%d) = %d\n", i, abs_val(i));
    printf("abs(%ld) = %ld\n", l, abs_val(l));
    printf("abs(%.2f) = %.2f\n", f, abs_val(f));
    printf("abs(%.2f) = %.2f\n", d, abs_val(d));

    printf("\nSquares:\n");
    printf("square(%d) = %d\n", i, square(i));
    printf("square(%.2f) = %.2f\n", f, square(f));
    printf("square(%.2f) = %.2f\n", d, square(d));

    return 0;
}

```

Alignment Support

C11 introduced alignment support through the `_Alignas` specifier and `_Alignof` operator, allowing for explicit control over memory alignment.

Alignment Specifier

```
#include <stdio.h>
#include <stdalign.h>

// Align structure to 16-byte boundary
struct aligned_struct {
    _Alignas(16) char data[16];
    int value;
};

// Using alignas macro (more portable)
struct portable_aligned {
    alignas(32) double matrix[4][4];
    int id;
};

int main() {
    struct aligned_struct s1;
    struct portable_aligned s2;

    printf("Size of aligned_struct: %zu bytes\n", sizeof(s1));
    printf("Alignment of aligned_struct: %zu\n", _Alignof(struct aligned_struct));
    printf("Address of s1.data: %p (aligned to 16: %s)\n",
        (void*)s1.data, ((uintptr_t)s1.data % 16 == 0) ? "Yes" : "No");

    printf("\nSize of portable_aligned: %zu bytes\n", sizeof(s2));
    printf("Alignment of portable_aligned: %zu\n", alignof(struct portable_aligned));
    printf("Address of s2.matrix: %p (aligned to 32: %s)\n",
        (void*)s2.matrix, ((uintptr_t)s2.matrix % 32 == 0) ? "Yes" : "No");

    return 0;
}
```

Practical Examples

Thread-Safe Logger with Static Assertions

```
#include <stdio.h>
#include <threads.h>
#include <time.h>
#include <stdarg.h>

// Static assertions for system requirements
```

```

_Static_assert(sizeof(void*) == 8, "64-bit system required");
_Static_assert(CHAR_BIT == 8, "8-bit char required");

// Thread-local logger instance
_Thread_local char log_buffer[1024];
_Thread_local int log_level = 0;

typedef enum {
    LOG_DEBUG,
    LOG_INFO,
    LOG_WARNING,
    LOG_ERROR
} LogLevel;

void log_message(LogLevel level, const char *format, ...) {
    const char *level_strings[] = {"DEBUG", "INFO", "WARNING", "ERROR"};

    if (level < log_level) return;

    // Get current time
    time_t now = time(NULL);
    struct tm *tm_info = localtime(&now);

    // Format timestamp
    strftime(log_buffer, sizeof(log_buffer), "%Y-%m-%d %H:%M:%S", tm_info);

    // Append log level and message
    int offset = strlen(log_buffer);
    snprintf(log_buffer + offset, sizeof(log_buffer) - offset,
             " [%s] ", level_strings[level]);

    // Format the actual message
    va_list args;
    va_start(args, format);
    offset = strlen(log_buffer);
    vsnprintf(log_buffer + offset, sizeof(log_buffer) - offset, format, args);
    va_end(args);

    printf("%s\n", log_buffer);
}

int logger_thread(void *arg) {
    int thread_id = *(int*)arg;

    // Set different log levels for different threads
    log_level = thread_id % 4;
}

```

```

    for (int i = 0; i < 3; i++) {
        log_message(LOG_DEBUG, "Thread %d debug message %d", thread_id, i);
        log_message(LOG_INFO, "Thread %d info message %d", thread_id, i);
        log_message(LOG_WARNING, "Thread %d warning message %d", thread_id, i);
        log_message(LOG_ERROR, "Thread %d error message %d", thread_id, i);
        thrd_sleep(&(struct timespec){.tv_sec=1}, NULL);
    }

    return 0;
}

int main() {
    thrd_t threads[3];
    int thread_ids[3] = {1, 2, 3};

    printf("Starting multi-threaded logger demo\n");

    // Create threads
    for (int i = 0; i < 3; i++) {
        if (thrd_create(&threads[i], logger_thread, &thread_ids[i]) != thrd_success) {
            printf("Error creating thread %d\n", i);
            return 1;
        }
    }

    // Wait for threads to complete
    for (int i = 0; i < 3; i++) {
        thrd_join(threads[i], NULL);
    }

    printf("Logger demo completed\n");
    return 0;
}

```

Summary

C11 introduced several important features that enhanced the safety, expressiveness, and concurrency support of the C programming language:

1. **Static Assertions:** Compile-time checks with `_Static_assert`
2. **Anonymous Structures/Unions:** Embed structures directly without member names
3. **Thread-Local Storage:** Per-thread variable instances with `_Thread_local`
4. **Bounds-Checking Functions:** Safer string and memory operations (Annex K)
5. **Generic Selection:** Type-generic programming with `_Generic`
6. **Alignment Support:** Explicit memory alignment control with `_Alignas` and `_Alignof`

These features make C11 a more robust and modern language while maintaining backward compatibility with earlier standards. Understanding and utilizing these features is essential for writing safe, efficient, and maintainable C code in contemporary development environments.

C17 and C23 Features

Introduction

The C17 standard (ISO/IEC 9899:2018) was primarily a technical corrigendum that fixed defects in the C11 standard without introducing significant new features. However, the upcoming C23 standard (ISO/IEC 9899:2024) represents a major evolution of the C language, introducing numerous modern features that make C more expressive, safer, and aligned with contemporary programming practices.

C23 features include `typeof` operators, binary literals, enhanced preprocessor directives, attribute syntax, and many other improvements that address long-standing limitations in the language. Understanding these features is crucial for writing modern C code.

C17: Technical Corrigendum

C17 was focused on fixing defects and inconsistencies in C11 rather than adding new features. Some notable changes include:

Defect Fixes

- Clarification of atomic operations semantics
- Fixes to alignment requirements
- Corrections to floating-point behavior
- Improvements to multi-threading support

Example of C17 Improvements

```
#include <stdio.h>
#include <stdalign.h>
#include <stdatomic.h>

// C17 clarified alignment requirements
_Alignas(16) struct aligned_data {
    int values[4];
};

// C17 improved atomic operations
```

```

_Atomic int atomic_counter = 0;

int main() {
    printf("Alignment of aligned_data: %zu\n", alignof(struct aligned_data));

    // Atomic operations with clearer semantics
    atomic_fetch_add(&atomic_counter, 1);
    printf("Atomic counter: %d\n", atomic_load(&atomic_counter));

    return 0;
}

```

C23: typeof Operators

One of the most significant additions in C23 is the `typeof` operator, which allows you to declare variables with the same type as an expression. This feature brings C closer to C++'s `decltype` and makes generic programming much easier.

Basic typeof Usage

```

#include <stdio.h>

#define max(a, b) ({ \
    typeof(a) _a = (a); \
    typeof(b) _b = (b); \
    _a > _b ? _a : _b; \
})

int main() {
    int x = 10, y = 20;
    double a = 3.14, b = 2.71;

    // typeof allows the macro to work with different types
    int max_int = max(x, y);
    double max_double = max(a, b);

    printf("Max of %d and %d: %d\n", x, y, max_int);
    printf("Max of %.2f and %.2f: %.2f\n", a, b, max_double);

    return 0;
}

```

typeof with Complex Expressions

```
#include <stdio.h>
#include <stdlib.h>

int main() {
    int array[10] = {1, 2, 3, 4, 5, 6, 7, 8, 9, 10};

    // Declare variable with same type as array element
    typeof(array[0]) element = 42;

    // Declare variable with same type as array (pointer decay)
    typeof(array) *ptr = &array;

    // Declare variable with same type as complex expression
    typeof((int)3.14 + (long)42) result = 0;

    printf("Element: %d\n", element);
    printf("First array element via pointer: %d\n", (*ptr)[0]);
    printf("Result type size: %zu bytes\n", sizeof(result));

    return 0;
}
```

typeof_unqual for Unqualified Types

```
#include <stdio.h>

int main() {
    const volatile int value = 42;

    // typeof preserves qualifiers
    typeof(value) qualified_var = 100; // Still const volatile

    // typeof_unqual removes all qualifiers
    typeof_unqual(value) unqualified_var = 200; // Just int

    printf("Qualified var type size: %zu\n", sizeof(qualified_var));
    printf("Unqualified var type size: %zu\n", sizeof(unqualified_var));

    return 0;
}
```

Binary Literals

C23 introduces binary literals, allowing you to write numbers in base-2 format using the `0b` or `0B` prefix.

Basic Binary Literals

```
#include <stdio.h>

int main() {
    // Binary literals
    int binary1 = 0b1010;           // 10 in decimal
    int binary2 = 0B11110000;       // 240 in decimal
    int binary3 = 0b1111'1111;      // 255 in decimal (with digit separator)

    printf("0b1010 = %d\n", binary1);
    printf("0B11110000 = %d\n", binary2);
    printf("0b1111'1111 = %d\n", binary3);

    // Useful for bit manipulation
    unsigned char flags = 0b1010'0101;
    printf("Flags: 0x%02X (binary: ", flags);
    for (int i = 7; i >= 0; i--) {
        printf("%d", (flags >> i) & 1);
    }
    printf(")\n");

    return 0;
}
```

Bit Manipulation with Binary Literals

```
#include <stdio.h>

// Bit manipulation macros using binary literals
#define BIT(n) (1U << (n))
#define SET_BIT(value, n) ((value) |= BIT(n))
#define CLEAR_BIT(value, n) ((value) &= ~BIT(n))
#define TOGGLE_BIT(value, n) ((value) ^= BIT(n))
#define CHECK_BIT(value, n) (((value) & BIT(n)) != 0)

int main() {
    unsigned char status = 0b0000'0000;
```

```

printf("Initial status: 0b");
for (int i = 7; i >= 0; i--) {
    printf("%d", (status >> i) & 1);
}
printf("\n");

// Set some bits
SET_BIT(status, 0); // Set bit 0
SET_BIT(status, 3); // Set bit 3
SET_BIT(status, 7); // Set bit 7

printf("After setting bits 0, 3, 7: 0b");
for (int i = 7; i >= 0; i--) {
    printf("%d", (status >> i) & 1);
}
printf("\n");

// Toggle a bit
TOGGLE_BIT(status, 3); // Toggle bit 3

printf("After toggling bit 3: 0b");
for (int i = 7; i >= 0; i--) {
    printf("%d", (status >> i) & 1);
}
printf("\n");

// Check specific bits
printf("Bit 0 is %s\n", CHECK_BIT(status, 0) ? "set" : "clear");
printf("Bit 3 is %s\n", CHECK_BIT(status, 3) ? "set" : "clear");
printf("Bit 5 is %s\n", CHECK_BIT(status, 5) ? "set" : "clear");

return 0;
}

```

Enhanced Preprocessor Directives

C23 introduces several new preprocessor directives that make conditional compilation more powerful and flexible.

#elifdef and #elifndef

```
#include <stdio.h>

#define DEBUG_MODE
// #define PRODUCTION_MODE

int main() {
    #if defined(DEBUG_MODE)
        printf("Debug mode enabled\n");
    #elifdef PRODUCTION_MODE
        printf("Production mode enabled\n");
    #elifndef TESTING_MODE
        printf("Testing mode not defined, using default mode\n");
    #else
        printf("Unknown mode\n");
    #endif

    return 0;
}
```

#warning Directive

```
#include <stdio.h>

// Compiler version check
#if __STDC_VERSION__ < 202311L
    #warning "This code is designed for C23 or later"
#endif

// Platform-specific warnings
#ifdef __APPLE__
    #warning "This code may need adjustments for macOS"
#endif

int main() {
    printf("Program running\n");

    #ifdef __GNUC__
        #warning "Consider using Clang for better C23 support"
    #endif

    return 0;
}
```

#embed Directive (Conceptual)

Note: The #embed directive is still being standardized and may not be available in all compilers yet.

```
// Example of proposed #embed directive
// #embed "shader.glsl" limit(1024) prefix(0xFF, 0xD8) suffix(0x00)

// This would embed the contents of shader.glsl as a byte array
// with a JPEG header prefix and null terminator suffix
```

Attribute Syntax

C23 introduces a standardized attribute syntax using double square brackets [[]], making it easier to specify compiler-specific features in a portable way.

Standard Attributes

```
#include <stdio.h>

// [[nodiscard]] - warn if return value is ignored
[[nodiscard]] int calculate_value(int x) {
    return x * x;
}

// [[maybe_unused]] - suppress unused variable warnings
void example_function([[maybe_unused]] int debug_param) {
    // debug_param might only be used in debug builds
#ifdef NDEBUG
    printf("Debug param: %d\n", debug_param);
#endif
}

// [[deprecated]] - mark functions as deprecated
[[deprecated("Use calculate_value instead")]]
int old_function(int x) {
    return x + x;
}

// [[fallthrough]] - indicate intentional fallthrough in switch
int process_state(int state) {
    switch (state) {
        case 1:
        case 2:
```



```

        printf("Low state\n");
        [[fallthrough]];
    case 3:
    case 4:
        printf("Medium state\n");
        [[fallthrough]];
    case 5:
        printf("High state\n");
        break;
    default:
        printf("Unknown state\n");
}
return 0;
}

int main() {
    // [[nodiscard]] example - compiler should warn if ignored
    calculate_value(5); // Warning: ignoring return value

    example_function(42);

    // [[deprecated]] example
    old_function(10); // Warning: deprecated function

    process_state(3);

    return 0;
}

```

Extended Integer Types

C23 introduces `_BitInt(N)` types for extended integer precision, allowing for integers with more than 64 bits.

`_BitInt` Example (Conceptual)

```

// Note: _BitInt is still being implemented in compilers
/*
#include <stdio.h>

int main() {
    // 128-bit integer
    _BitInt(128) big_number = 0;

```

```

    // 256-bit integer for cryptographic operations
    _BitInt(256) crypto_key = 0;

    printf("Size of 128-bit integer: %zu bytes\n", sizeof(big_number));
    printf("Size of 256-bit integer: %zu bytes\n", sizeof(crypto_key));

    return 0;
}
*/

```

Enhanced `_Static_assert`

C23 improves the `_Static_assert` macro by making the message parameter optional.

Enhanced Static Assertions

```

#include <stdio.h>
#include <stdint.h>

// Traditional static assertion with message
_Static_assert(sizeof(int) >= 4, "int must be at least 32 bits");

// C23 allows static assertion without message
_Static_assert(sizeof(long long) > sizeof(long));

// Complex static assertions
typedef struct {
    uint32_t header;
    char data[256];
    uint32_t checksum;
} Packet;

_Static_assert(sizeof(Packet) == 264, "Packet structure size must be exactly 264 bytes");
_Static_assert(_Alignof(Packet) >= 4, "Packet must be aligned to at least 4 bytes");

int main() {
    printf("All static assertions passed\n");
    printf("Size of Packet: %zu bytes\n", sizeof(Packet));
    printf("Alignment of Packet: %zu bytes\n", _Alignof(Packet));

    return 0;
}

```

New Standard Library Functions

C23 adds several new standard library functions, including improved string handling and memory allocation functions.

strdup and strndup

```
#define _GNU_SOURCE // For strdup in some implementations
#include <stdio.h>
#include <stdlib.h>
#include <string.h>

// Simple implementation of strdup if not available
char *my_strdup(const char *s) {
    size_t len = strlen(s) + 1;
    char *copy = malloc(len);
    if (copy) {
        memcpy(copy, s, len);
    }
    return copy;
}

int main() {
    const char *original = "Hello, World!";

    // Using strdup (if available in C23 implementation)
#ifdef _GNU_SOURCE
    char *copy1 = strdup(original);
    if (copy1) {
        printf("strdup copy: %s\n", copy1);
        free(copy1);
    }
#endif

    // Using custom implementation
    char *copy2 = my_strdup(original);
    if (copy2) {
        printf("my_strdup copy: %s\n", copy2);
        free(copy2);
    }

    // strndup example (copy at most n characters)
    char *copy3 = strndup(original, 5);
    if (copy3) {
        printf("strndup copy (5 chars): %s\n", copy3);
        free(copy3);
    }
}
```

```

    }

    return 0;
}

```

Enhanced aligned_alloc

```

#include <stdio.h>
#include <stdlib.h>

int main() {
    // Allocate 100 bytes aligned to 16-byte boundary
    void *ptr = aligned_alloc(16, 100);
    if (ptr) {
        printf("Allocated address: %p\n", ptr);
        printf("Alignment check: %s\n",
               ((uintptr_t)ptr % 16 == 0) ? "OK" : "FAIL");

        // Use the memory...

        free(ptr);
    } else {
        printf("Allocation failed\n");
    }

    return 0;
}

```

Practical Examples

Modern Configuration System

```

#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <stdbool.h>

// Using binary literals for configuration flags
typedef enum {
    CFG_DEBUG      = 0b0000'0001,
    CFG_LOGGING    = 0b0000'0010,
    CFG_NETWORK    = 0b0000'0100,
    CFG_DATABASE   = 0b0000'1000,
}

```

```

    CFG_SECURITY    = 0b0001'0000,
    CFG_COMPRESSION= 0b0010'0000
} ConfigFlags;

// Using typeof for generic configuration handling
#define CONFIG_GET(config, key) \
    typeof(((config)->key)) get_config_value(const Config *c, const char *k) { \
        if (strcmp(k, #key) == 0) return c->key; \
        /* Handle other keys */ \
        return (typeof(((config)->key)))0; \
    }

typedef struct {
    int max_connections;
    double timeout;
    bool enable_cache;
    ConfigFlags flags;
} Config;

// Using attributes for better code quality
[[nodiscard]] bool validate_config(const Config *config) {
    [[maybe_unused]] int debug_value = 42;

    if (config->max_connections <= 0) {
        return false;
    }

    if (config->timeout < 0.0) {
        return false;
    }

    return true;
}

[[deprecated("Use validate_config instead")]]
bool old_validate_config(Config *config) {
    return validate_config(config);
}

void print_config(const Config *config) {
    printf("Configuration:\n");
    printf("  Max Connections: %d\n", config->max_connections);
    printf("  Timeout: %.2f\n", config->timeout);
    printf("  Enable Cache: %s\n", config->enable_cache ? "Yes" : "No");
    printf("  Flags: 0b");
    for (int i = 7; i >= 0; i--) {
        printf("%d", (config->flags >> i) & 1);
    }
}

```

```

    }
    printf("\n");
}

int main() {
    // Initialize configuration with binary literals and designated initializers
    Config config = {
        .max_connections = 100,
        .timeout = 30.0,
        .enable_cache = true,
        .flags = CFG_DEBUG | CFG_LOGGING | CFG_NETWORK
    };

    // Static assertion to ensure configuration structure size
    _Static_assert(sizeof(Config) > 0, "Config structure must not be empty");

    if (validate_config(&config)) {
        print_config(&config);
    } else {
        printf("Invalid configuration\n");
    }

    // Deprecated function usage (should generate warning)
    old_validate_config(&config);

    return 0;
}

```

Summary

C17 and C23 represent important milestones in the evolution of the C programming language:

1. **C17**: Technical corrigendum fixing defects in C11
2. **C23**: Major feature release with modern language enhancements

Key C23 features include:

1. **typeof Operators**: Type introspection and generic programming
2. **Binary Literals**: More intuitive bit manipulation
3. **Enhanced Preprocessor**: More powerful conditional compilation
4. **Attribute Syntax**: Standardized compiler hints
5. **Extended Integer Types**: Support for arbitrary precision integers
6. **Enhanced Static Assertions**: More flexible compile-time checking
7. **New Standard Library Functions**: Improved string and memory operations

These features make modern C more expressive, safer, and better aligned with contemporary programming practices while maintaining the language's core principles of efficiency and low-level control.

Modern C Development Practices

Introduction

Modern C development has evolved significantly from traditional approaches, incorporating contemporary software engineering practices, DevOps methodologies, and cloud-native technologies. Today's C developers work with sophisticated toolchains, containerized environments, continuous integration systems, and modern deployment strategies.

This chapter explores the current landscape of C development, including DevOps integration, containerized development environments, WebAssembly compilation, cloud-native applications, and modern testing practices. Understanding these practices is essential for professional C development in 2026 and beyond.

DevOps and C: CI/CD Pipelines

Continuous Integration and Continuous Deployment (CI/CD) pipelines are crucial for modern software development, ensuring code quality, automated testing, and reliable deployments.

GitHub Actions for C Projects

```
# .github/workflows/build.yml
name: C CI

on:
  push:
    branches: [ main ]
  pull_request:
    branches: [ main ]

jobs:
  build:
    runs-on: ubuntu-latest

    steps:
      - uses: actions/checkout@v3

      - name: Install dependencies
        run: |
```



```

    sudo apt-get update
    sudo apt-get install -y build-essential cmake valgrind

- name: Configure CMake
  run: cmake -B ${github.workspace}}/build -DCMAKE_BUILD_TYPE=${env.BUILD_TYPE}}

- name: Build
  run: cmake --build ${github.workspace}}/build --config ${env.BUILD_TYPE}}

- name: Test
  working-directory: ${github.workspace}}/build
  run: ctest -C ${env.BUILD_TYPE}}

- name: Static Analysis
  run: |
    sudo apt-get install -y cppcheck
    cppcheck --enable=all --inconclusive src/

- name: Memory Check
  run: |
    cd ${github.workspace}}/build
    valgrind --leak-check=full --error-exitcode=1 ./my_program

```

GitLab CI for C Projects

```

# .gitlab-ci.yml
stages:
  - build
  - test
  - analyze
  - deploy

variables:
  BUILD_TYPE: Release

before_script:
  - apt-get update && apt-get install -y build-essential cmake

build_job:
  stage: build
  script:
    - cmake -B build -DCMAKE_BUILD_TYPE=$BUILD_TYPE
    - cmake --build build --config $BUILD_TYPE
  artifacts:
    paths:

```

```

        - build/

test_job:
  stage: test
  script:
    - cd build
    - ctest -C $BUILD_TYPE
  dependencies:
    - build_job

analyze_job:
  stage: analyze
  script:
    - apt-get install -y cppcheck
    - cppcheck --enable=all --inconclusive src/
  allow_failure: true

deploy_job:
  stage: deploy
  script:
    - echo "Deploying application..."
    - # Deployment commands here
  only:
    - main

```

Containerized Development

Containerization provides consistent development environments, simplifies dependency management, and enables reproducible builds.

Docker for C Development

```

# Dockerfile for C development environment
FROM ubuntu:22.04

# Install build tools and dependencies
RUN apt-get update && apt-get install -y \
    build-essential \
    cmake \
    gcc \
    g++ \
    gdb \
    valgrind \
    clang \

```

```

    clang-tools \
    cppcheck \
    git \
    vim \
    && rm -rf /var/lib/apt/lists/*

# Install additional tools
RUN apt-get update && apt-get install -y \
    lcov \
    doxygen \
    graphviz \
    && rm -rf /var/lib/apt/lists/*

# Set working directory
WORKDIR /app

# Copy project files
COPY . .

# Default command
CMD ["/bin/bash"]

```

Multi-stage Docker Build

```

# Multi-stage Dockerfile for C application
# Build stage
FROM ubuntu:22.04 AS builder

# Install build dependencies
RUN apt-get update && apt-get install -y \
    build-essential \
    cmake \
    && rm -rf /var/lib/apt/lists/*

# Set working directory
WORKDIR /app

# Copy source code
COPY src/ src/
COPY include/ include/
COPY CMakeLists.txt .

# Build the application
RUN cmake -B build -DCMAKE_BUILD_TYPE=Release
RUN cmake --build build --config Release

```

```

# Runtime stage
FROM ubuntu:22.04

# Install runtime dependencies only
RUN apt-get update && apt-get install -y \
    libssl-dev \
    && rm -rf /var/lib/apt/lists/*

# Copy built executable from builder stage
COPY --from=builder /app/build/myapp /usr/local/bin/myapp

# Create non-root user
RUN useradd -m -s /bin/bash appuser
USER appuser

# Expose port if needed
EXPOSE 8080

# Run the application
CMD ["myapp"]

```

Docker Compose for Development

```

# docker-compose.yml
version: '3.8'

services:
  app:
    build: .
    volumes:
      - ./app
      - /app/build
    working_dir: /app
    command: bash -c "cmake -B build && cmake --build build && ./build/myapp"
    environment:
      - DEBUG=1
    ports:
      - "8080:8080"
    depends_on:
      - database
      - cache

  database:
    image: postgres:15
    environment:

```

```

    POSTGRES_DB: myapp
    POSTGRES_USER: user
    POSTGRES_PASSWORD: password
volumes:
  - db_data:/var/lib/postgresql/data
ports:
  - "5432:5432"

cache:
  image: redis:7
  ports:
    - "6379:6379"

test:
  build: .
  volumes:
    - ./app
  working_dir: /app
  command: bash -c "cmake -B build && cmake --build build && ctest -V"
  depends_on:
    - database
    - cache

volumes:
  db_data:

```

Cross-compilation

Modern C development often requires building for multiple target platforms and architectures.

CMake Cross-compilation

```

# CMakeLists.txt with cross-compilation support
cmake_minimum_required(VERSION 3.20)
project(MyApp C)

# Set C standard
set(CMAKE_C_STANDARD 23)
set(CMAKE_C_STANDARD_REQUIRED ON)

# Compiler-specific flags
if(CMAKE_C_COMPILER_ID STREQUAL "GNU")
    add_compile_options(-Wall -Wextra -Wpedantic)
elseif(CMAKE_C_COMPILER_ID STREQUAL "Clang")

```

```

        add_compile_options(-Wall -Wextra -Wpedantic)
endif()

# Conditional compilation for different platforms
if(WIN32)
    add_compile_definitions(PLATFORM_WINDOWS)
elseif(APPLE)
    add_compile_definitions(PLATFORM_MACOS)
elseif(UNIX)
    add_compile_definitions(PLATFORM_LINUX)
endif()

# Source files
set(SOURCES
    src/main.c
    src/utils.c
    src/network.c
)

# Create executable
add_executable(myapp ${SOURCES})

# Link libraries based on platform
if(WIN32)
    target_link_libraries(myapp ws2_32)
elseif(UNIX)
    target_link_libraries(myapp pthread m)
endif()

# Installation rules
install(TARGETS myapp DESTINATION bin)

```

Cross-compilation Toolchain File

```

# arm-linux-gnueabi.cmake
set(CMAKE_SYSTEM_NAME Linux)
set(CMAKE_SYSTEM_PROCESSOR arm)

set(CMAKE_C_COMPILER arm-linux-gnueabi-gcc)
set(CMAKE_CXX_COMPILER arm-linux-gnueabi-g++)

set(CMAKE_FIND_ROOT_PATH /usr/arm-linux-gnueabi)
set(CMAKE_FIND_ROOT_PATH_MODE_PROGRAM NEVER)
set(CMAKE_FIND_ROOT_PATH_MODE_LIBRARY ONLY)
set(CMAKE_FIND_ROOT_PATH_MODE_INCLUDE ONLY)

```

```
set(CMAKE_FIND_ROOT_PATH_MODE_PACKAGE ONLY)
```

WebAssembly (WASM)

Compiling C to WebAssembly enables running C applications in web browsers and other WASM environments.

Emscripten Setup

```
# Install Emscripten
git clone https://github.com/emscripten-core/emsdk.git
cd emsdk
./emsdk install latest
./emsdk activate latest
source ./emsdk_env.sh
```

Simple C to WASM Example

```
// math_utils.c
#include <emscripten.h>

EMSCRIPTEN_KEEPALIVE
int add(int a, int b) {
    return a + b;
}

EMSCRIPTEN_KEEPALIVE
int fibonacci(int n) {
    if (n <= 1) return n;
    return fibonacci(n - 1) + fibonacci(n - 2);
}

EMSCRIPTEN_KEEPALIVE
double calculate_circle_area(double radius) {
    return 3.14159 * radius * radius;
}

# Compile to WebAssembly
emcc math_utils.c -o math_utils.js \
    -s EXPORTED_FUNCTIONS='["_add", "_fibonacci", "_calculate_circle_area"]' \
    -s EXPORTED_RUNTIME_METHODS='["ccall", "cwrap"]' \
    -s MODULARIZE=1 \
```

```
-s EXPORT_NAME="MathUtils"
```

HTML Integration

```
<!DOCTYPE html>
<html>
<head>
  <title>C WebAssembly Demo</title>
</head>
<body>
  <h1>C WebAssembly Math Demo</h1>
  <div id="output"></div>

  <script src="math_utils.js"></script>
  <script>
    MathUtils().then(function(Module) {
      // Wrap C functions
      const add = Module.cwrap('add', 'number', ['number', 'number']);
      const fibonacci = Module.cwrap('fibonacci', 'number', ['number']);
      const calculate_circle_area = Module.cwrap('calculate_circle_area', 'number', ['number', 'number']);

      // Use the functions
      document.getElementById('output').innerHTML = `
        <p>5 + 3 = ${add(5, 3)}</p>
        <p>Fibonacci(10) = ${fibonacci(10)}</p>
        <p>Circle area (radius=5) = ${calculate_circle_area(5)}</p>
      `;
    });
  </script>
</body>
</html>
```

Cloud-native C Applications

Cloud-native development involves building applications that are designed to run in cloud environments, leveraging microservices, containers, and orchestration platforms.

Simple HTTP Server in C

```
#define _GNU_SOURCE
#include <stdio.h>
#include <stdlib.h>
```



```

#include <string.h>
#include <unistd.h>
#include <sys/socket.h>
#include <netinet/in.h>
#include <arpa/inet.h>
#include <time.h>

#define PORT 8080
#define BUFFER_SIZE 1024

void handle_request(int client_fd) {
    char buffer[BUFFER_SIZE];
    ssize_t bytes_read;

    // Read HTTP request
    bytes_read = read(client_fd, buffer, BUFFER_SIZE - 1);
    if (bytes_read > 0) {
        buffer[bytes_read] = '\0';

        // Simple HTTP response
        const char *response =
            "HTTP/1.1 200 OK\r\n"
            "Content-Type: text/html\r\n"
            "Connection: close\r\n"
            "\r\n"
            "<html><body><h1>Hello from C!</h1>"
            "<p>Current time: %s</p>"
            "</body></html>";

        // Get current time
        time_t now = time(NULL);
        char *time_str = ctime(&now);
        time_str[strlen(time_str) - 1] = '\0'; // Remove newline

        // Send response
        char response_buffer[BUFFER_SIZE * 2];
        snprintf(response_buffer, sizeof(response_buffer), response, time_str);
        write(client_fd, response_buffer, strlen(response_buffer));
    }

    close(client_fd);
}

int main() {
    int server_fd, client_fd;
    struct sockaddr_in address;
    int opt = 1;

```

```

int addrlen = sizeof(address);

// Create socket
if ((server_fd = socket(AF_INET, SOCK_STREAM, 0)) == 0) {
    perror("socket failed");
    exit(EXIT_FAILURE);
}

// Set socket options
if (setsockopt(server_fd, SOL_SOCKET, SO_REUSEADDR | SO_REUSEPORT,
               &opt, sizeof(opt))) {
    perror("setsockopt");
    exit(EXIT_FAILURE);
}

// Configure address
address.sin_family = AF_INET;
address.sin_addr.s_addr = INADDR_ANY;
address.sin_port = htons(PORT);

// Bind socket
if (bind(server_fd, (struct sockaddr *)&address, sizeof(address)) < 0) {
    perror("bind failed");
    exit(EXIT_FAILURE);
}

// Listen for connections
if (listen(server_fd, 3) < 0) {
    perror("listen");
    exit(EXIT_FAILURE);
}

printf("Server listening on port %d\n", PORT);

// Main server loop
while (1) {
    // Accept connection
    if ((client_fd = accept(server_fd, (struct sockaddr *)&address,
                           (socklen_t *)&addrlen)) < 0) {
        perror("accept");
        continue;
    }

    // Handle request
    handle_request(client_fd);
}

```

```
    return 0;
}
```

Dockerfile for Cloud Deployment

```
# Dockerfile for cloud-native C application
FROM alpine:latest

# Install build tools
RUN apk add --no-cache build-base

# Set working directory
WORKDIR /app

# Copy source code
COPY server.c .

# Build the application
RUN gcc -o server server.c

# Create non-root user
RUN adduser -D -s /bin/sh appuser

# Switch to non-root user
USER appuser

# Expose port
EXPOSE 8080

# Health check
HEALTHCHECK --interval=30s --timeout=3s --start-period=5s --retries=3 \
    CMD wget --quiet --tries=1 --spider http://localhost:8080/ || exit 1

# Run the application
CMD ["/server"]
```

Supply Chain Security

Modern C development requires attention to supply chain security, dependency management, and vulnerability scanning.

Conan Package Manager

```
# conanfile.txt
[requires]
zlib/1.2.13
openssl/3.1.0
libcurl/8.0.1

[generators]
CMakeDeps
CMakeToolchain

[options]
openssl:shared=False
zlib:shared=False


# CMakeLists.txt with Conan integration
cmake_minimum_required(VERSION 3.20)
project(MyApp C)

# Include Conan-generated files
include(${CMAKE_BINARY_DIR}/conan_toolchain.cmake)
include(${CMAKE_BINARY_DIR}/conan_deps.cmake)

# Find packages
find_package(ZLIB REQUIRED)
find_package(OpenSSL REQUIRED)
find_package(CURL REQUIRED)

# Create executable
add_executable(myapp src/main.c)

# Link libraries
target_link_libraries(myapp
    ZLIB::ZLIB
    OpenSSL::SSL
    CURL::libcurl
)
```

Vulnerability Scanning

```
# Using OSV-Scanner for vulnerability detection
osv-scanner --lockfile=conan.lock

# Using Snyk for security scanning
```

```
snyk test --file=conanfile.txt

# Using GitHub Dependabot (configured in .github/dependabot.yml)

# .github/dependabot.yml
version: 2
updates:
  - package-ecosystem: "git submodule"
    directory: "/"
    schedule:
      interval: "weekly"

  - package-ecosystem: "docker"
    directory: "/"
    schedule:
      interval: "weekly"
```

Modern Testing Practices

Contemporary C development incorporates advanced testing methodologies including fuzzing, property-based testing, and mutation testing.

Unity Testing Framework

```
// test_math.c
#include "unity.h"
#include "math_utils.h"

void setUp(void) {
    // Set up before each test
}

void tearDown(void) {
    // Clean up after each test
}

void test_addition_positive_numbers(void) {
    TEST_ASSERT_EQUAL_INT(5, add(2, 3));
    TEST_ASSERT_EQUAL_INT(0, add(-5, 5));
}

void test_addition_negative_numbers(void) {
    TEST_ASSERT_EQUAL_INT(-5, add(-2, -3));
}
```

```

    TEST_ASSERT_EQUAL_INT(-1, add(2, -3));
}

void test_fibonacci_base_cases(void) {
    TEST_ASSERT_EQUAL_INT(0, fibonacci(0));
    TEST_ASSERT_EQUAL_INT(1, fibonacci(1));
}

void test_fibonacci_recursive_cases(void) {
    TEST_ASSERT_EQUAL_INT(1, fibonacci(2));
    TEST_ASSERT_EQUAL_INT(2, fibonacci(3));
    TEST_ASSERT_EQUAL_INT(3, fibonacci(4));
    TEST_ASSERT_EQUAL_INT(5, fibonacci(5));
}

int main(void) {
    UNITY_BEGIN();
    RUN_TEST(test_addition_positive_numbers);
    RUN_TEST(test_addition_negative_numbers);
    RUN_TEST(test_fibonacci_base_cases);
    RUN_TEST(test_fibonacci_recursive_cases);
    return UNITY_END();
}

```

Property-Based Testing with CMocka

```

// property_test.c
#include <stdarg.h>
#include <stddef.h>
#include <setjmp.h>
#include <cmocka.h>
#include <stdlib.h>

// Property: Addition is commutative
static void test_addition_commutative(void **state) {
    (void) state; // unused

    for (int i = 0; i < 100; i++) {
        int a = rand() % 1000 - 500; // -500 to 499
        int b = rand() % 1000 - 500; // -500 to 499

        assert_int_equal(add(a, b), add(b, a));
    }
}

```

```

// Property: Adding zero is identity
static void test_addition_identity(void **state) {
    (void) state; // unused

    for (int i = 0; i < 100; i++) {
        int a = rand() % 1000 - 500; // -500 to 499

        assert_int_equal(a, add(a, 0));
        assert_int_equal(a, add(0, a));
    }
}

int main(void) {
    const struct CMUnitTest tests[] = {
        cmocka_unit_test(test_addition_commutative),
        cmocka_unit_test(test_addition_identity),
    };

    return cmocka_run_group_tests(tests, NULL, NULL);
}

```

Fuzzing with AFL

```

// fuzz_target.c
#include <stdio.h>
#include <stdlib.h>
#include <string.h>

// Function to fuzz test
int parse_config(const char *input, size_t len) {
    if (len < 4) return -1;

    // Simple parsing logic
    if (input[0] == 'C' && input[1] == 'F' && input[2] == 'G') {
        int value = 0;
        for (size_t i = 3; i < len && i < 10; i++) {
            if (input[i] >= '0' && input[i] <= '9') {
                value = value * 10 + (input[i] - '0');
            } else {
                break;
            }
        }
        return value;
    }
}

```

```

    return -1;
}

#ifdef __AFL_HAVE_MANUAL_CONTROL
int main() {
    __AFL_INIT();

    unsigned char *input = __AFL_FUZZ_TESTCASE_BUF;
    while (__AFL_LOOP(10000)) {
        size_t len = __AFL_FUZZ_TESTCASE_LEN;
        parse_config((char*)input, len);
    }

    return 0;
}
#else
int main(int argc, char **argv) {
    if (argc != 2) {
        fprintf(stderr, "Usage: %s <input_file>\n", argv[0]);
        return 1;
    }

    FILE *fp = fopen(argv[1], "rb");
    if (!fp) {
        perror("fopen");
        return 1;
    }

    fseek(fp, 0, SEEK_END);
    long len = ftell(fp);
    fseek(fp, 0, SEEK_SET);

    char *input = malloc(len);
    if (!input) {
        fclose(fp);
        return 1;
    }

    fread(input, 1, len, fp);
    fclose(fp);

    parse_config(input, len);
    free(input);

    return 0;
}
#endif

```



```

# Build for fuzzing
AFL_USE_ASAN=1 afl-gcc -o fuzz_target fuzz_target.c

# Create input directory
mkdir -p input
echo "CFG123" > input/sample.cfg

# Run fuzzer
afl-fuzz -i input -o output -- ./fuzz_target @@

```

Documentation as Code

Modern C projects integrate documentation generation into the development workflow.

Doxygen Configuration

```

# Doxyfile
PROJECT_NAME = "My C Application"
PROJECT_VERSION = "1.0.0"
PROJECT_BRIEF = "A modern C application with comprehensive documentation"

INPUT = src include
RECURSIVE = YES

GENERATE_HTML = YES
GENERATE_LATEX = NO
GENERATE_MAN = NO

EXTRACT_ALL = YES
EXTRACT_PRIVATE = NO
EXTRACT_STATIC = YES

WARN_IF_UNDOCUMENTED = YES
WARN_IF_DOC_ERROR = YES

FILE_PATTERNS = *.c *.h
RECURSIVE = YES

HAVE_DOT = YES
CALL_GRAPH = YES
CALLER_GRAPH = YES

GENERATE_TREEVIEW = YES
DISABLE_INDEX = NO

```

```
HTML_OUTPUT = docs/html
HTML_FILE_EXTENSION = .html
```

Example Documentation

```
/**
 * @file math_utils.h
 * @brief Mathematical utility functions
 * @author John Doe
 * @date 2024
 */

#ifndef MATH_UTILS_H
#define MATH_UTILS_H

/**
 * @brief Adds two integers
 * @param a First integer
 * @param b Second integer
 * @return Sum of a and b
 * @note This function handles both positive and negative integers
 * @warning Integer overflow is not checked
 *
 * Example usage:
 * @code
 * int result = add(5, 3); // result = 8
 * @endcode
 */
int add(int a, int b);

/**
 * @brief Calculates the nth Fibonacci number
 * @param n The position in the Fibonacci sequence (0-indexed)
 * @return The nth Fibonacci number
 * @pre n >= 0
 * @post Return value >= 0
 *
 * @details This function uses a recursive approach for simplicity.
 * For large values of n, consider using an iterative implementation.
 *
 * Time complexity:  $O(2^n)$ 
 * Space complexity:  $O(n)$ 
 */
int fibonacci(int n);
```

```
#endif // MATH_UTILS_H
```

Practical Examples

Modern C Project Structure

```
my-c-project/  
  src/  
    main.c  
    utils.c  
    utils.h  
  include/  
    config.h  
  tests/  
    test_utils.c  
    CMakeLists.txt  
  docs/  
    Doxyfile  
    README.md  
  cmake/  
    FindMyLib.cmake  
  .github/  
    workflows/  
      ci.yml  
  docker/  
    Dockerfile  
    docker-compose.yml  
  CMakeLists.txt  
  conanfile.txt  
  .gitignore  
  README.md
```

Complete CMakeLists.txt

```
cmake_minimum_required(VERSION 3.20)  
project(MyCProject C)  
  
# Set C standard to C23  
set(CMAKE_C_STANDARD 23)  
set(CMAKE_C_STANDARD_REQUIRED ON)  
  
# Compiler options  
if(CMAKE_C_COMPILER_ID STREQUAL "GNU" OR CMAKE_C_COMPILER_ID STREQUAL "Clang")  
  add_compile_options(-Wall -Wextra -Wpedantic -Werror)
```

```

# Enable sanitizers in debug mode
if(CMAKE_BUILD_TYPE STREQUAL "Debug")
    add_compile_options(-fsanitize=address -fsanitize=undefined)
    add_link_options(-fsanitize=address -fsanitize=undefined)
endif()
endif()

# Include directories
include_directories(include)

# Source files
set(SOURCES
    src/main.c
    src/utils.c
)

# Create executable
add_executable(${PROJECT_NAME} ${SOURCES})

# Link libraries
target_link_libraries(${PROJECT_NAME} m)

# Enable testing
enable_testing()

# Add test subdirectory
add_subdirectory(tests)

# Installation
install(TARGETS ${PROJECT_NAME} DESTINATION bin)
install(DIRECTORY include/ DESTINATION include)

# Packaging
set(CPACK_PACKAGE_NAME "my-c-project")
set(CPACK_PACKAGE_VERSION "1.0.0")
set(CPACK_PACKAGE_DESCRIPTION "A modern C project")
set(CPACK_GENERATOR "TGZ;DEB;RPM")
include(CPack)

```

Summary

Modern C development practices encompass a wide range of contemporary software engineering methodologies:

1. **DevOps and CI/CD:** Automated building, testing, and deployment pipelines

2. **Containerized Development:** Consistent environments with Docker and containerization
3. **Cross-compilation:** Building for multiple target platforms and architectures
4. **WebAssembly:** Compiling C to run in web browsers and WASM environments
5. **Cloud-native Applications:** Designing applications for cloud deployment and orchestration
6. **Supply Chain Security:** Managing dependencies and scanning for vulnerabilities
7. **Modern Testing:** Fuzzing, property-based testing, and mutation testing
8. **Documentation as Code:** Integrated documentation generation and maintenance

These practices enable C developers to build robust, secure, and maintainable applications while leveraging modern development tools and methodologies. As C continues to evolve with new standards like C23, these practices become increasingly important for professional C development.

Advanced Preprocessor

Introduction

The C preprocessor is a powerful text processing tool that operates before the actual compilation begins. While often overlooked, mastering advanced preprocessor techniques can lead to more maintainable, efficient, and feature-rich C code. Modern C standards have enhanced preprocessor capabilities, making it even more versatile for metaprogramming and code generation.

This chapter explores advanced preprocessor features including macro definitions, conditional compilation, pragma directives, and token manipulation techniques that enable sophisticated code generation and platform-specific optimizations.

Macro Definitions and Function-like Macros

Macros are the most fundamental preprocessor feature, allowing for text substitution and code generation.

Advanced Macro Techniques

```
#include <stdio.h>
#include <stdlib.h>

// Stringification operator (#)
#define STRINGIFY(x) #x
#define PRINT_VAR(var) printf(#var " = %d\n", var)

// Token pasting operator (##)
#define CONCAT(a, b) a##b
#define DECLARE_VAR(type, name) type CONCAT(my_, name)

// Variadic macros (C99)
#define DEBUG_PRINT(fmt, ...) \
    fprintf(stderr, "[DEBUG] %s:%d: " fmt "\n", __FILE__, __LINE__, ##__VA_ARGS__)

#define LOG(level, fmt, ...) \
    printf("[ " #level " ] " fmt "\n", ##__VA_ARGS__)

// Macro overloading based on number of arguments
```

```

#define GET_MACRO(_1, _2, _3, NAME, ...) NAME
#define ADD2(a, b) ((a) + (b))
#define ADD3(a, b, c) ((a) + (b) + (c))
#define ADD(...) GET_MACRO(__VA_ARGS__, ADD3, ADD2)(__VA_ARGS__)

int main() {
    int x = 42;

    // Stringification
    printf("Value of x: %s\n", STRINGIFY(x));
    PRINT_VAR(x);

    // Token pasting
    DECLARE_VAR(int, counter) = 100;
    printf("my_counter = %d\n", my_counter);

    // Variadic macros
    DEBUG_PRINT("This is a debug message");
    DEBUG_PRINT("Value of x is %d", x);
    LOG(INFO, "Application started");
    LOG(ERROR, "Failed to open file: %s", "config.txt");

    // Macro overloading
    printf("ADD(1, 2) = %d\n", ADD(1, 2));
    printf("ADD(1, 2, 3) = %d\n", ADD(1, 2, 3));

    return 0;
}

```

Generic Programming with Macros

```

#include <stdio.h>
#include <stdlib.h>
#include <string.h>

// Generic swap macro
#define SWAP(a, b) do { \
    typeof(a) temp = (a); \
    (a) = (b); \
    (b) = temp; \
} while(0)

// Generic maximum macro
#define MAX(a, b) ({ \
    typeof(a) _a = (a); \

```

```

    typeof(b) _b = (b); \
    _a > _b ? _a : _b; \
})

// Generic container_of macro (similar to Linux kernel)
#define container_of(ptr, type, member) ({ \
    const typeof(((type *)0)->member) * __mptr = (ptr); \
    (type *)((char *)__mptr - offsetof(type, member)); \
})

#define offsetof(type, member) ((size_t) &((type *)0)->member)

typedef struct {
    int id;
    char name[50];
    double value;
} Item;

int main() {
    // Generic swap
    int x = 10, y = 20;
    printf("Before swap: x=%d, y=%d\n", x, y);
    SWAP(x, y);
    printf("After swap: x=%d, y=%d\n", x, y);

    double a = 3.14, b = 2.71;
    printf("Before swap: a=%.2f, b=%.2f\n", a, b);
    SWAP(a, b);
    printf("After swap: a=%.2f, b=%.2f\n", a, b);

    // Generic maximum
    printf("MAX(10, 20) = %d\n", MAX(10, 20));
    printf("MAX(3.14, 2.71) = %.2f\n", MAX(3.14, 2.71));

    // Container of example
    Item item = {1, "Test Item", 99.99};
    char *name_ptr = item.name;

    // Get the container from a member pointer
    Item *container = container_of(name_ptr, Item, name);
    printf("Container ID: %d, Name: %s, Value: %.2f\n",
        container->id, container->name, container->value);

    return 0;
}

```


Conditional Compilation

Conditional compilation allows you to include or exclude code based on preprocessor conditions, enabling platform-specific code and feature toggles.

Platform-Specific Compilation

```
#include <stdio.h>

// Platform detection
#if defined(_WIN32) || defined(_WIN64)
    #define PLATFORM_WINDOWS
    #include <windows.h>
#elif defined(__linux__)
    #define PLATFORM_LINUX
    #include <unistd.h>
#elif defined(__APPLE__)
    #define PLATFORM_MACOS
    #include <unistd.h>
#endif

// Compiler detection
#if defined(__GNUC__)
    #define COMPILER_GCC
    #define NORETURN __attribute__((noreturn))
    #define UNUSED __attribute__((unused))
#elif defined(_MSC_VER)
    #define COMPILER_MSVC
    #define NORETURN __declspec(noreturn)
    #define UNUSED
#endif

// Feature detection
#if __STDC_VERSION__ >= 202311L
    #define C23_AVAILABLE
    #define HAS_TYPEOF typeof
#elif __STDC_VERSION__ >= 201112L
    #define C11_AVAILABLE
    #define HAS_TYPEOF __typeof__
#else
    #define C99_AVAILABLE
#endif

// Debug vs Release builds
#ifdef DEBUG
    #define DBG_PRINT(fmt, ...) fprintf(stderr, "[DEBUG] " fmt "\n", ##__VA_ARGS__)
```

```

#define ASSERT(condition) do { \
    if (!(condition)) { \
        fprintf(stderr, "Assertion failed: %s at %s:%d\n", \
            #condition, __FILE__, __LINE__); \
        abort(); \
    } \
} while(0)
#else
#define DBG_PRINT(fmt, ...) do {} while(0)
#define ASSERT(condition) do {} while(0)
#endif

// Feature toggles
#define FEATURE_NETWORKING 1
#define FEATURE_DATABASE 0

#if FEATURE_NETWORKING
#include <sys/socket.h>
void init_networking(void) {
    printf("Networking initialized\n");
}
#else
void init_networking(void) {
    printf("Networking disabled\n");
}
#endif

#if FEATURE_DATABASE
#include <sqlite3.h>
void init_database(void) {
    printf("Database initialized\n");
}
#else
void init_database(void) {
    printf("Database disabled\n");
}
#endif

int main() {
    printf("Platform: ");
    #if defined(PLATFORM_WINDOWS)
        printf("Windows\n");
    #elif defined(PLATFORM_LINUX)
        printf("Linux\n");
    #elif defined(PLATFORM_MACOS)
        printf("macOS\n");
    #else

```

```

    printf("Unknown\n");
#endif

    printf("Compiler: ");
#ifdef COMPILER_GCC
    printf("GCC\n");
#elif defined(COMPILER_MSVC)
    printf("MSVC\n");
#else
    printf("Unknown\n");
#endif

    printf("C Standard: ");
#ifdef C23_AVAILABLE
    printf("C23\n");
#elif defined(C11_AVAILABLE)
    printf("C11\n");
#else
    printf("C99 or earlier\n");
#endif

    // Debug printing
    DBG_PRINT("Debug message");
    DBG_PRINT("Value of x: %d", 42);

    // Assertions
    int x = 10;
    ASSERT(x > 0);

    // Feature initialization
    init_networking();
    init_database();

    return 0;
}

```

Version-Specific Code

```

#include <stdio.h>

// Version checking macros
#define VERSION_CHECK(major, minor, patch) \
    ((major) * 10000 + (minor) * 100 + (patch))

#define LIBRARY_VERSION VERSION_CHECK(2, 3, 1)

```

```

// API version compatibility
#if LIBRARY_VERSION >= VERSION_CHECK(2, 0, 0)
    #define NEW_API_AVAILABLE
    void new_api_function(void) {
        printf("Using new API\n");
    }
#endif

#if LIBRARY_VERSION >= VERSION_CHECK(2, 3, 0)
    #define ENHANCED_API_AVAILABLE
    void enhanced_api_function(void) {
        printf("Using enhanced API\n");
    }
#endif

// Backward compatibility
#if LIBRARY_VERSION < VERSION_CHECK(2, 0, 0)
    void new_api_function(void) {
        printf("Using legacy implementation\n");
    }
#endif

int main() {
    printf("Library version: %d.%d.%d\n",
        LIBRARY_VERSION / 10000,
        (LIBRARY_VERSION % 10000) / 100,
        LIBRARY_VERSION % 100);

    #ifdef NEW_API_AVAILABLE
        new_api_function();
    #endif

    #ifdef ENHANCED_API_AVAILABLE
        enhanced_api_function();
    #endif

    return 0;
}

```

Pragma Directives

Pragma directives provide compiler-specific instructions and hints that can optimize code generation and control compilation behavior.

Common Pragma Directives

```
#include <stdio.h>
#include <stdlib.h>

// Region folding (supported by many IDEs)
#pragma region Utility Functions

static inline int max(int a, int b) {
    return a > b ? a : b;
}

static inline int min(int a, int b) {
    return a < b ? a : b;
}

#pragma endregion

// Optimization hints
#pragma GCC push_options
#pragma GCC optimize("O3")

// Performance-critical function
double compute_pi(int iterations) {
    double pi = 0.0;
    double denominator = 1.0;

    for (int i = 0; i < iterations; i++) {
        if (i % 2 == 0) {
            pi += 4.0 / denominator;
        } else {
            pi -= 4.0 / denominator;
        }
        denominator += 2.0;
    }

    return pi;
}

#pragma GCC pop_options

// Warning control
#pragma GCC diagnostic push
#pragma GCC diagnostic ignored "-Wunused-variable"

void example_with_warnings(void) {
```

```

    int unused_var = 42; // This warning is suppressed
    printf("Function with suppressed warnings\n");
}

#pragma GCC diagnostic pop

// Structure packing
#pragma pack(push, 1)
typedef struct {
    char a;        // 1 byte
    int b;         // 4 bytes (normally padded to offset 4)
    short c;       // 2 bytes (normally padded to offset 8)
} PackedStruct; // Total size: 7 bytes instead of 12
#pragma pack(pop)

// Once-only inclusion (alternative to include guards)
#pragma once

// Message directives (compiler-specific)
#pragma message("Compiling with advanced preprocessor features")

int main() {
    printf("Max of 10 and 20: %d\n", max(10, 20));
    printf("Min of 10 and 20: %d\n", min(10, 20));

    double pi = compute_pi(1000000);
    printf("Computed PI: %.10f\n", pi);

    example_with_warnings();

    printf("Size of PackedStruct: %zu bytes\n", sizeof(PackedStruct));

    return 0;
}

```

Compiler-Specific Pragmas

```

#include <stdio.h>

// GCC-specific pragmas
#ifdef __GNUC__
    #pragma GCC diagnostic push
    #pragma GCC diagnostic ignored "-Wdeprecated-declarations"

    // Function inlining hints

```

```

#pragma GCC push_options
#pragma GCC optimize("inline-functions-called-once")

static inline void optimized_function(void) {
    printf("This function may be inlined\n");
}

#pragma GCC pop_options
#pragma GCC diagnostic pop
#endif

// MSVC-specific pragmas
#ifdef _MSC_VER
#pragma warning(push)
#pragma warning(disable: 4996) // Disable deprecated function warnings

// Function inlining
#pragma inline_depth(255)
#pragma inline_recursion(on)

static __inline void msvc_optimized_function(void) {
    printf("MSVC optimized function\n");
}

#pragma warning(pop)
#endif

// Clang-specific pragmas
#ifdef __clang__
#pragma clang diagnostic push
#pragma clang diagnostic ignored "-Wdeprecated-declarations"

// Loop optimization hints
#pragma clang loop vectorize(enable)

void clang_optimized_loop(int *array, int size) {
    for (int i = 0; i < size; i++) {
        array[i] *= 2;
    }
}

#pragma clang diagnostic pop
#endif

int main() {
#ifdef __GNUC__
    optimized_function();

```

```

#endif

#ifdef _MSC_VER
    msvc_optimized_function();
#endif

#ifdef __clang__
    int array[10] = {1, 2, 3, 4, 5, 6, 7, 8, 9, 10};
    clang_optimized_loop(array, 10);
    printf("Array[0] after optimization: %d\n", array[0]);
#endif

    return 0;
}

```

Predefined Macros

C provides numerous predefined macros that offer information about the compilation environment.

System and Compiler Information

```

#include <stdio.h>
#include <time.h>

void print_compilation_info(void) {
    printf("=== Compilation Information ===\n");

    // Standard version
    printf("Standard version: %ld\n", __STDC_VERSION__);
#ifdef __STDC_VERSION__
    #if __STDC_VERSION__ >= 202311L
        printf("C Standard: C23\n");
    #elif __STDC_VERSION__ >= 201112L
        printf("C Standard: C11\n");
    #elif __STDC_VERSION__ >= 199901L
        printf("C Standard: C99\n");
    #else
        printf("C Standard: C89/C90\n");
    #endif
#endif
#endif

    // Compiler information
#ifdef __GNUC__

```



```

    printf("Compiler: GCC %d.%d.%d\n", __GNUC__, __GNUC_MINOR__, __GNUC_PATCHLEVEL__);
#endif

#ifdef __clang__
    printf("Compiler: Clang %d.%d.%d\n", __clang_major__, __clang_minor__, __clang_patchleve
#endif

#ifdef _MSC_VER
    printf("Compiler: MSVC %d\n", _MSC_VER);
#endif

    // Platform information
#ifdef __x86_64__
    printf("Architecture: x86_64\n");
#elif defined(__i386__)
    printf("Architecture: x86\n");
#elif defined(__aarch64__)
    printf("Architecture: ARM64\n");
#elif defined(__arm__)
    printf("Architecture: ARM\n");
#endif

#ifdef __LP64__
    printf("Data model: LP64\n");
#endif

    // Date and time
    printf("Compilation date: %s\n", __DATE__);
    printf("Compilation time: %s\n", __TIME__);

    // File and line information
    printf("Current file: %s\n", __FILE__);
    printf("Current line: %d\n", __LINE__);
    printf("Function name: %s\n", __func__);

    printf("\n");
}

// Using predefined macros for debugging
#ifdef DEBUG
    #define TRACE() printf("TRACE: %s() in %s:%d\n", __func__, __FILE__, __LINE__)
#else
    #define TRACE() do {} while(0)
#endif

void function_a(void) {
    TRACE();

```

```

    printf("Inside function_a\n");
}

void function_b(void) {
    TRACE();
    printf("Inside function_b\n");
    function_a();
}

int main() {
    print_compilation_info();

    function_b();

    // Counter macro (implementation-defined)
    printf("This is compilation #%d\n", __COUNTER__);
    printf("This is compilation #%d\n", __COUNTER__);

    return 0;
}

```

Token Manipulation and Advanced Techniques

Advanced preprocessor techniques involve complex token manipulation and metaprogramming.

Deferred Macro Expansion

```

#include <stdio.h>

// Deferred expansion technique
#define EMPTY()
#define DEFER(id) id EMPTY()
#define OBSTRUCT(...) __VA_ARGS__ DEFER(EMPTY)()
#define EXPAND(...) __VA_ARGS__

#define EVAL(...) EVAL1(EVAL1(EVAL1(__VA_ARGS__)))
#define EVAL1(...) EVAL2(EVAL2(EVAL2(__VA_ARGS__)))
#define EVAL2(...) EVAL3(EVAL3(EVAL3(__VA_ARGS__)))
#define EVAL3(...) EVAL4(EVAL4(EVAL4(__VA_ARGS__)))
#define EVAL4(...) __VA_ARGS__

// Recursive macro example
#define WHEN(c) IF_##c
#define IF_1(t, f) t

```

```

#define IF_0(t, f) f

#define COMMA() ,
#define SEMICOLON() ;

#define REPEAT_0(f, ...)
#define REPEAT_1(f, x, ...) f(x) REPEAT_0(f, __VA_ARGS__)
#define REPEAT_2(f, x, ...) f(x) REPEAT_1(f, __VA_ARGS__)
#define REPEAT_3(f, x, ...) f(x) REPEAT_2(f, __VA_ARGS__)
#define REPEAT_4(f, x, ...) f(x) REPEAT_3(f, __VA_ARGS__)

#define REPEAT_N(n, f, ...) EVAL(REPEAT_##n(f, __VA_ARGS__))

#define PRINT_ARG(x) printf("%s = %d\n", #x, x);

int main() {
    int a = 1, b = 2, c = 3, d = 4;

    printf("Repeating PRINT_ARG macro:\n");
    REPEAT_N(4, PRINT_ARG, a, b, c, d)

    return 0;
}

```

X-Macros Pattern

```

#include <stdio.h>
#include <string.h>

// X-Macros for enumerations and string tables
#define COLOR_TABLE \
    X(RED, "red", 0xFF0000) \
    X(GREEN, "green", 0x00FF00) \
    X(BLUE, "blue", 0x0000FF) \
    X(YELLOW, "yellow", 0xFFFF00) \
    X(PURPLE, "purple", 0xFF00FF) \
    X(CYAN, "cyan", 0x00FFFF) \
    X(WHITE, "white", 0xFFFFFF) \
    X(BLACK, "black", 0x000000)

// Generate enumeration
#define X(name, string, value) COLOR_##name,
typedef enum {
    COLOR_TABLE
    COLOR_COUNT
}

```

```

} Color;
#undef X

// Generate string array
#define X(name, string, value) string,
static const char *color_strings[] = {
    COLOR_TABLE
};
#undef X

// Generate value array
#define X(name, string, value) value,
static const int color_values[] = {
    COLOR_TABLE
};
#undef X

// Generate function to get string
const char *color_to_string(Color color) {
    if (color >= 0 && color < COLOR_COUNT) {
        return color_strings[color];
    }
    return "unknown";
}

// Generate function to get value
int color_to_value(Color color) {
    if (color >= 0 && color < COLOR_COUNT) {
        return color_values[color];
    }
    return 0;
}

// Generate lookup function
Color string_to_color(const char *str) {
    for (int i = 0; i < COLOR_COUNT; i++) {
        if (strcmp(str, color_strings[i]) == 0) {
            return (Color)i;
        }
    }
    return COLOR_COUNT; // Invalid color
}

int main() {
    printf("Color enumeration:\n");
    for (int i = 0; i < COLOR_COUNT; i++) {

```

```

        printf("  %d: %s (0x%06X)\n", i, color_to_string((Color)i), color_to_value((Color)i)
    }

    printf("\nString to color lookup:\n");
    const char *test_colors[] = {"red", "blue", "purple", "invalid"};
    for (int i = 0; i < 4; i++) {
        Color c = string_to_color(test_colors[i]);
        if (c < COLOR_COUNT) {
            printf("  %s -> %s (0x%06X)\n", test_colors[i], color_to_string(c), color_to_val
        } else {
            printf("  %s -> invalid color\n", test_colors[i]);
        }
    }

    return 0;
}

```

C23 Enhanced Preprocessor Features

C23 introduces several new preprocessor features that enhance metaprogramming capabilities.

#elifdef and #elifndef

```

#include <stdio.h>

// Define some macros for testing
#define DEBUG_MODE
// #define RELEASE_MODE
// #define TESTING_MODE

int main() {
    printf("Preprocessor conditional compilation:\n");

    #if defined(DEBUG_MODE)
        printf("  Debug mode is enabled\n");
    #elifdef RELEASE_MODE
        printf("  Release mode is enabled\n");
    #elifndef TESTING_MODE
        printf("  Testing mode is not defined\n");
    #else
        printf("  Unknown mode\n");
    #endif

    return 0;
}

```

```
}
```

#warning Directive

```
#include <stdio.h>

// Compiler version check
#if __STDC_VERSION__ < 202311L
    #warning "This code is designed for C23 or later"
#endif

// Platform-specific warnings
#ifdef __APPLE__
    #warning "This code may need adjustments for macOS"
#endif

// Feature availability warnings
#ifndef __STDC_LIB_EXT1__
    #warning "Bounds-checking functions not available"
#endif

int main() {
    printf("Program running with warnings (if any)\n");

#ifdef __GNUC__
    #warning "Consider using Clang for better C23 support"
#endif

    return 0;
}
```

Digit Separators in Preprocessor

```
#include <stdio.h>

// Using digit separators in preprocessor expressions
#define ONE_MILLION 1'000'000
#define MAX_BUFFER_SIZE 10'000
#define BINARY_PATTERN 0b1010'0101'1100'1111

int main() {
    printf("One million: %d\n", ONE_MILLION);
    printf("Max buffer size: %d\n", MAX_BUFFER_SIZE);
    printf("Binary pattern: 0x%X\n", BINARY_PATTERN);
}
```

```

    // Using in preprocessor conditionals
    #if ONE_MILLION > 500'000
        printf("One million is greater than 500 thousand\n");
    #endif

    return 0;
}

```

Practical Examples

Configuration System Using Preprocessor

```

#include <stdio.h>

// Configuration system using preprocessor
// Define default values
#ifndef CONFIG_MAX_CONNECTIONS
    #define CONFIG_MAX_CONNECTIONS 100
#endif

#ifndef CONFIG_TIMEOUT
    #define CONFIG_TIMEOUT 30
#endif

#ifndef CONFIG_LOG_LEVEL
    #define CONFIG_LOG_LEVEL 2
#endif

// Feature flags
#define FEATURE_NETWORKING 1
#define FEATURE_DATABASE 1
#define FEATURE_SECURITY 0

// Platform-specific configurations
#ifdef _WIN32
    #define PATH_SEPARATOR '\\'
    #define DEFAULT_CONFIG_FILE "config.ini"
#else
    #define PATH_SEPARATOR '/'
    #define DEFAULT_CONFIG_FILE "/etc/myapp/config.ini"
#endif

// Debug configuration
#ifdef DEBUG

```

```

#define LOG_LEVEL CONFIG_LOG_LEVEL
#define ENABLE_ASSERTIONS 1
#define VERBOSE_LOGGING 1
#else
#define LOG_LEVEL (CONFIG_LOG_LEVEL - 1)
#define ENABLE_ASSERTIONS 0
#define VERBOSE_LOGGING 0
#endif

// Conditional compilation based on features
#if FEATURE_NETWORKING
#include <sys/socket.h>
#define NETWORK_BUFFER_SIZE 8192

void init_networking(void) {
    printf("Networking initialized with buffer size: %d\n", NETWORK_BUFFER_SIZE);
}
#else
void init_networking(void) {
    printf("Networking disabled\n");
}
#endif

#if FEATURE_DATABASE
#include <sqlite3.h>
#define DB_CONNECTION_POOL_SIZE 10

void init_database(void) {
    printf("Database initialized with connection pool size: %d\n", DB_CONNECTION_POOL_SIZE);
}
#else
void init_database(void) {
    printf("Database disabled\n");
}
#endif

#if FEATURE_SECURITY
#include <openssl/ssl.h>
#define SSL_ENABLED 1

void init_security(void) {
    printf("Security initialized with SSL enabled\n");
}
#else
#define SSL_ENABLED 0
void init_security(void) {
    printf("Security disabled\n");
}

```



```

    }
#endif

void print_configuration(void) {
    printf("=== Application Configuration ===\n");
    printf("Max connections: %d\n", CONFIG_MAX_CONNECTIONS);
    printf("Timeout: %d seconds\n", CONFIG_TIMEOUT);
    printf("Log level: %d\n", LOG_LEVEL);
    printf("Path separator: '%c'\n", PATH_SEPARATOR);
    printf("Default config file: %s\n", DEFAULT_CONFIG_FILE);
    printf("Assertions: %s\n", ENABLE_ASSERTIONS ? "enabled" : "disabled");
    printf("Verbose logging: %s\n", VERBOSE_LOGGING ? "enabled" : "disabled");
    printf("\n");

    printf("=== Feature Status ===\n");
    printf("Networking: %s\n", FEATURE_NETWORKING ? "enabled" : "disabled");
    printf("Database: %s\n", FEATURE_DATABASE ? "enabled" : "disabled");
    printf("Security: %s\n", FEATURE_SECURITY ? "enabled" : "disabled");
    printf("SSL: %s\n", SSL_ENABLED ? "enabled" : "disabled");
    printf("\n");
}

int main() {
    print_configuration();

    init_networking();
    init_database();
    init_security();

    return 0;
}

```

Advanced Error Handling with Preprocessor

```

#include <stdio.h>
#include <stdlib.h>
#include <errno.h>
#include <string.h>

// Advanced error handling macros
#define STRINGIFY(x) #x
#define TOSTRING(x) STRINGIFY(x)

// Error checking with automatic cleanup
#define CHECK_ERRNO(call) do { \

```

```

    errno = 0; \
    call; \
    if (errno != 0) { \
        fprintf(stderr, "Error in %s at %s:%d: %s\n", \
            #call, __FILE__, __LINE__, strerror(errno)); \
        exit(EXIT_FAILURE); \
    } \
} while(0)

// Resource management with automatic cleanup
#define CLEANUP_VAR(type, name, cleanup_func) \
    type name = NULL; \
    defer_cleanup(cleanup_func, &name)

// Defer mechanism (simplified)
typedef struct defer_node {
    void (*func)(void*);
    void *arg;
    struct defer_node *next;
} defer_node_t;

static defer_node_t *defer_stack = NULL;

void defer_cleanup(void (*func)(void*), void *arg) {
    defer_node_t *node = malloc(sizeof(defer_node_t));
    node->func = func;
    node->arg = arg;
    node->next = defer_stack;
    defer_stack = node;
}

void execute_deferred(void) {
    while (defer_stack) {
        defer_node_t *node = defer_stack;
        defer_stack = node->next;
        node->func(node->arg);
        free(node);
    }
}

#define defer(func, arg) defer_cleanup(func, arg)

// Cleanup functions
void free_ptr(void *ptr) {
    free(*(void**)ptr);
}

```

```

void fclose_ptr(void *ptr) {
    if (*(FILE**)ptr) {
        fclose(*(FILE**)ptr);
    }
}

int main() {
    printf("Advanced error handling with preprocessor\n");

    // File operations with automatic error checking
    CLEANUP_VAR(FILE*, fp, fclose_ptr);
    fp = fopen("test.txt", "w");
    if (!fp) {
        perror("fopen");
        exit(EXIT_FAILURE);
    }

    CHECK_ERRNO(fprintf(fp, "Hello, World!\n"));
    CHECK_ERRNO(fflush(fp));

    // Memory allocation with automatic cleanup
    CLEANUP_VAR(char*, buffer, free_ptr);
    buffer = malloc(100);
    if (!buffer) {
        fprintf(stderr, "Memory allocation failed\n");
        exit(EXIT_FAILURE);
    }

    strcpy(buffer, "This is a test buffer");
    printf("Buffer content: %s\n", buffer);

    // Automatic cleanup when scope ends
    execute_deferred();

    return 0;
}

```

Summary

Advanced preprocessor techniques provide powerful metaprogramming capabilities in C:

1. **Macro Definitions:** Stringification, token pasting, variadic macros, and generic programming
2. **Conditional Compilation:** Platform-specific code, feature toggles, and version compatibility
3. **Pragma Directives:** Compiler hints, optimization control, and warning management

4. **Predefined Macros:** System information, compilation context, and debugging support
5. **Token Manipulation:** Deferred expansion, X-Macros, and complex metaprogramming
6. **C23 Enhancements:** Improved conditional directives and digit separators

Mastering these advanced preprocessor techniques enables C developers to write more maintainable, efficient, and feature-rich code while leveraging the full power of the C preprocessor for metaprogramming and code generation.

Module 9: Advanced C Features and Modern Standards

Overview

This module explores advanced features introduced in modern C standards (C99, C11, C17, and C23) as well as contemporary development practices. You'll learn about variable-length arrays, designated initializers, static assertions, thread-local storage, and many other features that make modern C more powerful and expressive. The module also covers modern development practices including DevOps integration, containerized development, and cloud-native C programming.

Learning Objectives

By the end of this module, you will be able to:

- Understand and apply C99 features including variable-length arrays and designated initializers
- Utilize C11 features such as static assertions and thread-local storage
- Implement C17/C23 features including typeof operators and binary literals
- Apply modern C development practices and tools
- Understand contemporary C ecosystem including WebAssembly and cloud-native development
- Work with advanced preprocessor features
- Implement generic programming techniques using modern C features
- Apply best practices for modern C development

Chapters

1. **C99 Features** - Variable-length arrays, designated initializers, compound literals, restrict keyword, inline functions
2. **C11 Features** - Static assertions, anonymous structures/unions, thread-local storage, bounds-checking functions
3. **C17 and C23 Features** - typeof operators, binary literals, enhanced preprocessor directives, attributes
4. **Modern C Development Practices** - DevOps integration, containerized development, WebAssembly, cloud-native C
5. **Advanced Preprocessor** - Macro definitions, conditional compilation, pragma directives, token manipulation

Key Concepts Covered

- Variable-length arrays (VLAs)
- Designated initializers
- Compound literals
- restrict keyword
- inline functions
- Static assertions (`_Static_assert`)
- Anonymous structures and unions
- Thread-local storage (`_Thread_local`)
- Bounds-checking functions (Annex K)
- `typeof` and `typeof_unqual` operators
- Binary integer constants (0b1010)
- Enhanced preprocessor directives (`#elifdef`, `#elifndef`, `#warning`, `#embed`)
- Attribute syntax (`[]`)
- Generic programming techniques
- DevOps and CI/CD for C projects
- Containerized C development with Docker
- WebAssembly compilation with Emscripten
- Cloud-native C applications
- Modern testing practices (fuzzing, property-based testing)
- Supply chain security and dependency management

Prerequisites

- Completion of Module 1: Foundations and Environment Setup
- Completion of Module 2: Data Types and Variables
- Completion of Module 3: Control Flow
- Completion of Module 4: Functions and Modular Programming
- Completion of Module 5: Arrays and Strings
- Completion of Module 6: Pointers and Memory Management
- Completion of Module 7: Structures and User-Defined Types
- Completion of Module 8: File I/O and System Programming
- Understanding of basic programming concepts
- Familiarity with C language fundamentals

Tools and Technologies

- Modern C compiler (GCC 14+, Clang 18+, MSVC 2022)
- Build systems (CMake, Meson)
- Package managers (Conan, vcpkg)
- Containerization tools (Docker)
- WebAssembly toolchains (Emscripten)
- CI/CD platforms (GitHub Actions, GitLab CI)
- Testing frameworks (Unity, CMocka)

- Static analysis tools (clang-analyzer, cppcheck)
- Fuzzing tools (AFL, libFuzzer)

Estimated Time to Complete

- Reading: 12-15 hours
- Exercises: 15-20 hours
- Projects: 10-12 hours

Assessment

- Complete all chapter exercises
- Successfully implement programs using modern C features
- Demonstrate understanding of contemporary C development practices
- Pass modern C features and practices quiz
- Submit a comprehensive project using modern C standards

Next Module

After completing this module, proceed to **Module 10: Data Structures and Algorithms** to learn about implementing classic data structures and algorithms in C.

Additional Resources

- “Modern C” by Jens Gustedt
- C standard documentation (ISO/IEC 9899:2024 for C23)
- Compiler documentation (GCC, Clang, MSVC)
- WebAssembly documentation
- Docker documentation
- CI/CD platform documentation
- Modern C development blogs and resources

Data Structures

Linear Data Structures

Introduction

Linear data structures are fundamental building blocks in computer science that organize data elements in a sequential manner. These structures include dynamic arrays, linked lists, stacks, queues, and their variations. Understanding linear data structures is crucial for efficient programming, as they form the foundation for more complex data structures and algorithms.

In this chapter, we'll explore the implementation and usage of various linear data structures in C, focusing on dynamic memory management, performance characteristics, and practical applications.

Dynamic Arrays (Vectors)

Dynamic arrays, also known as vectors, are resizable arrays that can grow or shrink during program execution. They provide the benefits of arrays (random access, cache locality) with the flexibility of dynamic sizing.

Basic Dynamic Array Implementation

```
#include <stdio.h>
#include <stdlib.h>
#include <string.h>

typedef struct {
    int *data;
    size_t size;
    size_t capacity;
} DynamicArray;

// Initialize dynamic array
DynamicArray* create_array(size_t initial_capacity) {
    DynamicArray *arr = malloc(sizeof(DynamicArray));
    if (arr == NULL) return NULL;

    arr->data = malloc(initial_capacity * sizeof(int));
    if (arr->data == NULL) {
        free(arr);
        return NULL;
    }
}
```

```

    }

    arr->size = 0;
    arr->capacity = initial_capacity;
    return arr;
}

// Free dynamic array
void free_array(DynamicArray *arr) {
    if (arr != NULL) {
        free(arr->data);
        free(arr);
    }
}

// Add element to end of array
void push_back(DynamicArray *arr, int value) {
    if (arr->size >= arr->capacity) {
        // Double the capacity
        size_t new_capacity = arr->capacity * 2;
        int *new_data = realloc(arr->data, new_capacity * sizeof(int));
        if (new_data == NULL) return; // Handle allocation failure

        arr->data = new_data;
        arr->capacity = new_capacity;
    }

    arr->data[arr->size] = value;
    arr->size++;
}

// Remove element from end of array
int pop_back(DynamicArray *arr) {
    if (arr->size == 0) return 0; // Handle empty array

    arr->size--;
    return arr->data[arr->size];
}

// Access element at index
int at(DynamicArray *arr, size_t index) {
    if (index >= arr->size) return 0; // Handle out of bounds
    return arr->data[index];
}

// Set element at index

```

```

void set(DynamicArray *arr, size_t index, int value) {
    if (index >= arr->size) return; // Handle out of bounds
    arr->data[index] = value;
}

// Insert element at index
void insert(DynamicArray *arr, size_t index, int value) {
    if (index > arr->size) return; // Handle invalid index

    if (arr->size >= arr->capacity) {
        size_t new_capacity = arr->capacity * 2;
        int *new_data = realloc(arr->data, new_capacity * sizeof(int));
        if (new_data == NULL) return;

        arr->data = new_data;
        arr->capacity = new_capacity;
    }

    // Shift elements to the right
    for (size_t i = arr->size; i > index; i--) {
        arr->data[i] = arr->data[i - 1];
    }

    arr->data[index] = value;
    arr->size++;
}

// Remove element at index
void erase(DynamicArray *arr, size_t index) {
    if (index >= arr->size) return; // Handle invalid index

    // Shift elements to the left
    for (size_t i = index; i < arr->size - 1; i++) {
        arr->data[i] = arr->data[i + 1];
    }

    arr->size--;
}

// Print array contents
void print_array(DynamicArray *arr) {
    printf("Array (size=%zu, capacity=%zu): ", arr->size, arr->capacity);
    for (size_t i = 0; i < arr->size; i++) {
        printf("%d ", arr->data[i]);
    }
    printf("\n");
}

```

```

int main() {
    DynamicArray *arr = create_array(2);
    if (arr == NULL) {
        printf("Failed to create array\n");
        return 1;
    }

    // Test basic operations
    push_back(arr, 10);
    push_back(arr, 20);
    print_array(arr);

    push_back(arr, 30); // This should trigger a resize
    print_array(arr);

    insert(arr, 1, 15);
    print_array(arr);

    set(arr, 0, 5);
    print_array(arr);

    printf("Element at index 2: %d\n", at(arr, 2));

    erase(arr, 1);
    print_array(arr);

    int popped = pop_back(arr);
    printf("Popped element: %d\n", popped);
    print_array(arr);

    free_array(arr);
    return 0;
}

```

Generic Dynamic Array

```

#include <stdio.h>
#include <stdlib.h>
#include <string.h>

typedef struct {
    void **data;
    size_t size;
    size_t capacity;
    size_t element_size;
}

```

```

} GenericArray;

// Create generic array
GenericArray* create_generic_array(size_t initial_capacity, size_t element_size) {
    GenericArray *arr = malloc(sizeof(GenericArray));
    if (arr == NULL) return NULL;

    arr->data = malloc(initial_capacity * sizeof(void*));
    if (arr->data == NULL) {
        free(arr);
        return NULL;
    }

    arr->size = 0;
    arr->capacity = initial_capacity;
    arr->element_size = element_size;
    return arr;
}

// Free generic array
void free_generic_array(GenericArray *arr) {
    if (arr != NULL) {
        for (size_t i = 0; i < arr->size; i++) {
            free(arr->data[i]);
        }
        free(arr->data);
        free(arr);
    }
}

// Add element to generic array
int push_generic(GenericArray *arr, const void *element) {
    if (arr->size >= arr->capacity) {
        size_t new_capacity = arr->capacity * 2;
        void **new_data = realloc(arr->data, new_capacity * sizeof(void*));
        if (new_data == NULL) return -1;

        arr->data = new_data;
        arr->capacity = new_capacity;
    }

    arr->data[arr->size] = malloc(arr->element_size);
    if (arr->data[arr->size] == NULL) return -1;

    memcpy(arr->data[arr->size], element, arr->element_size);
    arr->size++;
    return 0;
}

```

```

}

// Print generic array with custom print function
void print_generic_array(GenericArray *arr, void (*print_func)(const void*)) {
    printf("Generic Array (size=%zu, capacity=%zu): ", arr->size, arr->capacity);
    for (size_t i = 0; i < arr->size; i++) {
        print_func(arr->data[i]);
        printf(" ");
    }
    printf("\n");
}

// Example with strings
void print_string(const void *str) {
    printf("%s", (const char*)str);
}

int main() {
    // Create array of strings
    GenericArray *str_array = create_generic_array(2, sizeof(char*));
    if (str_array == NULL) {
        printf("Failed to create array\n");
        return 1;
    }

    const char *strings[] = {"Hello", "World", "Generic", "Array"};

    for (int i = 0; i < 4; i++) {
        push_generic(str_array, &strings[i]);
    }

    print_generic_array(str_array, print_string);

    free_generic_array(str_array);
    return 0;
}

```

Linked Lists

Linked lists are linear data structures where elements are stored in nodes, and each node contains a reference to the next node in the sequence. They provide efficient insertion and deletion operations but have slower access times compared to arrays.

Singly Linked List

```
#include <stdio.h>
#include <stdlib.h>

typedef struct Node {
    int data;
    struct Node *next;
} Node;

typedef struct {
    Node *head;
    size_t size;
} LinkedList;

// Create new node
Node* create_node(int data) {
    Node *node = malloc(sizeof(Node));
    if (node == NULL) return NULL;

    node->data = data;
    node->next = NULL;
    return node;
}

// Initialize linked list
LinkedList* create_list(void) {
    LinkedList *list = malloc(sizeof(LinkedList));
    if (list == NULL) return NULL;

    list->head = NULL;
    list->size = 0;
    return list;
}

// Free linked list
void free_list(LinkedList *list) {
    if (list == NULL) return;

    Node *current = list->head;
    while (current != NULL) {
        Node *next = current->next;
        free(current);
        current = next;
    }

    free(list);
}
```

```

}

// Insert at beginning
void push_front(LinkedList *list, int data) {
    Node *new_node = create_node(data);
    if (new_node == NULL) return;

    new_node->next = list->head;
    list->head = new_node;
    list->size++;
}

// Insert at end
void push_back(LinkedList *list, int data) {
    Node *new_node = create_node(data);
    if (new_node == NULL) return;

    if (list->head == NULL) {
        list->head = new_node;
    } else {
        Node *current = list->head;
        while (current->next != NULL) {
            current = current->next;
        }
        current->next = new_node;
    }
    list->size++;
}

// Remove from beginning
int pop_front(LinkedList *list) {
    if (list->head == NULL) return 0;

    Node *temp = list->head;
    int data = temp->data;
    list->head = list->head->next;
    free(temp);
    list->size--;

    return data;
}

// Find element
Node* find(LinkedList *list, int data) {
    Node *current = list->head;
    while (current != NULL) {
        if (current->data == data) {

```



```

        return current;
    }
    current = current->next;
}
return NULL;
}

// Insert after node
void insert_after(Node *node, int data) {
    if (node == NULL) return;

    Node *new_node = create_node(data);
    if (new_node == NULL) return;

    new_node->next = node->next;
    node->next = new_node;
}

// Delete node with specific data
void delete_node(LinkedList *list, int data) {
    if (list->head == NULL) return;

    // If head node holds the data
    if (list->head->data == data) {
        Node *temp = list->head;
        list->head = list->head->next;
        free(temp);
        list->size--;
        return;
    }

    // Search for the node to delete
    Node *current = list->head;
    while (current->next != NULL && current->next->data != data) {
        current = current->next;
    }

    // If node was found
    if (current->next != NULL) {
        Node *temp = current->next;
        current->next = current->next->next;
        free(temp);
        list->size--;
    }
}

// Print list

```

```

void print_list(LinkedList *list) {
    printf("List (size=%zu): ", list->size);
    Node *current = list->head;
    while (current != NULL) {
        printf("%d -> ", current->data);
        current = current->next;
    }
    printf("NULL\n");
}

int main() {
    LinkedList *list = create_list();
    if (list == NULL) {
        printf("Failed to create list\n");
        return 1;
    }

    // Test operations
    push_back(list, 10);
    push_back(list, 20);
    push_front(list, 5);
    print_list(list);

    insert_after(list->head, 7);
    print_list(list);

    Node *found = find(list, 20);
    if (found != NULL) {
        printf("Found node with data: %d\n", found->data);
    }

    delete_node(list, 7);
    print_list(list);

    int popped = pop_front(list);
    printf("Popped element: %d\n", popped);
    print_list(list);

    free_list(list);
    return 0;
}

```

Doubly Linked List

```
#include <stdio.h>
#include <stdlib.h>

typedef struct DNode {
    int data;
    struct DNode *next;
    struct DNode *prev;
} DNode;

typedef struct {
    DNode *head;
    DNode *tail;
    size_t size;
} DoublyLinkedList;

// Create new node
DNode* create_dnode(int data) {
    DNode *node = malloc(sizeof(DNode));
    if (node == NULL) return NULL;

    node->data = data;
    node->next = NULL;
    node->prev = NULL;
    return node;
}

// Initialize doubly linked list
DoublyLinkedList* create_dlist(void) {
    DoublyLinkedList *list = malloc(sizeof(DoublyLinkedList));
    if (list == NULL) return NULL;

    list->head = NULL;
    list->tail = NULL;
    list->size = 0;
    return list;
}

// Free doubly linked list
void free_dlist(DoublyLinkedList *list) {
    if (list == NULL) return;

    DNode *current = list->head;
    while (current != NULL) {
        DNode *next = current->next;
        free(current);
    }
```

```

        current = next;
    }

    free(list);
}

// Insert at beginning
void push_front_dlist(DoublyLinkedList *list, int data) {
    DNode *new_node = create_dnode(data);
    if (new_node == NULL) return;

    if (list->head == NULL) {
        list->head = list->tail = new_node;
    } else {
        new_node->next = list->head;
        list->head->prev = new_node;
        list->head = new_node;
    }
    list->size++;
}

// Insert at end
void push_back_dlist(DoublyLinkedList *list, int data) {
    DNode *new_node = create_dnode(data);
    if (new_node == NULL) return;

    if (list->tail == NULL) {
        list->head = list->tail = new_node;
    } else {
        new_node->prev = list->tail;
        list->tail->next = new_node;
        list->tail = new_node;
    }
    list->size++;
}

// Remove from beginning
int pop_front_dlist(DoublyLinkedList *list) {
    if (list->head == NULL) return 0;

    DNode *temp = list->head;
    int data = temp->data;

    if (list->head == list->tail) {
        list->head = list->tail = NULL;
    } else {
        list->head = list->head->next;
    }
}

```

```

        list->head->prev = NULL;
    }

    free(temp);
    list->size--;
    return data;
}

// Remove from end
int pop_back_dlist(DoublyLinkedList *list) {
    if (list->tail == NULL) return 0;

    DNode *temp = list->tail;
    int data = temp->data;

    if (list->head == list->tail) {
        list->head = list->tail = NULL;
    } else {
        list->tail = list->tail->prev;
        list->tail->next = NULL;
    }

    free(temp);
    list->size--;
    return data;
}

// Print list forward
void print_dlist_forward(DoublyLinkedList *list) {
    printf("Forward: ");
    DNode *current = list->head;
    while (current != NULL) {
        printf("%d ", current->data);
        current = current->next;
    }
    printf("\n");
}

// Print list backward
void print_dlist_backward(DoublyLinkedList *list) {
    printf("Backward: ");
    DNode *current = list->tail;
    while (current != NULL) {
        printf("%d ", current->data);
        current = current->prev;
    }
    printf("\n");
}

```

```

}

int main() {
    DoublyLinkedList *list = create_dlist();
    if (list == NULL) {
        printf("Failed to create list\n");
        return 1;
    }

    // Test operations
    push_back_dlist(list, 10);
    push_back_dlist(list, 20);
    push_front_dlist(list, 5);
    push_back_dlist(list, 30);

    print_dlist_forward(list);
    print_dlist_backward(list);

    int front = pop_front_dlist(list);
    printf("Popped from front: %d\n", front);

    int back = pop_back_dlist(list);
    printf("Popped from back: %d\n", back);

    print_dlist_forward(list);

    free_dlist(list);
    return 0;
}

```

Stacks

A stack is a Last-In-First-Out (LIFO) data structure where elements are added and removed from the same end, called the top.

Array-based Stack

```

#include <stdio.h>
#include <stdlib.h>
#include <stdbool.h>

#define STACK_SIZE 100

typedef struct {

```

```

    int data[STACK_SIZE];
    int top;
} ArrayStack;

// Initialize stack
void init_array_stack(ArrayStack *stack) {
    stack->top = -1;
}

// Check if stack is empty
bool is_empty_array(ArrayStack *stack) {
    return stack->top == -1;
}

// Check if stack is full
bool is_full_array(ArrayStack *stack) {
    return stack->top == STACK_SIZE - 1;
}

// Push element onto stack
bool push_array(ArrayStack *stack, int value) {
    if (is_full_array(stack)) {
        return false; // Stack overflow
    }

    stack->data[++stack->top] = value;
    return true;
}

// Pop element from stack
int pop_array(ArrayStack *stack) {
    if (is_empty_array(stack)) {
        return 0; // Stack underflow
    }

    return stack->data[stack->top--];
}

// Peek at top element
int peek_array(ArrayStack *stack) {
    if (is_empty_array(stack)) {
        return 0; // Stack is empty
    }

    return stack->data[stack->top];
}

```

```

// Print stack contents
void print_array_stack(ArrayStack *stack) {
    if (is_empty_array(stack)) {
        printf("Stack is empty\n");
        return;
    }

    printf("Stack (top to bottom): ");
    for (int i = stack->top; i >= 0; i--) {
        printf("%d ", stack->data[i]);
    }
    printf("\n");
}

int main() {
    ArrayStack stack;
    init_array_stack(&stack);

    // Test stack operations
    push_array(&stack, 10);
    push_array(&stack, 20);
    push_array(&stack, 30);

    print_array_stack(&stack);

    printf("Top element: %d\n", peek_array(&stack));

    int popped = pop_array(&stack);
    printf("Popped element: %d\n", popped);

    print_array_stack(&stack);

    return 0;
}

```

Linked List-based Stack

```

#include <stdio.h>
#include <stdlib.h>
#include <stdbool.h>

typedef struct StackNode {
    int data;
    struct StackNode *next;
} StackNode;

```



```

typedef struct {
    StackNode *top;
    size_t size;
} LinkedStack;

// Create new stack node
StackNode* create_stack_node(int data) {
    StackNode *node = malloc(sizeof(StackNode));
    if (node == NULL) return NULL;

    node->data = data;
    node->next = NULL;
    return node;
}

// Initialize stack
LinkedStack* create_linked_stack(void) {
    LinkedStack *stack = malloc(sizeof(LinkedStack));
    if (stack == NULL) return NULL;

    stack->top = NULL;
    stack->size = 0;
    return stack;
}

// Free stack
void free_linked_stack(LinkedStack *stack) {
    if (stack == NULL) return;

    StackNode *current = stack->top;
    while (current != NULL) {
        StackNode *next = current->next;
        free(current);
        current = next;
    }

    free(stack);
}

// Check if stack is empty
bool is_empty_linked(LinkedStack *stack) {
    return stack->top == NULL;
}

// Push element onto stack
bool push_linked(LinkedStack *stack, int value) {

```

```

    StackNode *new_node = create_stack_node(value);
    if (new_node == NULL) return false;

    new_node->next = stack->top;
    stack->top = new_node;
    stack->size++;

    return true;
}

// Pop element from stack
int pop_linked(LinkedStack *stack) {
    if (is_empty_linked(stack)) {
        return 0; // Stack underflow
    }

    StackNode *temp = stack->top;
    int data = temp->data;
    stack->top = stack->top->next;
    free(temp);
    stack->size--;

    return data;
}

// Peek at top element
int peek_linked(LinkedStack *stack) {
    if (is_empty_linked(stack)) {
        return 0; // Stack is empty
    }

    return stack->top->data;
}

// Print stack contents
void print_linked_stack(LinkedStack *stack) {
    if (is_empty_linked(stack)) {
        printf("Stack is empty\n");
        return;
    }

    printf("Stack (top to bottom): ");
    StackNode *current = stack->top;
    while (current != NULL) {
        printf("%d ", current->data);
        current = current->next;
    }
}

```

```

        printf("\n");
    }

    int main() {
        LinkedStack *stack = create_linked_stack();
        if (stack == NULL) {
            printf("Failed to create stack\n");
            return 1;
        }

        // Test stack operations
        push_linked(stack, 10);
        push_linked(stack, 20);
        push_linked(stack, 30);

        print_linked_stack(stack);

        printf("Top element: %d\n", peek_linked(stack));

        int popped = pop_linked(stack);
        printf("Popped element: %d\n", popped);

        print_linked_stack(stack);

        free_linked_stack(stack);
        return 0;
    }

```

Queues

A queue is a First-In-First-Out (FIFO) data structure where elements are added at the rear and removed from the front.

Circular Buffer Queue

```

#include <stdio.h>
#include <stdlib.h>
#include <stdbool.h>

#define QUEUE_SIZE 100

typedef struct {
    int data[QUEUE_SIZE];
    int front;

```

```

    int rear;
    int size;
} CircularQueue;

// Initialize queue
void init_circular_queue(CircularQueue *queue) {
    queue->front = 0;
    queue->rear = -1;
    queue->size = 0;
}

// Check if queue is empty
bool is_empty_queue(CircularQueue *queue) {
    return queue->size == 0;
}

// Check if queue is full
bool is_full_queue(CircularQueue *queue) {
    return queue->size == QUEUE_SIZE;
}

// Enqueue (add) element to queue
bool enqueue(CircularQueue *queue, int value) {
    if (is_full_queue(queue)) {
        return false; // Queue overflow
    }

    queue->rear = (queue->rear + 1) % QUEUE_SIZE;
    queue->data[queue->rear] = value;
    queue->size++;

    return true;
}

// Dequeue (remove) element from queue
int dequeue(CircularQueue *queue) {
    if (is_empty_queue(queue)) {
        return 0; // Queue underflow
    }

    int data = queue->data[queue->front];
    queue->front = (queue->front + 1) % QUEUE_SIZE;
    queue->size--;

    return data;
}

```

```

// Peek at front element
int front_queue(CircularQueue *queue) {
    if (is_empty_queue(queue)) {
        return 0; // Queue is empty
    }

    return queue->data[queue->front];
}

// Print queue contents
void print_queue(CircularQueue *queue) {
    if (is_empty_queue(queue)) {
        printf("Queue is empty\n");
        return;
    }

    printf("Queue (front to rear): ");
    int index = queue->front;
    for (int i = 0; i < queue->size; i++) {
        printf("%d ", queue->data[index]);
        index = (index + 1) % QUEUE_SIZE;
    }
    printf("\n");
}

int main() {
    CircularQueue queue;
    init_circular_queue(&queue);

    // Test queue operations
    enqueue(&queue, 10);
    enqueue(&queue, 20);
    enqueue(&queue, 30);

    print_queue(&queue);

    printf("Front element: %d\n", front_queue(&queue));

    int dequeued = dequeue(&queue);
    printf("Dequeued element: %d\n", dequeued);

    print_queue(&queue);

    return 0;
}

```

Linked List-based Queue

```
#include <stdio.h>
#include <stdlib.h>
#include <stdbool.h>

typedef struct QueueNode {
    int data;
    struct QueueNode *next;
} QueueNode;

typedef struct {
    QueueNode *front;
    QueueNode *rear;
    size_t size;
} LinkedQueue;

// Create new queue node
QueueNode* create_queue_node(int data) {
    QueueNode *node = malloc(sizeof(QueueNode));
    if (node == NULL) return NULL;

    node->data = data;
    node->next = NULL;
    return node;
}

// Initialize queue
LinkedQueue* create_linked_queue(void) {
    LinkedQueue *queue = malloc(sizeof(LinkedQueue));
    if (queue == NULL) return NULL;

    queue->front = NULL;
    queue->rear = NULL;
    queue->size = 0;
    return queue;
}

// Free queue
void free_linked_queue(LinkedQueue *queue) {
    if (queue == NULL) return;

    QueueNode *current = queue->front;
    while (current != NULL) {
        QueueNode *next = current->next;
        free(current);
        current = next;
    }
}
```

```

    }

    free(queue);
}

// Check if queue is empty
bool is_empty_linked_queue(LinkedException *queue) {
    return queue->front == NULL;
}

// Enqueue (add) element to queue
bool enqueue_linked(LinkedException *queue, int value) {
    QueueNode *new_node = create_queue_node(value);
    if (new_node == NULL) return false;

    if (is_empty_linked_queue(queue)) {
        queue->front = queue->rear = new_node;
    } else {
        queue->rear->next = new_node;
        queue->rear = new_node;
    }

    queue->size++;
    return true;
}

// Dequeue (remove) element from queue
int dequeue_linked(LinkedException *queue) {
    if (is_empty_linked_queue(queue)) {
        return 0; // Queue underflow
    }

    QueueNode *temp = queue->front;
    int data = temp->data;

    queue->front = queue->front->next;
    if (queue->front == NULL) {
        queue->rear = NULL; // Queue is now empty
    }

    free(temp);
    queue->size--;

    return data;
}

// Peek at front element

```

```

int front_linked_queue(LinkedList *queue) {
    if (is_empty_linked_queue(queue)) {
        return 0; // Queue is empty
    }

    return queue->front->data;
}

// Print queue contents
void print_linked_queue(LinkedList *queue) {
    if (is_empty_linked_queue(queue)) {
        printf("Queue is empty\n");
        return;
    }

    printf("Queue (front to rear): ");
    QueueNode *current = queue->front;
    while (current != NULL) {
        printf("%d ", current->data);
        current = current->next;
    }
    printf("\n");
}

int main() {
    LinkedList *queue = create_linked_queue();
    if (queue == NULL) {
        printf("Failed to create queue\n");
        return 1;
    }

    // Test queue operations
    enqueue_linked(queue, 10);
    enqueue_linked(queue, 20);
    enqueue_linked(queue, 30);

    print_linked_queue(queue);

    printf("Front element: %d\n", front_linked_queue(queue));

    int dequeued = dequeue_linked(queue);
    printf("Dequeued element: %d\n", dequeued);

    print_linked_queue(queue);

    free_linked_queue(queue);
    return 0;
}

```



```
}
```

Practical Examples

Expression Evaluation Using Stacks

```
#include <stdio.h>
#include <stdlib.h>
#include <stdbool.h>
#include <string.h>
#include <ctype.h>

#define STACK_SIZE 100

typedef struct {
    int data[STACK_SIZE];
    int top;
} Stack;

void init_stack(Stack *stack) {
    stack->top = -1;
}

bool is_empty(Stack *stack) {
    return stack->top == -1;
}

bool is_full(Stack *stack) {
    return stack->top == STACK_SIZE - 1;
}

bool push(Stack *stack, int value) {
    if (is_full(stack)) return false;
    stack->data[++stack->top] = value;
    return true;
}

int pop(Stack *stack) {
    if (is_empty(stack)) return 0;
    return stack->data[stack->top--];
}

int peek(Stack *stack) {
    if (is_empty(stack)) return 0;
```

```

    return stack->data[stack->top];
}

// Check if character is operator
bool is_operator(char c) {
    return (c == '+' || c == '-' || c == '*' || c == '/');
}

// Get precedence of operator
int precedence(char op) {
    switch (op) {
        case '+':
        case '-':
            return 1;
        case '*':
        case '/':
            return 2;
        default:
            return 0;
    }
}

// Convert infix to postfix
void infix_to_postfix(const char *infix, char *postfix) {
    Stack stack;
    init_stack(&stack);

    int i, j = 0;
    for (i = 0; infix[i] != '\0'; i++) {
        char c = infix[i];

        if (isdigit(c)) {
            postfix[j++] = c;
        } else if (c == '(') {
            push(&stack, c);
        } else if (c == ')') {
            while (!is_empty(&stack) && peek(&stack) != '(') {
                postfix[j++] = pop(&stack);
            }
            pop(&stack); // Pop '('
        } else if (is_operator(c)) {
            while (!is_empty(&stack) && precedence(peek(&stack)) >= precedence(c)) {
                postfix[j++] = pop(&stack);
            }
            push(&stack, c);
        }
    }
}

```

```

        while (!is_empty(&stack)) {
            postfix[j++] = pop(&stack);
        }

        postfix[j] = '\0';
    }

// Evaluate postfix expression
int evaluate_postfix(const char *postfix) {
    Stack stack;
    init_stack(&stack);

    for (int i = 0; postfix[i] != '\0'; i++) {
        char c = postfix[i];

        if (isdigit(c)) {
            push(&stack, c - '0');
        } else if (is_operator(c)) {
            int b = pop(&stack);
            int a = pop(&stack);

            switch (c) {
                case '+':
                    push(&stack, a + b);
                    break;
                case '-':
                    push(&stack, a - b);
                    break;
                case '*':
                    push(&stack, a * b);
                    break;
                case '/':
                    push(&stack, a / b);
                    break;
            }
        }
    }

    return pop(&stack);
}

int main() {
    char infix[] = "3+4*2/(1-5)";
    char postfix[100];

    printf("Infix expression: %s\n", infix);

```

```
    infix_to_postfix(infix, postfix);  
    printf("Postfix expression: %s\n", postfix);  
  
    int result = evaluate_postfix(postfix);  
    printf("Result: %d\n", result);  
  
    return 0;  
}
```

Summary

Linear data structures are fundamental to computer science and programming:

1. **Dynamic Arrays:** Provide random access with dynamic resizing capabilities
2. **Linked Lists:** Offer efficient insertion/deletion with sequential access
3. **Stacks:** Implement LIFO behavior for expression evaluation and backtracking
4. **Queues:** Implement FIFO behavior for scheduling and breadth-first processing

Each structure has its own trade-offs in terms of memory usage, access patterns, and performance characteristics. Understanding these trade-offs is crucial for selecting the appropriate data structure for specific applications.

Key considerations when choosing linear data structures: - **Access Pattern:** Random vs. sequential access - **Insertion/Deletion Frequency:** At ends vs. middle - **Memory Constraints:** Fixed vs. dynamic allocation - **Performance Requirements:** Time complexity for operations - **Implementation Complexity:** Simple vs. complex structures

Mastering linear data structures provides a solid foundation for understanding more complex data structures and algorithms.

Tree Structures

Introduction

Tree structures are hierarchical data structures that organize data in a parent-child relationship, where each node can have zero or more child nodes. Trees are fundamental in computer science and are used in various applications such as file systems, databases, expression parsing, and decision-making algorithms.

In this chapter, we'll explore different types of tree structures, from basic binary trees to more complex balanced trees like AVL and red-black trees. We'll implement these structures in C and examine their properties, operations, and performance characteristics.

Binary Trees

A binary tree is a tree data structure where each node has at most two children, referred to as the left child and the right child.

Basic Binary Tree Implementation

```
#include <stdio.h>
#include <stdlib.h>
#include <stdbool.h>

typedef struct TreeNode {
    int data;
    struct TreeNode *left;
    struct TreeNode *right;
} TreeNode;

// Create new tree node
TreeNode* create_node(int data) {
    TreeNode *node = malloc(sizeof(TreeNode));
    if (node == NULL) return NULL;

    node->data = data;
    node->left = NULL;
    node->right = NULL;
    return node;
}
```

```

}

// Free entire tree
void free_tree(TreeNode *root) {
    if (root == NULL) return;

    free_tree(root->left);
    free_tree(root->right);
    free(root);
}

// Insert node in binary tree (level-order insertion)
TreeNode* insert_level_order(TreeNode *root, int data) {
    TreeNode *new_node = create_node(data);
    if (new_node == NULL) return root;

    if (root == NULL) {
        return new_node;
    }

    // Use queue for level-order insertion
    TreeNode *queue[100];
    int front = 0, rear = 0;
    queue[rear++] = root;

    while (front < rear) {
        TreeNode *current = queue[front++];

        if (current->left == NULL) {
            current->left = new_node;
            return root;
        } else {
            queue[rear++] = current->left;
        }

        if (current->right == NULL) {
            current->right = new_node;
            return root;
        } else {
            queue[rear++] = current->right;
        }
    }

    return root;
}

// Tree traversals

```

```

void inorder_traversal(TreeNode *root) {
    if (root == NULL) return;

    inorder_traversal(root->left);
    printf("%d ", root->data);
    inorder_traversal(root->right);
}

void preorder_traversal(TreeNode *root) {
    if (root == NULL) return;

    printf("%d ", root->data);
    preorder_traversal(root->left);
    preorder_traversal(root->right);
}

void postorder_traversal(TreeNode *root) {
    if (root == NULL) return;

    postorder_traversal(root->left);
    postorder_traversal(root->right);
    printf("%d ", root->data);
}

// Level-order traversal (breadth-first)
void level_order_traversal(TreeNode *root) {
    if (root == NULL) return;

    TreeNode *queue[100];
    int front = 0, rear = 0;
    queue[rear++] = root;

    while (front < rear) {
        TreeNode *current = queue[front++];
        printf("%d ", current->data);

        if (current->left != NULL) {
            queue[rear++] = current->left;
        }

        if (current->right != NULL) {
            queue[rear++] = current->right;
        }
    }
}

```

```

// Tree height
int tree_height(TreeNode *root) {
    if (root == NULL) return 0;

    int left_height = tree_height(root->left);
    int right_height = tree_height(root->right);

    return 1 + (left_height > right_height ? left_height : right_height);
}

// Node count
int node_count(TreeNode *root) {
    if (root == NULL) return 0;

    return 1 + node_count(root->left) + node_count(root->right);
}

// Leaf count
int leaf_count(TreeNode *root) {
    if (root == NULL) return 0;

    if (root->left == NULL && root->right == NULL) {
        return 1;
    }

    return leaf_count(root->left) + leaf_count(root->right);
}

int main() {
    TreeNode *root = NULL;

    // Build tree: 1, 2, 3, 4, 5, 6, 7
    for (int i = 1; i <= 7; i++) {
        root = insert_level_order(root, i);
    }

    printf("Inorder traversal: ");
    inorder_traversal(root);
    printf("\n");

    printf("Preorder traversal: ");
    preorder_traversal(root);
    printf("\n");

    printf("Postorder traversal: ");
    postorder_traversal(root);
    printf("\n");
}

```



```

    printf("Level-order traversal: ");
    level_order_traversal(root);
    printf("\n");

    printf("Tree height: %d\n", tree_height(root));
    printf("Node count: %d\n", node_count(root));
    printf("Leaf count: %d\n", leaf_count(root));

    free_tree(root);
    return 0;
}

```

Binary Search Trees (BST)

A binary search tree is a binary tree where each node follows the property that all values in the left subtree are less than the node's value, and all values in the right subtree are greater than the node's value.

BST Implementation

```

#include <stdio.h>
#include <stdlib.h>
#include <stdbool.h>

typedef struct BSTNode {
    int data;
    struct BSTNode *left;
    struct BSTNode *right;
} BSTNode;

// Create new BST node
BSTNode* create_bst_node(int data) {
    BSTNode *node = malloc(sizeof(BSTNode));
    if (node == NULL) return NULL;

    node->data = data;
    node->left = NULL;
    node->right = NULL;
    return node;
}

// Free BST
void free_bst(BSTNode *root) {

```

```

        if (root == NULL) return;

        free_bst(root->left);
        free_bst(root->right);
        free(root);
    }

// Insert node in BST
BSTNode* insert_bst(BSTNode *root, int data) {
    if (root == NULL) {
        return create_bst_node(data);
    }

    if (data < root->data) {
        root->left = insert_bst(root->left, data);
    } else if (data > root->data) {
        root->right = insert_bst(root->right, data);
    }

    // If data == root->data, we don't insert duplicates

    return root;
}

// Search in BST
BSTNode* search_bst(BSTNode *root, int data) {
    if (root == NULL || root->data == data) {
        return root;
    }

    if (data < root->data) {
        return search_bst(root->left, data);
    } else {
        return search_bst(root->right, data);
    }
}

// Find minimum value node
BSTNode* find_min(BSTNode *root) {
    if (root == NULL) return NULL;

    while (root->left != NULL) {
        root = root->left;
    }

    return root;
}

```

```

// Find maximum value node
BSTNode* find_max(BSTNode *root) {
    if (root == NULL) return NULL;

    while (root->right != NULL) {
        root = root->right;
    }

    return root;
}

// Delete node from BST
BSTNode* delete_bst(BSTNode *root, int data) {
    if (root == NULL) return root;

    if (data < root->data) {
        root->left = delete_bst(root->left, data);
    } else if (data > root->data) {
        root->right = delete_bst(root->right, data);
    } else {
        // Node to be deleted found

        // Case 1: Node has no children
        if (root->left == NULL && root->right == NULL) {
            free(root);
            return NULL;
        }

        // Case 2: Node has one child
        if (root->left == NULL) {
            BSTNode *temp = root->right;
            free(root);
            return temp;
        }

        if (root->right == NULL) {
            BSTNode *temp = root->left;
            free(root);
            return temp;
        }

        // Case 3: Node has two children
        BSTNode *temp = find_min(root->right);
        root->data = temp->data;
        root->right = delete_bst(root->right, temp->data);
    }
}

```

```

    return root;
}

// Inorder traversal (gives sorted order for BST)
void inorder_bst(BSTNode *root) {
    if (root == NULL) return;

    inorder_bst(root->left);
    printf("%d ", root->data);
    inorder_bst(root->right);
}

// Check if tree is valid BST
bool is_valid_bst(BSTNode *root, long min_val, long max_val) {
    if (root == NULL) return true;

    if (root->data <= min_val || root->data >= max_val) {
        return false;
    }

    return is_valid_bst(root->left, min_val, root->data) &&
        is_valid_bst(root->right, root->data, max_val);
}

bool validate_bst(BSTNode *root) {
    return is_valid_bst(root, LONG_MIN, LONG_MAX);
}

int main() {
    BSTNode *root = NULL;

    // Insert nodes
    int values[] = {50, 30, 70, 20, 40, 60, 80, 10, 25, 35, 45};
    int n = sizeof(values) / sizeof(values[0]);

    for (int i = 0; i < n; i++) {
        root = insert_bst(root, values[i]);
    }

    printf("BST inorder traversal (sorted): ");
    inorder_bst(root);
    printf("\n");

    printf("Is valid BST: %s\n", validate_bst(root) ? "Yes" : "No");

    // Search operations
    int search_values[] = {25, 35, 90};

```

```

    for (int i = 0; i < 3; i++) {
        BSTNode *found = search_bst(root, search_values[i]);
        printf("Search %d: %s\n", search_values[i],
            found ? "Found" : "Not found");
    }

    // Find min and max
    BSTNode *min_node = find_min(root);
    BSTNode *max_node = find_max(root);
    printf("Minimum value: %d\n", min_node ? min_node->data : -1);
    printf("Maximum value: %d\n", max_node ? max_node->data : -1);

    // Delete operations
    printf("Deleting 20...\n");
    root = delete_bst(root, 20);
    printf("BST after deletion: ");
    inorder_bst(root);
    printf("\n");

    printf("Deleting 30...\n");
    root = delete_bst(root, 30);
    printf("BST after deletion: ");
    inorder_bst(root);
    printf("\n");

    printf("Deleting 50...\n");
    root = delete_bst(root, 50);
    printf("BST after deletion: ");
    inorder_bst(root);
    printf("\n");

    free_bst(root);
    return 0;
}

```

AVL Trees

AVL trees are self-balancing binary search trees where the heights of the two child subtrees of any node differ by at most one. This ensures $O(\log n)$ time complexity for search, insert, and delete operations.

AVL Tree Implementation

```
#include <stdio.h>
#include <stdlib.h>
#include <stdbool.h>
#include <math.h>

typedef struct AVLNode {
    int data;
    struct AVLNode *left;
    struct AVLNode *right;
    int height;
} AVLNode;

// Get height of node
int get_height(AVLNode *node) {
    return node ? node->height : 0;
}

// Get balance factor
int get_balance(AVLNode *node) {
    return node ? get_height(node->left) - get_height(node->right) : 0;
}

// Create new AVL node
AVLNode* create_avl_node(int data) {
    AVLNode *node = malloc(sizeof(AVLNode));
    if (node == NULL) return NULL;

    node->data = data;
    node->left = NULL;
    node->right = NULL;
    node->height = 1; // New node is initially at leaf level
    return node;
}

// Free AVL tree
void free_avl(AVLNode *root) {
    if (root == NULL) return;

    free_avl(root->left);
    free_avl(root->right);
    free(root);
}

// Update height of node
```

```

void update_height(AVLNode *node) {
    if (node != NULL) {
        node->height = 1 + fmax(get_height(node->left), get_height(node->right));
    }
}

// Right rotate
AVLNode* right_rotate(AVLNode *y) {
    AVLNode *x = y->left;
    AVLNode *T2 = x->right;

    // Perform rotation
    x->right = y;
    y->left = T2;

    // Update heights
    update_height(y);
    update_height(x);

    // Return new root
    return x;
}

// Left rotate
AVLNode* left_rotate(AVLNode *x) {
    AVLNode *y = x->right;
    AVLNode *T2 = y->left;

    // Perform rotation
    y->left = x;
    x->right = T2;

    // Update heights
    update_height(x);
    update_height(y);

    // Return new root
    return y;
}

// Insert node in AVL tree
AVLNode* insert_avl(AVLNode *root, int data) {
    // 1. Perform normal BST insertion
    if (root == NULL) {
        return create_avl_node(data);
    }
}

```

```

    if (data < root->data) {
        root->left = insert_avl(root->left, data);
    } else if (data > root->data) {
        root->right = insert_avl(root->right, data);
    } else {
        // Equal keys not allowed
        return root;
    }

    // 2. Update height of this ancestor node
    update_height(root);

    // 3. Get balance factor
    int balance = get_balance(root);

    // 4. If unbalanced, perform rotations

    // Left Left Case
    if (balance > 1 && data < root->left->data) {
        return right_rotate(root);
    }

    // Right Right Case
    if (balance < -1 && data > root->right->data) {
        return left_rotate(root);
    }

    // Left Right Case
    if (balance > 1 && data > root->left->data) {
        root->left = left_rotate(root->left);
        return right_rotate(root);
    }

    // Right Left Case
    if (balance < -1 && data < root->right->data) {
        root->right = right_rotate(root->right);
        return left_rotate(root);
    }

    // Return unchanged node pointer
    return root;
}

// Find node with minimum value
AVLNode* find_min_avl(AVLNode *root) {
    if (root == NULL) return NULL;

```



```

    while (root->left != NULL) {
        root = root->left;
    }

    return root;
}

// Delete node from AVL tree
AVLNode* delete_avl(AVLNode *root, int data) {
    // 1. Perform standard BST deletion
    if (root == NULL) return root;

    if (data < root->data) {
        root->left = delete_avl(root->left, data);
    } else if (data > root->data) {
        root->right = delete_avl(root->right, data);
    } else {
        // Node to be deleted found
        if (root->left == NULL || root->right == NULL) {
            AVLNode *temp = root->left ? root->left : root->right;

            if (temp == NULL) {
                temp = root;
                root = NULL;
            } else {
                *root = *temp;
            }

            free(temp);
        } else {
            AVLNode *temp = find_min_avl(root->right);
            root->data = temp->data;
            root->right = delete_avl(root->right, temp->data);
        }
    }

    if (root == NULL) return root;

    // 2. Update height of current node
    update_height(root);

    // 3. Get balance factor
    int balance = get_balance(root);

    // 4. If unbalanced, perform rotations

    // Left Left Case

```

```

    if (balance > 1 && get_balance(root->left) >= 0) {
        return right_rotate(root);
    }

    // Left Right Case
    if (balance > 1 && get_balance(root->left) < 0) {
        root->left = left_rotate(root->left);
        return right_rotate(root);
    }

    // Right Right Case
    if (balance < -1 && get_balance(root->right) <= 0) {
        return left_rotate(root);
    }

    // Right Left Case
    if (balance < -1 && get_balance(root->right) > 0) {
        root->right = right_rotate(root->right);
        return left_rotate(root);
    }

    return root;
}

// Search in AVL tree
AVLNode* search_avl(AVLNode *root, int data) {
    if (root == NULL || root->data == data) {
        return root;
    }

    if (data < root->data) {
        return search_avl(root->left, data);
    } else {
        return search_avl(root->right, data);
    }
}

// Inorder traversal
void inorder_avl(AVLNode *root) {
    if (root == NULL) return;

    inorder_avl(root->left);
    printf("%d ", root->data);
    inorder_avl(root->right);
}

// Check if tree is balanced

```

```

bool is_balanced(AVLNode *root) {
    if (root == NULL) return true;

    int balance = get_balance(root);
    if (abs(balance) > 1) return false;

    return is_balanced(root->left) && is_balanced(root->right);
}

int main() {
    AVLNode *root = NULL;

    // Insert nodes
    int values[] = {10, 20, 30, 40, 50, 25};
    int n = sizeof(values) / sizeof(values[0]);

    printf("Inserting values: ");
    for (int i = 0; i < n; i++) {
        printf("%d ", values[i]);
        root = insert_avl(root, values[i]);
    }
    printf("\n");

    printf("AVL tree inorder traversal: ");
    inorder_avl(root);
    printf("\n");

    printf("Is tree balanced: %s\n", is_balanced(root) ? "Yes" : "No");
    printf("Root height: %d\n", get_height(root));

    // Search operations
    int search_values[] = {25, 35, 60};
    for (int i = 0; i < 3; i++) {
        AVLNode *found = search_avl(root, search_values[i]);
        printf("Search %d: %s\n", search_values[i],
            found ? "Found" : "Not found");
    }

    // Delete operations
    printf("Deleting 30...\n");
    root = delete_avl(root, 30);
    printf("AVL tree after deletion: ");
    inorder_avl(root);
    printf("\n");

    printf("Is tree balanced: %s\n", is_balanced(root) ? "Yes" : "No");
}

```

```

    printf("Root height: %d\n", get_height(root));

    free_avl(root);
    return 0;
}

```

Red-Black Trees

Red-black trees are another type of self-balancing binary search tree with additional properties that ensure the tree remains approximately balanced during insertions and deletions.

Red-Black Tree Properties

1. Every node is either red or black
2. The root is black
3. All leaves (NIL nodes) are black
4. If a node is red, both its children are black
5. Every path from a node to its descendant leaves contains the same number of black nodes

```

#include <stdio.h>
#include <stdlib.h>
#include <stdbool.h>

typedef enum { RED, BLACK } Color;

typedef struct RBNode {
    int data;
    Color color;
    struct RBNode *left;
    struct RBNode *right;
    struct RBNode *parent;
} RBNode;

// Global NIL node (sentinel)
RBNode *NIL;

// Initialize NIL node
void init_nil(void) {
    NIL = malloc(sizeof(RBNode));
    NIL->color = BLACK;
    NIL->left = NIL->right = NIL->parent = NULL;
}

// Create new red-black node

```

```

RBNode* create_rb_node(int data) {
    RBNode *node = malloc(sizeof(RBNode));
    if (node == NULL) return NULL;

    node->data = data;
    node->color = RED; // New nodes are red by default
    node->left = node->right = node->parent = NIL;
    return node;
}

// Free red-black tree
void free_rb_tree(RBNode *root) {
    if (root == NULL || root == NIL) return;

    free_rb_tree(root->left);
    free_rb_tree(root->right);
    free(root);
}

// Left rotate
void left_rotate(RBNode **root, RBNode *x) {
    RBNode *y = x->right;

    // Turn y's left subtree into x's right subtree
    x->right = y->left;
    if (y->left != NIL) {
        y->left->parent = x;
    }

    // Link x's parent to y
    y->parent = x->parent;
    if (x->parent == NIL) {
        *root = y;
    } else if (x == x->parent->left) {
        x->parent->left = y;
    } else {
        x->parent->right = y;
    }

    // Put x on y's left
    y->left = x;
    x->parent = y;
}

// Right rotate
void right_rotate(RBNode **root, RBNode *y) {

```

```

RBNode *x = y->left;

// Turn x's right subtree into y's left subtree
y->left = x->right;
if (x->right != NIL) {
    x->right->parent = y;
}

// Link y's parent to x
x->parent = y->parent;
if (y->parent == NIL) {
    *root = x;
} else if (y == y->parent->right) {
    y->parent->right = x;
} else {
    y->parent->left = x;
}

// Put y on x's right
x->right = y;
y->parent = x;
}

// Fixup after insertion
void insert_fixup(RBNode **root, RBNode *z) {
    while (z->parent->color == RED) {
        if (z->parent == z->parent->parent->left) {
            RBNode *y = z->parent->parent->right;

            if (y->color == RED) {
                // Case 1: Uncle is red
                z->parent->color = BLACK;
                y->color = BLACK;
                z->parent->parent->color = RED;
                z = z->parent->parent;
            } else {
                if (z == z->parent->right) {
                    // Case 2: Uncle is black and z is right child
                    z = z->parent;
                    left_rotate(root, z);
                }
                // Case 3: Uncle is black and z is left child
                z->parent->color = BLACK;
                z->parent->parent->color = RED;
                right_rotate(root, z->parent->parent);
            }
        } else {

```

```

        RBNode *y = z->parent->parent->left;

        if (y->color == RED) {
            // Case 1: Uncle is red
            z->parent->color = BLACK;
            y->color = BLACK;
            z->parent->parent->color = RED;
            z = z->parent->parent;
        } else {
            if (z == z->parent->left) {
                // Case 2: Uncle is black and z is left child
                z = z->parent;
                right_rotate(root, z);
            }
            // Case 3: Uncle is black and z is right child
            z->parent->color = BLACK;
            z->parent->parent->color = RED;
            left_rotate(root, z->parent->parent);
        }
    }
}

(*root)->color = BLACK;
}

// Insert node in red-black tree
RBNode* insert_rb(RBNode **root, int data) {
    RBNode *z = create_rb_node(data);
    if (z == NULL) return NULL;

    RBNode *y = NIL;
    RBNode *x = *root;

    // Find insertion point
    while (x != NIL) {
        y = x;
        if (z->data < x->data) {
            x = x->left;
        } else {
            x = x->right;
        }
    }

    z->parent = y;
    if (y == NIL) {
        *root = z;
    } else if (z->data < y->data) {
        y->left = z;
    }
}

```

```

    } else {
        y->right = z;
    }

    // Fix red-black properties
    insert_fixup(root, z);

    return z;
}

// Inorder traversal
void inorder_rb(RBNode *root) {
    if (root == NULL || root == NIL) return;

    inorder_rb(root->left);
    printf("%d(%s) ", root->data, root->color == RED ? "R" : "B");
    inorder_rb(root->right);
}

int main() {
    init_nil();

    RBNode *root = NIL;

    // Insert nodes
    int values[] = {7, 3, 18, 10, 22, 8, 11, 26};
    int n = sizeof(values) / sizeof(values[0]);

    printf("Inserting values into red-black tree: ");
    for (int i = 0; i < n; i++) {
        printf("%d ", values[i]);
        insert_rb(&root, values[i]);
    }
    printf("\n");

    printf("Red-black tree inorder traversal: ");
    inorder_rb(root);
    printf("\n");

    free_rb_tree(root);
    free(NIL);
    return 0;
}

```


B-Trees

B-trees are balanced tree data structures that maintain sorted data and allow for efficient insertion, deletion, and search operations. They are commonly used in databases and file systems.

B-Tree Implementation (Simplified)

```
#include <stdio.h>
#include <stdlib.h>
#include <stdbool.h>

#define T 3 // Minimum degree

typedef struct BTreeNode {
    int n; // Number of keys
    int keys[2 * T - 1]; // Array of keys
    struct BTreeNode *children[2 * T]; // Array of child pointers
    bool leaf; // True if node is leaf
} BTreeNode;

// Create new B-tree node
BTreeNode* create_btree_node(bool leaf) {
    BTreeNode *node = malloc(sizeof(BTreeNode));
    if (node == NULL) return NULL;

    node->n = 0;
    node->leaf = leaf;

    for (int i = 0; i < 2 * T; i++) {
        node->children[i] = NULL;
    }

    return node;
}

// Search for key in B-tree
BTreeNode* search_btree(BTreeNode *root, int key) {
    int i = 0;
    while (i < root->n && key > root->keys[i]) {
        i++;
    }

    if (i < root->n && key == root->keys[i]) {
        return root;
    }
}
```

```

    if (root->leaf) {
        return NULL;
    }

    return search_btree(root->children[i], key);
}

// Split child node
void split_child(BTreeNode *parent, int i, BTreeNode *child) {
    BTreeNode *z = create_btree_node(child->leaf);
    z->n = T - 1;

    // Copy the last T-1 keys of child to z
    for (int j = 0; j < T - 1; j++) {
        z->keys[j] = child->keys[j + T];
    }

    // Copy the last T children of child to z
    if (!child->leaf) {
        for (int j = 0; j < T; j++) {
            z->children[j] = child->children[j + T];
        }
    }

    child->n = T - 1;

    // Shift parent's children to make space for z
    for (int j = parent->n; j >= i + 1; j--) {
        parent->children[j + 1] = parent->children[j];
    }

    parent->children[i + 1] = z;

    // Shift parent's keys to make space
    for (int j = parent->n - 1; j >= i; j--) {
        parent->keys[j + 1] = parent->keys[j];
    }

    parent->keys[i] = child->keys[T - 1];
    parent->n++;
}

// Insert non-full node
void insert_non_full(BTreeNode *node, int key) {
    int i = node->n - 1;

    if (node->leaf) {

```

```

    // Find location to insert key and shift keys
    while (i >= 0 && node->keys[i] > key) {
        node->keys[i + 1] = node->keys[i];
        i--;
    }

    node->keys[i + 1] = key;
    node->n++;
} else {
    // Find child to recurse on
    while (i >= 0 && node->keys[i] > key) {
        i--;
    }
    i++;

    // Split child if it's full
    if (node->children[i]->n == 2 * T - 1) {
        split_child(node, i, node->children[i]);

        if (node->keys[i] < key) {
            i++;
        }
    }

    insert_non_full(node->children[i], key);
}
}

// Insert key in B-tree
BTreeNode* insert_btree(BTreeNode *root, int key) {
    if (root == NULL) {
        root = create_btree_node(true);
        root->keys[0] = key;
        root->n = 1;
        return root;
    }

    // If root is full, split it
    if (root->n == 2 * T - 1) {
        BTreeNode *new_root = create_btree_node(false);
        new_root->children[0] = root;
        split_child(new_root, 0, root);

        int i = 0;
        if (new_root->keys[0] < key) {
            i++;
        }
    }

```

```

        insert_non_full(new_root->children[i], key);
        return new_root;
    } else {
        insert_non_full(root, key);
        return root;
    }
}

// Print B-tree
void print_btree(BTreeNode *root, int level) {
    if (root == NULL) return;

    for (int i = 0; i < level; i++) {
        printf(" ");
    }

    printf("Level %d: [", level);
    for (int i = 0; i < root->n; i++) {
        printf("%d ", root->keys[i]);
    }
    printf("]\n");

    if (!root->leaf) {
        for (int i = 0; i <= root->n; i++) {
            print_btree(root->children[i], level + 1);
        }
    }
}

int main() {
    BTreeNode *root = NULL;

    // Insert nodes
    int values[] = {10, 20, 5, 6, 12, 30, 7, 17};
    int n = sizeof(values) / sizeof(values[0]);

    printf("Inserting values into B-tree: ");
    for (int i = 0; i < n; i++) {
        printf("%d ", values[i]);
        root = insert_btree(root, values[i]);
    }
    printf("\n\n");

    printf("B-tree structure:\n");
    print_btree(root, 0);
}

```

```

// Search operations
int search_values[] = {6, 15, 17};
for (int i = 0; i < 3; i++) {
    BTreeNode *found = search_btree(root, search_values[i]);
    printf("Search %d: %s\n", search_values[i],
        found ? "Found" : "Not found");
}

return 0;
}

```

Practical Examples

File System Directory Tree

```

#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <stdbool.h>

typedef struct DirectoryNode {
    char name[100];
    bool is_directory;
    long size; // For files
    struct DirectoryNode *first_child;
    struct DirectoryNode *next_sibling;
} DirectoryNode;

// Create new directory node
DirectoryNode* create_directory_node(const char *name, bool is_dir, long size) {
    DirectoryNode *node = malloc(sizeof(DirectoryNode));
    if (node == NULL) return NULL;

    strncpy(node->name, name, sizeof(node->name) - 1);
    node->name[sizeof(node->name) - 1] = '\0';
    node->is_directory = is_dir;
    node->size = size;
    node->first_child = NULL;
    node->next_sibling = NULL;

    return node;
}

// Add child to directory

```

```

void add_child(DirectoryNode *parent, DirectoryNode *child) {
    if (parent->first_child == NULL) {
        parent->first_child = child;
    } else {
        DirectoryNode *current = parent->first_child;
        while (current->next_sibling != NULL) {
            current = current->next_sibling;
        }
        current->next_sibling = child;
    }
}

// Print directory tree
void print_directory_tree(DirectoryNode *root, int level) {
    if (root == NULL) return;

    for (int i = 0; i < level; i++) {
        printf("  ");
    }

    if (root->is_directory) {
        printf("[%s]\n", root->name);
    } else {
        printf("%s (%ld bytes)\n", root->name, root->size);
    }

    // Print children
    DirectoryNode *child = root->first_child;
    while (child != NULL) {
        print_directory_tree(child, level + 1);
        child = child->next_sibling;
    }
}

// Calculate directory size
long calculate_directory_size(DirectoryNode *root) {
    if (root == NULL) return 0;

    if (!root->is_directory) {
        return root->size;
    }

    long total_size = 0;
    DirectoryNode *child = root->first_child;
    while (child != NULL) {
        total_size += calculate_directory_size(child);
    }
}

```

```

        child = child->next_sibling;
    }

    return total_size;
}

// Find file by name
DirectoryNode* find_file(DirectoryNode *root, const char *name) {
    if (root == NULL) return NULL;

    if (strcmp(root->name, name) == 0) {
        return root;
    }

    // Search in children
    DirectoryNode *child = root->first_child;
    while (child != NULL) {
        DirectoryNode *found = find_file(child, name);
        if (found != NULL) {
            return found;
        }
        child = child->next_sibling;
    }

    return NULL;
}

int main() {
    // Create file system structure
    DirectoryNode *root = create_directory_node("root", true, 0);

    // Create directories
    DirectoryNode *home = create_directory_node("home", true, 0);
    DirectoryNode *usr = create_directory_node("usr", true, 0);
    DirectoryNode *john = create_directory_node("john", true, 0);
    DirectoryNode *documents = create_directory_node("documents", true, 0);
    DirectoryNode *pictures = create_directory_node("pictures", true, 0);

    // Create files
    DirectoryNode *file1 = create_directory_node("readme.txt", false, 1024);
    DirectoryNode *file2 = create_directory_node("report.pdf", false, 204800);
    DirectoryNode *file3 = create_directory_node("photo.jpg", false, 512000);
    DirectoryNode *file4 = create_directory_node("vacation.png", false, 1024000);

    // Build tree structure
    add_child(root, home);
    add_child(root, usr);

```

```

    add_child(home, john);
    add_child(john, documents);
    add_child(john, pictures);
    add_child(documents, file1);
    add_child(documents, file2);
    add_child(pictures, file3);
    add_child(pictures, file4);

    // Print directory tree
    printf("File System Directory Tree:\n");
    print_directory_tree(root, 0);

    // Calculate sizes
    printf("\nDirectory Sizes:\n");
    printf("Root directory size: %ld bytes\n", calculate_directory_size(root));
    printf("John's directory size: %ld bytes\n", calculate_directory_size(john));
    printf("Documents directory size: %ld bytes\n", calculate_directory_size(documents));

    // Find files
    printf("\nFile Search:\n");
    DirectoryNode *found = find_file(root, "photo.jpg");
    if (found) {
        printf("Found file: %s (%ld bytes)\n", found->name, found->size);
    } else {
        printf("File not found\n");
    }

    return 0;
}

```

Summary

Tree structures are essential data structures in computer science with various applications:

1. **Binary Trees:** Basic hierarchical structure with at most two children per node
2. **Binary Search Trees:** Ordered binary trees with efficient search, insert, and delete operations
3. **AVL Trees:** Self-balancing BSTs that maintain $O(\log n)$ operations
4. **Red-Black Trees:** Another self-balancing BST with relaxed balancing requirements
5. **B-Trees:** Multi-way trees optimized for disk storage and databases

Key properties and considerations:

- **Time Complexity:**
 - BST (unbalanced): $O(n)$ worst case, $O(\log n)$ average
 - Balanced trees (AVL, Red-Black): $O(\log n)$ guaranteed

- B-Trees: $O(\log n)$ with better cache performance
- **Space Complexity:** $O(n)$ for n nodes
- **Use Cases:**
 - BST: Expression trees, decision trees
 - AVL/Red-Black: In-memory databases, symbol tables
 - B-Trees: File systems, database indexes

Understanding tree structures and their variations is crucial for efficient data organization and retrieval in complex applications.

Hash Tables

Introduction

Hash tables (also known as hash maps) are data structures that store key-value pairs and provide efficient insertion, deletion, and lookup operations with an average time complexity of $O(1)$. They achieve this efficiency by using a hash function to compute an index into an array of buckets or slots, from which the desired value can be found.

Hash tables are fundamental in computer science and are used extensively in databases, caches, sets, and many other applications where fast data retrieval is essential.

Hash Function Design

A good hash function is crucial for the performance of a hash table. It should distribute keys uniformly across the hash table to minimize collisions and ensure efficient operations.

Basic Hash Functions

```
#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <stdbool.h>

// Simple hash function for integers
unsigned int hash_int(int key, int table_size) {
    return key % table_size;
}

// Simple hash function for strings (djb2 algorithm)
unsigned int hash_string(const char *str, int table_size) {
    unsigned long hash = 5381;
    int c;

    while ((c = *str++)) {
        hash = ((hash << 5) + hash) + c; // hash * 33 + c
    }

    return hash % table_size;
}
```

```

}

// Simple hash function for strings (sdbm algorithm)
unsigned int hash_string_sdbm(const char *str, int table_size) {
    unsigned long hash = 0;
    int c;

    while ((c = *str++)) {
        hash = c + (hash << 6) + (hash << 16) - hash;
    }

    return hash % table_size;
}

// Universal hash function
unsigned int hash_universal(unsigned int key, int table_size,
                           unsigned int a, unsigned int b, unsigned int p) {
    // Universal hashing: h(key) = ((a * key + b) mod p) mod table_size
    // where p is a prime number >= table_size
    return ((a * key + b) % p) % table_size;
}

int main() {
    int table_size = 100;

    // Test integer hash function
    printf("Hash values for integers:\n");
    for (int i = 1; i <= 10; i++) {
        printf("Key %d -> Hash %u\n", i, hash_int(i, table_size));
    }

    printf("\nHash values for strings (djb2):\n");
    const char *strings[] = {"hello", "world", "hash", "table", "example"};
    for (int i = 0; i < 5; i++) {
        printf("Key '%s' -> Hash %u\n", strings[i], hash_string(strings[i], table_size));
    }

    printf("\nHash values for strings (sdbm):\n");
    for (int i = 0; i < 5; i++) {
        printf("Key '%s' -> Hash %u\n", strings[i], hash_string_sdbm(strings[i], table_size));
    }

    return 0;
}

```

Collision Resolution Strategies

When two keys hash to the same index, a collision occurs. There are several strategies to handle collisions:

Chaining

In chaining, each bucket contains a linked list of elements that hash to the same index.

```
#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <stdbool.h>

#define TABLE_SIZE 10

typedef struct HashNode {
    char key[100];
    int value;
    struct HashNode *next;
} HashNode;

typedef struct {
    HashNode *buckets[TABLE_SIZE];
    int size;
} HashTable;

// Create new hash node
HashNode* create_hash_node(const char *key, int value) {
    HashNode *node = malloc(sizeof(HashNode));
    if (node == NULL) return NULL;

    strncpy(node->key, key, sizeof(node->key) - 1);
    node->key[sizeof(node->key) - 1] = '\0';
    node->value = value;
    node->next = NULL;

    return node;
}

// Initialize hash table
HashTable* create_hash_table(void) {
    HashTable *table = malloc(sizeof(HashTable));
    if (table == NULL) return NULL;

    for (int i = 0; i < TABLE_SIZE; i++) {
```

```

        table->buckets[i] = NULL;
    }

    table->size = 0;
    return table;
}

// Free hash table
void free_hash_table(HashTable *table) {
    if (table == NULL) return;

    for (int i = 0; i < TABLE_SIZE; i++) {
        HashNode *current = table->buckets[i];
        while (current != NULL) {
            HashNode *next = current->next;
            free(current);
            current = next;
        }
    }

    free(table);
}

// Hash function for strings
unsigned int hash(const char *key) {
    unsigned long hash = 5381;
    int c;

    while ((c = *key++)) {
        hash = ((hash << 5) + hash) + c;
    }

    return hash % TABLE_SIZE;
}

// Insert key-value pair
bool hash_table_insert(HashTable *table, const char *key, int value) {
    if (table == NULL || key == NULL) return false;

    unsigned int index = hash(key);

    // Check if key already exists
    HashNode *current = table->buckets[index];
    while (current != NULL) {
        if (strcmp(current->key, key) == 0) {
            current->value = value; // Update existing value
            return true;
        }
    }

```

```

    }
    current = current->next;
}

// Create new node and insert at beginning of chain
HashNode *new_node = create_hash_node(key, value);
if (new_node == NULL) return false;

new_node->next = table->buckets[index];
table->buckets[index] = new_node;
table->size++;

return true;
}

// Search for key
bool hash_table_search(HashTable *table, const char *key, int *value) {
    if (table == NULL || key == NULL) return false;

    unsigned int index = hash(key);

    HashNode *current = table->buckets[index];
    while (current != NULL) {
        if (strcmp(current->key, key) == 0) {
            if (value != NULL) {
                *value = current->value;
            }
            return true;
        }
        current = current->next;
    }

    return false;
}

// Delete key
bool hash_table_delete(HashTable *table, const char *key) {
    if (table == NULL || key == NULL) return false;

    unsigned int index = hash(key);

    HashNode *current = table->buckets[index];
    HashNode *prev = NULL;

    while (current != NULL) {
        if (strcmp(current->key, key) == 0) {
            if (prev == NULL) {

```

```

        // First node in chain
        table->buckets[index] = current->next;
    } else {
        prev->next = current->next;
    }

    free(current);
    table->size--;
    return true;
}

prev = current;
current = current->next;
}

return false;
}

// Print hash table
void print_hash_table(HashTable *table) {
    if (table == NULL) return;

    printf("Hash Table (size: %d):\n", table->size);
    for (int i = 0; i < TABLE_SIZE; i++) {
        printf("Bucket %d: ", i);

        HashNode *current = table->buckets[i];
        if (current == NULL) {
            printf("empty\n");
        } else {
            while (current != NULL) {
                printf("(%s: %d) -> ", current->key, current->value);
                current = current->next;
            }
            printf("NULL\n");
        }
    }
}

int main() {
    HashTable *table = create_hash_table();
    if (table == NULL) {
        printf("Failed to create hash table\n");
        return 1;
    }
}

```

```

// Insert key-value pairs
hash_table_insert(table, "apple", 5);
hash_table_insert(table, "banana", 3);
hash_table_insert(table, "orange", 8);
hash_table_insert(table, "grape", 12);
hash_table_insert(table, "kiwi", 7);

print_hash_table(table);

// Search for keys
int value;
if (hash_table_search(table, "apple", &value)) {
    printf("Found 'apple': %d\n", value);
}

if (hash_table_search(table, "mango", &value)) {
    printf("Found 'mango': %d\n", value);
} else {
    printf("'mango' not found\n");
}

// Update existing key
hash_table_insert(table, "apple", 10);
printf("After updating 'apple':\n");
if (hash_table_search(table, "apple", &value)) {
    printf("Found 'apple': %d\n", value);
}

// Delete key
if (hash_table_delete(table, "banana")) {
    printf("Deleted 'banana'\n");
}

print_hash_table(table);

free_hash_table(table);
return 0;
}

```

Open Addressing

In open addressing, all elements are stored in the hash table itself. When a collision occurs, alternative slots are probed until an empty slot is found.


```

#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <stdbool.h>

#define TABLE_SIZE 13
#define EMPTY_CELL -1
#define DELETED_CELL -2

typedef struct {
    char key[100];
    int value;
    int status; // EMPTY_CELL, DELETED_CELL, or occupied
} HashEntry;

typedef struct {
    HashEntry entries[TABLE_SIZE];
    int size;
} OpenHashTable;

// Initialize open hash table
OpenHashTable* create_open_hash_table(void) {
    OpenHashTable *table = malloc(sizeof(OpenHashTable));
    if (table == NULL) return NULL;

    for (int i = 0; i < TABLE_SIZE; i++) {
        table->entries[i].status = EMPTY_CELL;
    }

    table->size = 0;
    return table;
}

// Free open hash table
void free_open_hash_table(OpenHashTable *table) {
    free(table);
}

// Hash function
unsigned int hash_open(const char *key) {
    unsigned long hash = 5381;
    int c;

    while ((c = *key++)) {
        hash = ((hash << 5) + hash) + c;
    }
}

```

```

    return hash % TABLE_SIZE;
}

// Linear probing hash function
unsigned int hash_linear_probe(const char *key, int attempt) {
    return (hash_open(key) + attempt) % TABLE_SIZE;
}

// Quadratic probing hash function
unsigned int hash_quadratic_probe(const char *key, int attempt) {
    return (hash_open(key) + attempt * attempt) % TABLE_SIZE;
}

// Double hashing
unsigned int hash_double(const char *key, int attempt) {
    unsigned int hash1 = hash_open(key);
    unsigned int hash2 = 7 - (hash_open(key) % 7); // Second hash function
    return (hash1 + attempt * hash2) % TABLE_SIZE;
}

// Insert with linear probing
bool open_hash_table_insert(OpenHashTable *table, const char *key, int value) {
    if (table == NULL || key == NULL) return false;

    int attempt = 0;
    unsigned int index;

    do {
        index = hash_linear_probe(key, attempt);

        if (table->entries[index].status == EMPTY_CELL ||
            table->entries[index].status == DELETED_CELL) {
            // Found empty or deleted slot
            strncpy(table->entries[index].key, key, sizeof(table->entries[index].key) - 1);
            table->entries[index].key[sizeof(table->entries[index].key) - 1] = '\0';
            table->entries[index].value = value;
            table->entries[index].status = 0; // Occupied
            table->size++;
            return true;
        } else if (strcmp(table->entries[index].key, key) == 0) {
            // Key already exists, update value
            table->entries[index].value = value;
            return true;
        }

        attempt++;
    } while (attempt < TABLE_SIZE);
}

```

```

    // Hash table is full
    return false;
}

// Search with linear probing
bool open_hash_table_search(OpenHashTable *table, const char *key, int *value) {
    if (table == NULL || key == NULL) return false;

    int attempt = 0;
    unsigned int index;

    do {
        index = hash_linear_probe(key, attempt);

        if (table->entries[index].status == EMPTY_CELL) {
            // Key not found
            return false;
        } else if (table->entries[index].status != DELETED_CELL &&
            strcmp(table->entries[index].key, key) == 0) {
            // Key found
            if (value != NULL) {
                *value = table->entries[index].value;
            }
            return true;
        }

        attempt++;
    } while (attempt < TABLE_SIZE);

    // Key not found
    return false;
}

// Delete with linear probing
bool open_hash_table_delete(OpenHashTable *table, const char *key) {
    if (table == NULL || key == NULL) return false;

    int attempt = 0;
    unsigned int index;

    do {
        index = hash_linear_probe(key, attempt);

        if (table->entries[index].status == EMPTY_CELL) {
            // Key not found
            return false;
        }
    } while (attempt < TABLE_SIZE);

    // Key found
    table->entries[index].status = DELETED_CELL;
    return true;
}

```

```

    } else if (table->entries[index].status != DELETED_CELL &&
               strcmp(table->entries[index].key, key) == 0) {
        // Key found, mark as deleted
        table->entries[index].status = DELETED_CELL;
        table->size--;
        return true;
    }

    attempt++;
} while (attempt < TABLE_SIZE);

// Key not found
return false;
}

// Print open hash table
void print_open_hash_table(OpenHashTable *table) {
    if (table == NULL) return;

    printf("Open Hash Table (size: %d):\n", table->size);
    for (int i = 0; i < TABLE_SIZE; i++) {
        printf("Index %d: ", i);
        if (table->entries[i].status == EMPTY_CELL) {
            printf("EMPTY\n");
        } else if (table->entries[i].status == DELETED_CELL) {
            printf("DELETED\n");
        } else {
            printf("(%s: %d)\n", table->entries[i].key, table->entries[i].value);
        }
    }
}

int main() {
    OpenHashTable *table = create_open_hash_table();
    if (table == NULL) {
        printf("Failed to create hash table\n");
        return 1;
    }

    // Insert key-value pairs
    open_hash_table_insert(table, "apple", 5);
    open_hash_table_insert(table, "banana", 3);
    open_hash_table_insert(table, "orange", 8);
    open_hash_table_insert(table, "grape", 12);
    open_hash_table_insert(table, "kiwi", 7);

    print_open_hash_table(table);
}

```

```

// Search for keys
int value;
if (open_hash_table_search(table, "apple", &value)) {
    printf("Found 'apple': %d\n", value);
}

if (open_hash_table_search(table, "mango", &value)) {
    printf("Found 'mango': %d\n", value);
} else {
    printf("'mango' not found\n");
}

// Update existing key
open_hash_table_insert(table, "apple", 10);
printf("After updating 'apple':\n");
if (open_hash_table_search(table, "apple", &value)) {
    printf("Found 'apple': %d\n", value);
}

// Delete key
if (open_hash_table_delete(table, "banana")) {
    printf("Deleted 'banana'\n");
}

print_open_hash_table(table);

free_open_hash_table(table);
return 0;
}

```

Dynamic Hash Tables

Dynamic hash tables can resize themselves when the load factor becomes too high, ensuring efficient performance.

Resizable Hash Table with Chaining

```

#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <stdbool.h>

#define INITIAL_SIZE 8

```

```

#define LOAD_FACTOR_THRESHOLD 0.75

typedef struct DynamicHashNode {
    char key[100];
    int value;
    struct DynamicHashNode *next;
} DynamicHashNode;

typedef struct {
    DynamicHashNode **buckets;
    int size;
    int capacity;
} DynamicHashTable;

// Create new hash node
DynamicHashNode* create_dynamic_hash_node(const char *key, int value) {
    DynamicHashNode *node = malloc(sizeof(DynamicHashNode));
    if (node == NULL) return NULL;

    strncpy(node->key, key, sizeof(node->key) - 1);
    node->key[sizeof(node->key) - 1] = '\0';
    node->value = value;
    node->next = NULL;

    return node;
}

// Initialize dynamic hash table
DynamicHashTable* create_dynamic_hash_table(void) {
    DynamicHashTable *table = malloc(sizeof(DynamicHashTable));
    if (table == NULL) return NULL;

    table->buckets = calloc(INITIAL_SIZE, sizeof(DynamicHashNode*));
    if (table->buckets == NULL) {
        free(table);
        return NULL;
    }

    table->size = 0;
    table->capacity = INITIAL_SIZE;
    return table;
}

// Free dynamic hash table
void free_dynamic_hash_table(DynamicHashTable *table) {
    if (table == NULL) return;

```

```

    for (int i = 0; i < table->capacity; i++) {
        DynamicHashNode *current = table->buckets[i];
        while (current != NULL) {
            DynamicHashNode *next = current->next;
            free(current);
            current = next;
        }
    }

    free(table->buckets);
    free(table);
}

// Hash function
unsigned int dynamic_hash(const char *key, int capacity) {
    unsigned long hash = 5381;
    int c;

    while ((c = *key++)) {
        hash = ((hash << 5) + hash) + c;
    }

    return hash % capacity;
}

// Resize hash table
bool resize_hash_table(DynamicHashTable *table) {
    int old_capacity = table->capacity;
    int new_capacity = old_capacity * 2;

    DynamicHashNode **old_buckets = table->buckets;
    DynamicHashNode **new_buckets = calloc(new_capacity, sizeof(DynamicHashNode*));
    if (new_buckets == NULL) return false;

    table->buckets = new_buckets;
    table->capacity = new_capacity;

    // Rehash all elements
    for (int i = 0; i < old_capacity; i++) {
        DynamicHashNode *current = old_buckets[i];
        while (current != NULL) {
            DynamicHashNode *next = current->next;

            unsigned int new_index = dynamic_hash(current->key, new_capacity);
            current->next = table->buckets[new_index];
            table->buckets[new_index] = current;
        }
    }
}

```

```

        current = next;
    }
}

free(old_buckets);
return true;
}

// Insert key-value pair
bool dynamic_hash_table_insert(DynamicHashTable *table, const char *key, int value) {
    if (table == NULL || key == NULL) return false;

    // Check if resize is needed
    double load_factor = (double)table->size / table->capacity;
    if (load_factor > LOAD_FACTOR_THRESHOLD) {
        if (!resize_hash_table(table)) {
            return false;
        }
    }

    unsigned int index = dynamic_hash(key, table->capacity);

    // Check if key already exists
    DynamicHashNode *current = table->buckets[index];
    while (current != NULL) {
        if (strcmp(current->key, key) == 0) {
            current->value = value; // Update existing value
            return true;
        }
        current = current->next;
    }

    // Create new node and insert at beginning of chain
    DynamicHashNode *new_node = create_dynamic_hash_node(key, value);
    if (new_node == NULL) return false;

    new_node->next = table->buckets[index];
    table->buckets[index] = new_node;
    table->size++;

    return true;
}

// Search for key
bool dynamic_hash_table_search(DynamicHashTable *table, const char *key, int *value) {
    if (table == NULL || key == NULL) return false;

```



```

    unsigned int index = dynamic_hash(key, table->capacity);

    DynamicHashNode *current = table->buckets[index];
    while (current != NULL) {
        if (strcmp(current->key, key) == 0) {
            if (value != NULL) {
                *value = current->value;
            }
            return true;
        }
        current = current->next;
    }

    return false;
}

// Delete key
bool dynamic_hash_table_delete(DynamicHashTable *table, const char *key) {
    if (table == NULL || key == NULL) return false;

    unsigned int index = dynamic_hash(key, table->capacity);

    DynamicHashNode *current = table->buckets[index];
    DynamicHashNode *prev = NULL;

    while (current != NULL) {
        if (strcmp(current->key, key) == 0) {
            if (prev == NULL) {
                // First node in chain
                table->buckets[index] = current->next;
            } else {
                prev->next = current->next;
            }

            free(current);
            table->size--;
            return true;
        }

        prev = current;
        current = current->next;
    }

    return false;
}

```

```

// Print dynamic hash table
void print_dynamic_hash_table(DynamicHashTable *table) {
    if (table == NULL) return;

    printf("Dynamic Hash Table (size: %d, capacity: %d, load factor: %.2f):\n",
        table->size, table->capacity, (double)table->size / table->capacity);

    for (int i = 0; i < table->capacity; i++) {
        printf("Bucket %d: ", i);

        DynamicHashNode *current = table->buckets[i];
        if (current == NULL) {
            printf("empty\n");
        } else {
            while (current != NULL) {
                printf("(%s: %d) -> ", current->key, current->value);
                current = current->next;
            }
            printf("NULL\n");
        }
    }
}

int main() {
    DynamicHashTable *table = create_dynamic_hash_table();
    if (table == NULL) {
        printf("Failed to create hash table\n");
        return 1;
    }

    // Insert many key-value pairs to trigger resizing
    char key[20];
    for (int i = 0; i < 20; i++) {
        sprintf(key, "key%d", i);
        dynamic_hash_table_insert(table, key, i * 10);

        // Print status periodically
        if (i % 5 == 4) {
            printf("After inserting %d elements:\n", i + 1);
            print_dynamic_hash_table(table);
            printf("\n");
        }
    }

    // Search for some keys
    int value;

```

```

    if (dynamic_hash_table_search(table, "key5", &value)) {
        printf("Found 'key5': %d\n", value);
    }

    if (dynamic_hash_table_search(table, "key15", &value)) {
        printf("Found 'key15': %d\n", value);
    }

    // Delete some keys
    if (dynamic_hash_table_delete(table, "key10")) {
        printf("Deleted 'key10'\n");
    }

    printf("Final hash table:\n");
    print_dynamic_hash_table(table);

    free_dynamic_hash_table(table);
    return 0;
}

```

Hash Table Performance Analysis

Load Factor and Performance

```

#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <time.h>

#define TABLE_SIZE 1000

typedef struct PerfNode {
    int key;
    int value;
    struct PerfNode *next;
} PerfNode;

typedef struct {
    PerfNode *buckets[TABLE_SIZE];
    int size;
} PerfHashTable;

// Create performance hash table
PerfHashTable* create_perf_hash_table(void) {
    PerfHashTable *table = malloc(sizeof(PerfHashTable));

```

```

    if (table == NULL) return NULL;

    for (int i = 0; i < TABLE_SIZE; i++) {
        table->buckets[i] = NULL;
    }

    table->size = 0;
    return table;
}

// Free performance hash table
void free_perf_hash_table(PerfHashTable *table) {
    if (table == NULL) return;

    for (int i = 0; i < TABLE_SIZE; i++) {
        PerfNode *current = table->buckets[i];
        while (current != NULL) {
            PerfNode *next = current->next;
            free(current);
            current = next;
        }
    }

    free(table);
}

// Hash function
unsigned int perf_hash(int key) {
    return key % TABLE_SIZE;
}

// Insert with performance measurement
bool perf_hash_table_insert(PerfHashTable *table, int key, int value) {
    if (table == NULL) return false;

    unsigned int index = perf_hash(key);

    PerfNode *new_node = malloc(sizeof(PerfNode));
    if (new_node == NULL) return false;

    new_node->key = key;
    new_node->value = value;
    new_node->next = table->buckets[index];
    table->buckets[index] = new_node;
    table->size++;

    return true;
}

```

```

}

// Calculate average chain length
double calculate_average_chain_length(PerfHashTable *table) {
    if (table == NULL) return 0.0;

    int total_length = 0;
    int non_empty_buckets = 0;

    for (int i = 0; i < TABLE_SIZE; i++) {
        if (table->buckets[i] != NULL) {
            non_empty_buckets++;
            int length = 0;
            PerfNode *current = table->buckets[i];
            while (current != NULL) {
                length++;
                current = current->next;
            }
            total_length += length;
        }
    }

    return non_empty_buckets > 0 ? (double)total_length / non_empty_buckets : 0.0;
}

// Calculate load factor
double calculate_load_factor(PerfHashTable *table) {
    if (table == NULL) return 0.0;
    return (double)table->size / TABLE_SIZE;
}

// Find maximum chain length
int find_max_chain_length(PerfHashTable *table) {
    if (table == NULL) return 0;

    int max_length = 0;

    for (int i = 0; i < TABLE_SIZE; i++) {
        int length = 0;
        PerfNode *current = table->buckets[i];
        while (current != NULL) {
            length++;
            current = current->next;
        }
        if (length > max_length) {
            max_length = length;
        }
    }
}

```

```

    }

    return max_length;
}

int main() {
    PerfHashTable *table = create_perf_hash_table();
    if (table == NULL) {
        printf("Failed to create hash table\n");
        return 1;
    }

    // Insert elements with different patterns
    printf("Performance Analysis of Hash Table:\n");
    printf("=====\n\n");

    // Pattern 1: Sequential keys
    printf("Inserting 500 sequential keys (0-499):\n");
    clock_t start = clock();
    for (int i = 0; i < 500; i++) {
        perf_hash_table_insert(table, i, i * 2);
    }
    clock_t end = clock();

    printf("Time taken: %f seconds\n", ((double)(end - start)) / CLOCKS_PER_SEC);
    printf("Load factor: %.2f\n", calculate_load_factor(table));
    printf("Average chain length: %.2f\n", calculate_average_chain_length(table));
    printf("Maximum chain length: %d\n\n", find_max_chain_length(table));

    // Pattern 2: Random keys
    printf("Inserting 500 random keys:\n");
    srand(42); // Fixed seed for reproducibility
    start = clock();
    for (int i = 0; i < 500; i++) {
        int random_key = rand() % 10000;
        perf_hash_table_insert(table, random_key, random_key * 2);
    }
    end = clock();

    printf("Time taken: %f seconds\n", ((double)(end - start)) / CLOCKS_PER_SEC);
    printf("Load factor: %.2f\n", calculate_load_factor(table));
    printf("Average chain length: %.2f\n", calculate_average_chain_length(table));
    printf("Maximum chain length: %d\n\n", find_max_chain_length(table));

    free_perf_hash_table(table);
    return 0;
}

```

Practical Examples

Symbol Table Implementation

```
#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <stdbool.h>

#define SYMBOL_TABLE_SIZE 101

typedef enum {
    TYPE_INT,
    TYPE_FLOAT,
    TYPE_STRING,
    TYPE_FUNCTION
} SymbolType;

typedef struct Symbol {
    char name[100];
    SymbolType type;
    union {
        int int_val;
        float float_val;
        char *string_val;
        void (*func_ptr)(void);
    } value;
    struct Symbol *next;
} Symbol;

typedef struct {
    Symbol *buckets[SYMBOL_TABLE_SIZE];
} SymbolTable;

// Create symbol table
SymbolTable* create_symbol_table(void) {
    SymbolTable *table = malloc(sizeof(SymbolTable));
    if (table == NULL) return NULL;

    for (int i = 0; i < SYMBOL_TABLE_SIZE; i++) {
        table->buckets[i] = NULL;
    }

    return table;
}
```

```

// Free symbol table
void free_symbol_table(SymbolTable *table) {
    if (table == NULL) return;

    for (int i = 0; i < SYMBOL_TABLE_SIZE; i++) {
        Symbol *current = table->buckets[i];
        while (current != NULL) {
            Symbol *next = current->next;
            if (current->type == TYPE_STRING) {
                free(current->value.string_val);
            }
            free(current);
            current = next;
        }
    }

    free(table);
}

// Hash function for symbol names
unsigned int symbol_hash(const char *name) {
    unsigned long hash = 5381;
    int c;

    while ((c = *name++)) {
        hash = ((hash << 5) + hash) + c;
    }

    return hash % SYMBOL_TABLE_SIZE;
}

// Add integer symbol
bool add_integer_symbol(SymbolTable *table, const char *name, int value) {
    if (table == NULL || name == NULL) return false;

    unsigned int index = symbol_hash(name);

    Symbol *new_symbol = malloc(sizeof(Symbol));
    if (new_symbol == NULL) return false;

    strncpy(new_symbol->name, name, sizeof(new_symbol->name) - 1);
    new_symbol->name[sizeof(new_symbol->name) - 1] = '\0';
    new_symbol->type = TYPE_INT;
    new_symbol->value.int_val = value;
    new_symbol->next = table->buckets[index];
    table->buckets[index] = new_symbol;
}

```



```

    return true;
}

// Add float symbol
bool add_float_symbol(SymbolTable *table, const char *name, float value) {
    if (table == NULL || name == NULL) return false;

    unsigned int index = symbol_hash(name);

    Symbol *new_symbol = malloc(sizeof(Symbol));
    if (new_symbol == NULL) return false;

    strncpy(new_symbol->name, name, sizeof(new_symbol->name) - 1);
    new_symbol->name[sizeof(new_symbol->name) - 1] = '\0';
    new_symbol->type = TYPE_FLOAT;
    new_symbol->value.float_val = value;
    new_symbol->next = table->buckets[index];
    table->buckets[index] = new_symbol;

    return true;
}

// Add string symbol
bool add_string_symbol(SymbolTable *table, const char *name, const char *value) {
    if (table == NULL || name == NULL || value == NULL) return false;

    unsigned int index = symbol_hash(name);

    Symbol *new_symbol = malloc(sizeof(Symbol));
    if (new_symbol == NULL) return false;

    strncpy(new_symbol->name, name, sizeof(new_symbol->name) - 1);
    new_symbol->name[sizeof(new_symbol->name) - 1] = '\0';
    new_symbol->type = TYPE_STRING;
    new_symbol->value.string_val = malloc(strlen(value) + 1);
    if (new_symbol->value.string_val == NULL) {
        free(new_symbol);
        return false;
    }
    strcpy(new_symbol->value.string_val, value);
    new_symbol->next = table->buckets[index];
    table->buckets[index] = new_symbol;

    return true;
}

// Lookup symbol

```

```

Symbol* lookup_symbol(SymbolTable *table, const char *name) {
    if (table == NULL || name == NULL) return NULL;

    unsigned int index = symbol_hash(name);

    Symbol *current = table->buckets[index];
    while (current != NULL) {
        if (strcmp(current->name, name) == 0) {
            return current;
        }
        current = current->next;
    }

    return NULL;
}

// Print symbol table
void print_symbol_table(SymbolTable *table) {
    if (table == NULL) return;

    printf("Symbol Table:\n");
    printf("=====\n");

    for (int i = 0; i < SYMBOL_TABLE_SIZE; i++) {
        Symbol *current = table->buckets[i];
        while (current != NULL) {
            printf("Name: %-15s Type: ", current->name);

            switch (current->type) {
                case TYPE_INT:
                    printf("int      Value: %d\n", current->value.int_val);
                    break;
                case TYPE_FLOAT:
                    printf("float   Value: %.2f\n", current->value.float_val);
                    break;
                case TYPE_STRING:
                    printf("string  Value: %s\n", current->value.string_val);
                    break;
                case TYPE_FUNCTION:
                    printf("function\n");
                    break;
            }

            current = current->next;
        }
    }
}

```

```

int main() {
    SymbolTable *table = create_symbol_table();
    if (table == NULL) {
        printf("Failed to create symbol table\n");
        return 1;
    }

    // Add various symbols
    add_integer_symbol(table, "count", 42);
    add_integer_symbol(table, "max_size", 1000);
    add_float_symbol(table, "pi", 3.14159f);
    add_float_symbol(table, "e", 2.71828f);
    add_string_symbol(table, "name", "John Doe");
    add_string_symbol(table, "version", "1.0.0");

    printf("Symbol table contents:\n");
    print_symbol_table(table);

    // Lookup symbols
    printf("\nSymbol lookups:\n");
    printf("=====\n");

    Symbol *symbol = lookup_symbol(table, "count");
    if (symbol != NULL) {
        printf("Found 'count': %d\n", symbol->value.int_val);
    }

    symbol = lookup_symbol(table, "pi");
    if (symbol != NULL) {
        printf("Found 'pi': %.5f\n", symbol->value.float_val);
    }

    symbol = lookup_symbol(table, "name");
    if (symbol != NULL) {
        printf("Found 'name': %s\n", symbol->value.string_val);
    }

    symbol = lookup_symbol(table, "nonexistent");
    if (symbol != NULL) {
        printf("Found 'nonexistent'\n");
    } else {
        printf("'nonexistent' not found\n");
    }

    free_symbol_table(table);
    return 0;
}

```

```
}
```

Cache Implementation

```
#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <stdbool.h>
#include <time.h>

#define CACHE_SIZE 16
#define MAX_KEY_LENGTH 100
#define MAX_VALUE_LENGTH 1000

typedef struct CacheNode {
    char key[MAX_KEY_LENGTH];
    char value[MAX_VALUE_LENGTH];
    time_t timestamp;
    int access_count;
    struct CacheNode *next;
} CacheNode;

typedef struct {
    CacheNode *buckets[CACHE_SIZE];
    int size;
    int max_size;
} LRUCache;

// Create LRU cache
LRUCache* create_lru_cache(int max_size) {
    LRUCache *cache = malloc(sizeof(LRUCache));
    if (cache == NULL) return NULL;

    for (int i = 0; i < CACHE_SIZE; i++) {
        cache->buckets[i] = NULL;
    }

    cache->size = 0;
    cache->max_size = max_size;
    return cache;
}

// Free LRU cache
void free_lru_cache(LRUCache *cache) {
    if (cache == NULL) return;
```

```

    for (int i = 0; i < CACHE_SIZE; i++) {
        CacheNode *current = cache->buckets[i];
        while (current != NULL) {
            CacheNode *next = current->next;
            free(current);
            current = next;
        }
    }

    free(cache);
}

// Hash function
unsigned int cache_hash(const char *key) {
    unsigned long hash = 5381;
    int c;

    while ((c = *key++)) {
        hash = ((hash << 5) + hash) + c;
    }

    return hash % CACHE_SIZE;
}

// Remove least recently used item
void remove_lru_item(LRUCache *cache) {
    if (cache == NULL || cache->size == 0) return;

    time_t oldest_time = time(NULL);
    CacheNode *oldest_node = NULL;
    int oldest_bucket = -1;
    CacheNode *oldest_prev = NULL;

    // Find the oldest accessed item
    for (int i = 0; i < CACHE_SIZE; i++) {
        CacheNode *current = cache->buckets[i];
        CacheNode *prev = NULL;

        while (current != NULL) {
            if (current->timestamp < oldest_time) {
                oldest_time = current->timestamp;
                oldest_node = current;
                oldest_bucket = i;
                oldest_prev = prev;
            }
            prev = current;
            current = current->next;
        }

        // Remove the oldest node
        if (oldest_node != NULL) {
            if (oldest_prev != NULL)
                oldest_prev->next = oldest_node->next;
            else
                cache->buckets[oldest_bucket] = oldest_node->next;
            free(oldest_node);
        }
    }
    cache->size--;
}

```

```

        current = current->next;
    }
}

// Remove the oldest item
if (oldest_node != NULL) {
    if (oldest_prev == NULL) {
        cache->buckets[oldest_bucket] = oldest_node->next;
    } else {
        oldest_prev->next = oldest_node->next;
    }
    free(oldest_node);
    cache->size--;
}

}

// Put item in cache
bool cache_put(LRUCache *cache, const char *key, const char *value) {
    if (cache == NULL || key == NULL || value == NULL) return false;

    // Check if key already exists
    unsigned int index = cache_hash(key);
    CacheNode *current = cache->buckets[index];
    CacheNode *prev = NULL;

    while (current != NULL) {
        if (strcmp(current->key, key) == 0) {
            // Update existing item
            strncpy(current->value, value, MAX_VALUE_LENGTH - 1);
            current->value[MAX_VALUE_LENGTH - 1] = '\0';
            current->timestamp = time(NULL);
            current->access_count++;
            return true;
        }
        prev = current;
        current = current->next;
    }

    // Check if cache is full
    if (cache->size >= cache->max_size) {
        remove_lru_item(cache);
    }

    // Create new cache node
    CacheNode *new_node = malloc(sizeof(CacheNode));
    if (new_node == NULL) return false;

```

```

    strncpy(new_node->key, key, MAX_KEY_LENGTH - 1);
    new_node->key[MAX_KEY_LENGTH - 1] = '\0';
    strncpy(new_node->value, value, MAX_VALUE_LENGTH - 1);
    new_node->value[MAX_VALUE_LENGTH - 1] = '\0';
    new_node->timestamp = time(NULL);
    new_node->access_count = 1;
    new_node->next = cache->buckets[index];
    cache->buckets[index] = new_node;
    cache->size++;

    return true;
}

// Get item from cache
bool cache_get(LRUCache *cache, const char *key, char *value, size_t value_size) {
    if (cache == NULL || key == NULL || value == NULL) return false;

    unsigned int index = cache_hash(key);
    CacheNode *current = cache->buckets[index];

    while (current != NULL) {
        if (strcmp(current->key, key) == 0) {
            // Found item, update timestamp and access count
            strncpy(value, current->value, value_size - 1);
            value[value_size - 1] = '\0';
            current->timestamp = time(NULL);
            current->access_count++;
            return true;
        }
        current = current->next;
    }

    return false;
}

// Print cache statistics
void print_cache_stats(LRUCache *cache) {
    if (cache == NULL) return;

    printf("Cache Statistics:\n");
    printf("  Size: %d/%d\n", cache->size, cache->max_size);
    printf("  Buckets:\n");

    for (int i = 0; i < CACHE_SIZE; i++) {
        CacheNode *current = cache->buckets[i];
        if (current != NULL) {
            printf("    Bucket %d:\n", i);

```

```

        while (current != NULL) {
            printf("      Key: %-15s Value: %-20s Accesses: %d\n",
                current->key, current->value, current->access_count);
            current = current->next;
        }
    }
}

int main() {
    LRUCache *cache = create_lru_cache(5);
    if (cache == NULL) {
        printf("Failed to create cache\n");
        return 1;
    }

    // Add items to cache
    cache_put(cache, "user1", "John Doe");
    cache_put(cache, "user2", "Jane Smith");
    cache_put(cache, "config", "debug=true");
    cache_put(cache, "version", "1.2.3");
    cache_put(cache, "theme", "dark");

    printf("Initial cache state:\n");
    print_cache_stats(cache);

    // Access some items
    char value[100];
    if (cache_get(cache, "user1", value, sizeof(value))) {
        printf("\nRetrieved 'user1': %s\n", value);
    }

    if (cache_get(cache, "config", value, sizeof(value))) {
        printf("Retrieved 'config': %s\n", value);
    }

    // Add more items to trigger LRU eviction
    cache_put(cache, "new_item1", "value1");
    cache_put(cache, "new_item2", "value2");

    printf("\nCache after adding new items:\n");
    print_cache_stats(cache);

    // Try to access evicted item
    if (cache_get(cache, "user2", value, sizeof(value))) {
        printf("\nRetrieved 'user2': %s\n", value);
    } else {

```



```
        printf("\n'user2' was evicted from cache\n");
    }

    free_lru_cache(cache);
    return 0;
}
```

Summary

Hash tables are powerful data structures that provide efficient key-value storage and retrieval:

1. **Hash Functions:** Critical for distributing keys uniformly across the table
2. **Collision Resolution:**
 - Chaining: Store collided elements in linked lists
 - Open Addressing: Probe for alternative slots in the table
3. **Dynamic Resizing:** Maintain performance by resizing when load factor is too high
4. **Performance Characteristics:**
 - Average case: $O(1)$ for insert, delete, and search
 - Worst case: $O(n)$ when all keys collide
 - Space complexity: $O(n)$

Key considerations when implementing hash tables: - **Hash Function Quality:** Should minimize collisions and distribute keys uniformly - **Load Factor Management:** Resize when load factor exceeds threshold - **Collision Resolution Strategy:** Choose based on use case and performance requirements - **Memory Management:** Properly handle dynamic allocation and deallocation

Hash tables are widely used in: - Database indexing - Caching systems - Symbol tables in compilers - Sets and maps in programming languages - Network routers and switches - Cryptographic applications

Understanding hash table implementation and optimization techniques is essential for building high-performance applications that require fast data access.

Graph Algorithms

Introduction

Graphs are fundamental data structures used to model relationships between objects. A graph consists of vertices (nodes) connected by edges, and can represent various real-world scenarios such as social networks, transportation networks, computer networks, and dependency relationships.

In this chapter, we'll explore graph representations, traversal algorithms, shortest path algorithms, and minimum spanning tree algorithms. We'll implement these algorithms in C and examine their applications and performance characteristics.

Graph Representations

Graphs can be represented in several ways, each with its own advantages and disadvantages.

Adjacency Matrix

An adjacency matrix is a 2D array where the cell at position (i, j) indicates whether there is an edge from vertex i to vertex j.

```
#include <stdio.h>
#include <stdlib.h>
#include <stdbool.h>
#include <limits.h>

#define MAX_VERTICES 100

typedef struct {
    int vertices;
    int matrix[MAX_VERTICES][MAX_VERTICES];
} AdjMatrixGraph;

// Create adjacency matrix graph
AdjMatrixGraph* create_adj_matrix_graph(int vertices) {
    AdjMatrixGraph *graph = malloc(sizeof(AdjMatrixGraph));
    if (graph == NULL) return NULL;

    graph->vertices = vertices;
```

```

    // Initialize matrix with 0s (no edges)
    for (int i = 0; i < vertices; i++) {
        for (int j = 0; j < vertices; j++) {
            graph->matrix[i][j] = 0;
        }
    }

    return graph;
}

// Add edge to adjacency matrix graph
void add_edge_matrix(AdjMatrixGraph *graph, int src, int dest, int weight) {
    if (graph == NULL || src < 0 || src >= graph->vertices ||
        dest < 0 || dest >= graph->vertices) {
        return;
    }

    graph->matrix[src][dest] = weight;
    // For undirected graph, also add reverse edge
    graph->matrix[dest][src] = weight;
}

// Remove edge from adjacency matrix graph
void remove_edge_matrix(AdjMatrixGraph *graph, int src, int dest) {
    if (graph == NULL || src < 0 || src >= graph->vertices ||
        dest < 0 || dest >= graph->vertices) {
        return;
    }

    graph->matrix[src][dest] = 0;
    graph->matrix[dest][src] = 0;
}

// Check if edge exists
bool has_edge_matrix(AdjMatrixGraph *graph, int src, int dest) {
    if (graph == NULL || src < 0 || src >= graph->vertices ||
        dest < 0 || dest >= graph->vertices) {
        return false;
    }

    return graph->matrix[src][dest] != 0;
}

// Print adjacency matrix
void print_adj_matrix(AdjMatrixGraph *graph) {
    if (graph == NULL) return;

```

```

    printf("Adjacency Matrix:\n");
    printf("    ");
    for (int i = 0; i < graph->vertices; i++) {
        printf("%3d", i);
    }
    printf("\n");

    for (int i = 0; i < graph->vertices; i++) {
        printf("%3d ", i);
        for (int j = 0; j < graph->vertices; j++) {
            printf("%3d", graph->matrix[i][j]);
        }
        printf("\n");
    }
}

int main() {
    AdjMatrixGraph *graph = create_adj_matrix_graph(5);
    if (graph == NULL) {
        printf("Failed to create graph\n");
        return 1;
    }

    // Add edges
    add_edge_matrix(graph, 0, 1, 1);
    add_edge_matrix(graph, 0, 4, 1);
    add_edge_matrix(graph, 1, 2, 1);
    add_edge_matrix(graph, 1, 3, 1);
    add_edge_matrix(graph, 1, 4, 1);
    add_edge_matrix(graph, 2, 3, 1);
    add_edge_matrix(graph, 3, 4, 1);

    print_adj_matrix(graph);

    printf("\nEdge checks:\n");
    printf("Edge (0,1): %s\n", has_edge_matrix(graph, 0, 1) ? "Yes" : "No");
    printf("Edge (0,2): %s\n", has_edge_matrix(graph, 0, 2) ? "Yes" : "No");

    free(graph);
    return 0;
}

```

Adjacency List

An adjacency list represents a graph as an array of lists, where each list contains the neighbors of a vertex.

```

#include <stdio.h>
#include <stdlib.h>
#include <stdbool.h>

#define MAX_VERTICES 100

typedef struct AdjListNode {
    int vertex;
    int weight;
    struct AdjListNode *next;
} AdjListNode;

typedef struct {
    int vertices;
    AdjListNode *adj_list[MAX_VERTICES];
} AdjListGraph;

// Create new adjacency list node
AdjListNode* create_adj_list_node(int vertex, int weight) {
    AdjListNode *node = malloc(sizeof(AdjListNode));
    if (node == NULL) return NULL;

    node->vertex = vertex;
    node->weight = weight;
    node->next = NULL;
    return node;
}

// Create adjacency list graph
AdjListGraph* create_adj_list_graph(int vertices) {
    AdjListGraph *graph = malloc(sizeof(AdjListGraph));
    if (graph == NULL) return NULL;

    graph->vertices = vertices;

    // Initialize adjacency lists
    for (int i = 0; i < vertices; i++) {
        graph->adj_list[i] = NULL;
    }

    return graph;
}

// Free adjacency list graph
void free_adj_list_graph(AdjListGraph *graph) {
    if (graph == NULL) return;

```

```

    for (int i = 0; i < graph->vertices; i++) {
        AdjListNode *current = graph->adj_list[i];
        while (current != NULL) {
            AdjListNode *next = current->next;
            free(current);
            current = next;
        }
    }

    free(graph);
}

// Add edge to adjacency list graph
void add_edge_list(AdjListGraph *graph, int src, int dest, int weight) {
    if (graph == NULL || src < 0 || src >= graph->vertices ||
        dest < 0 || dest >= graph->vertices) {
        return;
    }

    // Add edge from src to dest
    AdjListNode *new_node = create_adj_list_node(dest, weight);
    new_node->next = graph->adj_list[src];
    graph->adj_list[src] = new_node;

    // For undirected graph, also add reverse edge
    new_node = create_adj_list_node(src, weight);
    new_node->next = graph->adj_list[dest];
    graph->adj_list[dest] = new_node;
}

// Remove edge from adjacency list graph
void remove_edge_list(AdjListGraph *graph, int src, int dest) {
    if (graph == NULL || src < 0 || src >= graph->vertices ||
        dest < 0 || dest >= graph->vertices) {
        return;
    }

    // Remove edge from src to dest
    AdjListNode *current = graph->adj_list[src];
    AdjListNode *prev = NULL;

    while (current != NULL) {
        if (current->vertex == dest) {
            if (prev == NULL) {
                graph->adj_list[src] = current->next;
            } else {
                prev->next = current->next;
            }
        }
        prev = current;
        current = current->next;
    }
}

```

```

        }
        free(current);
        break;
    }
    prev = current;
    current = current->next;
}

// Remove edge from dest to src
current = graph->adj_list[dest];
prev = NULL;

while (current != NULL) {
    if (current->vertex == src) {
        if (prev == NULL) {
            graph->adj_list[dest] = current->next;
        } else {
            prev->next = current->next;
        }
        free(current);
        break;
    }
    prev = current;
    current = current->next;
}

}

// Print adjacency list
void print_adj_list(AdjListGraph *graph) {
    if (graph == NULL) return;

    printf("Adjacency List:\n");
    for (int i = 0; i < graph->vertices; i++) {
        printf("%d: ", i);
        AdjListNode *current = graph->adj_list[i];
        while (current != NULL) {
            printf("(%d, %d) -> ", current->vertex, current->weight);
            current = current->next;
        }
        printf("NULL\n");
    }
}

int main() {
    AdjListGraph *graph = create_adj_list_graph(5);
    if (graph == NULL) {

```

```

        printf("Failed to create graph\n");
        return 1;
    }

    // Add edges
    add_edge_list(graph, 0, 1, 1);
    add_edge_list(graph, 0, 4, 1);
    add_edge_list(graph, 1, 2, 1);
    add_edge_list(graph, 1, 3, 1);
    add_edge_list(graph, 1, 4, 1);
    add_edge_list(graph, 2, 3, 1);
    add_edge_list(graph, 3, 4, 1);

    print_adj_list(graph);

    free_adj_list_graph(graph);
    return 0;
}

```

Graph Traversal Algorithms

Breadth-First Search (BFS)

BFS explores vertices level by level, visiting all neighbors of a vertex before moving to the next level.

```

#include <stdio.h>
#include <stdlib.h>
#include <stdbool.h>

#define MAX_VERTICES 100

typedef struct Queue {
    int items[MAX_VERTICES];
    int front;
    int rear;
} Queue;

// Initialize queue
Queue* create_queue(void) {
    Queue *queue = malloc(sizeof(Queue));
    if (queue == NULL) return NULL;

    queue->front = 0;
    queue->rear = -1;
}

```



```

    return queue;
}

// Check if queue is empty
bool is_queue_empty(Queue *queue) {
    return queue->rear < queue->front;
}

// Enqueue
void enqueue(Queue *queue, int item) {
    if (queue->rear < MAX_VERTICES - 1) {
        queue->items[++queue->rear] = item;
    }
}

// Dequeue
int dequeue(Queue *queue) {
    if (is_queue_empty(queue)) {
        return -1;
    }
    return queue->items[queue->front++];
}

typedef struct BFSGraph {
    int vertices;
    struct AdjListNode **adj_list;
} BFSGraph;

typedef struct AdjListNode {
    int vertex;
    struct AdjListNode *next;
} AdjListNode;

// Create new adjacency list node
AdjListNode* create_bfs_node(int vertex) {
    AdjListNode *node = malloc(sizeof(AdjListNode));
    if (node == NULL) return NULL;

    node->vertex = vertex;
    node->next = NULL;
    return node;
}

// Create BFS graph
BFSGraph* create_bfs_graph(int vertices) {
    BFSGraph *graph = malloc(sizeof(BFSGraph));
    if (graph == NULL) return NULL;

```

```

graph->vertices = vertices;
graph->adj_list = malloc(vertices * sizeof(AdjListNode*));
if (graph->adj_list == NULL) {
    free(graph);
    return NULL;
}

for (int i = 0; i < vertices; i++) {
    graph->adj_list[i] = NULL;
}

return graph;
}

// Free BFS graph
void free_bfs_graph(BFSGraph *graph) {
    if (graph == NULL) return;

    for (int i = 0; i < graph->vertices; i++) {
        AdjListNode *current = graph->adj_list[i];
        while (current != NULL) {
            AdjListNode *next = current->next;
            free(current);
            current = next;
        }
    }

    free(graph->adj_list);
    free(graph);
}

// Add edge to BFS graph
void add_bfs_edge(BFSGraph *graph, int src, int dest) {
    if (graph == NULL || src < 0 || src >= graph->vertices ||
        dest < 0 || dest >= graph->vertices) {
        return;
    }

    // Add edge from src to dest
    AdjListNode *new_node = create_bfs_node(dest);
    new_node->next = graph->adj_list[src];
    graph->adj_list[src] = new_node;

    // For undirected graph, also add reverse edge
    new_node = create_bfs_node(src);
    new_node->next = graph->adj_list[dest];

```

```

    graph->adj_list[dest] = new_node;
}

// BFS traversal
void bfs(BFSGraph *graph, int start_vertex) {
    if (graph == NULL || start_vertex < 0 || start_vertex >= graph->vertices) {
        return;
    }

    bool *visited = calloc(graph->vertices, sizeof(bool));
    if (visited == NULL) return;

    Queue *queue = create_queue();
    if (queue == NULL) {
        free(visited);
        return;
    }

    visited[start_vertex] = true;
    enqueue(queue, start_vertex);

    printf("BFS Traversal starting from vertex %d: ", start_vertex);

    while (!is_queue_empty(queue)) {
        int current_vertex = dequeue(queue);
        printf("%d ", current_vertex);

        AdjListNode *current = graph->adj_list[current_vertex];
        while (current != NULL) {
            int adj_vertex = current->vertex;
            if (!visited[adj_vertex]) {
                visited[adj_vertex] = true;
                enqueue(queue, adj_vertex);
            }
            current = current->next;
        }
    }

    printf("\n");

    free(visited);
    free(queue);
}

int main() {
    BFSGraph *graph = create_bfs_graph(6);

```

```

    if (graph == NULL) {
        printf("Failed to create graph\n");
        return 1;
    }

    // Add edges to create a sample graph
    add_bfs_edge(graph, 0, 1);
    add_bfs_edge(graph, 0, 2);
    add_bfs_edge(graph, 1, 3);
    add_bfs_edge(graph, 1, 4);
    add_bfs_edge(graph, 2, 4);
    add_bfs_edge(graph, 3, 5);
    add_bfs_edge(graph, 4, 5);

    bfs(graph, 0);
    bfs(graph, 2);

    free_bfs_graph(graph);
    return 0;
}

```

Depth-First Search (DFS)

DFS explores as far as possible along each branch before backtracking.

```

#include <stdio.h>
#include <stdlib.h>
#include <stdbool.h>

#define MAX_VERTICES 100

typedef struct DFSGraph {
    int vertices;
    struct AdjListNode **adj_list;
} DFSGraph;

typedef struct AdjListNode {
    int vertex;
    struct AdjListNode *next;
} AdjListNode;

// Create new adjacency list node
AdjListNode* create_dfs_node(int vertex) {
    AdjListNode *node = malloc(sizeof(AdjListNode));
    if (node == NULL) return NULL;

```

```

    node->vertex = vertex;
    node->next = NULL;
    return node;
}

// Create DFS graph
DFSGraph* create_dfs_graph(int vertices) {
    DFSGraph *graph = malloc(sizeof(DFSGraph));
    if (graph == NULL) return NULL;

    graph->vertices = vertices;
    graph->adj_list = malloc(vertices * sizeof(AdjListNode*));
    if (graph->adj_list == NULL) {
        free(graph);
        return NULL;
    }

    for (int i = 0; i < vertices; i++) {
        graph->adj_list[i] = NULL;
    }

    return graph;
}

// Free DFS graph
void free_dfs_graph(DFSGraph *graph) {
    if (graph == NULL) return;

    for (int i = 0; i < graph->vertices; i++) {
        AdjListNode *current = graph->adj_list[i];
        while (current != NULL) {
            AdjListNode *next = current->next;
            free(current);
            current = next;
        }
    }

    free(graph->adj_list);
    free(graph);
}

// Add edge to DFS graph
void add_dfs_edge(DFSGraph *graph, int src, int dest) {
    if (graph == NULL || src < 0 || src >= graph->vertices ||
        dest < 0 || dest >= graph->vertices) {
        return;
    }
}

```

```

// Add edge from src to dest
AdjListNode *new_node = create_dfs_node(dest);
new_node->next = graph->adj_list[src];
graph->adj_list[src] = new_node;

// For undirected graph, also add reverse edge
new_node = create_dfs_node(src);
new_node->next = graph->adj_list[dest];
graph->adj_list[dest] = new_node;
}

// Recursive DFS helper function
void dfs_recursive(DFSGraph *graph, int vertex, bool *visited) {
    visited[vertex] = true;
    printf("%d ", vertex);

    AdjListNode *current = graph->adj_list[vertex];
    while (current != NULL) {
        int adj_vertex = current->vertex;
        if (!visited[adj_vertex]) {
            dfs_recursive(graph, adj_vertex, visited);
        }
        current = current->next;
    }
}

// DFS traversal
void dfs(DFSGraph *graph, int start_vertex) {
    if (graph == NULL || start_vertex < 0 || start_vertex >= graph->vertices) {
        return;
    }

    bool *visited = calloc(graph->vertices, sizeof(bool));
    if (visited == NULL) return;

    printf("DFS Traversal starting from vertex %d: ", start_vertex);
    dfs_recursive(graph, start_vertex, visited);
    printf("\n");

    free(visited);
}

// Iterative DFS using stack
void dfs_iterative(DFSGraph *graph, int start_vertex) {
    if (graph == NULL || start_vertex < 0 || start_vertex >= graph->vertices) {
        return;
    }

```

```

}

bool *visited = calloc(graph->vertices, sizeof(bool));
if (visited == NULL) return;

int stack[MAX_VERTICES];
int top = -1;

stack[++top] = start_vertex;

printf("DFS Iterative Traversal starting from vertex %d: ", start_vertex);

while (top >= 0) {
    int current_vertex = stack[top--];

    if (!visited[current_vertex]) {
        visited[current_vertex] = true;
        printf("%d ", current_vertex);

        // Push all adjacent vertices to stack
        AdjListNode *current = graph->adj_list[current_vertex];
        while (current != NULL) {
            int adj_vertex = current->vertex;
            if (!visited[adj_vertex]) {
                stack[++top] = adj_vertex;
            }
            current = current->next;
        }
    }
}

printf("\n");

free(visited);
}

int main() {
    DFSGraph *graph = create_dfs_graph(6);
    if (graph == NULL) {
        printf("Failed to create graph\n");
        return 1;
    }

    // Add edges to create a sample graph
    add_dfs_edge(graph, 0, 1);
    add_dfs_edge(graph, 0, 2);

```

```

    add_dfs_edge(graph, 1, 3);
    add_dfs_edge(graph, 1, 4);
    add_dfs_edge(graph, 2, 4);
    add_dfs_edge(graph, 3, 5);
    add_dfs_edge(graph, 4, 5);

    dfs(graph, 0);
    dfs_iterative(graph, 0);

    free_dfs_graph(graph);
    return 0;
}

```

Shortest Path Algorithms

Dijkstra's Algorithm

Dijkstra's algorithm finds the shortest path from a source vertex to all other vertices in a weighted graph with non-negative edge weights.

```

#include <stdio.h>
#include <stdlib.h>
#include <stdbool.h>
#include <limits.h>

#define MAX_VERTICES 100

typedef struct DijkstraGraph {
    int vertices;
    struct DijkstraNode **adj_list;
} DijkstraGraph;

typedef struct DijkstraNode {
    int vertex;
    int weight;
    struct DijkstraNode *next;
} DijkstraNode;

// Create new adjacency list node
DijkstraNode* create_dijkstra_node(int vertex, int weight) {
    DijkstraNode *node = malloc(sizeof(DijkstraNode));
    if (node == NULL) return NULL;

    node->vertex = vertex;
    node->weight = weight;
}

```



```

    node->next = NULL;
    return node;
}

// Create Dijkstra graph
DijkstraGraph* create_dijkstra_graph(int vertices) {
    DijkstraGraph *graph = malloc(sizeof(DijkstraGraph));
    if (graph == NULL) return NULL;

    graph->vertices = vertices;
    graph->adj_list = malloc(vertices * sizeof(DijkstraNode*));
    if (graph->adj_list == NULL) {
        free(graph);
        return NULL;
    }

    for (int i = 0; i < vertices; i++) {
        graph->adj_list[i] = NULL;
    }

    return graph;
}

// Free Dijkstra graph
void free_dijkstra_graph(DijkstraGraph *graph) {
    if (graph == NULL) return;

    for (int i = 0; i < graph->vertices; i++) {
        DijkstraNode *current = graph->adj_list[i];
        while (current != NULL) {
            DijkstraNode *next = current->next;
            free(current);
            current = next;
        }
    }

    free(graph->adj_list);
    free(graph);
}

// Add edge to Dijkstra graph
void add_dijkstra_edge(DijkstraGraph *graph, int src, int dest, int weight) {
    if (graph == NULL || src < 0 || src >= graph->vertices ||
        dest < 0 || dest >= graph->vertices) {
        return;
    }
}

```

```

    // Add edge from src to dest
    DijkstraNode *new_node = create_dijkstra_node(dest, weight);
    new_node->next = graph->adj_list[src];
    graph->adj_list[src] = new_node;

    // For undirected graph, also add reverse edge
    new_node = create_dijkstra_node(src, weight);
    new_node->next = graph->adj_list[dest];
    graph->adj_list[dest] = new_node;
}

// Min heap implementation for Dijkstra's algorithm
typedef struct {
    int *vertices;
    int *distances;
    int size;
    int capacity;
    int *pos; // Position of vertices in heap
} MinHeap;

// Create min heap
MinHeap* create_min_heap(int capacity) {
    MinHeap *heap = malloc(sizeof(MinHeap));
    if (heap == NULL) return NULL;

    heap->capacity = capacity;
    heap->size = 0;
    heap->vertices = malloc(capacity * sizeof(int));
    heap->distances = malloc(capacity * sizeof(int));
    heap->pos = malloc(capacity * sizeof(int));

    if (heap->vertices == NULL || heap->distances == NULL || heap->pos == NULL) {
        free(heap->vertices);
        free(heap->distances);
        free(heap->pos);
        free(heap);
        return NULL;
    }

    return heap;
}

// Free min heap
void free_min_heap(MinHeap *heap) {
    if (heap == NULL) return;

    free(heap->vertices);

```

```

    free(heap->distances);
    free(heap->pos);
    free(heap);
}

// Check if heap is empty
bool is_heap_empty(MinHeap *heap) {
    return heap->size == 0;
}

// Min heapify operation
void min_heapify(MinHeap *heap, int idx) {
    int smallest = idx;
    int left = 2 * idx + 1;
    int right = 2 * idx + 2;

    if (left < heap->size && heap->distances[left] < heap->distances[smallest]) {
        smallest = left;
    }

    if (right < heap->size && heap->distances[right] < heap->distances[smallest]) {
        smallest = right;
    }

    if (smallest != idx) {
        // Swap positions
        int temp_pos = heap->pos[heap->vertices[smallest]];
        heap->pos[heap->vertices[smallest]] = heap->pos[heap->vertices[idx]];
        heap->pos[heap->vertices[idx]] = temp_pos;

        // Swap vertices
        int temp_vertex = heap->vertices[smallest];
        heap->vertices[smallest] = heap->vertices[idx];
        heap->vertices[idx] = temp_vertex;

        // Swap distances
        int temp_dist = heap->distances[smallest];
        heap->distances[smallest] = heap->distances[idx];
        heap->distances[idx] = temp_dist;

        min_heapify(heap, smallest);
    }
}

// Extract minimum distance vertex
int extract_min(MinHeap *heap) {

```

```

    if (is_heap_empty(heap)) return -1;

    int root_vertex = heap->vertices[0];

    // Replace root with last element
    int last_vertex = heap->vertices[heap->size - 1];
    heap->vertices[0] = last_vertex;

    // Update position of last vertex
    heap->pos[root_vertex] = heap->size - 1;
    heap->pos[last_vertex] = 0;

    // Reduce heap size and heapify root
    heap->size--;
    min_heapify(heap, 0);

    return root_vertex;
}

// Decrease key operation
void decrease_key(MinHeap *heap, int vertex, int new_dist) {
    int idx = heap->pos[vertex];
    heap->distances[idx] = new_dist;

    // Travel up while heap property is violated
    while (idx && heap->distances[idx] < heap->distances[(idx - 1) / 2]) {
        // Swap positions
        int temp_pos = heap->pos[heap->vertices[(idx - 1) / 2]];
        heap->pos[heap->vertices[(idx - 1) / 2]] = heap->pos[heap->vertices[idx]];
        heap->pos[heap->vertices[idx]] = temp_pos;

        // Swap vertices
        int temp_vertex = heap->vertices[(idx - 1) / 2];
        heap->vertices[(idx - 1) / 2] = heap->vertices[idx];
        heap->vertices[idx] = temp_vertex;

        // Swap distances
        int temp_dist = heap->distances[(idx - 1) / 2];
        heap->distances[(idx - 1) / 2] = heap->distances[idx];
        heap->distances[idx] = temp_dist;

        idx = (idx - 1) / 2;
    }
}

// Check if vertex is in min heap
bool is_in_min_heap(MinHeap *heap, int vertex) {

```

```

    return (heap->pos[vertex] < heap->size);
}

// Dijkstra's algorithm
void dijkstra(DijkstraGraph *graph, int src) {
    if (graph == NULL || src < 0 || src >= graph->vertices) {
        return;
    }

    int *distances = malloc(graph->vertices * sizeof(int));
    int *parent = malloc(graph->vertices * sizeof(int));
    bool *visited = calloc(graph->vertices, sizeof(bool));

    if (distances == NULL || parent == NULL || visited == NULL) {
        free(distances);
        free(parent);
        free(visited);
        return;
    }

    // Initialize distances and parent array
    for (int i = 0; i < graph->vertices; i++) {
        distances[i] = INT_MAX;
        parent[i] = -1;
    }

    distances[src] = 0;

    // Create min heap
    MinHeap *min_heap = create_min_heap(graph->vertices);
    if (min_heap == NULL) {
        free(distances);
        free(parent);
        free(visited);
        return;
    }

    // Initialize min heap with all vertices
    for (int i = 0; i < graph->vertices; i++) {
        min_heap->vertices[i] = i;
        min_heap->distances[i] = distances[i];
        min_heap->pos[i] = i;
    }

    min_heap->size = graph->vertices;

```

```

// Main loop of Dijkstra's algorithm
while (!is_heap_empty(min_heap)) {
    // Extract minimum distance vertex
    int u = extract_min(min_heap);
    visited[u] = true;

    // Update distance values of adjacent vertices
    DijkstraNode *current = graph->adj_list[u];
    while (current != NULL) {
        int v = current->vertex;

        // If shortest distance to v is not finalized yet,
        // and distance to v through u is less than its current distance
        if (!visited[v] && distances[u] != INT_MAX &&
            distances[u] + current->weight < distances[v]) {
            distances[v] = distances[u] + current->weight;
            parent[v] = u;

            // Update distance value in min heap
            decrease_key(min_heap, v, distances[v]);
        }

        current = current->next;
    }
}

// Print shortest distances
printf("Shortest distances from vertex %d:\n", src);
for (int i = 0; i < graph->vertices; i++) {
    if (distances[i] == INT_MAX) {
        printf("Vertex %d: INF\n", i);
    } else {
        printf("Vertex %d: %d\n", i, distances[i]);
    }
}

free(distances);
free(parent);
free(visited);
free_min_heap(min_heap);
}

int main() {
    DijkstraGraph *graph = create_dijkstra_graph(6);
    if (graph == NULL) {
        printf("Failed to create graph\n");
        return 1;
    }
}

```

```

    }

    // Add weighted edges
    add_dijkstra_edge(graph, 0, 1, 4);
    add_dijkstra_edge(graph, 0, 2, 3);
    add_dijkstra_edge(graph, 1, 2, 1);
    add_dijkstra_edge(graph, 1, 3, 2);
    add_dijkstra_edge(graph, 2, 3, 4);
    add_dijkstra_edge(graph, 3, 4, 2);
    add_dijkstra_edge(graph, 4, 5, 6);

    dijkstra(graph, 0);

    free_dijkstra_graph(graph);
    return 0;
}

```

Minimum Spanning Tree Algorithms

Prim's Algorithm

Prim's algorithm finds a minimum spanning tree for a weighted undirected graph.

```

#include <stdio.h>
#include <stdlib.h>
#include <stdbool.h>
#include <limits.h>

#define MAX_VERTICES 100

typedef struct PrimGraph {
    int vertices;
    struct PrimNode **adj_list;
} PrimGraph;

typedef struct PrimNode {
    int vertex;
    int weight;
    struct PrimNode *next;
} PrimNode;

// Create new adjacency list node
PrimNode* create_prim_node(int vertex, int weight) {
    PrimNode *node = malloc(sizeof(PrimNode));
    if (node == NULL) return NULL;

```

```

    node->vertex = vertex;
    node->weight = weight;
    node->next = NULL;
    return node;
}

// Create Prim graph
PrimGraph* create_prim_graph(int vertices) {
    PrimGraph *graph = malloc(sizeof(PrimGraph));
    if (graph == NULL) return NULL;

    graph->vertices = vertices;
    graph->adj_list = malloc(vertices * sizeof(PrimNode*));
    if (graph->adj_list == NULL) {
        free(graph);
        return NULL;
    }

    for (int i = 0; i < vertices; i++) {
        graph->adj_list[i] = NULL;
    }

    return graph;
}

// Free Prim graph
void free_prim_graph(PrimGraph *graph) {
    if (graph == NULL) return;

    for (int i = 0; i < graph->vertices; i++) {
        PrimNode *current = graph->adj_list[i];
        while (current != NULL) {
            PrimNode *next = current->next;
            free(current);
            current = next;
        }
    }

    free(graph->adj_list);
    free(graph);
}

// Add edge to Prim graph
void add_prim_edge(PrimGraph *graph, int src, int dest, int weight) {
    if (graph == NULL || src < 0 || src >= graph->vertices ||
        dest < 0 || dest >= graph->vertices) {

```



```

        return;
    }

    // Add edge from src to dest
    PrimNode *new_node = create_prim_node(dest, weight);
    new_node->next = graph->adj_list[src];
    graph->adj_list[src] = new_node;

    // For undirected graph, also add reverse edge
    new_node = create_prim_node(src, weight);
    new_node->next = graph->adj_list[dest];
    graph->adj_list[dest] = new_node;
}

// Min heap implementation for Prim's algorithm
typedef struct {
    int *vertices;
    int *keys;
    int *parent;
    int size;
    int capacity;
    int *pos; // Position of vertices in heap
} PrimHeap;

// Create min heap for Prim's algorithm
PrimHeap* create_prim_heap(int capacity) {
    PrimHeap *heap = malloc(sizeof(PrimHeap));
    if (heap == NULL) return NULL;

    heap->capacity = capacity;
    heap->size = 0;
    heap->vertices = malloc(capacity * sizeof(int));
    heap->keys = malloc(capacity * sizeof(int));
    heap->parent = malloc(capacity * sizeof(int));
    heap->pos = malloc(capacity * sizeof(int));

    if (heap->vertices == NULL || heap->keys == NULL ||
        heap->parent == NULL || heap->pos == NULL) {
        free(heap->vertices);
        free(heap->keys);
        free(heap->parent);
        free(heap->pos);
        free(heap);
        return NULL;
    }

    return heap;
}

```

```

}

// Free Prim heap
void free_prim_heap(PrimHeap *heap) {
    if (heap == NULL) return;

    free(heap->vertices);
    free(heap->keys);
    free(heap->parent);
    free(heap->pos);
    free(heap);
}

// Check if heap is empty
bool is_prim_heap_empty(PrimHeap *heap) {
    return heap->size == 0;
}

// Min heapify operation
void prim_heapify(PrimHeap *heap, int idx) {
    int smallest = idx;
    int left = 2 * idx + 1;
    int right = 2 * idx + 2;

    if (left < heap->size && heap->keys[left] < heap->keys[smallest]) {
        smallest = left;
    }

    if (right < heap->size && heap->keys[right] < heap->keys[smallest]) {
        smallest = right;
    }

    if (smallest != idx) {
        // Swap positions
        int temp_pos = heap->pos[heap->vertices[smallest]];
        heap->pos[heap->vertices[smallest]] = heap->pos[heap->vertices[idx]];
        heap->pos[heap->vertices[idx]] = temp_pos;

        // Swap vertices
        int temp_vertex = heap->vertices[smallest];
        heap->vertices[smallest] = heap->vertices[idx];
        heap->vertices[idx] = temp_vertex;

        // Swap keys
        int temp_key = heap->keys[smallest];
        heap->keys[smallest] = heap->keys[idx];
    }
}

```

```

        heap->keys[idx] = temp_key;

        prim_heapify(heap, smallest);
    }
}

// Extract minimum key vertex
int extract_min_prim(PrimHeap *heap) {
    if (is_prim_heap_empty(heap)) return -1;

    int root_vertex = heap->vertices[0];

    // Replace root with last element
    int last_vertex = heap->vertices[heap->size - 1];
    heap->vertices[0] = last_vertex;

    // Update position of last vertex
    heap->pos[root_vertex] = heap->size - 1;
    heap->pos[last_vertex] = 0;

    // Reduce heap size and heapify root
    heap->size--;
    prim_heapify(heap, 0);

    return root_vertex;
}

// Decrease key operation
void decrease_key_prim(PrimHeap *heap, int vertex, int new_key) {
    int idx = heap->pos[vertex];
    heap->keys[idx] = new_key;

    // Travel up while heap property is violated
    while (idx && heap->keys[idx] < heap->keys[(idx - 1) / 2]) {
        // Swap positions
        int temp_pos = heap->pos[heap->vertices[(idx - 1) / 2]];
        heap->pos[heap->vertices[(idx - 1) / 2]] = heap->pos[heap->vertices[idx]];
        heap->pos[heap->vertices[idx]] = temp_pos;

        // Swap vertices
        int temp_vertex = heap->vertices[(idx - 1) / 2];
        heap->vertices[(idx - 1) / 2] = heap->vertices[idx];
        heap->vertices[idx] = temp_vertex;

        // Swap keys
        int temp_key = heap->keys[(idx - 1) / 2];

```

```

        heap->keys[(idx - 1) / 2] = heap->keys[idx];
        heap->keys[idx] = temp_key;

        idx = (idx - 1) / 2;
    }
}

// Prim's algorithm for Minimum Spanning Tree
void prim_mst(PrimGraph *graph) {
    if (graph == NULL) return;

    int *parent = malloc(graph->vertices * sizeof(int));
    int *key = malloc(graph->vertices * sizeof(int));
    bool *in_mst = calloc(graph->vertices, sizeof(bool));

    if (parent == NULL || key == NULL || in_mst == NULL) {
        free(parent);
        free(key);
        free(in_mst);
        return;
    }

    // Create min heap
    PrimHeap *min_heap = create_prim_heap(graph->vertices);
    if (min_heap == NULL) {
        free(parent);
        free(key);
        free(in_mst);
        return;
    }

    // Initialize min heap with all vertices
    for (int i = 0; i < graph->vertices; i++) {
        parent[i] = -1;
        key[i] = INT_MAX;
        min_heap->vertices[i] = i;
        min_heap->keys[i] = key[i];
        min_heap->parent[i] = parent[i];
        min_heap->pos[i] = i;
    }

    // Make key value of 0th vertex as 0
    min_heap->keys[0] = 0;
    key[0] = 0;
    min_heap->parent[0] = -1;

```

```

min_heap->size = graph->vertices;

// Main loop of Prim's algorithm
while (!is_prim_heap_empty(min_heap)) {
    // Extract minimum key vertex
    int u = extract_min_prim(min_heap);
    in_mst[u] = true;

    // Update key values and parent indices of adjacent vertices
    PrimNode *current = graph->adj_list[u];
    while (current != NULL) {
        int v = current->vertex;

        // If v is not in MST and weight of (u,v) is smaller than current key of v
        if (!in_mst[v] && current->weight < key[v]) {
            key[v] = current->weight;
            parent[v] = u;
            decrease_key_prim(min_heap, v, key[v]);
        }

        current = current->next;
    }
}

// Print MST edges
printf("Minimum Spanning Tree Edges:\n");
int total_weight = 0;
for (int i = 1; i < graph->vertices; i++) {
    printf("%d - %d: %d\n", parent[i], i, key[i]);
    total_weight += key[i];
}
printf("Total weight: %d\n", total_weight);

free(parent);
free(key);
free(in_mst);
free_prim_heap(min_heap);
}

int main() {
    PrimGraph *graph = create_prim_graph(5);
    if (graph == NULL) {
        printf("Failed to create graph\n");
        return 1;
    }

    // Add weighted edges to create a sample graph

```

```

    add_prim_edge(graph, 0, 1, 2);
    add_prim_edge(graph, 0, 3, 6);
    add_prim_edge(graph, 1, 2, 3);
    add_prim_edge(graph, 1, 3, 8);
    add_prim_edge(graph, 1, 4, 5);
    add_prim_edge(graph, 2, 4, 7);
    add_prim_edge(graph, 3, 4, 9);

    prim_mst(graph);

    free_prim_graph(graph);
    return 0;
}

```

Practical Examples

Social Network Analysis

```

#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <stdbool.h>

#define MAX_USERS 100
#define MAX_NAME_LENGTH 50

typedef struct User {
    int id;
    char name[MAX_NAME_LENGTH];
    struct User **friends;
    int friend_count;
    int max_friends;
} User;

typedef struct SocialNetwork {
    User *users[MAX_USERS];
    int user_count;
} SocialNetwork;

// Create new user
User* create_user(int id, const char *name) {
    User *user = malloc(sizeof(User));
    if (user == NULL) return NULL;

```

```

    user->id = id;
    strncpy(user->name, name, MAX_NAME_LENGTH - 1);
    user->name[MAX_NAME_LENGTH - 1] = '\0';
    user->friend_count = 0;
    user->max_friends = 10;
    user->friends = malloc(user->max_friends * sizeof(User*));
    if (user->friends == NULL) {
        free(user);
        return NULL;
    }

    return user;
}

// Free user
void free_user(User *user) {
    if (user == NULL) return;
    free(user->friends);
    free(user);
}

// Create social network
SocialNetwork* create_social_network(void) {
    SocialNetwork *network = malloc(sizeof(SocialNetwork));
    if (network == NULL) return NULL;

    for (int i = 0; i < MAX_USERS; i++) {
        network->users[i] = NULL;
    }

    network->user_count = 0;
    return network;
}

// Free social network
void free_social_network(SocialNetwork *network) {
    if (network == NULL) return;

    for (int i = 0; i < MAX_USERS; i++) {
        if (network->users[i] != NULL) {
            free_user(network->users[i]);
        }
    }

    free(network);
}

```

```

// Add user to network
bool add_user(SocialNetwork *network, int id, const char *name) {
    if (network == NULL || id < 0 || id >= MAX_USERS || name == NULL) {
        return false;
    }

    if (network->users[id] != NULL) {
        return false; // User already exists
    }

    User *user = create_user(id, name);
    if (user == NULL) {
        return false;
    }

    network->users[id] = user;
    network->user_count++;
    return true;
}

// Add friendship between users
bool add_friendship(SocialNetwork *network, int user1_id, int user2_id) {
    if (network == NULL || user1_id < 0 || user1_id >= MAX_USERS ||
        user2_id < 0 || user2_id >= MAX_USERS) {
        return false;
    }

    User *user1 = network->users[user1_id];
    User *user2 = network->users[user2_id];

    if (user1 == NULL || user2 == NULL) {
        return false;
    }

    // Add user2 to user1's friend list
    if (user1->friend_count >= user1->max_friends) {
        user1->max_friends *= 2;
        User **temp = realloc(user1->friends, user1->max_friends * sizeof(User*));
        if (temp == NULL) {
            return false;
        }
        user1->friends = temp;
    }
    user1->friends[user1->friend_count++] = user2;

    // Add user1 to user2's friend list
    if (user2->friend_count >= user2->max_friends) {

```



```

        user2->max_friends *= 2;
        User **temp = realloc(user2->friends, user2->max_friends * sizeof(User*));
        if (temp == NULL) {
            return false;
        }
        user2->friends = temp;
    }
    user2->friends[user2->friend_count++] = user1;

    return true;
}

// Find mutual friends
void find_mutual_friends(SocialNetwork *network, int user1_id, int user2_id) {
    if (network == NULL || user1_id < 0 || user1_id >= MAX_USERS ||
        user2_id < 0 || user2_id >= MAX_USERS) {
        return;
    }

    User *user1 = network->users[user1_id];
    User *user2 = network->users[user2_id];

    if (user1 == NULL || user2 == NULL) {
        return;
    }

    printf("Mutual friends between %s and %s:\n", user1->name, user2->name);

    bool found = false;
    for (int i = 0; i < user1->friend_count; i++) {
        User *friend1 = user1->friends[i];
        for (int j = 0; j < user2->friend_count; j++) {
            if (user2->friends[j] == friend1) {
                printf("  %s\n", friend1->name);
                found = true;
                break;
            }
        }
    }

    if (!found) {
        printf("  No mutual friends found\n");
    }
}

// Calculate degree centrality (number of connections)
void calculate_degree_centrality(SocialNetwork *network) {

```

```

    if (network == NULL) return;

    printf("Degree Centrality (Number of Friends):\n");
    printf("=====\n");

    for (int i = 0; i < MAX_USERS; i++) {
        if (network->users[i] != NULL) {
            printf("%s: %d friends\n",
                network->users[i]->name,
                network->users[i]->friend_count);
        }
    }
}

// Print social network
void print_social_network(SocialNetwork *network) {
    if (network == NULL) return;

    printf("Social Network:\n");
    printf("=====\n");

    for (int i = 0; i < MAX_USERS; i++) {
        if (network->users[i] != NULL) {
            User *user = network->users[i];
            printf("%s (ID: %d): ", user->name, user->id);

            if (user->friend_count == 0) {
                printf("No friends\n");
            } else {
                printf("Friends: ");
                for (int j = 0; j < user->friend_count; j++) {
                    printf("%s", user->friends[j]->name);
                    if (j < user->friend_count - 1) {
                        printf(", ");
                    }
                }
                printf("\n");
            }
        }
    }
}

int main() {
    SocialNetwork *network = create_social_network();
    if (network == NULL) {
        printf("Failed to create social network\n");
        return 1;
    }
}

```

```

}

// Add users
add_user(network, 0, "Alice");
add_user(network, 1, "Bob");
add_user(network, 2, "Charlie");
add_user(network, 3, "Diana");
add_user(network, 4, "Eve");

// Add friendships
add_friendship(network, 0, 1); // Alice - Bob
add_friendship(network, 0, 2); // Alice - Charlie
add_friendship(network, 1, 2); // Bob - Charlie
add_friendship(network, 1, 3); // Bob - Diana
add_friendship(network, 2, 4); // Charlie - Eve
add_friendship(network, 3, 4); // Diana - Eve

print_social_network(network);

printf("\n");
find_mutual_friends(network, 0, 3); // Alice and Diana
find_mutual_friends(network, 1, 4); // Bob and Eve

printf("\n");
calculate_degree centrality(network);

free_social_network(network);
return 0;
}

```

Summary

Graph algorithms are essential for solving complex problems in computer science and real-world applications:

1. Graph Representations:

- Adjacency Matrix: Good for dense graphs, $O(1)$ edge lookup
- Adjacency List: Good for sparse graphs, space-efficient

2. Graph Traversal:

- BFS: Level-by-level exploration, useful for shortest path in unweighted graphs
- DFS: Deep exploration, useful for connectivity and cycle detection

3. Shortest Path Algorithms:

- Dijkstra's Algorithm: Finds shortest paths in weighted graphs with non-negative edges

- Time complexity: $O((V + E) \log V)$ with binary heap

4. Minimum Spanning Tree:

- Prim's Algorithm: Greedy approach to find MST
- Time complexity: $O((V + E) \log V)$ with binary heap

Key considerations when working with graph algorithms: - **Graph Density:** Choose appropriate representation (matrix vs. list) - **Edge Weights:** Consider whether negative weights are allowed - **Directed vs. Undirected:** Algorithm implementation varies - **Performance Requirements:** Balance time and space complexity

Graph algorithms have numerous applications: - Social network analysis - Network routing protocols - Transportation and logistics - Computer network design - Circuit design - Game development (pathfinding) - Recommendation systems - Dependency resolution

Understanding graph algorithms and their implementations is crucial for solving complex computational problems efficiently.

Sorting and Searching

Introduction

Sorting and searching are fundamental operations in computer science that are used extensively in software development. Sorting arranges data in a specific order, making it easier to search, analyze, and process. Searching finds specific elements within a dataset, which is essential for data retrieval and manipulation.

In this chapter, we'll explore various sorting algorithms including comparison-based and non-comparison-based approaches, as well as different searching techniques. We'll implement these algorithms in C and analyze their time and space complexity.

Comparison-Based Sorting Algorithms

Bubble Sort

Bubble sort repeatedly steps through the list, compares adjacent elements, and swaps them if they're in the wrong order.

```
#include <stdio.h>
#include <stdlib.h>
#include <stdbool.h>
#include <time.h>

// Bubble sort implementation
void bubble_sort(int arr[], int n) {
    for (int i = 0; i < n - 1; i++) {
        bool swapped = false;

        // Last i elements are already in place
        for (int j = 0; j < n - i - 1; j++) {
            if (arr[j] > arr[j + 1]) {
                // Swap elements
                int temp = arr[j];
                arr[j] = arr[j + 1];
                arr[j + 1] = temp;
                swapped = true;
            }
        }
    }
}
```

```

        // If no swapping occurred, array is sorted
        if (!swapped) {
            break;
        }
    }
}

// Print array
void print_array(int arr[], int n) {
    for (int i = 0; i < n; i++) {
        printf("%d ", arr[i]);
    }
    printf("\n");
}

int main() {
    int arr[] = {64, 34, 25, 12, 22, 11, 90};
    int n = sizeof(arr) / sizeof(arr[0]);

    printf("Original array: ");
    print_array(arr, n);

    bubble_sort(arr, n);

    printf("Sorted array: ");
    print_array(arr, n);

    return 0;
}

```

Selection Sort

Selection sort finds the minimum element from the unsorted part and places it at the beginning.

```

#include <stdio.h>

// Selection sort implementation
void selection_sort(int arr[], int n) {
    for (int i = 0; i < n - 1; i++) {
        // Find the minimum element in the remaining unsorted array
        int min_idx = i;
        for (int j = i + 1; j < n; j++) {
            if (arr[j] < arr[min_idx]) {
                min_idx = j;
            }
        }
    }
}

```

```

    }

    // Swap the found minimum element with the first element
    if (min_idx != i) {
        int temp = arr[min_idx];
        arr[min_idx] = arr[i];
        arr[i] = temp;
    }
}

// Print array
void print_array(int arr[], int n) {
    for (int i = 0; i < n; i++) {
        printf("%d ", arr[i]);
    }
    printf("\n");
}

int main() {
    int arr[] = {64, 25, 12, 22, 11, 90};
    int n = sizeof(arr) / sizeof(arr[0]);

    printf("Original array: ");
    print_array(arr, n);

    selection_sort(arr, n);

    printf("Sorted array: ");
    print_array(arr, n);

    return 0;
}

```

Insertion Sort

Insertion sort builds the final sorted array one item at a time by inserting each element into its correct position.

```

#include <stdio.h>

// Insertion sort implementation
void insertion_sort(int arr[], int n) {
    for (int i = 1; i < n; i++) {
        int key = arr[i];
        int j = i - 1;

```

```

        // Move elements greater than key one position ahead
        while (j >= 0 && arr[j] > key) {
            arr[j + 1] = arr[j];
            j--;
        }

        arr[j + 1] = key;
    }
}

// Print array
void print_array(int arr[], int n) {
    for (int i = 0; i < n; i++) {
        printf("%d ", arr[i]);
    }
    printf("\n");
}

int main() {
    int arr[] = {12, 11, 13, 5, 6};
    int n = sizeof(arr) / sizeof(arr[0]);

    printf("Original array: ");
    print_array(arr, n);

    insertion_sort(arr, n);

    printf("Sorted array: ");
    print_array(arr, n);

    return 0;
}

```

Merge Sort

Merge sort is a divide-and-conquer algorithm that divides the array into halves, sorts them, and then merges them.

```

#include <stdio.h>
#include <stdlib.h>

// Merge two sorted subarrays
void merge(int arr[], int left, int mid, int right) {
    int n1 = mid - left + 1;
    int n2 = right - mid;

```



```

// Create temporary arrays
int *left_arr = malloc(n1 * sizeof(int));
int *right_arr = malloc(n2 * sizeof(int));

if (left_arr == NULL || right_arr == NULL) {
    free(left_arr);
    free(right_arr);
    return;
}

// Copy data to temporary arrays
for (int i = 0; i < n1; i++) {
    left_arr[i] = arr[left + i];
}
for (int i = 0; i < n2; i++) {
    right_arr[i] = arr[mid + 1 + i];
}

// Merge the temporary arrays back into arr[left..right]
int i = 0, j = 0, k = left;

while (i < n1 && j < n2) {
    if (left_arr[i] <= right_arr[j]) {
        arr[k] = left_arr[i];
        i++;
    } else {
        arr[k] = right_arr[j];
        j++;
    }
    k++;
}

// Copy remaining elements of left_arr[], if any
while (i < n1) {
    arr[k] = left_arr[i];
    i++;
    k++;
}

// Copy remaining elements of right_arr[], if any
while (j < n2) {
    arr[k] = right_arr[j];
    j++;
    k++;
}

```

```

        free(left_arr);
        free(right_arr);
    }

    // Merge sort implementation
    void merge_sort(int arr[], int left, int right) {
        if (left < right) {
            int mid = left + (right - left) / 2;

            // Sort first and second halves
            merge_sort(arr, left, mid);
            merge_sort(arr, mid + 1, right);

            // Merge the sorted halves
            merge(arr, left, mid, right);
        }
    }

    // Print array
    void print_array(int arr[], int n) {
        for (int i = 0; i < n; i++) {
            printf("%d ", arr[i]);
        }
        printf("\n");
    }

    int main() {
        int arr[] = {12, 11, 13, 5, 6, 7};
        int n = sizeof(arr) / sizeof(arr[0]);

        printf("Original array: ");
        print_array(arr, n);

        merge_sort(arr, 0, n - 1);

        printf("Sorted array: ");
        print_array(arr, n);

        return 0;
    }

```

Quick Sort

Quick sort is a divide-and-conquer algorithm that picks an element as a pivot and partitions the array around it.

```

#include <stdio.h>

// Partition function for quick sort
int partition(int arr[], int low, int high) {
    int pivot = arr[high]; // Choose the rightmost element as pivot
    int i = low - 1;        // Index of smaller element

    for (int j = low; j < high; j++) {
        // If current element is smaller than or equal to pivot
        if (arr[j] <= pivot) {
            i++; // Increment index of smaller element
            // Swap arr[i] and arr[j]
            int temp = arr[i];
            arr[i] = arr[j];
            arr[j] = temp;
        }
    }

    // Swap arr[i+1] and arr[high] (pivot)
    int temp = arr[i + 1];
    arr[i + 1] = arr[high];
    arr[high] = temp;

    return i + 1;
}

// Quick sort implementation
void quick_sort(int arr[], int low, int high) {
    if (low < high) {
        // Partition the array and get pivot index
        int pi = partition(arr, low, high);

        // Recursively sort elements before and after partition
        quick_sort(arr, low, pi - 1);
        quick_sort(arr, pi + 1, high);
    }
}

// Print array
void print_array(int arr[], int n) {
    for (int i = 0; i < n; i++) {
        printf("%d ", arr[i]);
    }
    printf("\n");
}

int main() {

```

```

int arr[] = {10, 7, 8, 9, 1, 5};
int n = sizeof(arr) / sizeof(arr[0]);

printf("Original array: ");
print_array(arr, n);

quick_sort(arr, 0, n - 1);

printf("Sorted array: ");
print_array(arr, n);

return 0;
}

```

Heap Sort

Heap sort uses a binary heap data structure to sort elements efficiently.

```

#include <stdio.h>

// Heapify a subtree rooted at index i
void heapify(int arr[], int n, int i) {
    int largest = i;    // Initialize largest as root
    int left = 2 * i + 1;    // Left child
    int right = 2 * i + 2;   // Right child

    // If left child is larger than root
    if (left < n && arr[left] > arr[largest]) {
        largest = left;
    }

    // If right child is larger than largest so far
    if (right < n && arr[right] > arr[largest]) {
        largest = right;
    }

    // If largest is not root
    if (largest != i) {
        int temp = arr[i];
        arr[i] = arr[largest];
        arr[largest] = temp;

        // Recursively heapify the affected sub-tree
        heapify(arr, n, largest);
    }
}

```

```

// Heap sort implementation
void heap_sort(int arr[], int n) {
    // Build heap (rearrange array)
    for (int i = n / 2 - 1; i >= 0; i--) {
        heapify(arr, n, i);
    }

    // Extract elements from heap one by one
    for (int i = n - 1; i > 0; i--) {
        // Move current root to end
        int temp = arr[0];
        arr[0] = arr[i];
        arr[i] = temp;

        // Call heapify on the reduced heap
        heapify(arr, i, 0);
    }
}

// Print array
void print_array(int arr[], int n) {
    for (int i = 0; i < n; i++) {
        printf("%d ", arr[i]);
    }
    printf("\n");
}

int main() {
    int arr[] = {12, 11, 13, 5, 6, 7};
    int n = sizeof(arr) / sizeof(arr[0]);

    printf("Original array: ");
    print_array(arr, n);

    heap_sort(arr, n);

    printf("Sorted array: ");
    print_array(arr, n);

    return 0;
}

```

Non-Comparison Sorting Algorithms

Counting Sort

Counting sort counts the occurrences of each element and uses this information to place elements in the correct position.

```
#include <stdio.h>
#include <stdlib.h>
#include <string.h>

// Counting sort implementation
void counting_sort(int arr[], int n, int max_val) {
    // Create count array to store count of each element
    int *count = calloc(max_val + 1, sizeof(int));
    int *output = malloc(n * sizeof(int));

    if (count == NULL || output == NULL) {
        free(count);
        free(output);
        return;
    }

    // Store count of each element
    for (int i = 0; i < n; i++) {
        count[arr[i]]++;
    }

    // Change count[i] so that count[i] now contains actual
    // position of this element in output array
    for (int i = 1; i <= max_val; i++) {
        count[i] += count[i - 1];
    }

    // Build the output array
    for (int i = n - 1; i >= 0; i--) {
        output[count[arr[i]] - 1] = arr[i];
        count[arr[i]]--;
    }

    // Copy the output array to arr[], so that arr[] now
    // contains sorted elements
    for (int i = 0; i < n; i++) {
        arr[i] = output[i];
    }

    free(count);
}
```

```

        free(output);
    }

    // Print array
    void print_array(int arr[], int n) {
        for (int i = 0; i < n; i++) {
            printf("%d ", arr[i]);
        }
        printf("\n");
    }

    int main() {
        int arr[] = {4, 2, 2, 8, 3, 3, 1};
        int n = sizeof(arr) / sizeof(arr[0]);

        printf("Original array: ");
        print_array(arr, n);

        // Find maximum value in array
        int max_val = arr[0];
        for (int i = 1; i < n; i++) {
            if (arr[i] > max_val) {
                max_val = arr[i];
            }
        }

        counting_sort(arr, n, max_val);

        printf("Sorted array: ");
        print_array(arr, n);

        return 0;
    }

```

Radix Sort

Radix sort sorts numbers by processing individual digits, starting from the least significant digit to the most significant.

```

#include <stdio.h>
#include <stdlib.h>

// Get maximum value in array
int get_max(int arr[], int n) {
    int max = arr[0];
    for (int i = 1; i < n; i++) {

```

```

        if (arr[i] > max) {
            max = arr[i];
        }
    }
    return max;
}

// Counting sort for a particular digit (exp)
void counting_sort_for_radix(int arr[], int n, int exp) {
    int *output = malloc(n * sizeof(int));
    int count[10] = {0};

    if (output == NULL) return;

    // Store count of occurrences in count[]
    for (int i = 0; i < n; i++) {
        count[(arr[i] / exp) % 10]++;
    }

    // Change count[i] so that count[i] now contains actual
    // position of this digit in output[]
    for (int i = 1; i < 10; i++) {
        count[i] += count[i - 1];
    }

    // Build the output array
    for (int i = n - 1; i >= 0; i--) {
        output[count[(arr[i] / exp) % 10] - 1] = arr[i];
        count[(arr[i] / exp) % 10]--;
    }

    // Copy the output array to arr[], so that arr[] now
    // contains sorted numbers according to current digit
    for (int i = 0; i < n; i++) {
        arr[i] = output[i];
    }

    free(output);
}

// Radix sort implementation
void radix_sort(int arr[], int n) {
    // Find the maximum number to know number of digits
    int max = get_max(arr, n);

    // Do counting sort for every digit
    for (int exp = 1; max / exp > 0; exp *= 10) {

```



```

        counting_sort_for_radix(arr, n, exp);
    }
}

// Print array
void print_array(int arr[], int n) {
    for (int i = 0; i < n; i++) {
        printf("%d ", arr[i]);
    }
    printf("\n");
}

int main() {
    int arr[] = {170, 45, 75, 90, 2, 802, 24, 66};
    int n = sizeof(arr) / sizeof(arr[0]);

    printf("Original array: ");
    print_array(arr, n);

    radix_sort(arr, n);

    printf("Sorted array: ");
    print_array(arr, n);

    return 0;
}

```

Bucket Sort

Bucket sort distributes elements into buckets and then sorts each bucket individually.

```

#include <stdio.h>
#include <stdlib.h>
#include <string.h>

#define NUM_BUCKETS 10

typedef struct Node {
    float data;
    struct Node *next;
} Node;

// Create new node
Node* create_node(float data) {
    Node *node = malloc(sizeof(Node));
    if (node == NULL) return NULL;

```

```

    node->data = data;
    node->next = NULL;
    return node;
}

// Insert node in sorted order
void insert_sorted(Node **head, float data) {
    Node *new_node = create_node(data);
    if (new_node == NULL) return;

    // If list is empty or new node should be first
    if (*head == NULL || (*head)->data >= data) {
        new_node->next = *head;
        *head = new_node;
        return;
    }

    // Find the node before the point of insertion
    Node *current = *head;
    while (current->next != NULL && current->next->data < data) {
        current = current->next;
    }

    new_node->next = current->next;
    current->next = new_node;
}

// Bucket sort implementation
void bucket_sort(float arr[], int n) {
    // Create buckets
    Node *buckets[NUM_BUCKETS];
    for (int i = 0; i < NUM_BUCKETS; i++) {
        buckets[i] = NULL;
    }

    // Put array elements in different buckets
    for (int i = 0; i < n; i++) {
        int bucket_index = (int)(arr[i] * NUM_BUCKETS);
        // Handle edge case where arr[i] is 1.0
        if (bucket_index >= NUM_BUCKETS) {
            bucket_index = NUM_BUCKETS - 1;
        }
        insert_sorted(&buckets[bucket_index], arr[i]);
    }

    // Concatenate all buckets into arr[]

```

```

    int index = 0;
    for (int i = 0; i < NUM_BUCKETS; i++) {
        Node *current = buckets[i];
        while (current != NULL) {
            arr[index++] = current->data;
            Node *temp = current;
            current = current->next;
            free(temp);
        }
    }
}

// Print array
void print_array(float arr[], int n) {
    for (int i = 0; i < n; i++) {
        printf("%.2f ", arr[i]);
    }
    printf("\n");
}

int main() {
    float arr[] = {0.897, 0.565, 0.656, 0.1234, 0.665, 0.3434};
    int n = sizeof(arr) / sizeof(arr[0]);

    printf("Original array: ");
    print_array(arr, n);

    bucket_sort(arr, n);

    printf("Sorted array: ");
    print_array(arr, n);

    return 0;
}

```

Searching Algorithms

Linear Search

Linear search sequentially checks each element until the target is found or the end of the array is reached.

```

#include <stdio.h>

// Linear search implementation

```

```

int linear_search(int arr[], int n, int target) {
    for (int i = 0; i < n; i++) {
        if (arr[i] == target) {
            return i; // Return index of found element
        }
    }
    return -1; // Element not found
}

int main() {
    int arr[] = {10, 20, 80, 30, 60, 50, 110, 100, 130, 170};
    int n = sizeof(arr) / sizeof(arr[0]);
    int target = 110;

    int result = linear_search(arr, n, target);

    if (result != -1) {
        printf("Element %d found at index %d\n", target, result);
    } else {
        printf("Element %d not found in array\n", target);
    }

    return 0;
}

```

Binary Search

Binary search works on sorted arrays by repeatedly dividing the search interval in half.

```

#include <stdio.h>

// Iterative binary search
int binary_search_iterative(int arr[], int n, int target) {
    int left = 0;
    int right = n - 1;

    while (left <= right) {
        int mid = left + (right - left) / 2;

        if (arr[mid] == target) {
            return mid; // Element found
        } else if (arr[mid] < target) {
            left = mid + 1; // Search right half
        } else {
            right = mid - 1; // Search left half
        }
    }
}

```

```

    }

    return -1; // Element not found
}

// Recursive binary search
int binary_search_recursive(int arr[], int left, int right, int target) {
    if (left > right) {
        return -1; // Element not found
    }

    int mid = left + (right - left) / 2;

    if (arr[mid] == target) {
        return mid; // Element found
    } else if (arr[mid] < target) {
        return binary_search_recursive(arr, mid + 1, right, target);
    } else {
        return binary_search_recursive(arr, left, mid - 1, target);
    }
}

int main() {
    int arr[] = {2, 3, 4, 10, 40, 50, 80};
    int n = sizeof(arr) / sizeof(arr[0]);
    int target = 10;

    printf("Array: ");
    for (int i = 0; i < n; i++) {
        printf("%d ", arr[i]);
    }
    printf("\n");

    // Iterative binary search
    int result1 = binary_search_iterative(arr, n, target);
    if (result1 != -1) {
        printf("Element %d found at index %d (iterative)\n", target, result1);
    } else {
        printf("Element %d not found (iterative)\n", target);
    }

    // Recursive binary search
    int result2 = binary_search_recursive(arr, 0, n - 1, target);
    if (result2 != -1) {
        printf("Element %d found at index %d (recursive)\n", target, result2);
    } else {

```

```

        printf("Element %d not found (recursive)\n", target);
    }

    return 0;
}

```

Interpolation Search

Interpolation search is an improvement over binary search for uniformly distributed sorted arrays.

```

#include <stdio.h>

// Interpolation search implementation
int interpolation_search(int arr[], int n, int target) {
    int low = 0;
    int high = n - 1;

    while (low <= high && target >= arr[low] && target <= arr[high]) {
        // If array has only one element
        if (low == high) {
            if (arr[low] == target) return low;
            return -1;
        }

        // Probing the position with keeping uniform distribution in mind
        int pos = low + (((double)(high - low) / (arr[high] - arr[low])) * (target - arr[low]));

        // Condition of target found
        if (arr[pos] == target) {
            return pos;
        }

        // If target is larger, target is in upper part
        if (arr[pos] < target) {
            low = pos + 1;
        }
        // If target is smaller, target is in lower part
        else {
            high = pos - 1;
        }
    }

    return -1; // Element not found
}

int main() {

```

```

int arr[] = {10, 12, 13, 16, 18, 19, 20, 21, 22, 23, 24, 33, 35, 42, 47};
int n = sizeof(arr) / sizeof(arr[0]);
int target = 18;

printf("Array: ");
for (int i = 0; i < n; i++) {
    printf("%d ", arr[i]);
}
printf("\n");

int result = interpolation_search(arr, n, target);

if (result != -1) {
    printf("Element %d found at index %d\n", target, result);
} else {
    printf("Element %d not found\n", target);
}

return 0;
}

```

Advanced Searching Techniques

Exponential Search

Exponential search finds the range where the element might be present and then performs binary search in that range.

```

#include <stdio.h>

// Binary search helper function
int binary_search(int arr[], int left, int right, int target) {
    while (left <= right) {
        int mid = left + (right - left) / 2;

        if (arr[mid] == target) {
            return mid;
        } else if (arr[mid] < target) {
            left = mid + 1;
        } else {
            right = mid - 1;
        }
    }
    return -1;
}

```

```

// Exponential search implementation
int exponential_search(int arr[], int n, int target) {
    // If first element is target
    if (arr[0] == target) {
        return 0;
    }

    // Find range for binary search by repeated doubling
    int i = 1;
    while (i < n && arr[i] <= target) {
        i *= 2;
    }

    // Call binary search for the found range
    return binary_search(arr, i / 2, (i < n) ? i : n - 1, target);
}

int main() {
    int arr[] = {2, 3, 4, 10, 40, 50, 80, 100, 120, 150, 200};
    int n = sizeof(arr) / sizeof(arr[0]);
    int target = 100;

    printf("Array: ");
    for (int i = 0; i < n; i++) {
        printf("%d ", arr[i]);
    }
    printf("\n");

    int result = exponential_search(arr, n, target);

    if (result != -1) {
        printf("Element %d found at index %d\n", target, result);
    } else {
        printf("Element %d not found\n", target);
    }

    return 0;
}

```

Jump Search

Jump search creates blocks and jumps through them to find the target block, then performs linear search within that block.


```

#include <stdio.h>
#include <math.h>

// Jump search implementation
int jump_search(int arr[], int n, int target) {
    int step = sqrt(n);
    int prev = 0;

    // Finding the block where element is present (if it is present)
    while (arr[(int)fmin(step, n) - 1] < target) {
        prev = step;
        step += sqrt(n);
        if (prev >= n) {
            return -1;
        }
    }

    // Doing a linear search for target in block beginning with prev
    while (arr[prev] < target) {
        prev++;

        // If we reached next block or end of array, element is not present
        if (prev == fmin(step, n)) {
            return -1;
        }
    }

    // If element is found
    if (arr[prev] == target) {
        return prev;
    }

    return -1;
}

int main() {
    int arr[] = {0, 1, 1, 2, 3, 5, 8, 13, 21, 34, 55, 89, 144, 233, 377, 610};
    int n = sizeof(arr) / sizeof(arr[0]);
    int target = 55;

    printf("Array: ");
    for (int i = 0; i < n; i++) {
        printf("%d ", arr[i]);
    }
    printf("\n");

    int result = jump_search(arr, n, target);

```

```

    if (result != -1) {
        printf("Element %d found at index %d\n", target, result);
    } else {
        printf("Element %d not found\n", target);
    }

    return 0;
}

```

Selection Algorithms

Quickselect (Finding kth Smallest Element)

Quickselect finds the kth smallest element in an unordered list using a partitioning approach similar to quicksort.

```

#include <stdio.h>
#include <stdlib.h>

// Partition function (same as in quicksort)
int partition(int arr[], int low, int high) {
    int pivot = arr[high];
    int i = low - 1;

    for (int j = low; j < high; j++) {
        if (arr[j] <= pivot) {
            i++;
            int temp = arr[i];
            arr[i] = arr[j];
            arr[j] = temp;
        }
    }

    int temp = arr[i + 1];
    arr[i + 1] = arr[high];
    arr[high] = temp;

    return i + 1;
}

// Quickselect algorithm to find kth smallest element
int quickselect(int arr[], int low, int high, int k) {
    if (low <= high) {
        // Partition the array and get pivot index

```

```

    int pi = partition(arr, low, high);

    // If pivot is the kth element
    if (pi == k) {
        return arr[pi];
    }
    // If kth element is in left subarray
    else if (pi > k) {
        return quickselect(arr, low, pi - 1, k);
    }
    // If kth element is in right subarray
    else {
        return quickselect(arr, pi + 1, high, k);
    }
}

return -1; // Error case
}

// Find kth smallest element
int find_kth_smallest(int arr[], int n, int k) {
    if (k < 0 || k >= n) {
        return -1; // Invalid k
    }

    // Create a copy of array to avoid modifying original
    int *temp = malloc(n * sizeof(int));
    if (temp == NULL) return -1;

    for (int i = 0; i < n; i++) {
        temp[i] = arr[i];
    }

    int result = quickselect(temp, 0, n - 1, k);
    free(temp);

    return result;
}

int main() {
    int arr[] = {12, 3, 5, 7, 4, 19, 26};
    int n = sizeof(arr) / sizeof(arr[0]);
    int k = 3; // Find 4th smallest element (0-indexed)

    printf("Array: ");
    for (int i = 0; i < n; i++) {
        printf("%d ", arr[i]);
    }
}

```

```

    }
    printf("\n");

    int kth_smallest = find_kth_smallest(arr, n, k);

    if (kth_smallest != -1) {
        printf("The %dth smallest element is: %d\n", k + 1, kth_smallest);
    } else {
        printf("Invalid k value\n");
    }

    return 0;
}

```

Performance Comparison and Analysis

Sorting Algorithm Performance Test

```

#include <stdio.h>
#include <stdlib.h>
#include <time.h>
#include <string.h>

#define ARRAY_SIZE 10000

// Function prototypes
void bubble_sort(int arr[], int n);
void selection_sort(int arr[], int n);
void insertion_sort(int arr[], int n);
void merge_sort(int arr[], int left, int right);
void quick_sort(int arr[], int low, int high);
void heap_sort(int arr[], int n);

// Generate random array
void generate_random_array(int arr[], int n) {
    srand(time(NULL));
    for (int i = 0; i < n; i++) {
        arr[i] = rand() % 1000;
    }
}

// Copy array
void copy_array(int src[], int dest[], int n) {
    for (int i = 0; i < n; i++) {

```

```

        dest[i] = src[i];
    }
}

// Test sorting algorithm performance
double test_sorting_algorithm(void (*sort_func)(int[], int), int arr[], int n, const char *name)
{
    int *test_arr = malloc(n * sizeof(int));
    if (test_arr == NULL) return -1;

    copy_array(arr, test_arr, n);

    clock_t start = clock();
    sort_func(test_arr, n);
    clock_t end = clock();

    double time_taken = ((double)(end - start)) / CLOCKS_PER_SEC;

    printf("%-20s: %.6f seconds\n", name, time_taken);

    free(test_arr);
    return time_taken;
}

// Test sorting algorithm with range parameters
double test_sorting_algorithm_range(void (*sort_func)(int[], int, int), int arr[], int n, const char *name)
{
    int *test_arr = malloc(n * sizeof(int));
    if (test_arr == NULL) return -1;

    copy_array(arr, test_arr, n);

    clock_t start = clock();
    sort_func(test_arr, 0, n - 1);
    clock_t end = clock();

    double time_taken = ((double)(end - start)) / CLOCKS_PER_SEC;

    printf("%-20s: %.6f seconds\n", name, time_taken);

    free(test_arr);
    return time_taken;
}

// Bubble sort implementation
void bubble_sort(int arr[], int n) {
    for (int i = 0; i < n - 1; i++) {
        int swapped = 0;

```

```

        for (int j = 0; j < n - i - 1; j++) {
            if (arr[j] > arr[j + 1]) {
                int temp = arr[j];
                arr[j] = arr[j + 1];
                arr[j + 1] = temp;
                swapped = 1;
            }
        }
        if (!swapped) break;
    }
}

// Selection sort implementation
void selection_sort(int arr[], int n) {
    for (int i = 0; i < n - 1; i++) {
        int min_idx = i;
        for (int j = i + 1; j < n; j++) {
            if (arr[j] < arr[min_idx]) {
                min_idx = j;
            }
        }
        if (min_idx != i) {
            int temp = arr[min_idx];
            arr[min_idx] = arr[i];
            arr[i] = temp;
        }
    }
}

// Insertion sort implementation
void insertion_sort(int arr[], int n) {
    for (int i = 1; i < n; i++) {
        int key = arr[i];
        int j = i - 1;
        while (j >= 0 && arr[j] > key) {
            arr[j + 1] = arr[j];
            j--;
        }
        arr[j + 1] = key;
    }
}

// Merge function for merge sort
void merge(int arr[], int left, int mid, int right) {
    int n1 = mid - left + 1;
    int n2 = right - mid;

```

```

    int *left_arr = malloc(n1 * sizeof(int));
    int *right_arr = malloc(n2 * sizeof(int));

    for (int i = 0; i < n1; i++) {
        left_arr[i] = arr[left + i];
    }
    for (int i = 0; i < n2; i++) {
        right_arr[i] = arr[mid + 1 + i];
    }

    int i = 0, j = 0, k = left;
    while (i < n1 && j < n2) {
        if (left_arr[i] <= right_arr[j]) {
            arr[k] = left_arr[i];
            i++;
        } else {
            arr[k] = right_arr[j];
            j++;
        }
        k++;
    }

    while (i < n1) {
        arr[k] = left_arr[i];
        i++;
        k++;
    }

    while (j < n2) {
        arr[k] = right_arr[j];
        j++;
        k++;
    }

    free(left_arr);
    free(right_arr);
}

// Merge sort implementation
void merge_sort(int arr[], int left, int right) {
    if (left < right) {
        int mid = left + (right - left) / 2;
        merge_sort(arr, left, mid);
        merge_sort(arr, mid + 1, right);
        merge(arr, left, mid, right);
    }
}

```

```

// Partition function for quick sort
int partition(int arr[], int low, int high) {
    int pivot = arr[high];
    int i = low - 1;

    for (int j = low; j < high; j++) {
        if (arr[j] <= pivot) {
            i++;
            int temp = arr[i];
            arr[i] = arr[j];
            arr[j] = temp;
        }
    }

    int temp = arr[i + 1];
    arr[i + 1] = arr[high];
    arr[high] = temp;

    return i + 1;
}

// Quick sort implementation
void quick_sort(int arr[], int low, int high) {
    if (low < high) {
        int pi = partition(arr, low, high);
        quick_sort(arr, low, pi - 1);
        quick_sort(arr, pi + 1, high);
    }
}

// Heapify function for heap sort
void heapify(int arr[], int n, int i) {
    int largest = i;
    int left = 2 * i + 1;
    int right = 2 * i + 2;

    if (left < n && arr[left] > arr[largest]) {
        largest = left;
    }

    if (right < n && arr[right] > arr[largest]) {
        largest = right;
    }

    if (largest != i) {
        int temp = arr[i];

```



```

        arr[i] = arr[largest];
        arr[largest] = temp;
        heapify(arr, n, largest);
    }
}

// Heap sort implementation
void heap_sort(int arr[], int n) {
    for (int i = n / 2 - 1; i >= 0; i--) {
        heapify(arr, n, i);
    }

    for (int i = n - 1; i > 0; i--) {
        int temp = arr[0];
        arr[0] = arr[i];
        arr[i] = temp;
        heapify(arr, i, 0);
    }
}

int main() {
    int *original_arr = malloc(ARRAY_SIZE * sizeof(int));
    if (original_arr == NULL) {
        printf("Memory allocation failed\n");
        return 1;
    }

    printf("Generating random array of size %d...\n", ARRAY_SIZE);
    generate_random_array(original_arr, ARRAY_SIZE);

    printf("\nSorting Algorithm Performance Comparison:\n");
    printf("=====\n");

    // Test various sorting algorithms
    test_sorting_algorithm(bubble_sort, original_arr, ARRAY_SIZE, "Bubble Sort");
    test_sorting_algorithm(selection_sort, original_arr, ARRAY_SIZE, "Selection Sort");
    test_sorting_algorithm(insertion_sort, original_arr, ARRAY_SIZE, "Insertion Sort");
    test_sorting_algorithm_range(merge_sort, original_arr, ARRAY_SIZE, "Merge Sort");
    test_sorting_algorithm_range(quick_sort, original_arr, ARRAY_SIZE, "Quick Sort");
    test_sorting_algorithm(heap_sort, original_arr, ARRAY_SIZE, "Heap Sort");

    free(original_arr);
    return 0;
}

```

Practical Examples

Phone Book Implementation

```
#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <stdbool.h>

#define MAX_CONTACTS 1000
#define MAX_NAME_LENGTH 50
#define MAX_PHONE_LENGTH 20

typedef struct {
    char name[MAX_NAME_LENGTH];
    char phone[MAX_PHONE_LENGTH];
} Contact;

typedef struct {
    Contact contacts[MAX_CONTACTS];
    int count;
} PhoneBook;

// Create phone book
PhoneBook* create_phone_book(void) {
    PhoneBook *book = malloc(sizeof(PhoneBook));
    if (book == NULL) return NULL;

    book->count = 0;
    return book;
}

// Free phone book
void free_phone_book(PhoneBook *book) {
    free(book);
}

// Add contact to phone book
bool add_contact(PhoneBook *book, const char *name, const char *phone) {
    if (book == NULL || name == NULL || phone == NULL || book->count >= MAX_CONTACTS) {
        return false;
    }

    strncpy(book->contacts[book->count].name, name, MAX_NAME_LENGTH - 1);
    book->contacts[book->count].name[MAX_NAME_LENGTH - 1] = '\0';
```

```

        strncpy(book->contacts[book->count].phone, phone, MAX_PHONE_LENGTH - 1);
        book->contacts[book->count].phone[MAX_PHONE_LENGTH - 1] = '\0';

        book->count++;
        return true;
    }

    // Linear search for contact by name
    int search_contact_linear(PhoneBook *book, const char *name) {
        if (book == NULL || name == NULL) return -1;

        for (int i = 0; i < book->count; i++) {
            if (strcmp(book->contacts[i].name, name) == 0) {
                return i;
            }
        }

        return -1;
    }

    // Comparison function for sorting contacts by name
    int compare_contacts(const void *a, const void *b) {
        const Contact *contact_a = (const Contact *)a;
        const Contact *contact_b = (const Contact *)b;
        return strcmp(contact_a->name, contact_b->name);
    }

    // Sort phone book by name
    void sort_phone_book(PhoneBook *book) {
        if (book == NULL) return;

        qsort(book->contacts, book->count, sizeof(Contact), compare_contacts);
    }

    // Binary search for contact by name (requires sorted phone book)
    int search_contact_binary(PhoneBook *book, const char *name) {
        if (book == NULL || name == NULL) return -1;

        int left = 0;
        int right = book->count - 1;

        while (left <= right) {
            int mid = left + (right - left) / 2;
            int cmp = strcmp(book->contacts[mid].name, name);

            if (cmp == 0) {
                return mid;
            }
        }
    }

```

```

        } else if (cmp < 0) {
            left = mid + 1;
        } else {
            right = mid - 1;
        }
    }

    return -1;
}

// Print all contacts
void print_phone_book(PhoneBook *book) {
    if (book == NULL) return;

    printf("Phone Book (%d contacts):\n", book->count);
    printf("=====\n");

    for (int i = 0; i < book->count; i++) {
        printf("%-20s: %s\n", book->contacts[i].name, book->contacts[i].phone);
    }
    printf("\n");
}

// Print specific contact
void print_contact(PhoneBook *book, int index) {
    if (book == NULL || index < 0 || index >= book->count) return;

    printf("Name: %s\n", book->contacts[index].name);
    printf("Phone: %s\n", book->contacts[index].phone);
}

int main() {
    PhoneBook *book = create_phone_book();
    if (book == NULL) {
        printf("Failed to create phone book\n");
        return 1;
    }

    // Add contacts
    add_contact(book, "Alice Johnson", "555-1234");
    add_contact(book, "Bob Smith", "555-5678");
    add_contact(book, "Charlie Brown", "555-9012");
    add_contact(book, "Diana Prince", "555-3456");
    add_contact(book, "Eve Wilson", "555-7890");
    add_contact(book, "Frank Miller", "555-2345");
    add_contact(book, "Grace Lee", "555-6789");
}

```

```

printf("Unsorted phone book:\n");
print_phone_book(book);

// Linear search
printf("Linear search for 'Charlie Brown':\n");
int index = search_contact_linear(book, "Charlie Brown");
if (index != -1) {
    print_contact(book, index);
} else {
    printf("Contact not found\n");
}

// Sort phone book
printf("\nSorting phone book...\n");
sort_phone_book(book);

printf("\nSorted phone book:\n");
print_phone_book(book);

// Binary search
printf("Binary search for 'Eve Wilson':\n");
index = search_contact_binary(book, "Eve Wilson");
if (index != -1) {
    print_contact(book, index);
} else {
    printf("Contact not found\n");
}

free_phone_book(book);
return 0;
}

```

Summary

Sorting and searching algorithms are fundamental to computer science and software development:

1. Comparison-Based Sorting:

- Bubble Sort: $O(n^2)$ time, simple but inefficient
- Selection Sort: $O(n^2)$ time, minimal memory writes
- Insertion Sort: $O(n^2)$ time, efficient for small datasets
- Merge Sort: $O(n \log n)$ time, stable and consistent
- Quick Sort: $O(n \log n)$ average time, in-place sorting
- Heap Sort: $O(n \log n)$ time, in-place and consistent

2. Non-Comparison Sorting:

- Counting Sort: $O(n + k)$ time, for small range integers

- Radix Sort: $O(d \times (n + k))$ time, for fixed-width integers
- Bucket Sort: $O(n + k)$ average time, for uniformly distributed data

3. Searching Algorithms:

- Linear Search: $O(n)$ time, works on unsorted data
- Binary Search: $O(\log n)$ time, requires sorted data
- Interpolation Search: $O(\log \log n)$ average time, for uniformly distributed data
- Exponential Search: $O(\log n)$ time, for unbounded searches
- Jump Search: $O(\sqrt{n})$ time, alternative to binary search

4. Selection Algorithms:

- Quickselect: $O(n)$ average time, finds kth smallest element

Key considerations when choosing algorithms: - **Data Size:** Small datasets may benefit from simple algorithms - **Data Distribution:** Uniform distribution enables faster algorithms - **Memory Constraints:** In-place algorithms save memory - **Stability Requirements:** Stable sorts preserve relative order - **Performance Requirements:** Balance time and space complexity

Understanding these algorithms and their trade-offs is essential for: - Optimizing application performance - Choosing appropriate data structures - Implementing efficient algorithms - Solving complex computational problems - Passing technical interviews

Mastering sorting and searching techniques provides a strong foundation for advanced algorithmic thinking and problem-solving skills.

Module 10: Data Structures and Algorithms

Overview

This module explores fundamental data structures and algorithms essential for efficient C programming. You'll learn about linear data structures like dynamic arrays, linked lists, stacks, and queues, as well as more complex structures such as trees, hash tables, and graphs. The module also covers essential algorithms for sorting, searching, and graph traversal, providing a solid foundation for solving complex computational problems.

Learning Objectives

By the end of this module, you will be able to:

- Implement and use dynamic arrays (vectors) effectively
- Design and implement various types of linked lists
- Understand and apply stack and queue data structures
- Implement tree structures including binary trees and balanced trees
- Design and use hash tables for efficient data retrieval
- Apply graph algorithms for network and pathfinding problems
- Implement classic sorting and searching algorithms
- Analyze time and space complexity of algorithms
- Choose appropriate data structures for specific problems
- Apply best practices for data structure implementation in C

Chapters

1. **Linear Data Structures** - Dynamic arrays, linked lists, stacks, queues, deques, priority queues
2. **Tree Structures** - Binary trees, binary search trees, AVL trees, red-black trees, B-trees
3. **Hash Tables** - Hash function design, collision resolution, dynamic resizing, performance analysis
4. **Graph Algorithms** - Graph representations, traversal algorithms, shortest path algorithms, minimum spanning trees
5. **Sorting and Searching** - Comparison-based sorting, non-comparison sorting, binary search variations, selection algorithms

Key Concepts Covered

- Dynamic array implementation and resizing strategies
- Singly, doubly, and circular linked lists
- Stack implementation (array-based and linked-based)

- Queue implementation (circular buffer and linked-based)
- Deque and priority queue implementations
- Binary tree traversal (inorder, preorder, postorder)
- Binary search tree operations (insert, delete, search)
- AVL tree balancing and rotations
- Red-black tree properties and operations
- B-trees and B+ trees for database applications
- Hash function design and evaluation
- Collision resolution strategies (chaining, open addressing)
- Dynamic hash table resizing
- Graph representations (adjacency matrix, adjacency list)
- Breadth-first search (BFS) and depth-first search (DFS)
- Dijkstra's shortest path algorithm
- A* pathfinding algorithm
- Minimum spanning tree algorithms (Prim's, Kruskal's)
- Comparison-based sorting (quicksort, mergesort, heapsort)
- Non-comparison sorting (radix sort, counting sort, bucket sort)
- Binary search and its variations
- Selection algorithms (quickselect)
- Time and space complexity analysis
- Amortized analysis techniques
- Memory management for data structures
- Error handling in data structure operations

Prerequisites

- Completion of Module 1: Foundations and Environment Setup
- Completion of Module 2: Data Types and Variables
- Completion of Module 3: Control Flow
- Completion of Module 4: Functions and Modular Programming
- Completion of Module 5: Arrays and Strings
- Completion of Module 6: Pointers and Memory Management
- Completion of Module 7: Structures and User-Defined Types
- Completion of Module 8: File I/O and System Programming
- Completion of Module 9: Advanced C Features and Modern Standards
- Understanding of basic programming concepts
- Familiarity with C language fundamentals
- Knowledge of pointers and dynamic memory management

Tools and Technologies

- C compiler (GCC, Clang, or MSVC)
- Text editor or IDE
- Memory debugging tools (Valgrind, AddressSanitizer)
- Performance profiling tools (gprof, perf)

- Graph visualization tools (optional)
- Understanding of basic algorithms and data structures

Estimated Time to Complete

- Reading: 15-20 hours
- Exercises: 20-25 hours
- Projects: 15-20 hours

Assessment

- Complete all chapter exercises
- Successfully implement various data structures from scratch
- Demonstrate understanding of algorithm complexity analysis
- Pass data structures and algorithms quiz
- Submit comprehensive implementations of classic data structures

Next Module

After completing this module, proceed to **Module 11: Network Programming** to learn about socket programming and network communication.

Additional Resources

- “Introduction to Algorithms” by Cormen, Leiserson, Rivest, and Stein
- “Data Structures and Algorithms in C” by Adam Drozdek
- “The Algorithm Design Manual” by Steven Skiena
- Online algorithm visualization tools
- Competitive programming platforms (LeetCode, HackerRank)
- Open source data structure libraries

Network

Network Fundamentals

Network programming lets C programs exchange data across hosts. This chapter covers the models and addressing you need, then focuses on **real POSIX C**: endianness, `htons/htonl`, `inet_pton/inet_ntop`, filling `sockaddr_in`, and a complete **UDP echo client and server**.

Environment: Linux, `gcc -std=c17 -Wall -Wextra`. Link nothing extra for these examples (just `libc`).

```
gcc -std=c17 -Wall -Wextra -o prog prog.c
```

Computer Network Basics

A network is a set of interconnected devices that share a common communication path. Scope ranges from a USB gadget link to the global Internet.

Topologies (conceptual)

1. **Bus** — shared medium
2. **Star** — central switch/hub
3. **Ring** — circular path
4. **Mesh** — multiple redundant paths
5. **Tree** — hierarchical aggregation

Network types by scope

Type	Scope	Example
PAN	Person / desk	Bluetooth
LAN	Building / campus	Office Ethernet
MAN	City	Metro fiber
WAN	Large region	ISP backbone
Internet	Global	Public IP routing

OSI and TCP/IP Models

OSI (7 layers)

1. Physical — bits on wire/radio
2. Data Link — frames, MAC
3. Network — routing (IP)
4. Transport — end-to-end (TCP/UDP)
5. Session — dialogs
6. Presentation — encoding/encryption
7. Application — HTTP, DNS, ...

TCP/IP (practical 4 layers)

TCP/IP layer	Rough OSI	Examples
Link	1–2	Ethernet, Wi-Fi
Internet	3	IPv4, IPv6, ICMP
Transport	4	TCP, UDP
Application	5–7	HTTP, DNS, SSH

In socket programming you mostly touch **transport + addressing (IP + port)**.

IP Addressing and Ports

IPv4

32-bit addresses written dotted-decimal: 192.168.1.10.

Special addresses:

- 127.0.0.1 — loopback (this host)
- 0.0.0.0 — “any” address when binding a server
- 255.255.255.255 — limited broadcast (local link)

Historical classful ranges (A/B/C) matter less than **CIDR** (192.168.0.0/24), but you still see class language in older docs.

IPv6 (awareness)

128-bit addresses: `2001:db8::1`. Use `AF_INET6`, `struct sockaddr_in6`, and `inet_pton(AF_INET6, ...)`. This chapter's full programs use IPv4 for clarity; the conversion APIs are dual-stack ready.

Ports

16-bit service numbers (0–65535):

Range	Name	Examples
0–1023	Well-known	22 SSH, 80 HTTP, 443 HTTPS
1024–49151	Registered	Many app servers
49152–65535	Ephemeral	Client source ports

Binding ports below 1024 usually requires root/capabilities on Linux.

TCP vs UDP (what sockets expose)

	TCP (SOCK_STREAM)	UDP (SOCK_DGRAM)
Connection	Handshake (3-way)	None
Reliability	Retransmit, ordered	Best-effort
Message bounds	Byte stream	Datagram boundaries kept
Use cases	HTTP, SSH, databases	DNS, games, discovery, echo labs

This chapter ends with **UDP echo** because it is the smallest complete network program (no listen/accept).

Sockets and Address Structures

A **socket** is an endpoint (a file descriptor on Unix). You create one with `socket()`, then `bind/connect/sendto/recvfrom` as needed.

Headers

```
#include <sys/types.h>
#include <sys/socket.h>
#include <netinet/in.h>
#include <arpa/inet.h>
#include <unistd.h>
#include <netdb.h>
```

```
#include <errno.h>
#include <string.h>
#include <stdio.h>
#include <stdlib.h>
```

Structures

```
/* IPv4 socket address */
struct sockaddr_in {
    sa_family_t    sin_family; /* AF_INET */
    in_port_t      sin_port;   /* port in network byte order */
    struct in_addr sin_addr;    /* IPv4 address, network byte order */
    /* padding (sin_zero) often present */
};

struct in_addr {
    uint32_t s_addr;           /* address in network byte order */
};

/* Generic shape used by bind/connect (cast to/from) */
struct sockaddr {
    sa_family_t sa_family;
    char        sa_data[14];
};
```

Always set unused fields to zero (`memset` the whole `sockaddr_in`) before filling fields.

Endianness: Why Byte Order Matters

Multi-byte integers can be stored **little-endian** (least significant byte first — common on x86/x86-64) or **big-endian** (most significant byte first). **Network byte order is big-endian.**

If you put a host `uint16_t` port into a packet without conversion, peers on different architectures (or even the same machine's stack) will disagree about the value.

Program: inspect endianness and multi-byte layout

```
/* file: endian_demo.c */
#include <stdio.h>
#include <stdint.h>
#include <arpa/inet.h>
```

```

static void print_bytes(const char *label, const void *p, size_t n) {
    const unsigned char *b = p;
    size_t i;
    printf("%s:", label);
    for (i = 0; i < n; i++) {
        printf(" %02x", b[i]);
    }
    printf("\n");
}

int main(void) {
    uint32_t host_u32 = 0x12345678u;
    uint16_t host_u16 = 0xABCDu;
    uint32_t net_u32;
    uint16_t net_u16;

    {
        uint16_t probe = 0x0102;
        const unsigned char *p = (const unsigned char *)&probe;
        if (p[0] == 0x01 && p[1] == 0x02) {
            printf("This host appears big-endian\n");
        } else if (p[0] == 0x02 && p[1] == 0x01) {
            printf("This host appears little-endian\n");
        } else {
            printf("Unexpected byte order probe\n");
        }
    }

    print_bytes("host uint32 0x12345678", &host_u32, sizeof host_u32);
    print_bytes("host uint16 0xABCD", &host_u16, sizeof host_u16);

    net_u32 = htonl(host_u32);
    net_u16 = htons(host_u16);

    print_bytes("htonl(0x12345678)", &net_u32, sizeof net_u32);
    print_bytes("htons(0xABCD)", &net_u16, sizeof net_u16);

    printf("ntohl(htonl(x)) == x ? %s\n",
           ntohl(net_u32) == host_u32 ? "yes" : "no");
    printf("ntohs(htons(x)) == x ? %s\n",
           ntohs(net_u16) == host_u16 ? "yes" : "no");

    return 0;
}

```

```
gcc -std=c17 -Wall -Wextra -o endian_demo endian_demo.c
./endian_demo
```

On a little-endian host you typically see:

```
This host appears little-endian
host uint32 0x12345678: 78 56 34 12
htonl(0x12345678): 12 34 56 78
```

Network order always shows the “natural” human hex order in memory for that value’s big-endian encoding.

htons / htonl / ntohs / ntohl

```
#include <arpa/inet.h>

uint32_t htonl(uint32_t hostlong);
uint16_t htons(uint16_t hostshort);
uint32_t ntohl(uint32_t netlong);
uint16_t ntohs(uint16_t netshort);
```

Function	Direction	Width
htons	host → network	16-bit (ports)
htonl	host → network	32-bit (IPv4 s_addr pieces, some protocol fields)
ntohs	network → host	16-bit
ntohl	network → host	32-bit

Demo program

```
/* file: hton_demo.c */
#include <stdio.h>
#include <stdint.h>
#include <arpa/inet.h>

int main(void) {
    uint16_t port_host = 8080;
    uint16_t port_net = htons(port_host);
    uint32_t word_host = 0x12345678u;
```



```

printf("port host:  %u (0x%04x)\n", port_host, port_host);
printf("port net:   %u (0x%04x)  /* value after htons */\n",
      port_net, port_net);
printf("port back:  %u\n", ntohs(port_net));

printf("word host:  0x%08x\n", word_host);
printf("word net:   0x%08x\n", htonl(word_host));
printf("word back:  0x%08x\n", ntohl(htonl(word_host)));

/* Rule: store htons(port) into sin_port; print with ntohs(sin.sin_port) */
return 0;
}

gcc -std=c17 -Wall -Wextra -o hton_demo hton_demo.c
./hton_demo

```

Rule of thumb: every multi-byte field that crosses the wire (or sits in `sockaddr_in`) is converted to network order on write and to host order on read.

inet_pton and inet_ntop

Prefer these over older `inet_addr` / `inet_ntoa` (which are IPv4-only, weaker error reporting, and `inet_ntoa` is not thread-safe).

```

#include <arpa/inet.h>

int inet_pton(int af, const char *src, void *dst);
const char *inet_ntop(int af, const void *src, char *dst, socklen_t size);

```

- `inet_pton`: text $\rightarrow$ binary. Returns 1 success, 0 invalid text, -1 bad `af` / error.
- `inet_ntop`: binary $\rightarrow$ text. Returns `dst` on success, NULL on error.
- Buffer size: `INET_ADDRSTRLEN` (IPv4) or `INET6_ADDRSTRLEN` (IPv6).

Full conversion + `sockaddr_in` setup demo

```

/* file: addr_demo.c */
#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <arpa/inet.h>

```

```

#include <netinet/in.h>

int main(void) {
    const char *ip_text = "192.168.1.100";
    const uint16_t port = 9000;
    struct sockaddr_in addr;
    char ip_out[INET_ADDRSTRLEN];

    memset(&addr, 0, sizeof addr);
    addr.sin_family = AF_INET;
    addr.sin_port = htons(port);

    if (inet_pton(AF_INET, ip_text, &addr.sin_addr) != 1) {
        fprintf(stderr, "inet_pton failed for %s\n", ip_text);
        return EXIT_FAILURE;
    }

    if (inet_ntop(AF_INET, &addr.sin_addr, ip_out, sizeof ip_out) == NULL) {
        perror("inet_ntop");
        return EXIT_FAILURE;
    }

    printf("text in:  %s\n", ip_text);
    printf("text out: %s\n", ip_out);
    printf("port:      %u (wire field raw after htons: 0x%04x)\n",
           ntohs(addr.sin_port), addr.sin_port);
    printf("s_addr:    0x%08x (network order bits)\n",
           addr.sin_addr.s_addr);

    /* Bind-all pattern for servers */
    memset(&addr, 0, sizeof addr);
    addr.sin_family = AF_INET;
    addr.sin_addr.s_addr = htonl(INADDR_ANY); /* 0.0.0.0 */
    addr.sin_port = htons(9000);
    printf("server bind: 0.0.0.0:%u\n", ntohs(addr.sin_port));

    return 0;
}

```

```

gcc -std=c17 -Wall -Wextra -o addr_demo addr_demo.c
./addr_demo

```

Filling sockaddr_in checklist

```
struct sockaddr_in sa;
memset(&sa, 0, sizeof sa);
sa.sin_family = AF_INET;
sa.sin_port = htons(port);
inet_pton(AF_INET, "127.0.0.1", &sa.sin_addr);
/* or: sa.sin_addr.s_addr = htonl(INADDR_ANY); */
```

When calling `bind` / `sendto` / `recvfrom`, cast:

```
bind(fd, (struct sockaddr *)&sa, sizeof sa);
```

Minimal UDP Echo: Server and Client

UDP echo: client sends a datagram; server replies with the same payload to the source address.

Server (`udp_echo_server.c`)

```
/* file: udp_echo_server.c */
#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <unistd.h>
#include <errno.h>
#include <arpa/inet.h>
#include <netinet/in.h>
#include <sys/socket.h>

#define BUF_SIZE 2048

int main(int argc, char *argv[]) {
    int fd;
    uint16_t port;
    struct sockaddr_in srv;
    struct sockaddr_in cli;
    socklen_t cli_len;
    char buf[BUF_SIZE];
    char addrbuf[INET_ADDRSTRLEN];
    ssize_t n;

    if (argc != 2) {
```

```

    fprintf(stderr, "Usage: %s <port>\n", argv[0]);
    return EXIT_FAILURE;
}

port = (uint16_t)atoi(argv[1]);
if (port == 0) {
    fprintf(stderr, "invalid port\n");
    return EXIT_FAILURE;
}

fd = socket(AF_INET, SOCK_DGRAM, 0);
if (fd < 0) {
    perror("socket");
    return EXIT_FAILURE;
}

memset(&srv, 0, sizeof srv);
srv.sin_family = AF_INET;
srv.sin_addr.s_addr = htonl(INADDR_ANY);
srv.sin_port = htons(port);

if (bind(fd, (struct sockaddr *)&srv, sizeof srv) < 0) {
    perror("bind");
    close(fd);
    return EXIT_FAILURE;
}

printf("UDP echo server listening on port %u\n", port);

for (;;) {
    cli_len = sizeof cli;
    n = recvfrom(fd, buf, sizeof buf, 0,
                (struct sockaddr *)&cli, &cli_len);
    if (n < 0) {
        perror("recvfrom");
        continue;
    }

    if (inet_ntop(AF_INET, &cli.sin_addr, addrbuf, sizeof addrbuf) == NULL) {
        strncpy(addrbuf, "?", sizeof addrbuf);
        addrbuf[sizeof addrbuf - 1] = '\0';
    }
    printf("from %s:%u (%zd bytes)\n",
          addrbuf, ntohs(cli.sin_port), n);

    if (sendto(fd, buf, (size_t)n, 0,
               (struct sockaddr *)&cli, cli_len) < 0) {
        perror("sendto");
    }
}

```

```

    }
}

/* unreachable in this simple loop */
/* close(fd); */
/* return 0; */
}

```

Client (udp_echo_client.c)

```

/* file: udp_echo_client.c */
#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <unistd.h>
#include <arpa/inet.h>
#include <netinet/in.h>
#include <sys/socket.h>

#define BUF_SIZE 2048

int main(int argc, char *argv[]) {
    int fd;
    struct sockaddr_in srv;
    char buf[BUF_SIZE];
    ssize_t n;
    size_t len;

    if (argc != 4) {
        fprintf(stderr, "Usage: %s <server_ip> <port> <message>\n", argv[0]);
        return EXIT_FAILURE;
    }

    fd = socket(AF_INET, SOCK_DGRAM, 0);
    if (fd < 0) {
        perror("socket");
        return EXIT_FAILURE;
    }

    memset(&srv, 0, sizeof srv);
    srv.sin_family = AF_INET;
    srv.sin_port = htons((uint16_t)atoi(argv[2]));
    if (inet_pton(AF_INET, argv[1], &srv.sin_addr) != 1) {
        fprintf(stderr, "invalid server IP: %s\n", argv[1]);
        close(fd);
    }
}

```

```

        return EXIT_FAILURE;
    }

    len = strlen(argv[3]);
    if (len == 0 || len >= BUF_SIZE) {
        fprintf(stderr, "message empty or too long\n");
        close(fd);
        return EXIT_FAILURE;
    }

    if (sendto(fd, argv[3], len, 0,
               (struct sockaddr *)&srv, sizeof srv) < 0) {
        perror("sendto");
        close(fd);
        return EXIT_FAILURE;
    }

    n = recvfrom(fd, buf, sizeof buf - 1, 0, NULL, NULL);
    if (n < 0) {
        perror("recvfrom");
        close(fd);
        return EXIT_FAILURE;
    }
    buf[n] = '\0';
    printf("echo: %s\n", buf);

    close(fd);
    return 0;
}

```

Build and run (two terminals)

```

gcc -std=c17 -Wall -Wextra -o udp_echo_server udp_echo_server.c
gcc -std=c17 -Wall -Wextra -o udp_echo_client udp_echo_client.c

# terminal 1
./udp_echo_server 9000

# terminal 2
./udp_echo_client 127.0.0.1 9000 "hello UDP"
# echo: hello UDP

```

What just happened

1. Server socket + bind to 0.0.0.0:9000

2. Client `sendto` a datagram to 127.0.0.1:9000
3. Kernel delivers payload; `recvfrom` also fills **client address**
4. Server `sendto` same bytes back to that address
5. Client `recvfrom` prints the reply

No `listen`, no `accept` — that is TCP territory (next chapter).

Optional: see the packets

```
# while the exchange runs
sudo tcpdump -i lo -n udp port 9000
```

Network Errors (practical)

Always check return values:

```
if (fd < 0) {
    fprintf(stderr, "socket: %s\n", strerror(errno));
    exit(EXIT_FAILURE);
}
if (bind(fd, (struct sockaddr *)&srv, sizeof srv) < 0) {
    fprintf(stderr, "bind: %s\n", strerror(errno));
    close(fd);
    exit(EXIT_FAILURE);
}
```

Symptom	Typical cause
EADDRINUSE	Port already bound
EACCES	Port < 1024 without privilege
ECONNREFUSED	More common with TCP; UDP may get ICMP unreachable later
EINVAL	Wrong address length / family
ENETUNREACH / EHOSTUNREACH	Routing problem

UDP clients often hang in `recvfrom` if the server is down (no automatic error). Use `setsockopt` timeouts (`SO_RCVTIMEO`) for production clients.

```
struct timeval tv = {.tv_sec = 3, .tv_usec = 0};
setsockopt(fd, SOL_SOCKET, SO_RCVTIMEO, &tv, sizeof tv);
```

Hostname Lookup (brief, modern)

Prefer `getaddrinfo` over deprecated `gethostbyname`:

```
/* sketch - full dual-stack client belongs with TCP chapter */
struct addrinfo hints, *res;
memset(&hints, 0, sizeof hints);
hints.ai_family = AF_INET;
hints.ai_socktype = SOCK_DGRAM;
if (getaddrinfo("localhost", "9000", &hints, &res) != 0) {
    /* handle error */
}
/* use res->ai_addr with sendto; freeaddrinfo(res) later */
```

Protocol Quick Reference

- **TCP** — reliable stream; next chapter builds connect/accept.
 - **UDP** — datagrams; this chapter's echo lab.
 - **ICMP** — diagnostics (ping); raw sockets need privileges.
-

Exercises

Exercise 1 — Endianness report

Run `endian_demo` on your machine. Paste the byte listings for `0x12345678` before and after `htonl`. Explain which listing matches network order.

Exercise 2 — Port round-trip

Write a program that takes a port number as `argv[1]`, stores `htons` into a `sockaddr_in`, then prints `ntohs(sin_port)`. Confirm the printed port equals the argument for values 80, 8080, and 65535.

Exercise 3 — inet_pton validation

Feed `inet_pton` the strings `127.0.0.1`, `256.0.0.1`, `::1` (with `AF_INET`), and `not-an-ip`. Print the return code for each. Then convert `::1` with `AF_INET6` successfully.

Exercise 4 — Bind failure

Start `udp_echo_server 9000` twice. Capture the error from the second bind. Kill the first and confirm the second starts.

Exercise 5 — Echo timeout

Add `SO_RCVTIMEO` of 2 seconds to the client. Run the client without a server and show that it fails cleanly instead of hanging forever.

Exercise 6 — Multi-message server log

Modify the server to prefix each log line with a packet counter. Send five client messages and verify counters 1...5.

Exercise 7 — Binary payload

Send 8 raw bytes (not a C string) from a small client using an `unsigned char` array. Have the server reply with the same bytes. Confirm length stays 8 (no reliance on `\0`).

Summary

Concept	C takeaway
TCP/IP model	Sockets sit at transport + app
Endianness	Network = big-endian; use <code>hton*</code> / <code>ntoh*</code>
Text binary IP	<code>inet_pton</code> / <code>inet_ntop</code>
Address struct	Zero, set <code>AF_INET</code> , <code>htons</code> port, <code>inet_pton</code> or <code>INADDR_ANY</code>
UDP echo	<code>socket</code> → <code>bind/sendto/recvfrom</code>
Errors	Check every call; <code>strerror(errno)</code>

You now have the addressing and byte-order tools every C network program needs, plus a complete UDP exchange you can extend into real services. The next chapter builds **TCP** socket programming on the same foundations.

Socket Programming

Introduction

Socket programming is the foundation of network communication in C. Sockets provide an interface for applications to communicate over a network using various protocols such as TCP and UDP. This chapter covers the essential concepts and functions needed to create network applications using sockets.

Socket Basics

A socket is an endpoint for communication between two machines over a network. In Unix-like systems, sockets are treated as file descriptors, allowing the use of standard I/O functions for network communication.

Socket Creation

The first step in socket programming is creating a socket using the `socket()` function:

```
#include <sys/socket.h>

int socket(int domain, int type, int protocol);
```

Parameters: - domain: Communication domain (e.g., AF_INET for IPv4, AF_INET6 for IPv6)
- type: Communication type (e.g., SOCK_STREAM for TCP, SOCK_DGRAM for UDP) -
protocol: Specific protocol (usually 0 for default)

Return Value: - Socket file descriptor on success - -1 on error

Example Socket Creation

```
#include <stdio.h>
#include <stdlib.h>
#include <sys/socket.h>
#include <netinet/in.h>
#include <arpa/inet.h>

int main() {
```

```

int server_fd;

// Create TCP socket
server_fd = socket(AF_INET, SOCK_STREAM, 0);
if (server_fd == -1) {
    perror("Socket creation failed");
    exit(EXIT_FAILURE);
}

printf("Socket created successfully with file descriptor: %d\n", server_fd);

// Close socket
close(server_fd);
return 0;
}

```

Socket Address Structures

Different address structures are used for different address families:

IPv4 Address Structure

```

struct sockaddr_in {
    sa_family_t    sin_family; // Address family (AF_INET)
    in_port_t      sin_port;   // Port number (network byte order)
    struct in_addr sin_addr;    // IP address (network byte order)
    char           sin_zero[8]; // Padding to match sockaddr size
};

```

Generic Socket Address Structure

```

struct sockaddr {
    sa_family_t sa_family; // Address family
    char        sa_data[14]; // Address data
};

```

Server-Side Socket Programming

Creating a server involves several steps: socket creation, binding, listening, and accepting connections.

Binding a Socket

The `bind()` function associates a socket with a specific address and port:

```
#include <sys/socket.h>

int bind(int sockfd, const struct sockaddr *addr, socklen_t addrlen);
```

Listening for Connections

The `listen()` function makes a socket listen for incoming connections:

```
int listen(int sockfd, int backlog);
```

Accepting Connections

The `accept()` function accepts an incoming connection:

```
int accept(int sockfd, struct sockaddr *addr, socklen_t *addrlen);
```

Complete TCP Server Example

```
#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <unistd.h>
#include <sys/socket.h>
#include <netinet/in.h>
#include <arpa/inet.h>

#define PORT 8080
#define BUFFER_SIZE 1024

int main() {
    int server_fd, new_socket;
    struct sockaddr_in address;
    int opt = 1;
    int addrlen = sizeof(address);
    char buffer[BUFFER_SIZE] = {0};
    char *hello = "Hello from server";

    // Create socket file descriptor
    if ((server_fd = socket(AF_INET, SOCK_STREAM, 0)) == 0) {
```

```

        perror("socket failed");
        exit(EXIT_FAILURE);
    }

    // Forcefully attaching socket to the port
    if (setsockopt(server_fd, SOL_SOCKET, SO_REUSEADDR | SO_REUSEPORT,
        &opt, sizeof(opt))) {
        perror("setsockopt");
        exit(EXIT_FAILURE);
    }

    // Configure address structure
    address.sin_family = AF_INET;
    address.sin_addr.s_addr = INADDR_ANY;
    address.sin_port = htons(PORT);

    // Bind the socket to the network address and port
    if (bind(server_fd, (struct sockaddr *)&address, sizeof(address)) < 0) {
        perror("bind failed");
        exit(EXIT_FAILURE);
    }

    // Listen for incoming connections
    if (listen(server_fd, 3) < 0) {
        perror("listen");
        exit(EXIT_FAILURE);
    }

    printf("Server listening on port %d\n", PORT);

    // Accept incoming connection
    if ((new_socket = accept(server_fd, (struct sockaddr *)&address,
        (socklen_t *)&addrlen)) < 0) {
        perror("accept");
        exit(EXIT_FAILURE);
    }

    // Read message from client
    int valread = read(new_socket, buffer, BUFFER_SIZE);
    printf("Client: %s\n", buffer);

    // Send message to client
    send(new_socket, hello, strlen(hello), 0);
    printf("Hello message sent\n");

    // Close sockets
    close(new_socket);

```

```

        close(server_fd);

    return 0;
}

```

Client-Side Socket Programming

Creating a client involves socket creation and connecting to a server.

Connecting to a Server

The `connect()` function establishes a connection to a server:

```

#include <sys/socket.h>

int connect(int sockfd, const struct sockaddr *addr, socklen_t addrlen);

```

Complete TCP Client Example

```

#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <unistd.h>
#include <sys/socket.h>
#include <netinet/in.h>
#include <arpa/inet.h>

#define PORT 8080
#define BUFFER_SIZE 1024

int main() {
    int sock = 0;
    struct sockaddr_in serv_addr;
    char *hello = "Hello from client";
    char buffer[BUFFER_SIZE] = {0};

    // Create socket file descriptor
    if ((sock = socket(AF_INET, SOCK_STREAM, 0)) < 0) {
        printf("\n Socket creation error \n");
        return -1;
    }

    // Configure server address

```

```

serv_addr.sin_family = AF_INET;
serv_addr.sin_port = htons(PORT);

// Convert IPv4 and IPv6 addresses from text to binary form
if(inet_pton(AF_INET, "127.0.0.1", &serv_addr.sin_addr) <= 0) {
    printf("\nInvalid address/ Address not supported \n");
    return -1;
}

// Connect to server
if (connect(sock, (struct sockaddr *)&serv_addr, sizeof(serv_addr)) < 0) {
    printf("\nConnection Failed \n");
    return -1;
}

// Send message to server
send(sock, hello, strlen(hello), 0);
printf("Hello message sent\n");

// Read message from server
int valread = read(sock, buffer, BUFFER_SIZE);
printf("Server: %s\n", buffer);

// Close socket
close(sock);

return 0;
}

```

UDP Socket Programming

UDP (User Datagram Protocol) provides connectionless communication. Unlike TCP, there's no need for connection establishment.

UDP Server Example

```

#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <unistd.h>
#include <sys/socket.h>
#include <netinet/in.h>
#include <arpa/inet.h>

```

```

#define PORT 8080
#define BUFFER_SIZE 1024

int main() {
    int sockfd;
    char buffer[BUFFER_SIZE];
    struct sockaddr_in servaddr, cliaddr;

    // Create socket file descriptor
    if ((sockfd = socket(AF_INET, SOCK_DGRAM, 0)) < 0 ) {
        perror("socket creation failed");
        exit(EXIT_FAILURE);
    }

    // Clear address structures
    memset(&servaddr, 0, sizeof(servaddr));
    memset(&cliaddr, 0, sizeof(cliaddr));

    // Configure server address
    servaddr.sin_family = AF_INET;
    servaddr.sin_addr.s_addr = INADDR_ANY;
    servaddr.sin_port = htons(PORT);

    // Bind the socket with the server address
    if (bind(sockfd, (const struct sockaddr *)&servaddr, sizeof(servaddr)) < 0) {
        perror("bind failed");
        exit(EXIT_FAILURE);
    }

    int len, n;
    n = recvfrom(sockfd, (char *)buffer, BUFFER_SIZE, MSG_WAITALL,
        (struct sockaddr *) &cliaddr, &len);
    buffer[n] = '\0';
    printf("Client : %s\n", buffer);

    char *hello = "Hello from server";
    sendto(sockfd, (const char *)hello, strlen(hello), MSG_CONFIRM,
        (const struct sockaddr *) &cliaddr, len);
    printf("Hello message sent.\n");

    close(sockfd);
    return 0;
}

```


UDP Client Example

```
#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <unistd.h>
#include <sys/socket.h>
#include <netinet/in.h>
#include <arpa/inet.h>

#define PORT 8080
#define BUFFER_SIZE 1024

int main() {
    int sockfd;
    char buffer[BUFFER_SIZE];
    struct sockaddr_in servaddr;

    // Create socket file descriptor
    if ((sockfd = socket(AF_INET, SOCK_DGRAM, 0)) < 0) {
        perror("socket creation failed");
        exit(EXIT_FAILURE);
    }

    // Clear server address structure
    memset(&servaddr, 0, sizeof(servaddr));

    // Configure server address
    servaddr.sin_family = AF_INET;
    servaddr.sin_port = htons(PORT);
    servaddr.sin_addr.s_addr = INADDR_ANY;

    int n;
    char *hello = "Hello from client";
    sendto(sockfd, (const char *)hello, strlen(hello), MSG_CONFIRM,
        (const struct sockaddr *) &servaddr, sizeof(servaddr));
    printf("Hello message sent.\n");

    socklen_t len;
    n = recvfrom(sockfd, (char *)buffer, BUFFER_SIZE, MSG_WAITALL,
        (struct sockaddr *) &servaddr, &len);
    buffer[n] = '\0';
    printf("Server : %s\n", buffer);

    close(sockfd);
    return 0;
}
```

Advanced Socket Options

Socket Options

The `setsockopt()` and `getsockopt()` functions allow configuration of socket behavior:

```
#include <sys/socket.h>

int setsockopt(int sockfd, int level, int optname,
               const void *optval, socklen_t optlen);

int getsockopt(int sockfd, int level, int optname,
               void *optval, socklen_t *optlen);
```

Common Socket Options

1. **SO_REUSEADDR** - Allow reuse of local addresses
2. **SO_KEEPALIVE** - Enable keep-alive messages
3. **SO_LINGER** - Control behavior on close
4. **TCP_NODELAY** - Disable Nagle's algorithm

Example: Setting Socket Options

```
#include <stdio.h>
#include <stdlib.h>
#include <sys/socket.h>

int main() {
    int sockfd;
    int opt = 1;

    // Create socket
    sockfd = socket(AF_INET, SOCK_STREAM, 0);
    if (sockfd == -1) {
        perror("socket creation failed");
        exit(EXIT_FAILURE);
    }

    // Enable socket reuse
    if (setsockopt(sockfd, SOL_SOCKET, SO_REUSEADDR, &opt, sizeof(opt)) == -1) {
        perror("setsockopt SO_REUSEADDR failed");
        exit(EXIT_FAILURE);
    }
}
```

```

    // Enable keep-alive
    if (setsockopt(sockfd, SOL_SOCKET, SO_KEEPALIVE, &opt, sizeof(opt)) == -1) {
        perror("setsockopt SO_KEEPALIVE failed");
        exit(EXIT_FAILURE);
    }

    printf("Socket options set successfully\n");

    close(sockfd);
    return 0;
}

```

Error Handling in Socket Programming

Proper error handling is crucial in socket programming due to the unreliable nature of networks.

Common Socket Errors

1. **EACCES** - Permission denied
2. **EADDRINUSE** - Address already in use
3. **ECONNREFUSED** - Connection refused
4. **ETIMEDOUT** - Connection timed out
5. **EHOSTUNREACH** - No route to host

Robust Error Handling Example

```

#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <unistd.h>
#include <sys/socket.h>
#include <netinet/in.h>
#include <arpa/inet.h>
#include <errno.h>

int create_tcp_socket() {
    int sockfd;

    sockfd = socket(AF_INET, SOCK_STREAM, 0);
    if (sockfd == -1) {
        switch (errno) {
            case EACCES:
                fprintf(stderr, "Permission denied to create socket\n");

```

```

        break;
    case EMFILE:
        fprintf(stderr, "Process file descriptor limit reached\n");
        break;
    case ENFILE:
        fprintf(stderr, "System file descriptor limit reached\n");
        break;
    case ENOBUFFS:
        fprintf(stderr, "Insufficient buffer space\n");
        break;
    default:
        perror("socket creation failed");
    }
    return -1;
}

return sockfd;
}

int main() {
    int sockfd = create_tcp_socket();
    if (sockfd == -1) {
        exit(EXIT_FAILURE);
    }

    printf("Socket created successfully: %d\n", sockfd);
    close(sockfd);
    return 0;
}

```

Practical Examples

Simple Echo Server

```

#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <unistd.h>
#include <sys/socket.h>
#include <netinet/in.h>
#include <arpa/inet.h>
#include <time.h>

#define PORT 8080
#define BUFFER_SIZE 1024

```

```

void handle_client(int client_fd) {
    char buffer[BUFFER_SIZE];
    int bytes_read;

    while ((bytes_read = read(client_fd, buffer, BUFFER_SIZE - 1)) > 0) {
        buffer[bytes_read] = '\0';
        printf("Received: %s", buffer);

        // Echo back to client
        write(client_fd, buffer, bytes_read);

        // Exit if client sends "quit"
        if (strncmp(buffer, "quit", 4) == 0) {
            break;
        }
    }

    close(client_fd);
}

int main() {
    int server_fd, client_fd;
    struct sockaddr_in address;
    int opt = 1;
    int addrlen = sizeof(address);

    // Create socket
    if ((server_fd = socket(AF_INET, SOCK_STREAM, 0)) == 0) {
        perror("socket failed");
        exit(EXIT_FAILURE);
    }

    // Set socket options
    if (setsockopt(server_fd, SOL_SOCKET, SO_REUSEADDR | SO_REUSEPORT,
        &opt, sizeof(opt))) {
        perror("setsockopt");
        exit(EXIT_FAILURE);
    }

    // Configure address
    address.sin_family = AF_INET;
    address.sin_addr.s_addr = INADDR_ANY;
    address.sin_port = htons(PORT);

    // Bind socket
    if (bind(server_fd, (struct sockaddr *)&address, sizeof(address)) < 0) {

```

```

        perror("bind failed");
        exit(EXIT_FAILURE);
    }

    // Listen for connections
    if (listen(server_fd, 3) < 0) {
        perror("listen");
        exit(EXIT_FAILURE);
    }

    printf("Echo server listening on port %d\n", PORT);

    // Accept and handle clients
    while (1) {
        if ((client_fd = accept(server_fd, (struct sockaddr *)&address,
                                (socklen_t *)&addrlen)) < 0) {
            perror("accept");
            continue;
        }

        printf("Client connected\n");
        handle_client(client_fd);
        printf("Client disconnected\n");
    }

    close(server_fd);
    return 0;
}

```

Simple Echo Client

```

#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <unistd.h>
#include <sys/socket.h>
#include <netinet/in.h>
#include <arpa/inet.h>

#define PORT 8080
#define BUFFER_SIZE 1024

int main() {
    int sock = 0;
    struct sockaddr_in serv_addr;

```

```

char buffer[BUFFER_SIZE] = {0};

// Create socket
if ((sock = socket(AF_INET, SOCK_STREAM, 0)) < 0) {
    printf("Socket creation error\n");
    return -1;
}

// Configure server address
serv_addr.sin_family = AF_INET;
serv_addr.sin_port = htons(PORT);

// Convert IPv4 address
if (inet_pton(AF_INET, "127.0.0.1", &serv_addr.sin_addr) <= 0) {
    printf("Invalid address\n");
    return -1;
}

// Connect to server
if (connect(sock, (struct sockaddr *)&serv_addr, sizeof(serv_addr)) < 0) {
    printf("Connection failed\n");
    return -1;
}

printf("Connected to echo server. Type 'quit' to exit.\n");

// Main loop
while (1) {
    printf("Enter message: ");
    fgets(buffer, BUFFER_SIZE, stdin);

    // Remove newline
    buffer[strcspn(buffer, "\n")] = 0;

    // Send message
    send(sock, buffer, strlen(buffer), 0);

    // Exit if user types "quit"
    if (strcmp(buffer, "quit") == 0) {
        break;
    }

    // Receive echo
    int valread = read(sock, buffer, BUFFER_SIZE);
    buffer[valread] = '\0';
    printf("Echo: %s\n", buffer);
}

```

```
    close(sock);  
    return 0;  
}
```

Summary

Socket programming in C involves several key concepts:

1. **Socket Creation** - Using `socket()` to create endpoints for communication
2. **Server Programming** - Binding, listening, and accepting connections
3. **Client Programming** - Connecting to servers
4. **TCP vs UDP** - Connection-oriented vs connectionless communication
5. **Socket Options** - Configuring socket behavior with `setsockopt()`
6. **Error Handling** - Properly handling network errors and exceptions

These fundamentals form the basis for building more complex network applications. The next chapter will cover advanced networking concepts such as non-blocking I/O, select/poll, and multiplexing.

Advanced Networking

Introduction

Advanced networking concepts build upon the fundamentals of socket programming to create more sophisticated and efficient network applications. This chapter covers non-blocking I/O, multiplexing with select/poll, concurrent server architectures, and other advanced techniques that are essential for building scalable network applications.

Non-blocking I/O

Non-blocking I/O allows a program to continue execution while waiting for network operations to complete, rather than blocking until the operation finishes.

Setting Non-blocking Mode

There are two main ways to set a socket to non-blocking mode:

1. Using `fcntl()`:

```
#include <fcntl.h>

int flags = fcntl(sockfd, F_GETFL, 0);
fcntl(sockfd, F_SETFL, flags | O_NONBLOCK);
```

2. Using `ioctl()`:

```
#include <sys/ioctl.h>

int nonblocking = 1;
ioctl(sockfd, FIONBIO, &nonblocking);
```

Non-blocking Socket Example

```
#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <unistd.h>
```

```

#include <fcntl.h>
#include <errno.h>
#include <sys/socket.h>
#include <netinet/in.h>
#include <arpa/inet.h>

int main() {
    int sockfd;
    struct sockaddr_in server_addr;

    // Create socket
    sockfd = socket(AF_INET, SOCK_STREAM, 0);
    if (sockfd < 0) {
        perror("socket creation failed");
        exit(EXIT_FAILURE);
    }

    // Set socket to non-blocking mode
    int flags = fcntl(sockfd, F_GETFL, 0);
    if (fcntl(sockfd, F_SETFL, flags | O_NONBLOCK) < 0) {
        perror("fcntl failed");
        exit(EXIT_FAILURE);
    }

    // Configure server address
    memset(&server_addr, 0, sizeof(server_addr));
    server_addr.sin_family = AF_INET;
    server_addr.sin_port = htons(80);
    server_addr.sin_addr.s_addr = inet_addr("93.184.216.34"); // example.com

    // Attempt to connect (will likely return immediately with EINPROGRESS)
    int result = connect(sockfd, (struct sockaddr*)&server_addr, sizeof(server_addr));

    if (result < 0) {
        if (errno == EINPROGRESS) {
            printf("Connection in progress...\n");
            // Use select() or poll() to wait for connection completion
        } else {
            perror("connect failed");
            close(sockfd);
            exit(EXIT_FAILURE);
        }
    }

    close(sockfd);
    return 0;
}

```

I/O Multiplexing

I/O multiplexing allows a single process to monitor multiple file descriptors, waiting until one or more of them become “ready” for I/O operations.

select()

The `select()` function monitors multiple file descriptors to see if any of them are ready for I/O:

```
#include <sys/select.h>

int select(int nfd, fd_set *readfds, fd_set *writefds,
           fd_set *exceptfds, struct timeval *timeout);
```

Example: Using select() for Multiple Clients

```
#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <unistd.h>
#include <sys/select.h>
#include <sys/socket.h>
#include <netinet/in.h>
#include <arpa/inet.h>

#define PORT 8080
#define MAX_CLIENTS 30
#define BUFFER_SIZE 1024

int main() {
    int server_fd, new_socket, client_socket[MAX_CLIENTS], max_clients = MAX_CLIENTS;
    int activity, i, valread, sd;
    int max_sd;
    struct sockaddr_in address;

    char buffer[BUFFER_SIZE];

    fd_set readfds;

    // Initialize all client_socket[] to 0 (not valid)
    for (i = 0; i < max_clients; i++) {
        client_socket[i] = 0;
    }

    // Create socket
```

```

if ((server_fd = socket(AF_INET, SOCK_STREAM, 0)) == 0) {
    perror("socket failed");
    exit(EXIT_FAILURE);
}

// Set master socket to allow multiple connections
int opt = 1;
if (setsockopt(server_fd, SOL_SOCKET, SO_REUSEADDR, (char *)&opt, sizeof(opt)) < 0) {
    perror("setsockopt");
    exit(EXIT_FAILURE);
}

// Configure address
address.sin_family = AF_INET;
address.sin_addr.s_addr = INADDR_ANY;
address.sin_port = htons(PORT);

// Bind socket
if (bind(server_fd, (struct sockaddr *)&address, sizeof(address)) < 0) {
    perror("bind failed");
    exit(EXIT_FAILURE);
}

printf("Listener on port %d\n", PORT);

// Listen for connections
if (listen(server_fd, 3) < 0) {
    perror("listen");
    exit(EXIT_FAILURE);
}

// Accept connections
int addrlen = sizeof(address);
puts("Waiting for connections...");

while (1) {
    // Clear the socket set
    FD_ZERO(&readfds);

    // Add master socket to set
    FD_SET(server_fd, &readfds);
    max_sd = server_fd;

    // Add child sockets to set
    for (i = 0; i < max_clients; i++) {
        // Socket descriptor
        sd = client_socket[i];
    }
}

```

```

        // If valid socket descriptor then add to read list
        if (sd > 0)
            FD_SET(sd, &readfds);

        // Highest file descriptor number, needed for select()
        if (sd > max_sd)
            max_sd = sd;
    }

    // Wait for an activity on one of the sockets
    activity = select(max_sd + 1, &readfds, NULL, NULL, NULL);

    if ((activity < 0) && (errno != EINTR)) {
        printf("select error");
    }

    // If something happened on the master socket, then it's an incoming connection
    if (FD_ISSET(server_fd, &readfds)) {
        if ((new_socket = accept(server_fd, (struct sockaddr *)&address, (socklen_t *)&address_size)) > 0) {
            perror("accept");
            exit(EXIT_FAILURE);
        }

        // Inform user of socket number
        printf("New connection, socket fd is %d, ip is : %s, port : %d\n",
            new_socket, inet_ntoa(address.sin_addr), ntohs(address.sin_port));

        // Add new socket to array of sockets
        for (i = 0; i < max_clients; i++) {
            // If position is empty
            if (client_socket[i] == 0) {
                client_socket[i] = new_socket;
                printf("Adding to list of sockets as %d\n", i);
                break;
            }
        }
    }

    // Else it's some IO operation on some other socket
    for (i = 0; i < max_clients; i++) {
        sd = client_socket[i];

        if (FD_ISSET(sd, &readfds)) {
            // Check if it was for closing, and also read the incoming message
            if ((valread = read(sd, buffer, BUFFER_SIZE)) == 0) {
                // Somebody disconnected, get his details and print
            }
        }
    }
}

```

```

        getpeername(sd, (struct sockaddr*)&address, (socklen_t*)&addrlen);
        printf("Host disconnected, ip %s, port %d\n",
               inet_ntoa(address.sin_addr), ntohs(address.sin_port));

        // Close the socket and mark as 0 in list for reuse
        close(sd);
        client_socket[i] = 0;
    } else {
        // Set the string terminating NULL byte on the end of the data read
        buffer[valread] = '\0';
        send(sd, buffer, strlen(buffer), 0);
    }
}

}

}

return 0;
}

```

poll()

The `poll()` function provides similar functionality to `select()` but with some advantages:

```

#include <poll.h>

int poll(struct pollfd *fds, nfds_t nfds, int timeout);

```

Example: Using poll()

```

#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <unistd.h>
#include <poll.h>
#include <sys/socket.h>
#include <netinet/in.h>
#include <arpa/inet.h>

#define PORT 8080
#define MAX_CLIENTS 30
#define BUFFER_SIZE 1024

int main() {
    int server_fd, new_socket;

```

```

struct sockaddr_in address;
int addrlen = sizeof(address);

struct pollfd fds[MAX_CLIENTS + 1];
int nfds = 1; // Only server socket initially

// Initialize pollfd structure
for (int i = 0; i < MAX_CLIENTS + 1; i++) {
    fds[i].fd = -1;
    fds[i].events = POLLIN;
}

// Create server socket
if ((server_fd = socket(AF_INET, SOCK_STREAM, 0)) == 0) {
    perror("socket failed");
    exit(EXIT_FAILURE);
}

// Set socket options
int opt = 1;
if (setsockopt(server_fd, SOL_SOCKET, SO_REUSEADDR, (char *)&opt, sizeof(opt)) < 0) {
    perror("setsockopt");
    exit(EXIT_FAILURE);
}

// Configure address
address.sin_family = AF_INET;
address.sin_addr.s_addr = INADDR_ANY;
address.sin_port = htons(PORT);

// Bind and listen
if (bind(server_fd, (struct sockaddr *)&address, sizeof(address)) < 0) {
    perror("bind failed");
    exit(EXIT_FAILURE);
}

if (listen(server_fd, 3) < 0) {
    perror("listen");
    exit(EXIT_FAILURE);
}

// Add server socket to poll
fds[0].fd = server_fd;
fds[0].events = POLLIN;

printf("Server listening on port %d\n", PORT);

```

```

while (1) {
    int poll_count = poll(fds, nfds, -1);

    if (poll_count == -1) {
        perror("poll");
        exit(EXIT_FAILURE);
    }

    // Check server socket for new connections
    if (fds[0].revents & POLLIN) {
        if ((new_socket = accept(server_fd, (struct sockaddr *)&address, (socklen_t*)&address_size)) > 0) {
            perror("accept");
            exit(EXIT_FAILURE);
        }

        printf("New connection, socket fd is %d\n", new_socket);

        // Add new socket to poll array
        int i;
        for (i = 1; i < MAX_CLIENTS + 1; i++) {
            if (fds[i].fd == -1) {
                fds[i].fd = new_socket;
                nfds = (i + 1 > nfds) ? i + 1 : nfds;
                break;
            }
        }

        if (i == MAX_CLIENTS + 1) {
            printf("Too many clients\n");
            close(new_socket);
        }
    }

    // Check client sockets for data
    for (int i = 1; i < nfds; i++) {
        if (fds[i].fd != -1 && (fds[i].revents & POLLIN)) {
            char buffer[BUFFER_SIZE];
            int valread = read(fds[i].fd, buffer, BUFFER_SIZE);

            if (valread == 0) {
                // Client disconnected
                printf("Client disconnected, socket fd %d\n", fds[i].fd);
                close(fds[i].fd);
                fds[i].fd = -1;
            } else {
                // Echo data back
                buffer[valread] = '\0';

```



```

        printf("Received: %s", buffer);
        send(fds[i].fd, buffer, valread, 0);
    }
}

return 0;
}

```

Concurrent Server Architectures

Threading Approach

Using threads to handle multiple clients concurrently:

```

#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <unistd.h>
#include <pthread.h>
#include <sys/socket.h>
#include <netinet/in.h>
#include <arpa/inet.h>

#define PORT 8080
#define BUFFER_SIZE 1024

void *handle_client(void *arg) {
    int client_fd = *(int*)arg;
    char buffer[BUFFER_SIZE];
    int valread;

    // Detach thread to automatically clean up when finished
    pthread_detach(pthread_self());

    // Free the argument (allocated in main)
    free(arg);

    // Handle client communication
    while ((valread = read(client_fd, buffer, BUFFER_SIZE)) > 0) {
        buffer[valread] = '\0';
        printf("Received from client: %s", buffer);

        // Echo back to client
    }
}

```

```

        send(client_fd, buffer, valread, 0);

        // Exit if client sends "quit"
        if (strncmp(buffer, "quit", 4) == 0) {
            break;
        }
    }

    printf("Client disconnected\n");
    close(client_fd);
    return NULL;
}

int main() {
    int server_fd, client_fd;
    struct sockaddr_in address;
    int opt = 1;
    int addrlen = sizeof(address);

    // Create server socket
    if ((server_fd = socket(AF_INET, SOCK_STREAM, 0)) == 0) {
        perror("socket failed");
        exit(EXIT_FAILURE);
    }

    // Set socket options
    if (setsockopt(server_fd, SOL_SOCKET, SO_REUSEADDR | SO_REUSEPORT,
        &opt, sizeof(opt))) {
        perror("setsockopt");
        exit(EXIT_FAILURE);
    }

    // Configure address
    address.sin_family = AF_INET;
    address.sin_addr.s_addr = INADDR_ANY;
    address.sin_port = htons(PORT);

    // Bind socket
    if (bind(server_fd, (struct sockaddr *)&address, sizeof(address)) < 0) {
        perror("bind failed");
        exit(EXIT_FAILURE);
    }

    // Listen for connections
    if (listen(server_fd, 10) < 0) {
        perror("listen");
        exit(EXIT_FAILURE);
    }
}

```

```

}

printf("Server listening on port %d\n", PORT);

while (1) {
    // Accept connection
    if ((client_fd = accept(server_fd, (struct sockaddr *)&address,
                           (socklen_t *)&addrlen)) < 0) {
        perror("accept");
        continue;
    }

    printf("New client connected\n");

    // Create thread to handle client
    pthread_t thread_id;
    int *new_sock = malloc(sizeof(int));
    *new_sock = client_fd;

    if (pthread_create(&thread_id, NULL, handle_client, (void *)new_sock) < 0) {
        perror("could not create thread");
        free(new_sock);
        close(client_fd);
    }
}

close(server_fd);
return 0;
}

```

Forking Approach

Using processes to handle multiple clients:

```

#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <unistd.h>
#include <sys/socket.h>
#include <netinet/in.h>
#include <arpa/inet.h>
#include <sys/wait.h>

#define PORT 8080
#define BUFFER_SIZE 1024

```

```

void handle_client(int client_fd) {
    char buffer[BUFFER_SIZE];
    int valread;

    while ((valread = read(client_fd, buffer, BUFFER_SIZE)) > 0) {
        buffer[valread] = '\0';
        printf("Received from client: %s", buffer);

        // Echo back to client
        send(client_fd, buffer, valread, 0);

        // Exit if client sends "quit"
        if (strncmp(buffer, "quit", 4) == 0) {
            break;
        }
    }

    printf("Client disconnected\n");
    close(client_fd);
    exit(0);
}

int main() {
    int server_fd, client_fd;
    struct sockaddr_in address;
    int opt = 1;
    int addrlen = sizeof(address);

    // Create server socket
    if ((server_fd = socket(AF_INET, SOCK_STREAM, 0)) == 0) {
        perror("socket failed");
        exit(EXIT_FAILURE);
    }

    // Set socket options
    if (setsockopt(server_fd, SOL_SOCKET, SO_REUSEADDR | SO_REUSEPORT,
        &opt, sizeof(opt))) {
        perror("setsockopt");
        exit(EXIT_FAILURE);
    }

    // Configure address
    address.sin_family = AF_INET;
    address.sin_addr.s_addr = INADDR_ANY;
    address.sin_port = htons(PORT);

    // Bind socket

```

```

if (bind(server_fd, (struct sockaddr *)&address, sizeof(address)) < 0) {
    perror("bind failed");
    exit(EXIT_FAILURE);
}

// Listen for connections
if (listen(server_fd, 10) < 0) {
    perror("listen");
    exit(EXIT_FAILURE);
}

printf("Server listening on port %d\n", PORT);

while (1) {
    // Accept connection
    if ((client_fd = accept(server_fd, (struct sockaddr *)&address,
                           (socklen_t *)&addrlen)) < 0) {
        perror("accept");
        continue;
    }

    printf("New client connected\n");

    // Fork child process to handle client
    pid_t pid = fork();

    if (pid < 0) {
        perror("fork failed");
        close(client_fd);
    } else if (pid == 0) {
        // Child process
        close(server_fd); // Close listening socket in child
        handle_client(client_fd);
    } else {
        // Parent process
        close(client_fd); // Close client socket in parent
        // Clean up zombie processes
        while (waitpid(-1, NULL, WNOHANG) > 0);
    }
}

close(server_fd);
return 0;
}

```

Advanced Socket Options

Socket Buffer Sizes

Adjusting socket buffer sizes can improve performance:

```
#include <stdio.h>
#include <stdlib.h>
#include <sys/socket.h>

int main() {
    int sockfd;
    int buffer_size;
    socklen_t len = sizeof(buffer_size);

    // Create socket
    sockfd = socket(AF_INET, SOCK_STREAM, 0);
    if (sockfd < 0) {
        perror("socket creation failed");
        exit(EXIT_FAILURE);
    }

    // Get current receive buffer size
    if (getsockopt(sockfd, SOL_SOCKET, SO_RCVBUF, &buffer_size, &len) < 0) {
        perror("getsockopt failed");
        exit(EXIT_FAILURE);
    }
    printf("Default receive buffer size: %d bytes\n", buffer_size);

    // Get current send buffer size
    if (getsockopt(sockfd, SOL_SOCKET, SO_SNDBUF, &buffer_size, &len) < 0) {
        perror("getsockopt failed");
        exit(EXIT_FAILURE);
    }
    printf("Default send buffer size: %d bytes\n", buffer_size);

    // Set new buffer sizes
    buffer_size = 64 * 1024; // 64KB
    if (setsockopt(sockfd, SOL_SOCKET, SO_RCVBUF, &buffer_size, sizeof(buffer_size)) < 0) {
        perror("setsockopt SO_RCVBUF failed");
        exit(EXIT_FAILURE);
    }

    if (setsockopt(sockfd, SOL_SOCKET, SO_SNDBUF, &buffer_size, sizeof(buffer_size)) < 0) {
        perror("setsockopt SO_SNDBUF failed");
        exit(EXIT_FAILURE);
    }
}
```

```

    printf("New buffer sizes set to %d bytes\n", buffer_size);

    close(sockfd);
    return 0;
}

```

Timeout Options

Setting timeouts for socket operations:

```

#include <stdio.h>
#include <stdlib.h>
#include <sys/socket.h>
#include <netinet/in.h>
#include <time.h>

int main() {
    int sockfd;
    struct timeval timeout;

    // Create socket
    sockfd = socket(AF_INET, SOCK_STREAM, 0);
    if (sockfd < 0) {
        perror("socket creation failed");
        exit(EXIT_FAILURE);
    }

    // Set receive timeout
    timeout.tv_sec = 5; // 5 seconds
    timeout.tv_usec = 0;

    if (setsockopt(sockfd, SOL_SOCKET, SO_RCVTIMEO, &timeout, sizeof(timeout)) < 0) {
        perror("setsockopt SO_RCVTIMEO failed");
        exit(EXIT_FAILURE);
    }

    // Set send timeout
    timeout.tv_sec = 3; // 3 seconds

    if (setsockopt(sockfd, SOL_SOCKET, SO_SNDTIMEO, &timeout, sizeof(timeout)) < 0) {
        perror("setsockopt SO_SNDTIMEO failed");
        exit(EXIT_FAILURE);
    }

    printf("Socket timeouts set successfully\n");
}

```

```

        close(sockfd);
        return 0;
    }

```

Signal-Driven I/O

Signal-driven I/O allows a process to be notified when a socket is ready for I/O operations:

```

#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <unistd.h>
#include <signal.h>
#include <sys/socket.h>
#include <netinet/in.h>
#include <arpa/inet.h>
#include <fcntl.h>

int sockfd;

void signal_handler(int sig) {
    if (sig == SIGIO) {
        printf("Socket is ready for I/O\n");
        // Handle I/O operations here
    }
}

int main() {
    struct sockaddr_in address;

    // Create socket
    sockfd = socket(AF_INET, SOCK_STREAM, 0);
    if (sockfd < 0) {
        perror("socket creation failed");
        exit(EXIT_FAILURE);
    }

    // Set socket to non-blocking
    int flags = fcntl(sockfd, F_GETFL, 0);
    fcntl(sockfd, F_SETFL, flags | O_NONBLOCK);

    // Set up signal handler
    signal(SIGIO, signal_handler);

```



```

// Enable asynchronous I/O on socket
if (fcntl(sockfd, F_SETOWN, getpid()) < 0) {
    perror("fcntl F_SETOWN failed");
    exit(EXIT_FAILURE);
}

if (fcntl(sockfd, F_SETFL, flags | O_ASYNC) < 0) {
    perror("fcntl F_SETFL O_ASYNC failed");
    exit(EXIT_FAILURE);
}

// Configure address
memset(&address, 0, sizeof(address));
address.sin_family = AF_INET;
address.sin_port = htons(80);
address.sin_addr.s_addr = inet_addr("93.184.216.34"); // example.com

// Attempt to connect
connect(sockfd, (struct sockaddr*)&address, sizeof(address));

// Main loop
while (1) {
    pause(); // Wait for signals
}

close(sockfd);
return 0;
}

```

Practical Examples

High-Performance Echo Server

```

#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <unistd.h>
#include <sys/select.h>
#include <sys/socket.h>
#include <netinet/in.h>
#include <arpa/inet.h>
#include <errno.h>

#define PORT 8080

```

```

#define MAX_CLIENTS 100
#define BUFFER_SIZE 4096

typedef struct {
    int fd;
    char buffer[BUFFER_SIZE];
    int buffer_pos;
    int buffer_len;
} client_t;

int main() {
    int server_fd;
    struct sockaddr_in address;
    client_t clients[MAX_CLIENTS];
    int max_clients = MAX_CLIENTS;
    fd_set read_fds, write_fds;
    int max_fd;

    // Initialize client structures
    for (int i = 0; i < max_clients; i++) {
        clients[i].fd = -1;
        clients[i].buffer_pos = 0;
        clients[i].buffer_len = 0;
    }

    // Create server socket
    if ((server_fd = socket(AF_INET, SOCK_STREAM, 0)) == 0) {
        perror("socket failed");
        exit(EXIT_FAILURE);
    }

    // Set socket options
    int opt = 1;
    if (setsockopt(server_fd, SOL_SOCKET, SO_REUSEADDR, (char *)&opt, sizeof(opt)) < 0) {
        perror("setsockopt");
        exit(EXIT_FAILURE);
    }

    // Configure address
    address.sin_family = AF_INET;
    address.sin_addr.s_addr = INADDR_ANY;
    address.sin_port = htons(PORT);

    // Bind and listen
    if (bind(server_fd, (struct sockaddr *)&address, sizeof(address)) < 0) {
        perror("bind failed");
        exit(EXIT_FAILURE);
    }

```

```

}

if (listen(server_fd, 10) < 0) {
    perror("listen");
    exit(EXIT_FAILURE);
}

printf("High-performance echo server listening on port %d\n", PORT);

while (1) {
    // Clear file descriptor sets
    FD_ZERO(&read_fds);
    FD_ZERO(&write_fds);

    // Add server socket to read set
    FD_SET(server_fd, &read_fds);
    max_fd = server_fd;

    // Add client sockets to appropriate sets
    for (int i = 0; i < max_clients; i++) {
        if (clients[i].fd > 0) {
            FD_SET(clients[i].fd, &read_fds);

            // If we have data to send, add to write set
            if (clients[i].buffer_len > 0) {
                FD_SET(clients[i].fd, &write_fds);
            }

            if (clients[i].fd > max_fd) {
                max_fd = clients[i].fd;
            }
        }
    }

    // Wait for activity
    int activity = select(max_fd + 1, &read_fds, &write_fds, NULL, NULL);

    if (activity < 0 && errno != EINTR) {
        perror("select error");
        exit(EXIT_FAILURE);
    }

    // Check for new connections
    if (FD_ISSET(server_fd, &read_fds)) {
        int new_socket;
        int addrlen = sizeof(address);

```

```

if ((new_socket = accept(server_fd, (struct sockaddr *)&address, (socklen_t*)&ad
    perror("accept");
    continue;
}

// Add new socket to clients array
int i;
for (i = 0; i < max_clients; i++) {
    if (clients[i].fd == -1) {
        clients[i].fd = new_socket;
        clients[i].buffer_pos = 0;
        clients[i].buffer_len = 0;
        printf("New connection: socket %d\n", new_socket);
        break;
    }
}

if (i == max_clients) {
    printf("Too many clients, closing new connection\n");
    close(new_socket);
}
}

// Check for data from clients
for (int i = 0; i < max_clients; i++) {
    int sd = clients[i].fd;

    if (sd < 0) continue;

    // Check for data to read
    if (FD_ISSET(sd, &read_fds)) {
        int valread = read(sd, clients[i].buffer + clients[i].buffer_len,
                           BUFFER_SIZE - clients[i].buffer_len);

        if (valread <= 0) {
            // Client disconnected
            getpeername(sd, (struct sockaddr*)&address, (socklen_t*)&addrlen);
            printf("Client disconnected: %s:%d\n",
                   inet_ntoa(address.sin_addr), ntohs(address.sin_port));
            close(sd);
            clients[i].fd = -1;
            clients[i].buffer_pos = 0;
            clients[i].buffer_len = 0;
        } else {
            // Add data to buffer
            clients[i].buffer_len += valread;

```

```

        printf("Received %d bytes from client\n", valread);
    }
}

// Check for ability to write
if (sd > 0 && FD_ISSET(sd, &write_fds) && clients[i].buffer_len > 0) {
    int to_send = clients[i].buffer_len - clients[i].buffer_pos;
    int sent = write(sd, clients[i].buffer + clients[i].buffer_pos, to_send);

    if (sent < 0) {
        if (errno != EAGAIN && errno != EWOULDBLOCK) {
            printf("Write error, closing connection\n");
            close(sd);
            clients[i].fd = -1;
            clients[i].buffer_pos = 0;
            clients[i].buffer_len = 0;
        }
    } else {
        clients[i].buffer_pos += sent;

        // If we've sent all data, reset buffer
        if (clients[i].buffer_pos >= clients[i].buffer_len) {
            clients[i].buffer_pos = 0;
            clients[i].buffer_len = 0;
        }
    }
}

return 0;
}

```

Summary

Advanced networking concepts in C include:

1. **Non-blocking I/O** - Preventing programs from blocking on network operations
2. **I/O Multiplexing** - Using `select()` and `poll()` to handle multiple connections
3. **Concurrent Servers** - Implementing threading and forking approaches for handling multiple clients
4. **Advanced Socket Options** - Configuring buffer sizes, timeouts, and other socket parameters
5. **Signal-Driven I/O** - Using signals to notify when sockets are ready for I/O

These techniques are essential for building high-performance, scalable network applications. The next chapter will cover specific network protocols and their implementation in C.

Network Protocols

Introduction

Network protocols define the rules and conventions for communication between network devices. Understanding and implementing common network protocols is essential for developing robust network applications. This chapter covers the implementation of widely-used protocols such as HTTP, FTP, DNS, and others in C.

HTTP Protocol Implementation

HTTP (HyperText Transfer Protocol) is the foundation of data communication on the World Wide Web. Understanding HTTP is crucial for developing web applications and services.

HTTP Request Structure

An HTTP request consists of: 1. **Request Line** - Method, URI, and HTTP version 2. **Headers** - Key-value pairs providing metadata 3. **Body** - Optional data (for POST, PUT requests)

Example HTTP GET request:

```
GET /index.html HTTP/1.1
Host: www.example.com
User-Agent: Mozilla/5.0
Accept: text/html
Connection: close
```

HTTP Response Structure

An HTTP response consists of: 1. **Status Line** - HTTP version, status code, and reason phrase 2. **Headers** - Key-value pairs providing metadata 3. **Body** - Response content

Example HTTP response:

```
HTTP/1.1 200 OK
Content-Type: text/html
Content-Length: 1234
Connection: close
```

```
<HTML>...</HTML>
```

Simple HTTP Client

```
#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <unistd.h>
#include <sys/socket.h>
#include <netinet/in.h>
#include <arpa/inet.h>
#include <netdb.h>

#define BUFFER_SIZE 4096

int create_socket(const char *hostname, int port) {
    int sockfd;
    struct hostent *server;
    struct sockaddr_in serv_addr;

    // Create socket
    sockfd = socket(AF_INET, SOCK_STREAM, 0);
    if (sockfd < 0) {
        perror("ERROR opening socket");
        return -1;
    }

    // Get server IP address
    server = gethostbyname(hostname);
    if (server == NULL) {
        fprintf(stderr, "ERROR, no such host\n");
        return -1;
    }

    // Configure server address
    memset(&serv_addr, 0, sizeof(serv_addr));
    serv_addr.sin_family = AF_INET;
    memcpy(server->h_addr, &serv_addr.sin_addr.s_addr, server->h_length);
    serv_addr.sin_port = htons(port);

    // Connect to server
    if (connect(sockfd, (struct sockaddr *)&serv_addr, sizeof(serv_addr)) < 0) {
        perror("ERROR connecting");
        return -1;
    }

    return sockfd;
}
```

```

int http_get(const char *hostname, const char *path) {
    int sockfd;
    char request[BUFFER_SIZE];
    char response[BUFFER_SIZE];
    int bytes_sent, bytes_received;

    // Create socket and connect
    sockfd = create_socket(hostname, 80);
    if (sockfd < 0) {
        return -1;
    }

    // Create HTTP GET request
    snprintf(request, BUFFER_SIZE,
             "GET %s HTTP/1.1\r\n"
             "Host: %s\r\n"
             "Connection: close\r\n"
             "\r\n",
             path, hostname);

    // Send request
    bytes_sent = send(sockfd, request, strlen(request), 0);
    if (bytes_sent < 0) {
        perror("ERROR sending request");
        close(sockfd);
        return -1;
    }

    printf("HTTP Request sent:\n%s", request);

    // Receive response
    while ((bytes_received = recv(sockfd, response, BUFFER_SIZE - 1, 0)) > 0) {
        response[bytes_received] = '\0';
        printf("%s", response);
    }

    if (bytes_received < 0) {
        perror("ERROR receiving response");
        close(sockfd);
        return -1;
    }

    close(sockfd);
    return 0;
}

int main() {

```



```

    const char *hostname = "httpbin.org";
    const char *path = "/get";

    printf("Sending HTTP GET request to %s%s\n", hostname, path);
    http_get(hostname, path);

    return 0;
}

```

Simple HTTP Server

```

#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <unistd.h>
#include <sys/socket.h>
#include <netinet/in.h>
#include <arpa/inet.h>
#include <time.h>

#define PORT 8080
#define BUFFER_SIZE 4096

void send_http_response(int client_fd, const char *status, const char *content_type, const char *body) {
    char response[BUFFER_SIZE];
    time_t now;
    char time_str[100];

    // Get current time
    time(&now);
    strftime(time_str, sizeof(time_str), "%a, %d %b %Y %H:%M:%S GMT", gmtime(&now));

    // Create HTTP response
    int content_length = strlen(body);
    snprintf(response, BUFFER_SIZE,
        "HTTP/1.1 %s\r\n"
        "Date: %s\r\n"
        "Server: SimpleHTTPServer/1.0\r\n"
        "Content-Type: %s\r\n"
        "Content-Length: %d\r\n"
        "Connection: close\r\n"
        "\r\n"
        "%s",
        status, time_str, content_type, content_length, body);
}

```

```

    // Send response
    send(client_fd, response, strlen(response), 0);
}

void handle_request(int client_fd, const char *request) {
    printf("Received request:\n%s\n", request);

    // Parse request line (first line)
    char method[16], path[256], version[16];
    if (sscanf(request, "%15s %255s %15s", method, path, version) != 3) {
        send_http_response(client_fd, "400 Bad Request", "text/plain", "Bad Request");
        return;
    }

    // Handle different HTTP methods and paths
    if (strcmp(method, "GET") == 0) {
        if (strcmp(path, "/") == 0 || strcmp(path, "/index.html") == 0) {
            const char *html =
                "<!DOCTYPE html>\n"
                "<html>\n"
                "<head><title>Simple HTTP Server</title></head>\n"
                "<body><h1>Welcome to Simple HTTP Server</h1>\n"
                "<p>This is a simple HTTP server implemented in C.</p>\n"
                "</body>\n"
                "</html>";
            send_http_response(client_fd, "200 OK", "text/html", html);
        } else if (strcmp(path, "/api/time") == 0) {
            time_t now = time(NULL);
            char time_str[100];
            snprintf(time_str, sizeof(time_str), "{\"time\": \"%s\"}", ctime(&now));
            send_http_response(client_fd, "200 OK", "application/json", time_str);
        } else {
            const char *not_found =
                "<!DOCTYPE html>\n"
                "<html>\n"
                "<head><title>404 Not Found</title></head>\n"
                "<body><h1>404 Not Found</h1>\n"
                "<p>The requested resource was not found.</p>\n"
                "</body>\n"
                "</html>";
            send_http_response(client_fd, "404 Not Found", "text/html", not_found);
        }
    } else {
        send_http_response(client_fd, "405 Method Not Allowed", "text/plain", "Method Not Al
    }
}

```

```

int main() {
    int server_fd, client_fd;
    struct sockaddr_in address;
    int opt = 1;
    int addrlen = sizeof(address);
    char buffer[BUFFER_SIZE];

    // Create socket
    if ((server_fd = socket(AF_INET, SOCK_STREAM, 0)) == 0) {
        perror("socket failed");
        exit(EXIT_FAILURE);
    }

    // Set socket options
    if (setsockopt(server_fd, SOL_SOCKET, SO_REUSEADDR | SO_REUSEPORT,
        &opt, sizeof(opt))) {
        perror("setsockopt");
        exit(EXIT_FAILURE);
    }

    // Configure address
    address.sin_family = AF_INET;
    address.sin_addr.s_addr = INADDR_ANY;
    address.sin_port = htons(PORT);

    // Bind socket
    if (bind(server_fd, (struct sockaddr *)&address, sizeof(address)) < 0) {
        perror("bind failed");
        exit(EXIT_FAILURE);
    }

    // Listen for connections
    if (listen(server_fd, 10) < 0) {
        perror("listen");
        exit(EXIT_FAILURE);
    }

    printf("HTTP Server listening on port %d\n", PORT);

    while (1) {
        // Accept connection
        if ((client_fd = accept(server_fd, (struct sockaddr *)&address,
            (socklen_t*)&addrlen)) < 0) {
            perror("accept");
            continue;
        }
    }
}

```

```

        // Read request
        int bytes_read = read(client_fd, buffer, BUFFER_SIZE - 1);
        if (bytes_read > 0) {
            buffer[bytes_read] = '\0';
            handle_request(client_fd, buffer);
        }

        // Close client connection
        close(client_fd);
    }

    close(server_fd);
    return 0;
}

```

FTP Protocol Implementation

FTP (File Transfer Protocol) is a standard network protocol used for transferring files between a client and server on a computer network.

FTP Client Implementation

```

#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <unistd.h>
#include <sys/socket.h>
#include <netinet/in.h>
#include <arpa/inet.h>
#include <netdb.h>

#define BUFFER_SIZE 1024
#define FTP_PORT 21

int ftp_connect(const char *hostname, int port) {
    int sockfd;
    struct hostent *server;
    struct sockaddr_in serv_addr;

    // Create socket
    sockfd = socket(AF_INET, SOCK_STREAM, 0);
    if (sockfd < 0) {
        perror("ERROR opening socket");
        return -1;
    }
}

```

```

}

// Get server IP address
server = gethostbyname(hostname);
if (server == NULL) {
    fprintf(stderr, "ERROR, no such host\n");
    return -1;
}

// Configure server address
memset(&serv_addr, 0, sizeof(serv_addr));
serv_addr.sin_family = AF_INET;
memcpy(server->h_addr, &serv_addr.sin_addr.s_addr, server->h_length);
serv_addr.sin_port = htons(port);

// Connect to server
if (connect(sockfd, (struct sockaddr *)&serv_addr, sizeof(serv_addr)) < 0) {
    perror("ERROR connecting");
    return -1;
}

return sockfd;
}

int ftp_login(int sockfd, const char *username, const char *password) {
    char buffer[BUFFER_SIZE];
    int bytes_received;

    // Read welcome message
    bytes_received = recv(sockfd, buffer, BUFFER_SIZE - 1, 0);
    if (bytes_received > 0) {
        buffer[bytes_received] = '\0';
        printf("%s", buffer);
    }

    // Send USER command
    snprintf(buffer, BUFFER_SIZE, "USER %s\r\n", username);
    send(sockfd, buffer, strlen(buffer), 0);

    // Read response
    bytes_received = recv(sockfd, buffer, BUFFER_SIZE - 1, 0);
    if (bytes_received > 0) {
        buffer[bytes_received] = '\0';
        printf("%s", buffer);
    }

    // Send PASS command

```

```

    snprintf(buffer, BUFFER_SIZE, "PASS %s\r\n", password);
    send(sockfd, buffer, strlen(buffer), 0);

    // Read response
    bytes_received = recv(sockfd, buffer, BUFFER_SIZE - 1, 0);
    if (bytes_received > 0) {
        buffer[bytes_received] = '\0';
        printf("%s", buffer);

        // Check if login was successful
        if (strncmp(buffer, "230", 3) == 0) {
            return 0; // Success
        }
    }

    return -1; // Failure
}

int ftp_list(int sockfd) {
    char buffer[BUFFER_SIZE];

    // Send LIST command
    snprintf(buffer, BUFFER_SIZE, "LIST\r\n");
    send(sockfd, buffer, strlen(buffer), 0);

    // Read response
    int bytes_received = recv(sockfd, buffer, BUFFER_SIZE - 1, 0);
    if (bytes_received > 0) {
        buffer[bytes_received] = '\0';
        printf("%s", buffer);
        return 0;
    }

    return -1;
}

int main() {
    int sockfd;
    const char *hostname = "ftp.example.com"; // Replace with actual FTP server
    const char *username = "anonymous";
    const char *password = "user@example.com";

    printf("Connecting to FTP server: %s\n", hostname);

    // Connect to FTP server
    sockfd = ftp_connect(hostname, FTP_PORT);
    if (sockfd < 0) {

```

```

        exit(EXIT_FAILURE);
    }

    // Login to FTP server
    if (ftp_login(sockfd, username, password) < 0) {
        printf("FTP login failed\n");
        close(sockfd);
        exit(EXIT_FAILURE);
    }

    printf("FTP login successful\n");

    // List directory contents
    printf("Directory listing:\n");
    ftp_list(sockfd);

    // Close connection
    close(sockfd);
    return 0;
}

```

DNS Protocol Implementation

DNS (Domain Name System) translates domain names to IP addresses. Understanding DNS is important for network programming.

Simple DNS Resolver

```

#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <unistd.h>
#include <sys/socket.h>
#include <netinet/in.h>
#include <arpa/inet.h>
#include <netdb.h>

// DNS header structure
struct dns_header {
    unsigned short id;        // identification number
    unsigned short flags;     // flags
    unsigned short qdcount;   // number of question entries
    unsigned short ancount;   // number of answer entries
    unsigned short nscount;   // number of authority entries

```

```

    unsigned short arcount; // number of resource entries
};

// DNS question structure
struct dns_question {
    unsigned short qtype;
    unsigned short qclass;
};

// Function to convert domain name to DNS format
void dns_format(unsigned char* dns, unsigned char* host) {
    int lock = 0, i;
    strcat((char*)host, ".");

    for(i = 0; i < strlen((char*)host); i++) {
        if(host[i] == '.') {
            *dns++ = i - lock;
            for(; lock < i; lock++) {
                *dns++ = host[lock];
            }
            lock++;
        }
    }
    *dns++ = '\\0';
}

int dns_query(const char *hostname) {
    int sockfd;
    struct sockaddr_in dest;
    unsigned char buf[65536];
    struct dns_header *dns = NULL;
    struct dns_question *qinfo = NULL;

    // Create UDP socket
    sockfd = socket(AF_INET, SOCK_DGRAM, IPPROTO_UDP);
    if (sockfd < 0) {
        perror("socket creation failed");
        return -1;
    }

    // Configure destination address (Google DNS)
    dest.sin_family = AF_INET;
    dest.sin_port = htons(53);
    dest.sin_addr.s_addr = inet_addr("8.8.8.8");

    // Set up DNS header
    dns = (struct dns_header *)&buf;

```



```

dns->id = (unsigned short)htons(getpid());
dns->flags = htons(0x0100); // Standard query
dns->qdcount = htons(1);    // One question
dns->ancount = 0;
dns->nscount = 0;
dns->arcount = 0;

// Set up question section
unsigned char *qname = (unsigned char*)&buf[sizeof(struct dns_header)];
dns_format(qname, (unsigned char*)hostname);

qinfo = (struct dns_question*)&buf[sizeof(struct dns_header) + strlen((char*)qname) + 1];
qinfo->qtype = htons(1); // Type A (IPv4 address)
qinfo->qclass = htons(1); // Class IN (Internet)

// Send DNS query
int packet_size = sizeof(struct dns_header) + strlen((char*)qname) + 1 + sizeof(struct d
if (sendto(sockfd, (char*)buf, packet_size, 0, (struct sockaddr*)&dest, sizeof(dest)) <
    perror("sendto failed");
    close(sockfd);
    return -1;
}

// Receive DNS response
socklen_t dest_len = sizeof(dest);
int received_bytes = recvfrom(sockfd, (char*)buf, 65536, 0, (struct sockaddr*)&dest, &de
if (received_bytes < 0) {
    perror("recvfrom failed");
    close(sockfd);
    return -1;
}

// Parse DNS response
dns = (struct dns_header*)buf;
printf("DNS Response:\n");
printf("  ID: %d\n", ntohs(dns->id));
printf("  Questions: %d\n", ntohs(dns->qdcount));
printf("  Answers: %d\n", ntohs(dns->ancount));

// Extract IP addresses from answer section
unsigned char *reader = &buf[sizeof(struct dns_header) + strlen((char*)qname) + 1 + size

for(int i = 0; i < ntohs(dns->ancount); i++) {
    // Skip name (for simplicity, assuming it's a pointer)
    reader += 2; // Skip name pointer

```

```

    unsigned short type = ntohs(*(unsigned short*)reader);
    reader += 2;

    unsigned short class = ntohs(*(unsigned short*)reader);
    reader += 2;

    unsigned int ttl = ntohl(*(unsigned int*)reader);
    reader += 4;

    unsigned short data_len = ntohs(*(unsigned short*)reader);
    reader += 2;

    if (type == 1 && class == 1) { // Type A, Class IN
        unsigned char ip[4];
        for(int j = 0; j < 4; j++) {
            ip[j] = *reader++;
        }
        printf("  IP Address: %d.%d.%d.%d\n", ip[0], ip[1], ip[2], ip[3]);
    } else {
        reader += data_len;
    }
}

close(sockfd);
return 0;
}

int main() {
    const char *hostname = "google.com";

    printf("Resolving hostname: %s\n", hostname);
    dns_query(hostname);

    return 0;
}

```

SMTP Protocol Implementation

SMTP (Simple Mail Transfer Protocol) is used for sending emails. Here's a simple SMTP client implementation:

```

#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <unistd.h>

```

```

#include <sys/socket.h>
#include <netinet/in.h>
#include <arpa/inet.h>
#include <netdb.h>

#define BUFFER_SIZE 1024
#define SMTP_PORT 25

int smtp_connect(const char *smtp_server) {
    int sockfd;
    struct hostent *server;
    struct sockaddr_in serv_addr;

    // Create socket
    sockfd = socket(AF_INET, SOCK_STREAM, 0);
    if (sockfd < 0) {
        perror("ERROR opening socket");
        return -1;
    }

    // Get server IP address
    server = gethostbyname(smtp_server);
    if (server == NULL) {
        fprintf(stderr, "ERROR, no such host\n");
        return -1;
    }

    // Configure server address
    memset(&serv_addr, 0, sizeof(serv_addr));
    serv_addr.sin_family = AF_INET;
    memcpy(server->h_addr, &serv_addr.sin_addr.s_addr, server->h_length);
    serv_addr.sin_port = htons(SMTP_PORT);

    // Connect to server
    if (connect(sockfd, (struct sockaddr *)&serv_addr, sizeof(serv_addr)) < 0) {
        perror("ERROR connecting");
        return -1;
    }

    return sockfd;
}

int smtp_send_command(int sockfd, const char *command) {
    char buffer[BUFFER_SIZE];
    int bytes_sent, bytes_received;

    // Send command

```

```

bytes_sent = send(sockfd, command, strlen(command), 0);
if (bytes_sent < 0) {
    perror("ERROR sending command");
    return -1;
}

// Receive response
bytes_received = recv(sockfd, buffer, BUFFER_SIZE - 1, 0);
if (bytes_received > 0) {
    buffer[bytes_received] = '\0';
    printf("Server: %s", buffer);
    return 0;
}

return -1;
}

int smtp_send_email(const char *smtp_server, const char *from, const char *to,
                   const char *subject, const char *body) {
    int sockfd;
    char command[BUFFER_SIZE];

    // Connect to SMTP server
    sockfd = smtp_connect(smtp_server);
    if (sockfd < 0) {
        return -1;
    }

    // Read server greeting
    char buffer[BUFFER_SIZE];
    int bytes_received = recv(sockfd, buffer, BUFFER_SIZE - 1, 0);
    if (bytes_received > 0) {
        buffer[bytes_received] = '\0';
        printf("Server: %s", buffer);
    }

    // Send HELO command
    snprintf(command, BUFFER_SIZE, "HELO localhost\r\n");
    smtp_send_command(sockfd, command);

    // Send MAIL FROM command
    snprintf(command, BUFFER_SIZE, "MAIL FROM:<%s>\r\n", from);
    smtp_send_command(sockfd, command);

    // Send RCPT TO command
    snprintf(command, BUFFER_SIZE, "RCPT TO:<%s>\r\n", to);
    smtp_send_command(sockfd, command);
}

```

```

// Send DATA command
snprintf(command, BUFFER_SIZE, "DATA\r\n");
smtp_send_command(sockfd, command);

// Send email content
snprintf(command, BUFFER_SIZE,
    "From: %s\r\n"
    "To: %s\r\n"
    "Subject: %s\r\n"
    "\r\n"
    "%s\r\n"
    ".\r\n",
    from, to, subject, body);
send(sockfd, command, strlen(command), 0);

// Read response
bytes_received = recv(sockfd, buffer, BUFFER_SIZE - 1, 0);
if (bytes_received > 0) {
    buffer[bytes_received] = '\0';
    printf("Server: %s", buffer);
}

// Send QUIT command
snprintf(command, BUFFER_SIZE, "QUIT\r\n");
smtp_send_command(sockfd, command);

close(sockfd);
return 0;
}

int main() {
    const char *smtp_server = "smtp.example.com"; // Replace with actual SMTP server
    const char *from = "sender@example.com";
    const char *to = "recipient@example.com";
    const char *subject = "Test Email";
    const char *body = "This is a test email sent from a C program.";

    printf("Sending email via SMTP...\n");
    smtp_send_email(smtp_server, from, to, subject, body);

    return 0;
}

```

Practical Examples

Protocol Analyzer

```
#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <unistd.h>
#include <sys/socket.h>
#include <netinet/in.h>
#include <arpa/inet.h>
#include <netinet/ip.h>
#include <netinet/tcp.h>
#include <netinet/udp.h>
#include <netinet/ip_icmp.h>

void analyze_ip_packet(const unsigned char *buffer, int size) {
    struct iphdr *iph = (struct iphdr*)buffer;

    printf("IP Header\n");
    printf("    |-Version          : %d\n", (unsigned int)iph->version);
    printf("    |-Internet Header Length : %d DWORDS or %d Bytes\n",
        (unsigned int)iph->ihl, ((unsigned int)(iph->ihl))*4);
    printf("    |-Type Of Service   : %d\n", (unsigned int)iph->tos);
    printf("    |-Total Length      : %d Bytes\n", ntohs(iph->tot_len));
    printf("    |-Identification    : %d\n", ntohs(iph->id));
    printf("    |-Time To Live      : %d\n", (unsigned int)iph->ttl);
    printf("    |-Protocol          : %d\n", (unsigned int)iph->protocol);
    printf("    |-Header Checksum    : %d\n", ntohs(iph->check));
    printf("    |-Source IP         : %s\n", inet_ntoa(*(struct in_addr*)&iph->saddr));
    printf("    |-Destination IP    : %s\n", inet_ntoa(*(struct in_addr*)&iph->daddr));

    // Analyze protocol-specific headers
    switch (iph->protocol) {
        case 6: // TCP
            printf("TCP Protocol\n");
            break;
        case 17: // UDP
            printf("UDP Protocol\n");
            break;
        case 1: // ICMP
            printf("ICMP Protocol\n");
            break;
        default:
            printf("Other Protocol\n");
            break;
    }
}
```

```

    }
}

int main() {
    int sockfd;
    unsigned char buffer[65536];
    struct sockaddr_in dest;
    socklen_t dest_len = sizeof(dest);

    // Create raw socket (requires root privileges)
    sockfd = socket(AF_INET, SOCK_RAW, IPPROTO_TCP);
    if (sockfd < 0) {
        perror("socket creation failed (try running as root)");
        exit(EXIT_FAILURE);
    }

    printf("Protocol Analyzer started...\n");

    // Capture packets
    while (1) {
        int packet_size = recvfrom(sockfd, buffer, 65536, 0, (struct sockaddr*)&dest, &dest_len);
        if (packet_size < 0) {
            perror("recvfrom failed");
            continue;
        }

        printf("\n-----\n");
        analyze_ip_packet(buffer, packet_size);
        printf("-----\n");
    }

    close(sockfd);
    return 0;
}

```

Summary

Network protocol implementation in C involves understanding and working with various standard protocols:

1. **HTTP** - The foundation of web communication with request/response patterns
2. **FTP** - File transfer protocol for uploading and downloading files
3. **DNS** - Domain name resolution system
4. **SMTP** - Email transmission protocol
5. **Protocol Analysis** - Examining network packets at the protocol level

Key concepts for protocol implementation: - Understanding protocol message formats and structures
- Properly formatting requests and parsing responses - Handling different protocol states and error conditions - Managing connections and data flow - Implementing security considerations

These implementations form the basis for more complex network applications and services. The next chapter will cover network security considerations and best practices.

Network Security

Introduction

Network security is a critical aspect of modern network programming. As networks become more complex and interconnected, ensuring the confidentiality, integrity, and availability of data becomes increasingly important. This chapter covers essential network security concepts, encryption techniques, secure communication protocols, and best practices for developing secure network applications in C.

Security Fundamentals

CIA Triad

The foundation of information security is based on three core principles:

1. **Confidentiality** - Ensuring that information is accessible only to those authorized to have access
2. **Integrity** - Safeguarding the accuracy and completeness of information and processing methods
3. **Availability** - Ensuring that authorized users have access to information and associated assets when required

Common Security Threats

1. **Eavesdropping** - Unauthorized interception of network communications
2. **Man-in-the-Middle (MITM)** - Intercepting and potentially altering communication between two parties
3. **Denial of Service (DoS)** - Overwhelming a system to make it unavailable
4. **Injection Attacks** - Inserting malicious code into applications
5. **Buffer Overflows** - Exploiting memory management vulnerabilities

Cryptography Basics

Cryptography is the practice of securing communication from adversaries. Understanding basic cryptographic concepts is essential for implementing secure network applications.

Symmetric Encryption

Symmetric encryption uses the same key for both encryption and decryption.

```
#include <stdio.h>
#include <stdlib.h>
#include <string.h>

// Simple XOR cipher (for demonstration only - not secure)
void xor_encrypt_decrypt(char *data, size_t data_len, const char *key, size_t key_len) {
    for (size_t i = 0; i < data_len; i++) {
        data[i] ^= key[i % key_len];
    }
}

int main() {
    char message[] = "This is a secret message";
    char key[] = "mykey";
    size_t message_len = strlen(message);

    printf("Original message: %s\n", message);

    // Encrypt
    xor_encrypt_decrypt(message, message_len, key, strlen(key));
    printf("Encrypted message: ");
    for (size_t i = 0; i < message_len; i++) {
        printf("%02x ", (unsigned char)message[i]);
    }
    printf("\n");

    // Decrypt (XOR is symmetric)
    xor_encrypt_decrypt(message, message_len, key, strlen(key));
    printf("Decrypted message: %s\n", message);

    return 0;
}
```

Hash Functions

Hash functions create fixed-size digests of variable-length input data.

```
#include <stdio.h>
#include <stdlib.h>
#include <string.h>

// Simple hash function (for demonstration only - not cryptographically secure)
```

```

unsigned int simple_hash(const char *str) {
    unsigned int hash = 5381;
    int c;

    while ((c = *str++)) {
        hash = ((hash << 5) + hash) + c; // hash * 33 + c
    }

    return hash;
}

int main() {
    const char *message1 = "Hello, World!";
    const char *message2 = "Hello, World!";
    const char *message3 = "Hello, world!";

    printf("Hash of '%s': 0x%08x\n", message1, simple_hash(message1));
    printf("Hash of '%s': 0x%08x\n", message2, simple_hash(message2));
    printf("Hash of '%s': 0x%08x\n", message3, simple_hash(message3));

    return 0;
}

```

Secure Socket Layer (SSL/TLS)

SSL/TLS provides secure communication over computer networks. OpenSSL is a widely-used library for implementing SSL/TLS in C applications.

Installing OpenSSL

```

# Ubuntu/Debian
sudo apt-get install libssl-dev

# macOS (with Homebrew)
brew install openssl

# CentOS/RHEL
sudo yum install openssl-devel

```

Simple HTTPS Client with OpenSSL

```
#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <unistd.h>
#include <sys/socket.h>
#include <netinet/in.h>
#include <arpa/inet.h>
#include <netdb.h>
#include <openssl/ssl.h>
#include <openssl/err.h>

#define BUFFER_SIZE 4096

int create_socket(const char *hostname, int port) {
    int sockfd;
    struct hostent *server;
    struct sockaddr_in serv_addr;

    // Create socket
    sockfd = socket(AF_INET, SOCK_STREAM, 0);
    if (sockfd < 0) {
        perror("ERROR opening socket");
        return -1;
    }

    // Get server IP address
    server = gethostbyname(hostname);
    if (server == NULL) {
        fprintf(stderr, "ERROR, no such host\n");
        return -1;
    }

    // Configure server address
    memset(&serv_addr, 0, sizeof(serv_addr));
    serv_addr.sin_family = AF_INET;
    memcpy(server->h_addr, &serv_addr.sin_addr.s_addr, server->h_length);
    serv_addr.sin_port = htons(port);

    // Connect to server
    if (connect(sockfd, (struct sockaddr *)&serv_addr, sizeof(serv_addr)) < 0) {
        perror("ERROR connecting");
        return -1;
    }

    return sockfd;
}
```

```

}

SSL_CTX* init_ssl_ctx() {
    const SSL_METHOD *method;
    SSL_CTX *ctx;

    // Initialize OpenSSL
    SSL_library_init();
    SSL_load_error_strings();
    OpenSSL_add_all_algorithms();

    // Create SSL method
    method = TLS_client_method();

    // Create SSL context
    ctx = SSL_CTX_new(method);
    if (!ctx) {
        ERR_print_errors_fp(stderr);
        exit(EXIT_FAILURE);
    }

    return ctx;
}

int https_get(const char *hostname, const char *path) {
    int sockfd;
    SSL_CTX *ctx;
    SSL *ssl;
    char request[BUFFER_SIZE];
    char response[BUFFER_SIZE];
    int bytes_sent, bytes_received;

    // Initialize SSL context
    ctx = init_ssl_ctx();

    // Create socket and connect
    sockfd = create_socket(hostname, 443);
    if (sockfd < 0) {
        SSL_CTX_free(ctx);
        return -1;
    }

    // Create SSL structure
    ssl = SSL_new(ctx);
    SSL_set_fd(ssl, sockfd);

    // Perform SSL handshake

```

```

if (SSL_connect(ssl) <= 0) {
    ERR_print_errors_fp(stderr);
    SSL_free(ssl);
    close(sockfd);
    SSL_CTX_free(ctx);
    return -1;
}

printf("SSL connection established\n");
printf("SSL version: %s\n", SSL_get_version(ssl));
printf("Cipher: %s\n", SSL_get_cipher(ssl));

// Create HTTPS GET request
snprintf(request, BUFFER_SIZE,
    "GET %s HTTP/1.1\r\n"
    "Host: %s\r\n"
    "Connection: close\r\n"
    "\r\n",
    path, hostname);

// Send request over SSL
bytes_sent = SSL_write(ssl, request, strlen(request));
if (bytes_sent <= 0) {
    ERR_print_errors_fp(stderr);
    SSL_free(ssl);
    close(sockfd);
    SSL_CTX_free(ctx);
    return -1;
}

printf("HTTPS Request sent:\n%s", request);

// Receive response over SSL
while ((bytes_received = SSL_read(ssl, response, BUFFER_SIZE - 1)) > 0) {
    response[bytes_received] = '\0';
    printf("%s", response);
}

if (bytes_received < 0) {
    ERR_print_errors_fp(stderr);
    SSL_free(ssl);
    close(sockfd);
    SSL_CTX_free(ctx);
    return -1;
}

// Cleanup

```

```

        SSL_free(ssl);
        close(sockfd);
        SSL_CTX_free(ctx);

        return 0;
    }

    int main() {
        const char *hostname = "httpbin.org";
        const char *path = "/get";

        printf("Sending HTTPS GET request to %s%s\n", hostname, path);
        https_get(hostname, path);

        return 0;
    }

```

Authentication and Authorization

Authentication verifies the identity of users, while authorization determines what resources they can access.

Simple Authentication System

```

#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <unistd.h>
#include <crypt.h>

#define MAX_USERS 100
#define USERNAME_MAX 32
#define PASSWORD_MAX 128

typedef struct {
    char username[USERNAME_MAX];
    char password_hash[PASSWORD_MAX];
} User;

User users[MAX_USERS];
int user_count = 0;

// Add a new user
int add_user(const char *username, const char *password) {

```

```

    if (user_count >= MAX_USERS) {
        return -1;
    }

    // Hash the password (using crypt - not recommended for production)
    char *salt = "$6$salt$"; // SHA-512 with salt "salt"
    char *hash = crypt(password, salt);

    if (hash == NULL) {
        return -1;
    }

    strncpy(users[user_count].username, username, USERNAME_MAX - 1);
    strncpy(users[user_count].password_hash, hash, PASSWORD_MAX - 1);
    user_count++;

    return 0;
}

// Authenticate a user
int authenticate(const char *username, const char *password) {
    for (int i = 0; i < user_count; i++) {
        if (strcmp(users[i].username, username) == 0) {
            char *salt = users[i].password_hash;
            char *hash = crypt(password, salt);

            if (hash != NULL && strcmp(hash, users[i].password_hash) == 0) {
                return 1; // Authentication successful
            }
            return 0; // Wrong password
        }
    }
    return -1; // User not found
}

int main() {
    // Add some users
    add_user("admin", "admin123");
    add_user("user1", "password1");
    add_user("user2", "password2");

    // Test authentication
    const char *test_users[] = {"admin", "user1", "unknown"};
    const char *test_passwords[] = {"admin123", "wrongpass", "password1"};

    for (int i = 0; i < 3; i++) {

```



```

        printf("Authenticating user '%s' with password '%s': ",
               test_users[i], test_passwords[i]);

        int result = authenticate(test_users[i], test_passwords[i]);
        switch (result) {
            case 1:
                printf("SUCCESS\n");
                break;
            case 0:
                printf("FAILED (wrong password)\n");
                break;
            case -1:
                printf("FAILED (user not found)\n");
                break;
        }
    }

    return 0;
}

```

Secure Coding Practices

Input Validation

Proper input validation is crucial for preventing injection attacks and buffer overflows.

```

#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <ctype.h>

// Validate email address format
int validate_email(const char *email) {
    if (email == NULL) return 0;

    int len = strlen(email);
    if (len == 0 || len > 254) return 0;

    // Check for @ symbol
    char *at = strchr(email, '@');
    if (at == NULL || at == email || at == email + len - 1) return 0;

    // Check for domain part
    char *domain = at + 1;
    if (strlen(domain) == 0) return 0;
}

```

```

// Check for valid characters
for (int i = 0; i < len; i++) {
    if (!isalnum(email[i]) && email[i] != '@' && email[i] != '.' &&
        email[i] != '-' && email[i] != '_') {
        return 0;
    }
}

return 1;
}

// Safe string copy with bounds checking
char* safe_strncpy(char *dest, const char *src, size_t dest_size) {
    if (dest == NULL || src == NULL || dest_size == 0) {
        return NULL;
    }

    size_t src_len = strlen(src);
    size_t copy_len = (src_len < dest_size - 1) ? src_len : dest_size - 1;

    strncpy(dest, src, copy_len);
    dest[copy_len] = '\0';

    return dest;
}

int main() {
    // Test email validation
    const char *emails[] = {
        "user@example.com",
        "invalid.email",
        "@example.com",
        "user@",
        "user@.com",
        "valid.email@domain.co.uk"
    };

    for (int i = 0; i < 6; i++) {
        printf("Email '%s' is %s\n",
            emails[i],
            validate_email(emails[i]) ? "VALID" : "INVALID");
    }

    // Test safe string copy
    char dest[10];
    const char *src = "This is a very long string";

```

```

    printf("\nOriginal source: %s\n", src);
    if (safe_strncpy(dest, src, sizeof(dest)) != NULL) {
        printf("Safe copy result: %s\n", dest);
    } else {
        printf("Safe copy failed\n");
    }

    return 0;
}

```

Buffer Overflow Prevention

```

#include <stdio.h>
#include <stdlib.h>
#include <string.h>

// Vulnerable function (DO NOT USE)
void vulnerable_function(char *input) {
    char buffer[10];
    strcpy(buffer, input); // Buffer overflow risk
    printf("Buffer content: %s\n", buffer);
}

// Secure function
void secure_function(const char *input) {
    char buffer[10];
    // Use strncpy with proper bounds checking
    strncpy(buffer, input, sizeof(buffer) - 1);
    buffer[sizeof(buffer) - 1] = '\0'; // Ensure null termination
    printf("Buffer content: %s\n", buffer);
}

// Even more secure using snprintf
void very_secure_function(const char *input) {
    char buffer[10];
    // Use snprintf for guaranteed bounds checking
    snprintf(buffer, sizeof(buffer), "%s", input);
    printf("Buffer content: %s\n", buffer);
}

int main() {
    const char *safe_input = "Hello";
    const char *unsafe_input = "This is a very long string that will cause buffer overflow";

    printf("Testing secure function with safe input:\n");
}

```

```

    secure_function(safe_input);

    printf("\nTesting secure function with unsafe input:\n");
    secure_function(unsafe_input);

    printf("\nTesting very secure function with unsafe input:\n");
    very_secure_function(unsafe_input);

    return 0;
}

```

Firewall and Access Control

Simple Packet Filter

```

#include <stdio.h>
#include <stdlib.h>
#include <string.h>

#define MAX_RULES 100

typedef enum {
    ACTION_ALLOW,
    ACTION_DENY
} action_t;

typedef enum {
    PROTO_TCP,
    PROTO_UDP,
    PROTO_ANY
} protocol_t;

typedef struct {
    action_t action;
    protocol_t protocol;
    unsigned int src_ip;
    unsigned int dst_ip;
    unsigned short src_port;
    unsigned short dst_port;
} rule_t;

rule_t rules[MAX_RULES];
int rule_count = 0;

// Add a firewall rule

```

```

int add_rule(action_t action, protocol_t protocol,
             const char *src_ip, const char *dst_ip,
             unsigned short src_port, unsigned short dst_port) {
    if (rule_count >= MAX_RULES) {
        return -1;
    }

    rules[rule_count].action = action;
    rules[rule_count].protocol = protocol;
    // In a real implementation, convert IP strings to network byte order integers
    rules[rule_count].src_ip = 0; // Simplified
    rules[rule_count].dst_ip = 0; // Simplified
    rules[rule_count].src_port = src_port;
    rules[rule_count].dst_port = dst_port;

    rule_count++;
    return 0;
}

// Check if a packet should be allowed
int check_packet(protocol_t protocol, unsigned short src_port, unsigned short dst_port) {
    // Check rules in order (first match wins)
    for (int i = 0; i < rule_count; i++) {
        // Simplified matching - in reality, would check IP addresses too
        if ((rules[i].protocol == PROTO_ANY || rules[i].protocol == protocol) &&
            (rules[i].src_port == 0 || rules[i].src_port == src_port) &&
            (rules[i].dst_port == 0 || rules[i].dst_port == dst_port)) {

            return (rules[i].action == ACTION_ALLOW) ? 1 : 0;
        }
    }

    // Default deny
    return 0;
}

int main() {
    // Add some firewall rules
    add_rule(ACTION_DENY, PROTO_TCP, NULL, NULL, 0, 23); // Deny Telnet
    add_rule(ACTION_ALLOW, PROTO_TCP, NULL, NULL, 0, 80); // Allow HTTP
    add_rule(ACTION_ALLOW, PROTO_TCP, NULL, NULL, 0, 443); // Allow HTTPS
    add_rule(ACTION_DENY, PROTO_ANY, NULL, NULL, 0, 0); // Deny everything else

    // Test packets
    struct {
        protocol_t protocol;

```

```

        unsigned short src_port;
        unsigned short dst_port;
        const char *description;
    } test_packets[] = {
        {PROTO_TCP, 12345, 80, "HTTP request"},
        {PROTO_TCP, 12345, 443, "HTTPS request"},
        {PROTO_TCP, 12345, 23, "Telnet connection"},
        {PROTO_UDP, 12345, 53, "DNS query"}
    };

    for (int i = 0; i < 4; i++) {
        int allowed = check_packet(test_packets[i].protocol,
                                   test_packets[i].src_port,
                                   test_packets[i].dst_port);

        printf("Packet: %s -> %s\n",
               test_packets[i].description,
               allowed ? "ALLOWED" : "DENIED");
    }

    return 0;
}

```

Secure Communication Patterns

Mutual Authentication

```

#include <stdio.h>
#include <stdlib.h>
#include <string.h>

// Simple mutual authentication protocol
typedef struct {
    char client_id[32];
    char server_id[32];
    unsigned char nonce[16]; // Random number
    unsigned char session_key[32];
} auth_context_t;

// Generate random nonce (simplified)
void generate_nonce(unsigned char *nonce, size_t len) {
    // In a real implementation, use a cryptographically secure random number generator
    for (size_t i = 0; i < len; i++) {
        nonce[i] = rand() % 256;
    }
}

```

```

}

// Simple HMAC-like function (for demonstration only)
void simple_hmac(const char *key, const char *message, unsigned char *digest) {
    // In a real implementation, use a proper HMAC function
    unsigned int hash = 5381;
    const char *ptr = key;

    while (*ptr) {
        hash = ((hash << 5) + hash) + *ptr++;
    }

    ptr = message;
    while (*ptr) {
        hash = ((hash << 5) + hash) + *ptr++;
    }

    // Convert to byte array
    for (int i = 0; i < 4; i++) {
        digest[i] = (hash >> (i * 8)) & 0xFF;
    }
}

int mutual_authenticate_client(const char *client_id, const char *server_id,
                               const char *client_secret, const char *server_secret) {
    auth_context_t ctx;
    unsigned char client_hmac[16];
    unsigned char server_hmac[16];

    // Initialize context
    strncpy(ctx.client_id, client_id, sizeof(ctx.client_id) - 1);
    strncpy(ctx.server_id, server_id, sizeof(ctx.server_id) - 1);
    generate_nonce(ctx.nonce, sizeof(ctx.nonce));

    printf("Client: Starting mutual authentication with server %s\n", server_id);

    // Client sends client_id and nonce to server
    printf("Client: Sending client_id='%s' and nonce\n", client_id);

    // Server responds with its own nonce and HMAC
    generate_nonce(ctx.nonce, sizeof(ctx.nonce)); // Server's nonce
    simple_hmac(server_secret, ctx.client_id, server_hmac);

    printf("Server: Sending nonce and HMAC\n");

    // Client verifies server's HMAC and sends its own HMAC
    simple_hmac(client_secret, ctx.server_id, client_hmac);
}

```

```

    // Server verifies client's HMAC
    printf("Server: Verifying client HMAC\n");
    printf("Client: Verifying server HMAC\n");

    // If both verifications pass, establish session key
    generate_nonce(ctx.session_key, sizeof(ctx.session_key));
    printf("Mutual authentication successful. Session key established.\n");

    return 0;
}

int main() {
    srand(time(NULL));

    const char *client_id = "client123";
    const char *server_id = "server456";
    const char *client_secret = "client_secret_key";
    const char *server_secret = "server_secret_key";

    printf("Starting mutual authentication...\n");
    mutual_authenticate_client(client_id, server_id, client_secret, server_secret);

    return 0;
}

```

Practical Examples

Secure File Transfer

```

#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <unistd.h>
#include <sys/socket.h>
#include <netinet/in.h>
#include <arpa/inet.h>
#include <openssl/ssl.h>
#include <openssl/err.h>

#define BUFFER_SIZE 4096
#define PORT 8443

// Simple file encryption function
int encrypt_file(const char *input_file, const char *output_file, const char *key) {

```



```

FILE *in_fp, *out_fp;
char buffer[BUFFER_SIZE];
size_t bytes_read;
size_t key_len = strlen(key);
size_t key_index = 0;

in_fp = fopen(input_file, "rb");
if (!in_fp) {
    perror("Error opening input file");
    return -1;
}

out_fp = fopen(output_file, "wb");
if (!out_fp) {
    perror("Error opening output file");
    fclose(in_fp);
    return -1;
}

// XOR encryption (for demonstration only)
while ((bytes_read = fread(buffer, 1, BUFFER_SIZE, in_fp)) > 0) {
    for (size_t i = 0; i < bytes_read; i++) {
        buffer[i] ^= key[key_index++ % key_len];
    }
    fwrite(buffer, 1, bytes_read, out_fp);
}

fclose(in_fp);
fclose(out_fp);

return 0;
}

// Secure file transfer server
int secure_file_server() {
    int server_fd, client_fd;
    struct sockaddr_in address;
    int opt = 1;
    int addrlen = sizeof(address);

    // Create server socket
    if ((server_fd = socket(AF_INET, SOCK_STREAM, 0)) == 0) {
        perror("socket failed");
        exit(EXIT_FAILURE);
    }

    // Set socket options

```

```

if (setsockopt(server_fd, SOL_SOCKET, SO_REUSEADDR | SO_REUSEPORT,
               &opt, sizeof(opt))) {
    perror("setsockopt");
    exit(EXIT_FAILURE);
}

// Configure address
address.sin_family = AF_INET;
address.sin_addr.s_addr = INADDR_ANY;
address.sin_port = htons(PORT);

// Bind socket
if (bind(server_fd, (struct sockaddr *)&address, sizeof(address)) < 0) {
    perror("bind failed");
    exit(EXIT_FAILURE);
}

// Listen for connections
if (listen(server_fd, 3) < 0) {
    perror("listen");
    exit(EXIT_FAILURE);
}

printf("Secure file server listening on port %d\n", PORT);

while (1) {
    // Accept connection
    if ((client_fd = accept(server_fd, (struct sockaddr *)&address,
                           (socklen_t *)&addrlen)) < 0) {
        perror("accept");
        continue;
    }

    printf("Client connected\n");

    // In a real implementation, would establish SSL connection here
    // and transfer encrypted files

    close(client_fd);
}

close(server_fd);
return 0;
}

int main() {

```

```

printf("Network Security Examples\n");
printf("=====\n\n");

// Example: File encryption
const char *test_file = "test.txt";
const char *encrypted_file = "test.encrypted";
const char *key = "my_secret_key";

// Create a test file
FILE *fp = fopen(test_file, "w");
if (fp) {
    fprintf(fp, "This is a secret file that needs to be encrypted.\n");
    fprintf(fp, "It contains sensitive information.\n");
    fclose(fp);

    printf("Created test file: %s\n", test_file);

    // Encrypt the file
    if (encrypt_file(test_file, encrypted_file, key) == 0) {
        printf("File encrypted successfully: %s\n", encrypted_file);
    } else {
        printf("File encryption failed\n");
    }
}

return 0;
}

```

Summary

Network security in C programming involves several critical aspects:

1. **Cryptography Fundamentals** - Understanding encryption, hashing, and secure key management
2. **SSL/TLS Implementation** - Using OpenSSL for secure communication
3. **Authentication and Authorization** - Implementing secure user verification systems
4. **Secure Coding Practices** - Preventing common vulnerabilities like buffer overflows
5. **Firewall and Access Control** - Implementing network-level security controls
6. **Secure Communication Patterns** - Establishing trusted communication channels

Key principles for secure network programming: - Always validate and sanitize input - Use established cryptographic libraries rather than implementing your own - Implement proper error handling without revealing sensitive information - Follow the principle of least privilege - Keep software and libraries up to date - Use secure communication protocols (HTTPS, SSL/TLS) - Implement proper logging and monitoring

These security practices are essential for developing robust and trustworthy network applications that can withstand modern security threats.

Module 11: Network Programming

Overview

Network programming is a critical skill for developing modern applications that communicate over networks. This module covers the fundamentals of network programming in C, from basic concepts to advanced techniques. You'll learn how to create network applications that can communicate over local networks and the internet.

Learning Objectives

By the end of this module, you should be able to: - Understand fundamental networking concepts and protocols - Implement client-server applications using sockets - Work with different network protocols (TCP, UDP) - Handle network errors and implement robust communication - Implement advanced networking features like non-blocking I/O - Understand network security considerations

Chapters

1. **Network Fundamentals** - Basic networking concepts, OSI model, IP addressing, and network protocols
2. **Socket Programming** - Creating sockets, binding, connecting, and basic client-server communication
3. **Advanced Networking** - Non-blocking I/O, select/poll, multiplexing, and advanced socket options
4. **Network Protocols** - HTTP, FTP, DNS, and other common protocols implementation
5. **Network Security** - Secure communication, encryption, and network security best practices

Prerequisites

Before starting this module, you should have a solid understanding of: - C programming fundamentals - Pointers and memory management - File I/O operations - Basic understanding of computer networks

Key Concepts

- OSI and TCP/IP models
- IP addressing and subnetting
- Socket programming fundamentals
- Client-server architecture
- TCP and UDP protocols
- Network error handling
- Concurrent network programming
- Network security principles

Embedded

Embedded Basics

Introduction

Embedded systems are specialized computing systems designed to perform dedicated functions within larger mechanical or electrical systems. Unlike general-purpose computers, they are often resource-constrained and optimized for specific tasks. This chapter builds C-focused foundations you can practice on a Linux host: bit helpers, packed structs, `volatile`, register overlays over a buffer, a button-debounce state machine, and high-level cross-compile notes for `arm-none-eabi-gcc`.

Many examples are **host simulations** (plain `gcc` on Linux). They teach patterns you will use on bare metal without requiring a board.

Host compile convention:

```
gcc -std=c17 -Wall -Wextra -o program program.c
./program
```

What Are Embedded Systems?

An embedded system is a computer with a dedicated function inside a larger product, often with real-time constraints.

Characteristics

1. **Dedicated function** — not a general-purpose desktop OS workload
2. **Resource constraints** — limited RAM, flash, CPU, energy
3. **Real-time operation** — deadlines matter
4. **Reliability** — long uptime expectations
5. **Cost sensitivity** — BOM-driven design
6. **Power efficiency** — batteries and thermal limits

Types (rough scale)

Scale	Typical hardware	Examples
Small	8/16-bit MCU, KB of RAM	sensors, toys, simple appliances
Medium	32-bit MCU, tens–hundreds of KB	motor control, IoT nodes
Sophisticated	MPU + OS (Linux/RTOS), MB+	routers, automotive domains, phones

Microcontroller vs Microprocessor

Microcontroller (MCU): CPU + memory + peripherals on one chip; minimal external parts.

Microprocessor (MPU): CPU core; RAM, flash, and many peripherals are external; higher performance, more complex board design.

Typical Components

- Processor (CPU)
- Memory (Flash/ROM, RAM, sometimes EEPROM)
- GPIO and serial (UART, SPI, I²C)
- ADC/DAC, timers/counters
- Interrupt controller, clock tree, power management

Development Environments and Cross-Compilation

Embedded builds usually **cross-compile**: the compiler runs on a host (e.g. x86_64 Linux) and emits code for a different target (e.g. ARM Cortex-M).

High-Level Toolchain Notes (arm-none-eabi-gcc)

Piece	Role
arm-none-eabi-gcc	C compiler for bare-metal ARM (no Linux userspace ABI)
arm-none-eabi-as / ld	Assembler / linker (often driven by gcc)
arm-none-eabi-objcopy	ELF → binary/ihex for flashing
arm-none-eabi-objdump / size	Inspect sections and code size
arm-none-eabi-gdb	Debug via probe (OpenOCD, J-Link, etc.)

Piece	Role
Linker script (.ld)	Places <code>.text</code> , <code>.data</code> , <code>.bss</code> , stack at real addresses
Startup (.s / .c)	Reset handler, copy <code>.data</code> , zero <code>.bss</code> , call <code>main</code>

Typical flags (illustrative Cortex-M3):

```
# Install on Debian/Ubuntu-style hosts (example package name):
# sudo apt install gcc-arm-none-eabi binutils-arm-none-eabi

arm-none-eabi-gcc -std=c17 -Wall -Wextra \
  -mcpu=cortex-m3 -mthumb \
  -ffunction-sections -fdata-sections \
  -T linker_script.ld \
  -nostartfiles \
  -o firmware.elf startup.c main.c

arm-none-eabi-objcopy -O binary firmware.elf firmware.bin
arm-none-eabi-size firmware.elf
```

What those flags mean (high level):

- `-mcpu=cortex-m3 -mthumb` — generate Thumb code for that core
- `-ffunction-sections -fdata-sections` — enable linker garbage collection of unused sections (with `-Wl,--gc-sections`)
- `-T linker_script.ld` — memory map for flash/RAM
- `-nostartfiles` — you supply startup (or use a vendor CRT carefully)

You do **not** need this toolchain for the host simulation programs below; use native `gcc` on Linux for those.

Other Tools You Will Meet

- IDEs: vendor tools, VS Code + Cortex-Debug, CLion, etc.
- Debug probes: SWD/JTAG
- Simulators: QEMU (machine-dependent), Renode

Memory Organization

Region	Typical use
Flash / ROM	Code (<code>.text</code>), constants (<code>.rodata</code>), initializers
RAM	<code>.data</code> (copied at boot), <code>.bss</code> (zeroed), heap, stack
EEPROM / NVM	Configuration that survives power loss
MMIO registers	Peripheral control/status (volatile)

```
/* Conceptual Cortex-M style map (numbers vary by chip) */
/* 0x00000000 code flash
 * 0x20000000 SRAM
 * 0x40000000 peripherals
 * 0xE0000000 system (NVIC, SysTick, ...)
 */
```

Startup Sketch (bare metal idea)

```
/* Conceptual only - addresses come from the linker script */
extern unsigned int _etext, _data, _edata, _bss, _ebss;

void Reset_Handler(void) {
    unsigned int *src = &_etext;
    unsigned int *dst = &_data;
    while (dst < &_edata) {
        *dst++ = *src++;
    }
    dst = &_bss;
    while (dst < &_ebss) {
        *dst++ = 0;
    }
    main();
    for (;;) {
        /* halt if main returns */
    }
}
```

Section Placement Reminders

```
int initialized_var = 42;    /* .data */
int uninitialized_var;      /* .bss */
const int constant_var = 100; /* .rodata */
```

```

void example_function(void) {
    static int static_var = 10; /* .data */
    static int static_uninit;    /* .bss */
}

```

Bit Manipulation Helpers

Bit ops are daily drivers for GPIO, flags, and control registers. Prefer unsigned types and 1U shifts.

```

/* bit_helpers.c - host-runnable */
#include <stdio.h>
#include <stdint.h>
#include <stdbool.h>

static inline uint32_t bit_mask(unsigned bit) {
    return 1u << bit;
}

static inline void set_bit(uint32_t *reg, unsigned bit) {
    *reg |= bit_mask(bit);
}

static inline void clear_bit(uint32_t *reg, unsigned bit) {
    *reg &= ~bit_mask(bit);
}

static inline void toggle_bit(uint32_t *reg, unsigned bit) {
    *reg ^= bit_mask(bit);
}

static inline bool test_bit(uint32_t reg, unsigned bit) {
    return (reg & bit_mask(bit)) != 0u;
}

/* Multi-bit field helpers (field in [lsb, lsb+width)) */
static inline void write_field(uint32_t *reg, unsigned lsb, unsigned width,
                               uint32_t value) {
    uint32_t mask = ((1u << width) - 1u) << lsb;
    *reg = (*reg & ~mask) | ((value << lsb) & mask);
}

static inline uint32_t read_field(uint32_t reg, unsigned lsb, unsigned width) {

```

```

    return (reg >> lsb) & ((1u << width) - 1u);
}

int main(void) {
    uint32_t gpio = 0;

    set_bit(&gpio, 3);
    set_bit(&gpio, 1);
    printf("after set 3 and 1: 0x%08X\n", gpio);

    clear_bit(&gpio, 1);
    printf("after clear 1:      0x%08X\n", gpio);

    toggle_bit(&gpio, 3);
    toggle_bit(&gpio, 3);
    printf("after toggle 3 x2: 0x%08X (bit3=%d)\n",
           gpio, test_bit(gpio, 3));

    /* e.g. mode field bits [5:4] = 0b10 */
    write_field(&gpio, 4, 2, 0x2);
    printf("field[5:4]=%u reg=0x%08X\n",
           read_field(gpio, 4, 2), gpio);

    return 0;
}

gcc -std=c17 -Wall -Wextra -o bit_helpers bit_helpers.c
./bit_helpers

```

Macro form (common in headers; use carefully with side-effect arguments):

```

#define SET_BIT(reg, bit)    ((reg) |= (1U << (bit)))
#define CLEAR_BIT(reg, bit) ((reg) &= ~(1U << (bit)))
#define TOGGLE_BIT(reg, bit) ((reg) ^= (1U << (bit)))
#define CHECK_BIT(reg, bit) (((reg) >> (bit)) & 1U)

```

Shift safety: shifting by width of the type is undefined. Keep `bit < 32` for `uint32_t`.

volatile — Why Embedded Needs It

The compiler may cache memory values in registers or remove “useless” reads/writes. Memory-mapped registers and variables shared with ISRs **change outside the abstract machine**. Mark them `volatile` so each access in the source becomes a real access.

```

/* volatile_demo.c */
#include <stdio.h>
#include <stdint.h>

/* Simulated peripheral register (on bare metal this would be MMIO) */
static volatile uint32_t FAKE_UART_STATUS;
static volatile uint32_t FAKE_UART_DATA;

#define STATUS_RX_FULL (1u << 0)

static void hw_sim_inject(char c) {
    FAKE_UART_DATA = (uint32_t)(unsigned char)c;
    FAKE_UART_STATUS |= STATUS_RX_FULL;
}

static char uart_getc_blocking(void) {
    /* Without volatile on STATUS, an optimizer could read once and loop forever */
    while ((FAKE_UART_STATUS & STATUS_RX_FULL) == 0u) {
        /* spin - in real chips an IRQ or WFI would be better */
    }
    char c = (char)FAKE_UART_DATA;
    FAKE_UART_STATUS &= ~STATUS_RX_FULL;
    return c;
}

int main(void) {
    hw_sim_inject('Z');
    char c = uart_getc_blocking();
    printf("received: %c\n", c);
    return 0;
}

gcc -std=c17 -Wall -Wextra -O2 -o volatile_demo volatile_demo.c
./volatile_demo

```

Rules of thumb:

- MMIO registers: `volatile` (often via pointer typedef).
- Flags set in an ISR, read in main: `volatile`.
- `volatile` is **not** a substitute for mutexes on multi-core or for atomic RMW correctness across threads — it only limits *compiler* optimizations on that object.

Packed Structs

Compilers insert padding for alignment. Hardware register maps and wire formats often need **exact** layouts. Use packing attributes carefully (unaligned access cost/faults on some CPUs).

```
/* packed_structs.c */
#include <stdio.h>
#include <stdint.h>
#include <stddef.h>

struct Loose {
    uint8_t a;
    uint32_t b;
    uint8_t c;
};

struct Packed {
    uint8_t a;
    uint32_t b;
    uint8_t c;
} __attribute__((packed));

/* Bit-fields: handy for docs; packing/order is implementation-defined - verify */
struct StatusBits {
    uint8_t ready : 1;
    uint8_t error : 1;
    uint8_t mode : 2;
    uint8_t : 4; /* padding bits */
};

int main(void) {
    printf("sizeof(Loose) = %zu\n", sizeof(struct Loose));
    printf("sizeof(Packed) = %zu\n", sizeof(struct Packed));
    printf("offset b Loose = %zu\n", offsetof(struct Loose, b));
    printf("offset b Packed = %zu\n", offsetof(struct Packed, b));

    struct Packed p = { .a = 1, .b = 0x11223344u, .c = 0xFE };
    const unsigned char *bytes = (const unsigned char *)&p;
    printf("Packed bytes:");
    for (size_t i = 0; i < sizeof p; i++) {
        printf(" %02X", bytes[i]);
    }
    printf("\n");

    struct StatusBits st = { .ready = 1, .error = 0, .mode = 3 };
    printf("status ready=%u error=%u mode=%u\n", st.ready, st.error, st.mode);
    return 0;
}
```

```
}
```

```
gcc -std=c17 -Wall -Wextra -o packed_structs packed_structs.c  
./packed_structs
```

On a typical 64-bit Linux host, `Loose` is larger than 6 bytes due to padding; `Packed` is 6 bytes. Always check with `sizeof` / `offsetof` for your ABI.

Memory-Mapped Register Pattern as Struct Overlay

On hardware, a peripheral's registers sit at a fixed address. In C you often overlay a `struct` on that address. On the host, overlay the same struct on a **byte buffer** to practice without silicon.

```
/* mmio_overlay.c - simulate a tiny GPIO block in a buffer */  
#include <stdio.h>  
#include <stdint.h>  
#include <string.h>  
#include <stdbool.h>  
  
typedef struct {  
    volatile uint32_t DIR;    /* 0x00: 1 = output */  
    volatile uint32_t DATA;  /* 0x04: pin levels */  
    volatile uint32_t SET;    /* 0x08: write 1 to set bit in DATA */  
    volatile uint32_t CLEAR;  /* 0x0C: write 1 to clear bit in DATA */  
} GpioBlock;  
  
/* Host "device memory" */  
static uint8_t device_mem[sizeof(GpioBlock)];  
  
static GpioBlock *gpio_at_buffer(void) {  
    return (GpioBlock *) (void *) device_mem;  
}  
  
/* Emulate SET/CLEAR side effects a real peripheral would do in hardware */  
static void gpio_bus_write(GpioBlock *g, volatile uint32_t *reg, uint32_t value) {  
    if (reg == &g->SET) {  
        g->DATA |= value;  
    } else if (reg == &g->CLEAR) {  
        g->DATA &= ~value;  
    } else {  
        *reg = value;  
    }  
}
```



```

static void gpio_init_outputs(GpioBlock *g, uint32_t mask) {
    gpio_bus_write(g, &g->DIR, g->DIR | mask);
}

static void gpio_set_pins(GpioBlock *g, uint32_t mask) {
    gpio_bus_write(g, &g->SET, mask);
}

static void gpio_clear_pins(GpioBlock *g, uint32_t mask) {
    gpio_bus_write(g, &g->CLEAR, mask);
}

int main(void) {
    memset(device_mem, 0, sizeof device_mem);
    GpioBlock *gpio = gpio_at_buffer();

    const uint32_t LED = (1u << 0);
    gpio_init_outputs(gpio, LED);

    gpio_set_pins(gpio, LED);
    printf("DIR=0x%X DATA=0x%X (LED on)\n", gpio->DIR, gpio->DATA);

    gpio_clear_pins(gpio, LED);
    printf("DIR=0x%X DATA=0x%X (LED off)\n", gpio->DIR, gpio->DATA);

    /* On bare metal you would instead write:
     * #define GPIOA ((GpioBlock *)0x40020000u)
     * GPIOA->DIR |= LED;
     * with the real memory map - never do that to a random address on Linux.
     */
    return 0;
}

```

```

gcc -std=c17 -Wall -Wextra -o mmio_overlay mmio_overlay.c
./mmio_overlay

```

Bare-metal sketch (do not run as a normal Linux userspace poke):

```

#define GPIOA_BASE 0x40020000u
#define GPIOA ((volatile GpioBlock *)GPIOA_BASE)

```

Software Model: Button Debounce State Machine

Mechanical buttons bounce. A small state machine filters chatter before you act on a press/release.

```
/* button_debounce_fsm.c */
#include <stdio.h>
#include <stdint.h>
#include <stdbool.h>
#include <string.h>

typedef enum {
    BTN_IDLE,
    BTN_DEBOUNCE_PRESS,
    BTN_PRESSED,
    BTN_DEBOUNCE_RELEASE
} BtnState;

typedef struct {
    BtnState state;
    uint32_t ticks; /* time spent in current state */
    uint32_t threshold; /* ticks required to accept transition */
    bool stable_down; /* debounced level: true = pressed */
} BtnFsm;

static void btn_init(BtnFsm *b, uint32_t threshold) {
    memset(b, 0, sizeof *b);
    b->state = BTN_IDLE;
    b->threshold = threshold;
}

/* raw_down: instantaneous sample (active high here) */
static void btn_update(BtnFsm *b, bool raw_down) {
    switch (b->state) {
        case BTN_IDLE:
            if (raw_down) {
                b->state = BTN_DEBOUNCE_PRESS;
                b->ticks = 0;
            }
            break;

        case BTN_DEBOUNCE_PRESS:
            if (!raw_down) {
                b->state = BTN_IDLE;
                b->ticks = 0;
            } else if (++b->ticks >= b->threshold) {
                b->state = BTN_PRESSED;
            }
    }
}
```

```

        b->stable_down = true;
        printf("event: PRESS accepted\n");
    }
    break;

case BTN_PRESSED:
    if (!raw_down) {
        b->state = BTN_DEBOUNCE_RELEASE;
        b->ticks = 0;
    }
    break;

case BTN_DEBOUNCE_RELEASE:
    if (raw_down) {
        b->state = BTN_PRESSED;
        b->ticks = 0;
    } else if (++b->ticks >= b->threshold) {
        b->state = BTN_IDLE;
        b->stable_down = false;
        printf("event: RELEASE accepted\n");
    }
    break;
}
}

/* Simulated noisy waveform: 0=up, 1=down. Called once per "ms". */
int main(void) {
    BtnFsm btn;
    btn_init(&btn, 3); /* need 3 stable samples */

    /* bounce on press, solid hold, bounce on release */
    const char *wave =
        "000" /* idle */
        "10110111" /* noisy press */
        "11111111" /* held */
        "01001000" /* noisy release */
        "00000";

    printf("raw sequence (1=down): %s\n", wave);
    for (size_t i = 0; wave[i] != '\0'; i++) {
        bool raw = (wave[i] == '1');
        btn_update(&btn, raw);
        printf("t=%02zu raw=%d state=%d stable=%d\n",
            i, raw, (int)btn.state, btn.stable_down);
    }
    return 0;
}

```

```
gcc -std=c17 -Wall -Wextra -o button_debounce_fsm button_debounce_fsm.c
./button_debounce_fsm
```

Tune threshold to your sample rate (e.g. 20 ms at 1 kHz sampling → threshold 20).

Classic Patterns (Still Useful)

LED Blink Superloop (conceptual MMIO)

```
#include <stdint.h>

#define GPIO_PORT_ADDR 0x40020000u
#define GPIO_DIR (*(volatile uint32_t *) (GPIO_PORT_ADDR + 0x00))
#define GPIO_DATA (*(volatile uint32_t *) (GPIO_PORT_ADDR + 0x04))
#define LED_PIN 0

/* On host, prefer the buffer overlay example - these addresses are not yours. */

void delay(volatile uint32_t count) {
    while (count--) {
        __asm__ volatile("nop");
    }
}

void led_init(void) {
    GPIO_DIR |= (1u << LED_PIN);
}

int main(void) {
    led_init();
    for (;;) {
        GPIO_DATA |= (1u << LED_PIN);
        delay(1000000);
        GPIO_DATA &= ~(1u << LED_PIN);
        delay(1000000);
    }
}
```

Interrupt Sketch

```
void UART_IRQHandler(void) {
    uint32_t status = UART_STATUS_REG;
    if (status & UART_RX_INTERRUPT) {
        char data = (char)UART_DATA_REG;
        process_received_data(data);
        UART_STATUS_REG = UART_RX_INTERRUPT; /* clear */
    }
}
```

Keep ISRs short: set flags / push to a ring buffer; do heavy work in the main loop.

Timing

Busy-wait delays waste power and drift with optimization and clock changes. Prefer hardware timers or a SysTick counter for production firmware. Software `nop` loops are only for crude bring-up.

Practical Host Lab: Mini “Firmware” Module

Combine bits + overlay + FSM in one translation unit you can test on Linux.

```
/* mini_fw_sim.c */
#include <stdio.h>
#include <stdint.h>
#include <stdbool.h>
#include <string.h>

typedef struct {
    volatile uint32_t DIR;
    volatile uint32_t DATA;
} SimpleGpio;

static uint8_t ram[sizeof(SimpleGpio)];
static SimpleGpio *GPIO;

static inline void set_bit_u32(volatile uint32_t *r, unsigned b) {
    *r |= (1u << b);
}

static inline void clear_bit_u32(volatile uint32_t *r, unsigned b) {
    *r &= ~(1u << b);
}
```

```

#define LED 0u
#define BTN 1u

typedef enum { S_OFF, S_ON } LedMode;

int main(void) {
    memset(ram, 0, sizeof ram);
    GPIO = (SimpleGpio *) (void *) ram;

    set_bit_u32(&GPIO->DIR, LED);    /* LED out */
    clear_bit_u32(&GPIO->DIR, BTN); /* BTN in */

    LedMode mode = S_OFF;
    /* Pretend samples of button (active high) every step */
    const int samples[] = {0, 0, 1, 1, 1, 1, 0, 0};
    int stable = 0;
    int last = 0;

    for (size_t i = 0; i < sizeof samples / sizeof samples[0]; i++) {
        int raw = samples[i];
        if (raw == last) {
            stable++;
        } else {
            stable = 0;
            last = raw;
        }
        if (stable == 2 && raw == 1) {
            mode = (mode == S_OFF) ? S_ON : S_OFF;
            printf("toggle -> %s\n", mode == S_ON ? "ON" : "OFF");
        }
        if (mode == S_ON) {
            set_bit_u32(&GPIO->DATA, LED);
        } else {
            clear_bit_u32(&GPIO->DATA, LED);
        }
        printf("i=%zu raw=%d DATA=0x%X\n", i, raw, GPIO->DATA);
    }
    return 0;
}

```

```

gcc -std=c17 -Wall -Wextra -o mini_fw_sim mini_fw_sim.c
./mini_fw_sim

```

Exercises

Exercise 1 — Bit library

Implement `set`, `clear`, `toggle`, `test`, and `write_field` / `read_field` in a small header. Write tests that print expected masks for bits 0, 7, 31 on `uint32_t`. Reject bit `>= 32` safely.

```
gcc -std=c17 -Wall -Wextra -o bit_tests bit_tests.c
```

Exercise 2 — Packed vs unpacked

Define a packet:

```
u8 type; u16 length; u8 flags; u32 crc;
```

Print `sizeof` and offsets with and without `__attribute__((packed))`. Serialize to a `uint8_t` buffer with explicit byte writes (portable) and compare to a packed `memcpy` approach.

Exercise 3 — Overlay UART

Model a UART with `STATUS` (`RX_FULL`, `TX_EMPTY`) and `DATA` registers in a buffer. Write `uart_putc` / `uart_getc` that spin on flags. Inject bytes from `main` by writing the buffer as the “hardware” would.

Exercise 4 — Debounce FSM extension

Add a long-press state: if pressed longer than `L` ticks, emit `LONG_PRESS` once. Draw the state diagram in comments.

Exercise 5 — Cross-compile dry run

If `arm-none-eabi-gcc` is installed, compile a tiny `main` that only returns 0:

```
arm-none-eabi-gcc -std=c17 -Wall -Wextra -mcpu=cortex-m3 -mthumb \  
-c -o main.o main.c  
arm-none-eabi-size main.o
```

Note: without startup and a linker script you will not produce a full flashable image; the goal is to confirm the cross compiler runs and inspect section sizes.

Exercise 6 — Volatile experiment

Write a loop waiting on a non-volatile global that a second function would change (on one core with `-O2`, the wait may spin forever if the compiler caches the load). Then mark the flag `volatile` and show the wait completes when you flip the flag between calls (e.g. cooperative single-thread simulation).

Summary

Embedded C fundamentals covered here:

1. **Architecture** — MCU vs MPU, memory map, startup and sections
2. **Cross-compilation** — role of `arm-none-eabi-gcc`, flags, linker script, objcopy
3. **Bit helpers** — set/clear/toggle/test and field pack/unpack
4. **volatile** — MMIO and ISR-shared flags; not a lock
5. **Packed structs** — layout control and portability caveats
6. **Struct overlay** — register blocks on real addresses or host buffers
7. **Debounce FSM** — software model for noisy digital inputs
8. **Host labs** — practice patterns safely with Linux `gcc`

Next chapters go deeper into low-level techniques, embedded-specific C patterns, real-time structure, and hardware interfaces.

Low-Level Programming

Introduction

Low-level programming in embedded systems involves direct interaction with hardware components through memory-mapped I/O, bit manipulation, and precise control of system resources. This chapter explores advanced techniques for accessing hardware peripherals, managing memory at the register level, and optimizing code for resource-constrained environments.

Memory-Mapped I/O

Memory-mapped I/O allows software to interact with hardware peripherals by reading from and writing to specific memory addresses that correspond to hardware registers.

Direct Register Access

```
#include <stdint.h>

// Example register definitions for a microcontroller
#define PERIPH_BASE      0x40000000
#define GPIOA_BASE      (PERIPH_BASE + 0x00000)
#define USART1_BASE     (PERIPH_BASE + 0x10000)

// GPIOA registers
#define GPIOA_MODER      (*(volatile uint32_t*)(GPIOA_BASE + 0x00))
#define GPIOA_OTYPER     (*(volatile uint32_t*)(GPIOA_BASE + 0x04))
#define GPIOA_OSPEEDR    (*(volatile uint32_t*)(GPIOA_BASE + 0x08))
#define GPIOA_PUPDR      (*(volatile uint32_t*)(GPIOA_BASE + 0x0C))
#define GPIOA_IDR        (*(volatile uint32_t*)(GPIOA_BASE + 0x10))
#define GPIOA_ODR        (*(volatile uint32_t*)(GPIOA_BASE + 0x14))
#define GPIOA_BSRR       (*(volatile uint32_t*)(GPIOA_BASE + 0x18))
#define GPIOA_LCKR       (*(volatile uint32_t*)(GPIOA_BASE + 0x1C))

// USART1 registers
#define USART1_SR        (*(volatile uint32_t*)(USART1_BASE + 0x00))
#define USART1_DR        (*(volatile uint32_t*)(USART1_BASE + 0x04))
#define USART1_BRR       (*(volatile uint32_t*)(USART1_BASE + 0x08))
#define USART1_CR1       (*(volatile uint32_t*)(USART1_BASE + 0x0C))
```

```

#define USART1_CR2      (*(volatile uint32_t*)(USART1_BASE + 0x10))
#define USART1_CR3      (*(volatile uint32_t*)(USART1_BASE + 0x14))

// Configure GPIO pin as output
void gpio_config_output(uint8_t pin) {
    // Set pin mode to output (01)
    GPIOA_MODER &= ~(3UL << (pin * 2));    // Clear mode bits
    GPIOA_MODER |= (1UL << (pin * 2));    // Set output mode

    // Set output type to push-pull (0)
    GPIOA_OTYPER &= ~(1UL << pin);

    // Set output speed to high (11)
    GPIOA_OSPEEDR |= (3UL << (pin * 2));

    // Set pull-up/pull-down to no pull (00)
    GPIOA_PUPDR &= ~(3UL << (pin * 2));
}

// Set GPIO pin high
void gpio_set(uint8_t pin) {
    GPIOA_BSRR = (1UL << pin);
}

// Clear GPIO pin (set low)
void gpio_clear(uint8_t pin) {
    GPIOA_BSRR = (1UL << (pin + 16));
}

// Toggle GPIO pin
void gpio_toggle(uint8_t pin) {
    GPIOA_ODR ^= (1UL << pin);
}

// Read GPIO pin state
uint8_t gpio_read(uint8_t pin) {
    return (GPIOA_IDR & (1UL << pin)) ? 1 : 0;
}

```

Bit Manipulation Techniques

Bit manipulation is fundamental in embedded programming for efficient hardware control.

Advanced Bit Operations

```
#include <stdint.h>

// Bit manipulation macros
#define BIT(n) (1UL << (n))
#define BIT_MASK(len) (BIT(len) - 1)
#define BIT_FIELD(val, pos, len) (((val) & BIT_MASK(len)) << (pos))

#define SET_BIT(reg, bit) ((reg) |= BIT(bit))
#define CLEAR_BIT(reg, bit) ((reg) &= ~BIT(bit))
#define TOGGLE_BIT(reg, bit) ((reg) ^= BIT(bit))
#define CHECK_BIT(reg, bit) (((reg) & BIT(bit)) != 0)

#define SET_BITS(reg, pos, len, val) \
    ((reg) = ((reg) & ~(BIT_MASK(len) << (pos))) | BIT_FIELD(val, pos, len))

#define GET_BITS(reg, pos, len) \
    (((reg) >> (pos)) & BIT_MASK(len))

// Example usage
void example_bit_operations(void) {
    volatile uint32_t register_value = 0;

    // Set bit 5
    SET_BIT(register_value, 5);

    // Clear bits 2-3
    CLEAR_BITS(register_value, 2, 2);

    // Set bits 8-11 to value 0xA
    SET_BITS(register_value, 8, 4, 0xA);

    // Get value of bits 8-11
    uint32_t extracted_value = GET_BITS(register_value, 8, 4);

    // Check if bit 5 is set
    if (CHECK_BIT(register_value, 5)) {
        // Bit 5 is set
    }
}
```

Bit Fields

Bit fields provide a convenient way to access specific bits within a structure:

```

#include <stdint.h>

// GPIO register structure using bit fields
typedef struct {
    volatile uint32_t MODER    : 32; // Port mode register
    volatile uint32_t OTYPER   : 16; // Port output type register
    volatile uint32_t OSPEEDR  : 32; // Port output speed register
    volatile uint32_t PUPDR    : 32; // Port pull-up/pull-down register
    volatile uint32_t IDR      : 16; // Port input data register
    volatile uint32_t ODR      : 16; // Port output data register
    volatile uint32_t BSRR     : 32; // Port bit set/reset register
    volatile uint32_t LCKR     : 32; // Port configuration lock register
    // ... other registers
} GPIO_TypeDef;

// Access GPIO registers using bit fields
#define GPIOA_BASE_ADDR 0x40020000
#define GPIOA            ((GPIO_TypeDef *)GPIOA_BASE_ADDR)

void gpio_bitfield_example(void) {
    // Configure pin 0 as output using bit fields
    GPIOA->MODER &= ~(3UL << 0); // Clear mode bits for pin 0
    GPIOA->MODER |= (1UL << 0);  // Set output mode for pin 0

    // Set pin 0 high
    GPIOA->BSRR = (1UL << 0);
}

```

Inline Assembly

Inline assembly allows embedding assembly code directly within C functions for maximum performance and hardware control.

Basic Inline Assembly

```

#include <stdint.h>

// Simple inline assembly example
void nop_instruction(void) {
    __asm__ volatile ("nop");
}

// Inline assembly with input/output operands
uint32_t rotate_left(uint32_t value, uint32_t shift) {

```

```

uint32_t result;

__asm__ volatile (
    "ror %0, %1, %2"
    : "=r" (result)          // Output operand
    : "r" (value), "r" (shift) // Input operands
);

return result;
}

// Memory barrier using inline assembly
void memory_barrier(void) {
    __asm__ volatile (" ::: \"memory\");
}

```

Advanced Inline Assembly

```

// Atomic increment using inline assembly
int32_t atomic_increment(volatile int32_t *value) {
    int32_t result;

    __asm__ volatile (
        "1: ldrex %0, [%1]    \n" // Load exclusive
        "   add    %0, %0, #1  \n" // Increment
        "   strex  r1, %0, [%1] \n" // Store exclusive
        "   cmp    r1, #0      \n" // Check if successful
        "   bne    1b          \n" // Retry if not successful
        : "=&r" (result)
        : "r" (value)
        : "r1", "cc", "memory"
    );

    return result;
}

// Disable interrupts
void disable_interrupts(void) {
    __asm__ volatile (
        "cpsid i"
        ::: "memory"
    );
}

// Enable interrupts

```

```

void enable_interrupts(void) {
    __asm__ volatile (
        "cpsie i"
        ::: "memory"
    );
}

```

Memory Management

Efficient memory management is critical in resource-constrained embedded systems.

Custom Memory Allocator

```

#include <stdint.h>
#include <stddef.h>

#define HEAP_SIZE 4096
static uint8_t heap[HEAP_SIZE];
static size_t heap_top = 0;

// Simple memory allocator
void* simple_malloc(size_t size) {
    // Align to 4-byte boundary
    size = (size + 3) & ~3;

    if (heap_top + size > HEAP_SIZE) {
        return NULL; // Out of memory
    }

    void *ptr = &heap[heap_top];
    heap_top += size;

    return ptr;
}

// Simple memory deallocator (does nothing in this simple implementation)
void simple_free(void *ptr) {
    // In a real implementation, you would track allocated blocks
    // This simple version just ignores deallocation
    (void)ptr;
}

// Reset heap (for demonstration)
void heap_reset(void) {

```

```

    heap_top = 0;
}

```

Memory Sections and Linker Scripts

Understanding memory sections and linker scripts is crucial for embedded development:

```

// Example of placing variables in specific memory sections
__attribute__((section(".ramfunc")))
void fast_function(void) {
    // Function placed in RAM for faster execution
}

__attribute__((aligned(16)))
uint32_t aligned_buffer[32]; // 16-byte aligned buffer

// Place variable in specific memory region
__attribute__((section(".ccmram")))
uint32_t ccm_data[100]; // Placed in Core Coupled Memory

// Zero-initialize variable in .bss section
static uint32_t bss_variable;

// Initialize variable in .data section
static uint32_t data_variable = 0x12345678;

// Read-only variable in .rodata section
static const uint32_t rodata_variable = 0xABCDEF00;

```

Hardware Abstraction Layer (HAL)

Creating a hardware abstraction layer improves code portability and maintainability.

GPIO HAL Implementation

```

#include <stdint.h>

// GPIO port enumeration
typedef enum {
    GPIO_PORT_A,
    GPIO_PORT_B,
    GPIO_PORT_C,
    // ... other ports
}

```

```

} gpio_port_t;

// GPIO pin enumeration
typedef enum {
    GPIO_PIN_0 = 0,
    GPIO_PIN_1,
    GPIO_PIN_2,
    GPIO_PIN_3,
    GPIO_PIN_4,
    GPIO_PIN_5,
    GPIO_PIN_6,
    GPIO_PIN_7,
    GPIO_PIN_8,
    GPIO_PIN_9,
    GPIO_PIN_10,
    GPIO_PIN_11,
    GPIO_PIN_12,
    GPIO_PIN_13,
    GPIO_PIN_14,
    GPIO_PIN_15
} gpio_pin_t;

// GPIO mode enumeration
typedef enum {
    GPIO_MODE_INPUT = 0,
    GPIO_MODE_OUTPUT,
    GPIO_MODE_AF,
    GPIO_MODE_ANALOG
} gpio_mode_t;

// GPIO output type enumeration
typedef enum {
    GPIO_OTYPE_PP = 0, // Push-pull
    GPIO_OTYPE_OD      // Open-drain
} gpio_otype_t;

// GPIO speed enumeration
typedef enum {
    GPIO_SPEED_LOW = 0,
    GPIO_SPEED_MEDIUM,
    GPIO_SPEED_HIGH,
    GPIO_SPEED_VERY_HIGH
} gpio_speed_t;

// GPIO pull-up/pull-down enumeration
typedef enum {
    GPIO_PUPD_NONE = 0,

```



```

        GPIO_PUPD_UP,
        GPIO_PUPD_DOWN
    } gpio_pupd_t;

// GPIO configuration structure
typedef struct {
    gpio_mode_t mode;
    gpio_otype_t otype;
    gpio_speed_t speed;
    gpio_pupd_t pupd;
} gpio_config_t;

// Base addresses for GPIO ports
#define GPIOA_BASE 0x40020000
#define GPIOB_BASE 0x40020400
#define GPIOC_BASE 0x40020800

// GPIO register structure
typedef struct {
    volatile uint32_t MODER;    // GPIO port mode register
    volatile uint32_t OTYPER;   // GPIO port output type register
    volatile uint32_t OSPEEDR;  // GPIO port output speed register
    volatile uint32_t PUPDR;    // GPIO port pull-up/pull-down register
    volatile uint32_t IDR;      // GPIO port input data register
    volatile uint32_t ODR;      // GPIO port output data register
    volatile uint32_t BSRR;     // GPIO port bit set/reset register
    volatile uint32_t LCKR;     // GPIO port configuration lock register
    volatile uint32_t AFR[2];   // GPIO alternate function registers
} gpio_reg_t;

// Get GPIO register pointer
static gpio_reg_t* get_gpio_reg(gpio_port_t port) {
    switch (port) {
        case GPIO_PORT_A: return (gpio_reg_t*)GPIOA_BASE;
        case GPIO_PORT_B: return (gpio_reg_t*)GPIOB_BASE;
        case GPIO_PORT_C: return (gpio_reg_t*)GPIOC_BASE;
        default: return NULL;
    }
}

// Configure GPIO pin
int gpio_init(gpio_port_t port, gpio_pin_t pin, const gpio_config_t *config) {
    gpio_reg_t *gpio = get_gpio_reg(port);
    if (gpio == NULL || config == NULL) {
        return -1;
    }
}

```

```

uint32_t pin_pos = (uint32_t)pin * 2;

// Configure mode
gpio->MODER &= ~(3UL << pin_pos);
gpio->MODER |= ((uint32_t)config->mode << pin_pos);

// Configure output type (only for output modes)
if (config->mode == GPIO_MODE_OUTPUT || config->mode == GPIO_MODE_AF) {
    gpio->OTYPER &= ~(1UL << pin);
    gpio->OTYPER |= ((uint32_t)config->otype << pin);
}

// Configure speed
gpio->OSPEEDR &= ~(3UL << pin_pos);
gpio->OSPEEDR |= ((uint32_t)config->speed << pin_pos);

// Configure pull-up/pull-down
gpio->PUPDR &= ~(3UL << pin_pos);
gpio->PUPDR |= ((uint32_t)config->pupd << pin_pos);

return 0;
}

// Set GPIO pin high
void gpio_set(gpio_port_t port, gpio_pin_t pin) {
    gpio_reg_t *gpio = get_gpio_reg(port);
    if (gpio != NULL) {
        gpio->BSRR = (1UL << pin);
    }
}

// Clear GPIO pin (set low)
void gpio_clear(gpio_port_t port, gpio_pin_t pin) {
    gpio_reg_t *gpio = get_gpio_reg(port);
    if (gpio != NULL) {
        gpio->BSRR = (1UL << (pin + 16));
    }
}

// Toggle GPIO pin
void gpio_toggle(gpio_port_t port, gpio_pin_t pin) {
    gpio_reg_t *gpio = get_gpio_reg(port);
    if (gpio != NULL) {
        gpio->ODR ^= (1UL << pin);
    }
}

```

```

// Read GPIO pin state
int gpio_read(gpio_port_t port, gpio_pin_t pin) {
    gpio_reg_t *gpio = get_gpio_reg(port);
    if (gpio != NULL) {
        return (gpio->IDR & (1UL << pin)) ? 1 : 0;
    }
    return -1;
}

```

Interrupt Management

Proper interrupt management is essential for responsive embedded systems.

Nested Vectored Interrupt Controller (NVIC)

```

// NVIC registers
#define NVIC_BASE 0xE000E100
#define NVIC_ISER ((volatile uint32_t*)(NVIC_BASE + 0x000)) // Interrupt set-enable
#define NVIC_ICER ((volatile uint32_t*)(NVIC_BASE + 0x080)) // Interrupt clear-enable
#define NVIC_ISPR ((volatile uint32_t*)(NVIC_BASE + 0x100)) // Interrupt set-pending
#define NVIC_ICPR ((volatile uint32_t*)(NVIC_BASE + 0x180)) // Interrupt clear-pending
#define NVIC_IPR ((volatile uint8_t*)(NVIC_BASE + 0x300)) // Interrupt priority

// Enable interrupt
void nvic_enable_irq(uint8_t irq_number) {
    NVIC_ISER[irq_number >> 5] = (1UL << (irq_number & 0x1F));
}

// Disable interrupt
void nvic_disable_irq(uint8_t irq_number) {
    NVIC_ICER[irq_number >> 5] = (1UL << (irq_number & 0x1F));
}

// Set interrupt priority
void nvic_set_priority(uint8_t irq_number, uint8_t priority) {
    NVIC_IPR[irq_number] = (priority << 4) & 0xFF;
}

// Global interrupt enable/disable
void enable_irq(void) {
    __asm__ volatile ("cpsie i" ::: "memory");
}

void disable_irq(void) {

```

```

    __asm__ volatile ("cpsid i" ::: "memory");
}

```

Practical Examples

UART Driver Implementation

```

#include <stdint.h>

// UART registers
typedef struct {
    volatile uint32_t SR;    // Status register
    volatile uint32_t DR;    // Data register
    volatile uint32_t BRR;   // Baud rate register
    volatile uint32_t CR1;   // Control register 1
    volatile uint32_t CR2;   // Control register 2
    volatile uint32_t CR3;   // Control register 3
    volatile uint32_t GTPR;  // Guard time and prescaler register
} uart_reg_t;

#define USART1_BASE 0x40011000
#define USART1      ((uart_reg_t*)USART1_BASE)

// UART status flags
#define UART_FLAG_TXE (1UL << 7) // Transmit data register empty
#define UART_FLAG_RXNE (1UL << 5) // Read data register not empty
#define UART_FLAG_TC (1UL << 6) // Transmission complete

// UART control bits
#define UART_CR1_TE (1UL << 3) // Transmitter enable
#define UART_CR1_RE (1UL << 2) // Receiver enable
#define UART_CR1_UE (1UL << 0) // UART enable

// Initialize UART
int uart_init(uint32_t baudrate, uint32_t peripheral_clock) {
    // Enable UART clock (implementation-specific)
    // RCC->APB2ENR |= RCC_APB2ENR_USART1EN;

    // Configure baud rate
    uint32_t baud_div = (peripheral_clock + (baudrate / 2)) / baudrate;
    USART1->BRR = baud_div;

    // Configure UART (8N1, no flow control)
    USART1->CR1 = UART_CR1_TE | UART_CR1_RE | UART_CR1_UE;
}

```

```

    return 0;
}

// Send character
void uart_putc(char c) {
    // Wait for transmit data register to be empty
    while (!(USART1->SR & UART_FLAG_TXE));

    // Send character
    USART1->DR = c;
}

// Send string
void uart_puts(const char *str) {
    while (*str) {
        uart_putc(*str++);
    }
}

// Receive character
char uart_getc(void) {
    // Wait for receive data register to be not empty
    while (!(USART1->SR & UART_FLAG_RXNE));

    // Read character
    return (char)(USART1->DR & 0xFF);
}

// Check if data is available
int uart_available(void) {
    return (USART1->SR & UART_FLAG_RXNE) ? 1 : 0;
}

// Example usage
int main(void) {
    // Initialize UART for 115200 baud
    uart_init(115200, 16000000); // 16MHz peripheral clock

    uart_puts("Hello, Embedded World!\r\n");

    while (1) {
        if (uart_available()) {
            char c = uart_getc();
            uart_putc(c); // Echo character
        }
    }
}

```

```
    return 0;  
}
```

Summary

Low-level programming in embedded systems involves:

1. **Memory-Mapped I/O** - Direct hardware register access
2. **Bit Manipulation** - Efficient control of individual bits and bit fields
3. **Inline Assembly** - Maximum performance and hardware control
4. **Memory Management** - Custom allocators and memory section control
5. **Hardware Abstraction** - Portable and maintainable code
6. **Interrupt Management** - Responsive system behavior
7. **Peripheral Drivers** - Interface with hardware components

These techniques enable developers to create efficient, reliable embedded applications that make optimal use of limited system resources while maintaining precise control over hardware behavior.

Embedded C Techniques

Introduction

Embedded C programming requires specialized techniques and language extensions to efficiently utilize limited resources while maintaining code reliability and performance. This chapter explores advanced C programming techniques specifically tailored for embedded systems, including compiler extensions, optimization strategies, and specialized language features.

Compiler Extensions and Attributes

Modern embedded compilers provide extensions that enhance C's capabilities for embedded development.

GCC Attributes

GCC provides numerous attributes to control code generation and optimization:

```
#include <stdint.h>

// Align variable to specific boundary
__attribute__((aligned(16)))
uint32_t aligned_array[32];

// Place function in specific section
__attribute__((section(".ramfunc")))
void fast_function(void) {
    // Code placed in RAM for faster execution
}

// Never return function (noreturn)
__attribute__((noreturn))
void system_reset(void) {
    // Reset system
    // This function never returns
    while (1);
}

// Weak symbol (can be overridden)
```

```

__attribute__((weak))
void default_handler(void) {
    // Default interrupt handler
    while (1);
}

// Pack structure to minimize padding
__attribute__((packed))
struct packed_data {
    uint8_t byte1;
    uint16_t word1;
    uint8_t byte2;
};

// Always inline function
__attribute__((always_inline))
static inline uint32_t fast_multiply(uint32_t a, uint32_t b) {
    return a * b;
}

// Never inline function
__attribute__((noinline))
void critical_function(void) {
    // Function that should never be inlined
}

// Specify function call convention
__attribute__((long_call))
void far_function(void) {
    // Function that may be far away in memory
}

```

IAR and Keil Extensions

Different compilers provide their own extensions:

```

// IAR specific pragmas
#pragma location = "MY_SECTION"
__no_init uint32_t persistent_data; // No initialization

// Keil specific attributes
__align(16)
uint32_t keil_aligned_array[32];

// ARM-specific intrinsics
#include <arm_acle.h>

```



```

void example_intrinsics(void) {
    // Count leading zeros
    uint32_t value = 0x00100000;
    uint32_t clz = __clz(value); // Count leading zeros

    // Reverse bits
    uint32_t reversed = __rbit(value);

    // Byte-swap
    uint32_t swapped = __rev(value);
}

```

Optimization Techniques

Embedded systems require careful optimization to balance performance, size, and power consumption.

Compiler Optimization Levels

```

// Optimization levels (GCC examples)
// -O0: No optimization (default for debugging)
// -O1: Basic optimization
// -O2: Full optimization
// -O3: Aggressive optimization
// -Os: Optimize for size
// -Ofast: Fastest possible (may violate standards)

// Function-specific optimization
__attribute__((optimize("O3")))
void optimized_function(void) {
    // This function will be optimized with -O3 regardless of global settings
}

__attribute__((optimize("Os")))
void size_optimized_function(void) {
    // This function will be optimized for size
}

```

Manual Optimization Techniques

```
#include <stdint.h>

// Loop unrolling
void manual_unroll_example(uint32_t *array, uint32_t size) {
    uint32_t i;

    // Unrolled loop (process 4 elements per iteration)
    for (i = 0; i < size - 3; i += 4) {
        array[i] += 1;
        array[i+1] += 1;
        array[i+2] += 1;
        array[i+3] += 1;
    }

    // Handle remaining elements
    for (; i < size; i++) {
        array[i] += 1;
    }
}

// Strength reduction
uint32_t strength_reduction_example(uint32_t x, uint32_t iterations) {
    // Instead of multiplication in loop
    uint32_t result = 0;
    uint32_t multiplier = x;

    for (uint32_t i = 0; i < iterations; i++) {
        result += multiplier; // Addition instead of multiplication
        multiplier += x;      // Increment instead of multiplication
    }

    return result;
}

// Lookup tables for expensive calculations
static const uint8_t sine_table[256] = {
    // Pre-calculated sine values
    128, 131, 134, 137, 140, 143, 146, 149, 152, 155, 158, 161, 164, 167, 170, 173,
    176, 179, 182, 185, 188, 191, 194, 196, 199, 202, 205, 207, 210, 213, 215, 218,
    // ... more values
};

uint8_t fast_sine(uint8_t angle) {
    return sine_table[angle];
}
```

Fixed-Point Arithmetic

Floating-point operations are expensive on many embedded processors. Fixed-point arithmetic provides an efficient alternative.

```
#include <stdint.h>

// Fixed-point types (Q-format)
typedef int32_t q16_16; // 16 integer bits, 16 fractional bits
typedef int16_t q8_8;   // 8 integer bits, 8 fractional bits

// Convert between integer and fixed-point
#define INT_TO_Q16_16(x) ((q16_16)((x) << 16))
#define Q16_16_TO_INT(x) ((int32_t)((x) >> 16))

// Fixed-point multiplication (Q16.16)
q16_16 q16_16_mul(q16_16 a, q16_16 b) {
    // Multiply and shift to maintain Q16.16 format
    return (q16_16)(((int64_t)a * b) >> 16);
}

// Fixed-point division (Q16.16)
q16_16 q16_16_div(q16_16 a, q16_16 b) {
    // Divide and shift to maintain Q16.16 format
    return (q16_16)(((int64_t)a << 16) / b);
}

// Fixed-point addition and subtraction (same format)
q16_16 q16_16_add(q16_16 a, q16_16 b) {
    return a + b;
}

q16_16 q16_16_sub(q16_16 a, q16_16 b) {
    return a - b;
}

// Example usage
void fixed_point_example(void) {
    q16_16 a = INT_TO_Q16_16(3); // 3.0 in Q16.16
    q16_16 b = INT_TO_Q16_16(2) + (1 << 15); // 2.5 in Q16.16

    q16_16 result = q16_16_mul(a, b); // 3.0 * 2.5 = 7.5
    int32_t integer_result = Q16_16_TO_INT(result); // Convert to integer
}
```

Volatile Keyword and Memory-Mapped I/O

The volatile keyword is crucial for embedded programming to prevent compiler optimizations that could break hardware interaction.

```
#include <stdint.h>

// Memory-mapped hardware registers
#define GPIOA_BASE 0x40020000
#define GPIOA_ODR (*(volatile uint32_t*)(GPIOA_BASE + 0x14))
#define GPIOA_IDR (*(volatile uint32_t*)(GPIOA_BASE + 0x10))

// Without volatile, compiler might optimize away repeated reads
void non_volatile_example(void) {
    uint32_t reg_value = GPIOA_IDR; // Compiler might cache this value
    // ... some code ...
    reg_value = GPIOA_IDR; // Compiler might use cached value instead of re-reading
}

// With volatile, compiler will always read from hardware
void volatile_example(void) {
    volatile uint32_t reg_value = GPIOA_IDR; // Always read from hardware
    // ... some code ...
    reg_value = GPIOA_IDR; // Always read from hardware again
}

// Volatile pointers
volatile uint32_t * const gpio_odr = (volatile uint32_t*)(GPIOA_BASE + 0x14);

void pointer_example(void) {
    *gpio_odr = 0x00000001; // Set pin 0
    *gpio_odr = 0x00000002; // Set pin 1
}
```

Interrupt Service Routines

Properly written ISRs are critical for real-time embedded systems.

```
#include <stdint.h>

// Global variables used in ISR should be volatile
volatile uint32_t isr_counter = 0;
volatile uint8_t data_ready = 0;
volatile uint8_t received_data = 0;
```

```

// Interrupt service routine
void USART1_IRQHandler(void) __attribute__((interrupt));
void USART1_IRQHandler(void) {
    // Read status register
    uint32_t status = USART1->SR;

    // Check for receive interrupt
    if (status & USART_SR_RXNE) {
        // Read received data
        received_data = USART1->DR;
        data_ready = 1;
        isr_counter++;
    }

    // Check for transmit interrupt
    if (status & USART_SR_TXE) {
        // Disable transmit interrupt
        USART1->CR1 &= ~USART_CR1_TXEIE;
    }
}

// Critical section management
uint32_t enter_critical_section(void) {
    uint32_t primask;
    __asm__ volatile ("mrs %0, primask" : "=r" (primask));
    __asm__ volatile ("cpsid i");
    return primask;
}

void exit_critical_section(uint32_t primask) {
    __asm__ volatile ("msr primask, %0" : : "r" (primask));
}

// Safe access to shared variables
uint32_t get_isr_counter(void) {
    uint32_t counter;
    uint32_t primask = enter_critical_section();
    counter = isr_counter;
    exit_critical_section(primask);
    return counter;
}

```

Memory-Constrained Programming

Techniques for working with limited memory resources.

Stack Management

```
#include <stdint.h>

// Monitor stack usage
extern uint32_t _estack;      // End of stack (top)
extern uint32_t _stack_start; // Start of stack

// Stack pattern for usage monitoring
#define STACK_PATTERN 0xDEADBEEF

void init_stack_pattern(void) {
    uint32_t *stack_ptr = &_stack_start;
    while (stack_ptr < &_estack) {
        *stack_ptr++ = STACK_PATTERN;
    }
}

uint32_t get_stack_usage(void) {
    uint32_t *stack_ptr = &_stack_start;
    while (stack_ptr < &_estack && *stack_ptr == STACK_PATTERN) {
        stack_ptr++;
    }
    return (uint32_t)&_estack - (uint32_t)stack_ptr;
}

// Avoid large local variables
void bad_function(void) {
    uint8_t large_buffer[2048]; // Large stack allocation
    // ... use buffer ...
}

void good_function(void) {
    static uint8_t large_buffer[2048]; // Static allocation
    // ... use buffer ...
}

// Or use dynamic allocation with careful management
uint8_t *get_buffer(void) {
    static uint8_t buffer[2048];
    static uint8_t in_use = 0;

    if (!in_use) {
        in_use = 1;
        return buffer;
    }
    return NULL; // Buffer in use
}
```

```

}

void release_buffer(void) {
    // In a real implementation, you'd track which buffer is released
}

```

Code Size Optimization

```

#include <stdint.h>

// Use smaller data types when possible
void size_optimized_loop(uint8_t count) {
    for (uint8_t i = 0; i < count; i++) {
        // Process data
    }
}

// Instead of:
void size_unoptimized_loop(uint32_t count) {
    for (uint32_t i = 0; i < count; i++) {
        // Process data
    }
}

// Use function pointers for common operations
typedef void (*handler_func_t)(void);

void handler_a(void) { /* ... */ }
void handler_b(void) { /* ... */ }
void handler_c(void) { /* ... */ }

static const handler_func_t handlers[] = {
    handler_a,
    handler_b,
    handler_c
};

void call_handler(uint8_t index) {
    if (index < sizeof(handlers) / sizeof(handlers[0])) {
        handlers[index]();
    }
}

// Use lookup tables instead of complex calculations
static const uint8_t crc8_table[256] = {

```

```

    0x00, 0x07, 0x0E, 0x09, 0x1C, 0x1B, 0x12, 0x15,
    0x38, 0x3F, 0x36, 0x31, 0x24, 0x23, 0x2A, 0x2D,
    // ... table values
};

uint8_t crc8(uint8_t *data, uint32_t length) {
    uint8_t crc = 0;
    for (uint32_t i = 0; i < length; i++) {
        crc = crc8_table[crc ^ data[i]];
    }
    return crc;
}

```

Real-Time Considerations

Techniques for ensuring deterministic behavior in real-time systems.

Timing Analysis

```

#include <stdint.h>

// Measure execution time
uint32_t get_cycle_count(void) {
    uint32_t cycle_count;
    __asm__ volatile ("rdcycle %0" : "=r" (cycle_count));
    return cycle_count;
}

void timing_analysis_example(void) {
    uint32_t start_time = get_cycle_count();

    // Code to measure
    complex_calculation();

    uint32_t end_time = get_cycle_count();
    uint32_t execution_time = end_time - start_time;

    // Store or report execution time
}

// Avoid operations with variable execution time
void deterministic_delay(uint32_t cycles) {
    // Use timer instead of loop with unknown iterations
    uint32_t start = get_cycle_count();

```



```

    while ((get_cycle_count() - start) < cycles) {
        // Wait
    }
}

// Use constant-time operations when possible
uint32_t constant_time_compare(const uint8_t *a, const uint8_t *b, uint32_t len) {
    uint32_t result = 0;
    for (uint32_t i = 0; i < len; i++) {
        result |= a[i] ^ b[i]; // Always compare all bytes
    }
    return result; // 0 if equal, non-zero if different
}

```

Hardware-Specific Features

Leveraging processor-specific features for maximum efficiency.

ARM Cortex-M Specific Features

```

#include <stdint.h>

// Use Cortex-M specific instructions
void cortex_m_optimizations(void) {
    // Count leading zeros
    uint32_t value = 0x00100000;
    uint32_t clz;
    __asm__ volatile ("clz %0, %1" : "=r" (clz) : "r" (value));

    // Reverse bits
    uint32_t reversed;
    __asm__ volatile ("rbit %0, %1" : "=r" (reversed) : "r" (value));

    // Byte-reverse word
    uint32_t rev;
    __asm__ volatile ("rev %0, %1" : "=r" (rev) : "r" (value));
}

// Use bit-band region for atomic bit operations
#define BITBAND_SRAM_REF 0x20000000
#define BITBAND_SRAM_BASE 0x22000000
#define BITBAND_SRAM(a,b) ((BITBAND_SRAM_BASE + (a-BITBAND_SRAM_REF)*32 + (b*4)))

// Atomic bit manipulation using bit-band

```

```

void bitband_example(volatile uint32_t *reg, uint8_t bit) {
    // Create bit-band alias
    volatile uint32_t *bit_alias = (volatile uint32_t*)BITBAND_SRAM((uint32_t)reg, bit);

    *bit_alias = 1; // Atomically set bit
    *bit_alias = 0; // Atomically clear bit
}

```

Practical Examples

Efficient Sensor Data Processing

```

#include <stdint.h>

// Circular buffer for sensor data
#define BUFFER_SIZE 64

typedef struct {
    uint16_t buffer[BUFFER_SIZE];
    uint8_t head;
    uint8_t tail;
    uint8_t count;
} circular_buffer_t;

static circular_buffer_t sensor_buffer;

// Initialize circular buffer
void buffer_init(circular_buffer_t *buf) {
    buf->head = 0;
    buf->tail = 0;
    buf->count = 0;
}

// Add data to buffer
int buffer_put(circular_buffer_t *buf, uint16_t data) {
    if (buf->count >= BUFFER_SIZE) {
        return -1; // Buffer full
    }

    buf->buffer[buf->head] = data;
    buf->head = (buf->head + 1) % BUFFER_SIZE;
    buf->count++;

    return 0;
}

```

```

// Get data from buffer
int buffer_get(circular_buffer_t *buf, uint16_t *data) {
    if (buf->count == 0) {
        return -1; // Buffer empty
    }

    *data = buf->buffer[buf->tail];
    buf->tail = (buf->tail + 1) % BUFFER_SIZE;
    buf->count--;

    return 0;
}

// Moving average filter
#define AVERAGE_WINDOW 8

uint16_t moving_average(uint16_t new_value) {
    static uint16_t window[AVERAGE_WINDOW];
    static uint8_t index = 0;
    static uint32_t sum = 0;
    static uint8_t count = 0;

    // Remove oldest value from sum
    sum -= window[index];

    // Add new value to sum and window
    window[index] = new_value;
    sum += new_value;

    // Update index
    index = (index + 1) % AVERAGE_WINDOW;

    // Update count
    if (count < AVERAGE_WINDOW) {
        count++;
    }

    // Return average
    return (uint16_t)(sum / count);
}

```

Power Management

```
#include <stdint.h>

// Power management states
typedef enum {
    POWER_STATE_ACTIVE,
    POWER_STATE_SLEEP,
    POWER_STATE_DEEP_SLEEP
} power_state_t;

// Enter sleep mode
void enter_sleep_mode(void) {
    // Disable unnecessary peripherals
    // RCC->APB1ENR &= ~RCC_APB1ENR_TIM2EN; // Disable timer

    // Configure wake-up sources
    // EXTI->IMR |= EXTI_IMR_MR0; // Enable EXTI line 0

    // Enter sleep mode
    __asm__ volatile ("wfi"); // Wait for interrupt
}

// Enter deep sleep mode
void enter_deep_sleep_mode(void) {
    // Save critical data to retention memory

    // Configure deep sleep
    // SCB->SCR |= SCB_SCR_SLEEPDEEP_Msk;

    // Enter deep sleep
    __asm__ volatile ("wfi");
}

// Dynamic frequency scaling
void set_performance_level(uint8_t level) {
    switch (level) {
        case 0: // Low power
            // SystemClock_Config_LowPower();
            break;
        case 1: // Medium power
            // SystemClock_Config_Medium();
            break;
        case 2: // High performance
            // SystemClock_Config_High();
            break;
    }
}
```

```
}
```

Summary

Embedded C techniques involve specialized programming approaches to maximize efficiency in resource-constrained environments:

1. **Compiler Extensions** - Using attributes and pragmas for optimal code generation
2. **Optimization Techniques** - Balancing performance, size, and power consumption
3. **Fixed-Point Arithmetic** - Efficient alternative to floating-point operations
4. **Volatile Keyword** - Proper handling of hardware registers and shared variables
5. **Interrupt Management** - Writing efficient and safe ISRs
6. **Memory-Constrained Programming** - Techniques for limited RAM and Flash
7. **Real-Time Considerations** - Ensuring deterministic behavior
8. **Hardware-Specific Features** - Leveraging processor capabilities

These techniques enable developers to create efficient, reliable embedded applications that make optimal use of limited system resources while maintaining real-time performance requirements.

Real-time Programming

Introduction

Real-time programming in embedded systems requires deterministic behavior, precise timing control, and efficient resource management. This chapter explores the principles and techniques for developing real-time embedded applications, including scheduling strategies, timing analysis, and meeting strict temporal constraints.

Real-time System Characteristics

Real-time systems must respond to events within specified time constraints. Understanding these constraints is fundamental to successful real-time programming.

Types of Real-time Systems

1. **Hard Real-time Systems** - Missed deadlines can cause catastrophic failures
 - Examples: Aircraft control systems, medical devices, automotive safety systems
2. **Soft Real-time Systems** - Missed deadlines degrade performance but don't cause failures
 - Examples: Multimedia streaming, user interfaces, gaming
3. **Firm Real-time Systems** - Some deadlines can be missed with limited consequences
 - Examples: Industrial process control, data logging

Real-time Requirements

```
#include <stdint.h>
#include <stdbool.h>

// Real-time task structure
typedef struct {
    void (*function)(void);    // Task function
    uint32_t period;           // Task period in milliseconds
    uint32_t deadline;        // Task deadline in milliseconds
    uint32_t execution_time;   // Worst-case execution time
    uint32_t next_release;     // Next release time
    uint8_t priority;          // Task priority
}
```

```

    bool enabled;                // Task enabled flag
} rt_task_t;

// Real-time system parameters
typedef struct {
    uint32_t system_time;        // Current system time
    uint32_t max_response_time;  // Maximum observed response time
    uint32_t missed_deadlines;   // Count of missed deadlines
} rt_system_t;

static rt_system_t rt_system;

```

Task Scheduling

Effective task scheduling is crucial for meeting real-time constraints while maximizing system utilization.

Rate Monotonic Scheduling (RMS)

Rate Monotonic Scheduling assigns priorities based on task periods - shorter periods get higher priorities.

```

#include <stdint.h>
#include <stdbool.h>

#define MAX_TASKS 16

static rt_task_t task_list[MAX_TASKS];
static uint8_t task_count = 0;

// Add task to system
int rt_add_task(void (*function)(void), uint32_t period, uint32_t deadline,
                uint32_t execution_time, uint8_t priority) {
    if (task_count >= MAX_TASKS) {
        return -1; // Task list full
    }

    task_list[task_count].function = function;
    task_list[task_count].period = period;
    task_list[task_count].deadline = deadline;
    task_list[task_count].execution_time = execution_time;
    task_list[task_count].next_release = period;
    task_list[task_count].priority = priority;
    task_list[task_count].enabled = true;
}

```

```

    task_count++;
    return 0;
}

// Rate Monotonic priority assignment
void rt_assign_rm_priorities(void) {
    // Sort tasks by period (shortest first)
    for (int i = 0; i < task_count - 1; i++) {
        for (int j = 0; j < task_count - i - 1; j++) {
            if (task_list[j].period > task_list[j + 1].period) {
                // Swap tasks
                rt_task_t temp = task_list[j];
                task_list[j] = task_list[j + 1];
                task_list[j + 1] = temp;
            }
        }
    }

    // Assign priorities (higher number = higher priority)
    for (int i = 0; i < task_count; i++) {
        task_list[i].priority = task_count - i;
    }
}

// Check Rate Monotonic schedulability
bool rt_check_rm_schedulability(void) {
    // Calculate total utilization
    double total_utilization = 0.0;

    for (int i = 0; i < task_count; i++) {
        double utilization = (double)task_list[i].execution_time / task_list[i].period;
        total_utilization += utilization;
    }

    // Rate Monotonic bound:  $n * (2^{(1/n)} - 1)$ 
    double rm_bound = task_count * (pow(2.0, 1.0/task_count) - 1.0);

    return (total_utilization <= rm_bound);
}

```

Earliest Deadline First (EDF) Scheduling

EDF scheduling dynamically assigns priorities based on task deadlines.


```

// Find highest priority task (earliest deadline)
int rt_find_highest_priority_task(void) {
    int highest_priority_task = -1;
    uint32_t earliest_deadline = UINT32_MAX;

    for (int i = 0; i < task_count; i++) {
        if (task_list[i].enabled &&
            task_list[i].next_release <= rt_system.system_time &&
            task_list[i].deadline < earliest_deadline) {
            earliest_deadline = task_list[i].deadline;
            highest_priority_task = i;
        }
    }

    return highest_priority_task;
}

// Update task deadlines
void rt_update_deadlines(void) {
    for (int i = 0; i < task_count; i++) {
        if (task_list[i].enabled) {
            // If task is ready to run
            if (rt_system.system_time >= task_list[i].next_release) {
                // Set absolute deadline
                task_list[i].deadline = task_list[i].next_release +
                                         task_list[i].period;
            }
        }
    }
}

```

Timing and Synchronization

Precise timing control is essential for real-time systems.

Timer Management

```

#include <stdint.h>

// Timer structure
typedef struct {
    uint32_t reload_value;    // Timer reload value
    uint32_t current_value;   // Current timer value
    bool running;             // Timer running flag
}

```

```

    bool expired;                // Timer expired flag
    void (*callback)(void);      // Expiration callback
} timer_t;

#define MAX_TIMERS 8
static timer_t timers[MAX_TIMERS];

// Initialize timer system
void timer_init(void) {
    for (int i = 0; i < MAX_TIMERS; i++) {
        timers[i].running = false;
        timers[i].expired = false;
        timers[i].callback = NULL;
    }

    // Configure hardware timer
    // TIM2->PSC = SystemCoreClock / 1000000 - 1; // 1MHz timer
    // TIM2->ARR = 999; // 1ms interrupt
    // TIM2->DIER |= TIM_DIER_UIE; // Enable update interrupt
    // TIM2->CR1 |= TIM_CR1_CEN; // Start timer
}

// Create timer
int timer_create(uint32_t milliseconds, void (*callback)(void)) {
    for (int i = 0; i < MAX_TIMERS; i++) {
        if (!timers[i].running) {
            timers[i].reload_value = milliseconds;
            timers[i].current_value = milliseconds;
            timers[i].callback = callback;
            timers[i].running = true;
            timers[i].expired = false;
            return i;
        }
    }
    return -1; // No free timers
}

// Timer tick handler (called by timer interrupt)
void timer_tick_handler(void) {
    rt_system.system_time++;

    // Update all running timers
    for (int i = 0; i < MAX_TIMERS; i++) {
        if (timers[i].running) {
            if (timers[i].current_value > 0) {
                timers[i].current_value--;
            }
        }
    }
}

```

```

        if (timers[i].current_value == 0) {
            timers[i].expired = true;
            timers[i].running = false;

            if (timers[i].callback) {
                timers[i].callback();
            }
        }
    }

    // Update task deadlines
    rt_update_deadlines();
}

// Get system time
uint32_t rt_get_system_time(void) {
    return rt_system.system_time;
}

// Delay function (blocking)
void rt_delay_ms(uint32_t milliseconds) {
    uint32_t start_time = rt_system.system_time;
    while ((rt_system.system_time - start_time) < milliseconds) {
        // Wait
    }
}

```

Semaphore Implementation

Semaphores provide synchronization between tasks and ISRs.

```

#include <stdint.h>
#include <stdbool.h>

#define MAX_SEMAPHORES 16

typedef struct {
    volatile int32_t count;    // Semaphore count
    uint8_t max_count;        // Maximum count
    bool initialized;         // Initialization flag
} semaphore_t;

static semaphore_t semaphores[MAX_SEMAPHORES];

```

```

// Initialize semaphore
int sem_init(uint8_t sem_id, uint8_t initial_count, uint8_t max_count) {
    if (sem_id >= MAX_SEMAPHORES) {
        return -1;
    }

    semaphores[sem_id].count = initial_count;
    semaphores[sem_id].max_count = max_count;
    semaphores[sem_id].initialized = true;

    return 0;
}

// Wait on semaphore (blocking)
int sem_wait(uint8_t sem_id) {
    if (sem_id >= MAX_SEMAPHORES || !semaphores[sem_id].initialized) {
        return -1;
    }

    // Disable interrupts to ensure atomic operation
    uint32_t primask = __get_PRIMASK();
    __disable_irq();

    while (semaphores[sem_id].count <= 0) {
        // Re-enable interrupts and wait
        __set_PRIMASK(primask);

        // In a real RTOS, this would yield to other tasks
        // For simple implementation, just wait
        __asm__ volatile ("nop");

        // Disable interrupts again
        __disable_irq();
    }

    semaphores[sem_id].count--;

    // Restore interrupt state
    __set_PRIMASK(primask);

    return 0;
}

// Signal semaphore (non-blocking)
int sem_signal(uint8_t sem_id) {
    if (sem_id >= MAX_SEMAPHORES || !semaphores[sem_id].initialized) {
        return -1;
    }

```

```

    }

    // Disable interrupts to ensure atomic operation
    uint32_t primask = __get_PRIMASK();
    __disable_irq();

    if (semaphores[sem_id].count < semaphores[sem_id].max_count) {
        semaphores[sem_id].count++;
    }

    // Restore interrupt state
    __set_PRIMASK(primask);

    return 0;
}

// Try wait on semaphore (non-blocking)
int sem_try_wait(uint8_t sem_id) {
    if (sem_id >= MAX_SEMAPHORES || !semaphores[sem_id].initialized) {
        return -1;
    }

    // Disable interrupts to ensure atomic operation
    uint32_t primask = __get_PRIMASK();
    __disable_irq();

    int result = -1;
    if (semaphores[sem_id].count > 0) {
        semaphores[sem_id].count--;
        result = 0;
    }

    // Restore interrupt state
    __set_PRIMASK(primask);

    return result;
}

```

Interrupt Handling in Real-time Systems

Proper interrupt management is crucial for maintaining real-time performance.

Interrupt Latency Management

```
#include <stdint.h>

// Measure interrupt latency
static volatile uint32_t interrupt_timestamp = 0;
static volatile uint32_t interrupt_latency = 0;

// Interrupt service routine with latency measurement
void EXTIO_IRQHandler(void) {
    // Record interrupt entry time
    interrupt_timestamp = rt_system.system_time;

    // Clear interrupt flag
    // EXTI->PR = EXTI_PR_PRO;

    // Process interrupt
    process_external_interrupt();

    // Calculate latency
    interrupt_latency = rt_system.system_time - interrupt_timestamp;
}

// Critical section management
typedef struct {
    uint32_t primask;
    bool interrupts_disabled;
} critical_section_t;

// Enter critical section
critical_section_t enter_critical_section(void) {
    critical_section_t cs;
    cs.primask = __get_PRIMASK();
    __disable_irq();
    cs.interrupts_disabled = true;
    return cs;
}

// Exit critical section
void exit_critical_section(critical_section_t cs) {
    if (cs.interrupts_disabled) {
        __set_PRIMASK(cs.primask);
    }
}

// Deferred interrupt processing
#define INTERRUPT_QUEUE_SIZE 16
```

```

typedef enum {
    INTERRUPT_TYPE_GPIO,
    INTERRUPT_TYPE_UART,
    INTERRUPT_TYPE_TIMER
} interrupt_type_t;

typedef struct {
    interrupt_type_t type;
    uint32_t data;
    uint32_t timestamp;
} interrupt_event_t;

static interrupt_event_t interrupt_queue[INTERRUPT_QUEUE_SIZE];
static volatile uint8_t queue_head = 0;
static volatile uint8_t queue_tail = 0;
static volatile uint8_t queue_count = 0;

// Add interrupt event to queue
void queue_interrupt_event(interrupt_type_t type, uint32_t data) {
    critical_section_t cs = enter_critical_section();

    if (queue_count < INTERRUPT_QUEUE_SIZE) {
        interrupt_queue[queue_head].type = type;
        interrupt_queue[queue_head].data = data;
        interrupt_queue[queue_head].timestamp = rt_system.system_time;

        queue_head = (queue_head + 1) % INTERRUPT_QUEUE_SIZE;
        queue_count++;
    }

    exit_critical_section(cs);
}

// Process queued interrupts
void process_interrupt_queue(void) {
    while (queue_count > 0) {
        critical_section_t cs = enter_critical_section();

        interrupt_event_t event = interrupt_queue[queue_tail];
        queue_tail = (queue_tail + 1) % INTERRUPT_QUEUE_SIZE;
        queue_count--;

        exit_critical_section(cs);

        // Process the event
        switch (event.type) {

```

```

        case INTERRUPT_TYPE_GPIO:
            process_gpio_interrupt(event.data);
            break;
        case INTERRUPT_TYPE_UART:
            process_uart_interrupt(event.data);
            break;
        case INTERRUPT_TYPE_TIMER:
            process_timer_interrupt(event.data);
            break;
    }
}
}

```

Resource Management

Efficient resource management is essential for predictable real-time behavior.

Memory Pool Management

```

#include <stdint.h>
#include <stdbool.h>
#include <string.h>

#define MEMORY_POOL_SIZE 1024
#define BLOCK_SIZE 32
#define NUM_BLOCKS (MEMORY_POOL_SIZE / BLOCK_SIZE)

typedef struct {
    uint8_t pool[MEMORY_POOL_SIZE];
    bool block_used[NUM_BLOCKS];
    uint8_t free_blocks;
} memory_pool_t;

static memory_pool_t memory_pool;

// Initialize memory pool
void mem_pool_init(void) {
    memset(memory_pool.block_used, 0, sizeof(memory_pool.block_used));
    memory_pool.free_blocks = NUM_BLOCKS;
}

// Allocate memory block
void* mem_pool_alloc(void) {
    critical_section_t cs = enter_critical_section();

```



```

    if (memory_pool.free_blocks == 0) {
        exit_critical_section(cs);
        return NULL; // No free blocks
    }

    // Find free block
    for (int i = 0; i < NUM_BLOCKS; i++) {
        if (!memory_pool.block_used[i]) {
            memory_pool.block_used[i] = true;
            memory_pool.free_blocks--;
            exit_critical_section(cs);
            return &memory_pool.pool[i * BLOCK_SIZE];
        }
    }

    exit_critical_section(cs);
    return NULL; // Should not reach here
}

// Free memory block
void mem_pool_free(void* ptr) {
    if (ptr == NULL) return;

    // Calculate block index
    uintptr_t offset = (uintptr_t)ptr - (uintptr_t)memory_pool.pool;
    uint32_t block_index = offset / BLOCK_SIZE;

    if (block_index < NUM_BLOCKS) {
        critical_section_t cs = enter_critical_section();
        if (memory_pool.block_used[block_index]) {
            memory_pool.block_used[block_index] = false;
            memory_pool.free_blocks++;
        }
        exit_critical_section(cs);
    }
}

```

Priority Inheritance

Priority inheritance prevents priority inversion in resource sharing.

```

#include <stdint.h>
#include <stdbool.h>

#define MAX_PRIORITY 255

```

```

typedef struct {
    uint8_t original_priority;
    uint8_t current_priority;
    bool owned;
    uint8_t owner_task_id;
} resource_t;

#define MAX_RESOURCES 16
static resource_t resources[MAX_RESOURCES];

// Initialize resource
void resource_init(uint8_t resource_id) {
    if (resource_id < MAX_RESOURCES) {
        resources[resource_id].original_priority = 0;
        resources[resource_id].current_priority = 0;
        resources[resource_id].owned = false;
        resources[resource_id].owner_task_id = 0;
    }
}

// Acquire resource with priority inheritance
int resource_acquire(uint8_t resource_id, uint8_t task_id, uint8_t task_priority) {
    if (resource_id >= MAX_RESOURCES) {
        return -1;
    }

    critical_section_t cs = enter_critical_section();

    if (!resources[resource_id].owned) {
        // Resource is free, acquire it
        resources[resource_id].owned = true;
        resources[resource_id].owner_task_id = task_id;
        resources[resource_id].original_priority = task_priority;
        resources[resource_id].current_priority = task_priority;
    } else {
        // Resource is owned by another task
        uint8_t owner_id = resources[resource_id].owner_task_id;

        // Check if priority inversion would occur
        if (task_priority > resources[resource_id].current_priority) {
            // Apply priority inheritance
            resources[resource_id].current_priority = task_priority;
            // In a real system, you would boost the owner task's priority
            task_list[owner_id].priority = task_priority;
        }
    }
}

```

```

        exit_critical_section(cs);
        return -2; // Resource busy
    }

    exit_critical_section(cs);
    return 0; // Success
}

// Release resource
int resource_release(uint8_t resource_id, uint8_t task_id) {
    if (resource_id >= MAX_RESOURCES) {
        return -1;
    }

    critical_section_t cs = enter_critical_section();

    if (resources[resource_id].owned &&
        resources[resource_id].owner_task_id == task_id) {
        // Release resource
        resources[resource_id].owned = false;
        resources[resource_id].owner_task_id = 0;
        // Restore original priority
        // In a real system, you would restore the owner task's original priority
    } else {
        exit_critical_section(cs);
        return -2; // Not owner or not owned
    }

    exit_critical_section(cs);
    return 0; // Success
}

```

Timing Analysis and Deadline Monitoring

Monitoring and analyzing timing behavior ensures real-time constraints are met.

Deadline Monitoring

```

#include <stdint.h>
#include <stdbool.h>

// Task execution tracking
typedef struct {
    uint32_t start_time;

```

```

    uint32_t end_time;
    uint32_t execution_time;
    uint32_t deadline;
    bool deadline_missed;
} task_execution_t;

#define EXECUTION_HISTORY_SIZE 32
static task_execution_t execution_history[MAX_TASKS][EXECUTION_HISTORY_SIZE];
static uint8_t history_index[MAX_TASKS] = {0};

// Start task execution tracking
void task_start_execution(uint8_t task_id) {
    if (task_id < MAX_TASKS) {
        execution_history[task_id][history_index[task_id]].start_time =
            rt_system.system_time;
        execution_history[task_id][history_index[task_id]].deadline =
            task_list[task_id].deadline;
    }
}

// End task execution tracking
void task_end_execution(uint8_t task_id) {
    if (task_id < MAX_TASKS) {
        uint8_t idx = history_index[task_id];
        execution_history[task_id][idx].end_time = rt_system.system_time;
        execution_history[task_id][idx].execution_time =
            execution_history[task_id][idx].end_time -
            execution_history[task_id][idx].start_time;

        // Check for deadline miss
        if (rt_system.system_time > execution_history[task_id][idx].deadline) {
            execution_history[task_id][idx].deadline_missed = true;
            rt_system.missed_deadlines++;
        } else {
            execution_history[task_id][idx].deadline_missed = false;
        }

        // Update maximum response time
        if (execution_history[task_id][idx].execution_time >
            rt_system.max_response_time) {
            rt_system.max_response_time =
                execution_history[task_id][idx].execution_time;
        }

        // Advance history index
        history_index[task_id] = (history_index[task_id] + 1) % EXECUTION_HISTORY_SIZE;
    }
}

```

```

}

// Get task statistics
typedef struct {
    uint32_t avg_execution_time;
    uint32_t max_execution_time;
    uint32_t min_execution_time;
    uint32_t deadline_misses;
    double deadline_miss_rate;
} task_stats_t;

task_stats_t get_task_stats(uint8_t task_id) {
    task_stats_t stats = {0};

    if (task_id >= MAX_TASKS) {
        return stats;
    }

    uint32_t total_execution_time = 0;
    uint32_t valid_samples = 0;
    stats.max_execution_time = 0;
    stats.min_execution_time = UINT32_MAX;

    for (int i = 0; i < EXECUTION_HISTORY_SIZE; i++) {
        if (execution_history[task_id][i].end_time > 0) {
            uint32_t exec_time = execution_history[task_id][i].execution_time;
            total_execution_time += exec_time;
            valid_samples++;

            if (exec_time > stats.max_execution_time) {
                stats.max_execution_time = exec_time;
            }
            if (exec_time < stats.min_execution_time) {
                stats.min_execution_time = exec_time;
            }

            if (execution_history[task_id][i].deadline_missed) {
                stats.deadline_misses++;
            }
        }
    }

    if (valid_samples > 0) {
        stats.avg_execution_time = total_execution_time / valid_samples;
        stats.deadline_miss_rate = (double)stats.deadline_misses / valid_samples;
    }
}

```

```

    return stats;
}

```

Practical Examples

Real-time Control System

```

#include <stdint.h>
#include <math.h>

// PID Controller
typedef struct {
    float kp;           // Proportional gain
    float ki;           // Integral gain
    float kd;           // Derivative gain
    float setpoint;     // Desired value
    float integral;     // Integral term
    float previous_error; // Previous error for derivative
    float output_min;   // Minimum output
    float output_max;   // Maximum output
} pid_controller_t;

// Initialize PID controller
void pid_init(pid_controller_t *pid, float kp, float ki, float kd,
              float output_min, float output_max) {
    pid->kp = kp;
    pid->ki = ki;
    pid->kd = kd;
    pid->setpoint = 0;
    pid->integral = 0;
    pid->previous_error = 0;
    pid->output_min = output_min;
    pid->output_max = output_max;
}

// Compute PID output
float pid_compute(pid_controller_t *pid, float feedback, float dt) {
    // Calculate error
    float error = pid->setpoint - feedback;

    // Proportional term
    float proportional = pid->kp * error;

    // Integral term
    pid->integral += error * dt;
}

```

```

float integral = pid->ki * pid->integral;

// Derivative term
float derivative = pid->kd * (error - pid->previous_error) / dt;

// Calculate output
float output = proportional + integral + derivative;

// Clamp output to limits
if (output > pid->output_max) {
    output = pid->output_max;
    // Prevent integral windup
    pid->integral -= error * dt;
} else if (output < pid->output_min) {
    output = pid->output_min;
    // Prevent integral windup
    pid->integral -= error * dt;
}

// Save error for next derivative calculation
pid->previous_error = error;

return output;
}

// Set PID setpoint
void pid_set_setpoint(pid_controller_t *pid, float setpoint) {
    pid->setpoint = setpoint;
}

// Real-time motor control task
static pid_controller_t motor_pid;
static float motor_position = 0.0;
static float motor_speed = 0.0;

void motor_control_task(void) {
    static uint32_t last_time = 0;
    uint32_t current_time = rt_system.system_time;
    float dt = (current_time - last_time) / 1000.0; // Convert to seconds
    last_time = current_time;

    // Read motor position (simulated)
    // In real implementation: motor_position = read_encoder();

    // Compute PID output
    float control_output = pid_compute(&motor_pid, motor_position, dt);

```

```

    // Apply control output to motor
    // In real implementation: set_motor_pwm(control_output);

    // Update motor model (simulated)
    motor_speed = control_output;
    motor_position += motor_speed * dt;
}

// Initialize motor control system
void motor_control_init(void) {
    // Initialize PID controller
    pid_init(&motor_pid, 2.0, 1.0, 0.1, -100.0, 100.0);
    pid_set_setpoint(&motor_pid, 100.0); // Set target position

    // Add motor control task (10ms period)
    rt_add_task(motor_control_task, 10, 10, 2, 0);
}

```

Real-time Data Acquisition System

```

#include <stdint.h>

// Data acquisition system
#define ADC_CHANNELS 8
#define SAMPLE_BUFFER_SIZE 256

typedef struct {
    uint16_t samples[ADC_CHANNELS][SAMPLE_BUFFER_SIZE];
    uint16_t write_index[ADC_CHANNELS];
    uint16_t read_index[ADC_CHANNELS];
    uint32_t sample_count[ADC_CHANNELS];
    uint32_t overrun_count[ADC_CHANNELS];
} data_acquisition_t;

static data_acquisition_t data_acq;

// Initialize data acquisition system
void data_acq_init(void) {
    for (int i = 0; i < ADC_CHANNELS; i++) {
        data_acq.write_index[i] = 0;
        data_acq.read_index[i] = 0;
        data_acq.sample_count[i] = 0;
        data_acq.overrun_count[i] = 0;
    }
}

```



```

// ADC interrupt handler
void ADC_IRQHandler(void) {
    static uint8_t current_channel = 0;

    // Read ADC value (simulated)
    // uint16_t adc_value = ADC1->DR;
    uint16_t adc_value = 0; // Simulated value

    // Store sample
    uint16_t write_idx = data_acq.write_index[current_channel];
    data_acq.samples[current_channel][write_idx] = adc_value;

    // Update indices
    data_acq.write_index[current_channel] = (write_idx + 1) % SAMPLE_BUFFER_SIZE;
    data_acq.sample_count[current_channel]++;

    // Check for buffer overrun
    if (data_acq.write_index[current_channel] == data_acq.read_index[current_channel]) {
        data_acq.overrun_count[current_channel]++;
        // Handle overrun (e.g., increment read index to maintain FIFO behavior)
        data_acq.read_index[current_channel] =
            (data_acq.read_index[current_channel] + 1) % SAMPLE_BUFFER_SIZE;
    }

    // Move to next channel
    current_channel = (current_channel + 1) % ADC_CHANNELS;

    // Configure ADC for next channel
    // ADC1->SQR3 = current_channel; // Set channel in sequence register
}

// Read samples from buffer
int data_acq_read_samples(uint8_t channel, uint16_t *buffer, uint16_t max_samples) {
    if (channel >= ADC_CHANNELS) {
        return -1;
    }

    critical_section_t cs = enter_critical_section();

    uint16_t available_samples = 0;
    uint16_t read_idx = data_acq.read_index[channel];
    uint16_t write_idx = data_acq.write_index[channel];

    // Calculate available samples
    if (write_idx >= read_idx) {
        available_samples = write_idx - read_idx;
    }
}

```

```

    } else {
        available_samples = SAMPLE_BUFFER_SIZE - read_idx + write_idx;
    }

    // Limit to requested samples
    if (available_samples > max_samples) {
        available_samples = max_samples;
    }

    // Copy samples
    for (uint16_t i = 0; i < available_samples; i++) {
        buffer[i] = data_acq.samples[channel][read_idx];
        read_idx = (read_idx + 1) % SAMPLE_BUFFER_SIZE;
    }

    // Update read index
    data_acq.read_index[channel] = read_idx;

    exit_critical_section(cs);

    return available_samples;
}

// Data processing task
void data_processing_task(void) {
    static uint16_t sample_buffer[64];

    // Process samples from channel 0
    int samples_read = data_acq_read_samples(0, sample_buffer, 64);

    if (samples_read > 0) {
        // Perform data processing (e.g., filtering, FFT, etc.)
        // process_samples(sample_buffer, samples_read);
    }
}

// Initialize data acquisition system
void data_acquisition_system_init(void) {
    data_acq_init();

    // Add data processing task (50ms period)
    rt_add_task(data_processing_task, 50, 50, 10, 0);

    // Enable ADC interrupt
    // NVIC_EnableIRQ(ADC_IRQn);
}

```

Summary

Real-time programming in embedded systems requires careful consideration of:

1. **Scheduling Algorithms** - Rate Monotonic and Earliest Deadline First scheduling
2. **Timing Control** - Precise timer management and delay functions
3. **Synchronization** - Semaphores, mutexes, and critical sections
4. **Interrupt Handling** - Low-latency ISRs and deferred processing
5. **Resource Management** - Memory pools and priority inheritance
6. **Timing Analysis** - Deadline monitoring and performance measurement
7. **Control Systems** - PID controllers and real-time feedback loops

Key principles for successful real-time programming: - Minimize interrupt latency and execution time - Use deterministic algorithms and data structures - Carefully manage shared resources to avoid priority inversion - Monitor and analyze timing behavior to ensure constraints are met - Implement proper error handling for missed deadlines - Balance performance requirements with system complexity

These techniques enable developers to create reliable, predictable embedded systems that meet strict real-time requirements while efficiently utilizing limited system resources.

Hardware Interfaces

Introduction

Embedded systems interact with the physical world through various hardware interfaces. Understanding how to properly interface with different types of hardware components is essential for developing functional embedded applications. This chapter covers the most common hardware interfaces used in embedded systems, including GPIO, SPI, I2C, UART, and other specialized interfaces.

GPIO (General Purpose Input/Output)

GPIO pins are the most fundamental interface in embedded systems, providing direct digital input and output capabilities.

GPIO Configuration and Control

```
#include <stdint.h>
#include <stdbool.h>

// GPIO pin modes
typedef enum {
    GPIO_MODE_INPUT = 0,
    GPIO_MODE_OUTPUT,
    GPIO_MODE_AF,      // Alternate function
    GPIO_MODE_ANALOG
} gpio_mode_t;

// GPIO output types
typedef enum {
    GPIO_OTYPE_PP = 0,    // Push-pull
    GPIO_OTYPE_OD        // Open-drain
} gpio_otype_t;

// GPIO speed settings
typedef enum {
    GPIO_SPEED_LOW = 0,
    GPIO_SPEED_MEDIUM,
```

```

        GPIO_SPEED_HIGH,
        GPIO_SPEED_VERY_HIGH
    } gpio_speed_t;

// GPIO pull-up/pull-down settings
typedef enum {
    GPIO_PUPD_NONE = 0,
    GPIO_PUPD_UP,
    GPIO_PUPD_DOWN
} gpio_pupd_t;

// GPIO port definitions
typedef enum {
    GPIO_PORT_A = 0,
    GPIO_PORT_B,
    GPIO_PORT_C,
    GPIO_PORT_D,
    GPIO_PORT_E,
    GPIO_PORT_F,
    GPIO_PORT_G,
    GPIO_PORT_H
} gpio_port_t;

// GPIO pin definitions
typedef enum {
    GPIO_PIN_0 = 0,
    GPIO_PIN_1,
    GPIO_PIN_2,
    GPIO_PIN_3,
    GPIO_PIN_4,
    GPIO_PIN_5,
    GPIO_PIN_6,
    GPIO_PIN_7,
    GPIO_PIN_8,
    GPIO_PIN_9,
    GPIO_PIN_10,
    GPIO_PIN_11,
    GPIO_PIN_12,
    GPIO_PIN_13,
    GPIO_PIN_14,
    GPIO_PIN_15
} gpio_pin_t;

// GPIO register structure
typedef struct {
    volatile uint32_t MODER;    // Mode register
    volatile uint32_t OTYPER;  // Output type register

```

```

    volatile uint32_t OSPEEDR; // Output speed register
    volatile uint32_t PUPDR;   // Pull-up/pull-down register
    volatile uint32_t IDR;     // Input data register
    volatile uint32_t ODR;     // Output data register
    volatile uint32_t BSRR;    // Bit set/reset register
    volatile uint32_t LCKR;    // Configuration lock register
    volatile uint32_t AFR[2];  // Alternate function registers
} gpio_reg_t;

// Base addresses for GPIO ports
#define GPIOA_BASE 0x40020000
#define GPIOB_BASE 0x40020400
#define GPIOC_BASE 0x40020800
#define GPIOD_BASE 0x40020C00
#define GPIOE_BASE 0x40021000
#define GPIOF_BASE 0x40021400
#define GPIOG_BASE 0x40021800
#define GPIOH_BASE 0x40021C00

// Get GPIO register pointer
static gpio_reg_t* get_gpio_reg(gpio_port_t port) {
    switch (port) {
        case GPIO_PORT_A: return (gpio_reg_t*)GPIOA_BASE;
        case GPIO_PORT_B: return (gpio_reg_t*)GPIOB_BASE;
        case GPIO_PORT_C: return (gpio_reg_t*)GPIOC_BASE;
        case GPIO_PORT_D: return (gpio_reg_t*)GPIOD_BASE;
        case GPIO_PORT_E: return (gpio_reg_t*)GPIOE_BASE;
        case GPIO_PORT_F: return (gpio_reg_t*)GPIOF_BASE;
        case GPIO_PORT_G: return (gpio_reg_t*)GPIOG_BASE;
        case GPIO_PORT_H: return (gpio_reg_t*)GPIOH_BASE;
        default: return NULL;
    }
}

// Configure GPIO pin
int gpio_init(gpio_port_t port, gpio_pin_t pin, gpio_mode_t mode,
              gpio_otype_t otype, gpio_speed_t speed, gpio_pupd_t pupd) {
    gpio_reg_t *gpio = get_gpio_reg(port);
    if (gpio == NULL) {
        return -1;
    }

    uint32_t pin_pos = (uint32_t)pin * 2;

    // Configure mode
    gpio->MODER &= ~(3UL << pin_pos);

```

```

    gpio->MODER |= ((uint32_t)mode << pin_pos);

    // Configure output type (only for output modes)
    if (mode == GPIO_MODE_OUTPUT || mode == GPIO_MODE_AF) {
        gpio->OTYPER &= ~(1UL << pin);
        gpio->OTYPER |= ((uint32_t)otype << pin);
    }

    // Configure speed
    gpio->OSPEEDR &= ~(3UL << pin_pos);
    gpio->OSPEEDR |= ((uint32_t)speed << pin_pos);

    // Configure pull-up/pull-down
    gpio->PUPDR &= ~(3UL << pin_pos);
    gpio->PUPDR |= ((uint32_t)pupd << pin_pos);

    return 0;
}

// Set GPIO pin high
void gpio_set(gpio_port_t port, gpio_pin_t pin) {
    gpio_reg_t *gpio = get_gpio_reg(port);
    if (gpio != NULL) {
        gpio->BSRR = (1UL << pin);
    }
}

// Clear GPIO pin (set low)
void gpio_clear(gpio_port_t port, gpio_pin_t pin) {
    gpio_reg_t *gpio = get_gpio_reg(port);
    if (gpio != NULL) {
        gpio->BSRR = (1UL << (pin + 16));
    }
}

// Toggle GPIO pin
void gpio_toggle(gpio_port_t port, gpio_pin_t pin) {
    gpio_reg_t *gpio = get_gpio_reg(port);
    if (gpio != NULL) {
        gpio->ODR ^= (1UL << pin);
    }
}

// Read GPIO pin state
int gpio_read(gpio_port_t port, gpio_pin_t pin) {
    gpio_reg_t *gpio = get_gpio_reg(port);

```

```

    if (gpio != NULL) {
        return (gpio->IDR & (1UL << pin)) ? 1 : 0;
    }
    return -1;
}

// Configure GPIO pin as input with pull-up
int gpio_config_input_pullup(gpio_port_t port, gpio_pin_t pin) {
    return gpio_init(port, pin, GPIO_MODE_INPUT, GPIO_OTYPE_PP,
        GPIO_SPEED_LOW, GPIO_PUPD_UP);
}

// Configure GPIO pin as output push-pull
int gpio_config_output(gpio_port_t port, gpio_pin_t pin) {
    return gpio_init(port, pin, GPIO_MODE_OUTPUT, GPIO_OTYPE_PP,
        GPIO_SPEED_HIGH, GPIO_PUPD_NONE);
}

```

SPI (Serial Peripheral Interface)

SPI is a synchronous serial communication interface used for short-distance communication between microcontrollers and peripheral devices.

SPI Driver Implementation

```

#include <stdint.h>
#include <stdbool.h>

// SPI register structure
typedef struct {
    volatile uint32_t CR1;    // Control register 1
    volatile uint32_t CR2;    // Control register 2
    volatile uint32_t SR;     // Status register
    volatile uint32_t DR;     // Data register
    volatile uint32_t CRCPR;  // CRC polynomial register
    volatile uint32_t RXCRCR; // RX CRC register
    volatile uint32_t TXCRCR; // TX CRC register
    volatile uint32_t I2SCFGR; // I2S configuration register
    volatile uint32_t I2SPR;  // I2S prescaler register
} spi_reg_t;

// SPI base addresses
#define SPI1_BASE 0x40013000
#define SPI2_BASE 0x40003800

```



```

#define SPI3_BASE 0x40003C00

#define SPI1 ((spi_reg_t*)SPI1_BASE)
#define SPI2 ((spi_reg_t*)SPI2_BASE)
#define SPI3 ((spi_reg_t*)SPI3_BASE)

// SPI control register bits
#define SPI_CR1_CPHA (1UL << 0) // Clock phase
#define SPI_CR1_CPOL (1UL << 1) // Clock polarity
#define SPI_CR1_MSTR (1UL << 2) // Master selection
#define SPI_CR1_BR (7UL << 3) // Baud rate control
#define SPI_CR1_SPE (1UL << 6) // SPI enable
#define SPI_CR1_LSBFIRST (1UL << 7) // Frame format
#define SPI_CR1_SSI (1UL << 8) // Internal slave select
#define SPI_CR1_SSM (1UL << 9) // Software slave management
#define SPI_CR1_RXONLY (1UL << 10) // Receive only
#define SPI_CR1_DFF (1UL << 11) // Data frame format
#define SPI_CR1_CRCNEXT (1UL << 12) // CRC transfer next
#define SPI_CR1_CRCEN (1UL << 13) // Hardware CRC calculation
#define SPI_CR1_BIDIOE (1UL << 14) // Output enable in bidirectional mode
#define SPI_CR1_BIDIMODE (1UL << 15) // Bidirectional data mode enable

// SPI status register bits
#define SPI_SR_RXNE (1UL << 0) // Receive buffer not empty
#define SPI_SR_TXE (1UL << 1) // Transmit buffer empty
#define SPI_SR_CHSIDE (1UL << 2) // Channel side
#define SPI_SR_UDR (1UL << 3) // Underrun flag
#define SPI_SR_CRCERR (1UL << 4) // CRC error flag
#define SPI_SR_MODF (1UL << 5) // Mode fault
#define SPI_SR_OVR (1UL << 6) // Overrun flag
#define SPI_SR_BSY (1UL << 7) // Busy flag

// SPI configuration structure
typedef struct {
    uint32_t baud_rate; // Baud rate prescaler
    bool master; // Master mode
    bool cpol; // Clock polarity
    bool cpha; // Clock phase
    bool lsb_first; // LSB first
    bool bidirectional; // Bidirectional mode
    uint8_t data_size; // Data size (8 or 16 bits)
} spi_config_t;

// Initialize SPI peripheral
int spi_init(spi_reg_t *spi, const spi_config_t *config) {
    if (spi == NULL || config == NULL) {
        return -1;
    }

```

```

}

// Disable SPI
spi->CR1 &= ~SPI_CR1_SPE;

// Configure SPI
uint32_t cr1 = 0;

// Set baud rate
cr1 |= (config->baud_rate << 3) & SPI_CR1_BR;

// Set master/slave mode
if (config->master) {
    cr1 |= SPI_CR1_MSTR;
    cr1 |= SPI_CR1_SSI; // Set internal slave select
}

// Set clock polarity and phase
if (config->cpol) cr1 |= SPI_CR1_CPOL;
if (config->cpha) cr1 |= SPI_CR1_CPHA;

// Set data frame format
if (config->data_size == 16) {
    cr1 |= SPI_CR1_DFF;
}

// Set frame format
if (config->lsb_first) {
    cr1 |= SPI_CR1_LSBFIRST;
}

// Enable software slave management
cr1 |= SPI_CR1_SSM;

spi->CR1 = cr1;

// Configure bidirectional mode
if (config->bidirectional) {
    spi->CR1 |= SPI_CR1_BIDIMODE;
    spi->CR1 |= SPI_CR1_BIDIOE; // Enable output
}

// Enable SPI
spi->CR1 |= SPI_CR1_SPE;

return 0;
}

```

```

// SPI transmit function
int spi_transmit(spi_reg_t *spi, const uint8_t *data, uint32_t length) {
    if (spi == NULL || data == NULL || length == 0) {
        return -1;
    }

    for (uint32_t i = 0; i < length; i++) {
        // Wait for transmit buffer to be empty
        while (!(spi->SR & SPI_SR_TXE));

        // Send data
        spi->DR = data[i];

        // Wait for transmission to complete
        while (spi->SR & SPI_SR_BSY);
    }

    return 0;
}

// SPI receive function
int spi_receive(spi_reg_t *spi, uint8_t *data, uint32_t length) {
    if (spi == NULL || data == NULL || length == 0) {
        return -1;
    }

    for (uint32_t i = 0; i < length; i++) {
        // Wait for receive buffer to be empty
        while (!(spi->SR & SPI_SR_TXE));

        // Send dummy data to initiate reception
        spi->DR = 0xFF;

        // Wait for data to be received
        while (!(spi->SR & SPI_SR_RXNE));

        // Read received data
        data[i] = spi->DR;
    }

    return 0;
}

// SPI transmit and receive function (full duplex)
int spi_transceive(spi_reg_t *spi, const uint8_t *tx_data,
                   uint8_t *rx_data, uint32_t length) {

```

```

    if (spi == NULL || tx_data == NULL || rx_data == NULL || length == 0) {
        return -1;
    }

    for (uint32_t i = 0; i < length; i++) {
        // Wait for transmit buffer to be empty
        while (!(spi->SR & SPI_SR_TXE));

        // Send data
        spi->DR = tx_data[i];

        // Wait for data to be received
        while (!(spi->SR & SPI_SR_RXNE));

        // Read received data
        rx_data[i] = spi->DR;
    }

    return 0;
}

// SPI chip select control
void spi_select_device(gpio_port_t cs_port, gpio_pin_t cs_pin) {
    gpio_clear(cs_port, cs_pin); // Active low CS
}

void spi_deselect_device(gpio_port_t cs_port, gpio_pin_t cs_pin) {
    gpio_set(cs_port, cs_pin); // Inactive high CS
}

```

I2C (Inter-Integrated Circuit)

I2C is a multi-master, multi-slave, packet-switched, single-ended, serial communication bus used for short-distance communication.

I2C Driver Implementation

```

#include <stdint.h>
#include <stdbool.h>

// I2C register structure
typedef struct {
    volatile uint32_t CR1; // Control register 1
    volatile uint32_t CR2; // Control register 2
}

```

```

    volatile uint32_t OAR1;    // Own address register 1
    volatile uint32_t OAR2;    // Own address register 2
    volatile uint32_t DR;      // Data register
    volatile uint32_t SR1;     // Status register 1
    volatile uint32_t SR2;     // Status register 2
    volatile uint32_t CCR;     // Clock control register
    volatile uint32_t TRISE;   // TRISE register
    volatile uint32_t FLTR;    // FLTR register
} i2c_reg_t;

// I2C base addresses
#define I2C1_BASE 0x40005400
#define I2C2_BASE 0x40005800
#define I2C3_BASE 0x40005C00

#define I2C1 ((i2c_reg_t*)I2C1_BASE)
#define I2C2 ((i2c_reg_t*)I2C2_BASE)
#define I2C3 ((i2c_reg_t*)I2C3_BASE)

// I2C control register 1 bits
#define I2C_CR1_PE      (1UL << 0)    // Peripheral enable
#define I2C_CR1_SMBUS   (1UL << 1)    // SMBus mode
#define I2C_CR1_SMBTYPE (1UL << 3)    // SMBus type
#define I2C_CR1_ENARP   (1UL << 4)    // ARP enable
#define I2C_CR1_ENPEC   (1UL << 5)    // PEC enable
#define I2C_CR1_ENGC    (1UL << 6)    // General call enable
#define I2C_CR1_NOSTRETCH (1UL << 7) // Clock stretching disable
#define I2C_CR1_START   (1UL << 8)    // Start generation
#define I2C_CR1_STOP    (1UL << 9)    // Stop generation
#define I2C_CR1_ACK     (1UL << 10)   // Acknowledge enable
#define I2C_CR1_POS     (1UL << 11)   // Acknowledge/PEC position
#define I2C_CR1_PEC     (1UL << 12)   // Packet error checking
#define I2C_CR1_ALERT   (1UL << 13)   // SMBus alert
#define I2C_CR1_SWRST   (1UL << 15)   // Software reset

// I2C status register 1 bits
#define I2C_SR1_SB      (1UL << 0)    // Start bit
#define I2C_SR1_ADDR    (1UL << 1)    // Address sent/matched
#define I2C_SR1_BTF     (1UL << 2)    // Byte transfer finished
#define I2C_SR1_ADD10   (1UL << 3)    // 10-bit header sent
#define I2C_SR1_STOPF   (1UL << 4)    // Stop detection
#define I2C_SR1_RXNE    (1UL << 6)    // Data register not empty
#define I2C_SR1_TXE     (1UL << 7)    // Data register empty
#define I2C_SR1_BERR    (1UL << 8)    // Bus error
#define I2C_SR1_ARLO    (1UL << 9)    // Arbitration lost
#define I2C_SR1_AF      (1UL << 10)   // Acknowledge failure
#define I2C_SR1_OVR     (1UL << 11)   // Overrun/underrun

```

```

#define I2C_SR1_PECERR (1UL << 12) // PEC error
#define I2C_SR1_TIMEOUT (1UL << 14) // Timeout or Tlow error
#define I2C_SR1_SMBALERT (1UL << 15) // SMBus alert

// I2C configuration structure
typedef struct {
    uint32_t clock_speed; // I2C clock speed in Hz
    bool master; // Master mode
    uint16_t own_address; // Own address (slave mode)
    bool ack_enable; // Acknowledge enable
} i2c_config_t;

// Initialize I2C peripheral
int i2c_init(i2c_reg_t *i2c, const i2c_config_t *config) {
    if (i2c == NULL || config == NULL) {
        return -1;
    }

    // Disable I2C
    i2c->CR1 &= ~I2C_CR1_PE;

    // Configure I2C clock
    uint32_t freq_mhz = 0; // Peripheral clock frequency in MHz
    // freq_mhz = get_pclk1_frequency() / 1000000;

    // Configure CR2 register (frequency)
    i2c->CR2 = freq_mhz & 0x3F;

    // Configure CCR register (clock control)
    uint32_t ccr_value = 0;
    if (config->clock_speed <= 100000) {
        // Standard mode (up to 100kHz)
        ccr_value = (freq_mhz * 1000000) / (config->clock_speed * 2);
        if (ccr_value < 4) ccr_value = 4;
    } else {
        // Fast mode (up to 400kHz)
        ccr_value = (freq_mhz * 1000000) / (config->clock_speed * 3);
        if (ccr_value < 1) ccr_value = 1;
        ccr_value |= (1UL << 15); // Fast mode
    }
    i2c->CCR = ccr_value;

    // Configure TRISE register (rise time)
    if (config->clock_speed <= 100000) {
        // Standard mode
        i2c->TRISE = freq_mhz + 1;
    } else {

```

```

        // Fast mode
        i2c->TRISE = (freq_mhz * 300) / 1000 + 1;
    }

    // Configure own address (slave mode)
    if (!config->master) {
        i2c->OAR1 = (config->own_address << 1) | (1UL << 14);
    }

    // Enable acknowledge
    if (config->ack_enable) {
        i2c->CR1 |= I2C_CR1_ACK;
    }

    // Enable I2C
    i2c->CR1 |= I2C_CR1_PE;

    return 0;
}

// Generate start condition
static int i2c_start(i2c_reg_t *i2c) {
    // Generate start condition
    i2c->CR1 |= I2C_CR1_START;

    // Wait for start condition to be generated
    uint32_t timeout = 100000;
    while (!(i2c->SR1 & I2C_SR1_SB) && timeout--) {
        if (i2c->SR1 & (I2C_SR1_ARLO | I2C_SR1_BERR | I2C_SR1_AF)) {
            return -1; // Error occurred
        }
    }

    if (timeout == 0) {
        return -2; // Timeout
    }

    return 0;
}

// Generate stop condition
static void i2c_stop(i2c_reg_t *i2c) {
    i2c->CR1 |= I2C_CR1_STOP;
}

// Send address

```

```

static int i2c_send_address(i2c_reg_t *i2c, uint8_t address, bool read) {
    // Send address
    i2c->DR = (address << 1) | (read ? 1 : 0);

    // Wait for address to be sent
    uint32_t timeout = 100000;
    while (!(i2c->SR1 & I2C_SR1_ADDR) && timeout--) {
        if (i2c->SR1 & (I2C_SR1_ARLO | I2C_SR1_BERR | I2C_SR1_AF)) {
            return -1; // Error occurred
        }
    }

    if (timeout == 0) {
        return -2; // Timeout
    }

    // Clear ADDR flag by reading SR1 and SR2
    volatile uint32_t temp = i2c->SR1;
    temp = i2c->SR2;
    (void)temp; // Prevent compiler warning

    return 0;
}

// Send data
static int i2c_send_data(i2c_reg_t *i2c, const uint8_t *data, uint32_t length) {
    for (uint32_t i = 0; i < length; i++) {
        // Wait for transmit buffer to be empty
        uint32_t timeout = 100000;
        while (!(i2c->SR1 & I2C_SR1_TXE) && timeout--) {
            if (i2c->SR1 & (I2C_SR1_ARLO | I2C_SR1_BERR | I2C_SR1_AF | I2C_SR1_OVR)) {
                return -1; // Error occurred
            }
        }

        if (timeout == 0) {
            return -2; // Timeout
        }

        // Send data
        i2c->DR = data[i];
    }

    // Wait for byte transfer to finish
    uint32_t timeout = 100000;
    while (!(i2c->SR1 & I2C_SR1_BTF) && timeout--) {
        if (i2c->SR1 & (I2C_SR1_ARLO | I2C_SR1_BERR | I2C_SR1_AF | I2C_SR1_OVR)) {

```



```

        return -1; // Error occurred
    }
}

if (timeout == 0) {
    return -2; // Timeout
}

return 0;
}

// Receive data
static int i2c_receive_data(i2c_reg_t *i2c, uint8_t *data, uint32_t length) {
    // Enable acknowledge for all bytes except the last one
    if (length > 1) {
        i2c->CR1 |= I2C_CR1_ACK;
    } else {
        i2c->CR1 &= ~I2C_CR1_ACK;
    }

    for (uint32_t i = 0; i < length; i++) {
        // For the last byte, generate stop condition before reading
        if (i == length - 1) {
            i2c->CR1 &= ~I2C_CR1_ACK; // Disable acknowledge
            i2c_stop(i2c);           // Generate stop condition
        }

        // Wait for data to be received
        uint32_t timeout = 100000;
        while (!(i2c->SR1 & I2C_SR1_RXNE) && timeout--) {
            if (i2c->SR1 & (I2C_SR1_ARLO | I2C_SR1_BERR | I2C_SR1_AF | I2C_SR1_OVR)) {
                return -1; // Error occurred
            }
        }

        if (timeout == 0) {
            return -2; // Timeout
        }

        // Read data
        data[i] = i2c->DR;
    }

    return 0;
}

// I2C write function

```

```

int i2c_write(i2c_reg_t *i2c, uint8_t device_address,
              const uint8_t *data, uint32_t length) {
    if (i2c == NULL || data == NULL || length == 0) {
        return -1;
    }

    // Generate start condition
    if (i2c_start(i2c) != 0) {
        return -2;
    }

    // Send device address with write bit
    if (i2c_send_address(i2c, device_address, false) != 0) {
        i2c_stop(i2c);
        return -3;
    }

    // Send data
    if (i2c_send_data(i2c, data, length) != 0) {
        i2c_stop(i2c);
        return -4;
    }

    // Generate stop condition
    i2c_stop(i2c);

    return 0;
}

// I2C read function
int i2c_read(i2c_reg_t *i2c, uint8_t device_address,
             uint8_t *data, uint32_t length) {
    if (i2c == NULL || data == NULL || length == 0) {
        return -1;
    }

    // Generate start condition
    if (i2c_start(i2c) != 0) {
        return -2;
    }

    // Send device address with read bit
    if (i2c_send_address(i2c, device_address, true) != 0) {
        i2c_stop(i2c);
        return -3;
    }
}

```

```

// Receive data
if (i2c_receive_data(i2c, data, length) != 0) {
    return -4;
}

return 0;
}

// I2C write-read function (repeated start)
int i2c_write_read(i2c_reg_t *i2c, uint8_t device_address,
                  const uint8_t *write_data, uint32_t write_length,
                  uint8_t *read_data, uint32_t read_length) {
    if (i2c == NULL || write_data == NULL || read_data == NULL ||
        write_length == 0 || read_length == 0) {
        return -1;
    }

    // Generate start condition
    if (i2c_start(i2c) != 0) {
        return -2;
    }

    // Send device address with write bit
    if (i2c_send_address(i2c, device_address, false) != 0) {
        i2c_stop(i2c);
        return -3;
    }

    // Send write data
    if (i2c_send_data(i2c, write_data, write_length) != 0) {
        i2c_stop(i2c);
        return -4;
    }

    // Generate repeated start condition
    if (i2c_start(i2c) != 0) {
        i2c_stop(i2c);
        return -5;
    }

    // Send device address with read bit
    if (i2c_send_address(i2c, device_address, true) != 0) {
        i2c_stop(i2c);
        return -6;
    }

    // Receive read data

```

```

    if (i2c_receive_data(i2c, read_data, read_length) != 0) {
        return -7;
    }

    return 0;
}

```

UART (Universal Asynchronous Receiver-Transmitter)

UART is used for asynchronous serial communication between devices.

UART Driver Implementation

```

#include <stdint.h>
#include <stdbool.h>

// UART register structure
typedef struct {
    volatile uint32_t SR;        // Status register
    volatile uint32_t DR;        // Data register
    volatile uint32_t BRR;       // Baud rate register
    volatile uint32_t CR1;       // Control register 1
    volatile uint32_t CR2;       // Control register 2
    volatile uint32_t CR3;       // Control register 3
    volatile uint32_t GTPR;      // Guard time and prescaler register
} uart_reg_t;

// UART base addresses
#define USART1_BASE 0x40011000
#define USART2_BASE 0x40004400
#define USART3_BASE 0x40004800
#define UART4_BASE 0x40004C00
#define UART5_BASE 0x40005000

#define USART1 ((uart_reg_t*)USART1_BASE)
#define USART2 ((uart_reg_t*)USART2_BASE)
#define USART3 ((uart_reg_t*)USART3_BASE)
#define UART4 ((uart_reg_t*)UART4_BASE)
#define UART5 ((uart_reg_t*)UART5_BASE)

// UART status register bits
#define UART_SR_PE (1UL << 0) // Parity error
#define UART_SR_FE (1UL << 1) // Framing error
#define UART_SR_NE (1UL << 2) // Noise error

```

```

#define UART_SR_ORE      (1UL << 3)    // Overrun error
#define UART_SR_IDLE     (1UL << 4)    // IDLE line detected
#define UART_SR_RXNE     (1UL << 5)    // Read data register not empty
#define UART_SR_TC       (1UL << 6)    // Transmission complete
#define UART_SR_TXE       (1UL << 7)    // Transmit data register empty
#define UART_SR_LBD      (1UL << 8)    // LIN break detection flag
#define UART_SR_CTS      (1UL << 9)    // CTS flag

// UART control register 1 bits
#define UART_CR1_SBK      (1UL << 0)    // Send break
#define UART_CR1_RWU      (1UL << 1)    // Receiver wakeup
#define UART_CR1_RE        (1UL << 2)    // Receiver enable
#define UART_CR1_TE        (1UL << 3)    // Transmitter enable
#define UART_CR1_IDLEIE    (1UL << 4)    // IDLE interrupt enable
#define UART_CR1_RXNEIE    (1UL << 5)    // RXNE interrupt enable
#define UART_CR1_TCIE      (1UL << 6)    // Transmission complete interrupt enable
#define UART_CR1_TXEIE      (1UL << 7)    // TXE interrupt enable
#define UART_CR1_PEIE      (1UL << 8)    // PE interrupt enable
#define UART_CR1_PS        (1UL << 9)    // Parity selection
#define UART_CR1_PCE        (1UL << 10)    // Parity control enable
#define UART_CR1_WAKE      (1UL << 11)    // Wakeup method
#define UART_CR1_M          (1UL << 12)    // Word length
#define UART_CR1_UE         (1UL << 13)    // USART enable
#define UART_CR1_OVER8      (1UL << 15)    // Oversampling mode

// UART configuration structure
typedef struct {
    uint32_t baud_rate;        // Baud rate
    uint8_t data_bits;         // Data bits (8 or 9)
    uint8_t stop_bits;         // Stop bits (1 or 2)
    bool parity_enable;        // Parity enable
    bool parity_even;          // Even parity (true) or odd parity (false)
    uint32_t peripheral_clock; // Peripheral clock frequency
} uart_config_t;

// Initialize UART peripheral
int uart_init(uart_reg_t *uart, const uart_config_t *config) {
    if (uart == NULL || config == NULL) {
        return -1;
    }

    // Disable UART
    uart->CR1 &= ~UART_CR1_UE;

    // Configure baud rate
    uint32_t baud_div = 0;

```

```

    if (config->baud_rate != 0) {
        baud_div = (config->peripheral_clock + (config->baud_rate / 2)) / config->baud_rate;
    }
    uart->BRR = baud_div;

    // Configure data bits
    if (config->data_bits == 9) {
        uart->CR1 |= UART_CR1_M;
    } else {
        uart->CR1 &= ~UART_CR1_M;
    }

    // Configure parity
    if (config->parity_enable) {
        uart->CR1 |= UART_CR1_PCE;
        if (config->parity_even) {
            uart->CR1 &= ~UART_CR1_PS; // Even parity
        } else {
            uart->CR1 |= UART_CR1_PS; // Odd parity
        }
    } else {
        uart->CR1 &= ~UART_CR1_PCE;
    }

    // Configure stop bits
    uart->CR2 &= ~(3UL << 12); // Clear stop bits
    if (config->stop_bits == 2) {
        uart->CR2 |= (2UL << 12); // 2 stop bits
    } else {
        uart->CR2 |= (0UL << 12); // 1 stop bit
    }

    // Enable transmitter and receiver
    uart->CR1 |= UART_CR1_TE | UART_CR1_RE;

    // Enable UART
    uart->CR1 |= UART_CR1_UE;

    return 0;
}

// Send character
int uart_putc(uart_reg_t *uart, char c) {
    if (uart == NULL) {
        return -1;
    }
}

```

```

// Wait for transmit data register to be empty
uint32_t timeout = 100000;
while (!(uart->SR & UART_SR_TXE) && timeout--) {
    // Wait
}

if (timeout == 0) {
    return -2; // Timeout
}

// Send character
uart->DR = c;

return 0;
}

// Send string
int uart_puts(uart_reg_t *uart, const char *str) {
    if (uart == NULL || str == NULL) {
        return -1;
    }

    while (*str) {
        if (uart_putc(uart, *str++) != 0) {
            return -2;
        }
    }

    return 0;
}

// Receive character
int uart_getc(uart_reg_t *uart, char *c) {
    if (uart == NULL || c == NULL) {
        return -1;
    }

    // Check if data is available
    if (!(uart->SR & UART_SR_RXNE)) {
        return -2; // No data available
    }

    // Read character
    *c = (char)(uart->DR & 0xFF);

    // Check for errors
    if (uart->SR & (UART_SR_PE | UART_SR_FE | UART_SR_NE | UART_SR_ORE)) {

```

```

        return -3; // Error occurred
    }

    return 0;
}

// Check if data is available
bool uart_available(uart_reg_t *uart) {
    if (uart == NULL) {
        return false;
    }
    return (uart->SR & UART_SR_RXNE) ? true : false;
}

// Flush UART receive buffer
void uart_flush(uart_reg_t *uart) {
    if (uart == NULL) {
        return;
    }

    // Read all available data
    while (uart->SR & UART_SR_RXNE) {
        volatile uint32_t temp = uart->DR;
        (void)temp; // Prevent compiler warning
    }
}

```

Practical Examples

Sensor Interface Example

```

#include <stdint.h>
#include <stdbool.h>

// Temperature sensor (I2C) driver
#define TEMP_SENSOR_ADDR 0x48

typedef struct {
    i2c_reg_t *i2c;
    uint8_t address;
} temp_sensor_t;

// Initialize temperature sensor
int temp_sensor_init(temp_sensor_t *sensor, i2c_reg_t *i2c) {
    if (sensor == NULL || i2c == NULL) {

```



```

        return -1;
    }

    sensor->i2c = i2c;
    sensor->address = TEMP_SENSOR_ADDR;

    // Configure sensor (example configuration)
    uint8_t config_data[] = {0x01, 0x60}; // Configuration register and value
    return i2c_write(sensor->i2c, sensor->address, config_data, 2);
}

// Read temperature
float temp_sensor_read(temp_sensor_t *sensor) {
    if (sensor == NULL) {
        return -1.0;
    }

    // Read temperature register
    uint8_t reg_addr = 0x00; // Temperature register
    uint8_t data[2];

    if (i2c_write_read(sensor->i2c, sensor->address,
                        &reg_addr, 1, data, 2) != 0) {
        return -1.0;
    }

    // Convert to temperature value
    int16_t raw_temp = (data[0] << 8) | data[1];
    float temperature = raw_temp / 256.0;

    return temperature;
}

// EEPROM (I2C) driver
#define EEPROM_ADDR 0x50

typedef struct {
    i2c_reg_t *i2c;
    uint8_t address;
} eeprom_t;

// Initialize EEPROM
int eeprom_init(eeprom_t *eeprom, i2c_reg_t *i2c) {
    if (eeprom == NULL || i2c == NULL) {
        return -1;
    }
}

```

```

    eeprom->i2c = i2c;
    eeprom->address = EEPROM_ADDR;

    return 0;
}

// Write data to EEPROM
int eeprom_write(eeprom_t *eeprom, uint16_t address,
                 const uint8_t *data, uint32_t length) {
    if (eeprom == NULL || data == NULL || length == 0) {
        return -1;
    }

    // Write data in chunks (EEPROM page size is typically 16-64 bytes)
    uint32_t written = 0;
    while (written < length) {
        uint32_t chunk_size = (length - written > 16) ? 16 : (length - written);

        // Prepare data with address
        uint8_t write_data[18]; // 2 bytes address + 16 bytes data
        write_data[0] = (address + written) >> 8; // High byte of address
        write_data[1] = (address + written) & 0xFF; // Low byte of address
        for (uint32_t i = 0; i < chunk_size; i++) {
            write_data[2 + i] = data[written + i];
        }

        // Write chunk
        if (i2c_write(eeprom->i2c, eeprom->address,
                     write_data, 2 + chunk_size) != 0) {
            return -2;
        }

        written += chunk_size;

        // Wait for write to complete (EEPROM typical write time is 5ms)
        rt_delay_ms(5);
    }

    return 0;
}

// Read data from EEPROM
int eeprom_read(eeprom_t *eeprom, uint16_t address,
                uint8_t *data, uint32_t length) {
    if (eeprom == NULL || data == NULL || length == 0) {
        return -1;
    }
}

```

```

    // Send address
    uint8_t addr_data[2] = {(address >> 8) & 0xFF, address & 0xFF};
    if (i2c_write(eeprom->i2c, eeprom->address, addr_data, 2) != 0) {
        return -2;
    }

    // Read data
    return i2c_read(eeprom->i2c, eeprom->address, data, length);
}

// SPI Flash memory driver
typedef struct {
    spi_reg_t *spi;
    gpio_port_t cs_port;
    gpio_pin_t cs_pin;
} flash_t;

// Initialize flash memory
int flash_init(flash_t *flash, spi_reg_t *spi,
               gpio_port_t cs_port, gpio_pin_t cs_pin) {
    if (flash == NULL || spi == NULL) {
        return -1;
    }

    flash->spi = spi;
    flash->cs_port = cs_port;
    flash->cs_pin = cs_pin;

    // Configure CS pin as output
    gpio_config_output(cs_port, cs_pin);
    gpio_set(cs_port, cs_pin); // Deselect flash

    return 0;
}

// Send command to flash
static int flash_send_command(flash_t *flash, const uint8_t *cmd, uint32_t cmd_len,
                              const uint8_t *data_out, uint8_t *data_in, uint32_t data_len)
    if (flash == NULL || cmd == NULL || cmd_len == 0) {
        return -1;
    }

    // Select flash
    spi_select_device(flash->cs_port, flash->cs_pin);

    // Send command

```

```

if (spi_transmit(flash->spi, cmd, cmd_len) != 0) {
    spi_deselect_device(flash->cs_port, flash->cs_pin);
    return -2;
}

// Send/receive data
if (data_out != NULL && data_in != NULL) {
    if (spi_transceive(flash->spi, data_out, data_in, data_len) != 0) {
        spi_deselect_device(flash->cs_port, flash->cs_pin);
        return -3;
    }
} else if (data_out != NULL) {
    if (spi_transmit(flash->spi, data_out, data_len) != 0) {
        spi_deselect_device(flash->cs_port, flash->cs_pin);
        return -4;
    }
} else if (data_in != NULL) {
    if (spi_receive(flash->spi, data_in, data_len) != 0) {
        spi_deselect_device(flash->cs_port, flash->cs_pin);
        return -5;
    }
}

// Deselect flash
spi_deselect_device(flash->cs_port, flash->cs_pin);

return 0;
}

// Read flash ID
int flash_read_id(flash_t *flash, uint8_t *id, uint32_t id_len) {
    if (flash == NULL || id == NULL || id_len == 0) {
        return -1;
    }

    uint8_t cmd = 0x9F; // Read ID command
    return flash_send_command(flash, &cmd, 1, NULL, id, id_len);
}

// Read data from flash
int flash_read(flash_t *flash, uint32_t address, uint8_t *data, uint32_t length) {
    if (flash == NULL || data == NULL || length == 0) {
        return -1;
    }

    // Prepare read command with address
    uint8_t cmd[4] = {0x03, // Read command

```

```

        (address >> 16) & 0xFF, // Address high byte
        (address >> 8) & 0xFF,  // Address middle byte
        address & 0xFF};        // Address low byte

    return flash_send_command(flash, cmd, 4, NULL, data, length);
}

// Write data to flash (page program)
int flash_write(flash_t *flash, uint32_t address, const uint8_t *data, uint32_t length) {
    if (flash == NULL || data == NULL || length == 0) {
        return -1;
    }

    // Prepare write command with address
    uint8_t cmd[4] = {0x02, // Page program command
        (address >> 16) & 0xFF, // Address high byte
        (address >> 8) & 0xFF,  // Address middle byte
        address & 0xFF};        // Address low byte

    return flash_send_command(flash, cmd, 4, data, NULL, length);
}

```

Summary

Hardware interfaces are fundamental to embedded systems, enabling communication with various peripheral devices and sensors:

1. **GPIO** - Basic digital input/output for controlling LEDs, reading switches, etc.
2. **SPI** - High-speed synchronous communication for flash memory, displays, and sensors
3. **I2C** - Multi-master communication for EEPROMs, temperature sensors, and other devices
4. **UART** - Asynchronous serial communication for debug output and device communication

Key considerations for hardware interface implementation: - Proper initialization and configuration of peripheral registers - Error handling and timeout mechanisms - Efficient data transfer methods - Proper use of chip select signals for SPI devices - Addressing schemes for I2C devices - Baud rate configuration for UART communication - Interrupt-driven vs. polling-based approaches - Resource sharing and synchronization in multi-threaded environments

These interfaces form the foundation for connecting embedded systems to the physical world, enabling the development of sophisticated applications that interact with sensors, actuators, displays, storage devices, and communication modules.

Module 12: Embedded Systems Programming

Overview

Embedded systems programming is a specialized area of C development focused on creating software for dedicated computer systems that perform specific functions within larger mechanical or electrical systems. This module covers the unique challenges and techniques involved in programming for resource-constrained environments, real-time requirements, and direct hardware interaction.

Learning Objectives

By the end of this module, you should be able to: - Understand the fundamentals of embedded systems and their constraints - Program at the hardware level with direct memory-mapped I/O - Implement real-time programming techniques and scheduling - Work with embedded C extensions and specialized compilers - Interface with various hardware components and sensors - Debug and optimize code for embedded environments - Understand power management and resource optimization

Chapters

1. **Embedded Basics** - Introduction to embedded systems, architectures, and development environments
2. **Low-Level Programming** - Direct hardware access, memory-mapped I/O, and bit manipulation
3. **Embedded C Techniques** - Specialized C features for embedded systems, compiler extensions
4. **Real-time Programming** - Real-time constraints, scheduling, and timing considerations
5. **Hardware Interfaces** - Working with GPIO, SPI, I2C, UART, and other hardware interfaces

Prerequisites

Before starting this module, you should have a solid understanding of: - C programming fundamentals - Pointers and memory management - Basic electronics concepts - Computer architecture fundamentals

Key Concepts

- Resource-constrained programming
- Real-time constraints and deterministic behavior
- Direct hardware manipulation
- Memory-mapped I/O
- Bit manipulation and bit fields
- Interrupt handling
- Power management
- Cross-compilation for embedded targets
- Hardware-software interface design

Performance

Performance Analysis

Introduction

Performance analysis is the process of measuring, evaluating, and understanding the runtime behavior of programs to identify bottlenecks and optimization opportunities. In C programming, where performance is often critical, proper performance analysis is essential for creating efficient applications.

This chapter covers various techniques and tools for analyzing program performance, including profiling, benchmarking, and measurement methodologies.

Performance Metrics

Understanding the right metrics is crucial for effective performance analysis.

Time Complexity Metrics

```
#include <stdio.h>
#include <stdlib.h>
#include <time.h>

// Simple timing function
double get_time(void) {
    struct timespec ts;
    clock_gettime(CLOCK_MONOTONIC, &ts);
    return ts.tv_sec + ts.tv_nsec / 1e9;
}

// Example function to measure
void example_function(int n) {
    volatile int sum = 0;
    for (int i = 0; i < n; i++) {
        for (int j = 0; j < n; j++) {
            sum += i * j;
        }
    }
}
```

```

// Measure execution time
void measure_execution_time(void) {
    double start_time, end_time;

    start_time = get_time();
    example_function(1000);
    end_time = get_time();

    printf("Execution time: %.6f seconds\n", end_time - start_time);
}

```

Memory Usage Metrics

```

#include <stdio.h>
#include <stdlib.h>
#include <string.h>

// Memory tracking structure
typedef struct {
    size_t allocated_bytes;
    size_t freed_bytes;
    size_t current_usage;
    size_t peak_usage;
    size_t allocation_count;
    size_t deallocation_count;
} memory_stats_t;

static memory_stats_t mem_stats = {0};

// Custom malloc wrapper for tracking
void* tracked_malloc(size_t size) {
    void* ptr = malloc(size);
    if (ptr != NULL) {
        mem_stats.allocated_bytes += size;
        mem_stats.current_usage += size;
        mem_stats.allocation_count++;

        if (mem_stats.current_usage > mem_stats.peak_usage) {
            mem_stats.peak_usage = mem_stats.current_usage;
        }
    }
    return ptr;
}

// Custom free wrapper for tracking

```

```

void tracked_free(void* ptr, size_t size) {
    if (ptr != NULL) {
        free(ptr);
        mem_stats.freed_bytes += size;
        mem_stats.current_usage -= size;
        mem_stats.deallocation_count++;
    }
}

// Print memory statistics
void print_memory_stats(void) {
    printf("Memory Statistics:\n");
    printf("  Allocated: %zu bytes\n", mem_stats.allocated_bytes);
    printf("  Freed: %zu bytes\n", mem_stats.freed_bytes);
    printf("  Current Usage: %zu bytes\n", mem_stats.current_usage);
    printf("  Peak Usage: %zu bytes\n", mem_stats.peak_usage);
    printf("  Allocations: %zu\n", mem_stats.allocation_count);
    printf("  Deallocations: %zu\n", mem_stats.deallocation_count);
}

// Reset memory statistics
void reset_memory_stats(void) {
    memset(&mem_stats, 0, sizeof(mem_stats));
}

```

CPU Usage Metrics

```

#include <stdio.h>
#include <stdlib.h>
#include <sys/resource.h>

// Get CPU time usage
double get_cpu_time(void) {
    struct rusage usage;
    if (getrusage(RUSAGE_SELF, &usage) == 0) {
        return usage.ru_utime.tv_sec + usage.ru_utime.tv_usec / 1e6 +
            usage.ru_stime.tv_sec + usage.ru_stime.tv_usec / 1e6;
    }
    return -1.0;
}

// Example CPU-intensive function
void cpu_intensive_task(void) {
    volatile double result = 0.0;
    for (int i = 0; i < 1000000; i++) {

```

```

        result += i * 3.14159;
    }
}

// Measure CPU time
void measure_cpu_time(void) {
    double start_cpu, end_cpu;

    start_cpu = get_cpu_time();
    cpu_intensive_task();
    end_cpu = get_cpu_time();

    if (start_cpu >= 0 && end_cpu >= 0) {
        printf("CPU time used: %.6f seconds\n", end_cpu - start_cpu);
    }
}

```

Profiling Tools

Profiling tools provide detailed insights into program execution and resource usage.

GNU gprof

GNU gprof is a profiling tool that comes with GCC:

```

# Compile with profiling enabled
gcc -pg -o program program.c

# Run the program
./program

# Generate profiling report
gprof program gmon.out > analysis.txt

```

Example program for gprof profiling:

```

#include <stdio.h>
#include <stdlib.h>

// Function to profile
void function_a(void) {
    volatile int sum = 0;
    for (int i = 0; i < 1000000; i++) {
        sum += i;
    }
}

```

```

}

void function_b(void) {
    volatile int product = 1;
    for (int i = 1; i < 1000; i++) {
        product *= i;
    }
}

void function_c(void) {
    volatile double result = 0.0;
    for (int i = 0; i < 500000; i++) {
        result += i * 1.5;
    }
}

int main(void) {
    printf("Starting performance analysis...\n");

    // Call functions multiple times
    for (int i = 0; i < 100; i++) {
        function_a();
        if (i % 2 == 0) {
            function_b();
        }
        if (i % 5 == 0) {
            function_c();
        }
    }

    printf("Performance analysis complete.\n");
    return 0;
}

```

Valgrind Callgrind

Valgrind with Callgrind provides detailed call graph profiling:

```

# Compile normally
gcc -g -O2 -o program program.c

# Run with callgrind
valgrind --tool=callgrind ./program

# View results with kcachegrind
kcachegrind callgrind.out.*

```

perf (Linux Performance Events)

perf is a powerful performance analysis tool on Linux:

```
# Record performance data
perf record -g ./program

# View report
perf report

# Generate flame graph
perf script | flamegraph.pl > perf.svg
```

Custom Profiling Framework

Creating a custom profiling framework for specific needs:

```
#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <time.h>

#define MAX_PROFILERS 100

// Profiler structure
typedef struct {
    char name[64];
    struct timespec start_time;
    long long total_time_ns;
    int call_count;
    int is_active;
} profiler_t;

static profiler_t profilers[MAX_PROFILERS];
static int profiler_count = 0;

// Find or create profiler
static profiler_t* get_profiler(const char* name) {
    // Try to find existing profiler
    for (int i = 0; i < profiler_count; i++) {
        if (strcmp(profilers[i].name, name) == 0) {
            return &profilers[i];
        }
    }
}
```

```

    // Create new profiler if space available
    if (profiler_count < MAX_PROFILERS) {
        profiler_t* p = &profilers[profiler_count];
        strncpy(p->name, name, sizeof(p->name) - 1);
        p->name[sizeof(p->name) - 1] = '\0';
        p->total_time_ns = 0;
        p->call_count = 0;
        p->is_active = 0;
        profiler_count++;
        return p;
    }

    return NULL;
}

// Start profiling a section
void profile_start(const char* name) {
    profiler_t* p = get_profiler(name);
    if (p != NULL && !p->is_active) {
        clock_gettime(CLOCK_MONOTONIC, &p->start_time);
        p->is_active = 1;
    }
}

// Stop profiling a section
void profile_stop(const char* name) {
    profiler_t* p = get_profiler(name);
    if (p != NULL && p->is_active) {
        struct timespec end_time;
        clock_gettime(CLOCK_MONOTONIC, &end_time);

        long long elapsed_ns = (end_time.tv_sec - p->start_time.tv_sec) * 1000000000LL +
                                (end_time.tv_nsec - p->start_time.tv_nsec);

        p->total_time_ns += elapsed_ns;
        p->call_count++;
        p->is_active = 0;
    }
}

// Print profiling results
void profile_print_results(void) {
    printf("\n=== Profiling Results ===\n");
    printf("%-30s %12s %10s %15s\n", "Function", "Calls", "Time(ms)", "Avg Time( s)");
    printf("%-30s %12s %10s %15s\n", "-----", "-----", "-----", "-----");

    for (int i = 0; i < profiler_count; i++) {

```

```

        profiler_t* p = &profilers[i];
        double total_time_ms = p->total_time_ns / 1e6;
        double avg_time_us = (p->call_count > 0) ?
                               (p->total_time_ns / 1e3) / p->call_count : 0;

        printf("%-30s %12d %10.3f %15.3f\n",
               p->name, p->call_count, total_time_ms, avg_time_us);
    }
}

// Reset profiling data
void profile_reset(void) {
    profiler_count = 0;
    memset(profilers, 0, sizeof(profilers));
}

// Example usage
void example_function_1(void) {
    profile_start("example_function_1");

    // Simulate work
    volatile int sum = 0;
    for (int i = 0; i < 100000; i++) {
        sum += i;
    }

    profile_stop("example_function_1");
}

void example_function_2(void) {
    profile_start("example_function_2");

    // Simulate work
    volatile double product = 1.0;
    for (int i = 1; i < 1000; i++) {
        product *= i * 0.1;
    }

    profile_stop("example_function_2");
}

int main(void) {
    printf("Starting custom profiling example...\n");

    // Profile different functions
    for (int i = 0; i < 1000; i++) {
        example_function_1();
    }
}

```



```

        if (i % 2 == 0) {
            example_function_2();
        }
    }

    // Print results
    profile_print_results();

    return 0;
}

```

Benchmarking Framework

Creating a systematic benchmarking framework:

```

#include <stdio.h>
#include <stdlib.h>
#include <time.h>
#include <math.h>

// Benchmark result structure
typedef struct {
    double min_time;
    double max_time;
    double avg_time;
    double median_time;
    double stddev;
    int iterations;
} benchmark_result_t;

// Benchmark function type
typedef void (*benchmark_func_t)(void);

// Run benchmark
benchmark_result_t run_benchmark(benchmark_func_t func, int iterations) {
    benchmark_result_t result = {0};
    double* times = malloc(iterations * sizeof(double));

    if (times == NULL) {
        result.iterations = -1;
        return result;
    }

    struct timespec start, end;

```

```

// Warm up
for (int i = 0; i < 10; i++) {
    func();
}

// Run benchmark iterations
for (int i = 0; i < iterations; i++) {
    clock_gettime(CLOCK_MONOTONIC, &start);
    func();
    clock_gettime(CLOCK_MONOTONIC, &end);

    times[i] = (end.tv_sec - start.tv_sec) + (end.tv_nsec - start.tv_nsec) / 1e9;
}

// Calculate statistics
result.iterations = iterations;
result.min_time = times[0];
result.max_time = times[0];
double sum = 0.0;

for (int i = 0; i < iterations; i++) {
    if (times[i] < result.min_time) result.min_time = times[i];
    if (times[i] > result.max_time) result.max_time = times[i];
    sum += times[i];
}

result.avg_time = sum / iterations;

// Calculate median
// Sort times array (simplified bubble sort for small arrays)
for (int i = 0; i < iterations - 1; i++) {
    for (int j = 0; j < iterations - i - 1; j++) {
        if (times[j] > times[j + 1]) {
            double temp = times[j];
            times[j] = times[j + 1];
            times[j + 1] = temp;
        }
    }
}

if (iterations % 2 == 0) {
    result.median_time = (times[iterations/2 - 1] + times[iterations/2]) / 2.0;
} else {
    result.median_time = times[iterations/2];
}

// Calculate standard deviation

```

```

    double sum_sq_diff = 0.0;
    for (int i = 0; i < iterations; i++) {
        double diff = times[i] - result.avg_time;
        sum_sq_diff += diff * diff;
    }
    result.stddev = sqrt(sum_sq_diff / iterations);

    free(times);
    return result;
}

// Print benchmark results
void print_benchmark_result(const char* name, const benchmark_result_t* result) {
    if (result->iterations < 0) {
        printf("Benchmark '%s' failed\n", name);
        return;
    }

    printf("Benchmark: %s\n", name);
    printf("  Iterations: %d\n", result->iterations);
    printf("  Min Time: %.6f ms\n", result->min_time * 1000);
    printf("  Max Time: %.6f ms\n", result->max_time * 1000);
    printf("  Avg Time: %.6f ms\n", result->avg_time * 1000);
    printf("  Median Time: %.6f ms\n", result->median_time * 1000);
    printf("  Std Dev: %.6f ms\n", result->stddev * 1000);
    printf("\n");
}

// Example functions to benchmark
void benchmark_function_1(void) {
    volatile int sum = 0;
    for (int i = 0; i < 10000; i++) {
        sum += i;
    }
}

void benchmark_function_2(void) {
    volatile double product = 1.0;
    for (int i = 1; i < 1000; i++) {
        product *= (i * 0.01);
    }
}

int main(void) {
    printf("Starting benchmark tests...\n\n");

    // Run benchmarks

```

```

    benchmark_result_t result1 = run_benchmark(benchmark_function_1, 1000);
    print_benchmark_result("Sum Calculation", &result1);

    benchmark_result_t result2 = run_benchmark(benchmark_function_2, 1000);
    print_benchmark_result("Product Calculation", &result2);

    return 0;
}

```

Cache Performance Analysis

Understanding cache behavior is crucial for performance optimization:

```

#include <stdio.h>
#include <stdlib.h>
#include <time.h>

#define CACHE_LINE_SIZE 64
#define ARRAY_SIZE 1000000

// Test cache-friendly access pattern
void cache_friendly_access(int* array, int size) {
    for (int i = 0; i < size; i++) {
        array[i] += 1;
    }
}

// Test cache-unfriendly access pattern
void cache_unfriendly_access(int* array, int size) {
    for (int stride = 1; stride < size; stride *= 2) {
        for (int i = 0; i < size; i += stride) {
            array[i] += 1;
        }
    }
}

// Measure cache performance
void measure_cache_performance(void) {
    int* array = malloc(ARRAY_SIZE * sizeof(int));
    if (array == NULL) {
        printf("Failed to allocate memory\n");
        return;
    }

    struct timespec start, end;

```

```

// Initialize array
for (int i = 0; i < ARRAY_SIZE; i++) {
    array[i] = i;
}

// Measure cache-friendly access
clock_gettime(CLOCK_MONOTONIC, &start);
cache_friendly_access(array, ARRAY_SIZE);
clock_gettime(CLOCK_MONOTONIC, &end);

double friendly_time = (end.tv_sec - start.tv_sec) +
    (end.tv_nsec - start.tv_nsec) / 1e9;

// Reset array
for (int i = 0; i < ARRAY_SIZE; i++) {
    array[i] = i;
}

// Measure cache-unfriendly access
clock_gettime(CLOCK_MONOTONIC, &start);
cache_unfriendly_access(array, ARRAY_SIZE);
clock_gettime(CLOCK_MONOTONIC, &end);

double unfriendly_time = (end.tv_sec - start.tv_sec) +
    (end.tv_nsec - start.tv_nsec) / 1e9;

printf("Cache Performance Analysis:\n");
printf("  Cache-friendly access time: %.6f seconds\n", friendly_time);
printf("  Cache-unfriendly access time: %.6f seconds\n", unfriendly_time);
printf("  Performance ratio: %.2f\n", unfriendly_time / friendly_time);

free(array);
}

```

Memory Access Pattern Analysis

Analyzing memory access patterns for optimization:

```

#include <stdio.h>
#include <stdlib.h>
#include <time.h>

#define MATRIX_SIZE 1000

```

```

// Row-major access (cache-friendly)
void row_major_access(int matrix[MATRIX_SIZE][MATRIX_SIZE]) {
    for (int i = 0; i < MATRIX_SIZE; i++) {
        for (int j = 0; j < MATRIX_SIZE; j++) {
            matrix[i][j] += 1;
        }
    }
}

// Column-major access (cache-unfriendly)
void column_major_access(int matrix[MATRIX_SIZE][MATRIX_SIZE]) {
    for (int j = 0; j < MATRIX_SIZE; j++) {
        for (int i = 0; i < MATRIX_SIZE; i++) {
            matrix[i][j] += 1;
        }
    }
}

// Measure memory access patterns
void measure_memory_access_patterns(void) {
    int(*matrix)[MATRIX_SIZE] = malloc(MATRIX_SIZE * sizeof(*matrix));
    if (matrix == NULL) {
        printf("Failed to allocate memory\n");
        return;
    }

    struct timespec start, end;

    // Initialize matrix
    for (int i = 0; i < MATRIX_SIZE; i++) {
        for (int j = 0; j < MATRIX_SIZE; j++) {
            matrix[i][j] = i * MATRIX_SIZE + j;
        }
    }

    // Measure row-major access
    clock_gettime(CLOCK_MONOTONIC, &start);
    row_major_access(matrix);
    clock_gettime(CLOCK_MONOTONIC, &end);

    double row_time = (end.tv_sec - start.tv_sec) +
        (end.tv_nsec - start.tv_nsec) / 1e9;

    // Reset matrix
    for (int i = 0; i < MATRIX_SIZE; i++) {
        for (int j = 0; j < MATRIX_SIZE; j++) {

```

```

        matrix[i][j] = i * MATRIX_SIZE + j;
    }
}

// Measure column-major access
clock_gettime(CLOCK_MONOTONIC, &start);
column_major_access(matrix);
clock_gettime(CLOCK_MONOTONIC, &end);

double col_time = (end.tv_sec - start.tv_sec) +
    (end.tv_nsec - start.tv_nsec) / 1e9;

printf("Memory Access Pattern Analysis:\n");
printf("  Row-major access time: %.6f seconds\n", row_time);
printf("  Column-major access time: %.6f seconds\n", col_time);
printf("  Performance ratio: %.2f\n", col_time / row_time);

free(matrix);
}

```

Practical Examples

Comprehensive Performance Analysis Tool

```

#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <time.h>

// Performance analysis tool
typedef struct {
    char name[64];
    double execution_time;
    size_t memory_usage;
    int cpu_cycles;
} performance_data_t;

// Simple hash table implementation for performance data
#define HASH_TABLE_SIZE 101

typedef struct perf_node {
    performance_data_t data;
    struct perf_node* next;
} perf_node_t;

```

```

static perf_node_t* hash_table[HASH_TABLE_SIZE];

// Hash function
static unsigned int hash(const char* str) {
    unsigned int hash = 5381;
    int c;
    while ((c = *str++)) {
        hash = ((hash << 5) + hash) + c;
    }
    return hash % HASH_TABLE_SIZE;
}

// Add performance data
void add_performance_data(const char* name, double execution_time,
                          size_t memory_usage, int cpu_cycles) {
    unsigned int index = hash(name);

    perf_node_t* node = malloc(sizeof(perf_node_t));
    if (node == NULL) return;

    strncpy(node->data.name, name, sizeof(node->data.name) - 1);
    node->data.name[sizeof(node->data.name) - 1] = '\0';
    node->data.execution_time = execution_time;
    node->data.memory_usage = memory_usage;
    node->data.cpu_cycles = cpu_cycles;
    node->next = hash_table[index];
    hash_table[index] = node;
}

// Find performance data
performance_data_t* find_performance_data(const char* name) {
    unsigned int index = hash(name);
    perf_node_t* node = hash_table[index];

    while (node != NULL) {
        if (strcmp(node->data.name, name) == 0) {
            return &node->data;
        }
        node = node->next;
    }

    return NULL;
}

// Print all performance data
void print_all_performance_data(void) {

```



```

printf("\n=== Performance Analysis Report ===\n");
printf("%-30s %12s %15s %12s\n", "Function", "Time (ms)", "Memory (bytes)", "CPU Cycles");
printf("%-30s %12s %15s %12s\n", "-----", "-----", "-----", "-----");

for (int i = 0; i < HASH_TABLE_SIZE; i++) {
    perf_node_t* node = hash_table[i];
    while (node != NULL) {
        printf("%-30s %12.3f %15zu %12d\n",
            node->data.name,
            node->data.execution_time * 1000,
            node->data.memory_usage,
            node->data.cpu_cycles);
        node = node->next;
    }
}

// Clear performance data
void clear_performance_data(void) {
    for (int i = 0; i < HASH_TABLE_SIZE; i++) {
        perf_node_t* node = hash_table[i];
        while (node != NULL) {
            perf_node_t* next = node->next;
            free(node);
            node = next;
        }
        hash_table[i] = NULL;
    }
}

// Example performance test functions
void test_function_1(void) {
    volatile int sum = 0;
    for (int i = 0; i < 100000; i++) {
        sum += i;
    }
}

void test_function_2(void) {
    volatile double product = 1.0;
    for (int i = 1; i < 1000; i++) {
        product *= i * 0.1;
    }
}

void test_function_3(void) {
    char* buffer = malloc(1024);

```

```

    if (buffer != NULL) {
        memset(buffer, 0, 1024);
        free(buffer);
    }
}

int main(void) {
    struct timespec start, end;

    printf("Running comprehensive performance analysis...\n");

    // Test function 1
    clock_gettime(CLOCK_MONOTONIC, &start);
    test_function_1();
    clock_gettime(CLOCK_MONOTONIC, &end);
    double time1 = (end.tv_sec - start.tv_sec) + (end.tv_nsec - start.tv_nsec) / 1e9;
    add_performance_data("test_function_1", time1, 0, 0);

    // Test function 2
    clock_gettime(CLOCK_MONOTONIC, &start);
    test_function_2();
    clock_gettime(CLOCK_MONOTONIC, &end);
    double time2 = (end.tv_sec - start.tv_sec) + (end.tv_nsec - start.tv_nsec) / 1e9;
    add_performance_data("test_function_2", time2, 0, 0);

    // Test function 3
    clock_gettime(CLOCK_MONOTONIC, &start);
    test_function_3();
    clock_gettime(CLOCK_MONOTONIC, &end);
    double time3 = (end.tv_sec - start.tv_sec) + (end.tv_nsec - start.tv_nsec) / 1e9;
    add_performance_data("test_function_3", time3, 1024, 0);

    // Print results
    print_all_performance_data();

    // Clean up
    clear_performance_data();

    return 0;
}

```

Summary

Performance analysis is a critical skill for C programmers, involving:

1. **Metrics Collection** - Time, memory, and CPU usage measurements

2. **Profiling Tools** - gprof, Valgrind, perf, and custom profilers
3. **Benchmarking** - Systematic performance testing with statistical analysis
4. **Cache Analysis** - Understanding memory access patterns and cache behavior
5. **Memory Analysis** - Tracking allocations and identifying memory bottlenecks

Key principles for effective performance analysis: - Use appropriate tools for different types of analysis - Measure performance in realistic conditions - Collect sufficient data for statistical significance - Understand the trade-offs between accuracy and overhead - Focus on the most critical performance bottlenecks - Document performance characteristics for future reference - Continuously monitor performance during development

These techniques enable developers to create high-performance C applications that efficiently utilize system resources while meeting performance requirements.

Optimization Techniques

Introduction

Performance optimization in C programming involves applying various techniques to improve execution speed, reduce memory usage, and enhance overall efficiency. This chapter explores fundamental and advanced optimization strategies, including algorithmic improvements, loop optimizations, data structure selection, and compiler-assisted optimizations.

Algorithmic Optimization

The most significant performance improvements often come from choosing better algorithms and data structures.

Time Complexity Analysis

Understanding algorithmic complexity is fundamental to optimization:

```
#include <stdio.h>
#include <stdlib.h>
#include <time.h>

// O(n^2) bubble sort
void bubble_sort(int* arr, int n) {
    for (int i = 0; i < n - 1; i++) {
        for (int j = 0; j < n - i - 1; j++) {
            if (arr[j] > arr[j + 1]) {
                int temp = arr[j];
                arr[j] = arr[j + 1];
                arr[j + 1] = temp;
            }
        }
    }
}

// O(n log n) quick sort
void quick_sort(int* arr, int low, int high) {
    if (low < high) {
        int pi = partition(arr, low, high);
    }
}
```

```

        quick_sort(arr, low, pi - 1);
        quick_sort(arr, pi + 1, high);
    }
}

int partition(int* arr, int low, int high) {
    int pivot = arr[high];
    int i = (low - 1);

    for (int j = low; j <= high - 1; j++) {
        if (arr[j] < pivot) {
            i++;
            int temp = arr[i];
            arr[i] = arr[j];
            arr[j] = temp;
        }
    }

    int temp = arr[i + 1];
    arr[i + 1] = arr[high];
    arr[high] = temp;
    return (i + 1);
}

// O(n log n) merge sort
void merge_sort(int* arr, int l, int r) {
    if (l < r) {
        int m = l + (r - l) / 2;
        merge_sort(arr, l, m);
        merge_sort(arr, m + 1, r);
        merge(arr, l, m, r);
    }
}

void merge(int* arr, int l, int m, int r) {
    int i, j, k;
    int n1 = m - l + 1;
    int n2 = r - m;

    int* L = malloc(n1 * sizeof(int));
    int* R = malloc(n2 * sizeof(int));

    for (i = 0; i < n1; i++)
        L[i] = arr[l + i];
    for (j = 0; j < n2; j++)
        R[j] = arr[m + 1 + j];

    i = 0;

```

```

j = 0;
k = 1;

while (i < n1 && j < n2) {
    if (L[i] <= R[j]) {
        arr[k] = L[i];
        i++;
    } else {
        arr[k] = R[j];
        j++;
    }
    k++;
}

while (i < n1) {
    arr[k] = L[i];
    i++;
    k++;
}

while (j < n2) {
    arr[k] = R[j];
    j++;
    k++;
}

free(L);
free(R);
}

// Performance comparison
void compare_sorting_algorithms(void) {
    const int size = 10000;
    int* arr1 = malloc(size * sizeof(int));
    int* arr2 = malloc(size * sizeof(int));
    int* arr3 = malloc(size * sizeof(int));

    // Initialize arrays with random data
    for (int i = 0; i < size; i++) {
        arr1[i] = arr2[i] = arr3[i] = rand() % 10000;
    }

    struct timespec start, end;

    // Test bubble sort
    clock_gettime(CLOCK_MONOTONIC, &start);
    bubble_sort(arr1, size);

```

```

clock_gettime(CLOCK_MONOTONIC, &end);
double bubble_time = (end.tv_sec - start.tv_sec) + (end.tv_nsec - start.tv_nsec) / 1e9;

// Test quick sort
clock_gettime(CLOCK_MONOTONIC, &start);
quick_sort(arr2, 0, size - 1);
clock_gettime(CLOCK_MONOTONIC, &end);
double quick_time = (end.tv_sec - start.tv_sec) + (end.tv_nsec - start.tv_nsec) / 1e9;

// Test merge sort
clock_gettime(CLOCK_MONOTONIC, &start);
merge_sort(arr3, 0, size - 1);
clock_gettime(CLOCK_MONOTONIC, &end);
double merge_time = (end.tv_sec - start.tv_sec) + (end.tv_nsec - start.tv_nsec) / 1e9;

printf("Sorting Algorithm Performance (n=%d):\n", size);
printf("  Bubble Sort: %.6f seconds\n", bubble_time);
printf("  Quick Sort:  %.6f seconds\n", quick_time);
printf("  Merge Sort:  %.6f seconds\n", merge_time);

free(arr1);
free(arr2);
free(arr3);
}

```

Search Algorithm Optimization

Choosing the right search algorithm can dramatically improve performance:

```

#include <stdio.h>
#include <stdlib.h>
#include <time.h>

// O(n) linear search
int linear_search(int* arr, int n, int target) {
    for (int i = 0; i < n; i++) {
        if (arr[i] == target) {
            return i;
        }
    }
    return -1;
}

// O(log n) binary search (requires sorted array)
int binary_search(int* arr, int n, int target) {
    int left = 0;

```

```

    int right = n - 1;

    while (left <= right) {
        int mid = left + (right - left) / 2;

        if (arr[mid] == target) {
            return mid;
        }

        if (arr[mid] < target) {
            left = mid + 1;
        } else {
            right = mid - 1;
        }
    }

    return -1;
}

// Hash table implementation for O(1) average case search
#define HASH_TABLE_SIZE 1009

typedef struct hash_node {
    int key;
    int value;
    struct hash_node* next;
} hash_node_t;

static hash_node_t* hash_table[HASH_TABLE_SIZE];

unsigned int hash_function(int key) {
    return (unsigned int)(key % HASH_TABLE_SIZE);
}

void hash_table_insert(int key, int value) {
    unsigned int index = hash_function(key);

    hash_node_t* node = malloc(sizeof(hash_node_t));
    node->key = key;
    node->value = value;
    node->next = hash_table[index];
    hash_table[index] = node;
}

int hash_table_search(int key) {
    unsigned int index = hash_function(key);
    hash_node_t* node = hash_table[index];

```



```

    while (node != NULL) {
        if (node->key == key) {
            return node->value;
        }
        node = node->next;
    }

    return -1;
}

void hash_table_clear(void) {
    for (int i = 0; i < HASH_TABLE_SIZE; i++) {
        hash_node_t* node = hash_table[i];
        while (node != NULL) {
            hash_node_t* next = node->next;
            free(node);
            node = next;
        }
        hash_table[i] = NULL;
    }
}

// Performance comparison
void compare_search_algorithms(void) {
    const int size = 100000;
    int* sorted_array = malloc(size * sizeof(int));
    int* unsorted_array = malloc(size * sizeof(int));

    // Initialize arrays
    for (int i = 0; i < size; i++) {
        sorted_array[i] = i * 2; // Even numbers
        unsorted_array[i] = rand() % (size * 2);
        hash_table_insert(unsorted_array[i], i);
    }

    int target = sorted_array[size / 2]; // Target in middle

    struct timespec start, end;

    // Test linear search
    clock_gettime(CLOCK_MONOTONIC, &start);
    int linear_result = linear_search(unsorted_array, size, target);
    clock_gettime(CLOCK_MONOTONIC, &end);
    double linear_time = (end.tv_sec - start.tv_sec) + (end.tv_nsec - start.tv_nsec) / 1e9;

    // Test binary search

```

```

clock_gettime(CLOCK_MONOTONIC, &start);
int binary_result = binary_search(sorted_array, size, target);
clock_gettime(CLOCK_MONOTONIC, &end);
double binary_time = (end.tv_sec - start.tv_sec) + (end.tv_nsec - start.tv_nsec) / 1e9;

// Test hash table search
clock_gettime(CLOCK_MONOTONIC, &start);
int hash_result = hash_table_search(target);
clock_gettime(CLOCK_MONOTONIC, &end);
double hash_time = (end.tv_sec - start.tv_sec) + (end.tv_nsec - start.tv_nsec) / 1e9;

printf("Search Algorithm Performance (n=%d):\n", size);
printf("  Linear Search: %.9f seconds (result: %d)\n", linear_time, linear_result);
printf("  Binary Search: %.9f seconds (result: %d)\n", binary_time, binary_result);
printf("  Hash Search:   %.9f seconds (result: %d)\n", hash_time, hash_result);

free(sorted_array);
free(unsorted_array);
hash_table_clear();
}

```

Loop Optimization

Loop optimization techniques can significantly improve performance by reducing overhead and improving cache locality.

Loop Unrolling

Loop unrolling reduces loop overhead by processing multiple iterations in each loop cycle:

```

#include <stdio.h>
#include <stdlib.h>
#include <time.h>

// Standard loop
void standard_loop(int* arr, int n) {
    for (int i = 0; i < n; i++) {
        arr[i] = arr[i] * 2 + 1;
    }
}

// Unrolled loop (factor of 4)
void unrolled_loop(int* arr, int n) {
    int i;

```

```

// Process 4 elements per iteration
for (i = 0; i < n - 3; i += 4) {
    arr[i] = arr[i] * 2 + 1;
    arr[i + 1] = arr[i + 1] * 2 + 1;
    arr[i + 2] = arr[i + 2] * 2 + 1;
    arr[i + 3] = arr[i + 3] * 2 + 1;
}

// Handle remaining elements
for (; i < n; i++) {
    arr[i] = arr[i] * 2 + 1;
}

}

// Duff's device (advanced loop unrolling)
void duff_device(int* to, int* from, int count) {
    int n = (count + 7) / 8;
    switch (count % 8) {
        case 0: do { *to++ = *from++;
        case 7:      *to++ = *from++;
        case 6:      *to++ = *from++;
        case 5:      *to++ = *from++;
        case 4:      *to++ = *from++;
        case 3:      *to++ = *from++;
        case 2:      *to++ = *from++;
        case 1:      *to++ = *from++;
                    } while (--n > 0);
    }
}

// Performance comparison
void compare_loop_optimizations(void) {
    const int size = 1000000;
    int* arr1 = malloc(size * sizeof(int));
    int* arr2 = malloc(size * sizeof(int));

    // Initialize arrays
    for (int i = 0; i < size; i++) {
        arr1[i] = arr2[i] = i;
    }

    struct timespec start, end;

    // Test standard loop
    clock_gettime(CLOCK_MONOTONIC, &start);
    standard_loop(arr1, size);

```

```

clock_gettime(CLOCK_MONOTONIC, &end);
double standard_time = (end.tv_sec - start.tv_sec) + (end.tv_nsec - start.tv_nsec) / 1e9

// Test unrolled loop
clock_gettime(CLOCK_MONOTONIC, &start);
unrolled_loop(arr2, size);
clock_gettime(CLOCK_MONOTONIC, &end);
double unrolled_time = (end.tv_sec - start.tv_sec) + (end.tv_nsec - start.tv_nsec) / 1e9

printf("Loop Optimization Performance (n=%d):\n", size);
printf("  Standard Loop: %.6f seconds\n", standard_time);
printf("  Unrolled Loop: %.6f seconds\n", unrolled_time);
printf("  Speedup: %.2fx\n", standard_time / unrolled_time);

free(arr1);
free(arr2);
}

```

Loop Fusion and Fission

Combining or splitting loops can improve cache performance:

```

#include <stdio.h>
#include <stdlib.h>
#include <time.h>

#define SIZE 1000000

// Loop fusion example
void fused_loops(float* a, float* b, float* c, float* d, int n) {
    for (int i = 0; i < n; i++) {
        a[i] = b[i] + c[i]; // First operation
        d[i] = a[i] * 2.0f; // Second operation (depends on first)
    }
}

// Loop fission example
void fissioned_loops(float* a, float* b, float* c, float* d, int n) {
    // First loop
    for (int i = 0; i < n; i++) {
        a[i] = b[i] + c[i];
    }

    // Second loop
    for (int i = 0; i < n; i++) {
        d[i] = a[i] * 2.0f;
    }
}

```

```

    }
}

// Cache-friendly loop fusion
void cache_friendly_fusion(float* a, float* b, float* c, float* d, int n) {
    const int block_size = 1024; // Cache line friendly

    for (int block = 0; block < n; block += block_size) {
        int end = (block + block_size < n) ? block + block_size : n;

        // Process block
        for (int i = block; i < end; i++) {
            a[i] = b[i] + c[i];
            d[i] = a[i] * 2.0f;
        }
    }
}

// Performance comparison
void compare_loop_transformations(void) {
    float* a1 = malloc(SIZE * sizeof(float));
    float* b1 = malloc(SIZE * sizeof(float));
    float* c1 = malloc(SIZE * sizeof(float));
    float* d1 = malloc(SIZE * sizeof(float));

    float* a2 = malloc(SIZE * sizeof(float));
    float* b2 = malloc(SIZE * sizeof(float));
    float* c2 = malloc(SIZE * sizeof(float));
    float* d2 = malloc(SIZE * sizeof(float));

    float* a3 = malloc(SIZE * sizeof(float));
    float* b3 = malloc(SIZE * sizeof(float));
    float* c3 = malloc(SIZE * sizeof(float));
    float* d3 = malloc(SIZE * sizeof(float));

    // Initialize arrays
    for (int i = 0; i < SIZE; i++) {
        b1[i] = b2[i] = b3[i] = (float)i;
        c1[i] = c2[i] = c3[i] = (float)(i * 2);
    }

    struct timespec start, end;

    // Test fused loops
    clock_gettime(CLOCK_MONOTONIC, &start);
    fused_loops(a1, b1, c1, d1, SIZE);

```

```

clock_gettime(CLOCK_MONOTONIC, &end);
double fused_time = (end.tv_sec - start.tv_sec) + (end.tv_nsec - start.tv_nsec) / 1e9;

// Test fissioned loops
clock_gettime(CLOCK_MONOTONIC, &start);
fissioned_loops(a2, b2, c2, d2, SIZE);
clock_gettime(CLOCK_MONOTONIC, &end);
double fissioned_time = (end.tv_sec - start.tv_sec) + (end.tv_nsec - start.tv_nsec) / 1e9;

// Test cache-friendly fusion
clock_gettime(CLOCK_MONOTONIC, &start);
cache_friendly_fusion(a3, b3, c3, d3, SIZE);
clock_gettime(CLOCK_MONOTONIC, &end);
double cache_friendly_time = (end.tv_sec - start.tv_sec) + (end.tv_nsec - start.tv_nsec) / 1e9;

printf("Loop Transformation Performance (n=%d):\n", SIZE);
printf("  Fused Loops:          %.6f seconds\n", fused_time);
printf("  Fissioned Loops:       %.6f seconds\n", fissioned_time);
printf("  Cache-Friendly:        %.6f seconds\n", cache_friendly_time);

free(a1); free(b1); free(c1); free(d1);
free(a2); free(b2); free(c2); free(d2);
free(a3); free(b3); free(c3); free(d3);
}

```

Data Structure Optimization

Choosing the right data structures can have a dramatic impact on performance.

Array vs Linked List

Different data structures have different performance characteristics:

```

#include <stdio.h>
#include <stdlib.h>
#include <time.h>

// Array-based implementation
typedef struct {
    int* data;
    int size;
    int capacity;
} array_list_t;

```

```

array_list_t* array_list_create(int initial_capacity) {
    array_list_t* list = malloc(sizeof(array_list_t));
    list->data = malloc(initial_capacity * sizeof(int));
    list->size = 0;
    list->capacity = initial_capacity;
    return list;
}

void array_list_append(array_list_t* list, int value) {
    if (list->size >= list->capacity) {
        list->capacity *= 2;
        list->data = realloc(list->data, list->capacity * sizeof(int));
    }
    list->data[list->size++] = value;
}

int array_list_get(array_list_t* list, int index) {
    if (index >= 0 && index < list->size) {
        return list->data[index];
    }
    return -1;
}

void array_list_free(array_list_t* list) {
    free(list->data);
    free(list);
}

// Linked list implementation
typedef struct node {
    int data;
    struct node* next;
} node_t;

typedef struct {
    node_t* head;
    int size;
} linked_list_t;

linked_list_t* linked_list_create(void) {
    linked_list_t* list = malloc(sizeof(linked_list_t));
    list->head = NULL;
    list->size = 0;
    return list;
}

void linked_list_append(linked_list_t* list, int value) {

```

```

node_t* new_node = malloc(sizeof(node_t));
new_node->data = value;
new_node->next = NULL;

if (list->head == NULL) {
    list->head = new_node;
} else {
    node_t* current = list->head;
    while (current->next != NULL) {
        current = current->next;
    }
    current->next = new_node;
}
list->size++;
}

int linked_list_get(linked_list_t* list, int index) {
    node_t* current = list->head;
    for (int i = 0; i < index && current != NULL; i++) {
        current = current->next;
    }
    return (current != NULL) ? current->data : -1;
}

void linked_list_free(linked_list_t* list) {
    node_t* current = list->head;
    while (current != NULL) {
        node_t* next = current->next;
        free(current);
        current = next;
    }
    free(list);
}

// Performance comparison
void compare_data_structures(void) {
    const int size = 100000;

    struct timespec start, end;

    // Test array list
    clock_gettime(CLOCK_MONOTONIC, &start);
    array_list_t* arr_list = array_list_create(1000);
    for (int i = 0; i < size; i++) {
        array_list_append(arr_list, i);
    }
    for (int i = 0; i < size; i += 100) {

```



```

        array_list_get(arr_list, i);
    }
    clock_gettime(CLOCK_MONOTONIC, &end);
    double array_time = (end.tv_sec - start.tv_sec) + (end.tv_nsec - start.tv_nsec) / 1e9;

    // Test linked list
    clock_gettime(CLOCK_MONOTONIC, &start);
    linked_list_t* link_list = linked_list_create();
    for (int i = 0; i < size; i++) {
        linked_list_append(link_list, i);
    }
    for (int i = 0; i < size; i += 100) {
        linked_list_get(link_list, i);
    }
    clock_gettime(CLOCK_MONOTONIC, &end);
    double linked_time = (end.tv_sec - start.tv_sec) + (end.tv_nsec - start.tv_nsec) / 1e9;

    printf("Data Structure Performance (n=%d):\n", size);
    printf("  Array List:  %.6f seconds\n", array_time);
    printf("  Linked List:  %.6f seconds\n", linked_time);
    printf("  Ratio:  %.2fx\n", linked_time / array_time);

    array_list_free(arr_list);
    linked_list_free(link_list);
}

```

Compiler Optimization Techniques

Understanding and leveraging compiler optimizations is crucial for performance.

Compiler Flags and Pragmas

```

#include <stdio.h>
#include <stdlib.h>

// Function inlining hints
static inline int fast_add(int a, int b) {
    return a + b;
}

// Restrict keyword for pointer aliasing
void vector_add(float* restrict a, float* restrict b, float* restrict c, int n) {
    for (int i = 0; i < n; i++) {
        c[i] = a[i] + b[i];
    }
}

```

```

    }
}

// Loop unrolling pragma
void optimized_loop(float* arr, int n) {
    #pragma GCC unroll 4
    for (int i = 0; i < n; i++) {
        arr[i] = arr[i] * 2.0f + 1.0f;
    }
}

// Vectorization hints
void vectorized_operation(float* arr, int n) {
    #pragma GCC ivdep
    for (int i = 0; i < n; i++) {
        arr[i] = arr[i] * arr[i] + 2.0f * arr[i] + 1.0f;
    }
}

// Branch prediction hints
void branch_prediction_example(int* arr, int n) {
    for (int i = 0; i < n; i++) {
        if (__builtin_expect(arr[i] > 0, 1)) { // Likely true
            arr[i] *= 2;
        } else {
            arr[i] = 0;
        }
    }
}

// Memory alignment
struct aligned_struct {
    int a;
    double b;
    char c;
} __attribute__((aligned(16)));

// Example usage
void compiler_optimization_examples(void) {
    const int size = 1000000;
    float* a = malloc(size * sizeof(float));
    float* b = malloc(size * sizeof(float));
    float* c = malloc(size * sizeof(float));

    // Initialize arrays
    for (int i = 0; i < size; i++) {
        a[i] = b[i] = (float)i;
    }
}

```

```

    }

    // Test optimized operations
    vector_add(a, b, c, size);
    optimized_loop(c, size);
    vectorized_operation(c, size);

    printf("Compiler optimization examples completed\n");

    free(a);
    free(b);
    free(c);
}

```

Mathematical Optimization

Mathematical optimizations can significantly improve performance in numerical computations.

Fast Mathematical Operations

```

#include <stdio.h>
#include <stdlib.h>
#include <math.h>
#include <time.h>

// Fast integer square root (Newton's method)
int fast_sqrt(int x) {
    if (x == 0) return 0;

    int guess = x;
    int better_guess;

    do {
        better_guess = (guess + x / guess) / 2;
        if (better_guess >= guess) break;
        guess = better_guess;
    } while (1);

    return guess;
}

// Fast integer power
int fast_power(int base, int exp) {
    int result = 1;

```

```

    while (exp > 0) {
        if (exp & 1) {
            result *= base;
        }
        base *= base;
        exp >>= 1;
    }
    return result;
}

// Bit manipulation for multiplication by powers of 2
int multiply_by_power_of_2(int value, int power) {
    return value << power; // Equivalent to value * (2^power)
}

// Bit manipulation for division by powers of 2
int divide_by_power_of_2(int value, int power) {
    return value >> power; // Equivalent to value / (2^power) for positive numbers
}

// Fast absolute value
int fast_abs(int x) {
    int mask = x >> (sizeof(int) * 8 - 1);
    return (x + mask) ^ mask;
}

// Performance comparison
void compare_mathematical_optimizations(void) {
    const int iterations = 1000000;

    struct timespec start, end;

    // Test standard sqrt
    clock_gettime(CLOCK_MONOTONIC, &start);
    volatile double sum1 = 0;
    for (int i = 1; i <= iterations; i++) {
        sum1 += sqrt(i);
    }
    clock_gettime(CLOCK_MONOTONIC, &end);
    double sqrt_time = (end.tv_sec - start.tv_sec) + (end.tv_nsec - start.tv_nsec) / 1e9;

    // Test fast sqrt
    clock_gettime(CLOCK_MONOTONIC, &start);
    volatile int sum2 = 0;
    for (int i = 1; i <= iterations; i++) {
        sum2 += fast_sqrt(i);
    }
}

```

```

clock_gettime(CLOCK_MONOTONIC, &end);
double fast_sqrt_time = (end.tv_sec - start.tv_sec) + (end.tv_nsec - start.tv_nsec) / 1e9;

// Test standard power
clock_gettime(CLOCK_MONOTONIC, &start);
volatile int sum3 = 0;
for (int i = 1; i <= 1000; i++) {
    sum3 += (int)pow(2, i % 10);
}
clock_gettime(CLOCK_MONOTONIC, &end);
double power_time = (end.tv_sec - start.tv_sec) + (end.tv_nsec - start.tv_nsec) / 1e9;

// Test fast power
clock_gettime(CLOCK_MONOTONIC, &start);
volatile int sum4 = 0;
for (int i = 1; i <= 1000; i++) {
    sum4 += fast_power(2, i % 10);
}
clock_gettime(CLOCK_MONOTONIC, &end);
double fast_power_time = (end.tv_sec - start.tv_sec) + (end.tv_nsec - start.tv_nsec) / 1e9;

printf("Mathematical Optimization Performance:\n");
printf("  Standard sqrt: %.6f seconds\n", sqrt_time);
printf("  Fast sqrt:      %.6f seconds\n", fast_sqrt_time);
printf("  Standard power: %.6f seconds\n", power_time);
printf("  Fast power:     %.6f seconds\n", fast_power_time);
}

```

Practical Examples

Image Processing Optimization

```

#include <stdio.h>
#include <stdlib.h>
#include <time.h>

#define IMAGE_WIDTH 1920
#define IMAGE_HEIGHT 1080

// Simple image structure
typedef struct {
    unsigned char* data;
    int width;
    int height;
}

```

```

} image_t;

// Create image
image_t* create_image(int width, int height) {
    image_t* img = malloc(sizeof(image_t));
    img->data = malloc(width * height * sizeof(unsigned char));
    img->width = width;
    img->height = height;
    return img;
}

// Free image
void free_image(image_t* img) {
    free(img->data);
    free(img);
}

// Naive blur implementation
void naive_blur(image_t* src, image_t* dst) {
    for (int y = 1; y < src->height - 1; y++) {
        for (int x = 1; x < src->width - 1; x++) {
            int sum = 0;
            for (int ky = -1; ky <= 1; ky++) {
                for (int kx = -1; kx <= 1; kx++) {
                    sum += src->data[(y + ky) * src->width + (x + kx)];
                }
            }
            dst->data[y * src->width + x] = sum / 9;
        }
    }
}

// Optimized blur with loop unrolling
void optimized_blur(image_t* src, image_t* dst) {
    for (int y = 1; y < src->height - 1; y++) {
        for (int x = 1; x < src->width - 1; x++) {
            // Unroll the inner loops
            int sum = src->data[(y-1) * src->width + (x-1)] +
                src->data[(y-1) * src->width + x] +
                src->data[(y-1) * src->width + (x+1)] +
                src->data[y * src->width + (x-1)] +
                src->data[y * src->width + x] +
                src->data[y * src->width + (x+1)] +
                src->data[(y+1) * src->width + (x-1)] +
                src->data[(y+1) * src->width + x] +
                src->data[(y+1) * src->width + (x+1)];
        }
    }
}

```

```

        dst->data[y * src->width + x] = sum / 9;
    }
}

// SIMD-style optimization using multiple operations per loop
void simd_style_blur(image_t* src, image_t* dst) {
    for (int y = 1; y < src->height - 1; y++) {
        // Process 4 pixels at a time
        for (int x = 1; x < src->width - 4; x += 4) {
            // Process 4 pixels in parallel (simplified)
            for (int i = 0; i < 4; i++) {
                int sum = src->data[(y-1) * src->width + (x+i-1)] +
                    src->data[(y-1) * src->width + (x+i)] +
                    src->data[(y-1) * src->width + (x+i+1)] +
                    src->data[y * src->width + (x+i-1)] +
                    src->data[y * src->width + (x+i)] +
                    src->data[y * src->width + (x+i+1)] +
                    src->data[(y+1) * src->width + (x+i-1)] +
                    src->data[(y+1) * src->width + (x+i)] +
                    src->data[(y+1) * src->width + (x+i+1)];
                dst->data[y * src->width + (x+i)] = sum / 9;
            }
        }
    }
}

// Performance comparison
void compare_image_processing(void) {
    image_t* src = create_image(IMAGE_WIDTH, IMAGE_HEIGHT);
    image_t* dst1 = create_image(IMAGE_WIDTH, IMAGE_HEIGHT);
    image_t* dst2 = create_image(IMAGE_WIDTH, IMAGE_HEIGHT);
    image_t* dst3 = create_image(IMAGE_WIDTH, IMAGE_HEIGHT);

    // Initialize source image with random data
    for (int i = 0; i < IMAGE_WIDTH * IMAGE_HEIGHT; i++) {
        src->data[i] = rand() % 256;
    }

    struct timespec start, end;

    // Test naive blur
    clock_gettime(CLOCK_MONOTONIC, &start);
    naive_blur(src, dst1);
    clock_gettime(CLOCK_MONOTONIC, &end);
    double naive_time = (end.tv_sec - start.tv_sec) + (end.tv_nsec - start.tv_nsec) / 1e9;

```

```

// Test optimized blur
clock_gettime(CLOCK_MONOTONIC, &start);
optimized_blur(src, dst2);
clock_gettime(CLOCK_MONOTONIC, &end);
double optimized_time = (end.tv_sec - start.tv_sec) + (end.tv_nsec - start.tv_nsec) / 1e9;

// Test SIMD-style blur
clock_gettime(CLOCK_MONOTONIC, &start);
simd_style_blur(src, dst3);
clock_gettime(CLOCK_MONOTONIC, &end);
double simd_time = (end.tv_sec - start.tv_sec) + (end.tv_nsec - start.tv_nsec) / 1e9;

printf("Image Processing Performance (%dx%d):\n", IMAGE_WIDTH, IMAGE_HEIGHT);
printf("  Naive Blur:      %.6f seconds\n", naive_time);
printf("  Optimized Blur:    %.6f seconds\n", optimized_time);
printf("  SIMD-style:       %.6f seconds\n", simd_time);
printf("  Speedup:          %.2fx\n", naive_time / simd_time);

free_image(src);
free_image(dst1);
free_image(dst2);
free_image(dst3);
}

```

Summary

Optimization techniques in C programming encompass a wide range of approaches:

1. **Algorithmic Optimization** - Choosing better algorithms and data structures
2. **Loop Optimization** - Unrolling, fusion, and restructuring loops
3. **Data Structure Optimization** - Selecting appropriate data structures for the task
4. **Compiler Optimization** - Leveraging compiler features and flags
5. **Mathematical Optimization** - Fast mathematical operations and bit manipulation
6. **Memory Access Optimization** - Improving cache locality and reducing memory overhead

Key principles for effective optimization: - Profile first to identify actual bottlenecks - Focus on algorithmic improvements before micro-optimizations - Understand the trade-offs between performance and maintainability - Use appropriate compiler optimization flags - Consider the target hardware architecture - Test optimizations thoroughly to ensure correctness - Document optimization decisions for future maintenance

These techniques enable developers to create high-performance C applications that efficiently utilize system resources while maintaining code quality and readability.

Memory Optimization

Introduction

Memory optimization is a critical aspect of performance tuning in C programs, especially in resource-constrained environments and high-performance applications. Efficient memory management can significantly impact program performance through improved cache utilization, reduced allocation overhead, and better data locality. This chapter explores techniques for optimizing memory usage, from allocation strategies to cache-friendly data structures.

Memory Hierarchy and Performance

Understanding the memory hierarchy is fundamental to effective memory optimization. Modern computer systems have multiple levels of memory with vastly different access times.

Memory Latency Comparison

```
#include <stdio.h>
#include <time.h>
#include <stdlib.h>

// Example illustrating memory access time differences
void memory_latency_example(void) {
    printf("Typical Access Times:\n");
    printf("- CPU Register: ~0.3 ns\n");
    printf("- L1 Cache:      ~1 ns\n");
    printf("- L2 Cache:      ~4 ns\n");
    printf("- L3 Cache:      ~15 ns\n");
    printf("- Main Memory:   ~100 ns\n");
    printf("- SSD:          ~10 s\n");
    printf("- HDD:          ~10 ms\n");
}
```

Cache Behavior Impact

Cache misses can cause significant performance degradation. Understanding how data is accessed affects cache performance:

```

#include <stdio.h>
#include <stdlib.h>
#include <time.h>

#define MATRIX_SIZE 1024
#define ITERATIONS 100

// Row-major access (cache-friendly)
double row_major_sum(double matrix[MATRIX_SIZE][MATRIX_SIZE]) {
    double sum = 0.0;
    for (int i = 0; i < MATRIX_SIZE; i++) {
        for (int j = 0; j < MATRIX_SIZE; j++) {
            sum += matrix[i][j];
        }
    }
    return sum;
}

// Column-major access (cache-unfriendly)
double column_major_sum(double matrix[MATRIX_SIZE][MATRIX_SIZE]) {
    double sum = 0.0;
    for (int j = 0; j < MATRIX_SIZE; j++) {
        for (int i = 0; i < MATRIX_SIZE; i++) {
            sum += matrix[i][j];
        }
    }
    return sum;
}

void cache_performance_comparison(void) {
    double (*matrix)[MATRIX_SIZE] = malloc(MATRIX_SIZE * sizeof(*matrix));

    // Initialize matrix
    for (int i = 0; i < MATRIX_SIZE; i++) {
        for (int j = 0; j < MATRIX_SIZE; j++) {
            matrix[i][j] = (double)(i * MATRIX_SIZE + j);
        }
    }

    clock_t start, end;

    // Row-major timing
    start = clock();
    for (int i = 0; i < ITERATIONS; i++) {
        volatile double sum = row_major_sum(matrix);
    }
    end = clock();
}

```

```

double row_time = ((double)(end - start)) / CLOCKS_PER_SEC;

// Column-major timing
start = clock();
for (int i = 0; i < ITERATIONS; i++) {
    volatile double sum = column_major_sum(matrix);
}
end = clock();
double col_time = ((double)(end - start)) / CLOCKS_PER_SEC;

printf("Row-major time: %.4f seconds\n", row_time);
printf("Column-major time: %.4f seconds\n", col_time);
printf("Performance difference: %.2fx\n", col_time / row_time);

free(matrix);
}

```

Memory Allocation Optimization

Standard memory allocation functions like `malloc()` (see [system programming](#)) and `free()` (see [linear structures](#)) can be significant performance bottlenecks in performance-critical applications.

Custom Memory Allocators

For specific use cases, custom allocators can provide significant performance improvements:

```

#include <stdio.h>
#include <stdlib.h>
#include <stdint.h>
#include <string.h>

// Simple pool allocator
typedef struct {
    void *memory_pool;
    size_t pool_size;
    size_t block_size;
    size_t num_blocks;
    uint8_t *free_blocks;
    size_t next_free;
} pool_allocator_t;

// Initialize pool allocator
pool_allocator_t* pool_init(size_t block_size, size_t num_blocks) {
    pool_allocator_t *pool = malloc(sizeof(pool_allocator_t));

```

```

    if (!pool) return NULL;

    pool->pool_size = block_size * num_blocks;
    pool->memory_pool = malloc(pool->pool_size);
    if (!pool->memory_pool) {
        free(pool);
        return NULL;
    }

    pool->block_size = block_size;
    pool->num_blocks = num_blocks;
    pool->free_blocks = calloc(num_blocks, sizeof(uint8_t));
    pool->next_free = 0;

    return pool;
}

// Allocate block from pool
void* pool_alloc(pool_allocator_t *pool) {
    if (!pool || pool->next_free >= pool->num_blocks) {
        return NULL; // Pool exhausted
    }

    // Find free block
    for (size_t i = 0; i < pool->num_blocks; i++) {
        size_t index = (pool->next_free + i) % pool->num_blocks;
        if (!pool->free_blocks[index]) {
            pool->free_blocks[index] = 1;
            pool->next_free = (index + 1) % pool->num_blocks;
            return (char*)pool->memory_pool + (index * pool->block_size);
        }
    }

    return NULL; // Should not happen
}

// Free block back to pool
void pool_free(pool_allocator_t *pool, void *ptr) {
    if (!pool || !ptr) return;

    // Calculate block index
    ptrdiff_t offset = (char*)ptr - (char*)pool->memory_pool;
    if (offset < 0 || offset >= (ptrdiff_t)pool->pool_size) return;

    size_t index = offset / pool->block_size;
    if (index < pool->num_blocks) {
        pool->free_blocks[index] = 0;
    }
}

```

```

        pool->next_free = index; // Prefer this block for next allocation
    }
}

// Destroy pool allocator
void pool_destroy(pool_allocator_t *pool) {
    if (pool) {
        free(pool->memory_pool);
        free(pool->free_blocks);
        free(pool);
    }
}

// Example usage
void pool_allocator_example(void) {
    // Create pool for 100 blocks of 64 bytes each
    pool_allocator_t *pool = pool_init(64, 100);
    if (!pool) {
        printf("Failed to create pool\n");
        return;
    }

    // Allocate some blocks
    void *ptr1 = pool_alloc(pool);
    void *ptr2 = pool_alloc(pool);
    void *ptr3 = pool_alloc(pool);

    printf("Allocated 3 blocks\n");

    // Free some blocks
    pool_free(pool, ptr1);
    pool_free(pool, ptr3);

    printf("Freed 2 blocks\n");

    // Allocate more blocks (should reuse freed ones)
    void *ptr4 = pool_alloc(pool);
    void *ptr5 = pool_alloc(pool);

    printf("Allocated 2 more blocks\n");

    // Cleanup
    pool_destroy(pool);
}

```

Stack Allocation vs Heap Allocation

Stack allocation is significantly faster than heap allocation but has size and lifetime limitations:

```
#include <stdio.h>
#include <stdlib.h>
#include <time.h>

#define ARRAY_SIZE 10000
#define ITERATIONS 100000

// Stack allocation
void stack_allocation_test(void) {
    // Fast allocation on stack
    int array[ARRAY_SIZE];

    // Initialize array
    for (int i = 0; i < ARRAY_SIZE; i++) {
        array[i] = i;
    }
}

// Heap allocation
void heap_allocation_test(void) {
    // Slower allocation on heap
    int *array = malloc(ARRAY_SIZE * sizeof(int));
    if (array) {
        // Initialize array
        for (int i = 0; i < ARRAY_SIZE; i++) {
            array[i] = i;
        }
        free(array);
    }
}

void allocation_performance_comparison(void) {
    clock_t start, end;

    // Stack allocation timing
    start = clock();
    for (int i = 0; i < ITERATIONS; i++) {
        stack_allocation_test();
    }
    end = clock();
    double stack_time = ((double)(end - start)) / CLOCKS_PER_SEC;

    // Heap allocation timing
```

```

    start = clock();
    for (int i = 0; i < ITERATIONS; i++) {
        heap_allocation_test();
    }
    end = clock();
    double heap_time = ((double)(end - start)) / CLOCKS_PER_SEC;

    printf("Stack allocation time: %.4f seconds\n", stack_time);
    printf("Heap allocation time: %.4f seconds\n", heap_time);
    printf("Stack is %.2fx faster\n", heap_time / stack_time);
}

```

Data Structure Optimization

Choosing the right data structures and organizing them for cache efficiency can dramatically improve performance.

Structure of Arrays vs Array of Structures

The way data is organized in memory can significantly impact cache performance:

```

#include <stdio.h>
#include <stdlib.h>
#include <time.h>

#define NUM_POINTS 1000000

// Array of Structures (AoS) - cache-unfriendly for component-wise operations
typedef struct {
    float x, y, z;
} point_aos_t;

// Structure of Arrays (SoA) - cache-friendly for component-wise operations
typedef struct {
    float *x, *y, *z;
} point_soa_t;

// Initialize AoS
point_aos_t* init_aos(int count) {
    point_aos_t *points = malloc(count * sizeof(point_aos_t));
    for (int i = 0; i < count; i++) {
        points[i].x = (float)i;
        points[i].y = (float)(i * 2);
        points[i].z = (float)(i * 3);
    }
}

```

```

    }
    return points;
}

// Initialize SoA
point_soa_t* init_soa(int count) {
    point_soa_t *points = malloc(sizeof(point_soa_t));
    points->x = malloc(count * sizeof(float));
    points->y = malloc(count * sizeof(float));
    points->z = malloc(count * sizeof(float));

    for (int i = 0; i < count; i++) {
        points->x[i] = (float)i;
        points->y[i] = (float)(i * 2);
        points->z[i] = (float)(i * 3);
    }
    return points;
}

// Process all x coordinates (AoS version)
float process_x_aos(point_aos_t *points, int count) {
    float sum = 0.0f;
    for (int i = 0; i < count; i++) {
        sum += points[i].x * 2.0f;
    }
    return sum;
}

// Process all x coordinates (SoA version)
float process_x_soa(point_soa_t *points, int count) {
    float sum = 0.0f;
    for (int i = 0; i < count; i++) {
        sum += points->x[i] * 2.0f;
    }
    return sum;
}

void aos_vs_soa_comparison(void) {
    point_aos_t *aos_points = init_aos(NUM_POINTS);
    point_soa_t *soa_points = init_soa(NUM_POINTS);

    clock_t start, end;

    // AoS timing
    start = clock();
    volatile float aos_result = process_x_aos(aos_points, NUM_POINTS);
    end = clock();

```



```

double aos_time = ((double)(end - start)) / CLOCKS_PER_SEC;

// SoA timing
start = clock();
volatile float soa_result = process_x_soa(soa_points, NUM_POINTS);
end = clock();
double soa_time = ((double)(end - start)) / CLOCKS_PER_SEC;

printf("AoS processing time: %.4f seconds\n", aos_time);
printf("SoA processing time: %.4f seconds\n", soa_time);
printf("SoA is %.2fx faster\n", aos_time / soa_time);

// Cleanup
free(aos_points);
free(soa_points->x);
free(soa_points->y);
free(soa_points->z);
free(soa_points);
}

```

Memory Alignment

Proper memory alignment can improve performance by ensuring efficient memory access:

```

#include <stdio.h>
#include <stdlib.h>
#include <stdalign.h>
#include <time.h>

#define ARRAY_SIZE 1000000

// Unaligned structure
typedef struct {
    char a;        // 1 byte
    int b;         // 4 bytes (3 bytes padding before)
    short c;       // 2 bytes
    // 1 byte padding at end
} unaligned_struct_t;

// Aligned structure
typedef struct {
    alignas(16) int b; // 4 bytes, aligned to 16-byte boundary
    char a;           // 1 byte
    short c;          // 2 bytes
    // 9 bytes padding at end for 16-byte alignment
} aligned_struct_t;

```

```

void alignment_performance_test(void) {
    // Allocate arrays
    unaligned_struct_t *unaligned_array = malloc(ARRAY_SIZE * sizeof(unaligned_struct_t));
    aligned_struct_t *aligned_array = malloc(ARRAY_SIZE * sizeof(aligned_struct_t));

    clock_t start, end;

    // Test unaligned access
    start = clock();
    for (int i = 0; i < ARRAY_SIZE; i++) {
        unaligned_array[i].b = i * 2;
    }
    end = clock();
    double unaligned_time = ((double)(end - start)) / CLOCKS_PER_SEC;

    // Test aligned access
    start = clock();
    for (int i = 0; i < ARRAY_SIZE; i++) {
        aligned_array[i].b = i * 2;
    }
    end = clock();
    double aligned_time = ((double)(end - start)) / CLOCKS_PER_SEC;

    printf("Unaligned access time: %.4f seconds\n", unaligned_time);
    printf("Aligned access time: %.4f seconds\n", aligned_time);
    printf("Aligned is %.2fx faster\n", unaligned_time / aligned_time);

    // Print structure sizes
    printf("Unaligned struct size: %zu bytes\n", sizeof(unaligned_struct_t));
    printf("Aligned struct size: %zu bytes\n", sizeof(aligned_struct_t));

    free(unaligned_array);
    free(aligned_array);
}

```

Memory Layout Optimization

Understanding and optimizing memory layout can lead to significant performance improvements.

Data Locality

Keeping related data close together in memory improves cache performance:

```

#include <stdio.h>
#include <stdlib.h>
#include <time.h>

#define MATRIX_ROWS 1000
#define MATRIX_COLS 1000

// Poor data locality - accessing data with large strides
void poor_locality_example(int **matrix) {
    long sum = 0;
    // Access data column by column (poor cache usage)
    for (int col = 0; col < MATRIX_COLS; col++) {
        for (int row = 0; row < MATRIX_ROWS; row++) {
            sum += matrix[row][col];
        }
    }
    // Prevent optimization
    volatile long dummy = sum;
}

// Good data locality - accessing data sequentially
void good_locality_example(int **matrix) {
    long sum = 0;
    // Access data row by row (good cache usage)
    for (int row = 0; row < MATRIX_ROWS; row++) {
        for (int col = 0; col < MATRIX_COLS; col++) {
            sum += matrix[row][col];
        }
    }
    // Prevent optimization
    volatile long dummy = sum;
}

void data_locality_comparison(void) {
    // Allocate matrix
    int **matrix = malloc(MATRIX_ROWS * sizeof(int*));
    for (int i = 0; i < MATRIX_ROWS; i++) {
        matrix[i] = malloc(MATRIX_COLS * sizeof(int));
        for (int j = 0; j < MATRIX_COLS; j++) {
            matrix[i][j] = i * MATRIX_COLS + j;
        }
    }

    clock_t start, end;

    // Poor locality timing
    start = clock();

```

```

    poor_locality_example(matrix);
    end = clock();
    double poor_time = ((double)(end - start)) / CLOCKS_PER_SEC;

    // Good locality timing
    start = clock();
    good_locality_example(matrix);
    end = clock();
    double good_time = ((double)(end - start)) / CLOCKS_PER_SEC;

    printf("Poor locality time: %.4f seconds\n", poor_time);
    printf("Good locality time: %.4f seconds\n", good_time);
    printf("Good locality is %.2fx faster\n", poor_time / good_time);

    // Cleanup
    for (int i = 0; i < MATRIX_ROWS; i++) {
        free(matrix[i]);
    }
    free(matrix);
}

```

Memory Pooling

Memory pooling can reduce allocation overhead and improve data locality:

```

#include <stdio.h>
#include <stdlib.h>
#include <string.h>

#define POOL_SIZE 1024 * 1024 // 1MB pool
#define BLOCK_SIZE 256

typedef struct memory_block {
    struct memory_block *next;
    size_t size;
    int in_use;
    char data[1]; // Flexible array member
} memory_block_t;

typedef struct {
    char *pool_start;
    char *pool_end;
    memory_block_t *free_list;
    size_t total_size;
} memory_pool_t;

```

```

// Initialize memory pool
memory_pool_t* pool_create(size_t size) {
    memory_pool_t *pool = malloc(sizeof(memory_pool_t));
    if (!pool) return NULL;

    pool->pool_start = malloc(size);
    if (!pool->pool_start) {
        free(pool);
        return NULL;
    }

    pool->pool_end = pool->pool_start + size;
    pool->total_size = size;
    pool->free_list = (memory_block_t*)pool->pool_start;

    // Initialize first block to cover entire pool
    pool->free_list->next = NULL;
    pool->free_list->size = size - sizeof(memory_block_t);
    pool->free_list->in_use = 0;

    return pool;
}

// Allocate memory from pool
void* pool_malloc(memory_pool_t *pool, size_t size) {
    if (!pool || size == 0) return NULL;

    // Add space for block header
    size_t required_size = sizeof(memory_block_t) + size;

    // Find suitable block
    memory_block_t *prev = NULL;
    memory_block_t *current = pool->free_list;

    while (current) {
        if (!current->in_use && current->size >= size) {
            // Found suitable block
            current->in_use = 1;

            // Split block if large enough
            size_t remaining = current->size - size;
            if (remaining > sizeof(memory_block_t) + 32) { // Minimum block size
                memory_block_t *new_block = (memory_block_t*)
                    ((char*)current + sizeof(memory_block_t) + size);
                new_block->size = remaining - sizeof(memory_block_t);
                new_block->in_use = 0;
                new_block->next = current->next;
            }
        }
    }
}

```

```

        current->size = size;
        current->next = new_block;
    }

    return current->data;
}

prev = current;
current = current->next;
}

return NULL; // No suitable block found
}

// Free memory back to pool
void pool_free(memory_pool_t *pool, void *ptr) {
    if (!pool || !ptr) return;

    // Get block header
    memory_block_t *block = (memory_block_t*)((char*)ptr - sizeof(memory_block_t));
    block->in_use = 0;

    // Coalesce adjacent free blocks
    memory_block_t *current = pool->free_list;
    while (current) {
        if (!current->in_use) {
            // Check next block
            memory_block_t *next = current->next;
            if (next && !next->in_use) {
                // Merge with next block
                current->size += sizeof(memory_block_t) + next->size;
                current->next = next->next;
                continue; // Check same block again
            }
        }
        current = current->next;
    }
}

// Destroy memory pool
void pool_destroy(memory_pool_t *pool) {
    if (pool) {
        free(pool->pool_start);
        free(pool);
    }
}

```

```

// Example usage
void memory_pool_example(void) {
    memory_pool_t *pool = pool_create(PPOOL_SIZE);
    if (!pool) {
        printf("Failed to create memory pool\n");
        return;
    }

    printf("Created memory pool of %zu bytes\n", PPOOL_SIZE);

    // Allocate some blocks
    void *ptr1 = pool_malloc(pool, 128);
    void *ptr2 = pool_malloc(pool, 256);
    void *ptr3 = pool_malloc(pool, 64);

    printf("Allocated 3 blocks\n");

    // Free some blocks
    pool_free(pool, ptr1);
    pool_free(pool, ptr3);

    printf("Freed 2 blocks\n");

    // Allocate more blocks (should reuse freed space)
    void *ptr4 = pool_malloc(pool, 96);
    void *ptr5 = pool_malloc(pool, 160);

    printf("Allocated 2 more blocks\n");

    // Cleanup
    pool_destroy(pool);
}

```

Memory Access Patterns

Understanding and optimizing memory access patterns can significantly improve performance.

Sequential vs Random Access

Sequential memory access is much faster than random access due to cache behavior:

```

#include <stdio.h>
#include <stdlib.h>
#include <time.h>

```

```

#define ARRAY_SIZE 1000000
#define ACCESS_COUNT 1000000

void sequential_access_test(int *array) {
    long sum = 0;
    for (int i = 0; i < ARRAY_SIZE; i++) {
        sum += array[i];
    }
    // Prevent optimization
    volatile long dummy = sum;
}

void random_access_test(int *array) {
    long sum = 0;
    // Access elements in random order
    for (int i = 0; i < ARRAY_SIZE; i++) {
        int index = rand() % ARRAY_SIZE;
        sum += array[index];
    }
    // Prevent optimization
    volatile long dummy = sum;
}

void access_pattern_comparison(void) {
    int *array = malloc(ARRAY_SIZE * sizeof(int));
    if (!array) {
        printf("Failed to allocate array\n");
        return;
    }

    // Initialize array
    for (int i = 0; i < ARRAY_SIZE; i++) {
        array[i] = i;
    }

    clock_t start, end;

    // Sequential access timing
    start = clock();
    sequential_access_test(array);
    end = clock();
    double sequential_time = ((double)(end - start)) / CLOCKS_PER_SEC;

    // Random access timing
    srand(42); // Fixed seed for reproducible results
    start = clock();

```



```

    random_access_test(array);
    end = clock();
    double random_time = ((double)(end - start)) / CLOCKS_PER_SEC;

    printf("Sequential access time: %.4f seconds\n", sequential_time);
    printf("Random access time: %.4f seconds\n", random_time);
    printf("Sequential is %.2fx faster\n", random_time / sequential_time);

    free(array);
}

```

Practical Memory Optimization Techniques

Reducing Memory Fragmentation

Memory fragmentation can lead to inefficient memory usage and allocation failures:

```

#include <stdio.h>
#include <stdlib.h>

#define POOL_SIZE 10000

// Simple example showing fragmentation
void fragmentation_example(void) {
    void *ptrs[1000];

    printf("Demonstrating memory fragmentation:\n");

    // Allocate many small blocks
    for (int i = 0; i < 1000; i++) {
        ptrs[i] = malloc(10);
        if (!ptrs[i]) {
            printf("Allocation failed at %d\n", i);
            return;
        }
    }

    // Free every other block
    for (int i = 0; i < 1000; i += 2) {
        free(ptrs[i]);
        ptrs[i] = NULL;
    }

    // Try to allocate a larger block
    void *large_block = malloc(500);
}

```

```

    if (large_block) {
        printf("Large block allocation succeeded\n");
        free(large_block);
    } else {
        printf("Large block allocation failed due to fragmentation\n");
    }

    // Free remaining blocks
    for (int i = 1; i < 1000; i += 2) {
        if (ptrs[i]) {
            free(ptrs[i]);
        }
    }
}

```

Memory Prefetching

Modern processors support prefetching instructions to improve memory access performance:

```

#include <stdio.h>
#include <stdlib.h>
#include <time.h>

#define ARRAY_SIZE 1000000

// Manual prefetching example
void prefetch_example(int *array) {
    long sum = 0;

    for (int i = 0; i < ARRAY_SIZE; i++) {
        // Prefetch next cache line (typically 64 bytes)
        if (i + 16 < ARRAY_SIZE) {
            __builtin_prefetch(&array[i + 16], 0, 3);
        }

        sum += array[i];
    }

    // Prevent optimization
    volatile long dummy = sum;
}

// Without prefetching
void no_prefetch_example(int *array) {
    long sum = 0;

```

```

    for (int i = 0; i < ARRAY_SIZE; i++) {
        sum += array[i];
    }

    // Prevent optimization
    volatile long dummy = sum;
}

void prefetch_comparison(void) {
    int *array = malloc(ARRAY_SIZE * sizeof(int));
    if (!array) {
        printf("Failed to allocate array\n");
        return;
    }

    // Initialize array
    for (int i = 0; i < ARRAY_SIZE; i++) {
        array[i] = i;
    }

    clock_t start, end;

    // Without prefetching
    start = clock();
    no_prefetch_example(array);
    end = clock();
    double no_prefetch_time = ((double)(end - start)) / CLOCKS_PER_SEC;

    // With prefetching
    start = clock();
    prefetch_example(array);
    end = clock();
    double prefetch_time = ((double)(end - start)) / CLOCKS_PER_SEC;

    printf("Without prefetching: %.4f seconds\n", no_prefetch_time);
    printf("With prefetching: %.4f seconds\n", prefetch_time);
    printf("Prefetching improvement: %.2fx\n", no_prefetch_time / prefetch_time);

    free(array);
}

```

Memory Debugging and Analysis

Tools and techniques for analyzing memory usage and identifying optimization opportunities.

Simple Memory Tracker

A basic memory tracking system to monitor allocations:

```
#include <stdio.h>
#include <stdlib.h>
#include <string.h>

// Simple memory tracker
typedef struct mem_block {
    struct mem_block *next;
    void *ptr;
    size_t size;
    const char *file;
    int line;
} mem_block_t;

static mem_block_t *mem_list = NULL;
static size_t total_allocated = 0;
static size_t peak_allocated = 0;
static size_t allocation_count = 0;

// Track memory allocation
void* tracked_malloc(size_t size, const char *file, int line) {
    void *ptr = malloc(size);
    if (!ptr) return NULL;

    // Add to tracking list
    mem_block_t *block = malloc(sizeof(mem_block_t));
    if (!block) {
        free(ptr);
        return NULL;
    }

    block->ptr = ptr;
    block->size = size;
    block->file = file;
    block->line = line;
    block->next = mem_list;
    mem_list = block;

    total_allocated += size;
    allocation_count++;

    if (total_allocated > peak_allocated) {
        peak_allocated = total_allocated;
    }
}
```

```

    return ptr;
}

// Track memory deallocation
void tracked_free(void *ptr, const char *file, int line) {
    if (!ptr) return;

    // Find and remove from tracking list
    mem_block_t *prev = NULL;
    mem_block_t *current = mem_list;

    while (current) {
        if (current->ptr == ptr) {
            if (prev) {
                prev->next = current->next;
            } else {
                mem_list = current->next;
            }

            total_allocated -= current->size;
            free(current);
            free(ptr);
            return;
        }

        prev = current;
        current = current->next;
    }

    // Not found in tracking list - possible double free
    printf("Warning: Attempt to free untracked pointer at %s:%d\n", file, line);
    free(ptr);
}

// Print memory statistics
void print_memory_stats(void) {
    printf("\n=== Memory Statistics ===\n");
    printf("Total allocations: %zu\n", allocation_count);
    printf("Currently allocated: %zu bytes\n", total_allocated);
    printf("Peak allocation: %zu bytes\n", peak_allocated);

    if (mem_list) {
        printf("Leaked blocks:\n");
        mem_block_t *current = mem_list;
        while (current) {
            printf("  %p: %zu bytes at %s:%d\n",

```

```

        current->ptr, current->size, current->file, current->line);
        current = current->next;
    }
} else {
    printf("No memory leaks detected\n");
}
}

// Cleanup tracking (call at program exit)
void cleanup_memory_tracking(void) {
    mem_block_t *current = mem_list;
    while (current) {
        mem_block_t *next = current->next;
        free(current->ptr);
        free(current);
        current = next;
    }
    mem_list = NULL;
}

// Macros for easier usage
#define MALLOC(size) tracked_malloc(size, __FILE__, __LINE__)
#define FREE(ptr) tracked_free(ptr, __FILE__, __LINE__)

// Example usage
void memory_tracking_example(void) {
    printf("Memory tracking example:\n");

    // Allocate some memory
    int *array1 = (int*)MALLOC(100 * sizeof(int));
    char *string = (char*)MALLOC(50);

    // Use the memory
    for (int i = 0; i < 100; i++) {
        array1[i] = i;
    }
    strcpy(string, "Hello, World!");

    printf("Allocated array and string\n");
    print_memory_stats();

    // Free one block
    FREE(array1);
    printf("Freed array\n");
    print_memory_stats();
}

```

```
// Intentionally leak string to demonstrate leak detection
printf("Intentionally leaking string...\n");
print_memory_stats();

// Cleanup
cleanup_memory_tracking();
}
```

Summary

Memory optimization is a multifaceted discipline that requires understanding of:

1. **Memory Hierarchy** - How different levels of memory affect performance
2. **Allocation Strategies** - Custom allocators, stack vs heap allocation
3. **Data Structure Design** - Structure of Arrays vs Array of Structures, alignment
4. **Memory Layout** - Data locality, memory pooling
5. **Access Patterns** - Sequential vs random access, prefetching
6. **Fragmentation Management** - Reducing memory fragmentation
7. **Analysis Tools** - Memory tracking and debugging

Key principles for effective memory optimization:

- **Profile First** - Measure before and after optimization to verify improvements
- **Understand Your Data** - Know how your data is accessed and sized
- **Consider the Cache** - Design for cache-friendly access patterns
- **Minimize Allocations** - Reduce allocation frequency and overhead
- **Track Memory Usage** - Monitor for leaks and fragmentation
- **Use Appropriate Tools** - Leverage profilers and analysis tools

These techniques, when applied appropriately, can lead to significant performance improvements in C programs, especially in resource-constrained or performance-critical applications.

Advanced Optimization

Introduction

Advanced optimization techniques go beyond basic algorithmic improvements to leverage hardware-specific features, compiler optimizations, and parallel processing capabilities. This chapter explores sophisticated methods for maximizing performance in C programs, including vectorization, parallel processing, compiler-assisted optimizations, and specialized techniques for specific architectures.

Compiler Optimization

Modern compilers provide powerful optimization capabilities that can significantly improve performance with minimal programmer effort.

Optimization Levels

Different compiler optimization levels offer varying degrees of optimization:

```
// Example showing the impact of different optimization levels
// Compile with: gcc -O0, -O1, -O2, -O3, -Os

#include <stdio.h>
#include <time.h>

#define ITERATIONS 100000000

// Simple loop that can be optimized
int compute_sum(int n) {
    int sum = 0;
    for (int i = 0; i < n; i++) {
        sum += i * 2 + 1;
    }
    return sum;
}

void optimization_levels_demo(void) {
    clock_t start = clock();
    volatile int result = compute_sum(ITERATIONS);
    clock_t end = clock();
```



```

    double time_spent = ((double)(end - start)) / CLOCKS_PER_SEC;
    printf("Computation time: %.4f seconds\n", time_spent);
}

```

Compiler Hints and Pragmas

Compiler hints can guide optimization decisions:

```

#include <stdio.h>
#include <stdlib.h>

// Function likely to be called frequently
__attribute__((hot))
int hot_function(int x) {
    return x * x + 2 * x + 1;
}

// Function unlikely to be called
__attribute__((cold))
void error_handler(void) {
    printf("Error occurred!\n");
    exit(1);
}

// Loop unrolling hint
void loop_unrolling_example(void) {
    int array[1000];

    // Hint to unroll this loop
    #pragma GCC unroll 4
    for (int i = 0; i < 1000; i++) {
        array[i] = i * i;
    }
}

// Vectorization hints
void vectorization_example(float *a, float *b, float *c, int n) {
    // Tell compiler this loop can be vectorized
    #pragma GCC ivdep
    for (int i = 0; i < n; i++) {
        c[i] = a[i] + b[i];
    }
}

```

Profile-Guided Optimization (PGO)

PGO uses runtime profiling data to optimize code placement and branch prediction:

```
// Example of code that benefits from PGO
#include <stdio.h>

int process_data(int data) {
    // Branch that depends on input data
    if (data < 100) {
        // Common case
        return data * 2;
    } else if (data < 1000) {
        // Less common case
        return data + 100;
    } else {
        // Rare case
        return data / 2;
    }
}

void pgo_example(void) {
    // In a real PGO scenario, you would run the program with
    // representative data to generate profile data
    for (int i = 0; i < 10000; i++) {
        int result = process_data(i);
        // Use result to prevent optimization
        volatile int dummy = result;
    }
}
```

Vectorization

Vectorization allows processing multiple data elements simultaneously using SIMD (Single Instruction, Multiple Data) instructions.

Manual Vectorization with Intrinsics

Using compiler intrinsics for explicit vectorization:

```
#include <stdio.h>
#include <stdlib.h>
#include <time.h>

#ifdef __SSE__
```

```

#include <xmmintrin.h> // SSE
#endif

#ifdef __AVX__
#include <immintrin.h> // AVX
#endif

// Scalar addition
void scalar_add(float *a, float *b, float *c, int n) {
    for (int i = 0; i < n; i++) {
        c[i] = a[i] + b[i];
    }
}

#ifdef __SSE__
// SSE vectorized addition
void sse_add(float *a, float *b, float *c, int n) {
    int i = 0;

    // Process 4 elements at a time
    for (; i <= n - 4; i += 4) {
        __m128 va = _mm_load_ps(&a[i]);
        __m128 vb = _mm_load_ps(&b[i]);
        __m128 vc = _mm_add_ps(va, vb);
        _mm_store_ps(&c[i], vc);
    }

    // Handle remaining elements
    for (; i < n; i++) {
        c[i] = a[i] + b[i];
    }
}
#endif

#ifdef __AVX__
// AVX vectorized addition
void avx_add(float *a, float *b, float *c, int n) {
    int i = 0;

    // Process 8 elements at a time
    for (; i <= n - 8; i += 8) {
        __m256 va = _mm256_load_ps(&a[i]);
        __m256 vb = _mm256_load_ps(&b[i]);
        __m256 vc = _mm256_add_ps(va, vb);
        _mm256_store_ps(&c[i], vc);
    }
}

```

```

        // Handle remaining elements
        for (; i < n; i++) {
            c[i] = a[i] + b[i];
        }
    }
#endif

void vectorization_comparison(void) {
    const int size = 1000000;
    float *a = malloc(size * sizeof(float));
    float *b = malloc(size * sizeof(float));
    float *c = malloc(size * sizeof(float));

    // Initialize arrays
    for (int i = 0; i < size; i++) {
        a[i] = (float)i;
        b[i] = (float)(i * 2);
    }

    clock_t start, end;

    // Scalar version
    start = clock();
    scalar_add(a, b, c, size);
    end = clock();
    double scalar_time = ((double)(end - start)) / CLOCKS_PER_SEC;

#ifdef __SSE__
    // SSE version
    start = clock();
    sse_add(a, b, c, size);
    end = clock();
    double sse_time = ((double)(end - start)) / CLOCKS_PER_SEC;
    printf("SSE time: %.4f seconds (%.2fx speedup)\n",
           sse_time, scalar_time / sse_time);
#endif

#ifdef __AVX__
    // AVX version
    start = clock();
    avx_add(a, b, c, size);
    end = clock();
    double avx_time = ((double)(end - start)) / CLOCKS_PER_SEC;
    printf("AVX time: %.4f seconds (%.2fx speedup)\n",
           avx_time, scalar_time / avx_time);
#endif
}

```

```

    printf("Scalar time: %.4f seconds\n", scalar_time);

    free(a);
    free(b);
    free(c);
}

```

Auto-Vectorization

Modern compilers can automatically vectorize suitable loops:

```

#include <stdio.h>
#include <stdlib.h>
#include <time.h>

#define ARRAY_SIZE 1000000

// Compiler can auto-vectorize this function
void auto_vectorizable_loop(float *a, float *b, float *c, int n) {
    // Conditions for auto-vectorization:
    // 1. No dependencies between iterations
    // 2. Simple arithmetic operations
    // 3. Contiguous memory access
    // 4. Known loop bounds

    for (int i = 0; i < n; i++) {
        c[i] = a[i] * b[i] + 1.0f;
    }
}

// This loop cannot be auto-vectorized due to dependency
void non_vectorizable_loop(int *a, int n) {
    // Each iteration depends on the previous one
    for (int i = 1; i < n; i++) {
        a[i] = a[i-1] + i;
    }
}

void auto_vectorization_example(void) {
    float *a = malloc(ARRAY_SIZE * sizeof(float));
    float *b = malloc(ARRAY_SIZE * sizeof(float));
    float *c = malloc(ARRAY_SIZE * sizeof(float));

    // Initialize arrays
    for (int i = 0; i < ARRAY_SIZE; i++) {
        a[i] = (float)i;
    }
}

```

```

        b[i] = (float)(i * 2);
    }

    clock_t start = clock();
    auto_vectorizable_loop(a, b, c, ARRAY_SIZE);
    clock_t end = clock();

    double time_spent = ((double)(end - start)) / CLOCKS_PER_SEC;
    printf("Auto-vectorized loop time: %.4f seconds\n", time_spent);

    // Verify result (check first few elements)
    printf("First 5 results: ");
    for (int i = 0; i < 5; i++) {
        printf("%.0f ", c[i]);
    }
    printf("\n");

    free(a);
    free(b);
    free(c);
}

```

Parallel Processing

Utilizing multiple CPU cores can dramatically improve performance for suitable workloads.

OpenMP for Shared-Memory Parallelism

OpenMP provides a simple way to parallelize loops and sections:

```

#include <stdio.h>
#include <stdlib.h>
#include <time.h>
#include <omp.h>

#define ARRAY_SIZE 1000000

// Sequential matrix multiplication
void sequential_matrix_mult(int n, double *a, double *b, double *c) {
    for (int i = 0; i < n; i++) {
        for (int j = 0; j < n; j++) {
            double sum = 0.0;
            for (int k = 0; k < n; k++) {
                sum += a[i * n + k] * b[k * n + j];
            }
        }
    }
}

```

```

        }
        c[i * n + j] = sum;
    }
}

// Parallel matrix multiplication with OpenMP
void parallel_matrix_mult(int n, double *a, double *b, double *c) {
    #pragma omp parallel for
    for (int i = 0; i < n; i++) {
        for (int j = 0; j < n; j++) {
            double sum = 0.0;
            for (int k = 0; k < n; k++) {
                sum += a[i * n + k] * b[k * n + j];
            }
            c[i * n + j] = sum;
        }
    }
}

// Parallel reduction example
double parallel_sum(double *array, int n) {
    double sum = 0.0;

    #pragma omp parallel for reduction(+:sum)
    for (int i = 0; i < n; i++) {
        sum += array[i];
    }

    return sum;
}

void openmp_example(void) {
    const int size = 1000;
    double *a = malloc(size * size * sizeof(double));
    double *b = malloc(size * size * sizeof(double));
    double *c = malloc(size * size * sizeof(double));

    // Initialize matrices
    for (int i = 0; i < size * size; i++) {
        a[i] = (double)(i % 100);
        b[i] = (double)((i + 1) % 100);
    }

    // Measure sequential version
    clock_t start = clock();
    sequential_matrix_mult(size, a, b, c);

```

```

clock_t end = clock();
double sequential_time = ((double)(end - start)) / CLOCKS_PER_SEC;

// Measure parallel version
start = clock();
parallel_matrix_mult(size, a, b, c);
end = clock();
double parallel_time = ((double)(end - start)) / CLOCKS_PER_SEC;

printf("Sequential time: %.4f seconds\n", sequential_time);
printf("Parallel time: %.4f seconds\n", parallel_time);
printf("Speedup: %.2fx\n", sequential_time / parallel_time);

// Show number of threads
printf("Number of OpenMP threads: %d\n", omp_get_max_threads());

free(a);
free(b);
free(c);
}

```

Thread-Level Parallelism with Pthreads

Manual thread management for fine-grained control:

```

#include <stdio.h>
#include <stdlib.h>
#include <pthread.h>
#include <time.h>

#define NUM_THREADS 4
#define ARRAY_SIZE 1000000

typedef struct {
    int *array;
    int start;
    int end;
    long long sum;
} thread_data_t;

// Thread function to compute partial sum
void* compute_partial_sum(void *arg) {
    thread_data_t *data = (thread_data_t*)arg;

    data->sum = 0;
    for (int i = data->start; i < data->end; i++) {

```



```

        data->sum += data->array[i];
    }

    return NULL;
}

long long parallel_sum_pthreads(int *array, int n) {
    pthread_t threads[NUM_THREADS];
    thread_data_t thread_data[NUM_THREADS];

    // Divide work among threads
    int chunk_size = n / NUM_THREADS;

    for (int i = 0; i < NUM_THREADS; i++) {
        thread_data[i].array = array;
        thread_data[i].start = i * chunk_size;
        thread_data[i].end = (i == NUM_THREADS - 1) ? n : (i + 1) * chunk_size;

        pthread_create(&threads[i], NULL, compute_partial_sum, &thread_data[i]);
    }

    // Wait for all threads and accumulate results
    long long total_sum = 0;
    for (int i = 0; i < NUM_THREADS; i++) {
        pthread_join(threads[i], NULL);
        total_sum += thread_data[i].sum;
    }

    return total_sum;
}

long long sequential_sum(int *array, int n) {
    long long sum = 0;
    for (int i = 0; i < n; i++) {
        sum += array[i];
    }
    return sum;
}

void pthreads_example(void) {
    int *array = malloc(ARRAY_SIZE * sizeof(int));

    // Initialize array
    for (int i = 0; i < ARRAY_SIZE; i++) {
        array[i] = i % 1000;
    }
}

```

```

// Sequential sum
clock_t start = clock();
long long seq_sum = sequential_sum(array, ARRAY_SIZE);
clock_t end = clock();
double sequential_time = ((double)(end - start)) / CLOCKS_PER_SEC;

// Parallel sum
start = clock();
long long par_sum = parallel_sum_pthreads(array, ARRAY_SIZE);
end = clock();
double parallel_time = ((double)(end - start)) / CLOCKS_PER_SEC;

printf("Sequential sum: %lld (time: %.4f seconds)\n", seq_sum, sequential_time);
printf("Parallel sum: %lld (time: %.4f seconds)\n", par_sum, parallel_time);
printf("Speedup: %.2fx\n", sequential_time / parallel_time);

free(array);
}

```

Architecture-Specific Optimizations

Different CPU architectures have unique features that can be leveraged for performance.

Cache Optimization Techniques

Optimizing for cache hierarchy and behavior:

```

#include <stdio.h>
#include <stdlib.h>
#include <time.h>

#define MATRIX_SIZE 1024

// Cache-unfriendly matrix transpose
void transpose_unfriendly(int n, int *src, int *dst) {
    for (int i = 0; i < n; i++) {
        for (int j = 0; j < n; j++) {
            dst[j * n + i] = src[i * n + j];
        }
    }
}

// Cache-friendly matrix transpose with blocking
void transpose_blocked(int n, int *src, int *dst) {

```

```

const int block_size = 32; // Should match cache line size

for (int ii = 0; ii < n; ii += block_size) {
    for (int jj = 0; jj < n; jj += block_size) {
        for (int i = ii; i < ii + block_size && i < n; i++) {
            for (int j = jj; j < jj + block_size && j < n; j++) {
                dst[j * n + i] = src[i * n + j];
            }
        }
    }
}

}

void cache_optimization_example(void) {
    int *matrix = malloc(MATRIX_SIZE * MATRIX_SIZE * sizeof(int));
    int *transposed = malloc(MATRIX_SIZE * MATRIX_SIZE * sizeof(int));

    // Initialize matrix
    for (int i = 0; i < MATRIX_SIZE * MATRIX_SIZE; i++) {
        matrix[i] = i;
    }

    clock_t start, end;

    // Unfriendly transpose
    start = clock();
    transpose_unfriendly(MATRIX_SIZE, matrix, transposed);
    end = clock();
    double unfriendly_time = ((double)(end - start)) / CLOCKS_PER_SEC;

    // Blocked transpose
    start = clock();
    transpose_blocked(MATRIX_SIZE, matrix, transposed);
    end = clock();
    double blocked_time = ((double)(end - start)) / CLOCKS_PER_SEC;

    printf("Unfriendly transpose: %.4f seconds\n", unfriendly_time);
    printf("Blocked transpose: %.4f seconds\n", blocked_time);
    printf("Improvement: %.2fx\n", unfriendly_time / blocked_time);

    free(matrix);
    free(transposed);
}

```

Branch Prediction Optimization

Improving branch prediction accuracy for better performance:

```
#include <stdio.h>
#include <stdlib.h>
#include <time.h>

#define ARRAY_SIZE 1000000

// Function with unpredictable branches
void unpredictable_branches(int *array, int n) {
    for (int i = 0; i < n; i++) {
        // Random condition leads to poor branch prediction
        if (rand() % 2) {
            array[i] *= 2;
        } else {
            array[i] += 1;
        }
    }
}

// Function with predictable branches
void predictable_branches(int *array, int n) {
    // Process even indices differently from odd indices
    for (int i = 0; i < n; i++) {
        if (i % 2 == 0) {
            array[i] *= 2;
        } else {
            array[i] += 1;
        }
    }
}

// Branchless version using conditional moves
void branchless_version(int *array, int n) {
    for (int i = 0; i < n; i++) {
        int is_even = (i % 2 == 0);
        array[i] = array[i] * (1 + is_even) + (1 - is_even);
    }
}

void branch_prediction_example(void) {
    int *array1 = malloc(ARRAY_SIZE * sizeof(int));
    int *array2 = malloc(ARRAY_SIZE * sizeof(int));
    int *array3 = malloc(ARRAY_SIZE * sizeof(int));
```

```

// Initialize arrays
for (int i = 0; i < ARRAY_SIZE; i++) {
    array1[i] = array2[i] = array3[i] = i;
}

clock_t start, end;

// Unpredictable branches
srand(42);
start = clock();
unpredictable_branches(array1, ARRAY_SIZE);
end = clock();
double unpredictable_time = ((double)(end - start)) / CLOCKS_PER_SEC;

// Predictable branches
start = clock();
predictable_branches(array2, ARRAY_SIZE);
end = clock();
double predictable_time = ((double)(end - start)) / CLOCKS_PER_SEC;

// Branchless version
start = clock();
branchless_version(array3, ARRAY_SIZE);
end = clock();
double branchless_time = ((double)(end - start)) / CLOCKS_PER_SEC;

printf("Unpredictable branches: %.4f seconds\n", unpredictable_time);
printf("Predictable branches: %.4f seconds\n", predictable_time);
printf("Branchless version: %.4f seconds\n", branchless_time);
printf("Predictable improvement: %.2fx\n", unpredictable_time / predictable_time);
printf("Branchless improvement: %.2fx\n", unpredictable_time / branchless_time);

free(array1);
free(array2);
free(array3);
}

```

Memory-Level Parallelism

Exploiting memory-level parallelism to hide latency:

```

#include <stdio.h>
#include <stdlib.h>
#include <time.h>

```

```

#define ARRAY_SIZE 1000000

// Function that can benefit from memory-level parallelism
void memory_parallelism_example(void) {
    double *a = malloc(ARRAY_SIZE * sizeof(double));
    double *b = malloc(ARRAY_SIZE * sizeof(double));
    double *c = malloc(ARRAY_SIZE * sizeof(double));
    double *d = malloc(ARRAY_SIZE * sizeof(double));

    // Initialize arrays
    for (int i = 0; i < ARRAY_SIZE; i++) {
        a[i] = (double)i;
        b[i] = (double)(i * 2);
        c[i] = (double)(i * 3);
        d[i] = 0.0;
    }

    clock_t start = clock();

    // Independent operations that can be executed in parallel
    // Modern processors can execute multiple loads/stores simultaneously
    for (int i = 0; i < ARRAY_SIZE; i++) {
        d[i] = a[i] + b[i] + c[i];
    }

    clock_t end = clock();
    double time_spent = ((double)(end - start)) / CLOCKS_PER_SEC;

    printf("Memory parallelism example time: %.4f seconds\n", time_spent);
    printf("First few results: ");
    for (int i = 0; i < 5; i++) {
        printf("%.0f ", d[i]);
    }
    printf("\n");

    free(a);
    free(b);
    free(c);
    free(d);
}

```

Advanced Profiling and Analysis

Using sophisticated tools for performance analysis:

```

#include <stdio.h>
#include <stdlib.h>

#ifdef __linux__
#include <sys/resource.h>
#endif

// Function to measure memory usage
void print_memory_usage(void) {
#ifdef __linux__
    struct rusage usage;
    if (getrusage(RUSAGE_SELF, &usage) == 0) {
        printf("Maximum resident set size: %ld KB\n", usage.ru_maxrss);
        printf("Page reclaims: %ld\n", usage.ru_minflt);
        printf("Page faults: %ld\n", usage.ru_majflt);
    }
#endif
}

// Function to demonstrate performance counters (conceptual)
void performance_counter_example(void) {
    printf("Performance counter example:\n");
    printf("- Instructions per cycle (IPC)\n");
    printf("- Cache miss rates\n");
    printf("- Branch misprediction rate\n");
    printf("- Memory bandwidth utilization\n");

    // In practice, you would use tools like:
    // - perf (Linux)
    // - Intel VTune
    // - Apple Instruments
    // - GNU gprof
}

```

Optimization Decision Making

Making informed decisions about which optimizations to apply:

```

#include <stdio.h>
#include <time.h>

// Example showing when optimization matters
void optimization_decision_example(void) {
    printf("When to optimize:\n");
    printf("1. Profile first - identify actual bottlenecks\n");
}

```

```

printf("2. Consider algorithmic improvements before micro-optimizations\n");
printf("3. Optimize hot paths - functions called frequently\n");
printf("4. Balance performance with maintainability\n");
printf("5. Consider target platform and constraints\n");
printf("6. Measure impact of each optimization\n");

// Simple example of measuring optimization impact
const int iterations = 100000000;

clock_t start = clock();
volatile long long sum = 0;
for (int i = 0; i < iterations; i++) {
    sum += i;
}
clock_t end = clock();

double time_spent = ((double)(end - start)) / CLOCKS_PER_SEC;
printf("\nSimple loop time: %.4f seconds\n", time_spent);
}

```

Practical Optimization Workflow

A systematic approach to performance optimization:

```

#include <stdio.h>

void optimization_workflow(void) {
    printf("Performance Optimization Workflow:\n\n");

    printf("1. Profiling and Measurement\n");
    printf("    - Use profiling tools to identify bottlenecks\n");
    printf("    - Measure baseline performance\n");
    printf("    - Set performance goals\n\n");

    printf("2. Algorithmic Analysis\n");
    printf("    - Review algorithmic complexity\n");
    printf("    - Consider alternative algorithms\n");
    printf("    - Optimize data structures\n\n");

    printf("3. Implementation Optimization\n");
    printf("    - Apply compiler optimizations\n");
    printf("    - Use appropriate data layouts\n");
    printf("    - Minimize memory allocations\n\n");

    printf("4. Hardware-Specific Tuning\n");
}

```



```

printf("    - Leverage SIMD instructions\n");
printf("    - Utilize multiple cores\n");
printf("    - Optimize for cache hierarchy\n\n");

printf("5. Validation and Testing\n");
printf("    - Verify correctness after optimization\n");
printf("    - Measure performance improvements\n");
printf("    - Test on target hardware\n\n");

printf("6. Documentation and Maintenance\n");
printf("    - Document optimization decisions\n");
printf("    - Monitor performance over time\n");
printf("    - Maintain balance between performance and readability\n");
}

```

Summary

Advanced optimization techniques require a deep understanding of both the problem domain and the underlying hardware. Key takeaways include:

1. **Compiler Optimization** - Leverage compiler flags and hints for automatic improvements
2. **Vectorization** - Use SIMD instructions to process multiple data elements simultaneously
3. **Parallel Processing** - Distribute work across multiple CPU cores using OpenMP or pthreads
4. **Architecture-Specific Optimizations** - Tailor code for specific CPU features and cache behaviors
5. **Memory-Level Parallelism** - Exploit the ability to perform multiple memory operations simultaneously
6. **Branch Prediction** - Structure code to improve CPU branch prediction accuracy
7. **Systematic Approach** - Follow a methodical workflow for identifying and addressing performance bottlenecks

Important principles to remember:

- **Profile First** - Always measure before and after optimization
- **Algorithmic Improvements** - Often provide bigger gains than micro-optimizations
- **Platform Awareness** - Consider target hardware characteristics
- **Maintainability** - Balance performance with code readability and maintainability
- **Verification** - Ensure optimizations don't introduce bugs or change behavior

These advanced techniques, when applied appropriately, can yield substantial performance improvements in performance-critical C applications. However, they should be used judiciously and only after identifying actual bottlenecks through profiling.

Module 13: Performance Optimization

Overview

Performance optimization is a critical aspect of C programming, especially in resource-constrained environments and high-performance applications. This module covers techniques for analyzing, measuring, and improving the performance of C programs through profiling, algorithm optimization, memory management, and advanced optimization strategies.

Learning Objectives

By the end of this module, you should be able to: - Analyze program performance using profiling tools - Identify and eliminate performance bottlenecks - Optimize algorithms and data structures for better performance - Manage memory efficiently to reduce overhead - Apply compiler optimization techniques - Understand the trade-offs between performance and other factors - Implement caching strategies and algorithmic improvements - Measure and benchmark performance improvements

Chapters

1. **Performance Analysis** - Profiling tools, benchmarking techniques, and performance metrics
2. **Optimization Techniques** - Algorithm optimization, loop optimization, and code restructuring
3. **Memory Optimization** - Memory allocation strategies, cache optimization, and memory layout
4. **Advanced Optimization** - Compiler optimizations, vectorization, and parallel processing

Prerequisites

Before starting this module, you should have a solid understanding of: - C programming fundamentals - Data structures and algorithms - Memory management concepts - Basic understanding of computer architecture

Key Concepts

- Profiling and performance measurement
- Algorithmic complexity analysis
- Memory hierarchy and cache optimization
- Compiler optimization flags and techniques
- Loop optimization and vectorization
- Memory allocation strategies
- Performance benchmarking
- Trade-offs between performance and maintainability

Testing

Testing Methodologies

Testing ensures C code behaves correctly under normal use, edge cases, and failure conditions. This chapter focuses on **practical unit testing in C on Linux/POSIX**: small harnesses you can compile with `gcc`, table-driven tests, TDD, boundary analysis, coverage with `gcov/lcov`, and sanitizers.

Build flags used throughout this chapter:

```
gcc -std=c17 -Wall -Wextra -o program program.c
```

Add `-g` for debugging symbols, and sanitizer flags when hunting memory/UB bugs (covered below).

Why Testing Matters

- **Bug detection early**: cheaper than production incidents
- **Documentation**: tests show intended API contracts
- **Safe refactoring**: change internals without fear if tests pass
- **Regression control**: catch breakages when you touch old code

Types of Testing (C-focused)

Level	What you test	Typical tools in C
Unit	One function / small module	assert , custom macros, Unity, CUnit
Integration	Several .c files / libraries together	Same harness + real I/O stubs
System	Whole binary (CLI, server)	Shell scripts, golden files
Regression	“This used to work” cases	Any automated suite in CI

Unit testing characteristics

- **Small and focused** — one behavior per test when possible
- **Fast** — milliseconds, not minutes
- **Independent** — no order dependence

- **Repeatable** — same result every run
 - **Automated** — no human clicks
-

Style 1: `assert.h` Unit Tests

`<assert.h>` is enough for tiny libraries. Failed asserts abort the process (good for local runs; less pretty for large suites).

```
/* file: test_add.c */
#include <assert.h>
#include <stdio.h>

int add(int a, int b) {
    return a + b;
}

void test_add(void) {
    assert(add(2, 3) == 5);
    assert(add(-1, 1) == 0);
    assert(add(0, 0) == 0);
    assert(add(1000000, 2000000) == 3000000);
}

int main(void) {
    test_add();
    printf("all asserts passed\n");
    return 0;
}

gcc -std=c17 -Wall -Wextra -o test_add test_add.c
./test_add
```

Note: In release builds, `NDEBUG` disables asserts. Keep tests in a separate binary that never defines `NDEBUG`.

Style 2: Minimal Custom TEST Macro Harness

A few macros give you named tests, a pass/fail counter, and a non-zero exit code — without external libraries.

```

/* file: mini_test.h */
#ifndef MINI_TEST_H
#define MINI_TEST_H

#include <stdio.h>
#include <string.h>

static int g_tests_run;
static int g_tests_failed;

#define TEST_BEGIN() \
    do { g_tests_run = 0; g_tests_failed = 0; } while (0)

#define TEST_END() \
    do { \
        printf("\n%d tests, %d failed\n", g_tests_run, g_tests_failed); \
        return g_tests_failed ? 1 : 0; \
    } while (0)

#define RUN_TEST(fn) \
    do { \
        g_tests_run++; \
        printf("  RUN  %s ... ", #fn); \
        fflush(stdout); \
        if ((fn)() == 0) { \
            printf("OK\n"); \
        } else { \
            g_tests_failed++; \
            printf("FAIL\n"); \
        } \
    } while (0)

/* Return 0 on success, 1 on failure from these macros */
#define EXPECT_EQ_INT(actual, expected) \
    do { \
        int _a = (actual); \
        int _e = (expected); \
        if (_a != _e) { \
            fprintf(stderr, "\n    %s:%d: expected %d, got %d\n", \
                __FILE__, __LINE__, _e, _a); \
            return 1; \
        } \
    } while (0)

#define EXPECT_TRUE(cond) \
    do { \
        if (!(cond)) { \

```

```

        fprintf(stderr, "\n    %s:%d: condition false: %s\n",
                __FILE__, __LINE__, #cond);
        return 1;
    }
} while (0)

#define EXPECT_STREQ(actual, expected)
do {
    const char *_a = (actual);
    const char *_e = (expected);
    if (_a == NULL || _e == NULL || strcmp(_a, _e) != 0) {
        fprintf(stderr, "\n    %s:%d: expected \"%s\", got \"%s\"\n",
                __FILE__, __LINE__,
                _e ? _e : "(null)", _a ? _a : "(null)");
        return 1;
    }
} while (0)

#endif /* MINI_TEST_H */

/* file: test_clamp.c - uses mini_test.h */
#include "mini_test.h"

int clamp(int x, int lo, int hi) {
    if (x < lo) {
        return lo;
    }
    if (x > hi) {
        return hi;
    }
    return x;
}

static int test_clamp_inside(void) {
    EXPECT_EQ_INT(clamp(5, 0, 10), 5);
    return 0;
}

static int test_clamp_below(void) {
    EXPECT_EQ_INT(clamp(-3, 0, 10), 0);
    return 0;
}

static int test_clamp_above(void) {
    EXPECT_EQ_INT(clamp(99, 0, 10), 10);
    return 0;
}

```



```

}

int main(void) {
    TEST_BEGIN();
    RUN_TEST(test_clamp_inside);
    RUN_TEST(test_clamp_below);
    RUN_TEST(test_clamp_above);
    TEST_END();
}

gcc -std=c17 -Wall -Wextra -o test_clamp test_clamp.c
./test_clamp

```

Style 3: Unity-Style Suite That Actually Builds

Unity is a popular embedded-friendly C unit framework. You can either vendor its sources or use a **Unity-compatible subset** that builds as a single file for learning.

Option A — Self-contained “unity-like” runner (no download)

```

/* file: unity_lite.h - small subset of Unity API names */
#ifndef UNITY_LITE_H
#define UNITY_LITE_H

#include <stdio.h>
#include <string.h>

static int Unity_tests;
static int Unity_failures;
static const char *Unity_current;

#define UNITY_BEGIN() \
    do { Unity_tests = 0; Unity_failures = 0; printf("Unity-lite\n"); } while (0)

#define UNITY_END() \
    do { \
        printf("-----\n%d Tests %d Failures\n", \
            Unity_tests, Unity_failures); \
        return Unity_failures ? 1 : 0; \
    } while (0)

```

```

#define RUN_TEST(fn) \
do { \
    Unity_tests++; \
    Unity_current = #fn; \
    printf("TEST(%s)\n", Unity_current); \
    fn(); \
} while (0)

#define TEST_ASSERT_EQUAL_INT(expected, actual) \
do { \
    int _e = (int)(expected); \
    int _a = (int)(actual); \
    if (_e != _a) { \
        Unity_failures++; \
        printf("  FAIL: %s:%d expected %d was %d\n", \
            __FILE__, __LINE__, _e, _a); \
    } \
} while (0)

#define TEST_ASSERT_TRUE(cond) \
do { \
    if (!(cond)) { \
        Unity_failures++; \
        printf("  FAIL: %s:%d %s is false\n", \
            __FILE__, __LINE__, #cond); \
    } \
} while (0)

#define TEST_ASSERT_EQUAL_STRING(expected, actual) \
do { \
    const char *_e = (expected); \
    const char *_a = (actual); \
    if (_e == NULL || _a == NULL || strcmp(_e, _a) != 0) { \
        Unity_failures++; \
        printf("  FAIL: %s:%d expected \"%s\" was \"%s\"\n", \
            __FILE__, __LINE__, \
            _e ? _e : "(null)", _a ? _a : "(null)"); \
    } \
} while (0)

#endif

/* file: test_leap_unity.c */
#include "unity_lite.h"

/* Production code under test */

```

```

int is_leap_year(int year) {
    if (year % 400 == 0) {
        return 1;
    }
    if (year % 100 == 0) {
        return 0;
    }
    if (year % 4 == 0) {
        return 1;
    }
    return 0;
}

void test_common_year(void) {
    TEST_ASSERT_EQUAL_INT(0, is_leap_year(2019));
}

void test_leap_divisible_by_4(void) {
    TEST_ASSERT_EQUAL_INT(1, is_leap_year(2020));
}

void test_century_not_leap(void) {
    TEST_ASSERT_EQUAL_INT(0, is_leap_year(1900));
}

void test_400_year_leap(void) {
    TEST_ASSERT_EQUAL_INT(1, is_leap_year(2000));
}

int main(void) {
    UNITY_BEGIN();
    RUN_TEST(test_common_year);
    RUN_TEST(test_leap_divisible_by_4);
    RUN_TEST(test_century_not_leap);
    RUN_TEST(test_400_year_leap);
    return UNITY_END();
}

gcc -std=c17 -Wall -Wextra -o test_leap test_leap_unity.c
./test_leap

```

Option B — Real Unity (vendored)

If you clone Unity into `third_party/Unity`:

```
# Example layout
# third_party/Unity/src/unity.c
# third_party/Unity/src/unity.h
# third_party/Unity/src/unity_internals.h

gcc -std=c17 -Wall -Wextra \
  -I third_party/Unity/src \
  -o test_leap \
  test_leap_unity_real.c third_party/Unity/src/unity.c
```

Real Unity uses the same names (UNITY_BEGIN, RUN_TEST, TEST_ASSERT_EQUAL_INT, UNITY_END). Your tests stay almost identical to Option A.

Test-Driven Development (TDD) Walkthrough: is_leap_year()

TDD cycle: **Red** → **Green** → **Refactor**.

We implement Gregorian leap-year rules:

- Divisible by 400 → leap
- Else divisible by 100 → not leap
- Else divisible by 4 → leap
- Else → not leap

Step 0 — Scaffold (no real logic yet)

```
/* leap.c + test_leap_tdd.c evolution */

int is_leap_year(int year) {
    (void)year;
    return 0; /* deliberately wrong / incomplete */
}
```

Step 1 — RED: first failing test

```
/* test_leap_tdd.c */
#include "unity_lite.h"

int is_leap_year(int year); /* or #include "leap.h" */
```

```
void test_2020_is_leap(void) {
    TEST_ASSERT_EQUAL_INT(1, is_leap_year(2020));
}
```

```
int main(void) {
    UNITY_BEGIN();
    RUN_TEST(test_2020_is_leap);
    return UNITY_END();
}
```

```
gcc -std=c17 -Wall -Wextra -o test_leap_tdd test_leap_tdd.c leap.c
./test_leap_tdd
# FAIL: expected 1 was 0
```

Step 2 — GREEN: minimal code that passes

```
/* leap.c - first green */
int is_leap_year(int year) {
    return year % 4 == 0;
}
```

```
./test_leap_tdd
# OK for 2020
```

Step 3 — RED: century rule

```
void test_1900_not_leap(void) {
    TEST_ASSERT_EQUAL_INT(0, is_leap_year(1900));
}
```

```
/* main: also RUN_TEST(test_1900_not_leap); */
```

```
./test_leap_tdd
# FAIL for 1900 (1900 % 4 == 0 with naive rule)
```

Step 4 — GREEN: handle centuries

```
int is_leap_year(int year) {
    if (year % 100 == 0) {
        return 0;
    }
}
```

```

    }
    return year % 4 == 0;
}

```

Step 5 — RED: year 2000 must be leap

```

void test_2000_is_leap(void) {
    TEST_ASSERT_EQUAL_INT(1, is_leap_year(2000));
}

```

Step 6 — GREEN: full Gregorian rule

```

int is_leap_year(int year) {
    if (year % 400 == 0) {
        return 1;
    }
    if (year % 100 == 0) {
        return 0;
    }
    if (year % 4 == 0) {
        return 1;
    }
    return 0;
}

```

Step 7 — REFACTOR: keep tests green

```

int is_leap_year(int year) {
    return (year % 400 == 0) || (year % 4 == 0 && year % 100 != 0);
}

```

Re-run the suite after every refactor. All four (or more) tests should still pass.

Alternate TDD target: `clamp()`

Same cycle with fewer domain rules:

Phase	Test idea	Implementation
Red	<code>clamp(5,0,10)==5</code>	<code>return 0;</code> fails
Green	return x if in range	<code>return x;</code> (too simple)
Red	<code>clamp(-1,0,10)==0</code>	fails
Green	lower bound	<code>if (x < lo) return lo; return x;</code>

Phase	Test idea	Implementation
Red	<code>clamp(99,0,10)==10</code>	fails
Green	upper bound	full clamp
Refactor	<code>assert lo <= hi</code> or document	keep tests green

Boundary Value Testing

Bugs cluster at edges. Design cases for **just inside**, **on**, and **just outside** valid ranges.

Integers

```
#include "unity_lite.h"
#include <limits.h>

int clamp(int x, int lo, int hi) {
    if (x < lo) return lo;
    if (x > hi) return hi;
    return x;
}

void test_clamp_boundaries(void) {
    /* lo / hi themselves */
    TEST_ASSERT_EQUAL_INT(0, clamp(0, 0, 10));
    TEST_ASSERT_EQUAL_INT(10, clamp(10, 0, 10));

    /* just outside */
    TEST_ASSERT_EQUAL_INT(0, clamp(-1, 0, 10));
    TEST_ASSERT_EQUAL_INT(10, clamp(11, 0, 10));

    /* extreme ints */
    TEST_ASSERT_EQUAL_INT(INT_MIN, clamp(INT_MIN, INT_MIN, INT_MAX));
    TEST_ASSERT_EQUAL_INT(INT_MAX, clamp(INT_MAX, INT_MIN, INT_MAX));
    TEST_ASSERT_EQUAL_INT(0, clamp(INT_MIN, 0, 10));
    TEST_ASSERT_EQUAL_INT(10, clamp(INT_MAX, 0, 10));
}
```

Strings

```
#include <string.h>
#include <stddef.h>

/* Returns length, or (size_t)-1 on NULL; never walks past maxn */
size_t safe_strnlen(const char *s, size_t maxn) {
    size_t i;
    if (s == NULL) {
        return (size_t)-1;
    }
    for (i = 0; i < maxn; i++) {
        if (s[i] == '\0') {
            return i;
        }
    }
    return maxn;
}

void test_safe_strnlen_boundaries(void) {
    TEST_ASSERT_EQUAL_INT((int)(size_t)-1, (int)safe_strnlen(NULL, 10));
    TEST_ASSERT_EQUAL_INT(0, (int)safe_strnlen("", 10));
    TEST_ASSERT_EQUAL_INT(0, (int)safe_strnlen("abc", 0));
    TEST_ASSERT_EQUAL_INT(3, (int)safe_strnlen("abc", 3));
    TEST_ASSERT_EQUAL_INT(3, (int)safe_strnlen("abc", 100));
    TEST_ASSERT_EQUAL_INT(3, (int)safe_strnlen("abcdef", 3)); /* capped */
}
```

Boundary checklist for strings

- NULL
 - empty ""
 - length 1
 - exact buffer size
 - longer than limit
 - non-ASCII / embedded \0 if relevant
-

Table-Driven Tests in C

Encode many cases in an array of structs. One loop runs them all — easy to extend.

```
/* file: test_leap_table.c */
#include <stdio.h>

int is_leap_year(int year) {
    return (year % 400 == 0) || (year % 4 == 0 && year % 100 != 0);
}

struct leap_case {
    int year;
    int expected; /* 1 leap, 0 not */
    const char *note;
};

static const struct leap_case cases[] = {
    {2019, 0, "common"},
    {2020, 1, "div by 4"},
    {1900, 0, "century"},
    {2000, 1, "div by 400"},
    {1600, 1, "div by 400"},
    {2100, 0, "century"},
    {1, 0, "year 1"},
    {4, 1, "year 4"},
};

int main(void) {
    int failed = 0;
    size_t i;
    size_t n = sizeof cases / sizeof cases[0];

    for (i = 0; i < n; i++) {
        int got = is_leap_year(cases[i].year);
        if (got != cases[i].expected) {
            fprintf(stderr, "FAIL year=%d (%s): expected %d got %d\n",
                    cases[i].year, cases[i].note,
                    cases[i].expected, got);
            failed++;
        } else {
            printf("OK   year=%d (%s)\n", cases[i].year, cases[i].note);
        }
    }

    printf("%zu cases, %d failed\n", n, failed);
    return failed ? 1 : 0;
}
```

```
}
```

```
gcc -std=c17 -Wall -Wextra -o test_leap_table test_leap_table.c  
./test_leap_table
```

Table-driven style also works for `clamp`, parsers, and error-code matrices.

Code Coverage with `gcov` / `lcov`

Coverage shows which lines ran during tests — not that logic is correct, but that untested paths exist.

Instrument and run

```
# Compile with coverage (debug-friendly)  
gcc -std=c17 -Wall -Wextra -g --coverage \  
    -o test_leap test_leap_unity.c  
  
./test_leap  
  
# Per-file summary  
gcov test_leap_unity.c  
# creates test_leap_unity.c.gcov - open in an editor  
  
# Optional pretty HTML report (install lcov)  
lcov --capture --directory . --output-file coverage.info  
genhtml coverage.info --output-directory coverage_html  
# open coverage_html/index.html
```

Tips

- Clean between runs: `rm -f *.gcda *.gcno *.gcov coverage.info`
- Link with `--coverage` (or `-lgcov` on some toolchains) when splitting objects
- Aim for high branch coverage on pure logic; 100% line coverage bug-free

Example `.gcov` line prefixes:

- #####: — never executed
- 1: — executed once

- branch 0 taken 0% — branch never taken (needs another case)
-

Sanitizers: Address and Undefined Behavior

GCC/Clang sanitizers catch classes of bugs tests alone may miss.

AddressSanitizer (ASan)

Detects heap/stack buffer overflows, use-after-free, some leaks.

```
/* file: bug_overflow.c */
#include <stdio.h>
#include <stdlib.h>

int main(void) {
    int *p = malloc(4 * sizeof(int));
    if (!p) {
        return 1;
    }
    p[4] = 42; /* one past the end - undefined behavior */
    printf("%d\n", p[0]);
    free(p);
    return 0;
}
```

```
gcc -std=c17 -Wall -Wextra -g -fsanitize=address -o bug_overflow bug_overflow.c
./bug_overflow
# ASan prints stack of the bad write and the allocation site
```

UndefinedBehaviorSanitizer (UBSan)

```
/* file: bug_ub.c */
#include <limits.h>
#include <stdio.h>

int main(void) {
    int x = INT_MAX;
    x = x + 1; /* signed overflow - UB */
    printf("%d\n", x);
    return 0;
}
```

```
gcc -std=c17 -Wall -Wextra -g -fsanitize=undefined -o bug_ub bug_ub.c
./bug_ub
```

Combined (common for tests)

```
gcc -std=c17 -Wall -Wextra -g \
    -fsanitize=address,undefined \
    -o test_leap test_leap_unity.c
./test_leap
```

CI idea: run the unit suite under ASan+UBSan on every commit.

ThreadSanitizer (`-fsanitize=thread`) is for races; requires pthreads and is slower.

Integration and System Testing (brief)

- **Integration:** link `module_a.o` + `module_b.o` and exercise the boundary API.
- **System:** run the finished CLI with shell scripts:

```
#!/usr/bin/env bash
set -euo pipefail
./myprog --input sample.txt > out.txt
diff -u expected.txt out.txt
```

- **Regression:** keep golden files under `tests/fixtures/` and fail the pipeline on diff.

Black-Box vs White-Box (quick map)

Approach	You know	Techniques
Black box	Spec only	Equivalence classes, boundaries, decision tables
White box	Source	Statement/branch/path coverage, data-flow
Gray box	Partial	Spec + known module structure

In unit tests you usually do **white/gray box**; acceptance tests stay closer to black box.

FIRST and AAA

- **FIRST**: Fast, Independent, Repeatable, Self-validating, Timely
- **AAA**: Arrange (setup) → Act (call SUT) → Assert (check)

```
void test_clamp_above(void) {  
    /* Arrange */  
    int x = 99, lo = 0, hi = 10;  
    /* Act */  
    int got = clamp(x, lo, hi);  
    /* Assert */  
    TEST_ASSERT_EQUAL_INT(10, got);  
}
```

Best Practices

1. Keep production code free of test-only `#ifdef` noise when possible — separate `*_test.c` files.
2. Prefer pure functions for hard logic (easy to table-test).
3. One clear reason to fail per test (or clearly named multi-assert tests).
4. Run tests and sanitizers before every push.
5. Treat flaky tests as bugs (timing, order, shared files).
6. Do not commit coverage HTML; generate it in CI artifacts.

Exercises

Exercise 1 — Custom harness for clamp

Implement `int clamp(int x, int lo, int hi)` and a `mini_test.h` suite covering inside, below, above, and `INT_MIN/INT_MAX`. Build with:

```
gcc -std=c17 -Wall -Wextra -o test_clamp test_clamp.c
```

Exercise 2 — TDD `int max3(int a, int b, int c)`

Write failing tests first, implement, then refactor. Include cases where the max is in the first, second, and third argument, and all equal.

Exercise 3 — Table-driven string trim

Write `void trim_inplace(char *s)` that removes leading/trailing ASCII spaces. Drive tests from a table of {input, expected} pairs (use a writable buffer per case). Include empty string, all spaces, no spaces, and mixed.

Exercise 4 — Coverage gap hunt

Take `is_leap_year` and deliberately omit the year-2000 test. Run `gcov` and show which branch is under-tested. Add the missing case and re-check.

Exercise 5 — Sanitizer lab

Write a small program with a deliberate heap overflow. Confirm it “works” without sanitizers (undefined — may print garbage or crash). Rebuild with `-fsanitize=address` and capture the report. Fix the bug and re-run clean.

Exercise 6 — Boundary design

For a function `int parse_port(const char *s, int *out)` that accepts `"1"..."65535"`, list at least ten boundary/equivalence cases (NULL, "", "0", "1", "65535", "65536", leading zeros, non-digits, etc.) and implement tests.

Summary

Technique	When to use
<code>assert</code>	Tiny checks, throwaway scripts
Custom TEST macros	Portable unit suites with no deps
Unity-style runner	Familiar API / embeddable projects
TDD	New pure logic, clear contracts
Table-driven tests	Many similar cases
Boundary analysis	Ranges, sizes, parsers
<code>gcov/lcov</code>	Find untested lines/branches
ASan/UBSan	Memory safety and undefined behavior

Testing in C is mostly **discipline and tooling**, not language magic. A small harness plus tables, boundaries, coverage, and sanitizers will catch most defects before they leave your machine.

Debugging Techniques

Debugging is the systematic process of finding and fixing defects. On Linux, the core toolkit is **GDB**, **printf/logging macros**, **Valgrind Memcheck**, **strace/ltrace**, **core dumps**, and compiler sanitizers (see the testing chapter for ASan/UBSan).

Always compile with symbols when debugging:

```
gcc -std=c17 -Wall -Wextra -g -O0 -o program program.c
```

`-O0` keeps the mapping between source lines and machine code straightforward. Optimized builds (`-O2`) can reorder code and eliminate variables, which confuses stepping.

The Debugging Process

1. **Reproduce** — reliable steps beat intermittent luck
 2. **Isolate** — smallest input / fewest modules that fail
 3. **Hypothesize** — what would make the observed symptoms true?
 4. **Observe** — debugger, logs, sanitizers, traces
 5. **Fix** — change the root cause, not the symptom
 6. **Verify** — re-run failing case + nearby regression tests
-

GDB: Session-Style Walkthroughs

GDB (GNU Debugger) inspects a running or crashed process: breakpoints, stepping, printing, stack traces, and watchpoints.

Program under study

```
/* file: buggy_sum.c */
#include <stdio.h>

int sum_range(int n) {
    int total = 0;
    int i;
    /* Off-by-one: should be i <= n for inclusive sum 1..n */
    for (i = 1; i < n; i++) {
        total += i;
    }
    return total;
}

int main(void) {
    int n = 5;
    int s = sum_range(n);
    printf("sum 1..%d = %d (expected 15)\n", n, s);
    return 0;
}
```

```
gcc -std=c17 -Wall -Wextra -g -O0 -o buggy_sum buggy_sum.c
gdb ./buggy_sum
```

Transcript 1 — break, run, print, next, continue

```
$ gdb ./buggy_sum
(gdb) break main
Breakpoint 1 at 0x11a9: file buggy_sum.c, line 14.
(gdb) run
Starting program: .../buggy_sum

Breakpoint 1, main () at buggy_sum.c:14
14          int n = 5;
(gdb) next
15          int s = sum_range(n);
(gdb) print n
$1 = 5
(gdb) step
sum_range (n=5) at buggy_sum.c:5
5          int total = 0;
(gdb) next
7          for (i = 1; i < n; i++) {
```

```

(gdb) print i
$2 = 0
(gdb) next
8             total += i;
(gdb) print i
$3 = 1
(gdb) print total
$4 = 0
(gdb) next
7             for (i = 1; i < n; i++) {
(gdb) print total
$5 = 1
(gdb) break buggy_sum.c:9
Breakpoint 2 at 0x1178: file buggy_sum.c, line 9.
(gdb) continue
Continuing.

Breakpoint 2, sum_range (n=5) at buggy_sum.c:9
9             return total;
(gdb) print total
$6 = 10
(gdb) print n
$7 = 5
(gdb) quit

```

Observation: for $n = 5$, loop runs $i = 1..4$ so total is 10, not 15. Fix: $i \leq n$.

Transcript 2 — backtrace on a crash

```

/* file: crash_null.c */
#include <stdio.h>

void boom(int *p) {
    *p = 42; /* segfault if p is NULL */
}

void middle(void) {
    int *p = NULL;
    boom(p);
}

int main(void) {
    middle();
    return 0;
}

```

```
gcc -std=c17 -Wall -Wextra -g -O0 -o crash_null crash_null.c
gdb ./crash_null
```

```
(gdb) run
Program received signal SIGSEGV, Segmentation fault.
0x0000555555555155 in boom (p=0x0) at crash_null.c:5
5          *p = 42;
(gdb) backtrace
#0 boom (p=0x0) at crash_null.c:5
#1 0x0000555555555170 in middle () at crash_null.c:10
#2 0x0000555555555185 in main () at crash_null.c:14
(gdb) frame 1
#1 middle () at crash_null.c:10
10      boom(p);
(gdb) print p
$1 = (int *) 0x0
(gdb) info locals
p = 0x0
(gdb) list
```

Commands to memorize

Command	Purpose
break main / b file.c:42	Set breakpoint
break foo if x > 5	Conditional breakpoint
run / r	Start program (args: run arg1 arg2)
next / n	Step over
step / s	Step into
continue / c	Resume
print expr / p	Evaluate expression
backtrace / bt	Call stack
frame N	Select stack frame
info locals	Locals in current frame
list	Show source
quit	Exit

Transcript 3 — watchpoints

```
/* file: watch_demo.c */
#include <stdio.h>

int main(void) {
    int counter = 0;
    int i;
    for (i = 0; i < 5; i++) {
```

```

        counter += i;
    }
    printf("%d\n", counter);
    return 0;
}

(gdb) break main
(gdb) run
(gdb) next
(gdb) watch counter
Hardware watchpoint 2: counter
(gdb) continue
Hardware watchpoint 2: counter

Old value = 0
New value = 1
main () at watch_demo.c:8
(gdb) continue
...
(gdb) delete 2

```

Use `watch *pointer` for heap locations after they are allocated. Software watchpoints are slower; hardware watchpoints are limited in count but efficient.

Handy GDB extras

```

(gdb) break 10 if x > 5
(gdb) info breakpoints
(gdb) delete 1
(gdb) set var i = 99
(gdb) finish           # run until current function returns
(gdb) until 40         # continue until line 40 in this frame
(gdb) disassemble main
(gdb) info registers

```

TUI mode (source + assembly panes): start with `gdb -tui ./program` or press `C-x a` inside GDB.

printf Debugging and DEBUG_LOG Macros

Print debugging is fast for I/O-heavy paths and remote systems without an interactive debugger. Make it **toggleable** and **location-aware**.

```

/* file: debug_log.h */
#ifndef DEBUG_LOG_H
#define DEBUG_LOG_H

#include <stdio.h>

/* Compile with -DDEBUG=1 to enable, or #define DEBUG 1 before include */
#ifndef DEBUG
#define DEBUG 0
#endif

#if DEBUG
#define DEBUG_LOG(fmt, ...) \
    do { \
        fprintf(stderr, "[DEBUG] %s:%d:%s(): " fmt "\n", \
            __FILE__, __LINE__, __func__, ##__VA_ARGS__); \
    } while (0)
#else
#define DEBUG_LOG(fmt, ...) \
    do { /* no-op */ } while (0)
#endif

/* Always-on diagnostic (errors, rare paths) */
#define ERROR_LOG(fmt, ...) \
    do { \
        fprintf(stderr, "[ERROR] %s:%d:%s(): " fmt "\n", \
            __FILE__, __LINE__, __func__, ##__VA_ARGS__); \
    } while (0)

#endif

/* file: sum_debug.c */
#include "debug_log.h"

int sum_range(int n) {
    int total = 0;
    int i;
    DEBUG_LOG("enter n=%d", n);
    for (i = 1; i <= n; i++) {
        total += i;
        DEBUG_LOG("i=%d total=%d", i, total);
    }
    DEBUG_LOG("leave total=%d", total);
    return total;
}

```

```

int main(void) {
    int s = sum_range(5);
    printf("sum=%d\n", s);
    return 0;
}

# Logs on
gcc -std=c17 -Wall -Wextra -DDEBUG=1 -o sum_debug sum_debug.c
./sum_debug

# Logs compiled out
gcc -std=c17 -Wall -Wextra -o sum_debug sum_debug.c
./sum_debug

```

Patterns

- Log **entry/exit** of suspect functions with key arguments
- Log **loop indices** only for small **n** or every Nth iteration
- Prefer **stderr** so **stdout** can still be piped/compared
- Never leave secrets (passwords, tokens) in permanent logs
- For binary data, log lengths and a few hex bytes, not megabytes

```

/* Hex dump snippet */
static void dump_bytes(const unsigned char *p, size_t n) {
    size_t i;
    for (i = 0; i < n; i++) {
        fprintf(stderr, "%02x%s", p[i], ((i + 1) % 16 == 0) ? "\n" : " ");
    }
    if (n % 16) {
        fputc('\n', stderr);
    }
}

```

Rubber duck and binary-search debugging

- **Rubber duck**: explain each line aloud; contradictions surface quickly.
- **Binary search**: comment out half the work (or `git bisect`) until the bad change is isolated.

```

git bisect start
git bisect bad           # current broken revision
git bisect good v1.0     # known good tag

```

```
# build + test each checkout, then:
git bisect good    # or: git bisect bad
git bisect reset
```

Valgrind Memcheck: Intentional Leak and Reading the Report

Valgrind Memcheck detects invalid reads/writes, use of uninitialized values, double-free, and leaks.

Example program with a real leak and a bad read

```
/* file: leaky.c */
#include <stdio.h>
#include <stdlib.h>
#include <string.h>

char *make_greeting(const char *name) {
    size_t n = strlen(name);
    char *buf = malloc(n + 8); /* enough for "Hello, " + name + NUL roughly */
    if (buf == NULL) {
        return NULL;
    }
    /* Intentionally wrong size path for demo: we allocated n+8, use carefully */
    sprintf(buf, "Hello, %s", name);
    return buf; /* caller must free - we will "forget" */
}

int main(void) {
    char *msg = make_greeting("Ada");
    if (msg == NULL) {
        return 1;
    }
    printf("%s\n", msg);

    /* Invalid read demo (optional): read one past strlen */
    /* printf("past end byte maybe junk: %d\n", msg[strlen(msg) + 1]); */

    /* BUG: forgot free(msg) - definite leak */
    /* free(msg); */
    return 0;
}
```

```
gcc -std=c17 -Wall -Wextra -g -O0 -o leaky leaky.c
valgrind --tool=memcheck --leak-check=full --show-leak-kinds=all \
  --track-origins=yes ./leaky
```

How to read typical output

```
==12345== HEAP SUMMARY:
==12345==      in use at exit: 11 bytes in 1 blocks
==12345==    total heap usage: 1 allocs, 0 frees, 11 bytes allocated
==12345==
==12345== 11 bytes in 1 blocks are definitely lost in loss record 1 of 1
==12345==    at 0x...: malloc (vg_replace_malloc.c:...)
==12345==    by 0x...: make_greeting (leaky.c:8)
==12345==    by 0x...: main (leaky.c:18)
```

Phrase	Meaning
definitely lost	No pointer to the block remains — real leak
indirectly lost	Lost because a root pointer was lost (e.g. list head)
possibly lost	Interior pointer only (suspicious)
still reachable	Pointers exist at exit (globals, caches) — may be intentional
Invalid read/write	Access outside allocated region or after free
Conditional jump depends on uninitialised value	Branch on garbage
Invalid free / Mismatched free	Double free or <code>free</code> vs <code>delete</code> style issues

Fix for the demo: `free(msg);` before `return 0`.

Common Memcheck flags

```
valgrind --leak-check=full --show-leak-kinds=all \
  --track-origins=yes --errors-for-leak-kinds=all \
  ./program args...
```

Run under the same `-g` build you use for GDB. Suppressions (`.supp` files) hide known third-party noise — use sparingly.

strace and ltrace (brief)

strace — system calls

Shows what the kernel sees: open, read, write, connect, etc.

```
# Trace a simple program
strace -o strace.log ./buggy_sum
# Follow forks (servers, shells)
strace -f -o strace.log ./server
# Only file-related calls
strace -e trace=file ./myprog
# Network-related
strace -e trace=network ./client
```

Example snippet when a file is missing:

```
openat(AT_FDCWD, "config.txt", O_RDONLY) = -1 ENOENT (No such file or directory)
```

That single line often beats minutes of guessing about path logic.

ltrace — library calls

```
ltrace -o ltrace.log ./leaky
# Focus on malloc family
ltrace -e malloc+free+realloc ./leaky
```

Useful when you care about libc (malloc, printf, connect) rather than raw syscalls. Availability varies by distro/package.

Core Dumps: Capture and Analyze

When a process dies with a fatal signal (e.g. SIGSEGV), the kernel can write a **core** image of memory for post-mortem GDB.

Enable dumps

```
# Current shell only
ulimit -c unlimited

# Check limit
ulimit -c
```

```
# On many systems, core pattern is controlled by:
cat /proc/sys/kernel/core_pattern
# Example patterns: core, core.%p, or systemd-coredump pipes
```

If `core_pattern` pipes to `systemd-coredump`, use:

```
coredumpctl list
coredumpctl gdb ./crash_null
# or: coredumpctl dump -o core.crash
```

Produce a core

```
gcc -std=c17 -Wall -Wextra -g -O0 -o crash_null crash_null.c
ulimit -c unlimited
./crash_null
# may create file "core" or "core.<pid>" depending on system
ls -l core*
```

Analyze with GDB

```
gdb ./crash_null core
# or: gdb ./crash_null core.12345
```

```
(gdb) bt
#0 boom (p=0x0) at crash_null.c:5
#1 middle () at crash_null.c:10
#2 main () at crash_null.c:14
(gdb) print p
$1 = (int *) 0x0
(gdb) list
```

Same skills as live debugging: `bt`, `frame`, `print`, `info locals`.

Tips

- The binary path and **build ID** should match the one that crashed
- Keep `-g` builds (or separate debug packages) for production post-mortems
- Never ship secrets in cores; cores can contain full memory

Common Failure Modes

Segmentation faults

Typical causes:

- Null or wild pointer dereference
- Buffer overflow past array/heap end
- Use after free
- Stack overflow (deep recursion)

Workflow: GDB bt → inspect pointer → Valgrind or ASan if memory corruption is subtle.

Logic errors

Program runs, wrong answer. Prefer:

- Unit tests with known golden values
- GDB watchpoints on intermediate results
- Temporary `DEBUG_LOG` of invariants

Memory leaks and corruption

Valgrind / ASan first; then fix ownership (who allocates, who frees).

Race conditions (threads)

```
gcc -std=c17 -Wall -Wextra -g -fsanitize=thread -o races races.c -pthread
./races
```

Combine with careful locking; races are hard to reproduce with `printf` alone.

Static Analysis (quick)

```
gcc -std=c17 -Wall -Wextra -Wpedantic -Wshadow -Wconversion -c file.c
cppcheck --enable=all --inconclusive --std=c17 file.c
```

Static tools do not replace runtime debugging; they reduce the number of bugs that reach GDB.

Debugging Best Practices

1. **Reproduce first** — write a failing test or a minimal input file.
 2. **One change at a time** — random edits hide the real fix.
 3. **Trust tools more than memory** — print the actual value of `errno`, sizes, indices.
 4. **Keep a short log** — hypothesis, experiment, result.
 5. **Prefer `-g -O0` while hunting**, then re-test optimized builds.
 6. **Use version control** — `git bisect` for regressions.
 7. **Take breaks** — fatigue creates new bugs.
-

Exercises

Exercise 1 — GDB fix of `sum_range`

Use the buggy `sum_range` that uses `i < n`. Under GDB, break in the loop, print `i` and `total` for `n = 5`, confirm the off-by-one, fix the loop, and re-run outside GDB so the program prints 15.

Exercise 2 — `DEBUG_LOG` toggle

Add `debug_log.h` to a small multi-function program. Build once with `-DDEBUG=1` and once without. Confirm the non-debug binary produces no debug lines and is smaller / quieter.

Exercise 3 — Valgrind leak

Write a program that allocates a linked list of 10 nodes and “forgets” to free them. Run Memcheck with `--leak-check=full`. Then free the list correctly and show a clean report (All heap blocks were freed -- no leaks are possible).

Exercise 4 — Invalid write

Allocate `malloc(4)` and write 8 bytes. Compare:

1. Normal run (may appear to work — UB)
2. `valgrind ./prog`
3. `gcc -fsanitize=address`

Describe which tool gave the clearest diagnosis.

Exercise 5 — Core dump

Enable `ulimit -c unlimited`, run `crash_null`, open the core in GDB, and paste (or retype) the `bt` output. Identify the null pointer in which frame.

Exercise 6 — strace detective

Write a program that tries to open three paths in order until one succeeds. Hide the working path among missing ones. Use `strace -e openat` (or `trace=file`) to determine which path was used without reading the source.

Summary

Tool	Best for
GDB	Breakpoints, stack traces, watchpoints, live state
DEBUG_LOG	Fast insight, remote/headless runs
Valgrind Memcheck	Leaks, invalid access, uninit values
ASan/UBSan	Fast memory/UB checks (see testing chapter)
strace	“Why did open/connect fail?”
ltrace	Library call sequences
Core + GDB	Post-crash forensics

Debugging is a skill of **observation under control**. Master GDB and Memcheck early; use logs and traces when interactivity is limited; always finish with a regression test so the same bug stays fixed.

Error Handling

Error handling is a critical aspect of robust software development. In C programming, where memory management and low-level operations are common, proper error handling becomes even more important. This chapter explores various error handling strategies and techniques in C.

Introduction to Error Handling

Error handling refers to the process of anticipating, detecting, and resolving programming, application, or communication errors. In C, error handling is primarily manual since the language does not have built-in exception handling mechanisms like some other languages.

Why Error Handling is Important

- **Program Stability:** Prevents crashes and unexpected termination
- **Data Integrity:** Ensures data consistency and prevents corruption
- **User Experience:** Provides meaningful feedback to users
- **Debugging:** Helps identify and resolve issues quickly
- **Security:** Prevents vulnerabilities that could be exploited

Error Handling Strategies in C

Return Codes

Return codes are the most common error handling mechanism in C. Functions return specific values to indicate success or failure.

Basic Return Code Pattern

```
#include <stdio.h>
#include <stdlib.h>

// Function that returns 0 on success, non-zero on error
int divide(int a, int b, int *result) {
    if (b == 0) {
        return -1; // Error: division by zero
    }
}
```

```

        *result = a / b;
        return 0; // Success
    }

    int main() {
        int result;
        int error = divide(10, 2, &result);

        if (error != 0) {
            printf("Error occurred during division\n");
            return 1;
        }

        printf("Result: %d\n", result);
        return 0;
    }

```

Enumerated Return Codes

Using enums for return codes makes the code more readable and maintainable.

```

#include <stdio.h>

typedef enum {
    SUCCESS = 0,
    ERROR_INVALID_INPUT,
    ERROR_DIVISION_BY_ZERO,
    ERROR_MEMORY_ALLOCATION,
    ERROR_FILE_NOT_FOUND
} ErrorCode;

ErrorCode safe_divide(int a, int b, int *result) {
    if (a < 0 || b < 0) {
        return ERROR_INVALID_INPUT;
    }

    if (b == 0) {
        return ERROR_DIVISION_BY_ZERO;
    }

    *result = a / b;
    return SUCCESS;
}

int main() {
    int result;

```



```

    ErrorCode error = safe_divide(10, 2, &result);

    switch (error) {
        case SUCCESS:
            printf("Result: %d\n", result);
            break;
        case ERROR_INVALID_INPUT:
            printf("Error: Invalid input values\n");
            break;
        case ERROR_DIVISION_BY_ZERO:
            printf("Error: Division by zero\n");
            break;
        default:
            printf("Unknown error occurred\n");
            break;
    }

    return error;
}

```

Global Error Variables

C provides global variables like `errno` for system-level error reporting.

Using `errno`

```

#include <stdio.h>
#include <stdlib.h>
#include <errno.h>
#include <string.h>

int main() {
    FILE *file = fopen("nonexistent.txt", "r");

    if (file == NULL) {
        printf("Error opening file: %s\n", strerror(errno));
        return 1;
    }

    fclose(file);
    return 0;
}

```

Custom Error Variables

For application-specific errors, you can create your own global error variables.

```
#include <stdio.h>
#include <string.h>

// Global error variable
static int app_errno = 0;

// Error codes
#define APP_SUCCESS 0
#define APP_ERROR_INVALID_DATA 1
#define APP_ERROR_BUFFER_OVERFLOW 2
#define APP_ERROR_NETWORK_FAILURE 3

// Function to get error description
const char* app_strerror(int error_code) {
    switch (error_code) {
        case APP_SUCCESS:
            return "Success";
        case APP_ERROR_INVALID_DATA:
            return "Invalid data";
        case APP_ERROR_BUFFER_OVERFLOW:
            return "Buffer overflow";
        case APP_ERROR_NETWORK_FAILURE:
            return "Network failure";
        default:
            return "Unknown error";
    }
}

// Function that sets app_errno
int process_data(const char *data) {
    if (data == NULL) {
        app_errno = APP_ERROR_INVALID_DATA;
        return -1;
    }

    if (strlen(data) > 100) {
        app_errno = APP_ERROR_BUFFER_OVERFLOW;
        return -1;
    }

    // Process data...
    app_errno = APP_SUCCESS;
    return 0;
}
```

```

}

int main() {
    int result = process_data(NULL);

    if (result != 0) {
        printf("Error: %s\n", app_strerror(app_errno));
        return 1;
    }

    printf("Data processed successfully\n");
    return 0;
}

```

Advanced Error Handling Techniques

Setjmp and Longjmp

C provides `setjmp` and `longjmp` for non-local jumps, which can be used for exception-like behavior.

```

#include <stdio.h>
#include <stdlib.h>
#include <setjmp.h>
#include <string.h>

// Jump buffer for error handling
static jmp_buf error_buffer;

// Custom exception-like function
void throw_error(const char *message) {
    printf("Error: %s\n", message);
    longjmp(error_buffer, 1);
}

// Function that might fail
void risky_function(const char *data) {
    if (data == NULL) {
        throw_error("Null pointer passed to risky_function");
    }

    if (strlen(data) == 0) {
        throw_error("Empty string passed to risky_function");
    }
}

```

```

    printf("Processing data: %s\n", data);
}

int main() {
    // Set jump point
    if (setjmp(error_buffer) == 0) {
        // Normal execution path
        printf("Calling risky_function...\n");
        risky_function("Hello, World!");
        risky_function(NULL); // This will cause an error
    } else {
        // Error handling path
        printf("Error handled, continuing execution...\n");
    }

    printf("Program continues...\n");
    return 0;
}

```

Resource Acquisition Is Initialization (RAII)

While not native to C, RAII principles can be implemented using goto statements for cleanup.

```

#include <stdio.h>
#include <stdlib.h>
#include <string.h>

typedef struct {
    char *name;
    int *data;
    FILE *file;
} ResourceContainer;

int process_resources() {
    ResourceContainer container = {0};
    int result = -1; // Assume failure

    // Allocate resources
    container.name = malloc(100);
    if (container.name == NULL) {
        fprintf(stderr, "Failed to allocate name buffer\n");
        goto cleanup;
    }

    container.data = malloc(1000 * sizeof(int));
    if (container.data == NULL) {

```

```

        fprintf(stderr, "Failed to allocate data array\n");
        goto cleanup;
    }

    container.file = fopen("data.txt", "w");
    if (container.file == NULL) {
        fprintf(stderr, "Failed to open file\n");
        goto cleanup;
    }

    // Use resources
    strcpy(container.name, "Test Data");
    for (int i = 0; i < 1000; i++) {
        container.data[i] = i * i;
        fprintf(container.file, "%d\n", container.data[i]);
    }

    result = 0; // Success

cleanup:
    // Clean up resources in reverse order
    if (container.file != NULL) {
        fclose(container.file);
    }

    if (container.data != NULL) {
        free(container.data);
    }

    if (container.name != NULL) {
        free(container.name);
    }

    return result;
}

int main() {
    if (process_resources() != 0) {
        printf("Error occurred during resource processing\n");
        return 1;
    }

    printf("Resources processed successfully\n");
    return 0;
}

```

Error Handling Best Practices

Check All Return Values

Always check return values from functions, especially system calls and library functions.

```
#include <stdio.h>
#include <stdlib.h>

int safe_file_operations() {
    FILE *file = fopen("example.txt", "w");
    if (file == NULL) {
        perror("Failed to open file");
        return -1;
    }

    if (fprintf(file, "Hello, World!\n") < 0) {
        perror("Failed to write to file");
        fclose(file);
        return -1;
    }

    if (fclose(file) != 0) {
        perror("Failed to close file");
        return -1;
    }

    return 0;
}
```

Use Assertions for Debugging

Assertions help catch programming errors during development.

```
#include <stdio.h>
#include <assert.h>

int calculate_average(int *array, int size) {
    // Precondition checks
    assert(array != NULL);
    assert(size > 0);

    int sum = 0;
    for (int i = 0; i < size; i++) {
        sum += array[i];
    }
}
```

```

        return sum / size;
    }

    int main() {
        int numbers[] = {1, 2, 3, 4, 5};
        int average = calculate_average(numbers, 5);
        printf("Average: %d\n", average);

        // These would trigger assertions:
        // calculate_average(NULL, 5); // Assertion failed
        // calculate_average(numbers, 0); // Assertion failed

        return 0;
    }

```

Provide Meaningful Error Messages

Error messages should be clear and informative to help with debugging.

```

#include <stdio.h>
#include <stdlib.h>
#include <string.h>

typedef enum {
    ERR_SUCCESS = 0,
    ERR_NULL_POINTER,
    ERR_INVALID_SIZE,
    ERR_MEMORY_ALLOCATION,
    ERR_FILE_OPERATION
} ErrorType;

const char* get_error_message(ErrorType error, const char *context) {
    static char message[256];

    switch (error) {
        case ERR_SUCCESS:
            snprintf(message, sizeof(message), "Success in %s", context);
            break;
        case ERR_NULL_POINTER:
            snprintf(message, sizeof(message), "Null pointer error in %s", context);
            break;
        case ERR_INVALID_SIZE:
            snprintf(message, sizeof(message), "Invalid size parameter in %s", context);
            break;
        case ERR_MEMORY_ALLOCATION:

```

```

        snprintf(message, sizeof(message), "Memory allocation failed in %s", context);
        break;
    case ERR_FILE_OPERATION:
        snprintf(message, sizeof(message), "File operation failed in %s", context);
        break;
    default:
        snprintf(message, sizeof(message), "Unknown error in %s", context);
        break;
}

return message;
}

int safe_memory_copy(char *dest, const char *src, size_t size) {
    if (dest == NULL) {
        fprintf(stderr, "%s\n", get_error_message(ERR_NULL_POINTER, "safe_memory_copy"));
        return -1;
    }

    if (src == NULL) {
        fprintf(stderr, "%s\n", get_error_message(ERR_NULL_POINTER, "safe_memory_copy"));
        return -1;
    }

    if (size == 0) {
        fprintf(stderr, "%s\n", get_error_message(ERR_INVALID_SIZE, "safe_memory_copy"));
        return -1;
    }

    strncpy(dest, src, size - 1);
    dest[size - 1] = '\0';

    return 0;
}

```

Implement Error Logging

Logging errors helps with debugging and monitoring in production environments.

```

#include <stdio.h>
#include <time.h>
#include <stdarg.h>

typedef enum {
    LOG_DEBUG,
    LOG_INFO,

```



```

    LOG_WARNING,
    LOG_ERROR,
    LOG_CRITICAL
} LogLevel;

void log_message(LogLevel level, const char *format, ...) {
    const char *level_strings[] = {
        "DEBUG", "INFO", "WARNING", "ERROR", "CRITICAL"
    };

    // Get current time
    time_t now = time(NULL);
    struct tm *tm_info = localtime(&now);
    char timestamp[20];
    strftime(timestamp, sizeof(timestamp), "%Y-%m-%d %H:%M:%S", tm_info);

    // Print log level and timestamp
    printf("[%s] [%s] ", timestamp, level_strings[level]);

    // Print formatted message
    va_list args;
    va_start(args, format);
    vprintf(format, args);
    va_end(args);

    printf("\n");
}

int divide_with_logging(int a, int b, int *result) {
    log_message(LOG_DEBUG, "divide_with_logging called with a=%d, b=%d", a, b);

    if (b == 0) {
        log_message(LOG_ERROR, "Division by zero attempted");
        return -1;
    }

    *result = a / b;
    log_message(LOG_INFO, "Division successful: %d / %d = %d", a, b, *result);
    return 0;
}

```

Error Handling in Different Contexts

File I/O Error Handling

File operations require careful error handling due to various failure points.

```

#include <stdio.h>
#include <stdlib.h>
#include <errno.h>
#include <string.h>

typedef struct {
    int id;
    char name[50];
    double value;
} Record;

int read_records(const char *filename, Record *records, int max_records, int *record_count)
{
    FILE *file = fopen(filename, "rb");
    if (file == NULL) {
        fprintf(stderr, "Error opening file %s: %s\n", filename, strerror(errno));
        return -1;
    }

    *record_count = 0;
    Record temp;

    while (*record_count < max_records) {
        size_t elements_read = fread(&temp, sizeof(Record), 1, file);
        if (elements_read == 1) {
            records[*record_count] = temp;
            (*record_count)++;
        } else {
            if (feof(file)) {
                // End of file reached normally
                break;
            } else if (ferror(file)) {
                // Error occurred
                fprintf(stderr, "Error reading from file: %s\n", strerror(errno));
                fclose(file);
                return -1;
            }
        }
    }

    if (fclose(file) != 0) {
        fprintf(stderr, "Error closing file: %s\n", strerror(errno));
        return -1;
    }

    return 0;
}

```

Memory Management Error Handling

Proper memory management is crucial in C to prevent leaks and corruption.

```
#include <stdio.h>
#include <stdlib.h>
#include <string.h>

typedef struct {
    char *data;
    size_t size;
    size_t capacity;
} DynamicBuffer;

DynamicBuffer* create_buffer(size_t initial_capacity) {
    DynamicBuffer *buffer = malloc(sizeof(DynamicBuffer));
    if (buffer == NULL) {
        fprintf(stderr, "Failed to allocate buffer structure\n");
        return NULL;
    }

    buffer->data = malloc(initial_capacity);
    if (buffer->data == NULL) {
        fprintf(stderr, "Failed to allocate buffer data\n");
        free(buffer);
        return NULL;
    }

    buffer->size = 0;
    buffer->capacity = initial_capacity;
    return buffer;
}

int append_to_buffer(DynamicBuffer *buffer, const char *data, size_t data_size) {
    if (buffer == NULL || data == NULL) {
        return -1;
    }

    // Check if we need to resize
    if (buffer->size + data_size > buffer->capacity) {
        size_t new_capacity = buffer->capacity * 2;
        while (new_capacity < buffer->size + data_size) {
            new_capacity *= 2;
        }

        char *new_data = realloc(buffer->data, new_capacity);
        if (new_data == NULL) {
```

```

        fprintf(stderr, "Failed to resize buffer\n");
        return -1;
    }

    buffer->data = new_data;
    buffer->capacity = new_capacity;
}

// Copy data
memcpy(buffer->data + buffer->size, data, data_size);
buffer->size += data_size;

return 0;
}

void destroy_buffer(DynamicBuffer *buffer) {
    if (buffer != NULL) {
        free(buffer->data);
        free(buffer);
    }
}

```

Conclusion

Error handling is a fundamental aspect of writing robust C programs. By using appropriate error handling strategies, checking return values, providing meaningful error messages, and implementing proper cleanup procedures, developers can create more reliable and maintainable software. The key is to anticipate potential failure points and handle them gracefully, ensuring that programs can recover from errors or fail safely when recovery is not possible.

Code Quality

Code quality is a measure of how well software code meets specified requirements and satisfies user needs. High-quality code is not only functional but also maintainable, readable, and efficient. This chapter explores various aspects of code quality in C programming and best practices to achieve it — with **complete programs** you can compile under a strict warning set.

```
gcc -std=c17 -Wall -Wextra -Wpedantic -Werror -o prog prog.c
```

Treat `-Werror` as a quality gate once the code is past the first sketch.

Introduction to Code Quality

Code quality encompasses multiple dimensions that contribute to the overall excellence of software. It's not just about whether the code works, but also about how well it works, how easy it is to maintain, and how reliably it performs over time.

Dimensions of Code Quality

1. **Correctness:** The code produces the expected results
2. **Reliability:** The code performs consistently under various conditions
3. **Efficiency:** The code uses resources optimally
4. **Maintainability:** The code is easy to modify and extend
5. **Readability:** The code is easy to understand
6. **Testability:** The code can be easily tested
7. **Portability:** The code works across different platforms
8. **Security:** The code is resistant to vulnerabilities

Coding Standards and Style Guidelines

Consistent coding standards improve code readability and maintainability. They provide a common framework for developers to follow.

Naming Conventions

```
// Good naming conventions
int user_age;                // Descriptive variable names
double calculate_circle_area(double radius); // Descriptive function names
const int MAX_BUFFER_SIZE = 1024; // Constants in uppercase

// Avoid cryptic names
int x;                       // What does x represent?
int f(int a);                // What does f do?
```

Code Formatting

Consistent indentation and spacing improve code readability:

```
// Well-formatted code
int calculate_sum(int array[], int size) {
    int sum = 0;

    for (int i = 0; i < size; i++) {
        sum += array[i];
    }

    return sum;
}

// Poorly formatted code (avoid this)
int calculate_sum(int array[],int size){
int sum=0;
for(int i=0;i<size;i++){
sum+=array[i];}
return sum;}
```

Commenting Best Practices

Comments should explain the “why” rather than the “what”:

```
// Good comments - explain the reasoning
void initialize_buffer(Buffer *buf) {
    // Initialize to zero to prevent undefined behavior
    buf->size = 0;
    buf->capacity = INITIAL_CAPACITY;

    // Allocate memory with extra space to reduce reallocations
```

```

    buf->data = malloc(buf->capacity * sizeof(int));
}

// Avoid redundant comments
int x = 5; // Set x to 5 (obvious from code)

```

Code Review Practices

Code reviews are essential for maintaining code quality and sharing knowledge among team members.

Benefits of Code Reviews

- **Bug Detection:** Catch defects before they reach production
- **Knowledge Sharing:** Spread expertise across the team
- **Consistency:** Ensure adherence to coding standards
- **Mentoring:** Help junior developers improve
- **Design Feedback:** Get input on architectural decisions

Code Review Checklist

1. **Functionality:** Does the code meet requirements?
2. **Correctness:** Are there any logical errors?
3. **Readability:** Is the code easy to understand?
4. **Maintainability:** Is the code easy to modify?
5. **Performance:** Are there any efficiency concerns?
6. **Security:** Are there potential vulnerabilities?
7. **Standards:** Does the code follow established guidelines?

Example Code Review

```

// Before review
int process(int*a,int b){int c=0;for(int d=0;d<b;d++){c+=a[d];}return c;}

// After review
/**
 * Calculate the sum of integers in an array.
 *
 * @param array Array of integers to sum
 * @param size Number of elements in the array
 * @return Sum of all elements in the array
 */
int calculate_array_sum(int *array, int size) {

```

```

// Validate input parameters
if (array == NULL || size < 0) {
    return -1; // Error case
}

int sum = 0;
for (int i = 0; i < size; i++) {
    sum += array[i];
}

return sum;
}

```

Static Analysis Tools

Static analysis tools examine source code without executing it to find potential issues.

GCC with Warnings

GCC provides extensive warning options to catch common issues:

```

# Compile with comprehensive warnings
gcc -Wall -Wextra -Werror -pedantic -std=c99 program.c

# Explanation of flags:
# -Wall: Enable most warning messages
# -Wextra: Enable additional warnings
# -Werror: Treat warnings as errors
# -pedantic: Issue warnings for non-standard C
# -std=c99: Specify C standard

```

Cppcheck

Cppcheck is a static analysis tool specifically designed for C/C++:

```

# Basic usage
cppcheck program.c

# More thorough analysis
cppcheck --enable=all --inconclusive program.c

# Check specific directories
cppcheck --enable=all src/

```


Clang Static Analyzer

The Clang Static Analyzer detects bugs and potential issues:

```
# Analyze with Clang
clang --analyze program.c

# Generate HTML reports
scan-build make
```

Code Formatting and Style Tools

Automated tools help maintain consistent code style across projects.

Clang-Format

Clang-Format automatically formats C code according to specified rules:

```
# Format a single file
clang-format -i program.c

# Format multiple files
find . -name "*.c" -o -name "*.h" | xargs clang-format -i

# Use a specific style
clang-format -style=Google -i program.c
```

AStyle (Artistic Style)

AStyle is another popular code formatter:

```
# Format with Allman style
astyle --style=allman *.c *.h

# Format with K&R style
astyle --style=kr *.c *.h
```

Test Coverage Analysis

Test coverage measures how much of your code is exercised by tests.

GCC with Coverage

GCC can generate coverage information:

```
# Compile with coverage instrumentation
gcc -fprofile-arcs -ftest-coverage -o program program.c

# Run the program
./program

# Generate coverage report
gcov program.c

# Generate HTML report with lcov
lcov --capture --directory . --output-file coverage.info
genhtml coverage.info --output-directory coverage_report
```

Coverage Metrics

1. **Statement Coverage:** Percentage of statements executed
2. **Branch Coverage:** Percentage of branches taken
3. **Function Coverage:** Percentage of functions called
4. **Line Coverage:** Percentage of lines executed

Performance Analysis

Performance analysis helps identify bottlenecks and optimize code.

Profiling with gprof

gprof provides function-level profiling information:

```
# Compile with profiling enabled
gcc -pg -o program program.c

# Run the program
./program

# Generate profiling report
gprof program gmon.out > analysis.txt
```

Profiling with perf

perf is a powerful Linux profiling tool:

```
# Profile a program
perf record ./program

# Analyze results
perf report

# View call graph
perf record -g ./program
perf report -g
```

Security Best Practices

Security should be considered throughout the development process.

Input Validation

Always validate input to prevent buffer overflows and injection attacks:

```
#include <stdio.h>
#include <string.h>

// Unsafe function (vulnerable to buffer overflow)
void unsafe_copy(char *dest, char *src) {
    strcpy(dest, src); // Dangerous!
}

// Safe function with bounds checking
int safe_copy(char *dest, size_t dest_size, const char *src) {
    if (dest == NULL || src == NULL || dest_size == 0) {
        return -1; // Error
    }

    size_t src_len = strlen(src);
    if (src_len >= dest_size) {
        return -1; // Buffer too small
    }

    strcpy(dest, src);
    return 0; // Success
}
```

Memory Safety

Proper memory management prevents common security vulnerabilities:

```
#include <stdio.h>
#include <stdlib.h>
#include <string.h>

// Secure memory handling
char* duplicate_string(const char *source) {
    if (source == NULL) {
        return NULL;
    }

    size_t len = strlen(source);
    char *copy = malloc(len + 1);

    if (copy == NULL) {
        return NULL; // Allocation failed
    }

    strcpy(copy, source);
    return copy;
}

void free_string(char *str) {
    if (str != NULL) {
        free(str);
        str = NULL; // Prevent dangling pointer (local copy only)
    }
}
```

Secure Coding Practices

1. **Use Safe Functions:** Prefer `strncpy` over `strcpy`, `snprintf` over `sprintf`
2. **Validate Pointers:** Always check for NULL pointers
3. **Bounds Checking:** Verify array indices and buffer sizes
4. **Integer Overflow:** Check for arithmetic overflow
5. **Format Strings:** Use constant format strings with `printf/scanf`

Documentation and Code Comments

Good documentation is essential for code quality and maintainability.

Self-Documenting Code

Write code that is easy to understand without excessive comments:

```
// Self-documenting approach
typedef enum {
    USER_STATUS_ACTIVE,
    USER_STATUS_INACTIVE,
    USER_STATUS_SUSPENDED
} UserStatus;

typedef struct {
    char username[50];
    UserStatus status;
    time_t last_login;
} User;

int is_user_active(const User *user) {
    return user != NULL && user->status == USER_STATUS_ACTIVE;
}

// Avoid unclear code
int check_user(int a, int b) { // What do a and b represent?
    return a == 1 && b > 0;    // What does this check mean?
}
```

API Documentation

Document public APIs with clear descriptions and examples:

```
/**
 * @brief Calculate the factorial of a non-negative integer.
 *
 * This function calculates the factorial of a given non-negative integer
 * using an iterative approach. The factorial of n is the product of all
 * positive integers less than or equal to n.
 *
 * @param n The non-negative integer to calculate factorial for
 * @return The factorial of n, or -1 if n is negative
 *
 * @note This function can overflow for large values of n (n > 20)
 *
 * Example usage:
 * @code
 * int result = factorial(5); // Returns 120
 * @endcode
 */
```

```

    */
    int factorial(int n) {
        if (n < 0) {
            return -1; // Error case
        }

        if (n == 0 || n == 1) {
            return 1;
        }

        int result = 1;
        for (int i = 2; i <= n; i++) {
            result *= i;
        }

        return result;
    }

```

Refactoring and Code Improvement

Regular refactoring improves code quality over time.

Code Smells to Avoid

1. **Long Functions:** Functions that are too long and complex
2. **Duplicated Code:** Same or similar code in multiple places
3. **Complex Conditionals:** Overly complex if/else statements
4. **Magic Numbers:** Unexplained numeric constants in code
5. **Inappropriate Intimacy:** Functions that access internal data directly

Refactoring Example

```

// Before refactoring - complex and hard to understand
double calculate_discount(double price, int customer_type, int quantity) {
    if (customer_type == 1) {
        if (quantity > 100) {
            return price * 0.2;
        } else if (quantity > 50) {
            return price * 0.15;
        } else {
            return price * 0.1;
        }
    }
    } else if (customer_type == 2) {

```

```

        if (quantity > 100) {
            return price * 0.3;
        } else if (quantity > 50) {
            return price * 0.25;
        } else {
            return price * 0.2;
        }
    } else {
        if (quantity > 100) {
            return price * 0.1;
        } else if (quantity > 50) {
            return price * 0.05;
        } else {
            return 0;
        }
    }
}

// After refactoring - clean and maintainable
typedef enum {
    CUSTOMER_REGULAR,
    CUSTOMER_PREMIUM,
    CUSTOMER_VIP
} CustomerType;

typedef struct {
    double threshold_100;
    double threshold_50;
    double base_rate;
} DiscountRates;

static const DiscountRates DISCOUNT_RATES[] = {
    [CUSTOMER_REGULAR] = {0.2, 0.15, 0.1},
    [CUSTOMER_PREMIUM] = {0.3, 0.25, 0.2},
    [CUSTOMER_VIP] = {0.1, 0.05, 0.0}
};

double calculate_discount(double price, CustomerType customer_type, int quantity) {
    if (customer_type < 0 || customer_type >= sizeof(DISCOUNT_RATES)) {
        return 0; // Invalid customer type
    }

    const DiscountRates *rates = &DISCOUNT_RATES[customer_type];

    if (quantity > 100) {
        return price * rates->threshold_100;
    }

```

```

    } else if (quantity > 50) {
        return price * rates->threshold_50;
    } else {
        return price * rates->base_rate;
    }
}

```

Full Programs: Quality Gates You Can Run Locally

Program 1 — Strict-warning “clean” module (clamp_lib)

A small pure function plus a self-test binary. Clean under -Werror.

```

/* file: clamp.c */
#include <stdio.h>
#include <stdlib.h>
#include <limits.h>

/* Returns x clamped to [lo, hi]. If lo > hi, returns lo (documented choice). */
int clamp(int x, int lo, int hi) {
    if (lo > hi) {
        return lo;
    }
    if (x < lo) {
        return lo;
    }
    if (x > hi) {
        return hi;
    }
    return x;
}

#ifdef CLAMP_MAIN
int main(void) {
    struct {
        int x, lo, hi, want;
    } cases[] = {
        {5, 0, 10, 5},
        {-1, 0, 10, 0},
        {99, 0, 10, 10},
        {0, 0, 0, 0},
        {INT_MAX, 0, 10, 10},
        {INT_MIN, -10, 10, -10},
    };
    size_t i;

```



```

    int failed = 0;

    for (i = 0; i < sizeof cases / sizeof cases[0]; i++) {
        int got = clamp(cases[i].x, cases[i].lo, cases[i].hi);
        if (got != cases[i].want) {
            fprintf(stderr, "case %zu: clamp(%d,%d,%d)=%d want %d\n",
                i, cases[i].x, cases[i].lo, cases[i].hi, got, cases[i].want);
            failed = 1;
        }
    }
    if (failed) {
        return EXIT_FAILURE;
    }
    puts("clamp: all cases passed");
    return EXIT_SUCCESS;
}
#endif

gcc -std=c17 -Wall -Wextra -Wpedantic -Werror -DCLAMP_MAIN -o clamp_test clamp.c
./clamp_test

```

Program 2 — Warning archaeology (smell.c → clean.c)

Start with a deliberately smelly program, record warnings, then fix.

```

/* file: smell.c - DO NOT ship; for warning lab only */
#include <stdio.h>
#include <string.h>

int helper(); /* no prototype details */

int main() {
    int unused;
    char buf[8];
    char *p = 0;

    strcpy(buf, "toolongforbuffer"); /* overflow risk */
    if (p)
        printf("%s\n", p);
    printf("%d\n", helper(1, 2));
    return 0;
}

int helper(int a, int b) {
    return a + b;
}

```

```
}
```

```
gcc -std=c17 -Wall -Wextra -Wpedantic -c smell.c 2> smell_warnings.txt  
cat smell_warnings.txt
```

Clean rewrite:

```
/* file: clean.c */  
#include <stdio.h>  
#include <string.h>  
  
static int helper(int a, int b);  
  
int main(void) {  
    char buf[8];  
  
    /* bounded copy + explicit truncation detection */  
    if (snprintf(buf, sizeof buf, "%s", "toolongforbuffer") >= (int)sizeof buf) {  
        fprintf(stderr, "warning: truncated\n");  
    }  
    printf("buf=\"%s\"\n", buf);  
    printf("sum=%d\n", helper(1, 2));  
    return 0;  
}  
  
static int helper(int a, int b) {  
    return a + b;  
}  
  
gcc -std=c17 -Wall -Wextra -Wpedantic -Werror -o clean clean.c  
./clean
```

Program 3 — Complexity reduction: nested ifs → early return

```
/* file: access.c */  
#include <stdio.h>  
#include <stdbool.h>  
  
struct User {  
    bool active;  
    bool banned;  
    int age;  
    int role; /* 0 guest, 1 user, 2 admin */  
};
```

```

/* Opaque quality: one exit path per failure reason, easy to test. */
const char *can_edit(const struct User *u) {
    if (u == NULL) {
        return "null user";
    }
    if (!u->active) {
        return "inactive";
    }
    if (u->banned) {
        return "banned";
    }
    if (u->age < 13) {
        return "underage";
    }
    if (u->role < 1) {
        return "guest cannot edit";
    }
    return NULL; /* allowed */
}

int main(void) {
    struct User samples[] = {
        {true, false, 20, 1},
        {false, false, 20, 1},
        {true, true, 20, 2},
        {true, false, 10, 1},
        {true, false, 25, 0},
    };
    size_t i;

    for (i = 0; i < sizeof samples / sizeof samples[0]; i++) {
        const char *err = can_edit(&samples[i]);
        if (err == NULL) {
            printf("sample %zu: OK\n", i);
        } else {
            printf("sample %zu: DENY (%s)\n", i, err);
        }
    }
    return 0;
}

```

```

gcc -std=c17 -Wall -Wextra -Wpedantic -Werror -o access access.c
./access

```

Program 4 — Static assertions and size contracts

```
/* file: layout_check.c */
#include <stdio.h>
#include <stddef.h>
#include <stdint.h>
#include <assert.h>

struct PacketHdr {
    uint16_t type;
    uint16_t len;
    uint32_t id;
};

/* Catch ABI surprises at compile time */
_Static_assert(sizeof(struct PacketHdr) == 8, "PacketHdr must be 8 bytes");
_Static_assert(offsetof(struct PacketHdr, id) == 4, "id offset");

int main(void) {
    struct PacketHdr h = {1, 16, 42};
    printf("type=%u len=%u id=%u size=%zu\n",
          h.type, h.len, h.id, sizeof h);
    return 0;
}

gcc -std=c17 -Wall -Wextra -Wpedantic -Werror -o layout_check layout_check.c
./layout_check
```

Program 5 — Local quality script (check.sh + target)

```
/* file: avg.c */
#include <stdio.h>
#include <stdlib.h>

double average(const int *a, size_t n) {
    size_t i;
    long sum = 0;

    if (a == NULL || n == 0) {
        return 0.0;
    }
    for (i = 0; i < n; i++) {
        sum += a[i];
    }
}
```

```

        return (double)sum / (double)n;
    }

    int main(void) {
        int v[] = {1, 2, 3, 4};
        printf("%.2f\n", average(v, 4));
        return 0;
    }
}

# file: check.sh
#!/usr/bin/env bash
set -euo pipefail
CFLAGS=(-std=c17 -Wall -Wextra -Wpedantic -Werror -g)
gcc "${CFLAGS[@]}" -o avg avg.c
./avg
# optional tools if installed:
if command -v cppcheck >/dev/null; then
    cppcheck --enable=warning,style --error-exitcode=1 avg.c
fi
if command -v valgrind >/dev/null; then
    valgrind --error-exitcode=1 --leak-check=full ./avg
fi
echo "quality gate: OK"

chmod +x check.sh
./check.sh

```

Program 6 — Review checklist applied to a PR-sized change

Before you merge any non-trivial C change, walk this list (paste into PR template if you want):

```

[ ] Builds with -std=c17 -Wall -Wextra (and -Werror in CI)
[ ] No new compiler warnings
[ ] Every public function documents ownership (who frees?)
[ ] Error paths free / close what success paths free / close
[ ] No unbounded strcpy/sprintf/gets
[ ] Integer sizes / casts reviewed for truncation
[ ] Tests or a tiny main demo for the changed logic
[ ] Sanitizer or Valgrind run if memory/pointers touched

```

Example “before” (fails several checks):

```

char *load(const char *path) {
    FILE *fp = fopen(path, "r");
    char *buf = malloc(1024);

```

```

    fread(buf, 1, 1024, fp); /* no checks, no NUL, leak on error */
    return buf;
}

```

Example “after”:

```

/* file: load_text.c */
#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <errno.h>

/* Caller frees the returned buffer. Returns NULL on error. */
char *load_text(const char *path, size_t max_n) {
    FILE *fp;
    char *buf;
    size_t n;

    if (path == NULL || max_n == 0) {
        return NULL;
    }
    fp = fopen(path, "rb");
    if (fp == NULL) {
        return NULL;
    }
    buf = malloc(max_n + 1);
    if (buf == NULL) {
        fclose(fp);
        return NULL;
    }
    n = fread(buf, 1, max_n, fp);
    if (ferror(fp)) {
        free(buf);
        fclose(fp);
        return NULL;
    }
    buf[n] = '\0';
    fclose(fp);
    return buf;
}

int main(int argc, char *argv[]) {
    char *s;

    if (argc != 2) {
        fprintf(stderr, "Usage: %s <file>\n", argv[0]);
        return EXIT_FAILURE;
    }
}

```

```

    }
    s = load_text(argv[1], 4096);
    if (s == NULL) {
        fprintf(stderr, "load failed: %s\n", strerror(errno));
        return EXIT_FAILURE;
    }
    fputs(s, stdout);
    free(s);
    return EXIT_SUCCESS;
}

gcc -std=c17 -Wall -Wextra -Wpedantic -Werror -o load_text load_text.c
echo 'hello quality' > t.txt
./load_text t.txt

```

Continuous Integration and Quality Gates

Continuous Integration (CI) helps maintain code quality by automatically testing and validating changes.

CI Pipeline Components

1. **Automated Builds:** Compile code on every change
2. **Static Analysis:** Run code quality tools automatically
3. **Unit Tests:** Execute test suite on every change
4. **Code Coverage:** Measure test coverage
5. **Security Scans:** Check for vulnerabilities
6. **Deployment:** Automatically deploy passing builds

Example CI Configuration

```

# .github/workflows/ci.yml
name: C CI

on: [push, pull_request]

jobs:
  build:
    runs-on: ubuntu-latest

    steps:

```

```

- uses: actions/checkout@v4

- name: Install dependencies
  run: |
    sudo apt-get update
    sudo apt-get install -y gcc valgrind cppcheck

- name: Compile with warnings as errors
  run: gcc -std=c17 -Wall -Wextra -Wpedantic -Werror -o program program.c

- name: Static analysis
  run: cppcheck --enable=warning,style --error-exitcode=1 .

- name: Run tests
  run: ./run_tests.sh

- name: Memory check
  run: valgrind --error-exitcode=1 --leak-check=full ./program

```

Local gate CI gate

If it does not pass `check.sh` on your laptop, it should not pass in CI. Keep the same flags in both places.

Exercises

```
gcc -std=c17 -Wall -Wextra -Wpedantic -Werror -o exN exN.c
```

1. **Warning farm** — Write a 20-line file that triggers at least four distinct `-Wall -Wextra` diagnostics. Fix them one by one until `-Werror` is clean. Keep the smelly version in a comment.
2. **clamp table** — Add five more edge cases to `clamp_test` (including `lo > hi`) and make the policy explicit in a comment.
3. **Refactor nest** — Take a function with three nested `ifs` from your own code (or invent one) and rewrite with early returns. Confirm behavior with a small table of cases.
4. **_Static_assert** — Define a struct with intentional padding; assert `sizeof` and an `offsetof`. Change a field type and watch the assert fire.
5. **load_text growth** — Extend `load_text` to grow with `realloc` until EOF (cap at e.g. 1 MiB). Prove no leaks under Valgrind or ASan.

6. **check.sh for multi-file** — Point the script at `main.c + util.c`, fail the script if any command fails, and add an optional sanitizer build.
7. **Review role-play** — Review `smell.c` as if it were a PR. Write five bullet comments a reviewer should leave.
8. **Complexity budget** — Count decision points (`if/case/` or `||/for/while`) in one function; if over 10, split it.

Conclusion

Code quality is not a one-time achievement but an ongoing commitment. By following established coding standards, conducting regular code reviews, using automated tools, and continuously refactoring, developers can maintain high-quality codebases that are robust, maintainable, and secure.

Practical baseline for this book:

Gate	Command / habit
Warnings	<code>-Wall -Wextra -Wpedantic</code> (+ <code>-Werror</code> in CI)
Memory	ASan/UBSan or Valgrind on pointer-heavy changes
Structure	Early returns, clear ownership, small functions
Proof	Table-driven tests or a tiny self-check <code>main</code>
Process	Checklist + same flags locally and in CI

The investment in code quality pays dividends throughout the software lifecycle, reducing bugs, improving maintainability, and enhancing the overall developer experience.

Module 14: Testing

This module turns “it seems to work” into **evidence**. You will unit-test pure C with portable harnesses, drive design with TDD, hunt bugs with GDB and Valgrind, handle errors without exceptions, and raise code quality with warnings, sanitizers, and review habits.

Environment: Linux/POSIX, `gcc -std=c17 -Wall -Wextra`, plus optional `gcov/lcov`, Valgrind, and sanitizers.

Learning goals

By the end of this module you should be able to:

- Write **automated unit tests** with `assert`, a tiny custom macro harness, or a Unity-style runner
- Apply **TDD** (red → green → refactor) to small pure functions
- Design **boundary** and **table-driven** cases for integers and strings
- Measure **coverage** with `gcov/lcov` and interpret gaps
- Run tests under **AddressSanitizer** and **UndefinedBehaviorSanitizer**
- Drive **GDB** sessions (break, print, backtrace, watch) and read **core dumps**
- Detect leaks and invalid access with **Valgrind Memcheck**
- Use `strace/ltrace` and structured `DEBUG__LOG` macros
- Choose return codes, `errno`, and cleanup patterns for robust APIs
- Apply practical **quality** gates: `-Werror`, static analysis, review checklists, local `check.sh`

Chapters (map)

#	Chapter	What you will practice
1	Testing Methodologies	Unit harnesses, Unity-style suites, TDD (<code>is_leap_year / clamp</code>), boundaries, table-driven tests, <code>gcov/lcov</code> , ASan/UBSan
2	Debugging Techniques	GDB transcripts, <code>DEBUG_LOG</code> , Valgrind, <code>strace/ltrace</code> , core dumps, systematic isolation
3	Error Handling	Return codes, enums, <code>errno</code> , cleanup on failure, partial-success pitfalls
4	Code Quality	Warning archaeology, early-return design, <code>_Static_assert</code> , load/ownership patterns, CI-style local gates

Suggested order

1. **Testing Methodologies** first — build a harness and a few pure functions you trust.
2. **Debugging Techniques** — deliberately break code and recover with GDB/Valgrind.
3. **Error Handling** — make APIs fail loudly and cleanly.
4. **Code Quality** — lock in habits that keep the above sustainable (`-Werror`, checklists, scripts).

Tooling cheat sheet

```
# Unit test binary
gcc -std=c17 -Wall -Wextra -g -o test_foo test_foo.c
./test_foo

# Coverage
gcc -std=c17 -Wall -Wextra -g --coverage -o test_foo test_foo.c
./test_foo && gcov test_foo.c

# Sanitizers
gcc -std=c17 -Wall -Wextra -g -fsanitize=address,undefined -o test_foo test_foo.c

# Strict quality compile
gcc -std=c17 -Wall -Wextra -Wpedantic -Werror -o prog prog.c

# Debug
gcc -std=c17 -Wall -Wextra -g -O0 -o prog prog.c
gdb ./prog
valgrind --leak-check=full ./prog
```

How this module connects

- **Earlier modules** (functions, pointers, memory, file I/O) supply the bugs and APIs you will test.
- **Later practice** (projects, networking) should inherit the same habit: small tests for pure logic, integration checks for I/O and sockets.
- Treat Module 14 as a **skill overlay**: re-test older chapters' utilities under a harness once you have one.

Module self-check

Each chapter ends with exercises. A solid bar for Module 14:

- ☐ A multi-test binary exits non-zero on failure
- ☐ At least one function was written strictly with TDD
- ☐ A table of 8 cases drives a pure function
- ☐ `gcov` (or HTML via `lcov`) reviewed once
- ☐ ASan or Valgrind used to find a planted memory bug
- ☐ A GDB `backtrace` obtained from a segfault or core file
- ☐ A “smelly” file cleaned until `-Werror` is silent
- ☐ One PR-style checklist applied to a small change

Work through the chapters in order, keep every example compiling with `-Wall -Wextra`, and prefer **complete programs** you can re-run when you change the code under test.

Estimated time

- Reading: 5–8 hours
- Labs (tests + GDB + sanitizers + quality gates): 8–12 hours

Next

Apply these habits to **Module 15 (Advanced)** projects and to any workbook project: ship tests and a quality script with the feature, not after.

Advanced

Module 15.1: Multithreading in C

Introduction

Multithreading is a powerful technique that allows programs to perform multiple tasks concurrently within a single process. In C programming, multithreading can significantly improve performance for CPU-intensive applications and enable responsive user interfaces. This module explores the fundamentals of multithreading using the POSIX Threads (pthreads) library, which is the standard threading API on Unix-like systems.

Learning Objectives

By the end of this module, you will be able to:

- Understand the concepts of processes and threads
- Create and manage threads using the pthread library
- Implement thread synchronization mechanisms
- Handle thread communication and data sharing
- Debug and optimize multithreaded applications
- Apply best practices for thread safety

1. Fundamentals of Multithreading

1.1 Processes vs Threads

A **process** is an independent program execution unit that has its own memory space, file descriptors, and system resources. A **thread** is a lightweight subprocess that exists within a process and shares the process's resources.

Key differences:

- Processes have separate memory spaces; threads share memory
- Process creation is expensive; thread creation is relatively cheap
- Inter-process communication requires special mechanisms; threads can communicate directly through shared memory
- Processes are isolated; threads can affect each other if not properly synchronized

1.2 Thread Lifecycle

A thread goes through several states during its lifetime:

1. **New:** Thread is being created
2. **Runnable:** Thread is ready to run
3. **Running:** Thread is executing
4. **Blocked:** Thread is waiting for a resource
5. **Terminated:** Thread has finished execution

1.3 Thread Attributes

Threads can have various attributes that control their behavior: - **Detached vs Joinable**: Detached threads clean up automatically; joinable threads require explicit joining - **Scheduling policy**: Determines how the thread is scheduled (SCHED_FIFO, SCHED_RR, SCHED_OTHER) - **Stack size**: Controls the amount of stack memory allocated to the thread

2. Creating and Managing Threads

2.1 Basic Thread Creation

The pthread library provides the `pthread_create()` function to create threads:

```
#include <pthread.h>
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>

// Thread function
void* thread_function(void* arg) {
    int thread_id = *(int*)arg;
    printf("Hello from thread %d\n", thread_id);

    // Simulate some work
    sleep(1);

    printf("Thread %d finishing\n", thread_id);
    return NULL;
}

int main() {
    pthread_t thread1, thread2;
    int id1 = 1, id2 = 2;

    // Create threads
    if (pthread_create(&thread1, NULL, thread_function, &id1) != 0) {
        perror("pthread_create failed");
        exit(EXIT_FAILURE);
    }

    if (pthread_create(&thread2, NULL, thread_function, &id2) != 0) {
        perror("pthread_create failed");
        exit(EXIT_FAILURE);
    }

    // Wait for threads to complete
```

```

    pthread_join(thread1, NULL);
    pthread_join(thread2, NULL);

    printf("All threads completed\n");
    return 0;
}

```

2.2 Thread Arguments and Return Values

Threads can receive arguments and return values:

```

#include <pthread.h>
#include <stdio.h>
#include <stdlib.h>

typedef struct {
    int start;
    int end;
    long long sum;
} thread_data_t;

void* calculate_sum(void* arg) {
    thread_data_t* data = (thread_data_t*)arg;

    data->sum = 0;
    for (int i = data->start; i <= data->end; i++) {
        data->sum += i;
    }

    return NULL;
}

int main() {
    pthread_t thread1, thread2;
    thread_data_t data1 = {1, 50000000, 0};
    thread_data_t data2 = {50000001, 100000000, 0};

    // Create threads
    pthread_create(&thread1, NULL, calculate_sum, &data1);
    pthread_create(&thread2, NULL, calculate_sum, &data2);

    // Wait for completion
    pthread_join(thread1, NULL);
    pthread_join(thread2, NULL);

    long long total_sum = data1.sum + data2.sum;
}

```



```

    printf("Sum of 1 to 100000000: %lld\n", total_sum);

    return 0;
}

```

3. Thread Synchronization

3.1 Race Conditions

A race condition occurs when multiple threads access shared data concurrently, and the outcome depends on the timing of their execution. Consider this problematic code:

```

#include <pthread.h>
#include <stdio.h>
#include <stdlib.h>

int shared_counter = 0;

void* increment_counter(void* arg) {
    for (int i = 0; i < 1000000; i++) {
        shared_counter++; // Race condition here!
    }
    return NULL;
}

int main() {
    pthread_t thread1, thread2;

    pthread_create(&thread1, NULL, increment_counter, NULL);
    pthread_create(&thread2, NULL, increment_counter, NULL);

    pthread_join(thread1, NULL);
    pthread_join(thread2, NULL);

    printf("Expected: 2000000, Actual: %d\n", shared_counter);
    return 0;
}

```

3.2 Mutexes

Mutexes (mutual exclusions) are the most common synchronization primitive:

```

#include <pthread.h>
#include <stdio.h>

```

```

#include <stdlib.h>

int shared_counter = 0;
pthread_mutex_t counter_mutex = PTHREAD_MUTEX_INITIALIZER;

void* increment_counter(void* arg) {
    for (int i = 0; i < 1000000; i++) {
        pthread_mutex_lock(&counter_mutex);
        shared_counter++;
        pthread_mutex_unlock(&counter_mutex);
    }
    return NULL;
}

int main() {
    pthread_t thread1, thread2;

    pthread_create(&thread1, NULL, increment_counter, NULL);
    pthread_create(&thread2, NULL, increment_counter, NULL);

    pthread_join(thread1, NULL);
    pthread_join(thread2, NULL);

    printf("Expected: 2000000, Actual: %d\n", shared_counter);

    pthread_mutex_destroy(&counter_mutex);
    return 0;
}

```

3.3 Condition Variables

Condition variables allow threads to wait for specific conditions:

```

#include <pthread.h>
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>

int buffer = -1;
int count = 0;
pthread_mutex_t buffer_mutex = PTHREAD_MUTEX_INITIALIZER;
pthread_cond_t buffer_not_empty = PTHREAD_COND_INITIALIZER;
pthread_cond_t buffer_not_full = PTHREAD_COND_INITIALIZER;

void* producer(void* arg) {
    for (int i = 0; i < 10; i++) {

```

```

    pthread_mutex_lock(&buffer_mutex);

    // Wait while buffer is full
    while (count == 1) {
        pthread_cond_wait(&buffer_not_full, &buffer_mutex);
    }

    buffer = i;
    count = 1;
    printf("Produced: %d\n", i);

    // Signal consumer that buffer is not empty
    pthread_cond_signal(&buffer_not_empty);
    pthread_mutex_unlock(&buffer_mutex);

    sleep(1); // Simulate work
}
return NULL;
}

void* consumer(void* arg) {
    for (int i = 0; i < 10; i++) {
        pthread_mutex_lock(&buffer_mutex);

        // Wait while buffer is empty
        while (count == 0) {
            pthread_cond_wait(&buffer_not_empty, &buffer_mutex);
        }

        int item = buffer;
        count = 0;
        printf("Consumed: %d\n", item);

        // Signal producer that buffer is not full
        pthread_cond_signal(&buffer_not_full);
        pthread_mutex_unlock(&buffer_mutex);

        sleep(1); // Simulate work
    }
    return NULL;
}

int main() {
    pthread_t producer_thread, consumer_thread;

    pthread_create(&producer_thread, NULL, producer, NULL);
    pthread_create(&consumer_thread, NULL, consumer, NULL);

```

```

    pthread_join(producer_thread, NULL);
    pthread_join(consumer_thread, NULL);

    pthread_mutex_destroy(&buffer_mutex);
    pthread_cond_destroy(&buffer_not_empty);
    pthread_cond_destroy(&buffer_not_full);

    return 0;
}

```

4. Advanced Threading Concepts

4.1 Thread-Specific Data

Thread-specific data allows each thread to have its own copy of a variable:

```

#include <pthread.h>
#include <stdio.h>
#include <stdlib.h>

pthread_key_t thread_key;

void destructor(void* value) {
    printf("Destroying thread-specific data: %d\n", *(int*)value);
    free(value);
}

void* thread_function(void* arg) {
    int* thread_data = malloc(sizeof(int));
    *thread_data = *(int*)arg;

    pthread_setspecific(thread_key, thread_data);

    int* retrieved_data = (int*)pthread_getspecific(thread_key);
    printf("Thread %d has data: %d\n", *(int*)arg, *retrieved_data);

    return NULL;
}

int main() {
    pthread_t thread1, thread2;
    int data1 = 100, data2 = 200;

    pthread_key_create(&thread_key, destructor);
}

```

```

pthread_create(&thread1, NULL, thread_function, &data1);
pthread_create(&thread2, NULL, thread_function, &data2);

pthread_join(thread1, NULL);
pthread_join(thread2, NULL);

pthread_key_delete(thread_key);
return 0;
}

```

4.2 Thread Cancellation

Threads can be cancelled by other threads:

```

#include <pthread.h>
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>

void cleanup_handler(void* arg) {
    printf("Cleanup handler called with argument: %d\n", *(int*)arg);
    free(arg);
}

void* cancellable_thread(void* arg) {
    int* data = malloc(sizeof(int));
    *data = 42;

    // Install cleanup handler
    pthread_cleanup_push(cleanup_handler, data);

    // Enable cancellation
    pthread_setcancelstate(PTHREAD_CANCEL_ENABLE, NULL);
    pthread_setcanceltype(PTHREAD_CANCEL_DEFERRED, NULL);

    while (1) {
        printf("Working...\n");
        sleep(1);

        // Cancellation point
        pthread_testcancel();
    }

    // This won't be reached due to infinite loop
    pthread_cleanup_pop(0);
    return NULL;
}

```

```

}

int main() {
    pthread_t thread;

    pthread_create(&thread, NULL, cancellable_thread, NULL);

    sleep(3); // Let thread work for a while

    printf("Cancelling thread...\n");
    pthread_cancel(thread);

    pthread_join(thread, NULL);
    printf("Thread cancelled and joined\n");

    return 0;
}

```

5. Performance Considerations

5.1 Thread Pool Implementation

A thread pool can improve performance by reusing threads:

```

#include <pthread.h>
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>

#define NUM_THREADS 4
#define QUEUE_SIZE 10

typedef struct {
    void (*function)(void* arg);
    void* arg;
} task_t;

typedef struct {
    task_t task_queue[QUEUE_SIZE];
    int front, rear, count;
    pthread_mutex_t queue_mutex;
    pthread_cond_t queue_not_empty;
    pthread_cond_t queue_not_full;
    int shutdown;
    pthread_t threads[NUM_THREADS];
}

```

```

} threadpool_t;

threadpool_t pool = {0};

void* worker_thread(void* arg) {
    while (1) {
        pthread_mutex_lock(&pool.queue_mutex);

        // Wait for tasks or shutdown
        while (pool.count == 0 && !pool.shutdown) {
            pthread_cond_wait(&pool.queue_not_empty, &pool.queue_mutex);
        }

        if (pool.shutdown) {
            pthread_mutex_unlock(&pool.queue_mutex);
            break;
        }

        // Get task from queue
        task_t task = pool.task_queue[pool.front];
        pool.front = (pool.front + 1) % QUEUE_SIZE;
        pool.count--;

        pthread_mutex_unlock(&pool.queue_mutex);
        pthread_cond_signal(&pool.queue_not_full);

        // Execute task
        task.function(task.arg);
    }
    return NULL;
}

int threadpool_init() {
    pthread_mutex_init(&pool.queue_mutex, NULL);
    pthread_cond_init(&pool.queue_not_empty, NULL);
    pthread_cond_init(&pool.queue_not_full, NULL);

    for (int i = 0; i < NUM_THREADS; i++) {
        pthread_create(&pool.threads[i], NULL, worker_thread, NULL);
    }

    return 0;
}

int threadpool_add_task(void (*function)(void*), void* arg) {
    pthread_mutex_lock(&pool.queue_mutex);

```

```

// Wait if queue is full
while (pool.count == QUEUE_SIZE) {
    pthread_cond_wait(&pool.queue_not_full, &pool.queue_mutex);
}

// Add task to queue
pool.task_queue[pool.rear].function = function;
pool.task_queue[pool.rear].arg = arg;
pool.rear = (pool.rear + 1) % QUEUE_SIZE;
pool.count++;

pthread_mutex_unlock(&pool.queue_mutex);
pthread_cond_signal(&pool.queue_not_empty);

return 0;
}

int threadpool_shutdown() {
    pthread_mutex_lock(&pool.queue_mutex);
    pool.shutdown = 1;
    pthread_mutex_unlock(&pool.queue_mutex);

    pthread_cond_broadcast(&pool.queue_not_empty);

    for (int i = 0; i < NUM_THREADS; i++) {
        pthread_join(pool.threads[i], NULL);
    }

    pthread_mutex_destroy(&pool.queue_mutex);
    pthread_cond_destroy(&pool.queue_not_empty);
    pthread_cond_destroy(&pool.queue_not_full);

    return 0;
}

// Example task function
void example_task(void* arg) {
    int task_id = *(int*)arg;
    printf("Executing task %d\n", task_id);
    sleep(1);
    printf("Task %d completed\n", task_id);
}

int main() {
    threadpool_init();

    // Add tasks to thread pool

```



```

    for (int i = 0; i < 10; i++) {
        int* task_id = malloc(sizeof(int));
        *task_id = i;
        threadpool_add_task(example_task, task_id);
    }

    sleep(5); // Let tasks execute

    threadpool_shutdown();
    printf("Thread pool shutdown complete\n");

    return 0;
}

```

6. Debugging Multithreaded Applications

6.1 Common Issues

1. **Deadlocks:** Threads waiting for each other indefinitely
2. **Race conditions:** Unpredictable behavior due to timing
3. **Priority inversion:** High-priority thread blocked by low-priority thread
4. **Thread starvation:** Thread never gets CPU time

6.2 Debugging Tools

- **Valgrind with Helgrind:** Detects race conditions and threading errors
- **GDB:** Can debug individual threads
- **ThreadSanitizer:** Compiler-based tool for detecting threading bugs

Example Valgrind command:

```
valgrind --tool=helgrind ./multithreaded_program
```

7. Best Practices

7.1 Design Principles

1. **Minimize shared data:** Reduce the need for synchronization
2. **Use immutable data:** Avoid race conditions by not modifying shared data
3. **Keep critical sections small:** Reduce lock contention
4. **Avoid nested locks:** Prevent deadlocks
5. **Use RAII for resource management:** Automatic cleanup

7.2 Performance Optimization

1. **Profile before optimizing:** Identify actual bottlenecks
2. **Consider Amdahl's Law:** Understand the limits of parallelization
3. **Use appropriate synchronization primitives:** Mutex for mutual exclusion, condition variables for signaling
4. **Avoid false sharing:** Ensure cache lines aren't shared unnecessarily
5. **Consider lock-free data structures:** For high-performance scenarios

Summary

Multithreading in C provides powerful capabilities for concurrent programming but requires careful design and implementation to avoid common pitfalls. Key concepts include:

- Understanding the difference between processes and threads
- Properly creating and managing threads with `pthread_create` and `pthread_join`
- Implementing synchronization with mutexes and condition variables
- Handling thread-specific data and cancellation
- Optimizing performance with thread pools and careful resource management

Mastering multithreading requires practice and understanding of both the technical details and the underlying concepts. Always prioritize correctness over performance and use appropriate debugging tools to identify issues in multithreaded code.

Module 15.2: System Programming in C

Introduction

System programming involves writing software that interacts directly with the operating system and hardware. In C, system programming encompasses a wide range of topics including process management, file system operations, inter-process communication, and low-level system interfaces. This module explores the essential concepts and techniques for writing system-level code in C.

Learning Objectives

By the end of this module, you will be able to: - Understand system calls and their role in system programming - Manage processes and implement process control - Perform advanced file system operations - Implement inter-process communication mechanisms - Work with signals and signal handling - Understand memory management at the system level - Apply best practices for secure and robust system programming

1. System Calls and APIs

1.1 What are System Calls?

System calls are the interface between user-space programs and the kernel. They provide controlled access to system resources and services. In C, system calls are typically accessed through wrapper functions in the standard library.

1.2 Common System Call Categories

1. **Process Control:** `fork()`, `exec()`, `wait()`, `exit()`
2. **File Management:** `open()`, `read()`, `write()`, `close()`, `stat()`
3. **Device Management:** `ioctl()`
4. **Information Maintenance:** `getpid()`, `getuid()`, `gettimeofday()`
5. **Communication:** `pipe()`, `socket()`, `send()`, `receive()`

1.3 Error Handling in System Calls

System calls return -1 on error and set the global `errno` variable:

```

#include <stdio.h>
#include <unistd.h>
#include <errno.h>
#include <string.h>

int main() {
    if (close(-1) == -1) {
        fprintf(stderr, "Error: %s\n", strerror(errno));
        // Error: Bad file descriptor
    }
    return 0;
}

```

2. Process Management

2.1 Process Creation with fork()

The `fork()` system call creates a new process by duplicating the calling process:

```

#include <stdio.h>
#include <unistd.h>
#include <sys/wait.h>
#include <stdlib.h>

int main() {
    pid_t pid = fork();

    if (pid == -1) {
        perror("fork failed");
        exit(EXIT_FAILURE);
    } else if (pid == 0) {
        // Child process
        printf("Child process (PID: %d)\n", getpid());
        exit(EXIT_SUCCESS);
    } else {
        // Parent process
        printf("Parent process (PID: %d), Child PID: %d\n", getpid(), pid);
        wait(NULL); // Wait for child to complete
        printf("Child process completed\n");
    }

    return 0;
}

```

2.2 Process Execution with exec()

The `exec()` family of functions replaces the current process image:

```
#include <stdio.h>
#include <unistd.h>
#include <stdlib.h>

int main() {
    pid_t pid = fork();

    if (pid == 0) {
        // Child process - execute ls command
        char *args[] = {"ls", "-l", "/home", NULL};
        execvp("ls", args);

        // If execvp returns, it failed
        perror("execvp failed");
        exit(EXIT_FAILURE);
    } else if (pid > 0) {
        // Parent process
        wait(NULL);
        printf("Child process completed\n");
    } else {
        perror("fork failed");
        exit(EXIT_FAILURE);
    }

    return 0;
}
```

2.3 Process Groups and Sessions

Process groups and sessions help organize related processes:

```
#include <stdio.h>
#include <unistd.h>
#include <sys/wait.h>
#include <signal.h>

int main() {
    pid_t pid = fork();

    if (pid == 0) {
        // Create new session
        if (setsid() == -1) {
```

```

        perror("setsid failed");
        exit(EXIT_FAILURE);
    }

    printf("New session leader (PID: %d)\n", getpid());
    printf("Session ID: %d\n", getsid(0));

    // Simulate daemon work
    sleep(10);
    exit(EXIT_SUCCESS);
} else if (pid > 0) {
    printf("Parent process (PID: %d)\n", getpid());
    wait(NULL);
}

return 0;
}

```

3. Advanced File System Operations

3.1 File Metadata with stat()

The `stat()` family of functions retrieves file metadata:

```

#include <stdio.h>
#include <sys/stat.h>
#include <time.h>
#include <pwd.h>
#include <grp.h>

void print_file_info(const char *filename) {
    struct stat file_stat;

    if (stat(filename, &file_stat) == -1) {
        perror("stat failed");
        return;
    }

    printf("File: %s\n", filename);
    printf("Size: %ld bytes\n", file_stat.st_size);
    printf("Permissions: %o\n", file_stat.st_mode & 0777);
    printf("Inode: %ld\n", file_stat.st_ino);
    printf("Links: %d\n", file_stat.st_nlink);

    // Convert timestamps to readable format

```

```

printf("Last access: %s", ctime(&file_stat.st_atime));
printf("Last modification: %s", ctime(&file_stat.st_mtime));
printf("Last status change: %s", ctime(&file_stat.st_ctime));

// Get user and group names
struct passwd *pwd = getpwuid(file_stat.st_uid);
struct group *grp = getgrgid(file_stat.st_gid);

if (pwd) printf("Owner: %s\n", pwd->pw_name);
if (grp) printf("Group: %s\n", grp->gr_name);
}

int main(int argc, char *argv[]) {
    if (argc != 2) {
        fprintf(stderr, "Usage: %s <filename>\n", argv[0]);
        return 1;
    }

    print_file_info(argv[1]);
    return 0;
}

```

3.2 Directory Operations

Working with directories using `opendir()`, `readdir()`, and `closedir()`:

```

#include <stdio.h>
#include <dirent.h>
#include <sys/stat.h>
#include <string.h>

void list_directory(const char *path) {
    DIR *dir = opendir(path);
    if (!dir) {
        perror("opendir failed");
        return;
    }

    struct dirent *entry;
    while ((entry = readdir(dir)) != NULL) {
        // Skip current and parent directory entries
        if (strcmp(entry->d_name, ".") == 0 || strcmp(entry->d_name, "..") == 0) {
            continue;
        }

        printf("%s\n", entry->d_name);
    }
}

```

```

// Get file type information
switch (entry->d_type) {
    case DT_REG:
        printf(" (regular file)\n");
        break;
    case DT_DIR:
        printf(" (directory)\n");
        break;
    case DT_LNK:
        printf(" (symbolic link)\n");
        break;
    case DT_CHR:
        printf(" (character device)\n");
        break;
    case DT_BLK:
        printf(" (block device)\n");
        break;
    case DT_FIFO:
        printf(" (named pipe)\n");
        break;
    case DT SOCK:
        printf(" (socket)\n");
        break;
    default:
        printf(" (unknown)\n");
        break;
}
}

closedir(dir);
}

int main(int argc, char *argv[]) {
    const char *path = (argc > 1) ? argv[1] : ".";
    list_directory(path);
    return 0;
}

```

3.3 Hard Links and Symbolic Links

Creating and working with links:

```

#include <stdio.h>
#include <unistd.h>
#include <sys/stat.h>

```



```

int main() {
    const char *original = "original.txt";
    const char *hard_link = "hard_link.txt";
    const char *sym_link = "sym_link.txt";

    // Create original file
    FILE *file = fopen(original, "w");
    if (file) {
        fprintf(file, "This is the original file.\n");
        fclose(file);
    }

    // Create hard link
    if (link(original, hard_link) == -1) {
        perror("link failed");
    } else {
        printf("Hard link created: %s\n", hard_link);
    }

    // Create symbolic link
    if (symlink(original, sym_link) == -1) {
        perror("symlink failed");
    } else {
        printf("Symbolic link created: %s\n", sym_link);
    }

    // Show link information
    struct stat stat_buf;
    if (stat(original, &stat_buf) == 0) {
        printf("Original file link count: %d\n", stat_buf.st_nlink);
    }

    return 0;
}

```

4. Inter-Process Communication (IPC)

4.1 Pipes

Creating anonymous pipes for communication between related processes:

```

#include <stdio.h>
#include <unistd.h>
#include <sys/wait.h>

```

```

#include <string.h>
#include <stdlib.h>

int main() {
    int pipefd[2];
    pid_t pid;
    char buffer[100];

    // Create pipe
    if (pipe(pipefd) == -1) {
        perror("pipe failed");
        exit(EXIT_FAILURE);
    }

    pid = fork();
    if (pid == -1) {
        perror("fork failed");
        exit(EXIT_FAILURE);
    }

    if (pid == 0) {
        // Child process - write to pipe
        close(pipefd[0]); // Close read end

        const char *message = "Hello from child process!";
        write(pipefd[1], message, strlen(message) + 1);

        close(pipefd[1]); // Close write end
        exit(EXIT_SUCCESS);
    } else {
        // Parent process - read from pipe
        close(pipefd[1]); // Close write end

        ssize_t bytes_read = read(pipefd[0], buffer, sizeof(buffer));
        if (bytes_read > 0) {
            printf("Parent received: %s\n", buffer);
        }

        close(pipefd[0]); // Close read end
        wait(NULL); // Wait for child to complete
    }

    return 0;
}

```

4.2 Named Pipes (FIFOs)

Creating named pipes for communication between unrelated processes:

```
#include <stdio.h>
#include <sys/stat.h>
#include <fcntl.h>
#include <unistd.h>
#include <string.h>

#define FIFO_NAME "/tmp/my_fifo"

int main() {
    // Create named pipe
    if (mkfifo(FIFO_NAME, 0666) == -1) {
        perror("mkfifo failed");
        return 1;
    }

    pid_t pid = fork();
    if (pid == 0) {
        // Child process - write to FIFO
        int fd = open(FIFO_NAME, O_WRONLY);
        if (fd == -1) {
            perror("open failed");
            exit(EXIT_FAILURE);
        }

        const char *message = "Message through named pipe";
        write(fd, message, strlen(message) + 1);
        close(fd);
        exit(EXIT_SUCCESS);
    } else {
        // Parent process - read from FIFO
        int fd = open(FIFO_NAME, O_RDONLY);
        if (fd == -1) {
            perror("open failed");
            return 1;
        }

        char buffer[100];
        ssize_t bytes_read = read(fd, buffer, sizeof(buffer));
        if (bytes_read > 0) {
            printf("Received: %s\n", buffer);
        }

        close(fd);
    }
}
```

```

        wait(NULL);
        unlink(FIFO_NAME); // Remove FIFO
    }

    return 0;
}

```

4.3 Shared Memory

Using shared memory for efficient data sharing:

```

#include <stdio.h>
#include <sys/shm.h>
#include <sys/stat.h>
#include <unistd.h>
#include <sys/wait.h>
#include <string.h>

int main() {
    const int SHM_SIZE = 1024;
    const key_t shm_key = 0x1234;

    // Create shared memory segment
    int shmid = shmget(shm_key, SHM_SIZE, IPC_CREAT | 0666);
    if (shmid == -1) {
        perror("shmget failed");
        return 1;
    }

    pid_t pid = fork();
    if (pid == 0) {
        // Child process - write to shared memory
        char *shm_ptr = (char*)shmat(shmid, NULL, 0);
        if (shm_ptr == (char*)-1) {
            perror("shmat failed");
            exit(EXIT_FAILURE);
        }

        strcpy(shm_ptr, "Hello from shared memory!");
        shmdt(shm_ptr);
        exit(EXIT_SUCCESS);
    } else {
        // Parent process - read from shared memory
        sleep(1); // Wait for child to write

        char *shm_ptr = (char*)shmat(shmid, NULL, 0);
    }
}

```

```

    if (shm_ptr == (char*)-1) {
        perror("shmat failed");
        return 1;
    }

    printf("Received: %s\n", shm_ptr);

    shmdt(shm_ptr);
    shmctl(shmid, IPC_RMID, NULL); // Remove shared memory
    wait(NULL);
}

return 0;
}

```

5. Signal Handling

5.1 Basic Signal Handling

Handling signals with `signal()` or `sigaction()`:

```

#include <stdio.h>
#include <signal.h>
#include <unistd.h>
#include <stdlib.h>

void signal_handler(int sig) {
    switch (sig) {
        case SIGINT:
            printf("\nReceived SIGINT (Ctrl+C)\n");
            break;
        case SIGTERM:
            printf("\nReceived SIGTERM\n");
            break;
        default:
            printf("\nReceived signal %d\n", sig);
            break;
    }
}

int main() {
    // Register signal handlers
    signal(SIGINT, signal_handler);
    signal(SIGTERM, signal_handler);
}

```

```

printf("Process PID: %d\n", getpid());
printf("Send signals to test handling\n");
printf("Press Ctrl+C to send SIGINT\n");

// Infinite loop - wait for signals
while (1) {
    printf("Working...\n");
    sleep(2);
}

return 0;
}

```

5.2 Advanced Signal Handling with sigaction()

More robust signal handling using sigaction():

```

#include <stdio.h>
#include <signal.h>
#include <unistd.h>
#include <stdlib.h>
#include <string.h>

volatile sig_atomic_t signal_received = 0;

void sigint_handler(int sig, siginfo_t *info, void *context) {
    printf("\nReceived SIGINT from PID %d\n", info->si_pid);
    signal_received = 1;
}

void sigusr1_handler(int sig, siginfo_t *info, void *context) {
    printf("\nReceived SIGUSR1\n");
    signal_received = 2;
}

int main() {
    struct sigaction sa_int, sa_usr1;

    // Setup SIGINT handler
    memset(&sa_int, 0, sizeof(sa_int));
    sa_int.sa_sigaction = sigint_handler;
    sa_int.sa_flags = SA_SIGINFO;
    sigaction(SIGINT, &sa_int, NULL);

    // Setup SIGUSR1 handler
    memset(&sa_usr1, 0, sizeof(sa_usr1));

```

```

sa_usr1.sa_sigaction = sigusr1_handler;
sa_usr1.sa_flags = SA_SIGINFO;
sigaction(SIGUSR1, &sa_usr1, NULL);

printf("Process PID: %d\n", getpid());
printf("Send SIGINT (Ctrl+C) or SIGUSR1 to test\n");

while (!signal_received) {
    printf("Working... (signal_received = %d)\n", signal_received);
    sleep(2);
}

printf("Exiting due to signal %d\n", signal_received);
return 0;
}

```

6. Memory Management

6.1 Memory Mapping with mmap()

Using mmap() for file mapping and shared memory:

```

#include <stdio.h>
#include <sys/mman.h>
#include <sys/stat.h>
#include <fcntl.h>
#include <unistd.h>
#include <string.h>

int main() {
    const char *filename = "mapped_file.txt";
    const char *content = "This content is mapped to memory.\n";

    // Create and write to file
    int fd = open(filename, O_CREAT | O_RDWR, 0666);
    if (fd == -1) {
        perror("open failed");
        return 1;
    }

    write(fd, content, strlen(content));

    // Map file to memory
    struct stat sb;
    fstat(fd, &sb);

```

```

char *mapped = mmap(NULL, sb.st_size, PROT_READ | PROT_WRITE, MAP_SHARED, fd, 0);
if (mapped == MAP_FAILED) {
    perror("mmap failed");
    close(fd);
    return 1;
}

printf("Original content: %s", mapped);

// Modify mapped content
strcpy(mapped, "Modified content through memory mapping!\n");

// Unmap and cleanup
munmap(mapped, sb.st_size);
close(fd);

// Verify changes
fd = open(filename, O_RDONLY);
char buffer[100];
ssize_t bytes_read = read(fd, buffer, sizeof(buffer) - 1);
if (bytes_read > 0) {
    buffer[bytes_read] = '\0';
    printf("File content after modification: %s", buffer);
}

close(fd);
unlink(filename);
return 0;
}

```

6.2 Memory Protection

Setting memory protection with `mprotect()`:

```

#include <stdio.h>
#include <sys/mman.h>
#include <stdlib.h>
#include <signal.h>
#include <setjmp.h>

static jmp_buf jmp_env;

void segv_handler(int sig) {
    printf("Segmentation fault caught!\n");
    longjmp(jmp_env, 1);
}

```



```

}

int main() {
    // Allocate memory with mmap
    size_t page_size = getpagesize();
    char *memory = mmap(NULL, page_size, PROT_READ | PROT_WRITE,
                        MAP_PRIVATE | MAP_ANONYMOUS, -1, 0);

    if (memory == MAP_FAILED) {
        perror("mmap failed");
        return 1;
    }

    // Set up signal handler
    signal(SIGSEGV, segv_handler);

    // Write to memory (should work)
    strcpy(memory, "Hello, protected memory!");
    printf("Before protection: %s\n", memory);

    // Set memory to read-only
    if (mprotect(memory, page_size, PROT_READ) == -1) {
        perror("mprotect failed");
        munmap(memory, page_size);
        return 1;
    }

    // Try to write to protected memory
    if (setjmp(jmp_env) == 0) {
        strcpy(memory, "This should fail!");
        printf("Write succeeded (unexpected)\n");
    } else {
        printf("Write to protected memory failed as expected\n");
    }

    printf("After protection: %s\n", memory);

    // Cleanup
    munmap(memory, page_size);
    return 0;
}

```

7. Security Considerations

7.1 Secure Programming Practices

1. **Input Validation:** Always validate and sanitize input
2. **Buffer Overflow Prevention:** Use safe string functions
3. **Privilege Management:** Drop privileges when not needed
4. **Secure File Operations:** Set appropriate permissions
5. **Environment Sanitization:** Clear sensitive environment variables

7.2 Example of Secure File Handling

```
#include <stdio.h>
#include <fcntl.h>
#include <unistd.h>
#include <sys/stat.h>
#include <errno.h>
#include <string.h>

int create_secure_file(const char *filename) {
    // Create file with restrictive permissions
    int fd = open(filename, O_CREAT | O_WRONLY | O_TRUNC, 0600);
    if (fd == -1) {
        return -1;
    }

    // Verify file was created with correct permissions
    struct stat st;
    if (fstat(fd, &st) == -1) {
        close(fd);
        unlink(filename);
        return -1;
    }

    if ((st.st_mode & 0777) != 0600) {
        close(fd);
        unlink(filename);
        errno = EACCES;
        return -1;
    }

    return fd;
}

int main() {
    const char *filename = "secure_file.txt";
```

```

int fd = create_secure_file(filename);

if (fd == -1) {
    perror("Failed to create secure file");
    return 1;
}

const char *data = "Sensitive data";
write(fd, data, strlen(data));
close(fd);

printf("Secure file created successfully\n");

// Verify permissions
struct stat st;
if (stat(filename, &st) == 0) {
    printf("File permissions: %o\n", st.st_mode & 0777);
}

unlink(filename);
return 0;
}

```

8. Best Practices

8.1 Error Handling

1. Check all system call return values
2. Use errno for detailed error information
3. Implement proper cleanup on error
4. Log errors appropriately

8.2 Resource Management

1. Always close file descriptors
2. Free allocated memory
3. Remove temporary files
4. Detach shared memory segments

8.3 Portability

1. Use feature test macros
2. Handle platform differences
3. Avoid non-standard extensions

4. Test on multiple platforms

Summary

System programming in C requires a deep understanding of operating system concepts and careful attention to detail. Key areas covered in this module include:

- Process management with `fork()`, `exec()`, and `wait()`
- Advanced file system operations with `stat()` and directory functions
- Inter-process communication through pipes, FIFOs, and shared memory
- Signal handling for process control
- Memory management with `mmap()` and `mprotect()`
- Security considerations for robust system code

Mastering system programming enables developers to create efficient, low-level applications that interact directly with the operating system. Always prioritize security, robustness, and proper error handling when writing system-level code.

Module 15.3: Domain-Specific Applications in C

Introduction

Domain-specific applications are software solutions designed to address particular problem domains with specialized requirements. In C programming, developing domain-specific applications requires leveraging the language's low-level capabilities while implementing domain-specific logic and algorithms. This module explores the development of applications in various domains including embedded systems, scientific computing, networking, and real-time systems.

Learning Objectives

By the end of this module, you will be able to: - Understand the characteristics of domain-specific applications - Design and implement applications for specific domains - Apply domain-specific algorithms and data structures - Optimize performance for domain requirements - Handle domain-specific constraints and requirements - Integrate with domain-specific hardware or protocols

1. Characteristics of Domain-Specific Applications

1.1 Domain Requirements

Domain-specific applications have unique requirements that distinguish them from general-purpose software:

1. **Performance Constraints:** Real-time systems require deterministic response times
2. **Resource Limitations:** Embedded systems often have limited memory and processing power
3. **Reliability Requirements:** Safety-critical systems demand high reliability
4. **Domain-Specific Logic:** Specialized algorithms and data structures
5. **Integration Needs:** Interface with specific hardware or protocols

1.2 Design Considerations

When designing domain-specific applications in C:

1. **Modularity:** Separate domain logic from infrastructure code
2. **Abstraction:** Hide domain complexity behind clean interfaces
3. **Optimization:** Focus on domain-critical performance paths
4. **Testing:** Implement domain-specific test scenarios
5. **Documentation:** Provide domain-specific context and examples

2. Embedded Systems Applications

2.1 Microcontroller Programming

Programming for microcontrollers requires careful resource management:

```
#include <stdint.h>
#include <stdbool.h>

// Hardware register definitions (example for ARM Cortex-M)
#define GPIO_BASE_ADDR    0x40020000
#define GPIO_MODER_REG    (*(volatile uint32_t*)(GPIO_BASE_ADDR + 0x00))
#define GPIO_ODR_REG      (*(volatile uint32_t*)(GPIO_BASE_ADDR + 0x14))

// GPIO pin configuration
typedef enum {
    GPIO_MODE_INPUT = 0,
    GPIO_MODE_OUTPUT = 1,
    GPIO_MODE_AF = 2,
    GPIO_MODE_ANALOG = 3
} gpio_mode_t;

typedef enum {
    GPIO_SPEED_LOW = 0,
    GPIO_SPEED_MEDIUM = 1,
    GPIO_SPEED_HIGH = 2,
    GPIO_SPEED_VERY_HIGH = 3
} gpio_speed_t;

typedef struct {
    uint8_t pin;
    gpio_mode_t mode;
    gpio_speed_t speed;
    bool pull_up;
    bool pull_down;
} gpio_config_t;

// GPIO driver functions
void gpio_init(const gpio_config_t *config) {
    // Configure pin mode
    uint32_t mode_mask = 0x3 << (config->pin * 2);
    GPIO_MODER_REG = (GPIO_MODER_REG & ~mode_mask) |
        ((config->mode & 0x3) << (config->pin * 2));

    // Configure output speed (simplified)
    // Additional configuration for pull-up/pull-down would go here
}
```

```

void gpio_write(uint8_t pin, bool value) {
    if (value) {
        GPIO_ODR_REG |= (1 << pin);
    } else {
        GPIO_ODR_REG &= ~(1 << pin);
    }
}

bool gpio_read(uint8_t pin) {
    return (GPIO_ODR_REG & (1 << pin)) != 0;
}

// Example application: LED blinker
void led_blinker_task(void) {
    static uint32_t counter = 0;
    const uint32_t BLINK_PERIOD = 1000;    // ms

    gpio_config_t led_config = {
        .pin = 13,    // Assuming LED on pin 13
        .mode = GPIO_MODE_OUTPUT,
        .speed = GPIO_SPEED_LOW,
        .pull_up = false,
        .pull_down = false
    };

    gpio_init(&led_config);

    // Main loop (typically called by RTOS or timer interrupt)
    while (1) {
        if (counter % BLINK_PERIOD == 0) {
            static bool led_state = false;
            led_state = !led_state;
            gpio_write(13, led_state);
        }
        counter++;

        // In a real system, this would be a delay function
        // or the function would be called periodically
    }
}

```

2.2 Real-Time Operating System Integration

Integrating with RTOS for deterministic behavior:

```

#include <stdint.h>
#include <stddef.h>

// Simplified RTOS API (similar to FreeRTOS)
typedef void* TaskHandle_t;
typedef void* QueueHandle_t;
typedef void* SemaphoreHandle_t;

// Task function prototype
typedef void (*TaskFunction_t)(void* pvParameters);

// RTOS function declarations
TaskHandle_t xTaskCreate(TaskFunction_t pvTaskCode,
                        const char * const pcName,
                        const uint16_t usStackDepth,
                        void * const pvParameters,
                        uint32_t uxPriority,
                        TaskHandle_t * const pxCreatedTask);

QueueHandle_t xQueueCreate(uint32_t uxQueueLength,
                          uint32_t uxItemSize);

uint32_t xQueueSend(QueueHandle_t xQueue,
                   const void * const pvItemToQueue,
                   uint32_t xTicksToWait);

uint32_t xQueueReceive(QueueHandle_t xQueue,
                     void * const pvBuffer,
                     uint32_t xTicksToWait);

SemaphoreHandle_t xSemaphoreCreateBinary(void);
void xSemaphoreTake(SemaphoreHandle_t xSemaphore, uint32_t xBlockTime);
void xSemaphoreGive(SemaphoreHandle_t xSemaphore);

// Domain-specific application: Sensor data processing
typedef struct {
    uint16_t temperature;
    uint16_t humidity;
    uint32_t timestamp;
} sensor_data_t;

// Task handles
TaskHandle_t sensor_task_handle;
TaskHandle_t processing_task_handle;
TaskHandle_t communication_task_handle;

// Communication queue

```



```

QueueHandle_t sensor_queue;

// Synchronization semaphore
SemaphoreHandle_t data_ready_semaphore;

// Sensor reading task
void sensor_task(void* pvParameters) {
    sensor_data_t sensor_data;

    while (1) {
        // Simulate sensor reading
        sensor_data.temperature = 250 + (rand() % 20); // 25.0°C ± 1.0°C
        sensor_data.humidity = 500 + (rand() % 100); // 50.0% ± 5.0%
        sensor_data.timestamp = 0; // In real system, use actual timestamp

        // Send data to processing task
        if (xQueueSend(sensor_queue, &sensor_data, 100) != 0) {
            // Handle queue full error
        }

        // Signal that data is ready
        xSemaphoreGive(data_ready_semaphore);

        // Delay for 1 second (1000ms)
        // In real RTOS, use vTaskDelay() or similar
    }
}

// Data processing task
void processing_task(void* pvParameters) {
    sensor_data_t sensor_data;

    while (1) {
        // Wait for data to be ready
        xSemaphoreTake(data_ready_semaphore, 0xFFFFFFFF); // Wait indefinitely

        // Receive sensor data
        if (xQueueReceive(sensor_queue, &sensor_data, 0) == 0) {
            // Process sensor data
            float temp_celsius = sensor_data.temperature / 10.0f;
            float humidity_percent = sensor_data.humidity / 10.0f;

            // Apply calibration or filtering if needed
            // ... processing logic ...

            // Forward processed data to communication task
            // ... communication logic ...
        }
    }
}

```

```

    }
}

// Application initialization
void application_init(void) {
    // Create communication queue
    sensor_queue = xQueueCreate(10, sizeof(sensor_data_t));

    // Create synchronization semaphore
    data_ready_semaphore = xSemaphoreCreateBinary();

    // Create tasks
    sensor_task_handle = xTaskCreate(sensor_task,
                                    "SensorTask",
                                    256,    // Stack size
                                    NULL,   // Parameters
                                    1,      // Priority
                                    NULL);  // Task handle

    processing_task_handle = xTaskCreate(processing_task,
                                        "ProcessingTask",
                                        512,    // Stack size
                                        NULL,   // Parameters
                                        2,      // Priority
                                        NULL);  // Task handle

    // In a real system, you would also create communication task
    // and start the scheduler
}

```

3. Scientific Computing Applications

3.1 Numerical Libraries Integration

Integrating with numerical libraries for scientific computing:

```

#include <stdio.h>
#include <stdlib.h>
#include <math.h>
#include <string.h>

// Matrix operations for scientific computing
typedef struct {
    size_t rows;

```

```

    size_t cols;
    double *data;
} matrix_t;

// Create matrix
matrix_t* matrix_create(size_t rows, size_t cols) {
    matrix_t *mat = malloc(sizeof(matrix_t));
    if (!mat) return NULL;

    mat->rows = rows;
    mat->cols = cols;
    mat->data = calloc(rows * cols, sizeof(double));

    if (!mat->data) {
        free(mat);
        return NULL;
    }

    return mat;
}

// Free matrix
void matrix_destroy(matrix_t *mat) {
    if (mat) {
        free(mat->data);
        free(mat);
    }
}

// Get element at (row, col)
double matrix_get(const matrix_t *mat, size_t row, size_t col) {
    if (row >= mat->rows || col >= mat->cols) return 0.0;
    return mat->data[row * mat->cols + col];
}

// Set element at (row, col)
void matrix_set(matrix_t *mat, size_t row, size_t col, double value) {
    if (row < mat->rows && col < mat->cols) {
        mat->data[row * mat->cols + col] = value;
    }
}

// Matrix multiplication
matrix_t* matrix_multiply(const matrix_t *a, const matrix_t *b) {
    if (a->cols != b->rows) return NULL;

    matrix_t *result = matrix_create(a->rows, b->cols);

```

```

    if (!result) return NULL;

    for (size_t i = 0; i < a->rows; i++) {
        for (size_t j = 0; j < b->cols; j++) {
            double sum = 0.0;
            for (size_t k = 0; k < a->cols; k++) {
                sum += matrix_get(a, i, k) * matrix_get(b, k, j);
            }
            matrix_set(result, i, j, sum);
        }
    }

    return result;
}

// Solve linear system using Gaussian elimination
int matrix_solve_linear_system(matrix_t *a, matrix_t *b) {
    if (a->rows != a->cols || a->rows != b->rows || b->cols != 1) {
        return -1; // Invalid dimensions
    }

    size_t n = a->rows;

    // Forward elimination
    for (size_t i = 0; i < n; i++) {
        // Find pivot
        size_t max_row = i;
        for (size_t k = i + 1; k < n; k++) {
            if (fabs(matrix_get(a, k, i)) > fabs(matrix_get(a, max_row, i))) {
                max_row = k;
            }
        }

        // Swap rows if needed
        if (max_row != i) {
            for (size_t j = 0; j < n; j++) {
                double temp = matrix_get(a, i, j);
                matrix_set(a, i, j, matrix_get(a, max_row, j));
                matrix_set(a, max_row, j, temp);
            }
            double temp = matrix_get(b, i, 0);
            matrix_set(b, i, 0, matrix_get(b, max_row, 0));
            matrix_set(b, max_row, 0, temp);
        }

        // Check for singular matrix

```

```

    if (fabs(matrix_get(a, i, i)) < 1e-10) {
        return -2; // Singular matrix
    }

    // Eliminate column
    for (size_t k = i + 1; k < n; k++) {
        double factor = matrix_get(a, k, i) / matrix_get(a, i, i);
        for (size_t j = i; j < n; j++) {
            double new_value = matrix_get(a, k, j) - factor * matrix_get(a, i, j);
            matrix_set(a, k, j, new_value);
        }
        double new_b = matrix_get(b, k, 0) - factor * matrix_get(b, i, 0);
        matrix_set(b, k, 0, new_b);
    }
}

// Back substitution
for (int i = n - 1; i >= 0; i--) {
    double sum = matrix_get(b, i, 0);
    for (size_t j = i + 1; j < n; j++) {
        sum -= matrix_get(a, i, j) * matrix_get(b, j, 0);
    }
    matrix_set(b, i, 0, sum / matrix_get(a, i, i));
}

return 0; // Success
}

// Example: Solve a system of linear equations
void example_linear_system(void) {
    // Solve:
    // 2x + 3y + z = 1
    // x + 4y + 2z = 2
    // 3x + y + 5z = 3

    matrix_t *a = matrix_create(3, 3);
    matrix_t *b = matrix_create(3, 1);

    if (!a || !b) {
        printf("Memory allocation failed\n");
        return;
    }

    // Coefficient matrix A
    matrix_set(a, 0, 0, 2.0); matrix_set(a, 0, 1, 3.0); matrix_set(a, 0, 2, 1.0);
    matrix_set(a, 1, 0, 1.0); matrix_set(a, 1, 1, 4.0); matrix_set(a, 1, 2, 2.0);
    matrix_set(a, 2, 0, 3.0); matrix_set(a, 2, 1, 1.0); matrix_set(a, 2, 2, 5.0);

```

```

    // Constants vector B
    matrix_set(b, 0, 0, 1.0);
    matrix_set(b, 1, 0, 2.0);
    matrix_set(b, 2, 0, 3.0);

    printf("Solving linear system:\n");
    printf("2x + 3y + z = 1\n");
    printf("x + 4y + 2z = 2\n");
    printf("3x + y + 5z = 3\n\n");

    int result = matrix_solve_linear_system(a, b);

    if (result == 0) {
        printf("Solution:\n");
        printf("x = %.6f\n", matrix_get(b, 0, 0));
        printf("y = %.6f\n", matrix_get(b, 1, 0));
        printf("z = %.6f\n", matrix_get(b, 2, 0));
    } else {
        printf("Failed to solve system (error code: %d)\n", result);
    }

    matrix_destroy(a);
    matrix_destroy(b);
}

```

3.2 Fast Fourier Transform (FFT)

Implementing FFT for signal processing:

```

#include <stdio.h>
#include <stdlib.h>
#include <math.h>
#include <complex.h>

#ifndef M_PI
#define M_PI 3.14159265358979323846
#endif

// Complex number type
typedef double complex complex_t;

// FFT implementation using Cooley-Tukey algorithm
void fft_recursive(complex_t *x, size_t n, size_t step) {
    if (step < n) {
        fft_recursive(x, n, step * 2);
    }
}

```

```

        fft_recursive(x + step, n, step * 2);

    for (size_t i = 0; i < n; i += step * 2) {
        complex_t t = cexp(-I * M_PI * i / n) * x[i + step];
        x[i + step] = x[i] - t;
        x[i] += t;
    }
}

}

void fft(complex_t *x, size_t n) {
    // Make sure n is a power of 2
    if (n & (n - 1)) {
        fprintf(stderr, "FFT size must be a power of 2\n");
        return;
    }

    // Bit-reversal permutation
    size_t j = 0;
    for (size_t i = 0; i < n; i++) {
        if (i < j) {
            complex_t temp = x[i];
            x[i] = x[j];
            x[j] = temp;
        }

        size_t k = n >> 1;
        while (k <= j) {
            j -= k;
            k >>= 1;
        }
        j += k;
    }

    // Recursive FFT
    fft_recursive(x, n, 1);
}

// Inverse FFT
void ifft(complex_t *x, size_t n) {
    // Conjugate the input
    for (size_t i = 0; i < n; i++) {
        x[i] = conj(x[i]);
    }

    // Apply FFT
    fft(x, n);
}

```

```

    // Conjugate and normalize
    for (size_t i = 0; i < n; i++) {
        x[i] = conj(x[i]) / n;
    }
}

// Generate a test signal
void generate_test_signal(complex_t *signal, size_t n, double frequency) {
    for (size_t i = 0; i < n; i++) {
        double t = (double)i / n;
        // Composite signal: sum of sine waves
        double value = sin(2 * M_PI * frequency * t) +
            0.5 * sin(2 * M_PI * 3 * frequency * t) +
            0.3 * sin(2 * M_PI * 5 * frequency * t);
        signal[i] = value + 0.0 * I; // Real signal
    }
}

// Example: FFT analysis of a signal
void example_fft_analysis(void) {
    const size_t N = 64; // Must be power of 2
    const double sample_rate = 1000.0; // Hz
    const double signal_frequency = 50.0; // Hz

    complex_t *signal = malloc(N * sizeof(complex_t));
    if (!signal) {
        fprintf(stderr, "Memory allocation failed\n");
        return;
    }

    // Generate test signal
    generate_test_signal(signal, N, signal_frequency);

    printf("Original signal samples:\n");
    for (size_t i = 0; i < 8; i++) { // Show first 8 samples
        printf("Sample[%zu] = %.3f\n", i, creal(signal[i]));
    }

    // Apply FFT
    fft(signal, N);

    printf("\nFrequency domain representation:\n");
    printf("Frequency (Hz)\tMagnitude\tPhase (rad)\n");

    // Analyze frequency components
    for (size_t i = 0; i <= N/2; i++) {

```



```

    double frequency = i * sample_rate / N;
    double magnitude = cabs(signal[i]) / N;
    double phase = carg(signal[i]);

    // Only show significant components
    if (magnitude > 0.01) {
        printf("%.1f\t\t%.3f\t\t%.3f\n", frequency, magnitude, phase);
    }
}

// Apply inverse FFT to verify
ifft(signal, N);

printf("\nReconstructed signal (first 8 samples):\n");
for (size_t i = 0; i < 8; i++) {
    printf("Sample[%zu] = %.3f\n", i, creal(signal[i]));
}

free(signal);
}

```

4. Networking Applications

4.1 Protocol Implementation

Implementing custom network protocols:

```

#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <stdint.h>
#include <arpa/inet.h>

// Custom protocol packet structure
typedef struct {
    uint16_t magic;           // Protocol identifier
    uint16_t version;        // Protocol version
    uint32_t sequence;       // Packet sequence number
    uint32_t timestamp;      // Timestamp (seconds since epoch)
    uint16_t type;           // Packet type
    uint16_t length;         // Payload length
    uint8_t payload[0];      // Variable length payload
} __attribute__((packed)) protocol_packet_t;

#define PROTOCOL_MAGIC 0xCAFE

```

```

#define PROTOCOL_VERSION 1

// Packet types
#define PACKET_TYPE_DATA 1
#define PACKET_TYPE_ACK 2
#define PACKET_TYPE_NACK 3
#define PACKET_TYPE_HEARTBEAT 4

// Protocol error codes
typedef enum {
    PROTOCOL_SUCCESS = 0,
    PROTOCOL_ERROR_INVALID_MAGIC = -1,
    PROTOCOL_ERROR_INVALID_VERSION = -2,
    PROTOCOL_ERROR_BUFFER_TOO_SMALL = -3,
    PROTOCOL_ERROR_INVALID_LENGTH = -4
} protocol_error_t;

// Create a protocol packet
protocol_packet_t* protocol_create_packet(uint16_t type,
                                          const void *payload,
                                          uint16_t payload_length) {
    size_t packet_size = sizeof(protocol_packet_t) + payload_length;
    protocol_packet_t *packet = malloc(packet_size);

    if (!packet) {
        return NULL;
    }

    // Initialize packet header
    packet->magic = htons(PROTOCOL_MAGIC);
    packet->version = htons(PROTOCOL_VERSION);
    packet->sequence = 0; // Would be set by sequence manager
    packet->timestamp = htonl((uint32_t)time(NULL));
    packet->type = htons(type);
    packet->length = htons(payload_length);

    // Copy payload
    if (payload && payload_length > 0) {
        memcpy(packet->payload, payload, payload_length);
    }

    return packet;
}

// Validate a received packet
protocol_error_t protocol_validate_packet(const protocol_packet_t *packet,
                                          size_t buffer_size) {

```

```

// Check buffer size
if (buffer_size < sizeof(protocol_packet_t)) {
    return PROTOCOL_ERROR_BUFFER_TOO_SMALL;
}

// Check magic number
if (ntohs(packet->magic) != PROTOCOL_MAGIC) {
    return PROTOCOL_ERROR_INVALID_MAGIC;
}

// Check version
if (ntohs(packet->version) != PROTOCOL_VERSION) {
    return PROTOCOL_ERROR_INVALID_VERSION;
}

// Check payload length
uint16_t payload_length = ntohs(packet->length);
size_t total_size = sizeof(protocol_packet_t) + payload_length;

if (buffer_size < total_size) {
    return PROTOCOL_ERROR_BUFFER_TOO_SMALL;
}

return PROTOCOL_SUCCESS;
}

// Parse packet payload based on type
void protocol_parse_payload(const protocol_packet_t *packet) {
    uint16_t type = ntohs(packet->type);
    uint16_t length = ntohs(packet->length);

    printf("Packet Type: %u, Length: %u\n", type, length);

    switch (type) {
        case PACKET_TYPE_DATA:
            printf("Data packet received: ");
            if (length > 0) {
                // Print first 32 bytes of payload as hex
                for (int i = 0; i < length && i < 32; i++) {
                    printf("%02x ", packet->payload[i]);
                }
                if (length > 32) printf("...");
            }
            printf("\n");
            break;

        case PACKET_TYPE_ACK:

```

```

        printf("Acknowledgment packet\n");
        break;

    case PACKET_TYPE_NACK:
        printf("Negative acknowledgment packet\n");
        break;

    case PACKET_TYPE_HEARTBEAT:
        printf("Heartbeat packet\n");
        break;

    default:
        printf("Unknown packet type\n");
        break;
    }
}

// Example: Protocol usage
void example_protocol_usage(void) {
    printf("Creating protocol packets...\n");

    // Create data packet
    const char *message = "Hello, Protocol World!";
    protocol_packet_t *data_packet = protocol_create_packet(
        PACKET_TYPE_DATA, message, strlen(message));

    if (data_packet) {
        printf("Created data packet:\n");
        protocol_parse_payload(data_packet);

        // Validate the packet
        protocol_error_t result = protocol_validate_packet(
            data_packet, sizeof(protocol_packet_t) + strlen(message));

        if (result == PROTOCOL_SUCCESS) {
            printf("Packet validation successful\n");
        } else {
            printf("Packet validation failed: %d\n", result);
        }

        free(data_packet);
    }

    // Create heartbeat packet
    protocol_packet_t *heartbeat_packet = protocol_create_packet(
        PACKET_TYPE_HEARTBEAT, NULL, 0);

```

```

    if (heartbeat_packet) {
        printf("\nCreated heartbeat packet:\n");
        protocol_parse_payload(heartbeat_packet);
        free(heartbeat_packet);
    }
}

```

4.2 HTTP Client Implementation

Simple HTTP client for web services:

```

#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <unistd.h>
#include <sys/socket.h>
#include <netinet/in.h>
#include <arpa/inet.h>
#include <netdb.h>

// HTTP request structure
typedef struct {
    char method[16];
    char host[256];
    int port;
    char path[512];
    char *headers;
    char *body;
    size_t body_length;
} http_request_t;

// HTTP response structure
typedef struct {
    int status_code;
    char status_text[64];
    char *headers;
    char *body;
    size_t body_length;
} http_response_t;

// Create HTTP request
http_request_t* http_create_request(const char *method,
                                    const char *host,
                                    int port,
                                    const char *path) {
    http_request_t *req = malloc(sizeof(http_request_t));

```

```

    if (!req) return NULL;

    strncpy(req->method, method, sizeof(req->method) - 1);
    strncpy(req->host, host, sizeof(req->host) - 1);
    req->port = port;
    strncpy(req->path, path, sizeof(req->path) - 1);

    req->headers = NULL;
    req->body = NULL;
    req->body_length = 0;

    return req;
}

// Add header to request
void http_add_header(http_request_t *req, const char *header) {
    size_t current_len = req->headers ? strlen(req->headers) : 0;
    size_t header_len = strlen(header);
    size_t new_len = current_len + header_len + 3; // +3 for \r\n\0

    char *new_headers = realloc(req->headers, new_len);
    if (!new_headers) return;

    req->headers = new_headers;

    if (current_len == 0) {
        strcpy(req->headers, header);
    } else {
        strcat(req->headers, "\r\n");
        strcat(req->headers, header);
    }
}

// Set request body
void http_set_body(http_request_t *req, const char *body, size_t length) {
    req->body = malloc(length);
    if (req->body) {
        memcpy(req->body, body, length);
        req->body_length = length;
    }
}

// Build HTTP request string
char* http_build_request_string(const http_request_t *req) {
    size_t req_size = strlen(req->method) + strlen(req->path) + 16; // Method + Path + HTTP
    req_size += strlen(req->host) + 16; // Host header

```

```

    if (req->headers) {
        req_size += strlen(req->headers) + 2; // Headers + \r\n
    }

    if (req->body) {
        req_size += req->body_length + 16; // Content-Length header
    }

    req_size += 32; // Extra space for safety

    char *request = malloc(req_size);
    if (!request) return NULL;

    // Build request line
    snprintf(request, req_size, "%s %s HTTP/1.1\r\n", req->method, req->path);

    // Add Host header
    char host_header[300];
    snprintf(host_header, sizeof(host_header), "Host: %s:%d", req->host, req->port);
    strcat(request, host_header);
    strcat(request, "\r\n");

    // Add custom headers
    if (req->headers) {
        strcat(request, req->headers);
        strcat(request, "\r\n");
    }

    // Add Content-Length if body exists
    if (req->body) {
        char content_length[32];
        snprintf(content_length, sizeof(content_length), "Content-Length: %zu\r\n", req->body_length);
        strcat(request, content_length);
    }

    // End headers
    strcat(request, "\r\n");

    // Add body
    if (req->body) {
        size_t current_len = strlen(request);
        memcpy(request + current_len, req->body, req->body_length);
        request[current_len + req->body_length] = '\0';
    }

    return request;
}

```

```

// Simple HTTP GET request
int http_get(const char *url, char **response_body, size_t *response_length) {
    // Parse URL (simplified - assumes http://host:port/path format)
    if (strncmp(url, "http://", 7) != 0) {
        fprintf(stderr, "Only HTTP URLs supported\n");
        return -1;
    }

    const char *host_start = url + 7;
    const char *port_start = strchr(host_start, ':');
    const char *path_start = strchr(host_start, '/');

    if (!path_start) {
        fprintf(stderr, "Invalid URL format\n");
        return -1;
    }

    char host[256];
    int port = 80;

    if (port_start && port_start < path_start) {
        // Extract host
        size_t host_len = port_start - host_start;
        if (host_len >= sizeof(host)) {
            fprintf(stderr, "Host name too long\n");
            return -1;
        }
        strncpy(host, host_start, host_len);
        host[host_len] = '\0';

        // Extract port
        port = atoi(port_start + 1);
    } else {
        // Extract host only
        size_t host_len = path_start - host_start;
        if (host_len >= sizeof(host)) {
            fprintf(stderr, "Host name too long\n");
            return -1;
        }
        strncpy(host, host_start, host_len);
        host[host_len] = '\0';
    }

    const char *path = path_start;

    // Create socket

```



```

int sock = socket(AF_INET, SOCK_STREAM, 0);
if (sock < 0) {
    perror("socket creation failed");
    return -1;
}

// Resolve hostname
struct hostent *server = gethostbyname(host);
if (!server) {
    fprintf(stderr, "Failed to resolve hostname: %s\n", host);
    close(sock);
    return -1;
}

// Connect to server
struct sockaddr_in server_addr;
memset(&server_addr, 0, sizeof(server_addr));
server_addr.sin_family = AF_INET;
memcpy(&server_addr.sin_addr.s_addr, server->h_addr, server->h_length);
server_addr.sin_port = htons(port);

if (connect(sock, (struct sockaddr*)&server_addr, sizeof(server_addr)) < 0) {
    perror("connection failed");
    close(sock);
    return -1;
}

// Create HTTP request
http_request_t *req = http_create_request("GET", host, port, path);
if (!req) {
    close(sock);
    return -1;
}

http_add_header(req, "Connection: close");
http_add_header(req, "User-Agent: C-HTTP-Client/1.0");

char *request_string = http_build_request_string(req);
if (!request_string) {
    free(req);
    close(sock);
    return -1;
}

// Send request
if (send(sock, request_string, strlen(request_string), 0) < 0) {
    perror("send failed");
}

```

```

        free(request_string);
        free(req);
        close(sock);
        return -1;
    }

    // Receive response
    char buffer[4096];
    ssize_t bytes_received;
    char *full_response = NULL;
    size_t total_received = 0;

    while ((bytes_received = recv(sock, buffer, sizeof(buffer) - 1, 0)) > 0) {
        char *new_response = realloc(full_response, total_received + bytes_received + 1);
        if (!new_response) {
            free(full_response);
            free(request_string);
            free(req);
            close(sock);
            return -1;
        }

        full_response = new_response;
        memcpy(full_response + total_received, buffer, bytes_received);
        total_received += bytes_received;
        full_response[total_received] = '\0';
    }

    close(sock);
    free(request_string);
    free(req);

    if (bytes_received < 0) {
        perror("recv failed");
        free(full_response);
        return -1;
    }

    // Parse response (simplified)
    *response_body = full_response;
    *response_length = total_received;

    return 0;
}

// Example: HTTP client usage
void example_http_client(void) {

```

```

printf("Making HTTP GET request...\n");

char *response_body;
size_t response_length;

// Make a simple HTTP request
int result = http_get("http://httpbin.org/get", &response_body, &response_length);

if (result == 0) {
    printf("HTTP request successful\n");
    printf("Response length: %zu bytes\n", response_length);

    // Print first 500 characters of response
    size_t print_length = (response_length < 500) ? response_length : 500;
    printf("Response (first %zu chars):\n", print_length);
    fwrite(response_body, 1, print_length, stdout);
    printf("\n");

    free(response_body);
} else {
    printf("HTTP request failed\n");
}
}

```

5. Best Practices for Domain-Specific Applications

5.1 Domain Analysis

1. **Requirements Gathering:** Understand domain-specific constraints and requirements
2. **Performance Analysis:** Identify critical performance paths
3. **Resource Constraints:** Account for memory, CPU, and I/O limitations
4. **Safety Considerations:** Implement appropriate error handling and recovery

5.2 Design Patterns

1. **Modular Architecture:** Separate domain logic from infrastructure
2. **Abstraction Layers:** Hide domain complexity behind clean interfaces
3. **Configuration Management:** Support domain-specific configuration
4. **Extensibility:** Design for future domain requirements

5.3 Testing Strategies

1. **Domain-Specific Tests:** Implement tests that validate domain requirements
2. **Edge Case Testing:** Test boundary conditions specific to the domain

3. **Performance Testing:** Validate performance under domain-specific loads
4. **Integration Testing:** Test integration with domain-specific systems

Summary

Domain-specific applications in C require a deep understanding of both the programming language and the target domain. Key concepts covered in this module include:

- Characteristics and requirements of domain-specific applications
- Embedded systems programming with microcontroller integration
- Scientific computing with numerical libraries and algorithms
- Networking applications with custom protocols and HTTP clients
- Best practices for design, implementation, and testing

Developing domain-specific applications in C allows for highly optimized, efficient solutions that can take full advantage of system resources while meeting specialized requirements. Success in this area requires both strong C programming skills and domain expertise.

Module 15.4: Capstone Project

Introduction

The capstone project represents the culmination of your C programming journey, integrating all the concepts, techniques, and best practices you've learned throughout this course. This project challenges you to design, implement, and document a substantial application that demonstrates mastery of advanced C programming concepts while solving a real-world problem.

Learning Objectives

By completing this capstone project, you will demonstrate the ability to:

- Design and implement a complex, multi-module C application
- Apply advanced programming techniques including multi-threading, memory management, and system programming
- Integrate domain-specific knowledge with C programming skills
- Follow professional software development practices including documentation, testing, and version control
- Debug and optimize complex applications
- Present technical work effectively

Project Requirements

1. Technical Requirements

Your capstone project must include:

1. **Multi-file Structure:** Organized into logical modules with proper header files
2. **Advanced C Features:** Use of at least three advanced C concepts (e.g., multithreading, memory mapping, signal handling, etc.)
3. **Error Handling:** Comprehensive error handling with appropriate recovery mechanisms
4. **Memory Management:** Proper dynamic memory allocation and deallocation
5. **Input/Output:** Robust file I/O and/or network communication
6. **Data Structures:** Implementation of appropriate data structures for the problem domain
7. **Testing:** Unit tests and integration tests for critical components
8. **Documentation:** Complete documentation including user guide and API reference

2. Project Scope

The project should be substantial enough to demonstrate advanced skills but manageable within the timeframe. Consider projects that:

- Solve a real-world problem or fill a gap in existing tools
- Demonstrate integration of multiple concepts from the course
- Include both algorithmic complexity and systems programming aspects
- Have clear requirements and measurable outcomes

3. Deliverables

1. **Source Code:** Complete, well-organized C source code
2. **Documentation:**
 - README with project overview, installation, and usage instructions
 - Technical documentation explaining design decisions
 - API documentation for any libraries or modules created
3. **Testing:** Test suite with clear test cases and expected outcomes
4. **Presentation:** A presentation (written or oral) explaining the project
5. **Reflection:** A written reflection on the development process and lessons learned

Project Ideas

1. System Monitoring Tool

Create a comprehensive system monitoring application that collects and displays system metrics:

```
#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <unistd.h>
#include <pthread.h>
#include <time.h>
#include <sys/sysinfo.h>
#include <sys/statvfs.h>

// Data structures for system metrics
typedef struct {
    double cpu_usage;
    unsigned long memory_total;
    unsigned long memory_free;
    unsigned long memory_available;
    double load_average[3];
    time_t timestamp;
} system_metrics_t;
```

```

typedef struct {
    char filesystem[256];
    unsigned long total_space;
    unsigned long free_space;
    unsigned long available_space;
} disk_metrics_t;

// Thread-safe metrics storage
typedef struct {
    system_metrics_t *metrics;
    size_t count;
    size_t capacity;
    pthread_mutex_t mutex;
} metrics_storage_t;

// Global metrics storage
static metrics_storage_t g_metrics = {
    .metrics = NULL,
    .count = 0,
    .capacity = 0,
    .mutex = PTHREAD_MUTEX_INITIALIZER
};

// Function to collect CPU usage
double get_cpu_usage() {
    static unsigned long long prev_idle = 0;
    static unsigned long long prev_total = 0;

    FILE *fp = fopen("/proc/stat", "r");
    if (!fp) return -1.0;

    char buffer[1024];
    fgets(buffer, sizeof(buffer), fp);
    fclose(fp);

    unsigned long long user, nice, system, idle, iowait, irq, softirq, steal;
    sscanf(buffer, "cpu %llu %llu %llu %llu %llu %llu %llu %llu",
           &user, &nice, &system, &idle, &iowait, &irq, &softirq, &steal);

    unsigned long long total = user + nice + system + idle + iowait + irq + softirq + steal;
    unsigned long long total_idle = idle + iowait;

    double cpu_usage = 0.0;
    if (prev_total != 0) {
        unsigned long long total_diff = total - prev_total;
        unsigned long long idle_diff = total_idle - prev_idle;
    }

```

```

        cpu_usage = 100.0 * (total_diff - idle_diff) / total_diff;
    }

    prev_idle = total_idle;
    prev_total = total;

    return cpu_usage;
}

// Function to collect memory information
void get_memory_info(unsigned long *total, unsigned long *free, unsigned long *available) {
    struct sysinfo si;
    if (sysinfo(&si) == 0) {
        *total = si.totalram * si.mem_unit;
        *free = si.freeram * si.mem_unit;
        *available = si.freeram * si.mem_unit + si.bufferram * si.mem_unit;
    }
}

// Function to collect disk space information
int get_disk_info(const char *path, disk_metrics_t *disk) {
    struct statvfs buf;
    if (statvfs(path, &buf) != 0) {
        return -1;
    }

    strncpy(disk->filesystem, path, sizeof(disk->filesystem) - 1);
    disk->total_space = buf.f_blocks * buf.f_frsize;
    disk->free_space = buf.f_bfree * buf.f_frsize;
    disk->available_space = buf.f_bavail * buf.f_frsize;

    return 0;
}

// Function to collect system load average
void get_load_average(double loadavg[3]) {
    FILE *fp = fopen("/proc/loadavg", "r");
    if (fp) {
        fscanf(fp, "%lf %lf %lf", &loadavg[0], &loadavg[1], &loadavg[2]);
        fclose(fp);
    }
}

// Metrics collection thread function
void* metrics_collection_thread(void* arg) {
    while (1) {

```



```

    system_metrics_t metrics = {0};
    metrics.timestamp = time(NULL);

    // Collect metrics
    metrics.cpu_usage = get_cpu_usage();
    get_memory_info(&metrics.memory_total, &metrics.memory_free, &metrics.memory_available);
    get_load_average(metrics.load_average);

    // Store metrics safely
    pthread_mutex_lock(&g_metrics.mutex);

    // Expand storage if needed
    if (g_metrics.count >= g_metrics.capacity) {
        size_t new_capacity = (g_metrics.capacity == 0) ? 10 : g_metrics.capacity * 2;
        system_metrics_t *new_metrics = realloc(g_metrics.metrics,
                                                new_capacity * sizeof(system_metrics_t));

        if (new_metrics) {
            g_metrics.metrics = new_metrics;
            g_metrics.capacity = new_capacity;
        }
    }

    // Add new metrics
    if (g_metrics.count < g_metrics.capacity) {
        g_metrics.metrics[g_metrics.count] = metrics;
        g_metrics.count++;
    }

    pthread_mutex_unlock(&g_metrics.mutex);

    // Wait for 5 seconds before next collection
    sleep(5);
}

return NULL;
}

// Function to display current metrics
void display_current_metrics() {
    pthread_mutex_lock(&g_metrics.mutex);

    if (g_metrics.count > 0) {
        system_metrics_t *latest = &g_metrics.metrics[g_metrics.count - 1];

        printf("=== System Metrics ===\n");
        printf("Timestamp: %s", ctime(&latest->timestamp));
    }
}

```

```

    printf("CPU Usage: %.2f%%\n", latest->cpu_usage);
    printf("Memory Total: %lu MB\n", latest->memory_total / (1024 * 1024));
    printf("Memory Free: %lu MB\n", latest->memory_free / (1024 * 1024));
    printf("Memory Available: %lu MB\n", latest->memory_available / (1024 * 1024));
    printf("Load Average: %.2f, %.2f, %.2f\n",
           latest->load_average[0], latest->load_average[1], latest->load_average[2]);

    // Display disk information
    disk_metrics_t disk;
    if (get_disk_info("/", &disk) == 0) {
        printf("Disk (/): %.2f GB total, %.2f GB free, %.2f GB available\n",
               disk.total_space / (1024.0 * 1024 * 1024),
               disk.free_space / (1024.0 * 1024 * 1024),
               disk.available_space / (1024.0 * 1024 * 1024));
    }
}

pthread_mutex_unlock(&g_metrics.mutex);
}

// Main function for system monitor
int main_system_monitor() {
    printf("Starting System Monitor...\n");

    // Create metrics collection thread
    pthread_t collector_thread;
    if (pthread_create(&collector_thread, NULL, metrics_collection_thread, NULL) != 0) {
        perror("Failed to create metrics collection thread");
        return 1;
    }

    // Main loop - display metrics and handle user input
    char command[10];
    while (1) {
        printf("\nCommands: [d]isplay metrics, [q]uit: ");
        if (fgets(command, sizeof(command), stdin)) {
            if (command[0] == 'd' || command[0] == 'D') {
                display_current_metrics();
            } else if (command[0] == 'q' || command[0] == 'Q') {
                break;
            }
        }
    }
}

// Clean up
pthread_cancel(collector_thread);

```

```

pthread_join(collector_thread, NULL);

if (g_metrics.metrics) {
    free(g_metrics.metrics);
}

printf("System Monitor stopped.\n");
return 0;
}

```

2. Custom Database Engine

Implement a simple embedded database with basic SQL-like functionality:

```

#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <stdint.h>
#include <stdbool.h>

// Data types
typedef enum {
    TYPE_INTEGER,
    TYPE_STRING,
    TYPE_FLOAT
} column_type_t;

// Column definition
typedef struct {
    char name[64];
    column_type_t type;
    size_t offset;
} column_t;

// Table definition
typedef struct {
    char name[64];
    column_t *columns;
    size_t column_count;
    size_t row_size;
    void *data;
    size_t row_count;
    size_t capacity;
} table_t;

// Database structure

```

```

typedef struct {
    table_t *tables;
    size_t table_count;
    size_t capacity;
} database_t;

// Global database instance
static database_t g_db = {0};

// Create a new table
table_t* db_create_table(const char *name, size_t column_count) {
    // Expand tables array if needed
    if (g_db.table_count >= g_db.capacity) {
        size_t new_capacity = (g_db.capacity == 0) ? 4 : g_db.capacity * 2;
        table_t *new_tables = realloc(g_db.tables, new_capacity * sizeof(table_t));
        if (!new_tables) return NULL;

        g_db.tables = new_tables;
        g_db.capacity = new_capacity;
    }

    // Initialize new table
    table_t *table = &g_db.tables[g_db.table_count];
    strncpy(table->name, name, sizeof(table->name) - 1);
    table->columns = calloc(column_count, sizeof(column_t));
    table->column_count = column_count;
    table->row_size = 0;
    table->data = NULL;
    table->row_count = 0;
    table->capacity = 0;

    g_db.table_count++;
    return table;
}

// Add a column to a table
int table_add_column(table_t *table, const char *name, column_type_t type) {
    if (table->column_count >= 64) return -1; // Max columns reached

    column_t *column = &table->columns[table->column_count - 1];
    strncpy(column->name, name, sizeof(column->name) - 1);

    // Calculate offset based on previous columns
    size_t offset = 0;
    for (size_t i = 0; i < table->column_count - 1; i++) {
        switch (table->columns[i].type) {
            case TYPE_INTEGER:

```

```

        offset = (offset + sizeof(int) - 1) & ~(sizeof(int) - 1);
        offset += sizeof(int);
        break;
    case TYPE_FLOAT:
        offset = (offset + sizeof(double) - 1) & ~(sizeof(double) - 1);
        offset += sizeof(double);
        break;
    case TYPE_STRING:
        offset = (offset + sizeof(char*) - 1) & ~(sizeof(char*) - 1);
        offset += sizeof(char*);
        break;
    }
}

column->type = type;
column->offset = offset;

// Update row size
switch (type) {
    case TYPE_INTEGER:
        table->row_size = (offset + sizeof(int) - 1) & ~(sizeof(int) - 1);
        table->row_size += sizeof(int);
        break;
    case TYPE_FLOAT:
        table->row_size = (offset + sizeof(double) - 1) & ~(sizeof(double) - 1);
        table->row_size += sizeof(double);
        break;
    case TYPE_STRING:
        table->row_size = (offset + sizeof(char*) - 1) & ~(sizeof(char*) - 1);
        table->row_size += sizeof(char*);
        break;
}

return 0;
}

// Insert a row into a table
int table_insert_row(table_t *table, void *row_data) {
    // Expand data array if needed
    if (table->row_count >= table->capacity) {
        size_t new_capacity = (table->capacity == 0) ? 16 : table->capacity * 2;
        void *new_data = realloc(table->data, new_capacity * table->row_size);
        if (!new_data) return -1;

        table->data = new_data;
        table->capacity = new_capacity;
    }
}

```

```

    // Copy row data
    void *dest = (char*)table->data + table->row_count * table->row_size;
    memcpy(dest, row_data, table->row_size);
    table->row_count++;

    return 0;
}

// Simple query execution
void db_execute_query(const char *query) {
    printf("Executing query: %s\n", query);

    // Simple parser for demonstration
    if (strncmp(query, "CREATE TABLE", 12) == 0) {
        // Parse table creation
        char table_name[64];
        if (sscanf(query, "CREATE TABLE %63s", table_name) == 1) {
            table_t *table = db_create_table(table_name, 4); // Max 4 columns for demo
            if (table) {
                printf("Table '%s' created successfully\n", table_name);
            } else {
                printf("Failed to create table '%s'\n", table_name);
            }
        }
    } else if (strncmp(query, "INSERT INTO", 11) == 0) {
        printf("Insert operation executed\n");
    } else if (strncmp(query, "SELECT", 6) == 0) {
        printf("Select query executed\n");
    } else {
        printf("Unsupported query\n");
    }
}

// Database initialization
void db_init() {
    g_db.tables = NULL;
    g_db.table_count = 0;
    g_db.capacity = 0;
    printf("Database initialized\n");
}

// Database cleanup
void db_cleanup() {
    for (size_t i = 0; i < g_db.table_count; i++) {
        table_t *table = &g_db.tables[i];
        free(table->columns);
    }
}

```

```

        free(table->data);
    }
    free(g_db.tables);
    g_db.tables = NULL;
    g_db.table_count = 0;
    g_db.capacity = 0;
    printf("Database cleaned up\n");
}

// Main function for database engine
int main_database_engine() {
    db_init();

    printf("Custom Database Engine\n");
    printf("Enter SQL-like commands (type 'quit' to exit):\n");

    char query[256];
    while (1) {
        printf("> ");
        if (!fgets(query, sizeof(query), stdin)) break;

        // Remove newline
        query[strcspn(query, "\n")] = 0;

        if (strcmp(query, "quit") == 0) break;

        db_execute_query(query);
    }

    db_cleanup();
    return 0;
}

```

3. Network File Transfer System

Create a secure file transfer system with client-server architecture:

```

#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <unistd.h>
#include <sys/socket.h>
#include <netinet/in.h>
#include <arpa/inet.h>
#include <sys/stat.h>
#include <fcntl.h>

```

```

#include <pthread.h>

#define PORT 8080
#define BUFFER_SIZE 4096
#define MAX_FILENAME 256

// Protocol commands
#define CMD_UPLOAD 1
#define CMD_DOWNLOAD 2
#define CMD_LIST 3
#define CMD_QUIT 4

// Protocol packet structure
typedef struct {
    uint32_t command;
    uint32_t data_length;
    char filename[MAX_FILENAME];
} __attribute__((packed)) protocol_header_t;

// Thread arguments for client handling
typedef struct {
    int client_socket;
    struct sockaddr_in client_addr;
} client_args_t;

// Function to send a file to client
int send_file(int client_socket, const char *filename) {
    int file_fd = open(filename, O_RDONLY);
    if (file_fd < 0) {
        perror("Failed to open file");
        return -1;
    }

    // Get file size
    struct stat st;
    if (fstat(file_fd, &st) < 0) {
        perror("Failed to get file stats");
        close(file_fd);
        return -1;
    }

    // Send file size to client
    uint32_t file_size = htonl(st.st_size);
    if (send(client_socket, &file_size, sizeof(file_size), 0) < 0) {
        perror("Failed to send file size");
        close(file_fd);
        return -1;
    }

```



```

    }

    // Send file data
    char buffer[BUFFER_SIZE];
    ssize_t bytes_read;
    while ((bytes_read = read(file_fd, buffer, BUFFER_SIZE)) > 0) {
        if (send(client_socket, buffer, bytes_read, 0) < 0) {
            perror("Failed to send file data");
            close(file_fd);
            return -1;
        }
    }

    close(file_fd);
    return 0;
}

// Function to receive a file from client
int receive_file(int client_socket, const char *filename) {
    // Receive file size
    uint32_t file_size;
    if (recv(client_socket, &file_size, sizeof(file_size), 0) < 0) {
        perror("Failed to receive file size");
        return -1;
    }

    file_size = ntohl(file_size);

    // Create file
    int file_fd = open(filename, O_CREAT | O_WRONLY | O_TRUNC, 0644);
    if (file_fd < 0) {
        perror("Failed to create file");
        return -1;
    }

    // Receive file data
    char buffer[BUFFER_SIZE];
    size_t total_received = 0;

    while (total_received < file_size) {
        size_t to_receive = (file_size - total_received < BUFFER_SIZE) ?
            (file_size - total_received) : BUFFER_SIZE;

        ssize_t bytes_received = recv(client_socket, buffer, to_receive, 0);
        if (bytes_received < 0) {
            perror("Failed to receive file data");
            close(file_fd);

```

```

        return -1;
    }

    if (write(file_fd, buffer, bytes_received) < 0) {
        perror("Failed to write file data");
        close(file_fd);
        return -1;
    }

    total_received += bytes_received;
}

close(file_fd);
return 0;
}

// Function to list files in directory
void list_files(int client_socket) {
    FILE *fp = popen("ls -l", "r");
    if (!fp) {
        perror("Failed to execute ls command");
        return;
    }

    char buffer[BUFFER_SIZE];
    size_t total_size = 0;

    // First pass: calculate total size
    while (fgets(buffer, sizeof(buffer), fp)) {
        total_size += strlen(buffer);
    }

    // Send total size
    uint32_t size = htonl(total_size);
    send(client_socket, &size, sizeof(size), 0);

    // Second pass: send file list
    rewind(fp);
    while (fgets(buffer, sizeof(buffer), fp)) {
        send(client_socket, buffer, strlen(buffer), 0);
    }

    pclose(fp);
}

// Client handling thread
void* handle_client(void* arg) {

```

```

client_args_t *client_args = (client_args_t*)arg;
int client_socket = client_args->client_socket;

printf("Client connected: %s:%d\n",
       inet_ntoa(client_args->client_addr.sin_addr),
       ntohs(client_args->client_addr.sin_port));

free(client_args);

// Handle client requests
while (1) {
    protocol_header_t header;
    ssize_t bytes_received = recv(client_socket, &header, sizeof(header), 0);

    if (bytes_received <= 0) {
        break; // Client disconnected or error
    }

    header.command = ntohl(header.command);
    header.data_length = ntohl(header.data_length);

    switch (header.command) {
        case CMD_UPLOAD:
            printf("Upload request for file: %s\n", header.filename);
            if (receive_file(client_socket, header.filename) == 0) {
                printf("File uploaded successfully: %s\n", header.filename);
            } else {
                printf("Failed to upload file: %s\n", header.filename);
            }
            break;

        case CMD_DOWNLOAD:
            printf("Download request for file: %s\n", header.filename);
            if (send_file(client_socket, header.filename) == 0) {
                printf("File sent successfully: %s\n", header.filename);
            } else {
                printf("Failed to send file: %s\n", header.filename);
            }
            break;

        case CMD_LIST:
            printf("File list request\n");
            list_files(client_socket);
            break;

        case CMD_QUIT:
            printf("Client requested to quit\n");

```

```

        goto cleanup;

    default:
        printf("Unknown command: %u\n", header.command);
        break;
    }
}

cleanup:
    close(client_socket);
    printf("Client disconnected\n");
    return NULL;
}

// Server main function
int main_file_transfer_server() {
    int server_fd;
    struct sockaddr_in address;
    int opt = 1;

    // Create socket
    if ((server_fd = socket(AF_INET, SOCK_STREAM, 0)) == 0) {
        perror("socket failed");
        exit(EXIT_FAILURE);
    }

    // Set socket options
    if (setsockopt(server_fd, SOL_SOCKET, SO_REUSEADDR | SO_REUSEPORT,
        &opt, sizeof(opt))) {
        perror("setsockopt failed");
        exit(EXIT_FAILURE);
    }

    address.sin_family = AF_INET;
    address.sin_addr.s_addr = INADDR_ANY;
    address.sin_port = htons(PORT);

    // Bind socket
    if (bind(server_fd, (struct sockaddr *)&address, sizeof(address)) < 0) {
        perror("bind failed");
        exit(EXIT_FAILURE);
    }

    // Listen for connections
    if (listen(server_fd, 3) < 0) {
        perror("listen failed");
        exit(EXIT_FAILURE);
    }
}

```

```

}

printf("File transfer server listening on port %d\n", PORT);

// Accept connections
while (1) {
    int client_socket;
    struct sockaddr_in client_addr;
    socklen_t client_addr_len = sizeof(client_addr);

    if ((client_socket = accept(server_fd, (struct sockaddr *)&client_addr,
                                &client_addr_len)) < 0) {
        perror("accept failed");
        continue;
    }

    // Create thread to handle client
    client_args_t *client_args = malloc(sizeof(client_args_t));
    client_args->client_socket = client_socket;
    client_args->client_addr = client_addr;

    pthread_t client_thread;
    if (pthread_create(&client_thread, NULL, handle_client, client_args) != 0) {
        perror("Failed to create client thread");
        free(client_args);
        close(client_socket);
    } else {
        pthread_detach(client_thread); // Automatically clean up thread
    }
}

close(server_fd);
return 0;
}

// Client main function
int main_file_transfer_client(const char *server_ip) {
    int client_fd;
    struct sockaddr_in serv_addr;

    // Create socket
    if ((client_fd = socket(AF_INET, SOCK_STREAM, 0)) < 0) {
        perror("Socket creation error");
        return -1;
    }
}

```

```

serv_addr.sin_family = AF_INET;
serv_addr.sin_port = htons(PORT);

// Convert IPv4 address from text to binary
if (inet_pton(AF_INET, server_ip, &serv_addr.sin_addr) <= 0) {
    printf("Invalid address/ Address not supported\n");
    return -1;
}

// Connect to server
if (connect(client_fd, (struct sockaddr *)&serv_addr, sizeof(serv_addr)) < 0) {
    perror("Connection Failed");
    return -1;
}

printf("Connected to server %s:%d\n", server_ip, PORT);

// Simple client interface
char command[10];
while (1) {
    printf("\nCommands: [u]pload, [d]ownload, [l]ist, [q]uit: ");
    if (fgets(command, sizeof(command), stdin)) {
        if (command[0] == 'u' || command[0] == 'U') {
            char filename[MAX_FILENAME];
            printf("Enter filename to upload: ");
            if (fgets(filename, sizeof(filename), stdin)) {
                filename[strlen(filename) - 1] = 0; // Remove newline

                protocol_header_t header = {0};
                header.command = htonl(CMD_UPLOAD);
                strncpy(header.filename, filename, sizeof(header.filename) - 1);

                send(client_fd, &header, sizeof(header), 0);
                if (receive_file(client_fd, filename) == 0) {
                    printf("File uploaded successfully\n");
                } else {
                    printf("Failed to upload file\n");
                }
            }
        }
        else if (command[0] == 'd' || command[0] == 'D') {
            char filename[MAX_FILENAME];
            printf("Enter filename to download: ");
            if (fgets(filename, sizeof(filename), stdin)) {
                filename[strlen(filename) - 1] = 0; // Remove newline

                protocol_header_t header = {0};
                header.command = htonl(CMD_DOWNLOAD);

```

```

        strncpy(header.filename, filename, sizeof(header.filename) - 1);

        send(client_fd, &header, sizeof(header), 0);
        if (send_file(client_fd, filename) == 0) {
            printf("File downloaded successfully\n");
        } else {
            printf("Failed to download file\n");
        }
    }
} else if (command[0] == 'l' || command[0] == 'L') {
    protocol_header_t header = {0};
    header.command = htonl(CMD_LIST);
    send(client_fd, &header, sizeof(header), 0);

    // Receive file list
    uint32_t list_size;
    recv(client_fd, &list_size, sizeof(list_size), 0);
    list_size = ntohl(list_size);

    char *list_buffer = malloc(list_size + 1);
    if (list_buffer) {
        recv(client_fd, list_buffer, list_size, 0);
        list_buffer[list_size] = '\0';
        printf("Files on server:\n%s", list_buffer);
        free(list_buffer);
    }
} else if (command[0] == 'q' || command[0] == 'Q') {
    protocol_header_t header = {0};
    header.command = htonl(CMD_QUIT);
    send(client_fd, &header, sizeof(header), 0);
    break;
}
}

close(client_fd);
return 0;
}

```

Project Development Process

1. Planning Phase

1. **Project Selection:** Choose a project that interests you and matches your skill level
2. **Requirements Analysis:** Define clear, measurable requirements
3. **Design Documentation:** Create architecture and design documents

4. **Technology Selection:** Choose appropriate libraries, tools, and frameworks
5. **Timeline Planning:** Create a realistic development timeline

2. Implementation Phase

1. **Environment Setup:** Configure development environment and tools
2. **Version Control:** Initialize Git repository and establish branching strategy
3. **Modular Development:** Implement features incrementally
4. **Testing:** Write and execute tests for each module
5. **Documentation:** Maintain documentation throughout development

3. Testing and Debugging Phase

1. **Unit Testing:** Test individual functions and modules
2. **Integration Testing:** Test module interactions
3. **System Testing:** Test complete system functionality
4. **Performance Testing:** Validate performance requirements
5. **Debugging:** Identify and fix issues

4. Deployment and Presentation Phase

1. **Final Testing:** Conduct comprehensive testing
2. **Documentation Finalization:** Complete all documentation
3. **Packaging:** Package application for distribution
4. **Presentation Preparation:** Prepare presentation materials
5. **Project Submission:** Submit all deliverables

Evaluation Criteria

Your capstone project will be evaluated based on:

Technical Excellence (40%)

- Code quality and organization
- Proper use of C programming concepts
- Error handling and robustness
- Performance optimization
- Security considerations

Design and Architecture (25%)

- System design quality
- Modularity and maintainability
- Appropriate use of data structures and algorithms
- Scalability considerations

Documentation (20%)

- Code comments and readability
- User documentation
- Technical documentation
- API documentation

Testing (10%)

- Test coverage
- Test quality
- Bug identification and resolution

Presentation (5%)

- Clarity of presentation
- Technical depth
- Professionalism

Best Practices

1. Code Quality

- Follow consistent coding standards
- Use meaningful variable and function names
- Write modular, reusable code
- Include comprehensive error handling
- Document complex algorithms and decisions

2. Project Management

- Use version control effectively
- Break project into manageable milestones
- Track progress and adjust plans as needed
- Seek feedback regularly
- Maintain regular backups

3. Collaboration

- If working in a team, establish clear roles
- Use collaborative tools effectively
- Communicate regularly with team members
- Review each other's code
- Resolve conflicts constructively

Resources and Support

Tools and Libraries

- **Development Environment:** GCC, Clang, or preferred C compiler
- **Debugging:** GDB, Valgrind
- **Version Control:** Git
- **Documentation:** Doxygen for API documentation
- **Testing:** CUnit, Unity, or custom testing framework

Learning Resources

- **Books:** “The C Programming Language” by Kernighan and Ritchie
- **Online:** C standard documentation, man pages
- **Communities:** Stack Overflow, C programming forums
- **Tutorials:** System programming tutorials, networking guides

Conclusion

The capstone project is your opportunity to demonstrate mastery of C programming and apply it to solve real-world problems. Choose a project that challenges you but is achievable within the timeframe. Focus on writing clean, well-documented code and following professional software development practices.

Remember that the goal is not just to complete a project, but to learn and grow as a developer. Take the time to understand the concepts you're implementing, and don't hesitate to explore beyond the basic requirements. Good luck with your capstone project!

Workbook

Examples

Foundations

Data Types

Control Flow

Functions

Arrays and Strings

Pointers

Structures

File I/O

Modern C

Data Structures

Network

Embedded

Performance

Testing

Advanced

Exercises

Foundations

Module 1: Foundations Exercises

Exercise 1: Basic Program Structure

Write a C program that prints the following output:

```
*****
*   My First C Program   *
*****
```

Requirements: - Use proper program structure with `#include` and `main()` function - Include comments explaining each part of your code - Compile and run your program successfully

Exercise 2: Personal Information Display

Create a program that displays your personal information in a formatted way: - Name - Age - City/Country - Favorite programming language - Years of programming experience

Requirements: - Use at least 3 different `printf` format specifiers - Use both single-line and multi-line comments - Make the output visually appealing with proper formatting

Exercise 3: Simple Calculator Input

Write a program that asks the user for two numbers and then displays: - The sum - The difference - The product - The quotient (handle division by zero)

Requirements: - Use `scanf()` to get input from the user - Check the return value of `scanf()` for error handling - Use proper variable declarations - Format the output clearly

Exercise 4: Character Analysis

Create a program that asks the user to enter a character and then displays: - The character itself - Its ASCII value - Whether it's a vowel (a, e, i, o, u) - Whether it's uppercase or lowercase (if it's a letter)

Requirements: - Use `getchar()` for character input - Handle whitespace properly - Use conditional statements (will be covered in detail in Module 3) - Display results in a formatted manner

Exercise 5: Temperature Converter

Write a program that converts temperatures between Celsius and Fahrenheit: - Ask the user for a temperature value - Ask the user to specify C for Celsius or F for Fahrenheit - Convert and display the equivalent temperature in the other unit

Formulas: - $F = (C \times 9/5) + 32$ - $C = (F - 32) \times 5/9$

Requirements: - Use proper input validation - Format output to 2 decimal places - Handle both uppercase and lowercase input for unit specification - Include clear prompts and instructions

Exercise 6: String Formatting Practice

Create a program that demonstrates various printf formatting options: - Display integers with different width specifications - Display floating-point numbers with different precision - Display strings with width and alignment options - Use escape sequences for formatting

Requirements: - Demonstrate at least 5 different printf format options - Include examples of left and right alignment - Use at least 3 different data types - Comment your code explaining each formatting choice

Exercise 7: Input Buffer Management

Write a program that demonstrates proper input buffer management: - Read a string with spaces using `fgets()` - Read an integer after the string - Read a character after the integer - Show how to properly handle the input buffer between different input types

Requirements: - Avoid common scanf pitfalls with buffer management - Use appropriate buffer sizes to prevent overflow - Handle input errors gracefully - Include comments explaining buffer management techniques

Exercise 8: Comprehensive Demo Program

Create a comprehensive program that integrates all Module 1 concepts: - Proper program structure with includes and main function - Multiple function calls (at least 2 custom functions) - Various printf formatting options - Multiple scanf inputs with error checking - Character I/O demonstrations - Well-commented code with both comment types

Requirements: - At least 50 lines of code - Use all major concepts from Module 1 - Include a README section in comments explaining your program - Compile without warnings (use `-Wall -Wextra` flags)

Solutions and Tips

Exercise 1 Solution Example:

```
#include <stdio.h>

int main() {
    // Print a decorative header
    printf("*****\n");
    printf("*   My First C Program   *\n");
    printf("*****\n");
    return 0;
}
```

Exercise 3 Solution Example:

```
#include <stdio.h>

int main() {
    float num1, num2;

    printf("Enter first number: ");
    if (scanf("%f", &num1) != 1) {
        printf("Invalid input!\n");
        return 1;
    }

    printf("Enter second number: ");
    if (scanf("%f", &num2) != 1) {
        printf("Invalid input!\n");
        return 1;
    }

    printf("\nResults:\n");
    printf("%.2f + %.2f = %.2f\n", num1, num2, num1 + num2);
    printf("%.2f - %.2f = %.2f\n", num1, num2, num1 - num2);
    printf("%.2f * %.2f = %.2f\n", num1, num2, num1 * num2);

    if (num2 != 0) {
        printf("%.2f / %.2f = %.2f\n", num1, num2, num1 / num2);
    } else {
        printf("Cannot divide by zero!\n");
    }

    return 0;
}
```



```
}
```

Common Pitfalls to Avoid:

1. **Missing & in scanf:** Always use & for non-string variables
2. **Buffer overflow:** Limit string input lengths
3. **Whitespace issues:** Handle newlines properly with character input
4. **Missing return 0:** Always return 0 from main function
5. **Forgetting headers:** Include <stdio.h> for I/O functions

Compilation Tips:

```
# Basic compilation
gcc program.c -o program

# With warnings enabled
gcc -Wall -Wextra program.c -o program

# With debugging information
gcc -g program.c -o program

# With C standard specification
gcc -std=c99 program.c -o program
```

Complete these exercises to solidify your understanding of Module 1 concepts. Each exercise builds upon the previous ones, gradually increasing in complexity.

Data Types

Module 2: Data Types and Variables Exercises

Exercise 1: Data Type Explorer

Write a program that displays the size and range of all fundamental data types on your system. Include:

- All integer types (char, short, int, long, long long) and their unsigned variants
- All floating-point types (float, double, long double)
- Fixed-width integer types (int8_t, int16_t, int32_t, int64_t)
- Boolean type

Requirements:

- Use `sizeof` operator to determine sizes
- Use `limits.h`, `float.h`, and `stdint.h` headers for range information
- Format output clearly with labels
- Include both signed and unsigned variants where applicable

Exercise 2: Variable Declaration Challenge

Create a program that declares variables using different naming conventions and demonstrates proper initialization:

- Declare 5 variables using snake_case convention
- Declare 5 variables using camelCase convention
- Initialize all variables at declaration
- Demonstrate the use of const variables
- Show examples of multiple variable declarations in single statements

Requirements:

- Use at least 3 different data types
- Include both global and local variables
- Demonstrate scope differences
- Use meaningful variable names related to a real-world scenario (e.g., student records, inventory system)

Exercise 3: Constants Comparison

Write a program that demonstrates the three methods of creating constants in C:

- Using `#define` preprocessor directive
- Using `const` qualifier
- Using `enum` enumeration

Create constants for a real-world application (e.g., a calendar system with months, days, or a game with difficulty levels).

Requirements:

- Implement all three constant creation methods
- Show advantages and disadvantages of each method
- Demonstrate usage in calculations or conditional statements
- Include string constants where appropriate

Exercise 4: Storage Class Demonstrator

Create a program that demonstrates the behavior of different storage classes: - **auto** (automatic) - **static** (static) - **register** (register) - **extern** (external)

For **extern**, create two source files to demonstrate external linkage.

Requirements: - Show how static variables retain values between function calls - Demonstrate the difference between local and global scope - Show how register variables might be optimized - Implement proper extern usage with header files

Exercise 5: Type Conversion Safety

Write a program that safely handles type conversions between different data types: - Integer to floating-point and vice versa - Large integer to smaller integer (handle overflow) - Signed to unsigned and vice versa - Character to integer conversions

Include error checking and validation for unsafe conversions.

Requirements: - Check for potential overflow before conversion - Handle truncation appropriately - Use standard library functions for string conversion (strtol, strtod) - Provide meaningful error messages for failed conversions

Exercise 6: Operator Expression Builder

Create a program that evaluates complex mathematical expressions using various operators: - Arithmetic operators with mixed data types - Relational and logical operators in compound expressions - Bitwise operators for bit manipulation - Assignment operators for cumulative calculations - Conditional operators for decision making

Requirements: - Include expressions with different precedence levels - Use parentheses to clarify intent - Demonstrate short-circuit evaluation - Show the difference between pre and post increment/decrement

Exercise 7: Bit Manipulation Toolkit

Develop a set of functions that perform common bit manipulation operations: - Function to check if a specific bit is set - Function to set a specific bit - Function to clear a specific bit - Function to toggle a specific bit - Function to count the number of set bits - Function to reverse all bits in a byte

Requirements: - Implement all functions using bitwise operators - Include parameter validation - Provide clear documentation for each function - Test with various input values

Exercise 8: Data Type Selection Advisor

Create a program that helps users select appropriate data types based on their requirements: - Input: Range of values, precision requirements, memory constraints - Output: Recommended data type with justification - Include considerations for portability and performance

Requirements: - Consider both value range and memory usage - Account for signed vs unsigned requirements - Include fixed-width type recommendations - Provide warnings for potential overflow situations

Exercise 9: Comprehensive Calculator

Implement a calculator that demonstrates advanced use of operators and type handling: - Basic arithmetic operations - Bitwise operations - Logical operations - Type-safe conversions - Error handling for invalid operations (division by zero, etc.)

Requirements: - Support both integer and floating-point modes - Handle operator precedence correctly - Include memory functions (store/recall) - Provide clear error messages - Support both interactive and batch modes

Exercise 10: Real-World Application

Design and implement a complete program that integrates all Module 2 concepts: - A student grade management system - An inventory tracking system - A simple banking application - A measurement conversion tool

Requirements: - Use appropriate data types for all values - Implement proper variable scoping - Include constants for fixed values - Handle type conversions safely - Use operators effectively for calculations - Apply storage classes where appropriate - Include comprehensive error checking

Solutions and Tips

Exercise 1 Solution Example:

```
#include <stdio.h>
#include <limits.h>
#include <float.h>
#include <stdint.h>

int main() {
    printf("=== Data Type Explorer ===\n\n");

    // Integer types
    printf("Integer Types:\n");
```

```

printf("char: %zu bytes, range %d to %d\n", sizeof(char), CHAR_MIN, CHAR_MAX);
printf("unsigned char: %zu bytes, range 0 to %u\n", sizeof(unsigned char), UCHAR_MAX);
printf("short: %zu bytes, range %d to %d\n", sizeof(short), SHRT_MIN, SHRT_MAX);
printf("unsigned short: %zu bytes, range 0 to %u\n", sizeof(unsigned short), USHRT_MAX);
// ... continue for other types

return 0;
}

```

Exercise 3 Solution Example:

```

#include <stdio.h>

// Using #define
#define MAX_STUDENTS 100
#define UNIVERSITY_NAME "Tech University"

// Using const
const float PI = 3.14159f;
const int PASSING_GRADE = 60;

// Using enum
enum Difficulty {
    EASY = 1,
    MEDIUM = 2,
    HARD = 3
};

enum Days {
    MONDAY = 1,
    TUESDAY,
    WEDNESDAY,
    THURSDAY,
    FRIDAY,
    SATURDAY,
    SUNDAY
};

int main() {
    printf("University: %s\n", UNIVERSITY_NAME);
    printf("Max students: %d\n", MAX_STUDENTS);
    printf("PI: %.5f\n", PI);
    printf("Passing grade: %d\n", PASSING_GRADE);
    printf("Difficulty level: %d\n", MEDIUM);
    printf("Day of week: %d\n", WEDNESDAY);
}

```

```

    return 0;
}

```

Exercise 6 Solution Example:

```

#include <stdio.h>

int main() {
    int a = 10, b = 3;
    float x = 15.7f, y = 4.2f;

    // Mixed arithmetic with precedence
    int result1 = a + b * 2;           // 10 + (3 * 2) = 16
    float result2 = (a + b) * x / y;   // ((10 + 3) * 15.7) / 4.2

    printf("Integer expression: %d + %d * 2 = %d\n", a, b, result1);
    printf("Float expression: (%d + %d) * %.1f / %.1f = %.2f\n", a, b, x, y, result2);

    // Logical expressions with short-circuit
    if (b != 0 && a / b > 1) {
        printf("Division is safe and result > 1\n");
    }

    // Conditional operator
    char *status = (a > b) ? "Greater" : "Less or Equal";
    printf("Status: %s\n", status);

    return 0;
}

```

Exercise 7 Solution Example:

```

#include <stdio.h>

// Check if bit at position is set
int is_bit_set(unsigned int num, int position) {
    return (num >> position) & 1;
}

// Set bit at position
unsigned int set_bit(unsigned int num, int position) {
    return num | (1 << position);
}

```

```

// Clear bit at position
unsigned int clear_bit(unsigned int num, int position) {
    return num & ~(1 << position);
}

// Toggle bit at position
unsigned int toggle_bit(unsigned int num, int position) {
    return num ^ (1 << position);
}

int main() {
    unsigned int number = 0b10101010;

    printf("Original: %08b\n", number);
    printf("Bit 3 set? %d\n", is_bit_set(number, 3));

    number = set_bit(number, 2);
    printf("After setting bit 2: %08b\n", number);

    number = clear_bit(number, 7);
    printf("After clearing bit 7: %08b\n", number);

    number = toggle_bit(number, 5);
    printf("After toggling bit 5: %08b\n", number);

    return 0;
}

```

Common Pitfalls to Avoid

1. **Uninitialized Variables:** Always initialize variables at declaration
2. **Integer Overflow:** Check ranges before conversions
3. **Division by Zero:** Always validate divisors
4. **Implicit Conversions:** Be aware of precision loss
5. **Operator Precedence:** Use parentheses for clarity
6. **Side Effects in Expressions:** Avoid undefined behavior with ++/-- operators
7. **Signed/Unsigned Comparisons:** Be careful when mixing signed and unsigned types
8. **Bitwise vs Logical Operators:** Use &/| for bit operations, &&/|| for logical operations

Compilation Tips

```
# Basic compilation
gcc program.c -o program

# With warnings for type safety
gcc -Wall -Wextra program.c -o program

# With standard specification
gcc -std=c99 program.c -o program

# With debugging information
gcc -g -Wall program.c -o program
```

Complete these exercises to solidify your understanding of Module 2 concepts. Each exercise builds upon the previous ones, gradually increasing in complexity.

Control Flow

Module 3: Control Flow Exercises

Exercise 1: Grade Calculator with Multiple Conditions

Write a program that calculates letter grades based on numerical scores using nested if-else statements. Include: - Input validation for scores (0-100) - Different grading scales for different courses - Bonus points for perfect attendance - Extra credit opportunities

Requirements: - Use nested conditionals for complex grading logic - Implement multiple grading scales (e.g., strict, lenient, standard) - Include bonus point calculations - Provide detailed feedback for each grade range

Exercise 2: Pattern Generator

Create a program that generates various number and character patterns using nested loops: - Right-angled triangles - Pyramids - Diamond patterns - Floyd's triangle - Pascal's triangle (advanced)

Requirements: - Use nested loops for pattern generation - Allow user to select pattern type - Implement input validation for pattern size - Include both number and character patterns

Exercise 3: Menu-Driven Calculator

Develop a comprehensive calculator using do-while loops and switch statements: - Basic arithmetic operations - Scientific functions (sin, cos, tan, log, sqrt) - Memory functions (store, recall, clear) - History of calculations - Unit conversions

Requirements: - Use do-while for main menu loop - Implement switch for operation selection - Include error handling for invalid inputs - Provide clear exit mechanism

Exercise 4: Prime Number Analyzer

Write a program that finds and analyzes prime numbers using various loop constructs: - Prime number detection - Prime factorization - Twin prime identification - Prime number distribution analysis - Sieve of Eratosthenes implementation

Requirements: - Use efficient algorithms for prime detection - Implement multiple loop types (for, while) - Include performance timing - Provide statistical analysis of results

Exercise 5: Resource Manager with goto

Create a resource management system that demonstrates proper use of goto for error handling:
- File operations - Memory allocation - Network connections (simulated) - Database connections (simulated) - Cleanup procedures

Requirements: - Use goto for structured error handling - Implement proper resource cleanup - Include error simulation - Demonstrate resource leak prevention

Exercise 6: Function Pointer Calculator

Develop a calculator that uses function pointers for operation selection:
- Basic arithmetic operations
- Advanced mathematical functions - Custom function registration - Dynamic operation selection - Plugin architecture simulation

Requirements: - Use function pointers for operation dispatch - Implement callback mechanisms - Include function registration system - Provide extensibility examples

Exercise 7: State Machine Implementation

Create a state machine for a real-world system:
- Vending machine controller - Traffic light system
- Bank account management - Game character states - Network protocol handler

Requirements: - Use enum for state definitions - Implement state transition logic - Include event handling - Provide state visualization

Exercise 8: Error Handling Framework

Design an error handling system that demonstrates various error management techniques:
- Return code patterns - errno simulation - Exception-like behavior with setjmp/longjmp - Logging mechanisms - Recovery procedures

Requirements: - Implement multiple error handling approaches - Include error code definitions - Provide error message system - Demonstrate error recovery

Exercise 9: Complex Loop Optimization

Write programs that demonstrate loop optimization techniques:
- Loop unrolling - Loop fusion - Loop invariant code motion - Strength reduction - Cache-friendly iterations

Requirements: - Compare performance of optimized vs. unoptimized loops - Include timing measurements - Provide analysis of optimization effects - Demonstrate when optimizations are beneficial

Exercise 10: Comprehensive Control Flow Application

Design and implement a complete application that integrates all control flow concepts: - A simple game (tic-tac-toe, snake, etc.) - A data processing system - A simulation program - A configuration tool - A monitoring application

Requirements: - Use all control flow constructs appropriately - Include complex decision-making logic - Implement proper error handling - Provide user-friendly interface - Include comprehensive documentation

Solutions and Tips

Exercise 1 Solution Example:

```
#include <stdio.h>

int main() {
    float score;
    int attendance;
    int extra_credit;

    printf("Enter score (0-100): ");
    scanf("%f", &score);

    printf("Enter attendance (0-100%%): ");
    scanf("%d", &attendance);

    printf("Enter extra credit points (0-10): ");
    scanf("%d", &extra_credit);

    // Input validation
    if (score < 0 || score > 100) {
        printf("Invalid score!\n");
        return 1;
    }

    if (attendance < 0 || attendance > 100) {
        printf("Invalid attendance!\n");
        return 1;
    }

    if (extra_credit < 0 || extra_credit > 10) {
        printf("Invalid extra credit!\n");
        return 1;
    }
}
```

```

// Apply bonus for perfect attendance
if (attendance == 100) {
    score += 2.0f;
    if (score > 100) score = 100;
}

// Add extra credit
score += extra_credit;
if (score > 100) score = 100;

// Determine grade
if (score >= 97) {
    printf("Grade: A+ (Excellent)\n");
} else if (score >= 93) {
    printf("Grade: A (Outstanding)\n");
} else if (score >= 90) {
    printf("Grade: A- (Very Good)\n");
} else if (score >= 87) {
    printf("Grade: B+ (Good)\n");
} else if (score >= 83) {
    printf("Grade: B (Satisfactory)\n");
} else if (score >= 80) {
    printf("Grade: B- (Above Average)\n");
} else if (score >= 77) {
    printf("Grade: C+ (Average)\n");
} else if (score >= 73) {
    printf("Grade: C (Below Average)\n");
} else if (score >= 70) {
    printf("Grade: C- (Poor)\n");
} else if (score >= 60) {
    printf("Grade: D (Very Poor)\n");
} else {
    printf("Grade: F (Fail)\n");
}

printf("Final score: %.1f\n", score);

return 0;
}

```

Exercise 2 Solution Example:

```
#include <stdio.h>

void print_pyramid(int rows) {
    for (int i = 1; i <= rows; i++) {
        // Print spaces
        for (int j = 1; j <= rows - i; j++) {
            printf(" ");
        }
        // Print stars
        for (int k = 1; k <= 2 * i - 1; k++) {
            printf("*");
        }
        printf("\n");
    }
}

void print_floyds_triangle(int rows) {
    int num = 1;
    for (int i = 1; i <= rows; i++) {
        for (int j = 1; j <= i; j++) {
            printf("%d ", num++);
        }
        printf("\n");
    }
}

int main() {
    int choice, rows;

    printf("Pattern Generator\n");
    printf("1. Pyramid\n");
    printf("2. Floyd's Triangle\n");
    printf("Enter choice: ");
    scanf("%d", &choice);

    printf("Enter number of rows: ");
    scanf("%d", &rows);

    if (rows <= 0 || rows > 20) {
        printf("Invalid number of rows!\n");
        return 1;
    }

    switch (choice) {
        case 1:
```

```

        print_pyramid(rows);
        break;
    case 2:
        print_floyds_triangle(rows);
        break;
    default:
        printf("Invalid choice!\n");
}

return 0;
}

```

Exercise 3 Solution Example:

```

#include <stdio.h>
#include <math.h>

double memory = 0.0;
double history[100];
int history_count = 0;

void add_to_history(double result) {
    if (history_count < 100) {
        history[history_count++] = result;
    }
}

void show_history() {
    printf("\nCalculation History:\n");
    for (int i = 0; i < history_count && i < 10; i++) {
        printf("%d. %.2f\n", i + 1, history[i]);
    }
}

int main() {
    int choice;
    double num1, num2, result;

    do {
        printf("\n=== Scientific Calculator ===\n");
        printf("1. Addition\n");
        printf("2. Subtraction\n");
        printf("3. Multiplication\n");
        printf("4. Division\n");
        printf("5. Sine\n");
    } while (1);
}

```



```

printf("6. Cosine\n");
printf("7. Square Root\n");
printf("8. Memory Store\n");
printf("9. Memory Recall\n");
printf("10. Show History\n");
printf("0. Exit\n");
printf("Enter your choice: ");

scanf("%d", &choice);

switch (choice) {
    case 1:
        printf("Enter two numbers: ");
        scanf("%lf %lf", &num1, &num2);
        result = num1 + num2;
        printf("Result: %.2f\n", result);
        add_to_history(result);
        break;
    case 2:
        printf("Enter two numbers: ");
        scanf("%lf %lf", &num1, &num2);
        result = num1 - num2;
        printf("Result: %.2f\n", result);
        add_to_history(result);
        break;
    case 3:
        printf("Enter two numbers: ");
        scanf("%lf %lf", &num1, &num2);
        result = num1 * num2;
        printf("Result: %.2f\n", result);
        add_to_history(result);
        break;
    case 4:
        printf("Enter two numbers: ");
        scanf("%lf %lf", &num1, &num2);
        if (num2 != 0) {
            result = num1 / num2;
            printf("Result: %.2f\n", result);
            add_to_history(result);
        } else {
            printf("Error: Division by zero!\n");
        }
        break;
    case 5:
        printf("Enter angle in radians: ");
        scanf("%lf", &num1);
        result = sin(num1);

```

```

        printf("Result: %.2f\n", result);
        add_to_history(result);
        break;
    case 6:
        printf("Enter angle in radians: ");
        scanf("%lf", &num1);
        result = cos(num1);
        printf("Result: %.2f\n", result);
        add_to_history(result);
        break;
    case 7:
        printf("Enter number: ");
        scanf("%lf", &num1);
        if (num1 >= 0) {
            result = sqrt(num1);
            printf("Result: %.2f\n", result);
            add_to_history(result);
        } else {
            printf("Error: Cannot calculate square root of negative number!\n");
        }
        break;
    case 8:
        printf("Enter number to store: ");
        scanf("%lf", &memory);
        printf("Stored %.2f in memory\n", memory);
        break;
    case 9:
        printf("Memory: %.2f\n", memory);
        break;
    case 10:
        show_history();
        break;
    case 0:
        printf("Thank you for using the calculator!\n");
        break;
    default:
        printf("Invalid choice! Please try again.\n");
    }
} while (choice != 0);

return 0;
}

```

Exercise 5 Solution Example:

```
#include <stdio.h>
#include <stdlib.h>
#include <string.h>

typedef struct {
    char *filename;
    FILE *file;
    char *buffer;
    int *data;
} ResourceManager;

int initialize_resources(ResourceManager *rm) {
    // Allocate filename
    rm->filename = malloc(100);
    if (!rm->filename) {
        goto cleanup;
    }
    strcpy(rm->filename, "data.txt");

    // Open file
    rm->file = fopen(rm->filename, "w");
    if (!rm->file) {
        printf("Error opening file\n");
        goto cleanup;
    }

    // Allocate buffer
    rm->buffer = malloc(1024);
    if (!rm->buffer) {
        printf("Error allocating buffer\n");
        goto cleanup;
    }

    // Allocate data
    rm->data = malloc(100 * sizeof(int));
    if (!rm->data) {
        printf("Error allocating data\n");
        goto cleanup;
    }

    printf("All resources initialized successfully\n");
    return 0; // Success

cleanup:
```

```

    // Cleanup in reverse order
    if (rm->data) {
        free(rm->data);
        rm->data = NULL;
    }
    if (rm->buffer) {
        free(rm->buffer);
        rm->buffer = NULL;
    }
    if (rm->file) {
        fclose(rm->file);
        rm->file = NULL;
    }
    if (rm->filename) {
        free(rm->filename);
        rm->filename = NULL;
    }
    return -1; // Error
}

void cleanup_resources(ResourceManager *rm) {
    if (rm->data) {
        free(rm->data);
        rm->data = NULL;
    }
    if (rm->buffer) {
        free(rm->buffer);
        rm->buffer = NULL;
    }
    if (rm->file) {
        fclose(rm->file);
        rm->file = NULL;
    }
    if (rm->filename) {
        free(rm->filename);
        rm->filename = NULL;
    }
    printf("All resources cleaned up\n");
}

int main() {
    ResourceManager rm = {0};

    if (initialize_resources(&rm) == 0) {
        printf("Resource management successful\n");
        // Use resources here
        cleanup_resources(&rm);
    }
}

```

```

    } else {
        printf("Resource management failed\n");
        return 1;
    }

    return 0;
}

```

Common Pitfalls to Avoid

1. **Infinite Loops:** Always ensure loop termination conditions
2. **Off-by-One Errors:** Carefully check loop bounds
3. **goto Overuse:** Use goto only for structured error handling
4. **Deep Nesting:** Use early returns to reduce nesting levels
5. **Missing Break Statements:** Always include break in switch cases
6. **Floating-Point Loop Variables:** Avoid using floats as loop counters
7. **Resource Leaks:** Always clean up allocated resources
8. **Uninitialized Variables:** Initialize all variables before use

Compilation Tips

```

# Basic compilation
gcc program.c -o program

# With math library for scientific functions
gcc program.c -lm -o program

# With warnings enabled
gcc -Wall -Wextra program.c -o program

# With debugging information
gcc -g -Wall program.c -o program

# With optimization
gcc -O2 program.c -o program

```

Complete these exercises to solidify your understanding of Module 3 concepts. Each exercise builds upon the previous ones, gradually increasing in complexity.

Functions

Module 4: Functions and Modular Programming Exercises

Exercise 1: Basic Function Implementation

Write a program that implements several mathematical functions: - A function to calculate the area of different shapes (circle, rectangle, triangle) - A function to convert temperatures between Celsius and Fahrenheit - A function to calculate compound interest - A function to determine if a year is a leap year

Requirements: - Use proper function declarations and definitions - Include appropriate header files - Handle edge cases and invalid inputs - Provide clear function documentation - Use const parameters where appropriate

Exercise 2: Recursive Functions

Create a program that implements various recursive algorithms: - Factorial calculation with error handling - Fibonacci sequence with memoization - Binary search on a sorted array - Tower of Hanoi solution - Greatest Common Divisor (GCD) using Euclidean algorithm

Requirements: - Implement proper base cases to prevent infinite recursion - Handle negative inputs appropriately - Include performance timing for comparison with iterative versions - Document recursion depth limitations - Provide clear output showing the recursive process

Exercise 3: Variable Argument Functions

Develop a program that demonstrates variable argument functions: - A function that calculates the average of a variable number of integers - A function that finds the maximum value among variable arguments - A formatted print function similar to printf - A function that concatenates a variable number of strings

Requirements: - Use stdarg.h for variable argument handling - Implement proper argument counting and type checking - Handle memory allocation for string concatenation - Include error handling for invalid arguments - Provide examples with different numbers of arguments

Exercise 4: Function Pointers and Callbacks

Create a program that uses function pointers for various purposes: - A calculator that uses function pointers for operations - A sorting function that accepts a comparison function pointer - A filter function that applies a predicate to an array - A state machine implementation using function pointers

Requirements: - Implement function pointer arrays for dispatch tables - Create callback mechanisms for event handling - Demonstrate function pointer casting - Include error handling for null function pointers - Provide clear examples of different use cases

Exercise 5: Modular Programming with Multiple Files

Design and implement a complete module system: - Create a statistics library with header and source files - Implement a string processing module - Create a configuration management module - Develop a logging system module - Write a main program that uses all modules

Requirements: - Use proper include guards in header files - Separate interface from implementation - Demonstrate static and external linkage - Include Makefile or build script for compilation - Document module dependencies and usage

Exercise 6: Inline Functions and Macro Functions

Write a program that compares inline functions with macro functions: - Implement mathematical operations as both inline functions and macros - Create performance benchmarks for both approaches - Demonstrate side effects with macros - Show type safety differences between approaches - Implement conditional compilation for different optimization levels

Requirements: - Use appropriate compiler flags for inline function support - Include timing measurements for performance comparison - Demonstrate macro pitfalls with complex expressions - Show how inline functions provide type checking - Document when to use each approach

Exercise 7: Standard Library Function Practice

Create a comprehensive program that uses various standard library functions: - String manipulation with functions from `string.h` - Mathematical calculations with functions from `math.h` - Time and date processing with functions from `time.h` - Memory management with functions from `stdlib.h` - Character classification with functions from `ctype.h`

Requirements: - Demonstrate proper error handling for library functions - Include examples of common pitfalls and how to avoid them - Show memory management best practices - Implement proper buffer size checking - Provide clear output showing function results

Exercise 8: Advanced Modular Programming

Design a complete application using advanced modular programming techniques: - Implement an opaque pointer design pattern - Create a plugin system using dynamic loading (if supported) - Develop a configuration system with multiple backends - Implement a thread-safe module with proper synchronization - Create a unit testing framework for your modules

Requirements: - Use advanced features like function pointers for polymorphism - Implement proper error handling and resource cleanup - Include documentation for module interfaces - Demonstrate proper separation of concerns - Provide examples of module extensibility

Solutions and Tips

Exercise 1 Solution Example:

```
#include <stdio.h>
#include <math.h>

#define PI 3.14159265359

// Function to calculate area of circle
double circle_area(double radius) {
    if (radius < 0) {
        printf("Error: Radius cannot be negative\n");
        return -1;
    }
    return PI * radius * radius;
}

// Function to calculate area of rectangle
double rectangle_area(double length, double width) {
    if (length < 0 || width < 0) {
        printf("Error: Dimensions cannot be negative\n");
        return -1;
    }
    return length * width;
}

// Function to calculate area of triangle
double triangle_area(double base, double height) {
    if (base < 0 || height < 0) {
        printf("Error: Dimensions cannot be negative\n");
        return -1;
    }
    return 0.5 * base * height;
}
```

```

// Function to convert Celsius to Fahrenheit
double celsius_to_fahrenheit(double celsius) {
    return (celsius * 9.0 / 5.0) + 32.0;
}

// Function to convert Fahrenheit to Celsius
double fahrenheit_to_celsius(double fahrenheit) {
    return (fahrenheit - 32.0) * 5.0 / 9.0;
}

// Function to calculate compound interest
double compound_interest(double principal, double rate, int time) {
    if (principal < 0 || rate < 0 || time < 0) {
        printf("Error: Invalid parameters\n");
        return -1;
    }
    return principal * pow(1 + rate / 100.0, time);
}

// Function to check if a year is a leap year
int is_leap_year(int year) {
    if (year % 400 == 0) return 1;
    if (year % 100 == 0) return 0;
    if (year % 4 == 0) return 1;
    return 0;
}

int main() {
    // Test area calculations
    printf("Circle area (radius 5): %.2f\n", circle_area(5));
    printf("Rectangle area (5x3): %.2f\n", rectangle_area(5, 3));
    printf("Triangle area (base 4, height 6): %.2f\n", triangle_area(4, 6));

    // Test temperature conversion
    printf("25°C = %.1f°F\n", celsius_to_fahrenheit(25));
    printf("77°F = %.1f°C\n", fahrenheit_to_celsius(77));

    // Test compound interest
    printf("Compound interest ($1000, 5%%, 10 years): $%.2f\n",
        compound_interest(1000, 5, 10));

    // Test leap year
    printf("2024 is %s a leap year\n", is_leap_year(2024) ? "" : "not");
    printf("1900 is %s a leap year\n", is_leap_year(1900) ? "" : "not");

    return 0;
}

```

```
}
```

Exercise 2 Solution Example:

```
#include <stdio.h>
#include <time.h>

// Simple recursive factorial
long long factorial(int n) {
    if (n < 0) {
        printf("Error: Factorial of negative number\n");
        return -1;
    }
    if (n == 0 || n == 1) return 1;
    return n * factorial(n - 1);
}

// Fibonacci with memoization
long long fib_memo[100] = {0};

long long fibonacci_memo(int n) {
    if (n < 0) {
        printf("Error: Invalid Fibonacci index\n");
        return -1;
    }
    if (n <= 1) return n;

    if (fib_memo[n] != 0) return fib_memo[n];

    fib_memo[n] = fibonacci_memo(n - 1) + fibonacci_memo(n - 2);
    return fib_memo[n];
}

// Binary search recursive implementation
int binary_search_recursive(int arr[], int left, int right, int target) {
    if (left > right) return -1; // Not found

    int mid = left + (right - left) / 2;

    if (arr[mid] == target) return mid;
    if (arr[mid] > target) {
        return binary_search_recursive(arr, left, mid - 1, target);
    } else {
        return binary_search_recursive(arr, mid + 1, right, target);
    }
}
```

```

}

// Tower of Hanoi
void tower_of_hanoi(int n, char from, char to, char aux) {
    if (n == 1) {
        printf("Move disk 1 from %c to %c\n", from, to);
        return;
    }

    tower_of_hanoi(n - 1, from, aux, to);
    printf("Move disk %d from %c to %c\n", n, from, to);
    tower_of_hanoi(n - 1, aux, to, from);
}

// GCD using Euclidean algorithm
int gcd_recursive(int a, int b) {
    if (b == 0) return a;
    return gcd_recursive(b, a % b);
}

int main() {
    // Test factorial
    printf("Factorial of 5: %lld\n", factorial(5));

    // Test Fibonacci with memoization
    for (int i = 0; i < 10; i++) {
        printf("Fibonacci(%d) = %lld\n", i, fibonacci_memo(i));
    }

    // Test binary search
    int arr[] = {1, 3, 5, 7, 9, 11, 13, 15};
    int target = 7;
    int index = binary_search_recursive(arr, 0, 7, target);
    if (index != -1) {
        printf("Found %d at index %d\n", target, index);
    } else {
        printf("%d not found\n", target);
    }

    // Test Tower of Hanoi
    printf("\nTower of Hanoi (3 disks):\n");
    tower_of_hanoi(3, 'A', 'C', 'B');

    // Test GCD
    printf("\nGCD of 48 and 18: %d\n", gcd_recursive(48, 18));

    return 0;
}

```

```
}
```

Exercise 3 Solution Example:

```
#include <stdio.h>
#include <stdarg.h>
#include <string.h>
#include <stdlib.h>

// Function to calculate average of variable integers
double average_ints(int count, ...) {
    if (count <= 0) {
        printf("Error: Invalid count\n");
        return 0;
    }

    va_list args;
    va_start(args, count);

    int sum = 0;
    for (int i = 0; i < count; i++) {
        sum += va_arg(args, int);
    }

    va_end(args);
    return (double)sum / count;
}

// Function to find maximum among variable arguments
int max_ints(int count, ...) {
    if (count <= 0) {
        printf("Error: Invalid count\n");
        return 0;
    }

    va_list args;
    va_start(args, count);

    int max = va_arg(args, int);
    for (int i = 1; i < count; i++) {
        int current = va_arg(args, int);
        if (current > max) max = current;
    }

    va_end(args);
}
```

```

    return max;
}

// Function to concatenate variable strings
char* concat_strings(int count, ...) {
    if (count <= 0) {
        printf("Error: Invalid count\n");
        return NULL;
    }

    // First pass: calculate total length
    va_list args;
    va_start(args, count);

    int total_length = 0;
    for (int i = 0; i < count; i++) {
        const char* str = va_arg(args, const char*);
        if (str) total_length += strlen(str);
    }

    va_end(args);

    // Allocate memory for result
    char* result = (char*)malloc(total_length + 1);
    if (!result) {
        printf("Error: Memory allocation failed\n");
        return NULL;
    }

    // Second pass: concatenate strings
    result[0] = '\0';
    va_start(args, count);

    for (int i = 0; i < count; i++) {
        const char* str = va_arg(args, const char*);
        if (str) strcat(result, str);
    }

    va_end(args);
    return result;
}

int main() {
    // Test average function
    printf("Average of 1, 2, 3, 4, 5: %.2f\n",
        average_ints(5, 1, 2, 3, 4, 5));
}

```

```

// Test maximum function
printf("Maximum of 10, 25, 7, 33, 15: %d\n",
      max_ints(5, 10, 25, 7, 33, 15));

// Test string concatenation
char* result = concat_strings(4, "Hello", " ", "World", "!");
if (result) {
    printf("Concatenated string: %s\n", result);
    free(result);
}

return 0;
}

```

Exercise 4 Solution Example:

```

#include <stdio.h>
#include <stdlib.h>

// Calculator operations
int add(int a, int b) { return a + b; }
int subtract(int a, int b) { return a - b; }
int multiply(int a, int b) { return a * b; }
int divide(int a, int b) { return b != 0 ? a / b : 0; }

// Calculator using function pointers
int calculate(int a, int b, int (*operation)(int, int)) {
    if (operation == NULL) {
        printf("Error: Invalid operation\n");
        return 0;
    }
    return operation(a, b);
}

// Comparison functions for sorting
int compare_ascending(const void *a, const void *b) {
    int int_a = *(const int*)a;
    int int_b = *(const int*)b;
    return (int_a > int_b) - (int_a < int_b);
}

int compare_descending(const void *a, const void *b) {
    int int_a = *(const int*)a;
    int int_b = *(const int*)b;
    return (int_b > int_a) - (int_b < int_a);
}

```

```

}

// Generic sorting function with comparison function pointer
void sort_array(int *arr, int size, int (*compare)(const void*, const void*)) {
    if (arr == NULL || compare == NULL) {
        printf("Error: Invalid parameters\n");
        return;
    }
    qsort(arr, size, sizeof(int), compare);
}

// Filter function with predicate
int* filter_array(int *arr, int size, int (*predicate)(int), int *result_size) {
    if (arr == NULL || predicate == NULL || result_size == NULL) {
        printf("Error: Invalid parameters\n");
        return NULL;
    }

    int *result = (int*)malloc(size * sizeof(int));
    if (!result) {
        printf("Error: Memory allocation failed\n");
        return NULL;
    }

    *result_size = 0;
    for (int i = 0; i < size; i++) {
        if (predicate(arr[i])) {
            result[*result_size] = arr[i];
            (*result_size)++;
        }
    }

    return result;
}

// Predicate functions
int is_even(int n) { return n % 2 == 0; }
int is_positive(int n) { return n > 0; }
int is_greater_than_10(int n) { return n > 10; }

int main() {
    // Test calculator with function pointers
    printf("Calculator using function pointers:\n");
    printf("10 + 5 = %d\n", calculate(10, 5, add));
    printf("10 - 5 = %d\n", calculate(10, 5, subtract));
    printf("10 * 5 = %d\n", calculate(10, 5, multiply));
    printf("10 / 5 = %d\n", calculate(10, 5, divide));
}

```



```

// Test sorting with function pointers
int arr[] = {5, 2, 8, 1, 9, 3};
int size = sizeof(arr) / sizeof(arr[0]);

printf("\nOriginal array: ");
for (int i = 0; i < size; i++) {
    printf("%d ", arr[i]);
}

sort_array(arr, size, compare_ascending);
printf("\nSorted ascending: ");
for (int i = 0; i < size; i++) {
    printf("%d ", arr[i]);
}

sort_array(arr, size, compare_descending);
printf("\nSorted descending: ");
for (int i = 0; i < size; i++) {
    printf("%d ", arr[i]);
}

// Test filter function
int numbers[] = {1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12};
int num_size = sizeof(numbers) / sizeof(numbers[0]);
int filtered_size;

int *evens = filter_array(numbers, num_size, is_even, &filtered_size);
if (evens) {
    printf("\n\nEven numbers: ");
    for (int i = 0; i < filtered_size; i++) {
        printf("%d ", evens[i]);
    }
    free(evens);
}

int *greater_than_10 = filter_array(numbers, num_size, is_greater_than_10, &filtered_size);
if (greater_than_10) {
    printf("\n\nNumbers > 10: ");
    for (int i = 0; i < filtered_size; i++) {
        printf("%d ", greater_than_10[i]);
    }
    free(greater_than_10);
}

printf("\n");
return 0;

```

```
}
```

Common Pitfalls to Avoid

1. **Infinite Recursion:** Always ensure proper base cases
2. **Memory Leaks:** Free dynamically allocated memory
3. **Null Pointer Dereference:** Always check function pointer validity
4. **Buffer Overflows:** Ensure adequate buffer sizes for string operations
5. **Variable Argument Misuse:** Properly count and type-check variable arguments
6. **Linkage Issues:** Understand static vs. external linkage
7. **Include Guard Problems:** Use proper include guards to prevent multiple inclusions
8. **Macro Side Effects:** Be careful with complex expressions in macros
9. **Integer Overflow:** Check for overflow in mathematical operations
10. **Resource Cleanup:** Always clean up resources in error paths

Compilation Tips

```
# Basic compilation with multiple source files
gcc main.c math_utils.c string_utils.c -o program

# With math library for mathematical functions
gcc main.c math_utils.c -lm -o program

# With warnings enabled
gcc -Wall -Wextra main.c math_utils.c -o program

# With debugging information
gcc -g -Wall main.c math_utils.c -o program

# With optimization
gcc -O2 main.c math_utils.c -o program

# Creating object files separately
gcc -c math_utils.c -o math_utils.o
gcc -c string_utils.c -o string_utils.o
gcc main.c math_utils.o string_utils.o -o program
```

Complete these exercises to solidify your understanding of Module 4 concepts. Each exercise builds upon the previous ones, gradually increasing in complexity.

Arrays and Strings

Module 5: Arrays and Strings Exercises

Exercise 1: Array Manipulation Functions

Write a program that implements various array manipulation functions: - A function to find the maximum and minimum values in an array - A function to calculate the average of array elements - A function to reverse an array in place - A function to check if an array is sorted (ascending or descending) - A function to remove duplicates from an array

Requirements: - Use proper function declarations with array parameters - Handle edge cases (empty arrays, single element arrays) - Include appropriate header files - Provide clear function documentation - Test with different array sizes and data types

Exercise 2: Matrix Operations

Create a program that performs various matrix operations: - Matrix addition and subtraction - Matrix multiplication - Transpose of a matrix - Determinant calculation for 2x2 and 3x3 matrices - Check if a matrix is symmetric

Requirements: - Use 2D arrays for matrix representation - Implement proper bounds checking - Handle memory allocation for dynamic matrices (bonus) - Include error handling for incompatible matrix operations - Provide clear output formatting for matrices

Exercise 3: String Processing Library

Develop a comprehensive string processing library: - Function to count words in a string - Function to count vowels and consonants - Function to reverse words in a string (maintain word order) - Function to remove extra whitespace - Function to check if a string is a palindrome

Requirements: - Use proper string handling functions from string.h - Implement your own versions of common string functions - Handle edge cases (empty strings, NULL pointers) - Include memory safety considerations - Provide examples with various string inputs

Exercise 4: Dynamic Array Implementation

Create a dynamic array (vector-like) implementation:

- Function to initialize a dynamic array
- Function to add elements (with automatic resizing)
- Function to remove elements by index
- Function to search for elements
- Function to free the dynamic array

Requirements:

- Use malloc, realloc, and free for memory management
- Implement proper error handling for memory allocation
- Include capacity and size tracking
- Handle memory leaks properly
- Provide a complete API with clear documentation

Exercise 5: Multi-dimensional Array Challenges

Solve complex problems using multi-dimensional arrays:

- Implement a simple image processing function (2D array of pixels)
- Create a tic-tac-toe game board with win detection
- Implement a simple maze solver using backtracking
- Create a program to process student grades in a 2D array
- Implement Conway's Game of Life

Requirements:

- Use appropriate data structures for each problem
- Include proper initialization and cleanup
- Handle boundary conditions correctly
- Provide clear visualization of results where applicable
- Include performance considerations for large arrays

Exercise 6: String Search and Replace

Write a program that implements advanced string operations:

- Function to find all occurrences of a substring
- Function to replace all occurrences of a substring
- Function to tokenize a string based on multiple delimiters
- Function to implement a simple pattern matching algorithm
- Function to perform case-insensitive string operations

Requirements:

- Handle overlapping matches correctly
- Manage memory allocation for result strings
- Include error checking for invalid inputs
- Optimize for performance where possible
- Provide comprehensive test cases

Exercise 7: Command Line Argument Processing

Create a program that processes command line arguments effectively:

- Parse command line options and flags
- Handle positional arguments
- Implement a help system with usage information
- Validate argument types and ranges
- Provide meaningful error messages for invalid input

Requirements:

- Use argc and argv properly
- Implement both short (-h) and long (--help) options
- Handle combined flags (-abc equivalent to -a -b -c)
- Include default values for optional arguments
- Provide examples of usage

Exercise 8: Advanced Array and String Integration

Design a complete application that integrates arrays and strings: - Implement a simple database system using arrays of structures - Create a text analysis tool that processes multiple files - Develop a word frequency counter with sorting capabilities - Implement a simple encryption/decryption system - Create a program that generates reports from structured data

Requirements: - Use modular design with separate functions for each major task - Include proper error handling and recovery - Handle file I/O for data persistence - Implement sorting and searching algorithms - Provide clear user interface and output formatting

Solutions and Tips

Exercise 1 Solution Example:

```
#include <stdio.h>
#include <limits.h>

// Function to find maximum value in array
int find_max(int arr[], int size) {
    if (size <= 0) return INT_MIN;

    int max = arr[0];
    for (int i = 1; i < size; i++) {
        if (arr[i] > max) {
            max = arr[i];
        }
    }
    return max;
}

// Function to find minimum value in array
int find_min(int arr[], int size) {
    if (size <= 0) return INT_MAX;

    int min = arr[0];
    for (int i = 1; i < size; i++) {
        if (arr[i] < min) {
            min = arr[i];
        }
    }
    return min;
}

// Function to calculate average
```

```

double calculate_average(int arr[], int size) {
    if (size <= 0) return 0.0;

    long long sum = 0;
    for (int i = 0; i < size; i++) {
        sum += arr[i];
    }
    return (double)sum / size;
}

int main() {
    int numbers[] = {5, 2, 8, 1, 9, 3};
    int size = sizeof(numbers) / sizeof(numbers[0]);

    printf("Array: ");
    for (int i = 0; i < size; i++) {
        printf("%d ", numbers[i]);
    }
    printf("\n");

    printf("Maximum: %d\n", find_max(numbers, size));
    printf("Minimum: %d\n", find_min(numbers, size));
    printf("Average: %.2f\n", calculate_average(numbers, size));

    return 0;
}

```

Exercise 3 Solution Example:

```

#include <stdio.h>
#include <string.h>
#include <ctype.h>

// Function to count words in a string
int count_words(const char *str) {
    if (str == NULL) return 0;

    int count = 0;
    int in_word = 0;

    for (int i = 0; str[i] != '\0'; i++) {
        if (isspace(str[i])) {
            in_word = 0;
        } else if (!in_word) {
            in_word = 1;
        }
    }
    return count;
}

```

```

        count++;
    }
}
return count;
}

// Function to check if string is palindrome
int is_palindrome(const char *str) {
    if (str == NULL) return 0;

    int len = strlen(str);
    for (int i = 0; i < len / 2; i++) {
        if (tolower(str[i]) != tolower(str[len - 1 - i])) {
            return 0;
        }
    }
    return 1;
}

int main() {
    char text1[] = "Hello world, this is a test string";
    char text2[] = "A man a plan a canal Panama";

    printf("Text 1: \"%s\\n\", text1);
    printf("Word count: %d\\n", count_words(text1));

    printf("\\nText 2: \"%s\\n\", text2);
    printf("Is palindrome: %s\\n", is_palindrome(text2) ? "Yes" : "No");

    return 0;
}

```

Common Pitfalls to Avoid:

1. **Array bounds:** Always check array indices before accessing elements
2. **String termination:** Ensure strings are properly null-terminated
3. **Memory allocation:** Check return values from malloc/realloc
4. **Pointer validation:** Always check for NULL pointers
5. **Buffer overflow:** Limit string input sizes to prevent overflow

Performance Tips:

1. **Pass by reference:** Use pointers for large arrays to avoid copying
2. **Avoid repeated strlen:** Store string length in a variable when used multiple times
3. **Use const:** Mark string parameters as const when not modifying them
4. **Pre-allocate:** When size is known, pre-allocate arrays to avoid repeated resizing

5. **Cache results:** For expensive operations, cache results when possible

Complete these exercises to solidify your understanding of arrays and strings in C. Each exercise builds upon the previous ones, gradually increasing in complexity.

Pointers

Module 6: Pointers Exercises

Exercise 1: Basic Pointer Operations

Write a program that demonstrates fundamental pointer operations: - Declare and initialize pointers to different data types - Perform pointer arithmetic (increment, decrement, addition, subtraction) - Dereference pointers to modify values - Swap two variables using pointers - Calculate the distance between two pointers

Requirements: - Use proper pointer declarations and initialization - Demonstrate pointer arithmetic with different data types - Include bounds checking to prevent invalid memory access - Provide clear output showing address and value changes - Handle NULL pointer checks appropriately

Exercise 2: Pointers and Arrays

Create a program that explores the relationship between pointers and arrays: - Access array elements using pointer notation - Pass arrays to functions using pointer parameters - Implement array traversal using pointers - Demonstrate pointer arithmetic as array indexing - Compare performance of array indexing vs pointer arithmetic

Requirements: - Implement functions that accept array parameters as pointers - Show equivalent operations using array and pointer notation - Include examples with multi-dimensional arrays - Handle array bounds properly - Provide timing comparisons where relevant

Exercise 3: Dynamic Memory Allocation

Develop a program that demonstrates dynamic memory management: - Allocate memory for single variables using malloc - Allocate memory for arrays using malloc and calloc - Resize allocated memory using realloc - Free allocated memory properly - Handle memory allocation failures gracefully

Requirements: - Check return values from all memory allocation functions - Initialize allocated memory appropriately - Avoid memory leaks by freeing all allocated memory - Handle reallocation failures properly - Include memory usage tracking (bonus)

Exercise 4: Pointer to Pointers

Write a program that works with pointers to pointers: - Declare and initialize pointers to pointers - Modify values through multiple levels of indirection - Implement a function that allocates memory and returns it through a pointer parameter - Create a dynamic 2D array using pointers to pointers - Demonstrate proper cleanup of multi-level allocations

Requirements: - Handle multiple levels of pointer dereferencing correctly - Include proper error checking at each level - Free memory in the correct order to prevent leaks - Document the relationship between levels clearly - Provide examples of practical use cases

Exercise 5: Function Pointers

Create a program that demonstrates function pointer usage: - Declare and initialize function pointers - Call functions through function pointers - Implement a simple calculator using function pointers - Create an array of function pointers for dispatch - Pass function pointers as parameters to other functions

Requirements: - Use proper function pointer syntax and declarations - Include error checking for NULL function pointers - Demonstrate different ways to assign functions to pointers - Implement callback mechanisms using function pointers - Provide clear examples of practical applications

Exercise 6: Advanced Pointer Techniques

Write a program that implements advanced pointer concepts: - Implement a generic swap function using void pointers - Create a flexible data structure using void pointers - Demonstrate pointer casting between different types - Implement a simple memory pool allocator - Create a basic garbage collection simulation (bonus)

Requirements: - Handle type safety when working with void pointers - Include proper casting and alignment considerations - Implement error checking for invalid operations - Document memory management strategies clearly - Provide examples of real-world applications

Exercise 7: Pointer Debugging and Safety

Create a program that demonstrates safe pointer practices: - Implement a memory debugging system with allocation tracking - Create functions to detect common pointer errors - Demonstrate proper error handling for pointer operations - Implement a simple smart pointer-like mechanism - Show techniques for preventing dangling pointers

Requirements: - Include comprehensive error checking - Handle edge cases and invalid operations gracefully - Provide meaningful error messages - Implement cleanup procedures for all resources - Document safety practices and best approaches

Exercise 8: Comprehensive Pointer Application

Design a complete application that integrates all pointer concepts: - Implement a dynamic data structure (linked list, tree, etc.) - Create a memory management system with custom allocators - Develop a plugin system using function pointers - Implement a callback-based event system - Include comprehensive error handling and resource management

Requirements: - Use modular design with separate functions for each major component - Include proper documentation for all pointer-related functions - Handle all memory allocation and deallocation properly - Implement robust error handling throughout - Provide clear examples and test cases

Solutions and Tips

Exercise 1 Solution Example:

```
#include <stdio.h>

// Function to swap two integers using pointers
void swap_int(int *a, int *b) {
    if (a == NULL || b == NULL) {
        printf("Error: NULL pointer passed to swap_int\n");
        return;
    }

    int temp = *a;
    *a = *b;
    *b = temp;
}

int main() {
    int x = 10, y = 20;
    int *ptr_x = &x;
    int *ptr_y = &y;

    printf("Before swap: x = %d, y = %d\n", x, y);
    printf("Addresses: &x = %p, &y = %p\n", (void*)&x, (void*)&y);
    printf("Pointer values: ptr_x = %p, ptr_y = %p\n", (void*)ptr_x, (void*)ptr_y);

    swap_int(ptr_x, ptr_y);

    printf("After swap: x = %d, y = %d\n", x, y);

    // Pointer arithmetic demonstration
    int arr[] = {1, 2, 3, 4, 5};
```

```

    int *ptr = arr;

    printf("\nArray elements using pointer arithmetic:\n");
    for (int i = 0; i < 5; i++) {
        printf("Element %d: value = %d, address = %p\n", i, *ptr, (void*)ptr);
        ptr++;
    }

    return 0;
}

```

Exercise 3 Solution Example:

```

#include <stdio.h>
#include <stdlib.h>

// Function to safely allocate memory
int* safe_malloc_int(int count) {
    if (count <= 0) {
        printf("Error: Invalid count %d\n", count);
        return NULL;
    }

    int *ptr = malloc(count * sizeof(int));
    if (ptr == NULL) {
        printf("Error: Memory allocation failed\n");
        return NULL;
    }

    // Initialize allocated memory
    for (int i = 0; i < count; i++) {
        ptr[i] = 0;
    }

    return ptr;
}

int main() {
    // Allocate memory for 5 integers
    int *dynamic_array = safe_malloc_int(5);
    if (dynamic_array == NULL) {
        return 1;
    }

    // Use the allocated memory

```

```

for (int i = 0; i < 5; i++) {
    dynamic_array[i] = (i + 1) * 10;
}

printf("Dynamic array contents: ");
for (int i = 0; i < 5; i++) {
    printf("%d ", dynamic_array[i]);
}
printf("\n");

// Resize the array
int *resized_array = realloc(dynamic_array, 8 * sizeof(int));
if (resized_array == NULL) {
    printf("Error: Memory reallocation failed\n");
    free(dynamic_array);
    return 1;
}

dynamic_array = resized_array;

// Initialize new elements
for (int i = 5; i < 8; i++) {
    dynamic_array[i] = (i + 1) * 10;
}

printf("Resized array contents: ");
for (int i = 0; i < 8; i++) {
    printf("%d ", dynamic_array[i]);
}
printf("\n");

// Free allocated memory
free(dynamic_array);
dynamic_array = NULL; // Prevent dangling pointer

return 0;
}

```

Common Pitfalls to Avoid:

1. **Uninitialized pointers:** Always initialize pointers to NULL or valid addresses
2. **Dangling pointers:** Set pointers to NULL after freeing memory
3. **Memory leaks:** Free all allocated memory exactly once
4. **Buffer overruns:** Check bounds before pointer arithmetic
5. **Type mismatches:** Use proper casting when changing pointer types

Best Practices:

1. **NULL checking:** Always check pointers before dereferencing
2. **Consistent naming:** Use clear naming conventions for pointer variables
3. **Documentation:** Comment pointer usage and ownership clearly
4. **Error handling:** Handle allocation failures gracefully
5. **Resource management:** Follow RAII-like patterns for cleanup

Complete these exercises to solidify your understanding of pointers in C. Each exercise builds upon the previous ones, gradually increasing in complexity.

Structures

Module 7: Structures, Unions, and Enums

Exercises

Exercise 1: Basic Structure Implementation

Write a program that defines and uses various structures: - Define a structure to represent a date (day, month, year) - Define a structure to represent a person (name, age, date of birth) - Define a structure to represent a book (title, author, ISBN, price) - Create functions to initialize and print these structures - Demonstrate structure assignment and comparison

Requirements: - Use proper structure definitions with appropriate data types - Implement initialization functions for each structure - Create print functions that display structure contents - Handle string copying safely to prevent buffer overflows - Include examples of nested structures

Exercise 2: Array of Structures

Create a program that manages collections using arrays of structures: - Implement a student database using an array of student structures - Create functions to add, search, and remove students - Implement sorting functions for the student array - Add functionality to calculate class averages and statistics - Include file I/O to save and load student data

Requirements: - Use dynamic memory allocation for flexible array sizing - Implement proper error handling for array bounds - Include search functions with different criteria - Handle memory management properly to prevent leaks - Provide clear user interface for database operations

Exercise 3: Linked List Implementation

Develop a program that implements a complete linked list: - Define a node structure for the linked list - Implement functions to insert, delete, and search nodes - Create functions to traverse and display the list - Implement sorting and merging of linked lists - Include memory management for all list operations

Requirements: - Handle edge cases (empty list, single node, etc.) - Implement proper memory allocation and deallocation - Include error checking for all operations - Provide both iterative and recursive implementations (bonus) - Document time and space complexity of operations

Exercise 4: Unions and Bit Fields

Write a program that demonstrates advanced data types: - Define and use unions for memory-efficient data storage - Implement structures with bit fields for flags and settings - Create a program that uses unions for type-generic operations - Demonstrate memory layout differences between structures and unions - Implement a simple variant type using unions

Requirements: - Include proper documentation for union usage - Handle type safety when working with unions - Demonstrate bit field packing and alignment - Include examples of practical union applications - Provide clear explanations of memory usage differences

Exercise 5: Enumerations and Type Definitions

Create a program that uses enumerations effectively: - Define enumerations for different categories (status, types, etc.) - Use typedef to create aliases for complex structure types - Implement state machines using enumerations - Create lookup tables using enumerations as indices - Demonstrate enum-to-string conversion functions

Requirements: - Use meaningful names for enumeration constants - Include bounds checking for enumeration values - Implement proper error handling for invalid enum values - Provide examples of enum-based switch statements - Document the advantages of using enumerations

Exercise 6: Advanced Structure Techniques

Write a program that implements advanced structure concepts: - Create structures with flexible array members (C99) - Implement opaque pointer design pattern - Demonstrate structure packing and alignment - Create self-referential structures for trees and graphs - Implement structure serialization and deserialization

Requirements: - Handle flexible array members correctly with dynamic allocation - Include proper memory management for opaque structures - Demonstrate compiler-specific packing directives - Implement proper traversal algorithms for tree structures - Include error checking for serialization operations

Exercise 7: Structure Memory Management

Create a program that demonstrates safe structure handling: - Implement a memory pool for structure allocation - Create functions to track structure usage and leaks - Demonstrate proper deep copying of structures - Implement reference counting for shared structures - Include debugging features for structure validation

Requirements: - Include comprehensive error checking - Handle all memory allocation failures gracefully - Provide clear documentation for memory management functions - Implement proper cleanup procedures - Include performance measurements for different approaches

Exercise 8: Comprehensive Structure Application

Design a complete application that integrates all structure concepts: - Implement a complete data management system - Create a configuration system using nested structures - Develop a plugin architecture using function pointers in structures - Implement a serialization system for data persistence - Include comprehensive error handling and resource management

Requirements: - Use modular design with clear separation of concerns - Include proper documentation for all components - Handle all memory allocation and deallocation properly - Implement robust error handling throughout - Provide clear examples and test cases

Solutions and Tips

Exercise 1 Solution Example:

```
#include <stdio.h>
#include <string.h>

// Structure definitions
typedef struct {
    int day;
    int month;
    int year;
} Date;

typedef struct {
    char name[50];
    int age;
    Date birth_date;
} Person;

typedef struct {
    char title[100];
    char author[50];
    char isbn[20];
    double price;
} Book;

// Function to initialize a date
Date create_date(int day, int month, int year) {
    Date d = {day, month, year};
    return d;
}

// Function to print a date
```

```

void print_date(const Date *d) {
    if (d == NULL) return;
    printf("%02d/%02d/%04d", d->day, d->month, d->year);
}

// Function to initialize a person
Person create_person(const char *name, int age, Date birth_date) {
    Person p;
    strncpy(p.name, name, sizeof(p.name) - 1);
    p.name[sizeof(p.name) - 1] = '\0';
    p.age = age;
    p.birth_date = birth_date;
    return p;
}

// Function to print a person
void print_person(const Person *p) {
    if (p == NULL) return;
    printf("Name: %s\n", p->name);
    printf("Age: %d\n", p->age);
    printf("Birth Date: ");
    print_date(&p->birth_date);
    printf("\n");
}

int main() {
    Date birth = create_date(15, 6, 1990);
    Person person = create_person("John Doe", 33, birth);

    printf("Person Information:\n");
    print_person(&person);

    return 0;
}

```

Exercise 3 Solution Example:

```

#include <stdio.h>
#include <stdlib.h>
#include <string.h>

// Node structure for linked list
typedef struct Node {
    int data;
    struct Node *next;
}

```

```

} Node;

// Function to create a new node
Node* create_node(int data) {
    Node *new_node = malloc(sizeof(Node));
    if (new_node == NULL) {
        printf("Error: Memory allocation failed\n");
        return NULL;
    }

    new_node->data = data;
    new_node->next = NULL;
    return new_node;
}

// Function to insert at beginning of list
void insert_at_beginning(Node **head, int data) {
    if (head == NULL) return;

    Node *new_node = create_node(data);
    if (new_node == NULL) return;

    new_node->next = *head;
    *head = new_node;
}

// Function to print the list
void print_list(Node *head) {
    printf("List: ");
    Node *current = head;

    while (current != NULL) {
        printf("%d ", current->data);
        current = current->next;
    }
    printf("\n");
}

// Function to free the entire list
void free_list(Node *head) {
    Node *current = head;
    Node *next;

    while (current != NULL) {
        next = current->next;
        free(current);
        current = next;
    }
}

```

```

    }
}

int main() {
    Node *head = NULL;

    // Insert elements
    insert_at_beginning(&head, 10);
    insert_at_beginning(&head, 20);
    insert_at_beginning(&head, 30);

    printf("Linked list after insertions:\n");
    print_list(head);

    // Free memory
    free_list(head);

    return 0;
}

```

Common Pitfalls to Avoid:

1. **Shallow copying:** Always implement deep copy for structures with pointers
2. **Memory leaks:** Free all allocated memory in structures properly
3. **Buffer overflows:** Check string lengths when copying to fixed-size arrays
4. **Uninitialized structures:** Initialize all structure members properly
5. **Flexible array misuse:** Handle flexible array members with dynamic allocation

Best Practices:

1. **Consistent naming:** Use clear, consistent names for structures and members
2. **Documentation:** Comment structure purposes and member meanings
3. **Encapsulation:** Hide implementation details using opaque pointers when appropriate
4. **Error handling:** Check all memory allocations and pointer operations
5. **Memory management:** Follow consistent patterns for allocation and deallocation

Complete these exercises to solidify your understanding of structures, unions, and enums in C. Each exercise builds upon the previous ones, gradually increasing in complexity.

File I/O

Module 8: File I/O Exercises

Exercise 1: Basic File Operations

Write a program that performs fundamental file operations: - Create and write text to a file - Read and display the contents of a file - Append new content to an existing file - Copy one file to another - Check if a file exists and get its properties

Requirements: - Use proper file handling functions (fopen, fclose, etc.) - Include error checking for all file operations - Handle different file modes appropriately - Implement proper resource cleanup - Provide clear feedback for operation results

Exercise 2: Text File Processing

Create a program that processes text files in various ways: - Count lines, words, and characters in a text file - Search for specific patterns or words in a file - Replace occurrences of a word with another word - Extract specific lines or sections from a file - Format and restructure text data

Requirements: - Use appropriate text file I/O functions - Handle large files efficiently - Include proper buffer management - Implement case-sensitive and case-insensitive searches - Provide statistics and progress feedback

Exercise 3: Binary File Operations

Develop a program that works with binary files: - Create and read binary data files - Implement a simple database using binary files - Serialize and deserialize structures to/from binary files - Handle endianness when working with binary data - Include data validation and integrity checking

Requirements: - Use binary file modes appropriately - Handle structure padding and alignment issues - Implement proper error checking for binary operations - Include byte order conversion functions if needed - Provide examples of practical binary file applications

Exercise 4: File System Operations

Write a program that performs file system level operations: - List files and directories in a given path - Create, rename, and delete files and directories - Check file permissions and attributes - Implement a simple file browser - Handle file system errors gracefully

Requirements: - Use platform-appropriate file system functions - Include proper error handling for system calls - Handle cross-platform compatibility issues - Implement recursive directory operations - Provide clear user interface for file system operations

Exercise 5: Random Access Files

Create a program that demonstrates random access file operations: - Implement direct record access in a data file - Create an index for fast data retrieval - Update specific records without reading entire file - Handle file positioning with `fseek` and `ftell` - Implement a simple database with random access

Requirements: - Use `fseek`, `ftell`, and `rewind` appropriately - Handle fixed and variable record lengths - Include proper data serialization for records - Implement error checking for positioning operations - Provide efficient access patterns for large datasets

Exercise 6: Advanced File I/O Techniques

Write a program that implements advanced file handling concepts: - Use file buffering and flushing strategies - Implement non-blocking I/O operations - Create temporary files and handle automatic cleanup - Work with file locking for concurrent access - Implement file compression and decompression (bonus)

Requirements: - Include proper buffer management and flushing - Handle concurrent access scenarios safely - Implement proper cleanup for temporary files - Include error recovery mechanisms - Document performance considerations for different approaches

Exercise 7: Error Handling and Recovery

Create a program that demonstrates robust file error handling: - Implement comprehensive error checking for all file operations - Create recovery mechanisms for interrupted operations - Handle disk full and other system-level errors - Implement logging for file operations - Include graceful degradation for non-critical failures

Requirements: - Check return values from all file functions - Provide meaningful error messages for different failure scenarios - Implement retry mechanisms for transient errors - Include proper cleanup in error paths - Document error handling strategies clearly

Exercise 8: Comprehensive File Management System

Design a complete application that integrates all file I/O concepts: - Implement a full-featured file manager with GUI or CLI interface - Create backup and restore functionality - Implement file synchronization between directories - Include search and filtering capabilities - Provide comprehensive error handling and logging

Requirements: - Use modular design with clear separation of concerns - Include proper documentation for all components - Handle all file operations safely and efficiently - Implement robust error handling throughout - Provide clear examples and test cases

Solutions and Tips

Exercise 1 Solution Example:

```
#include <stdio.h>
#include <stdlib.h>
#include <string.h>

// Function to write text to a file
int write_text_file(const char *filename, const char *text) {
    FILE *file = fopen(filename, "w");
    if (file == NULL) {
        printf("Error: Could not open file %s for writing\n", filename);
        return -1;
    }

    if (fputs(text, file) == EOF) {
        printf("Error: Failed to write to file %s\n", filename);
        fclose(file);
        return -1;
    }

    fclose(file);
    return 0;
}

// Function to read and display file contents
int read_text_file(const char *filename) {
    FILE *file = fopen(filename, "r");
    if (file == NULL) {
        printf("Error: Could not open file %s for reading\n", filename);
        return -1;
    }

    char buffer[1024];
    while (fgets(buffer, sizeof(buffer), file) != NULL) {
        printf("%s", buffer);
    }

    if (ferror(file)) {
        printf("Error: Failed to read from file %s\n", filename);
    }
}
```

```

        fclose(file);
        return -1;
    }

    fclose(file);
    return 0;
}

// Function to append text to a file
int append_text_file(const char *filename, const char *text) {
    FILE *file = fopen(filename, "a");
    if (file == NULL) {
        printf("Error: Could not open file %s for appending\n", filename);
        return -1;
    }

    if (fputs(text, file) == EOF) {
        printf("Error: Failed to append to file %s\n", filename);
        fclose(file);
        return -1;
    }

    fclose(file);
    return 0;
}

int main() {
    const char *filename = "test.txt";
    const char *initial_text = "This is the initial content.\nLine 2\nLine 3\n";
    const char *append_text = "This line was appended.\n";

    // Write initial content
    if (write_text_file(filename, initial_text) != 0) {
        return 1;
    }
    printf("Wrote initial content to %s\n", filename);

    // Read and display content
    printf("\nFile contents:\n");
    if (read_text_file(filename) != 0) {
        return 1;
    }

    // Append additional content
    if (append_text_file(filename, append_text) != 0) {
        return 1;
    }
}

```

```

    printf("\nAppended content to %s\n", filename);

    // Read and display updated content
    printf("\nUpdated file contents:\n");
    if (read_text_file(filename) != 0) {
        return 1;
    }

    return 0;
}

```

Exercise 3 Solution Example:

```

#include <stdio.h>
#include <stdlib.h>
#include <string.h>

// Structure for binary data
typedef struct {
    int id;
    char name[32];
    double value;
} Record;

// Function to write records to binary file
int write_records(const char *filename, const Record *records, int count) {
    FILE *file = fopen(filename, "wb");
    if (file == NULL) {
        printf("Error: Could not open file %s for writing\n", filename);
        return -1;
    }

    if (fwrite(records, sizeof(Record), count, file) != (size_t)count) {
        printf("Error: Failed to write all records to file %s\n", filename);
        fclose(file);
        return -1;
    }

    fclose(file);
    return 0;
}

// Function to read records from binary file
int read_records(const char *filename, Record *records, int max_count) {
    FILE *file = fopen(filename, "rb");

```

```

    if (file == NULL) {
        printf("Error: Could not open file %s for reading\n", filename);
        return -1;
    }

    size_t records_read = fread(records, sizeof(Record), max_count, file);
    if (ferror(file)) {
        printf("Error: Failed to read records from file %s\n", filename);
        fclose(file);
        return -1;
    }

    fclose(file);
    return (int)records_read;
}

int main() {
    // Create sample records
    Record records[] = {
        {1, "Record One", 100.50},
        {2, "Record Two", 200.75},
        {3, "Record Three", 300.25}
    };
    int record_count = sizeof(records) / sizeof(records[0]);

    const char *filename = "records.dat";

    // Write records to binary file
    if (write_records(filename, records, record_count) != 0) {
        return 1;
    }
    printf("Wrote %d records to %s\n", record_count, filename);

    // Read records from binary file
    Record read_records[10];
    int read_count = read_records(filename, read_records, 10);
    if (read_count < 0) {
        return 1;
    }

    printf("\nRead %d records from %s:\n", read_count, filename);
    for (int i = 0; i < read_count; i++) {
        printf("ID: %d, Name: %s, Value: %.2f\n",
            read_records[i].id, read_records[i].name, read_records[i].value);
    }

    return 0;
}

```

```
}
```

Common Pitfalls to Avoid:

1. **File handle leaks:** Always close files with `fclose`
2. **Buffer overflows:** Check buffer sizes when reading text
3. **Binary mode issues:** Use “rb”/“wb” for binary files on all platforms
4. **Error checking:** Always check return values from file functions
5. **Path separators:** Handle cross-platform path differences properly

Best Practices:

1. **Resource management:** Use RAII-like patterns for file handles
2. **Error handling:** Implement comprehensive error checking
3. **Buffer management:** Use appropriate buffer sizes for efficiency
4. **File modes:** Choose correct file modes for intended operations
5. **Data integrity:** Include validation and checksums for critical data

Complete these exercises to solidify your understanding of file I/O in C. Each exercise builds upon the previous ones, gradually increasing in complexity.

Modern C

Module 9: Modern C Features Exercises

Exercise 1: C99 Features Implementation

Write a program that demonstrates key C99 features: - Use mixed declarations and code within blocks - Implement designated initializers for structures - Create functions using inline keyword for performance - Use compound literals for temporary objects - Demonstrate variable-length arrays (VLAs)

Requirements: - Use proper C99 compilation flags - Include examples of improved flexibility with mixed declarations - Demonstrate the benefits of designated initializers - Show performance comparisons with inline functions - Handle VLA size checking to prevent stack overflow

Exercise 2: C11 Standard Features

Create a program that utilizes C11 features: - Implement static assertions for compile-time checks - Use generic selections for type-generic programming - Create programs with anonymous structures and unions - Demonstrate alignment specifications - Use thread-local storage for multi-threading

Requirements: - Use appropriate C11 compilation flags - Include comprehensive static assertion examples - Implement generic functions using `_Generic` - Show practical uses of anonymous structures - Document thread-local storage behavior

Exercise 3: C17/C23 New Features

Develop a program that explores recent C standard features: - Use UTF-8 string literals for international text - Implement binary literals and digit separators - Create programs with enhanced enumeration features - Use `nullptr` constant (C23) for safer pointer handling - Demonstrate standard attributes like `[[deprecated]]`

Requirements: - Use appropriate C17/C23 compilation flags - Include examples of UTF-8 handling - Show benefits of binary literals and digit separators - Implement proper `nullptr` usage patterns - Document attribute-based code annotations

Exercise 4: Type Safety and Modern Practices

Write a program that emphasizes modern C safety practices: - Use `stdbool.h` for boolean types - Implement fixed-width integer types from `stdint.h` - Create safe string handling functions - Use `restrict` keyword for optimization hints - Demonstrate proper `const` correctness

Requirements: - Include comprehensive examples of type-safe programming - Show performance benefits of `restrict` keyword - Implement bounds-checked string functions - Document `const` correctness principles - Provide examples of portable integer types

Exercise 5: Advanced Preprocessor Techniques

Create a program that demonstrates modern preprocessor capabilities: - Implement complex macro functions with error checking - Use variadic macros for flexible logging - Create X-macros for data-driven code generation - Demonstrate conditional compilation with modern features - Implement include guards and `pragma once`

Requirements: - Include proper macro hygiene practices - Show benefits of variadic macros - Document X-macro patterns and use cases - Implement cross-platform conditional compilation - Provide examples of modern header organization

Exercise 6: Memory Management and Allocation

Write a program that implements modern memory management: - Use `aligned_alloc` for specific alignment requirements - Implement custom allocators with modern features - Create memory debugging tools with modern C - Demonstrate safe memory handling practices - Include performance profiling for allocation strategies

Requirements: - Include proper error checking for allocation functions - Show benefits of aligned memory access - Implement comprehensive memory tracking - Document safe memory handling patterns - Provide performance comparison data

Exercise 7: Error Handling and Diagnostics

Create a program that demonstrates modern error handling: - Implement detailed error reporting with `errno.h` - Use static assertions for design-by-contract - Create comprehensive logging systems - Demonstrate proper resource cleanup with modern techniques - Include debugging aids and diagnostic tools

Requirements: - Include comprehensive error classification - Show benefits of static assertions for validation - Implement structured logging with timestamps - Document proper cleanup patterns - Provide debugging support functions

Exercise 8: Comprehensive Modern C Application

Design a complete application that integrates all modern C features: - Implement a data processing pipeline with modern C - Create a configuration system using modern features - Develop a plugin architecture with dynamic loading - Include comprehensive testing and validation - Provide detailed documentation and examples

Requirements: - Use modular design with clear separation of concerns - Include proper documentation for all components - Handle all resource management properly - Implement robust error handling throughout - Provide clear examples and test cases

Solutions and Tips

Exercise 1 Solution Example:

```
#include <stdio.h>
#include <stdlib.h>
#include <stdbool.h>

// Structure with designated initializers
typedef struct {
    int id;
    char name[32];
    double value;
    bool active;
} Item;

// Inline function for performance
inline int max_int(int a, int b) {
    return (a > b) ? a : b;
}

int main() {
    // Mixed declarations and code (C99)
    for (int i = 0; i < 5; i++) {
        int square = i * i; // Declaration inside block
        printf("Square of %d is %d\n", i, square);
    }

    // Designated initializers (C99)
    Item item = {
        .id = 100,
        .name = "Sample Item",
        .value = 99.99,
        .active = true
    }
```

```

};

printf("Item: ID=%d, Name=%s, Value=%.2f, Active=%s\n",
      item.id, item.name, item.value,
      item.active ? "true" : "false");

// Compound literals (C99)
int *arr = (int[]){1, 2, 3, 4, 5};
printf("Array from compound literal: ");
for (int i = 0; i < 5; i++) {
    printf("%d ", arr[i]);
}
printf("\n");

// Variable-length array (C99)
int size;
printf("Enter array size: ");
scanf("%d", &size);

if (size > 0 && size <= 1000) {
    int vla[size]; // VLA declaration
    for (int i = 0; i < size; i++) {
        vla[i] = i * i;
    }

    printf("VLA contents: ");
    for (int i = 0; i < size && i < 10; i++) { // Limit output
        printf("%d ", vla[i]);
    }
    printf("\n");
}

return 0;
}

```

Exercise 2 Solution Example:

```

#include <stdio.h>
#include <stdlib.h>
#include <stdbool.h>
#include <stdalign.h>
#include <assert.h>

// Static assertion example (C11)
_Static_assert(sizeof(int) >= 4, "int must be at least 4 bytes");

```

```

// Generic selection example (C11)
#define MAX_GENERIC(x, y) _Generic((x), \
    int: max_int, \
    float: fmaxf, \
    double: fmax \
)(x, y)

inline int max_int(int a, int b) {
    return (a > b) ? a : b;
}

// Anonymous union example (C11)
typedef struct {
    enum { INT_TYPE, FLOAT_TYPE, STRING_TYPE } type;
    union {
        int int_value;
        float float_value;
        char string_value[32];
    }; // Anonymous union
} Variant;

int main() {
    // Generic selection usage
    int a = 10, b = 20;
    float x = 1.5f, y = 2.5f;

    printf("Max of %d and %d: %d\n", a, b, MAX_GENERIC(a, b));
    printf("Max of %.1f and %.1f: %.1f\n", x, y, MAX_GENERIC(x, y));

    // Anonymous union usage
    Variant v;

    v.type = INT_TYPE;
    v.int_value = 42;
    printf("Integer variant: %d\n", v.int_value);

    v.type = STRING_TYPE;
    snprintf(v.string_value, sizeof(v.string_value), "Hello, C11!");
    printf("String variant: %s\n", v.string_value);

    // Alignment specification (C11)
    alignas(16) char aligned_buffer[64];
    printf("Aligned buffer address: %p\n", (void*)aligned_buffer);
    printf("Alignment check: %s\n",
        ((uintptr_t)aligned_buffer % 16 == 0) ? "Aligned" : "Not aligned");
}

```

```
    return 0;  
}
```

Common Pitfalls to Avoid:

1. **Compiler compatibility:** Not all compilers support all modern features
2. **Feature detection:** Always check for feature availability before use
3. **Performance assumptions:** Measure actual performance benefits
4. **Portability issues:** Consider cross-platform compatibility
5. **Complexity creep:** Don't over-engineer with modern features

Best Practices:

1. **Feature detection:** Use preprocessor checks for feature availability
2. **Gradual adoption:** Introduce modern features incrementally
3. **Documentation:** Comment modern feature usage clearly
4. **Testing:** Verify behavior across different compilers
5. **Fallbacks:** Provide alternatives for older compiler versions

Complete these exercises to solidify your understanding of modern C features. Each exercise builds upon the previous ones, gradually increasing in complexity.

Data Structures

Module 10: Data Structures Exercises

Exercise 1: Linear Data Structures Implementation

Write a program that implements fundamental linear data structures: - Implement a complete stack with push, pop, and peek operations - Create a queue with enqueue and dequeue functions - Develop a deque (double-ended queue) with operations on both ends - Implement a circular buffer for efficient memory usage - Include comprehensive error handling and boundary checks

Requirements: - Use dynamic memory allocation for flexible sizing - Include proper initialization and cleanup functions - Implement comprehensive error checking - Provide performance analysis for operations - Include examples of practical applications

Exercise 2: Linked List Variations

Create a program that implements various linked list types: - Singly linked list with insert, delete, and search operations - Doubly linked list with forward and backward traversal - Circular linked list with proper termination handling - Implement sorted insertion and merge operations - Include memory management and leak prevention

Requirements: - Handle all edge cases (empty list, single node, etc.) - Implement both iterative and recursive algorithms - Include performance comparisons between list types - Provide comprehensive test cases - Document time and space complexity

Exercise 3: Tree Data Structures

Develop a program that implements tree structures: - Binary search tree with insert, delete, and search operations - AVL tree with automatic balancing - Heap implementation (min-heap and max-heap) - Trie for string storage and retrieval - Include traversal algorithms (inorder, preorder, postorder, level-order)

Requirements: - Implement proper balancing algorithms for AVL trees - Include visualization functions for tree structures - Handle memory management correctly - Provide performance analysis for operations - Include comprehensive error handling

Exercise 4: Hash Tables and Associative Arrays

Write a program that implements hash-based data structures: - Hash table with collision resolution (chaining or open addressing) - Implement different hash functions and compare their performance - Create an associative array (dictionary) interface - Include dynamic resizing based on load factor - Provide statistics on collision rates and performance

Requirements: - Implement multiple collision resolution strategies - Include proper hash function selection and testing - Handle dynamic resizing efficiently - Provide performance profiling tools - Include comprehensive error handling

Exercise 5: Graph Representations and Algorithms

Create a program that works with graph data structures: - Implement adjacency list and adjacency matrix representations - Create graph traversal algorithms (BFS and DFS) - Implement shortest path algorithms (Dijkstra, Bellman-Ford) - Develop minimum spanning tree algorithms (Prim, Kruskal) - Include cycle detection and topological sorting

Requirements: - Handle both directed and undirected graphs - Implement weighted and unweighted graph versions - Include proper memory management for graph structures - Provide visualization tools for graph algorithms - Document algorithm complexity and use cases

Exercise 6: Advanced Data Structures

Write a program that implements specialized data structures: - Implement a priority queue using heap - Create a disjoint-set (union-find) data structure - Develop a segment tree for range queries - Implement a B-tree for database indexing - Include a bloom filter for probabilistic set membership

Requirements: - Include comprehensive documentation for each structure - Provide performance analysis and use case examples - Handle edge cases and error conditions properly - Include testing frameworks for verification - Document implementation details and optimizations

Exercise 7: Data Structure Performance Analysis

Create a program that analyzes and compares data structures: - Implement benchmarking tools for different operations - Create performance profiling for various data structures - Compare memory usage and cache efficiency - Analyze time complexity with empirical data - Include visualization of performance characteristics

Requirements: - Include comprehensive benchmarking framework - Provide statistical analysis of results - Handle large datasets for meaningful comparisons - Include cross-platform performance considerations - Document methodology and findings clearly

Exercise 8: Comprehensive Data Structure Application

Design a complete application that integrates multiple data structures: - Implement a database system using various data structures - Create a file system simulator with directory structures - Develop a network routing system with graph algorithms - Include a caching system with hash tables and LRU eviction - Provide comprehensive testing and validation

Requirements: - Use modular design with clear separation of concerns - Include proper documentation for all components - Handle all resource management properly - Implement robust error handling throughout - Provide clear examples and test cases

Solutions and Tips

Exercise 1 Solution Example:

```
#include <stdio.h>
#include <stdlib.h>
#include <stdbool.h>
#include <limits.h>

// Stack implementation
typedef struct {
    int *data;
    int top;
    int capacity;
} Stack;

// Function to create a stack
Stack* create_stack(int capacity) {
    Stack *stack = malloc(sizeof(Stack));
    if (stack == NULL) return NULL;

    stack->data = malloc(capacity * sizeof(int));
    if (stack->data == NULL) {
        free(stack);
        return NULL;
    }

    stack->top = -1;
    stack->capacity = capacity;
    return stack;
}

// Function to check if stack is empty
bool is_stack_empty(Stack *stack) {
```

```

    return stack == NULL || stack->top == -1;
}

// Function to check if stack is full
bool is_stack_full(Stack *stack) {
    return stack != NULL && stack->top == stack->capacity - 1;
}

// Function to push element onto stack
bool push(Stack *stack, int value) {
    if (is_stack_full(stack)) {
        printf("Error: Stack overflow\n");
        return false;
    }

    stack->data[++stack->top] = value;
    return true;
}

// Function to pop element from stack
int pop(Stack *stack) {
    if (is_stack_empty(stack)) {
        printf("Error: Stack underflow\n");
        return INT_MIN;
    }

    return stack->data[stack->top--];
}

// Function to peek at top element
int peek(Stack *stack) {
    if (is_stack_empty(stack)) {
        printf("Error: Stack is empty\n");
        return INT_MIN;
    }

    return stack->data[stack->top];
}

// Function to free stack memory
void free_stack(Stack *stack) {
    if (stack != NULL) {
        free(stack->data);
        free(stack);
    }
}

```

```

int main() {
    Stack *stack = create_stack(5);
    if (stack == NULL) {
        printf("Error: Failed to create stack\n");
        return 1;
    }

    // Test stack operations
    printf("Testing stack operations:\n");

    // Push elements
    for (int i = 1; i <= 5; i++) {
        if (push(stack, i * 10)) {
            printf("Pushed %d onto stack\n", i * 10);
        }
    }

    // Try to push to full stack
    if (!push(stack, 60)) {
        printf("Successfully handled stack overflow\n");
    }

    // Peek at top element
    printf("Top element: %d\n", peek(stack));

    // Pop elements
    while (!is_stack_empty(stack)) {
        int value = pop(stack);
        printf("Popped %d from stack\n", value);
    }

    // Try to pop from empty stack
    int empty_pop = pop(stack);
    if (empty_pop == INT_MIN) {
        printf("Successfully handled stack underflow\n");
    }

    free_stack(stack);
    return 0;
}

```

Exercise 3 Solution Example:

```
#include <stdio.h>
#include <stdlib.h>
#include <stdbool.h>

// Binary Search Tree Node
typedef struct TreeNode {
    int data;
    struct TreeNode *left;
    struct TreeNode *right;
} TreeNode;

// Function to create a new tree node
TreeNode* create_node(int data) {
    TreeNode *node = malloc(sizeof(TreeNode));
    if (node == NULL) return NULL;

    node->data = data;
    node->left = NULL;
    node->right = NULL;
    return node;
}

// Function to insert a node into BST
TreeNode* insert_bst(TreeNode *root, int data) {
    if (root == NULL) {
        return create_node(data);
    }

    if (data < root->data) {
        root->left = insert_bst(root->left, data);
    } else if (data > root->data) {
        root->right = insert_bst(root->right, data);
    }

    // If data == root->data, we don't insert duplicates

    return root;
}

// Function for inorder traversal
void inorder_traversal(TreeNode *root) {
    if (root != NULL) {
        inorder_traversal(root->left);
        printf("%d ", root->data);
        inorder_traversal(root->right);
    }
}
```

```

}

// Function to search for a value in BST
TreeNode* search_bst(TreeNode *root, int data) {
    if (root == NULL || root->data == data) {
        return root;
    }

    if (data < root->data) {
        return search_bst(root->left, data);
    } else {
        return search_bst(root->right, data);
    }
}

// Function to find minimum value node
TreeNode* find_min(TreeNode *node) {
    while (node && node->left != NULL) {
        node = node->left;
    }
    return node;
}

// Function to delete a node from BST
TreeNode* delete_bst(TreeNode *root, int data) {
    if (root == NULL) return root;

    if (data < root->data) {
        root->left = delete_bst(root->left, data);
    } else if (data > root->data) {
        root->right = delete_bst(root->right, data);
    } else {
        // Node to be deleted found
        if (root->left == NULL) {
            TreeNode *temp = root->right;
            free(root);
            return temp;
        } else if (root->right == NULL) {
            TreeNode *temp = root->left;
            free(root);
            return temp;
        }

        // Node with two children
        TreeNode *temp = find_min(root->right);
        root->data = temp->data;
    }
}

```

```

        root->right = delete_bst(root->right, temp->data);
    }

    return root;
}

// Function to free entire tree
void free_tree(TreeNode *root) {
    if (root != NULL) {
        free_tree(root->left);
        free_tree(root->right);
        free(root);
    }
}

int main() {
    TreeNode *root = NULL;

    // Insert elements into BST
    int values[] = {50, 30, 70, 20, 40, 60, 80};
    int n = sizeof(values) / sizeof(values[0]);

    printf("Inserting values into BST: ");
    for (int i = 0; i < n; i++) {
        printf("%d ", values[i]);
        root = insert_bst(root, values[i]);
    }
    printf("\n");

    // Inorder traversal (should be sorted)
    printf("Inorder traversal: ");
    inorder_traversal(root);
    printf("\n");

    // Search for values
    int search_values[] = {40, 90};
    for (int i = 0; i < 2; i++) {
        TreeNode *result = search_bst(root, search_values[i]);
        if (result != NULL) {
            printf("Found %d in BST\n", search_values[i]);
        } else {
            printf("%d not found in BST\n", search_values[i]);
        }
    }

    // Delete a node

```

```

    printf("Deleting 30 from BST\n");
    root = delete_bst(root, 30);

    printf("Inorder traversal after deletion: ");
    inorder_traversal(root);
    printf("\n");

    // Free memory
    free_tree(root);
    return 0;
}

```

Common Pitfalls to Avoid:

1. **Memory leaks:** Always free allocated memory for dynamic structures
2. **Null pointer dereference:** Check pointers before dereferencing
3. **Boundary conditions:** Handle empty structures and edge cases properly
4. **Infinite recursion:** Ensure proper base cases for recursive algorithms
5. **Performance degradation:** Choose appropriate data structures for use cases

Best Practices:

1. **Modular design:** Implement each data structure as separate modules
2. **Error handling:** Include comprehensive error checking and reporting
3. **Documentation:** Comment interface functions and implementation details
4. **Testing:** Provide thorough test cases for all operations
5. **Performance analysis:** Include complexity analysis and benchmarking

Complete these exercises to solidify your understanding of data structures in C. Each exercise builds upon the previous ones, gradually increasing in complexity.

Network

Module 11: Network Programming Exercises

Exercise 1: Basic Socket Programming

Write a program that demonstrates fundamental socket operations: - Create and configure TCP and UDP sockets - Implement basic client-server communication - Handle socket binding, listening, and connection acceptance - Send and receive data through sockets - Properly close and cleanup socket resources

Requirements: - Implement both TCP and UDP examples - Include proper error handling for all socket operations - Handle cross-platform differences (Windows vs Unix) - Implement timeout handling for socket operations - Provide clear examples of client-server interaction

Exercise 2: HTTP Client Implementation

Create a program that implements a simple HTTP client: - Parse URLs and extract host, port, and path components - Establish HTTP connections to web servers - Send HTTP GET and POST requests - Parse HTTP response headers and status codes - Handle chunked transfer encoding and content length

Requirements: - Implement proper HTTP/1.1 protocol compliance - Include support for common HTTP headers - Handle redirects and authentication (basic) - Implement connection reuse for multiple requests - Include proper error handling for network issues

Exercise 3: Multi-client Server

Develop a program that implements a server handling multiple clients: - Use `fork()` or threading to handle concurrent clients - Implement connection pooling for efficient resource usage - Handle client disconnections gracefully - Include client session management - Implement basic rate limiting and security measures

Requirements: - Handle race conditions with proper synchronization - Include timeout mechanisms for inactive clients - Implement proper resource cleanup for terminated clients - Provide logging and monitoring capabilities - Include comprehensive error handling

Exercise 4: Network Protocol Implementation

Write a program that implements custom network protocols: - Design and implement a simple chat protocol - Create a file transfer protocol with checksum verification - Implement a simple remote command execution protocol - Include protocol versioning and compatibility - Handle protocol state management

Requirements: - Define clear protocol specifications - Include proper message framing and parsing - Implement error recovery mechanisms - Handle network byte order conversion - Provide protocol documentation and examples

Exercise 5: Advanced Networking Concepts

Create a program that demonstrates advanced networking features: - Implement non-blocking I/O with select() or poll() - Use epoll() or kqueue() for high-performance I/O (Linux/BSD) - Implement socket timeouts and keep-alive mechanisms - Handle network address resolution and DNS lookups - Include IPv6 support alongside IPv4

Requirements: - Include performance comparisons between different I/O models - Handle partial reads and writes correctly - Implement proper signal handling for network applications - Include network interface enumeration and monitoring - Provide cross-platform compatibility

Exercise 6: Network Security Implementation

Write a program that implements basic network security: - Implement simple encryption/decryption functions - Create secure communication channels with basic cryptography - Handle certificate validation and SSL/TLS (bonus) - Implement basic authentication mechanisms - Include secure key exchange protocols

Requirements: - Use established cryptographic libraries (OpenSSL, etc.) - Include proper random number generation - Handle key management and storage securely - Implement secure session establishment - Provide security best practices documentation

Exercise 7: Network Debugging and Monitoring

Create a program that provides network debugging capabilities: - Implement packet capture and analysis (bonus) - Create network performance monitoring tools - Develop connection state tracking and logging - Include bandwidth measurement and reporting - Implement network error diagnosis tools

Requirements: - Include comprehensive logging and tracing - Provide real-time monitoring capabilities - Handle large volumes of network data efficiently - Include visualization of network statistics - Document debugging techniques and tools

Exercise 8: Comprehensive Network Application

Design a complete network application that integrates all concepts: - Implement a distributed system with multiple network components - Create a peer-to-peer file sharing network - Develop a simple web server with dynamic content - Include network service discovery and registration - Provide comprehensive testing and validation

Requirements: - Use modular design with clear separation of concerns - Include proper documentation for all components - Handle all network resource management properly - Implement robust error handling throughout - Provide clear examples and test cases

Solutions and Tips

Exercise 1 Solution Example:

```
#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <unistd.h>

#ifdef _WIN32
    #include <winsock2.h>
    #include <ws2tcpip.h>
    #pragma comment(lib, "ws2_32.lib")
#else
    #include <sys/socket.h>
    #include <arpa/inet.h>
    #include <netinet/in.h>
    #include <netdb.h>
#endif

// Cross-platform socket definitions
#ifdef _WIN32
    typedef SOCKET socket_t;
    #define CLOSE_SOCKET closesocket
    #define SOCK_ERR SOCKET_ERROR
#else
    typedef int socket_t;
    #define CLOSE_SOCKET close
    #define SOCK_ERR -1
#endif

// Function to initialize sockets (Windows specific)
int initialize_sockets(void) {
#ifdef _WIN32
```

```

WSADATA wsaData;
int result = WSASStartup(MAKEWORD(2, 2), &wsaData);
if (result != 0) {
    printf("WSAStartup failed: %d\n", result);
    return 0;
}
#endif
return 1;
}

// Function to cleanup sockets (Windows specific)
void cleanup_sockets(void) {
#ifdef _WIN32
    WSACleanup();
#endif
}

// Function to create TCP socket
socket_t create_tcp_socket(void) {
    socket_t sock = socket(AF_INET, SOCK_STREAM, 0);
    if (sock == SOCK_ERR) {
        perror("Socket creation failed");
    }
    return sock;
}

// Function to create UDP socket
socket_t create_udp_socket(void) {
    socket_t sock = socket(AF_INET, SOCK_DGRAM, 0);
    if (sock == SOCK_ERR) {
        perror("UDP Socket creation failed");
    }
    return sock;
}

// Simple TCP server example
int tcp_server_example(int port) {
    if (!initialize_sockets()) return -1;

    socket_t server_sock = create_tcp_socket();
    if (server_sock == SOCK_ERR) {
        cleanup_sockets();
        return -1;
    }

    struct sockaddr_in server_addr;
    memset(&server_addr, 0, sizeof(server_addr));

```

```

server_addr.sin_family = AF_INET;
server_addr.sin_addr.s_addr = INADDR_ANY;
server_addr.sin_port = htons(port);

if (bind(server_sock, (struct sockaddr*)&server_addr, sizeof(server_addr)) == SOCK_ERR)
    perror("Bind failed");
    CLOSE_SOCKET(server_sock);
    cleanup_sockets();
    return -1;
}

if (listen(server_sock, 5) == SOCK_ERR) {
    perror("Listen failed");
    CLOSE_SOCKET(server_sock);
    cleanup_sockets();
    return -1;
}

printf("TCP Server listening on port %d\n", port);

struct sockaddr_in client_addr;
socklen_t client_len = sizeof(client_addr);
socket_t client_sock = accept(server_sock, (struct sockaddr*)&client_addr, &client_len);

if (client_sock != SOCK_ERR) {
    char buffer[1024];
    int bytes_received = recv(client_sock, buffer, sizeof(buffer) - 1, 0);
    if (bytes_received > 0) {
        buffer[bytes_received] = '\0';
        printf("Received: %s\n", buffer);

        const char *response = "Hello from TCP server!";
        send(client_sock, response, strlen(response), 0);
    }
    CLOSE_SOCKET(client_sock);
}

CLOSE_SOCKET(server_sock);
cleanup_sockets();
return 0;
}

// Simple TCP client example
int tcp_client_example(const char *server_ip, int port) {
    if (!initialize_sockets()) return -1;

    socket_t client_sock = create_tcp_socket();

```

```

    if (client_sock == SOCK_ERR) {
        cleanup_sockets();
        return -1;
    }

    struct sockaddr_in server_addr;
    memset(&server_addr, 0, sizeof(server_addr));
    server_addr.sin_family = AF_INET;
    server_addr.sin_port = htons(port);

    if (inet_pton(AF_INET, server_ip, &server_addr.sin_addr) <= 0) {
        printf("Invalid address\n");
        CLOSE_SOCKET(client_sock);
        cleanup_sockets();
        return -1;
    }

    if (connect(client_sock, (struct sockaddr*)&server_addr, sizeof(server_addr)) == SOCK_ER

        perror("Connection failed");
        CLOSE_SOCKET(client_sock);
        cleanup_sockets();
        return -1;
    }

    const char *message = "Hello from TCP client!";
    send(client_sock, message, strlen(message), 0);

    char buffer[1024];
    int bytes_received = recv(client_sock, buffer, sizeof(buffer) - 1, 0);
    if (bytes_received > 0) {
        buffer[bytes_received] = '\0';
        printf("Received: %s\n", buffer);
    }

    CLOSE_SOCKET(client_sock);
    cleanup_sockets();
    return 0;
}

int main() {
    // Note: For testing, you would run the server and client separately
    // This is just an example of how to structure the code

    printf("Network Programming Examples\n");
    printf("Run tcp_server_example() in one terminal\n");
    printf("Run tcp_client_example() in another terminal\n");

```

```

    return 0;
}

```

Exercise 2 Solution Example:

```

#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <unistd.h>

#ifdef _WIN32
    #include <winsock2.h>
    #include <ws2tcpip.h>
    #pragma comment(lib, "ws2_32.lib")
#else
    #include <sys/socket.h>
    #include <arpa/inet.h>
    #include <netinet/in.h>
    #include <netdb.h>
#endif

#ifdef _WIN32
    typedef SOCKET socket_t;
    #define CLOSE_SOCKET closesocket
    #define SOCK_ERR SOCKET_ERROR
#else
    typedef int socket_t;
    #define CLOSE_SOCKET close
    #define SOCK_ERR -1
#endif

// URL parsing structure
typedef struct {
    char protocol[16];
    char host[256];
    int port;
    char path[512];
} URL;

// Function to parse URL
int parse_url(const char *url_str, URL *url) {
    // Simple URL parsing (supports http://host:port/path)
    if (sscanf(url_str, "%15[^:]:/%255[^:]:%d%511s",
               url->protocol, url->host, &url->port, url->path) == 4) {
        return 0;
    }
}

```



```

    }

    // Try without port
    if (sscanf(url_str, "%15[^:]://%255[~/]%511s",
        url->protocol, url->host, url->path) == 3) {
        url->port = (strcmp(url->protocol, "https") == 0) ? 443 : 80;
        return 0;
    }

    return -1;
}

// Simple HTTP GET request
int http_get_request(const char *url_str) {
#ifdef _WIN32
    WSADATA wsaData;
    if (WSAStartup(MAKEWORD(2, 2), &wsaData) != 0) {
        return -1;
    }
#endif

    URL url;
    if (parse_url(url_str, &url) != 0) {
        printf("Failed to parse URL: %s\n", url_str);
#ifdef _WIN32
        WSACleanup();
#endif
        return -1;
    }

    // Create socket
    socket_t sock = socket(AF_INET, SOCK_STREAM, 0);
    if (sock == SOCK_ERR) {
        perror("Socket creation failed");
#ifdef _WIN32
        WSACleanup();
#endif
        return -1;
    }

    // Resolve hostname
    struct hostent *host = gethostbyname(url.host);
    if (host == NULL) {
        printf("Failed to resolve hostname: %s\n", url.host);
        CLOSE_SOCKET(sock);
#ifdef _WIN32
        WSACleanup();
#endif
    }

```

```

#endif
    return -1;
}

// Connect to server
struct sockaddr_in server_addr;
memset(&server_addr, 0, sizeof(server_addr));
server_addr.sin_family = AF_INET;
server_addr.sin_port = htons(url.port);
memcpy(&server_addr.sin_addr, host->h_addr_list[0], host->h_length);

if (connect(sock, (struct sockaddr*)&server_addr, sizeof(server_addr)) == SOCK_ERR) {
    perror("Connection failed");
    CLOSE_SOCKET(sock);
#ifdef _WIN32
    WSACleanup();
#endif
    return -1;
}

// Send HTTP GET request
char request[1024];
snprintf(request, sizeof(request),
    "GET %s HTTP/1.1\r\n"
    "Host: %s:%d\r\n"
    "Connection: close\r\n"
    "\r\n",
    url.path, url.host, url.port);

if (send(sock, request, strlen(request), 0) == SOCK_ERR) {
    perror("Send failed");
    CLOSE_SOCKET(sock);
#ifdef _WIN32
    WSACleanup();
#endif
    return -1;
}

// Receive response
char buffer[4096];
int bytes_received;
while ((bytes_received = recv(sock, buffer, sizeof(buffer) - 1, 0)) > 0) {
    buffer[bytes_received] = '\0';
    printf("%s", buffer);
}

CLOSE_SOCKET(sock);

```

```

#ifdef _WIN32
    WSACleanup();
#endif
return 0;
}

int main() {
    const char *url = "http://httpbin.org/get";
    printf("Making HTTP GET request to: %s\n\n", url);

    if (http_get_request(url) != 0) {
        printf("HTTP request failed\n");
        return 1;
    }

    return 0;
}

```

Common Pitfalls to Avoid:

1. **Platform differences:** Handle Windows and Unix socket API differences
2. **Resource leaks:** Always close sockets and cleanup network libraries
3. **Buffer overflows:** Check buffer sizes when receiving network data
4. **Blocking operations:** Handle timeouts to prevent hanging applications
5. **Error handling:** Check return values from all network functions

Best Practices:

1. **Cross-platform compatibility:** Use conditional compilation for platform differences
2. **Resource management:** Implement proper cleanup in error paths
3. **Security considerations:** Validate all network input and use secure protocols
4. **Performance optimization:** Use appropriate I/O models for scalability
5. **Error recovery:** Implement robust error handling and recovery mechanisms

Complete these exercises to solidify your understanding of network programming in C. Each exercise builds upon the previous ones, gradually increasing in complexity.

Embedded

Module 12: Embedded Systems Programming Exercises

Exercise 1: Bit Manipulation and Register Operations

Write a program that demonstrates fundamental bit manipulation operations commonly used in embedded systems: - Implement functions for setting, clearing, and toggling individual bits - Create macros for bit manipulation operations - Demonstrate bit field usage for hardware register access - Implement functions for reading and writing hardware registers - Show how to handle endianness in register operations

Requirements: - Implement bit manipulation functions without using built-in bit operations - Create reusable macros for common bit operations - Demonstrate proper handling of volatile variables for hardware registers - Include examples of bit field structures for hardware register mapping - Provide clear documentation for all bit manipulation functions

Exercise 2: Memory-Mapped I/O Implementation

Create a program that simulates memory-mapped I/O operations: - Implement functions for reading and writing to memory-mapped registers - Create a simple GPIO (General Purpose Input/Output) simulation - Demonstrate proper use of volatile pointers for hardware access - Implement interrupt handling for I/O operations - Show how to handle memory alignment requirements

Requirements: - Use appropriate data types for memory-mapped registers - Implement proper error checking for I/O operations - Include examples of direct memory access patterns - Demonstrate proper handling of memory barriers if needed - Provide clear examples of GPIO configuration and control

Exercise 3: Interrupt Service Routines

Develop a program that implements interrupt service routines (ISRs): - Create a framework for registering and handling interrupts - Implement a simple timer interrupt simulation - Demonstrate proper context saving and restoration - Show how to handle nested interrupts - Implement interrupt prioritization mechanisms

Requirements: - Follow proper ISR coding practices (minimal code, no blocking operations) - Include examples of interrupt enable/disable functions - Demonstrate proper use of atomic operations in ISRs - Implement interrupt vector tables or similar dispatch mechanisms - Provide clear documentation of interrupt handling flow

Exercise 4: Real-Time Constraints Implementation

Write a program that demonstrates handling real-time constraints: - Implement a simple real-time scheduler - Create functions with guaranteed execution time limits - Demonstrate proper handling of timing-critical sections - Implement watchdog timer functionality - Show how to measure and optimize code execution time

Requirements: - Use appropriate timing functions for the target platform - Include examples of critical section protection - Demonstrate proper use of hardware timers - Implement timeout mechanisms for blocking operations - Provide clear examples of real-time performance measurement

Exercise 5: Low-Power Programming Techniques

Create a program that implements low-power programming techniques: - Implement sleep and power-down modes - Create functions for dynamic clock scaling - Demonstrate proper handling of wake-up sources - Show how to optimize code for power consumption - Implement power management state machines

Requirements: - Include examples of different power-saving modes - Demonstrate proper handling of sleep/wake cycles - Show how to preserve state during power transitions - Implement power consumption measurement and reporting - Provide clear documentation of power management strategies

Exercise 6: Hardware Abstraction Layer (HAL)

Write a program that implements a simple Hardware Abstraction Layer: - Create a unified interface for different hardware platforms - Implement device driver interfaces for common peripherals - Demonstrate proper error handling in HAL functions - Show how to handle hardware initialization and configuration - Implement version control for HAL interfaces

Requirements: - Use consistent naming conventions across all HAL functions - Include proper error codes and return value handling - Demonstrate proper initialization and cleanup sequences - Provide clear documentation for all HAL interfaces - Include examples of platform-specific implementations

Exercise 7: Embedded Debugging and Monitoring

Create a program that provides debugging and monitoring capabilities for embedded systems: - Implement a simple logging system with minimal resource usage - Create functions for runtime system monitoring - Demonstrate proper use of assertions in embedded environments - Show how to implement non-intrusive debugging techniques - Implement memory usage tracking and reporting

Requirements: - Include configurable logging levels and output destinations - Demonstrate proper handling of limited memory resources - Show how to implement circular buffers for logging - Provide examples of runtime error detection and reporting - Include clear documentation of debugging techniques

Exercise 8: Comprehensive Embedded Application

Design a complete embedded application that integrates all concepts: - Implement a simple embedded system with multiple components - Create a state machine for system control - Demonstrate proper handling of hardware and software components - Include comprehensive error handling and recovery mechanisms - Provide clear documentation and testing procedures

Requirements: - Use modular design with clear separation of concerns - Include proper initialization and shutdown sequences - Demonstrate proper resource management throughout the application - Implement robust error handling and recovery mechanisms - Provide clear examples and test cases for all components

Solutions and Tips

Exercise 1 Solution Example:

```
#include <stdio.h>
#include <stdint.h>

// Bit manipulation macros
#define SET_BIT(reg, bit)      ((reg) |= (1U << (bit)))
#define CLEAR_BIT(reg, bit)   ((reg) &= ~(1U << (bit)))
#define TOGGLE_BIT(reg, bit)  ((reg) ^= (1U << (bit)))
#define CHECK_BIT(reg, bit)   (((reg) >> (bit)) & 1U)

// Bit manipulation functions
void set_bit(uint32_t *reg, int bit) {
    *reg |= (1U << bit);
}

void clear_bit(uint32_t *reg, int bit) {
    *reg &= ~(1U << bit);
}

void toggle_bit(uint32_t *reg, int bit) {
    *reg ^= (1U << bit);
}

int check_bit(uint32_t reg, int bit) {
    return (reg >> bit) & 1U;
}

// Bit field structure for hardware register
typedef struct {
    uint32_t enable    : 1; // Bit 0
```

```

    uint32_t direction : 1; // Bit 1
    uint32_t reserved  : 2; // Bits 2-3
    uint32_t mode       : 4; // Bits 4-7
    uint32_t config     : 24; // Bits 8-31
} gpio_register_t;

// Volatile pointer for hardware register access
volatile gpio_register_t *gpio_reg = (volatile gpio_register_t *)0x40020000;

int main() {
    uint32_t test_reg = 0;

    // Test bit manipulation functions
    printf("Initial register value: 0x%08X\n", test_reg);

    set_bit(&test_reg, 5);
    printf("After setting bit 5: 0x%08X\n", test_reg);

    toggle_bit(&test_reg, 5);
    printf("After toggling bit 5: 0x%08X\n", test_reg);

    clear_bit(&test_reg, 5);
    printf("After clearing bit 5: 0x%08X\n", test_reg);

    // Test bit manipulation macros
    SET_BIT(test_reg, 3);
    printf("After setting bit 3 with macro: 0x%08X\n", test_reg);

    if (CHECK_BIT(test_reg, 3)) {
        printf("Bit 3 is set\n");
    }

    // Test bit field structure
    gpio_reg->enable = 1;
    gpio_reg->direction = 1;
    gpio_reg->mode = 0xF;

    printf("GPIO Register - Enable: %d, Direction: %d, Mode: 0x%X\n",
        gpio_reg->enable, gpio_reg->direction, gpio_reg->mode);

    return 0;
}

```


Exercise 2 Solution Example:

```
#include <stdio.h>
#include <stdint.h>
#include <stdlib.h>
#include <string.h>

// Memory-mapped I/O simulation
#define GPIO_BASE_ADDR 0x40020000
#define GPIO_DIR_REG    (*(volatile uint32_t *) (GPIO_BASE_ADDR + 0x00))
#define GPIO_DATA_REG   (*(volatile uint32_t *) (GPIO_BASE_ADDR + 0x04))
#define GPIO_SET_REG    (*(volatile uint32_t *) (GPIO_BASE_ADDR + 0x08))
#define GPIO_CLR_REG    (*(volatile uint32_t *) (GPIO_BASE_ADDR + 0x0C))

// GPIO pin definitions
#define GPIO_PIN_0  (1U << 0)
#define GPIO_PIN_1  (1U << 1)
#define GPIO_PIN_2  (1U << 2)
#define GPIO_PIN_3  (1U << 3)

// Simple GPIO simulation structure
typedef struct {
    uint32_t direction; // 0 = input, 1 = output
    uint32_t data;      // Current pin states
} gpio_sim_t;

// Global GPIO simulation instance
static gpio_sim_t gpio_sim = {0};

// Function to initialize GPIO
void gpio_init(void) {
    gpio_sim.direction = 0;
    gpio_sim.data = 0;
    printf("GPIO initialized\n");
}

// Function to set GPIO direction
void gpio_set_direction(uint32_t pins, int direction) {
    if (direction) {
        gpio_sim.direction |= pins; // Set as output
    } else {
        gpio_sim.direction &= ~pins; // Set as input
    }

    // Update hardware register
    GPIO_DIR_REG = gpio_sim.direction;
}
```

```

        printf("GPIO direction set: 0x%08X (pins: 0x%08X, direction: %s)\n",
               gpio_sim.direction, pins, direction ? "output" : "input");
    }

// Function to write to GPIO pins
void gpio_write(uint32_t pins, int value) {
    if (value) {
        gpio_sim.data |= pins;
        GPIO_SET_REG = pins; // Hardware register for setting pins
    } else {
        gpio_sim.data &= ~pins;
        GPIO_CLR_REG = pins; // Hardware register for clearing pins
    }

    GPIO_DATA_REG = gpio_sim.data; // Update data register
    printf("GPIO write: pins 0x%08X set to %d\n", pins, value);
}

// Function to read from GPIO pins
uint32_t gpio_read(uint32_t pins) {
    uint32_t value = GPIO_DATA_REG & pins; // Read from hardware register
    printf("GPIO read: pins 0x%08X = 0x%08X\n", pins, value);
    return value;
}

// Function to simulate external input change
void gpio_simulate_input(uint32_t pins, int value) {
    if (value) {
        gpio_sim.data |= pins;
    } else {
        gpio_sim.data &= ~pins;
    }
    printf("Simulated input change: pins 0x%08X set to %d\n", pins, value);
}

int main() {
    // Initialize GPIO
    gpio_init();

    // Configure pins 0-1 as outputs, pins 2-3 as inputs
    gpio_set_direction(GPIO_PIN_0 | GPIO_PIN_1, 1); // Output
    gpio_set_direction(GPIO_PIN_2 | GPIO_PIN_3, 0); // Input

    // Write to output pins
    gpio_write(GPIO_PIN_0, 1);
    gpio_write(GPIO_PIN_1, 0);

```

```

// Simulate input changes
gpio_simulate_input(GPIO_PIN_2, 1);
gpio_simulate_input(GPIO_PIN_3, 0);

// Read input pins
uint32_t input_value = gpio_read(GPIO_PIN_2 | GPIO_PIN_3);
printf("Input pins value: 0x%08X\n", input_value);

return 0;
}

```

Common Pitfalls to Avoid:

1. **Volatile keyword misuse:** Always use volatile for hardware registers
2. **Endianness issues:** Handle byte order correctly in multi-byte registers
3. **Timing constraints:** Avoid blocking operations in time-critical code
4. **Memory alignment:** Ensure proper alignment for memory-mapped I/O
5. **Interrupt safety:** Protect critical sections in interrupt-driven code

Best Practices:

1. **Hardware abstraction:** Use consistent interfaces for hardware access
2. **Error handling:** Implement proper error checking for all operations
3. **Resource management:** Carefully manage memory and other resources
4. **Documentation:** Provide clear documentation for hardware interactions
5. **Testing:** Include comprehensive testing for embedded functionality

Complete these exercises to solidify your understanding of embedded systems programming in C. Each exercise builds upon the previous ones, gradually increasing in complexity.

Performance

Module 13: Performance Optimization Exercises

Exercise 1: Algorithm Complexity Analysis

Write a program that demonstrates algorithm complexity analysis and optimization: - Implement multiple versions of common algorithms with different complexities - Create functions to measure execution time of algorithms - Demonstrate the impact of algorithmic complexity on performance - Show how to profile code to identify bottlenecks - Implement optimized versions of inefficient algorithms

Requirements: - Include examples of $O(1)$, $O(\log n)$, $O(n)$, $O(n \log n)$, and $O(n^2)$ algorithms - Implement accurate timing functions for performance measurement - Provide clear visualization of performance differences - Include examples of common algorithmic optimizations - Document the analysis process and findings

Exercise 2: Memory Access Patterns Optimization

Create a program that demonstrates memory access pattern optimization: - Implement functions with different memory access patterns - Show the impact of cache locality on performance - Demonstrate proper use of memory alignment - Implement techniques for reducing cache misses - Show how to optimize data structure layouts

Requirements: - Include examples of sequential, random, and strided memory access - Implement cache-friendly data structures - Demonstrate proper memory alignment techniques - Show the impact of data structure padding and packing - Provide clear performance comparisons between different approaches

Exercise 3: Compiler Optimization Techniques

Develop a program that demonstrates compiler optimization techniques: - Implement code that showcases different compiler optimization levels - Show how to use compiler hints and attributes effectively - Demonstrate proper use of inline functions and macros - Implement examples of loop optimization techniques - Show how to prevent unwanted compiler optimizations for testing

Requirements: - Include examples of common compiler optimization flags - Demonstrate proper use of `restrict` keyword and `volatile` - Show how to use compiler-specific attributes and pragmas - Implement examples of manual loop unrolling and vectorization - Provide clear documentation of optimization techniques

Exercise 4: Data Structure Optimization

Write a program that demonstrates data structure optimization: - Implement multiple versions of common data structures - Show the impact of different data structure choices on performance - Demonstrate proper use of memory pools and object pools - Implement custom allocators for specific use cases - Show how to optimize data structures for specific access patterns

Requirements: - Include examples of array vs linked list performance trade-offs - Implement cache-oblivious data structures - Demonstrate proper use of memory preallocation - Show how to optimize data structures for concurrent access - Provide clear performance comparisons and analysis

Exercise 5: Parallel Processing Implementation

Create a program that implements parallel processing techniques: - Implement multi-threaded versions of compute-intensive tasks - Show how to use OpenMP or pthreads effectively - Demonstrate proper synchronization techniques - Implement lock-free data structures (bonus) - Show how to measure and optimize parallel performance

Requirements: - Include examples of task parallelism and data parallelism - Demonstrate proper use of thread pools - Show how to avoid common parallel programming pitfalls - Implement proper load balancing techniques - Provide clear scalability analysis and measurements

Exercise 6: SIMD and Vectorization

Write a program that demonstrates SIMD (Single Instruction, Multiple Data) optimization: - Implement vectorized versions of compute-intensive functions - Show how to use SSE, AVX, or other SIMD instruction sets - Demonstrate proper data alignment for SIMD operations - Implement fallback implementations for different architectures - Show how to measure SIMD performance improvements

Requirements: - Include examples of manual vectorization and compiler auto-vectorization - Demonstrate proper handling of data alignment requirements - Show how to handle different vector widths and instruction sets - Implement proper error handling for SIMD operations - Provide clear performance comparisons between scalar and vector versions

Exercise 7: Memory Management Optimization

Create a program that demonstrates memory management optimization: - Implement custom memory allocators for specific use cases - Show how to reduce memory fragmentation - Demonstrate proper use of memory pools and arenas - Implement garbage collection-like mechanisms (bonus) - Show how to profile and optimize memory usage

Requirements: - Include examples of stack vs heap allocation performance - Demonstrate proper use of memory preallocation techniques - Show how to implement efficient memory pooling strategies - Implement proper memory leak detection and prevention - Provide clear memory usage analysis and optimization

Exercise 8: Comprehensive Performance Optimization

Design a complete application that demonstrates all performance optimization techniques: - Implement a compute-intensive application with multiple components - Apply all optimization techniques learned in previous exercises - Include comprehensive performance profiling and analysis - Demonstrate proper trade-offs between different optimization approaches - Provide clear documentation and performance results

Requirements: - Use modular design with clear performance boundaries - Include comprehensive profiling and benchmarking - Demonstrate proper optimization prioritization - Implement robust performance monitoring throughout the application - Provide clear examples and performance comparisons

Solutions and Tips

Exercise 1 Solution Example:

```
#include <stdio.h>
#include <stdlib.h>
#include <time.h>
#include <string.h>

// Timing function
double get_time_diff(struct timespec start, struct timespec end) {
    return (end.tv_sec - start.tv_sec) + (end.tv_nsec - start.tv_nsec) / 1e9;
}

// O(n²) bubble sort
void bubble_sort(int arr[], int n) {
    for (int i = 0; i < n - 1; i++) {
        for (int j = 0; j < n - i - 1; j++) {
            if (arr[j] > arr[j + 1]) {
                int temp = arr[j];
                arr[j] = arr[j + 1];
                arr[j + 1] = temp;
            }
        }
    }
}

// O(n log n) quick sort
void quick_sort(int arr[], int low, int high) {
    if (low < high) {
        int pi = partition(arr, low, high);
        quick_sort(arr, low, pi - 1);
    }
}
```

```

        quick_sort(arr, pi + 1, high);
    }
}

int partition(int arr[], int low, int high) {
    int pivot = arr[high];
    int i = (low - 1);

    for (int j = low; j <= high - 1; j++) {
        if (arr[j] < pivot) {
            i++;
            int temp = arr[i];
            arr[i] = arr[j];
            arr[j] = temp;
        }
    }
    int temp = arr[i + 1];
    arr[i + 1] = arr[high];
    arr[high] = temp;
    return (i + 1);
}

// O(n) linear search
int linear_search(int arr[], int n, int key) {
    for (int i = 0; i < n; i++) {
        if (arr[i] == key) {
            return i;
        }
    }
    return -1;
}

// O(log n) binary search
int binary_search(int arr[], int low, int high, int key) {
    if (high >= low) {
        int mid = low + (high - low) / 2;
        if (arr[mid] == key) {
            return mid;
        }
        if (arr[mid] > key) {
            return binary_search(arr, low, mid - 1, key);
        }
        return binary_search(arr, mid + 1, high, key);
    }
    return -1;
}

```



```

// Performance testing function
void test_algorithms() {
    const int sizes[] = {1000, 5000, 10000, 20000};
    const int num_sizes = sizeof(sizes) / sizeof(sizes[0]);

    for (int s = 0; s < num_sizes; s++) {
        int n = sizes[s];
        int *arr1 = malloc(n * sizeof(int));
        int *arr2 = malloc(n * sizeof(int));

        // Fill arrays with random data
        for (int i = 0; i < n; i++) {
            arr1[i] = rand() % 10000;
            arr2[i] = arr1[i];
        }

        struct timespec start, end;

        // Test bubble sort
        clock_gettime(CLOCK_MONOTONIC, &start);
        bubble_sort(arr1, n);
        clock_gettime(CLOCK_MONOTONIC, &end);
        double bubble_time = get_time_diff(start, end);

        // Test quick sort
        clock_gettime(CLOCK_MONOTONIC, &start);
        quick_sort(arr2, 0, n - 1);
        clock_gettime(CLOCK_MONOTONIC, &end);
        double quick_time = get_time_diff(start, end);

        printf("Array size: %d\n", n);
        printf("Bubble sort time: %.6f seconds\n", bubble_time);
        printf("Quick sort time: %.6f seconds\n", quick_time);
        printf("Speedup: %.2fx\n\n", bubble_time / quick_time);

        free(arr1);
        free(arr2);
    }
}

int main() {
    srand(time(NULL));
    test_algorithms();
    return 0;
}

```

Exercise 2 Solution Example:

```
#include <stdio.h>
#include <stdlib.h>
#include <time.h>
#include <string.h>

#define MATRIX_SIZE 1000
#define CACHE_LINE_SIZE 64

// Matrix structure with proper alignment
typedef struct {
    int rows;
    int cols;
    int *data;
} matrix_t;

// Create matrix with specified alignment
matrix_t* create_matrix(int rows, int cols) {
    matrix_t *mat = malloc(sizeof(matrix_t));
    mat->rows = rows;
    mat->cols = cols;
    // Allocate aligned memory for better cache performance
    posix_memalign((void**)&mat->data, CACHE_LINE_SIZE, rows * cols * sizeof(int));
    return mat;
}

void free_matrix(matrix_t *mat) {
    free(mat->data);
    free(mat);
}

// Initialize matrix with random values
void init_matrix(matrix_t *mat) {
    for (int i = 0; i < mat->rows * mat->cols; i++) {
        mat->data[i] = rand() % 100;
    }
}

// Matrix multiplication - ijk order (poor cache performance)
void matrix_multiply_ijk(matrix_t *a, matrix_t *b, matrix_t *c) {
    for (int i = 0; i < a->rows; i++) {
        for (int j = 0; j < b->cols; j++) {
            c->data[i * c->cols + j] = 0;
            for (int k = 0; k < a->cols; k++) {
                c->data[i * c->cols + j] +=
```

```

        a->data[i * a->cols + k] * b->data[k * b->cols + j];
    }
}
}

// Matrix multiplication - ikj order (better cache performance)
void matrix_multiply_ikj(matrix_t *a, matrix_t *b, matrix_t *c) {
    // Initialize result matrix
    memset(c->data, 0, c->rows * c->cols * sizeof(int));

    for (int i = 0; i < a->rows; i++) {
        for (int k = 0; k < a->cols; k++) {
            int temp = a->data[i * a->cols + k];
            for (int j = 0; j < b->cols; j++) {
                c->data[i * c->cols + j] += temp * b->data[k * b->cols + j];
            }
        }
    }
}

// Performance comparison
void compare_matrix_multiplication() {
    matrix_t *a = create_matrix(MATRIX_SIZE, MATRIX_SIZE);
    matrix_t *b = create_matrix(MATRIX_SIZE, MATRIX_SIZE);
    matrix_t *c1 = create_matrix(MATRIX_SIZE, MATRIX_SIZE);
    matrix_t *c2 = create_matrix(MATRIX_SIZE, MATRIX_SIZE);

    init_matrix(a);
    init_matrix(b);

    struct timespec start, end;

    // Test ijk order
    clock_gettime(CLOCK_MONOTONIC, &start);
    matrix_multiply_ijk(a, b, c1);
    clock_gettime(CLOCK_MONOTONIC, &end);
    double time_ijk = get_time_diff(start, end);

    // Test ikj order
    clock_gettime(CLOCK_MONOTONIC, &start);
    matrix_multiply_ikj(a, b, c2);
    clock_gettime(CLOCK_MONOTONIC, &end);
    double time_ikj = get_time_diff(start, end);

    printf("Matrix multiplication performance comparison:\n");
    printf("Matrix size: %dx%d\n", MATRIX_SIZE, MATRIX_SIZE);

```

```

printf("ijk order time: %.6f seconds\n", time_ijk);
printf("ikj order time: %.6f seconds\n", time_ikj);
printf("Performance improvement: %.2fx\n", time_ijk / time_ikj);

// Verify results are the same
int diff_count = 0;
for (int i = 0; i < c1->rows * c1->cols; i++) {
    if (c1->data[i] != c2->data[i]) {
        diff_count++;
    }
}
printf("Results match: %s\n", diff_count == 0 ? "Yes" : "No");

free_matrix(a);
free_matrix(b);
free_matrix(c1);
free_matrix(c2);
}

double get_time_diff(struct timespec start, struct timespec end) {
    return (end.tv_sec - start.tv_sec) + (end.tv_nsec - start.tv_nsec) / 1e9;
}

int main() {
    srand(time(NULL));
    compare_matrix_multiplication();
    return 0;
}

```

Common Pitfalls to Avoid:

1. **Premature optimization:** Always profile before optimizing
2. **Algorithmic vs implementation optimization:** Focus on algorithmic improvements first
3. **Micro-optimizations:** Avoid micro-optimizations that hurt readability without significant gains
4. **Platform-specific optimizations:** Consider portability when optimizing
5. **Over-optimization:** Balance performance with maintainability and correctness

Best Practices:

1. **Profile-driven optimization:** Use profiling tools to identify real bottlenecks
2. **Algorithmic improvements:** Focus on better algorithms before low-level optimizations
3. **Memory locality:** Optimize for cache-friendly access patterns
4. **Compiler assistance:** Use compiler optimization flags and hints effectively
5. **Measurement and validation:** Always measure performance improvements and verify correctness

Complete these exercises to solidify your understanding of performance optimization in C. Each exercise builds upon the previous ones, gradually increasing in complexity.

Testing

Module 14: Testing and Quality Assurance Exercises

Exercise 1: Unit Testing Framework Implementation

Write a program that implements a simple unit testing framework: - Create basic assertion functions for different data types - Implement a test runner that executes test cases - Develop a reporting system for test results - Show how to organize test cases into test suites - Demonstrate proper test fixture setup and teardown

Requirements: - Implement assertions for common data types (int, float, string, etc.) - Create a simple test registration and execution system - Include detailed reporting with pass/fail counts and timing - Show how to handle expected failures and skip conditions - Provide clear examples of test organization

Exercise 2: Mocking and Stubbing Techniques

Create a program that demonstrates mocking and stubbing techniques: - Implement simple mock objects for function dependencies - Show how to replace real implementations with test doubles - Demonstrate proper use of function pointers for mocking - Implement a simple mocking framework - Show how to verify mock interactions

Requirements: - Include examples of state-based and behavior-based testing - Demonstrate proper isolation of units under test - Show how to handle complex dependencies with mocks - Implement proper mock verification mechanisms - Provide clear documentation of mocking techniques

Exercise 3: Integration Testing Implementation

Develop a program that implements integration testing: - Create tests that verify interactions between multiple components - Show how to set up and tear down integration test environments - Demonstrate proper test data management for integration tests - Implement tests for API interactions and data flow - Show how to handle external dependencies in integration tests

Requirements: - Include examples of component integration testing - Demonstrate proper test environment setup and cleanup - Show how to handle test data dependencies and isolation - Implement tests for complex system interactions - Provide clear documentation of integration testing strategies

Exercise 4: Code Coverage Analysis

Write a program that demonstrates code coverage analysis: - Implement instrumentation to track code execution - Create a coverage reporting system - Show how to measure different types of coverage (line, branch, path) - Demonstrate proper coverage goal setting and tracking - Show how to identify untested code paths

Requirements: - Include examples of coverage instrumentation techniques - Implement proper coverage measurement and reporting - Show how to handle complex control flow in coverage analysis - Demonstrate coverage-guided test development - Provide clear documentation of coverage analysis results

Exercise 5: Static Analysis and Linting

Create a program that demonstrates static analysis techniques: - Implement simple static analysis checks for common issues - Show how to integrate static analysis tools into the build process - Demonstrate proper configuration of static analysis rules - Implement custom static analysis checks - Show how to handle static analysis findings

Requirements: - Include examples of common static analysis checks - Demonstrate proper integration with development workflow - Show how to configure and customize analysis rules - Implement proper handling of false positives - Provide clear documentation of static analysis results

Exercise 6: Fuzz Testing Implementation

Write a program that implements fuzz testing techniques: - Create a simple fuzz testing framework - Implement input generation strategies for fuzz testing - Show how to monitor for crashes and unexpected behavior - Demonstrate proper fuzz test case reduction - Show how to integrate fuzz testing into the testing pipeline

Requirements: - Include examples of different fuzzing strategies - Implement proper crash detection and reporting - Show how to generate meaningful test inputs - Demonstrate proper handling of fuzz test results - Provide clear documentation of fuzz testing techniques

Exercise 7: Performance and Stress Testing

Create a program that implements performance and stress testing: - Develop tests that verify performance requirements - Implement stress tests that push system limits - Show how to measure and report performance metrics - Demonstrate proper test environment setup for performance testing - Show how to identify performance bottlenecks

Requirements: - Include examples of performance benchmarking - Implement proper metrics collection and reporting - Show how to handle performance test data analysis - Demonstrate proper test environment configuration - Provide clear documentation of performance test results

Exercise 8: Comprehensive Testing Strategy

Design a complete testing strategy that integrates all testing techniques: - Implement a comprehensive test suite covering all testing types - Create a testing pipeline that executes different test categories - Demonstrate proper test prioritization and execution ordering - Show how to maintain test quality and reliability - Provide clear documentation and reporting

Requirements: - Use modular design with clear test organization - Include comprehensive test coverage across all areas - Demonstrate proper test maintenance and evolution - Implement robust error handling throughout the testing system - Provide clear examples and test documentation

Solutions and Tips

Exercise 1 Solution Example:

```
#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <time.h>

// Simple unit testing framework

// Test result structure
typedef struct {
    int passed;
    int failed;
    int skipped;
    double total_time;
} test_result_t;

// Global test result tracker
static test_result_t g_test_result = {0};

// Assertion functions
#define ASSERT_TRUE(condition) \
    do { \
        if (!(condition)) { \
            printf("ASSERT_TRUE failed at %s:%d\n", __FILE__, __LINE__); \
            g_test_result.failed++; \
            return; \
        } \
        g_test_result.passed++; \
    } while(0)

#define ASSERT_FALSE(condition) \
    do { \
```

```

        if (condition) { \
            printf("ASSERT_FALSE failed at %s:%d\n", __FILE__, __LINE__); \
            g_test_result.failed++; \
            return; \
        } \
        g_test_result.passed++; \
    } while(0)

#define ASSERT_EQUAL(expected, actual) \
    do { \
        if ((expected) != (actual)) { \
            printf("ASSERT_EQUAL failed at %s:%d: expected %d, got %d\n", \
                __FILE__, __LINE__, (expected), (actual)); \
            g_test_result.failed++; \
            return; \
        } \
        g_test_result.passed++; \
    } while(0)

#define ASSERT_STRING_EQUAL(expected, actual) \
    do { \
        if (strcmp((expected), (actual)) != 0) { \
            printf("ASSERT_STRING_EQUAL failed at %s:%d: expected '%s', got '%s'\n", \
                __FILE__, __LINE__, (expected), (actual)); \
            g_test_result.failed++; \
            return; \
        } \
        g_test_result.passed++; \
    } while(0)

// Test function type
typedef void (*test_func_t)(void);

// Test case structure
typedef struct {
    const char *name;
    test_func_t function;
} test_case_t;

// Example functions to test
int add(int a, int b) {
    return a + b;
}

int subtract(int a, int b) {
    return a - b;
}

```

```

const char* get_hello_message(void) {
    return "Hello, World!";
}

// Test cases
void test_add_positive(void) {
    ASSERT_EQUAL(5, add(2, 3));
}

void test_add_negative(void) {
    ASSERT_EQUAL(-1, add(2, -3));
}

void test_subtract(void) {
    ASSERT_EQUAL(2, subtract(5, 3));
}

void test_string_message(void) {
    ASSERT_STRING_EQUAL("Hello, World!", get_hello_message());
}

// Test suite
test_case_t test_suite[] = {
    {"test_add_positive", test_add_positive},
    {"test_add_negative", test_add_negative},
    {"test_subtract", test_subtract},
    {"test_string_message", test_string_message},
    {NULL, NULL} // Sentinel
};

// Test runner
void run_tests(void) {
    printf("Running tests...\n");

    struct timespec start, end;
    clock_gettime(CLOCK_MONOTONIC, &start);

    for (int i = 0; test_suite[i].name != NULL; i++) {
        printf("Running %s... ", test_suite[i].name);

        // Reset test counters for each test
        int initial_passed = g_test_result.passed;
        int initial_failed = g_test_result.failed;

        test_suite[i].function();
    }
}

```

```

        if (g_test_result.failed > initial_failed) {
            printf("FAILED\n");
        } else {
            printf("PASSED\n");
        }
    }

    clock_gettime(CLOCK_MONOTONIC, &end);
    g_test_result.total_time = (end.tv_sec - start.tv_sec) +
                              (end.tv_nsec - start.tv_nsec) / 1e9;

    // Print summary
    printf("\nTest Results:\n");
    printf("Passed: %d\n", g_test_result.passed);
    printf("Failed: %d\n", g_test_result.failed);
    printf("Skipped: %d\n", g_test_result.skipped);
    printf("Total time: %.6f seconds\n", g_test_result.total_time);

    if (g_test_result.failed == 0) {
        printf("All tests passed!\n");
    } else {
        printf("Some tests failed!\n");
    }
}

int main() {
    run_tests();
    return g_test_result.failed > 0 ? 1 : 0;
}

```

Exercise 2 Solution Example:

```

#include <stdio.h>
#include <stdlib.h>
#include <string.h>

// Mocking framework example

// Real database interface
typedef struct {
    int (*connect)(const char *host, int port);
    int (*query)(const char *sql, char **result);
    void (*disconnect)(void);
} database_interface_t;

```

```

// Real database implementation
static int real_db_connect(const char *host, int port) {
    printf("Connecting to database at %s:%d\n", host, port);
    return 1; // Success
}

static int real_db_query(const char *sql, char **result) {
    printf("Executing query: %s\n", sql);
    *result = strdup("Mock query result");
    return 1; // Success
}

static void real_db_disconnect(void) {
    printf("Disconnecting from database\n");
}

database_interface_t real_database = {
    .connect = real_db_connect,
    .query = real_db_query,
    .disconnect = real_db_disconnect
};

// Mock database implementation
static int mock_connect_calls = 0;
static int mock_query_calls = 0;
static int mock_disconnect_calls = 0;
static int mock_connect_result = 1;
static int mock_query_result = 1;

static int mock_db_connect(const char *host, int port) {
    mock_connect_calls++;
    printf("Mock connect called (%d times)\n", mock_connect_calls);
    return mock_connect_result;
}

static int mock_db_query(const char *sql, char **result) {
    mock_query_calls++;
    printf("Mock query called (%d times): %s\n", mock_query_calls, sql);
    *result = strdup("Mock result");
    return mock_query_result;
}

static void mock_db_disconnect(void) {
    mock_disconnect_calls++;
    printf("Mock disconnect called (%d times)\n", mock_disconnect_calls);
}

```

```

database_interface_t mock_database = {
    .connect = mock_db_connect,
    .query = mock_db_query,
    .disconnect = mock_db_disconnect
};

// Function under test
int fetch_user_data(database_interface_t *db, const char *user_id, char **data) {
    if (!db->connect("localhost", 5432)) {
        return 0;
    }

    char sql[256];
    snprintf(sql, sizeof(sql), "SELECT * FROM users WHERE id = '%s'", user_id);

    int result = db->query(sql, data);
    db->disconnect();

    return result;
}

// Test with mock
void test_fetch_user_data_success(void) {
    // Reset mock counters
    mock_connect_calls = 0;
    mock_query_calls = 0;
    mock_disconnect_calls = 0;
    mock_connect_result = 1;
    mock_query_result = 1;

    char *data = NULL;
    int result = fetch_user_data(&mock_database, "123", &data);

    // Verify results
    if (result != 1) {
        printf("Test failed: Expected success, got failure\n");
        return;
    }

    if (mock_connect_calls != 1) {
        printf("Test failed: Expected 1 connect call, got %d\n", mock_connect_calls);
        return;
    }

    if (mock_query_calls != 1) {
        printf("Test failed: Expected 1 query call, got %d\n", mock_query_calls);
        return;
    }
}

```

```

    }

    if (mock_disconnect_calls != 1) {
        printf("Test failed: Expected 1 disconnect call, got %d\n", mock_disconnect_calls);
        return;
    }

    printf("Test passed: All mock interactions verified\n");
    free(data);
}

void test_fetch_user_data_connection_failure(void) {
    // Reset mock counters
    mock_connect_calls = 0;
    mock_query_calls = 0;
    mock_disconnect_calls = 0;
    mock_connect_result = 0; // Simulate connection failure

    char *data = NULL;
    int result = fetch_user_data(&mock_database, "123", &data);

    // Verify results
    if (result != 0) {
        printf("Test failed: Expected failure, got success\n");
        return;
    }

    if (mock_connect_calls != 1) {
        printf("Test failed: Expected 1 connect call, got %d\n", mock_connect_calls);
        return;
    }

    // Should not call query or disconnect on connection failure
    if (mock_query_calls != 0) {
        printf("Test failed: Expected 0 query calls, got %d\n", mock_query_calls);
        return;
    }

    if (mock_disconnect_calls != 0) {
        printf("Test failed: Expected 0 disconnect calls, got %d\n", mock_disconnect_calls);
        return;
    }

    printf("Test passed: Connection failure handled correctly\n");
}

int main() {

```

```

    printf("Testing with mock database...\n");
    test_fetch_user_data_success();
    test_fetch_user_data_connection_failure();

    printf("\nFor comparison, running with real database:\n");
    char *data = NULL;
    fetch_user_data(&real_database, "123", &data);
    free(data);

    return 0;
}

```

Common Pitfalls to Avoid:

1. **Testing implementation details:** Focus on behavior rather than implementation
2. **Insufficient test coverage:** Ensure adequate coverage of edge cases and error conditions
3. **Flaky tests:** Write deterministic tests that produce consistent results
4. **Over-mocking:** Don't mock everything; focus on complex or external dependencies
5. **Ignoring test maintenance:** Keep tests up-to-date with code changes

Best Practices:

1. **Test pyramid:** Follow the testing pyramid (unit > integration > end-to-end)
2. **Descriptive test names:** Use clear, descriptive names for test cases
3. **Isolated tests:** Ensure tests don't depend on each other or shared state
4. **Fast feedback:** Keep tests fast to encourage frequent execution
5. **Continuous testing:** Integrate testing into the development workflow

Complete these exercises to solidify your understanding of testing and quality assurance in C. Each exercise builds upon the previous ones, gradually increasing in complexity.

Advanced

Module 15: Advanced Topics Exercises

Exercise 1: Metaprogramming with Macros

Write a program that demonstrates advanced macro techniques: - Implement complex macros with conditional compilation - Create macros that generate code based on parameters - Demonstrate proper use of token concatenation and stringification - Show how to implement compile-time assertions - Implement a simple domain-specific language using macros

Requirements: - Include examples of advanced preprocessor features - Demonstrate proper macro hygiene and scoping - Show how to handle macro expansion edge cases - Implement useful utility macros for common patterns - Provide clear documentation of macro behavior

Exercise 2: Generic Programming Techniques

Create a program that demonstrates generic programming in C: - Implement type-generic data structures using void pointers - Show how to use function pointers for polymorphic behavior - Demonstrate proper type safety in generic code - Implement a simple container library with generic operations - Show how to handle memory management in generic code

Requirements: - Include examples of generic containers (lists, trees, etc.) - Demonstrate proper error handling in generic functions - Show how to implement type-safe generic interfaces - Implement proper memory management for generic data - Provide clear documentation of generic programming techniques

Exercise 3: Coroutines and Cooperative Multitasking

Develop a program that implements coroutines and cooperative multitasking: - Create a simple coroutine framework - Implement yield and resume mechanisms - Show how to manage coroutine state and context - Demonstrate proper coroutine scheduling - Implement a simple cooperative multitasking system

Requirements: - Include examples of coroutine usage patterns - Demonstrate proper context switching mechanisms - Show how to handle coroutine lifecycle management - Implement proper error handling in coroutine systems - Provide clear documentation of coroutine concepts

Exercise 4: Reflection and Introspection

Write a program that demonstrates reflection and introspection techniques: - Implement a simple type system with runtime type information - Create functions for inspecting data structure layouts - Show how to implement serialization and deserialization - Demonstrate proper metadata generation and usage - Implement a simple object model with reflection capabilities

Requirements: - Include examples of runtime type inspection - Demonstrate proper serialization techniques - Show how to implement metadata-driven programming - Implement proper error handling for reflection operations - Provide clear documentation of reflection mechanisms

Exercise 5: Template Metaprogramming Simulation

Create a program that simulates template metaprogramming techniques: - Implement compile-time computation and type manipulation - Show how to use preprocessor metaprogramming - Demonstrate proper use of constexpr-like constructs - Implement a simple compile-time data structure library - Show how to optimize code through compile-time evaluation

Requirements: - Include examples of compile-time algorithms - Demonstrate proper use of preprocessor for metaprogramming - Show how to implement type traits and compile-time checks - Implement proper error reporting for metaprogramming errors - Provide clear documentation of metaprogramming techniques

Exercise 6: Memory Model and Concurrency

Write a program that demonstrates advanced memory model and concurrency concepts: - Implement lock-free data structures - Show how to use atomic operations effectively - Demonstrate proper memory ordering and synchronization - Implement a simple memory model simulator - Show how to handle race conditions and data races

Requirements: - Include examples of atomic operations and memory barriers - Demonstrate proper use of concurrent data structures - Show how to implement lock-free algorithms - Implement proper testing for concurrent code - Provide clear documentation of memory model concepts

Exercise 7: Language Interoperability

Create a program that demonstrates language interoperability: - Implement a C library that can be called from other languages - Show how to interface with C++ code - Demonstrate proper use of foreign function interfaces - Implement a simple binding generator (bonus) - Show how to handle different calling conventions

Requirements: - Include examples of C API design for language interoperability - Demonstrate proper handling of data type conversions - Show how to implement proper error handling across language boundaries - Implement proper resource management for interop scenarios - Provide clear documentation of interoperability techniques

Exercise 8: Comprehensive Advanced Application

Design a complete application that integrates all advanced concepts: - Implement a sophisticated system using multiple advanced techniques - Create a domain-specific language or framework - Demonstrate proper architecture and design patterns - Include comprehensive testing and validation - Provide clear documentation and examples

Requirements: - Use modular design with clear separation of concerns - Include comprehensive error handling and recovery mechanisms - Demonstrate proper resource management throughout the application - Implement robust testing and validation procedures - Provide clear examples and documentation

Solutions and Tips

Exercise 1 Solution Example:

```
#include <stdio.h>
#include <stdlib.h>
#include <string.h>

// Advanced macro techniques

// Stringification macro
#define STRINGIFY(x) #x
#define TOSTRING(x) STRINGIFY(x)

// Token concatenation macros
#define CONCAT(a, b) a##b
#define PASTE(a, b) CONCAT(a, b)

// Conditional compilation based on parameters
#define DEBUG_LEVEL 2

#if DEBUG_LEVEL >= 2
    #define DEBUG_PRINT(fmt, ...) \
        fprintf(stderr, "[DEBUG] %s:%d: " fmt "\n", __FILE__, __LINE__, ##__VA_ARGS__)
#elif DEBUG_LEVEL == 1
    #define DEBUG_PRINT(fmt, ...) \
        fprintf(stderr, "[INFO] " fmt "\n", ##__VA_ARGS__)
#else
    #define DEBUG_PRINT(fmt, ...) do {} while(0)
#endif

// Compile-time assertion
#define STATIC_ASSERT(condition, message) \
```

```

typedef char PASTE(static_assert_, __LINE__)[(condition) ? 1 : -1]

// Macro for generating function declarations
#define DECLARE_FUNCTION(return_type, name, ...) \
    return_type name(__VA_ARGS__);

// Macro for generating function definitions with logging
#define DEFINE_FUNCTION(return_type, name, params, body) \
    return_type name params { \
        DEBUG_PRINT("Entering function %s", TOSTRING(name)); \
        body \
        DEBUG_PRINT("Exiting function %s", TOSTRING(name)); \
    }

// Macro for creating a simple DSL for defining structures
#define BEGIN_STRUCT(name) typedef struct {
#define FIELD(type, name) type name;
#define END_STRUCT(name) } name##_t;

// Example usage of the DSL
BEGIN_STRUCT(person)
    FIELD(char*, name)
    FIELD(int, age)
    FIELD(float, height)
END_STRUCT(person)

// Generic max macro
#define MAX(a, b) ((a) > (b) ? (a) : (b))
#define MAX3(a, b, c) MAX(MAX(a, b), c)

// Type-generic print macro
#define PRINT_VALUE(x) _Generic((x), \
    int: printf("int: %d\n", (x)), \
    float: printf("float: %f\n", (x)), \
    double: printf("double: %f\n", (x)), \
    char*: printf("string: %s\n", (x)), \
    default: printf("unknown type\n"))

// Compile-time assertions
STATIC_ASSERT(sizeof(int) >= 4, "int must be at least 4 bytes");
STATIC_ASSERT(sizeof(void*) == 8 || sizeof(void*) == 4, "Pointer size must be 4 or 8 bytes")

int main() {
    // Test stringification
    printf("File: %s\n", __FILE__);
    printf("Line: %d\n", __LINE__);
    printf("Function: %s\n", TOSTRING(main));
}

```

```

// Test debug printing
DEBUG_PRINT("This is a debug message with value %d", 42);

// Test generic max
int max_int = MAX3(10, 20, 15);
printf("Max of 10, 20, 15: %d\n", max_int);

// Test type-generic print
int a = 42;
float b = 3.14f;
char *c = "Hello";
PRINT_VALUE(a);
PRINT_VALUE(b);
PRINT_VALUE(c);

// Test DSL-generated structure
person_t person = {"Alice", 30, 5.5f};
printf("Person: %s, Age: %d, Height: %.1f\n",
       person.name, person.age, person.height);

return 0;
}

```

Exercise 2 Solution Example:

```

#include <stdio.h>
#include <stdlib.h>
#include <string.h>

// Generic programming techniques

// Generic list node
typedef struct list_node {
    void *data;
    struct list_node *next;
} list_node_t;

// Generic list structure
typedef struct {
    list_node_t *head;
    size_t data_size;
    int (*compare)(const void *a, const void *b);
    void (*print)(const void *data);
} generic_list_t;

```

```

// Create a new generic list
generic_list_t* list_create(size_t data_size,
                           int (*compare)(const void *a, const void *b),
                           void (*print)(const void *data)) {
    generic_list_t *list = malloc(sizeof(generic_list_t));
    if (list) {
        list->head = NULL;
        list->data_size = data_size;
        list->compare = compare;
        list->print = print;
    }
    return list;
}

// Add an element to the list
int list_add(generic_list_t *list, const void *data) {
    list_node_t *node = malloc(sizeof(list_node_t));
    if (!node) return 0;

    node->data = malloc(list->data_size);
    if (!node->data) {
        free(node);
        return 0;
    }

    memcpy(node->data, data, list->data_size);
    node->next = list->head;
    list->head = node;
    return 1;
}

// Find an element in the list
void* list_find(generic_list_t *list, const void *key) {
    for (list_node_t *node = list->head; node; node = node->next) {
        if (list->compare(node->data, key) == 0) {
            return node->data;
        }
    }
    return NULL;
}

// Print all elements in the list
void list_print(generic_list_t *list) {
    for (list_node_t *node = list->head; node; node = node->next) {
        list->print(node->data);
    }
}

```

```

// Free the list
void list_destroy(generic_list_t *list) {
    list_node_t *node = list->head;
    while (node) {
        list_node_t *next = node->next;
        free(node->data);
        free(node);
        node = next;
    }
    free(list);
}

// Comparison functions for different types
int int_compare(const void *a, const void *b) {
    int ia = *(const int*)a;
    int ib = *(const int*)b;
    return (ia > ib) - (ia < ib);
}

int string_compare(const void *a, const void *b) {
    return strcmp(*(const char**)a, *(const char**)b);
}

// Print functions for different types
void int_print(const void *data) {
    printf("%d ", *(const int*)data);
}

void string_print(const void *data) {
    printf("%s ", *(const char**)data);
}

int main() {
    // Test with integers
    printf("Testing integer list:\n");
    generic_list_t *int_list = list_create(sizeof(int), int_compare, int_print);

    int values[] = {10, 20, 30, 40, 50};
    for (int i = 0; i < 5; i++) {
        list_add(int_list, &values[i]);
    }

    printf("List contents: ");
    list_print(int_list);
    printf("\n");
}

```



```

    int search_key = 30;
    int *found = (int*)list_find(int_list, &search_key);
    if (found) {
        printf("Found: %d\n", *found);
    } else {
        printf("Not found: %d\n", search_key);
    }

    list_destroy(int_list);

    // Test with strings
    printf("\nTesting string list:\n");
    generic_list_t *string_list = list_create(sizeof(char*), string_compare, string_print);

    const char *strings[] = {"apple", "banana", "cherry", "date", "elderberry"};
    for (int i = 0; i < 5; i++) {
        list_add(string_list, &strings[i]);
    }

    printf("List contents: ");
    list_print(string_list);
    printf("\n");

    const char *search_string = "cherry";
    char **found_string = (char**)list_find(string_list, &search_string);
    if (found_string) {
        printf("Found: %s\n", *found_string);
    } else {
        printf("Not found: %s\n", search_string);
    }

    list_destroy(string_list);

    return 0;
}

```

Common Pitfalls to Avoid:

1. **Macro complexity:** Avoid overly complex macros that hurt readability
2. **Type safety:** Maintain type safety when using generic programming techniques
3. **Memory management:** Properly handle memory in generic and advanced code
4. **Portability:** Consider platform differences in advanced features
5. **Debugging difficulty:** Advanced techniques can make debugging more challenging

Best Practices:

1. **Clear documentation:** Document advanced techniques thoroughly
2. **Gradual complexity:** Introduce advanced concepts incrementally
3. **Error handling:** Implement robust error handling for advanced features
4. **Testing:** Include comprehensive tests for advanced functionality
5. **Maintainability:** Balance advanced features with code maintainability

Complete these exercises to solidify your understanding of advanced topics in C. Each exercise builds upon the previous ones, gradually increasing in complexity.

Projects

Advanced

Advanced Projects

These projects are designed for experienced C programmers who want to tackle complex challenges involving system programming, networking, multithreading, and advanced data structures.

Project 1: Concurrent Web Server

Description

Create a multithreaded web server that can handle multiple client requests concurrently and serve static files.

Learning Objectives

- Network programming with sockets
- Multithreading and synchronization
- HTTP protocol implementation
- File I/O operations
- Error handling and logging

Requirements

1. Handle multiple concurrent client connections
2. Serve static HTML, CSS, JavaScript, and image files
3. Implement basic HTTP/1.1 protocol
4. Support GET and POST requests
5. Handle 404 errors gracefully
6. Log requests and responses
7. Support basic authentication (optional)
8. Implement connection pooling

Implementation Steps

1. Set up socket programming for server
2. Implement HTTP request parsing
3. Create thread pool for handling connections
4. Add file serving functionality
5. Implement error handling and logging

6. Add support for different HTTP methods
7. Create configuration system
8. Test with various clients

Sample Code Structure

```
#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <unistd.h>
#include <sys/socket.h>
#include <netinet/in.h>
#include <arpa/inet.h>
#include <pthread.h>
#include <time.h>
#include <fcntl.h>
#include <sys/stat.h>
#include <dirent.h>

#define PORT 8080
#define MAX_CLIENTS 100
#define BUFFER_SIZE 4096
#define THREAD_POOL_SIZE 10

typedef struct {
    int socket;
    struct sockaddr_in address;
} client_info_t;

typedef struct {
    pthread_t threads[THREAD_POOL_SIZE];
    int thread_count;
    pthread_mutex_t queue_mutex;
    pthread_cond_t queue_cond;
    client_info_t *queue;
    int queue_size;
    int queue_front;
    int queue_rear;
    int shutdown;
} thread_pool_t;

// Function prototypes
void init_thread_pool(thread_pool_t *pool);
void destroy_thread_pool(thread_pool_t *pool);
void *worker_thread(void *arg);
void enqueue_client(thread_pool_t *pool, client_info_t client);
```

```

client_info_t dequeue_client(thread_pool_t *pool);
void handle_client(client_info_t client);
void send_response(int client_socket, const char *status, const char *content_type, const ch
void send_file(int client_socket, const char *filepath);
void log_request(const char *method, const char *path, int status_code);
int parse_http_request(const char *request, char *method, char *path, char *version);
char* get_content_type(const char *filepath);
void *server_thread(void *arg);

int main() {
    int server_fd;
    struct sockaddr_in address;
    int opt = 1;
    int addrlen = sizeof(address);
    thread_pool_t pool;

    // Initialize thread pool
    init_thread_pool(&pool);

    // Create socket
    if ((server_fd = socket(AF_INET, SOCK_STREAM, 0)) == 0) {
        perror("socket failed");
        exit(EXIT_FAILURE);
    }

    // Set socket options
    if (setsockopt(server_fd, SOL_SOCKET, SO_REUSEADDR | SO_REUSEPORT,
        &opt, sizeof(opt))) {
        perror("setsockopt failed");
        exit(EXIT_FAILURE);
    }

    address.sin_family = AF_INET;
    address.sin_addr.s_addr = INADDR_ANY;
    address.sin_port = htons(PORT);

    // Bind socket
    if (bind(server_fd, (struct sockaddr *)&address, sizeof(address)) < 0) {
        perror("bind failed");
        exit(EXIT_FAILURE);
    }

    // Listen for connections
    if (listen(server_fd, 10) < 0) {
        perror("listen failed");
        exit(EXIT_FAILURE);
    }
}

```

```

printf("Server listening on port %d\n", PORT);

// Main server loop
while (1) {
    client_info_t client;
    if ((client.socket = accept(server_fd, (struct sockaddr *)&client.address,
                                (socklen_t*)&addrlen)) < 0) {
        perror("accept failed");
        continue;
    }

    // Add client to thread pool queue
    enqueue_client(&pool, client);
}

// Cleanup
destroy_thread_pool(&pool);
close(server_fd);
return 0;
}

// Implement all functions here

```

Common Pitfalls to Avoid

1. Thread safety issues with shared resources
2. Memory leaks from improper resource management
3. Not handling partial reads/writes correctly
4. Buffer overflows in request parsing
5. Not properly closing socket connections
6. Race conditions in thread pool implementation

Best Practices

1. Use proper synchronization primitives (mutexes, condition variables)
2. Always validate input and handle errors gracefully
3. Implement proper resource cleanup
4. Use connection pooling for efficiency
5. Log important events for debugging
6. Handle edge cases in HTTP protocol implementation
7. Use appropriate buffer sizes to prevent overflows

Project 2: Memory Allocator

Description

Implement a custom memory allocator that provides dynamic memory management similar to malloc/free.

Learning Objectives

- Low-level memory management
- Pointer arithmetic
- Data structure design
- Performance optimization
- System programming concepts

Requirements

1. Implement malloc, free, calloc, and realloc functions
2. Handle memory fragmentation
3. Support different allocation strategies (first-fit, best-fit, worst-fit)
4. Provide memory usage statistics
5. Handle alignment requirements
6. Detect memory leaks and corruption (optional)
7. Support multi-threading (optional)
8. Implement garbage collection (optional)

Implementation Steps

1. Design internal data structures for memory blocks
2. Implement basic allocation and deallocation
3. Add different allocation strategies
4. Handle memory fragmentation
5. Implement memory statistics
6. Add alignment support
7. Test with various allocation patterns
8. Optimize for performance

Sample Code Structure

```
#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <unistd.h>
#include <pthread.h>
```

```

#define ALIGNMENT 8
#define ALIGN(size) (((size) + (ALIGNMENT-1)) & ~(ALIGNMENT-1))
#define BLOCK_SIZE sizeof(block_header_t)

typedef enum {
    ALLOC_FIRST_FIT,
    ALLOC_BEST_FIT,
    ALLOC_WORST_FIT
} alloc_strategy_t;

typedef struct block_header {
    size_t size;
    int free;
    struct block_header *next;
    struct block_header *prev;
} block_header_t;

typedef struct {
    block_header_t *head;
    size_t total_size;
    size_t used_size;
    size_t free_size;
    alloc_strategy_t strategy;
    pthread_mutex_t mutex;
} memory_pool_t;

static memory_pool_t g_pool = {0};

// Function prototypes
void init_memory_pool(size_t size, alloc_strategy_t strategy);
void destroy_memory_pool();
void *my_malloc(size_t size);
void my_free(void *ptr);
void *my_calloc(size_t nmemb, size_t size);
void *my_realloc(void *ptr, size_t size);
void print_memory_stats();
block_header_t *find_free_block(size_t size);
block_header_t *split_block(block_header_t *block, size_t size);
void coalesce_blocks();
size_t align_size(size_t size);

void *my_malloc(size_t size) {
    if (size == 0) return NULL;

    pthread_mutex_lock(&g_pool.mutex);

```

```

size_t aligned_size = align_size(size);
block_header_t *block = find_free_block(aligned_size);

if (block) {
    // Split block if it's larger than needed
    if (block->size > aligned_size + BLOCK_SIZE + ALIGNMENT) {
        block = split_block(block, aligned_size);
    }

    block->free = 0;
    g_pool.used_size += block->size;
    g_pool.free_size -= block->size;

    pthread_mutex_unlock(&g_pool.mutex);
    return (void*)(block + 1);
}

pthread_mutex_unlock(&g_pool.mutex);
return NULL; // Out of memory
}

void my_free(void *ptr) {
    if (!ptr) return;

    pthread_mutex_lock(&g_pool.mutex);

    block_header_t *block = (block_header_t*)ptr - 1;
    block->free = 1;
    g_pool.used_size -= block->size;
    g_pool.free_size += block->size;

    // Coalesce adjacent free blocks
    coalesce_blocks();

    pthread_mutex_unlock(&g_pool.mutex);
}

// Implement all other functions here

```

Common Pitfalls to Avoid

1. Memory leaks from improper block management
2. Buffer overflows in block headers
3. Not handling alignment requirements
4. Incorrect coalescing of free blocks
5. Thread safety issues in multi-threaded version

6. Not validating pointers in `free()` function

Best Practices

1. Use proper data structures for efficient block management
2. Handle edge cases in allocation and deallocation
3. Implement proper alignment for performance
4. Use synchronization for thread safety
5. Validate all input pointers
6. Provide detailed memory statistics
7. Test with various allocation patterns

Project 3: Real-time Data Processing System

Description

Create a system that processes real-time data streams with configurable filters and transformations.

Learning Objectives

- Multithreading and inter-thread communication
- Data pipeline design
- Performance optimization
- Configuration management
- Error handling in real-time systems

Requirements

1. Process multiple data streams concurrently
2. Apply configurable filters and transformations
3. Support real-time data ingestion
4. Provide monitoring and statistics
5. Handle backpressure and buffering
6. Support plugin architecture for custom processors
7. Implement fault tolerance
8. Provide REST API for configuration (optional)

Implementation Steps

1. Design data pipeline architecture
2. Implement data source connectors
3. Create processor framework
4. Add filtering and transformation capabilities
5. Implement threading and synchronization
6. Add monitoring and statistics
7. Handle backpressure and buffering
8. Test with real-world data scenarios

Sample Code Structure

```
#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <pthread.h>
#include <unistd.h>
#include <time.h>
#include <sys/queue.h>

#define MAX_PROCESSORS 100
#define BUFFER_SIZE 1024
#define MAX_PIPELINE_DEPTH 10

typedef struct {
    char *data;
    size_t size;
    time_t timestamp;
    int priority;
} data_packet_t;

typedef enum {
    FILTER_TYPE_NONE,
    FILTER_TYPE_REGEX,
    FILTER_TYPE_RANGE,
    FILTER_TYPE_CUSTOM
} filter_type_t;

typedef struct {
    filter_type_t type;
    void *config;
    int (*apply)(data_packet_t *packet, void *config);
} filter_t;

typedef struct {
```

```

    char name[50];
    int (*process)(data_packet_t *input, data_packet_t *output, void *config);
    void *config;
    filter_t *filters;
    int filter_count;
} processor_t;

typedef struct pipeline_node {
    processor_t *processor;
    struct pipeline_node *next;
} pipeline_node_t;

typedef struct {
    pipeline_node_t *head;
    int depth;
    pthread_mutex_t mutex;
    pthread_cond_t data_available;
    data_packet_t buffer[BUFFER_SIZE];
    int head_idx;
    int tail_idx;
    int count;
} pipeline_t;

typedef struct {
    char source_name[50];
    pthread_t thread;
    int running;
    pipeline_t *pipeline;
    void *source_config;
    int (*read_data)(data_packet_t *packet, void *config);
} data_source_t;

// Function prototypes
pipeline_t* create_pipeline();
void destroy_pipeline(pipeline_t *pipeline);
int add_processor(pipeline_t *pipeline, processor_t *processor);
int remove_processor(pipeline_t *pipeline, const char *name);
void process_data(pipeline_t *pipeline, data_packet_t *input);
data_source_t* create_data_source(const char *name,
                                int (*read_func)(data_packet_t*, void*),
                                void *config);
void start_data_source(data_source_t *source, pipeline_t *pipeline);
void stop_data_source(data_source_t *source);
void *data_source_thread(void *arg);
int apply_filters(data_packet_t *packet, filter_t *filters, int filter_count);
void monitor_pipeline(pipeline_t *pipeline);

```

```

int main() {
    // Create pipeline
    pipeline_t *pipeline = create_pipeline();
    if (!pipeline) {
        fprintf(stderr, "Failed to create pipeline\n");
        return 1;
    }

    // Add processors to pipeline
    // ... add processors ...

    // Create data sources
    data_source_t *source1 = create_data_source("sensor1", read_sensor_data, NULL);
    data_source_t *source2 = create_data_source("network", read_network_data, NULL);

    // Start data sources
    start_data_source(source1, pipeline);
    start_data_source(source2, pipeline);

    // Monitor pipeline
    while (1) {
        monitor_pipeline(pipeline);
        sleep(5);
    }

    // Cleanup
    stop_data_source(source1);
    stop_data_source(source2);
    destroy_pipeline(pipeline);

    return 0;
}

// Implement all functions here

```

Common Pitfalls to Avoid

1. Deadlocks in threading implementation
2. Buffer overflows in data packets
3. Not handling backpressure properly
4. Memory leaks in dynamic data structures
5. Race conditions in shared data access
6. Not validating data packet integrity

Best Practices

1. Use proper synchronization primitives
2. Implement efficient buffering strategies
3. Handle backpressure with appropriate mechanisms
4. Validate all data packets
5. Provide detailed monitoring and logging
6. Use configuration files for flexibility
7. Implement graceful shutdown procedures

Project 4: Embedded System Simulator

Description

Create a simulator for an embedded system with hardware components like GPIO, UART, timers, and interrupts.

Learning Objectives

- Embedded systems programming concepts
- Hardware simulation
- Interrupt handling
- Real-time constraints
- Memory-mapped I/O

Requirements

1. Simulate GPIO pins with input/output capabilities
2. Implement UART communication
3. Simulate timer functionality
4. Handle interrupts and interrupt priorities
5. Support memory-mapped I/O
6. Provide debugging interface
7. Simulate real-time constraints
8. Support multiple peripheral devices

Implementation Steps

1. Design hardware simulation architecture
2. Implement GPIO simulation
3. Add UART communication simulation
4. Create timer simulation
5. Implement interrupt controller
6. Add memory-mapped I/O support

7. Create debugging interface
8. Test with sample embedded applications

Sample Code Structure

```
#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <pthread.h>
#include <unistd.h>
#include <time.h>

#define GPIO_PINS 32
#define UART_BUFFER_SIZE 256
#define MAX_INTERRUPTS 256

// Memory-mapped addresses
#define GPIO_BASE_ADDR 0x40000000
#define UART_BASE_ADDR 0x40001000
#define TIMER_BASE_ADDR 0x40002000
#define INTERRUPT_ADDR 0x40003000

typedef enum {
    GPIO_MODE_INPUT,
    GPIO_MODE_OUTPUT,
    GPIO_MODE_AF,
    GPIO_MODE_ANALOG
} gpio_mode_t;

typedef struct {
    unsigned int moder;        // Mode register
    unsigned int otyper;       // Output type register
    unsigned int ospeedr;      // Output speed register
    unsigned int pupdr;        // Pull-up/pull-down register
    unsigned int idr;          // Input data register
    unsigned int odr;          // Output data register
    unsigned int bsrr;         // Bit set/reset register
    unsigned int lckr;         // Configuration lock register
    unsigned int afr[2];       // Alternate function registers
} gpio_registers_t;

typedef struct {
    unsigned int sr;          // Status register
    unsigned int dr;          // Data register
    unsigned int brr;         // Baud rate register
    unsigned int cr1;         // Control register 1
```

```

    unsigned int cr2;        // Control register 2
    unsigned int cr3;        // Control register 3
    unsigned int gtp;        // Guard time and prescaler register
} uart_registers_t;

typedef struct {
    unsigned int cr1;        // Control register 1
    unsigned int cr2;        // Control register 2
    unsigned int smcr;       // Slave mode control register
    unsigned int dier;       // DMA/interrupt enable register
    unsigned int sr;         // Status register
    unsigned int egr;        // Event generation register
    unsigned int ccmr1;      // Capture/compare mode register 1
    unsigned int ccmr2;      // Capture/compare mode register 2
    unsigned int ccer;       // Capture/compare enable register
    unsigned int cnt;        // Counter
    unsigned int psc;        // Prescaler
    unsigned int arr;        // Auto-reload register
} timer_registers_t;

typedef struct {
    unsigned int isr;        // Interrupt set register
    unsigned int ier;        // Interrupt enable register
    unsigned int icr;        // Interrupt clear register
    unsigned int ipr[8];     // Interrupt priority registers
} interrupt_registers_t;

typedef struct {
    gpio_registers_t gpio;
    uart_registers_t uart;
    timer_registers_t timer;
    interrupt_registers_t interrupt;
    unsigned char *memory;
    size_t memory_size;
    pthread_t cpu_thread;
    int running;
    unsigned long long cycles;
    time_t start_time;
} embedded_system_t;

// Function prototypes
embedded_system_t* create_embedded_system(size_t memory_size);
void destroy_embedded_system(embedded_system_t *system);
void* cpu_thread(void *arg);
unsigned int read_memory(embedded_system_t *system, unsigned int addr);
void write_memory(embedded_system_t *system, unsigned int addr, unsigned int value);

```

```

void gpio_write_pin(embedded_system_t *system, int pin, int value);
int gpio_read_pin(embedded_system_t *system, int pin);
void uart_transmit(embedded_system_t *system, unsigned char data);
unsigned char uart_receive(embedded_system_t *system);
void timer_start(embedded_system_t *system);
void timer_stop(embedded_system_t *system);
void interrupt_set(embedded_system_t *system, int irq);
void interrupt_clear(embedded_system_t *system, int irq);
void handle_interrupt(embedded_system_t *system, int irq);
void load_program(embedded_system_t *system, const char *filename);
void run_system(embedded_system_t *system);
void stop_system(embedded_system_t *system);

int main() {
    // Create embedded system
    embedded_system_t *system = create_embedded_system(64 * 1024); // 64KB memory
    if (!system) {
        fprintf(stderr, "Failed to create embedded system\n");
        return 1;
    }

    // Load sample program
    load_program(system, "sample_program.bin");

    // Run system
    run_system(system);

    // Let it run for a while
    sleep(10);

    // Stop and cleanup
    stop_system(system);
    destroy_embedded_system(system);

    return 0;
}

// Implement all functions here

```

Common Pitfalls to Avoid

1. Incorrect memory-mapped I/O implementation
2. Race conditions in interrupt handling
3. Not simulating real-time constraints properly
4. Memory leaks in dynamic structures
5. Buffer overflows in peripheral simulations

6. Incorrect bit manipulation in hardware registers

Best Practices

1. Use proper data structures for hardware registers
2. Implement accurate timing simulation
3. Handle interrupts with proper priority levels
4. Validate all memory accesses
5. Provide debugging interfaces for development
6. Use bit manipulation macros for register access
7. Implement proper error handling for hardware operations

Tips for Success

1. **Understand the Domain:** Research the problem domain thoroughly before implementation
2. **Design First:** Create detailed designs and architectures before coding
3. **Incremental Development:** Build and test components incrementally
4. **Performance Considerations:** Profile and optimize critical paths
5. **Error Handling:** Implement comprehensive error handling and recovery
6. **Testing:** Create thorough test cases including edge cases
7. **Documentation:** Document complex algorithms and architectures
8. **Security:** Consider security implications in networked applications
9. **Scalability:** Design for scalability from the beginning
10. **Maintainability:** Write clean, modular code that's easy to maintain

These advanced projects will challenge you to apply sophisticated C programming techniques while solving complex real-world problems. Focus on robustness, performance, and maintainability in your implementations.

Beginner

Beginner Projects

These projects are designed for those new to C programming. They focus on fundamental concepts and provide a gentle introduction to building complete applications.

Project 1: Personal Information Manager

Description

Create a simple console application that allows users to store and manage personal contact information.

Learning Objectives

- Basic input/output operations
- Working with strings and arrays
- Simple file I/O operations
- Menu-driven interface design

Requirements

1. Store contact information (name, phone number, email address)
2. Add new contacts
3. View all contacts
4. Search for contacts by name
5. Delete contacts
6. Save contacts to a file
7. Load contacts from a file
8. Implement a simple menu system

Implementation Steps

1. Design a structure to hold contact information
2. Implement functions for each menu option
3. Create a menu system using a loop and switch statement
4. Implement file I/O for saving/loading contacts
5. Add error handling for invalid inputs

Sample Code Structure

```
#include <stdio.h>
#include <stdlib.h>
#include <string.h>

#define MAX_CONTACTS 100
#define NAME_LENGTH 50
#define PHONE_LENGTH 15
#define EMAIL_LENGTH 50

typedef struct {
    char name[NAME_LENGTH];
    char phone[PHONE_LENGTH];
    char email[EMAIL_LENGTH];
} Contact;

// Function prototypes
void add_contact(Contact contacts[], int *count);
void view_contacts(Contact contacts[], int count);
void search_contact(Contact contacts[], int count);
void delete_contact(Contact contacts[], int *count);
void save_contacts(Contact contacts[], int count);
int load_contacts(Contact contacts[]);
void display_menu();

int main() {
    Contact contacts[MAX_CONTACTS];
    int count = 0;

    // Load existing contacts
    count = load_contacts(contacts);

    int choice;
    do {
        display_menu();
        printf("Enter your choice: ");
        scanf("%d", &choice);

        switch (choice) {
            case 1:
                add_contact(contacts, &count);
                break;
            case 2:
                view_contacts(contacts, count);
                break;
            case 3:
```

```

        search_contact(contacts, count);
        break;
    case 4:
        delete_contact(contacts, &count);
        break;
    case 5:
        save_contacts(contacts, count);
        printf("Contacts saved successfully!\n");
        break;
    case 6:
        printf("Goodbye!\n");
        break;
    default:
        printf("Invalid choice. Please try again.\n");
    }
} while (choice != 6);

return 0;
}

// Implement all functions here

```

Common Pitfalls to Avoid

1. Buffer overflows when reading strings
2. Not checking return values of file operations
3. Memory leaks (though not common in this project)
4. Not handling invalid user input properly

Best Practices

1. Use `fgets()` instead of `gets()` for string input
2. Always validate user input
3. Check return values of file operations
4. Use constants for array sizes
5. Comment your code appropriately

Project 2: Simple Calculator

Description

Create a command-line calculator that can perform basic arithmetic operations and some advanced functions.

Learning Objectives

- Working with floating-point numbers
- Implementing mathematical functions
- Error handling
- Menu systems

Requirements

1. Perform basic operations: addition, subtraction, multiplication, division
2. Perform advanced operations: power, square root, logarithm, trigonometric functions
3. Handle invalid inputs gracefully
4. Support both interactive and command-line argument modes
5. Display help information
6. Clear operation history

Implementation Steps

1. Design the main calculator loop
2. Implement basic arithmetic functions
3. Add advanced mathematical functions
4. Handle command-line arguments
5. Implement error handling
6. Create a clean user interface

Sample Code Structure

```
#include <stdio.h>
#include <stdlib.h>
#include <math.h>
#include <string.h>

// Function prototypes
double add(double a, double b);
double subtract(double a, double b);
double multiply(double a, double b);
double divide(double a, double b);
double power(double base, double exponent);
double square_root(double x);
double logarithm(double x);
double sine(double x);
double cosine(double x);

void print_menu();
void interactive_mode();
```

```

void command_line_mode(int argc, char *argv[]);
int parse_operation(const char *op);

int main(int argc, char *argv[]) {
    if (argc > 1) {
        command_line_mode(argc, argv);
    } else {
        interactive_mode();
    }
    return 0;
}

// Implement all functions here

```

Common Pitfalls to Avoid

1. Division by zero errors
2. Invalid input for mathematical functions (e.g., square root of negative numbers)
3. Not handling command-line argument parsing correctly
4. Precision issues with floating-point arithmetic

Best Practices

1. Always check for division by zero
2. Validate inputs before performing operations
3. Use appropriate data types for precision requirements
4. Provide clear error messages
5. Implement proper input validation

Project 3: Number Guessing Game

Description

Create a number guessing game where the computer generates a random number and the player tries to guess it with feedback on whether their guess is too high or too low.

Learning Objectives

- Working with random numbers
- Loop control structures
- Conditional statements
- User interaction

Requirements

1. Generate a random number within a specified range
2. Allow the user to make guesses
3. Provide feedback on each guess (too high, too low, correct)
4. Count the number of attempts
5. Allow the user to play multiple rounds
6. Keep track of high scores
7. Provide difficulty levels

Implementation Steps

1. Implement random number generation
2. Create the main game loop
3. Add difficulty levels
4. Implement scoring system
5. Add high score tracking
6. Create a user-friendly interface

Sample Code Structure

```
#include <stdio.h>
#include <stdlib.h>
#include <time.h>

#define EASY_MAX 50
#define MEDIUM_MAX 100
#define HARD_MAX 200

// Function prototypes
int generate_random_number(int max);
void play_game(int max_number);
void display_high_scores();
void save_high_score(int attempts, int max_number);
void display_menu();
int get_difficulty();

int main() {
    srand(time(NULL)); // Seed the random number generator

    int choice;
    do {
        display_menu();
        printf("Enter your choice: ");
        scanf("%d", &choice);
```

```

switch (choice) {
    case 1: {
        int difficulty = get_difficulty();
        int max_number;
        switch (difficulty) {
            case 1: max_number = EASY_MAX; break;
            case 2: max_number = MEDIUM_MAX; break;
            case 3: max_number = HARD_MAX; break;
            default: max_number = MEDIUM_MAX;
        }
        play_game(max_number);
        break;
    }
    case 2:
        display_high_scores();
        break;
    case 3:
        printf("Thanks for playing!\n");
        break;
    default:
        printf("Invalid choice. Please try again.\n");
}
} while (choice != 3);

return 0;
}

// Implement all functions here

```

Common Pitfalls to Avoid

1. Not seeding the random number generator
2. Not validating user input
3. Integer overflow issues
4. Not handling invalid menu choices

Best Practices

1. Always seed the random number generator with `srand(time(NULL))`
2. Validate all user input
3. Provide clear instructions and feedback
4. Use constants for magic numbers
5. Implement proper error handling

Project 4: Temperature Converter

Description

Create a temperature conversion tool that can convert between Celsius, Fahrenheit, and Kelvin.

Learning Objectives

- Working with mathematical formulas
- User interface design
- Data validation
- File I/O for saving conversion history

Requirements

1. Convert between Celsius, Fahrenheit, and Kelvin
2. Support batch conversions from files
3. Save conversion history to a file
4. Load previous conversion history
5. Provide a clean command-line interface
6. Handle invalid temperature values

Implementation Steps

1. Implement temperature conversion formulas
2. Create a menu-driven interface
3. Add file I/O for history management
4. Implement batch conversion feature
5. Add data validation
6. Create a user-friendly interface

Sample Code Structure

```
#include <stdio.h>
#include <stdlib.h>

// Conversion formulas
double celsius_to_fahrenheit(double celsius);
double celsius_to_kelvin(double celsius);
double fahrenheit_to_celsius(double fahrenheit);
double fahrenheit_to_kelvin(double fahrenheit);
double kelvin_to_celsius(double kelvin);
double kelvin_to_fahrenheit(double kelvin);
```

```

// Function prototypes
void single_conversion();
void batch_conversion();
void save_history();
void load_history();
void display_menu();

int main() {
    int choice;
    do {
        display_menu();
        printf("Enter your choice: ");
        scanf("%d", &choice);

        switch (choice) {
            case 1:
                single_conversion();
                break;
            case 2:
                batch_conversion();
                break;
            case 3:
                save_history();
                break;
            case 4:
                load_history();
                break;
            case 5:
                printf("Goodbye!\n");
                break;
            default:
                printf("Invalid choice. Please try again.\n");
        }
    } while (choice != 5);

    return 0;
}

// Implement all functions here

```

Common Pitfalls to Avoid

1. Not handling invalid temperature values (e.g., below absolute zero)
2. Precision issues with floating-point arithmetic
3. Not validating file operations
4. Buffer overflows with string inputs

Best Practices

1. Validate temperature values (e.g., Kelvin cannot be negative)
2. Use appropriate precision for floating-point numbers
3. Check file operation return values
4. Provide clear error messages
5. Use constants for conversion formulas

Tips for Success

1. **Start Simple:** Begin with basic functionality and gradually add features
2. **Test Frequently:** Test your code after implementing each feature
3. **Handle Errors:** Always consider what could go wrong and handle it gracefully
4. **Use Functions:** Break your code into logical functions for better organization
5. **Comment Your Code:** Explain complex logic and algorithms
6. **Validate Input:** Never trust user input; always validate it
7. **Save Often:** Use version control or save your work frequently
8. **Ask for Help:** Don't hesitate to seek help when you're stuck

These beginner projects will help you build confidence in C programming while practicing essential concepts. Remember to focus on writing clean, readable code and handling errors appropriately.

Capstone

Capstone Projects

These capstone projects are designed to integrate all the skills and knowledge you've gained throughout the C programming course. They represent substantial, real-world applications that demonstrate mastery of advanced C programming concepts.

Project 1: Distributed Chat Application

Description

Create a distributed chat application with multiple servers, clients, and advanced features like encryption, file sharing, and user authentication.

Learning Objectives

- Network programming with TCP/IP and UDP
- Multithreading and concurrency
- Cryptography and security
- Distributed systems concepts
- Protocol design and implementation
- Database integration
- User interface design

Requirements

1. Server Components:

- Multiple chat servers with load balancing
- User authentication and session management
- Message routing between servers
- Persistent message storage
- Administration interface

2. Client Features:

- User registration and login
- Real-time messaging with multiple users
- Private messaging
- Group chats
- File sharing capability
- Message history

- Presence indication (online/offline)
- Emoticons and rich text support

3. Advanced Features:

- End-to-end encryption
- Message queuing for offline users
- Voice messaging (optional)
- Video calling (optional)
- Plugin architecture for extensions
- Mobile client compatibility

Implementation Steps

1. Design system architecture and communication protocols
2. Implement core server infrastructure
3. Create client application with user interface
4. Add authentication and security features
5. Implement message routing and storage
6. Add advanced features like file sharing
7. Create administration tools
8. Test with multiple users and servers
9. Optimize for performance and scalability

Sample Code Structure

```
#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <unistd.h>
#include <sys/socket.h>
#include <netinet/in.h>
#include <arpa/inet.h>
#include <pthread.h>
#include <openssl/ssl.h>
#include <openssl/err.h>
#include <sqlite3.h>

#define MAX_CLIENTS 1000
#define MAX_MESSAGE_LENGTH 1024
#define SERVER_PORT 8080
#define DATABASE_FILE "chat.db"

// Message types
typedef enum {
    MSG_TEXT,
    MSG_FILE,
```

```

    MSG_JOIN,
    MSG_LEAVE,
    MSG_AUTH,
    MSG_ERROR
} message_type_t;

// User structure
typedef struct {
    int id;
    char username[50];
    char password_hash[64];
    int socket_fd;
    SSL *ssl;
    int authenticated;
    time_t last_activity;
} user_t;

// Message structure
typedef struct {
    message_type_t type;
    int sender_id;
    int receiver_id; // 0 for broadcast
    char content[MAX_MESSAGE_LENGTH];
    time_t timestamp;
    int encrypted;
} message_t;

// Server structure
typedef struct {
    int server_fd;
    SSL_CTX *ssl_ctx;
    sqlite3 *db;
    user_t clients[MAX_CLIENTS];
    int client_count;
    pthread_mutex_t clients_mutex;
    pthread_t worker_threads[10];
    int thread_count;
    int running;
} chat_server_t;

// Client structure
typedef struct {
    int socket_fd;
    SSL *ssl;
    char username[50];
    int authenticated;
    pthread_t receive_thread;

```

```

    pthread_t send_thread;
    int connected;
} chat_client_t;

// Function prototypes
chat_server_t* create_chat_server(int port);
void destroy_chat_server(chat_server_t *server);
int start_server(chat_server_t *server);
void stop_server(chat_server_t *server);
void* client_handler(void *arg);
int authenticate_user(chat_server_t *server, const char *username, const char *password);
int register_user(chat_server_t *server, const char *username, const char *password);
int send_message(chat_server_t *server, message_t *msg);
int broadcast_message(chat_server_t *server, message_t *msg);
int store_message(chat_server_t *server, message_t *msg);
message_t* retrieve_message_history(chat_server_t *server, int user_id, int limit);
int init_database(chat_server_t *server);
SSL_CTX* create_ssl_context();
int create_user_table(chat_server_t *server);
void* server_worker_thread(void *arg);
void log_message(const char *format, ...);

// Client functions
chat_client_t* create_chat_client(const char *server_ip, int port);
void destroy_chat_client(chat_client_t *client);
int connect_to_server(chat_client_t *client);
void disconnect_from_server(chat_client_t *client);
int login(chat_client_t *client, const char *username, const char *password);
int send_chat_message(chat_client_t *client, const char *message, int receiver_id);
void* receive_messages(void *arg);
void* send_messages(void *arg);

int main(int argc, char *argv[]) {
    if (argc < 2) {
        fprintf(stderr, "Usage: %s [server|client]\n", argv[0]);
        return 1;
    }

    if (strcmp(argv[1], "server") == 0) {
        chat_server_t *server = create_chat_server(SERVER_PORT);
        if (!server) {
            fprintf(stderr, "Failed to create server\n");
            return 1;
        }

        if (start_server(server) != 0) {
            fprintf(stderr, "Failed to start server\n");

```

```

        destroy_chat_server(server);
        return 1;
    }

    printf("Chat server started on port %d\n", SERVER_PORT);

    // Keep server running
    while (1) {
        sleep(1);
    }

    stop_server(server);
    destroy_chat_server(server);
} else if (strcmp(argv[1], "client") == 0) {
    chat_client_t *client = create_chat_client("127.0.0.1", SERVER_PORT);
    if (!client) {
        fprintf(stderr, "Failed to create client\n");
        return 1;
    }

    if (connect_to_server(client) != 0) {
        fprintf(stderr, "Failed to connect to server\n");
        destroy_chat_client(client);
        return 1;
    }

    // Login process
    char username[50], password[50];
    printf("Enter username: ");
    scanf("%s", username);
    printf("Enter password: ");
    scanf("%s", password);

    if (login(client, username, password) != 0) {
        fprintf(stderr, "Login failed\n");
        disconnect_from_server(client);
        destroy_chat_client(client);
        return 1;
    }

    printf("Login successful! Starting chat...\n");

    // Start message threads
    pthread_create(&client->receive_thread, NULL, receive_messages, client);
    pthread_create(&client->send_thread, NULL, send_messages, client);

    // Wait for threads

```

```

        pthread_join(client->receive_thread, NULL);
        pthread_join(client->send_thread, NULL);

        disconnect_from_server(client);
        destroy_chat_client(client);
    } else {
        fprintf(stderr, "Invalid mode. Use 'server' or 'client'\n");
        return 1;
    }

    return 0;
}

// Implement all functions here

```

Key Challenges

1. **Concurrency:** Managing multiple clients and threads safely
2. **Security:** Implementing robust encryption and authentication
3. **Scalability:** Designing for thousands of concurrent users
4. **Reliability:** Ensuring message delivery and system stability
5. **Performance:** Optimizing database queries and network operations
6. **Protocol Design:** Creating efficient communication protocols

Evaluation Criteria

- **Functionality:** All required features implemented correctly
- **Security:** Proper encryption and authentication mechanisms
- **Performance:** Efficient handling of multiple connections
- **Robustness:** Error handling and recovery mechanisms
- **Code Quality:** Clean, well-documented, maintainable code
- **Testing:** Comprehensive test suite with edge cases
- **Documentation:** Clear user and developer documentation

Project 2: Operating System Kernel

Description

Implement a basic operating system kernel that can boot, manage processes, handle interrupts, and provide system calls.

Learning Objectives

- Low-level system programming
- Memory management
- Process scheduling
- Interrupt handling
- Device drivers
- File system implementation
- Boot process understanding

Requirements

1. Boot Process:

- Custom bootloader
- Kernel loading and initialization
- Basic hardware detection

2. Memory Management:

- Physical memory management
- Virtual memory with paging
- Heap allocation
- Memory protection

3. Process Management:

- Process creation and termination
- Process scheduling (round-robin, priority-based)
- Inter-process communication
- Process synchronization

4. Interrupt Handling:

- Interrupt descriptor table (IDT)
- Exception handling
- Hardware interrupt processing
- System call interface

5. Device Drivers:

- Keyboard driver
- Display driver (VGA text mode)
- Timer driver
- Storage driver (optional)

6. File System:

- Simple file system implementation
- File operations (create, read, write, delete)
- Directory management

Implementation Steps

1. Set up development environment with cross-compiler
2. Implement bootloader and kernel entry point
3. Set up GDT, IDT, and basic interrupt handling
4. Implement memory management system
5. Create process management and scheduling
6. Add device drivers for basic I/O
7. Implement file system
8. Create system call interface
9. Develop user-space programs
10. Test and debug the complete system

Sample Code Structure

```
// kernel.h
#ifndef KERNEL_H
#define KERNEL_H

#include <stdint.h>
#include <stddef.h>

// Basic types
typedef uint8_t  u8;
typedef uint16_t u16;
typedef uint32_t u32;
typedef uint64_t u64;

// Kernel entry point
void kernel_main();

// Hardware abstraction
void outb(u16 port, u8 val);
u8 inb(u16 port);
void outw(u16 port, u16 val);
u16 inw(u16 port);

#endif

// gdt.h
#ifndef GDT_H
#define GDT_H

#include "kernel.h"

typedef struct {
```



```

    u16 limit_low;
    u16 base_low;
    u8 base_middle;
    u8 access;
    u8 granularity;
    u8 base_high;
} __attribute__((packed)) gdt_entry_t;

typedef struct {
    u16 limit;
    u32 base;
} __attribute__((packed)) gdt_ptr_t;

void init_gdt();
void gdt_set_gate(int num, u32 base, u32 limit, u8 access, u8 gran);

#endif

// idt.h
#ifndef IDT_H
#define IDT_H

#include "kernel.h"

typedef struct {
    u16 base_low;
    u16 sel;
    u8 always0;
    u8 flags;
    u16 base_high;
} __attribute__((packed)) idt_entry_t;

typedef struct {
    u16 limit;
    u32 base;
} __attribute__((packed)) idt_ptr_t;

void init_idt();
void idt_set_gate(u8 num, u32 base, u16 sel, u8 flags);

#endif

// memory.h
#ifndef MEMORY_H
#define MEMORY_H

#include "kernel.h"

```

```

#define PAGE_SIZE 4096

typedef struct {
    u32 present      : 1;
    u32 rw            : 1;
    u32 user         : 1;
    u32 writethru    : 1;
    u32 cached       : 1;
    u32 accessed     : 1;
    u32 dirty        : 1;
    u32 pat          : 1;
    u32 global       : 1;
    u32 available    : 3;
    u32 frame        : 20;
} __attribute__((packed)) page_t;

typedef struct {
    page_t pages[1024];
} __attribute__((packed)) page_table_t;

typedef struct {
    page_table_t *tables[1024];
    u32 tables_physical[1024];
    u32 physical_address;
} __attribute__((packed)) page_directory_t;

void init_paging();
void switch_page_directory(page_directory_t *dir);
page_t *get_page(u32 address, int make, page_directory_t *dir);
void alloc_frame(page_t *page, int is_kernel, int is_writeable);
void free_frame(page_t *page);

#endif

// process.h
#ifndef PROCESS_H
#define PROCESS_H

#include "kernel.h"
#include "memory.h"

#define MAX_PROCESSES 256
#define KERNEL_STACK_SIZE 4096

typedef enum {
    PROCESS_RUNNING,

```

```

    PROCESS_READY,
    PROCESS_BLOCKED,
    PROCESS_TERMINATED
} process_state_t;

typedef struct {
    u32 esp, ebp, esi, edi, eax, ebx, ecx, edx;
    u32 eip, eflags;
    u32 cr3; // Page directory
    u32 kernel_stack;
    process_state_t state;
    u32 pid;
    char name[32];
} __attribute__((packed)) process_t;

typedef struct {
    process_t *processes[MAX_PROCESSES];
    u32 current_process;
    u32 process_count;
} process_manager_t;

void init_process_manager();
process_t *create_process(const char *name, u32 entry_point);
void switch_to_process(process_t *process);
void schedule();
void yield();

#endif

// syscall.h
#ifdef SYSCALL_H
#define SYSCALL_H

#include "kernel.h"

#define SYSCALL_WRITE 1
#define SYSCALL_READ 2
#define SYSCALL_EXIT 3
#define SYSCALL_FORK 4

void init_syscalls();
void syscall_handler(registers_t *regs);

#endif

// main kernel implementation
void kernel_main() {

```

```

// Initialize hardware
init_gdt();
init_idt();

// Initialize memory management
init_paging();

// Initialize process management
init_process_manager();

// Initialize system calls
init_syscalls();

// Create initial processes
process_t *init = create_process("init", (u32)init_process);
switch_to_process(init);

// Start scheduler
while (1) {
    schedule();
    __asm__ __volatile__("hlt");
}
}

// Implement all other components here

```

Key Challenges

1. **Low-level Programming:** Working directly with hardware
2. **Memory Management:** Implementing virtual memory and paging
3. **Concurrency:** Managing processes and threads safely
4. **Interrupt Handling:** Properly handling hardware interrupts
5. **Boot Process:** Understanding and implementing the boot sequence
6. **Debugging:** Debugging at the kernel level

Evaluation Criteria

- **Boot Process:** Successful kernel loading and initialization
- **Memory Management:** Correct implementation of paging and allocation
- **Process Management:** Functional process creation and scheduling
- **Interrupt Handling:** Proper exception and interrupt processing
- **System Calls:** Working system call interface
- **Stability:** Kernel stability under various conditions
- **Documentation:** Clear explanation of kernel components

Project 3: Database Management System

Description

Create a full-featured relational database management system with SQL support, transactions, indexing, and optimization.

Learning Objectives

- Data storage and retrieval
- Query processing and optimization
- Transaction management
- Concurrency control
- Indexing and search algorithms
- File system integration
- SQL parser and executor

Requirements

1. Storage Engine:

- Page-based storage system
- Buffer pool management
- Transaction log for recovery
- Checkpointing mechanism

2. SQL Support:

- DDL (CREATE, ALTER, DROP)
- DML (SELECT, INSERT, UPDATE, DELETE)
- Query optimization
- Aggregation functions
- Joins (INNER, LEFT, RIGHT, FULL)

3. Transaction Management:

- ACID properties
- Lock-based concurrency control
- Deadlock detection and resolution
- Transaction isolation levels

4. Indexing:

- B+ tree implementation
- Hash index support
- Composite index support
- Index maintenance

5. Administration:

- Database creation and management
- User authentication and permissions
- Backup and recovery
- Performance monitoring

Implementation Steps

1. Design database architecture and storage format
2. Implement storage engine with page management
3. Create buffer pool and cache system
4. Develop SQL parser and query planner
5. Implement B+ tree and other index structures
6. Add transaction management and concurrency control
7. Create administration tools and utilities
8. Test with standard SQL benchmarks
9. Optimize for performance
10. Document the system architecture

Sample Code Structure

```
#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <pthread.h>
#include <unistd.h>
#include <fcntl.h>
#include <sys/stat.h>

#define PAGE_SIZE 4096
#define MAX_DATABASES 100
#define MAX_TABLES 1000
#define MAX_COLUMNS 50
#define MAX_INDEXES 100

// Data types
typedef enum {
    TYPE_INT,
    TYPE_BIGINT,
    TYPE_FLOAT,
    TYPE_DOUBLE,
    TYPE_VARCHAR,
    TYPE_TEXT,
    TYPE_DATE,
    TYPE_DATETIME
} data_type_t;
```

```

// Page structure
typedef struct {
    u32 page_id;
    u32 next_page;
    u32 prev_page;
    u16 free_space;
    u8 page_type;
    u8 data[PAGE_SIZE - sizeof(u32) * 3 - sizeof(u16) - sizeof(u8)];
} __attribute__((packed)) page_t;

// Table structure
typedef struct {
    u32 table_id;
    char name[64];
    u32 page_id; // First page of table data
    u32 row_count;
    u32 column_count;
    struct {
        char name[32];
        data_type_t type;
        u16 size;
        u8 nullable;
    } columns[MAX_COLUMNS];
} table_t;

// Index structure
typedef struct {
    u32 index_id;
    char name[64];
    u32 table_id;
    u32 column_count;
    u32 columns[MAX_COLUMNS];
    u8 index_type; // 0 = B+ tree, 1 = Hash
    u32 root_page;
} index_t;

// Transaction structure
typedef enum {
    TXN_ACTIVE,
    TXN_COMMITTED,
    TXN_ABORTED
} txn_state_t;

typedef struct {
    u32 txn_id;
    txn_state_t state;

```

```

    u64 start_time;
    u64 commit_time;
    pthread_mutex_t lock;
    // Undo/redo logs
    struct log_entry *undo_log;
    struct log_entry *redo_log;
} transaction_t;

// Database structure
typedef struct {
    char name[64];
    int fd; // File descriptor
    u32 page_count;
    table_t tables[MAX_TABLES];
    u32 table_count;
    index_t indexes[MAX_INDEXES];
    u32 index_count;
    pthread_mutex_t lock;
    pthread_rwlock_t rw_lock;
} database_t;

// Buffer pool
typedef struct {
    page_t *pages;
    u32 *page_ids;
    u8 *dirty_flags;
    u64 *access_times;
    u32 capacity;
    u32 count;
    pthread_mutex_t lock;
} buffer_pool_t;

// SQL Parser structures
typedef enum {
    SQL_SELECT,
    SQL_INSERT,
    SQL_UPDATE,
    SQL_DELETE,
    SQL_CREATE_TABLE,
    SQL_DROP_TABLE,
    SQL_CREATE_INDEX,
    SQL_BEGIN_TRANSACTION,
    SQL_COMMIT,
    SQL_ROLLBACK
} sql_statement_type_t;

typedef struct {

```



```

    sql_statement_type_t type;
    char *table_name;
    struct {
        char *column_name;
        char *value;
    } *columns;
    u32 column_count;
    char *where_clause;
    char **order_by;
    u32 order_by_count;
    u32 limit;
} sql_statement_t;

// Query plan
typedef enum {
    PLAN_SCAN,
    PLAN_INDEX_SCAN,
    PLAN_JOIN,
    PLAN_SORT,
    PLAN_AGGREGATE
} plan_node_type_t;

typedef struct plan_node {
    plan_node_type_t type;
    struct plan_node *left;
    struct plan_node *right;
    void *data; // Node-specific data
} plan_node_t;

// Database manager
typedef struct {
    database_t databases[MAX_DATABASES];
    u32 database_count;
    buffer_pool_t buffer_pool;
    transaction_t *active_transactions;
    u32 transaction_count;
    pthread_mutex_t txn_lock;
} db_manager_t;

// Function prototypes
db_manager_t* create_db_manager(u32 buffer_pool_size);
void destroy_db_manager(db_manager_t *manager);
database_t* create_database(db_manager_t *manager, const char *name);
int drop_database(db_manager_t *manager, const char *name);
table_t* create_table(database_t *db, const char *name, /* column definitions */);
int drop_table(database_t *db, const char *name);
index_t* create_index(database_t *db, const char *name, const char *table_name, /* column li

```

```

int drop_index(database_t *db, const char *name);
transaction_t* begin_transaction(db_manager_t *manager);
int commit_transaction(db_manager_t *manager, transaction_t *txn);
int rollback_transaction(db_manager_t *manager, transaction_t *txn);
sql_statement_t* parse_sql(const char *sql);
plan_node_t* create_query_plan(sql_statement_t *stmt, database_t *db);
int execute_query_plan(plan_node_t *plan, transaction_t *txn, /* result set */);
page_t* get_page(database_t *db, u32 page_id);
int write_page(database_t *db, page_t *page);
void* btree_insert(index_t *index, void *key, u32 page_id);
void* btree_search(index_t *index, void *key);
int btree_delete(index_t *index, void *key);
void lock_table(database_t *db, table_t *table, int exclusive);
void unlock_table(database_t *db, table_t *table);
void lock_row(database_t *db, table_t *table, u32 row_id, int exclusive);
void unlock_row(database_t *db, table_t *table, u32 row_id);

int main() {
    // Create database manager
    db_manager_t *manager = create_db_manager(1000); // 1000 page buffer pool
    if (!manager) {
        fprintf(stderr, "Failed to create database manager\n");
        return 1;
    }

    // Create database
    database_t *db = create_database(manager, "testdb");
    if (!db) {
        fprintf(stderr, "Failed to create database\n");
        destroy_db_manager(manager);
        return 1;
    }

    // Example SQL execution
    const char *sql = "CREATE TABLE users (id INT PRIMARY KEY, name VARCHAR(50), email VARCHAR(50))";
    sql_statement_t *stmt = parse_sql(sql);
    if (stmt) {
        plan_node_t *plan = create_query_plan(stmt, db);
        if (plan) {
            transaction_t *txn = begin_transaction(manager);
            execute_query_plan(plan, txn, NULL);
            commit_transaction(manager, txn);
            // Free plan
        }
        // Free statement
    }
}

```

```
        // Cleanup
        destroy_db_manager(manager);
        return 0;
    }

    // Implement all functions here
```

Key Challenges

1. **Data Consistency:** Ensuring ACID properties in transactions
2. **Concurrency Control:** Managing concurrent access to data
3. **Query Optimization:** Creating efficient execution plans
4. **Storage Management:** Efficiently managing disk and memory resources
5. **Index Implementation:** Implementing complex data structures like B+ trees
6. **SQL Parsing:** Creating a robust SQL parser

Evaluation Criteria

- **SQL Compliance:** Support for standard SQL features
- **Performance:** Efficient query execution and optimization
- **Reliability:** Transaction safety and recovery mechanisms
- **Scalability:** Ability to handle large datasets
- **Concurrency:** Proper handling of concurrent operations
- **Code Quality:** Well-structured, maintainable code
- **Documentation:** Comprehensive system documentation

Project 4: Game Engine

Description

Create a 2D game engine with rendering, physics, audio, input handling, and scripting capabilities.

Learning Objectives

- Graphics programming with OpenGL/DirectX
- Physics simulation
- Audio processing
- Input handling
- Entity-component system
- Resource management
- Cross-platform development

Requirements

1. Graphics System:

- 2D rendering pipeline
- Sprite management
- Animation system
- Particle effects
- Camera system

2. Physics Engine:

- Collision detection (AABB, circle, polygon)
- Collision response
- Rigidbody physics
- Constraints and joints

3. Audio System:

- Sound loading and playback
- Music streaming
- 3D audio positioning
- Audio effects

4. Input Handling:

- Keyboard, mouse, gamepad support
- Input mapping system
- Gesture recognition

5. Game Object System:

- Entity-component architecture
- Scene management
- Serialization system
- Scripting interface

6. Resource Management:

- Asset loading and caching
- Texture compression
- Memory pooling
- Async loading

Implementation Steps

1. Set up graphics rendering framework
2. Implement entity-component system
3. Create physics simulation engine
4. Add audio processing capabilities
5. Develop input handling system
6. Build resource management system

7. Create scripting interface
8. Develop tools and editor
9. Test with sample games
10. Optimize for performance

Sample Code Structure

```
#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <math.h>
#include <pthread.h>
#ifdef _WIN32
    #include <windows.h>
    #include <GL/gl.h>
#else
    #include <GL/gl.h>
    #include <unistd.h>
#endif

#define MAX_ENTITIES 10000
#define MAX_COMPONENTS 32
#define MAX_SYSTEMS 64

// Math structures
typedef struct {
    float x, y;
} vec2_t;

typedef struct {
    float x, y, z;
} vec3_t;

typedef struct {
    float x, y, z, w;
} vec4_t;

typedef struct {
    float m[16]; // 4x4 matrix
} mat4_t;

// Component types
typedef enum {
    COMPONENT_TRANSFORM,
    COMPONENT_SPRITE,
    COMPONENT_RIGIDBODY,
```

```

        COMPONENT_COLLIDER,
        COMPONENT_SCRIPT,
        COMPONENT_CAMERA,
        COMPONENT_AUDIO_SOURCE,
        COMPONENT_PARTICLE_SYSTEM
    } component_type_t;

    // Transform component
    typedef struct {
        vec3_t position;
        vec3_t rotation;
        vec3_t scale;
        mat4_t world_matrix;
    } transform_component_t;

    // Sprite component
    typedef struct {
        u32 texture_id;
        vec2_t size;
        vec4_t color;
        int visible;
    } sprite_component_t;

    // Rigidbody component
    typedef struct {
        vec2_t velocity;
        vec2_t acceleration;
        float mass;
        float drag;
        int use_gravity;
    } rigidbody_component_t;

    // Collider component
    typedef enum {
        COLLIDER_AABB,
        COLLIDER_CIRCLE,
        COLLIDER_POLYGON
    } collider_type_t;

    typedef struct {
        collider_type_t type;
        union {
            struct { vec2_t min, max; } aabb;
            struct { vec2_t center; float radius; } circle;
            struct { vec2_t *vertices; int vertex_count; } polygon;
        } shape;
        int is_trigger;
    }

```

```

} collider_component_t;

// Entity structure
typedef struct {
    u32 id;
    u64 component_mask; // Bitmask of components
    char name[64];
    int active;
} entity_t;

// Component arrays (structure of arrays for cache efficiency)
typedef struct {
    transform_component_t transforms[MAX_ENTITIES];
    sprite_component_t sprites[MAX_ENTITIES];
    rigidbody_component_t rigidbodies[MAX_ENTITIES];
    collider_component_t colliders[MAX_ENTITIES];
    // ... other component arrays
} component_arrays_t;

// System base structure
typedef struct {
    char name[32];
    void (*update)(float delta_time);
    void (*render)();
    int enabled;
} system_t;

// Renderer system
typedef struct {
    system_t base;
    u32 shader_program;
    u32 vao, vbo;
    mat4_t projection_matrix;
    mat4_t view_matrix;
} renderer_system_t;

// Physics system
typedef struct {
    system_t base;
    vec2_t gravity;
    float time_scale;
} physics_system_t;

// Collision system
typedef struct {
    system_t base;
    struct {

```

```

        u32 entity_a, entity_b;
        vec2_t contact_point;
        vec2_t normal;
        float penetration;
    } *contacts;
    u32 contact_count;
    u32 contact_capacity;
} collision_system_t;

// Scene structure
typedef struct {
    entity_t entities[MAX_ENTITIES];
    u32 entity_count;
    component_arrays_t components;
    system_t *systems[MAX_SYSTEMS];
    u32 system_count;
    float time_accumulator;
} scene_t;

// Game engine structure
typedef struct {
    scene_t *current_scene;
    int running;
    float delta_time;
    u64 frame_count;
    struct {
        int width, height;
        const char *title;
        void *window_handle;
    } window;
    struct {
        float master_volume;
        float sfx_volume;
        float music_volume;
    } audio_settings;
} game_engine_t;

// Function prototypes
game_engine_t* create_game_engine(const char *title, int width, int height);
void destroy_game_engine(game_engine_t *engine);
int initialize_engine(game_engine_t *engine);
void run_game_loop(game_engine_t *engine);
void shutdown_engine(game_engine_t *engine);

// Entity management
entity_t* create_entity(scene_t *scene, const char *name);
void destroy_entity(scene_t *scene, u32 entity_id);

```



```

int add_component(scene_t *scene, u32 entity_id, component_type_t type);
int remove_component(scene_t *scene, u32 entity_id, component_type_t type);
int has_component(scene_t *scene, u32 entity_id, component_type_t type);
void* get_component(scene_t *scene, u32 entity_id, component_type_t type);

// System management
int register_system(scene_t *scene, system_t *system);
int unregister_system(scene_t *scene, const char *system_name);
system_t* get_system(scene_t *scene, const char *system_name);

// Math functions
vec2_t vec2_add(vec2_t a, vec2_t b);
vec2_t vec2_subtract(vec2_t a, vec2_t b);
vec2_t vec2_multiply(vec2_t v, float scalar);
float vec2_length(vec2_t v);
vec2_t vec2_normalize(vec2_t v);
mat4_t mat4_identity();
mat4_t mat4_orthographic(float left, float right, float bottom, float top, float near, float
mat4_t mat4_transform(vec3_t position, vec3_t rotation, vec3_t scale);

// Renderer functions
renderer_system_t* create_renderer_system();
void destroy_renderer_system(renderer_system_t *renderer);
void renderer_update(float delta_time);
void renderer_render();
void renderer_draw_sprite(sprite_component_t *sprite, transform_component_t *transform);

// Physics functions
physics_system_t* create_physics_system();
void destroy_physics_system(physics_system_t *physics);
void physics_update(float delta_time);
void integrate_rigidbody(rigidbody_component_t *rb, transform_component_t *transform, float

// Collision functions
collision_system_t* create_collision_system();
void destroy_collision_system(collision_system_t *collision);
void collision_update(float delta_time);
int check_collision_aabb_circle(vec2_t aabb_min, vec2_t aabb_max, vec2_t circle_center, float
int check_collision_circle_circle(vec2_t center1, float radius1, vec2_t center2, float radiu
void resolve_collision(u32 entity_a, u32 entity_b, vec2_t contact_point, vec2_t normal, float

// Input functions
int is_key_pressed(int key);
int is_key_released(int key);
int is_key_held(int key);
vec2_t get_mouse_position();
int is_mouse_button_pressed(int button);

```

```

// Resource management
u32 load_texture(const char *filename);
void unload_texture(u32 texture_id);
void* load_shader(const char *vertex_shader, const char *fragment_shader);
void unload_shader(u32 shader_program);

int main() {
    // Create game engine
    game_engine_t *engine = create_game_engine("My Game Engine", 1024, 768);
    if (!engine) {
        fprintf(stderr, "Failed to create game engine\n");
        return 1;
    }

    // Initialize engine
    if (initialize_engine(engine) != 0) {
        fprintf(stderr, "Failed to initialize engine\n");
        destroy_game_engine(engine);
        return 1;
    }

    // Create sample scene
    scene_t *scene = engine->current_scene;

    // Create player entity
    entity_t *player = create_entity(scene, "Player");
    add_component(scene, player->id, COMPONENT_TRANSFORM);
    add_component(scene, player->id, COMPONENT_SPRITE);
    add_component(scene, player->id, COMPONENT_RIGIDBODY);
    add_component(scene, player->id, COMPONENT_COLLIDER);

    // Configure player components
    transform_component_t *player_transform = get_component(scene, player->id, COMPONENT_TRANSFORM);
    player_transform->position.x = 400;
    player_transform->position.y = 300;
    player_transform->scale.x = player_transform->scale.y = 2.0f;

    sprite_component_t *player_sprite = get_component(scene, player->id, COMPONENT_SPRITE);
    player_sprite->texture_id = load_texture("player.png");
    player_sprite->size.x = 32;
    player_sprite->size.y = 32;

    // Run game loop
    run_game_loop(engine);

    // Cleanup

```

```
    shutdown_engine(engine);  
    destroy_game_engine(engine);  
  
    return 0;  
}  
  
// Implement all functions here
```

Key Challenges

1. **Performance:** Optimizing rendering and physics for real-time performance
2. **Memory Management:** Efficiently managing game assets and entities
3. **Cross-platform:** Supporting multiple operating systems and graphics APIs
4. **Architecture:** Designing a flexible, extensible system
5. **Mathematics:** Implementing complex mathematical operations for graphics and physics
6. **Resource Management:** Loading and managing game assets efficiently

Evaluation Criteria

- **Functionality:** All required systems implemented and working
- **Performance:** Smooth gameplay with good frame rates
- **Flexibility:** Extensible architecture for adding new features
- **Code Quality:** Well-organized, maintainable code
- **Documentation:** Clear documentation of engine architecture
- **Sample Games:** Demonstrating engine capabilities with sample games
- **Tools:** Development tools for content creation

Tips for Capstone Success

1. **Choose Wisely:** Select a project that matches your interests and skill level
2. **Plan Extensively:** Create detailed designs and specifications before coding
3. **Iterative Development:** Build and test components incrementally
4. **Focus on Core Features:** Implement essential functionality first
5. **Document Everything:** Maintain comprehensive documentation throughout
6. **Test Thoroughly:** Create extensive test cases including edge cases
7. **Seek Feedback:** Get regular feedback from peers and mentors
8. **Manage Time:** Create realistic timelines and stick to them
9. **Learn Continuously:** Research and learn new techniques as needed
10. **Enjoy the Process:** Take pride in creating something substantial and meaningful

These capstone projects represent the culmination of your C programming education. They require integrating multiple concepts and techniques while solving complex, real-world problems. Choose a project that excites you and demonstrates your mastery of C programming.

Intermediate

Intermediate Projects

These projects are designed for programmers with basic C knowledge who want to advance their skills. They involve more complex data structures, file handling, and problem-solving techniques.

Project 1: Student Grade Management System

Description

Create a comprehensive system to manage student grades, calculate statistics, and generate reports.

Learning Objectives

- Working with dynamic data structures
- File I/O operations with structured data
- Data analysis and statistics
- Memory management
- Modular programming

Requirements

1. Store student information (ID, name, grades for multiple subjects)
2. Add, edit, and delete student records
3. Calculate average grades for students and subjects
4. Generate class statistics (highest, lowest, median grades)
5. Save and load data from files
6. Search and sort student records
7. Generate detailed reports
8. Handle multiple classes/sections

Implementation Steps

1. Design data structures for students and grades
2. Implement dynamic memory allocation for records
3. Create functions for CRUD operations
4. Implement statistical analysis functions
5. Add file I/O for data persistence

6. Create sorting and searching algorithms
7. Develop report generation features
8. Build a menu-driven interface

Sample Code Structure

```
#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <math.h>

#define MAX_NAME_LENGTH 50
#define MAX_SUBJECTS 10
#define FILENAME "students.dat"

typedef struct {
    int id;
    char name[MAX_NAME_LENGTH];
    int num_subjects;
    char subjects[MAX_SUBJECTS][MAX_NAME_LENGTH];
    double grades[MAX_SUBJECTS];
} Student;

typedef struct {
    Student *students;
    int count;
    int capacity;
} StudentManager;

// Function prototypes
StudentManager* create_student_manager(int initial_capacity);
void destroy_student_manager(StudentManager *manager);
int add_student(StudentManager *manager, Student student);
int edit_student(StudentManager *manager, int id, Student new_student);
int delete_student(StudentManager *manager, int id);
Student* find_student(StudentManager *manager, int id);
void sort_students_by_name(StudentManager *manager);
void sort_students_by_average(StudentManager *manager);
double calculate_student_average(Student *student);
double calculate_subject_average(StudentManager *manager, const char *subject);
void generate_class_report(StudentManager *manager);
void save_to_file(StudentManager *manager, const char *filename);
int load_from_file(StudentManager *manager, const char *filename);
void display_menu();

int main() {
```

```

StudentManager *manager = create_student_manager(10);
if (!manager) {
    fprintf(stderr, "Failed to create student manager\n");
    return 1;
}

// Load existing data
load_from_file(manager, FILENAME);

int choice;
do {
    display_menu();
    printf("Enter your choice: ");
    scanf("%d", &choice);

    switch (choice) {
        case 1: {
            Student student;
            printf("Enter student ID: ");
            scanf("%d", &student.id);
            printf("Enter student name: ");
            scanf(" %[^\\n]", student.name);
            // Add more input handling
            add_student(manager, student);
            break;
        }
        // Implement other cases
        case 8:
            save_to_file(manager, FILENAME);
            printf("Data saved successfully!\\n");
            break;
        case 9:
            printf("Goodbye!\\n");
            break;
        default:
            printf("Invalid choice. Please try again.\\n");
    }
} while (choice != 9);

destroy_student_manager(manager);
return 0;
}

// Implement all functions here

```

Common Pitfalls to Avoid

1. Memory leaks from dynamic allocation
2. Buffer overflows with string inputs
3. Not checking return values of file operations
4. Incorrect implementation of sorting algorithms
5. Not handling edge cases in statistical calculations

Best Practices

1. Always free dynamically allocated memory
2. Check for NULL pointers before dereferencing
3. Validate all user inputs
4. Use appropriate data structures for efficiency
5. Implement proper error handling
6. Comment complex algorithms
7. Use constants for array sizes and limits

Project 2: Simple Text Editor

Description

Create a basic text editor that can create, edit, and save text files with basic editing features.

Learning Objectives

- Working with dynamic strings
- File I/O operations
- Text processing algorithms
- Memory management
- User interface design

Requirements

1. Create new text files
2. Open and edit existing text files
3. Save files to disk
4. Basic editing operations (insert, delete, replace)
5. Search and replace functionality
6. Undo/redo functionality
7. Line numbering
8. Status bar with file information

Implementation Steps

1. Design data structures for text storage
2. Implement file I/O operations
3. Create basic editing functions
4. Add search and replace features
5. Implement undo/redo functionality
6. Add line numbering
7. Create user interface
8. Add status bar and file information

Sample Code Structure

```
#include <stdio.h>
#include <stdlib.h>
#include <string.h>

#define MAX_LINE_LENGTH 1000
#define INITIAL_LINES 100

typedef struct {
    char **lines;
    int line_count;
    int capacity;
    char filename[256];
    int modified;
} TextEditor;

// Function prototypes
TextEditor* create_editor();
void destroy_editor(TextEditor *editor);
int load_file(TextEditor *editor, const char *filename);
int save_file(TextEditor *editor, const char *filename);
void insert_line(TextEditor *editor, int line_num, const char *text);
void delete_line(TextEditor *editor, int line_num);
void replace_line(TextEditor *editor, int line_num, const char *text);
int search_text(TextEditor *editor, const char *text);
void replace_all(TextEditor *editor, const char *old_text, const char *new_text);
void display_content(TextEditor *editor);
void display_status(TextEditor *editor);
void expand_lines(TextEditor *editor);

int main() {
    TextEditor *editor = create_editor();
    if (!editor) {
        fprintf(stderr, "Failed to create text editor\n");
    }
}
```

```

        return 1;
    }

    char filename[256];
    printf("Enter filename to open (or new filename): ");
    scanf(" %[^\\n]", filename);

    if (load_file(editor, filename)) {
        printf("File loaded successfully\\n");
    } else {
        printf("Creating new file: %s\\n", filename);
        strcpy(editor->filename, filename);
    }

    int choice;
    do {
        display_content(editor);
        display_status(editor);
        // Add menu options
        printf("1. Insert line\\n");
        printf("2. Delete line\\n");
        printf("3. Replace line\\n");
        printf("4. Search text\\n");
        printf("5. Save file\\n");
        printf("6. Exit\\n");
        printf("Enter your choice: ");
        scanf("%d", &choice);

        // Implement menu options
    } while (choice != 6);

    destroy_editor(editor);
    return 0;
}

// Implement all functions here

```

Common Pitfalls to Avoid

1. Memory leaks from dynamic string allocation
2. Buffer overflows when reading input
3. Not handling file operation errors
4. Incorrect implementation of undo/redo functionality
5. Performance issues with large files

Best Practices

1. Use dynamic memory allocation carefully
2. Validate all file operations
3. Implement proper error handling
4. Use efficient data structures for text storage
5. Provide clear user feedback
6. Handle edge cases in editing operations

Project 3: Library Management System

Description

Create a system to manage a library's collection of books, borrowers, and transactions.

Learning Objectives

- Working with complex data relationships
- File I/O with structured data
- Data validation and integrity
- Search and sorting algorithms
- Transaction management

Requirements

1. Manage book inventory (title, author, ISBN, quantity)
2. Manage borrower information (ID, name, contact info)
3. Handle book borrowing and returning
4. Track overdue books
5. Generate reports (borrowed books, overdue books, popular books)
6. Save and load data from files
7. Search books by various criteria
8. Handle multiple copies of the same book

Implementation Steps

1. Design data structures for books, borrowers, and transactions
2. Implement CRUD operations for books and borrowers
3. Create borrowing and returning functionality
4. Add overdue tracking
5. Implement file I/O for data persistence
6. Create search and reporting features
7. Build a menu-driven interface
8. Add data validation

Sample Code Structure

```
#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <time.h>

#define MAX_TITLE_LENGTH 100
#define MAX_AUTHOR_LENGTH 50
#define MAX_BORROWERS 1000
#define MAX_BOOKS 1000

typedef struct {
    int isbn;
    char title[MAX_TITLE_LENGTH];
    char author[MAX_AUTHOR_LENGTH];
    int total_copies;
    int available_copies;
} Book;

typedef struct {
    int id;
    char name[50];
    char contact[50];
    int borrowed_count;
} Borrower;

typedef struct {
    int book_isbn;
    int borrower_id;
    time_t borrow_date;
    time_t due_date;
    time_t return_date; // 0 if not returned
} Transaction;

typedef struct {
    Book books[MAX_BOOKS];
    int book_count;

    Borrower borrowers[MAX_BORROWERS];
    int borrower_count;

    Transaction transactions[10000];
    int transaction_count;
} Library;
```

```

// Function prototypes
Library* create_library();
void destroy_library(Library *library);
int add_book(Library *library, Book book);
int remove_book(Library *library, int isbn);
int add_borrower(Library *library, Borrower borrower);
int remove_borrower(Library *library, int id);
int borrow_book(Library *library, int isbn, int borrower_id);
int return_book(Library *library, int isbn, int borrower_id);
Book* find_book(Library *library, int isbn);
Borrower* find_borrower(Library *library, int id);
void search_books(Library *library, const char *query);
void generate_overdue_report(Library *library);
void save_library(Library *library, const char *filename);
int load_library(Library *library, const char *filename);
void display_menu();

int main() {
    Library *library = create_library();
    if (!library) {
        fprintf(stderr, "Failed to create library\n");
        return 1;
    }

    // Load existing data
    load_library(library, "library.dat");

    int choice;
    do {
        display_menu();
        printf("Enter your choice: ");
        scanf("%d", &choice);

        switch (choice) {
            case 1: {
                Book book;
                printf("Enter ISBN: ");
                scanf("%d", &book.isbn);
                printf("Enter title: ");
                scanf(" %[^\n]", book.title);
                printf("Enter author: ");
                scanf(" %[^\n]", book.author);
                printf("Enter total copies: ");
                scanf("%d", &book.total_copies);
                book.available_copies = book.total_copies;
                add_book(library, book);
                break;
            }
        }
    } while (choice != 0);
}

```

```

    }
    // Implement other cases
    case 9:
        save_library(library, "library.dat");
        printf("Library data saved successfully!\n");
        break;
    case 10:
        printf("Goodbye!\n");
        break;
    default:
        printf("Invalid choice. Please try again.\n");
    }
} while (choice != 10);

destroy_library(library);
return 0;
}

// Implement all functions here

```

Common Pitfalls to Avoid

1. Data inconsistency between related entities
2. Not handling date/time calculations correctly
3. Memory leaks in dynamic structures
4. Not validating transaction integrity
5. Performance issues with large datasets

Best Practices

1. Maintain data integrity between related entities
2. Use proper date/time functions
3. Implement efficient search algorithms
4. Validate all transactions
5. Handle edge cases in borrowing logic
6. Provide clear error messages
7. Use appropriate data structures for performance

Project 4: Simple Database Engine

Description

Create a basic database engine that can store, retrieve, and manipulate tabular data.

Learning Objectives

- Working with tabular data structures
- File I/O with binary data
- Query processing
- Indexing and search optimization
- Memory management

Requirements

1. Create and manage database tables
2. Define table schemas (column names and types)
3. Insert, update, and delete records
4. Query data with simple WHERE clauses
5. Save and load databases from files
6. Support basic data types (int, float, string)
7. Implement indexing for faster searches
8. Export data to CSV format

Implementation Steps

1. Design data structures for tables and records
2. Implement table creation and schema management
3. Create CRUD operations for records
4. Add query processing functionality
5. Implement file I/O for persistence
6. Add indexing for performance
7. Create CSV export functionality
8. Build a command-line interface

Sample Code Structure

```
#include <stdio.h>
#include <stdlib.h>
#include <string.h>

#define MAX_TABLES 100
#define MAX_COLUMNS 50
#define MAX_RECORDS 10000
#define MAX_STRING_LENGTH 256

typedef enum {
    TYPE_INT,
    TYPE_FLOAT,
    TYPE_STRING
}
```

```

} DataType;

typedef struct {
    char name[50];
    DataType type;
    int size;
} Column;

typedef struct {
    int id;
    void *data; // Pointer to actual data
} Record;

typedef struct {
    char name[50];
    Column columns[MAX_COLUMNS];
    int column_count;
    Record records[MAX_RECORDS];
    int record_count;
    int *index; // Simple index array
    int index_count;
} Table;

typedef struct {
    Table tables[MAX_TABLES];
    int table_count;
} Database;

// Function prototypes
Database* create_database();
void destroy_database(Database *db);
Table* create_table(Database *db, const char *name);
int add_column(Table *table, const char *name, DataType type, int size);
int insert_record(Table *table, void *data[]);
Record* find_record(Table *table, int record_id);
int update_record(Table *table, int record_id, void *data[]);
int delete_record(Table *table, int record_id);
Record** query_table(Table *table, const char *where_clause, int *result_count);
void save_database(Database *db, const char *filename);
int load_database(Database *db, const char *filename);
void export_to_csv(Table *table, const char *filename);
void display_menu();

int main() {
    Database *db = create_database();
    if (!db) {
        fprintf(stderr, "Failed to create database\n");
    }
}

```



```

        return 1;
    }

    // Load existing database
    load_database(db, "database.dat");

    int choice;
    do {
        display_menu();
        printf("Enter your choice: ");
        scanf("%d", &choice);

        switch (choice) {
            case 1: {
                char table_name[50];
                printf("Enter table name: ");
                scanf(" %[^\\n]", table_name);
                create_table(db, table_name);
                break;
            }
            // Implement other cases
            case 7:
                save_database(db, "database.dat");
                printf("Database saved successfully!\\n");
                break;
            case 8:
                printf("Goodbye!\\n");
                break;
            default:
                printf("Invalid choice. Please try again.\\n");
        }
    } while (choice != 8);

    destroy_database(db);
    return 0;
}

// Implement all functions here

```

Common Pitfalls to Avoid

1. Memory leaks from dynamic data allocation
2. Data type mismatches
3. Not handling file format compatibility
4. Performance issues with large datasets
5. Not implementing proper indexing strategies

Best Practices

1. Use appropriate data structures for different data types
2. Validate data types during operations
3. Implement efficient indexing strategies
4. Handle file format versioning
5. Provide clear error messages for query syntax
6. Use constants for size limits
7. Implement proper memory management

Tips for Success

1. **Plan Before Coding:** Design your data structures and algorithms before implementation
2. **Modular Approach:** Break complex problems into smaller, manageable functions
3. **Error Handling:** Always consider what could go wrong and handle it gracefully
4. **Memory Management:** Be careful with dynamic allocation and always free memory
5. **Testing:** Test each function individually before integrating
6. **Performance:** Consider the efficiency of your algorithms, especially for large datasets
7. **Documentation:** Comment your code, especially complex algorithms
8. **Version Control:** Use version control to track changes and experiment safely

These intermediate projects will help you develop more advanced C programming skills while working on realistic applications. Focus on writing efficient, maintainable code with proper error handling.

Appendices

Appendix A: Additional Resources

Books and Publications

Core C Programming

1. **“The C Programming Language”** by Brian W. Kernighan and Dennis M. Ritchie
 - Often referred to as K&R, this is the definitive reference for C programming
 - Written by the creators of the C language
 - Essential for understanding the foundations of C
2. **“C Programming: A Modern Approach”** by K. N. King
 - Comprehensive guide covering both basic and advanced C concepts
 - Includes modern C standards (C99, C11)
 - Excellent for self-study with exercises and solutions
3. **“Expert C Programming: Deep C Secrets”** by Peter van der Linden
 - Advanced C programming techniques and lesser-known features
 - Covers complex topics like pointers, memory management, and optimization
 - Provides insights into the internals of C implementations

Systems Programming

4. **“Advanced Programming in the UNIX Environment”** by W. Richard Stevens
 - Essential reference for UNIX/Linux systems programming
 - Covers file I/O, process control, inter-process communication
 - Detailed explanations of system calls and library functions
5. **“UNIX Network Programming”** by W. Richard Stevens
 - Comprehensive guide to network programming in C
 - Covers sockets, protocols, and network applications
 - Invaluable for developing networked applications

Data Structures and Algorithms

6. **“Algorithms in C”** by Robert Sedgewick
 - Implementation of fundamental algorithms in C
 - Covers sorting, searching, graph algorithms, and more

- Excellent reference for algorithmic problem-solving

7. **“Data Structures and Algorithms in C”** by Michael T. Goodrich

- Comprehensive coverage of data structures with C implementations
- Includes analysis of time and space complexity
- Practical approach to algorithm design and implementation

Online Resources

Official Documentation

1. **ISO C Standard Documents**

- C17/C18: ISO/IEC 9899:2018
- C11: ISO/IEC 9899:2011
- C99: ISO/IEC 9899:1999
- Available from ISO for purchase, but drafts are often available

2. **GNU C Library Documentation**

- Comprehensive reference for glibc functions
- Available at: <https://www.gnu.org/software/libc/manual/>
- Detailed explanations of standard library functions

3. **POSIX Standards**

- Portable Operating System Interface standards
- Defines system interfaces and command-line shells
- Available at: <https://pubs.opengroup.org/onlinepubs/9699919799/>

Tutorials and Learning Platforms

4. **GeeksforGeeks C Programming**

- Extensive collection of C programming tutorials
- Covers basic to advanced topics with examples
- Available at: <https://www.geeksforgeeks.org/c-programming-language/>

5. **C Programming.com**

- Beginner-friendly tutorials and references
- Interactive compiler for testing code examples
- Available at: <https://www.cprogramming.com/>

6. **Learn-C.org**

- Free interactive C tutorial
- Browser-based coding environment
- Available at: <https://www.learn-c.org/>

Reference Sites

7. **cppreference.com**

- Comprehensive C and C++ reference
- Detailed documentation of standard library functions
- Available at: <https://en.cppreference.com/w/c>

8. **Linux Man Pages**

- System documentation for UNIX/Linux commands and functions
- Accessible via **man** command or online at: <https://man7.org/linux/man-pages/>
- Essential for systems programming

Development Tools

Compilers

1. **GCC (GNU Compiler Collection)**

- The most widely used C compiler
- Supports all C standards and many extensions
- Available at: <https://gcc.gnu.org/>

2. **Clang/LLVM**

- Modern C compiler with excellent diagnostics
- Fast compilation and good optimization
- Available at: <https://clang.llvm.org/>

3. **Microsoft Visual C++**

- Microsoft's C compiler for Windows
- Integrated with Visual Studio IDE
- Part of Visual Studio Community (free)

Debugging Tools

4. **GDB (GNU Debugger)**

- Powerful command-line debugger for C programs
- Supports breakpoints, watchpoints, and stack tracing
- Available at: <https://www.gnu.org/software/gdb/>

5. **Valgrind**

- Memory debugging and profiling tool
- Detects memory leaks, invalid memory access, and more
- Available at: <https://valgrind.org/>

6. **AddressSanitizer**

- Fast memory error detector
- Part of GCC and Clang compilers
- Detects buffer overflows, use-after-free, and other errors

Profiling Tools

7. gprof

- GNU profiler for analyzing program performance
- Shows function call counts and execution times
- Part of GNU binutils

8. perf

- Linux performance analysis tool
- Hardware performance counter support
- Available on most Linux distributions

Static Analysis Tools

9. Cppcheck

- Static analysis tool for C/C++ code
- Detects bugs, memory leaks, and style issues
- Available at: <http://cppcheck.sourceforge.net/>

10. Clang Static Analyzer

- Source code analysis tool for C/C++
- Finds bugs and potential security issues
- Part of Clang/LLVM project

Libraries and Frameworks

Standard Libraries

1. glibc (GNU C Library)

- The standard C library on most Linux systems
- Comprehensive implementation of ISO C standard
- Available at: <https://www.gnu.org/software/libc/>

2. musl libc

- Lightweight, standards-conforming C library
- Designed for static linking and embedded systems
- Available at: <https://musl.libc.org/>

Networking Libraries

3. libcurl

- Multiprotocol file transfer library
- Supports HTTP, FTP, SMTP, and many other protocols
- Available at: <https://curl.se/libcurl/>

4. libevent

- Event notification library for developing scalable network applications
- Supports event-driven programming model
- Available at: <http://libevent.org/>

Data Structure Libraries

5. GLib

- General-purpose utility library from GNOME project
- Provides data structures, string utilities, and more
- Available at: <https://developer.gnome.org/glib/>

6. Apache Portable Runtime (APR)

- Cross-platform library for system programming
- Provides memory management, file I/O, and threading utilities
- Available at: <https://apr.apache.org/>

Communities and Forums

Online Communities

1. Stack Overflow

- Large community of developers answering programming questions
- Extensive C programming tag with thousands of questions
- Available at: <https://stackoverflow.com/questions/tagged/c>

2. Reddit r/C_Programming

- Community discussion forum for C programmers
- News, discussions, and code reviews
- Available at: https://www.reddit.com/r/C_Programming/

3. C Board

- Dedicated C programming forum
- Friendly community for beginners and experts
- Available at: <http://cboard.cprogramming.com/>

Mailing Lists

4. **comp.lang.c**

- Usenet newsgroup for C programming discussions
- Long-running community with expert participants
- Available through Google Groups

Conferences and Events

Major Conferences

1. **ACM SIGPLAN Conference on Programming Language Design and Implementation (PLDI)**

- Premier conference for programming language research
- Includes C-related research and developments

2. **USENIX Annual Technical Conference**

- Focuses on systems implementation and applications
- Relevant for systems programming in C

3. **Linux Plumbers Conference**

- Focuses on Linux kernel and systems development
- Many sessions relevant to C systems programming

Open Source Projects

Learning from Real Code

1. **Linux Kernel**

- Largest open-source C project
- Excellent example of systems programming
- Available at: <https://github.com/torvalds/linux>

2. **Git**

- Distributed version control system written in C
- Well-designed codebase with good documentation
- Available at: <https://github.com/git/git>

3. **Redis**

- In-memory data structure store written in C
- Clean, well-documented codebase
- Available at: <https://github.com/redis/redis>

4. **Nginx**

- High-performance web server written in C
- Excellent example of network programming
- Available at: <https://github.com/nginx/nginx>

Practice Platforms

Coding Challenges

1. LeetCode

- Programming practice platform with C support
- Thousands of coding problems with solutions
- Available at: <https://leetcode.com/>

2. HackerRank

- Coding practice and competition platform
- Dedicated C programming section
- Available at: <https://www.hackerrank.com/>

3. Project Euler

- Mathematical and computational programming challenges
- Excellent for practicing algorithms and optimization
- Available at: <https://projecteuler.net/>

Standards Organizations

Key Organizations

1. ISO/IEC JTC1/SC22/WG14

- International standardization working group for C
- Responsible for the ISO C standard
- Website: <https://www.open-std.org/jtc1/sc22/wg14/>

2. ANSI C Standards Committee

- American National Standards Institute committee
- Works with ISO on C standard development

Additional Learning Resources

Video Courses

1. CS50 (Harvard)

- Introduction to Computer Science with C programming

- Free online course with video lectures
- Available at: <https://cs50.harvard.edu/>

2. MIT OpenCourseWare

- Free lecture videos and materials for computer science courses
- Several courses include C programming content
- Available at: <https://ocw.mit.edu/>

Podcasts

1. C Developers Journal

- Podcast focused on C programming and systems development
- Interviews with C experts and discussions of best practices

2. Software Engineering Daily

- Daily interviews with software engineers
- Many episodes relevant to systems programming in C

This appendix provides a comprehensive list of resources to help you continue your journey in C programming. Whether you're looking to deepen your understanding of specific topics, find tools to improve your development workflow, or connect with the C programming community, these resources will serve as valuable references.

Appendix B: C Language Reference

C Keywords

Standard C Keywords (C17/C18)

```
auto      break   case    char    const   continue default do
double    else    enum    extern  float   for      goto   if
inline    int     long    register return  short   signed  sizeof
static    struct  switch  typedef union   unsigned void    volatile
while     _Alignas _Alignof _Atomic _Bool   _Complex _Generic
_Imaginary _Noreturn _Static_assert _Thread_local
```

Deprecated Keywords

- restrict (C99-C17, still available but not commonly used)

Data Types and Sizes

Basic Data Types

Type	Typical Size (64-bit)	Range
char	1 byte	-128 to 127 or 0 to 255
unsigned char	1 byte	0 to 255
short	2 bytes	-32,768 to 32,767
unsigned short	2 bytes	0 to 65,535
int	4 bytes	-2,147,483,648 to 2,147,483,647
unsigned int	4 bytes	0 to 4,294,967,295
long	8 bytes	-9,223,372,036,854,775,808 to 9,223,372,036,854,775,807
unsigned long	8 bytes	0 to 18,446,744,073,709,551,615

Type	Typical Size (64-bit)	Range
float	4 bytes	1.2E-38 to 3.4E+38 (6 decimal places)
double	8 bytes	2.3E-308 to 1.7E+308 (15 decimal places)
long double	16 bytes	3.4E-4932 to 1.1E+4932 (19 decimal places)

Type Specifiers

- **signed**: Explicitly signed integers
- **unsigned**: Non-negative integers
- **short**: Shorter integers
- **long**: Longer integers
- **long long**: Even longer integers (C99)

Type Qualifiers

- **const**: Read-only data
- **volatile**: Data that may change unexpectedly
- **restrict**: Pointer optimization hint (C99)
- **_Atomic**: Atomic operations (C11)

Operators

Arithmetic Operators

Operator	Description	Example
+	Addition	a + b
-	Subtraction	a - b
*	Multiplication	a * b
/	Division	a / b
%	Modulus (remainder)	a % b
++	Increment	++a or a++
-	Decrement	-a or a--

Relational Operators

Operator	Description	Example
==	Equal to	a == b
!=	Not equal to	a != b
>	Greater than	a > b
<	Less than	a < b
>=	Greater than or equal to	a >= b
<=	Less than or equal to	a <= b

Logical Operators

Operator	Description	Example
&&	Logical AND	a && b
	Logical OR	a b
!	Logical NOT	!a

Bitwise Operators

Operator	Description	Example
&	Bitwise AND	a & b
	Bitwise OR	a b
^	Bitwise XOR	a ^ b
~	Bitwise NOT (one's complement)	~a
«	Left shift	a « b
»	Right shift	a » b

Assignment Operators

Operator	Description	Example	Equivalent
=	Simple assignment	a = b	a = b
+=	Add and assign	a += b	a = a + b
-=	Subtract and assign	a -= b	a = a - b
*=	Multiply and assign	a *= b	a = a * b
/=	Divide and assign	a /= b	a = a / b
%=	Modulus and assign	a %= b	a = a % b
&=	Bitwise AND and assign	a &= b	a = a & b
=	Bitwise OR and assign	a = b	a = a b
^=	Bitwise XOR and assign	a ^= b	a = a ^ b
«=	Left shift and assign	a «= b	a = a « b
»=	Right shift and assign	a »= b	a = a » b

Other Operators

Operator	Description	Example
sizeof	Size of data type or variable	sizeof(int)
&	Address of	&a
*	Value at address (dereference)	*ptr
?:	Conditional (ternary)	condition ? a : b
,	Comma	a, b
.	Member access	struct.member
->	Pointer member access	ptr->member
[]	Array subscript	arr[i]
()	Function call	func()

Standard Library Functions

Input/Output Functions (<stdio.h>)

```
// Formatted I/O
int printf(const char *format, ...);
int scanf(const char *format, ...);
int fprintf(FILE *stream, const char *format, ...);
int fscanf(FILE *stream, const char *format, ...);
int sprintf(char *str, const char *format, ...);
int snprintf(char *str, size_t size, const char *format, ...);

// Character I/O
int getchar(void);
int putchar(int c);
int fgetc(FILE *stream);
int fputc(int c, FILE *stream);

// String I/O
char *fgets(char *str, int n, FILE *stream);
int fputs(const char *str, FILE *stream);

// File operations
FILE *fopen(const char *filename, const char *mode);
int fclose(FILE *stream);
size_t fread(void *ptr, size_t size, size_t nmemb, FILE *stream);
size_t fwrite(const void *ptr, size_t size, size_t nmemb, FILE *stream);
int fseek(FILE *stream, long offset, int whence);
long ftell(FILE *stream);
void rewind(FILE *stream);
```

String Functions (<string.h>)

```
// String copying
char *strcpy(char *dest, const char *src);
char *strncpy(char *dest, const char *src, size_t n);

// String concatenation
char *strcat(char *dest, const char *src);
char *strncat(char *dest, const char *src, size_t n);

// String comparison
int strcmp(const char *s1, const char *s2);
int strncmp(const char *s1, const char *s2, size_t n);

// String searching
char *strchr(const char *s, int c);
char *strrchr(const char *s, int c);
char *strstr(const char *haystack, const char *needle);
size_t strlen(const char *s);

// Memory operations
void *memcpy(void *dest, const void *src, size_t n);
void *memmove(void *dest, const void *src, size_t n);
int memcmp(const void *s1, const void *s2, size_t n);
void *memset(void *s, int c, size_t n);
```

Memory Management Functions (<stdlib.h>)

```
void *malloc(size_t size);
void *calloc(size_t nmemb, size_t size);
void *realloc(void *ptr, size_t size);
void free(void *ptr);
```

Mathematical Functions (<math.h>)

```
// Trigonometric functions
double sin(double x);
double cos(double x);
double tan(double x);
double asin(double x);
double acos(double x);
double atan(double x);
double atan2(double y, double x);
```



```

// Hyperbolic functions
double sinh(double x);
double cosh(double x);
double tanh(double x);

// Exponential and logarithmic functions
double exp(double x);
double log(double x);    // Natural logarithm
double log10(double x);  // Base-10 logarithm
double pow(double x, double y);
double sqrt(double x);

// Rounding and absolute value
double ceil(double x);
double floor(double x);
double fabs(double x);
double fmod(double x, double y);

```

Utility Functions (<stdlib.h>)

```

// Conversion functions
int atoi(const char *nptr);
long atol(const char *nptr);
double atof(const char *nptr);
long int strtol(const char *nptr, char **endptr, int base);
unsigned long int strtoul(const char *nptr, char **endptr, int base);

// Random number generation
int rand(void);
void srand(unsigned int seed);

// Searching and sorting
void *bsearch(const void *key, const void *base,
              size_t nmemb, size_t size,
              int (*compar)(const void *, const void *));
void qsort(void *base, size_t nmemb, size_t size,
           int (*compar)(const void *, const void *));

// Process control
void exit(int status);
void abort(void);
int atexit(void (*func)(void));

```

Time Functions (<time.h>)

```
time_t time(time_t *tloc);
char *ctime(const time_t *timep);
struct tm *localtime(const time_t *timep);
struct tm *gmtime(const time_t *timep);
time_t mktime(struct tm *tm);
size_t strftime(char *s, size_t max, const char *format,
                const struct tm *tm);
```

Preprocessor Directives

Common Directives

```
#define MACRO_NAME value
#undef MACRO_NAME
#include <header.h>
#include "file.h"
#if condition
#ifdef MACRO
#ifndef MACRO
#else
#elif condition
#endif
#error message
#pragma directive
#line number "filename"
```

Standard Macros

```
__LINE__      // Current line number
__FILE__      // Current filename
__DATE__      // Compilation date
__TIME__      // Compilation time
__STDC__      // Standard C compliance
__STDC_VERSION__ // C standard version
__STDC_HOSTED__ // Hosted implementation
```

Format Specifiers

printf Format Specifiers

Specifier	Description	Example
%d, %i	Signed decimal integer	printf(“%d”, 42);
%u	Unsigned decimal integer	printf(“%u”, 42u);
%o	Octal integer	printf(“%o”, 42);
%x, %X	Hexadecimal integer	printf(“%x”, 42);
%f, %F	Decimal floating point	printf(“%f”, 3.14);
%e, %E	Scientific notation	printf(“%e”, 3.14);
%g, %G	Use %f or %e, whichever is shorter	printf(“%g”, 3.14);
%c	Character	printf(“%c”, ‘A’);
%s	String	printf(“%s”, “Hello”);
%p	Pointer address	printf(“%p”, ptr);
%n	Number of characters written so far	printf(“%n”, &count);

scanf Format Specifiers

Specifier	Description	Example
%d	Signed decimal integer	scanf(“%d”, &num);
%u	Unsigned decimal integer	scanf(“%u”, &num);
%o	Octal integer	scanf(“%o”, &num);
%x	Hexadecimal integer	scanf(“%x”, &num);
%f	Floating point	scanf(“%f”, &num);
%e	Floating point	scanf(“%e”, &num);
%g	Floating point	scanf(“%g”, &num);
%c	Character	scanf(“%c”, &ch);
%s	String	scanf(“%s”, str);
%p	Pointer	scanf(“%p”, &ptr);
%[...]	Character set	scanf(“%[aeiou]”, str);

Format Modifiers

Modifier	Description	Example
hh	char (char)	printf(“%hhd”, c);
h	short (short)	printf(“%hd”, s);
l	long (long)	printf(“%ld”, l);
ll	long long (long long)	printf(“%lld”, ll);
j	intmax_t	printf(“%jd”, imax);
z	size_t	printf(“%zd”, sz);
t	ptrdiff_t	printf(“%td”, ptrdiff);
L	long double	printf(“%Lf”, ld);

Escape Sequences

Standard Escape Sequences

Sequence	Description	ASCII Value
\a	Alert (bell)	007
\b	Backspace	008
\f	Form feed	012
\n	New line	010
\r	Carriage return	013
\t	Horizontal tab	009
\v	Vertical tab	011
\\	Backslash	092
\'	Single quote	039
\"	Double quote	034
\?	Question mark	063
\0	Null character	000

Octal and Hexadecimal Escapes

Sequence	Description	Example
\\ooo	Octal character	\\040 (space)
\\xhh	Hexadecimal character	\\x20 (space)

Storage Classes

Storage Class Specifiers

Specifier	Storage Duration	Linkage	Scope
auto	Automatic	None	Block
register	Automatic	None	Block
static	Static	Internal	File/Block
extern	Static	External	File
typedef	N/A	N/A	File

Error Handling

errno Values (<errno.h>)

```
EDOM    // Domain error
ERANGE  // Range error
EILSEQ  // Illegal byte sequence
```

Assertions (<assert.h>)

```
void assert(int expression);
```

Signal Handling (<signal.h>)

Signal Types

```
SIGABRT // Abnormal termination
SIGFPE  // Floating-point error
SIGILL  // Illegal instruction
SIGINT  // Interactive attention signal
SIGSEGV // Invalid memory access
SIGTERM // Termination request
```

Signal Functions

```
void (*signal(int sig, void (*func)(int)))(int);
int raise(int sig);
```

File Operations Modes

fopen Modes

Mode	Description	File Position	File Creation
"r"	Read only	Beginning	No
"w"	Write only	Beginning	Yes (truncates)
"a"	Append only	End	Yes
"r+"	Read/Write	Beginning	No
"w+"	Read/Write	Beginning	Yes (truncates)

Mode	Description	File Position	File Creation
"a+"	Read/Append	End	Yes

Binary Mode Suffix

Add 'b' to any mode for binary files (e.g., "rb", "wb", "ab")

Standard Streams

Predefined Streams

```
stdin    // Standard input (keyboard)
stdout   // Standard output (screen)
stderr   // Standard error (screen)
```

Limits and Constants

Implementation Limits (<limits.h>)

```
CHAR_BIT    // Bits in a char (8)
CHAR_MAX    // Maximum char value
CHAR_MIN    // Minimum char value
INT_MAX     // Maximum int value
INT_MIN     // Minimum int value
LONG_MAX    // Maximum long value
LONG_MIN    // Minimum long value
SCHAR_MAX   // Maximum signed char value
SCHAR_MIN   // Minimum signed char value
SHRT_MAX    // Maximum short value
SHRT_MIN    // Minimum short value
UCHAR_MAX   // Maximum unsigned char value
UINT_MAX    // Maximum unsigned int value
ULONG_MAX   // Maximum unsigned long value
USHRT_MAX   // Maximum unsigned short value
```

Floating-Point Limits (<float.h>)

```
FLT_DIG      // Decimal digits of precision
FLT_EPSILON  // Smallest x where 1.0 + x != 1.0
FLT_MANT_DIG // Number of digits in mantissa
FLT_MAX      // Maximum float value
FLT_MAX_10_EXP // Maximum decimal exponent
FLT_MIN      // Minimum float value
FLT_MIN_10_EXP // Minimum decimal exponent
DBL_DIG      // Decimal digits of precision (double)
```

Thread Support (C11) (<threads.h>)

Thread Functions

```
int thrd_create(thrd_t *thr, thrd_start_t func, void *arg);
int thrd_equal(thrd_t thr0, thrd_t thr1);
thrd_t thrd_current(void);
int thrd_sleep(const struct timespec *duration, struct timespec *remaining);
void thrd_yield(void);
_Noreturn void thrd_exit(int res);
int thrd_detach(thrd_t thr);
int thrd_join(thrd_t thr, int *res);
```

Mutex Functions

```
int mtx_init(mtx_t *mtx, int type);
void mtx_destroy(mtx_t *mtx);
int mtx_lock(mtx_t *mtx);
int mtx_timedlock(mtx_t *mtx, const struct timespec *ts);
int mtx_trylock(mtx_t *mtx);
int mtx_unlock(mtx_t *mtx);
```

Condition Variable Functions

```
int cnd_init(cnd_t *cnd);
void cnd_destroy(cnd_t *cnd);
int cnd_signal(cnd_t *cnd);
int cnd_broadcast(cnd_t *cnd);
int cnd_wait(cnd_t *cnd, mtx_t *mtx);
int cnd_timedwait(cnd_t *cnd, mtx_t *mtx, const struct timespec *ts);
```

Atomic Operations (C11) (<stdatomic.h>)

Atomic Types

```
_Atomic bool
_Atomic char
_Atomic short
_Atomic int
_Atomic long
_Atomic long long
_Atomic unsigned char
_Atomic unsigned short
_Atomic unsigned int
_Atomic unsigned long
_Atomic unsigned long long
```

Atomic Operations

```
memory_order memory_order_relaxed;
memory_order memory_order_consume;
memory_order memory_order_acquire;
memory_order memory_order_release;
memory_order memory_order_acq_rel;
memory_order memory_order_seq_cst;

#define ATOMIC_VAR_INIT(value)
#define ATOMIC_FLAG_INIT

bool atomic_flag_test_and_set(volatile atomic_flag *object);
void atomic_flag_clear(volatile atomic_flag *object);
```

Generic Selection (C11) (<stdnoreturn.h>)

_Generic Keyword

```
#define MAX(x, y) _Generic((x), \
    int: max_int, \
    float: max_float, \
    double: max_double \
)(x, y)
```


Alignment (C11) (<stdalign.h>)

Alignment Specifiers

```
_Alignas(expression)  
alignas(type)  
_Alignof(type)  
alignof(type)
```

This reference appendix provides a comprehensive overview of the C programming language, covering keywords, operators, standard library functions, and other essential elements. It serves as a quick reference for C programmers of all levels.

Appendix C: Exercise Solutions

This appendix provides solutions to selected exercises from each module of the course. These solutions are meant to serve as examples of good programming practices and to help you understand the concepts covered in each module.

Module 1: Foundations

Exercise 1.1: Hello World Variations

```
#include <stdio.h>

int main() {
    // Basic Hello World
    printf("Hello, World!\n");

    // Personalized greeting
    printf("Hello, C Programmer!\n");

    // Multi-line message
    printf("Welcome to C programming.\n");
    printf("This is the beginning of your journey.\n");

    return 0;
}
```

Exercise 1.2: Basic Calculator

```
#include <stdio.h>

int main() {
    double num1, num2;
    char operator;

    printf("Enter first number: ");
    scanf("%lf", &num1);

    printf("Enter an operator (+, -, *, /): ");
```

```

scanf(" %c", &operator);

printf("Enter second number: ");
scanf("%lf", &num2);

switch (operator) {
    case '+':
        printf("%.2lf + %.2lf = %.2lf\n", num1, num2, num1 + num2);
        break;
    case '-':
        printf("%.2lf - %.2lf = %.2lf\n", num1, num2, num1 - num2);
        break;
    case '*':
        printf("%.2lf * %.2lf = %.2lf\n", num1, num2, num1 * num2);
        break;
    case '/':
        if (num2 != 0) {
            printf("%.2lf / %.2lf = %.2lf\n", num1, num2, num1 / num2);
        } else {
            printf("Error: Division by zero\n");
        }
        break;
    default:
        printf("Error: Invalid operator\n");
}

return 0;
}

```

Module 2: Data Types

Exercise 2.1: Size and Range Calculator

```

#include <stdio.h>
#include <limits.h>
#include <float.h>

int main() {
    printf("Data Type Sizes and Ranges:\n");
    printf("=====\n");

    printf("char: %zu bytes\n", sizeof(char));
    printf("  Range: %d to %d\n", CHAR_MIN, CHAR_MAX);
}

```

```

printf("unsigned char: %zu bytes\n", sizeof(unsigned char));
printf("  Range: 0 to %u\n", UCHAR_MAX);

printf("short: %zu bytes\n", sizeof(short));
printf("  Range: %d to %d\n", SHRT_MIN, SHRT_MAX);

printf("unsigned short: %zu bytes\n", sizeof(unsigned short));
printf("  Range: 0 to %u\n", USHRT_MAX);

printf("int: %zu bytes\n", sizeof(int));
printf("  Range: %d to %d\n", INT_MIN, INT_MAX);

printf("unsigned int: %zu bytes\n", sizeof(unsigned int));
printf("  Range: 0 to %u\n", UINT_MAX);

printf("long: %zu bytes\n", sizeof(long));
printf("  Range: %ld to %ld\n", LONG_MIN, LONG_MAX);

printf("unsigned long: %zu bytes\n", sizeof(unsigned long));
printf("  Range: 0 to %lu\n", ULONG_MAX);

printf("float: %zu bytes\n", sizeof(float));
printf("  Range: %e to %e\n", FLT_MIN, FLT_MAX);
printf("  Precision: %d digits\n", FLT_DIG);

printf("double: %zu bytes\n", sizeof(double));
printf("  Range: %e to %e\n", DBL_MIN, DBL_MAX);
printf("  Precision: %d digits\n", DBL_DIG);

return 0;
}

```

Exercise 2.2: Temperature Converter

```

#include <stdio.h>

double celsius_to_fahrenheit(double celsius) {
    return (celsius * 9.0 / 5.0) + 32.0;
}

double fahrenheit_to_celsius(double fahrenheit) {
    return (fahrenheit - 32.0) * 5.0 / 9.0;
}

int main() {

```

```

double temperature;
int choice;

printf("Temperature Converter\n");
printf("1. Celsius to Fahrenheit\n");
printf("2. Fahrenheit to Celsius\n");
printf("Enter your choice (1 or 2): ");
scanf("%d", &choice);

if (choice == 1) {
    printf("Enter temperature in Celsius: ");
    scanf("%lf", &temperature);
    printf("%.2lf°C = %.2lf°F\n", temperature, celsius_to_fahrenheit(temperature));
} else if (choice == 2) {
    printf("Enter temperature in Fahrenheit: ");
    scanf("%lf", &temperature);
    printf("%.2lf°F = %.2lf°C\n", temperature, fahrenheit_to_celsius(temperature));
} else {
    printf("Invalid choice\n");
}

return 0;
}

```

Module 3: Control Flow

Exercise 3.1: Prime Number Checker

```

#include <stdio.h>
#include <stdbool.h>

bool is_prime(int n) {
    if (n <= 1) return false;
    if (n <= 3) return true;
    if (n % 2 == 0 || n % 3 == 0) return false;

    for (int i = 5; i * i <= n; i += 6) {
        if (n % i == 0 || n % (i + 2) == 0) {
            return false;
        }
    }
    return true;
}

int main() {

```

```

    int number;

    printf("Enter a number: ");
    scanf("%d", &number);

    if (is_prime(number)) {
        printf("%d is a prime number\n", number);
    } else {
        printf("%d is not a prime number\n", number);
    }

    return 0;
}

```

Exercise 3.2: Pattern Printer

```

#include <stdio.h>

void print_triangle(int n) {
    for (int i = 1; i <= n; i++) {
        // Print spaces
        for (int j = 1; j <= n - i; j++) {
            printf(" ");
        }
        // Print stars
        for (int j = 1; j <= 2 * i - 1; j++) {
            printf("*");
        }
        printf("\n");
    }
}

void print_diamond(int n) {
    // Upper half (including middle)
    for (int i = 1; i <= n; i++) {
        for (int j = 1; j <= n - i; j++) {
            printf(" ");
        }
        for (int j = 1; j <= 2 * i - 1; j++) {
            printf("*");
        }
        printf("\n");
    }

    // Lower half

```

```

        for (int i = n - 1; i >= 1; i--) {
            for (int j = 1; j <= n - i; j++) {
                printf(" ");
            }
            for (int j = 1; j <= 2 * i - 1; j++) {
                printf("*");
            }
            printf("\n");
        }
    }

int main() {
    int n;
    int choice;

    printf("Pattern Printer\n");
    printf("1. Triangle\n");
    printf("2. Diamond\n");
    printf("Enter your choice (1 or 2): ");
    scanf("%d", &choice);

    printf("Enter the size: ");
    scanf("%d", &n);

    if (choice == 1) {
        print_triangle(n);
    } else if (choice == 2) {
        print_diamond(n);
    } else {
        printf("Invalid choice\n");
    }

    return 0;
}

```

Module 4: Functions

Exercise 4.1: Mathematical Functions Library

```

#include <stdio.h>
#include <math.h>

// Function to calculate factorial
long long factorial(int n) {

```

```

    if (n < 0) return -1; // Error case
    if (n == 0 || n == 1) return 1;

    long long result = 1;
    for (int i = 2; i <= n; i++) {
        result *= i;
    }
    return result;
}

// Function to calculate power
double power(double base, int exponent) {
    if (exponent == 0) return 1.0;

    double result = 1.0;
    int abs_exp = (exponent < 0) ? -exponent : exponent;

    for (int i = 0; i < abs_exp; i++) {
        result *= base;
    }

    return (exponent < 0) ? 1.0 / result : result;
}

// Function to calculate GCD using Euclidean algorithm
int gcd(int a, int b) {
    while (b != 0) {
        int temp = b;
        b = a % b;
        a = temp;
    }
    return a;
}

// Function to calculate LCM
int lcm(int a, int b) {
    return (a * b) / gcd(a, b);
}

int main() {
    int choice, a, b;
    double x;
    int n;

    printf("Mathematical Functions Library\n");
    printf("1. Factorial\n");
    printf("2. Power\n");

```



```

printf("3. GCD\n");
printf("4. LCM\n");
printf("Enter your choice (1-4): ");
scanf("%d", &choice);

switch (choice) {
    case 1:
        printf("Enter a number: ");
        scanf("%d", &n);
        if (n >= 0) {
            printf("Factorial of %d = %lld\n", n, factorial(n));
        } else {
            printf("Factorial is not defined for negative numbers\n");
        }
        break;

    case 2:
        printf("Enter base: ");
        scanf("%lf", &x);
        printf("Enter exponent: ");
        scanf("%d", &n);
        printf("%.2lf^%d = %.2lf\n", x, n, power(x, n));
        break;

    case 3:
        printf("Enter first number: ");
        scanf("%d", &a);
        printf("Enter second number: ");
        scanf("%d", &b);
        printf("GCD of %d and %d = %d\n", a, b, gcd(a, b));
        break;

    case 4:
        printf("Enter first number: ");
        scanf("%d", &a);
        printf("Enter second number: ");
        scanf("%d", &b);
        printf("LCM of %d and %d = %d\n", a, b, lcm(a, b));
        break;

    default:
        printf("Invalid choice\n");
}

return 0;
}

```

Exercise 4.2: Recursive Functions

```
#include <stdio.h>

// Recursive factorial
long long recursive_factorial(int n) {
    if (n < 0) return -1;
    if (n == 0 || n == 1) return 1;
    return n * recursive_factorial(n - 1);
}

// Recursive Fibonacci
int recursive_fibonacci(int n) {
    if (n <= 0) return 0;
    if (n == 1) return 1;
    return recursive_fibonacci(n - 1) + recursive_fibonacci(n - 2);
}

// Recursive power
double recursive_power(double base, int exponent) {
    if (exponent == 0) return 1.0;
    if (exponent > 0) {
        return base * recursive_power(base, exponent - 1);
    } else {
        return 1.0 / base * recursive_power(base, exponent + 1);
    }
}

int main() {
    int choice, n;
    double x;
    int exp;

    printf("Recursive Functions\n");
    printf("1. Factorial\n");
    printf("2. Fibonacci\n");
    printf("3. Power\n");
    printf("Enter your choice (1-3): ");
    scanf("%d", &choice);

    switch (choice) {
        case 1:
            printf("Enter a number: ");
            scanf("%d", &n);
            if (n >= 0) {
                printf("Factorial of %d = %lld\n", n, recursive_factorial(n));
            } else {

```

```

        printf("Factorial is not defined for negative numbers\n");
    }
    break;

case 2:
    printf("Enter position: ");
    scanf("%d", &n);
    if (n >= 0) {
        printf("Fibonacci number at position %d = %d\n", n, recursive_fibonacci(n));
    } else {
        printf("Invalid position\n");
    }
    break;

case 3:
    printf("Enter base: ");
    scanf("%lf", &x);
    printf("Enter exponent: ");
    scanf("%d", &exp);
    printf("%.2lf^%d = %.2lf\n", x, exp, recursive_power(x, exp));
    break;

default:
    printf("Invalid choice\n");
}

return 0;
}

```

Module 5: Arrays and Strings

Exercise 5.1: Array Operations

```

#include <stdio.h>
#include <limits.h>

// Function to find maximum element in array
int find_max(int arr[], int size) {
    int max = arr[0];
    for (int i = 1; i < size; i++) {
        if (arr[i] > max) {
            max = arr[i];
        }
    }
}

```

```

        return max;
    }

    // Function to find minimum element in array
    int find_min(int arr[], int size) {
        int min = arr[0];
        for (int i = 1; i < size; i++) {
            if (arr[i] < min) {
                min = arr[i];
            }
        }
        return min;
    }

    // Function to calculate average
    double calculate_average(int arr[], int size) {
        long long sum = 0;
        for (int i = 0; i < size; i++) {
            sum += arr[i];
        }
        return (double)sum / size;
    }

    // Function to reverse array
    void reverse_array(int arr[], int size) {
        for (int i = 0; i < size / 2; i++) {
            int temp = arr[i];
            arr[i] = arr[size - 1 - i];
            arr[size - 1 - i] = temp;
        }
    }

    // Function to print array
    void print_array(int arr[], int size) {
        printf("[");
        for (int i = 0; i < size; i++) {
            printf("%d", arr[i]);
            if (i < size - 1) printf(", ");
        }
        printf("]\n");
    }

    int main() {
        int arr[] = {64, 34, 25, 12, 22, 11, 90};
        int size = sizeof(arr) / sizeof(arr[0]);

        printf("Original array: ");
    }

```

```

    print_array(arr, size);

    printf("Maximum element: %d\n", find_max(arr, size));
    printf("Minimum element: %d\n", find_min(arr, size));
    printf("Average: %.2f\n", calculate_average(arr, size));

    reverse_array(arr, size);
    printf("Reversed array: ");
    print_array(arr, size);

    return 0;
}

```

Exercise 5.2: String Manipulation

```

#include <stdio.h>
#include <string.h>
#include <ctype.h>

// Function to count vowels
int count_vowels(const char *str) {
    int count = 0;
    for (int i = 0; str[i] != '\0'; i++) {
        char ch = tolower(str[i]);
        if (ch == 'a' || ch == 'e' || ch == 'i' || ch == 'o' || ch == 'u') {
            count++;
        }
    }
    return count;
}

// Function to reverse string
void reverse_string(char *str) {
    int len = strlen(str);
    for (int i = 0; i < len / 2; i++) {
        char temp = str[i];
        str[i] = str[len - 1 - i];
        str[len - 1 - i] = temp;
    }
}

// Function to check palindrome
int is_palindrome(const char *str) {
    int len = strlen(str);
    for (int i = 0; i < len / 2; i++) {

```

```

        if (tolower(str[i]) != tolower(str[len - 1 - i])) {
            return 0;
        }
    }
    return 1;
}

// Function to remove spaces
void remove_spaces(char *str) {
    int i, j = 0;
    for (i = 0; str[i] != '\0'; i++) {
        if (str[i] != ' ') {
            str[j++] = str[i];
        }
    }
    str[j] = '\0';
}

int main() {
    char str[100];

    printf("Enter a string: ");
    fgets(str, sizeof(str), stdin);

    // Remove newline character if present
    str[strcspn(str, "\n")] = 0;

    printf("Original string: %s\n", str);
    printf("Number of vowels: %d\n", count_vowels(str));

    char str_copy[100];
    strcpy(str_copy, str);
    reverse_string(str_copy);
    printf("Reversed string: %s\n", str_copy);

    if (is_palindrome(str)) {
        printf("The string is a palindrome\n");
    } else {
        printf("The string is not a palindrome\n");
    }

    strcpy(str_copy, str);
    remove_spaces(str_copy);
    printf("String without spaces: %s\n", str_copy);

    return 0;
}

```

Module 6: Pointers

Exercise 6.1: Pointer Arithmetic

```
#include <stdio.h>

// Function to swap two integers using pointers
void swap(int *a, int *b) {
    int temp = *a;
    *a = *b;
    *b = temp;
}

// Function to find maximum in array using pointers
int* find_max_ptr(int *arr, int size) {
    int *max = arr;
    for (int i = 1; i < size; i++) {
        if (*(arr + i) > *max) {
            max = arr + i;
        }
    }
    return max;
}

// Function to reverse array using pointers
void reverse_array_ptr(int *arr, int size) {
    int *start = arr;
    int *end = arr + size - 1;

    while (start < end) {
        int temp = *start;
        *start = *end;
        *end = temp;
        start++;
        end--;
    }
}

// Function to print array using pointers
void print_array_ptr(int *arr, int size) {
    printf("[");
    for (int i = 0; i < size; i++) {
        printf("%d", *(arr + i));
        if (i < size - 1) printf(", ");
    }
    printf("]\n");
}
```

```

}

int main() {
    int arr[] = {64, 34, 25, 12, 22, 11, 90};
    int size = sizeof(arr) / sizeof(arr[0]);

    printf("Original array: ");
    print_array_ptr(arr, size);

    int a = 10, b = 20;
    printf("Before swap: a = %d, b = %d\n", a, b);
    swap(&a, &b);
    printf("After swap: a = %d, b = %d\n", a, b);

    int *max_ptr = find_max_ptr(arr, size);
    printf("Maximum element: %d at address %p\n", *max_ptr, (void*)max_ptr);

    reverse_array_ptr(arr, size);
    printf("Reversed array: ");
    print_array_ptr(arr, size);

    return 0;
}

```

Exercise 6.2: Dynamic Memory Allocation

```

#include <stdio.h>
#include <stdlib.h>

// Function to create dynamic array
int* create_array(int size) {
    int *arr = (int*)malloc(size * sizeof(int));
    if (arr == NULL) {
        printf("Memory allocation failed\n");
        return NULL;
    }
    return arr;
}

// Function to initialize array with random values
void initialize_array(int *arr, int size) {
    for (int i = 0; i < size; i++) {
        arr[i] = rand() % 100; // Random numbers 0-99
    }
}

```



```

// Function to resize array
int* resize_array(int *arr, int old_size, int new_size) {
    int *new_arr = (int*)realloc(arr, new_size * sizeof(int));
    if (new_arr == NULL) {
        printf("Memory reallocation failed\n");
        return arr; // Return original array
    }

    // Initialize new elements to 0
    if (new_size > old_size) {
        for (int i = old_size; i < new_size; i++) {
            new_arr[i] = 0;
        }
    }

    return new_arr;
}

// Function to free array
void free_array(int *arr) {
    free(arr);
}

// Function to print array
void print_array(int *arr, int size) {
    printf("[");
    for (int i = 0; i < size; i++) {
        printf("%d", arr[i]);
        if (i < size - 1) printf(", ");
    }
    printf("]\n");
}

int main() {
    int size = 5;
    int *arr = create_array(size);

    if (arr != NULL) {
        initialize_array(arr, size);
        printf("Original array: ");
        print_array(arr, size);

        // Resize array
        int new_size = 8;
        arr = resize_array(arr, size, new_size);
        if (arr != NULL) {

```

```

        printf("Resized array: ");
        print_array(arr, new_size);
    }

    free_array(arr);
    printf("Memory freed successfully\n");
}

return 0;
}

```

Module 7: Structures

Exercise 7.1: Student Database

```

#include <stdio.h>
#include <string.h>
#include <stdlib.h>

#define MAX_STUDENTS 100

typedef struct {
    int id;
    char name[50];
    int age;
    float gpa;
} Student;

// Function to add student
void add_student(Student students[], int *count) {
    if (*count >= MAX_STUDENTS) {
        printf("Database is full\n");
        return;
    }

    Student *s = &students[*count];

    printf("Enter student ID: ");
    scanf("%d", &s->id);

    printf("Enter student name: ");
    scanf(" %[^\\n]", s->name);

    printf("Enter student age: ");
    scanf("%d", &s->age);
}

```

```

    printf("Enter student GPA: ");
    scanf("%f", &s->gpa);

    (*count)++;
    printf("Student added successfully\n");
}

// Function to display all students
void display_students(Student students[], int count) {
    if (count == 0) {
        printf("No students in database\n");
        return;
    }

    printf("\nStudent Database:\n");
    printf("ID\tName\t\tAge\tGPA\n");
    printf("-----\n");

    for (int i = 0; i < count; i++) {
        Student *s = &students[i];
        printf("%d\t%-20s\t%d\t%.2f\n", s->id, s->name, s->age, s->gpa);
    }
}

// Function to find student by ID
Student* find_student(Student students[], int count, int id) {
    for (int i = 0; i < count; i++) {
        if (students[i].id == id) {
            return &students[i];
        }
    }
    return NULL;
}

// Function to calculate average GPA
float calculate_average_gpa(Student students[], int count) {
    if (count == 0) return 0.0;

    float sum = 0.0;
    for (int i = 0; i < count; i++) {
        sum += students[i].gpa;
    }
    return sum / count;
}

int main() {

```

```

Student students[MAX_STUDENTS];
int count = 0;
int choice, id;

while (1) {
    printf("\nStudent Database Menu:\n");
    printf("1. Add Student\n");
    printf("2. Display All Students\n");
    printf("3. Find Student by ID\n");
    printf("4. Calculate Average GPA\n");
    printf("5. Exit\n");
    printf("Enter your choice: ");
    scanf("%d", &choice);

    switch (choice) {
        case 1:
            add_student(students, &count);
            break;

        case 2:
            display_students(students, count);
            break;

        case 3:
            printf("Enter student ID: ");
            scanf("%d", &id);
            Student *s = find_student(students, count, id);
            if (s != NULL) {
                printf("Student found:\n");
                printf("ID: %d\n", s->id);
                printf("Name: %s\n", s->name);
                printf("Age: %d\n", s->age);
                printf("GPA: %.2f\n", s->gpa);
            } else {
                printf("Student not found\n");
            }
            break;

        case 4:
            printf("Average GPA: %.2f\n", calculate_average_gpa(students, count));
            break;

        case 5:
            printf("Exiting program\n");
            exit(0);

        default:

```

```

        printf("Invalid choice\n");
    }
}

return 0;
}

```

Exercise 7.2: Complex Number Operations

```

#include <stdio.h>
#include <math.h>

typedef struct {
    double real;
    double imag;
} Complex;

// Function to create complex number
Complex create_complex(double real, double imag) {
    Complex c = {real, imag};
    return c;
}

// Function to add two complex numbers
Complex add_complex(Complex a, Complex b) {
    Complex result;
    result.real = a.real + b.real;
    result.imag = a.imag + b.imag;
    return result;
}

// Function to subtract two complex numbers
Complex subtract_complex(Complex a, Complex b) {
    Complex result;
    result.real = a.real - b.real;
    result.imag = a.imag - b.imag;
    return result;
}

// Function to multiply two complex numbers
Complex multiply_complex(Complex a, Complex b) {
    Complex result;
    result.real = a.real * b.real - a.imag * b.imag;
    result.imag = a.real * b.imag + a.imag * b.real;
    return result;
}

```

```

}

// Function to calculate magnitude of complex number
double magnitude_complex(Complex c) {
    return sqrt(c.real * c.real + c.imag * c.imag);
}

// Function to print complex number
void print_complex(Complex c) {
    if (c.imag >= 0) {
        printf("%.2f + %.2fi", c.real, c.imag);
    } else {
        printf("%.2f - %.2fi", c.real, -c.imag);
    }
}

int main() {
    Complex a, b, result;
    int choice;

    printf("Enter first complex number:\n");
    printf("Real part: ");
    scanf("%lf", &a.real);
    printf("Imaginary part: ");
    scanf("%lf", &a.imag);

    printf("Enter second complex number:\n");
    printf("Real part: ");
    scanf("%lf", &b.real);
    printf("Imaginary part: ");
    scanf("%lf", &b.imag);

    while (1) {
        printf("\nComplex Number Operations:\n");
        printf("1. Add\n");
        printf("2. Subtract\n");
        printf("3. Multiply\n");
        printf("4. Magnitude of first number\n");
        printf("5. Exit\n");
        printf("Enter your choice: ");
        scanf("%d", &choice);

        switch (choice) {
            case 1:
                result = add_complex(a, b);
                printf("Result: ");
                print_complex(a);

```

```

        printf(" + ");
        print_complex(b);
        printf(" = ");
        print_complex(result);
        printf("\n");
        break;

    case 2:
        result = subtract_complex(a, b);
        printf("Result: ");
        print_complex(a);
        printf(" - ");
        print_complex(b);
        printf(" = ");
        print_complex(result);
        printf("\n");
        break;

    case 3:
        result = multiply_complex(a, b);
        printf("Result: ");
        print_complex(a);
        printf(" * ");
        print_complex(b);
        printf(" = ");
        print_complex(result);
        printf("\n");
        break;

    case 4:
        printf("Magnitude of ");
        print_complex(a);
        printf(" = %.2f\n", magnitude_complex(a));
        break;

    case 5:
        printf("Exiting program\n");
        return 0;

    default:
        printf("Invalid choice\n");
    }
}

return 0;
}

```

Module 8: File I/O

Exercise 8.1: Text File Processor

```
#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <ctype.h>

// Function to count lines, words, and characters
void analyze_file(const char *filename) {
    FILE *file = fopen(filename, "r");
    if (file == NULL) {
        perror("Error opening file");
        return;
    }

    int lines = 0, words = 0, characters = 0;
    int ch, in_word = 0;

    while ((ch = fgetc(file)) != EOF) {
        characters++;

        if (ch == '\n') {
            lines++;
        }

        if (isspace(ch)) {
            in_word = 0;
        } else if (!in_word) {
            in_word = 1;
            words++;
        }
    }

    // Handle case where file doesn't end with newline
    if (characters > 0 && fseek(file, -1, SEEK_END) == 0) {
        if (fgetc(file) != '\n') {
            lines++;
        }
    }

    fclose(file);

    printf("File Analysis for '%s':\n", filename);
    printf("Lines: %d\n", lines);
}
```



```

    printf("Words: %d\n", words);
    printf("Characters: %d\n", characters);
}

// Function to copy file
void copy_file(const char *source, const char *destination) {
    FILE *src = fopen(source, "r");
    if (src == NULL) {
        perror("Error opening source file");
        return;
    }

    FILE *dest = fopen(destination, "w");
    if (dest == NULL) {
        perror("Error opening destination file");
        fclose(src);
        return;
    }

    int ch;
    while ((ch = fgetc(src)) != EOF) {
        fputc(ch, dest);
    }

    fclose(src);
    fclose(dest);

    printf("File copied successfully from '%s' to '%s'\n", source, destination);
}

// Function to convert text to uppercase
void convert_to_uppercase(const char *input, const char *output) {
    FILE *in = fopen(input, "r");
    if (in == NULL) {
        perror("Error opening input file");
        return;
    }

    FILE *out = fopen(output, "w");
    if (out == NULL) {
        perror("Error opening output file");
        fclose(in);
        return;
    }

    int ch;

```

```

while ((ch = fgetc(in)) != EOF) {
    fputc(toupper(ch), out);
}

fclose(in);
fclose(out);

printf("Text converted to uppercase and saved to '%s'\n", output);
}

int main() {
    int choice;
    char filename1[100], filename2[100];

    while (1) {
        printf("\nText File Processor:\n");
        printf("1. Analyze file\n");
        printf("2. Copy file\n");
        printf("3. Convert to uppercase\n");
        printf("4. Exit\n");
        printf("Enter your choice: ");
        scanf("%d", &choice);

        switch (choice) {
            case 1:
                printf("Enter filename: ");
                scanf("%s", filename1);
                analyze_file(filename1);
                break;

            case 2:
                printf("Enter source filename: ");
                scanf("%s", filename1);
                printf("Enter destination filename: ");
                scanf("%s", filename2);
                copy_file(filename1, filename2);
                break;

            case 3:
                printf("Enter input filename: ");
                scanf("%s", filename1);
                printf("Enter output filename: ");
                scanf("%s", filename2);
                convert_to_uppercase(filename1, filename2);
                break;

            case 4:

```

```

        printf("Exiting program\n");
        return 0;

    default:
        printf("Invalid choice\n");
    }
}

return 0;
}

```

Exercise 8.2: Binary File Operations

```

#include <stdio.h>
#include <stdlib.h>
#include <string.h>

#define MAX_RECORDS 100

typedef struct {
    int id;
    char name[50];
    float salary;
} Employee;

// Function to write employee data to binary file
void write_employees(const char *filename, Employee employees[], int count) {
    FILE *file = fopen(filename, "wb");
    if (file == NULL) {
        perror("Error opening file for writing");
        return;
    }

    fwrite(&count, sizeof(int), 1, file);
    fwrite(employees, sizeof(Employee), count, file);

    fclose(file);
    printf("Employee data written to '%s'\n", filename);
}

// Function to read employee data from binary file
int read_employees(const char *filename, Employee employees[]) {
    FILE *file = fopen(filename, "rb");
    if (file == NULL) {
        perror("Error opening file for reading");
    }
}

```

```

        return 0;
    }

    int count;
    fread(&count, sizeof(int), 1, file);

    if (count > MAX_RECORDS) {
        printf("Too many records in file\n");
        fclose(file);
        return 0;
    }

    fread(employees, sizeof(Employee), count, file);
    fclose(file);

    printf("Employee data read from '%s'\n", filename);
    return count;
}

// Function to add employee
void add_employee(Employee employees[], int *count) {
    if (*count >= MAX_RECORDS) {
        printf("Maximum number of employees reached\n");
        return;
    }

    Employee *emp = &employees[*count];

    printf("Enter employee ID: ");
    scanf("%d", &emp->id);

    printf("Enter employee name: ");
    scanf(" %[^\\n]", emp->name);

    printf("Enter employee salary: ");
    scanf("%f", &emp->salary);

    (*count)++;
    printf("Employee added successfully\n");
}

// Function to display all employees
void display_employees(Employee employees[], int count) {
    if (count == 0) {
        printf("No employees in database\n");
        return;
    }
}

```

```

printf("\nEmployee Database:\n");
printf("ID\tName\t\t\tSalary\n");
printf("-----\n");

for (int i = 0; i < count; i++) {
    Employee *emp = &employees[i];
    printf("%d\t%-20s\t$%.2f\n", emp->id, emp->name, emp->salary);
}

}

int main() {
    Employee employees[MAX_RECORDS];
    int count = 0;
    int choice;
    char filename[100];

    while (1) {
        printf("\nBinary File Employee Database:\n");
        printf("1. Add Employee\n");
        printf("2. Display Employees\n");
        printf("3. Save to File\n");
        printf("4. Load from File\n");
        printf("5. Exit\n");
        printf("Enter your choice: ");
        scanf("%d", &choice);

        switch (choice) {
            case 1:
                add_employee(employees, &count);
                break;

            case 2:
                display_employees(employees, count);
                break;

            case 3:
                printf("Enter filename: ");
                scanf("%s", filename);
                write_employees(filename, employees, count);
                break;

            case 4:
                printf("Enter filename: ");
                scanf("%s", filename);
                count = read_employees(filename, employees);
                break;
        }
    }
}

```

```

        case 5:
            printf("Exiting program\n");
            return 0;

        default:
            printf("Invalid choice\n");
    }
}

return 0;
}

```

Module 9: Modern C Features

Exercise 9.1: C99 Features Demonstration

```

#include <stdio.h>
#include <stdbool.h>
#include <stdint.h>
#include <inttypes.h>

// Variable-length array example
void vla_example(int size) {
    int arr[size]; // VLA - size determined at runtime

    printf("VLA of size %d:\n", size);
    for (int i = 0; i < size; i++) {
        arr[i] = i * i;
        printf("arr[%d] = %d\n", i, arr[i]);
    }
}

// Inline function example
inline int max(int a, int b) {
    return (a > b) ? a : b;
}

// _Bool example
void bool_example(void) {
    bool flag1 = true;
    bool flag2 = false;

    printf("Boolean values: flag1 = %s, flag2 = %s\n",
        flag1 ? "true" : "false",

```

```

        flag2 ? "true" : "false");

    bool result = flag1 && !flag2;
    printf("flag1 AND NOT flag2 = %s\n", result ? "true" : "false");
}

int main() {
    printf("C99 Features Demonstration\n");
    printf("=====\n");

    // Variable-length array
    int size;
    printf("Enter size for VLA: ");
    scanf("%d", &size);
    vla_example(size);

    // Inline function
    int a = 15, b = 25;
    printf("Max of %d and %d = %d\n", a, b, max(a, b));

    // Boolean example
    bool_example();

    // Fixed-width integers
    int32_t fixed_int = 123456789;
    uint64_t big_uint = 123456789012345ULL;

    printf("Fixed-width integers:\n");
    printf("int32_t: %" PRId32 "\n", fixed_int);
    printf("uint64_t: %" PRIu64 "\n", big_uint);

    return 0;
}

```

Exercise 9.2: C11 Features Demonstration

```

#include <stdio.h>
#include <stdlib.h>
#include <stdalign.h>
#include <stdnoreturn.h>
#include <threads.h>
#include <stdatomic.h>

// Static assertion example
#define ARRAY_SIZE 10

```

```

_Static_assert(ARRAY_SIZE > 0, "Array size must be positive");

// Generic selection example
#define PRINT_GENERIC(x) _Generic((x), \
    int: printf("Integer: %d\n"), \
    float: printf("Float: %f\n"), \
    double: printf("Double: %f\n"), \
    char*: printf("String: %s\n"), \
    default: printf("Unknown type\n") \
)(x)

// Noreturn function example
noreturn void fatal_error(const char *msg) {
    printf("Fatal error: %s\n", msg);
    exit(EXIT_FAILURE);
}

// Alignment example
struct aligned_struct {
    alignas(16) double x;
    int y;
};

// Atomic operations example
atomic_int global_counter = ATOMIC_VAR_INIT(0);

int increment_counter(void *arg) {
    for (int i = 0; i < 1000; i++) {
        atomic_fetch_add(&global_counter, 1);
    }
    return 0;
}

int main() {
    printf("C11 Features Demonstration\n");
    printf("=====\n");

    // Generic selection
    int i = 42;
    float f = 3.14f;
    double d = 2.718;
    char *s = "Hello, C11!";

    PRINT_GENERIC(i);
    PRINT_GENERIC(f);
    PRINT_GENERIC(d);
    PRINT_GENERIC(s);
}

```



```

// Alignment
printf("Alignment of aligned_struct: %zu\n", alignof(struct aligned_struct));
printf("Size of aligned_struct: %zu\n", sizeof(struct aligned_struct));

// Atomic operations with threads
thrd_t thread1, thread2;

thrd_create(&thread1, increment_counter, NULL);
thrd_create(&thread2, increment_counter, NULL);

thrd_join(thread1, NULL);
thrd_join(thread2, NULL);

printf("Final counter value: %d\n", atomic_load(&global_counter));

// Static assertion (compile-time check)
int arr[ARRAY_SIZE];
printf("Array of size %d created successfully\n", ARRAY_SIZE);

return 0;
}

```

Module 10: Data Structures

Exercise 10.1: Linked List Implementation

```

#include <stdio.h>
#include <stdlib.h>

typedef struct Node {
    int data;
    struct Node* next;
} Node;

// Function to create new node
Node* create_node(int data) {
    Node* new_node = (Node*)malloc(sizeof(Node));
    if (new_node == NULL) {
        printf("Memory allocation failed\n");
        return NULL;
    }
    new_node->data = data;
    new_node->next = NULL;
    return new_node;
}

```

```

}

// Function to insert at beginning
void insert_at_beginning(Node** head, int data) {
    Node* new_node = create_node(data);
    if (new_node == NULL) return;

    new_node->next = *head;
    *head = new_node;
}

// Function to insert at end
void insert_at_end(Node** head, int data) {
    Node* new_node = create_node(data);
    if (new_node == NULL) return;

    if (*head == NULL) {
        *head = new_node;
        return;
    }

    Node* current = *head;
    while (current->next != NULL) {
        current = current->next;
    }
    current->next = new_node;
}

// Function to delete node
void delete_node(Node** head, int key) {
    if (*head == NULL) return;

    // If head node holds the key
    if ((*head)->data == key) {
        Node* temp = *head;
        *head = (*head)->next;
        free(temp);
        return;
    }

    // Search for the key to be deleted
    Node* current = *head;
    while (current->next != NULL && current->next->data != key) {
        current = current->next;
    }
}

```

```

    // If key was not present
    if (current->next == NULL) return;

    // Remove the node
    Node* temp = current->next;
    current->next = current->next->next;
    free(temp);
}

// Function to search for a node
Node* search_node(Node* head, int key) {
    Node* current = head;
    while (current != NULL) {
        if (current->data == key) {
            return current;
        }
        current = current->next;
    }
    return NULL;
}

// Function to display the list
void display_list(Node* head) {
    if (head == NULL) {
        printf("List is empty\n");
        return;
    }

    printf("List: ");
    Node* current = head;
    while (current != NULL) {
        printf("%d", current->data);
        if (current->next != NULL) printf(" -> ");
        current = current->next;
    }
    printf(" -> NULL\n");
}

// Function to free the entire list
void free_list(Node* head) {
    Node* current = head;
    while (current != NULL) {
        Node* next = current->next;
        free(current);
        current = next;
    }
}

```

```

int main() {
    Node* head = NULL;
    int choice, data;

    while (1) {
        printf("\nLinked List Operations:\n");
        printf("1. Insert at beginning\n");
        printf("2. Insert at end\n");
        printf("3. Delete node\n");
        printf("4. Search node\n");
        printf("5. Display list\n");
        printf("6. Exit\n");
        printf("Enter your choice: ");
        scanf("%d", &choice);

        switch (choice) {
            case 1:
                printf("Enter data: ");
                scanf("%d", &data);
                insert_at_beginning(&head, data);
                break;

            case 2:
                printf("Enter data: ");
                scanf("%d", &data);
                insert_at_end(&head, data);
                break;

            case 3:
                printf("Enter data to delete: ");
                scanf("%d", &data);
                delete_node(&head, data);
                break;

            case 4:
                printf("Enter data to search: ");
                scanf("%d", &data);
                Node* found = search_node(head, data);
                if (found != NULL) {
                    printf("Node found with data: %d\n", found->data);
                } else {
                    printf("Node not found\n");
                }
                break;

            case 5:

```

```

        display_list(head);
        break;

    case 6:
        free_list(head);
        printf("Exiting program\n");
        return 0;

    default:
        printf("Invalid choice\n");
    }
}

return 0;
}

```

Exercise 10.2: Binary Search Tree Implementation

```

#include <stdio.h>
#include <stdlib.h>

typedef struct TreeNode {
    int data;
    struct TreeNode* left;
    struct TreeNode* right;
} TreeNode;

// Function to create new node
TreeNode* create_node(int data) {
    TreeNode* new_node = (TreeNode*)malloc(sizeof(TreeNode));
    if (new_node == NULL) {
        printf("Memory allocation failed\n");
        return NULL;
    }
    new_node->data = data;
    new_node->left = NULL;
    new_node->right = NULL;
    return new_node;
}

// Function to insert node
TreeNode* insert_node(TreeNode* root, int data) {
    if (root == NULL) {
        return create_node(data);
    }
}

```

```

        if (data < root->data) {
            root->left = insert_node(root->left, data);
        } else if (data > root->data) {
            root->right = insert_node(root->right, data);
        }
        // If data is equal, we don't insert (no duplicates)

        return root;
    }

    // Function for inorder traversal
    void inorder_traversal(TreeNode* root) {
        if (root != NULL) {
            inorder_traversal(root->left);
            printf("%d ", root->data);
            inorder_traversal(root->right);
        }
    }

    // Function for preorder traversal
    void preorder_traversal(TreeNode* root) {
        if (root != NULL) {
            printf("%d ", root->data);
            preorder_traversal(root->left);
            preorder_traversal(root->right);
        }
    }

    // Function for postorder traversal
    void postorder_traversal(TreeNode* root) {
        if (root != NULL) {
            postorder_traversal(root->left);
            postorder_traversal(root->right);
            printf("%d ", root->data);
        }
    }

    // Function to search for a node
    TreeNode* search_node(TreeNode* root, int data) {
        if (root == NULL || root->data == data) {
            return root;
        }

        if (data < root->data) {
            return search_node(root->left, data);
        } else {

```

```

        return search_node(root->right, data);
    }
}

// Function to find minimum value node
TreeNode* find_min(TreeNode* root) {
    while (root && root->left != NULL) {
        root = root->left;
    }
    return root;
}

// Function to delete node
TreeNode* delete_node(TreeNode* root, int data) {
    if (root == NULL) return root;

    if (data < root->data) {
        root->left = delete_node(root->left, data);
    } else if (data > root->data) {
        root->right = delete_node(root->right, data);
    } else {
        // Node to be deleted found
        if (root->left == NULL) {
            TreeNode* temp = root->right;
            free(root);
            return temp;
        } else if (root->right == NULL) {
            TreeNode* temp = root->left;
            free(root);
            return temp;
        }

        // Node with two children
        TreeNode* temp = find_min(root->right);
        root->data = temp->data;
        root->right = delete_node(root->right, temp->data);
    }
    return root;
}

// Function to free the entire tree
void free_tree(TreeNode* root) {
    if (root != NULL) {
        free_tree(root->left);
        free_tree(root->right);
        free(root);
    }
}

```

```

}

int main() {
    TreeNode* root = NULL;
    int choice, data;

    while (1) {
        printf("\nBinary Search Tree Operations:\n");
        printf("1. Insert node\n");
        printf("2. Delete node\n");
        printf("3. Search node\n");
        printf("4. Inorder traversal\n");
        printf("5. Preorder traversal\n");
        printf("6. Postorder traversal\n");
        printf("7. Exit\n");
        printf("Enter your choice: ");
        scanf("%d", &choice);

        switch (choice) {
            case 1:
                printf("Enter data: ");
                scanf("%d", &data);
                root = insert_node(root, data);
                break;

            case 2:
                printf("Enter data to delete: ");
                scanf("%d", &data);
                root = delete_node(root, data);
                break;

            case 3:
                printf("Enter data to search: ");
                scanf("%d", &data);
                TreeNode* found = search_node(root, data);
                if (found != NULL) {
                    printf("Node found with data: %d\n", found->data);
                } else {
                    printf("Node not found\n");
                }
                break;

            case 4:
                printf("Inorder traversal: ");
                inorder_traversal(root);
                printf("\n");
                break;
        }
    }
}

```



```

        case 5:
            printf("Preorder traversal: ");
            preorder_traversal(root);
            printf("\n");
            break;

        case 6:
            printf("Postorder traversal: ");
            postorder_traversal(root);
            printf("\n");
            break;

        case 7:
            free_tree(root);
            printf("Exiting program\n");
            return 0;

        default:
            printf("Invalid choice\n");
    }
}

return 0;
}

```

Module 11: Network Programming

Exercise 11.1: Simple TCP Client-Server

```

// Server code (server.c)
#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <unistd.h>
#include <sys/socket.h>
#include <netinet/in.h>
#include <arpa/inet.h>

#define PORT 8080
#define BUFFER_SIZE 1024

int main_server() {
    int server_fd, new_socket;
    struct sockaddr_in address;

```

```

int opt = 1;
int addrlen = sizeof(address);
char buffer[BUFFER_SIZE] = {0};

// Create socket
if ((server_fd = socket(AF_INET, SOCK_STREAM, 0)) == 0) {
    perror("socket failed");
    exit(EXIT_FAILURE);
}

// Set socket options
if (setsockopt(server_fd, SOL_SOCKET, SO_REUSEADDR | SO_REUSEPORT,
               &opt, sizeof(opt))) {
    perror("setsockopt failed");
    exit(EXIT_FAILURE);
}

address.sin_family = AF_INET;
address.sin_addr.s_addr = INADDR_ANY;
address.sin_port = htons(PORT);

// Bind socket
if (bind(server_fd, (struct sockaddr *)&address, sizeof(address)) < 0) {
    perror("bind failed");
    exit(EXIT_FAILURE);
}

// Listen for connections
if (listen(server_fd, 3) < 0) {
    perror("listen failed");
    exit(EXIT_FAILURE);
}

printf("Server listening on port %d\n", PORT);

// Accept connection
if ((new_socket = accept(server_fd, (struct sockaddr *)&address,
                        (socklen_t *)&addrlen)) < 0) {
    perror("accept failed");
    exit(EXIT_FAILURE);
}

// Receive message
int valread = read(new_socket, buffer, BUFFER_SIZE);
printf("Client: %s\n", buffer);

// Send response

```

```

    const char *response = "Hello from server";
    send(new_socket, response, strlen(response), 0);
    printf("Response sent\n");

    close(new_socket);
    close(server_fd);
    return 0;
}

// Client code (client.c)
int main_client() {
    int sock = 0;
    struct sockaddr_in serv_addr;
    char *hello = "Hello from client";
    char buffer[1024] = {0};

    // Create socket
    if ((sock = socket(AF_INET, SOCK_STREAM, 0)) < 0) {
        printf("Socket creation error\n");
        return -1;
    }

    serv_addr.sin_family = AF_INET;
    serv_addr.sin_port = htons(PORT);

    // Convert IPv4 address from text to binary
    if (inet_pton(AF_INET, "127.0.0.1", &serv_addr.sin_addr) <= 0) {
        printf("Invalid address/ Address not supported\n");
        return -1;
    }

    // Connect to server
    if (connect(sock, (struct sockaddr *)&serv_addr, sizeof(serv_addr)) < 0) {
        printf("Connection Failed\n");
        return -1;
    }

    // Send message
    send(sock, hello, strlen(hello), 0);
    printf("Hello message sent\n");

    // Receive response
    int valread = read(sock, buffer, 1024);
    printf("Server: %s\n", buffer);

    close(sock);
    return 0;
}

```

```

}

int main() {
    int choice;
    printf("1. Run as server\n");
    printf("2. Run as client\n");
    printf("Enter choice: ");
    scanf("%d", &choice);

    if (choice == 1) {
        return main_server();
    } else if (choice == 2) {
        return main_client();
    } else {
        printf("Invalid choice\n");
        return 1;
    }
}

```

Exercise 11.2: HTTP Client

```

#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <unistd.h>
#include <sys/socket.h>
#include <netinet/in.h>
#include <arpa/inet.h>
#include <netdb.h>

int http_get_request(const char *hostname, const char *path) {
    int sock;
    struct hostent *server;
    struct sockaddr_in serv_addr;
    char request[1024];
    char response[4096];

    // Create socket
    sock = socket(AF_INET, SOCK_STREAM, 0);
    if (sock < 0) {
        perror("socket creation failed");
        return 1;
    }

    // Get server IP address

```

```

server = gethostbyname(hostname);
if (server == NULL) {
    fprintf(stderr, "Error, no such host\n");
    close(sock);
    return 1;
}

// Set up server address
memset(&serv_addr, 0, sizeof(serv_addr));
serv_addr.sin_family = AF_INET;
memcpy(server->h_addr, &serv_addr.sin_addr.s_addr, server->h_length);
serv_addr.sin_port = htons(80);

// Connect to server
if (connect(sock, (struct sockaddr*)&serv_addr, sizeof(serv_addr)) < 0) {
    perror("connection failed");
    close(sock);
    return 1;
}

// Create HTTP GET request
snprintf(request, sizeof(request),
    "GET %s HTTP/1.1\r\n"
    "Host: %s\r\n"
    "Connection: close\r\n\r\n",
    path, hostname);

// Send request
if (send(sock, request, strlen(request), 0) < 0) {
    perror("send failed");
    close(sock);
    return 1;
}

// Receive response
int bytes_received;
while ((bytes_received = recv(sock, response, sizeof(response) - 1, 0)) > 0) {
    response[bytes_received] = '\0';
    printf("%s", response);
}

close(sock);
return 0;
}

int main() {
    char hostname[256];

```

```

    char path[256];

    printf("Enter hostname (e.g., httpbin.org): ");
    scanf("%s", hostname);

    printf("Enter path (e.g., /get): ");
    scanf("%s", path);

    return http_get_request(hostname, path);
}

```

Module 12: Embedded Systems Programming

Exercise 12.1: Bit Manipulation Library

```

#include <stdio.h>
#include <stdint.h>

// Bit manipulation macros
#define SET_BIT(reg, bit)      ((reg) |= (1U << (bit)))
#define CLEAR_BIT(reg, bit)   ((reg) &= ~(1U << (bit)))
#define TOGGLE_BIT(reg, bit)  ((reg) ^= (1U << (bit)))
#define CHECK_BIT(reg, bit)   (((reg) >> (bit)) & 1U)

// Bit field structure for hardware register
typedef struct {
    uint32_t enable      : 1; // Bit 0
    uint32_t direction   : 1; // Bit 1
    uint32_t reserved    : 2; // Bits 2-3
    uint32_t mode         : 4; // Bits 4-7
    uint32_t config       : 24; // Bits 8-31
} gpio_register_t;

// Function implementations
void set_bit(uint32_t *reg, int bit) {
    *reg |= (1U << bit);
}

void clear_bit(uint32_t *reg, int bit) {
    *reg &= ~(1U << bit);
}

void toggle_bit(uint32_t *reg, int bit) {
    *reg ^= (1U << bit);
}

```

```

int check_bit(uint32_t reg, int bit) {
    return (reg >> bit) & 1U;
}

int main() {
    uint32_t test_reg = 0;

    printf("Initial register value: 0x%08X\n", test_reg);

    // Test bit manipulation functions
    set_bit(&test_reg, 5);
    printf("After setting bit 5: 0x%08X\n", test_reg);

    toggle_bit(&test_reg, 5);
    printf("After toggling bit 5: 0x%08X\n", test_reg);

    clear_bit(&test_reg, 5);
    printf("After clearing bit 5: 0x%08X\n", test_reg);

    // Test bit manipulation macros
    SET_BIT(test_reg, 3);
    printf("After setting bit 3 with macro: 0x%08X\n", test_reg);

    if (CHECK_BIT(test_reg, 3)) {
        printf("Bit 3 is set\n");
    }

    // Test bit field structure
    gpio_register_t gpio_reg = {0};
    gpio_reg.enable = 1;
    gpio_reg.direction = 1;
    gpio_reg.mode = 0xF;

    printf("GPIO Register - Enable: %d, Direction: %d, Mode: 0x%X\n",
        gpio_reg.enable, gpio_reg.direction, gpio_reg.mode);

    return 0;
}

```

Exercise 12.2: Memory-Mapped I/O Simulation

```

#include <stdio.h>
#include <stdint.h>
#include <stdlib.h>
#include <string.h>

```

```

// Memory-mapped I/O simulation
#define GPIO_BASE_ADDR 0x40020000
#define GPIO_DIR_REG    (*(volatile uint32_t *) (GPIO_BASE_ADDR + 0x00))
#define GPIO_DATA_REG    (*(volatile uint32_t *) (GPIO_BASE_ADDR + 0x04))
#define GPIO_SET_REG     (*(volatile uint32_t *) (GPIO_BASE_ADDR + 0x08))
#define GPIO_CLR_REG     (*(volatile uint32_t *) (GPIO_BASE_ADDR + 0x0C))

// GPIO pin definitions
#define GPIO_PIN_0  (1U << 0)
#define GPIO_PIN_1  (1U << 1)
#define GPIO_PIN_2  (1U << 2)
#define GPIO_PIN_3  (1U << 3)

// Simple GPIO simulation structure
typedef struct {
    uint32_t direction; // 0 = input, 1 = output
    uint32_t data;       // Current pin states
} gpio_sim_t;

// Global GPIO simulation instance
static gpio_sim_t gpio_sim = {0};

// Function to initialize GPIO
void gpio_init(void) {
    gpio_sim.direction = 0;
    gpio_sim.data = 0;
    printf("GPIO initialized\n");
}

// Function to set GPIO direction
void gpio_set_direction(uint32_t pins, int direction) {
    if (direction) {
        gpio_sim.direction |= pins; // Set as output
    } else {
        gpio_sim.direction &= ~pins; // Set as input
    }

    // Update hardware register
    GPIO_DIR_REG = gpio_sim.direction;
    printf("GPIO direction set: 0x%08X (pins: 0x%08X, direction: %s)\n",
        gpio_sim.direction, pins, direction ? "output" : "input");
}

// Function to write to GPIO pins
void gpio_write(uint32_t pins, int value) {
    if (value) {

```



```

        gpio_sim.data |= pins;
        GPIO_SET_REG = pins; // Hardware register for setting pins
    } else {
        gpio_sim.data &= ~pins;
        GPIO_CLR_REG = pins; // Hardware register for clearing pins
    }

    GPIO_DATA_REG = gpio_sim.data; // Update data register
    printf("GPIO write: pins 0x%08X set to %d\n", pins, value);
}

// Function to read from GPIO pins
uint32_t gpio_read(uint32_t pins) {
    uint32_t value = GPIO_DATA_REG & pins; // Read from hardware register
    printf("GPIO read: pins 0x%08X = 0x%08X\n", pins, value);
    return value;
}

// Function to simulate external input change
void gpio_simulate_input(uint32_t pins, int value) {
    if (value) {
        gpio_sim.data |= pins;
    } else {
        gpio_sim.data &= ~pins;
    }
    printf("Simulated input change: pins 0x%08X set to %d\n", pins, value);
}

int main() {
    // Initialize GPIO
    gpio_init();

    // Configure pins 0-1 as outputs, pins 2-3 as inputs
    gpio_set_direction(GPIO_PIN_0 | GPIO_PIN_1, 1); // Output
    gpio_set_direction(GPIO_PIN_2 | GPIO_PIN_3, 0); // Input

    // Write to output pins
    gpio_write(GPIO_PIN_0, 1);
    gpio_write(GPIO_PIN_1, 0);

    // Simulate input changes
    gpio_simulate_input(GPIO_PIN_2, 1);
    gpio_simulate_input(GPIO_PIN_3, 0);

    // Read input pins
    uint32_t input_value = gpio_read(GPIO_PIN_2 | GPIO_PIN_3);
    printf("Input pins value: 0x%08X\n", input_value);
}

```

```
    return 0;
}
```

Module 13: Performance Optimization

Exercise 13.1: Algorithm Complexity Analysis

```
#include <stdio.h>
#include <stdlib.h>
#include <time.h>
#include <string.h>

// Timing function
double get_time_diff(struct timespec start, struct timespec end) {
    return (end.tv_sec - start.tv_sec) + (end.tv_nsec - start.tv_nsec) / 1e9;
}

// O(n2) bubble sort
void bubble_sort(int arr[], int n) {
    for (int i = 0; i < n - 1; i++) {
        for (int j = 0; j < n - i - 1; j++) {
            if (arr[j] > arr[j + 1]) {
                int temp = arr[j];
                arr[j] = arr[j + 1];
                arr[j + 1] = temp;
            }
        }
    }
}

// O(n log n) quick sort
void quick_sort(int arr[], int low, int high) {
    if (low < high) {
        int pi = partition(arr, low, high);
        quick_sort(arr, low, pi - 1);
        quick_sort(arr, pi + 1, high);
    }
}

int partition(int arr[], int low, int high) {
    int pivot = arr[high];
    int i = (low - 1);

    for (int j = low; j <= high - 1; j++) {
```

```

        if (arr[j] < pivot) {
            i++;
            int temp = arr[i];
            arr[i] = arr[j];
            arr[j] = temp;
        }
    }
    int temp = arr[i + 1];
    arr[i + 1] = arr[high];
    arr[high] = temp;
    return (i + 1);
}

// O(n) linear search
int linear_search(int arr[], int n, int key) {
    for (int i = 0; i < n; i++) {
        if (arr[i] == key) {
            return i;
        }
    }
    return -1;
}

// O(log n) binary search
int binary_search(int arr[], int low, int high, int key) {
    if (high >= low) {
        int mid = low + (high - low) / 2;
        if (arr[mid] == key) {
            return mid;
        }
        if (arr[mid] > key) {
            return binary_search(arr, low, mid - 1, key);
        }
        return binary_search(arr, mid + 1, high, key);
    }
    return -1;
}

// Performance testing function
void test_algorithms() {
    const int sizes[] = {1000, 5000, 10000, 20000};
    const int num_sizes = sizeof(sizes) / sizeof(sizes[0]);

    for (int s = 0; s < num_sizes; s++) {
        int n = sizes[s];
        int *arr1 = malloc(n * sizeof(int));
        int *arr2 = malloc(n * sizeof(int));
    }
}

```

```

// Fill arrays with random data
for (int i = 0; i < n; i++) {
    arr1[i] = rand() % 10000;
    arr2[i] = arr1[i];
}

struct timespec start, end;

// Test bubble sort
clock_gettime(CLOCK_MONOTONIC, &start);
bubble_sort(arr1, n);
clock_gettime(CLOCK_MONOTONIC, &end);
double bubble_time = get_time_diff(start, end);

// Test quick sort
clock_gettime(CLOCK_MONOTONIC, &start);
quick_sort(arr2, 0, n - 1);
clock_gettime(CLOCK_MONOTONIC, &end);
double quick_time = get_time_diff(start, end);

printf("Array size: %d\n", n);
printf("Bubble sort time: %.6f seconds\n", bubble_time);
printf("Quick sort time: %.6f seconds\n", quick_time);
printf("Speedup: %.2fx\n\n", bubble_time / quick_time);

free(arr1);
free(arr2);
}
}

int main() {
    srand(time(NULL));
    test_algorithms();
    return 0;
}

```

Exercise 13.2: Memory Access Patterns

```

#include <stdio.h>
#include <stdlib.h>
#include <time.h>
#include <string.h>

#define MATRIX_SIZE 1000
#define CACHE_LINE_SIZE 64

```

```

// Matrix structure with proper alignment
typedef struct {
    int rows;
    int cols;
    int *data;
} matrix_t;

// Create matrix with specified alignment
matrix_t* create_matrix(int rows, int cols) {
    matrix_t *mat = malloc(sizeof(matrix_t));
    mat->rows = rows;
    mat->cols = cols;
    // Allocate aligned memory for better cache performance
    posix_memalign((void**)&mat->data, CACHE_LINE_SIZE, rows * cols * sizeof(int));
    return mat;
}

void free_matrix(matrix_t *mat) {
    free(mat->data);
    free(mat);
}

// Initialize matrix with random values
void init_matrix(matrix_t *mat) {
    for (int i = 0; i < mat->rows * mat->cols; i++) {
        mat->data[i] = rand() % 100;
    }
}

// Matrix multiplication - ijk order (poor cache performance)
void matrix_multiply_ijk(matrix_t *a, matrix_t *b, matrix_t *c) {
    for (int i = 0; i < a->rows; i++) {
        for (int j = 0; j < b->cols; j++) {
            c->data[i * c->cols + j] = 0;
            for (int k = 0; k < a->cols; k++) {
                c->data[i * c->cols + j] +=
                    a->data[i * a->cols + k] * b->data[k * b->cols + j];
            }
        }
    }
}

// Matrix multiplication - ikj order (better cache performance)
void matrix_multiply_ikj(matrix_t *a, matrix_t *b, matrix_t *c) {
    // Initialize result matrix
    memset(c->data, 0, c->rows * c->cols * sizeof(int));

```

```

    for (int i = 0; i < a->rows; i++) {
        for (int k = 0; k < a->cols; k++) {
            int temp = a->data[i * a->cols + k];
            for (int j = 0; j < b->cols; j++) {
                c->data[i * c->cols + j] += temp * b->data[k * b->cols + j];
            }
        }
    }
}

// Performance comparison
void compare_matrix_multiplication() {
    matrix_t *a = create_matrix(MATRIX_SIZE, MATRIX_SIZE);
    matrix_t *b = create_matrix(MATRIX_SIZE, MATRIX_SIZE);
    matrix_t *c1 = create_matrix(MATRIX_SIZE, MATRIX_SIZE);
    matrix_t *c2 = create_matrix(MATRIX_SIZE, MATRIX_SIZE);

    init_matrix(a);
    init_matrix(b);

    struct timespec start, end;

    // Test ijk order
    clock_gettime(CLOCK_MONOTONIC, &start);
    matrix_multiply_ijk(a, b, c1);
    clock_gettime(CLOCK_MONOTONIC, &end);
    double time_ijk = get_time_diff(start, end);

    // Test ikj order
    clock_gettime(CLOCK_MONOTONIC, &start);
    matrix_multiply_ikj(a, b, c2);
    clock_gettime(CLOCK_MONOTONIC, &end);
    double time_ikj = get_time_diff(start, end);

    printf("Matrix multiplication performance comparison:\n");
    printf("Matrix size: %dx%d\n", MATRIX_SIZE, MATRIX_SIZE);
    printf("ijk order time: %.6f seconds\n", time_ijk);
    printf("ikj order time: %.6f seconds\n", time_ikj);
    printf("Performance improvement: %.2fx\n", time_ijk / time_ikj);

    // Verify results are the same
    int diff_count = 0;
    for (int i = 0; i < c1->rows * c1->cols; i++) {
        if (c1->data[i] != c2->data[i]) {
            diff_count++;
        }
    }
}

```

```

    }
    printf("Results match: %s\n", diff_count == 0 ? "Yes" : "No");

    free_matrix(a);
    free_matrix(b);
    free_matrix(c1);
    free_matrix(c2);
}

double get_time_diff(struct timespec start, struct timespec end) {
    return (end.tv_sec - start.tv_sec) + (end.tv_nsec - start.tv_nsec) / 1e9;
}

int main() {
    srand(time(NULL));
    compare_matrix_multiplication();
    return 0;
}

```

Module 14: Testing and Quality Assurance

Exercise 14.1: Unit Testing Framework

```

#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <time.h>

// Simple unit testing framework

// Test result structure
typedef struct {
    int passed;
    int failed;
    int skipped;
    double total_time;
} test_result_t;

// Global test result tracker
static test_result_t g_test_result = {0};

// Assertion functions
#define ASSERT_TRUE(condition) \
    do { \
        if (!(condition)) { \

```

```

        printf("ASSERT_TRUE failed at %s:%d\n", __FILE__, __LINE__); \
        g_test_result.failed++; \
        return; \
    } \
    g_test_result.passed++; \
} while(0)

#define ASSERT_FALSE(condition) \
do { \
    if (condition) { \
        printf("ASSERT_FALSE failed at %s:%d\n", __FILE__, __LINE__); \
        g_test_result.failed++; \
        return; \
    } \
    g_test_result.passed++; \
} while(0)

#define ASSERT_EQUAL(expected, actual) \
do { \
    if ((expected) != (actual)) { \
        printf("ASSERT_EQUAL failed at %s:%d: expected %d, got %d\n", \
            __FILE__, __LINE__, (expected), (actual)); \
        g_test_result.failed++; \
        return; \
    } \
    g_test_result.passed++; \
} while(0)

#define ASSERT_STRING_EQUAL(expected, actual) \
do { \
    if (strcmp((expected), (actual)) != 0) { \
        printf("ASSERT_STRING_EQUAL failed at %s:%d: expected '%s', got '%s'\n", \
            __FILE__, __LINE__, (expected), (actual)); \
        g_test_result.failed++; \
        return; \
    } \
    g_test_result.passed++; \
} while(0)

// Test function type
typedef void (*test_func_t)(void);

// Test case structure
typedef struct {
    const char *name;
    test_func_t function;
} test_case_t;

```



```

// Example functions to test
int add(int a, int b) {
    return a + b;
}

int subtract(int a, int b) {
    return a - b;
}

const char* get_hello_message(void) {
    return "Hello, World!";
}

// Test cases
void test_add_positive(void) {
    ASSERT_EQUAL(5, add(2, 3));
}

void test_add_negative(void) {
    ASSERT_EQUAL(-1, add(2, -3));
}

void test_subtract(void) {
    ASSERT_EQUAL(2, subtract(5, 3));
}

void test_string_message(void) {
    ASSERT_STRING_EQUAL("Hello, World!", get_hello_message());
}

// Test suite
test_case_t test_suite[] = {
    {"test_add_positive", test_add_positive},
    {"test_add_negative", test_add_negative},
    {"test_subtract", test_subtract},
    {"test_string_message", test_string_message},
    {NULL, NULL} // Sentinel
};

// Test runner
void run_tests(void) {
    printf("Running tests...\n");

    struct timespec start, end;
    clock_gettime(CLOCK_MONOTONIC, &start);

```

```

for (int i = 0; test_suite[i].name != NULL; i++) {
    printf("Running %s... ", test_suite[i].name);

    // Reset test counters for each test
    int initial_passed = g_test_result.passed;
    int initial_failed = g_test_result.failed;

    test_suite[i].function();

    if (g_test_result.failed > initial_failed) {
        printf("FAILED\n");
    } else {
        printf("PASSED\n");
    }
}

clock_gettime(CLOCK_MONOTONIC, &end);
g_test_result.total_time = (end.tv_sec - start.tv_sec) +
    (end.tv_nsec - start.tv_nsec) / 1e9;

// Print summary
printf("\nTest Results:\n");
printf("Passed: %d\n", g_test_result.passed);
printf("Failed: %d\n", g_test_result.failed);
printf("Skipped: %d\n", g_test_result.skipped);
printf("Total time: %.6f seconds\n", g_test_result.total_time);

if (g_test_result.failed == 0) {
    printf("All tests passed!\n");
} else {
    printf("Some tests failed!\n");
}
}

int main() {
    run_tests();
    return g_test_result.failed > 0 ? 1 : 0;
}

```

Exercise 14.2: Mocking Framework

```

#include <stdio.h>
#include <stdlib.h>
#include <string.h>

```

```

// Mocking framework example

// Real database interface
typedef struct {
    int (*connect)(const char *host, int port);
    int (*query)(const char *sql, char **result);
    void (*disconnect)(void);
} database_interface_t;

// Real database implementation
static int real_db_connect(const char *host, int port) {
    printf("Connecting to database at %s:%d\n", host, port);
    return 1; // Success
}

static int real_db_query(const char *sql, char **result) {
    printf("Executing query: %s\n", sql);
    *result = strdup("Mock query result");
    return 1; // Success
}

static void real_db_disconnect(void) {
    printf("Disconnecting from database\n");
}

database_interface_t real_database = {
    .connect = real_db_connect,
    .query = real_db_query,
    .disconnect = real_db_disconnect
};

// Mock database implementation
static int mock_connect_calls = 0;
static int mock_query_calls = 0;
static int mock_disconnect_calls = 0;
static int mock_connect_result = 1;
static int mock_query_result = 1;

static int mock_db_connect(const char *host, int port) {
    mock_connect_calls++;
    printf("Mock connect called (%d times)\n", mock_connect_calls);
    return mock_connect_result;
}

static int mock_db_query(const char *sql, char **result) {
    mock_query_calls++;

```

```

    printf("Mock query called (%d times): %s\n", mock_query_calls, sql);
    *result = strdup("Mock result");
    return mock_query_result;
}

static void mock_db_disconnect(void) {
    mock_disconnect_calls++;
    printf("Mock disconnect called (%d times)\n", mock_disconnect_calls);
}

database_interface_t mock_database = {
    .connect = mock_db_connect,
    .query = mock_db_query,
    .disconnect = mock_db_disconnect
};

// Function under test
int fetch_user_data(database_interface_t *db, const char *user_id, char **data) {
    if (!db->connect("localhost", 5432)) {
        return 0;
    }

    char sql[256];
    snprintf(sql, sizeof(sql), "SELECT * FROM users WHERE id = '%s'", user_id);

    int result = db->query(sql, data);
    db->disconnect();

    return result;
}

// Test with mock
void test_fetch_user_data_success(void) {
    // Reset mock counters
    mock_connect_calls = 0;
    mock_query_calls = 0;
    mock_disconnect_calls = 0;
    mock_connect_result = 1;
    mock_query_result = 1;

    char *data = NULL;
    int result = fetch_user_data(&mock_database, "123", &data);

    // Verify results
    if (result != 1) {
        printf("Test failed: Expected success, got failure\n");
        return;
    }
}

```

```

    }

    if (mock_connect_calls != 1) {
        printf("Test failed: Expected 1 connect call, got %d\n", mock_connect_calls);
        return;
    }

    if (mock_query_calls != 1) {
        printf("Test failed: Expected 1 query call, got %d\n", mock_query_calls);
        return;
    }

    if (mock_disconnect_calls != 1) {
        printf("Test failed: Expected 1 disconnect call, got %d\n", mock_disconnect_calls);
        return;
    }

    printf("Test passed: All mock interactions verified\n");
    free(data);
}

void test_fetch_user_data_connection_failure(void) {
    // Reset mock counters
    mock_connect_calls = 0;
    mock_query_calls = 0;
    mock_disconnect_calls = 0;
    mock_connect_result = 0; // Simulate connection failure

    char *data = NULL;
    int result = fetch_user_data(&mock_database, "123", &data);

    // Verify results
    if (result != 0) {
        printf("Test failed: Expected failure, got success\n");
        return;
    }

    if (mock_connect_calls != 1) {
        printf("Test failed: Expected 1 connect call, got %d\n", mock_connect_calls);
        return;
    }

    // Should not call query or disconnect on connection failure
    if (mock_query_calls != 0) {
        printf("Test failed: Expected 0 query calls, got %d\n", mock_query_calls);
        return;
    }
}

```

```

        if (mock_disconnect_calls != 0) {
            printf("Test failed: Expected 0 disconnect calls, got %d\n", mock_disconnect_calls);
            return;
        }

        printf("Test passed: Connection failure handled correctly\n");
    }

int main() {
    printf("Testing with mock database...\n");
    test_fetch_user_data_success();
    test_fetch_user_data_connection_failure();

    printf("\nFor comparison, running with real database:\n");
    char *data = NULL;
    fetch_user_data(&real_database, "123", &data);
    free(data);

    return 0;
}

```

Module 15: Advanced Topics

Exercise 15.1: Multithreading Example

```

#include <stdio.h>
#include <stdlib.h>
#include <pthread.h>
#include <unistd.h>

// Thread function
void* thread_function(void* arg) {
    int thread_id = *(int*)arg;
    printf("Hello from thread %d\n", thread_id);

    // Simulate some work
    sleep(1);

    printf("Thread %d finishing\n", thread_id);
    return NULL;
}

int main() {
    pthread_t thread1, thread2;

```

```

int id1 = 1, id2 = 2;

// Create threads
if (pthread_create(&thread1, NULL, thread_function, &id1) != 0) {
    perror("pthread_create failed");
    exit(EXIT_FAILURE);
}

if (pthread_create(&thread2, NULL, thread_function, &id2) != 0) {
    perror("pthread_create failed");
    exit(EXIT_FAILURE);
}

// Wait for threads to complete
pthread_join(thread1, NULL);
pthread_join(thread2, NULL);

printf("All threads completed\n");
return 0;
}

```

Exercise 15.2: System Programming Example

```

#include <stdio.h>
#include <unistd.h>
#include <sys/wait.h>
#include <stdlib.h>

int main() {
    pid_t pid = fork();

    if (pid == -1) {
        perror("fork failed");
        exit(EXIT_FAILURE);
    } else if (pid == 0) {
        // Child process
        printf("Child process (PID: %d)\n", getpid());
        exit(EXIT_SUCCESS);
    } else {
        // Parent process
        printf("Parent process (PID: %d), Child PID: %d\n", getpid(), pid);
        wait(NULL); // Wait for child to complete
        printf("Child process completed\n");
    }
}

```

```
    return 0;  
}
```

This appendix provides solutions to selected exercises from each module of the course. These solutions demonstrate good programming practices and illustrate the concepts covered in each module. Remember that there are often multiple ways to solve a problem, and these solutions are meant to serve as examples rather than the only correct approach.