

Writing Kernel with AI

K19G

2026-10-04

Table of contents

Writing Kernel with AI	7
CPU to AI-Written Kernels	7
Welcome to Kernel Engineering	7
The Trajectory of This Book	7
Who This Is For	8
Formats	8
 Part I: Parallelism, from the CPU up	 9
Parallelism on the CPU	10
Parallelism on the CPU	11
1. Instruction-Level Parallelism (ILP) and Microarchitecture	11
The Out-of-Order Engine	12
2. SIMD: Vectorized Data Parallelism	12
Vector Widths Across Hardware Architectures	12
Explicit AVX2 Vector Operations in C++	13
Alignment and Vector Loads	13
3. Thread-Level Parallelism: Heavyweight vs Lightweight	14
The Cost of OS Threads	14
User-Space Concurrency (M:N Models)	14
4. Cache Lines and False Sharing	14
The False Sharing Trap	14
Remediation via Padding	15
5. Architectural Contrast: CPU vs GPU Design Philosophy	15
 Roofline Lab: Predict Then Measure	 17
Roofline Lab: Predict Then Measure	18
1. The Roofline Model Formulation	18
Definitions	18
The Ridge Point (I^*)	19
2. Hardware Ridge Points Across Generations	19
3. Dissecting Common Machine Learning Operations	20
Case 1: Vector Addition ($y = x + z$)	20
Case 2: LayerNorm / RMSNorm	20
Case 3: Dense Matrix Multiplication (GEMM)	21
4. The Multi-Tier Roofline (Hierarchy of Limits)	21
5. Profiling Roofline Performance with NVIDIA Nsight Compute	21
Key NCU Metrics to Inspect:	22

The CUDA Programming Model & GPU Puzzles	23
The CUDA Programming Model & GPU Puzzles	24
1. The Logical vs Physical Hierarchy	24
The Dimensions	24
2. Multi-Dimensional Index Calculations	25
1D Global Index Calculation	25
2D Global Index Calculation (Matrix Coordinates)	25
3. Warp Divergence and the SIMT Execution Mask	25
The Cost of Divergence	26
Best Practices to Eliminate Divergence	26
4. Complete Worked Example: Vectorized FMA with Error Handling	26
5. Inspecting Intermediate PTX and Machine SASS	28
The Memory Hierarchy in Anger: Matrix Transpose	29
The Memory Hierarchy in Anger: Matrix Transpose	30
1. The GPU Memory Hierarchy	30
Memory Space Properties	31
2. Global Memory Coalescing: The 10x Lever	31
The Golden Rule of Coalescing	31
3. Shared Memory Bank Conflicts	32
The Conflict Condition	32
4. Case Study: Step-by-Step Optimization of Matrix Transpose	32
Step 1: Naive Transpose (Uncoalesced Writes)	32
Step 2: Coalesced Transpose via Shared Memory (With Bank Conflicts)	33
Step 3: Conflict-Free Transpose with Shared Memory Padding	34
Full Production Code	34
5. Performance Comparison	35
Part II: The GEMM Worklog: Naive to cuBLAS	36
GEMM Worklog I: Naive to Tiled	37
GEMM Worklog I: Naive to Tiled	38
1. Dissecting the Naive GEMM: Why is it Slow?	38
The Naive CUDA Kernel	38
Performance Autopsy	39
2. The Breakthrough: Shared Memory 2D Block Tiling	39
Tile Dimensions and Reuse Factor	40
3. The 2D Tiled CUDA Kernel	40
4. Benchmark Results on NVIDIA A100	41
Why Did Performance Jump by 27x?	42
GEMM Worklog II: Registers & Warp Tiling	43
GEMM Worklog II: Registers & Warp Tiling	44
Micro-Tiling in Registers	44

Vectorized Memory Loads with <code>float4</code>	44
The Warptiling Hierarchy	44
Tensor Cores & WMMA: The Second Worklog	45
Tensor Cores & WMMA: The Second Worklog	46
The Tensor Core Advantage	46
Direct Programming via WMMA (Warp Matrix Multiply Accumulate)	46
 Part III: Attention & Profiling	 47
Profiling & Debugging Like a Pro	48
Profiling & Debugging Like a Pro	49
The Professional Diagnostic Stack	49
Debugging Sabotaged Kernels: 3 Classic Incident Scenarios	49
Attention: The Kernel That Ate the World & FlashAttention v1	50
Attention: The Kernel That Ate the World & FlashAttention v1	51
The Breakthrough: Online Softmax & Tiling	51
The Online Softmax Formula	51
 Part IV: The Modern Frontier	 52
FlashAttention from Scratch (FA1 → FA2 → FA3)	53
FlashAttention from Scratch (FA1 → FA2 → FA3)	54
FlashAttention Evolution	54
Beating cuBLAS on an H100	55
Beating cuBLAS on an H100	56
1. Tensor Memory Accelerator (TMA)	56
2. Warp Specialization	56
Triton → CUTLASS → CuTe-DSL	57
Triton → CUTLASS → CuTe-DSL	58
OpenAI Triton	58
CUTLASS & CuTe	58
Inference-Serving Kernels	59
Inference-Serving Kernels	60
PagedAttention (vLLM)	60
Blackwell & NVFP4	61

Blackwell & NVFP4	62
Blackwell Architectural Highlights	62
Flash Attention 4	63
Flash Attention 4	64
The 2026 Frontier: Breaking the 1 PFLOPS Single-GPU Barrier	64
PagedAttention and the KV Cache	65
PagedAttention and the KV Cache	66
Mental model	66
Why contiguous KV wastes memory	66
Capacity sketch	67
KV cache in Transformers (reminder)	67
Block structure	67
Worked layout numbers	68
Block table and allocator	68
Prefix sharing (copy-on-write intuition)	68
Prefill vs decode under paging	68
Implication	69
Paged KV layouts (kernel writer view)	69
Illustrative host-side table (not a production kernel)	70
Implications for custom kernels	71
Verification and benchmark notes	71
Functional checks	71
Performance / capacity checks	71
Profiling tips	72
Pitfalls	72
Checkpoint	72
Summary	73
Sources (for further reading)	73
 Part V: AI-Written Kernels	 74
DeepSeek: FlashMLA & DeepGEMM	75
DeepSeek: FlashMLA & DeepGEMM	76
FlashMLA: Multi-Head Latent Attention	76
DeepGEMM: Clean Architecture over Brute Force	76
LLM-Generated Kernels: The Agent + Profiler Loop	77
LLM-Generated Kernels: The Agent + Profiler Loop	78
The Agentic Optimization Loop	78
Google DeepMind AlphaEvolve & KernelBench	78
The Role of the Human Kernel Engineer	78

Part VI: Kernel Engineering Capstone Projects	79
Crusoe Capstone Specifications & Problem Tracks	80
Crusoe Capstone Specifications & Problem Tracks	81
Problem Track 1: Fused Sparse Attention	81
Problem Track 2: Quantized FP4/FP8 GEMM with Dynamic Dequantization	81
Problem Track 3: High-Throughput Speculative Decoding Verification Kernel	81
Verification, Profiling, and Benchmarking Harness	82
Verification, Profiling, and Benchmarking Harness	83
1. Correctness Verification	83
2. CUDA Event Benchmarking	83
Demo Day Presentation & Portfolio	84
Demo Day Presentation & Portfolio	85
Deliverables Checklist	85

Writing Kernel with AI

CPU to AI-Written Kernels

A comprehensive, production-grade guide to high-performance kernel engineering: from the fundamentals of CPU instruction parallelism to hand-tuned CUDA, FlashAttention, modern Hopper/Blackwell hardware architectures, and AI-driven kernel optimization loops.

Welcome to Kernel Engineering

Accelerated computing and modern artificial intelligence rest on a single engineering discipline: **extracting peak theoretical throughput from hardware.**

Whether training frontier foundation models or serving millions of concurrent inference tokens, the software stack converges onto specialized computational kernels. Writing fast kernels is not a matter of guesswork or trial-and-error—it is an empirical science governed by the physics of computation: arithmetic intensity, memory bandwidth, latency hiding, and hardware execution hierarchies.

The Trajectory of This Book

This curriculum spans six rigorous parts designed to take you from first principles to the cutting edge:

1. **Part I: Parallelism, from the CPU up:** Understand how modern CPUs achieve massive throughput via SIMD vectorization and multi-threading before stepping into the GPU execution hierarchy: warps, streaming multiprocessors, memory coalescing, and the Roofline model.
2. **Part II: The GEMM Worklog: Naive → cuBLAS:** Master General Matrix Multiply (GEMM)—the operation that comprises over 90% of deep learning FLOPs. Step by step, elevate a naive kernel achieving 1.3% of GPU capacity to a hierarchical warp-tiled and Tensor Core (WMMA) implementation rivaling NVIDIA’s cuBLAS library.
3. **Part III: Attention & Profiling:** Diagnose hardware bottlenecks using NVIDIA Nsight Compute like a compiler engineer. Construct FlashAttention v1 from scratch, utilizing the online softmax algorithm to eliminate high-bandwidth memory bottlenecks.
4. **Part IV: The Modern Frontier:** Explore modern production kernels running on Hopper (H100) and Blackwell (B200) architectures: FlashAttention-2/3/4, TMA asynchronous copy, WGMMMA, Triton, CUTLASS, CuTe-DSL, PagedAttention, and NVFP4.

5. **Part V: AI-Written Kernels:** Analyze DeepSeek’s ultra-compact, high-performance kernels (FlashMLA and DeepGEMM). Step into the future of kernel development: using autonomous LLM agents (e.g., DeepMind AlphaEvolve, KernelBench) with profiler-in-the-loop guidance.
 6. **Part VI: Kernel Engineering Capstone Projects:** Put your skills into practice with industry-grade capstone specifications developed with Crusoe, complete with verification harnesses, performance benchmarks, and demo-day technical presentations.
-

Who This Is For

- **ML Infrastructure & Systems Engineers:** Developers optimizing training clusters and inference backends (vLLM, SGLang, TensorRT-LLM).
 - **GPU Programmers & Performance Engineers:** Engineers seeking mastery of modern GPU architectures (Hopper, Blackwell), memory coalescing, and async execution pipelines.
 - **Software Engineers Transitioning to Accelerated Computing:** Developers wanting deep understanding of what happens when high-level PyTorch or JAX code hits the silicon.
-

Formats

This book is available in multiple formats via the library portal: - **HTML** — responsive interactive reading experience with syntax highlighting and light/dark theme toggle. - **PDF** — complete typeset document for offline review and reference. - **EPUB** — eBook format optimized for tablets and mobile e-readers.

Part I: Parallelism, from the CPU up

Parallelism on the CPU


Parallelism on the CPU

Before programming thousands of lockstep cores on a GPU, we must build a mechanical understanding of how modern CPUs execute concurrent and parallel work. High-performance computing is fundamentally an exercise in orchestrating physical execution units while feeding them data fast enough to prevent pipeline stalls.

Modern CPUs extract performance across three distinct tiers of parallelism: 1. **Instruction-Level Parallelism (ILP)**: Pipelining, superscalar dispatch, out-of-order execution, and branch prediction within a single core. 2. **Data-Level Parallelism (SIMD)**: Single Instruction, Multiple Data execution where one machine instruction operates over a register packed with multiple independent values. 3. **Thread-Level Parallelism (TLP)**: Multiple independent execution streams running concurrently across physical and logical processor cores.

1. Instruction-Level Parallelism (ILP) and Microarchitecture

A modern CPU core is not a simple sequential state machine. It is an asynchronous processing pipeline designed to retire multiple micro-operations (uops) per clock cycle.



The Out-of-Order Engine

1. **Instruction Fetch & Branch Prediction:** The CPU fetches instructions ahead of the current program counter using branch target buffers (BTB) and direction predictors. When a branch is mispredicted, the entire speculative pipeline must be flushed, costing 15 to 20 cycles of wasted energy.
 2. **Decode and Micro-op Generation:** Complex x86 instructions are broken down into simpler, RISC-like micro-operations (uops).
 3. **Register Renaming & Reservation Stations:** To break false register dependencies (Write-After-Read and Write-After-Write hazards), the CPU maps architectural registers (e.g., `rax`, `ymm0`) to a much larger pool of physical registers.
 4. **Out-of-Order (OoO) Dispatch:** Independent instructions whose operands are ready in the cache or register file execute out of program order across parallel execution ports.
 5. **In-Order Retirement:** The Reorder Buffer (ROB) ensures that memory state updates and architectural register commits happen in strict program order, guaranteeing consistent exception semantics.
-

2. SIMD: Vectorized Data Parallelism

While ILP extracts parallelism implicitly from scalar instructions, SIMD (Single Instruction, Multiple Data) makes data parallelism explicit. Instead of performing arithmetic on a single 32-bit floating-point value, SIMD instructions load, operate upon, and store entire vectors of values simultaneously.

Vector Widths Across Hardware Architectures

Vector Extension	Vector Width	32-bit Floats per Vector	Target Architecture
x86 SSE2	128-bit	4	Legacy x86-64
x86 AVX2	256-bit	8	Intel Haswell+ / AMD Zen2+
x86 AVX-512	512-bit	16	Intel Skylake-SP / Emerald Rapids, AMD Zen4+
x86 AVX10	256 / 512-bit	8 or 16	Modern Intel Architecture
ARM Neon	128-bit	4	Apple Silicon, ARM Neoverse V1/V2, Cortex-A
ARM SVE / SVE2	Scalable (128–2048b)	Variable	Fujitsu A64FX, AWS Graviton3/4

Explicit AVX2 Vector Operations in C++

When performance is paramount, reliance on compiler autovectorization is risky. Compilers frequently fail to vectorize loops due to pointer aliasing concerns or unknown loop boundary alignments. Engineers write architecture-specific intrinsic functions:

```
#include <immintrin.h>
#include <iostream>
#include <vector>

void vector_fma_avx2(const float* __restrict__ a,
                    const float* __restrict__ b,
                    const float* __restrict__ c,
                    float* __restrict__ out,
                    size_t n) {
    size_t i = 0;

    // Process 8 floats per iteration using 256-bit AVX2 registers
    for (; i + 8 <= n; i += 8) {
        // Load vectors from memory
        __m256 va = _mm256_loadu_ps(&a[i]);
        __m256 vb = _mm256_loadu_ps(&b[i]);
        __m256 vc = _mm256_loadu_ps(&c[i]);

        // Fused Multiply-Add: out = (va * vb) + vc
        // Computes 8 FMAs (16 floating point operations) in 4-5 cycles latency
        __m256 vres = _mm256_fmadd_ps(va, vb, vc);

        // Store back to destination
        _mm256_storeu_ps(&out[i], vres);
    }

    // Scalar cleanup tail loop
    for (; i < n; ++i) {
        out[i] = a[i] * b[i] + c[i];
    }
}
```

Alignment and Vector Loads

- `_mm256_loadu_ps`: Unaligned load. Safe across any memory address, but incurs small penalty if crossing a 64-byte cache line boundary.
- `_mm256_load_ps`: Aligned load. Requires the source address to be strictly 32-byte aligned (`alignas(32)` or `posix_memalign`). A non-aligned pointer will trigger a hardware general protection fault (`SIGSEGV`).

3. Thread-Level Parallelism: Heavyweight vs Lightweight

To leverage multi-core CPUs, applications distribute work across threads. How threads are modeled dramatically impacts context switching overhead and memory footprint.

```
OS Heavyweight Threads:
```

```
[Kernel Scheduler] ---> [Core 0: OS Thread (2-8 MB Stack)] (Switch: ~1-3 us)
                        ---> [Core 1: OS Thread (2-8 MB Stack)]
```

```
Lightweight User-Space Green Threads / Fibers:
```

```
[OS Thread / SM]      ---> [Fiber A (2 KB)] -> [Fiber B (2 KB)] -> [Fiber C (2 KB)]
```

The Cost of OS Threads

POSIX threads (`pthread_create`) or OpenMP workers (`#pragma omp parallel for`) allocate their own kernel task structure and a fixed memory stack (typically 2 MB to 8 MB). - When an OS thread context switches, the CPU must switch privilege rings into the kernel, save and restore architectural register state (including wide AVX-512 state), and suffer TLB and cache pollution. - Total latency: **1,000 to 3,000 CPU cycles** (1–3 microseconds).

User-Space Concurrency (M:N Models)


Modern programming environments decouple logical concurrency from physical OS threads: - **Go Goroutines**: Allocate a dynamic 2 KB stack that grows on demand. The runtime multiplexes thousands of goroutines across a fixed thread pool (`GOMAXPROCS`), switching in user space within ~10 to 50 nanoseconds. - **Java Virtual Threads (Project Loom)**: Suspendable continuation objects stored on the garbage-collected heap. - **Rust Tokio**: Futures compiled into zero-allocation state machines driven by a lock-free work-stealing executor.

4. Cache Lines and False Sharing

CPU memory controllers do not fetch individual 4-byte floating-point numbers from system DRAM. All transfers between main memory and CPU caches occur in fixed chunks of **64 bytes**, called **cache lines**.

The False Sharing Trap

When two distinct threads running on separate CPU cores write to independent variables that happen to sit within the same 64-byte cache line, the CPU hardware coherence engine (e.g., MESI protocol) must invalidate that entire cache line across both cores back and forth:



This phenomenon—**false sharing**—can slow down multi-threaded code by a factor of 50x to 100x, reducing parallel throughput below that of a single-threaded loop.

Remediation via Padding

```
// Bad: Threads write to adjacent elements in an array
struct BadThreadState {
    uint64_t counter; // 8 bytes -> adjacent elements share a 64B line!
};

// Good: Align to 64 bytes (the hardware cache line size)
struct alignas(64) GoodThreadState {
    uint64_t counter;
    // 56 bytes of implicit padding inserted by the compiler
};
```

5. Architectural Contrast: CPU vs GPU Design Philosophy

Why do GPUs dominate machine learning and scientific compute when CPUs have such sophisticated out-of-order execution engines?

```
CPU Architecture (Latency-Optimized):
+-----+
| [Control: Branch Predictor, Out-of-Order Engine] |
| [Large L1 / L2 / L3 Caches (Tens of Megabytes)] |
| [A Few Powerful ALU / FMA Execution Cores (8-64 Cores)] |
+-----+

GPU Architecture (Throughput-Optimized):
+-----+
| [Control] [Control] [Control] [Control] |
| [Thousands of Compact ALU / Tensor Cores] |
| [High-Bandwidth HBM3 / GDDR Memory Interface] |
+-----+
```

+-----+

Dimension	Modern Server CPU (e.g. AMD Genoa / Intel Xeon)	Modern Data Center GPU (e.g. NVIDIA H100)
Optimization Target	Minimum latency for a single sequence	Maximum aggregate throughput
Physical Cores	64 – 128 complex cores	132 – 144 Streaming Multiprocessors (16,896 cores)
Control Logic	Heavy: speculative, out-of-order	Light: SIMT lockstep warps, in-order dispatch
Cache Hierarchy	Massive: ~384 MB L3 cache	Compact: 50 MB L2 cache, 228 KB SM SRAM
Memory Bandwidth	300 – 460 GB/s (DDR5)	3,350 GB/s (HBM3)
FP32 Compute	4 – 8 TFLOP/s	67 TFLOP/s (Vector) / 989 TFLOP/s (Tensor)

Understanding this trade-off is the foundation of accelerated computing: **CPUs excel at complex branching, low latency, and irregular memory traversals. GPUs excel at massive arithmetic intensity and regular, uniform matrix operations.**

Roofline Lab: Predict Then Measure

Roofline Lab: Predict Then Measure

Every high-performance kernel is bounded by one of two physical constraints: 1. **Compute Bound:** The execution units (ALUs, FPUs, Tensor Cores) are saturated executing instructions; the memory interface has idle bandwidth. 2. **Memory Bound:** The execution units starve waiting for data to arrive from off-chip DRAM (HBM or GDDR); arithmetic throughput is a fraction of hardware peak.

Before writing a single line of CUDA, a kernel engineer must predict which regime the workload inhabits. Writing faster assembly or unrolling loops inside a memory-bound kernel achieves zero speedup. Conversely, optimizing memory tiling for a kernel already compute-bound yields diminishing returns.


1. The Roofline Model Formulation

The Roofline Model, formulated by Williams et al. (Berkeley), visualizes maximum achievable performance as a function of **Operational Intensity** (also called Arithmetic Intensity).

Definitions

$$\text{Operational Intensity } (I) = \frac{\text{Total Arithmetic Operations (FLOPs)}}{\text{Total Bytes Transferred from DRAM (Bytes)}}$$

$$\text{Attainable Performance } (P) = \min(\text{Peak Hardware Compute } (P_{\text{peak}}), \text{Peak Memory Bandwidth } (B) \times I)$$





The Ridge Point (I^*)

The intersection between the memory bandwidth slope and the peak compute roof defines the **Ridge Point** (I^*):

$$I^* = \frac{P_{\text{peak}}}{B_{\text{peak}}}$$

- If a workload's intensity $I < I^*$: The workload is **strictly memory-bound**. Maximum possible throughput is capped at $B_{\text{peak}} \times I$. To go faster, you must reduce memory traffic (e.g. through kernel fusion or caching).
- If a workload's intensity $I \geq I^*$: The workload is **compute-bound**. Maximum possible throughput is capped at P_{peak} . To go faster, you must optimize instruction throughput, eliminate warp divergence, or switch to lower-precision numeric formats (e.g. FP16, BF16, FP8).

2. Hardware Ridge Points Across Generations

Notice how rapidly the Ridge Point has shifted over successive GPU generations. Compute capability has grown far faster than memory bandwidth:

GPU Model	FP32 Peak Compute	HBM / GDDR Bandwidth	Vector Ridge Point (I_{FP32}^*)	FP16 Tensor Compute	Tensor Ridge Point (I_{Tensor}^*)
Tesla V100 (PCIe)	14.0 TFLOP/s	900 GB/s	15.5 FLOPs/B	112 TFLOP/s	124.4 FLOPs/B
Ampere A100 (SXM4)	19.5 TFLOP/s	2,039 GB/s	9.6 FLOPs/B	312 TFLOP/s	153.0 FLOPs/B
Hopper H100 (SXM5)	67.0 TFLOP/s	3,350 GB/s	20.0 FLOPs/B	989 TFLOP/s	295.2 FLOPs/B
Blackwell B200	90.0 TFLOP/s	8,000 GB/s	11.25 FLOPs/B	2,250 TFLOP/s (FP16)	281.2 FLOPs/B

Key Takeaway: On an NVIDIA H100 running Tensor Core operations, an algorithm must execute nearly **300 floating-point operations for every single byte loaded from HBM** before it can even theoretically achieve peak compute saturation.

3. Dissecting Common Machine Learning Operations

Let us calculate the operational intensity for four fundamental deep learning operations executed on an NVIDIA A100 (SXM4: $B = 2,039$ GB/s, $P_{\text{vector}} = 19.5$ TFLOP/s).

Case 1: Vector Addition ($y = x + z$)

Given two vectors of length $N = 10^7$ FP32 elements: - **Operations:** N additions = 10^7 FLOPs. - **Memory Traffic:** - Read x : $10^7 \times 4$ B = 40 MB. - Read z : $10^7 \times 4$ B = 40 MB. - Write y : $10^7 \times 4$ B = 40 MB. - Total Memory = 120 MB = 1.2×10^8 Bytes. - **Operational Intensity:**

$$I = \frac{10^7 \text{ FLOPs}}{1.2 \times 10^8 \text{ Bytes}} \approx 0.0833 \text{ FLOPs/Byte}$$

- **Analysis:** $I = 0.0833 \ll I^* = 9.6$. Strictly memory bound. - **Maximum Theoretical Throughput:**

$$\text{Throughput} = 2,039 \text{ GB/s} \times 0.0833 \text{ FLOPs/Byte} \approx 169.8 \text{ GFLOP/s}$$

The GPU achieves less than **0.9%** of its FP32 compute capability. The ALU cores spend 99% of their time stalled waiting for memory controllers.

Case 2: LayerNorm / RMSNorm

RMSNorm computes: $y_i = \frac{x_i}{\sqrt{\frac{1}{d} \sum_{j=1}^d x_j^2 + \epsilon}} \cdot \gamma_i$. For a hidden dimension $d = 4096$: - **Operations per token:** - Square: d FLOPs. - Sum reduction: $d - 1$ FLOPs. - Mean, add ϵ , reciprocal square root (rsqrt): ~ 4 FLOPs. - Scale by γ and write: $2d$ FLOPs. - Total $\approx 4d \approx 16,384$ FLOPs. - **Memory per token (unfused naive PyTorch):** - Read input x : $d \times 4$ B = 16 KB. - Intermediates written and re-read from global memory: $\sim 3 \times 16$ KB = 48 KB. - Write output y : 16 KB. - Total ~ 80 KB. - **Operational Intensity:**

$$I = \frac{16,384 \text{ FLOPs}}{81,920 \text{ Bytes}} \approx 0.20 \text{ FLOPs/Byte}$$

- **Analysis:** Completely memory bound. This is why fused normalization kernels (which keep intermediates in registers and shared memory) deliver massive 5x to 10x speedups.

Case 3: Dense Matrix Multiplication (GEMM)


Compute $C = A \times B$ where $A, B, C \in \mathbb{R}^{M \times N}$, with $M = K = N = 4096$: - **Operations:** $2 \times M \times N \times K = 2 \times 4096^3 \approx 137.4 \times 10^9$ FLOPs (137.4 GFLOPs). - **Minimum Algorithmic Memory (I/O):** - Read A: $4096^2 \times 4 \text{ B} = 67.1 \text{ MB}$. - Read B: $4096^2 \times 4 \text{ B} = 67.1 \text{ MB}$. - Write C: $4096^2 \times 4 \text{ B} = 67.1 \text{ MB}$. - Total $\approx 201.3 \text{ MB}$. - **Operational Intensity (with perfect reuse):**

$$I_{\text{optimal}} = \frac{1.374 \times 10^{11} \text{ FLOPs}}{2.013 \times 10^8 \text{ Bytes}} \approx 682.6 \text{ FLOPs/Byte}$$

- **Analysis:** $I = 682.6 > I^* = 9.6$ (and even $> I_{\text{Tensor}}^* = 153.0$). Dense GEMM is **thoroughly compute-bound** when memory reuse is leveraged via shared memory tiling!

4. The Multi-Tier Roofline (Hierarchy of Limits)

Modern GPUs have multiple memory tiers. A kernel that exhausts L2 cache bandwidth or shared memory bandwidth hits an intermediate roofline before reaching DRAM:



1. **HBM / DRAM Bandwidth:** The slowest and largest layer. Optimization goal: load each byte from DRAM exactly once into on-chip memory.
2. **L2 Cache Bandwidth:** Over 2x to 3x faster than DRAM. Shared across all SMs.
3. **Shared Memory / L1 SRAM Bandwidth:** Over 5x to 8x faster than DRAM (~19 TB/s aggregate on H100). Serves warps inside a single SM.
4. **Register File Bandwidth:** Over 30 TB/s aggregate bandwidth. The only memory tier that can feed execution units at full single-cycle latency.

5. Profiling Roofline Performance with NVIDIA Nsight Compute

To verify where a kernel lands on the physical Roofline curve, run NVIDIA Nsight Compute (`ncu`):

```
# Profile Speed of Light and Roofline metrics
ncu --set full \
    --section SpeedOfLight \
    --section SpeedOfLight_HierarchicalTensorRooflineChart \
    -o profile_result \
    ./my_cuda_binary
```

Key NCU Metrics to Inspect:

- `dram__throughput.avg.pct_of_peak_sustained_elapsed`: High percentage (>80%) indicates saturation of off-chip DRAM bandwidth (Memory Bound).
- `sm__throughput.avg.pct_of_peak_sustained_elapsed`: High percentage (>80%) indicates saturation of SM computation pipelines (Compute Bound).
- `gpu__time_duration.sum`: Total kernel wall clock time.

By comparing these two percentages, you instantly know whether to optimize memory access patterns or computational instructions.

The CUDA Programming Model & GPU Puzzles


The CUDA Programming Model & GPU Puzzles

Modern GPUs execute programs using the **SIMT** (**Single Instruction, Multiple Threads**) execution model. While SIMD on a CPU executes a single vector instruction over vector registers on a single core, SIMT executes hundreds of independent scalar threads grouped into lockstep units called **warps**.

Writing effective CUDA kernels requires understanding both the logical programming abstractions and how the hardware scheduler physically maps them onto silicon.

1. The Logical vs Physical Hierarchy

CUDA defines a hierarchical programming model that scales seamlessly from a single-SM embedded chip (e.g. Jetson) to a massive data-center accelerator with 144 SMs (e.g. Hopper H100).



The Dimensions

1. **Thread:** The finest unit of execution. Each thread has its own program counter, private registers, and local memory stack.
2. **Warp:** A contiguous group of **32 threads**. The warp is the fundamental unit of scheduling on the GPU. All 32 threads in a warp execute the same instruction simultaneously.
3. **Thread Block (CTA - Cooperative Thread Array):** A group of up to 1,024 threads. Threads within the same block reside on the **same physical SM**, share fast on-chip Shared Memory (SRAM), and can synchronize using barrier primitives (`__syncthreads()`).
4. **Grid:** The full set of thread blocks dispatched to execute a kernel. Thread blocks within a grid execute independently; there is no guaranteed execution order across blocks.

2. Multi-Dimensional Index Calculations

CUDA allows Grids and Blocks to be configured in 1D, 2D, or 3D using `dim3`. The hardware provides built-in read-only variables: - `threadIdx.{x,y,z}`: Thread coordinates within its block. - `blockIdx.{x,y,z}`: Block coordinates within the grid. - `blockDim.{x,y,z}`: Dimensions (thread count) of the block. - `gridDim.{x,y,z}`: Dimensions (block count) of the grid.

1D Global Index Calculation

For a 1D grid processing a flat vector of length N :

$$\text{Global ID} = \text{blockIdx.x} \times \text{blockDim.x} + \text{threadIdx.x}$$

```
int idx = blockIdx.x * blockDim.x + threadIdx.x;
if (idx < n) {
    // Process element idx
}
```

2D Global Index Calculation (Matrix Coordinates)

For a 2D matrix with width W (columns) and height H (rows):

```
// Map block and thread coordinates to row and column
int col = blockIdx.x * blockDim.x + threadIdx.x;
int row = blockIdx.y * blockDim.y + threadIdx.y;

// Linear memory index (Row-Major layout)
int linear_idx = row * width + col;

if (row < height && col < width) {
    matrix[linear_idx] = 42.0f;
}
```

3. Warp Divergence and the SIMT Execution Mask

All 32 threads in a warp share a single instruction unit. If threads within a warp take different execution paths at a conditional branch (`if / else`), the hardware executes every branch path sequentially, disabling threads that are not on that path via a hardware **execution mask**.

```
Instruction Stream:
Thread 0-15:    Takes IF branch (condition is true)
```

```
Thread 16-31: Takes ELSE branch (condition is false)
```

Execution Timeline:

```
Cycle 1-10: [Thread 0-15 Active] [Thread 16-31 MASKED OFF] -> Executes IF body
```

```
Cycle 11-20: [Thread 0-15 MASKED OFF] [Thread 16-31 Active] -> Executes ELSE body
```

```
Reconverge: [All Threads 0-31 Active]
```

The Cost of Divergence

- When a warp diverges 2-ways, execution time roughly doubles.
- When 32 threads take 32 distinct branch paths (e.g. `switch(threadIdx.x % 32)`), the warp serializes into 32 separate execution passes, dropping warp computational throughput to **3.1%** of peak.

Best Practices to Eliminate Divergence

1. **Align branches by warp boundary:** Branches that depend strictly on `blockIdx` or `threadIdx.x / 32` cause zero divergence because all 32 threads in the warp evaluate to the exact same condition.
2. **Predicated Instructions:** For tiny branches (1-2 instructions), the compiler automatically uses predicated instructions instead of jumping.
3. **Use Intrinsic Math Primitives:** Replace `if (a < b) val = a; else val = b;` with `fminf(a, b)`, which compiles to a single non-branching instruction.

4. Complete Worked Example: Vectorized FMA with Error Handling

A robust production CUDA implementation requires careful host allocation, device memory management, asynchronous stream synchronization, and error checking macros.

```
#include <cuda_runtime.h>
#include <stdio.h>
#include <stdlib.h>

// Macro for robust CUDA API error checking
#define CUDA_CHECK(call) \
    do { \
        cudaError_t err = call; \
        if (err != cudaSuccess) { \
            fprintf(stderr, "CUDA error at %s:%d: %s\n", __FILE__, __LINE__, \
                cudaGetErrorString(err)); \
            exit(EXIT_FAILURE); \
        } \
    }
```

```

    } while (0)

// Fused Multiply-Add Kernel: d[i] = a[i] * b[i] + c[i]
__global__ void fma_kernel(const float* __restrict__ a,
                          const float* __restrict__ b,
                          const float* __restrict__ c,
                          float* __restrict__ d,
                          int n) {
    int idx = blockIdx.x * blockDim.x + threadIdx.x;
    if (idx < n) {
        // Generates single hardware FFMA instruction
        d[idx] = a[idx] * b[idx] + c[idx];
    }
}

int main(void) {
    const int N = 1 << 22; // ~4.19 million elements
    const size_t bytes = N * sizeof(float);

    // 1. Allocate Host Memory
    float *h_a = (float*)malloc(bytes);
    float *h_b = (float*)malloc(bytes);
    float *h_c = (float*)malloc(bytes);
    float *h_d = (float*)malloc(bytes);

    for (int i = 0; i < N; ++i) {
        h_a[i] = 1.0f;
        h_b[i] = 2.0f;
        h_c[i] = 3.0f;
    }

    // 2. Allocate Device Memory
    float *d_a, *d_b, *d_c, *d_d;
    CUDA_CHECK(cudaMalloc(&d_a, bytes));
    CUDA_CHECK(cudaMalloc(&d_b, bytes));
    CUDA_CHECK(cudaMalloc(&d_c, bytes));
    CUDA_CHECK(cudaMalloc(&d_d, bytes));

    // 3. Copy inputs to Device
    CUDA_CHECK(cudaMemcpy(d_a, h_a, bytes, cudaMemcpyHostToDevice));
    CUDA_CHECK(cudaMemcpy(d_b, h_b, bytes, cudaMemcpyHostToDevice));
    CUDA_CHECK(cudaMemcpy(d_c, h_c, bytes, cudaMemcpyHostToDevice));

    // 4. Configure Grid & Launch Kernel
    const int threads_per_block = 256;
    const int blocks_per_grid = (N + threads_per_block - 1) / threads_per_block;

```

```

fma_kernel<<<blocks_per_grid, threads_per_block>>>(d_a, d_b, d_c, d_d, N);

// Check for launch errors and synchronize
CUDA_CHECK(cudaGetLastError());
CUDA_CHECK(cudaDeviceSynchronize());

// 5. Copy results back
CUDA_CHECK(cudaMemcpy(h_d, d_d, bytes, cudaMemcpyDeviceToHost));

// Verify result: 1.0 * 2.0 + 3.0 = 5.0
printf("Result verification: d[0] = %f, d[N-1] = %f\n", h_d[0], h_d[N-1]);

// 6. Cleanup
CUDA_CHECK(cudaFree(d_a));
CUDA_CHECK(cudaFree(d_b));
CUDA_CHECK(cudaFree(d_c));
CUDA_CHECK(cudaFree(d_d));
free(h_a); free(h_b); free(h_c); free(h_d);

printf("Kernel execution completed successfully.\n");
return 0;
}

```

5. Inspecting Intermediate PTX and Machine SASS

To understand what the GPU physically executes, compile to PTX (virtual assembly) and SASS (machine assembly):

```

# Generate PTX virtual assembly
nvcc -ptx -arch=sm_90 -O3 fma.cu -o fma.ptx

# Disassemble binary machine SASS
nvcc -arch=sm_90 -O3 fma.cu -o fma
cuobjdump --dump-sass fma > fma.sass

```

Inside `fma.sass`, locate the inner loop body:

```

/*0030*/ LDG.E.SYS R2, [R4.64] ;      // Load 32-bit float from global memory
/*0040*/ LDG.E.SYS R3, [R6.64] ;      // Load 32-bit float from global memory
/*0050*/ FFMA R0, R2, R3, R0 ;        // Single-cycle Fused Multiply-Add instruction!
/*0060*/ STG.E.SYS [R8.64], R0 ;      // Store result back to global memory

```

Notice **FFMA**: The hardware fused multiply and addition into a single instruction without an intermediate rounding step, preserving numeric precision while doubling arithmetic throughput.

The Memory Hierarchy in Anger: Matrix Transpose

The Memory Hierarchy in Anger: Matrix Transpose

Writing fast GPU kernels is fundamentally about **choreographing data movement through the memory hierarchy**. While modern GPUs boast tens of teraflops of arithmetic power, every floating-point operation is useless if data cannot reach the functional units.

In this chapter, we dissect the GPU memory hierarchy, master global memory coalescing, resolve shared memory bank conflicts, and walk through the canonical optimization benchmark: **Matrix Transpose**.

1. The GPU Memory Hierarchy

A modern GPU features a steeply tiered memory architecture. Each level represents a compromise between capacity, latency, and aggregate bandwidth:

```
+-----+
| Registers: 64K x 32-bit registers per SM (~0-1 cycle latency)|
| Bandwidth: >30 TB/s aggregate across all SMs                |
+-----+
|
+-----+
| Shared Memory / L1 Data Cache: 64 - 228 KB per SM (~19-32 cyc)|
| Bandwidth: ~19 TB/s aggregate (Software-managed SRAM)        |
+-----+
|
+-----+
| L2 Cache: 40 - 50 MB on-chip shared across all SMs (~200 cyc)|
| Bandwidth: ~5.2 TB/s aggregate                               |
+-----+
|
+-----+
| Global Memory (HBM3 / GDDR6): 16 - 96 GB off-chip (~400 cyc)|
| Bandwidth: 1.5 - 3.35 TB/s                                    |
+-----+
```

Memory Space Properties

Memory Space	Scope	Lifetime	Physical Location	Speed
Registers	Single Thread	Thread duration	On-chip Register File	Single-cycle (~0 ns)
Shared Memory	Single Thread	Block duration	On-chip SRAM	~15–30 cycles (~1.5 ns)
L1 Cache	Streaming Multiprocessor	Hardware managed	On-chip SRAM	~15–30 cycles
L2 Cache	All SMs on GPU	Hardware managed	On-chip Crossbar	~150–200 cycles (~15 ns)
Global Memory	All threads + Host	Application duration	Off-chip DRAM (HBM/GDDR)	~400+ cycles (~100 ns)

2. Global Memory Coalescing: The 10x Lever

When threads in a warp issue a read or write instruction to global memory, the memory controller does not process 32 separate individual 4-byte requests. Instead, it aggregates adjacent addresses into **coalesced 32-byte or 128-byte memory transactions**.

The Golden Rule of Coalescing

Consecutive threads in a warp (threadIdx.x from 0 to 31) must access consecutive memory addresses (address + 0, address + 4, ..., address + 124).

Coalesced Access (Consecutive):

```
Threads:      T0    T1    T2    T3    ...    T31
              |     |     |     |           |
              |     |     |     |           |
```

```
Memory Addr: 0x00 0x04 0x08 0x0C    ...    0x7C
```

Result: ONE single 128-byte transaction! (100% bus utilization)

Uncoalesced Strided Access (Column-wise):

```
Threads:      T0          T1          T2    ...    T31
              |           |           |           |
              |           |           |           |
```

```
Memory Addr: 0x0000          0x1000          0x2000 ... 0x1F000 (Stride = N elements)
```

Result: 32 SEPARATE 32-byte transactions! (3.1% bus utilization!)

In the uncoalesced case, the GPU fetches 32×32 Bytes = 1024 Bytes over the physical memory bus to retrieve only 32×4 Bytes = 128 Bytes of useful data. **90% of your memory bandwidth is wasted.**

3. Shared Memory Bank Conflicts

Shared memory is organized into **32 independent memory banks**, each with a word width of **4 bytes (32 bits)**.

Successive 32-bit words are assigned to successive banks in round-robin fashion: - `bank_index = (byte_address / 4) % 32`

```
Bank 0:  Word 0,  Word 32, Word 64 ...
Bank 1:  Word 1,  Word 33, Word 65 ...
...
Bank 31: Word 31, Word 63, Word 95 ...
```

The Conflict Condition

- **Conflict-Free:** All 32 threads in a warp access **different banks** (or all threads broadcast-read the exact same word). The memory request completes in **1 cycle**.
- **N-Way Bank Conflict:** If N threads in a warp request different addresses that map to the **same bank**, the hardware must **serialize** the requests into N successive single-bank accesses. An 8-way conflict reduces shared memory throughput by $8\times$.

4. Case Study: Step-by-Step Optimization of Matrix Transpose

Matrix transpose ($B_{j,i} = A_{i,j}$) is the ultimate memory benchmark: if you read A row-wise (coalesced), you must write B column-wise (uncoalesced), or vice-versa.

Achieving near-peak memory bandwidth requires staging tiles through shared memory.

Step 1: Naive Transpose (Uncoalesced Writes)

```
__global__ void transpose_naive(const float* __restrict__ in,
                                float* __restrict__ out,
                                int width, int height) {
    int x = blockIdx.x * blockDim.x + threadIdx.x; // column
    int y = blockIdx.y * blockDim.y + threadIdx.y; // row

    if (x < width && y < height) {
        // Coalesced Read: Adjacent threads have adjacent x (consecutive in memory)
        float val = in[y * width + x];

        // UNCOALESCEDED Write: Adjacent threads have adjacent y, strided by height!
        out[x * height + y] = val;
    }
}
```

```

    }
}

```

Benchmark Result on A100: Reaches only ~18% of peak global memory bandwidth due to serialized uncoalesced writes.

Step 2: Coalesced Transpose via Shared Memory (With Bank Conflicts)

We load a 2D tile from global memory into shared memory coalesced, synchronize threads, and then write the transposed tile out to global memory coalesced:

```

#define TILE_DIM 32
#define BLOCK_ROWS 8

__global__ void transpose_coalesced_conflicted(const float* __restrict__ in,
                                              float* __restrict__ out,
                                              int width, int height) {
    __shared__ float tile[TILE_DIM][TILE_DIM];

    int x = blockIdx.x * TILE_DIM + threadIdx.x;
    int y = blockIdx.y * TILE_DIM + threadIdx.y;

    // 1. Read input tile coalesced into shared memory
    for (int j = 0; j < TILE_DIM; j += BLOCK_ROWS) {
        if (x < width && (y + j) < height) {
            tile[threadIdx.y + j][threadIdx.x] = in[(y + j) * width + x];
        }
    }

    __syncthreads();

    // Transpose block coordinates for writing
    x = blockIdx.y * TILE_DIM + threadIdx.x;
    y = blockIdx.x * TILE_DIM + threadIdx.y;

    // 2. Write transposed tile coalesced to global memory
    for (int j = 0; j < TILE_DIM; j += BLOCK_ROWS) {
        if (x < height && (y + j) < width) {
            // Reading tile[threadIdx.x][threadIdx.y + j] triggers a 32-WAY BANK CONFLICT!
            out[(y + j) * height + x] = tile[threadIdx.x][threadIdx.y + j];
        }
    }
}

```

The Problem: While global memory reads and writes are now 100% coalesced, reading `tile[threadIdx.x][...]` causes all 32 threads in a warp to access the **same bank simultaneously** (bank index = $(\text{threadIdx.x} * 32 + \dots) \% 32 = 0$), creating a catastrophic **32-way shared memory bank conflict!**

Step 3: Conflict-Free Transpose with Shared Memory Padding

We eliminate bank conflicts with a single architectural stroke: **add one element of padding to each row of the shared memory tile:**

```
// Add +1 column padding: shifts each row by 1 bank!
__shared__ float tile[TILE_DIM][TILE_DIM + 1];
```

With `[TILE_DIM][TILE_DIM + 1]`: - Row 0 starts at Bank 0 ($0 \% 32$). - Row 1 starts at Bank 1 ($33 \% 32 = 1$). - Row 2 starts at Bank 2 ($66 \% 32 = 2$). - Reading column `threadIdx.x` now distributes the 32 threads across **32 unique banks**. Bank conflicts drop to **zero!**

Full Production Code

```
#include <cuda_runtime.h>
#include <stdio.h>

#define TILE_DIM 32
#define BLOCK_ROWS 8

__global__ void transpose_optimal(const float* __restrict__ in,
                                   float* __restrict__ out,
                                   int width, int height) {
    // 32 x 33 shared memory array: +1 padding eliminates all bank conflicts
    __shared__ float tile[TILE_DIM][TILE_DIM + 1];

    int x = blockIdx.x * TILE_DIM + threadIdx.x;
    int y = blockIdx.y * TILE_DIM + threadIdx.y;

    // 1. Coalesced Global Read -> Conflict-Free Shared Write
    #pragma unroll
    for (int j = 0; j < TILE_DIM; j += BLOCK_ROWS) {
        if (x < width && (y + j) < height) {
            tile[threadIdx.y + j][threadIdx.x] = in[(y + j) * width + x];
        }
    }

    __syncthreads();
}
```

```

// Transpose tile destination mapping
x = blockIdx.y * TILE_DIM + threadIdx.x;
y = blockIdx.x * TILE_DIM + threadIdx.y;

// 2. Conflict-Free Shared Read -> Coalesced Global Write
#pragma unroll
for (int j = 0; j < TILE_DIM; j += BLOCK_ROWS) {
    if (x < height && (y + j) < width) {
        out[(y + j) * height + x] = tile[threadIdx.x][threadIdx.y + j];
    }
}
}

```

5. Performance Comparison

Benchmarked on an NVIDIA A100 GPU (4096×4096 FP32 Matrix, Peak HBM Bandwidth = 2,039 GB/s):

Implementation	Effective Bandwidth	% of Peak HBM Bandwidth	Bottleneck / Limiter
Naive Transpose	358 GB/s	17.5%	Severe uncoalesced global write fragmentation
Coalesced with Bank Conflicts	1,220 GB/s	59.8%	32-way shared memory bank conflict serialization
Padded Conflict-Free Transpose	1,845 GB/s	90.5%	Near-hardware speed of light saturation

To verify bank conflicts in Nsight Compute:

```
ncu --metrics l1tex__data_bank_conflicts_pipe_lsu_mem_shared.sum ./transpose_test
```

For the padded kernel, the conflict count drops from hundreds of thousands to **0**.

Part II: The GEMM Worklog: Naive to cuBLAS

GEMM Worklog I: Naive to Tiled

GEMM Worklog I: Naive to Tiled

General Matrix Multiply (**GEMM**) is the computational engine of modern artificial intelligence. Fully connected layers, multi-head attention projections, convolutional layers (via `im2col`), and mixture-of-experts routers all compile down to GEMM:


$$C = \alpha(A \times B) + \beta C$$

where $A \in \mathbb{R}^{M \times K}$, $B \in \mathbb{R}^{K \times N}$, and $C \in \mathbb{R}^{M \times N}$.

In this worklog, we embark on the canonical kernel engineering journey: elevating a naive GEMM achieving less than **1.5%** of GPU capacity to a shared-memory tiled kernel achieving over **36%** of theoretical peak compute.

1. Dissecting the Naive GEMM: Why is it Slow?

In a naive GEMM implementation, each thread computes a single output element $C_{i,j}$ by performing a dot product between row i of matrix A and column j of matrix B .



The Naive CUDA Kernel

```
__global__ void gemm_naive(const float* __restrict__ A,
                          const float* __restrict__ B,
                          float* __restrict__ C,
                          int M, int N, int K) {
    // Determine row and column for this thread
    int row = blockIdx.y * blockDim.y + threadIdx.y;
    int col = blockIdx.x * blockDim.x + threadIdx.x;
```

```

    if (row < M && col < N) {
        float sum = 0.0f;
        // Compute dot product along dimension K
        for (int k = 0; k < K; ++k) {
            sum += A[row * K + k] * B[k * N + col];
        }
        C[row * N + col] = sum;
    }
}

```

Performance Autopsy

Consider running this kernel on an NVIDIA A100 ($M = N = K = 4096$): 1. **Flop Count:** $2 \times M \times N \times K = 2 \times 4096^3 \approx 137.4 \times 10^9$ FLOPs. 2. **Global Memory Loads:** - For every step of loop k , thread (i, j) fetches $A_{i,k}$ (4 bytes) and $B_{k,j}$ (4 bytes) from off-chip DRAM. - Total bytes loaded from DRAM across all threads = $M \times N \times K \times 2 \times 4 \text{ B} \approx 549.7 \times 10^9$ Bytes (~**550 GB**). 3. **Operational Intensity:**

$$I_{\text{naive}} = \frac{2 \text{ FLOPs}}{8 \text{ Bytes}} = 0.25 \text{ FLOPs/Byte}$$

4. **The Bottleneck:** - On the A100, the ridge point is $I^* = 9.6$ FLOPs/Byte. - At $I = 0.25$, the kernel is trapped deep within the **Memory-Bound regime**. - Maximum attainable performance:

$$P_{\text{naive}} = 2,039 \text{ GB/s} \times 0.25 \text{ FLOPs/Byte} \approx 509.7 \text{ GFLOP/s}$$

- **Result:** 509.7 GFLOP/s/19,500 GFLOP/s = **2.6%** of theoretical peak. In reality, uncoalesced memory traffic on matrix B drags actual performance down to ^{**}_{1.3%**} (260 GFLOP/s).

2. The Breakthrough: Shared Memory 2D Block Tiling

To escape the memory bottleneck, we must **reuse loaded data**.

Notice: - All threads in the same row i need the exact same value $A_{i,k}$. - All threads in the same column j need the exact same value $B_{k,j}$.

Instead of streaming from global memory for every multiply-add, thread blocks load a 2D tile of A and a 2D tile of B into ultra-fast on-chip **Shared Memory (SRAM)**, synchronize, and then perform multiple arithmetic operations from SRAM.

```

Global Memory (Slow HBM)
|
(Coalesced Block Load)
v
Shared Memory Tile

```

```

    s_a [TILE x TILE]    s_b [TILE x TILE]
      |
(All threads reuse data TILE times!)
      v
    Registers & ALUs

```

Tile Dimensions and Reuse Factor

With a tile size of $TILE = 32$: - Each thread loads **1 element** of A and **1 element** of B into shared memory. - During the inner multiplication loop, each loaded element is read and reused by **32 separate threads!** - Global memory traffic drops by a factor of $32\times$. - Operational intensity jumps from 0.25 FLOPs/Byte to:

$$I_{\text{tiled}} = 0.25 \times 32 = \mathbf{8.0 \text{ FLOPs/Byte}}$$

At 8.0 FLOPs/Byte, we are right next to the hardware Ridge Point ($I^* = 9.6$).

3. The 2D Tiled CUDA Kernel

```

#include <cuda_runtime.h>
#include <stdio.h>

#define TILE_WIDTH 32

__global__ void gemm_tiled_shared_mem(const float* __restrict__ A,
                                       const float* __restrict__ B,
                                       float* __restrict__ C,
                                       int M, int N, int K) {

    // 2D shared memory cache tiles
    __shared__ float s_A[TILE_WIDTH][TILE_WIDTH];
    __shared__ float s_B[TILE_WIDTH][TILE_WIDTH];

    int tx = threadIdx.x;
    int ty = threadIdx.y;

    // Output coordinate in matrix C
    int row = blockIdx.y * TILE_WIDTH + ty;
    int col = blockIdx.x * TILE_WIDTH + tx;

    float acc = 0.0f;

    // Number of phase tiles needed along dimension K

```

```

int num_tiles = (K + TILE_WIDTH - 1) / TILE_WIDTH;

for (int t = 0; t < num_tiles; ++t) {
    // 1. Collaborative loading of Tile A into shared memory
    int a_col = t * TILE_WIDTH + tx;
    if (row < M && a_col < K) {
        s_A[ty][tx] = A[row * K + a_col];
    } else {
        s_A[ty][tx] = 0.0f;
    }

    // 2. Collaborative loading of Tile B into shared memory
    int b_row = t * TILE_WIDTH + ty;
    if (b_row < K && col < N) {
        s_B[ty][tx] = B[b_row * N + col];
    } else {
        s_B[ty][tx] = 0.0f;
    }

    // 3. Barrier synchronization: ensure entire tile is present in SRAM
    __syncthreads();

    // 4. Accumulate partial dot product from shared memory
    #pragma unroll
    for (int k = 0; k < TILE_WIDTH; ++k) {
        acc += s_A[ty][k] * s_B[k][tx];
    }


    // 5. Barrier synchronization: ensure computation finishes before next tile load
    __syncthreads();
}

// Write final accumulated result to global memory
if (row < M && col < N) {
    C[row * N + col] = acc;
}
}

```

4. Benchmark Results on NVIDIA A100

Evaluating on square FP32 matrices ($M = N = K = 4096$):



Implementation	Execution Time ($M = N = K = 4096$)	Measured TFLOP/s	% of Theoretical Peak Compute
Naive GEMM	528.4 ms	0.26 TFLOP/s	1.3%
2D Tiled Shared Memory GEMM	19.3 ms	7.12 TFLOP/s	36.5%
NVIDIA cuBLAS (Closed Source)	7.5 ms	18.32 TFLOP/s	93.9%

Why Did Performance Jump by 27x?

1. **Global Memory Bandwidth Slashed:** Memory traffic dropped from 550 GB to ~17 GB.
2. **Coalesced Reads:** Both A and B tiles are fetched using contiguous thread indexing (`tx`).
3. **Latency Hiding:** Staged data in shared memory is accessed in ~20 cycles instead of ~400 cycles.

In the next chapter, we push further into the micro-architecture: register micro-tiling, `float4` vectorization, and warp-level hierarchical tiling to close the gap to cuBLAS.

GEMM Worklog II: Registers & Warp Tiling

GEMM Worklog II: Registers & Warp Tiling

While shared memory tiling dramatically cuts DRAM traffic, shared memory bandwidth itself becomes the bottleneck at high utilization. Reaching near-cuBLAS performance requires **register tiling** and **vectorized memory instructions**.

Micro-Tiling in Registers

Registers provide ~20 TB/s aggregate bandwidth. Instead of each thread calculating a single output element $C_{i,j}$, each thread computes a $TMimesTN$ block of elements (e.g. $8imes8 = 64$ values per thread).

Vectorized Memory Loads with float4

Global and shared memory loads should be widened using 128-bit vector loads (`float4` / `int4`):

```
// Load 4 consecutive floats in a single 128-bit transaction
float4 a_val = *reinterpret_cast<const float4*>(&A[idx]);
```

The Warptiling Hierarchy

Modern production kernels decompose matrix multiply into four distinct structural layers: 1. **Grid Tile:** Handled by the entire GPU grid across all SMs. 2. **Block Tile:** Handled by a single thread block using shared memory. 3. **Warp Tile:** Handled by a 32-thread warp executing synchronized sub-tiles. 4. **Thread Tile:** Handled by individual thread registers.

By structuring the kernel into a $128imes128$ block tile and $64imes64$ warp tiles with double buffering, manual register allocation reaches **93.7% of cuBLAS speed** on NVIDIA Ampere hardware.

Tensor Cores & WMMA: The Second Worklog

Tensor Cores & WMMA: The Second Worklog

Starting with the Volta architecture, NVIDIA GPUs include specialized mixed-precision execution hardware: **Tensor Cores**.

The Tensor Core Advantage

- Conventional CUDA cores execute 1 FMA (Fused Multiply-Add = 2 FLOPs) per cycle.
- Tensor Cores compute full matrix multiply-accumulate operations ($D = A \times B + C$) on small matrix fragments (e.g. $16 \times 16 \times 16$) in hardware every cycle.
- Speedup: ~5-10x higher theoretical compute density compared to FP32 CUDA cores.

Direct Programming via WMMA (Warp Matrix Multiply Accumulate)

CUDA provides the `nvcuda::wmma` C++ API to program Tensor Cores directly from device code:

```
#include <mma.h>
using namespace nvcuda;

// Cooperative 16x16x16 matrix multiply across 32 threads in a warp
__global__ void wmma_gemm_kernel(const half* A, const half* B, float* C, int M, int N, int K)
{
    // Declare WMMA fragments
    wmma::fragment<wmma::matrix_a, 16, 16, 16, half, wmma::row_major> a_frag;
    wmma::fragment<wmma::matrix_b, 16, 16, 16, half, wmma::col_major> b_frag;
    wmma::fragment<wmma::accumulator, 16, 16, 16, float> c_frag;

    wmma::fill_fragment(c_frag, 0.0f);

    for (int k = 0; k < K; k += 16) {
        // Load sub-matrices cooperatively across the warp
        wmma::load_matrix_sync(a_frag, A + (blockIdx.y * 16 * K) + k, K);
        wmma::load_matrix_sync(b_frag, B + (blockIdx.x * 16 * K) + k, K);

        // Hardware Tensor Core Matrix Multiply
        wmma::mma_sync(c_frag, a_frag, b_frag, c_frag);
    }

    // Store accumulator fragment back to global memory
    wmma::store_matrix_sync(C + (blockIdx.y * 16 * N) + (blockIdx.x * 16), c_frag, N, wmma::);
}
```

Part III: Attention & Profiling

Profiling & Debugging Like a Pro

Profiling & Debugging Like a Pro

Optimization without profiling is speculation. Top kernel engineers diagnose bottlenecks systematically using the NVIDIA toolchain:

The Professional Diagnostic Stack

1. **compute-sanitizer:** Traps out-of-bounds memory accesses, misaligned memory addresses, uninitialized shared memory, and warp synchronization race conditions.
2. **NVIDIA Nsight Systems (NSYS):** Timeline visualization across CPU threads, CUDA streams, memory transfers, and kernel durations.
3. **NVIDIA Nsight Compute (NCU):** Deep architectural profiling inside individual kernels—analyzing warp stall reasons, SM issue pipe utilization, memory chart traffic, and register spills.

Debugging Sabotaged Kernels: 3 Classic Incident Scenarios

- **Incident 1: Silent Race Condition in Shared Memory:** Threads write to `s_data[threadIdx.x]` and read from `s_data[threadIdx.x + 1]` without an intervening `__syncthreads()`. *Detection:* `compute-sanitizer --tool racecheck ./my_kernel`
- **Incident 2: Shared Memory Bank Conflicts:** All 32 threads in a warp access different rows in the same shared memory bank simultaneously. *Detection:* NCU metric `l1tex__data_bank_conflicts_pipe_lsu_mem_shared.sum`. *Resolution:* Stride padding (`__shared__ float smem[32][33]`).
- **Incident 3: Catastrophic Register Spilling:** Kernel uses 80 registers per thread when the hardware limit per thread at target occupancy is 64. Excessive variables spill to local DRAM memory, increasing latency 100x. *Detection:* Compiler flag `--ptxas-options=-v` reporting `spill stores / spill loads`. *Resolution:* Kernel launch bound annotation `__launch_bounds__(256, 2)`.

Attention: The Kernel That Ate the World & FlashAttention v1

Attention: The Kernel That Ate the World & FlashAttention v1

Standard Transformer Attention computes:

$$\text{Attention}(Q, K, V) = \text{Softmax} \left(\frac{QK^T}{\sqrt{d_k}} \right) V$$

For a sequence length N and head dimension d ($N \gg d$): - $S = QK^T$ generates an $N \times N$ attention score matrix. - Traditional PyTorch implementations materialize S and $P = \text{Softmax}(S)$ in global HBM memory ($O(N^2)$ memory reads and writes). - For long contexts ($N = 32k$ or $128k$), intermediate matrices consume gigabytes of memory and choke HBM bandwidth.

The Breakthrough: Online Softmax & Tiling

Tri Dao et al. introduced **FlashAttention**: compute the exact attention output without ever writing the $N \times N$ intermediate attention matrix to global memory.

Traditional Attention (Memory-Bound):

Q, K -> [Global HBM Write: S] -> [Global HBM Read: S] -> Softmax -> [Global HBM Write: P]

FlashAttention (Fused into SRAM):

Load Blocks of Q, K, V into On-chip SRAM -> Compute Local Gemm -> Online Softmax Update ->

The Online Softmax Formula

Traditional softmax requires two passes over the sequence: pass 1 to find the maximum $m = \max(x_i)$, and pass 2 to compute $\sum e^{x_i - m}$. Online softmax updates the running maximum m and normalization sum l incrementally as blocks stream through shared memory:

$$m_{\text{new}} = \max(m_{\text{prev}}, m_{\text{block}})$$

$$l_{\text{new}} = e^{m_{\text{prev}} - m_{\text{new}}} l_{\text{prev}} + e^{m_{\text{block}} - m_{\text{new}}} l_{\text{block}}$$

$$O_{\text{new}} = e^{m_{\text{prev}} - m_{\text{new}}} O_{\text{prev}} + P_{\text{block}} V_{\text{block}}$$

Part IV: The Modern Frontier

FlashAttention from Scratch (FA1 $\rightarrow$ FA2 $\rightarrow$ FA3)

FlashAttention from Scratch (FA1 → FA2 → FA3)

FlashAttention Evolution

- **FlashAttention-1 (2022):** Proved that fusing attention and computing online softmax in SRAM yields 2-4x wall-clock speedups while cutting memory footprint from $O(N^2)$ to $O(N)$.
- **FlashAttention-2 (2023):**
 - Swapped the outer and inner loops (looping over blocks of Q in the outer loop, streaming K, V in the inner loop).
 - Reduced non-matmul FLOPs by deferring output scaling until the very end.
 - Introduced sequence parallel work distribution across thread blocks.
- **FlashAttention-3 (2024):**
 - Architected specifically for NVIDIA Hopper (H100) hardware.
 - Utilizes asynchronous Tensor Memory Accelerator (TMA) hardware pipelines.
 - Leverages Hopper WGMMMA (Warp Group Matrix Multiply and Accumulate) instructions.
 - Implements FP8 mixed precision with low quantization error.

Beating cuBLAS on an H100

Beating cuBLAS on an H100

NVIDIA's Hopper H100 microarchitecture introduced two revolutionary hardware subsystems:

1. Tensor Memory Accelerator (TMA)

In pre-Hopper architectures, threads in a warp explicitly load data from global memory into registers, then write from registers into shared memory. TMA executes multidimensional tensor transfers (e.g. 2D/3D tiles) directly between global HBM and shared memory **asynchronously in hardware**, completely bypassing the register file and consuming zero SM ALU cycles.

2. Warp Specialization

Instead of every thread in a block performing both memory loading and arithmetic: - **Producer Warps (1-2 warps)**: Issue asynchronous TMA copy instructions to fetch upcoming tiles into double-buffered shared memory. - **Consumer Warps (Warp Groups)**: Execute continuous WGMMMA matrix multiply instructions without ever pausing for memory latency. - Hardware barrier arrive/wait primitives (`cuda::barrier`) synchronize producers and consumers with cycle-level precision.

Triton → CUTLASS → CuTe-DSL

Triton → CUTLASS → CuTe-DSL

Kernel engineering in 2025-2026 has transitioned from hand-written CUDA C++ toward declarative and high-level DSL frameworks.

OpenAI Triton

Triton allows writing high-performance GPU kernels in idiomatic Python. The programmer writes code at the granularity of block tiles rather than individual threads. The Triton compiler handles intra-block scheduling, memory coalescing, and shared memory allocation automatically.

```
import triton
import triton.language as tl

@triton.jit
def add_kernel(x_ptr, y_ptr, output_ptr, n_elements, BLOCK_SIZE: tl.constexpr):
    pid = tl.program_id(axis=0)
    block_start = pid * BLOCK_SIZE
    offsets = block_start + tl.arange(0, BLOCK_SIZE)
    mask = offsets < n_elements
    x = tl.load(x_ptr + offsets, mask=mask)
    y = tl.load(y_ptr + offsets, mask=mask)
    tl.store(output_ptr + offsets, x + y, mask=mask)
```

CUTLASS & CuTe

NVIDIA CUTLASS 3.x is structured around **CuTe**: a C++ template library that unifies data layouts, coordinates, and hierarchical tensor tiling through mathematical layout algebra:

$$extLayout = (extShape, extStride)$$

Inference-Serving Kernels

Inference-Serving Kernels

LLM inference splits into two distinct execution regimes with radically different performance characteristics: 1. **Prefill Phase:** The prompt is processed in parallel. High operational intensity, strictly compute-bound. 2. **Decode Phase:** Tokens are generated autoregressively one by one. The model loads billions of weights to compute a single new vector. Strictly memory-bound!

PagedAttention (vLLM)

In standard inference, KV caches must be allocated as contiguous memory buffers. This leads to 60-80% memory waste due to internal and external memory fragmentation. PagedAttention treats KV cache memory like virtual memory pages in an operating system: - KV vectors are stored in non-contiguous 16-token page blocks. - A page table maps logical token positions to physical GPU memory addresses. - Result: Memory waste drops to near 0%, enabling 2-4x higher serving concurrency on the same hardware.

Blackwell & NVFP4

Blackwell & NVFP4

NVIDIA's Blackwell (B200 / GB200) microarchitecture represents the state of the art in frontier hardware (2025-2026).

Blackwell Architectural Highlights

- **tcgen05 Tensor Cores:** 5th generation Tensor Cores delivering 4x higher FP4 training and inference density compared to Hopper FP8.
- **Tensor Memory (TMEM):** A dedicated ultra-fast on-chip memory structure decoupled from shared memory, specifically optimized for multi-tensor staging and GEMM accumulation.
- **NVFP4 Microscopic Scaling:** 4-bit floating point format with micro-block scaling factors that preserves accuracy across large foundation models while doubling token throughput.

Flash Attention 4

Flash Attention 4

The 2026 Frontier: Breaking the 1 PFLOPS Single-GPU Barrier

On Blackwell hardware, Tensor Core matrix multiply units run so extraordinarily fast that the normalization step in Softmax—historically a minor fraction of runtime—becomes the dominant bottleneck.

FlashAttention-4 introduces **Adaptive Softmax**: - Dynamically checks divergence between running block maximums and candidate maximums. - Skips full re-normalization when numerical differences fall below precision thresholds. - Overlaps Softmax correction entirely behind asynchronous WGMMMA execution pipelines. - Written using modern Python CuTe-DSL, FA4 is the first kernel to achieve over 1 PFLOPS of sustained attention compute on a single processor.

PagedAttention and the KV Cache

PagedAttention and the KV Cache

Autoregressive LLM serving is memory-bound in decode and **capacity-bound** in concurrency: each active sequence carries a growing key/value (KV) cache. Contiguous per-sequence KV allocations waste large fractions of GPU DRAM to fragmentation and over-reservation. **PagedAttention** (vLLM) borrows the OS virtual-memory idea—fixed-size blocks, a block table, allocate-on-demand—so attention kernels read logically contiguous tokens from physically non-contiguous KV blocks.

This chapter goes deeper than a one-paragraph mention in inference-serving notes: why contiguous KV fails, how block tables work, prefill vs decode implications, layout choices for custom kernels, and how to verify/benchmark paging behavior.

Mental model

OS analogy	LLM serving
Process virtual address space	Logical token positions $0..T - 1$ in a sequence
Pages	KV blocks of B tokens
Page table	Per-sequence block table (logical block $\rightarrow$ physical block)
Physical frames	Slots in a GPU KV pool
Internal fragmentation	At most one partially filled tail block
External fragmentation	Largely eliminated (any free block can be used)

```
logical tokens:  [t0 t1 ... t15 | t16 ... t31 | t32 ... t40____]
                  block 0          block 1          block 2 (partial)
```

```
block table:     L0  $\rightarrow$  P17
                  L1  $\rightarrow$  P3
                  L2  $\rightarrow$  P9
```

```
physical pool:   P0 P1 P2 P3[K/V] ... P9[K/V] ... P17[K/V] ...
```

Attention still needs Q against all prior K , and weighted sum over V —but the kernel **gathers** K, V through the block table instead of assuming one contiguous pointer.

Why contiguous KV wastes memory

For a sequence with maximum length $L_{\max}$ and current length t :

1. **Reservation waste:** many systems allocate KV for $L_{\max}$ up front even while $t \ll L_{\max}$.
2. **External fragmentation:** variable-length contiguous regions cannot pack tightly as sequences finish at different times.
3. **Batching rigidity:** contiguous layouts make sharing prefixes or compacting live sequences expensive.

Empirical reports in the PagedAttention / vLLM line of work put waste often in the **60–80%** range under contiguous reservation—memory that cannot become additional concurrent sequences. Paging caps internal waste at roughly one block (B tokens) per sequence (per managed cache partition) and removes most external fragmentation.

Capacity sketch

Let one token’s KV footprint (all layers, all heads, K and V) be b bytes. Contiguous reservation for N sequences costs $\sim NL_{\max}b$. Paged on-demand costs $\sim b \sum_i t_i$ plus $< NBb$ partial-block slack (order-of-magnitude). The ratio $L_{\max}/\bar{t}$ dominates the win when outputs are much shorter than the configured max.

KV cache in Transformers (reminder)

For each layer ℓ and head h , token position i stores vectors $k_i^{(\ell,h)}, v_i^{(\ell,h)} \in \mathbb{R}^d$. Causality at decode step t uses

$$\alpha_{t,i} = \text{softmax}_i \left(\frac{q_t^\top k_i}{\sqrt{d}} \right), \quad o_t = \sum_{i=1}^t \alpha_{t,i} v_i.$$

Prefill computes many q rows in parallel against the prompt’s K, V ; decode appends one new (k_t, v_t) and attends with a single q_t (or a small speculative set).

Serving reality: weights are shared; **KV grows per request**. Throughput (tokens/s/GPU) often rises when more sequences fit—hence KV packing is a first-class systems problem, not a minor allocator tweak.

Block structure

Choose block size B (vLLM commonly discusses values like 16). Logical block j holds tokens $(j-1)B+1, \dots, jB$ (or 0-based equivalents):

$$K_j = (k_{(j-1)B+1}, \dots, k_{jB}), \quad V_j = (v_{(j-1)B+1}, \dots, v_{jB}).$$

Attention becomes a loop over logical blocks: load the physical block pointed by the table, compute partial scores/outputs, advance. The last block carries a **filled-token count** $1..B$.

Worked layout numbers

$B = 16$, head dim $d = 128$, dtype FP16: one head's K block stores 16×128 elements = 2048 halves ≈ 4 KiB (order of magnitude), and V similarly. Multi-layer multi-head caches scale linearly; production engines may use separate block tables per layer/head partition or pack carefully for coalescing—design choice, not a change to the paging idea.

Block table and allocator

Per sequence (and often per KV partition):

Field	Role
Logical index j	Token range $[jB, (j+1)B)$
Physical block id	Index into GPU (or CPU swap) pool
Filled count	Valid tokens in the tail block
Refcount (advanced)	Prefix sharing / copy-on-write

Allocate on demand: when decode would write past the filled count, pull a free physical block, append a table entry. **Free on finish:** return blocks to the pool (and decrement refcounts if shared).

```
decode step:
  if filled == B:
    phys = alloc_block()
    table.append(phys)
    filled = 0
  write k_t, v_t into table[-1] at offset filled
  filled += 1
  run paged_attention(q_t, table)
```

Prefix sharing (copy-on-write intuition)

If two requests share a prompt prefix, map their early logical blocks to the **same** physical blocks with elevated refcounts. Divergence (different generated tokens) allocates private blocks from the fork point. Automatic prefix caching builds on the same indirection: hash logical block contents (+prefix identity) to reuse physical KV.

Prefill vs decode under paging

Phase	Compute character	Paging interaction
Prefill	Compute-heavy; many tokens written	Allocate enough blocks for prompt; write K/V densely left-to-right; attention over prompt blocks
Decode	Memory-bound; one token at a time	Append into tail block; attention gathers all prior blocks for one q

Custom kernels must handle **variable block lists per sequence in a batch**: sequence A may have 3 blocks, sequence B 40. Batching strategies (continuous batching) reshape which sequences sit in a wave while the block tables remain the source of truth for KV addresses.

Implication

A FlashAttention-style tiled softmax over a **contiguous** K, V buffer is not drop-in compatible. You need a kernel that:

1. Accepts block tables + pool base pointers, or
2. Materializes contiguous KV (defeats paging), or
3. Uses a hybrid (e.g., paged decode kernels + specialized prefill paths).

vLLM's historical PagedAttention CUDA kernels (`attention_kernels`-style designs) are built for case (1): thread blocks iterate paged K/V vectors with careful vectorized loads.

Paged KV layouts (kernel writer view)

Exact layouts evolve; treat the following as a **mental checklist** when reading kernels or writing prototypes:

```
Pool: physical_block[p] → storage for B tokens × head_dim (K or V)
```

```
K access for logical token t:
```

```
  j = t / B
```

```
  o = t % B
```

```
  p = block_table[seq][j]
```

```
  addr = base_K + p * stride_block + o * stride_token + head_offset
```

```
V access: similar with base_V / possibly different vectorization shape
```

K vs V vectorization: kernels often pack K for efficient $Q \cdot K$ dots (e.g., threads cooperating on a head vector) while V layouts favor weighting and accumulating into the output. Do not assume K and V share identical lane mapping.

Separate K and V caches: common—two pools, parallel tables or mirrored indices.

Illustrative host-side table (not a production kernel)

```
# file: toy_block_table.py
# Toy model of logical→physical mapping (host-side). Not CUDA.

from dataclasses import dataclass, field

B = 16

@dataclass
class SequenceKV:
    table: list[int] = field(default_factory=list) # physical ids
    filled: int = 0 # tokens used in last block

class BlockPool:
    def __init__(self, n_blocks: int):
        self.free = list(range(n_blocks))

    def alloc(self) -> int:
        if not self.free:
            raise MemoryError("KV pool exhausted")
        return self.free.pop()

    def free_block(self, p: int) -> None:
        self.free.append(p)

def append_token(seq: SequenceKV, pool: BlockPool) -> None:
    if not seq.table or seq.filled == B:
        seq.table.append(pool.alloc())
        seq.filled = 0
    # device would write K/V into pool[seq.table[-1]][seq.filled]
    seq.filled += 1

def release(seq: SequenceKV, pool: BlockPool) -> None:
    for p in seq.table:
        pool.free_block(p)
    seq.table.clear()
    seq.filled = 0
```

Use this only to reason about allocation traffic; real engines track layers/heads, CUDA streams, and fragmentation across devices.

Implications for custom kernels

1. **Indirection tax:** every K/V load pays a table lookup (often cached in shared memory / registers for the active block list).
2. **Tail blocks:** masks for `filled < B` are mandatory—same family of bugs as ragged sequence lengths.
3. **Batching:** one wave may include many tables; prefer grouped launches or paged-friendly attention APIs over naively padding to $L_{\max}$ contiguous.
4. **Fusion:** paging composes with continuous batching, speculative decoding (extra parallel tokens), and quantization of KV (INT8/FP8 caches)—each changes strides and scales, not the block-table story.
5. **Swap / multi-tier:** physical blocks may live in CPU RAM; kernels and schedulers must not assume all table entries are GPU-resident without a residency check.

When generating kernels with AI assistants: specify **block size**, **K/V separation**, **dtype**, **head dim**, and **tail masking** explicitly. Ambiguous “write attention” prompts usually emit contiguous FlashAttention clones that cannot consume a paged pool.

Verification and benchmark notes

Functional checks

Check	Idea
Match contiguous golden	For small T , gather paged KV to contiguous and compare attention outs vs a reference softmax
Tail mask	Lengths $B - 1$, B , $B + 1$ across block boundaries
Multi-sequence batch	Different table lengths in one launch
Prefix share	RefCount: mutate after fork must not clobber peer

Performance / capacity checks

Metric	Why
Max concurrent sequences at fixed $L_{\max}$	Primary paging win
KV bytes reserved vs bytes touched	Fragmentation / over-reserve detector

Metric	Why
Decode tokens/s vs batch size	Memory bandwidth + kernel efficiency
Allocation failures under load	Pool sizing / preemption / swap policy

Microbenchmarks that only time a contiguous GEMM miss the point—page-table attention is about **fitting more live KV** without catastrophic gather overhead. Profile both **allocator** and **attention kernel**.

Profiling tips

- Trace block alloc/free counts per request; spikes may indicate missing prefix sharing.
- Compare achieved occupancy / bytes loaded for paged vs contiguous decode on identical batch shapes.
- Watch for CPU-sync on block-table uploads each step—pin tables or update incrementally.

Pitfalls

Pitfall	Result
Assuming contiguous FlashAttention works unchanged	Wrong answers or illegal loads
Ignoring partial-block fill	Garbage tokens in softmax
Oversized B	More internal waste; less flexible packing
Undersized pool	Preemption thrash or request rejects
Forgetting refcount on shared prefixes	Silent cross-request corruption
Benchmarking only prefill TFLOPs	Misses decode memory + paging behavior

Checkpoint

- Explain why contiguous $L_{\max}$ reservation destroys concurrency.
- Map OS paging terms onto block tables and KV pools.
- Describe allocate-on-demand during decode and tail **filled** counts.
- List what a paged attention kernel must accept beyond Q, K, V pointers.
- Name two functional tests and one capacity metric for paging.

Summary

PagedAttention makes KV cache **virtually contiguous and physically sparse**. The win is utilization—more sequences per GPU—at the cost of attention kernels that understand block tables, tail masks, and pool lifetimes. Prefill still wants heavy compute paths; decode wants lean paged gathers. For custom or AI-written kernels, treat the block table as part of the ABI, verify against contiguous goldens at block boundaries, and benchmark concurrency and KV bytes—not only matmul TFLOPs.

Sources (for further reading)

- Kwon et al., “Efficient Memory Management for Large Language Model Serving with PagedAttention” (SOSP 2023; arXiv:2309.06180).
- vLLM design docs on PagedAttention / paged KV layouts.
- vLLM blog: “Easy, Fast, and Cheap LLM Serving with PagedAttention.”

Part V: AI-Written Kernels

DeepSeek: FlashMLA & DeepGEMM

DeepSeek: FlashMLA & DeepGEMM

DeepSeek stunned the AI industry by open-sourcing compact, ultra-efficient kernels that outcompeted monolithic enterprise libraries.

FlashMLA: Multi-Head Latent Attention

DeepSeek-V3 and R1 utilize Multi-Head Latent Attention (MLA) to compress KV caches into low-dimensional latent spaces. Standard attention kernels cannot handle this asymmetrical projection efficiently. FlashMLA is custom-tuned to decode latent vectors directly into shared memory, cutting KV memory bandwidth consumption by over 70%.

DeepGEMM: Clean Architecture over Brute Force

In fewer than 300 lines of readable, modular CUDA code, **DeepGEMM** implements fine-grained FP8 matrix multiplication with grouped scaling. By avoiding deeply nested template metaprogramming and focusing purely on optimal software pipelining and TMA scheduling, DeepGEMM matches and often beats commercial closed-source GEMM kernels.


LLM-Generated Kernels: The Agent + Profiler Loop

LLM-Generated Kernels: The Agent + Profiler Loop

We are entering a transformative era where AI models generate, optimize, and verify GPU kernels autonomously.

The Agentic Optimization Loop

[Prompt / Specification]



Google DeepMind AlphaEvolve & KernelBench

- **KernelBench:** An automated evaluation benchmark testing whether language models can generate numerically correct and high-performance CUDA/Triton kernels across realistic neural network primitives.
- **DeepMind AlphaEvolve:** An evolutionary coding agent combining LLM code mutations with real hardware profiler feedback. AlphaEvolve discovered non-trivial memory layout schedules that sped up production attention kernels by ~32%.

The Role of the Human Kernel Engineer

The engineer's job is not becoming obsolete; it is elevating: - Defining invariants, numerical bounds, and correctness testbenches. - Providing architectural guidance and identifying hardware bottlenecks. - Interpreting profiler telemetry to steer the agent's exploration away from local minima.

Part VI: Kernel Engineering Capstone Projects

Crusoe Capstone Specifications & Problem Tracks

Crusoe Capstone Specifications & Problem Tracks

The capstone project requires building a production-ready GPU kernel addressing genuine infrastructure challenges.

Problem Track 1: Fused Sparse Attention

Implement a fused attention kernel supporting block-sparse patterns (e.g. BigBird or sliding window + global tokens). The kernel must: - Skip computation of zero-weight blocks entirely. - Maintain online softmax correctness across sparse blocks. - Outperform standard dense PyTorch attention by at least 2x on sequences >16k.

Problem Track 2: Quantized FP4/FP8 GEMM with Dynamic Dequantization

Build a high-throughput matrix multiply kernel that ingests FP8 or INT4 quantized weights and dequantizes on-the-fly inside SM registers before execution.

Problem Track 3: High-Throughput Speculative Decoding Verification Kernel

Design a batched verification kernel for speculative decoding that verifies K draft tokens concurrently against a target model distribution in a single fused pass.

Verification, Profiling, and Benchmarking Harness

Verification, Profiling, and Benchmarking Harness

Every kernel submission must pass a three-tiered validation suite:

1. Correctness Verification

Compare kernel output tensors against double-precision PyTorch / FP64 reference outputs across edge-case dimensions, non-multiple-of-32 matrix sizes, and extreme numerical scales (10^{-5} to 10^5). Maximum allowable relative error:

$$\frac{\|Y_{\text{kernel}} - Y_{\text{ref}}\|_2}{\|Y_{\text{ref}}\|_2} < 10^{-3}$$

2. CUDA Event Benchmarking

Measure kernel execution time using `cudaEventRecord` over 1,000 warm iterations, discarding outliers:

```
cudaEvent_t start, stop;
cudaEventCreate(&start);
cudaEventCreate(&stop);

cudaEventRecord(start);
for (int i = 0; i < 1000; ++i) {
    launch_kernel<<<grid, block>>>(...);
}
cudaEventRecord(stop);
cudaEventSynchronize(stop);

float ms = 0;
cudaEventElapsedTime(&ms, start, stop);
printf("Average execution time: %f ms\n", ms / 1000.0f);
```

Demo Day Presentation & Portfolio

Demo Day Presentation & Portfolio

Presenting your work effectively is as critical as writing fast code.

Deliverables Checklist

1. **GitHub Repository:** Clean, modular repository containing your kernel source, CMake / setup.py build configuration, and automated test runners.
2. **Performance Roofline Graph:** A generated plot showing where your kernel sits relative to theoretical memory and compute ceilings across varying matrix dimensions.
3. **Speedup Comparison Table:** Side-by-side throughput (TFLOP/s) comparisons against naive CUDA, PyTorch `torch.compile`, and official vendor libraries (cuBLAS / cuDNN).
4. **Technical Defense:** Prepare concise explanations for architectural decisions: tile sizes, shared memory allocations, register counts per thread, and how AI-assisted iteration shaped the final solution.